
Amazon Relational Database Service

User Guide



Amazon Relational Database Service: User Guide

Copyright © 2019 Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

Table of Contents

What Is Amazon RDS?	1
Overview	1
DB Instances	1
AWS Regions and Availability Zones	2
Security	2
Monitoring an Amazon RDS DB Instance	3
How to Work with Amazon RDS	3
AWS Management Console	3
Command Line Interface	3
Programming with Amazon RDS	3
How You Are Charged for Amazon RDS	3
What's Next?	4
Getting Started	4
Database Engine-Specific Topics	4
Setting Up	5
Sign Up for AWS	5
Create an IAM User	5
Determine Requirements	7
Provide Access to Your DB Instance in Your VPC by Creating a Security Group	8
Getting Started	11
Creating a MariaDB DB Instance and Connecting to a Database	11
Creating a MariaDB Instance	11
Connecting to a Database on a DB Instance Running MariaDB	20
Deleting a DB Instance	22
Creating a Microsoft SQL Server DB Instance and Connecting to a DB Instance	23
Creating a Sample SQL Server DB Instance	23
Connecting to Your Sample DB Instance	30
Exploring Your Sample DB Instance	31
Deleting Your Sample DB Instance	32
Creating a MySQL DB Instance and Connecting to a Database	33
Creating a MySQL DB Instance	33
Connecting to a Database on a DB Instance Running MySQL	42
Deleting a DB Instance	44
Creating an Oracle DB Instance and Connecting to a Database	45
Creating a Sample Oracle DB Instance	45
Connecting to Your Sample DB Instance	53
Deleting Your Sample DB Instance	55
Creating a PostgreSQL DB Instance and Connecting to a Database	55
Creating a PostgreSQL DB Instance	55
Connecting to a PostgreSQL DB Instance	64
Deleting a DB Instance	68
Tutorial: Create a Web Server and an Amazon RDS Database	69
Step 1: Create a DB Instance	70
Next Step	81
Step 2: Create a Web Server	81
Tutorials	93
Best Practices for Amazon RDS	94
Amazon RDS Basic Operational Guidelines	94
DB Instance RAM Recommendations	95
Using Enhanced Monitoring to Identify Operating System Issues	95
Using Metrics to Identify Performance Issues	95
Viewing Performance Metrics	95
Evaluating Performance Metrics	98
Tuning Queries	99

Best Practices for Working with MySQL Storage Engines	100
Best Practices for Working with MariaDB Storage Engines	101
Best Practices for Working with Oracle	101
Best Practices for Working with PostgreSQL	101
Loading Data into a PostgreSQL DB Instance	101
Working with the fsync and full_page_writes database parameters	102
Working with the PostgreSQL Autovacuum Feature	102
Best Practices for Working with SQL Server	103
Working with DB Parameter Groups	103
Amazon RDS Best Practices Presentation Video	104
DB Instances	105
Choosing the DB Instance Class	107
DB Instance Class Types	107
Terminology	108
Hardware Specifications	109
Supported DB Engines	112
Changing Your DB Instance Class	116
Configuring the Processor	117
DB Instance Status	129
Regions and Availability Zones	132
.....	132
DB Instance Storage	135
Storage Types	135
General Purpose SSD Storage	135
Provisioned IOPS Storage	137
Magnetic storage	138
Monitoring storage performance	139
Factors That Affect Storage Performance	139
High Availability (Multi-AZ)	144
Modifying a DB Instance to be a Multi-AZ Deployment	145
Failover Process for Amazon RDS	145
Related Topics	146
DB Instance Lifecycle	147
Creating a DB Instance	148
Connecting to a DB Instance	149
Modifying a DB Instance	150
Maintaining a DB Instance	152
Upgrading the Engine Version	159
Renaming a DB Instance	163
Rebooting a DB Instance	166
Stopping a DB Instance	169
Starting a DB Instance	172
Deleting a DB Instance	174
Tagging RDS Resources	177
Overview	177
Working with Read Replicas	182
Overview	182
Creating a Read Replica	184
Promoting a Read Replica	186
Creating a Read Replica in a Different AWS Region	188
Monitoring Read Replication	194
Working with Option Groups	196
Option Groups Overview	196
Creating an Option Group	197
Making a Copy of an Option Group	199
Adding an Option to an Option Group	200
Listing the Options and Option Settings for an Option Group	203

Modifying an Option Setting	204
Removing an Option from an Option Group	207
Deleting an Option Group	208
Working with Parameter Groups	211
Creating a DB Parameter Group	212
Modifying Parameters in a DB Parameter Group	213
Copying a DB Parameter Group	217
Listing DB Parameter Groups	218
Viewing Parameter Values for a DB Parameter Group	219
Comparing DB Parameter Groups	220
DB Parameter Values	220
Working with ARNs	224
Constructing an ARN	224
Getting an Existing ARN	227
Working with Storage	231
Increasing DB Instance Storage Capacity	231
Managing Capacity Automatically with Storage Autoscaling	232
Modifying Provisioned IOPS	236
DB Instance Billing for Amazon RDS	239
On-Demand DB Instances	240
Reserved DB Instances	241
Backing Up and Restoring	253
Working With Backups	254
Backup Storage	254
Backup Window	254
Backup Retention Period	256
Disabling Automated Backups	256
Enabling Automated Backups	257
Retaining Automated Backups	259
Automated Backups with Unsupported MySQL Storage Engines	262
Automated Backups with Unsupported MariaDB Storage Engines	263
.....	263
Creating a DB Snapshot	264
Restoring from a DB Snapshot	266
Parameter Groups	266
Security Groups	266
Option Groups	266
Microsoft SQL Server	267
Oracle	267
Restoring from a Snapshot	267
Copying a Snapshot	269
Limitations	269
Snapshot Retention	269
Shared Snapshots	269
Encryption	269
Copying Snapshots Across AWS Regions	270
Option Groups	270
Parameter Groups	270
Copying a DB Snapshot	271
Sharing a Snapshot	278
Sharing an Encrypted Snapshot	279
Sharing a Snapshot	281
Point-in-Time Recovery	285
Deleting a Snapshot	287
Deleting a DB Snapshot	287
Tutorial: Restore a DB Instance from a DB Snapshot	289
Prerequisites for Restoring a DB Instance from a DB Snapshot	289

Restoring a DB Instance from a DB Snapshot	290
Modifying a Restored DB Instance	291
Monitoring	294
Overview of Monitoring	295
Monitoring Tools	296
Monitoring with CloudWatch	297
Publishing to CloudWatch Logs	303
Enhanced Monitoring	307
Enhanced Monitoring Availability	307
Differences Between CloudWatch and Enhanced Monitoring Metrics	307
Setting Up for and Enabling Enhanced Monitoring	307
Viewing Enhanced Monitoring	309
Viewing Enhanced Monitoring by Using CloudWatch Logs	313
Performance Insights	319
Enabling Performance Insights	320
Access Control for Performance Insights	324
Using the Performance Insights Dashboard	326
Additional User Interface Features	337
Performance Insights API	338
Metrics Published to CloudWatch	351
Performance Insights Counters	352
Logging Performance Insights Calls by Using AWS CloudTrail	362
Using Amazon RDS Recommendations	365
Responding to Recommendations	367
Using Amazon RDS Event Notification	370
Amazon RDS Event Categories and Event Messages	371
Subscribing to Amazon RDS Event Notification	378
Listing Your Amazon RDS Event Notification Subscriptions	380
Modifying an Amazon RDS Event Notification Subscription	382
Adding a Source Identifier to an Amazon RDS Event Notification Subscription	384
Removing a Source Identifier from an Amazon RDS Event Notification Subscription	385
Listing the Amazon RDS Event Notification Categories	386
Deleting an Amazon RDS Event Notification Subscription	387
Viewing Amazon RDS Events	388
Console	388
AWS CLI	388
API	388
.....	389
Database Log Files	390
Viewing and Listing Database Log Files	390
Downloading a Database Log File	390
Watching a Database Log File	392
Publishing to CloudWatch Logs	392
Reading Log File Contents Using REST	392
MariaDB Database Log Files	394
Microsoft SQL Server Database Log Files	402
MySQL Database Log Files	405
Oracle Database Log Files	413
PostgreSQL Database Log Files	419
Logging Amazon RDS API Calls with AWS CloudTrail	423
Amazon RDS Information in CloudTrail	423
Understanding Amazon RDS Log File Entries	423
Security	427
Data Protection	428
Data Encryption	429
Internetwork Traffic Privacy	444
Identity and Access Management	445

Audience	445
Authenticating With Identities	445
Managing Access Using Policies	447
How Amazon RDS Works with IAM	449
Identity-Based Policy Examples	451
IAM Database Authentication for MySQL and PostgreSQL	461
Troubleshooting	478
Kerberos Authentication	480
Logging and Monitoring	480
Compliance Validation	483
Resilience	484
Backup and Restore	484
Replication	484
Failover	484
Infrastructure Security	485
Security Groups	485
Public Accessibility	485
Security Best Practices	485
Controlling Access with Security Groups	486
DB Security Groups	486
VPC Security Groups	486
DB Security Groups vs. VPC Security Groups	487
Security Group Scenario	487
Creating a VPC Security Group	488
Associating with a DB Instance	488
Deleting DB VPC Security Groups	488
DB Security Groups on EC2-Classic	491
Master User Account Privileges	499
Service-Linked Roles	501
Service-Linked Role Permissions for Amazon RDS	501
Creating a Service-Linked Role for Amazon RDS	503
Editing a Service-Linked Role for Amazon RDS	503
Deleting a Service-Linked Role for Amazon RDS	503
Using Amazon RDS with Amazon VPC	505
Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform	506
Scenarios for Accessing a DB Instance in a VPC	507
Working with a DB Instance in a VPC	514
Updating the VPC for a DB Instance	519
Tutorial: Create an Amazon VPC for Use with a DB Instance	521
Amazon RDS on AWS Outposts (Preview)	526
Prerequisites	527
Limitations	527
Supported DB Instance Classes	527
Creating DB Instances	528
MariaDB on Amazon RDS	532
Common Management Tasks	532
MariaDB Versions	534
Version and Feature Support	535
MariaDB 10.3 Support	535
MariaDB 10.2 Support	535
MariaDB 10.1 Support	536
MariaDB 10.0 Support	536
Features Not Supported	536
Supported Storage Engines	537
MariaDB Security	537
SSL Support	539
Cache Warming	540

Dumping and Loading the Buffer Pool on Demand	541
Database Parameters	541
Common DBA Tasks	541
Local Time Zone	541
Creating a DB Instance Running MariaDB	543
Available Settings	551
Connecting to a DB Instance Running MariaDB	555
Connecting from the mysql Utility	557
Connecting with SSL	557
Maximum MariaDB Connections	558
Related Topics	558
Updating Applications for New SSL/TLS Certificates	560
Determining Whether a Client Requires Certificate Verification in Order to Connect	560
Updating Your Application Trust Store	561
Example Java Code for Establishing SSL Connections	562
Modifying a DB Instance Running MariaDB	564
Available Settings	565
Upgrading the MariaDB DB Engine	573
Overview	573
Upgrading a MariaDB DB Instance	574
Migrating Data from a MySQL DB Snapshot to a MariaDB DB Instance	575
Incompatibilities Between MariaDB and MySQL	575
Console	575
AWS CLI	577
API	578
Working with MariaDB Replication	579
Working with MariaDB Read Replicas	579
Configuring GTID-Based Replication	582
Importing Data into a MariaDB DB Instance	586
Options for MariaDB	587
MariaDB Audit Plugin Support	587
Parameters for MariaDB	590
MariaDB on Amazon RDS SQL Reference	595
mysql.rds_set_external_master_gtid	595
mysql.rds_kill_query_id	597
Microsoft SQL Server on Amazon RDS	598
Common Management Tasks	598
Limits	600
DB Instance Class Support	602
Security	603
Compliance Programs	604
HIPAA	604
SSL Support	605
Version and Feature Support	605
SQL Server 2017 Support	605
SQL Server 2016 Support	605
SQL Server 2014 Support	606
SQL Server 2012 Support on Amazon RDS	607
SQL Server 2008 R2 Deprecated on Amazon RDS	608
MSSQL Version Management	608
Database Engine Patches and Versions	608
Deprecation Schedule	608
CDC Support	609
Features Not Supported and Features with Limited Support	609
Multi-AZ Deployments	610
Using TDE	610
Local Time Zone	610

Supported Time Zones	611
Licensing SQL Server on Amazon RDS	615
Restoring License-Terminated DB Instances	615
Using SQL Server Developer Edition on AWS	615
Related Topics	615
Creating a DB Instance Running SQL Server	616
Available Settings	625
Connecting to a DB Instance Running SQL Server	630
Connecting to Your DB Instance with SSMS	630
Connecting to Your DB Instance with SQL Workbench/J	634
Security Group Considerations	636
Troubleshooting	637
.....	638
Updating Applications for New SSL/TLS Certificates	639
Determining Whether Any Applications are Connecting to Your Microsoft SQL Server DB Instance Using SSL	639
Determining Whether a Client Requires Certificate Verification in Order to Connect	640
Updating Your Application Trust Store	641
Modifying a DB Instance Running SQL Server	643
Available Settings	644
Related Topics	651
Upgrading the SQL Server DB Engine	653
Overview	653
Major Version Upgrades	654
Multi-AZ and In-Memory Optimization Considerations	655
Option and Parameter Group Considerations	655
Testing an Upgrade	655
Upgrading a SQL Server DB Instance	656
Upgrading Deprecated DB Instances Before Support Ends	656
Importing and Exporting SQL Server Databases	657
Limitations and Recommendations	657
Setting Up	658
Using Native Backup and Restore	660
Compressing Backup Files	669
Troubleshooting	670
.....	671
Importing and Exporting SQL Server Data Using Other Methods	672
Multi-AZ for SQL Server	681
Adding Multi-AZ to a SQL Server DB Instance	681
Notes and Recommendations	682
Determining the Location of the Secondary	684
Migrating to Always On AGs	685
Additional Features for SQL Server	686
Using SSL with a SQL Server DB Instance	687
Using Windows Authentication with a SQL Server DB Instance	691
Amazon S3 Integration	700
Options for SQL Server	712
Native Backup and Restore	712
Transparent Data Encryption	714
SQL Server Audit	717
Common DBA Tasks for SQL Server	723
Accessing the tempdb Database	724
Analyzing Your Database Workload with SQL Server Tuning Advisor	726
Collations and Character Sets	728
Determining a Recovery Model	730
Dropping a SQL Server Database	731
Using CDC	731

Renaming a Multi-AZ Database	733
Resetting the db_owner Role Password	734
Restoring License-Terminated DB Instances	734
Transitioning a Database from OFFLINE to ONLINE	735
Using SQL Server Agent	735
Working with SQL Server Logs	736
Working with Trace and Dump Files	737
MySQL on Amazon RDS	739
Common Management Tasks	739
MySQL Versions	742
MySQL Features Not Supported By Amazon RDS	743
Supported Storage Engines	744
MySQL Security	744
Password Validation Plugin	746
SSL Support	746
Using memcached and Other Options with MySQL	748
InnoDB Cache Warming	748
Dumping and Loading the Buffer Pool on Demand	749
Local Time Zone	749
Known Issues and Limitations	750
Creating a DB Instance Running MySQL	751
Available Settings	758
Connecting to a DB Instance Running MySQL	764
Connecting from the MySQL Client	766
Connecting with SSL	766
Connecting from MySQL Workbench	767
Maximum MySQL connections	769
Managing Connections with Amazon RDS Proxy (Preview)	770
.....	796
Updating Applications for New SSL/TLS Certificates	797
Determining Whether any Applications are Connecting to Your MySQL DB Instance Using SSL ...	797
Determining Whether a Client Requires Certificate Verification to Connect	798
Updating Your Application Trust Store	799
Example Java Code for Establishing SSL Connections	800
Modifying a DB Instance Running MySQL	802
Available Settings	803
Upgrading the MySQL DB Engine	812
Overview	812
Major Version Upgrades	812
Testing an Upgrade	815
Upgrading a MySQL DB Instance	816
Upgrading a MySQL Database with Reduced Downtime	816
Upgrading a MySQL DB Snapshot	818
Upgrading a MySQL DB Snapshot	818
AWS CLI	818
RDS API	819
Importing Data into a MySQL DB Instance	820
Overview	820
Importing Data Considerations	822
Restoring a Backup into an Amazon RDS MySQL DB Instance	826
Importing Data from a MySQL or MariaDB DB to an Amazon RDS MySQL or MariaDB DB	
Instance	833
Importing Data to an Amazon RDS MySQL or MariaDB DB Instance with Reduced Downtime	835
Importing Data From Any Source to a MySQL or MariaDB DB Instance	847
Working with MySQL Replication	852
Working with MySQL Read Replicas	852
Using GTID-Based Replication	858

Replication with a MySQL or MariaDB Instance Running External to Amazon RDS	862
Exporting Data From a MySQL DB Instance	869
Prepare an Instance of MySQL External to Amazon RDS	869
Prepare the Replication Source	870
Copy the Database	870
Complete the Export	872
Related Topics	872
Options for MySQL	873
MariaDB Audit Plugin	874
memcached	877
Common DBA Tasks for MySQL	881
Killing a Session or Query	881
Skipping the Current Replication Error	881
Working with InnoDB Tablespaces to Improve Crash Recovery Times	882
Managing the Global Status History	883
Known Issues and Limitations	885
Inconsistent InnoDB Buffer Pool Size	885
Index Merge Optimization Returns Wrong Results	885
Log File Size	886
MySQL Parameter Exceptions for Amazon RDS DB Instances	886
MySQL File Size Limits	887
MySQL Keyring Plugin Not Supported	888
MySQL on Amazon RDS SQL Reference	889
Overview	889
SQL Reference Conventions	890
mysql.rds_set_master_auto_position	890
mysql.rds_set_external_master	891
mysql.rds_set_external_master_with_delay	893
mysql.rds_set_external_master_with_auto_position	895
mysql.rds_reset_external_master	897
mysql.rds_import_binlog_ssl_material	898
mysql.rds_remove_binlog_ssl_material	899
mysql.rds_set_source_delay	900
mysql.rds_start_replication	901
mysql.rds_start_replication_until	901
mysql.rds_start_replication_until_gtid	902
mysql.rds_stop_replication	903
mysql.rds_skip_transaction_with_gtid	904
mysql.rds_skip_repl_error	904
mysql.rds_next_master_log	905
mysql.rds_innodb_buffer_pool_dump_now	907
mysql.rds_innodb_buffer_pool_load_now	907
mysql.rds_innodb_buffer_pool_load_abort	907
mysql.rds_set_configuration	908
mysql.rds_show_configuration	909
mysql.rds_kill	910
mysql.rds_kill_query	910
mysql.rds_rotate_general_log	911
mysql.rds_rotate_slow_log	911
mysql.rds_enable_gsh_collector	911
mysql.rds_set_gsh_collector	912
mysql.rds_disable_gsh_collector	912
mysql.rds_collect_global_status_history	912
mysql.rds_enable_gsh_rotation	912
mysql.rds_set_gsh_rotation	913
mysql.rds_disable_gsh_rotation	913
mysql.rds_rotate_global_status_history	913

Oracle on Amazon RDS	914
Common Management Tasks	914
Licensing	916
License Included	916
Bring Your Own License (BYOL)	916
Licensing Oracle Multi-AZ Deployments	917
Migrating Between Oracle Editions	917
DB Instance Class Support	917
Deprecation of db.t2 DB Instance Classes	919
Deprecated db.m3 and db.r3 DB Instance Classes	919
Deprecated db.m1 and db.m2 DB Instance Classes	920
Security	920
SSL Support	921
Oracle 19c	921
Amazon RDS Parameter Changes for Oracle 19c Version 19.0.0.0	921
Oracle 18c	922
Amazon RDS Parameter Changes for Oracle 18c Version 18.0.0.0	922
Oracle 12c	923
Oracle 12c Version 12.2.0.1	923
Oracle 12c Version 12.1.0.2	926
Oracle Database Feature Support	932
Oracle Database Parameter Support	934
Engine Version Management	934
Deprecation of Oracle 11.2.0.2	934
Deprecation of Oracle 11.2.0.3	935
Deprecation of Oracle 12.1.0.1	935
Using Huge Pages	935
Using utl_http, utl_tcp, and utl_smtp	938
Using OEM, APEX, TDE, and Other Options	940
Using Extended Data Types	940
Enabling Extended Data Types for a New DB Instance	940
Enabling Extended Data Types for an Existing DB Instance	941
Public Synonyms	941
Creating a DB Instance Running Oracle	942
Available Settings	949
Connecting to a DB Instance Running Oracle	955
Finding the Endpoint	955
SQL Developer	957
SQL*Plus	959
Security Group Considerations	960
Dedicated and Shared Server Processes	960
Troubleshooting	961
.....	961
Updating Applications for New SSL/TLS Certificates	962
Determining Whether Any Applications Are Connecting to Your Oracle DB Instance Using SSL ...	962
Updating Your Application Trust Store	963
Example Java Code for Establishing SSL Connections	964
Modifying a DB Instance Running Oracle	966
Available Settings	967
Modifying Oracle sqlnet.ora Parameters	976
Upgrading the Oracle DB Engine	979
Overview	979
Major Version Upgrades	980
Minor Version Upgrades	981
SE2 Upgrade Paths	981
Option and Parameter Group Considerations	982
Testing an Upgrade	982

Upgrading an Oracle DB Instance	983
Upgrading an Oracle DB Snapshot	984
Console	984
AWS CLI	984
RDS API	985
Related Topics	985
Importing Data into Oracle	987
Importing Using Oracle SQL Developer	987
Importing Using Oracle Data Pump	987
Oracle Export/Import Utilities	995
Oracle SQL*Loader	995
Oracle Materialized Views	996
Working with Oracle Read Replicas	998
Configuring Read Replicas for Oracle	998
Read Replica Limitations with Oracle	999
Troubleshooting an Oracle Read Replica Problem	1000
Oracle Character Sets	1001
Options for Oracle	1004
Amazon S3 Integration	1005
Application Express (APEX)	1016
Enterprise Manager	1024
Java Virtual Machine (JVM)	1037
Label Security	1040
Locator	1043
Multimedia	1046
Native Network Encryption (NNE)	1049
Secure Sockets Layer (SSL)	1051
Spatial	1060
SQLT	1063
Statspack	1068
Time Zone	1071
Transparent Data Encryption (TDE)	1074
UTL_MAIL	1076
XML DB	1078
Common DBA Tasks for Oracle	1080
System Tasks	1085
Database Tasks	1095
Log Tasks	1105
RMAN Tasks	1112
Oracle Scheduler Tasks	1128
Miscellaneous Tasks	1131
Using Kerberos Authentication	1134
Setting Up	1135
Managing a DB Instance	1144
Connecting with Kerberos Authentication	1145
Tools and Third-Party Software for Oracle	1147
Setting Up	1147
Using Oracle GoldenGate	1153
Using the Oracle Repository Creation Utility	1165
Installing a Siebel Database on Oracle on Amazon RDS	1170
Oracle Database Engine Release Notes	1173
Oracle Versions 19.0.0, 18.0.0, and 12.2.0.1	1173
Oracle Versions 12.1.0.2 and 11.2.0.4	1173
Database Engine: 19.0.0.0	1174
Database Engine: 18.0.0.0	1176
Database Engine: 12.2.0.1	1182
Database Engine: 12.1.0.2	1195

Database Engine: 11.2.0.4	1225
PostgreSQL on Amazon RDS	1253
Common Management Tasks for PostgreSQL on Amazon RDS	1254
Creating a DB Instance Running PostgreSQL	1257
Available Settings	1265
Connecting to a DB Instance Running the PostgreSQL Database Engine	1270
Using pgAdmin to Connect to a PostgreSQL DB Instance	1270
Using psql to Connect to a PostgreSQL DB Instance	1272
Troubleshooting Connection Issues	1273
Updating Applications for New SSL/TLS Certificates	1275
Determining Whether Applications Are Connecting to PostgreSQL DB Instances Using SSL	1275
Determining Whether a Client Requires Certificate Verification in Order to Connect	1276
Updating Your Application Trust Store	1277
Using SSL/TLS Connections for Different Types of Applications	1278
Modifying a DB Instance Running PostgreSQL	1279
Available Settings	1280
Upgrading the PostgreSQL DB Engine	1289
Overview	1289
Major Version Upgrades	1290
Automatic Minor Version Upgrades	1294
Upgrading a PostgreSQL DB Snapshot	1296
Working with PostgreSQL Read Replicas	1298
Read Replica Configuration with PostgreSQL	1298
Monitoring PostgreSQL Read Replicas	1299
Read Replica Limitations with PostgreSQL	1299
Replication Interruptions with PostgreSQL Read Replicas	1299
Troubleshooting a PostgreSQL Read Replica Problem	1299
Importing Data into PostgreSQL on Amazon RDS	1302
Importing a PostgreSQL Database from an Amazon EC2 Instance	1303
Using the \copy Command to Import Data to a Table on a PostgreSQL DB Instance	1305
Importing S3 Data into RDS for PostgreSQL	1305
Transporting PostgreSQL Databases Between DB Instances	1316
Common DBA Tasks for PostgreSQL	1322
Creating Roles	1322
Managing PostgreSQL Database Access	1323
Working with PostgreSQL Parameters	1323
Working with PostgreSQL Autovacuum	1332
Audit Logging for a PostgreSQL DB Instance	1340
Working with the pgaudit Extension	1341
Working with the pg_repack Extension	1342
Working with PostGIS	1343
Using pgBadger for Log Analysis with PostgreSQL	1345
Viewing the Contents of pg_config	1346
Working with the orafce Extension	1346
Accessing External Data with the postgres_fdw Extension	1347
Using a Custom DNS Server for Outbound Network Access	1348
Restricting Password Management	1349
Using Kerberos Authentication	1351
Overview	1351
Setting Up	1352
Managing a DB Instance	1360
Connecting with Kerberos Authentication	1361
Working with the Database Preview Environment	1362
Features Not Supported in the Preview Environment	1362
PostgreSQL Extensions Supported in the Preview Environment	1362
Creating a New DB Instance in the Preview Environment	1364
PostgreSQL Versions and Extensions	1364

Supported PostgreSQL Database Versions	1365
Supported Features and Extensions	1383
Limits	1414
Limits in Amazon RDS	1414
Naming Constraints in Amazon RDS	1415
File Size Limits in Amazon RDS	1417
MySQL File Size Limits in Amazon RDS	1417
MariaDB File Size Limits in Amazon RDS	1418
Troubleshooting	1420
Cannot Connect to DB Instance	1420
Testing the DB Instance Connection	1420
Troubleshooting Connection Authentication	1421
Security Issues	1421
Error Message "Failed to retrieve account attributes, certain console functions may be impaired."	1421
Resetting the DB Instance Owner Role Password	1422
DB Instance Outage or Reboot	1422
Parameter Changes Not Taking Effect	1423
DB Instance Out of Storage	1423
Insufficient DB Instance Capacity	1424
MySQL Issues	1425
Index Merge Optimization Returns Wrong Results	1425
Diagnosing and Resolving Lag Between Read Replicas	1425
Diagnosing and Resolving a MySQL or MariaDB Read Replication Failure	1427
Creating Triggers with Binary Logging Enabled Requires SUPER Privilege	1428
Diagnosing and Resolving Point-In-Time Restore Failures	1429
Slave Down or Disabled Error	1430
Read Replica Create Fails or Replication Breaks With Fatal Error 1236	1430
Oracle GoldenGate Issues	1430
Retaining Logs for Sufficient Time	1430
Cannot Connect to SQL Server DB Instance	1431
Cannot Connect to PostgreSQL DB Instance	1431
Cannot set backup retention to 0	1432
Amazon RDS API Reference	1433
Using the Query API	1433
Query Parameters	1433
Query Request Authentication	1433
Troubleshooting Applications	1434
Retrieving Errors	1434
Troubleshooting Tips	1434
Document History	1435
Earlier Updates	1453
AWS Glossary	1474

What Is Amazon Relational Database Service (Amazon RDS)?

Amazon Relational Database Service (Amazon RDS) is a web service that makes it easier to set up, operate, and scale a relational database in the AWS Cloud. It provides cost-efficient, resizable capacity for an industry-standard relational database and manages common database administration tasks.

Note

This guide covers Amazon RDS database engines other than Amazon Aurora. For information about using Amazon Aurora, see the [Amazon Aurora User Guide](#).

This guide covers using Amazon RDS in the AWS Cloud. For information about using Amazon RDS in on-premises VMware environments, see the [Amazon RDS on VMware User Guide](#).

Overview of Amazon RDS

Why do you want a managed relational database service? Because Amazon RDS takes over many of the difficult or tedious management tasks of a relational database:

- When you buy a server, you get CPU, memory, storage, and IOPS, all bundled together. With Amazon RDS, these are split apart so that you can scale them independently. If you need more CPU, less IOPS, or more storage, you can easily allocate them.
- Amazon RDS manages backups, software patching, automatic failure detection, and recovery.
- To deliver a managed service experience, Amazon RDS doesn't provide shell access to DB instances. It also restricts access to certain system procedures and tables that require advanced privileges.
- You can have automated backups performed when you need them, or manually create your own backup snapshot. You can use these backups to restore a database. The Amazon RDS restore process works reliably and efficiently.
- You can get high availability with a primary instance and a synchronous secondary instance that you can fail over to when problems occur. You can also use MySQL, MariaDB, or PostgreSQL Read Replicas to increase read scaling.
- You can use the database products you are already familiar with: MySQL, MariaDB, PostgreSQL, Oracle, Microsoft SQL Server.
- In addition to the security in your database package, you can help control who can access your RDS databases by using AWS Identity and Access Management (IAM) to define users and permissions. You can also help protect your databases by putting them in a virtual private cloud.

If you are new to AWS products and services, begin learning more with the following resources:

- For an overview of all AWS products, see [What is Cloud Computing?](#)
- Amazon Web Services provides a number of database services. For guidance on which service is best for your environment, see [Running Databases on AWS](#).

DB Instances

The basic building block of Amazon RDS is the DB instance. A *DB instance* is an isolated database environment in the AWS Cloud. Your DB instance can contain multiple user-created databases. You can access your DB instance by using the same tools and applications that you use with a standalone

database instance. You can create and modify a DB instance by using the AWS Command Line Interface, the Amazon RDS API, or the AWS Management Console.

Each DB instance runs a *DB engine*. Amazon RDS currently supports the MySQL, MariaDB, PostgreSQL, Oracle, and Microsoft SQL Server DB engines. Each DB engine has its own supported features, and each version of a DB engine may include specific features. Additionally, each DB engine has a set of parameters in a DB parameter group that control the behavior of the databases that it manages.

The computation and memory capacity of a DB instance is determined by its *DB instance class*. You can select the DB instance that best meets your needs. If your needs change over time, you can change DB instances. For information, see [Choosing the DB Instance Class \(p. 107\)](#).

Note

For pricing information on DB instance classes, see the Pricing section of the [Amazon RDS](#) product page.

DB instance storage comes in three types: Magnetic, General Purpose (SSD), and Provisioned IOPS (PIOPS). They differ in performance characteristics and price, allowing you to tailor your storage performance and cost to the needs of your database. Each DB instance has minimum and maximum storage requirements depending on the storage type and the database engine it supports. It's important to have sufficient storage so that your databases have room to grow. Also, sufficient storage makes sure that features for the DB engine have room to write content or log entries. For more information, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

You can run a DB instance on a virtual private cloud (VPC) using the Amazon Virtual Private Cloud (Amazon VPC) service. When you use a VPC, you have control over your virtual networking environment. You can choose your own IP address range, create subnets, and configure routing and access control lists. The basic functionality of Amazon RDS is the same whether it's running in a VPC or not. Amazon RDS manages backups, software patching, automatic failure detection, and recovery. There's no additional cost to run your DB instance in a VPC. For more information on using Amazon VPC with RDS, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).

Amazon RDS uses Network Time Protocol (NTP) to synchronize the time on DB Instances.

AWS Regions and Availability Zones

Amazon cloud computing resources are housed in highly available data center facilities in different areas of the world (for example, North America, Europe, or Asia). Each data center location is called an AWS Region.

Each AWS Region contains multiple distinct locations called Availability Zones, or AZs. Each Availability Zone is engineered to be isolated from failures in other Availability Zones. Each is engineered to provide inexpensive, low-latency network connectivity to other Availability Zones in the same AWS Region. By launching instances in separate Availability Zones, you can protect your applications from the failure of a single location. For more information, see [Regions and Availability Zones \(p. 132\)](#).

You can run your DB instance in several Availability Zones, an option called a Multi-AZ deployment. When you choose this option, Amazon automatically provisions and maintains a secondary standby DB instance in a different Availability Zone. Your primary DB instance is synchronously replicated across Availability Zones to the secondary instance. This approach helps provide data redundancy and failover support, eliminate I/O freezes, and minimize latency spikes during system backups. For more information, see [High Availability \(Multi-AZ\) for Amazon RDS \(p. 144\)](#).

Security

A *security group* controls the access to a DB instance. It does so by allowing access to IP address ranges or Amazon EC2 instances that you specify.

Amazon RDS uses DB security groups, VPC security groups, and EC2 security groups. In simple terms, you can look at it like this:

- A DB security group controls access to a DB instance that is not in a VPC.
- A VPC security group controls access to a DB instance inside a VPC.
- An EC2 security group controls access to an EC2 instance and can be used with a DB instance.

For more information about security groups, see [Security in Amazon RDS \(p. 427\)](#).

Monitoring an Amazon RDS DB Instance

There are several ways that you can track the performance and health of a DB instance. You can use the Amazon CloudWatch service to monitor the performance and health of a DB instance. CloudWatch performance charts are shown in the Amazon RDS console. You can also subscribe to Amazon RDS events to be notified about changes to a DB instance, DB snapshot, DB parameter group, or DB security group. For more information, see [Monitoring Amazon RDS \(p. 294\)](#).

How to Work with Amazon RDS

There are several ways that you can interact with Amazon RDS.

AWS Management Console

The AWS Management Console is a simple web-based user interface. You can manage your DB instances from the console with no programming required. To access the Amazon RDS console, sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.

Command Line Interface

You can use the AWS Command Line Interface (AWS CLI) to access the Amazon RDS API interactively. To install the AWS CLI, see [Installing the AWS Command Line Interface](#). To begin using the AWS CLI for RDS, see [AWS Command Line Interface Reference for Amazon RDS](#).

Programming with Amazon RDS

If you are a developer, you can access the Amazon RDS programmatically. For more information, see [Amazon RDS Application Programming Interface \(API\) Reference \(p. 1433\)](#).

For application development, we recommend that you use one of the AWS Software Development Kits (SDKs). The AWS SDKs handle low-level details such as authentication, retry logic, and error handling, so that you can focus on your application logic. AWS SDKs are available for a wide variety of languages. For more information, see [Tools for Amazon Web Services](#).

AWS also provides libraries, sample code, tutorials, and other resources to help you get started more easily. For more information, see [Sample Code & Libraries](#).

How You Are Charged for Amazon RDS

When you use Amazon RDS, you can choose to use on-demand DB instances or reserved DB instances. For more information, see [DB Instance Billing for Amazon RDS \(p. 239\)](#).

For Amazon RDS pricing information, see the [Amazon RDS product page](#).

What's Next?

The preceding section introduced you to the basic infrastructure components that RDS offers. What should you do next?

Getting Started

Create a DB instance using instructions in [Getting Started with Amazon RDS \(p. 11\)](#).

Database Engine–Specific Topics

You can review information specific to a particular DB engine in the following sections:

- [MariaDB on Amazon RDS \(p. 532\)](#)
- [Microsoft SQL Server on Amazon RDS \(p. 598\)](#)
- [MySQL on Amazon RDS \(p. 739\)](#)
- [Oracle on Amazon RDS \(p. 914\)](#)
- [PostgreSQL on Amazon RDS \(p. 1253\)](#)

Setting Up for Amazon RDS

Following, you can find how to set up Amazon Relational Database Service (Amazon RDS) for the first time. If you already have an AWS account, know your Amazon RDS requirements, and prefer to use the defaults for IAM and VPC security groups, skip ahead to [Getting Started \(p. 4\)](#).

A couple things you should know about Amazon Web Services (AWS):

- When you sign up for AWS, your AWS account automatically has access to all services in AWS, including Amazon RDS. However, you are charged only for the services that you use.
- With Amazon RDS, you pay only for the RDS instances that are active. The Amazon RDS DB instance that you create is live (not running in a sandbox). You incur the standard Amazon RDS usage fees for the instance until you terminate it. For more information about Amazon RDS usage rates, see the [Amazon RDS product page](#).

Topics

- [Sign Up for AWS \(p. 5\)](#)
- [Create an IAM User \(p. 5\)](#)
- [Determine Requirements \(p. 7\)](#)
- [Provide Access to Your DB Instance in Your VPC by Creating a Security Group \(p. 8\)](#)

Sign Up for AWS

If you have an AWS account already, skip to the next section, [Create an IAM User \(p. 5\)](#).

If you don't have an AWS account, you can use the following procedure to create one. If you are a new AWS customer, you can get started with Amazon RDS for free; for more information, see [AWS Free Usage Tier](#).

To create a new AWS account

1. Open <https://portal.aws.amazon.com/billing/signup>.
2. Follow the online instructions.

Part of the sign-up procedure involves receiving a phone call and entering a verification code on the phone keypad.

Create an IAM User

After you create an AWS account and successfully connect to the AWS Management Console, you can create an AWS Identity and Access Management (IAM) user. Instead of signing in with your AWS root account, we recommend that you use an IAM administrative user with Amazon RDS.

One way to do this is to create a new IAM user and grant it administrator permissions. Alternatively, you can add an existing IAM user to an IAM group with Amazon RDS administrative permissions. You can then access AWS from a special URL using the credentials for the IAM user.

If you signed up for AWS but haven't created an IAM user for yourself, you can create one using the IAM console.

To create an administrator user for yourself and add the user to an administrators group (console)

1. Use your AWS account email address and password to sign in as the [AWS account root user](#) to the IAM console at <https://console.aws.amazon.com/iam/>.

Note

We strongly recommend that you adhere to the best practice of using the **Administrator** IAM user below and securely lock away the root user credentials. Sign in as the root user only to perform a few [account and service management tasks](#).

2. In the navigation pane, choose **Users** and then choose **Add user**.
3. For **User name**, enter **Administrator**.
4. Select the check box next to **AWS Management Console access**. Then select **Custom password**, and then enter your new password in the text box.
5. (Optional) By default, AWS requires the new user to create a new password when first signing in. You can clear the check box next to **User must create a new password at next sign-in** to allow the new user to reset their password after they sign in.
6. Choose **Next: Permissions**.
7. Under **Set permissions**, choose **Add user to group**.
8. Choose **Create group**.
9. In the **Create group** dialog box, for **Group name** enter **Administrators**.
10. Choose **Filter policies**, and then select **AWS managed -job function** to filter the table contents.
11. In the policy list, select the check box for **AdministratorAccess**. Then choose **Create group**.

Note

You must activate IAM user and role access to Billing before you can use the **AdministratorAccess** permissions to access the AWS Billing and Cost Management console. To do this, follow the instructions in [step 1 of the tutorial about delegating access to the billing console](#).

12. Back in the list of groups, select the check box for your new group. Choose **Refresh** if necessary to see the group in the list.
13. Choose **Next: Tags**.
14. (Optional) Add metadata to the user by attaching tags as key-value pairs. For more information about using tags in IAM, see [Tagging IAM Entities](#) in the *IAM User Guide*.
15. Choose **Next: Review** to see the list of group memberships to be added to the new user. When you are ready to proceed, choose **Create user**.

You can use this same process to create more groups and users and to give your users access to your AWS account resources. To learn about using policies that restrict user permissions to specific AWS resources, see [Access Management](#) and [Example Policies](#).

To sign in as the new IAM user, first sign out of the AWS Management Console. Then use the following URL, where **your_aws_account_id** is your AWS account number without the hyphens. For example, if your AWS account number is 1234-5678-9012, your AWS account ID is 123456789012.

`https://your_aws_account_id.signin.aws.amazon.com/console/`

Type the IAM user name and password that you just created. When you're signed in, the navigation bar displays "your_user_name @ your_aws_account_id".

If you don't want the URL for your sign-in page to contain your AWS account ID, you can create an account alias. From the IAM dashboard, choose **Customize** and type an alias, such as your company name. To sign in after you create an account alias, use the following URL.

`https://your_account_alias.signin.aws.amazon.com/console/`

To verify the sign-in link for IAM users for your account, open the IAM console and check under **AWS Account Alias** on the dashboard.

You can also create access keys for your AWS account. These access keys can be used to access AWS through the AWS Command Line Interface (AWS CLI) or through the Amazon RDS API. For more information, see [Managing Access Keys for Your AWS Account](#), [Installing the AWS Command Line Interface](#), and the [Amazon RDS API Reference](#).

Determine Requirements

The basic building block of Amazon RDS is the DB instance. In a DB instance, you create your databases. A DB instance provides a network address called an *endpoint*. Your applications use this endpoint to connect to your DB instance. When you create a DB instance, you specify details like storage, memory, database engine and version, network configuration, security, and maintenance periods. You control network access to a DB instance through a security group.

Before you create a DB instance and a security group, you must know your DB instance and network needs. Here are some important things to consider:

- **Resource requirements** – What are the memory and processor requirements for your application or service? You use these settings to help you determine what DB instance class to use. For specifications about DB instance classes, see [Choosing the DB Instance Class \(p. 107\)](#).
- **VPC, subnet, and security group** – Your DB instance is most likely in a virtual private cloud (VPC). To connect to your DB instance, you need to set up security group rules. These rules are set up differently depending on what kind of VPC you use and how you use it: in a default VPC, in a user-defined VPC, or outside of a VPC.

Note

Some legacy accounts don't use a VPC. If you are accessing a new AWS Region or you are a new RDS user (after 2013), you are most likely creating a DB instance inside a VPC.

For information on how to determine if your account has a default VPC in a particular AWS Region, see [Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform \(p. 506\)](#).

The following list describes the rules for each VPC option:

- **Default VPC** – If your AWS account has a default VPC in the current AWS Region, that VPC is configured to support DB instances. If you specify the default VPC when you create the DB instance, do the following:
 - Create a *VPC security group* that authorizes connections from the application or service to the Amazon RDS DB instance with the database. Use the [Amazon EC2 API](#) or the **Security Group** option on the VPC console to create VPC security groups. For information, see [Step 4: Create a VPC Security Group \(p. 518\)](#).
 - Specify the default DB subnet group. If this is the first DB instance you have created in this AWS Region, Amazon RDS creates the default DB subnet group when it creates the DB instance.
- **User-defined VPC** – If you want to specify a user-defined VPC when you create a DB instance, be aware of the following:
 - Make sure to create a *VPC security group* that authorizes connections from the application or service to the Amazon RDS DB instance with the database. Use the [Amazon EC2 API](#) or the **Security Group** option on the VPC console to create VPC security groups. For information, see [Step 4: Create a VPC Security Group \(p. 518\)](#).
 - The VPC must meet certain requirements in order to host DB instances, such as having at least two subnets, each in a separate availability zone. For information, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).

- Make sure to specify a DB subnet group that defines which subnets in that VPC can be used by the DB instance. For information, see the DB subnet group section in [Working with a DB Instance in a VPC \(p. 514\)](#).
- **No VPC** – If your AWS account doesn't have a default VPC and you don't specify a user-defined VPC, create a DB security group. A *DB security group* authorizes connections from the devices and Amazon RDS instances running the applications or utilities to access the databases in the DB instance. For more information, see [Working with DB Security Groups \(EC2-Classic Platform\) \(p. 491\)](#).
- **High availability:** Do you need failover support? On Amazon RDS, a Multi-AZ deployment creates a primary DB instance and a secondary standby DB instance in another Availability Zone for failover support. We recommend Multi-AZ deployments for production workloads to maintain high availability. For development and test purposes, you can use a deployment that isn't Multi-AZ. For more information, see [High Availability \(Multi-AZ\) for Amazon RDS \(p. 144\)](#).
- **IAM policies:** Does your AWS account have policies that grant the permissions needed to perform Amazon RDS operations? If you are connecting to AWS using IAM credentials, your IAM account must have IAM policies that grant the permissions required to perform Amazon RDS operations. For more information, see [Identity and Access Management in Amazon RDS \(p. 445\)](#).
- **Open ports:** What TCP/IP port does your database listen on? The firewall at some companies might block connections to the default port for your database engine. If your company firewall blocks the default port, choose another port for the new DB instance. When you create a DB instance that listens on a port you specify, you can change the port by modifying the DB instance.
- **AWS Region:** What AWS Region do you want your database in? Having your database in close proximity to your application or web service can reduce network latency.
- **DB disk subsystem:** What are your storage requirements? Amazon RDS provides three storage types:
 - Magnetic (Standard Storage)
 - General Purpose (SSD)
 - Provisioned IOPS (PIOPS)

Magnetic storage offers cost-effective storage that is ideal for applications with light or burst I/O requirements. General purpose, SSD-backed storage, also called *gp2*, can provide faster access than disk-based storage. Provisioned IOPS storage is designed to meet the needs of I/O-intensive workloads, particularly database workloads, which are sensitive to storage performance and consistency in random access I/O throughput. For more information on Amazon RDS storage, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

When you have the information you need to create the security group and the DB instance, continue to the next step.

Provide Access to Your DB Instance in Your VPC by Creating a Security Group

VPC security groups provide access to DB instances in a VPC. They act as a firewall for the associated DB instance, controlling both inbound and outbound traffic at the instance level. DB instances are created by default with a firewall and a default security group that protect the DB instance.

Before you can connect to your DB instance, you must add rules to security group that enable you to connect. Use your network and configuration information to create rules to allow access to your DB instance.

Note

If your legacy DB instance was created before March 2013 and isn't in a VPC, it might not have associated security groups. If your DB instance was created after this date, it might be inside a default VPC.

For example, suppose that you have an application that accesses a database on your DB instance in a VPC. In this case, you must add a custom TCP rule that specifies the port range and IP addresses that your application uses to access the database. If you have an application on an Amazon EC2 instance, you can use the VPC or EC2 security group that you set up for the Amazon EC2 instance.

To create a VPC security group

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc>.
2. In the top right corner of the AWS Management Console, choose the AWS Region where you want to create your VPC security group and DB instance. In the list of Amazon VPC resources for that AWS Region, you should see at least one VPC and several subnets. If you don't, you don't have a default VPC in that AWS Region.
3. In the navigation pane, choose **Security Groups**.
4. Choose **Create Security Group**.
5. In the **Create Security Group** window, type **Name tag**, **Group name**, and **Description** values for your security group. For **VPC**, choose the VPC that you want to create your DB instance in. Choose **Yes, Create**.
6. The VPC security group that you created should still be selected. If not, locate it in the list, and choose it. The details pane at the bottom of the console window displays the details for the security group, and tabs for working with inbound and outbound rules. Choose the **Inbound Rules** tab.
7. On the **Inbound Rules** tab, choose **Edit**.
 - a. For **Type**, choose **Custom TCP Rule**.
 - b. For **Port Range**, type the port value to use for your DB instance.
 - c. For **Source**, choose a security group name or type the IP address range (CIDR value) from where you access the instance.
8. Choose **Add another rule** if you need to add more IP addresses or different port ranges.
9. (Optional) Use the **Outbound Rules** tab to add rules for outbound traffic. By default, all outbound traffic is allowed.

You can use the VPC security group that you just created as the security group for your DB instance when you create it. If your DB instance isn't going to be in a VPC, see [Working with DB Security Groups \(EC2-Classic Platform\) \(p. 491\)](#) to create a DB security group to use when you create your DB instance.

Note

If you use a default VPC, a default subnet group spanning all of the VPC's subnets is created for you. When you create a DB instance, you can select the default VPC and use **default** for **DB Subnet Group**.

Once you have completed the setup requirements, you can launch a DB instance using your requirements and security group. For information on creating a DB instance, see the relevant documentation in the following table.

Database Engine	Documentation
MariaDB	Creating a MariaDB DB Instance and Connecting to a Database on a MariaDB DB Instance (p. 11)
Microsoft SQL Server	Creating a Microsoft SQL Server DB Instance and Connecting to a DB Instance (p. 23)
MySQL	Creating a MySQL DB Instance and Connecting to a Database on a MySQL DB Instance (p. 33)

Database Engine	Documentation
Oracle	Creating an Oracle DB Instance and Connecting to a Database on an Oracle DB Instance (p. 45)
PostgreSQL	Creating a PostgreSQL DB Instance and Connecting to a Database on a PostgreSQL DB Instance (p. 55)

Getting Started with Amazon RDS

In the following section, you can find out how to create and connect to a DB instance using Amazon Relational Database Service (Amazon RDS). You can create a DB instance that uses MariaDB, MySQL, Microsoft SQL Server, Oracle, or PostgreSQL.

Important

Before you can create or connect to a DB instance, you must complete the tasks in [Setting Up for Amazon RDS \(p. 5\)](#).

Creating a DB instance and connecting to a database on a DB instance is slightly different for each of the DB engines. Choose the DB engine following that you want to use for detailed information on creating and connecting to the DB instance. After you have created and connected to your DB instance, there are instructions to help you delete the DB instance.

Topics

- [Creating a MariaDB DB Instance and Connecting to a Database on a MariaDB DB Instance \(p. 11\)](#)
- [Creating a Microsoft SQL Server DB Instance and Connecting to a DB Instance \(p. 23\)](#)
- [Creating a MySQL DB Instance and Connecting to a Database on a MySQL DB Instance \(p. 33\)](#)
- [Creating an Oracle DB Instance and Connecting to a Database on an Oracle DB Instance \(p. 45\)](#)
- [Creating a PostgreSQL DB Instance and Connecting to a Database on a PostgreSQL DB Instance \(p. 55\)](#)
- [Tutorial: Create a Web Server and an Amazon RDS Database \(p. 69\)](#)

Creating a MariaDB DB Instance and Connecting to a Database on a MariaDB DB Instance

The easiest way to create a MariaDB DB instance is to use the Amazon RDS console. After you create the DB instance, you can use command line tools such as mysql or standard graphical tools such as HeidiSQL to connect to a database on the DB instance.

Important

Before you can create or connect to a DB instance, you must complete the tasks in [Setting Up for Amazon RDS \(p. 5\)](#).

Topics

- [Creating a MariaDB Instance \(p. 11\)](#)
- [Connecting to a Database on a DB Instance Running the MariaDB Database Engine \(p. 20\)](#)
- [Deleting a DB Instance \(p. 22\)](#)

Creating a MariaDB Instance

The basic building block of Amazon RDS is the DB instance. This environment is where you run your MariaDB databases.

Note

A new console interface is available for database creation. Choose either the **New Console** or the **Original Console** instructions based on the console that you are using. The **New Console** instructions are open by default.

New Console

You can create a DB instance running MariaDB with the AWS Management Console with **Easy Create** enabled or not enabled. With **Easy Create** enabled, you specify only the DB engine type, DB instance size, and DB instance identifier. **Easy Create** uses the default setting for other configuration options. With **Easy Create** not enabled, you specify more configuration options when you create a database, including ones for availability, security, backups, and maintenance.

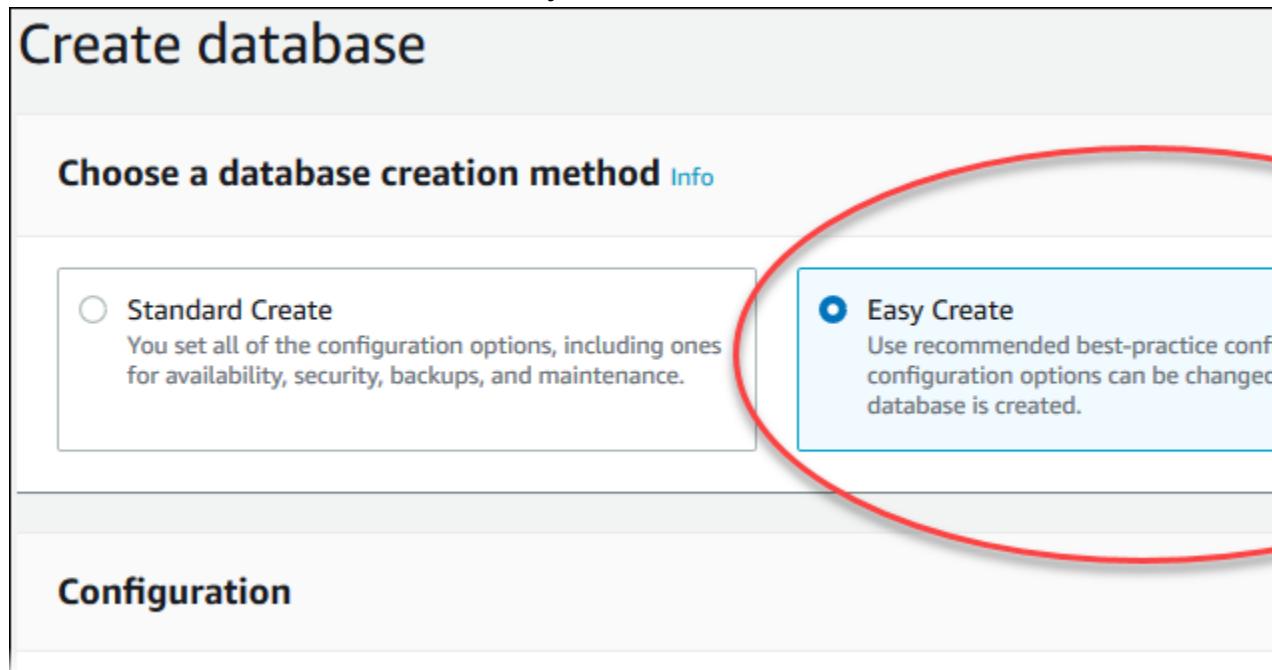
In this example, you use **Easy Create** to create a DB instance running the MariaDB database engine with a db.t2.micro DB instance class.

Note

For information about creating a MariaDB DB instance with **Easy Create** not enabled, see [Creating a DB Instance Running the MariaDB Database Engine \(p. 543\)](#).

To create a MariaDB DB instance with **Easy Create** enabled

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the Amazon RDS console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
4. Choose **Create database** and make sure that **Easy Create** is chosen.



5. In **Configuration**, choose **MariaDB**.
6. For **DB instance size**, choose **Free tier**.
7. For **DB instance identifier**, enter a name for the DB instance, or leave the default name.
8. For **Master username**, enter a name for the master user, or leave the default name.

The **Create database** page should look similar to the following image.

Create database

Choose a database creation method [Info](#)

Standard Create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy Create

Use recommended best-practice configuration options can be changed after the database is created.

Configuration

Engine type [Info](#)

Amazon Aurora



MySQL



MariaDB



PostgreSQL



Oracle

ORACLE®

Microsoft SQL Server



DB instance size

Production

db.r4.xlarge
4 vCPUs
30.5 GiB RAM
500 GiB

Dev/Test

db.r4.large
2 vCPUs
15.25 GiB RAM
100 GiB

Free tier

db.t2.micro
1 vCPUs
1 GiB RAM
20 GiB

DB instance identifier

Type a name for your DB instance. The name must be unique across all DB instances owned by your AWS account in this Region.

13

database-1

The DB instance identifier is case-insensitive, but is stored as all lowercase (as in "mydbinstance"). Constraints: 1 to 63 characters, lowercase (A to Z), SQL Server: First character must be letter. Constraints to ensure the identifier is unique across all DB instances owned by your AWS account in this Region.

9. To use an automatically generated master password for the DB instance, make sure that the **Auto generate a password** check box is chosen.

To enter your master password, clear the **Auto generate a password** check box, and then enter the same password in **Master password** and **Confirm password**.

10. (Optional) Open **View default settings for Easy create**.

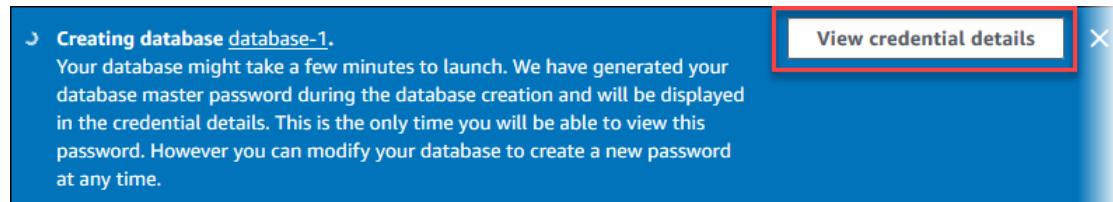
▼ View default settings for Easy create		
Configuration	Value	Editable after database is created
Database Location	Regional	No
Database Features	provisioned	No
Automatic Backups	Enabled	No

You can examine the default settings used when **Easy Create** is enabled. If you want to change one or more settings during database creation, choose **Standard Create** to set them. The **Editable after database creation** column shows which options you can change after database creation. To change a setting with **No** in that column, use **Standard Create**. For settings with **Yes** in that column, you can either use **Standard Create** or modify the DB instance after it's created to change the setting.

11. Choose **Create database**.

If you chose to use an automatically generated password, the **View credential details** button appears on the **Databases** page.

To view the master user name and password for the DB instance, choose **View credential details**.



To connect to the DB instance as the master user, use the user name and password that appear.

Important

You can't view the master user password again. If you don't record it, you might have to change it. If you need to change the master user password after the DB instance is available, you can modify the DB instance to do so. For more information about modifying a DB instance, see [Modifying a DB Instance Running the MariaDB Database Engine \(p. 564\)](#).

12. For **Databases**, choose the name of the new MariaDB DB instance.

On the RDS console, the details for new DB instance appear. The DB instance has a status of **creating** until the DB instance is ready to use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and the amount of storage, it can take up to 20 minutes before the new instance is available.

The screenshot shows the 'database-1' instance details in the Amazon RDS console. The 'Info' column displays 'Creating' with a red circle around it, indicating the status of the database instance.

DB identifier	CPU	Info	Class
database-1		Creating	db.t2.micro
Role	Current activity	Engine	Region & AZ
Instance		MariaDB	-

Below the table, there are tabs for Connectivity & security, Monitoring, Logs & events, Configuration, Maintenance & backups, and Tags.

Original Console

In this example, you create a DB instance running the MariaDB database engine called *mariadb-instance1*, with a *db.t2.small* DB instance class, 20 GiB of storage, and automated backups enabled with a retention period of one day.

To create a MariaDB DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the Amazon RDS console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
If the navigation pane is closed, choose the menu icon at the upper left to open it.
4. Choose **Create database**. The **Select engine** page opens.

Select engine

Engine options

Amazon Aurora **Amazon Aurora**

MySQL 

MariaDB 

PostgreSQL 

Oracle **ORACLE**

Microsoft SQL Server 

MariaDB
MariaDB Community Edition is a MySQL-compatible database with strong support from the open source community, and extra features and performance optimizations.

- Supports database size up to 16 TB.
- Instances offer up to 32 vCPUs and 244 GiB Memory.
- Supports automated backup and point-in-time recovery.
- Supports cross-region read replicas.
- Supports global transaction ID (GTID) and thread pooling.
- Developed and supported by the MariaDB open source community.

Only enable options eligible for RDS Free Usage Tier [info](#)

[Cancel](#) **Next**

5. Choose **MariaDB**, and then choose **Next**.
6. The **Choose use case** page asks if you plan to use the DB instance you are creating for production. Because this is an example instance, choose **Dev/Test - MariaDB**. Then choose **Next**.

Note

If you create a production instance, you typically choose **Production - MariaDB** on this page to enable the failover option Multi-AZ and the Provisioned IOPS storage option.

7. On the **Specify DB details** page, specify your DB instance information. The following table shows settings for an example DB instance. When the settings are as you want them, choose **Next**.

For This Parameter	Do This
License model	Choose the default, general-public-license , to use the GNU General Public License, version 2 for MariaDB. MariaDB has only one license model.
DB engine version	Choose the version of MariaDB that you want to use.

For This Parameter	Do This
DB instance class	Choose db.t2.small for a configuration that equates to 2 GiB memory, 1 ECU (1 virtual core with 1 ECU), 64-bit platform, and moderate I/O capacity.
Multi-AZ deployment	<p>Choose Create replica in a different zone to have a standby replica of your DB instance created in another Availability Zone for failover support. We recommend Multi-AZ for production workloads to maintain high availability. For development and testing, you can choose No.</p> <p>For more information, see High Availability (Multi-AZ) for Amazon RDS (p. 144).</p>
Storage type	Choose the storage type General Purpose (SSD) . For more information about storage, see Amazon RDS DB Instance Storage (p. 135) .
Allocated storage	Enter 20 to allocate 20 GiB of storage for your database. In some cases, allocating a higher amount of storage for your DB instance than the size of your database can improve I/O performance. For more information about storage allocation, see Amazon RDS Features .
DB instance identifier	Enter a name for the DB instance that is unique for your account in the AWS Region you chose. You can add some intelligence to the name, such as including the AWS Region and DB engine you chose, for example mariadb-instance1 .
Master username	Enter a name using 1–16 alphanumeric characters to use as the master user name to log on to your DB instance. You use this user name to log on to your database on the DB instance for the first time.
Master password and Confirm password	Enter a password that contains from 8–41 printable ASCII characters (excluding /, ", and @) for your master user password. You use this password with the user name when you log on to your database. Enter the password again for Confirm Password .

Specify DB details

Instance specifications
Estimate your monthly costs for the DB Instance using the [AWS Simple Monthly Calculator](#) 

DB engine
MariaDB Community Edition

License model [Info](#)
▼

DB engine version [Info](#)
▼

Free tier
The Amazon RDS Free Tier provides a single db.t2.micro instance as well as up to 20 GiB of storage, allowing new AWS customers to gain hands-on experience with Amazon RDS. Learn more about the RDS Free Tier and the instance restrictions [here](#).
 Only enable options eligible for RDS Free Usage Tier [Info](#)

DB instance class [Info](#)
▼

8. On the **Configure advanced settings** page, provide additional information that RDS needs to launch the MariaDB DB instance. The table shows settings for an example DB instance. Specify your DB instance information, then choose **Create database**.

For This Parameter	Do This
Virtual Private Cloud (VPC)	Choose the name of the Amazon Virtual Private Cloud (Amazon VPC) to host your MariaDB DB instance. For more information about using VPC, see Amazon Virtual Private Cloud VPCs and Amazon RDS (p. 505) .
Subnet group	Choose Create new DB subnet group .
Public accessibility	Choose Yes .
Availability zone	Determine if you want to specify a particular availability zone. For more information about Availability Zones, see Regions and Availability Zones (p. 132) .
VPC security groups	Choose Create new VPC security group .
Database name	Enter a name for your default database that is 1–64 alphanumeric characters. If you don't provide a name,

For This Parameter	Do This
	Amazon RDS doesn't automatically create a database on the DB instance you are creating. To create additional databases, connect to the DB instance and use the SQL command <code>CREATE DATABASE</code> . For more information about connecting to the DB instance, see Connecting to a DB Instance Running the MariaDB Database Engine (p. 555) .
Database port	Keep the default value of 3306 unless you have a specific port you want to access the database through. MariaDB installations default to port 3306.
DB parameter group	Accept the default value of default.mariadb10.0 unless you created your own DB parameter group. For more information about parameter groups, see Working with DB Parameter Groups (p. 211) .
Option group	Accept the default value.
Copy tags to snapshots	Choose this option to have any DB instance tags copied to a DB snapshot when you create a snapshot. For more information, see Tagging Amazon RDS Resources (p. 177) .
Encryption	Choose Disable encryption . Note You usually choose Enable encryption for production instances to enable encryption at rest for this DB instance. For more information, see Encrypting Amazon RDS Resources (p. 429) .
Backup retention period	Set the number of days you want automatic backups of your database to be retained. For testing purposes, you can set this value to 1 day .
Backup window	Unless you have a specific time that you want to have your database back up, use the default of No Preference .
Enhanced Monitoring	Unless you want to enable gathering metrics in real time for the operating system that your DB instance runs on, use the default of Disable enhanced monitoring .
Log exports	Select General log . For more information, see MariaDB Database Log Files (p. 394) .
Auto minor version upgrade	Choose Enable auto minor version upgrade to enable your DB instance to receive preferred minor DB engine version upgrades automatically when they become available.
Maintenance window	Choose the 30-minute window in which pending modifications to your DB instance are applied. If the time period doesn't matter, choose No preference .

9. Choose **Create database**.
10. Choose **View DB instance details**.

On the RDS console, the details for new DB instance appear. The DB instance has a status of **creating** until the DB instance is ready to use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and the amount of storage, it can take up to 20 minutes before the new instance is available.

The screenshot shows the 'mariadb-instance1' summary page. The 'Summary' section displays the following details:

Engine	DB instance class info	DB instance status
MariaDB 10.1.26	db.t2.small	creating

Below the summary, there is a 'CloudWatch (17)' section with a legend indicating the instance name 'mariadb-instance1'. A search bar is also present in this section.

Connecting to a Database on a DB Instance Running the MariaDB Database Engine

After Amazon RDS provisions your DB instance, you can use any standard SQL client application to connect to a database on the DB instance. In this example, you connect to a database on a MariaDB DB instance using the mysql command-line tool. One GUI-based application you can use to connect is HeidiSQL. For more information, see the [Download HeidiSQL](#) page. For more information on using MariaDB, see the [MariaDB documentation](#).

To connect to a database on a DB instance using the mysql command-line tool

1. Find the endpoint (DNS name) and port number for your DB instance.
 - a. Open the RDS console and then choose **Databases** to display a list of your DB instances.
 - b. Choose the MariaDB DB instance name to display its details.
 - c. On the **Connectivity & security** tab, copy the endpoint. Also, note the port number. You need both the endpoint and the port number to connect to the DB instance.

Summary

DB identifier database-1

Role Instance

CPU

Conn...

Logs & events

Configuration

Connectivity & security

Endpoint & port

Endpoint database-1. [REDACTED].us-east-2.rds.amazonaws.com

Port 3306

2. Enter the following command at a command prompt on a client computer to connect to a database on a MariaDB DB instance. Substitute the DNS name (endpoint) for your DB instance for <endpoint>, the master user name you used for <mymasteruser>, and provide the master password you used when prompted for a password.

```
PROMPT> mysql -h <endpoint> -P 3306 -u <mymasteruser> -p
```

After you enter the password for the user, you should see output similar to the following.

```
Welcome to the MySQL monitor. Commands end with ; or \g.  
Your MySQL connection id is 272  
Server version: 5.5.5-10.0.17-MariaDB-log MariaDB Server  
Copyright (c) 2000, 2015, Oracle and/or its affiliates. All rights reserved.  
Oracle is a registered trademark of Oracle Corporation and/or its
```

```
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql >
```

Deleting a DB Instance

After you have connected to the sample DB instance that you created, you should delete the DB instance so you are no longer charged for it.

To delete a DB instance with no final DB snapshot

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the DB instance you want to delete.
4. For **Actions**, choose **Delete**.
5. For **Create final snapshot?**, choose **No**, and select the acknowledgment.
6. Choose **Delete**.

Creating a Microsoft SQL Server DB Instance and Connecting to a DB Instance

The basic building block of Amazon RDS is the DB instance. Your Amazon RDS DB instance is similar to your on-premises Microsoft SQL Server. After you create your SQL Server DB instance, you can add one or more custom databases to it.

Important

You must have an AWS account before you can create a DB instance. If you don't have an AWS account, open <https://aws.amazon.com/>, and then choose **Create an AWS Account**.

In this topic, you create a sample SQL Server DB instance. You then connect to the DB instance and run a simple query. Finally, you delete the sample DB instance.

Creating a Sample SQL Server DB Instance

The basic building block of Amazon RDS is the DB instance. This environment is where you run your Microsoft SQL Server databases.

Note

A new console interface is available for database creation. Choose either the **New Console** or the **Original Console** instructions based on the console that you are using. The **New Console** instructions are open by default.

New Console

You can create a DB instance running Microsoft SQL Server with the AWS Management Console with **Easy create** enabled or not enabled. With **Easy Create** enabled, you specify only the DB engine type, DB instance size, and DB instance identifier. **Easy Create** uses the default setting for other configuration options. With **Easy Create** not enabled, you specify more configuration options when you create a database, including ones for availability, security, backups, and maintenance.

For this example, you use **Easy Create** to create a DB instance running the Microsoft SQL Server database engine with a db.t2.micro DB instance class.

Note

For information about creating a Microsoft SQL Server DB instance with **Easy Create** not enabled, see [Creating a DB Instance Running the Microsoft SQL Server Database Engine \(p. 616\)](#).

To create a Microsoft SQL Server DB instance with Easy Create enabled

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the Amazon RDS console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
4. Choose **Create database** and make sure that **Easy Create** is chosen.

Create database

Choose a database creation method Info

Standard Create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy Create

Use recommended best-practice configuration options can be changed after the database is created.

Configuration

5. In **Configuration**, choose **Microsoft SQL Server**.
6. For **DB instance size**, choose **Free tier**.
7. For **DB instance identifier**, enter a name for the DB instance, or leave the default name.
8. For **Master username**, enter a name for the master user, or leave the default name.

The **Create database** page should look similar to the following image.

Create database

Choose a database creation method Info

Standard Create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy Create

Use recommended best-practice configurations. Some configuration options can be changed after the database is created.

Configuration

Engine type Info

Amazon Aurora



MySQL



MariaDB

PostgreSQL



Oracle

ORACLE®

Microsoft SQL Server



DB instance size

Production

db.r4.large
2 vCPUs
15.25 GiB RAM
500 GiB

Dev/Test

db.m4.large
2 vCPUs
8 GiB RAM
100 GiB

Free tier

db.t2.micro
1 vCPUs
1 GiB RAM
20 GiB

DB instance identifier

Type a name for your DB instance. The name must be unique across all DB instances owned by your AWS account.
Region.

25

database-1

The DB instance identifier is case-insensitive, but is stored as all lowercase (as in "mydbinstance"). Constraints: 1

9. To use an automatically generated master password for the DB instance, make sure that the **Auto generate a password** check box is chosen.

To enter your master password, clear the **Auto generate a password** check box, and then enter the same password in **Master password** and **Confirm password**.

10. (Optional) Open **View default settings for Easy create**.

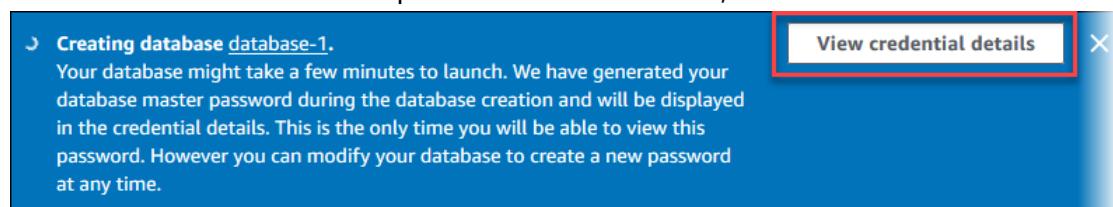
Configuration	Value	Editable after database is created
Database Location	Regional	No
Database Features	provisioned	No
Automatic Backups	Enabled	No

You can examine the default settings used when **Easy Create** is enabled. If you want to change one or more settings during database creation, choose **Standard Create** to set them. The **Editable after database creation** column shows which options you can change after database creation. To change a setting with **No** in that column, use **Standard Create**. For settings with **Yes** in that column, you can either use **Standard Create** or modify the DB instance after it's created to change the setting.

11. Choose **Create database**.

If you chose to use an automatically generated password, the **View credential details** button appears on the **Databases** page.

To view the master user name and password for the DB instance, choose **View credential details**.



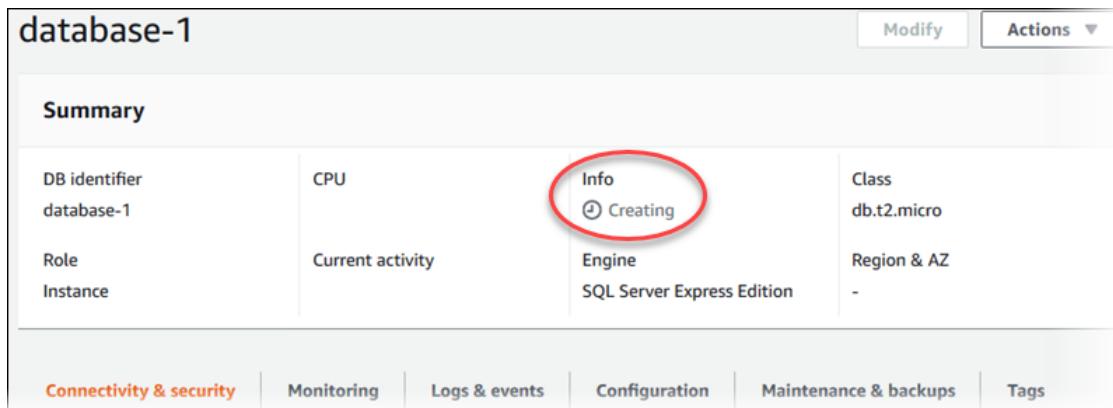
To connect to the DB instance as the master user, use the user name and password that appear.

Important

You can't view the master user password again. If you don't record it, you might have to change it. If you need to change the master user password after the DB instance is available, you can modify the DB instance to do so. For more information about modifying a DB instance, see [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#).

12. For **Databases**, choose the name of the new Microsoft SQL Server DB instance.

On the RDS console, the details for new DB instance appear. The DB instance has a status of **creating** until the DB instance is ready to use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and the amount of storage, it can take up to 20 minutes before the new instance is available.



Original Console

In this procedure you use the AWS Management Console to create a sample DB instance. Because you are only creating a sample DB instance, each setting is not fully explained. For a full explanation of each setting, see [Creating a DB Instance Running the Microsoft SQL Server Database Engine \(p. 616\)](#).

Before you begin, you must have an AWS account before you can create a DB instance. If you don't have an AWS account, open <https://aws.amazon.com/>, and then choose **Create an AWS Account**.

To create a DB instance running the Microsoft SQL Server DB engine

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the Amazon RDS console, choose the AWS Region where you want to create the DB instance.
3. Locate the **Create database** card below the **Resources** card. If you don't see it, scroll down or choose **Databases** from the navigation pane at left.

Choose **Create database**.

4. On the **Select engine** page, choose **Microsoft SQL Server**.
5. At the bottom of the page, check the box for **Only enable options eligible for RDS Free Usage Tier**. Doing this simplifies the setup process. For more information, see [Amazon RDS Free Tier](#).
6. Choose for the **SQL Server Express Edition**. This is the only SQL Server option that is eligible for RDS free Usage Tier.
7. On the **Instance specifications** page on the **Specify DB Details** page, provide the information for your DB instance as shown in the following table.

For This Parameter	Skip all of these. Provided here for reference only.
License Model	Skip. There is only one option available: license-included .
DB Engine Version	Skip. This defaults to the most recent version of SQL Server available in the list.
DB Instance Class	Skip. This defaults to the Free Tier eligible option, the smallest DB instance class available. This instance class is appropriate for testing.
Time Zone	Skip. If you don't choose a different time zone, your DB instance uses the default time zone.

For This Parameter	Skip all of these. Provided here for reference only.
Storage Type	Skip. This defaults to General Purpose (SSD) .
Allocated Storage	Skip. This defaults to 20 to allocate 20 GiB of storage for your database.
Storage autoscaling	Skip. Use the default settings, which enable storage autoscaling and set the maximum storage threshold to 1000 GB .

- On the **Settings** card at the bottom of the **Specify DB Details** page, provide the information for your DB instance as shown in the following table.

For This Parameter	Enter the Following
DB Instance Identifier	Enter sample-instance .
Username	Enter a name to use as the primary user name. The "master" user name is a login with all database privileges. It uses SQL Server Authentication.
Password and Confirm Password	Enter a password. It must contain between 8–128 printable ASCII characters (excluding /, ", and @).

Choose **Next** to continue.

- The **Configure Advanced Settings** page appears.

On the **Network & Security** card on the **Configure Advanced Settings** page, provide the information for your DB instance as shown in the following table.

For This Parameter	Do This
Virtual Private Cloud (VPC)	Choose Create new VPC .
Subnet Group	Choose Create new DB Subnet Group .
Publicly Accessible	Choose No .
Availability Zone	Choose No Preference .
VPC Security Group	Choose Create new Security Group .

- On the **Microsoft SQL Server Windows Authentication** card on the **Configure Advanced Settings** page, provide the information for your DB instance as shown in the following table.

For This Parameter	Do This
Directory	Skip. This defaults to None .

- On the **Database Options** card on the **Configure Advanced Settings** page, provide the information for your DB instance as shown in the following table.

For This Parameter	Do This
Database Port	Skip. This defaults to 1433 . SQL Server installations default to port 1433, but in some cases a firewall might block this port. If in doubt, ask your network administrator what port you should use.
DB Parameter Group	Skip. Keep the default value.
Option Group	Skip. Keep the default value.

12. On the **Encryption** card on the **Configure Advanced Settings** provide the information for your DB instance as shown in the following table.

For This Parameter	Do This
Encryption	Choose Disable encryption .

13. On the **Backup** card on the **Configure Advanced Settings** page, provide the information for your DB instance as shown in the following table.

For This Parameter	Do This
Backup Retention Period	Skip. Defaults to 7 days .
Backup Window	Skip. Defaults to No Preference .
Copy Tags To Snapshots	Skip. Defaults to Copy tags to snapshots . However, you can clear the check box if you prefer to do so.

14. On the **Monitoring** card on the **Configure Advanced Settings** page, provide the information for your DB instance as shown in the following table.

For This Parameter	Do This
Enhanced Monitoring	Choose Disable enhanced monitoring .

15. On the **Performance Insights** card on the **Configure Advanced Settings** page, provide the information for your DB instance as shown in the following table.

For This Parameter	Do This
Performance Insights	Choose Disable Performance Insights

16. On the **Maintenance** card on the **Configure Advanced Settings** page, provide the information for your DB instance as shown in the following table.

For This Parameter	Do This
Auto Minor Version Upgrade	Choose Disable auto minor version upgrade .
Maintenance Window	Choose No Preference .

17. On the **Deletion protection** card on the **Configure Advanced Settings** page, provide the information for your DB instance as shown in the following table.

For This Parameter	Do This
Deletion protection	Clear the box for Enable deletion protection , to make it easier to delete this DB instance when you are done with it.

18. Choose **Create database** to confirm your choices and create the DB instance.
19. After you create the database, choose **View Your DB Instances**.

On the RDS console, the new DB instance appears in the list of DB instances. The DB instance has a status of **Creating** until the DB instance is ready to use. When the state changes to **Available**, you can connect to the DB instance.

Connecting to Your Sample SQL Server DB Instance

In this procedure you connect to your sample DB instance by using Microsoft SQL Server Management Studio (SSMS).

Before you begin, your database should have a status of **Available**. If it has a status of **Creating** or **Backup-up**, wait until it's **Available**. The status updates without requiring you to refresh the page. This process can take up to 20 minutes.

Also, make sure you have SSMS installed. If you can also connect to SQL Server on RDS by using a different tools, such as an add-in for your development environment or some other database tool. However, this tutorial only covers using SSMS. To download a stand-alone version of this SSMS, see [Download SQL Server Management Studio \(SSMS\)](#) in the Microsoft documentation.

To connect to a DB Instance using SSMS

1. Find the DNS name and port number for your DB Instance.
 - a. Open the RDS console and then choose **Databases** to display a list of your DB instances.
 - b. Hover your mouse cursor over the name **sample-instance**, which is blue. When you do this, the mouse cursor changes into a selection icon (for example, a pointing hand). Also, the DB instance name, becomes underlined.

Click on the DB instance name to choose it. The screen changes to display the information for the DB instance you choose.

 - c. On the **Connectivity** tab, which opens by default, copy the endpoint. The **Endpoint** looks something like this: `sample-instance.abc2defghije.us-west-2.rds.amazonaws.com`. Also, take note of the port number. The default port for SQL Server is 1433. If yours is different, write it down.
 2. Start SQL Server Management Studio.
- The **Connect to Server** dialog box appears.
3. Provide the information for your sample DB instance.
 - a. For **Server type**, choose **Database Engine**.
 - b. For **Server name**, enter the DNS name, followed by a comma and the port number (the default port is 1433). For example, your server name should look like the following.

`sample-instance.abc2defghije.us-west-2.rds.amazonaws.com,1433`

- c. For **Authentication**, choose **SQL Server Authentication**.

- d. For **Login**, enter the user name that you chose to use for your sample DB instance. This is also known as the master user name.
 - e. For **Password**, enter the password that you chose earlier for your sample DB instance. This is also known as the master user password.
4. Choose **Connect**.

After a few moments, SSMS connects to your DB instance.

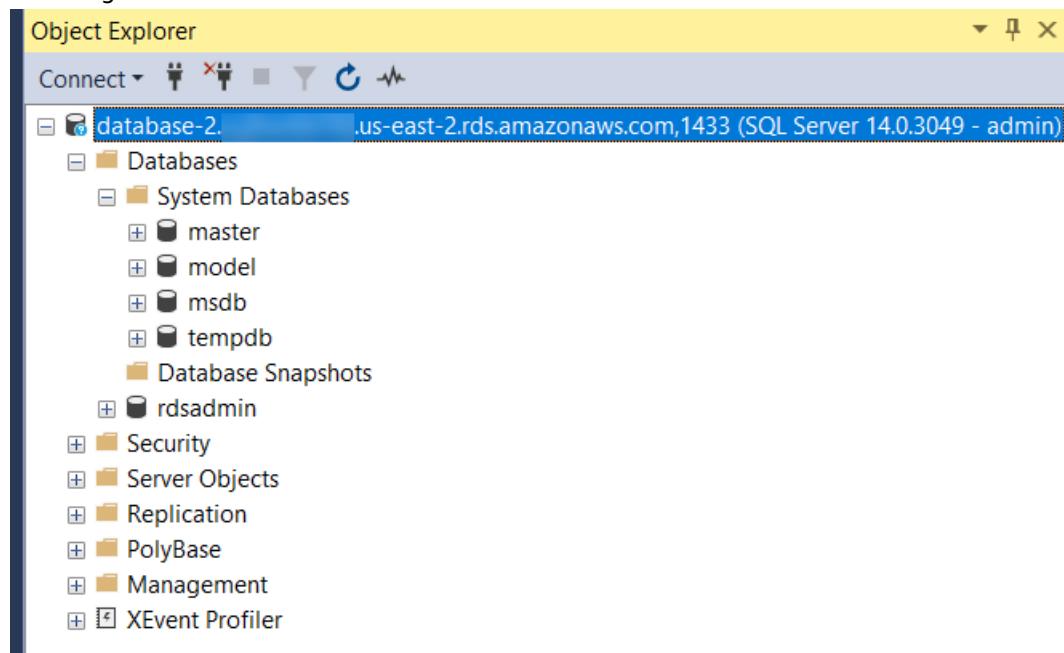
If you can't connect to your DB instance, see [Troubleshooting the Connection to Your SQL Server DB Instance \(p. 637\)](#).

Exploring Your Sample SQL Server DB Instance

In this procedure you continue the previous procedure and explore your sample DB instance by using Microsoft SQL Server Management Studio (SSMS).

To explore a DB Instance using SSMS

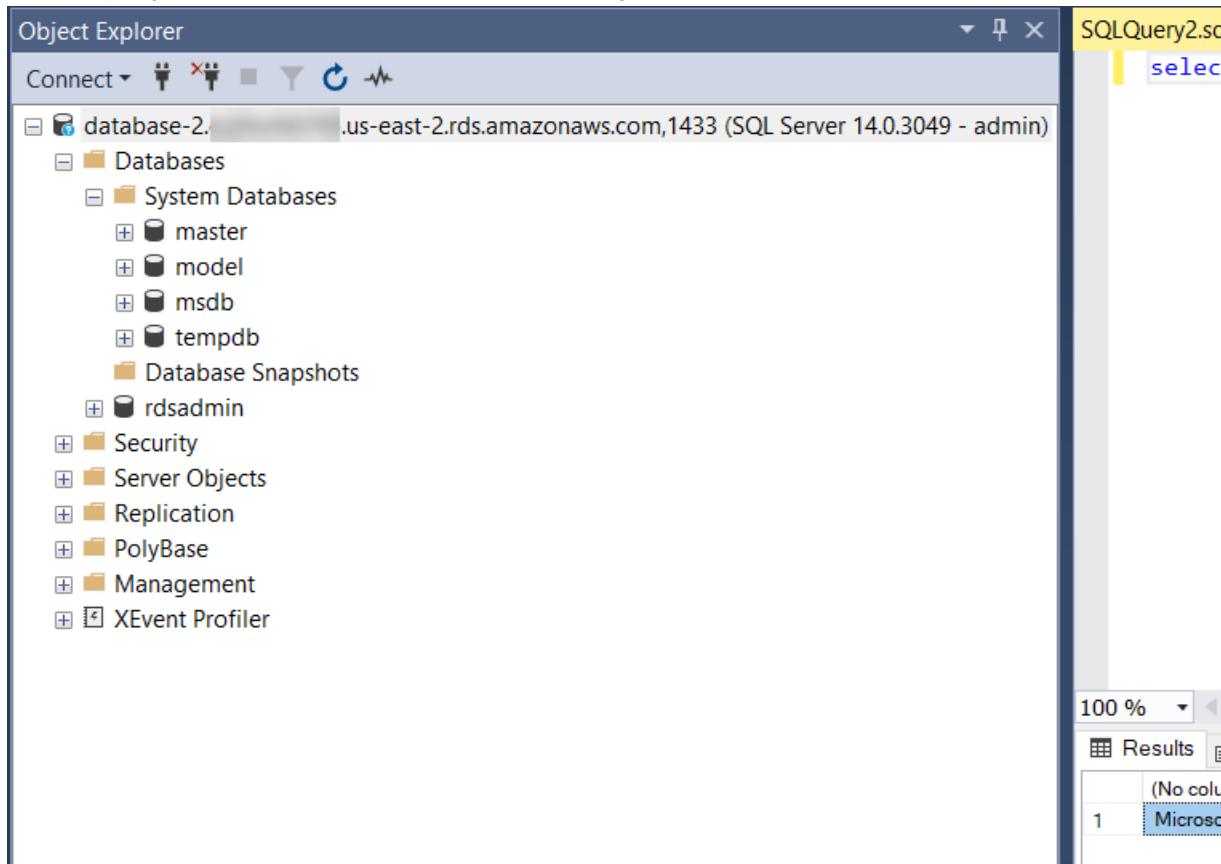
1. Your SQL Server DB instance comes with SQL Server's standard built-in system databases (master, model, msdb, and tempdb). To explore the system databases, do the following:
 - a. In SSMS, on the **View** menu, choose **Object Explorer**.
 - b. Expand your DB instance, expand **Databases**, and then expand **System Databases** as shown following.



2. Your SQL Server DB instance also comes with a database named `rdsadmin`. Amazon RDS uses this database to store the objects that it uses to manage your database. The `rdsadmin` database also includes stored procedures that you can run to perform advanced tasks.
3. You can now start creating your own databases and running queries against your DB instance and databases as usual. To run a test query against your sample DB instance, do the following:
 - a. In SSMS, on the **File** menu point to **New** and then choose **Query with Current Connection**.
 - b. Enter the following SQL query.

```
select @@VERSION
```

- c. Run the query. SSMS returns the SQL Server version of your Amazon RDS DB instance.



Deleting Your Sample DB Instance

After you are done exploring the sample DB instance that you created, you should delete the DB instance so that you are no longer charged for it.

To delete a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the button next to **sample-instance**, or whatever you named your sample DB instance.
4. Choose **Delete** from the **Actions** above (it's to the left of the orange **Create database** button).

If you see a message that says **This database has deletion protection option enabled**, follow these steps:

- Choose **Modify**.
- On the **Deletion protection** card (near the bottom of the page), clear the box next to **Enable deletion protection**. Then choose **Continue**.

- On the **Scheduling of modifications** card, choose **Apply immediately**. Then choose **Modify DB Instance**.
 - Try again to delete the instance by choosing **Delete** from the **Actions** menu.
5. Clear the box for **Create final snapshot**. Because this isn't a production database, you don't need to save a copy of it.
 6. Verify that you selected the correct database to delete. The name "sample-instance" displays in the title of the screen: **Delete sample-instance instance?**

If you don't recognize the name of your sample instance in the title, choose **Cancel** and start over.
 7. To confirm that you want to permanently delete the database that is displayed in the title of this screen, do the following:
 - Check the box to confirm: **I acknowledge that upon instance deletion, automated backups, including system snapshots and point-in-time recovery, will no longer be available.**
 - Type "**delete me**" into the box **To confirm deletion, type delete me into the field**.
 - Choose **Delete**. This action can't be undone.

The database shows a status of **Deleting** until deletion is complete.

Creating a MySQL DB Instance and Connecting to a Database on a MySQL DB Instance

The easiest way to create a DB instance is to use the AWS Management Console. After you have created the DB instance, you can use standard MySQL utilities such as MySQL Workbench to connect to a database on the DB instance.

Important

Before you can create or connect to a DB instance, you must complete the tasks in [Setting Up for Amazon RDS \(p. 5\)](#).

Topics

- [Creating a MySQL DB Instance \(p. 33\)](#)
- [Connecting to a Database on a DB Instance Running the MySQL Database Engine \(p. 42\)](#)
- [Deleting a DB Instance \(p. 44\)](#)

Creating a MySQL DB Instance

The basic building block of Amazon RDS is the DB instance. This environment is where you run your MySQL databases.

Note

A new console interface is available for database creation. Choose either the **New Console** or the **Original Console** instructions based on the console that you are using. The **New Console** instructions are open by default.

New Console

You can create a DB instance running MySQL with the AWS Management Console with **Easy Create** enabled or disabled. With **Easy Create** enabled, you specify only the DB engine type, DB instance size,

and DB instance identifier. **Easy Create** uses the default setting for other configuration options. With **Easy Create** not enabled, you specify more configuration options when you create a database, including ones for availability, security, backups, and maintenance.

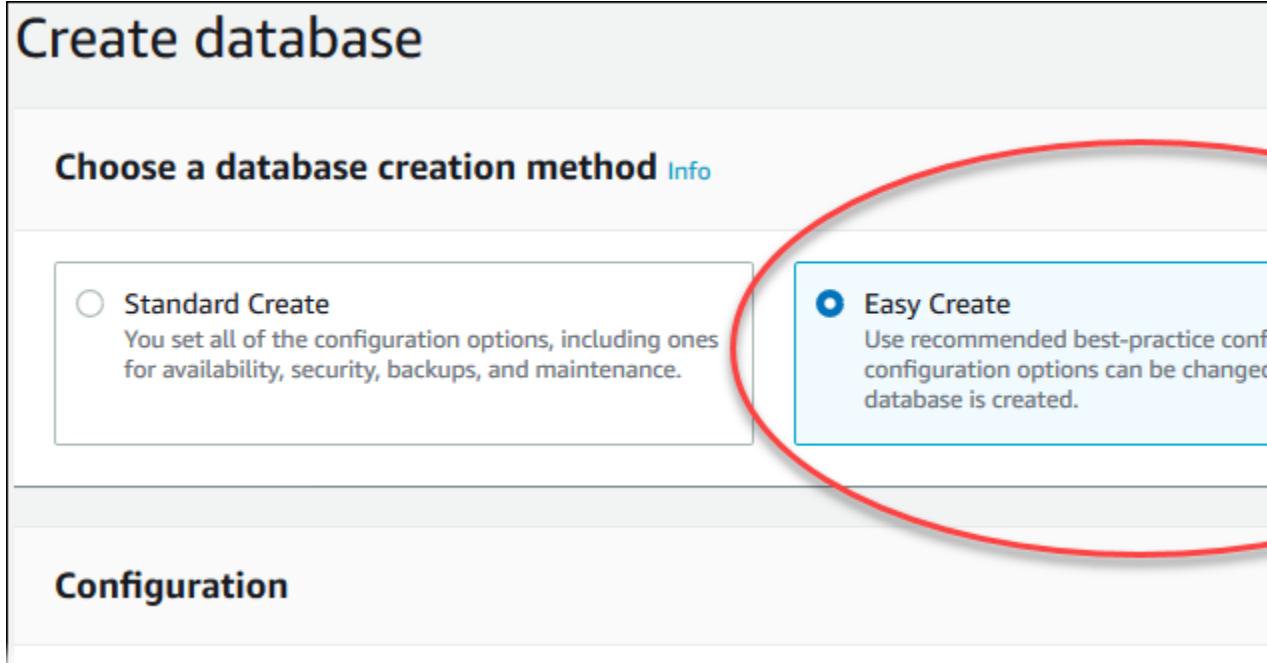
In this example, you use **Easy Create** to create a DB instance running the MySQL database engine with a db.t2.micro DB instance class.

Note

For information about creating a MySQL DB instance with **Easy Create** not enabled, see [Creating a DB Instance Running the MySQL Database Engine \(p. 751\)](#).

To create a MySQL DB instance with Easy Create enabled

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the Amazon RDS console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
4. Choose **Create database** and make sure that **Easy Create** is chosen.



5. In **Configuration**, choose **MySQL**.
6. For **DB instance size**, choose **Free tier**.
7. For **DB instance identifier**, enter a name for the DB instance, or leave the default name.
8. For **Master username**, enter a name for the master user, or leave the default name.

The **Create database** page should look similar to the following image.

Create database

Choose a database creation method Info

Standard Create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy Create

Use recommended best-practice configuration options can be changed after the database is created.

Configuration

Engine type Info

Amazon Aurora



MySQL



MariaDB

PostgreSQL



Oracle

ORACLE®

Microsoft SQL Server



DB instance size

Production

db.r5.xlarge
4 vCPUs
32 GiB RAM
500 GiB

Dev/Test

db.r5.large
2 vCPUs
16 GiB RAM
100 GiB

Free tier

db.t2.micro
1 vCPUs
1 GiB RAM
20 GiB

DB instance identifier

Type a name for your DB instance. The name must be unique across all DB instances owned by your AWS account in the specified Region.

35

database-1

The DB instance identifier is case-insensitive, but is stored as all lowercase (as in "mydbinstance"). Constraints: 1–63 characters, must begin with a letter, and can contain numbers, underscores, and hyphens.

9. To use an automatically generated master password for the DB instance, enable **Auto generate a password**.

To enter your master password, disable **Auto generate a password**, and then enter the same password in **Master password** and **Confirm password**.

10. (Optional) Open **View default settings for Easy create**.

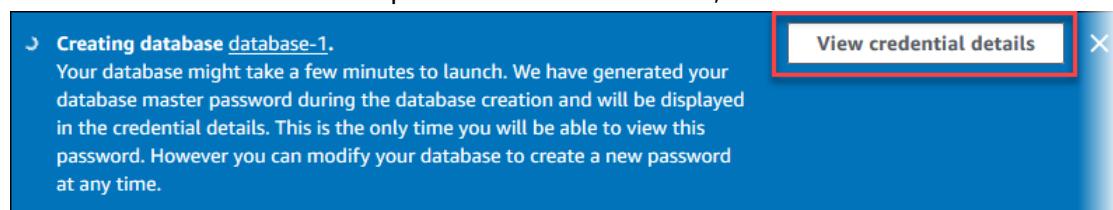
Configuration	Value	Editable after database is created
Database Location	Regional	No
Database Features	provisioned	No
Automatic Backups	Enabled	No

You can examine the default settings used when **Easy Create** is enabled. If you want to change one or more settings during database creation, choose **Standard Create** to set them. The **Editable after database creation** column shows which options you can change after database creation. To change a setting with **No** in that column, use **Standard Create**. For settings with **Yes** in that column, you can either use **Standard Create** or modify the DB instance after it is created to change the setting.

11. Choose **Create database**.

If you chose to use an automatically generated password, the **View credential details** button appears on the **Databases** page.

To view the master username and password for the DB instance, choose **View credential details**.



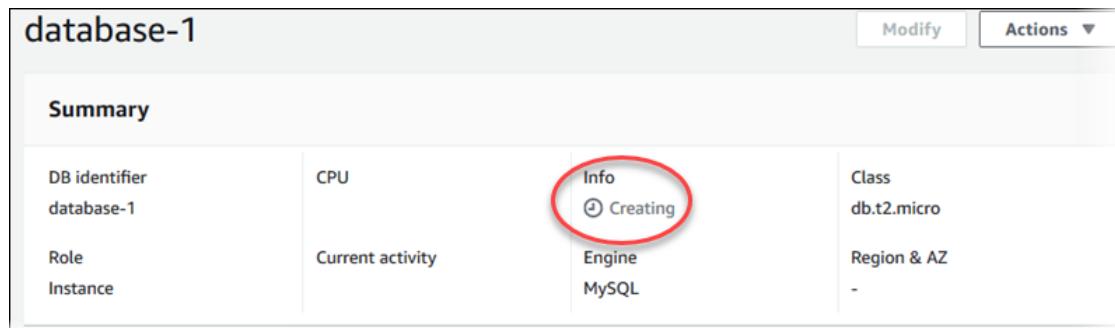
You can use the username and password that appears to connect to the DB instance as the master user.

Important

You won't be able to view master user password again. If you don't record it, you might have to change it. If you need to change the master user password after the DB instance is available, you can modify the DB instance to do so. For more information about modifying a DB instance, see [Modifying a DB Instance Running the MySQL Database Engine \(p. 802\)](#).

12. In the **Databases** list, choose the name of the new MySQL DB instance.

On the RDS console, the details for new DB instance appear. The DB instance has a status of **creating** until the DB instance is ready to use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and the amount of storage, it can take up to 20 minutes before the new instance is available.



Original Console

In this example, you create a DB instance running the MySQL database engine called *mysql-instance1*, with a *db.m1.small* DB instance class, 20 GiB of storage, and automated backups enabled with a retention period of one day.

To create a MySQL DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the top right corner of the Amazon RDS console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
If the navigation pane is closed, choose the menu icon at the top left to open it.
4. Choose **Create database**. The **Select engine** page opens.

Select engine

Engine options

- Amazon Aurora
Amazon Aurora
- MySQL

- MariaDB

- PostgreSQL

- Oracle

- Microsoft SQL Server


MySQL

MySQL is the most popular open source database in the world. MySQL on RDS offers the rich features of the MySQL community edition with the flexibility to easily scale compute resources or storage capacity for your database.

- Supports database size up to 16 TB.
- Instances offer up to 32 vCPUs and 244 GiB Memory.
- Supports automated backup and point-in-time recovery.
- Supports cross-region read replicas.

Only enable options eligible for RDS Free Usage Tier [info](#)

[Cancel](#) [Next](#)

5. Choose **MySQL**, and then choose **Next**.
6. The **Choose use case** page asks if you are planning to use the DB instance you are creating for production. Choose **Dev/Test** and then choose **Next**.
7. On the **Specify DB Details** page, specify your DB instance information. The following table shows settings for an example DB instance. When the settings are as you want them, choose **Next**.

For This Parameter	Do This
License model	Choose the default, general-public-license , to use the general license agreement for MySQL. MySQL has only one license model.
DB engine version	Choose the default version of MySQL. Amazon RDS supports multiple versions of MySQL in some AWS Regions.
DB instance class	Choose db.m1.small .
Multi-AZ deployment	Choose Yes to have a standby replica of your DB instance created in another Availability Zone for failover support. We recommend Multi-AZ for production workloads to

For This Parameter	Do This
	<p>maintain high availability. For development and testing, you can choose No.</p> <p>For more information, see High Availability (Multi-AZ) for Amazon RDS (p. 144).</p>
Storage type	Choose the storage type General Purpose (SSD) . For more information about storage, see Amazon RDS DB Instance Storage (p. 135) .
Allocated storage	Enter 20 to allocate 20 GiB of storage for your database. In some cases, allocating a higher amount of storage for your DB instance than the size of your database can improve I/O performance. For more information about storage allocation, see Amazon RDS Features .
DB instance identifier	Enter a name for the DB instance that is unique for your account in the AWS Region you chose. You can add some intelligence to the name, such as including the AWS Region and DB engine you chose, for example mysql-instance1 .
Master username	Enter a name using alphanumeric characters to use as the master user name to log on to your DB instance. This is the user name you use to log on to your database on the DB instance for the first time.
Master password and Confirm password	Enter a password that contains from 8 to 41 printable ASCII characters (excluding /, ", and @) for your master user password. This is the password to use when you use the user name to log on to your database. Then type the password again in the Confirm Password box.

Specify DB details

Instance specifications

Estimate your monthly costs for the DB Instance using the [AWS Simple Monthly Calculator](#)

DB engine

MySQL Community Edition

License model [Info](#)

general-public-license

DB engine version [Info](#)

MySQL 5.6.40



Known Issues/Limitations

Review the [Known Issues/Limitations](#) to learn about potential compatibility issues with specific database versions.



Free tier

The Amazon RDS Free Tier provides a single db.t2.micro instance as well as up to 20 GiB of storage, allowing new AWS customers to gain hands-on experience with Amazon RDS. Learn more about the RDS Free Tier and the instance restrictions [here](#).

Only enable options eligible for RDS Free Usage Tier [Info](#)

DB instance class [Info](#)

db.r4.xlarge — 4 vCPU, 30.5 GiB RAM

Multi-AZ deployment [Info](#)

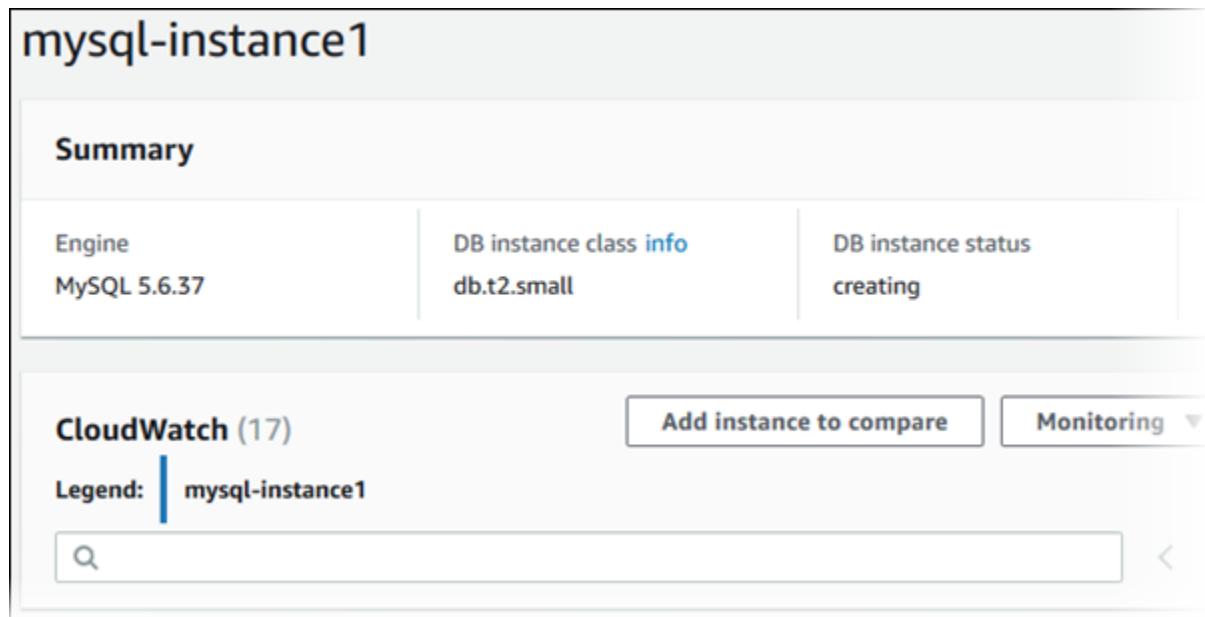
8. Choose **Next**.
9. On the **Configure advanced settings** page, provide additional information that RDS needs to launch the MySQL DB instance. The table shows settings for an example DB instance. Specify your DB instance information, then choose **Create database**.

For This Parameter	Do This
Virtual Private Cloud (VPC)	Choose Create new VPC .
Subnet group	Choose Create new DB subnet group .
Public accessibility	Choose Yes .
Availability zone	Choose No Preference .

For This Parameter	Do This
VPC security groups	Choose Create new VPC security group .
Database name	<p>Enter a name for your default database that is 1 to 64 alpha-numeric characters. If you don't provide a name, Amazon RDS doesn't automatically create a database on the DB instance you are creating.</p> <p>To create additional databases, connect to the DB instance and use the SQL command <code>CREATE DATABASE</code>. For more information about connecting to the DB instance, see Connecting to a DB Instance Running the MySQL Database Engine (p. 764).</p>
Database port	Leave the default value of 3306 unless you have a specific port you want to access the database through. MySQL installations default to port 3306.
DB parameter group	Leave the default value unless you created your own DB parameter group. For more information about parameter groups, see Working with DB Parameter Groups (p. 211) .
Option group	Choose the default value because this option group is used with the MySQL version you chose on the previous page.
Copy tags To snapshots	Choose this option to have any DB instance tags copied to a DB snapshot when you create a snapshot. For more information, see Tagging Amazon RDS Resources (p. 177) .
IAM DB authentication	Choose No . For more information, see Identity and Access Management in Amazon RDS (p. 445) .
Encryption	Choose Enable encryption to enable encryption at rest for this DB instance. For more information, see Encrypting Amazon RDS Resources (p. 429) .
Backup retention period	Set the number of days you want automatic backups of your database to be retained. For testing purposes, you can set this value to 1 .
Backup window	Unless you have a specific time that you want to have your database backup, use the default of No Preference .
Enhanced monitoring	Unless you want to enable gathering metrics in real time for the operating system that your DB instance runs on, use the default of Disable enhanced monitoring .
Log exports	Select General log . For more information, see MySQL Database Log Files (p. 405) .
Auto minor version upgrade	Choose Enable auto minor version upgrade to enable your DB instance to receive preferred minor DB engine version upgrades automatically when they become available.
Maintenance window	Choose No preference .

10. Choose **Create database**.
11. Choose **View DB instance details**.

On the RDS console, the details for new DB instance appear. The DB instance has a status of **creating** until the DB instance is ready to use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and the amount of storage, it can take up to 20 minutes before the new instance is available.



Connecting to a Database on a DB Instance Running the MySQL Database Engine

Once Amazon RDS provisions your DB instance, you can use any standard SQL client application to connect to a database on the DB instance. In this example, you connect to a database on a MySQL DB instance using MySQL monitor commands. One GUI-based application you can use to connect is MySQL Workbench; for more information, go to the [Download MySQL Workbench](#) page. For more information on using MySQL, go to the [MySQL documentation](#). For information about installing MySQL (including the MySQL client), see [Installing and Upgrading MySQL](#).

To connect to a database on a DB instance using MySQL monitor

1. Find the endpoint (DNS name) and port number for your DB Instance.
 - a. Open the RDS console and then choose **Databases** to display a list of your DB instances.
 - b. Choose the MySQL DB instance name to display its details.
 - c. On the **Connectivity & security** tab, copy the endpoint. Also, note the port number. You need both the endpoint and the port number to connect to the DB instance.

Summary

DB identifier
database-1

Role
Instance

CPU

Memory

Logs & events

Configuration

Connectivity & security

Endpoint & port

Endpoint
database-1. [REDACTED].us-east-2.rds.amazonaws.com

Port
3306

2. Download a SQL client that you can use to connect to the DB instance.

You can connect to an Amazon RDS MySQL DB instance by using tools like the MySQL command line utility. For more information on using the MySQL client, go to [mysql - The MySQL Command Line Tool](#) in the MySQL documentation. One GUI-based application you can use to connect is MySQL Workbench. For more information, go to the [Download MySQL Workbench](#) page.

3. Connect to the a database on a MySQL DB instance. For example, enter the following command at a command prompt on a client computer to connect to a database on a MySQL DB instance using the MySQL client. Substitute the DNS name for your DB instance for `<endpoint>`, the master user name you used for `<mymasteruser>`, and provide the master password you used when prompted for a password.

```
PROMPT> mysql -h <endpoint> -P 3306 -u <mymasteruser> -p
```

After you enter the password for the user, you should see output similar to the following.

```
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 350
Server version: 5.6.40-log MySQL Community Server (GPL)

Type 'help;' or '\h' for help. Type '\c' to clear the buffer.

mysql>
```

If you can't connect to your MySQL DB instance, two common causes of connection failures to a new DB instance are:

- The DB instance was created using a security group that does not authorize connections from the device or Amazon EC2 instance where the MySQL application or utility is running. If the DB instance was created in a VPC, it must have a VPC security group that authorizes the connections. If the DB instance was created outside of a VPC, it must have a DB security group that authorizes the connections. For more information, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).
- The DB instance was created using the default port of 3306, and your company has firewall rules blocking connections to that port from devices in your company network. To fix this failure, recreate the instance with a different port.

For more information about connecting to a MySQL DB instance, see [Connecting to a DB Instance Running the MySQL Database Engine \(p. 764\)](#).

Deleting a DB Instance

After you have connected to the sample DB instance that you created, you should delete the DB instance so you are no longer charged for it.

To delete a DB instance with no final DB snapshot

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the DB instance that you want to delete.
4. For **Actions**, choose **Delete**.
5. For **Create final snapshot?**, choose **No**, and select the acknowledgment.
6. Choose **Delete**.

Creating an Oracle DB Instance and Connecting to a Database on an Oracle DB Instance

The basic building block of Amazon RDS is the DB instance. Your Amazon RDS DB instance is similar to your on-premises Oracle database.

Important

You must have an AWS account before you can create a DB instance. If you don't have an AWS account, open <https://aws.amazon.com/>, and then choose **Create an AWS Account**.

In this topic, you create a sample Oracle DB instance. You then connect to the DB instance and run a simple query. Finally, you delete the sample DB instance.

Creating a Sample Oracle DB Instance

The basic building block of Amazon RDS is the DB instance. This environment is where you run your Oracle databases.

Note

A new console interface is available for database creation. Choose either the **New Console** or the **Original Console** instructions based on the console that you are using. The **New Console** instructions are open by default.

New Console

You can create a DB instance running Oracle with the AWS Management Console with **Easy create** enabled or not enabled. With **Easy Create** enabled, you specify only the DB engine type, DB instance size, and DB instance identifier. **Easy Create** uses the default setting for other configuration options. With **Easy Create** not enabled, you specify more configuration options when you create a database, including ones for availability, security, backups, and maintenance.

For this example, you use **Easy Create** to create a DB instance running the Oracle database engine with a db.t2.micro DB instance class.

Note

For information about creating an Oracle DB instance with **Easy Create** not enabled, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).

To create an Oracle DB instance with Easy Create enabled

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the Amazon RDS console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
4. Choose **Create database** and ensure that **Easy Create** is chosen.

Create database

Choose a database creation method Info

Standard Create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy Create

Use recommended best-practice configuration options can be changed after the database is created.

Configuration

5. In **Configuration**, choose **Oracle**.
6. For **DB instance size**, choose **Free tier**. If **Free tier** isn't available, choose **Dev/Test**.
7. For **DB instance identifier**, enter a name for the DB instance, or leave the default name.
8. For **Master username**, enter a name for the master user, or leave the default name.

The **Create database** page should look similar to the following image.

Create database

Choose a database creation method Info

Standard Create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy Create

Use recommended best-practice configuration options can be changed after the database is created.

Configuration

Engine type Info

Amazon Aurora



MySQL



MariaDB



PostgreSQL



Oracle

ORACLE®

Microsoft SQL Server



DB instance size

Production

db.r4.large
2 vCPUs
15.25 GiB RAM
500 GiB

Dev/Test

db.m4.large
2 vCPUs
8 GiB RAM
100 GiB

DB instance identifier

Type a name for your DB instance. The name must be unique across all DB instances owned by your AWS account.

Region:
database-1

9. To use an automatically generated master password for the DB instance, make sure that the **Auto generate a password** check box is chosen.

To enter your master password, clear the **Auto generate a password** check box, and then enter the same password in **Master password** and **Confirm password**.

10. (Optional) Open **View default settings for Easy create**.

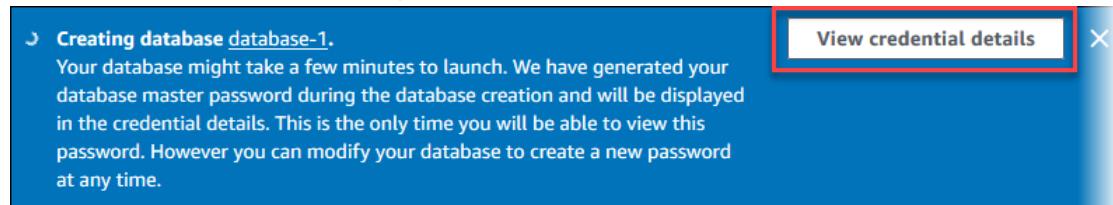
Configuration	Value	Editable after database is created
Database Location	Regional	No
Database Features	provisioned	No
Automatic Backups	Enabled	No

You can examine the default settings that are used when **Easy Create** is enabled. If you want to change one or more settings during database creation, choose **Standard Create** to set them. The **Editable after database creation** column shows which options you can change after database creation. To change a setting with **No** in that column, use **Standard Create**. For settings with **Yes** in that column, you can either use **Standard Create** or modify the DB instance after it's created to change the setting.

11. Choose **Create database**.

If you used an automatically generated password, the **View credential details** button appears on the **Databases** page.

To view the master user name and password for the DB instance, choose **View credential details**.



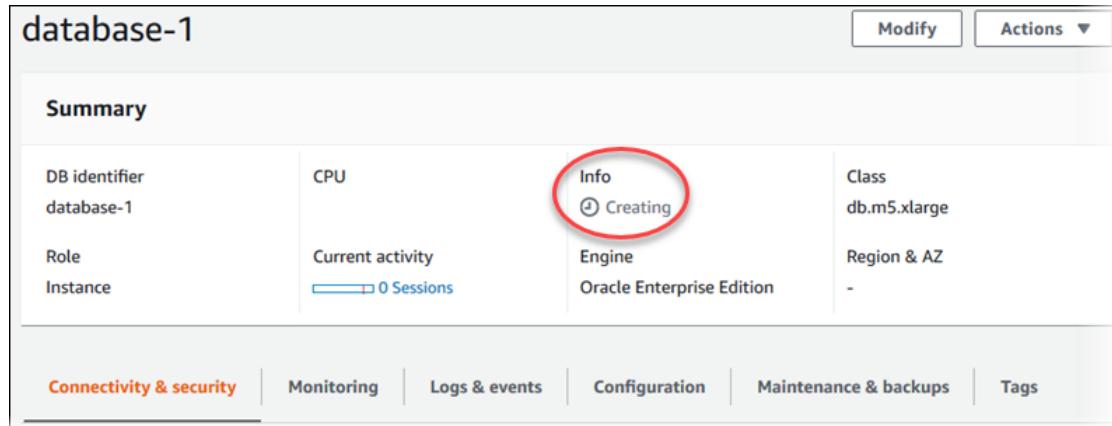
To connect to the DB instance as the master user, use the user name and password that appear.

Important

> You can't view the master user password again. If you don't record it, you might have to change it. If you need to change the master user password after the DB instance is available, you can modify the DB instance to do so. For more information about modifying a DB instance, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

12. For **Databases**, choose the name of the new Oracle DB instance.

On the RDS console, the details for new DB instance appear. The DB instance has a status of **creating** until the DB instance is ready to use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and the amount of storage, it can take up to 20 minutes before the new instance is available.



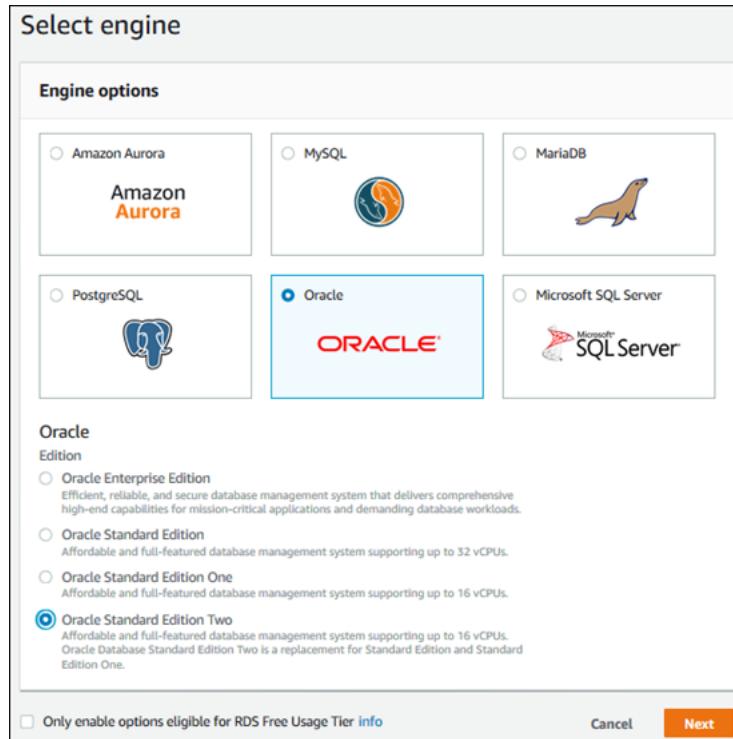
Original Console

In this procedure you use the AWS Management Console to create a sample DB instance. Since you are only creating a sample DB instance, each setting is not fully explained. For a full explanation of each setting, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).

To create a DB instance running the Oracle database engine

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the top right corner of the Amazon RDS console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
4. Choose **Create database**.

The **Select engine** page appears.



5. Choose the Oracle icon, and then choose **Select** for the **Oracle Standard Edition Two** edition.
6. The **Choose use case** page asks if you are planning to use the DB instance you are creating for production. Choose **Dev/Test** and then choose **Next**.

The **Specify DB details** page appears.

Specify DB details

Instance specifications
Estimate your monthly costs for the DB Instance using the [AWS Simple Monthly Calculator](#).

DB engine
Oracle Database Standard Edition Two

License model info

DB engine version info

Free tier
The Amazon RDS Free Tier provides a single db.t2.micro instance as well as up to 20 GB of storage, allowing new AWS customers to gain hands-on experience with Amazon RDS. Learn more about the RDS Free Tier and the instance restrictions [here](#).

Only enable options eligible for RDS Free Usage Tier info

DB instance class info

Multi-AZ deployment info
 Create replica in different zone
Creates a replica in a different Availability Zone (AZ) to provide data redundancy, eliminate I/O freezes, and minimize latency spikes during system backups.

 No

Storage type info

Allocated storage
 GB
(Minimum: 20 GB, Maximum: 6144 GB) Higher allocated storage may improve IOOPS performance.

Provisioning less than 100 GB of General Purpose (SSD) storage for high throughput workloads could result in higher latencies upon exhaustion of the initial General Purpose (SSD) IO credit balance. Click here for more details.

7. On the **Specify DB details** page, provide the information for your DB instance as shown in the following table.

For This Parameter	Do This
License model	Choose license-included to use the general license agreement for Oracle.
DB engine version	Choose the most recent version of Oracle available in the list.
DB instance class	Choose db.t2.small . This instance class is appropriate for testing.
Multi-AZ deployment	For development and testing, choose No .
Storage type	Choose the storage type General Purpose (SSD) .
Allocated storage	Enter 20 to allocate 20 GiB of storage for your database.
DB instance identifier	Enter oracle-instance1 .

For This Parameter	Do This
Master username	Enter a name to use as the master user name to log on to your DB instance with all database privileges. The master user name is a SQL Server Authentication login.
Master password and confirm Password	Enter a password for your master user password. It must contain 8–128 printable ASCII characters (excluding /, ", and @).

- Choose **Next** to continue.

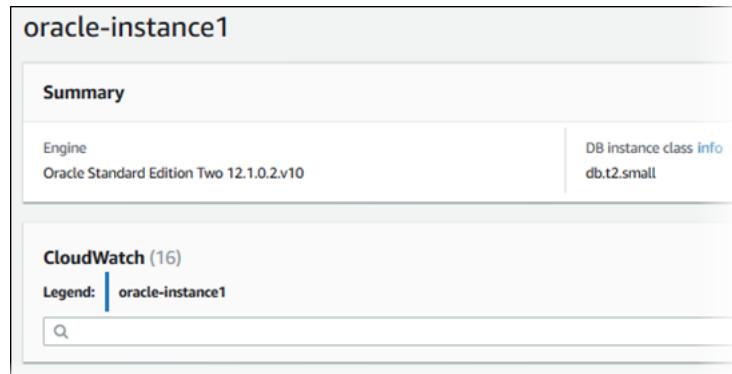
The **Configure Advanced Settings** page appears.

- On the **Configure advanced settings** page, provide the information for your DB instance as shown in the following table.

For This Parameter	Do This
Virtual Private Cloud (VPC)	Choose Create new VPC .
Subnet group	Choose Create new DB subnet group .
Public accessibility	Choose Yes .
Availability zone	Choose No Preference .
VPC security groups	Choose Create new VPC security group .
Database name	Enter ORCL .
Database port	Keep the default value of 1521 unless you have a specific port you want to access the database through. Oracle installations default to port 1521, but in some cases a firewall might block this port. If in doubt, ask your network administrator what port you should use.
DB parameter group	Keep the default value.
Option group	Keep the default value.
Copy tags to snapshots	Don't choose this setting; keep it clear.
Character set name	Choose the default value of AL32UTF8 for the Unicode 5.0 UTF-8 Universal character set.
Enable encryption	Choose No to enable encryption at rest for this DB instance.
Backup retention period	Choose 7 days .
Backup window	Choose No preference .
Enhanced monitoring	Choose Disable enhanced monitoring .
Auto minor version upgrade	Choose Enable auto minor version upgrade to enable your DB instance to receive preferred minor DB engine version upgrades automatically when they become available.
Maintenance window	Choose No preference .

10. Choose **Create database**.
11. Choose **View DB instance details**.

On the RDS console, the details for new DB instance appear. The DB instance has a status of **creating** until the DB instance is ready to use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and the amount of storage, it can take up to 20 minutes before the new instance is available.



Connecting to Your Sample Oracle DB Instance

After Amazon RDS provisions your DB instance, you can use any standard SQL client application to connect to the DB instance. In this procedure you connect to your sample DB instance by using the Oracle *sqlplus* command line utility. To download a stand-alone version of this utility, see [SQL*Plus User's Guide and Reference](#).

To connect to a DB Instance using SQL*Plus

1. Find the endpoint (DNS name) and port number for your DB Instance.
 - a. Open the RDS console and then choose **Databases** to display a list of your DB instances.
 - b. Choose the Oracle DB instance name to display its details.
 - c. On the **Connectivity & security** tab, copy the endpoint. Also, note the port number. You need both the endpoint and the port number to connect to the DB instance.

The screenshot shows the AWS RDS console for a database named 'database-1'. The 'Summary' section displays the DB identifier as 'database-1' and the instance role as 'Instance'. Below this, the 'Connectivity & security' tab is active, showing the endpoint details. A red oval highlights the 'Endpoint' field, which contains the value 'database-1.us-west-2.rds.amazonaws.com'. Other visible fields include the port number '1521'.

2. Enter the following command on one line at a command prompt to connect to your DB instance by using the sqlplus utility. The value for Host is the endpoint for your DB instance, and the value for Port is the port you assigned the DB instance. The value for the Oracle SID is the name of the DB instance's database that you specified when you created the DB instance, not the name of the DB instance.

```
PROMPT>sqlplus 'mydbusr@(DESCRIPTION=(ADDRESS=(PROTOCOL=TCP)(HOST=endpoint)(PORT=1521))(CONNECT_DATA=(SID=ORCL)))'
```

You should see output similar to the following.

```
SQL*Plus: Release 11.1.0.7.0 - Production on Wed May 25 15:13:59 2011
```

SQL>

For more information about connecting to an Oracle DB instance, see [Connecting to a DB Instance Running the Oracle Database Engine \(p. 955\)](#).

Deleting Your Sample DB Instance

After you are done exploring the sample DB instance that you created, you should delete the DB instance so that you are no longer charged for it.

To delete a DB instance with no final DB snapshot

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the DB instance that you want to delete.
4. For **Actions**, choose **Delete**.
5. For **Create final snapshot?**, choose **No**, and choose the acknowledgment.
6. Choose **Delete**.

Creating a PostgreSQL DB Instance and Connecting to a Database on a PostgreSQL DB Instance

The easiest way to create a DB instance is to use the RDS console. After you have created the DB instance, you can use standard SQL client utilities to connect to the DB instance such as the pgAdmin utility. In this example, you create a DB instance running the PostgreSQL database engine called west2-postgres1, with a db.m1.small DB instance class, 10 GiB of storage, and automated backups enabled with a retention period of one day.

Important

Before you can create or connect to a DB instance, you must complete the tasks in [Setting Up for Amazon RDS \(p. 5\)](#).

Topics

- [Creating a PostgreSQL DB Instance \(p. 55\)](#)
- [Connecting to a PostgreSQL DB Instance \(p. 64\)](#)
- [Deleting a DB Instance \(p. 68\)](#)

Creating a PostgreSQL DB Instance

The basic building block of Amazon RDS is the DB instance. This environment is where you run your PostgreSQL databases.

Note

A new console interface is available for database creation. Choose either the **New Console** or the **Original Console** instructions based on the console that you are using. The **New Console** instructions are open by default.

New Console

You can create a DB instance running PostgreSQL with the AWS Management Console with **Easy Create** enabled or disabled. With **Easy Create** enabled, you specify only the DB engine type, DB instance size, and DB instance identifier. **Easy Create** uses the default setting for other configuration options. With **Easy Create** not enabled, you specify more configuration options when you create a database, including ones for availability, security, backups, and maintenance.

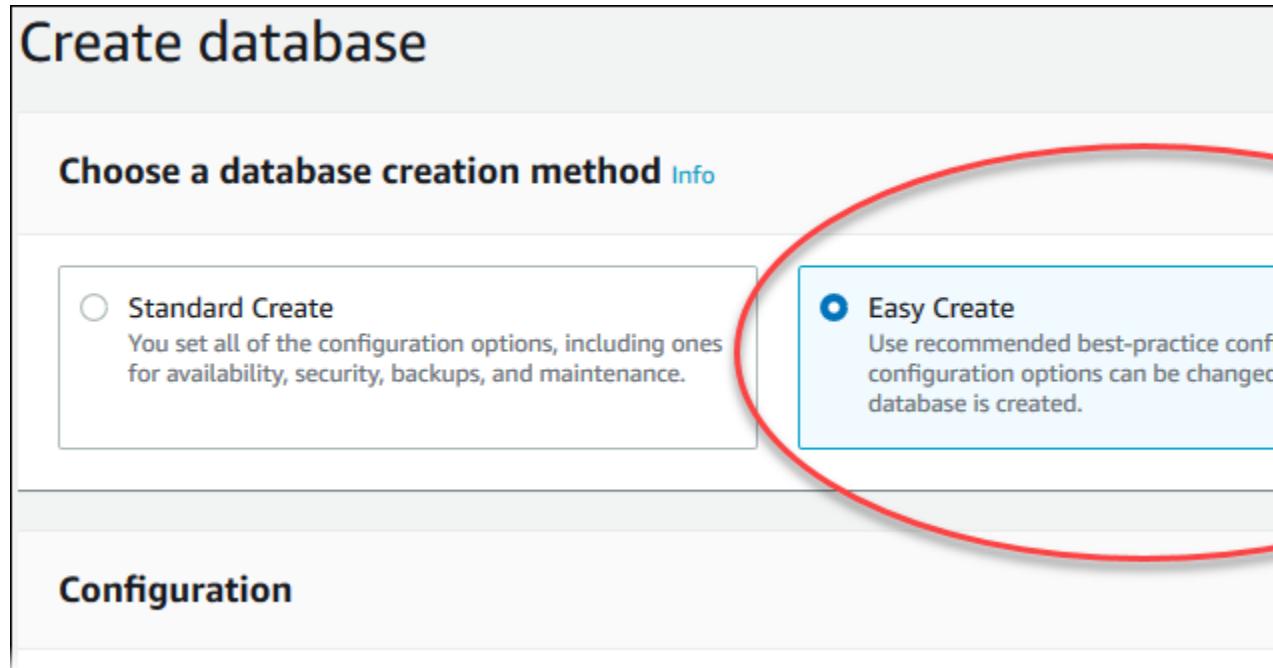
In this example, you use **Easy Create** to create a DB instance running the PostgreSQL database engine with a db.t2.micro DB instance class.

Note

For information about creating a PostgreSQL DB instance with **Easy Create** not enabled, see [Creating a DB Instance Running the PostgreSQL Database Engine \(p. 1257\)](#).

To create a PostgreSQL DB instance with Easy Create enabled

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the Amazon RDS console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
4. Choose **Create database** and make sure that **Easy Create** is chosen.



5. In **Configuration**, choose **PostgreSQL**.
6. For **DB instance size**, choose **Free tier**.
7. For **DB instance identifier**, enter a name for the DB instance, or leave the default name.
8. For **Master username**, enter a name for the master user, or leave the default name.

The **Create database** page should look similar to the following image.

Create database

Choose a database creation method [Info](#)

Standard Create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy Create

Use recommended best-practice configuration options can be changed after the database is created.

Configuration

Engine type [Info](#)

Amazon Aurora



MySQL



MariaDB



PostgreSQL



Oracle

ORACLE®

Microsoft SQL Server



DB instance size

Production

db.r4.large
2 vCPUs
15.25 GiB RAM
500 GiB

Dev/Test

db.m4.large
2 vCPUs
8 GiB RAM
100 GiB

Free tier

db.t2.micro
1 vCPUs
1 GiB RAM
20 GiB

DB instance identifier

Type a name for your DB instance. The name must be unique across all DB instances owned by your AWS account in this Region.

9. To use an automatically generated master password for the DB instance, make sure that the **Auto generate a password** check box is chosen.

To enter your master password, clear the **Auto generate a password** check box, and then enter the same password in **Master password** and **Confirm password**.

10. (Optional) Open **View default settings for Easy create**.

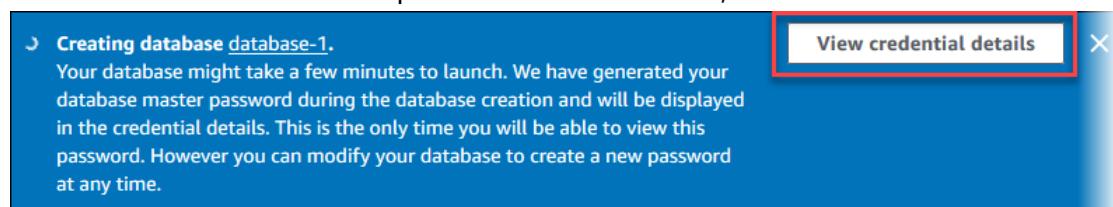
▼ View default settings for Easy create		
Configuration	Value	Editable after database is created
Database Location	Regional	No
Database Features	provisioned	No
Automatic Backups	Enabled	No

You can examine the default settings used when **Easy Create** is enabled. If you want to change one or more settings during database creation, choose **Standard Create** to set them. The **Editable after database creation** column shows which options you can change after database creation. To change a setting with **No** in that column, use **Standard Create**. For settings with **Yes** in that column, you can either use **Standard Create** or modify the DB instance after it's created to change the setting.

11. Choose **Create database**.

If you chose to use an automatically generated password, the **View credential details** button appears on the **Databases** page.

To view the master user name and password for the DB instance, choose **View credential details**.



To connect to the DB instance as the master user, use the user name and password that appear.

Important

You can't view the master user password again. If you don't record it, you might have to change it. If you need to change the master user password after the DB instance is available, you can modify the DB instance to do so. For more information about modifying a DB instance, see [Modifying a DB Instance Running the PostgreSQL Database Engine \(p. 1279\)](#).

12. For **Databases**, choose the name of the new PostgreSQL DB instance.

On the RDS console, the details for new DB instance appear. The DB instance has a status of **creating** until the DB instance is ready to use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and the amount of storage, it can take up to 20 minutes before the new instance is available.

The screenshot shows the 'Summary' tab for a database instance named 'database-1'. The 'Info' column displays the status as 'Creating'. A red circle highlights this status. The 'Engine' column shows 'PostgreSQL'. Other details include 'DB identifier: database-1', 'CPU', 'Class: db.t2.micro', 'Role: Instance', 'Current activity', 'Region & AZ', and a blank 'Tags' section. Below the summary, there are tabs for 'Connectivity & security', 'Monitoring', 'Logs & events', 'Configuration', 'Maintenance & backups', and 'Tags'.

Original Console

To create a DB Instance Running the PostgreSQL DB Engine

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the top right corner of the AWS Management Console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
If the navigation pane is closed, choose the menu icon at the top left to open it.
4. Choose **Create database** to open the **Select engine** page.

Select engine

Engine options

Amazon Aurora
Amazon Aurora

MySQL


MariaDB


PostgreSQL


Oracle
ORACLE

Microsoft SQL Server


PostgreSQL

PostgreSQL is a powerful, open-source object-relational database system with a strong reputation of reliability, stability, and correctness.

- High reliability and stability in a variety of workloads.
- Advanced features to perform in high-volume environments.
- Vibrant open-source community that releases new features multiple times per year.
- Supports multiple extensions that add even more functionality to the database.
- Supports up to 5 Read Replicas per instance, within a single Region or cross-region.
- The most Oracle-compatible open-source database.

Info If you want to create PostgreSQL 11 in the Preview environment, click [here](#)

Aurora global database feature is now available.
This feature is now available in our new database creation flow.
[Try it now](#)

Only enable options eligible for RDS Free Usage Tier [Info](#)

[Cancel](#) **Next**

5. On the **Select engine** page, choose the PostgreSQL icon, and then choose **Next**.
6. Next, the **Use case** page asks if you are planning to use the DB instance you are creating for production. If you are, choose **Production**. If you choose this option, the failover option **Multi-AZ** and the **Provisioned IOPS** storage options are preselected in the following step. Choose **Next** when you are finished.

7. On the **Specify DB Details** page, specify your DB instance information. Choose **Next** when you are finished.

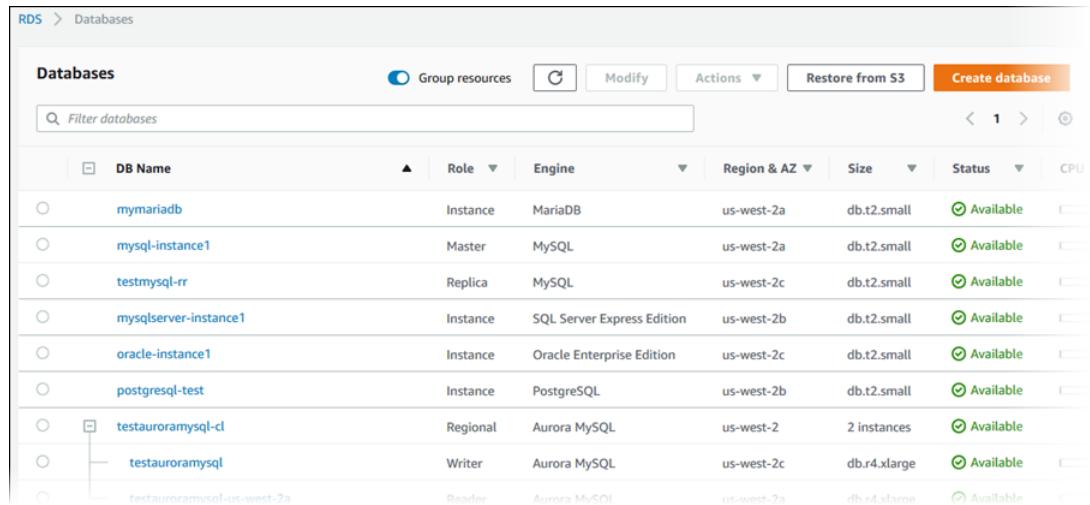
For This Parameter	Do This
License Model	PostgreSQL has only one license model. Choose postgresql-license to use the general license agreement for PostgreSQL.
DB Engine Version	Choose the version of PostgreSQL you want to use.
DB Instance Class	Choose db.t2.small for a configuration that equates to 2 GiB memory, 1 ECU (1 virtual core with 1 ECU), 64-bit platform, and moderate I/O capacity. For more information about all the DB instance class options, see Choosing the DB Instance Class (p. 107) .
Multi-AZ Deployment	Choose Yes to have a standby replica of your DB instance created in another Availability Zone for failover support. We recommend Multi-AZ for production workloads to maintain high availability. For development and testing, you can choose No . For more information, see High Availability (Multi-AZ) for Amazon RDS (p. 144) .
Storage Type	Choose the storage type General Purpose (SSD) . For more information about storage, see Amazon RDS DB Instance Storage (p. 135) .
Allocated Storage	Enter 20 to allocate 20 GiB of storage for your database. In some cases, allocating a higher amount of storage for your DB instance than the size of your database can improve I/O performance. For more information about storage allocation, see Amazon RDS Features .
DB Instance Identifier	Enter a name for the DB instance that is unique for your account in the AWS Region you chose. You can add some intelligence to the name, such as including the AWS Region and DB engine you chose, for example postgreSQL-test .
Master Username	Enter a name using alphanumeric characters to use as the master user name to log on to your DB instance. For information on the default privileges granted to the master user name, see Amazon RDS for PostgreSQL Versions and Extensions (p. 1364)
Master Password and Confirm Password	Enter a password that contains from 8 to 128 printable ASCII characters (excluding /, ", and @) for your master password, then type the password again in the Confirm Password box.

8. On the **Configure Advanced Settings** page, provide additional information that RDS needs to launch the PostgreSQL DB instance. The table shows settings for an example DB instance. Specify your DB instance information, then choose **Create database**.

For This Parameter	Do This
VPC	This setting depends on the platform you are on. If you are a new customer to AWS, choose the default VPC shown. If you are creating a DB instance on the previous E2-Classic platform that does not use a VPC, choose Not in VPC . For more information about VPC, see Amazon Virtual Private Cloud VPCs and Amazon RDS (p. 505) .
Subnet Group	This setting depends on the platform you are on. If you are a new customer to AWS, choose default , which is the default DB subnet group that was created for your account. If you are creating a DB instance on the previous E2-Classic platform and you want your DB instance in a specific VPC, choose the DB subnet group you created for that VPC. For more information about VPC, see Amazon Virtual Private Cloud VPCs and Amazon RDS (p. 505) .
Publicly Accessible	Choose Yes to give the DB instance a public IP address, meaning that it is accessible outside the VPC; otherwise, choose No , so the DB instance is only accessible from inside the VPC. For more information about hiding DB instances from public access, see Hiding a DB Instance in a VPC from the Internet (p. 515) .
Availability Zone	Use the default value of No Preference unless you want to specify an Availability Zone.
VPC Security Group	If you are a new customer to AWS, choose the default VPC. If you created a VPC security group, choose the VPC security group you previously created.
Database Name	Enter a name for your database of up to 63 alpha-numeric characters. If you do not provide a name, the default "postgres" database is created. To create additional databases, connect to the DB instance and use the SQL command <code>CREATE DATABASE</code> . For more information about connecting to the DB instance, see Connecting to a DB Instance Running the PostgreSQL Database Engine (p. 1270) .
Database Port	Specify a port you want to use to access the database. PostgreSQL installations default to port 5432.
DB Parameter Group	Use the default value unless you have created your own parameter group.
Option Group	Use the default value unless you have created your own option group.
Copy Tags To Snapshots	Choose this option to have any DB instance tags copied to a DB snapshot when you create a snapshot. For more information, see Tagging Amazon RDS Resources (p. 177) .

For This Parameter	Do This
Enable Encryption	Choose Yes to enable encryption at rest for this DB instance. For more information, see Encrypting Amazon RDS Resources (p. 429) .
Backup Retention Period	Set the number of days you want automatic backups of your database to be retained. For testing purposes, you can set this value to 1 .
Backup Window	Unless you have a specific time that you want to have your database backup, use the default of No Preference .
Enable Enhanced Monitoring	Choose Yes to enable real-time OS monitoring. Amazon RDS provides metrics in real time for the operating system (OS) that your DB instance runs on. You are only charged for Enhanced Monitoring that exceeds the free tier provided by Amazon CloudWatch Logs.
Monitoring Role	Choose Default to use the default IAM role.
Granularity	Choose 60 to monitor the instance every minute.
Auto Minor Version Upgrade	Choose Enable auto minor version upgrade to enable your DB instance to receive preferred minor DB engine version upgrades automatically when they become available.
Maintenance Window	Choose the 30-minute window in which pending modifications to your DB instance are applied. If the time period doesn't matter, choose No Preference .

- On the final page, choose **Create database**.
- On the Amazon RDS console, the new DB instance appears in the list of DB instances. The DB instance has a status of **creating** until the DB instance is created and ready for use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and store allocated, it could take several minutes for the new instance to be available.



The screenshot shows the 'Databases' section of the Amazon RDS console. The table lists various DB instances with columns for DB Name, Role, Engine, Region & AZ, Size, Status, and CPU. One instance, 'testauroramysql-cl', is selected, indicated by a blue outline around its row. The 'Status' column for this instance shows a green checkmark and the word 'Available'.

DB Name	Role	Engine	Region & AZ	Size	Status	CPU
mymariadb	Instance	MariaDB	us-west-2a	db.t2.small	Available	
mysql-instance1	Master	MySQL	us-west-2a	db.t2.small	Available	
testmysql-rr	Replica	MySQL	us-west-2c	db.t2.small	Available	
mysqlserver-instance1	Instance	SQL Server Express Edition	us-west-2b	db.t2.small	Available	
oracle-instance1	Instance	Oracle Enterprise Edition	us-west-2c	db.t2.small	Available	
postgresql-test	Instance	PostgreSQL	us-west-2b	db.t2.small	Available	
testauroramysql-cl	Regional	Aurora MySQL	us-west-2	2 instances	Available	
testauroramysql	Writer	Aurora MySQL	us-west-2c	db.r4.xlarge	Available	
testauroramysql-us-west-2a	Reader	Aurora MySQL	us-west-2a	db.r4.xlarge	Available	

Connecting to a PostgreSQL DB Instance

After Amazon RDS provisions your DB instance, you can use any standard SQL client application to connect to the instance. The security group that you assigned to the DB instance when you created it must allow access to the DB instance. If you have difficulty connecting to the DB instance, the problem is most often with the access rules you set up in the security group you assigned to the DB instance.

This section shows two ways to connect to a PostgreSQL DB instance. The first example uses pgAdmin, a popular open-source administration and development tool for PostgreSQL. You can download and use pgAdmin without having a local instance of PostgreSQL on your client computer. The second example uses psql, a command line utility that is part of a PostgreSQL installation. To use psql, you must have a PostgreSQL installed on your client computer or have installed the psql client on your machine.

In this example, you connect to a PostgreSQL DB instance using pgAdmin.

Using pgAdmin to Connect to a PostgreSQL DB Instance

To connect to a PostgreSQL DB instance using pgAdmin

1. Find the endpoint (DNS name) and port number for your DB Instance.
 - a. Open the RDS console and then choose **Databases** to display a list of your DB instances.
 - b. Choose the PostgreSQL DB instance name to display its details.
 - c. On the **Connectivity & security** tab, copy the endpoint. Also, note the port number. You need both the endpoint and the port number to connect to the DB instance.

database-1

Summary

DB identifier
database-1

Role
Instance

Connectivity & security Monitoring Logs & events

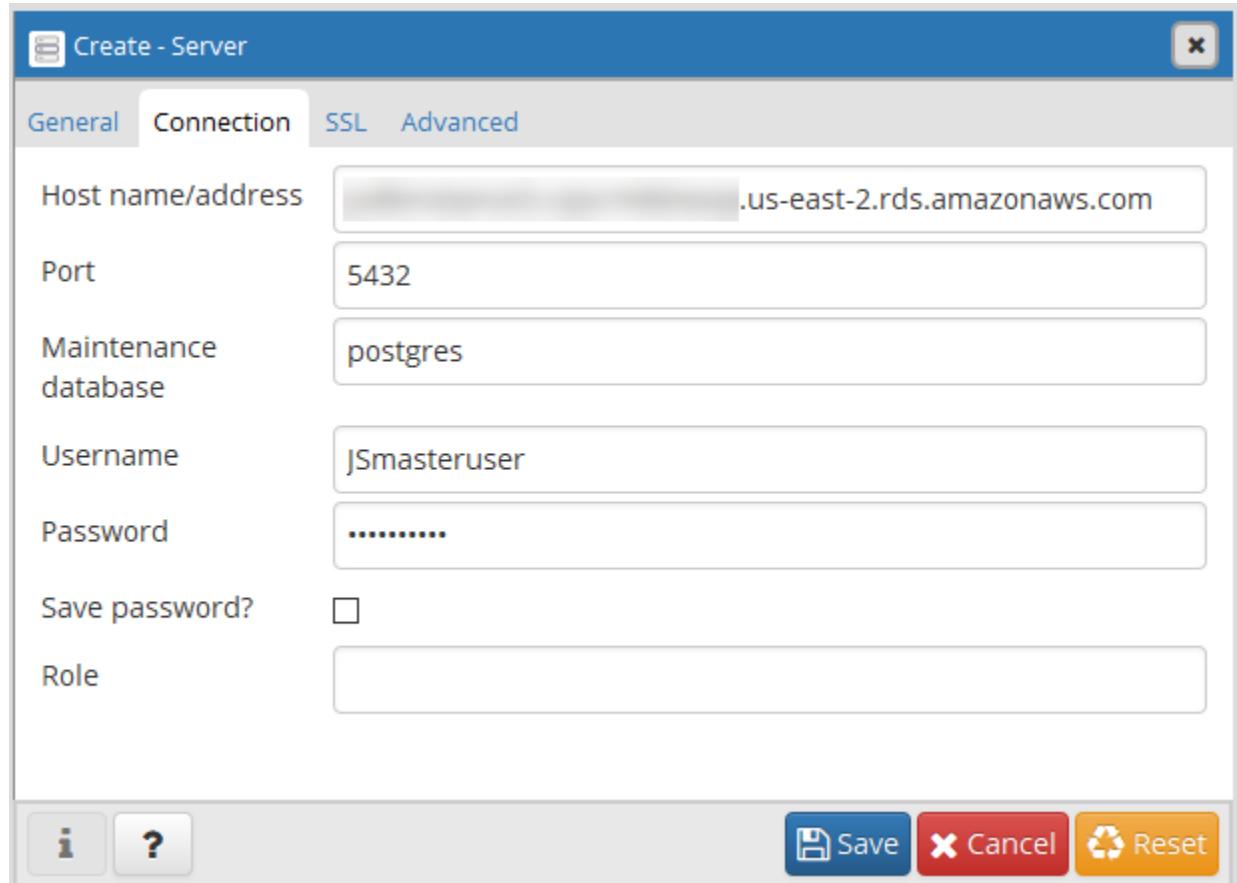
Connectivity & security

Endpoint
database-1.us-west-1.rds.amazonaws.com

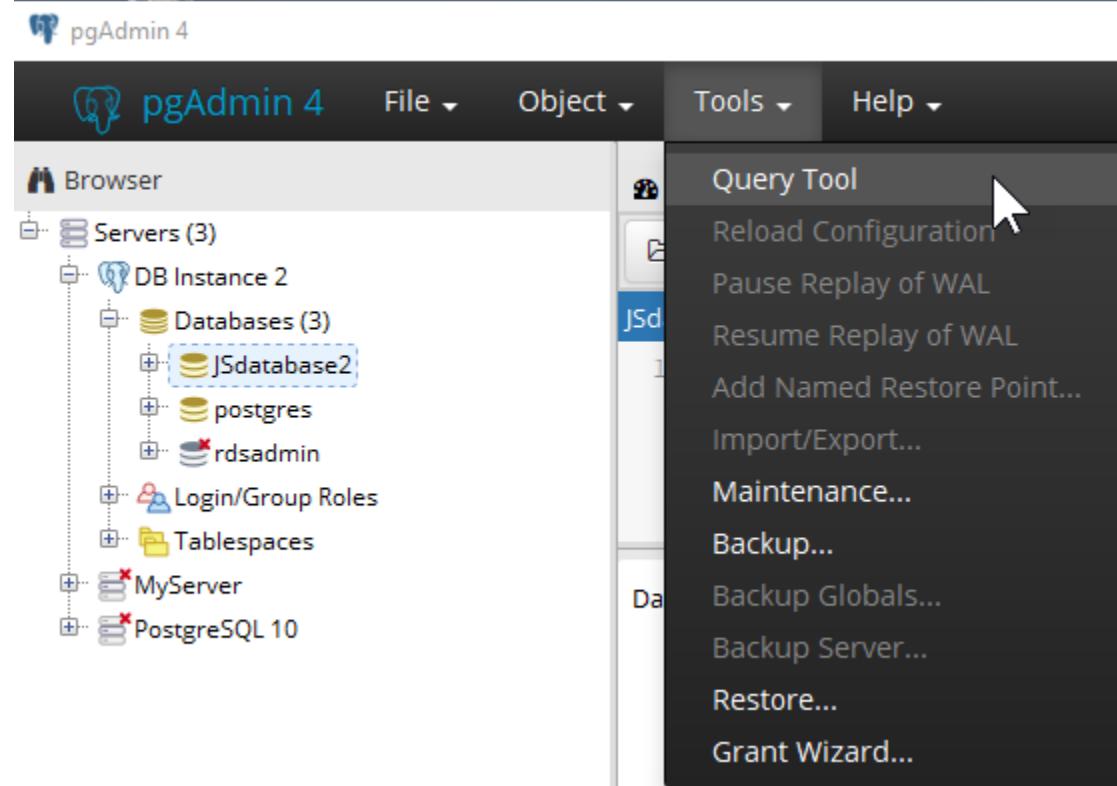
Port
5432

2. Install pgAdmin from <http://www.pgadmin.org/>. You can download and use pgAdmin without having a local instance of PostgreSQL on your client computer.
3. Launch the pgAdmin application on your client computer.
4. Choose **Add Server** from the **File** menu.
5. In the **New Server Registration** dialog box, enter the DB instance endpoint (for example, `mypostgresql.c6c8dntfzzhgv0.us-west-2.rds.amazonaws.com`) in the **Host** box. Don't include the colon or port number as shown on the Amazon RDS console (`mypostgresql.c6c8dntfzzhgv0.us-west-2.rds.amazonaws.com:5432`).

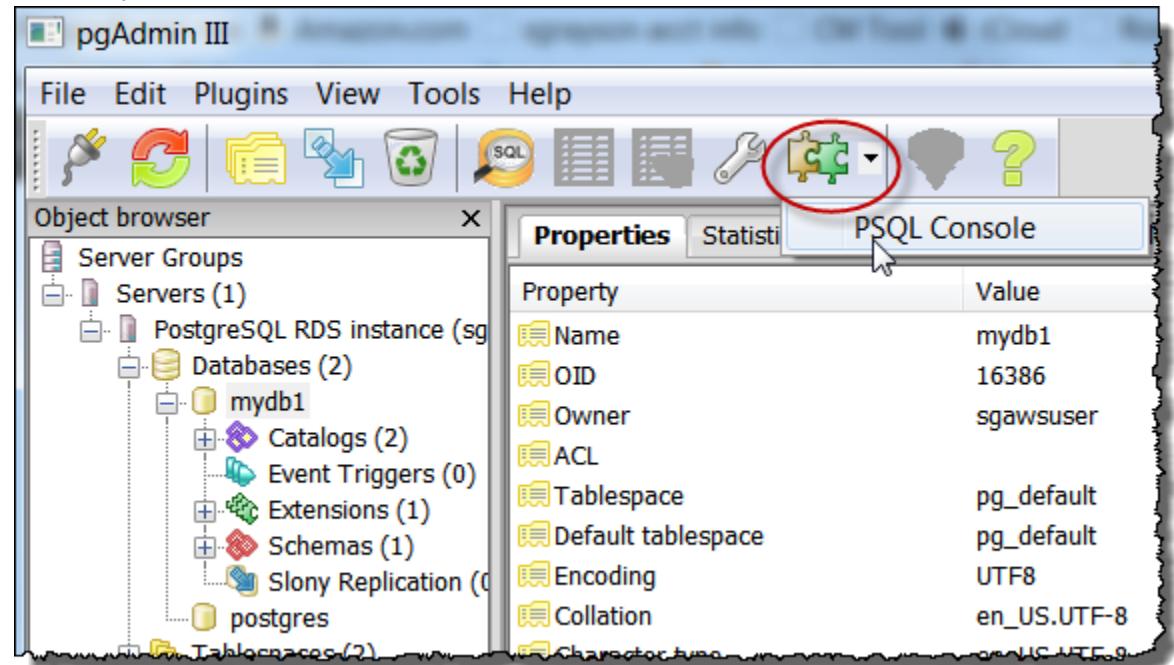
Enter the port you assigned to the DB instance for **Port**. Enter the user name and user password that you entered when you created the DB instance for **Username** and **Password**.



6. Choose **OK**.
7. In the Object browser, expand **Server Groups**. Choose the server (the DB instance) you created, and then choose the database name.



8. Choose the plugin icon and choose **PSQL Console**. The psql command window opens for the default database you created.



9. Use the command window to enter SQL or psql commands. Enter \q to close the window.

Using psql to Connect to a PostgreSQL DB Instance

If your client computer has PostgreSQL installed, you can use a local instance of psql to connect to a PostgreSQL DB instance. To connect to your PostgreSQL DB instance using psql, provide host information and access credentials.

The following format is used to connect to a PostgreSQL DB instance on Amazon RDS.

```
psql --host=<DB instance endpoint> --port=<port> --username=<master user name> --password  
--dbname=<database name>
```

For example, the following command connects to a database called mypgdb on a PostgreSQL DB instance called mypostgresql using fictitious credentials.

```
psql --host=mypostgresql.c6c8mwvfdgv0.us-west-2.rds.amazonaws.com --port=5432 --  
username=awsuser --password --dbname=mypgdb
```

Troubleshooting Connection Issues

By far the most common problem that occurs when attempting to connect to a database on a DB instance is the access rules in the security group assigned to the DB instance. If you used the default DB security group when you created the DB instance, chances are good that the security group did not have the rules that enable you to access the instance. For more information about Amazon RDS security groups, see [Controlling Access with Security Groups \(p. 486\)](#)

The most common error is could not connect to server: Connection timed out. If you receive this error, check that the host name is the DB instance endpoint and that the port number is correct. Check that the security group assigned to the DB instance has the necessary rules to allow access through any firewall your connection may be going through.

Deleting a DB Instance

After you have connected to the sample DB instance that you created, you should delete the DB instance so you are no longer charged for it.

To delete a DB instance with no final DB snapshot

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the DB instance that you want to delete.
4. For **Actions**, choose **Delete**.
5. For **Create final snapshot?**, choose **No**, and select the acknowledgment.
6. Choose **Delete**.

Tutorial: Create a Web Server and an Amazon RDS Database

This tutorial helps you install an Apache web server with PHP, and create a MySQL database. The web server runs on an Amazon EC2 instance using Amazon Linux, and the MySQL database is an Amazon RDS MySQL DB instance. Both the Amazon EC2 instance and the Amazon RDS DB instance run in a virtual private cloud (VPC) based on the Amazon VPC service.

Note

This tutorial works with Amazon Linux and might not work for other versions of Linux such as Ubuntu.

Before you begin this tutorial, you must have a VPC with both public and private subnets, and corresponding security groups. If you don't have these, complete the following tasks in [Tutorial: Create an Amazon VPC for Use with a DB Instance \(p. 521\)](#):

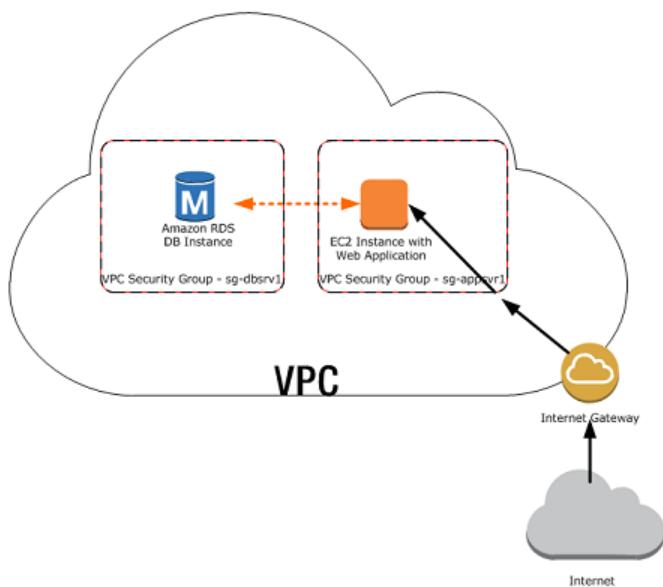
1. [Create a VPC with Private and Public Subnets \(p. 521\)](#)
2. [Create Additional Subnets \(p. 522\)](#)
3. [Create a VPC Security Group for a Public Web Server \(p. 523\)](#)
4. [Create a VPC Security Group for a Private DB Instance \(p. 524\)](#)
5. [Create a DB Subnet Group \(p. 524\)](#)

In the tutorial that follows, you specify the VPC, subnets, and security groups when you create the DB instance. You also specify them when you create the EC2 instance that will host your web server. The VPC, subnets, and security groups are required for the DB instance and the web server to communicate. After the VPC is set up, this tutorial shows you how to create the DB instance and install the web server. You connect your web server to your RDS DB instance in the VPC using the DB instance endpoint.

In this tutorial, you perform the following procedures:

- [Step 1: Create an RDS DB Instance \(p. 70\)](#)
- [Step 2: Create an EC2 Instance and Install a Web Server \(p. 81\)](#)

The following diagram shows the configuration when the tutorial is complete.



Step 1: Create an RDS DB Instance

In this step, you create an Amazon RDS MySQL DB instance that maintains the data used by a web application.

Important

Before you begin this step, you must have a VPC with both public and private subnets, and corresponding security groups. If you don't have these, see [Tutorial: Create an Amazon VPC for Use with a DB Instance \(p. 521\)](#). Complete the steps in [Create a VPC with Private and Public Subnets \(p. 521\)](#), [Create Additional Subnets \(p. 522\)](#), [Create a VPC Security Group for a Public Web Server \(p. 523\)](#), and [Create a VPC Security Group for a Private DB Instance \(p. 524\)](#).

Note

A new console interface is available for database creation. Choose either the **New Console** or the **Original Console** instructions based on the console that you are using. The **New Console** instructions are open by default.

New Console

To launch a MySQL DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the AWS Management Console, choose the AWS Region in which you want to create the DB instance. This example uses the US West (Oregon) Region.
3. In the navigation pane, choose **Databases**.
4. Choose **Create database**.
5. On the **Create database** page, shown following, make sure that the **Standard Create** option is chosen, and then choose **MySQL**.

Create database

Choose a database creation method [Info](#)

Standard Create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy Create

Use recommended best-practice configurations. Some configuration options can be changed after the database is created.

Engine options

Engine type [Info](#)

Amazon Aurora



MySQL



MariaDB

PostgreSQL



Oracle

ORACLE®

Microsoft SQL Server



Edition

MySQL Community

Version [Info](#)

MySQL 5.7.20

6. In the **Templates** section, choose **Dev/Test**.
7. In the **Settings** section, set these values:
 - **DB instance identifier – tutorial-db-instance**

- **Master username** – `tutorial_user`
- **Auto generate a password** – Disable the option
- **Master password** – Choose a password.
- **Confirm password** – Retype the password.

Settings

DB instance identifier [Info](#)
Type a name for your DB instance. The name must be unique cross all DB instances owned by your AWS account in the current AWS Region.

`tutorial-db-instance`

The DB instance identifier is case-insensitive, but is stored as all lowercase (as in "mydbinstance"). Constraints: 1 to 60 alphanumeric characters or hyphens (1 to 15 for SQL Server). First character must be a letter. Can't contain two consecutive hyphens. Can't end with a hyphen.

Credentials Settings

Master username [Info](#)
Type a login ID for the master user of your DB instance.

`tutorial_user`

1 to 16 alphanumeric characters. First character must be a letter

Auto generate a password
Amazon RDS can generate a password for you, or you can specify your own password

Master password [Info](#)
`*****`

Constraints: At least 8 printable ASCII characters. Can't contain any of the following: / (slash), "(double quote) and @ (at sign).

Confirm password [Info](#)
`*****`

8. In the **DB instance size** section, set these values:

- **DB instance performance type** – **Burstable**
- **DB instance class** – **db.t2.small**

The screenshot shows the 'DB instance size' configuration section. Under 'DB instance performance type', 'Burstable' is selected. Under 'DB instance class', 'db.t2.small' is chosen, showing details: 1 vCPUs, 2 GiB RAM, and Not EBS Optimized. There is also an option to 'Include previous generation classes'.

9. In the **Storage** and **Availability & durability** sections, use the default values.
10. In the **Connectivity** section, open **Additional connectivity configuration** and set these values:
 - **Virtual Private Cloud (VPC)** – Choose an existing VPC with both public and private subnets, such as the `tutorial-vpc` (`vpc-identifier`) created in [Create a VPC with Private and Public Subnets \(p. 521\)](#)

Note

The VPC must have subnets in different Availability Zones.

- **Subnet group** – The DB subnet group for the VPC, such as the `tutorial-db-subnet-group` created in [Create a DB Subnet Group \(p. 524\)](#)
- **Publicly accessible** – No
- **VPC security groups** – Choose an existing VPC security group that is configured for private access, such as the `tutorial-db-securitygroup` created in [Create a VPC Security Group for a Private DB Instance \(p. 524\)](#).

Remove other security groups, such as the default security group, by choosing the X associated with each.

- **Availability zone – No Preference**
- **Database port – 3306**

Connectivity

Virtual Private Cloud (VPC) [Info](#)
VPC that defines the virtual networking environment for this DB instance.

tutorial-vpc (vpc-) ▾
Only VPCs with a corresponding DB subnet group are listed.

i After a database is created, you can't change the VPC selection.

Subnet group [Info](#)
DB subnet group that defines which subnets and IP ranges the DB instance can use in the VPC you selected.

tutorial-db-subnet-group ▾

Publicly accessible [Info](#)

Yes
Amazon EC2 instances and devices outside the VPC can connect to your database. Choose one or more VPC security groups that specify which EC2 instances and devices inside the VPC can connect to the database.

No
RDS will not assign a public IP address to the database. Only Amazon EC2 instances and devices inside the VPC can connect to your database.

VPC security group
Choose one or more RDS security groups to allow access to your database. Ensure that the security group rules allow incoming traffic from EC2 instances and devices outside your VPC. (Security groups are required for publicly accessible databases.)

Choose existing
Choose existing VPC security groups

Create new
Create new VPC security group

Existing VPC security groups

Choose VPC security groups ▾
tutorial-db-securitygroup X

Availability zone [Info](#)

No preference ▾

Database port [Info](#)
TCP/IP port the database will use for application connections.

3306 ▾

11. Open the **Additional configuration** section, and enter **sample** for **Initial database name**. Keep the default settings for the other options.
 12. To create your Amazon RDS MySQL DB instance, choose **Create database**.
- Your new DB instance appears in the **Databases** list with the status **Creating**.
13. Wait for the **Status** of your new DB instance to show as **Available**. Then choose the DB instance name to show its details.
 14. In the **Connectivity & security** section, view the **Endpoint** and **Port** of the DB instance.

tutorial-db-instance

Summary

DB identifier	CPU
tutorial-db-instance	
Role	Current activity
Instance	 0 Connections

Connectivity & security Monitoring Logs & events Configuration

Connectivity & security

Endpoint & port

Endpoint	tutorial-db-instance.  .us-west-2.rds.amazonaws.com
Port	3306

Note the endpoint and port for your DB instance. You use this information to connect your web server to your RDS DB instance.

To make sure your RDS MySQL DB instance is as secure as possible, verify that sources outside of the VPC can't connect to your RDS MySQL DB instance.

Original Console

To launch a MySQL DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.

2. In the top-right corner of the AWS Management Console, choose the AWS Region in which you want to create the DB instance. This example uses the US West (Oregon) region.
3. In the navigation pane, choose **Databases**.
If the navigation pane is closed, choose the menu icon at the top left to open it.
4. Choose **Create database** to open the **Select engine** page.
5. On the **Select engine** page, shown following, choose **MySQL**, and then choose **Next**.

The screenshot shows the 'Select engine' page with the title 'Select engine' at the top. Below it is a section titled 'Engine options' containing six database engine choices: Amazon Aurora, MySQL, MariaDB, PostgreSQL, Oracle, and Microsoft SQL Server. The MySQL option is selected, indicated by a blue border around its box and a checked radio button. The MySQL logo, featuring a stylized orange and blue circular design, is displayed next to the name. Below the engine options, there is a detailed description of MySQL, highlighting its popularity as an open-source database and its features like automated backup and cross-region read replicas. At the bottom of the page are two buttons: 'Cancel' and a prominent orange 'Next' button.

6. On the **Choose use case** page, choose **Dev/Test – MySQL**, and then choose **Next**.
7. On the **Specify DB details** page, shown following, set these values:
 - **License model:** Use the default value.
 - **DB engine version:** Use the default value.
 - **DB instance class:** db.t2.small
 - **Multi-AZ deployment:** No
 - **Storage type:** General Purpose (SSD)
 - **Allocated storage:** 20 GiB
 - **DB instance identifier:** tutorial-db-instance
 - **Master username:** tutorial_user
 - **Master password:** Choose a password.

- **Confirm password:** Retype the password.

Specify DB details

Instance specifications
Estimate your monthly costs for the DB Instance using the AWS Simple Monthly Calculator.

DB engine: MySQL Community Edition

License model info: general-public-license

DB engine version info: mysql 5.6.37

Known Issues/Limitations
Review the Known Issues/Limitations to learn about potential compatibility issues with specific database versions.

Free tier
The Amazon RDS Free Tier provides a single db.t2.micro instance as well as up to 20 GB of storage, allowing new AWS customers to gain hands-on experience with Amazon RDS. Learn more about the RDS Free Tier and the instance restrictions [here](#).

Only enable options eligible for RDS Free Usage Tier info

DB instance class info: db.t2.small — 1 vCPU, 2 GiB RAM

Multi-AZ deployment info:
 Create replica in different zone: Creates a replica in a different Availability Zone (AZ) to provide data redundancy, eliminate I/O freezes, and minimize latency spikes during system backups.
 No

Storage type info: General Purpose (SSD)

Allocated storage: 20 GB
(Minimum: 20 GB, Maximum: 16384 GB) Higher allocated storage may improve IOPS performance.

Settings

DB instance identifier info: Specify a name that is unique for all DB instances owned by your AWS account in the current region.
tutorial-db-instance

DB instance identifier is case insensitive, but stored as all lower-case, as in "mydbinstance".

Master username info: Specify an alphanumeric string that defines the login ID for the master user.
tutorial_user

Master Username must start with a letter.

Master password info: Confirm password info:
Master Password must be at least eight characters long, as in "mypassword".

Cancel **Previous** **Next**

8. Choose **Next** and set the following values in the **Configure advanced settings** page:

- **Virtual Private Cloud (VPC):** Choose an existing VPC with both public and private subnets, such as the `tutorial-vpc` (`vpc-identifier`) created in [Create a VPC with Private and Public Subnets \(p. 521\)](#)

Note

The VPC must have subnets in different Availability Zones.

- **Subnet group:** The DB subnet group for the VPC, such as the `tutorial-db-subnet-group` created in [Create a DB Subnet Group \(p. 524\)](#)
- **Public accessibility:** No
- **Availability zone:** No Preference
- **VPC security groups:** Choose an existing VPC security group that is configured for private access, such as the `tutorial-db-securitygroup` created in [Create a VPC Security Group for a Private DB Instance \(p. 524\)](#).

Remove other security groups, such as the default security group, by choosing the **X** associated with each.

- **Database name:** sample

Leave the default settings for the other options.

Configure advanced settings

Network & Security

Virtual Private Cloud (VPC) [Info](#)

VPC defines the virtual networking environment for this DB instance.

tutorial-vpc ([REDACTED])



Only VPCs with a corresponding DB subnet group are listed.

Subnet group [Info](#)

DB subnet group that defines which subnets and IP ranges the DB instance can use in the VPC you selected.

tutorial-db-subnet-group



Public accessibility [Info](#)

Yes

EC2 instances and devices outside of the VPC hosting the DB instance will connect to the DB instances. You must also select one or more VPC security groups that specify which EC2 instances and devices can connect to the DB instance.

No

DB instance will not have a public IP address assigned. No EC2 instance or devices outside of the VPC will be able to connect.

Availability zone [Info](#)

No preference



VPC security groups

Security groups have rules authorizing connections from all the EC2 instances and devices that need to access the DB instance.

Create new VPC security group

Choose existing VPC security groups

Choose VPC security groups



tutorial-db-securitygroup X

Database options

Database name

sample

Note: if no database name is specified then no initial MySQL database will be created on the DB Instance.

9. To create your Amazon RDS MySQL DB instance, choose **Create database**.
10. On the next page, choose **View DB instances details** to view your RDS MySQL DB instance.
11. Wait for the **DB instance status** of your new DB instance to show as **available**. Then scroll to the **Connect** section, shown following.

The screenshot shows the 'Connect' section of the Amazon RDS console. It displays the endpoint as 'tutorial-db-instance.rds.amazonaws.com' and the port as '3306'. The 'Publicly accessible' option is set to 'No'. Below this, the 'Security group rules (2)' section is shown, listing two inbound rules from the 'tutorial-db-securitygroup'. The first rule is a 'Security Group - Inbound' and the second is a 'CIDR/IP - Outbound' rule for 0.0.0.0/0.

Make note of the endpoint and port for your DB instance. You will use this information to connect your web server to your RDS DB instance.

To make sure your RDS MySQL DB instance is as secure as possible, verify that sources outside of the VPC cannot connect to your RDS MySQL DB instance.

Next Step

[Step 2: Create an EC2 Instance and Install a Web Server \(p. 81\)](#)

Step 2: Create an EC2 Instance and Install a Web Server

In this step you create a web server to connect to the Amazon RDS DB instance that you created in [Step 1: Create an RDS DB Instance \(p. 70\)](#).

Launch an EC2 Instance

First you create an Amazon EC2 instance in the public subnet of your VPC.

To launch an EC2 instance

1. Sign in to the AWS Management Console and open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. Choose **EC2 Dashboard**, and then choose **Launch Instance**, as shown following.

Resources

You are using the following Amazon EC2 resources in the US East (N. Virginia) region:

2 Running Instances
0 Dedicated Hosts
2 Volumes
2 Key Pairs
0 Placement Groups

1 Elastic IPs
0 Snapshots
0 Load Balancers
24 Security Groups

EC2 Spot. Save up to 90% off On-Demand Prices. Turbo Boost your Workloads. [Get started with Amazon EC2 Spot Instances.](#)

Create Instance

To start using Amazon EC2 you will want to launch a virtual server, known as an Amazon EC2 instance.

[Launch Instance](#)

Note: Your instances will launch in the US East (N. Virginia) region

Service Health

Scheduled Events

3. Choose the **Amazon Linux AMI**, as shown following.

Step 1: Choose an Amazon Machine Image (AMI)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. You can select an AMI provided by AWS, our user community, or the AWS Marketplace; or you can select one of your own AMIs.

Cancel and Exit

Search for an AMI by entering a search term e.g. "Windows"

Quick Start

My AMIs

AWS Marketplace

Community AMIs

Free tier only ⓘ

Amazon Linux 2 AMI (HVM), SSD Volume Type - ami-061392db613a6357b (64-bit x86) / ami-062ce7f8c1e7ffd3c (64-bit Arm)

Amazon Linux 2 comes with five years support. It provides Linux kernel 4.14 tuned for optimal performance on Amazon EC2, systemd 219, GCC 7.3, Glibc 2.26, Binutils 2.29.1, and the latest software packages through extras.

Root device type: ebs Virtualization type: hvm ENA Enabled: Yes

Select

64-bit (x86)

 64-bit (Arm)

Amazon Linux AMI 2018.03.0 (HVM), SSD Volume Type - ami-01e24be29428c15b2

The Amazon Linux AMI is an EBS-backed, AWS-supported image. The default image includes AWS command line tools, Python, Ruby, Perl, and Java. The repositories include Docker, PHP, MySQL, PostgreSQL, and other packages.

Root device type: ebs Virtualization type: hvm ENA Enabled: Yes

Select

64-bit (x86)

 64-bit (Arm)

Red Hat Enterprise Linux 8 (HVM), SSD Volume Type - ami-079596bf7a949ddf8 (64-bit x86) / ami-07a968a2c17fb48b (64-bit Arm)

Red Hat

Select

64-bit (x86)

 64-bit (Arm)

Important

Don't choose **Amazon Linux 2 AMI** because it doesn't have the software packages required for this tutorial.

4. Choose the **t2.small** instance type, as shown following, and then choose **Next: Configure Instance Details**.

Step 2: Choose an Instance Type

Amazon EC2 provides a wide selection of instance types optimized to fit different use cases. Instances are virtual servers that can run applications. They have varying combinations of CPU, memory, storage, and networking capacity, and give you the flexibility to choose the appropriate mix of resources for your applications. [Learn more](#) about instance types and how they can meet your computing needs.

Currently selected: t2.small (Variable ECUs, 1 vCPUs, 2.5 GHz, Intel Xeon Family, 2 GiB memory, EBS only)								
Family	Type	vCPUs	Memory (GiB)	Instance Storage (GB)	EBS-Optimized Available	Network Performance	IPv6 Support	
General purpose	t2.nano	1	0.5	EBS only	-	Low to Moderate	Yes	
General purpose	t2.micro <small>Free tier eligible</small>	1	1	EBS only	-	Low to Moderate	Yes	
<input checked="" type="checkbox"/> General purpose	t2.small	1	2	EBS only	-	Low to Moderate	Yes	
General purpose	t2.medium	2	4	EBS only	-	Low to Moderate	Yes	
General purpose	t2.large	2	8	EBS only	-	Low to Moderate	Yes	
General purpose	t2.xlarge	4	16	EBS only	-	Moderate	Yes	
General purpose	t2.2xlarge	8	32	EBS only	-	Moderate	Yes	
General purpose	t3a.nano	2	0.5	EBS only	Yes	Up to 5 Gigabit	Yes	
General purpose	t3a.micro	2	1	EBS only	Yes	Up to 5 Gigabit	Yes	
General purpose	t3a.small	2	2	EBS only	Yes	Up to 5 Gigabit	Yes	
General purpose	t3a.medium	2	4	EBS only	Yes	Up to 5 Gigabit	Yes	
General purpose	t3a.large	2	8	EBS only	Yes	Up to 5 Gigabit	Yes	

[Cancel](#) [Previous](#) **Review and Launch** [Next: Configure Instance Details](#)

5. On the **Configure Instance Details** page, shown following, set these values and leave the other values as their defaults:

- **Network:** Choose the VPC with both public and private subnets that you chose for the DB instance, such as the `tutorial-vpc` (`vpc-identifier`) created in [Create a VPC with Private and Public Subnets \(p. 521\)](#).
- **Subnet:** Choose an existing public subnet, such as `subnet-identifier | Tutorial public | us-west-2a` created in [Create a VPC Security Group for a Public Web Server \(p. 523\)](#).
- **Auto-assign Public IP:** Choose **Enable**.

Step 3: Configure Instance Details

Configure the instance to suit your requirements. You can launch multiple instances from the same AMI, request Spot instances to take advantage of the lower pricing, assign an access management role to the instance, and more.

Number of Instances: 1 [Launch into Auto Scaling Group](#)

Purchasing option: Request Spot Instances

Network: vpc- tutorial-vpc [Create new VPC](#)
Subnet: subnet- Tutorial public | us-west-2a [Create new subnet](#)
 249 IP Addresses available

Auto-assign Public IP: Enable

Placement group: Add instance to placement group

Capacity Reservation: Open [Create new Capacity Reservation](#)

IAM role: None [Create new IAM role](#)

Shutdown behavior: Stop

Enable termination protection: Protect against accidental termination

Monitoring: Enable CloudWatch detailed monitoring
Additional charges apply.

Tenancy: Shared - Run a shared hardware instance
Additional charges will apply for dedicated tenancy.

Elastic Inference: Add an Elastic Inference accelerator
Additional charges apply.

T2/T3 Unlimited: Enable
Additional charges may apply

[Cancel](#) [Previous](#) [Review and Launch](#) [Next: Add Storage](#)

6. Choose **Next: Add Storage**.
7. On the **Add Storage** page, keep the default values and choose **Next: Add Tags**.
8. On the **Add Tags** page, shown following, choose **Add Tag**, then enter **Name** for **Key** and enter **tutorial-web-server** for **Value**.

Step 5: Add Tags

A tag consists of a case-sensitive key-value pair. For example, you could define a tag with key = Name and value = Webserver.
 A copy of a tag can be applied to volumes, instances or both.
 Tags will be applied to all instances and volumes. [Learn more](#) about tagging your Amazon EC2 resources.

Key	(127 characters maximum)	Value	(255 characters maximum)	Instances	Volumes
Name	tutorial-web-server	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
Add another tag (Up to 50 tags maximum)					

[Cancel](#) [Previous](#) [Review and Launch](#) [Next: Configure Security Group](#)

9. Choose **Next: Configure Security Group**.
10. On the **Configure Security Group** page, shown following, choose **Select an existing security group**, and then choose an existing security group, such as the **tutorial-securitygroup** created in [Create a VPC Security Group for a Public Web Server \(p. 523\)](#). The security group must include inbound rules for SSH and HTTP access.

Step 6: Configure Security Group

A security group is a set of firewall rules that control the traffic for your instance. On this page, you can add rules to allow specific traffic to reach your instance. For example, if you want to set up a web server and allow Internet traffic to reach your instance, add rules that allow unrestricted access to the HTTP and HTTPS ports. You can create a new security group or select from an existing one below. [Learn more](#) about Amazon EC2 security groups.

Assign a security group: Create a new security group
 Select an existing security group

Security Group ID	Name	Description	Actions
sg-[REDACTED]	default	default VPC security group	Copy to new
sg-[REDACTED]	tutorial-db-securitygroup	Tutorial DB Instance Security Group	Copy to new
<input checked="" type="checkbox"/> sg-[REDACTED]	tutorial-securitygroup	Tutorial Security Group	Copy to new

Inbound rules for sg-0ef508f81f84a5764 (Selected security groups: sg-0ef508f81f84a5764)				
Type	Protocol	Port Range	Source	Description
HTTP	TCP	80	[REDACTED]	
SSH	TCP	22	[REDACTED]	

[Cancel](#) [Previous](#) [Review and Launch](#)

11. Choose Review and Launch.

12. On the Review Instance Launch page, shown following, verify your settings and then choose Launch.

Step 7: Review Instance Launch

Please review your instance launch details. You can go back to edit changes for each section. Click **Launch** to assign a key pair to your instance and complete the launch process.

⚠ Your instance configuration is not eligible for the free usage tier

To launch an instance that's eligible for the free usage tier, check your AMI selection, instance type, configuration options, or storage devices. Learn more about [free usage tier](#) eligibility and usage restrictions.

[Don't show me this again](#)

AMI Details [Edit AMI](#)

	Amazon Linux 2 AMI (HVM), SSD Volume Type - ami-061392db613a6357b
Free tier eligible	Amazon Linux 2 comes with five years support. It provides Linux kernel 4.14 tuned for optimal performance on Amazon EC2, systemd 219, GCC 7.3, Glibc 2.26, Binutils 2.29.1, and the latest software packages through extras.
Root Device Type:	ebs
Virtualization type:	hvm

Instance Type [Edit instance type](#)

Instance Type	ECUs	vCPUs	Memory (GiB)	Instance Storage (GB)	EBS-Optimized Available	Network Performance
t2.small	Variable	1	2	EBS only	-	Low to Moderate

Security Groups [Edit security groups](#)

Security Group ID	Name	Description
sg-[REDACTED]	tutorial-securitygroup	Tutorial Security Group

All selected security groups inbound rules

Type	Protocol	Port Range	Source	Description
HTTP	TCP	80	[REDACTED]	
SSH	TCP	22	[REDACTED]	

Instance Details [Edit instance details](#)

AMI	Amazon Linux 2 AMI (HVM), SSD Volume Type - ami-061392db613a6357b
Instance Type	t2.small
Security Group	tutorial-securitygroup
Key Name	[REDACTED]
Block Device Mapping	[REDACTED]
Network Interfaces	[REDACTED]
Monitoring	Enabled
Tags	[REDACTED]

[Cancel](#) [Previous](#) [Launch](#)

13. On the Select an existing key pair or create a new key pair page, shown following, choose Create a new key pair and set Key pair name to tutorial-key-pair. Choose Download Key Pair, and then save the key pair file on your local machine. You use this key pair file to connect to your EC2 instance.

Select an existing key pair or create a new key pair X

A key pair consists of a **public key** that AWS stores, and a **private key file** that you store. Together, they allow you to connect to your instance securely. For Windows AMIs, the private key file is required to obtain the password used to log into your instance. For Linux AMIs, the private key file allows you to securely SSH into your instance.

Note: The selected key pair will be added to the set of keys authorized for this instance. Learn more about [removing existing key pairs from a public AMI](#).

Key pair name

You have to download the **private key file** (*.pem file) before you can continue.
Store it in a secure and accessible location. You will not be able to download the file again after it's created.

[Cancel](#) Launch Instances

14. To launch your EC2 instance, choose **Launch Instances**. On the **Launch Status** page, shown following, note the identifier for your new EC2 instance, for example: i-0288d65fd4470b6a9.

Launch Status

Your instances are now launching
The following instance launches have been initiated: **I-0288d65fd4470b6a9** [View launch log](#)

Get notified of estimated charges
Create billing alerts to get an email notification when estimated charges on your AWS bill exceed an amount you define (for example, if you exceed the free usage tier).

How to connect to your instances

Your instances are launching, and it may take a few minutes until they are in the **running** state, when they will be ready for you to use. Usage hours on your new instances will start immediately and continue to accrue until you stop or terminate your instances.

Click [View Instances](#) to monitor your instances' status. Once your instances are in the **running** state, you can [connect](#) to them from the Instances screen. [Find out](#) how to connect to your instances.

▼ Here are some helpful resources to get you started

- [How to connect to your Linux instance](#)
- [Learn about AWS Free Usage Tier](#)
- [Amazon EC2: User Guide](#)
- [Amazon EC2: Discussion Forum](#)

While your instances are launching you can also

- [Create status check alarms](#) to be notified when these instances fail status checks. (Additional charges may apply)
- [Create and attach additional EBS volumes](#) (Additional charges may apply)
- [Manage security groups](#)

[View Instances](#)

15. To find your instance, choose **[View Instances](#)**.
16. Wait until **Instance Status** for your instance reads as **running** before continuing.

Install an Apache Web Server with PHP

Next you connect to your EC2 instance and install the web server.

To connect to your EC2 instance and install the Apache web server with PHP

1. To connect to the EC2 instance that you created earlier, follow the steps in [Connect to Your Linux Instance](#).
2. To get the latest bug fixes and security updates, update the software on your EC2 instance by using the following command:

Note

The `-y` option installs the updates without asking for confirmation. To examine updates before installing, omit this option.

```
[ec2-user ~]$ sudo yum update -y
```

3. After the updates complete, install the Apache web server with the PHP software package using the **yum install** command, which installs multiple software packages and related dependencies at the same time.

```
[ec2-user ~]$ sudo yum install -y httpd24 php56 php56-mysqlnd
```

Note

If you receive the error `No package package-name available`, then your instance was not launched with the Amazon Linux AMI (perhaps you are using the Amazon Linux 2 AMI instead). You can view your version of Amazon Linux with the following command.

```
cat /etc/system-release
```

For more information, see [Updating Instance Software](#).

4. Start the web server with the command shown following.

```
[ec2-user ~]$ sudo service httpd start
```

You can test that your web server is properly installed and started by entering the public DNS name of your EC2 instance in the address bar of a web browser, for example: `http://ec2-42-8-168-21.us-west-1.compute.amazonaws.com`. If your web server is running, then you see the Apache test page. If you don't see the Apache test page, then verify that your inbound rules for the VPC security group that you created in [Tutorial: Create an Amazon VPC for Use with a DB Instance \(p. 521\)](#) include a rule allowing HTTP (port 80) access for the IP address you use to connect to the web server.

Note

The Apache test page appears only when there is no content in the document root directory, `/var/www/html`. After you add content to the document root directory, your content appears at the public DNS address of your EC2 instance instead of the Apache test page.

5. Configure the web server to start with each system boot using the `chkconfig` command.

```
[ec2-user ~]$ sudo chkconfig httpd on
```

To allow `ec2-user` to manage files in the default root directory for your Apache web server, you need to modify the ownership and permissions of the `/var/www` directory. In this tutorial, you add a group named `www` to your EC2 instance, and then you give that group ownership of the `/var/www` directory and add write permissions for the group. Any members of that group can then add, delete, and modify files for the web server.

To set file permissions for the Apache web server

1. Add the `www` group to your EC2 instance with the following command.

```
[ec2-user ~]$ sudo groupadd www
```

2. Add the `ec2-user` user to the `www` group.

```
[ec2-user ~]$ sudo usermod -a -G www ec2-user
```

3. To refresh your permissions and include the new www group, log out.

```
[ec2-user ~]$ exit
```

4. Log back in again and verify that the www group exists with the groups command.

```
[ec2-user ~]$ groups  
ec2-user wheel www
```

5. Change the group ownership of the /var/www directory and its contents to the www group.

```
[ec2-user ~]$ sudo chgrp -R www /var/www
```

6. Change the directory permissions of /var/www and its subdirectories to add group write permissions and set the group ID on subdirectories created in the future.

```
[ec2-user ~]$ sudo chmod 2775 /var/www  
[ec2-user ~]$ find /var/www -type d -exec sudo chmod 2775 {} +
```

7. Recursively change the permissions for files in the /var/www directory and its subdirectories to add group write permissions.

```
[ec2-user ~]$ find /var/www -type f -exec sudo chmod 0664 {} +
```

Connect your Apache web server to your RDS DB instance

Next, you add content to your Apache web server that connects to your Amazon RDS DB instance.

To add content to the Apache web server that connects to your RDS DB instance

1. While still connected to your EC2 instance, change the directory to /var/www and create a new subdirectory named inc.

```
[ec2-user ~]$ cd /var/www  
[ec2-user ~]$ mkdir inc  
[ec2-user ~]$ cd inc
```

2. Create a new file in the inc directory named dbinfo.inc, and then edit the file by calling nano (or the editor of your choice).

```
[ec2-user ~]$ >dbinfo.inc  
[ec2-user ~]$ nano dbinfo.inc
```

3. Add the following contents to the dbinfo.inc file, where *db_instance_endpoint* is the endpoint of your RDS MySQL DB instance, without the port, and *master password* is the master password for your RDS MySQL DB instance.

Note

Placing the user name and password information in a folder that is not part of the document root for your web server reduces the possibility of your security information being exposed.

```
<?php

define('DB_SERVER', 'db_instance_endpoint');
define('DB_USERNAME', 'tutorial_user');
define('DB_PASSWORD', 'master password');
define('DB_DATABASE', 'sample');

?>
```

4. Save and close the dbinfo.inc file.
5. Change the directory to /var/www/html.

```
[ec2-user ~]$ cd /var/www/html
```

6. Create a new file in the html directory named SamplePage.php, and then edit the file by calling nano (or the editor of your choice).

```
[ec2-user ~]$ >SamplePage.php
[ec2-user ~]$ nano SamplePage.php
```

7. Add the following contents to the SamplePage.php file:

Note

Placing the user name and password information in a folder that is not part of the document root for your web server reduces the possibility of your security information being exposed.

```
<?php include "../inc/dbinfo.inc"; ?>
<html>
<body>
<h1>Sample page</h1>
<?php

/* Connect to MySQL and select the database. */
$connection = mysqli_connect(DB_SERVER, DB_USERNAME, DB_PASSWORD);

if (mysqli_connect_errno()) echo "Failed to connect to MySQL: " .
mysqli_connect_error();

$dbname = mysqli_select_db($connection, DB_DATABASE);

/* Ensure that the EMPLOYEES table exists. */
VerifyEmployeesTable($connection, DB_DATABASE);

/* If input fields are populated, add a row to the EMPLOYEES table. */
$employee_name = htmlentities($_POST['NAME']);
$employee_address = htmlentities($_POST['ADDRESS']);

if (strlen($employee_name) || strlen($employee_address)) {
    AddEmployee($connection, $employee_name, $employee_address);
}
?>
```

```

<!-- Input form -->
<form action=<?PHP echo $_SERVER['SCRIPT_NAME'] ?>" method="POST">
  <table border="0">
    <tr>
      <td>NAME</td>
      <td>ADDRESS</td>
    </tr>
    <tr>
      <td>
        <input type="text" name="NAME" maxlength="45" size="30" />
      </td>
      <td>
        <input type="text" name="ADDRESS" maxlength="90" size="60" />
      </td>
      <td>
        <input type="submit" value="Add Data" />
      </td>
    </tr>
  </table>
</form>

<!-- Display table data. -->
<table border="1" cellpadding="2" cellspacing="2">
  <tr>
    <td>ID</td>
    <td>NAME</td>
    <td>ADDRESS</td>
  </tr>
<?php

$result = mysqli_query($connection, "SELECT * FROM EMPLOYEES");

while($query_data = mysqli_fetch_row($result)) {
  echo "<tr>";
  echo "<td>",$query_data[0], "</td>",
        "<td>",$query_data[1], "</td>",
        "<td>",$query_data[2], "</td>";
  echo "</tr>";
}
?>

</table>

<!-- Clean up. -->
<?php

mysqli_free_result($result);
mysqli_close($connection);

?>

</body>
</html>

<?php

/* Add an employee to the table. */
function AddEmployee($connection, $name, $address) {
  $n = mysqli_real_escape_string($connection, $name);
  $a = mysqli_real_escape_string($connection, $address);

$query = "INSERT INTO EMPLOYEES (NAME, ADDRESS) VALUES ('$n', '$a');";

```

```
        if(!mysqli_query($connection, $query)) echo("<p>Error adding employee data.</p>");

/* Check whether the table exists and, if not, create it. */
function VerifyEmployeesTable($connection, $dbName) {
    if(!TableExists("EMPLOYEES", $connection, $dbName))
    {
        $query = "CREATE TABLE EMPLOYEES (
            ID int(11) UNSIGNED AUTO_INCREMENT PRIMARY KEY,
            NAME VARCHAR(45),
            ADDRESS VARCHAR(90)
        )";

        if(!mysqli_query($connection, $query)) echo("<p>Error creating table.</p>");
    }
}

/* Check for the existence of a table. */
function TableExists($tableName, $connection, $dbName) {
    $t = mysqli_real_escape_string($connection, $tableName);
    $d = mysqli_real_escape_string($connection, $dbName);

    $checktable = mysqli_query($connection,
        "SELECT TABLE_NAME FROM information_schema.TABLES WHERE TABLE_NAME = '$t' AND
        TABLE_SCHEMA = '$d'");

    if(mysqli_num_rows($checktable) > 0) return true;

    return false;
}
?>
```

8. Save and close the SamplePage.php file.
9. Verify that your web server successfully connects to your RDS MySQL DB instance by opening a web browser and browsing to <http://EC2 instance endpoint>/SamplePage.php, for example: <http://ec2-55-122-41-31.us-west-2.compute.amazonaws.com/SamplePage.php>.

You can use SamplePage.php to add data to your RDS MySQL DB instance. The data that you add is then displayed on the page.

To make sure your RDS MySQL DB instance is as secure as possible, verify that sources outside of the VPC cannot connect to your RDS MySQL DB instance.

Tutorials

The following tutorials show you how to perform common tasks that use Amazon RDS:

- [Tutorial: Create an Amazon VPC for Use with a DB Instance \(p. 521\)](#)
- [Tutorial: Create a Web Server and an Amazon RDS Database \(p. 69\)](#)
- [Tutorial: Restore a DB Instance from a DB Snapshot \(p. 289\)](#)

For videos, see [AWS Instructional Videos and Labs](#).

Best Practices for Amazon RDS

Learn best practices for working with Amazon RDS. As new best practices are identified, we will keep this section up to date.

Topics

- [Amazon RDS Basic Operational Guidelines \(p. 94\)](#)
- [DB Instance RAM Recommendations \(p. 95\)](#)
- [Using Enhanced Monitoring to Identify Operating System Issues \(p. 95\)](#)
- [Using Metrics to Identify Performance Issues \(p. 95\)](#)
- [Best Practices for Working with MySQL Storage Engines \(p. 100\)](#)
- [Best Practices for Working with MariaDB Storage Engines \(p. 101\)](#)
- [Best Practices for Working with Oracle \(p. 101\)](#)
- [Best Practices for Working with PostgreSQL \(p. 101\)](#)
- [Best Practices for Working with SQL Server \(p. 103\)](#)
- [Working with DB Parameter Groups \(p. 103\)](#)
- [Amazon RDS Best Practices Presentation Video \(p. 104\)](#)

Note

For common recommendations for Amazon RDS, see [Using Amazon RDS Recommendations \(p. 365\)](#).

Amazon RDS Basic Operational Guidelines

The following are basic operational guidelines that everyone should follow when working with Amazon RDS. Note that the Amazon RDS Service Level Agreement requires that you follow these guidelines:

- Monitor your memory, CPU, and storage usage. Amazon CloudWatch can be set up to notify you when usage patterns change or when you approach the capacity of your deployment, so that you can maintain system performance and availability.
- Scale up your DB instance when you are approaching storage capacity limits. You should have some buffer in storage and memory to accommodate unforeseen increases in demand from your applications.
- Enable automatic backups and set the backup window to occur during the daily low in write IOPS.
- If your database workload requires more I/O than you have provisioned, recovery after a failover or database failure will be slow. To increase the I/O capacity of a DB instance, do any or all of the following:
 - Migrate to a DB instance class with High I/O capacity.
 - Convert from standard storage to either General Purpose or Provisioned IOPS storage, depending on how much of an increase you need. For information on available storage types, see [Amazon RDS Storage Types \(p. 135\)](#).

If you convert to Provisioned IOPS storage, make sure you also use a DB instance class that is optimized for Provisioned IOPS. For information on Provisioned IOPS, see [Provisioned IOPS SSD Storage \(p. 137\)](#).

- If you are already using Provisioned IOPS storage, provision additional throughput capacity.

- If your client application is caching the Domain Name Service (DNS) data of your DB instances, set a time-to-live (TTL) value of less than 30 seconds. Because the underlying IP address of a DB instance can change after a failover, caching the DNS data for an extended time can lead to connection failures if your application tries to connect to an IP address that no longer is in service.
- Test failover for your DB instance to understand how long the process takes for your use case and to ensure that the application that accesses your DB instance can automatically connect to the new DB instance after failover.

DB Instance RAM Recommendations

An Amazon RDS performance best practice is to allocate enough RAM so that your *working set* resides almost completely in memory. The working set is the data and indexes that are frequently in use on your instance. The more you use the DB instance, the more the working set will grow.

To tell if your working set is almost all in memory, check the ReadIOPS metric (using Amazon CloudWatch) while the DB instance is under load. The value of ReadIOPS should be small and stable. If scaling up the DB instance class—to a class with more RAM—results in a dramatic drop in ReadIOPS, your working set was not almost completely in memory. Continue to scale up until ReadIOPS no longer drops dramatically after a scaling operation, or ReadIOPS is reduced to a very small amount. For information on monitoring a DB instance's metrics, see [Viewing DB Instance Metrics \(p. 304\)](#).

Using Enhanced Monitoring to Identify Operating System Issues

Amazon RDS provides metrics in real time for the operating system (OS) that your DB instance runs on. You can view the metrics for your DB instance using the console, or consume the Enhanced Monitoring JSON output from Amazon CloudWatch Logs in a monitoring system of your choice. For more information about Enhanced Monitoring, see [Enhanced Monitoring \(p. 307\)](#)

Enhanced Monitoring is available for the following database engines:

- MariaDB
- Microsoft SQL Server
- MySQL version 5.5 or later
- Oracle
- PostgreSQL

Enhanced monitoring is available for all DB instance classes except for db.m1.small. Enhanced Monitoring is available in all regions except for AWS GovCloud (US-West).

Using Metrics to Identify Performance Issues

To identify performance issues caused by insufficient resources and other common bottlenecks, you can monitor the metrics available for your Amazon RDS DB instance.

Viewing Performance Metrics

You should monitor performance metrics on a regular basis to see the average, maximum, and minimum values for a variety of time ranges. If you do so, you can identify when performance is degraded. You

can also set Amazon CloudWatch alarms for particular metric thresholds so you are alerted if they are reached.

To troubleshoot performance issues, it's important to understand the baseline performance of the system. When you set up a new DB instance and get it running with a typical workload, you should capture the average, maximum, and minimum values of all of the performance metrics at a number of different intervals (for example, one hour, 24 hours, one week, two weeks) to get an idea of what is normal. It helps to get comparisons for both peak and off-peak hours of operation. You can then use this information to identify when performance is dropping below standard levels.

To view performance metrics

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose a DB instance.
3. Choose **Monitoring**. The first eight performance metrics display. The metrics default to showing information for the current day.
4. Use the numbered buttons at top right to page through the additional metrics, or choose adjust the settings to see more metrics.
5. Choose a performance metric to adjust the time range in order to see data for other than the current day. You can change the **Statistic**, **Time Range**, and **Period** values to adjust the information displayed. For example, to see the peak values for a metric for each day of the last two weeks, set **Statistic** to **Maximum**, **Time Range** to **Last 2 Weeks**, and **Period** to **Day**.

Note

Changing the **Statistic**, **Time Range**, and **Period** values changes them for all metrics. The updated values persist for the remainder of your session or until you change them again.

You can also view performance metrics using the CLI or API. For more information, see [Viewing DB Instance Metrics \(p. 304\)](#).

To set a CloudWatch alarm

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose a DB instance.
3. Choose **Logs & events**.
4. In the **CloudWatch alarms** section, choose **Create alarm**.

Create alarm

You can use CloudWatch alarms to be notified automatically whenever metric data reaches a level you define.

Settings

[Refresh](#)

Send notifications

- Yes
 No

Send notifications to

- ARN
 New email or SMS topic

Topic name

Name of the topic.

Manually enter a topic name...

With these recipients

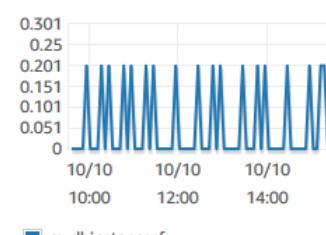
Email addresses or phone numbers of SMS enabled devices to send the notifications to

awsAccount@domain.com

Metric

Average ▼ of CPU Utilization ▼

CPU Utilization Percent



Threshold

>= ▼ ▲ Percent

Evaluation period

1 ▲ consecutive period(s) of 5 Minutes ▼

Name of alarm

awsrds-mydbinstancecf-High-CPU-Utilization

[Cancel](#)

[Create alarm](#)

5. For **Send notifications**, choose **Yes**, and for **Send notifications to**, choose **New email or SMS topic**.
6. For **Topic name**, enter a name for the notification, and for **With these recipients**, enter a comma-separated list of email addresses and phone numbers.
7. For **Metric**, choose the alarm statistic and metric to set.
8. For **Threshold**, specify whether the metric must be greater than, less than, or equal to the threshold, and specify the threshold value.
9. For **Evaluation period**, choose the evaluation period for the alarm, and for **consecutive period(s) of**, choose the period during which the threshold must have been reached in order to trigger the alarm.
10. For **Name of alarm**, enter a name for the alarm.
11. Choose **Create Alarm**.

The alarm appears in the **CloudWatch alarms** section.

Evaluating Performance Metrics

A DB instance has a number of different categories of metrics, and how to determine acceptable values depends on the metric.

CPU

- CPU Utilization – Percentage of computer processing capacity used.

Memory

- Freeable Memory – How much RAM is available on the DB instance, in megabytes. The red line in the Monitoring tab metrics is marked at 75% for CPU, Memory and Storage Metrics. If instance memory consumption frequently crosses that line, then this indicates that you should check your workload or upgrade your instance.
- Swap Usage – How much swap space is used by the DB instance, in megabytes.

Disk space

- Free Storage Space – How much disk space is not currently being used by the DB instance, in megabytes.

Input/output operations

- Read IOPS, Write IOPS – The average number of disk read or write operations per second.
- Read Latency, Write Latency – The average time for a read or write operation in milliseconds.
- Read Throughput, Write Throughput – The average number of megabytes read from or written to disk per second.
- Queue Depth – The number of I/O operations that are waiting to be written to or read from disk.

Network traffic

- Network Receive Throughput, Network Transmit Throughput – The rate of network traffic to and from the DB instance in megabytes per second.

Database connections

- DB Connections – The number of client sessions that are connected to the DB instance.

For more detailed individual descriptions of the performance metrics available, see [Amazon RDS Dimensions and Metrics](#).

Generally speaking, acceptable values for performance metrics depend on what your baseline looks like and what your application is doing. Investigate consistent or trending variances from your baseline. Advice about specific types of metrics follows:

- **High CPU or RAM consumption** – High values for CPU or RAM consumption might be appropriate, provided that they are in keeping with your goals for your application (like throughput or concurrency) and are expected.

- **Disk space consumption** – Investigate disk space consumption if space used is consistently at or above 85 percent of the total disk space. See if it is possible to delete data from the instance or archive data to a different system to free up space.
- **Network traffic** – For network traffic, talk with your system administrator to understand what expected throughput is for your domain network and Internet connection. Investigate network traffic if throughput is consistently lower than expected.
- **Database connections** – Consider constraining database connections if you see high numbers of user connections in conjunction with decreases in instance performance and response time. The best number of user connections for your DB instance will vary based on your instance class and the complexity of the operations being performed. You can determine the number of database connections by associating your DB instance with a parameter group where the *User Connections* parameter is set to other than 0 (unlimited). You can either use an existing parameter group or create a new one. For more information, see [Working with DB Parameter Groups \(p. 211\)](#).
- **IOPS metrics** – The expected values for IOPS metrics depend on disk specification and server configuration, so use your baseline to know what is typical. Investigate if values are consistently different than your baseline. For best IOPS performance, make sure your typical working set will fit into memory to minimize read and write operations.

For issues with any performance metrics, one of the first things you can do to improve performance is tune the most used and most expensive queries to see if that lowers the pressure on system resources. For more information, see [Tuning Queries \(p. 99\)](#)

If your queries are tuned and an issue persists, consider upgrading your Amazon RDS [Choosing the DB Instance Class \(p. 107\)](#) to one with more of the resource (CPU, RAM, disk space, network bandwidth, I/O capacity) that is related to the issue you are experiencing.

Tuning Queries

One of the best ways to improve DB instance performance is to tune your most commonly used and most resource-intensive queries to make them less expensive to run.

MySQL Query Tuning

Go to [Optimizing SELECT Statements](#) in the MySQL documentation for more information on writing queries for better performance. You can also go to [MySQL Performance Tuning and Optimization Resources](#) for additional query tuning resources.

Oracle Query Tuning

Go to the [Database SQL Tuning Guide](#) in the Oracle documentation for more information on writing and analyzing queries for better performance.

SQL Server Query Tuning

Go to [Analyzing a Query](#) in the SQL Server documentation to improve queries for SQL Server DB instances. You can also use the execution-, index- and I/O-related data management views (DMVs) described in the [Dynamic Management Views and Functions](#) documentation to troubleshoot SQL Server query issues.

A common aspect of query tuning is creating effective indexes. You can use the [Database Engine Tuning Advisor](#) to get potential index improvements for your DB instance. For more information, see [Analyzing Your Database Workload on an Amazon RDS DB Instance with SQL Server Tuning Advisor \(p. 726\)](#).

PostgreSQL Query Tuning

Go to [Using EXPLAIN](#) in the PostgreSQL documentation to learn how to analyze a query plan. You can use this information to modify a query or underlying tables in order to improve query performance. You

can also go to [Controlling the Planner with Explicit JOIN Clauses](#) to get tips about how to specify joins in your query for the best performance.

MariaDB Query Tuning

Go to [Query Optimizations](#) in the MariaDB documentation for more information on writing queries for better performance.

Best Practices for Working with MySQL Storage Engines

On a MySQL DB instance, observe the following table creation limits:

- You're limited to 10,000 tables if you are either using Provisioned IOPS storage, or using General Purpose storage and the DB instance is 200 GiB or larger in size.
- You're limited to 1000 tables if you are either using standard storage, or using General Purpose storage and the DB instance is less than 200 GiB in size.

We recommend these limits because having large numbers of tables significantly increases database recovery time after a failover or database crash. If you need to create more tables than recommended, set the `innodb_file_per_table` parameter to 0. For more information, see [Working with InnoDB Tablespaces to Improve Crash Recovery Times \(p. 882\)](#) and [Working with DB Parameter Groups \(p. 211\)](#).

For MySQL DB instances that use version 5.7 or later, you can exceed these table creation limits due to improvements in InnoDB crash recovery. However, we still recommend that you take caution due to the potential performance impact of creating very large numbers of tables.

On a MySQL DB instance, avoid tables in your database growing too large. Although the general storage limit is 64 TiB, provisioned storage limits restrict the maximum size of a MySQL table file to 16 TiB. Partition your large tables so that file sizes are well under the 16 TiB limit. This approach can also improve performance and recovery time. For more information, see [MySQL File Size Limits \(p. 887\)](#).

The Point-In-Time Restore and snapshot restore features of Amazon RDS for MySQL require a crash-recoverable storage engine and are supported for the InnoDB storage engine only. Although MySQL supports multiple storage engines with varying capabilities, not all of them are optimized for crash recovery and data durability. For example, the MyISAM storage engine does not support reliable crash recovery and might prevent a Point-In-Time Restore or snapshot restore from working as intended. This might result in lost or corrupt data when MySQL is restarted after a crash.

InnoDB is the recommended and supported storage engine for MySQL DB instances on Amazon RDS. InnoDB instances can also be migrated to Aurora, while MyISAM instances can't be migrated. However, MyISAM performs better than InnoDB if you require intense, full-text search capability. If you still choose to use MyISAM with Amazon RDS, following the steps outlined in [Automated Backups with Unsupported MySQL Storage Engines \(p. 262\)](#) can be helpful in certain scenarios for snapshot restore functionality.

If you want to convert existing MyISAM tables to InnoDB tables, you can use the process outlined in the [MySQL documentation](#). MyISAM and InnoDB have different strengths and weaknesses, so you should fully evaluate the impact of making this switch on your applications before doing so.

In addition, Federated Storage Engine is currently not supported by Amazon RDS for MySQL.

Best Practices for Working with MariaDB Storage Engines

The point-in-time restore and snapshot restore features of Amazon RDS for MariaDB require a crash-recoverable storage engine. Although MariaDB supports multiple storage engines with varying capabilities, not all of them are optimized for crash recovery and data durability. For example, although Aria is a crash-safe replacement for MyISAM, it might still prevent a point-in-time restore or snapshot restore from working as intended. This might result in lost or corrupt data when MariaDB is restarted after a crash. InnoDB (for version 10.2 and higher) and XtraDB (for version 10.0 and 10.1) are the recommended and supported storage engines for MariaDB DB instances on Amazon RDS. If you still choose to use Aria with Amazon RDS, following the steps outlined in [Automated Backups with Unsupported MariaDB Storage Engines \(p. 263\)](#) can be helpful in certain scenarios for snapshot restore functionality.

Best Practices for Working with Oracle

For information about best practices for working with Amazon RDS for Oracle, see [Best Practices for Running Oracle Database on Amazon Web Services](#) and the video [Running Oracle Databases on Amazon RDS](#).

Best Practices for Working with PostgreSQL

Two important areas where you can improve performance with PostgreSQL on Amazon RDS are when loading data into a DB instance and when using the PostgreSQL autovacuum feature. The following sections cover some of the practices we recommend for these areas.

Loading Data into a PostgreSQL DB Instance

When loading data into an Amazon RDS PostgreSQL DB instance, you should modify your DB instance settings and your DB parameter group values to allow for the most efficient importing of data into your DB instance.

Modify your DB instance settings to the following:

- Disable DB instance backups (set `backup_retention` to 0)
- Disable Multi-AZ

Modify your DB parameter group to include the following settings. You should test the parameter settings to find the most efficient settings for your DB instance:

- Increase the value of the `maintenance_work_mem` parameter. For more information about PostgreSQL resource consumption parameters, see the [PostgreSQL documentation](#).
- Increase the value of the `checkpoint_segments` and `checkpoint_timeout` parameters to reduce the number of writes to the wal log.
- Disable the `synchronous_commit` parameter (do not turn off FSYNC).
- Disable the PostgreSQL autovacuum parameter.
- Make sure none of the tables you are importing are unlogged. Data stored in unlogged tables can be lost during a failover. For more information see, [CREATE TABLE UNLOGGED](#)

Use the `pg_dump -Fc` (compressed) or `pg_restore -j` (parallel) commands with these settings.

Working with the fsync and full_page_writes database parameters

In PostgreSQL 9.4.1 on Amazon RDS, the `fsync` and `full_page_writes` database parameters are not modifiable. Disabling the `fsync` and `full_page_writes` database parameters can lead to data corruption, so we have enabled them for you. We recommend that customers with other 9.3 DB engine versions of PostgreSQL not disable the `fsync` and `full_page_writes` parameters.

Working with the PostgreSQL Autovacuum Feature

The autovacuum feature for PostgreSQL databases is a feature that we strongly recommend you use to maintain the health of your PostgreSQL DB instance. Autovacuum automates the execution of the `VACUUM` and `ANALYZE` command; using autovacuum is required by PostgreSQL, not imposed by Amazon RDS, and its use is critical to good performance. The feature is enabled by default for all new Amazon RDS PostgreSQL DB instances, and the related configuration parameters are appropriately set by default.

Your database administrator needs to know and understand this maintenance operation. For the PostgreSQL documentation on autovacuum, see <http://www.postgresql.org/docs/current/static/routine-vacuuming.html#AUTOVACUUM>.

Autovacuum is not a “resource free” operation, but it works in the background and yields to user operations as much as possible. When enabled, autovacuum checks for tables that have had a large number of updated or deleted tuples. It also protects against loss of very old data due to [transaction ID wraparound](#).

Autovacuum should not be thought of as a high-overhead operation that can be reduced to gain better performance. On the contrary, tables that have a high velocity of updates and deletes will quickly deteriorate over time if autovacuum is not run.

Important

Not running autovacuum can result in an eventual required outage to perform a much more intrusive vacuum operation. When an Amazon RDS PostgreSQL DB instance becomes unavailable because of an over conservative use of autovacuum, the PostgreSQL database will shut down to protect itself. At that point, Amazon RDS must perform a single-user-mode full vacuum directly on the DB instance , which can result in a multi-hour outage. Thus, we strongly recommend that you do not turn off autovacuum, which is enabled by default.

The autovacuum parameters determine when and how hard autovacuum works. The `autovacuum_vacuum_threshold` and `autovacuum_vacuum_scale_factor` parameters determine when autovacuum is run. The `autovacuum_max_workers`, `autovacuum_nap_time`, `autovacuum_cost_limit`, and `autovacuum_cost_delay` parameters determine how hard autovacuum works. For more information about autovacuum, when it runs, and what parameters are required, see the [PostgreSQL documentation](#).

The following query shows the number of "dead" tuples in a table named `table1` :

```
PROMPT> select relname, n_dead_tup, last_vacuum, last_autovacuum from pg_catalog.pg_stat_all_tables where n_dead_tup > 0 and relname = 'table1' order by n_dead_tup desc;
```

The results of the query will resemble the following:

<code>relname n_dead_tup last_vacuum last_autovacuum</code>

tasks		81430522	
(1 row)			

Best Practices for Working with SQL Server

Best practices for a Multi-AZ deployment with a SQL Server DB instance include the following:

- Use Amazon RDS DB events to monitor failovers. For example, you can be notified by text message or email when a DB instance fails over. For more information about Amazon RDS events, see [Using Amazon RDS Event Notification \(p. 370\)](#).
- If your application caches DNS values, set time to live (TTL) to less than 30 seconds. Setting TTL as so is a good practice in case there is a failover, where the IP address might change and the cached value might no longer be in service.
- We recommend that you *do not* enable the following modes because they turn off transaction logging, which is required for Multi-AZ:
 - Simple recover mode
 - Offline mode
 - Read-only mode
- Test to determine how long it takes for your DB instance to failover. Failover time can vary due to the type of database, the instance class, and the storage type you use. You should also test your application's ability to continue working if a failover occurs.
- To shorten failover time, you should do the following:
 - Ensure that you have sufficient Provisioned IOPS allocated for your workload. Inadequate I/O can lengthen failover times. Database recovery requires I/O.
 - Use smaller transactions. Database recovery relies on transactions, so if you can break up large transactions into multiple smaller transactions, your failover time should be shorter.
- Take into consideration that during a failover, there will be elevated latencies. As part of the failover process, Amazon RDS automatically replicates your data to a new standby instance. This replication means that new data is being committed to two different DB instances, so there might be some latency until the standby DB instance has caught up to the new primary DB instance.
- Deploy your applications in all Availability Zones. If an Availability Zone does go down, your applications in the other Availability Zones will still be available.

When working with a Multi-AZ deployment of SQL Server, remember that Amazon RDS creates replicas for all SQL Server databases on your instance. If you don't want specific databases to have secondary replicas, set up a separate DB instance that doesn't use Multi-AZ for those databases.

Working with DB Parameter Groups

We recommend that you try out DB parameter group changes on a test DB instance before applying parameter group changes to your production DB instances. Improperly setting DB engine parameters in a DB parameter group can have unintended adverse effects, including degraded performance and system instability. Always exercise caution when modifying DB engine parameters and back up your DB instance before modifying a DB parameter group.

For information about backing up your DB instance, see [Backing Up and Restoring Amazon RDS DB Instances \(p. 253\)](#).

Amazon RDS Best Practices Presentation Video

The 2016 AWS Summit conference in Chicago included a presentation on best practices for creating and configuring a secure, highly available database instance using Amazon RDS. A video of the presentation is available [here](#).

Amazon RDS DB Instances

A *DB instance* is an isolated database environment running in the cloud. It is the basic building block of Amazon RDS. A DB instance can contain multiple user-created databases, and can be accessed using the same client tools and applications you might use to access a standalone database instance. DB instances are simple to create and modify with the Amazon AWS command line tools, Amazon RDS API operations, or the AWS Management Console.

Note

Amazon RDS supports access to databases using any standard SQL client application. Amazon RDS does not allow direct host access.

You can have up to 40 Amazon RDS DB instances:

- 10 for each SQL Server edition (Enterprise, Standard, Web, and Express) under the "license-included" model
- 40 for MySQL, MariaDB, or PostgreSQL
- 40 for Oracle under the "bring-your-own-license" (BYOL) licensing model

Note

If your application requires more DB instances, you can request additional DB instances using the form at <https://console.aws.amazon.com/support/home#/case/create?issueType=service-limit-increase&limitType=service-code-rds-instances>.

Each DB instance has a DB instance identifier. This customer-supplied name uniquely identifies the DB instance when interacting with the Amazon RDS API and AWS CLI commands. The DB instance identifier must be unique for that customer in an AWS Region.

The identifier is used as part of the DNS hostname allocated to your instance by RDS. For example, if you specify db1 as the DB instance identifier, then RDS will automatically allocate a DNS endpoint for your instance, such as db1.123456789012.us-east-1.rds.amazonaws.com, where 123456789012 is the fixed identifier for a specific region for your account.

Each DB instance supports a database engine. Amazon RDS currently supports MySQL, MariaDB, PostgreSQL, Oracle, Microsoft SQL Server, and Amazon Aurora database engines.

When creating a DB instance, some database engines require that a database name be specified. A DB instance can host multiple databases, or a single Oracle database with multiple schemas. The database name value depends on the database engine:

- For the MySQL and MariaDB database engines, the database name is the name of a database hosted in your DB instance. Databases hosted by the same DB instance must have a unique name within that instance.
- For the Oracle database engine, database name is used to set the value of ORACLE_SID, which must be supplied when connecting to the Oracle RDS instance.
- For the Microsoft SQL Server database engine, database name is not a supported parameter.
- For the PostgreSQL database engine, the database name is the name of a database hosted in your DB instance. A database name is not required when creating a DB instance. Databases hosted by the same DB instance must have a unique name within that instance.

Amazon RDS creates a master user account for your DB instance as part of the creation process. This master user has permissions to create databases and to perform create, delete, select, update, and insert operations on tables the master user creates. You must set the master user password when you create a

DB instance, but you can change it at any time using the Amazon AWS command line tools, Amazon RDS API operations, or the AWS Management Console. You can also change the master user password and manage users using standard SQL commands.

Note

This guide covers non-Aurora Amazon RDS database engines. For information about using Amazon Aurora, see the [Amazon Aurora User Guide](#).

Topics

- [Choosing the DB Instance Class \(p. 107\)](#)
- [DB Instance Status \(p. 129\)](#)
- [Regions and Availability Zones \(p. 132\)](#)
- [Amazon RDS DB Instance Storage \(p. 135\)](#)
- [High Availability \(Multi-AZ\) for Amazon RDS \(p. 144\)](#)
- [Amazon RDS DB Instance Lifecycle \(p. 147\)](#)
- [Tagging Amazon RDS Resources \(p. 177\)](#)
- [Working with Read Replicas \(p. 182\)](#)
- [Working with Option Groups \(p. 196\)](#)
- [Working with DB Parameter Groups \(p. 211\)](#)
- [Working with Amazon Resource Names \(ARNs\) in Amazon RDS \(p. 224\)](#)
- [Working with Storage for Amazon RDS DB Instances \(p. 231\)](#)
- [DB Instance Billing for Amazon RDS \(p. 239\)](#)

Choosing the DB Instance Class

The DB instance class determines the computation and memory capacity of an Amazon RDS DB instance. The DB instance class you need depends on your processing power and memory requirements.

For more information about instance class pricing, see [Amazon RDS Pricing](#).

Topics

- [DB Instance Class Types \(p. 107\)](#)
- [Terminology for DB Instance Class Hardware Specifications \(p. 108\)](#)
- [Hardware Specifications for All Available DB Instance Classes \(p. 109\)](#)
- [Supported DB Engines for All Available DB Instance Classes \(p. 112\)](#)
- [Changing Your DB Instance Class \(p. 116\)](#)
- [Configuring the Processor for a DB Instance Class \(p. 117\)](#)

DB Instance Class Types

Amazon RDS supports three types of instance classes: Standard, Memory Optimized, and Burstable Performance. For more information about Amazon EC2 instance types, see [Instance Type](#) in the Amazon EC2 documentation.

The following are the Standard DB instance classes available:

- **db.m5** – Latest-generation general-purpose instance classes that provide a balance of compute, memory, and network resources, and are a good choice for many applications. The db.m5 instance classes provide more computing capacity than the previous db.m4 instance classes. They are powered by the AWS Nitro System, a combination of dedicated hardware and lightweight hypervisor.
- **db.m4** – Current-generation general-purpose instance classes that provide more computing capacity than the previous db.m3 instance classes.
- **db.m3** – Previous-generation general-purpose instance classes that provide more computing capacity than the previous db.m1 instance classes.
- **db.m1** – Previous-generation general-purpose instance classes.

The following are the Memory Optimized DB instance classes available:

- **db.z1d** – Latest-generation instance classes optimized for memory-intensive applications. These offer both high compute capacity and a high memory footprint. High frequency z1d instances deliver a sustained all core frequency of up to 4.0 GHz.
- **db.x1e** – Latest-generation instance classes optimized for memory-intensive applications. These offer one of the lowest price per GiB of RAM among the DB instance classes and up to 3,904 GiB of DRAM-based instance memory. The db.x1e instance classes are available only in the following regions: US East (N. Virginia), US West (Oregon), Europe (Ireland), Asia Pacific (Tokyo), and Asia Pacific (Sydney).
- **db.x1** – Current-generation instance classes optimized for memory-intensive applications. These offer one of the lowest price per GiB of RAM among the DB instance classes and up to 1,952 GiB of DRAM-based instance memory.
- **db.r5** – Latest-generation instance classes optimized for memory-intensive applications. These offer improved networking and Amazon Elastic Block Store (Amazon EBS) performance. They are powered by the AWS Nitro System, a combination of dedicated hardware and lightweight hypervisor.
- **db.r4** – Current-generation instance classes optimized for memory-intensive applications. These offer improved networking and Amazon EBS performance.

- **db.r3** – Previous-generation instance classes that provide memory optimization. The db.r3 instances classes are not available in the Europe (Paris) region and the South America (São Paulo) region.
- **db.m2** – Previous-generation memory-optimized instance classes.

The following are the Burstable Performance DB instance classes available:

- **db.t3** – Latest-generation instance classes that provide a baseline performance level, with the ability to burst to full CPU usage. These instance classes provide more computing capacity than the previous db.t2 instance classes. They are powered by the AWS Nitro System, a combination of dedicated hardware and lightweight hypervisor.
- **db.t2** – Current-generation instance classes that provide a baseline performance level, with the ability to burst to full CPU usage.

Note

The DB instance classes that use the AWS Nitro System (db.m5, db.r5, db.t3) are throttled on combined read plus write workload.

Terminology for DB Instance Class Hardware Specifications

The following terminology is used to describe hardware specifications for DB instance classes:

- **vCPU** – The number of virtual central processing units (CPUs). A *virtual CPU* is a unit of capacity that you can use to compare DB instance classes. Instead of purchasing or leasing a particular processor to use for several months or years, you are renting capacity by the hour. Our goal is to make a consistent and specific amount of CPU capacity available, within the limits of the actual underlying hardware.
- **ECU** – The relative measure of the integer processing power of an Amazon EC2 instance. To make it easy for developers to compare CPU capacity between different instance classes, we have defined an Amazon EC2 Compute Unit. The amount of CPU that is allocated to a particular instance is expressed in terms of these EC2 Compute Units. One ECU currently provides CPU capacity equivalent to a 1.0–1.2 GHz 2007 Opteron or 2007 Xeon processor.
- **Memory (GiB)** – The RAM, in gibibytes, allocated to the DB instance. There is often a consistent ratio between memory and vCPU. As an example, take the db.r4 instance class, which has a memory to vCPU ratio similar to the db.r5 instance class. However, for most use cases the db.r5 instance class provides better, more consistent performance than the db.r4 instance class.
- **VPC Only** – The instance class is supported only for DB instances that are in a VPC based on the Amazon VPC service. In some cases, you might want to use an instance class that requires a VPC but your current DB instance isn't in a VPC. In these cases, start by moving your DB instance into a VPC. For more information, see [Moving a DB Instance Not in a VPC into a VPC \(p. 520\)](#).
- **EBS-Optimized** – The DB instance uses an optimized configuration stack and provides additional, dedicated capacity for I/O. This optimization provides the best performance by minimizing contention between I/O and other traffic from your instance. For more information about Amazon EBS-optimized instances, see [Amazon EBS-Optimized Instances](#) in the *Amazon EC2 User Guide for Linux Instances*.
- **Max. Bandwidth (Mbps)** – The maximum bandwidth in megabits per second. Divide by 8 to get the expected throughput in megabytes per second.

Important

General Purpose SSD (gp2) volumes for Amazon RDS DB instances have a throughput limit of 250 MiB/s in most cases. However, the throughput limit can vary depending on volume size. For more information, see [Amazon EBS Volume Types](#) in the *Amazon EC2 User Guide for Linux Instances*. For information on estimating bandwidth for gp2 storage, see [General Purpose SSD Storage \(p. 135\)](#).

- **Network Performance** – The network speed relative to other DB instance classes.

Hardware Specifications for All Available DB Instance Classes

In the following table, you can find details about the Amazon RDS DB instance classes. For a more detailed explanation of the table column terminology, see [Terminology for DB Instance Class Hardware Specifications \(p. 108\)](#). For information about Amazon RDS DB engine support for each DB instance class, see [Supported DB Engines for All Available DB Instance Classes \(p. 112\)](#).

Instance Class	vCPU	ECU	Memory (GiB)	VPC Only	EBS Optimized	Max. Bandwidth (Mbps)	Network Performance
db.m5 – Latest Generation Standard Instance Classes							
db.m5.24xlarge	96	345	384	Yes	Yes	19,000	25 Gbps
db.m5.16xlarge	64	262	256	Yes	Yes	13,600	20 Gbps
db.m5.12xlarge	48	173	192	Yes	Yes	9,500	10 Gbps
db.m5.8xlarge	32	131	128	Yes	Yes	6,800	10 Gbps
db.m5.4xlarge	16	61	64	Yes	Yes	4,750	Up to 10 Gbps
db.m5.2xlarge*	8	31	32	Yes	Yes	Up to 4,750	Up to 10 Gbps
db.m5.xlarge*	4	15	16	Yes	Yes	Up to 4,750	Up to 10 Gbps
db.m5.large*	2	10	8	Yes	Yes	Up to 4,750	Up to 10 Gbps
db.m4 – Current Generation Standard Instance Classes							
db.m4.16xlarge	64	188	256	Yes	Yes	10,000	25 Gbps
db.m4.10xlarge	40	124.5	160	Yes	Yes	4,000	10 Gbps
db.m4.4xlarge	16	53.5	64	Yes	Yes	2,000	High
db.m4.2xlarge	8	25.5	32	Yes	Yes	1,000	High
db.m4.xlarge	4	13	16	Yes	Yes	750	High
db.m4.large	2	6.5	8	Yes	Yes	450	Moderate
db.m3 – Previous Generation Standard Instance Classes							
db.m3.2xlarge	8	26	30	No	Yes	1,000	High
db.m3.xlarge	4	13	15	No	Yes	500	High
db.m3.large	2	6.5	7.5	No	No	—	Moderate
db.m3.medium	1	3	3.75	No	No	—	Moderate
db.m1 – Previous Generation Standard Instance Classes							

db.m1.xlarge	4	4	15	No	Yes	450	High
db.m1.large	2	2	7.5	No	Yes	450	Moderate
db.m1.medium	1	1	3.75	No	No	—	Moderate
db.m1.small	1	1	1.7	No	No	—	Very Low
Instance Class	vCPU	ECU	Memory (GiB)	VPC Only	EBS Optimized	Max. Bandwidth (Mbps)	Network Performance

db.z1d – Latest Generation Memory Optimized Instance Classes

db.z1d.12xlarge	48	271	384	Yes	Yes	14,000	25 Gbps
db.z1d.6xlarge	24	134	192	Yes	Yes	7,000	10 Gbps
db.z1d.3xlarge	12	75	96	Yes	Yes	3,500	Up to 10 Gbps
db.z1d.2xlarge	8	53	64	Yes	Yes	2,333	Up to 10 Gbps
db.z1d.xlarge*	4	28	32	Yes	Yes	Up to 2,333	Up to 10 Gbps
db.z1d.large*	2	15	16	Yes	Yes	Up to 2,333	Up to 10 Gbps
Instance Class	vCPU	ECU	Memory (GiB)	VPC Only	EBS Optimized	Max. Bandwidth (Mbps)	Network Performance

db.x1e – Latest Generation Memory Optimized Instance Classes

db.x1e.32xlarge	128	340	3,904	Yes	Yes	14,000	25 Gbps
db.x1e.16xlarge	64	179	1,952	Yes	Yes	7,000	10 Gbps
db.x1e.8xlarge	32	91	976	Yes	Yes	3,500	Up to 10 Gbps
db.x1e.4xlarge	16	47	488	Yes	Yes	1,750	Up to 10 Gbps
db.x1e.2xlarge	8	23	244	Yes	Yes	1,000	Up to 10 Gbps
db.x1e.xlarge	4	12	122	Yes	Yes	500	Up to 10 Gbps

db.x1 – Current Generation Memory Optimized Instance Classes

db.x1.32xlarge	128	349	1,952	Yes	Yes	14,000	25 Gbps
db.x1.16xlarge	64	174.5	976	Yes	Yes	7,000	10 Gbps
Instance Class	vCPU	ECU	Memory (GiB)	VPC Only	EBS Optimized	Max. Bandwidth (Mbps)	Network Performance

db.r5 – Latest Generation Memory Optimized Instance Classes

db.r5.24xlarge	96	347	768	Yes	Yes	19,000	25 Gbps
db.r5.16xlarge	64	264	512	Yes	Yes	13,600	20 Gbps
db.r5.12xlarge	48	173	384	Yes	Yes	9,500	10 Gbps
db.r5.8xlarge	32	132	256	Yes	Yes	6,800	10 Gbps
db.r5.4xlarge	16	71	128	Yes	Yes	4,750	Up to 10 Gbps
db.r5.2xlarge*	8	38	64	Yes	Yes	Up to 4,750	Up to 10 Gbps
db.r5.xlarge*	4	19	32	Yes	Yes	Up to 4,750	Up to 10 Gbps
db.r5.large*	2	10	16	Yes	Yes	Up to 4,750	Up to 10 Gbps

db.r4 – Current Generation Memory Optimized Instance Classes

db.r4.16xlarge	64	195	488	Yes	Yes	14,000	25 Gbps
db.r4.8xlarge	32	99	244	Yes	Yes	7,000	10 Gbps
db.r4.4xlarge	16	53	122	Yes	Yes	3,500	Up to 10 Gbps
db.r4.2xlarge	8	27	61	Yes	Yes	1,700	Up to 10 Gbps
db.r4.xlarge	4	13.5	30.5	Yes	Yes	850	Up to 10 Gbps
db.r4.large	2	7	15.25	Yes	Yes	425	Up to 10 Gbps

db.r3 – Previous Generation Memory Optimized Instance Classes

db.r3.8xlarge	32	104	244	No	No	—	10 Gbps
db.r3.4xlarge	16	52	122	No	Yes	2,000	High
db.r3.2xlarge	8	26	61	No	Yes	1,000	High
db.r3.xlarge	4	13	30.5	No	Yes	500	Moderate
db.r3.large	2	6.5	15.25	No	No	—	Moderate
Instance Class	vCPU	ECU	Memory (GiB)	VPC Only	EBS Optimized	Max. Bandwidth (Mbps)	Network Performance

db.m2 – Previous Generation Memory Optimized Instance Classes

db.m2.4xlarge	8	26	68.4	No	Yes	1,000	High
db.m2.2xlarge	4	13	34.2	No	Yes	500	Moderate

db.m2.xlarge	2	6.5	17.1	No	No	—	Moderate
Instance Class	vCPU	ECU	Memory (GiB)	VPC Only	EBS Optimized	Max. Bandwidth (Mbps)	Network Performance
db.t3 – Latest Generation Burstable Performance Instance Classes							
db.t3.2xlarge*	8	Variable	32	Yes	Yes	Up to 2,048	Up to 5 Gbps
db.t3.xlarge*	4	Variable	16	Yes	Yes	Up to 2,048	Up to 5 Gbps
db.t3.large*	2	Variable	8	Yes	Yes	Up to 2,048	Up to 5 Gbps
db.t3.medium*	2	Variable	4	Yes	Yes	Up to 1,536	Up to 5 Gbps
db.t3.small*	2	Variable	2	Yes	Yes	Up to 1,536	Up to 5 Gbps
db.t3.micro*	2	Variable	1	Yes	Yes	Up to 1,536	Up to 5 Gbps
db.t2 – Current Generation Burstable Performance Instance Classes							
db.t2.2xlarge	8	Variable	32	Yes	No	—	Moderate
db.t2.xlarge	4	Variable	16	Yes	No	—	Moderate
db.t2.large	2	Variable	8	Yes	No	—	Moderate
db.t2.medium	2	Variable	4	Yes	No	—	Moderate
db.t2.small	1	Variable	2	Yes	No	—	Low
db.t2.micro	1	Variable	1	Yes	No	—	Low

* These DB instance classes can support maximum performance for 30 minutes at least once every 24 hours. For more information on baseline performance of these instance types, see [Amazon EBS-Optimized Instances](#) in the *Amazon EC2 User Guide for Linux Instances*.

Supported DB Engines for All Available DB Instance Classes

In the following table, you can find details about supported Amazon RDS DB instance classes for each Amazon RDS DB engine. For DB instance class specifications, see [Hardware Specifications for All Available DB Instance Classes \(p. 109\)](#).

The following are DB engine considerations for DB instance classes:

- **Microsoft SQL Server** – Instance class support varies according to the version and edition of SQL Server. For instance class support by version and edition, see [DB Instance Class Support for Microsoft SQL Server \(p. 602\)](#).
- **Oracle** – Instance class support varies according to the version and edition of Oracle. For instance class support by version and edition, see [DB Instance Class Support for Oracle \(p. 917\)](#).

- **PostgreSQL** – The db.m5, db.r5, and db.t3 DB instance classes are supported for the following Amazon RDS PostgreSQL versions:
 - PostgreSQL 9.6.9 and higher 9.6 versions
 - PostgreSQL 10.4 and higher 10 versions
 - PostgreSQL 11.1 and higher 11 versions

Instance Class	MariaDB	Microsoft SQL Server	MySQL	Oracle	PostgreSQL
db.m5 – Latest Generation Standard Instance Classes					
db.m5.24xlarge	Yes	Yes	Yes	Yes	Yes
db.m5.16xlarge	No	Yes	No	Yes	No
db.m5.12xlarge	Yes	Yes	Yes	Yes	Yes
db.m5.8xlarge	No	Yes	No	Yes	No
db.m5.4xlarge	Yes	Yes	Yes	Yes	Yes
db.m5.2xlarge	Yes	Yes	Yes	Yes	Yes
db.m5.xlarge	Yes	Yes	Yes	Yes	Yes
db.m5.large	Yes	Yes	Yes	Yes	Yes
db.m4 – Current Generation Standard Instance Classes					
db.m4.16xlarge	Yes	Yes	MySQL 8.0, 5.7, 5.6	Yes	Yes
db.m4.10xlarge	Yes	Yes	Yes	Yes	Yes
db.m4.4xlarge	Yes	Yes	Yes	Yes	Yes
db.m4.2xlarge	Yes	Yes	Yes	Yes	Yes
db.m4.xlarge	Yes	Yes	Yes	Yes	Yes
db.m4.large	Yes	Yes	Yes	Yes	Yes
db.m3 – Previous Generation Standard Instance Classes					
db.m3.2xlarge	No	Yes	Yes	Deprecated	Yes
db.m3.xlarge	No	Yes	Yes	Deprecated	Yes
db.m3.large	No	Yes	Yes	Deprecated	Yes
db.m3.medium	No	Yes	Yes	Deprecated	Yes
db.m1 – Previous Generation Standard Instance Classes					
db.m1.xlarge	No	Yes	Deprecated	Deprecated	PostgreSQL 9.4, 9.3

db.m1.large	No	Yes	Deprecated	Deprecated	PostgreSQL 9.4, 9.3
db.m1.medium	No	Yes	Deprecated	Deprecated	PostgreSQL 9.4, 9.3
db.m1.small	No	Yes	Deprecated	Deprecated	PostgreSQL 9.4, 9.3
Instance Class	MariaDB	Microsoft SQL Server	MySQL	Oracle	PostgreSQL
db.z1d – Latest Generation Memory Optimized Instance Classes					
db.z1d.12xlarge	No	No	No	Yes	No
db.z1d.6xlarge	No	No	No	Yes	No
db.z1d.3xlarge	No	No	No	Yes	No
db.z1d.2xlarge	No	No	No	Yes	No
db.z1d.xlarge	No	No	No	Yes	No
db.z1d.large	No	No	No	Yes	No
Instance Class	MariaDB	Microsoft SQL Server	MySQL	Oracle	PostgreSQL
db.x1e – Latest Generation Memory Optimized Instance Classes					
db.x1e.32xlarge	No	Yes	No	Yes	No
db.x1e.16xlarge	No	Yes	No	Yes	No
db.x1e.8xlarge	No	Yes	No	Yes	No
db.x1e.4xlarge	No	Yes	No	Yes	No
db.x1e.2xlarge	No	Yes	No	Yes	No
db.x1e.xlarge	No	Yes	No	Yes	No
db.x1 – Current Generation Memory Optimized Instance Classes					
db.x1.32xlarge	No	Yes	No	Yes	No
db.x1.16xlarge	No	Yes	No	Yes	No
Instance Class	MariaDB	Microsoft SQL Server	MySQL	Oracle	PostgreSQL
db.r5 – Latest Generation Memory Optimized Instance Classes					
db.r5.24xlarge	Yes	Yes	Yes	Yes	Yes
db.r5.16xlarge	No	Yes	No	Yes	No
db.r5.12xlarge	Yes	Yes	Yes	Yes	Yes

db.r5.8xlarge	No	Yes	No	Yes	No
db.r5.4xlarge	Yes	Yes	Yes	Yes	Yes
db.r5.2xlarge	Yes	Yes	Yes	Yes	Yes
db.r5.xlarge	Yes	Yes	Yes	Yes	Yes
db.r5.large	Yes	Yes	Yes	Yes	Yes
db.r4 – Current Generation Memory Optimized Instance Classes					
db.r4.16xlarge	Yes	Yes	MySQL 8.0, 5.7, 5.6	Yes	Yes
db.r4.8xlarge	Yes	Yes	MySQL 8.0, 5.7, 5.6	Yes	Yes
db.r4.4xlarge	Yes	Yes	MySQL 8.0, 5.7, 5.6	Yes	Yes
db.r4.2xlarge	Yes	Yes	MySQL 8.0, 5.7, 5.6	Yes	Yes
db.r4.xlarge	Yes	Yes	MySQL 8.0, 5.7, 5.6	Yes	Yes
db.r4.large	Yes	Yes	MySQL 8.0, 5.7, 5.6	Yes	Yes
db.r3 – Previous Generation Memory Optimized Instance Classes					
db.r3.8xlarge	Yes	Yes	Yes	Deprecated	Yes
db.r3.4xlarge	Yes	Yes	Yes	Deprecated	Yes
db.r3.2xlarge	Yes	Yes	Yes	Deprecated	Yes
db.r3.xlarge	Yes	Yes	Yes	Deprecated	Yes
db.r3.large	Yes	Yes	Yes	Deprecated	Yes
Instance Class	MariaDB	Microsoft SQL Server	MySQL	Oracle	PostgreSQL
db.m2 – Previous Generation Memory Optimized Instance Classes					
db.m2.4xlarge	No	Yes	Deprecated	Deprecated	PostgreSQL 9.4, 9.3
db.m2.2xlarge	No	Yes	Deprecated	Deprecated	PostgreSQL 9.4, 9.3

db.m2.xlarge	No	Yes	Deprecated	Deprecated	PostgreSQL 9.4, 9.3
Instance Class	MariaDB	Microsoft SQL Server	MySQL	Oracle	PostgreSQL
db.t3 – Latest Generation Burstable Performance Instance Classes					
db.t3.2xlarge	Yes	Yes	Yes	Yes	Yes
db.t3.xlarge	Yes	Yes	Yes	Yes	Yes
db.t3.large	Yes	Yes	Yes	Yes	Yes
db.t3.medium	Yes	Yes	Yes	Yes	Yes
db.t3.small	Yes	Yes	Yes	Yes	Yes
db.t3.micro	Yes	No	Yes	Yes	Yes
db.t2 – Current Generation Burstable Performance Instance Classes					
db.t2.2xlarge	Yes	No	MySQL 8.0, 5.7, 5.6	Deprecated	PostgreSQL 9.6, 9.5, 9.4
db.t2.xlarge	Yes	No	MySQL 8.0, 5.7, 5.6	Deprecated	PostgreSQL 9.6, 9.5, 9.4
db.t2.large	Yes	Yes	Yes	Deprecated	Yes
db.t2.medium	Yes	Yes	Yes	Deprecated	Yes
db.t2.small	Yes	Yes	Yes	Deprecated	Yes
db.t2.micro	Yes	Yes	Yes	Deprecated	Yes

Changing Your DB Instance Class

You can change the CPU and memory available to a DB instance by changing its DB instance class. To change the DB instance class, modify your DB instance by following the instructions for your specific database engine.

- [Modifying a DB Instance Running the MariaDB Database Engine \(p. 564\)](#)
- [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#)
- [Modifying a DB Instance Running the MySQL Database Engine \(p. 802\)](#)
- [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#)
- [Modifying a DB Instance Running the PostgreSQL Database Engine \(p. 1279\)](#)

Some instance classes require that your DB instance is in a VPC. If your current DB instance isn't in a VPC, and you want to use an instance class that requires one, first move your DB instance into a VPC. For more information, see [Moving a DB Instance Not in a VPC into a VPC \(p. 520\)](#).

Configuring the Processor for a DB Instance Class

Amazon RDS DB instance classes support Intel Hyper-Threading Technology, which enables multiple threads to run concurrently on a single Intel Xeon CPU core. Each thread is represented as a virtual CPU (vCPU) on the DB instance. A DB instance has a default number of CPU cores, which varies according to DB instance type. For example, a db.m4.xlarge DB instance type has two CPU cores and two threads per core by default—four vCPUs in total.

Note

Each vCPU is a hyperthread of an Intel Xeon CPU core.

Topics

- [Overview of Configuring the Processor \(p. 117\)](#)
- [CPU Cores and Threads Per CPU Core Per DB Instance Class \(p. 117\)](#)
- [Setting the CPU Cores and Threads per CPU Core for a DB Instance Class \(p. 120\)](#)

Overview of Configuring the Processor

In most cases, you can find a DB instance class that has a combination of memory and number of vCPUs to suit your workloads. However, you can also specify the following processor features to optimize your DB instance for specific workloads or business needs:

- **Number of CPU cores** – You can customize the number of CPU cores for the DB instance. You might do this to potentially optimize the licensing costs of your software with a DB instance that has sufficient amounts of RAM for memory-intensive workloads but fewer CPU cores.
- **Threads per core** – You can disable Intel Hyper-Threading Technology by specifying a single thread per CPU core. You might do this for certain workloads, such as high-performance computing (HPC) workloads.

You can control the number of CPU cores and threads for each core separately. You can set one or both in a request. After a setting is associated with a DB instance, the setting persists until you change it.

The processor settings for a DB instance are associated with snapshots of the DB instance. When a snapshot is restored, its restored DB instance uses the processor feature settings used when the snapshot was taken.

If you modify the DB instance class for a DB instance with nondefault processor settings, either specify default processor settings or explicitly specify processor settings at modification. This requirement ensures that you are aware of the third-party licensing costs that might be incurred when you modify the DB instance.

There is no additional or reduced charge for specifying processor features on an Amazon RDS DB instance. You're charged the same as for DB instances that are launched with default CPU configurations.

CPU Cores and Threads Per CPU Core Per DB Instance Class

In the following table, you can find the DB instance classes that support setting a number of CPU cores and CPU threads per core. You can also find the default value and the valid values for the number of CPU cores and CPU threads per core for each DB instance class.

DB Instance Class	Default vCPUs	Default CPU Cores	Default Threads per Core	Valid Number of CPU Cores	Valid Number of Threads per Core
db.m5.large	2	1	2	1	1, 2

DB Instance Class	Default vCPUs	Default CPU Cores	Default Threads per Core	Valid Number of CPU Cores	Valid Number of Threads per Core
db.m5.xlarge	4	2	2	2	1, 2
db.m5.2xlarge	8	4	2	2, 4	1, 2
db.m5.4xlarge	16	8	2	2, 4, 6, 8	1, 2
db.m5.8xlarge	32	16	2	2, 4, 6, 8, 10, 12, 14, 16	1, 2
db.m5.12xlarge	48	24	2	2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24	1, 2
db.m5.16xlarge	64	32	2	2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32	1, 2
db.m5.24xlarge	96	48	2	4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32, 34, 36, 38, 40, 42, 44, 46, 48	1, 2
db.m4.10xlarge	40	20	2	2, 4, 6, 8, 10, 12, 14, 16, 18, 20	1, 2
db.m4.16xlarge	64	32	2	2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32	1, 2
db.r3.large	2	1	2	1	1, 2
db.r3.xlarge	4	2	2	1, 2	1, 2
db.r3.2xlarge	8	4	2	1, 2, 3, 4	1, 2
db.r3.4xlarge	16	8	2	1, 2, 3, 4, 5, 6, 7, 8	1, 2
db.r3.8xlarge	32	16	2	2, 4, 6, 8, 10, 12, 14, 16	1, 2
db.r5.large	2	1	2	1	1, 2
db.r5.xlarge	4	2	2	2	1, 2
db.r5.2xlarge	8	4	2	2, 4	1, 2
db.r5.4xlarge	16	8	2	2, 4, 6, 8	1, 2
db.r5.8xlarge	32	16	2	2, 4, 6, 8, 10, 12, 14, 16	1, 2

DB Instance Class	Default vCPUs	Default CPU Cores	Default Threads per Core	Valid Number of CPU Cores	Valid Number of Threads per Core
db.r5.12xlarge	48	24	2	2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24	1, 2
db.r5.16xlarge	64	32	2	2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32	1, 2
db.r5.24xlarge	96	48	2	4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32, 34, 36, 38, 40, 42, 44, 46, 48	1, 2
db.r4.large	2	1	2	1	1, 2
db.r4.xlarge	4	2	2	1, 2	1, 2
db.r4.2xlarge	8	4	2	1, 2, 3, 4	1, 2
db.r4.4xlarge	16	8	2	1, 2, 3, 4, 5, 6, 7, 8	1, 2
db.r4.8xlarge	32	16	2	1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16	1, 2
db.r4.16xlarge	64	32	2	2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32	1, 2
db.x1.16xlarge	64	32	2	2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32	1, 2
db.x1.32xlarge	128	64	2	4, 8, 12, 16, 20, 24, 28, 32, 36, 40, 44, 48, 52, 56, 60, 64	1, 2
db.x1e.xlarge	4	2	2	1, 2	1, 2
db.x1e.2xlarge	8	4	2	1, 2, 3, 4	1, 2
db.x1e.4xlarge	16	8	2	1, 2, 3, 4, 5, 6, 7, 8	1, 2
db.x1e.8xlarge	32	16	2	1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16	1, 2

DB Instance Class	Default vCPUs	Default CPU Cores	Default Threads per Core	Valid Number of CPU Cores	Valid Number of Threads per Core
db.x1e.16xlarge	64	32	2	2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32	1, 2
db.x1e.32xlarge	128	64	2	4, 8, 12, 16, 20, 24, 28, 32, 36, 40, 44, 48, 52, 56, 60, 64	1, 2
db.z1d.large	2	1	2	1	1, 2
db.z1d.xlarge	4	2	2	2	1, 2
db.z1d.2xlarge	8	4	2	2, 4	1, 2
db.z1d.3xlarge	12	6	2	2, 4, 6	1, 2
db.z1d.6xlarge	24	12	2	2, 4, 6, 8, 10, 12	1, 2
db.z1d.12xlarge	48	24	2	4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24	1, 2

Note

Currently, you can configure the number of CPU cores and threads per core only for Oracle DB instances. For information about the DB instance classes supported by different Oracle database editions, see [DB Instance Class Support for Oracle \(p. 917\)](#).

For Oracle DB instances, configuring the number of CPU cores and threads per core is only supported with the Bring Your Own License (BYOL) licensing option. For more information about Oracle licensing options, see [Oracle Licensing \(p. 916\)](#).

You can use AWS CloudTrail to monitor and audit changes to the process configuration of Amazon RDS for Oracle DB instances. For more information about using CloudTrail, see [Logging Amazon RDS API Calls with AWS CloudTrail \(p. 423\)](#).

Setting the CPU Cores and Threads per CPU Core for a DB Instance Class

You can configure the number of CPU cores and threads per core for the DB instance class when you perform the following operations:

- [Creating an Amazon RDS DB Instance \(p. 148\)](#)
- [Modifying an Amazon RDS DB Instance \(p. 150\)](#)
- [Restoring from a DB Snapshot \(p. 266\)](#)
- [Restoring a DB Instance to a Specified Time \(p. 285\)](#)

Note

When you modify a DB instance to configure the number of CPU cores or threads per core, there is a brief DB instance outage.

You can set the CPU cores and the threads per CPU core for a DB instance class using the AWS Management Console, the AWS CLI, or the RDS API.

[Console](#)

When you are creating, modifying, or restoring a DB instance, you set the DB instance class in the AWS Management Console. The **Instance specifications** section shows options for the processor. The following image shows the processor features options.

Instance specifications

Estimate your monthly costs for the DB Instance using the [AWS Simple Monthly Calculator](#)

DB engine
Oracle Database Enterprise Edition

License model [Info](#)

DB engine version [Info](#)

DB instance class [Info](#)

Multi-AZ deployment [Info](#)

Create replica in different zone
Creates a replica in a different Availability Zone (AZ) to provide data redundancy, eliminate I/O freezes, and minimize latency spikes during system backups.

No

Storage type [Info](#)

Allocated storage
 GiB
(Minimum: 100 GiB, Maximum: 16384 GiB)

Provisioned IOPS [Info](#)

▼ Additional configuration

Processor features

Override default values
You can change the number of CPU cores and threads per core on the DB instance class.

Core count [Info](#)

Threads per core [Info](#)

Set the following options to the appropriate values for your DB instance class under **Processor features**:

- **Core count** – Set the number of CPU cores using this option. The value must be equal to or less than the maximum number of CPU cores for the DB instance class.
- **Threads per core** – Specify **2** to enable multiple threads per core, or specify **1** to disable multiple threads per core.

When you modify or restore a DB instance, you can also set the CPU cores and the threads per CPU core to the defaults for the instance class.

When you view the details for a DB instance in the console, you can view the processor information for its DB instance class on the **Configuration** tab. The following image shows a DB instance class with one CPU core and multiple threads per core enabled.

Instance and IOPS	
Instance Class	db.r4.large
Core count	1
Threads per core	2
vCPU enabled	2
Storage Type	Provisioned IOPS (SSD)
IOPS	1000
Storage	100 GiB

For Oracle DB instances, the processor information only appears for Bring Your Own License (BYOL) DB instances.

AWS CLI

You can set the processor features for a DB instance when you run one of the following AWS CLI commands:

- [create-db-instance](#)

- [modify-db-instance](#)
- [restore-db-instance-from-db-snapshot](#)
- [restore-db-instance-from-s3](#)
- [restore-db-instance-to-point-in-time](#)

To configure the processor of a DB instance class for a DB instance by using the AWS CLI, include the `--processor-features` option in the command. Specify the number of CPU cores with the `coreCount` feature name, and specify whether multiple threads per core are enabled with the `threadsPerCore` feature name.

The option has the following syntax.

```
--processor-features "Name=coreCount,Value=<value>" "Name=threadsPerCore,Value=<value>"
```

The following are examples that configure the processor:

Examples

- [Setting the Number of CPU Cores for a DB Instance \(p. 124\)](#)
- [Setting the Number of CPU Cores and Disabling Multiple Threads for a DB Instance \(p. 124\)](#)
- [Viewing the Valid Processor Values for a DB Instance Class \(p. 125\)](#)
- [Returning to Default Processor Settings for a DB Instance \(p. 126\)](#)
- [Returning to the Default Number of CPU Cores for a DB Instance \(p. 126\)](#)
- [Returning to the Default Number of Threads Per Core for a DB Instance \(p. 127\)](#)

Setting the Number of CPU Cores for a DB Instance

Example

The following example modifies `mydbinstance` by setting the number of CPU cores to 4. The changes are applied immediately by using `--apply-immediately`. If you want to apply the changes during the next scheduled maintenance window, omit the `--apply-immediately` option.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
  --processor-features "Name=coreCount,Value=4" \
  --apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
  --processor-features "Name=coreCount,Value=4" ^
  --apply-immediately
```

Setting the Number of CPU Cores and Disabling Multiple Threads for a DB Instance

Example

The following example modifies `mydbinstance` by setting the number of CPU cores to 4 and disabling multiple threads per core. The changes are applied immediately by using `--apply-immediately`. If

you want to apply the changes during the next scheduled maintenance window, omit the `--apply-immediately` option.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--processor-features "Name=coreCount,Value=4" "Name=threadsPerCore,Value=1" \
--apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
--processor-features "Name=coreCount,Value=4" "Name=threadsPerCore,Value=1" ^
--apply-immediately
```

Viewing the Valid Processor Values for a DB Instance Class

Example

You can view the valid processor values for a particular DB instance class by running the [describe-orderable-db-instance-options](#) command and specifying the instance class for the `--db-instance-class` option. For example, the output for the following command shows the processor options for the db.r3.large instance class.

```
aws rds describe-orderable-db-instance-options --engine oracle-ee --db-instance-class db.r3.large
```

Following is sample output for the command in JSON format.

```
{
    "SupportsIops": true,
    "MaxIopsPerGib": 50.0,
    "LicenseModel": "bring-your-own-license",
    "DBInstanceClass": "db.r3.large",
    "SupportsSIAMDatabaseAuthentication": false,
    "MinStorageSize": 100,
    "AvailabilityZones": [
        {
            "Name": "us-west-2a"
        },
        {
            "Name": "us-west-2b"
        },
        {
            "Name": "us-west-2c"
        }
    ],
    "EngineVersion": "12.1.0.2.v2",
    "MaxStorageSize": 32768,
    "MinIopsPerGib": 1.0,
    "MaxIopsPerDbInstance": 40000,
    "ReadReplicaCapable": false,
    "AvailableProcessorFeatures": [
        {
            "Name": "coreCount",
            "DefaultValue": "1",
            "AllowedValues": "1"
        },
        ...
    ]
}
```

```
{  
    "Name": "threadsPerCore",  
    "DefaultValue": "2",  
    "AllowedValues": "1,2"  
}  
],  
"SupportsEnhancedMonitoring": true,  
"SupportsPerformanceInsights": false,  
"MinIopsPerDbInstance": 1000,  
"StorageType": "io1",  
"Vpc": false,  
"SupportsStorageEncryption": true,  
"Engine": "oracle-ee",  
"MultiAZCapable": true  
}
```

In addition, you can run the following commands for DB instance class processor information:

- [describe-db-instances](#) – Shows the processor information for the specified DB instance.
- [describe-db-snapshots](#) – Shows the processor information for the specified DB snapshot.
- [describe-valid-db-instance-modifications](#) – Shows the valid modifications to the processor for the specified DB instance.

Returning to Default Processor Settings for a DB Instance

Example

The following example modifies `mydbinstance` by returning its DB instance class to the default processor values for it. The changes are applied immediately by using `--apply-immediately`. If you want to apply the changes during the next scheduled maintenance window, omit the `--apply-immediately` option.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \  
  --use-default-processor-features \  
  --apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^  
  --use-default-processor-features ^  
  --apply-immediately
```

Returning to the Default Number of CPU Cores for a DB Instance

Example

The following example modifies `mydbinstance` by returning its DB instance class to the default number of CPU cores for it. The threads per core setting isn't changed. The changes are applied immediately by using `--apply-immediately`. If you want to apply the changes during the next scheduled maintenance window, omit the `--apply-immediately` option.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \  
  --processor-features "Name=coreCount,Value=DEFAULT" \  
  --apply-immediately
```

```
--apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
--processor-features "Name=coreCount,Value=DEFAULT" ^
--apply-immediately
```

Returning to the Default Number of Threads Per Core for a DB Instance

Example

The following example modifies `mydbinstance` by returning its DB instance class to the default number of threads per core for it. The number of CPU cores setting isn't changed. The changes are applied immediately by using `--apply-immediately`. If you want to apply the changes during the next scheduled maintenance window, omit the `--apply-immediately` option.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--processor-features "Name=threadsPerCore,Value=DEFAULT" \
--apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
--processor-features "Name=threadsPerCore,Value=DEFAULT" ^
--apply-immediately
```

RDS API

You can set the processor features for a DB instance when you call one of the following Amazon RDS API operations:

- [CreateDBInstance](#)
- [ModifyDBInstance](#)
- [RestoreDBInstanceFromDBSnapshot](#)
- [RestoreDBInstanceFromS3](#)
- [RestoreDBInstanceToPointInTime](#)

To configure the processor features of a DB instance class for a DB instance by using the Amazon RDS API, include the `ProcessFeatures` parameter in the call.

The parameter has the following syntax.

```
ProcessFeatures "Name=coreCount,Value=<value>" "Name=threadsPerCore,Value=<value>"
```

Specify the number of CPU cores with the `coreCount` feature name, and specify whether multiple threads per core are enabled with the `threadsPerCore` feature name.

You can view the valid processor values for a particular instance class by running the [DescribeOrderableDBInstanceOptions](#) operation and specifying the instance class for the `DBInstanceClass` parameter.

In addition, you can use the following actions for DB instance class processor information:

- [DescribeDBInstances](#) – Shows the processor information for the specified DB instance.
- [DescribeDBSnapshots](#) – Shows the processor information for the specified DB snapshot.
- [DescribeValidDBInstanceModifications](#) – Shows the valid modifications to the processor for the specified DB instance.

DB Instance Status

The status of a DB instance indicates the health of the DB instance. You can view the status of a DB instance by using the Amazon RDS console, the AWS CLI command [describe-db-instances](#), or the API operation [DescribeDBInstances](#).

Note

Amazon RDS also uses another status called *maintenance status*, which is shown in the **Maintenance** column of the Amazon RDS console. This value indicates the status of any maintenance patches that need to be applied to a DB instance. Maintenance status is independent of DB instance status. For more information on *maintenance status*, see [Applying Updates for a DB Instance \(p. 154\)](#).

Find the possible status values for DB instances in the following table, which also shows how you are billed for each status. It shows if you will be billed for the DB instance and storage, billed only for storage, or not billed. For all DB instance statuses, you are always billed for backup usage.

DB Instance Status	Billed	Description
available	Billed	The DB instance is healthy and available.
backing-up	Billed	The DB instance is currently being backed up.
backtracking	Billed	The DB instance is currently being backtracked. This status only applies to Aurora MySQL.
configuring-enhanced-monitoring	Billed	Enhanced Monitoring is being enabled or disabled for this DB instance.
configuring-iam-database-auth	Billed	AWS Identity and Access Management (IAM) database authentication is being enabled or disabled for this DB instance.
configuring-log-exports	Billed	Publishing log files to Amazon CloudWatch Logs is being enabled or disabled for this DB instance.
converting-to-vpc	Billed	The DB instance is being converted from a DB instance that is not in an Amazon Virtual Private Cloud (Amazon VPC) to a DB instance that is in an Amazon VPC.
creating	Not billed	The DB instance is being created. The DB instance is inaccessible while it is being created.
deleting	Not billed	The DB instance is being deleted.
failed	Not billed	The DB instance has failed and Amazon RDS can't recover it. Perform a point-in-time restore to the latest restorable time of the DB instance to recover the data.
inaccessible-encryption-credentials	Not billed	The AWS KMS key used to encrypt or decrypt the DB instance can't be accessed.
incompatible-network	Not billed	Amazon RDS is attempting to perform a recovery action on a DB instance but can't do so because the VPC is in a state that prevents the action from being completed. This status can occur if, for example, all available IP addresses in a subnet are in use and Amazon RDS can't get an IP address for the DB instance.

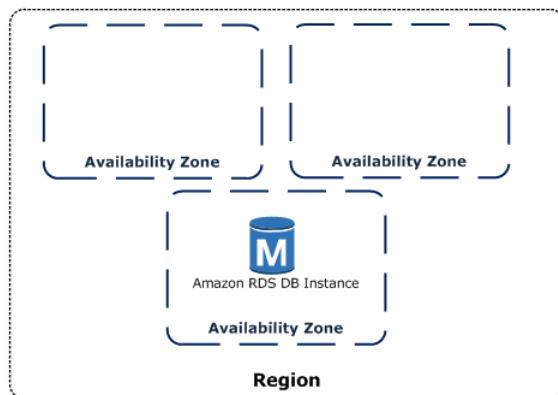
DB Instance Status	Billed	Description
incompatible-option-group	Billed	Amazon RDS attempted to apply an option group change but can't do so, and Amazon RDS can't roll back to the previous option group state. For more information, check the Recent Events list for the DB instance. This status can occur if, for example, the option group contains an option such as TDE and the DB instance doesn't contain encrypted information.
incompatible-parameters	Billed	Amazon RDS can't start the DB instance because the parameters specified in the DB instance's DB parameter group aren't compatible with the DB instance. Revert the parameter changes or make them compatible with the DB instance to regain access to your DB instance. For more information about the incompatible parameters, check the Recent Events list for the DB instance.
incompatible-restore	Not billed	Amazon RDS can't do a point-in-time restore. Common causes for this status include using temp tables, using MyISAM tables with MySQL, or using Aria tables with MariaDB.
maintenance	Billed	Amazon RDS is applying a maintenance update to the DB instance. This status is used for instance-level maintenance that RDS schedules well in advance.
modifying	Billed	The DB instance is being modified because of a customer request to modify the DB instance.
moving-to-vpc	Billed	The DB instance is being moved to a new Amazon Virtual Private Cloud (Amazon VPC).
rebooting	Billed	The DB instance is being rebooted because of a customer request or an Amazon RDS process that requires the rebooting of the DB instance.
renaming	Billed	The DB instance is being renamed because of a customer request to rename it.
resetting-master-credentials	Billed	The master credentials for the DB instance are being reset because of a customer request to reset them.
restore-error	Billed	The DB instance encountered an error attempting to restore to a point-in-time or from a snapshot.
starting	Billed for storage	The DB instance is starting.
stopped	Billed for storage	The DB instance is stopped.
stopping	Billed for storage	The DB instance is being stopped.

DB Instance Status	Billed	Description
storage-full	Billed	The DB instance has reached its storage capacity allocation. This is a critical status, and we recommend that you fix this issue immediately. To do so, scale up your storage by modifying the DB instance. To avoid this situation, set Amazon CloudWatch alarms to warn you when storage space is getting low.
storage-optimization	Billed	Your DB instance is being modified to change the storage size or type. The DB instance is fully operational. However, while the status of your DB instance is storage-optimization , you can't request any changes to the storage of your DB instance. The storage optimization process is usually short, but can sometimes take up to and even beyond 24 hours.
upgrading	Billed	The database engine version is being upgraded.

Regions and Availability Zones

Amazon cloud computing resources are hosted in multiple locations world-wide. These locations are composed of AWS Regions and Availability Zones. Each *AWS Region* is a separate geographic area. Each AWS Region has multiple, isolated locations known as *Availability Zones*. Amazon RDS provides you the ability to place resources, such as instances, and data in multiple locations. Resources aren't replicated across AWS Regions unless you do so specifically.

Amazon operates state-of-the-art, highly-available data centers. Although rare, failures can occur that affect the availability of instances that are in the same location. If you host all your instances in a single location that is affected by such a failure, none of your instances would be available.



It is important to remember that each AWS Region is completely independent. Any Amazon RDS activity you initiate (for example, creating database instances or listing available database instances) runs only in your current default AWS Region. The default AWS Region can be changed in the console, by setting the EC2_REGION environment variable, or it can be overridden by using the --region parameter with the AWS Command Line Interface. See [Configuring the AWS Command Line Interface](#), specifically, the sections on Environment Variables and Command Line Options for more information.

Amazon RDS supports a special AWS Region called AWS GovCloud (US-West) that is designed to allow US government agencies and customers to move more sensitive workloads into the cloud. AWS GovCloud (US-West) addresses the US government's specific regulatory and compliance requirements. For more information about AWS GovCloud (US-West), see [What Is AWS GovCloud \(US-West\)?](#)

To create or work with an Amazon RDS DB instance in a specific AWS Region, use the corresponding regional service endpoint.

Amazon RDS supports the endpoints listed in the following table.

Region Name	Region	Endpoint	Protocol	
US East (Ohio)	us-east-2	rds.us-east-2.amazonaws.com	HTTPS	
US East (N. Virginia)	us-east-1	rds.us-east-1.amazonaws.com	HTTPS	
US West (N. California)	us-west-1	rds.us-west-1.amazonaws.com	HTTPS	

Region Name	Region	Endpoint	Protocol	
US West (Oregon)	us-west-2	rds.us-west-2.amazonaws.com	HTTPS	
Asia Pacific (Hong Kong)	ap-east-1	rds.ap-east-1.amazonaws.com	HTTPS	
Asia Pacific (Mumbai)	ap-south-1	rds.ap-south-1.amazonaws.com	HTTPS	
Asia Pacific (Osaka-Local)	ap-northeast-3	rds.ap-northeast-3.amazonaws.com	HTTPS	
Asia Pacific (Seoul)	ap-northeast-2	rds.ap-northeast-2.amazonaws.com	HTTPS	
Asia Pacific (Singapore)	ap-southeast-1	rds.ap-southeast-1.amazonaws.com	HTTPS	
Asia Pacific (Sydney)	ap-southeast-2	rds.ap-southeast-2.amazonaws.com	HTTPS	
Asia Pacific (Tokyo)	ap-northeast-1	rds.ap-northeast-1.amazonaws.com	HTTPS	
Canada (Central)	ca-central-1	rds.ca-central-1.amazonaws.com	HTTPS	
China (Beijing)	cn-north-1	rds.cn-north-1.amazonaws.com.cn	HTTPS	
China (Ningxia)	cn-northwest-1	rds.cn-northwest-1.amazonaws.com.cn	HTTPS	
EU (Frankfurt)	eu-central-1	rds.eu-central-1.amazonaws.com	HTTPS	
EU (Ireland)	eu-west-1	rds.eu-west-1.amazonaws.com	HTTPS	
EU (London)	eu-west-2	rds.eu-west-2.amazonaws.com	HTTPS	
EU (Paris)	eu-west-3	rds.eu-west-3.amazonaws.com	HTTPS	
EU (Stockholm)	eu-north-1	rds.eu-north-1.amazonaws.com	HTTPS	

Region Name	Region	Endpoint	Protocol	
Middle East (Bahrain)	me-south-1	rds.me-south-1.amazonaws.com	HTTPS	
South America (Sao Paulo)	sa-east-1	rds.sa-east-1.amazonaws.com	HTTPS	
AWS GovCloud (US-East)	us-gov-east-1	rds.us-gov-east-1.amazonaws.com	HTTPS	
AWS GovCloud (US-West)	us-gov-west-1	rds.us-gov-west-1.amazonaws.com	HTTPS	

If you do not explicitly specify an endpoint, the US West (Oregon) endpoint is the default.

Amazon RDS DB Instance Storage

DB instances for Amazon RDS for MySQL, MariaDB, PostgreSQL, Oracle, and Microsoft SQL Server use Amazon Elastic Block Store (Amazon EBS) volumes for database and log storage. Depending on the amount of storage requested, Amazon RDS automatically stripes across multiple Amazon EBS volumes to enhance performance.

Amazon RDS Storage Types

Amazon RDS provides three storage types: General Purpose SSD (also known as gp2), Provisioned IOPS SSD (also known as io1), and magnetic. They differ in performance characteristics and price, which means that you can tailor your storage performance and cost to the needs of your database workload. You can create MySQL, MariaDB, and PostgreSQL RDS DB instances with up to 64 TiB of storage. You can create Oracle RDS DB instances with up to 64 TiB of storage. You can create SQL Server RDS DB instances with up to 16 TiB of storage. For this amount of storage, use the Provisioned IOPS SSD and General Purpose SSD storage types.

The following list briefly describes the three storage types:

- **General Purpose SSD** – General Purpose SSD volumes offer cost-effective storage that is ideal for a broad range of workloads. These volumes deliver single-digit millisecond latencies and the ability to burst to 3,000 IOPS for extended periods of time. Baseline performance for these volumes is determined by the volume's size.

For more information about General Purpose SSD storage, including the storage size ranges, see [General Purpose SSD Storage \(p. 135\)](#).

- **Provisioned IOPS** – Provisioned IOPS storage is designed to meet the needs of I/O-intensive workloads, particularly database workloads, that require low I/O latency and consistent I/O throughput.

For more information about provisioned IOPS storage, including the storage size ranges, see [Provisioned IOPS SSD Storage \(p. 137\)](#).

- **Magnetic** – Amazon RDS also supports magnetic storage for backward compatibility. We recommend that you use General Purpose SSD or Provisioned IOPS for any new storage needs. The maximum amount of storage allowed for DB instances on magnetic storage is less than that of the other storage types. For more information, see [Magnetic storage \(p. 138\)](#).

Several factors can affect the performance of Amazon EBS volumes, such as instance configuration, I/O characteristics, and workload demand. For more information about getting the most out of your Provisioned IOPS volumes, see [Amazon EBS Volume Performance](#).

General Purpose SSD Storage

General Purpose SSD storage offers cost-effective storage that is acceptable for most database workloads. The following are the storage size ranges for General Purpose SSD DB instances:

- MariaDB, MySQL and PostgreSQL database instances: 20 GiB–64 TiB
- SQL Server for Enterprise, Standard, Web, and Express editions: 20 GiB–16 TiB
- Oracle instances: 20 GiB– 64 TiB

Baseline I/O performance for General Purpose SSD storage is 3 IOPS for each GiB. This relationship means that larger volumes have better performance. For example, baseline performance for a 100-GiB volume is 300 IOPS. Baseline performance for a 1-TiB volume is 3,000 IOPS. And baseline performance for a 5.34-TiB volume is 16,000 IOPS.

Volumes below 1 TiB in size also have ability to burst to 3,000 IOPS for extended periods of time. Burst is not relevant for volumes above 1 TiB. Instance I/O credit balance determines burst performance. For more information about instance I/O credits see, [I/O Credits and Burst Performance \(p. 136\)](#).

Many workloads never deplete the burst balance, making General Purpose SSD an ideal storage choice for many workloads. However, some workloads can exhaust the 3,000 IOPS burst storage credit balance, so you should plan your storage capacity to meet the needs of your workloads.

I/O Credits and Burst Performance

General Purpose SSD storage performance is governed by volume size, which dictates the base performance level of the volume and how quickly it accumulates I/O credits. Larger volumes have higher base performance levels and accumulate I/O credits faster. *I/O credits* represent the available bandwidth that your General Purpose SSD storage can use to burst large amounts of I/O when more than the base level of performance is needed. The more I/O credits your storage has for I/O, the more time it can burst beyond its base performance level and the better it performs when your workload requires more performance.

When using General Purpose SSD storage, your DB instance receives an initial I/O credit balance of 5.4 million I/O credits. This initial credit balance is enough to sustain a burst performance of 3,000 IOPS for 30 minutes. This balance is designed to provide a fast initial boot cycle for boot volumes and to provide a good bootstrapping experience for other applications. Volumes earn I/O credits at the baseline performance rate of 3 IOPS for each GiB of volume size. For example, a 100-GiB SSD volume has a baseline performance of 300 IOPS.

When your storage requires more than the base performance I/O level, it uses I/O credits in the I/O credit balance to burst to the required performance level. Such a burst goes to a maximum of 3,000 IOPS. Storage larger than 1,000 GiB has a base performance that is equal or greater than the maximum burst performance. When your storage uses fewer I/O credits than it earns in a second, unused I/O credits are added to the I/O credit balance. The maximum I/O credit balance for a DB instance using General Purpose SSD storage is equal to the initial I/O credit balance (5.4 million I/O credits).

Suppose that your storage uses all of its I/O credit balance. If so, its maximum performance remains at the base performance level until I/O demand drops below the base level and unused I/O credits are added to the I/O credit balance. (The *base performance level* is the rate at which your storage earns I/O credits.) The more storage, the greater the base performance is and the faster it replenishes the I/O credit balance.

Note

Storage conversions between magnetic storage and General Purpose SSD storage can potentially deplete your I/O credit balance, resulting in longer conversion times. For more information about scaling storage, see [Working with Storage for Amazon RDS DB Instances \(p. 231\)](#).

The following table lists several storage sizes. For each storage size, it lists the associated base performance of the storage, which is also the rate at which it accumulates I/O credits. The table also lists the burst duration at the 3,000 IOPS maximum, when starting with a full I/O credit balance. In addition, the table lists the time in seconds that the storage takes to refill an empty I/O credit balance.

Note

The IOPS figure reaches its maximum value at a volume storage size of 5,334 GiB.

Storage size (GiB)	Base Performance (IOPS)	Maximum Burst Duration at 3,000 IOPS (Seconds)	Seconds to Fill Empty I/O Credit Balance
1	100	1,862	54,000
100	300	2,000	18,000

Storage size (GiB)	Base Performance (IOPS)	Maximum Burst Duration at 3,000 IOPS (Seconds)	Seconds to Fill Empty I/O Credit Balance
250	750	2,400	7,200
500	1,500	3,600	3,600
750	2,250	7,200	2,400
1,000	3,000	Infinite	N/A
5,334	16,000	N/A	N/A

The burst duration of your storage depends on the size of the storage, the burst IOPS required, and the I/O credit balance when the burst begins. This relationship is shown in the equation following.

$$\text{Burst duration} = \frac{(\text{Credit balance})}{(\text{Burst IOPS}) - 3(\text{Storage size in GiB})}$$

You might notice that your storage performance is frequently limited to the base level due to an empty I/O credit balance. If so, consider allocating more General Purpose SSD storage with a higher base performance level. Alternatively, you can switch to Provisioned IOPS storage for workloads that require sustained IOPS performance.

For workloads with steady state I/O requirements, provisioning less than 100 GiB of General Purpose SSD storage might result in higher latencies if you exhaust your I/O credit balance.

Note

In general, most workloads never exceed the I/O credit balance.

For a more detailed description of how baseline performance and I/O credit balance affect performance see [Understanding Burst vs. Baseline Performance with Amazon RDS and GP2](#).

Provisioned IOPS SSD Storage

For production application that requires fast and consistent I/O performance, we recommend Provisioned IOPS (input/output operations per second) storage. Provisioned IOPS storage is a storage type that delivers predictable performance, and consistently low latency. Provisioned IOPS storage is optimized for online transaction processing (OLTP) workloads that have consistent performance requirements. Provisioned IOPS helps performance tuning of these workloads.

When you create a DB instance, you specify an IOPS rate and the size of the volume. Amazon RDS provides that IOPS rate for the DB instance until you change it.

Note

Your database workload might not be able to achieve 100 percent of the IOPS that you have provisioned.

The following table shows the range of Provisioned IOPS and storage size range for each database engine.

Database Engine	Range of Provisioned IOPS	Range of Storage
MariaDB	1,000–80,000 IOPS	100 GiB–64 TiB

Database Engine	Range of Provisioned IOPS	Range of Storage
SQL Server, Enterprise and Standard editions	1,000–64,000 IOPS*	20 GiB–16 TiB
SQL Server, Web and Express editions	1,000–64,000 IOPS*	100 GiB–16 TiB
MySQL	1,000–80,000 IOPS	100 GiB–64 TiB
Oracle	1,000–80,000 IOPS	100 GiB–64 TiB
PostgreSQL	1,000–80,000 IOPS	100 GiB–64 TiB

* Maximum IOPS of 64,000 is guaranteed only on [Nitro-based instances](#) that are on m5 instance types. Other instance families guarantee performance up to 32,000 IOPS.

Combining Provisioned IOPS Storage with Multi-AZ deployments, or Read Replicas

For production OLTP use cases, we recommend that you use Multi-AZ deployments for enhanced fault tolerance with Provisioned IOPS storage for fast and predictable performance.

You can also use Provisioned IOPS SSD storage with Read Replicas for MySQL, MariaDB or PostgreSQL. The type of storage for a Read Replica is independent of that on the master DB instance. For example, you might use General Purpose SSD for Read Replicas with a master DB instance that uses Provisioned IOPS SSD storage to reduce costs. However, your Read Replicas performance in this case might differ from that of a configuration where both the master DB instance and the Read Replicas use Provisioned IOPS SSD storage.

Provisioned IOPS Storage Costs

With Provisioned IOPS storage, you are charged for the provisioned resources whether or not you use them in a given month.

For more information about pricing, see [Amazon RDS Pricing](#).

Getting the most out of Amazon RDS Provisioned IOPS SSD storage

If your workload is I/O constrained, using Provisioned IOPS SSD storage can increase the number of I/O requests that the system can process concurrently. Increased concurrency allows for decreased latency because I/O requests spend less time in a queue. Decreased latency allows for faster database commits, which improves response time and allows for higher database throughput.

Provisioned IOPS SSD storage provides a way to reserve I/O capacity by specifying IOPS. However, as with any other system capacity attribute, its maximum throughput under load is constrained by the resource that is consumed first. That resource might be network bandwidth, CPU, memory, or database internal resources.

Magnetic storage

Amazon RDS also supports magnetic storage for backward compatibility. We recommend that you use General Purpose SSD or Provisioned IOPS SSD for any new storage needs. The following are some limitations for magnetic storage:

- Doesn't allow you to scale storage when using the SQL Server database engine.
- Doesn't support elastic volumes.
- Limited to a maximum size of 3 TiB.
- Limited to a maximum of 1,000 IOPS.

Monitoring storage performance

Amazon RDS provides several metrics that you can use to determine how your DB instance is performing. You can view the metrics on the summary page for your instance in Amazon RDS Management Console. You can also use Amazon CloudWatch to monitor these metrics. For more information, see [Viewing DB Instance Metrics \(p. 304\)](#). Enhanced Monitoring provides more detailed I/O metrics; for more information, see [Enhanced Monitoring \(p. 307\)](#).

The following metrics are useful for monitoring storage for your DB instance:

- **IOPS** – The number of I/O operations completed each second. This metric is reported as the average IOPS for a given time interval. Amazon RDS reports read and write IOPS separately on 1-minute intervals. Total IOPS is the sum of the read and write IOPS. Typical values for IOPS range from zero to tens of thousands per second.
- **Latency** – The elapsed time between the submission of an I/O request and its completion. This metric is reported as the average latency for a given time interval. Amazon RDS reports read and write latency separately on 1-minute intervals in units of seconds. Typical values for latency are in the millisecond (ms). For example, Amazon RDS reports 2 ms as 0.002 seconds.
- **Throughput** – The number of bytes each second that are transferred to or from disk. This metric is reported as the average throughput for a given time interval. Amazon RDS reports read and write throughput separately on 1-minute intervals using units of megabytes per second (MB/s). Typical values for throughput range from zero to the I/O channel's maximum bandwidth.
- **Queue Depth** – The number of I/O requests in the queue waiting to be serviced. These are I/O requests that have been submitted by the application but have not been sent to the device because the device is busy servicing other I/O requests. Time spent waiting in the queue is a component of latency and service time (not available as a metric). This metric is reported as the average queue depth for a given time interval. Amazon RDS reports queue depth in 1-minute intervals. Typical values for queue depth range from zero to several hundred.

Measured IOPS values are independent of the size of the individual I/O operation. This means that when you measure I/O performance, you should look at the throughput of the instance, not simply the number of I/O operations.

Factors That Affect Storage Performance

Both system activities and database workload can affect storage performance.

System activities

The following system-related activities consume I/O capacity and might reduce database instance performance while in progress:

- Multi-AZ standby creation
- Read replica creation
- Changing storage types

Database workload

In some cases your database or application design results in concurrency issues, locking, or other forms of database contention. In these cases, you might not be able to use all the provisioned bandwidth directly. In addition, you may encounter the following workload-related situations:

- The throughput limit of the underlying instance type is reached.
- Queue depth is consistently less than 1 because your application is not driving enough I/O operations.
- You experience query contention in the database even though some I/O capacity is unused.

If there isn't at least one system resource that is at or near a limit, and adding threads doesn't increase the database transaction rate, the bottleneck is most likely contention in the database. The most common forms are row lock and index page lock contention, but there are many other possibilities. If this is your situation, you should seek the advice of a database performance tuning expert.

DB instance class

To get the most performance out of your Amazon RDS database instance, choose a current generation instance type with enough bandwidth to support your storage type. For example, you can choose EBS-optimized instances and instances with 10-gigabit network connectivity.

For the full list of Amazon EC2 instance types that support EBS optimization, see [Instance types that support EBS optimization](#).

We encourage you to use the latest generation of instances to get the best performance. Previous generation DB instances have a lower instance storage limit. The following table shows the maximum storage that each DB instance class can scale to for each database engine. All values are in tebibytes (TiB).

Instance Class	Mar	Mic	MyS	Ora	Pos									
db.m5 – Latest Generation Standard Instance Classes														
db.m5.24xlarge	64	16	64	64	64									
db.m5.12xlarge	64	16	64	64	64									
db.m5.4xlarge	64	16	64	64	64									
db.m5.2xlarge	64	16	64	64	64									
db.m5.xlarge	64	16	64	64	64									
db.m5.large	64	16	64	64	64									
db.m4 – Current Generation Standard Instance Classes														
db.m4.16xlarge	64	16	64	64	64									
db.m4.10xlarge	64	16	64	64	64									
db.m4.4xlarge	64	16	64	64	64									
db.m4.2xlarge	64	16	64	64	64									
db.m4.xlarge	64	16	64	64	64									
db.m4.large	64	16	64	64	64									

Instance Class	Mari	Mic	MyS	Ora	Pos							
Instance Class	Mari	DB2	MySQL	Oracle	PostgreSQL							
db.m3 – Previous Generation Standard Instance Classes												
db.m3.2xlarge	6	16	6	6	6							
db.m3.xlarge	6	16	6	6	6							
db.m3.large	6	16	6	6	6							
db.m3.medium	32	16	32	32	32							
Instance Class	Mari	DB2	MySQL	Oracle	PostgreSQL							
db.r5 – Latest Generation Memory Optimized Instance Classes												
db.r5.24xlarge	16	16	64	64								
db.r5.12xlarge	16	16	64	64								
db.r5.4xlarge	16	16	64	64								
db.r5.2xlarge	16	16	64	64								
db.r5.xlarge	16	16	16	64	64							
db.r5.large	16	16	16	64	64							
db.r4 – Current Generation Memory Optimized Instance Classes												
db.r4.16xlarge	64	16	64	64	64							
db.r4.8xlarge	64	16	64	64	64							
db.r4.4xlarge	64	16	64	64	64							
db.r4.2xlarge	64	16	64	64	64							
db.r4.xlarge	64	16	64	64	64							
db.r4.large	64	16	64	64	64							
db.r3 – Previous Generation Memory Optimized Instance Classes												
db.r3.8xlarge	64	16	64	64	64							
db.r3.4xlarge	64	16	64	64	64							
db.r3.2xlarge	64	16	64	64	64							
db.r3.xlarge	64	16	64	64	64							

Instance Class	Mar	Mic	MySQL	Ora	PostgreSQL							
Instance Class	MariaDB	Oracle	MySQL	Oracle	PostgreSQL							
db.r3.large	64	16	64	64	64							
Instance Class	MariaDB	Oracle	MySQL	Oracle	PostgreSQL							
db.t3 – Latest Generation Burstable Performance Instance Classes												
db.t3.2xlarge	16	16	16	64	64							
db.t3.xlarge	16	16	16	64	64							
db.t3.large	16	16	16	64	64							
db.t3.medium	16	16	32	32	32							
db.t3.small	16	16	16	32	16							
db.t3.micro	16	16	16	32	16							
db.t2 – Current Generation Burstable Performance Instance Classes												
db.t2.2xlarge	64	16	64	64	64							
db.t2.xlarge	64	16	64	64	64							
db.t2.large	64	16	64	64	64							
db.t2.medium	32	16	32	32	32							
db.t2.small	16	16	16	16	16							
db.t2.micro	16	16	16	16	16							
Instance Class	MariaDB	Oracle	MySQL	Oracle	PostgreSQL							
db.x1e – Latest Generation Memory Optimized Instance Classes												
db.x1e.32xlarge		16		64								
db.x1e.16xlarge		16		64								
db.x1e.8xlarge		16		64								
db.x1e.4xlarge		16		64								
db.x1e.2xlarge		16		64								
db.x1e.xlarge		16		64								

Instance Class	Mar	Mic	MyS	Ora	Pos								
db.x1 – Current Generation Memory Optimized Instance Classes													
db.x1.32xlarge	16		64										
db.x1.16xlarge	16		64										

For Oracle, scaling up to 80,000 IOPS is only supported on the following instance classes.

- db.m5.24xlarge
- db.r5.24xlarge
- db.x1.32xlarge
- db.x1e.32xlarge

For more details on all instance classes supported, see [Previous Generation DB Instances](#).

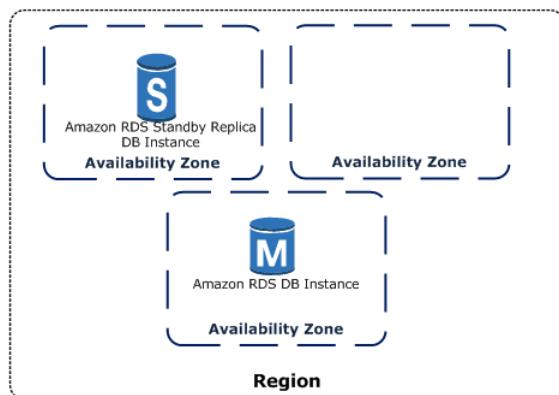
High Availability (Multi-AZ) for Amazon RDS

Amazon RDS provides high availability and failover support for DB instances using Multi-AZ deployments. Amazon RDS uses several different technologies to provide failover support. Multi-AZ deployments for Oracle, PostgreSQL, MySQL, and MariaDB DB instances use Amazon's failover technology. SQL Server DB instances use SQL Server Database Mirroring (DBM).

In a Multi-AZ deployment, Amazon RDS automatically provisions and maintains a synchronous standby replica in a different Availability Zone. The primary DB instance is synchronously replicated across Availability Zones to a standby replica to provide data redundancy, eliminate I/O freezes, and minimize latency spikes during system backups. Running a DB instance with high availability can enhance availability during planned system maintenance, and help protect your databases against DB instance failure and Availability Zone disruption. For more information on Availability Zones, see [Regions and Availability Zones \(p. 132\)](#).

Note

The high-availability feature is not a scaling solution for read-only scenarios; you cannot use a standby replica to serve read traffic. To service read-only traffic, you should use a Read Replica. For more information, see [Working with Read Replicas \(p. 182\)](#).



Using the RDS console, you can create a Multi-AZ deployment by simply specifying Multi-AZ when creating a DB instance. You can use the console to convert existing DB instances to Multi-AZ deployments by modifying the DB instance and specifying the Multi-AZ option. You can also specify a Multi-AZ deployment with the AWS CLI or Amazon RDS API. Use the [create-db-instance](#) or [modify-db-instance](#) CLI command, or the [CreateDBInstance](#) or [ModifyDBInstance](#) API operation.

The RDS console shows the Availability Zone of the standby replica (called the secondary AZ). You can also use the [describe-db-instances](#) CLI command or the [DescribeDBInstances](#) API operation to find the secondary AZ.

DB instances using Multi-AZ deployments may have increased write and commit latency compared to a Single-AZ deployment, due to the synchronous data replication that occurs. You may have a change in latency if your deployment fails over to the standby replica, although AWS is engineered with low-latency network connectivity between Availability Zones. For production workloads, we recommend that you use Provisioned IOPS and DB instance classes that are optimized for Provisioned IOPS for fast, consistent performance. For more information about DB instance classes, see [Choosing the DB Instance Class \(p. 107\)](#).

Modifying a DB Instance to Be a Multi-AZ Deployment

If you have a DB instance in a Single-AZ deployment and you modify it to be a Multi-AZ deployment (for engines other than SQL Server or Amazon Aurora), Amazon RDS takes several steps. First, Amazon RDS takes a snapshot of the primary DB instance from your deployment and then restores the snapshot into another Availability Zone. Amazon RDS then sets up synchronous replication between your primary DB instance and the new instance. This action avoids downtime when you convert from Single-AZ to Multi-AZ, but you can experience a significant performance impact when first converting to Multi-AZ. This impact is more noticeable for large and write-intensive DB instances.

Once the modification is complete, Amazon RDS triggers an event (RDS-EVENT-0025) that indicates the process is complete. You can monitor Amazon RDS events; for more information about events, see [Using Amazon RDS Event Notification \(p. 370\)](#).

Failover Process for Amazon RDS

In the event of a planned or unplanned outage of your DB instance, Amazon RDS automatically switches to a standby replica in another Availability Zone if you have enabled Multi-AZ. The time it takes for the failover to complete depends on the database activity and other conditions at the time the primary DB instance became unavailable. Failover times are typically 60-120 seconds. However, large transactions or a lengthy recovery process can increase failover time. When the failover is complete, it can take additional time for the RDS console UI to reflect the new Availability Zone.

The failover mechanism automatically changes the DNS record of the DB instance to point to the standby DB instance. As a result, you need to re-establish any existing connections to your DB instance. Due to how the Java DNS caching mechanism works, you may need to reconfigure your JVM environment. For more information on how to manage a Java application that caches DNS values in the case of a failover, see the [AWS SDK for Java](#).

Note

You can force a failover manually when you reboot a DB instance. For more information, see [Rebooting a DB Instance \(p. 166\)](#).

Amazon RDS handles failovers automatically so you can resume database operations as quickly as possible without administrative intervention. The primary DB instance switches over automatically to the standby replica if any of the following conditions occur:

- An Availability Zone outage
- The primary DB instance fails
- The DB instance's server type is changed
- The operating system of the DB instance is undergoing software patching
- A manual failover of the DB instance was initiated using **Reboot with failover**

There are several ways to determine if your Multi-AZ DB instance has failed over:

- DB event subscriptions can be setup to notify you via email or SMS that a failover has been initiated. For more information about events, see [Using Amazon RDS Event Notification \(p. 370\)](#)
- You can view your DB events by using the Amazon RDS console or API operations.
- You can view the current state of your Multi-AZ deployment by using the Amazon RDS console and API operations.

For information on how you can respond to failovers, reduce recovery time, and other best practices for Amazon RDS, see [Best Practices for Amazon RDS \(p. 94\)](#).

Related Topics

- [Multi-AZ Deployments for Microsoft SQL Server \(p. 681\)](#)
- [Licensing Oracle Multi-AZ Deployments \(p. 917\)](#)

Amazon RDS DB Instance Lifecycle

The lifecycle of an Amazon RDS DB instance includes creating, modifying, maintaining and upgrading, performing backups and restores, rebooting, and deleting the instance. This section provides information on and links to more about these processes.

Topics

- [Creating an Amazon RDS DB Instance \(p. 148\)](#)
- [Connecting to an Amazon RDS DB Instance \(p. 149\)](#)
- [Modifying an Amazon RDS DB Instance \(p. 150\)](#)
- [Maintaining a DB Instance \(p. 152\)](#)
- [Upgrading a DB Instance Engine Version \(p. 159\)](#)
- [Renaming a DB Instance \(p. 163\)](#)
- [Rebooting a DB Instance \(p. 166\)](#)
- [Stopping an Amazon RDS DB Instance Temporarily \(p. 169\)](#)
- [Starting an Amazon RDS DB Instance That Was Previously Stopped \(p. 172\)](#)
- [Deleting a DB Instance \(p. 174\)](#)

Creating an Amazon RDS DB Instance

The basic building block of Amazon RDS is the DB instance. You choose the engine-specific characteristics of the DB instance when you create it. You also choose the capacity of storage, CPU, memory, and so on of the AWS instance on which the database server runs.

To create an Amazon RDS DB instance, follow the instructions for your specific database engine:

- [Creating a DB Instance Running the MariaDB Database Engine \(p. 543\)](#)
- [Creating a DB Instance Running the Microsoft SQL Server Database Engine \(p. 616\)](#)
- [Creating a DB Instance Running the MySQL Database Engine \(p. 751\)](#)
- [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#)
- [Creating a DB Instance Running the PostgreSQL Database Engine \(p. 1257\)](#)

Connecting to an Amazon RDS DB Instance

After you create an Amazon RDS DB instance, you can use any standard SQL client application to connect to the DB instance. To connect to an Amazon RDS DB instance, follow the instructions for your specific database engine.

- [Connecting to a DB Instance Running the MariaDB Database Engine \(p. 555\)](#)
- [Connecting to a DB Instance Running the Microsoft SQL Server Database Engine \(p. 630\)](#)
- [Connecting to a DB Instance Running the MySQL Database Engine \(p. 764\)](#)
- [Connecting to a DB Instance Running the Oracle Database Engine \(p. 955\)](#)
- [Connecting to a DB Instance Running the PostgreSQL Database Engine \(p. 1270\)](#)

Modifying an Amazon RDS DB Instance

Most modifications to a DB instance can be applied immediately or deferred until the next maintenance window. Some modifications, such as parameter group changes, require that you manually reboot your DB instance for the change to take effect.

Important

Some modifications result in an outage because Amazon RDS must reboot your DB instance for the change to take effect. Review the impact to your database and applications before modifying your DB instance settings.

To modify an Amazon RDS DB instance, follow the instructions for your specific database engine.

- [Modifying a DB Instance Running the MariaDB Database Engine \(p. 564\)](#)
- [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#)
- [Modifying a DB Instance Running the MySQL Database Engine \(p. 802\)](#)
- [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#)
- [Modifying a DB Instance Running the PostgreSQL Database Engine \(p. 1279\)](#)

Using the Apply Immediately Parameter

When you modify a DB instance, you can apply the changes immediately. To apply changes immediately, you select the **Apply Immediately** option in the AWS Management Console, you use the `--apply-immediately` parameter when calling the AWS CLI, or you set the `ApplyImmediately` parameter to `true` when using the Amazon RDS API.

If you don't choose to apply changes immediately, the changes are put into the pending modifications queue. During the next maintenance window, any pending changes in the queue are applied. If you choose to apply changes immediately, your new changes and any changes in the pending modifications queue are applied.

Important

If any of the pending modifications require downtime, choosing apply immediately can cause unexpected downtime.

When you choose to apply a change immediately, any pending modifications are also applied immediately, instead of during the next maintenance window.

If you don't want a pending change to be applied in the next maintenance window, you can modify the DB instance to revert the change using the AWS CLI and specify the `--apply-immediately` option.

Changes to some database settings are applied immediately, even if you choose to defer your changes. To see how the different database settings interact with the apply immediately setting, see the settings for your specific database engine.

- [Settings for MariaDB DB Instances \(p. 565\)](#)
- [Settings for Microsoft SQL Server DB Instances \(p. 644\)](#)
- [Settings for MySQL DB Instances \(p. 803\)](#)
- [Settings for Oracle DB Instances \(p. 967\)](#)
- [Settings for PostgreSQL DB Instances \(p. 1280\)](#)

Related Topics

- [Renaming a DB Instance \(p. 163\)](#)
- [Rebooting a DB Instance \(p. 166\)](#)

- [Stopping an Amazon RDS DB Instance Temporarily \(p. 169\)](#)
- [modify-db-instance](#)
- [ModifyDBInstance](#)

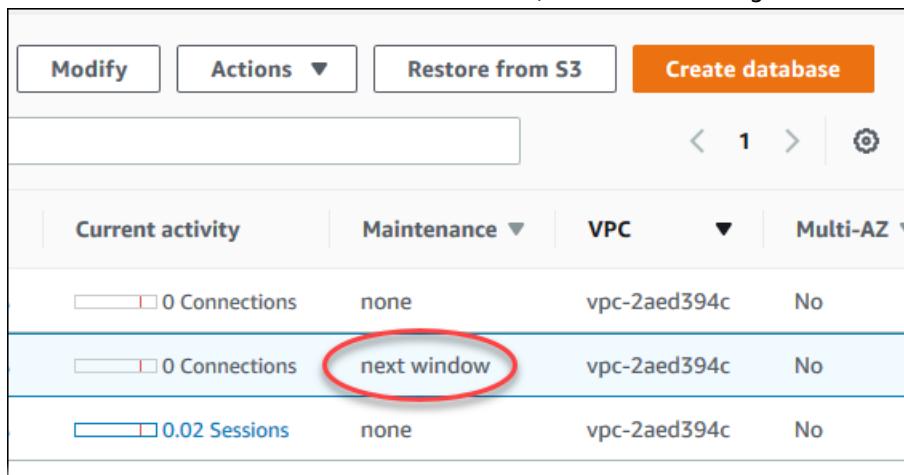
Maintaining a DB Instance

Periodically, Amazon RDS performs maintenance on Amazon RDS resources. Maintenance most often involves updates to the DB instance's underlying hardware, underlying operating system (OS), or database engine version. Updates to the operating system most often occur for security issues and should be done as soon as possible.

Some maintenance items require that Amazon RDS take your DB instance offline for a short time. Maintenance items that require a resource to be offline include required operating system or database patching. Required patching is automatically scheduled only for patches that are related to security and instance reliability. Such patching occurs infrequently (typically once every few months) and seldom requires more than a fraction of your maintenance window.

Deferred DB instance modifications that you have chosen not to apply immediately are applied during the maintenance window. For example, you may choose to change the DB instance class or parameter group during the maintenance window. For information about modifying a DB instance, see [Modifying an Amazon RDS DB Instance \(p. 150\)](#).

You can view whether a maintenance update is available for your DB instance by using the RDS console, the AWS CLI, or the Amazon RDS API. If an update is available, it is indicated in the **Maintenance** column for the DB instance on the Amazon RDS console, as shown following.



Current activity	Maintenance	VPC	Multi-AZ
0 Connections	none	vpc-2aed394c	No
0 Connections	next window	vpc-2aed394c	No
0.02 Sessions	none	vpc-2aed394c	No

If no maintenance update is available for a DB instance, the column value is **none** for it.

If a maintenance update is available for a DB instance, the following column values are possible:

- **required** – The maintenance action will be applied to the resource and can't be deferred.
- **available** – The maintenance action is available, but it will not be applied to the resource automatically. You can apply it manually.
- **next window** – The maintenance action will be applied to the resource during the next maintenance window.
- **In progress** – The maintenance action is in the process of being applied to the resource.

If an update is available, you can take one of the actions:

- If the maintenance value is **next window**, defer the maintenance items by choosing **Defer upgrade** from **Actions**. You can't defer a maintenance action if it has already started.
- Apply the maintenance items immediately.
- Schedule the maintenance items to start during your next maintenance window.

- Take no action.

Note

Certain OS updates are marked as **required**. If you defer a required update, you get a notice from Amazon RDS indicating when the update will be performed. Other updates are marked as **available**, and these you can defer indefinitely.

To take an action, choose the DB instance to show its details, then choose **Maintenance & backups**. The pending maintenance items appear.

The screenshot shows the 'Maintenance & backups' tab selected in the top navigation bar. Below it, a section titled 'Maintenance' shows the status of 'Auto minor version upgrade' as 'Enabled'. Under 'Pending maintenance', there is one item listed:

Description	Type	Status	Apply date
Automatic minor version upgrade to postgres 9.6.11	db-upgrade	next window	February 25th 2019, 3:28:00 am UTC-8 (local)

Buttons for 'Apply now' and 'Apply at next maintenance window' are present above the table, along with a search bar and pagination controls.

The maintenance window determines when pending operations start, but doesn't limit the total execution time of these operations. Maintenance operations aren't guaranteed to finish before the maintenance window ends, and can continue beyond the specified end time. For more information, see [The Amazon RDS Maintenance Window \(p. 156\)](#).

Applying Updates for a DB Instance

With Amazon RDS, you can choose when to apply maintenance operations. You can decide when Amazon RDS applies updates by using the RDS console, AWS Command Line Interface (AWS CLI), or RDS API.

Console

To manage an update for a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the DB instance that has a required update.
4. For **Actions**, choose one of the following:
 - **Upgrade now**
 - **Upgrade at next window**

Note

If you choose **Upgrade at next window** and later want to delay the update, you can choose **Defer upgrade**. You can't defer a maintenance action if it has already started.

AWS CLI

To apply a pending update to a DB instance, use the [apply-pending-maintenance-action](#) AWS CLI command.

Example

For Linux, OS X, or Unix:

```
aws rds apply-pending-maintenance-action \
--resource-identifier arn:aws:rds:us-west-2:001234567890:db:mysql-db \
--apply-action system-update \
--opt-in-type immediate
```

For Windows:

```
aws rds apply-pending-maintenance-action ^
--resource-identifier arn:aws:rds:us-west-2:001234567890:db:mysql-db ^
--apply-action system-update ^
--opt-in-type immediate
```

Note

You can defer a maintenance action by specifying `undo-opt-in` for `--opt-in-type`. You can't specify `undo-opt-in` for `--opt-in-type` if the maintenance action has already started.

To return a list of resources that have at least one pending update, use the [describe-pending-maintenance-actions](#) AWS CLI command.

Example

For Linux, OS X, or Unix:

```
aws rds describe-pending-maintenance-actions \
--resource-identifier arn:aws:rds:us-west-2:001234567890:db:mysql-db
```

For Windows:

```
aws rds describe-pending-maintenance-actions ^  
--resource-identifier arn:aws:rds:us-west-2:001234567890:db:mysql-db
```

You can also return a list of resources for a DB instance by specifying the `--filters` parameter of the `describe-pending-maintenance-actions` AWS CLI command. The format for the `--filters` command is `Name=filter-name,Value=resource-id,...`.

The following are the accepted values for the `Name` parameter of a filter:

- `db-instance-id` – Accepts a list of DB instance identifiers or Amazon Resource Names (ARNs). The returned list only includes pending maintenance actions for the DB instances identified by these identifiers or ARNs.
- `db-cluster-id` – Accepts a list of DB cluster identifiers or ARNs for Amazon Aurora. The returned list only includes pending maintenance actions for the DB clusters identified by these identifiers or ARNs.

For example, the following example returns the pending maintenance actions for the `sample-instance1` and `sample-instance2` DB instances.

Example

For Linux, OS X, or Unix:

```
aws rds describe-pending-maintenance-actions \  
--filters Name=db-instance-id,Values=sample-instance1,sample-instance2
```

For Windows:

```
aws rds describe-pending-maintenance-actions ^  
--filters Name=db-instance-id,Values=sample-instance1,sample-instance2
```

RDS API

To apply an update to a DB instance, call the Amazon RDS API [ApplyPendingMaintenanceAction](#) operation.

To return a list of resources that have at least one pending update, call the Amazon RDS API [DescribePendingMaintenanceActions](#) operation.

Maintenance for Multi-AZ Deployments

Running a DB instance as a Multi-AZ deployment can further reduce the impact of a maintenance event, because Amazon RDS applies operating system updates by following these steps:

1. Perform maintenance on the standby.
2. Promote the standby to primary.
3. Perform maintenance on the old primary, which becomes the new standby.

When you modify the database engine for your DB instance in a Multi-AZ deployment, then Amazon RDS upgrades both the primary and secondary DB instances at the same time. In this case, the database engine for the entire Multi-AZ deployment is shut down during the upgrade.

For more information on Multi-AZ deployments, see [High Availability \(Multi-AZ\) for Amazon RDS \(p. 144\)](#).

The Amazon RDS Maintenance Window

Every DB instance has a weekly maintenance window during which any system changes are applied. You can think of the maintenance window as an opportunity to control when modifications and software patching occur, in the event either are requested or required. If a maintenance event is scheduled for a given week, it is initiated during the 30-minute maintenance window you identify. Most maintenance events also complete during the 30-minute maintenance window, although larger maintenance events may take more than 30 minutes to complete.

The 30-minute maintenance window is selected at random from an 8-hour block of time per region. If you don't specify a preferred maintenance window when you create the DB instance, then Amazon RDS assigns a 30-minute maintenance window on a randomly selected day of the week.

RDS will consume some of the resources on your DB instance while maintenance is being applied. You might observe a minimal effect on performance. For a DB instance, on rare occasions, a Multi-AZ failover might be required for a maintenance update to complete.

Following, you can find the time blocks for each region from which default maintenance windows are assigned.

Region Name	Region	Time Block
US East (Ohio)	us-east-2	03:00–11:00 UTC
US East (N. Virginia)	us-east-1	03:00–11:00 UTC
US West (N. California)	us-west-1	06:00–14:00 UTC
US West (Oregon)	us-west-2	06:00–14:00 UTC
Asia Pacific (Hong Kong)	ap-east-1	06:00–14:00 UTC
Asia Pacific (Mumbai)	ap-south-1	17:30–01:30 UTC
Asia Pacific (Osaka-Local)	ap-northeast-3	22:00–23:59 UTC
Asia Pacific (Seoul)	ap-northeast-2	13:00–21:00 UTC
Asia Pacific (Singapore)	ap-southeast-1	14:00–22:00 UTC
Asia Pacific (Sydney)	ap-southeast-2	12:00–20:00 UTC
Asia Pacific (Tokyo)	ap-northeast-1	13:00–21:00 UTC
Canada (Central)	ca-central-1	03:00–11:00 UTC
China (Beijing)	cn-north-1	06:00–14:00 UTC
China (Ningxia)	cn-northwest-1	06:00–14:00 UTC
Europe (Frankfurt)	eu-central-1	23:00–07:00 UTC
Europe (Ireland)	eu-west-1	22:00–06:00 UTC
Europe (London)	eu-west-2	22:00–06:00 UTC
Europe (Paris)	eu-west-3	23:59–07:29 UTC
Europe (Stockholm)	eu-north-1	23:00–07:00 UTC

Region Name	Region	Time Block
Middle East (Bahrain)	me-south-1	06:00–14:00 UTC
South America (São Paulo)	sa-east-1	00:00–08:00 UTC
AWS GovCloud (US-West)	us-gov-west-1	06:00–14:00 UTC

Adjusting the Preferred DB Instance Maintenance Window

The maintenance window should fall at the time of lowest usage and thus might need modification from time to time. Your DB instance will only be unavailable during this time if the system changes, such as a change in DB instance class, are being applied and require an outage, and only for the minimum amount of time required to make the necessary changes.

In the following example, you adjust the preferred maintenance window for a DB instance.

For the purpose of this example, we assume that the DB instance named *mydbinstance* exists and has a preferred maintenance window of "Sun:05:00-Sun:06:00" UTC.

Console

To adjust the preferred maintenance window

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then select the DB instance that you want to modify.
3. Choose **Modify**. The **Modify DB Instance** page appears.
4. In the **Maintenance** section, update the maintenance window.

Note

The maintenance window and the backup window for the DB instance cannot overlap. If you enter a value for the maintenance window that overlaps the backup window, an error message appears.

5. Choose **Continue**.
On the confirmation page, review your changes.
6. To apply the changes to the maintenance window immediately, select **Apply immediately**.
7. Choose **Modify DB Instance** to save your changes.

Alternatively, choose **Back** to edit your changes, or choose **Cancel** to cancel your changes.

AWS CLI

To adjust the preferred maintenance window, use the AWS CLI `modify-db-instance` command with the following parameters:

- `--db-instance-identifier`
- `--preferred-maintenance-window`

Example

The following code example sets the maintenance window to Tuesdays from 4:00-4:30AM UTC.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier mydbinstance \
--preferred-maintenance-window Tue:04:00-Tue:04:30
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier mydbinstance ^
--preferred-maintenance-window Tue:04:00-Tue:04:30
```

RDS API

To adjust the preferred maintenance window, use the Amazon RDS API [ModifyDBInstance](#) operation with the following parameters:

- `DBInstanceIdentifier = mydbinstance`
- `PreferredMaintenanceWindow = Tue:04:00-Tue:04:30`

Example

The following code example sets the maintenance window to Tuesdays from 4:00-4:30AM UTC.

```
https://rds.us-west-2.amazonaws.com/
?Action=ModifyDBInstance
&DBInstanceIdentifier=mydbinstance
&PreferredMaintenanceWindow=Tue:04:00-Tue:04:30
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2014-09-01
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140425/us-east-1/rds/aws4_request
&X-Amz-Date=20140425T192732Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=1dc9dd716f4855e9bdf188c70f1cf9f6251b070b68b81103b59ec70c3e7854b3
```

Upgrading a DB Instance Engine Version

Amazon RDS provides newer versions of each supported database engine so you can keep your DB instance up-to-date. Newer versions can include bug fixes, security enhancements, and other improvements for the database engine. When Amazon RDS supports a new version of a database engine, you can choose how and when to upgrade your database DB instances.

There are two kinds of upgrades: major version upgrades and minor version upgrades. In general, a *major engine version upgrade* can introduce changes that are not compatible with existing applications. In contrast, a *minor version upgrade* includes only changes that are backward-compatible with existing applications.

The version numbering sequence is specific to each database engine. For example, Amazon RDS MySQL 5.7 and 8.0 are major engine versions and upgrading from any 5.7 version to any 8.0 version is a major version upgrade. Amazon RDS MySQL version 5.7.22 and 5.7.23 are minor versions and upgrading from 5.7.22 to 5.7.23 is a minor version upgrade.

For more information about major and minor version upgrades for a specific DB engine, see the following documentation for your DB engine:

- [Upgrading the MariaDB DB Engine \(p. 573\)](#)
- [Upgrading the Microsoft SQL Server DB Engine \(p. 653\)](#)
- [Upgrading the MySQL DB Engine \(p. 812\)](#)
- [Upgrading the Oracle DB Engine \(p. 979\)](#)
- [Upgrading the PostgreSQL DB Engine for Amazon RDS \(p. 1289\)](#)

For major version upgrades, you must manually modify the DB engine version through the AWS Management Console, AWS CLI, or RDS API. For minor version upgrades, you can manually modify the engine version, or you can choose to enable auto minor version upgrades.

Topics

- [Manually Upgrading the Engine Version \(p. 159\)](#)
- [Automatically Upgrading the Minor Engine Version \(p. 160\)](#)

Manually Upgrading the Engine Version

To manually upgrade the engine version of a DB instance, you can use the AWS Management Console, the AWS CLI, or the RDS API.

Console

To upgrade the engine version of a DB instance by using the console

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to upgrade.
3. Choose **Modify**. The **Modify DB Instance** page appears.
4. For **DB engine version**, choose the new version.
5. Choose **Continue** and check the summary of modifications.
6. To apply the changes immediately, choose **Apply immediately**. Choosing this option can cause an outage in some cases. For more information, see [Using the Apply Immediately Parameter \(p. 150\)](#).

7. On the confirmation page, review your changes. If they are correct, choose **Modify DB Instance** to save your changes.

Alternatively, choose **Back** to edit your changes, or choose **Cancel** to cancel your changes.

AWS CLI

To upgrade the engine version of a DB instance, use the CLI [modify-db-instance](#) command. Specify the following parameters:

- **--db-instance-identifier** – the name of the DB instance.
- **--engine-version** – the version number of the database engine to upgrade to.
For information about valid engine versions, use the AWS CLI [describe-db-engine-versions](#) command.
- **--allow-major-version-upgrade** – to upgrade the major version.
- **--no-apply-immediately** – to apply changes during the next maintenance window. To apply changes immediately, use **--apply-immediately**.

Example

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
  --db-instance-identifier mydbinstance \
  --engine-version new_version \
  --allow-major-version-upgrade \
  --no-apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
  --db-instance-identifier mydbinstance ^
  --engine-version new_version ^
  --allow-major-version-upgrade ^
  --no-apply-immediately
```

RDS API

To upgrade the engine version of a DB instance, use the [ModifyDBInstance](#) action. Specify the following parameters:

- **DBInstanceIdentifier** – the name of the DB instance, for example *mydbinstance*.
- **EngineVersion** – the version number of the database engine to upgrade to. For information about valid engine versions, use the [DescribeDBEngineVersions](#) operation.
- **AllowMajorVersionUpgrade** – whether to allow a major version upgrade. To do so, set the value to **true**.
- **ApplyImmediately** – whether to apply changes immediately or during the next maintenance window. To apply changes immediately, set the value to **true**. To apply changes during the next maintenance window, set the value to **false**.

Automatically Upgrading the Minor Engine Version

A *minor engine version* is an update to a DB engine version within a major engine version. For example, a major engine version might be 9.6 with the minor engine versions 9.6.11 and 9.6.12 within it.

If you want Amazon RDS to upgrade the DB engine version of a database automatically, you can enable auto minor version upgrades for the database. When a minor engine version is designated as the preferred minor engine version, each database that meets both of the following conditions is upgraded to the minor engine version automatically:

- The database is running a minor version of the DB engine that is lower than the preferred minor engine version.
- The database has auto minor version upgrade enabled.

You can control whether auto minor version upgrade is enabled for a DB instance when you perform the following tasks:

- [Creating a DB instance \(p. 148\)](#)
- [Modifying a DB instance \(p. 150\)](#)
- [Creating a Read Replica \(p. 184\)](#)
- [Restoring a DB instance from a snapshot \(p. 266\)](#)
- [Restoring a DB instance to a specific time \(p. 285\)](#)
- [Importing a DB instance from Amazon S3 \(p. 826\) \(for a MySQL backup on Amazon S3\)](#)

When you perform these tasks, you can control whether auto minor version upgrade is enabled for the DB instance in the following ways:

- Using the console, set the **Auto minor version upgrade** option.
- Using the AWS CLI, set the `--auto-minor-version-upgrade` | `--no-auto-minor-version-upgrade` option.
- Using the RDS API, set the `AutoMinorVersionUpgrade` parameter.

To determine whether a maintenance update, such as a DB engine version upgrade, is available for your DB instance, you can use the console, AWS CLI, or RDS API. You can also upgrade the DB engine version manually and adjust the maintenance window. For more information, see [Maintaining a DB Instance \(p. 152\)](#).

A PostgreSQL engine upgrade doesn't upgrade any PostgreSQL extensions. To update an extension after an engine upgrade, use the `ALTER EXTENSION UPDATE` command.

Note

If you are running the PostGIS extension in your Amazon RDS PostgreSQL DB instance, make sure that you follow the [PostGIS upgrade instructions](#) in the PostGIS documentation before you upgrade the extension.

To upgrade an extension, use the following command.

```
ALTER EXTENSION extension_name UPDATE TO 'new_version';
```

To list your currently installed extensions, use the PostgreSQL `pg_extension` catalog in the following command:

```
SELECT * FROM pg_extension;
```

To view a list of the specific extension versions that are available for your installation, use the PostgreSQL `pg_available_extension_versions` view in the following command:

```
SELECT * FROM pg_available_extension_versions;
```


Renaming a DB Instance

You can rename a DB instance by using the AWS Management Console, the AWS CLI `modify-db-instance` command, or the Amazon RDS API `ModifyDBInstance` action. Renaming a DB instance can have far-reaching effects; the following is a list of things you should know before you rename a DB instance.

- When you rename a DB instance, the endpoint for the DB instance changes, because the URL includes the name you assigned to the DB instance. You should always redirect traffic from the old URL to the new one.
- When you rename a DB instance, the old DNS name that was used by the DB instance is immediately deleted, although it could remain cached for a few minutes. The new DNS name for the renamed DB instance becomes effective in about 10 minutes. The renamed DB instance is not available until the new name becomes effective.
- You cannot use an existing DB instance name when renaming an instance.
- All read replicas associated with a DB instance remain associated with that instance after it is renamed. For example, suppose you have a DB instance that serves your production database and the instance has several associated read replicas. If you rename the DB instance and then replace it in the production environment with a DB snapshot, the DB instance that you renamed will still have the read replicas associated with it.
- Metrics and events associated with the name of a DB instance are maintained if you reuse a DB instance name. For example, if you promote a Read Replica and rename it to be the name of the previous master, the events and metrics associated with the master are associated with the renamed instance.
- DB instance tags remain with the DB instance, regardless of renaming.
- DB snapshots are retained for a renamed DB instance.

Renaming to Replace an Existing DB Instance

The most common reasons for renaming a DB instance are that you are promoting a Read Replica or you are restoring data from a DB snapshot or PITR. By renaming the database, you can replace the DB instance without having to change any application code that references the DB instance. In these cases, you would do the following:

1. Stop all traffic going to the master DB instance. This can involve redirecting traffic from accessing the databases on the DB instance or some other way you want to use to prevent traffic from accessing your databases on the DB instance.
2. Rename the master DB instance to a name that indicates it is no longer the master as described later in this topic.
3. Create a new master DB instance by restoring from a DB snapshot or by promoting a read replica, and then give the new instance the name of the previous master DB instance.
4. Associate any read replicas with the new master DB instance.

If you delete the old master DB instance, you are responsible for deleting any unwanted DB snapshots of the old master instance.

For information about promoting a Read Replica, see [Promoting a Read Replica to Be a Standalone DB Instance \(p. 186\)](#).

Console

To rename a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the DB instance that you want to rename.
4. Choose **Modify**.
5. In **Settings**, enter a new name for **DB instance identifier**.
6. Choose **Continue**.
7. To apply the changes immediately, choose **Apply immediately**. Choosing this option can cause an outage in some cases. For more information, see [Using the Apply Immediately Parameter \(p. 150\)](#).
8. On the confirmation page, review your changes. If they are correct, choose **Modify DB Instance** to save your changes.

Alternatively, choose **Back** to edit your changes, or choose **Cancel** to cancel your changes.

AWS CLI

To rename a DB instance, use the AWS CLI command `modify-db-instance`. Provide the current `--db-instance-identifier` value and `--new-db-instance-identifier` parameter with the new name of the DB instance.

Example

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
  --db-instance-identifier DBInstanceIdentifier \
  --new-db-instance-identifier NewDBInstanceIdentifier
```

For Windows:

```
aws rds modify-db-instance ^
  --db-instance-identifier DBInstanceIdentifier ^
  --new-db-instance-identifier NewDBInstanceIdentifier
```

RDS API

To rename a DB instance, call Amazon RDS API function `ModifyDBInstance` with the following parameters:

- `DBInstanceIdentifier` = existing name for the instance
- `NewDBInstanceIdentifier` = new name for the instance

```
https://rds.amazonaws.com/
?Action=ModifyDBInstance
&DBInstanceIdentifier=mydbinstance
&NewDBInstanceIdentifier=mynewdbinstanceidentifier
&Version=2012-01-15
&SignatureVersion=2
&SignatureMethod=HmacSHA256
```

```
&Timestamp=2012-01-20T22%3A06%3A23.624Z  
&AWSAccessKeyId=<AWS Access Key ID>  
&Signature=<Signature>
```

Rebooting a DB Instance

You might need to reboot your DB instance, usually for maintenance reasons. For example, if you make certain modifications, or if you change the DB parameter group associated with the DB instance , you must reboot the instance for the changes to take effect.

Note

If a DB instance isn't using the latest changes to its associated DB parameter group, the AWS Management Console shows the DB parameter group with a status of **pending-reboot**. The **pending-reboot** parameter groups status doesn't result in an automatic reboot during the next maintenance window. To apply the latest parameter changes to that DB instance, manually reboot the DB instance. For more information about parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#).

Rebooting a DB instance restarts the database engine service. Rebooting a DB instance results in a momentary outage, during which the DB instance status is set to *rebooting*.

If the Amazon RDS instance is configured for Multi-AZ, you can perform the reboot with a failover. An Amazon RDS event is created when the reboot is completed. If your DB instance is a Multi-AZ deployment, you can force a failover from one Availability Zone (AZ) to another when you reboot. When you force a failover of your DB instance, Amazon RDS automatically switches to a standby replica in another Availability Zone, and updates the DNS record for the DB instance to point to the standby DB instance. As a result, you need to clean up and re-establish any existing connections to your DB instance. Rebooting with failover is beneficial when you want to simulate a failure of a DB instance for testing, or restore operations to the original AZ after a failover occurs. For more information, see [High Availability \(Multi-AZ\) for Amazon RDS \(p. 144\)](#).

Note

When you force a failover from one Availability Zone to another when you reboot, the Availability Zone change might not be reflected in the AWS Management Console, and in calls to the AWS CLI and RDS API, for several minutes.

You can't reboot your DB instance if it is not in the available state. Your database can be unavailable for several reasons, such as an in-progress backup, a previously requested modification, or a maintenance-window action.

The time required to reboot your DB instance depends on the crash recovery process of your specific database engine. To improve the reboot time, we recommend that you reduce database activity as much as possible during the reboot process. Reducing database activity reduces rollback activity for in-transit transactions.

Console

To reboot a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to reboot.
3. For **Actions**, choose **Reboot**.

The **Reboot DB Instance** page appears.

4. (Optional) Choose **Reboot with failover?** to force a failover from one AZ to another.
5. Choose **Reboot** to reboot your DB instance.

Alternatively, choose **Cancel**.

AWS CLI

To reboot a DB instance by using the AWS CLI, call the `reboot-db-instance` command.

Example Simple Reboot

For Linux, OS X, or Unix:

```
aws rds reboot-db-instance \
--db-instance-identifier mydbinstance
```

For Windows:

```
aws rds reboot-db-instance ^
--db-instance-identifier mydbinstance
```

Example Reboot with Failover

To force a failover from one AZ to the other, use the `--force-failover` parameter.

For Linux, OS X, or Unix:

```
aws rds reboot-db-instance \
--db-instance-identifier mydbinstance \
--force-failover
```

For Windows:

```
aws rds reboot-db-instance ^
--db-instance-identifier mydbinstance ^
--force-failover
```

RDS API

To reboot a DB instance by using the Amazon RDS API, call the `RebootDBInstance` action.

Example Simple Reboot

```
https://rds.amazonaws.com/
?Action=RebootDBInstance
&DBInstanceIdentifier=mydbinstance
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20131016/us-west-1/rds/aws4_request
&X-Amz-Date=20131016T233051Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=087a8eb41cb1ab5f99e81575f23e73757ffc6a1e42d7d2b30b9cc0be988cff97
```

Example Reboot with Failover

To force a failover from one AZ to the other, set the `ForceFailover` parameter to `true`.

```
https://rds.amazonaws.com/
?Action=RebootDBInstance
&DBInstanceIdentifier=mydbinstance
&ForceFailover=true
&Version=2014-10-31
```

```
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20131016/us-west-1/rds/aws4_request
&X-Amz-Date=20131016T233051Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=087a8eb41cb1ab5f99e81575f23e73757ffc6a1e42d7d2b30b9cc0be988cff97
```

Stopping an Amazon RDS DB Instance Temporarily

If you use a DB instance intermittently, for temporary testing, or for a daily development activity, you can stop your Amazon RDS DB instance temporarily to save money. While your DB instance is stopped, you are charged for provisioned storage (including Provisioned IOPS) and backup storage (including manual snapshots and automated backups within your specified retention window), but not for DB instance hours. For more information, see [Billing FAQs](#).

Note

In some cases, a large amount of time is required to stop a DB instance. If you want to stop your DB instance and restart it immediately, you can reboot the DB instance. For information about rebooting a DB instance, see [Rebooting a DB Instance \(p. 166\)](#).

You can stop and start DB instances that are running the following engines:

- MariaDB
- Microsoft SQL Server
- MySQL
- Oracle
- PostgreSQL

Stopping and starting a DB instance is supported for all DB instance classes, and in all AWS Regions.

You can stop and start a DB instance whether it is configured for a single Availability Zone or for Multi-AZ, for database engines that support Multi-AZ deployments. You can't stop an Amazon RDS for SQL Server DB instance in a Multi-AZ configuration.

When you stop a DB instance, the DB instance performs a normal shutdown and stops running. The status of the DB instance changes to stopping and then stopped. Any storage volumes remain attached to the DB instance, and their data is kept. Any data stored in the RAM of the DB instance is deleted.

Stopping a DB instance removes pending actions, except for pending actions for the DB instance's option group or DB parameter group.

Important

You can stop a DB instance for up to seven days. If you don't manually start your DB instance after seven days, your DB instance is automatically started so that it doesn't fall behind any required maintenance updates.

Benefits

Stopping and starting a DB instance is faster than creating a DB snapshot, and then restoring the snapshot.

When you stop a DB instance it retains its ID, Domain Name Server (DNS) endpoint, parameter group, security group, and option group. When you start a DB instance, it has the same configuration as when you stopped it. In addition, if you stop a DB instance, Amazon RDS retains the Amazon Simple Storage Service (Amazon S3) transaction logs so you can do a point-in-time restore if necessary.

Limitations

The following are some limitations to stopping and starting a DB instance:

- You can't stop a DB instance that has a Read Replica, or that is a Read Replica.
- You can't stop an Amazon RDS for SQL Server DB instance in a Multi-AZ configuration.

- You can't modify a stopped DB instance.
- You can't delete an option group that is associated with a stopped DB instance.
- You can't delete a DB parameter group that is associated with a stopped DB instance.

Option and Parameter Group Considerations

You can't remove persistent options (including permanent options) from an option group if there are DB instances associated with that option group. This functionality is also true of any DB instance with a state of stopping, stopped, or starting.

You can change the option group or DB parameter group that is associated with a stopped DB instance, but the change does not occur until the next time you start the DB instance. If you chose to apply changes immediately, the change occurs when you start the DB instance. Otherwise the changes occurs during the next maintenance window after you start the DB instance.

VPC Considerations

When you stop a DB instance it retains its DNS endpoint. If you stop a DB instance that is not in an Amazon Virtual Private Cloud (Amazon VPC), Amazon RDS releases the IP addresses of the DB instance. If you stop a DB instance that is in a VPC, the DB instance retains its IP addresses.

Note

You should always connect to a DB instance using the DNS endpoint, not the IP address.

Console

To stop a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to stop.
3. For **Actions**, choose **Stop**.
4. (Optional) In the **Stop DB Instance** window, choose **Yes** for **Create Snapshot?** and enter the snapshot name for **Snapshot name**. Choose **Yes** if you want to create a snapshot of the DB instance before stopping it.
5. Choose **Yes, Stop Now** to stop the DB instance, or choose **Cancel** to cancel the operation.

AWS CLI

To stop a DB instance by using the AWS CLI, call the `stop-db-instance` command with the following parameters:

- `--db-instance-identifier` – the name of the DB instance.

Example

```
stop-db-instance --db-instance-identifier mydbinstance
```

RDS API

To stop a DB instance by using the Amazon RDS API, call the `StopDBInstance` operation with the following parameter:

- **DBInstanceIdentifier** – the name of the DB instance.

Starting an Amazon RDS DB Instance That Was Previously Stopped

You can stop your Amazon RDS DB instance temporarily to save money. After you stop your DB instance, you can restart it to begin using it again. For more details about stopping and starting DB instances, see [Stopping an Amazon RDS DB Instance Temporarily \(p. 169\)](#).

When you start a DB instance that you previously stopped, the DB instance retains the ID, Domain Name Server (DNS) endpoint, parameter group, security group, and option group. When you start a stopped instance, you are charged a full instance hour.

Console

To start a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to start.
3. For **Actions**, choose **Start**.

AWS CLI

To start a DB instance by using the AWS CLI, call the `start-db-instance` command with the following parameters:

- `--db-instance-identifier` – the name of the DB instance.

Example

```
start-db-instance --db-instance-identifier mydbinstance
```

RDS API

To start a DB instance by using the Amazon RDS API, call the `StartDBInstance` operation with the following parameters:

- `DBInstanceIdentifier` – the name of the DB instance.

Example

```
https://rds.amazonaws.com/
    ?Action=StartDBInstance
    &DBInstanceIdentifier=mydbinstance
    &SignatureMethod=HmacSHA256
    &SignatureVersion=4
    &Version=2014-10-31
    &X-Amz-Algorithm=AWS4-HMAC-SHA256
    &X-Amz-Credential=AKIADQKE4SARGYLE/20131016/us-west-1/rds/aws4_request
    &X-Amz-Date=20131016T233051Z
    &X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
    &X-Amz-Signature=087a8eb41cb1ab5f99e81575f23e73757ffc6a1e42d7d2b30b9cc0be988cff97
```

Related Topics

- [Deleting a DB Instance \(p. 174\)](#)
- [Rebooting a DB Instance \(p. 166\)](#)

Deleting a DB Instance

To delete a DB instance, you must do the following:

- Provide the name of the instance
- Enable or disable the option to take a final DB snapshot of the instance
- Enable or disable the option to retain automated backups

If the DB instance that you want to delete has a Read Replica, you should either promote the Read Replica or delete it. For more information, see [Promoting a Read Replica to Be a Standalone DB Instance \(p. 186\)](#).

Deletion Protection

You can only delete instances that don't have deletion protection enabled. When you create or modify a DB instance, you have the option to enable deletion protection so that users can't delete the DB instance. Deletion protection is disabled by default for you when you use AWS CLI and API commands. Deletion protection is enabled for you when you use the AWS Management Console to create a production DB instance. However, Amazon RDS enforces deletion protection when you use the console, the CLI, or the API to delete a DB instance. To delete a DB instance that has deletion protection enabled, first modify the instance and disable deletion protection. Enabling or disabling deletion protection doesn't cause an outage.

Creating a Final Snapshot and Retaining Automated Backups

When you delete a DB instance, you can choose whether to create a final snapshot of the DB instance. You can also choose to retain automated backups after the DB instance is deleted. To be able to restore the DB instance later, create a final snapshot, retain automated backups, or both.

When you delete a DB instance, you have the following choices:

- Create a final DB snapshot.

To be able to restore your deleted DB instance later with a final DB snapshot, create a final DB snapshot.

To delete a DB instance quickly, you can skip creating a final DB snapshot.

Important

If you skip the final DB snapshot, to restore your DB instance you need one of the following:

- Use an earlier manual snapshot of the DB instance to restore the DB instance to that DB snapshot's point in time.
- Retain automated backups. You can use those to restore your DB instance during your retention period, but not after your retention period has expired.

You can't create a final DB snapshot of your DB instance if it has the status `creating`, `failed`, `incompatible-restore`, or `incompatible-network`. For more information about DB instance statuses, see [DB Instance Status \(p. 129\)](#).

- Retain automated backups.

You can choose to retain automated backups when you delete a DB instance. These backups are still subject to the retention period of the DB instance and age out the same way systems snapshots do.

When you choose to retain automated backups, your automated backups are retained for a set period of time. This set retention period occurs regardless of whether you chose to create a final DB snapshot.

Your automated backups are retained for the retention period that was set on the DB instance at the time that you deleted it.

To delete a retained automated backup, follow the instructions in [Deleting Retained Automated Backups \(p. 260\)](#).

If you don't choose to retain automated backups, your automated backups are deleted and can't be recovered when you delete a DB instance. You typically don't need to retain automated backups if you create a final DB snapshot.

Note

Regardless of your choice, earlier manual DB snapshots aren't deleted.

Deleting a DB Instance by Using the Console, CLI, and API

You can delete a DB instance using the AWS Management Console, the AWS CLI, or the RDS API.

Console

To delete a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to delete.
3. For **Actions**, choose **Delete**.
4. To create a final DB snapshot for the DB instance, enable **Create final snapshot?**.
5. If you enabled **Create final snapshot?** in the previous step, for **Final snapshot name** enter the name of your final DB snapshot.
6. To retain automated backups, choose **Retain automated backups**.
7. Enter **delete me** in the box.
8. Choose **Delete**.

AWS CLI

To delete a DB instance by using the AWS CLI, call the `delete-db-instance` command with the following options:

- `--db-instance-identifier`
- `--final-db-snapshot-identifier` or `--skip-final-snapshot`

Example With a final snapshot and no retained automated backups

For Linux, OS X, or Unix:

```
aws rds delete-db-instance \
  --db-instance-identifier mydbinstance \
  --final-db-snapshot-identifier mydbinstancefinalsnapshot \
  --delete-automated-backups
```

For Windows:

```
aws rds delete-db-instance ^
  --db-instance-identifier mydbinstance ^
  --final-db-snapshot-identifier mydbinstancefinalsnapshot ^
```

```
--delete-automated-backups
```

Example With retained automated backups and no final snapshot

For Linux, OS X, or Unix:

```
aws rds delete-db-instance \
--db-instance-identifier mydbinstance \
--skip-final-snapshot \
--no-delete-automated-backups
```

For Windows:

```
aws rds delete-db-instance ^
--db-instance-identifier mydbinstance ^
--skip-final-snapshot ^
--no-delete-automated-backups
```

RDS API

To delete a DB instance by using the Amazon RDS API, call the [DeleteDBInstance](#) operation with the following parameters:

- `DBInstanceIdentifier`
- `FinalDBSnapshotIdentifier` or `SkipFinalSnapshot`

Example With a final snapshot and no retained automated backups

```
https://rds.amazonaws.com/
?Action=DeleteDBInstance
&DBInstanceIdentifier=mydbinstance
&FinalDBSnapshotIdentifier=mydbinstancefinalsnapshot
&DeleteAutomatedBackups=true
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140305/us-west-1/rds/aws4_request
&X-Amz-Date=20140305T185838Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=b441901545441d3c7a48f63b5b1522c5b2b37c137500c93c45e209d4b3a064a3
```

Example With retained automated backups and no final snapshot

```
https://rds.amazonaws.com/
?Action=DeleteDBInstance
&DBInstanceIdentifier=mydbinstance
&SkipFinalSnapshot=true
&DeleteAutomatedBackups=false
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140305/us-west-1/rds/aws4_request
&X-Amz-Date=20140305T185838Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=b441901545441d3c7a48f63b5b1522c5b2b37c137500c93c45e209d4b3a064a3
```

Tagging Amazon RDS Resources

You can use Amazon RDS tags to add metadata to your Amazon RDS resources. In addition, these tags can be used with IAM policies to manage access to Amazon RDS resources and to control what actions can be applied to the Amazon RDS resources. Finally, these tags can be used to track costs by grouping expenses for similarly tagged resources.

All Amazon RDS resources can be tagged

- DB instances
- DB clusters
- Read Replicas
- DB snapshots
- DB cluster snapshots
- Reserved DB instances
- Event subscriptions
- DB option groups
- DB parameter groups
- DB cluster parameter groups
- DB security groups
- DB subnet groups

For information on managing access to tagged resources with IAM policies, see [Identity and Access Management in Amazon RDS \(p. 445\)](#).

Overview of Amazon RDS Resource Tags

An Amazon RDS tag is a name-value pair that you define and associate with an Amazon RDS resource. The name is referred to as the key. Supplying a value for the key is optional. You can use tags to assign arbitrary information to an Amazon RDS resource. You can use a tag key, for example, to define a category, and the tag value might be an item in that category. For example, you might define a tag key of "project" and a tag value of "Salix," indicating that the Amazon RDS resource is assigned to the Salix project. You can also use tags to designate Amazon RDS resources as being used for test or production by using a key such as environment=test or environment=production. We recommend that you use a consistent set of tag keys to make it easier to track metadata associated with Amazon RDS resources.

Use tags to organize your AWS bill to reflect your own cost structure. To do this, sign up to get your AWS account bill with tag key values included. Then, to see the cost of combined resources, organize your billing information according to resources with the same tag key values. For example, you can tag several resources with a specific application name, and then organize your billing information to see the total cost of that application across several services. For more information, see [Cost Allocation and Tagging in About AWS Billing and Cost Management](#).

Each Amazon RDS resource has a tag set, which contains all the tags that are assigned to that Amazon RDS resource. A tag set can contain as many as 50 tags, or it can be empty. If you add a tag to an Amazon RDS resource that has the same key as an existing tag on resource, the new value overwrites the old value.

AWS does not apply any semantic meaning to your tags; tags are interpreted strictly as character strings. Amazon RDS can set tags on a DB instance or other Amazon RDS resources, depending on the settings that you use when you create the resource. For example, Amazon RDS might add a tag indicating that a DB instance is for production or for testing.

- The tag key is the required name of the tag. The string value can be from 1 to 128 Unicode characters in length and cannot be prefixed with "aws:" or "rds:". The string can contain only the set of Unicode letters, digits, white-space, '_', ':', ';', '/', '=', '+', '-', '@' (Java regex: "^([\u0000-\uFFFF]\u0000-\uFFFF]*\$").
 - The tag value is an optional string value of the tag. The string value can be from 1 to 256 Unicode characters in length and cannot be prefixed with "aws:". The string can contain only the set of Unicode letters, digits, white-space, '_', ':', ';', '/', '=', '+', '-', '@' (Java regex: "^([\u0000-\uFFFF]\u0000-\uFFFF]*\$").
- Values do not have to be unique in a tag set and can be null. For example, you can have a key-value pair in a tag set of project/Trinity and cost-center/Trinity.

Note

You can add a tag to a snapshot, however, your bill will not reflect this grouping.

You can use the AWS Management Console, the command line interface, or the Amazon RDS API to add, list, and delete tags on Amazon RDS resources. When using the command line interface or the Amazon RDS API, you must provide the Amazon Resource Name (ARN) for the Amazon RDS resource you want to work with. For more information about constructing an ARN, see [Constructing an ARN for Amazon RDS \(p. 224\)](#).

Tags are cached for authorization purposes. Because of this, additions and updates to tags on Amazon RDS resources can take several minutes before they are available.

Copying Tags

When you create or restore a DB instance, you can specify that the tags from the DB instance are copied to snapshots of the DB instance. Copying tags ensures that the metadata for the DB snapshots matches that of the source DB instance and any access policies for the DB snapshot also match those of the source DB instance. Tags are not copied by default.

You can specify that tags are copied to DB snapshots for the following actions:

- Creating a DB instance.
- Restoring a DB instance.
- Creating a Read Replica.
- Copying a DB snapshot.

Note

If you include a value for the `--tag-key` parameter of the [create-db-snapshot](#) AWS CLI command (or supply at least one tag to the [CreateDBSnapshot](#) API operation) then RDS doesn't copy tags from the source DB instance to the new DB snapshot. This functionality applies even if the source DB instance has the `--copy-tags-to-snapshot` (`CopyTagsToSnapshot`) option enabled. If you take this approach, you can create a copy of a DB instance from a DB snapshot and avoid adding tags that don't apply to the new DB instance. Once you have created your DB snapshot using the AWS CLI `create-db-snapshot` command (or the [CreateDBSnapshot](#) Amazon RDS API operation) you can then add tags as described later in this topic.

Console

The process to tag an Amazon RDS resource is similar for all resources. The following procedure shows how to tag an Amazon RDS DB instance.

To add a tag to a DB instance

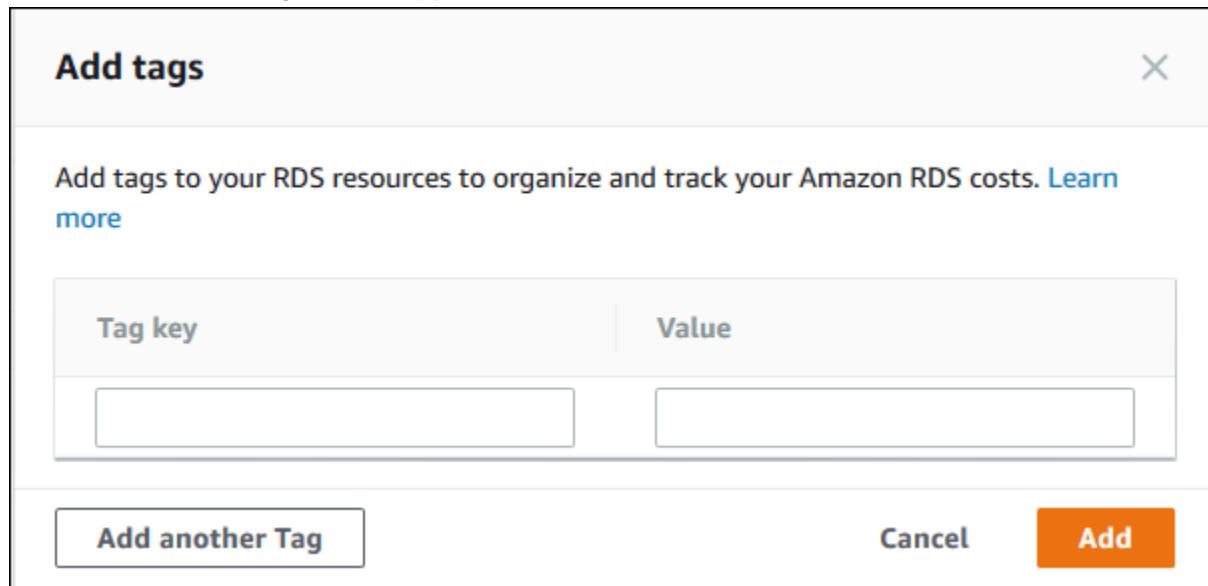
1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.

2. In the navigation pane, choose **Databases**.

Note

To filter the list of DB instances in the **Databases** pane, enter a text string for **Filter databases**. Only DB instances that contain the string appear.

3. Choose the name of the DB instance that you want to tag to show its details.
4. In the details section, scroll down to the **Tags** section.
5. Choose **Add**. The **Add tags** window appears.



6. Enter a value for **Tag key** and **Value**.
7. To add another tag, you can choose **Add another Tag** and enter a value for its **Tag key** and **Value**.
Repeat this step as many times as necessary.
8. Choose **Add**.

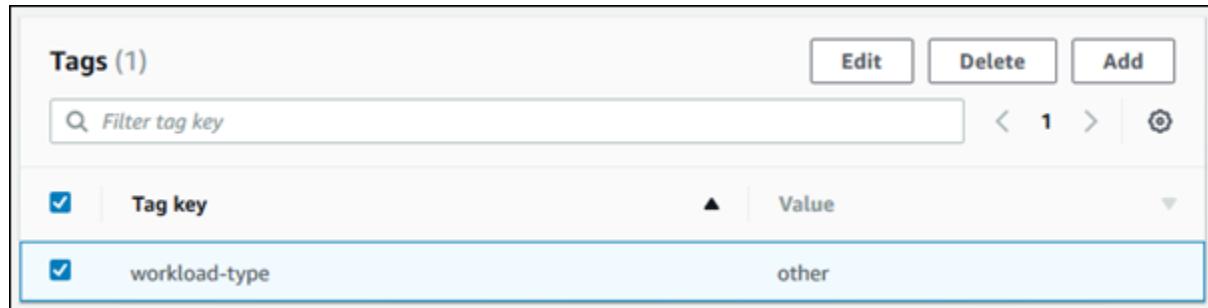
To delete a tag from a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.

Note

To filter the list of DB instances in the **Databases** pane, enter a text string in the **Filter databases** box. Only DB instances that contain the string appear.

3. Choose the name of the DB instance to show its details.
4. In the details section, scroll down to the **Tags** section.
5. Choose the tag you want to delete.



- Choose **Delete**, and then choose **Delete** in the **Delete tags** window.

AWS CLI

You can add, list, or remove tags for a DB instance using the AWS CLI.

- To add one or more tags to an Amazon RDS resource, use the AWS CLI command [add-tags-to-resource](#).
- To list the tags on an Amazon RDS resource, use the AWS CLI command [list-tags-for-resource](#).
- To remove one or more tags from an Amazon RDS resource, use the AWS CLI command [remove-tags-from-resource](#).

To learn more about how to construct the required ARN, see [Constructing an ARN for Amazon RDS \(p. 224\)](#).

RDS API

You can add, list, or remove tags for a DB instance using the Amazon RDS API.

- To add a tag to an Amazon RDS resource, use the [AddTagsToResource](#) operation.
- To list tags that are assigned to an Amazon RDS resource, use the [ListTagsForResource](#).
- To remove tags from an Amazon RDS resource, use the [RemoveTagsFromResource](#) operation.

To learn more about how to construct the required ARN, see [Constructing an ARN for Amazon RDS \(p. 224\)](#).

When working with XML using the Amazon RDS API, tags use the following schema:

```
<Tagging>
  <TagSet>
    <Tag>
      <Key>Project</Key>
      <Value>Trinity</Value>
    </Tag>
    <Tag>
      <Key>User</Key>
      <Value>Jones</Value>
    </Tag>
  </TagSet>
</Tagging>
```

The following table provides a list of the allowed XML tags and their characteristics. Values for Key and Value are case-dependent. For example, project=Trinity and PROJECT=Trinity are two distinct tags.

Tagging Element	Description
TagSet	A tag set is a container for all tags assigned to an Amazon RDS resource. There can be only one tag set per resource. You work with a TagSet only through the Amazon RDS API.
Tag	A tag is a user-defined key-value pair. There can be from 1 to 50 tags in a tag set.
Key	<p>A key is the required name of the tag. The string value can be from 1 to 128 Unicode characters in length and cannot be prefixed with "rds:" or "aws:". The string can only contain only the set of Unicode letters, digits, white-space, '_', '.', '/', '=', '+', '-' (Java regex: "<code>^([\\p{L}\\p{Z}\\p{N}_.:/=-\\-]*\$</code>").</p> <p>Keys must be unique to a tag set. For example, you cannot have a key-pair in a tag set with the key the same but with different values, such as project/Trinity and project/Xanadu.</p>
Value	<p>A value is the optional value of the tag. The string value can be from 1 to 256 Unicode characters in length and cannot be prefixed with "rds:" or "aws:". The string can only contain only the set of Unicode letters, digits, white-space, '_', '.', '/', '=', '+', '-' (Java regex: "<code>^([\\p{L}\\p{Z}\\p{N}_.:/=-\\-]*\$</code>").</p> <p>Values do not have to be unique in a tag set and can be null. For example, you can have a key-value pair in a tag set of project/Trinity and cost-center/Trinity.</p>

Working with Read Replicas

Amazon RDS uses the MariaDB, MySQL, Oracle, and PostgreSQL DB engines' built-in replication functionality to create a special type of DB instance called a Read Replica from a source DB instance. Updates made to the source DB instance are asynchronously copied to the Read Replica. You can reduce the load on your source DB instance by routing read queries from your applications to the Read Replica. Using Read Replicas, you can elastically scale out beyond the capacity constraints of a single DB instance for read-heavy database workloads.

Note

The information following applies to creating Amazon RDS Read Replicas either in the same AWS Region as the source DB instance, or in a separate AWS Region. The information following doesn't apply to setting up replication with an instance that is running on an Amazon EC2 instance or that is on-premises.

When you create a Read Replica, you first specify an existing DB instance as the source. Then Amazon RDS takes a snapshot of the source instance and creates a read-only instance from the snapshot. Amazon RDS then uses the asynchronous replication method for the DB engine to update the Read Replica whenever there is a change to the source DB instance. The Read Replica operates as a DB instance that allows only read-only connections. Applications connect to a Read Replica the same way they do to any DB instance. Amazon RDS replicates all databases in the source DB instance.

In some cases, a Read Replica resides in a different AWS Region than its source DB instance. In these cases, Amazon RDS sets up a secure communications channel between the source DB instance and the Read Replica. Amazon RDS establishes any AWS security configurations needed to enable the secure channel, such as adding security group entries.

Read Replicas are supported by the MariaDB, MySQL, Oracle, and PostgreSQL engines. In this section, you can find general information about using Read Replicas with all of these engines. For information about using Read Replicas with a specific engine, see the following sections:

- [Working with MySQL Read Replicas \(p. 852\)](#)
- [Working with MariaDB Read Replicas \(p. 579\)](#)
- [Working with Oracle Read Replicas for Amazon RDS \(p. 998\)](#)
- [Working with PostgreSQL Read Replicas \(p. 1298\)](#)

Overview of Amazon RDS Read Replicas

Deploying one or more Read Replicas for a given source DB instance might make sense in a variety of scenarios, including the following:

- Scaling beyond the compute or I/O capacity of a single DB instance for read-heavy database workloads. You can direct this excess read traffic to one or more Read Replicas.
- Serving read traffic while the source DB instance is unavailable. In some cases, your source DB instance might not be able to take I/O requests, for example due to I/O suspension for backups or scheduled maintenance. In these cases, you can direct read traffic to your Read Replicas. For this use case, keep in mind that the data on the Read Replica might be "stale" because the source DB instance is unavailable.
- Business reporting or data warehousing scenarios where you might want business reporting queries to run against a Read Replica, rather than your primary, production DB instance.
- Implementing disaster recovery. You can promote a Read Replica to a standalone instance as a disaster recovery solution if the source DB instance fails.

By default, a Read Replica is created with the same storage type as the source DB instance. However, you can create a Read Replica that has a different storage type from the source DB instance based on the options listed in the following table.

Source DB Instance Storage Type	Source DB Instance Storage Allocation	Read Replica Storage Type Options
PIOPS	100 GiB–32 TiB	PIOPS, GP2, Standard
GP2	100 GiB–32 TiB	PIOPS, GP2, Standard
GP2	Less than 100 GiB	GP2, Standard
Standard	100 GiB 6 TiB	PIOPS, GP2, Standard
Standard	Less than 100 GiB	GP2, Standard

Amazon RDS doesn't support circular replication. You can't configure a DB instance to serve as a replication source for an existing DB instance. You can only create a new Read Replica from an existing DB instance. For example, if `MyDBInstance` replicates to `ReadReplica1`, you can't configure `ReadReplica1` to replicate back to `MyDBInstance`. For MariaDB, MySQL, and PostgreSQL, you can create a Read Replica from an existing Read Replica. For example, from `ReadReplica1`, you can create a new Read Replica, such as `ReadReplica2`. For Oracle, you can't create a Read Replica from an existing Read Replica.

Differences Between Read Replicas for Different DB Engines

Because Amazon RDS DB engines implement replication differently, there are several significant differences you should know about, as shown in the following table.

Feature or Behavior	MySQL and MariaDB	Oracle	PostgreSQL
What is the replication method?	Logical replication.	Physical replication.	Physical replication.
How are transaction logs purged?	RDS MySQL and RDS MariaDB keep any binary logs that haven't been applied.	If a source DB instance has no cross-region Read Replicas, Amazon RDS for Oracle keeps a minimum of two hours of transaction logs on the source DB instance. Logs are purged from the source DB instance after two hours or after the archive log retention hours setting has passed, whichever is longer. Logs are purged from the Read Replica after the archive log retention hours setting has passed only if they have been successfully applied to the database.	PostgreSQL has the parameter <code>wal_keep_segments</code> that dictates how many write ahead log (WAL) files are kept to provide data to the Read Replicas. The parameter value specifies the number of logs to keep.

Feature or Behavior	MySQL and MariaDB	Oracle	PostgreSQL
		<p>In some cases, a source DB instance might have one or more cross-region Read Replicas. If so, Amazon RDS for Oracle keeps the transaction logs on the source DB instance until they have been transmitted and applied to all cross-region Read Replicas.</p> <p>For information about setting archivelog retention hours, see Retaining Archived Redo Logs (p. 1110).</p>	
Can a replica be made writable?	Yes. You can enable the MySQL or MariaDB Read Replica to be writable.	No. An Oracle Read Replica is a physical copy, and Oracle doesn't allow for writes in a Read Replica. You can promote the Read Replica to make it writable. The promoted Read Replica has the replicated data to the point when the request was made to promote it.	No. A PostgreSQL Read Replica is a physical copy, and PostgreSQL doesn't allow for a Read Replica to be made writable.
Can backups be performed on the replica?	Yes. You can enable automatic backups on a MySQL or MariaDB Read Replica.	No. You can't create manual snapshots of Amazon RDS for Oracle Read Replicas or enable automatic backups for them.	Yes, you can create a manual snapshot of a PostgreSQL Read Replica, but you can't enable automatic backups.
Can you use parallel replication?	Yes. MySQL version 5.6 and later and all supported MariaDB versions allow for parallel replication threads.	Yes. Redo log data is always transmitted in parallel from the source database to all of its Read Replicas.	No. PostgreSQL has a single process handling replication.

Creating a Read Replica

You can create a Read Replica from an existing MySQL, MariaDB, Oracle, or PostgreSQL DB instance using the AWS Management Console, AWS CLI, or AWS API. You create a Read Replica by specifying the `SourceDBInstanceIdentifier`, which is the DB instance identifier of the source DB instance from which you wish to replicate.

When you create a Read Replica, Amazon RDS takes a DB snapshot of your source DB instance and begins replication. As a result, you experience a brief I/O suspension on your source DB instance while the DB snapshot occurs. The I/O suspension typically lasts about one minute. You can avoid the I/O suspension if the source DB instance is a Multi-AZ deployment, because in that case the snapshot is taken from the

secondary DB instance. An active, long-running transaction can slow the process of creating the Read Replica. We recommend that you wait for long-running transactions to complete before creating a Read Replica. If you create multiple Read Replicas in parallel from the same source DB instance, Amazon RDS takes only one snapshot at the start of the first create action.

When creating a Read Replica, there are a few things to consider. First, you must enable automatic backups on the source DB instance by setting the backup retention period to a value other than 0. This requirement also applies to a Read Replica that is the source DB instance for another Read Replica. For MySQL DB instances, automatic backups are supported only for Read Replicas running MySQL 5.6 and later, but not for MySQL versions 5.5. To enable automatic backups on an Amazon RDS MySQL version 5.6 and later Read Replica, first create the Read Replica, then modify the Read Replica to enable automatic backups.

Note

Within an AWS Region, all Read Replicas must be created in the same Amazon VPC as the source DB instance, even if VPC peering is configured in the AWS Region.

Console

To create a Read Replica from a source MySQL, MariaDB, Oracle, or PostgreSQL DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the MySQL, MariaDB, Oracle, or PostgreSQL DB instance that you want to use as the source for a Read Replica.
4. For **Actions**, choose **Create read replica**.
5. Choose the instance specifications that you want to use. We recommend that you use the same DB instance class and storage type as the source DB instance for the Read Replica. For **Multi-AZ deployment**, choose **Yes** to create a standby of your replica in another Availability Zone for failover support for the replica. Creating your Read Replica as a Multi-AZ DB instance is independent of whether the source database is a Multi-AZ DB instance.
6. Choose the settings that you want to use. For **DB instance identifier**, enter a name for the Read Replica. Adjust other settings as needed.
7. Choose the other settings that you want to use.
8. Choose **Create read replica**.

AWS CLI

To create a Read Replica from a source MySQL, MariaDB, Oracle, or PostgreSQL DB instance, use the AWS CLI command `create-db-instance-read-replica`.

Example

For Linux, OS X, or Unix:

```
aws rds create-db-instance-read-replica \
--db-instance-identifier myreadreplica \
--source-db-instance-identifier mydbinstance
```

For Windows:

```
aws rds create-db-instance-read-replica ^
--db-instance-identifier myreadreplica ^
--source-db-instance-identifier mydbinstance
```

RDS API

To create a Read Replica from a source MySQL, MariaDB, Oracle, or PostgreSQL DB instance, call the Amazon RDS API function [CreateDBInstanceReadReplica](#).

```
https://rds.amazonaws.com/  
?Action=CreateDBInstanceReadReplica  
&DBInstanceIdentifier=myreadreplica  
&SourceDBInstanceIdentifier=mydbinstance  
&Version=2012-01-15  
&SignatureVersion=2  
&SignatureMethod=HmacSHA256  
&Timestamp=2012-01-20T22%3A06%3A23.624Z  
&AWSAccessKeyId=<AWS Access Key ID>  
&Signature=<Signature>
```

Promoting a Read Replica to Be a Standalone DB Instance

You can promote a MySQL, MariaDB, Oracle, or PostgreSQL Read Replica into a standalone DB instance. When you promote a Read Replica, the DB instance is rebooted before it becomes available.

There are several reasons you might want to promote a Read Replica to a standalone DB instance:

- **Performing DDL operations (MySQL and MariaDB only)** – DDL operations, such as creating or rebuilding indexes, can take time and impose a significant performance penalty on your DB instance. You can perform these operations on a MySQL or MariaDB Read Replica once the Read Replica is in sync with its source DB instance. Then you can promote the Read Replica and direct your applications to use the promoted instance.
- **Sharding** – Sharding embodies the "share-nothing" architecture and essentially involves breaking a large database into several smaller databases. One common way to split a database is splitting tables that are not joined in the same query onto different hosts. Another method is duplicating a table across multiple hosts and then using a hashing algorithm to determine which host receives a given update. You can create Read Replicas corresponding to each of your shards (smaller databases) and promote them when you decide to convert them into standalone shards. You can then carve out the key space (if you are splitting rows) or distribution of tables for each of the shards depending on your requirements.
- **Implementing failure recovery** – You can use Read Replica promotion as a data recovery scheme if the source DB instance fails. This approach complements synchronous replication, automatic failure detection, and failover.

If you are aware of the ramifications and limitations of asynchronous replication and you still want to use Read Replica promotion for data recovery, you can do so. To do this, first create a Read Replica and then monitor the source DB instance for failures. In the event of a failure, do the following:

1. Promote the Read Replica.
2. Direct database traffic to the promoted DB instance.
3. Create a replacement Read Replica with the promoted DB instance as its source.

When you promote a Read Replica, the new DB instance that is created retains the backup retention period, the backup window, the option group, and the parameter group of the former Read Replica source. The promotion process can take several minutes or longer to complete, depending on the size of the Read Replica. Once you promote the Read Replica to a new DB instance, it's just like any other DB instance. For example, you can create Read Replicas from the new DB instance and perform point-in-time restore operations. Because the promoted DB instance is no longer a Read Replica, you can't use it as a

replication target. If a source DB instance has several Read Replicas, promoting one of the Read Replicas to a DB instance has no effect on the other replicas.

Backup duration is a function of the number of changes to the database since the previous backup. If you plan to promote a Read Replica to a standalone instance, we recommend that you enable backups and complete at least one backup prior to promotion. In addition, a Read Replica cannot be promoted to a standalone instance when it is in the backing-up status. If you have enabled backups on your Read Replica, configure the automated backup window so that daily backups do not interfere with Read Replica promotion.

The following steps show the general process for promoting a Read Replica to a DB instance:

1. Stop any transactions from being written to the Read Replica source DB instance, and then wait for all updates to be made to the Read Replica. Database updates occur on the Read Replica after they have occurred on the source DB instance, and this replication lag can vary significantly. Use the [Replica Lag](#) metric to determine when all updates have been made to the Read Replica.
2. For MySQL and MariaDB only: If you need to make changes to the MySQL or MariaDB Read Replica, you must set the `read_only` parameter to `0` in the DB parameter group for the Read Replica. You can then perform all needed DDL operations, such as creating indexes, on the Read Replica. Actions taken on the Read Replica don't affect the performance of the source DB instance.
3. Promote the Read Replica by using the **Promote Read Replica** option on the Amazon RDS console, the AWS CLI command [promote-read-replica](#), or the [PromoteReadReplica](#) Amazon RDS API operation.

Note

The promotion process takes a few minutes to complete. When you promote a Read Replica, replication is stopped and the Read Replica is rebooted. When the reboot is complete, the Read Replica is available as a new DB instance.

4. (Optional) Modify the new DB instance to be a Multi-AZ deployment. For more information, see [Modifying an Amazon RDS DB Instance \(p. 150\)](#) and [High Availability \(Multi-AZ\) for Amazon RDS \(p. 144\)](#).

Console

To promote a Read Replica to a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the Amazon RDS console, choose **Databases**.
The **Databases** pane appears. Each Read Replica shows **Replica** in the **Role** column.
3. Choose the Read Replica that you want to promote.
4. For **Actions**, choose **Promote read replica**.
5. On the **Promote Read Replica** page, enter the backup retention period and the backup window for the new promoted DB instance.
6. When the settings are as you want them, choose **Continue**.
7. On the acknowledgment page, choose **Promote Read Replica**.

AWS CLI

To promote a Read Replica to a DB instance, use the AWS CLI [promote-read-replica](#) command.

Example

For Linux, OS X, or Unix:

```
aws rds promote-read-replica \
--db-instance-identifier myreadreplica
```

For Windows:

```
aws rds promote-read-replica ^
--db-instance-identifier myreadreplica
```

RDS API

To promote a Read Replica to a DB instance, call [PromoteReadReplica](#).

```
https://rds.amazonaws.com/
?Action=PromoteReadReplica
&DBInstanceIdentifier=myreadreplica
&Version=2012-01-15
&SignatureVersion=2
&SignatureMethod=HmacSHA256
&Timestamp=2012-01-20T22%3A06%3A23.624Z
&AWSAccessKeyId=<AWS Access Key ID>
&Signature=<Signature>
```

Creating a Read Replica in a Different AWS Region

With Amazon RDS, you can create a MariaDB, MySQL, Oracle, or PostgreSQL Read Replica in a different AWS Region than the source DB instance. You create a Read Replica to do the following:

- Improve your disaster recovery capabilities.
- Scale read operations into an AWS Region closer to your users.
- Make it easier to migrate from a data center in one AWS Region to a data center in another AWS Region.

Creating a Read Replica in a different AWS Region than the source instance is similar to creating a replica in the same AWS Region. To create a Read Replica across AWS Regions, you can use the AWS Management Console, run the [create-db-instance-read-replica](#) command, or call the [CreateDBInstanceReadReplica](#) API operation.

To create an encrypted Read Replica in a different AWS Region than the source DB instance, the source DB instance must be encrypted.

Following, you can find information on how to create a Read Replica from a source MariaDB, MySQL, Oracle, or PostgreSQL DB instance in a different AWS Region.

Console

You can create a Read Replica across AWS Regions using the AWS Management Console.

To create a Read Replica across AWS Regions with the console

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the MariaDB, MySQL, Oracle, or PostgreSQL DB instance that you want to use as the source for a Read Replica. For **Actions**, choose **Create read replica**. To create an encrypted Read Replica, the source DB instance must be encrypted. To learn more about encrypting the source DB instance, see [Encrypting Amazon RDS Resources \(p. 429\)](#).

4. Choose the instance specifications you want to use. We recommend that you use the same DB instance class and storage type for the Read Replica.
5. Choose the other settings you want to use:
 - For **DB instance identifier**, enter a name for the Read Replica.
 - In the **Network & Security** section, choose a value for **Destination region** and **Destination DB subnet group**.
 - To create an encrypted Read Replica in another AWS Region, choose **Enable Encryption**, and then choose the **Master key**. For the **Master key**, choose the AWS Key Management Service (AWS KMS) key identifier of the destination AWS Region.
 - Choose the other settings that you want to use.
6. Choose **Create read replica**.

AWS CLI

To create a Read Replica from a source MySQL, MariaDB, Oracle, or PostgreSQL DB instance in a different AWS Region, you can use the [create-db-instance-read-replica](#) command. In this case, you use [create-db-instance-read-replica](#) from the AWS Region where you want the Read Replica and specify the Amazon Resource Name (ARN) for the source DB instance. An ARN uniquely identifies a resource created in Amazon Web Services.

For example, if your source DB instance is in the US East (N. Virginia) Region, the ARN looks similar to the following.

```
arn:aws:rds:us-east-1:123456789012:db:my-mysql-instance
```

For information about ARNs, see [Working with Amazon Resource Names \(ARNs\) in Amazon RDS \(p. 224\)](#).

To create an encrypted Read Replica in a different AWS Region than the source DB instance, you can use the AWS CLI [create-db-instance-read-replica](#) command from the destination AWS Region. The following parameters are used to create an encrypted Read Replica in another AWS Region:

- **--source-region** — The AWS Region that the encrypted Read Replica is created in. If **source-region** is not specified, you must specify a **pre-signed-url** value. A **pre-signed-url** is an URL that contains a Signature Version 4 signed request for the **CreateDBInstanceReadReplica** operation that is called in the source AWS Region that the Read Replica is created from. To learn more about **pre-signed-url**, see [CreateDBInstanceReadReplica](#).
- **--source-db-instance-identifier** — The DB instance identifier for the encrypted Read Replica that is created. This identifier must be in the ARN format for the source AWS Region. The AWS Region specified in **source-db-instance-identifier** must match the AWS Region specified as **source-region**.
- **--db-instance-identifier** — The identifier for the encrypted Read Replica in the destination AWS Region.
- **--kms-key-id** — The AWS KMS key identifier for the key to use to encrypt the Read Replica in the destination AWS Region.

The following code creates a Read Replica in the **us-west-2** Region.

Example

For Linux, OS X, or Unix:

```
aws rds create-db-instance-read-replica \
--db-instance-identifier DBInstanceIdentifier \
```

```
--region us-west-2 \
--source-db-instance-identifier arn:aws:rds:us-east-1:123456789012:db:my-mysql-instance
```

For Windows:

```
aws rds create-db-instance-read-replica ^
--db-instance-identifier DBInstanceIdentifier ^
--region us-west-2 ^
--source-db-instance-identifier arn:aws:rds:us-east-1:123456789012:db:my-mysql-instance
```

The following code creates a Read Replica in a different AWS Region than the source DB instance. The AWS Region where you call the `create-db-instance-read-replica` command is the destination AWS Region for the encrypted Read Replica.

Example

For Linux, OS X, or Unix:

```
aws rds create-db-instance-read-replica \
--db-instance-identifier DBInstanceIdentifier \
--region us-west-2 \
--source-db-instance-identifier arn:aws:rds:us-east-1:123456789012:db:my-mysql-instance \
--source-region us-east-1 \
--kms-key-id my-us-east-1-key
```

For Windows:

```
aws rds create-db-instance-read-replica ^
--db-instance-identifier DBInstanceIdentifier ^
--region us-west-2 ^
--source-db-instance-identifier arn:aws:rds:us-east-1:123456789012:db:my-mysql-instance ^
--source-region us-east-1 ^
--kms-key-id my-us-east-1-key
```

RDS API

To create a Read Replica from a source MySQL, MariaDB, Oracle, or PostgreSQL DB instance in a different AWS Region, you can call the Amazon RDS API function [CreateDBInstanceReadReplica](#). In this case, you call [CreateDBInstanceReadReplica](#) from the AWS Region where you want the Read Replica and specify the Amazon Resource Name (ARN) for the source DB instance. An ARN uniquely identifies a resource created in Amazon Web Services.

To create an encrypted Read Replica in a different AWS Region than the source DB instance, you can use the Amazon RDS API [CreateDBInstanceReadReplica](#) action from the destination AWS Region. To create an encrypted Read Replica in another AWS Region, you must specify a value for `PreSignedURL`. `PreSignedURL` should contain a request for the [CreateDBInstanceReadReplica](#) action to call in the source AWS Region where the Read Replica is created in. To learn more about `PreSignedUrl`, see [CreateDBInstanceReadReplica](#).

For example, if your source DB instance is in the US East (N. Virginia) Region, the ARN looks similar to the following.

```
arn:aws:rds:us-east-1:123456789012:db:my-mysql-instance
```

For information about ARNs, see [Working with Amazon Resource Names \(ARNs\) in Amazon RDS \(p. 224\)](#).

Example

```
https://us-west-2.rds.amazonaws.com/
?Action=CreateDBInstanceReadReplica
&KmsKeyId=my-us-east-1-key
&PreSignedUrl=https%253A%252Frds.us-west-2.amazonaws.com%252F
%253FAction%253D CreateDBInstanceReadReplica
%2526DestinationRegion%253Dus-east-1
%2526KmsKeyId%253Dmy-us-east-1-key
%2526SourceDBIdentifier%253Darn%25253Aaws%25253Ards%25253Aus-
west-2%1234567890
12%25253Adb%25253Amy-mysql-instance
%2526SignatureMethod%253DHmacSHA256
%2526SignatureVersion%253D4%2526SourceDBIdentifier%253Darn%25253Aaws
%25253Ards%25253Aus-west-2%25253A123456789012%25253Ainstance%25253Amysql-instance1-
instance-20161115
%2526Version%253D2014-10-31
%2526X-Amz-Algorithm%253DAWS4-HMAC-SHA256
%2526X-Amz-Credential%253DAKIADQKE4SARGYLE%252F20161117%252Fus-west-2%252Frds
%252Faws4_request
%2526X-Amz-Date%253D20161117T215409Z
%2526X-Amz-Expires%253D3600
%2526X-Amz-SignedHeaders%253Dcontent-type%253Bhost%253Buser-agent%253Bx-amz-
content-sha256%253Bx-amz-date
%2526X-Amz-Signature
%253D255a0f17b4e717d3b67fad163c3ec26573b882c03a65523522cf890a67fca613
&DBIdentifier=myreadreplica
&SourceDBIdentifier=arn:aws:rds:us-east-1:123456789012:db:my-mysql-instance
&Version=2012-01-15
&SignatureVersion=2
&SignatureMethod=HmacSHA256
&Timestamp=2012-01-20T22%3A06%3A23.624Z
&AWSAccessKeyId=<AWS Access Key ID>
&Signature=<Signature>
```

Cross-Region Replication Considerations

All of the considerations for performing replication within an AWS Region apply to cross-region replication. The following extra considerations apply when replicating between AWS Regions:

- You can only replicate between AWS Regions when using the following Amazon RDS DB instances:
 - MariaDB (all versions).
 - MySQL version 5.6 and later.
 - Oracle Enterprise Edition (EE) engine version 12.1.0.2.v10 and higher 12.1 versions, all versions of 12.2, and all versions of 18.0.

An Active Data Guard license is required. For information about limitations for Oracle cross-region Read Replicas, see [Read Replica Limitations with Oracle \(p. 999\)](#).

- PostgreSQL version 9.4.7 and later 9.4 versions, and version 9.5.2 and later 9.5 versions.
- A source DB instance can have cross-region Read Replicas in multiple AWS Regions.
- You can only create a cross-region Amazon RDS Read Replica from a source Amazon RDS DB instance that is not a Read Replica of another Amazon RDS DB instance.
- You can't set up a replication channel into or out of the AWS GovCloud (US-West) Region.
- You can expect to see a higher level of lag time for any Read Replica that is in a different AWS Region than the source instance. This lag time comes from the longer network channels between regional data centers.

- Within an AWS Region, all cross-region Read Replicas created from the same source DB instance must either be in the same Amazon VPC or be outside of a VPC. For cross-region Read Replicas, any of the create Read Replica commands that specify the `--db-subnet-group-name` parameter must specify a DB subnet group from the same VPC.
- You can create a cross-region Read Replica in a VPC from a source DB instance that is in a VPC in another AWS Region. You can also create a cross-region Read Replica in a VPC from a source DB instance that is not in a VPC. You can also create a cross-region Read Replica that is not in a VPC from a source DB instance that is in a VPC.
- Due to the limit on the number of access control list (ACL) entries for a VPC, we can't guarantee more than five cross-region Read Replica instances.
- The Read Replica uses the default DB parameter group for the specified DB engine.
- The Read Replica uses the default security group.
- For Oracle DB instances, when the source for a cross-region Read Replica is deleted, the Read Replica is promoted. For MariaDB, MySQL, and PostgreSQL DB instances, when the source for a cross-region Read Replica is deleted, the replication status of the Read Replica is set to terminated. However, the Read Replica isn't promoted.

Cross-Region Replication Costs

The data transferred for cross-region replication incurs Amazon RDS data transfer charges. These cross-region replication actions generate charges for the data transferred out of the source AWS Region:

- When you create a Read Replica, Amazon RDS takes a snapshot of the source instance and transfers the snapshot to the Read Replica AWS Region.
- For each data modification made in the source databases, Amazon RDS transfers data from the source AWS Region to the Read Replica AWS Region.

For more information about data transfer pricing, see [Amazon RDS Pricing](#).

For MySQL and MariaDB instances, you can reduce your data transfer costs by reducing the number of cross-region Read Replicas that you create. For example, suppose that you have a source DB instance in one AWS Region and want to have three Read Replicas in another AWS Region. In this case, you create only one of the Read Replicas from the source DB instance. You create the other two replicas from the first Read Replica instead of the source DB instance.

For example, if you have `source-instance-1` in one AWS Region, you can do the following:

- Create `read-replica-1` in the new AWS Region, specifying `source-instance-1` as the source.
- Create `read-replica-2` from `read-replica-1`.
- Create `read-replica-3` from `read-replica-1`.

In this example, you are only charged for the data transferred from `source-instance-1` to `read-replica-1`. You are not charged for the data transferred from `read-replica-1` to the other two replicas because they are all in the same AWS Region. If you create all three replicas directly from `source-instance-1`, you are charged for the data transfers to all three replicas.

How Amazon RDS Does Cross-Region Replication

Amazon RDS uses the following process to create a cross-region Read Replica. Depending on the AWS Regions involved and the amount of data in the databases, this process can take hours to complete. You can use this information to determine how far the process has proceeded when you create a cross-region Read Replica:

1. Amazon RDS begins configuring the source DB instance as a replication source and sets the status to *modifying*.
2. Amazon RDS begins setting up the specified Read Replica in the destination AWS Region and sets the status to *creating*.
3. Amazon RDS creates an automated DB snapshot of the source DB instance in the source AWS Region. The format of the DB snapshot name is `rds:<InstanceID>-<timestamp>`, where `<InstanceID>` is the identifier of the source instance, and `<timestamp>` is the date and time the copy started. For example, `rds:mysourceinstance-2013-11-14-09-24` was created from the instance `mysourceinstance` at `2013-11-14-09-24`. During the creation of an automated DB snapshot, the source DB instance status remains *modifying*, the Read Replica status remains *creating*, and the DB snapshot status is *creating*. The progress column of the DB snapshot page in the console reports how far the DB snapshot creation has progressed. When the DB snapshot is complete, the status of both the DB snapshot and source DB instance are set to *available*.
4. Amazon RDS begins a cross-region snapshot copy for the initial data transfer. The snapshot copy is listed as an automated snapshot in the destination AWS Region with a status of *creating*. It has the same name as the source DB snapshot. The progress column of the DB snapshot display indicates how far the copy has progressed. When the copy is complete, the status of the DB snapshot copy is set to *available*.
5. Amazon RDS then uses the copied DB snapshot for the initial data load on the Read Replica. During this phase, the Read Replica is in the list of DB instances in the destination, with a status of *creating*. When the load is complete, the Read Replica status is set to *available*, and the DB snapshot copy is deleted.
6. When the Read Replica reaches the *available* status, Amazon RDS starts by replicating the changes made to the source instance since the start of the create Read Replica operation. During this phase, the replication lag time for the Read Replica is greater than 0.

For information about replication lag time, see [Monitoring Read Replication \(p. 194\)](#).

Cross-Region Replication Examples

Example Create a Cross-Region Read Replica Outside of Any VPC

The following example creates a Read Replica in us-west-2 from a source DB instance in us-east-1. The Read Replica is created outside of a VPC:

For Linux, OS X, or Unix:

```
aws rds create-db-instance-read-replica \
    --db-instance-identifier SimCoProd01Replica01 \
    --region us-west-2 \
    --source-db-instance-identifier arn:aws:rds:us-east-1:123456789012:db:SimcoProd01
```

For Windows:

```
aws rds create-db-instance-read-replica ^
    --db-instance-identifier SimCoProd01Replica01 ^
    --region us-west-2 ^
    --source-db-instance-identifier arn:aws:rds:us-east-1:123456789012:db:SimcoProd01
```

Example Create Cross-Region Read Replica in a VPC

This example creates a Read Replica in us-west-2 from a source DB instance in us-east-1. The Read Replica is created in the VPC associated with the specified DB subnet group:

For Linux, OS X, or Unix:

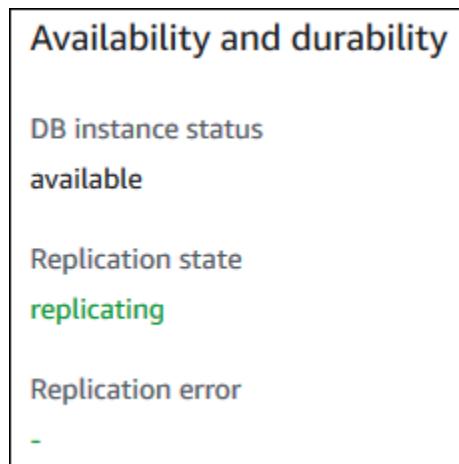
```
aws rds create-db-instance-read-replica \
--db-instance-identifier SimCoProd01Replica01 \
--region us-west-2 \
--db-subnet-group-name my-us-west-2-subnet \
--source-db-instance-identifier arn:aws:rds:us-east-1:123456789012:db:SimcoProd01
```

For Windows:

```
aws rds create-db-instance-read-replica ^
--db-instance-identifier SimCoProd01Replica01 ^
--region us-west-2 ^
--db-subnet-group-name my-us-west-2-subnet ^
--source-db-instance-identifier arn:aws:rds:us-east-1:123456789012:db:SimcoProd01
```

Monitoring Read Replication

You can monitor the status of a Read Replica in several ways. The Amazon RDS console shows the status of a Read Replica in the **Availability and durability** section of the Read Replica details. To view the details for a Read Replica, choose the name of the Read Replica in the list of instances in the Amazon RDS console.



You can also see the status of a Read Replica using the AWS CLI `describe-db-instances` command or the Amazon RDS API `DescribeDBInstances` action.

The status of a Read Replica can be one of the following:

- **replicating**—The Read Replica is replicating successfully.
- **error**—An error has occurred with the replication. Check the **Replication Error** field in the Amazon RDS console or the event log to determine the exact error. For more information about troubleshooting a replication error, see [Troubleshooting a MySQL Read Replica Problem \(p. 856\)](#).
- **terminated (MariaDB, MySQL, or PostgreSQL only)**—Replication is terminated. This occurs if replication is stopped for more than 30 consecutive days, either manually or due to a replication error. In this case, Amazon RDS terminates replication between the source DB instance and all Read Replicas. Amazon RDS does this to prevent increased storage requirements on the source DB instance and long failover times.

Broken replication can affect storage because the logs can grow in size and number due to the high volume of errors messages being written to the log. Broken replication can also affect failure recovery due to the time Amazon RDS requires to maintain and process the large number of logs during recovery.

- **stopped (MariaDB or MySQL only)**—Replication has stopped because of a customer initiated request.
- **replication stop point set (MySQL only)**—A customer initiated stop point was set using the [mysql.rds_start_replication_until \(p. 901\)](#) stored procedure and the replication is in progress.
- **replication stop point reached (MySQL only)**—A customer initiated stop point was set using the [mysql.rds_start_replication_until \(p. 901\)](#) stored procedure and replication is stopped because the stop point was reached.

You can monitor replication lag in Amazon CloudWatch by viewing the Amazon RDS ReplicaLag metric. For MySQL and MariaDB, the ReplicaLag metric reports the value of the Seconds_Behind_Master field of the SHOW SLAVE STATUS command. For PostgreSQL, the ReplicaLag metric reports the value of `SELECT extract(epoch from now() - pg_last_xact_replay_timestamp()) AS slave_lag`.

Common causes for replication lag for MySQL and MariaDB are the following:

- A network outage.
- Writing to tables with indexes on a Read Replica. If the `read_only` parameter is not set to 0 on the Read Replica, it can break replication.
- Using a nontransactional storage engine such as MyISAM. Replication is only supported for the InnoDB storage engine on MySQL and the XtraDB storage engine on MariaDB.

When the ReplicaLag metric reaches 0, the replica has caught up to the source DB instance. If the ReplicaLag metric returns -1, then replication is currently not active. `ReplicaLag = -1` is equivalent to `Seconds_Behind_Master = NULL`.

For Oracle, the ReplicaLag metric is the sum of the `Apply_Lag` value and the difference between the current time and the apply lag's `DATUM_TIME` value. The `DATUM_TIME` value is the last time the Read Replica received data from its source DB instance. For more information, see [V\\$DATAGUARD_STATS](#) in the Oracle documentation.

PostgreSQL versions 9.4.7 and 9.5.2 and later use physical replication slots to manage write ahead log (WAL) retention on the source instance. For each cross-region Read Replica instance, Amazon RDS creates a physical replication slot and associates it with the instance. Two Amazon CloudWatch metrics, Oldest Replication Slot Lag and Transaction Logs Disk Usage, show how far behind the most lagging replica is in terms of WAL data received and how much storage is being used for WAL data. The Transaction Logs Disk Usage value can substantially increase when a cross-region Read Replica is lagging significantly.

For more information about monitoring a DB instance with CloudWatch, see [Monitoring with Amazon CloudWatch \(p. 297\)](#).

Working with Option Groups

Some DB engines offer additional features that make it easier to manage data and databases, and to provide additional security for your database. Amazon RDS uses option groups to enable and configure these features. An *option group* can specify features, called options, that are available for a particular Amazon RDS DB instance. Options can have settings that specify how the option works. When you associate a DB instance with an option group, the specified options and option settings are enabled for that DB instance.

Amazon RDS supports options for the following database engines:

Database Engine	Relevant Documentation
MariaDB	Options for MariaDB Database Engine (p. 587)
Microsoft SQL Server	Options for the Microsoft SQL Server Database Engine (p. 712)
MySQL	Options for MySQL DB Instances (p. 873)
Oracle	Options for Oracle DB Instances (p. 1004)

Option Groups Overview

Amazon RDS provides an empty default option group for each new DB instance. You cannot modify this default option group, but any new option group that you create derives its settings from the default option group. To apply an option to a DB instance, you must do the following:

1. Create a new option group, or copy or modify an existing option group.
2. Add one or more options to the option group.
3. Associate the option group with the DB instance.

Both DB instances and DB snapshots can be associated with an option group. In some cases, you might restore from a DB snapshot or perform a point-in-time restore for a DB instance. In these cases, the option group associated with the DB snapshot or DB instance is, by default, associated with the restored DB instance. You can associate a different option group with a restored DB instance. However, the new option group must contain any persistent or permanent options that were included in the original option group. Persistent and permanent options are described following.

Options require additional memory to run on a DB instance. Thus, you might need to launch a larger instance to use them, depending on your current use of your DB instance. For example, Oracle Enterprise Manager Database Control uses about 300 MB of RAM. If you enable this option for a small DB instance, you might encounter performance problems or out-of-memory errors.

Persistent and Permanent Options

Two types of options, persistent and permanent, require special consideration when you add them to an option group.

Persistent options can't be removed from an option group while DB instances are associated with the option group. An example of a persistent option is the TDE option for Microsoft SQL Server transparent data encryption (TDE). You must disassociate all DB instances from the option group before a persistent option can be removed from the option group. In some cases, you might restore or perform a point-in-time restore from a DB snapshot. In these cases, if the option group associated with that DB snapshot contains a persistent option, you can only associate the restored DB instance with that option group.

Permanent options, such as the TDE option for Oracle Advanced Security TDE, can never be removed from an option group. You can change the option group of a DB instance that is using the permanent option. However, the option group associated with the DB instance must include the same permanent option. In some cases, you might restore or perform a point-in-time restore from a DB snapshot. In these cases, if the option group associated with that DB snapshot contains a permanent option, you can only associate the restored DB instance with an option group with that permanent option.

For Oracle DB instances, you can copy shared DB snapshots that have the options `Timezone` or `OLS` (or both). To do so, specify a target option group that includes these options when you copy the DB snapshot. The OLS option is permanent and persistent only for Oracle DB instances running Oracle version 12.2 or higher. For more information about these options, see [Oracle Time Zone \(p. 1071\)](#) and [Oracle Label Security \(p. 1040\)](#).

VPC and Platform Considerations

When an option group is assigned to a DB instance, it is linked to the platform that the DB instance is on. That platform can either be a VPC supported by the Amazon VPC service, or EC2-Classic (non-VPC) supported by the Amazon EC2 service. For details on these two platforms, see [Amazon EC2](#) and [Amazon Virtual Private Cloud](#).

If a DB instance is in a VPC, the option group associated with the instance is linked to that VPC. This means that you can't use the option group assigned to a DB instance if you try to restore the instance to a different VPC or a different platform. If you restore a DB instance to a different VPC or a different platform, you can do one of the following:

- Assign the default option group to the DB instance.
- Assign an option group that is linked to that VPC or platform.
- Create a new option group and assign it to the DB instance.

With persistent or permanent options, such as Oracle TDE, you must create a new option group that includes the persistent or permanent option when restoring a DB instance into a different VPC.

Option settings control the behavior of an option. For example, the Oracle Advanced Security option `NATIVE_NETWORK_ENCRYPTION` has a setting that you can use to specify the encryption algorithm for network traffic to and from the DB instance. Some options settings are optimized for use with Amazon RDS and cannot be changed.

Mutually Exclusive Options

Some options are mutually exclusive. You can use one or the other, but not both at the same time. The following options are mutually exclusive:

- [Oracle Enterprise Manager Database Express \(p. 1025\)](#) and [Oracle Management Agent for Enterprise Manager Cloud Control \(p. 1028\)](#).
- [Oracle Native Network Encryption \(p. 1049\)](#) and [Oracle Secure Sockets Layer \(p. 1051\)](#).

Creating an Option Group

You can create a new option group that derives its settings from the default option group, and then add one or more options to the new option group. Alternatively, if you already have an existing option group, you can copy that option group with all of its options to a new option group. For more information, see [Making a Copy of an Option Group \(p. 199\)](#).

After you create a new option group, it has no options. To learn how to add options to the option group, see [Adding an Option to an Option Group \(p. 200\)](#). After you have added the options you want, you

can then associate the option group with a DB instance so that the options become available on the DB instance. For information about associating an option group with a DB instance, see the documentation for your specific engine listed at [Working with Option Groups \(p. 196\)](#).

Console

One way of creating an option group is by using the AWS Management Console.

To create a new option group by using the console

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Option groups**.
3. Choose **Create group**.
4. In the **Create option group** window, do the following:
 - a. For **Name**, type a name for the option group that is unique within your AWS account. The name can contain only letters, digits, and hyphens.
 - b. For **Description**, type a brief description of the option group. The description is used for display purposes.
 - c. For **Engine**, choose the DB engine that you want.
 - d. For **Major engine version**, choose the major version of the DB engine that you want.
5. To continue, choose **Create**. To cancel the operation instead, choose **Cancel**.

AWS CLI

To create an option group, use the AWS CLI `create-option-group` command with the following required parameters.

- `--option-group-name`
- `--engine-name`
- `--major-engine-version`
- `--option-group-description`

Example

The following example creates an option group named `testoptiongroup`, which is associated with the Oracle Enterprise Edition DB engine. The description is enclosed in quotation marks.

For Linux, OS X, or Unix:

```
aws rds create-option-group \
--option-group-name testoptiongroup \
--engine-name oracle-ee \
--major-engine-version 12.1 \
--option-group-description "Test option group"
```

For Windows:

```
aws rds create-option-group ^
--option-group-name testoptiongroup ^
--engine-name oracle-ee ^
```

```
--major-engine-version 12.1 ^
--option-group-description "Test option group"
```

RDS API

To create an option group, call the Amazon RDS API [CreateOptionGroup](#) operation. Include the following parameters:

- OptionGroupName
- EngineName
- MajorEngineVersion
- OptionGroupDescription

Making a Copy of an Option Group

You can use the AWS CLI or the Amazon RDS API to make a copy of an option group. Copying an option group is convenient when you have an existing option group and you want to include most of its custom parameters and values in a new option group. You can also make a copy of an option group that you use in production and then modify the copy to test other option settings.

AWS CLI

To copy an option group, use the AWS CLI [copy-option-group](#) command. Include the following required parameters:

- --source-option-group-identifier
- --target-option-group-identifier
- --target-option-group-description

Example

The following example creates an option group named `new-local-option-group`, which is a local copy of the option group `my-remote-option-group`.

For Linux, OS X, or Unix:

```
aws rds copy-option-group \
  --source-option-group-identifier arn:aws:rds:us-west-2:123456789012:og:my-remote-
option-group \
  --target-option-group-identifier new-local-option-group \
  --target-option-group-description "Option group 2"
```

For Windows:

```
aws rds copy-option-group ^
  --source-option-group-identifier arn:aws:rds:us-west-2:123456789012:og:my-remote-
option-group ^
  --target-option-group-identifier new-local-option-group ^
  --target-option-group-description "Option group 2"
```

RDS API

To copy an option group, call the Amazon RDS API [CopyOptionGroup](#) operation. Include the following required parameters.

- `SourceOptionGroupIdentifier`
- `TargetOptionGroupIdentifier`
- `TargetOptionGroupDescription`

Adding an Option to an Option Group

You can add an option to an existing option group. After you have added the options you want, you can then associate the option group with a DB instance so that the options become available on the DB instance. For information about associating an option group with a DB instance, see the documentation for your specific DB engine listed at [Working with Option Groups \(p. 196\)](#).

Option group changes must be applied immediately in two cases:

- When you add an option that adds or updates a port value, such as the `OEM` option.
- When you add or remove an option group with an option that includes a port value.

In these cases, choose the **Apply Immediately** option in the console. Or you can include the `--apply-immediately` option when using the AWS CLI or set the `ApplyImmediately` parameter to `true` when using the Amazon RDS API. Options that don't include port values can be applied immediately, or can be applied during the next maintenance window for the DB instance.

Console

You can use the AWS Management Console to add an option to an option group.

To add an option to an option group by using the console

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Option groups**.
3. Select the option group that you want to modify, and then choose **Add Option**.

Option groups (9)		
<input type="text"/> Filter subnet groups		
<input type="checkbox"/>	Name	Description
<input type="checkbox"/>	carpcmysql	carpcmysql
<input checked="" type="checkbox"/>	carporacle	carporacle
<input type="checkbox"/>	default:mysql-5-5	Default option group for mysql 5.5
<input type="checkbox"/>	default:mysql-5-6	Default option group for mysql 5.6
<input type="checkbox"/>	default:mysql-5-7	Default option group for mysql 5.7

4. In the **Add option** window, do the following:

- a. Choose the option that you want to add. You might need to provide additional values, depending on the option that you select. For example, when you choose the `OEM` option, you must also type a port value and specify a DB security group.
- b. To enable the option on all associated DB instances as soon as you add it, for **Apply Immediately**, choose **Yes**. If you choose **No** (the default), the option is enabled for each associated DB instance during its next maintenance window.

Add Option

Option details

Option group name
carpcoracle

Option
Name of Option you want to add to this group

Port
The port number, if applicable, to use when connecting to the Option

Security Groups
A list of VPC or DB Security Groups for which this Option is enabled

default

Apply Immediately [info](#)
 Yes
 No

- When the settings are as you want them, choose **Add Option**.

AWS CLI

To add an option to an option group, run the AWS CLI [add-option-to-option-group](#) command with the option that you want to add. To enable the new option immediately on all associated DB instances, include the `--apply-immediately` parameter. By default, the option is enabled for each associated DB instance during its next maintenance window. Include the following required parameter:

- `--option-group-name`

Example

The following example adds the Oracle Enterprise Manager Database Control (OEM) option to an option group named `testoptiongroup` and immediately enables it. Even if you use the default security group, you must specify that security group.

For Linux, OS X, or Unix:

```
aws rds add-option-to-option-group \
```

```
--option-group-name testoptiongroup \
--options OptionName=OEM,Port=5500,DBSecurityGroupMemberships=default \
--apply-immediately
```

For Windows:

```
aws rds add-option-to-option-group ^
--option-group-name testoptiongroup ^
--options OptionName=OEM,Port=5500,DBSecurityGroupMemberships=default ^
--apply-immediately
```

Command output is similar to the following:

```
OPTIONGROUP  False  oracle-ee  12.1  arn:aws:rds:us-east-1:1234567890:og:testoptiongroup
Test Option Group  testoptiongroup default
OPTIONS Oracle 12c EM Express  OEM      False    False   5500
DBSECURITYGROUPEMEMBERSHIPS  default authorized
```

Example

The following example adds the Oracle OEM option to an option group. It also specifies a custom port and a pair of Amazon EC2 VPC security groups to use for that port.

For Linux, OS X, or Unix:

```
aws rds add-option-to-option-group \
--option-group-name testoptiongroup \
--options OptionName=OEM,Port=5500,VpcSecurityGroupMemberships="sg-test1,sg-test2" \
--apply-immediately
```

For Windows:

```
aws rds add-option-to-option-group ^
--option-group-name testoptiongroup ^
--options OptionName=OEM,Port=5500,VpcSecurityGroupMemberships="sg-test1,sg-test2" ^
--apply-immediately
```

Command output is similar to the following:

```
OPTIONGROUP  False  oracle-ee  12.1  arn:aws:rds:us-east-1:1234567890:og:testoptiongroup
Test Option Group  testoptiongroup vpc-test
OPTIONS Oracle 12c EM Express  OEM      False    False   5500
VPCSECURITYGROUPEMEMBERSHIPS  active  sg-test1
VPCSECURITYGROUPEMEMBERSHIPS  active  sg-test2
```

Example

The following example adds the Oracle option NATIVE_NETWORK_ENCRYPTION to an option group and specifies the option settings. If no option settings are specified, default values are used.

For Linux, OS X, or Unix:

```
aws rds add-option-to-option-group \
--option-group-name testoptiongroup \
--options '[{"OptionSettings": [{"Name": "SQLNET.ENCRYPTION_SERVER", "Value": "REQUIRED"}, {"Name": "SQLNET.ENCRYPTION_TYPES_SERVER", "Value": "AES256,AES192,DES"}]}, {"OptionName": "NATIVE_NETWORK_ENCRYPTION", "Value": "AES256,AES192,DES"}]' \
--apply-immediately
```

For Windows:

```
aws rds add-option-to-option-group ^
--option-group-name testoptiongroup ^
--options "OptionSettings=[{ "Name": "SQLNET.ENCRYPTION_SERVER", "Value": "REQUIRED"}, { "Name": "SQLNET.ENCRYPTION_TYPES_SERVER", "Value": "AES256\,AES192\,DES"}], "OptionName": "NATIVE_NETWORK_ENCRYPTION", ^
--apply-immediately
```

Command output is similar to the following:

```
OPTIONGROUP False oracle-ee 12.1 arn:aws:rds:us-east-1:1234567890:og:testoptiongroup
Test Option Group testoptiongroup
OPTIONS Oracle Advanced Security - Native Network Encryption      NATIVE_NETWORK_ENCRYPTION
      False False
OPTIONSETTINGS
      RC4_256,AES256,AES192,3DES168,RC4_128,AES128,3DES112,RC4_56,DES,RC4_40,DES40
          STATIC STRING
      RC4_256,AES256,AES192,3DES168,RC4_128,AES128,3DES112,RC4_56,DES,RC4_40,DES40      Specifies
          list of encryption algorithms in order of intended use
      True True SQLNET.ENCRYPTION_TYPES_SERVER AES256,AES192,DES
OPTIONSETTINGS ACCEPTED,REJECTED,REQUESTED,REQUIRED STATIC STRING REQUESTED
          Specifies the desired encryption behavior False True SQLNET.ENCRYPTION_SERVER
          REQUIRED
OPTIONSETTINGS SHA1,MD5 STATIC STRING SHA1,MD5 Specifies list of checksumming
          algorithms in order of intended use True True SQLNET.CRYPTO_CHECKSUM_TYPES_SERVER
          SHA1,MD5
```

RDS API

To add an option to an option group using the Amazon RDS API, call the [ModifyOptionGroup](#) operation with the option that you want to add. To enable the new option immediately on all associated DB instances, include the `ApplyImmediately` parameter and set it to `true`. By default, the option is enabled for each associated DB instance during its next maintenance window. Include the following required parameter:

- `OptionGroupName`

Listing the Options and Option Settings for an Option Group

You can list all the options and option settings for an option group.

Console

You can use the AWS Management Console to list all of the options and option settings for an option group.

To list the options and option settings for an option group

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Option groups**.
3. Choose the name of the option group to display its details. The options and option settings in the option group are listed.

AWS CLI

To list the options and option settings for an option group, use the AWS CLI [describe-option-groups](#) command. Specify the name of the option group whose options and settings you want to view. If you don't specify an option group name, all option groups are described.

Example

The following example lists the options and option settings for all option groups.

```
aws rds describe-option-groups
```

Example

The following example lists the options and option settings for an option group named `testoptiongroup`.

```
aws rds describe-option-groups --option-group-name testoptiongroup
```

RDS API

To list the options and option settings for an option group, use the Amazon RDS API [DescribeOptionGroups](#) operation. Specify the name of the option group whose options and settings you want to view. If you don't specify an option group name, all option groups are described.

Modifying an Option Setting

After you have added an option that has modifiable option settings, you can modify the settings at any time. If you change options or option settings in an option group, those changes are applied to all DB instances that are associated with that option group. For more information on what settings are available for the various options, see the documentation for your specific engine listed at [Working with Option Groups \(p. 196\)](#).

Option group changes must be applied immediately in two cases:

- When you add an option that adds or updates a port value, such as the `OEM` option.
- When you add or remove an option group with an option that includes a port value.

In these cases, choose the **Apply Immediately** option in the console. Or you can include the `--apply-immediately` option when using the AWS CLI or set the `ApplyImmediately` parameter to `true` when using the RDS API. Options that don't include port values can be applied immediately, or can be applied during the next maintenance window for the DB instance.

Console

You can use the AWS Management Console to modify an option setting.

To modify an option setting by using the console

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Option groups**.
3. Select the option group whose option that you want to modify, and then choose **Modify option**.
4. In the **Modify option** window, from **Installed Options**, choose the option whose setting you want to modify. Make the changes that you want.
5. To enable the option as soon as you add it, for **Apply Immediately**, choose **Yes**. If you choose **No** (the default), the option is enabled for each associated DB instance during its next maintenance window.
6. When the settings are as you want them, choose **Modify Option**.

AWS CLI

To modify an option setting, use the AWS CLI `add-option-to-option-group` command with the option group and option that you want to modify. By default, the option is enabled for each associated DB instance during its next maintenance window. To apply the change immediately to all associated DB instances, include the `--apply-immediately` parameter. To modify an option setting, use the `--settings` argument.

Example

The following example modifies the port that the Oracle Enterprise Manager Database Control (OEM) uses in an option group named `testoptiongroup` and immediately applies the change.

For Linux, OS X, or Unix:

```
aws rds add-option-to-option-group \
--option-group-name testoptiongroup \
--options OptionName=OEM,Port=5432,DBSecurityGroupMemberships=default \
--apply-immediately
```

For Windows:

```
aws rds add-option-to-option-group ^
--option-group-name testoptiongroup ^
--options OptionName=OEM,Port=5432,DBSecurityGroupMemberships=default ^
--apply-immediately
```

Command output is similar to the following:

```
OPTIONGROUP  False  oracle-ee  12.1  arn:aws:rds:us-east-1:1234567890:og:testoptiongroup
Test Option Group      testoptiongroup
OPTIONS Oracle 12c EM Express OEM      False    False   5432
DBSECURITYGROUPMEMBERSHIPS default authorized
```

Example

The following example modifies the Oracle option NATIVE_NETWORK_ENCRYPTION and changes the option settings.

For Linux, OS X, or Unix:

```
aws rds add-option-to-option-group \
--option-group-name testoptiongroup \
--options '[{"OptionSettings": [{"Name": "SQLNET.ENCRYPTION_SERVER", "Value": "REQUIRED"}, {"Name": "SQLNET.ENCRYPTION_TYPES_SERVER", "Value": "AES256,AES192,DES,RC4_256"}]}, {"OptionName": "NATIVE_NETWORK_ENCRYPTION", "Value": "AES256,AES192,DES,RC4_256"}]' \
--apply-immediately
```

For Windows:

```
aws rds add-option-to-option-group ^
--option-group-name testoptiongroup ^
--options "OptionSettings=[{"Name": "SQLNET.ENCRYPTION_SERVER", "Value": "REQUIRED"}, {"Name": "SQLNET.ENCRYPTION_TYPES_SERVER", "Value": "AES256\,AES192\,DES\,RC4_256"}], "OptionName": "NATIVE_NETWORK_ENCRYPTION" ^
--apply-immediately
```

Command output is similar to the following:

```
OPTIONGROUP  False  oracle-ee 12.1  arn:aws:rds:us-east-1:1234567890:og:testoptiongroup
Test Option Group      testoptiongroup
OPTIONS Oracle Advanced Security - Native Network Encryption      NATIVE_NETWORK_ENCRYPTION
      False    False
OPTIONSETTINGS
      RC4_256,AES256,AES192,3DES168,RC4_128,AES128,3DES112,RC4_56,DES,RC4_40,DES40  STATIC
      STRING
      RC4_256,AES256,AES192,3DES168,RC4_128,AES128,3DES112,RC4_56,DES,RC4_40,DES40
      Specifies list of encryption algorithms in order of intended use
      True    True  SQLNET.ENCRYPTION_TYPES_SERVER  AES256,AES192,DES,RC4_256
OPTIONSETTINGS ACCEPTED,REJECTED,REQUESTED,REQUIRED  STATIC STRING REQUESTED
      Specifies the desired encryption behavior  False    True  SQLNET.ENCRYPTION_SERVER
      REQUIRED
OPTIONSETTINGS SHA1,MD5  STATIC STRING SHA1,MD5      Specifies list of
      checksumming algorithms in order of intended use      True    True
      SQLNET.CRYPTO_CHECKSUM_TYPES_SERVER  SHA1,MD5
OPTIONSETTINGS ACCEPTED,REJECTED,REQUESTED,REQUIRED  STATIC STRING
      REQUESTED      Specifies the desired data integrity behavior  False    True
      SQLNET.CRYPTO_CHECKSUM_SERVER  REQUESTED
```

RDS API

To modify an option setting, use the Amazon RDS API [ModifyOptionGroup](#) command with the option group and option that you want to modify. By default, the option is enabled for each associated DB instance during its next maintenance window. To apply the change immediately to all associated DB instances, include the `ApplyImmediately` parameter and set it to `true`.

Removing an Option from an Option Group

Some options can be removed from an option group, and some cannot. A persistent option cannot be removed from an option group until all DB instances associated with that option group are disassociated. A permanent option can never be removed from an option group. For more information about what options are removable, see the documentation for your specific engine listed at [Working with Option Groups \(p. 196\)](#).

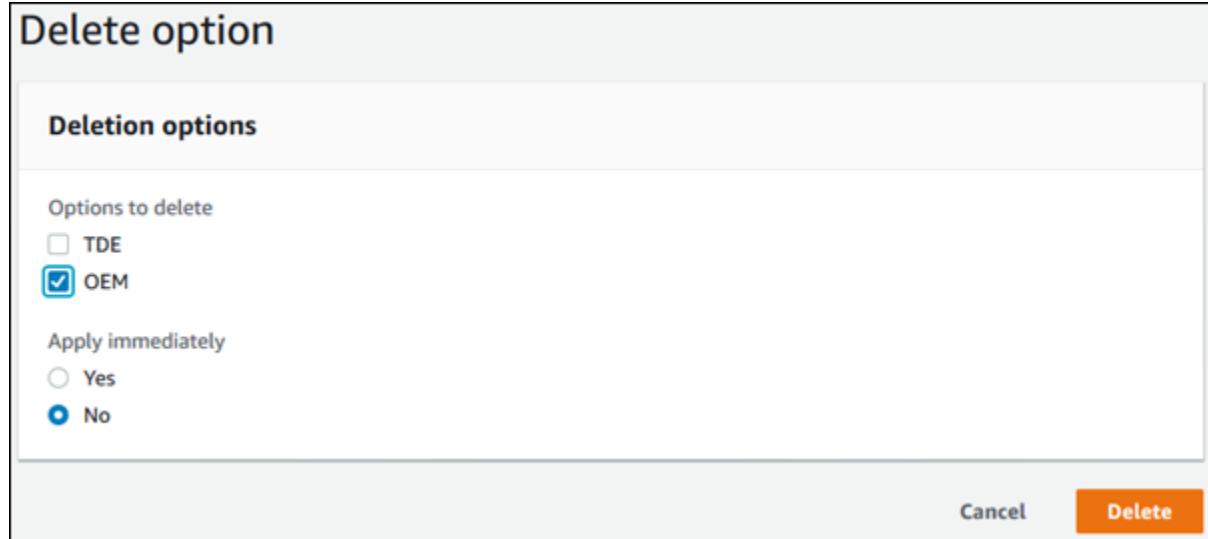
If you remove all options from an option group, Amazon RDS doesn't delete the option group. DB instances that are associated with the empty option group continue to be associated with it; they just won't have any active options. Alternatively, to remove all options from a DB instance, you can associate the DB instance with the default (empty) option group.

Console

You can use the AWS Management Console to remove an option from an option group.

To remove an option from an option group by using the console

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Option groups**.
3. Select the option group whose option you want to remove, and then choose **Delete option**.
4. In the **Delete option** window, do the following:
 - Select the check box for the option that you want to delete.
 - For the deletion to take effect as soon as you make it, for **Apply immediately**, choose **Yes**. If you choose **No** (the default), the option is deleted for each associated DB instance during its next maintenance window.



5. When the settings are as you want them, choose **Yes, Delete**.

AWS CLI

To remove an option from an option group, use the AWS CLI `remove-option-from-option-group` command with the option that you want to delete. By default, the option is removed from each

associated DB instance during its next maintenance window. To apply the change immediately, include the --apply-immediately parameter.

Example

The following example removes the Oracle Enterprise Manager Database Control (OEM) option from an option group named `testoptiongroup` and immediately applies the change.

For Linux, OS X, or Unix:

```
aws rds remove-option-from-option-group \
--option-group-name testoptiongroup \
--options OEM \
--apply-immediately
```

For Windows:

```
aws rds remove-option-from-option-group ^
--option-group-name testoptiongroup ^
--options OEM ^
--apply-immediately
```

Command output is similar to the following:

```
OPTIONGROUP      testoptiongroup oracle-ee    12.1      Test option group
```

RDS API

To remove an option from an option group, use the Amazon RDS API [ModifyOptionGroup](#) action. By default, the option is removed from each associated DB instance during its next maintenance window. To apply the change immediately, include the `ApplyImmediately` parameter and set it to `true`.

Include the following parameters:

- `OptionGroupName`
- `OptionsToRemove.OptionName`

Deleting an Option Group

You can delete an option group that is not associated with any Amazon RDS resource. An option group can be associated with a DB instance, a manual DB snapshot, or an automated DB snapshot.

If you try to delete an option group that is associated with an Amazon RDS resource, an error similar to the following is returned.

```
An error occurred (InvalidOptionGroupStateException) when calling the DeleteOptionGroup
operation: The option group 'optionGroupName' cannot be deleted because it is in use.
```

To find the Amazon RDS resources associated with an option group

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Option groups**.
3. Choose the name of the option group to show its details.
4. Check the **Associated Instances and Snapshots** section for the associated Amazon RDS resources.

If a DB instance is associated with the option group, modify the DB instance to use a different option group. For more information, see [Modifying an Amazon RDS DB Instance \(p. 150\)](#).

If a manual DB snapshot is associated with the option group, modify the DB snapshot to use a different option group using the AWS CLI `modify-db-snapshot` command.

Note

You can't modify the option group of an automated DB snapshot.

Console

One way of deleting an option group is by using the AWS Management Console.

To delete an option group by using the console

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Option groups**.
3. Choose the option group.
4. Choose **Delete group**.
5. On the confirmation page, choose **Delete** to finish deleting the option group, or choose **Cancel** to cancel the deletion.

AWS CLI

To delete an option group, use the AWS CLI `delete-option-group` command with the following required parameter.

- `--option-group-name`

Example

The following example deletes an option group named `testoptiongroup`.

For Linux, OS X, or Unix:

```
aws rds delete-option-group \
--option-group-name testoptiongroup
```

For Windows:

```
aws rds delete-option-group ^
--option-group-name testoptiongroup
```

RDS API

To delete an option group, call the Amazon RDS API [DeleteOptionGroup](#) operation. Include the following parameter:

- `OptionGroupName`

Working with DB Parameter Groups

You manage your DB engine configuration by associating your DB instances with parameter groups. Amazon RDS defines parameter groups with default settings that apply to newly created DB instances. You can define your own parameter groups with customized settings. Then you can modify your DB instances to use your own parameter groups.

A *DB parameter group* acts as a container for engine configuration values that are applied to one or more DB instances.

If you create a DB instance without specifying a DB parameter group, the DB instance uses a default DB parameter group. Each default DB parameter group contains database engine defaults and Amazon RDS system defaults based on the engine, compute class, and allocated storage of the instance. You can't modify the parameter settings of a default parameter group. Instead, you create your own parameter group where you choose your own parameter settings. Not all DB engine parameters can be changed in a parameter group that you create.

If you want to use your own parameter group, you create a new parameter group and modify the parameters that you want to. You then modify your DB instance to use the new parameter group. If you update parameters within a DB parameter group, the changes apply to all DB instances that are associated with that parameter group.

You can copy an existing DB parameter group with the AWS CLI [copy-db-parameter-group](#) command. Copying a parameter group can be convenient when you want to include most of an existing DB parameter group's custom parameters and values in a new DB parameter group.

Here are some important points about working with parameters in a DB parameter group:

- When you change a dynamic parameter and save the DB parameter group, the change is applied immediately regardless of the **Apply Immediately** setting. When you change a static parameter and save the DB parameter group, the parameter change takes effect after you manually reboot the DB instance. You can reboot a DB instance using the RDS console or by explicitly calling the `RebootDBInstance` API operation (without failover, if the DB instance is in a Multi-AZ deployment). The requirement to reboot the associated DB instance after a static parameter change helps mitigate the risk of a parameter misconfiguration affecting an API call, such as calling `ModifyDBInstance` to change DB instance class or scale storage.

If a DB instance isn't using the latest changes to its associated DB parameter group, the AWS Management Console shows the DB parameter group with a status of **pending-reboot**. The **pending-reboot** parameter groups status doesn't result in an automatic reboot during the next maintenance window. To apply the latest parameter changes to that DB instance, manually reboot the DB instance.

- When you change the DB parameter group associated with a DB instance, you must manually reboot the instance before the new DB parameter group is used by the DB instance.
- You can specify the value for a DB parameter as an integer or as an integer expression built from formulas, variables, functions, and operators. Functions can include a mathematical log expression. For more information, see [DB Parameter Values \(p. 220\)](#).
- Set any parameters that relate to the character set or collation of your database in your parameter group before creating the DB instance and before you create a database in your DB instance. This ensures that the default database and new databases in your DB instance use the character set and collation values that you specify. If you change character set or collation parameters for your DB instance, the parameter changes are not applied to existing databases.

You can change character set or collation values for an existing database using the `ALTER DATABASE` command, for example:

```
ALTER DATABASE database_name CHARACTER SET character_set_name COLLATE collation;
```

- Improperly setting parameters in a DB parameter group can have unintended adverse effects, including degraded performance and system instability. Always exercise caution when modifying database parameters and back up your data before modifying a DB parameter group. Try out parameter group setting changes on a test DB instance before applying those parameter group changes to a production DB instance.
- To determine the supported parameters for your DB engine, you can view the parameters in the DB parameter group used by the DB instance. For more information, see [Viewing Parameter Values for a DB Parameter Group \(p. 219\)](#).

Topics

- [Creating a DB Parameter Group \(p. 212\)](#)
- [Modifying Parameters in a DB Parameter Group \(p. 213\)](#)
- [Copying a DB Parameter Group \(p. 217\)](#)
- [Listing DB Parameter Groups \(p. 218\)](#)
- [Viewing Parameter Values for a DB Parameter Group \(p. 219\)](#)
- [Comparing DB Parameter Groups \(p. 220\)](#)
- [DB Parameter Values \(p. 220\)](#)

Creating a DB Parameter Group

You can create a new DB parameter group using the AWS Management Console, the AWS CLI, or the RDS API.

Console

To create a DB parameter group

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Parameter groups**.
3. Choose **Create parameter group**.

The **Create parameter group** window appears.

4. In the **Parameter group family** list, select a DB parameter group family.
5. In the **Type** list, select **DB Parameter Group**.
6. In the **Group name** box, enter the name of the new DB parameter group.
7. In the **Description** box, enter a description for the new DB parameter group.
8. Choose **Create**.

AWS CLI

To create a DB parameter group, use the AWS CLI `create-db-parameter-group` command. The following example creates a DB parameter group named *mydbparametergroup* for MySQL version 5.6 with a description of "My new parameter group."

Include the following required parameters:

- `--db-parameter-group-name`
- `--db-parameter-group-family`
- `--description`

To list all of the available parameter group families, use the following command:

```
aws rds describe-db-engine-versions --query "DBEngineVersions[].DBParameterGroupFamily"
```

Note

The output contains duplicates.

Example

For Linux, OS X, or Unix:

```
aws rds create-db-parameter-group \
--db-parameter-group-name mydbparametergroup \
--db-parameter-group-family MySQL5.6 \
--description "My new parameter group"
```

For Windows:

```
aws rds create-db-parameter-group ^
--db-parameter-group-name mydbparametergroup ^
--db-parameter-group-family MySQL5.6 ^
--description "My new parameter group"
```

This command produces output similar to the following:

```
DBPARAMETERGROUP mydbparametergroup mysql5.6 My new parameter group
```

RDS API

To create a DB parameter group, use the RDS API [CreateDBParameterGroup](#) operation.

Include the following required parameters:

- `DBParameterGroupName`
- `DBParameterGroupFamily`
- `Description`

Modifying Parameters in a DB Parameter Group

You can modify parameter values in a customer-created DB parameter group; you can't change the parameter values in a default DB parameter group. Changes to parameters in a customer-created DB parameter group are applied to all DB instances that are associated with the DB parameter group.

If you change a parameter value, when the change is applied is determined by the type of parameter. Changes to dynamic parameters are applied immediately. Changes to static parameters require that

the DB instance associated with DB parameter group be rebooted before the change takes effect. To determine the type of a parameter, list the parameters in a parameter group using one of the procedures shown in the section [Listing DB Parameter Groups \(p. 218\)](#).

The RDS console shows the status of the DB parameter group associated with a DB instance on the **Configuration** tab. For example, if the DB instance isn't using the latest changes to its associated DB parameter group, the RDS console shows the DB parameter group with a status of **pending-reboot**. To apply the latest parameter changes to that DB instance, manually reboot the DB instance.

Connectivity Monitoring Logs & events Configuration

Instance

Configuration

DB instance id
oracle-instance1

Engine version
12.1.0.2.v14

Storage type
General Purpose (SSD)

IOPS
-

Storage
20 GiB

DB name
ORCL

License model
Bring Your Own License

Character set
AL32UTF8

Option groups
default:oracle-ee-12-1

ARN
[REDACTED]

Resource id
[REDACTED]

Created time
Fri Nov 30 2018 15:41:20 GMT-0800 (Pacific Standard Time)

Parameter group
[testsqlnet \(pending-reboot\)](#)

Deletion protection
Disabled

Console

To modify a DB parameter group

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Parameter groups**.
3. In the list, choose the parameter group that you want to modify.
4. For **Parameter group actions**, choose **Edit**.
5. Change the values of the parameters that you want to modify. You can scroll through the parameters using the arrow keys at the top right of the dialog box.

You can't change values in a default parameter group.
6. Choose **Save changes**.

AWS CLI

To modify a DB parameter group, use the AWS CLI `modify-db-parameter-group` command with the following required parameters:

- `--db-parameter-group-name`
- `--parameters`

The following example modifies the `max_connections` and `max_allowed_packet` values in the DB parameter group named *mydbparametergroup*.

Example

For Linux, OS X, or Unix:

```
aws rds modify-db-parameter-group \
  --db-parameter-group-name mydbparametergroup \
  --parameters "ParameterName=max_connections,ParameterValue=250,ApplyMethod=immediate" \
  "ParameterName=max_allowed_packet,ParameterValue=1024,ApplyMethod=immediate"
```

For Windows:

```
aws rds modify-db-parameter-group ^
  --db-parameter-group-name mydbparametergroup ^
  --parameters "ParameterName=max_connections,ParameterValue=250,ApplyMethod=immediate" ^
  "ParameterName=max_allowed_packet,ParameterValue=1024,ApplyMethod=immediate"
```

The command produces output like the following:

```
DBPARAMETERGROUP  mydbparametergroup
```

RDS API

To modify a DB parameter group, use the RDS API `ModifyDBParameterGroup` command with the following required parameters:

- `DBParameterGroupName`
- `Parameters`

Copying a DB Parameter Group

You can copy custom DB parameter groups that you create. Copying a parameter group is a convenient solution when you have already created a DB parameter group and you want to include most of the custom parameters and values from that group in a new DB parameter group. You can copy a DB parameter group by using the AWS CLI [copy-db-parameter-group](#) command or the RDS API [CopyDBParameterGroup](#) operation.

After you copy a DB parameter group, wait at least 5 minutes before creating your first DB instance that uses that DB parameter group as the default parameter group. Doing this allows Amazon RDS to fully complete the copy action before the parameter group is used. This is especially important for parameters that are critical when creating the default database for a DB instance. An example is the character set for the default database defined by the `character_set_database` parameter. Use the **Parameter Groups** option of the [Amazon RDS console](#) or the [describe-db-parameters](#) command to verify that your DB parameter group is created.

Note

You can't copy a default parameter group. However, you can create a new parameter group that is based on a default parameter group.

Console

To copy a DB parameter group

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Parameter groups**.
3. In the list, choose the custom parameter group that you want to copy.
4. For **Parameter group actions**, choose **Copy**.
5. In **New DB parameter group identifier**, enter a name for the new parameter group.
6. In **Description**, enter a description for the new parameter group.
7. Choose **Copy**.

AWS CLI

To copy a DB parameter group, use the AWS CLI [copy-db-parameter-group](#) command with the following required parameters:

- `--source-db-parameter-group-identifier`
- `--target-db-parameter-group-identifier`
- `--target-db-parameter-group-description`

The following example creates a new DB parameter group named `mygroup2` that is a copy of the DB parameter group `mygroup1`.

Example

For Linux, OS X, or Unix:

```
aws rds copy-db-parameter-group \
--source-db-parameter-group-identifier mygroup1 \
--target-db-parameter-group-identifier mygroup2 \
--target-db-parameter-group-description "DB parameter group 2"
```

For Windows:

```
aws rds copy-db-parameter-group ^
--source-db-parameter-group-identifier mygroup1 ^
--target-db-parameter-group-identifier mygroup2 ^
--target-db-parameter-group-description "DB parameter group 2"
```

RDS API

To copy a DB parameter group, use the RDS API [CopyDBParameterGroup](#) operation with the following required parameters:

- `SourceDBParameterGroupIdentifier`
- `TargetDBParameterGroupIdentifier`
- `TargetDBParameterGroupDescription`

List DB Parameter Groups

You can list the DB parameter groups you've created for your AWS account.

Note

Default parameter groups are automatically created from a default parameter template when you create a DB instance for a particular DB engine and version. These default parameter groups contain preferred parameter settings and can't be modified. When you create a custom parameter group, you can modify parameter settings.

Console

To list all DB parameter groups for an AWS account

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Parameter groups**.

The DB parameter groups appear in a list.

AWS CLI

To list all DB parameter groups for an AWS account, use the AWS CLI [describe-db-parameter-groups](#) command.

Example

The following example lists all available DB parameter groups for an AWS account.

```
aws rds describe-db-parameter-groups
```

The command returns a response like the following:

DBPARAMETERGROUP	default.mysql5.5	mysql5.5	Default parameter group for MySQL5.5
DBPARAMETERGROUP	default.mysql5.6	mysql5.6	Default parameter group for MySQL5.6
DBPARAMETERGROUP	mydbparametergroup	mysql5.6	My new parameter group

The following example describes the *mydbparamgroup1* parameter group.

For Linux, OS X, or Unix:

```
aws rds describe-db-parameter-groups \
--db-parameter-group-name mydbparamgroup1
```

For Windows:

```
aws rds describe-db-parameter-groups ^
--db-parameter-group-name mydbparamgroup1
```

The command returns a response like the following:

```
DBPARAMETERGROUP mydbparametergroup1 mysql5.5 My new parameter group
```

RDS API

To list all DB parameter groups for an AWS account, use the RDS API [DescribeDBParameterGroups](#) operation.

Viewing Parameter Values for a DB Parameter Group

You can get a list of all parameters in a DB parameter group and their values.

Console

To view the parameter values for a DB parameter group

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Parameter groups**.
The DB parameter groups appear in a list.
3. Choose the name of the parameter group to see its list of parameters.

AWS CLI

To view the parameter values for a DB parameter group, use the AWS CLI [describe-db-parameters](#) command with the following required parameter.

- `--db-parameter-group-name`

Example

The following example lists the parameters and parameter values for a DB parameter group named `mydbparametergroup`.

```
aws rds describe-db-parameters --db-parameter-group-name mydbparametergroup
```

The command returns a response like the following:

DBPARAMETER Type	Parameter Name Is Modifiable	Parameter Value	Source	Data Type	Apply
DBPARAMETER	allow-suspicious-udfs		engine-default	boolean	static
DBPARAMETER	auto_increment_increment	true	engine-default	integer	dynamic

DBPARAMETER	auto_increment_offset		engine-default	integer	dynamic
DBPARAMETER	true				
DBPARAMETER	binlog_cache_size	32768	system	integer	dynamic
DBPARAMETER	true				
DBPARAMETER	socket	/tmp/mysql.sock	system	string	static
	false				

RDS API

To view the parameter values for a DB parameter group, use the RDS API [DescribeDBParameters](#) command with the following required parameter.

- `DBParameterGroupName`

Comparing DB Parameter Groups

You can use the AWS Management Console to view the differences between two parameter groups for the same DB engine and version.

To compare two parameter groups

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Parameter groups**.
3. In the list, choose the two parameter groups that you want to compare.
4. For **Parameter group actions**, choose **Compare**.

Note

If the items you selected aren't equivalent, you can't choose **Compare**. For example, you can't compare a MySQL 5.6 and a MySQL 5.7 parameter group. You can't compare a DB parameter group and an Aurora DB cluster parameter group.

DB Parameter Values

You can specify the value for a DB parameter as any of the following:

- An integer constant
- A DB parameter formula
- A DB parameter function
- A character string constant
- A log expression (the log function represents log base 2), such as
`value={log(DBInstanceClassMemory/8187281418)*1000}`

DB Parameter Formulas

A DB parameter formula is an expression that resolves to an integer value or a Boolean value, and is enclosed in braces: `{}`. You can specify formulas for either a DB parameter value or as an argument to a DB parameter function.

Syntax

```
{FormulaVariable}  
{FormulaVariable*Integer}  
{FormulaVariable*Integer/Integer}  
{FormulaVariable/Integer}
```

DB Parameter Formula Variables

Each formula variable returns integer or a Boolean value. The names of the variables are case-sensitive.

AllocatedStorage

Returns the size, in bytes, of the data volume.

DBInstanceClassMemory

Returns the number of bytes of memory allocated to the DB instance class associated with the current DB instance, less the memory used by the Amazon RDS processes that manage the instance.

EndPointPort

Returns the number of the port used when connecting to the DB instance.

DBInstanceClassHugePagesDefault

Returns a Boolean value. Currently, it is only supported for Oracle engines.

For more information, see [Using Huge Pages with an Oracle DB Instance \(p. 935\)](#).

DB Parameter Formula Operators

DB parameter formulas support two operators: division and multiplication.

Division Operator: /

Divides the dividend by the divisor, returning an integer quotient. Decimals in the quotient are truncated, not rounded.

Syntax

```
dividend / divisor
```

The dividend and divisor arguments must be integer expressions.

Multiplication Operator: *

Multiplies the expressions, returning the product of the expressions. Decimals in the expressions are truncated, not rounded.

Syntax

```
expression * expression
```

Both expressions must be integers.

DB Parameter Functions

The parameter arguments can be specified as either integers or formulas. Each function must have at least one argument. Multiple arguments can be specified as a comma-separated list. The list can't have any empty members, such as *argument1,,argument3*. Function names are case-insensitive.

Note

DB Parameter functions are not currently supported in the AWS CLI.

IF()

Returns an argument.

Currently, it is only supported for Oracle engines, and the only supported first argument is `{DBInstanceClassHugePagesDefault}`. For more information, see [Using Huge Pages with an Oracle DB Instance \(p. 935\)](#).

Syntax

```
IF(argument1, argument2, argument3)
```

Returns the second argument if the first argument evaluates to true. Returns the third argument otherwise.

GREATEST()

Returns the largest value from a list of integers or parameter formulas.

Syntax

```
GREATEST(argument1, argument2,...argumentn)
```

Returns an integer.

LEAST()

Returns the smallest value from a list of integers or parameter formulas.

Syntax

```
LEAST(argument1, argument2,...argumentn)
```

Returns an integer.

SUM()

Adds the values of the specified integers or parameter formulas.

Syntax

```
SUM(argument1, argument2,...argumentn)
```

Returns an integer.

DB Parameter Value Examples

These examples show using formulas and functions in the values for DB parameters.

Warning

Improperly setting parameters in a DB parameter group can have unintended adverse effects, including degraded performance and system instability. Always exercise caution when modifying database parameters and back up your data before modifying your DB parameter group. Try out parameter group changes on a test DB instances, created using point-in-time-restores, before applying those parameter group changes to your production DB instances.

You can specify the `GREATEST` function in an Oracle processes parameter to set the number of user processes to the larger of either 80 or `DBInstanceClassMemory` divided by 9,868,951.

```
GREATEST({DBInstanceClassMemory/9868951},80)
```

You can specify the `LEAST()` function in a MySQL `max_binlog_cache_size` parameter value to set the maximum cache size a transaction can use in a MySQL instance to the lesser of 1 MB or `DBInstanceClass/256`.

```
LEAST({DBInstanceClassMemory/256},10485760)
```

Working with Amazon Resource Names (ARNs) in Amazon RDS

Resources created in Amazon Web Services are each uniquely identified with an Amazon Resource Name (ARN). For certain Amazon RDS operations, you must uniquely identify an Amazon RDS resource by specifying its ARN. For example, when you create an RDS DB instance Read Replica, you must supply the ARN for the source DB instance.

Constructing an ARN for Amazon RDS

Resources created in Amazon Web Services are each uniquely identified with an Amazon Resource Name (ARN). You can construct an ARN for an Amazon RDS resource using the following syntax.

`arn:aws:rds:<region>:<account number>:<resourcetype>:<name>`

Region Name	Region	Endpoint	Protocol	
US East (Ohio)	us-east-2	rds.us-east-2.amazonaws.com	HTTPS	
US East (N. Virginia)	us-east-1	rds.us-east-1.amazonaws.com	HTTPS	
US West (N. California)	us-west-1	rds.us-west-1.amazonaws.com	HTTPS	
US West (Oregon)	us-west-2	rds.us-west-2.amazonaws.com	HTTPS	
Asia Pacific (Hong Kong)	ap-east-1	rds.ap-east-1.amazonaws.com	HTTPS	
Asia Pacific (Mumbai)	ap-south-1	rds.ap-south-1.amazonaws.com	HTTPS	
Asia Pacific (Osaka-Local)	ap-northeast-3	rds.ap-northeast-3.amazonaws.com	HTTPS	
Asia Pacific (Seoul)	ap-northeast-2	rds.ap-northeast-2.amazonaws.com	HTTPS	
Asia Pacific (Singapore)	ap-southeast-1	rds.ap-southeast-1.amazonaws.com	HTTPS	
Asia Pacific (Sydney)	ap-southeast-2	rds.ap-southeast-2.amazonaws.com	HTTPS	

Region Name	Region	Endpoint	Protocol	
Asia Pacific (Tokyo)	ap-northeast-1	rds.ap-northeast-1.amazonaws.com	HTTPS	
Canada (Central)	ca-central-1	rds.ca-central-1.amazonaws.com	HTTPS	
China (Beijing)	cn-north-1	rds.cn-north-1.amazonaws.com.cn	HTTPS	
China (Ningxia)	cn-northwest-1	rds.cn-northwest-1.amazonaws.com.cn	HTTPS	
EU (Frankfurt)	eu-central-1	rds.eu-central-1.amazonaws.com	HTTPS	
EU (Ireland)	eu-west-1	rds.eu-west-1.amazonaws.com	HTTPS	
EU (London)	eu-west-2	rds.eu-west-2.amazonaws.com	HTTPS	
EU (Paris)	eu-west-3	rds.eu-west-3.amazonaws.com	HTTPS	
EU (Stockholm)	eu-north-1	rds.eu-north-1.amazonaws.com	HTTPS	
Middle East (Bahrain)	me-south-1	rds.me-south-1.amazonaws.com	HTTPS	
South America (Sao Paulo)	sa-east-1	rds.sa-east-1.amazonaws.com	HTTPS	
AWS GovCloud (US-East)	us-gov-east-1	rds.us-gov-east-1.amazonaws.com	HTTPS	
AWS GovCloud (US-West)	us-gov-west-1	rds.us-gov-west-1.amazonaws.com	HTTPS	

The following table shows the format that you should use when constructing an ARN for a particular Amazon RDS resource type.

Resource Type	ARN Format
DB instance	<p>arn:aws:rds:<region>:<account>:db:<name></p> <p>For example:</p> <pre>arn:aws:rds:us-east-2:123456789012:db:my-mysql-instance-1</pre>

Resource Type	ARN Format
Event subscription	<p>arn:aws:rds:<region>:<account>:es:<name></p> <p>For example:</p> <pre>arn:aws:rds:us-east-2:123456789012:es:my-subscription</pre>
DB option group	<p>arn:aws:rds:<region>:<account>:og:<name></p> <p>For example:</p> <pre>arn:aws:rds:us-east-2:123456789012:og:my-og</pre>
DB parameter group	<p>arn:aws:rds:<region>:<account>:pg:<name></p> <p>For example:</p> <pre>arn:aws:rds:us-east-2:123456789012:pg:my-param-enable-logs</pre>
Reserved DB instance	<p>arn:aws:rds:<region>:<account>:ri:<name></p> <p>For example:</p> <pre>arn:aws:rds:us-east-2:123456789012:ri:my-reserved-postgresql</pre>
DB security group	<p>arn:aws:rds:<region>:<account>:secgrp:<name></p> <p>For example:</p> <pre>arn:aws:rds:us-east-2:123456789012:secgrp:my-public</pre>
Automated DB snapshot	<p>arn:aws:rds:<region>:<account>:snapshot:rds:<name></p> <p>For example:</p> <pre>arn:aws:rds:us-east-2:123456789012:snapshot:rds:my-mysql-db-2019-07-22-07-23</pre>
Manual DB snapshot	<p>arn:aws:rds:<region>:<account>:snapshot:<name></p> <p>For example:</p> <pre>arn:aws:rds:us-east-2:123456789012:snapshot:my-mysql-db-snap</pre>
DB subnet group	<p>arn:aws:rds:<region>:<account>:subgrp:<name></p> <p>For example:</p> <pre>arn:aws:rds:us-east-2:123456789012:subgrp:my-subnet-10</pre>

Getting an Existing ARN

You can get the ARN of an RDS resource by using the AWS Management Console, AWS Command Line Interface (AWS CLI), or RDS API.

Console

To get an ARN from the AWS Management Console, navigate to the resource you want an ARN for, and view the details for that resource. For example, you can get the ARN for a DB instance from the **Configuration** tab of the DB instance details, as shown following.

The screenshot shows the 'Configuration' tab selected in the top navigation bar. Below it, the 'Instance' section displays various configuration details. A red oval highlights the 'Configuration' tab. Another red oval highlights the 'ARN' field, which contains the value 'arn:aws:rds:us-west-2:XXXXXXXXXX:db:oracle-instance1'. The 'Resource id' field below it contains the value 'oracle-instance1'.

Connectivity Monitoring Logs & events Configuration

Instance

Configuration

DB instance id
oracle-instance1

Engine version
12.1.0.2.v14

Storage type
General Purpose (SSD)

IOPS
-

Storage
20 GiB

DB name
ORCL

License model
Bring Your Own License

Character set
AL32UTF8

Option groups
default:oracle-ee-12-1

ARN
arn:aws:rds:us-west-2:XXXXXXXXXX:db:oracle-instance1

Resource id
oracle-instance1

AWS CLI

To get an ARN from the AWS CLI for a particular RDS resource, you use the `describe` command for that resource. The following table shows each AWS CLI command, and the ARN property used with the command to get an ARN.

AWS CLI Command	ARN Property
<code>describe-event-subscriptions</code>	<code>EventSubscriptionArn</code>
<code>describe-certificates</code>	<code>CertificateArn</code>
<code>describe-db-parameter-groups</code>	<code>DBParameterGroupArn</code>
<code>describe-db-instances</code>	<code>DBInstanceArn</code>
<code>describe-db-security-groups</code>	<code>DBSecurityGroupArn</code>
<code>describe-db-snapshots</code>	<code>DBSnapshotArn</code>
<code>describe-events</code>	<code>SourceArn</code>
<code>describe-reserved-db-instances</code>	<code>ReservedDBInstanceArn</code>
<code>describe-db-subnet-groups</code>	<code>DBSubnetGroupArn</code>
<code>describe-option-groups</code>	<code>OptionGroupArn</code>

For example, the following AWS CLI command gets the ARN for a DB instance.

Example

For Linux, OS X, or Unix:

```
aws rds describe-db-instances \
--db-instance-identifier DBInstanceIdentifier \
--region us-west-2
```

For Windows:

```
aws rds describe-db-instances ^
--db-instance-identifier DBInstanceIdentifier ^
--region us-west-2
```

RDS API

To get an ARN for a particular RDS resource, you can call the following RDS API operations and use the ARN properties shown following.

RDS API Operation	ARN Property
<code>DescribeEventSubscriptions</code>	<code>EventSubscriptionArn</code>
<code>DescribeCertificates</code>	<code>CertificateArn</code>
<code>DescribeDBParameterGroups</code>	<code>DBParameterGroupArn</code>

RDS API Operation	ARN Property
DescribeDBInstances	DBInstanceArn
DescribeDBSecurityGroups	DBSecurityGroupArn
DescribeDBSchemas	DBSnapshotArn
DescribeEvents	SourceArn
DescribeReservedDBInstances	ReservedDBInstanceArn
DescribeDBSubnetGroups	DBSubnetGroupArn
DescribeOptionGroups	OptionGroupArn

Working with Storage for Amazon RDS DB Instances

To specify how you want your data stored in Amazon RDS, choose a storage type and provide a storage size when you create or modify a DB instance. Later, you can increase the amount or change the type of storage by modifying the DB instance. For more information about which storage type to use for your workload, see [Amazon RDS Storage Types \(p. 135\)](#).

Topics

- [Increasing DB Instance Storage Capacity \(p. 231\)](#)
- [Managing Capacity Automatically with Amazon RDS Storage Autoscaling \(p. 232\)](#)
- [Modifying SSD Storage Settings for Provisioned IOPS \(p. 236\)](#)

Increasing DB Instance Storage Capacity

If you need space for additional data, you can scale up the storage of an existing DB instance. To do so, you can use the Amazon RDS Management Console, the Amazon RDS API, or the AWS Command Line Interface (AWS CLI). For information about storage limits, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

Note

Scaling storage for Amazon RDS for Microsoft SQL Server DB instances is supported only for General Purpose SSD or Provisioned IOPS SSD storage types.

To monitor the amount of free storage for your DB instance so you can respond when necessary, we recommend that you create an Amazon CloudWatch alarm. For more information on setting CloudWatch alarms, see [Using CloudWatch Alarms](#).

In most cases, scaling storage doesn't require any outage and doesn't degrade performance of the server. After you modify the storage size for a DB instance, the status of the DB instance is **storage-optimization**. The DB instance is fully operational after a storage modification.

Note

You can't make further storage modifications until six (6) hours after storage optimization has completed on the instance.

However, a special case is if you have a SQL Server DB instance and haven't modified the storage configuration since November 2017. In this case, you might experience a short outage of a few minutes when you modify your DB instance to increase the allocated storage. After the outage, the DB instance is online but in the **storage-optimization** state. Performance might be degraded during storage optimization.

Note

You can't reduce the amount of storage for a DB instance after storage has been allocated.

Console

To increase storage for a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the DB instance that you want to modify.
4. Choose **Modify**.

5. Enter a new value for **Allocated Storage**. It must be greater than the current value.

The screenshot shows the AWS RDS console's storage configuration page. At the top, it says "Storage type: General Purpose (SSD)". Below that, the "Allocated storage" field is set to "16384 GiB". A note below the field states: "This instance supports multiple storage ranges between 20 and 16384 GiB. [See all](#)". A warning box contains the following text:
⚠ Scaling your instance storage can:

- Deplete the initial General Purpose (SSD) I/O credits, leading to longer conversion times. [Learn more](#)
- Impact instance performance until operation completes. [Learn more](#)

Note

When you increase **Allocated Storage**, it must be by at least 10 percent. If you try to increase the value by less than 10 percent, you see an error.

6. Choose **Continue** to move to the next screen.
7. Choose **Apply immediately** in the **Scheduling of modifications** section to apply the storage changes to the DB instance immediately. Or choose **Apply during the next scheduled maintenance window** to apply the changes during the next maintenance window.
8. When the settings are as you want them, choose **Modify DB instance**.

AWS CLI

To increase the storage for a DB instance, use the AWS CLI command `modify-db-instance`. Set the following parameters:

- `--allocated-storage` – Amount of storage to be allocated for the DB instance, in gibibytes.
- `--apply-immediately` – Use `--apply-immediately` to change to the new storage type immediately. Or use `--no-apply-immediately` (the default) to apply storage changes during the next maintenance window. An immediate outage occurs when the changes are applied.

For more information about storage, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

Amazon RDS API

To increase storage for a DB instance, use the Amazon RDS API operation `ModifyDBInstance`. Set the following parameters:

- `AllocatedStorage` – Amount of storage to be allocated for the DB instance, in gibibytes.
- `ApplyImmediately` – Set this option to `True` to apply scaling changes immediately. Set this option to `False` (the default) to apply scaling changes during the next maintenance window. An immediate outage occurs when the changes are applied.

For more information about storage, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

Managing Capacity Automatically with Amazon RDS Storage Autoscaling

If your workload is unpredictable, you can enable storage autoscaling for an Amazon RDS DB instance. To do so, you can use the Amazon RDS console, the Amazon RDS API, or the AWS CLI.

For example, you might use this feature for a new mobile gaming application that users are adopting rapidly. In this case, a rapidly increasing workload might exceed the available database storage. To avoid having to manually scale up database storage, you can use Amazon RDS storage autoscaling.

With storage autoscaling enabled, when Amazon RDS detects that you are running out of free database space it automatically scales up your storage. Amazon RDS starts a storage modification for an autoscaling-enabled DB instance when these factors apply:

- Free available space is less than 10 percent of the allocated storage.
- The low-storage condition lasts at least five minutes.
- At least six hours have passed since the last storage modification.

The additional storage is in increments of whichever is greater, 5 GiB or 12% of currently allocated storage. The maximum storage threshold is the limit to which the DB instance can be autoscaled. You can't set the maximum storage threshold for autoscaling-enabled instances to a value greater than the maximum allocated storage.

For example, SQL Server Standard Edition on db.m5.xlarge has a default allocated storage for the instance of 20 GiB (the minimum) and a maximum allocated storage of 16,384 GiB. The default maximum storage threshold for autoscaling is 1,000 GiB. If you use this default, the instance doesn't autoscale above 1,000 GiB. This is true even though the maximum allocated storage for the instance is 16,384 GiB.

Note

Autoscaling doesn't occur if the maximum storage threshold would be exceeded by the storage increment.

Autoscaling can't completely prevent storage-full situations for large data loads, because further storage modifications can't be made until six hours after storage optimization has completed on the instance. If you perform a large data load, and autoscaling doesn't provide enough space, the database might remain in the storage-full state for several hours. This can harm the database.

If you start a storage scaling operation at the same time that Amazon RDS starts an autoscaling operation, your storage modification takes precedence. The autoscaling operation is canceled.

Although automatic scaling helps you to increase storage on your Amazon RDS DB instance dynamically, you should still configure the initial storage for your DB instance to an appropriate size for your typical workload.

Enabling Storage Autoscaling for a New DB Instance

When you create a new Amazon RDS DB instance, you can choose whether to enable storage autoscaling. You can also set an upper limit on the storage that Amazon RDS can allocate for the DB instance.

Note

When you clone an Amazon RDS DB instance that has storage autoscaling enabled, that setting isn't automatically inherited by the cloned instance. The new DB instance has the same amount of allocated storage as the original instance. You can turn storage autoscaling on again for the new instance if the cloned instance continues to increase its storage requirements.

Console

To enable storage autoscaling for a new DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the Amazon RDS console, choose the AWS Region where you want to create the DB instance.

3. In the navigation pane, choose **Databases**.
4. Choose **Create database**. On the **Select engine** page, choose your database engine and specify your DB instance information as described in [Getting Started with Amazon RDS \(p. 11\)](#).
5. In the **Storage autoscaling** section, set the **Maximum storage threshold** value for the DB instance.
6. Specify the rest of your DB instance information as described in [Getting Started with Amazon RDS \(p. 11\)](#).

AWS CLI

To enable storage autoscaling for a new DB instance, use the AWS CLI command `create-db-instance`. Set the following parameter:

- `--max-allocated-storage` – Turns on storage autoscaling and sets the upper limit on storage size, in gibibytes.

To verify that Amazon RDS storage autoscaling is available for your DB instance, use the AWS CLI `describe-valid-db-instance-modifications` command. To check based on the instance class before creating an instance, use the `describe-orderable-db-instance-options` command. Check the following field in the return value:

- `SupportsStorageAutoscaling` – Indicates whether the DB instance or instance class supports storage autoscaling.

For more information about storage, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

Amazon RDS API

To enable storage autoscaling for a new DB instance, use the Amazon RDS API operation `CreateDBInstance`. Set the following parameter:

- `MaxAllocatedStorage` – Turns on Amazon RDS storage autoscaling and sets the upper limit on storage size, in gibibytes.

To verify that Amazon RDS storage autoscaling is available for your DB instance, use the Amazon RDS API `DescribeValidDbInstanceModifications` operation for an existing instance, or the `DescribeOrderableDBInstanceOptions` operation before creating an instance. Check the following field in the return value:

- `SupportsStorageAutoscaling` – Indicates whether the DB instance supports storage autoscaling.

For more information about storage, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

Changing the Storage Autoscaling Settings for a DB Instance

You can turn storage autoscaling on for an existing Amazon RDS DB instance. You can also change the upper limit on the storage that Amazon RDS can allocate for the DB instance.

Console

To change the storage autoscaling settings for a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.

2. In the navigation pane, choose **Databases**.
3. Choose the DB instance that you want to modify, and choose **Modify**. The **Modify DB instance** page appears.
4. Change the storage limit in the **Autoscaling** section. For more information, see [Modifying an Amazon RDS DB Instance \(p. 150\)](#).
5. When all the changes are as you want them, choose **Continue** and check your modifications.
6. On the confirmation page, review your changes. If they're correct, choose **Modify DB Instance** to save your changes. If they aren't correct, choose **Back** to edit your changes or **Cancel** to cancel your changes.

Changing the storage autoscaling limit occurs immediately. This setting ignores the **Apply immediately** setting.

AWS CLI

To change the storage autoscaling settings for a DB instance, use the AWS CLI command `modify-db-instance`. Set the following parameter:

- `--max-allocated-storage` – Sets the upper limit on storage size, in gibibytes. If the value is greater than the `--allocated-storage` parameter, storage autoscaling is turned on. If the value is the same as the `--allocated-storage` parameter, storage autoscaling is turned off.

To verify that Amazon RDS storage autoscaling is available for your DB instance, use the AWS CLI `describe-valid-db-instance-modifications` command. To check based on the instance class before creating an instance, use the `describe-orderable-db-instance-options` command. Check the following field in the return value:

- `SupportsStorageAutoscaling` – Indicates whether the DB instance supports storage autoscaling.

For more information about storage, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

Amazon RDS API

To change the storage autoscaling settings for a DB instance, use the Amazon RDS API operation `ModifyDBInstance`. Set the following parameter:

- `MaxAllocatedStorage` – Sets the upper limit on storage size, in gibibytes.

To verify that Amazon RDS storage autoscaling is available for your DB instance, use the Amazon RDS API `DescribeValidDbInstanceModifications` operation for an existing instance, or the `DescribeOrderableDBInstanceOptions` operation before creating an instance. Check the following field in the return value:

- `SupportsStorageAutoscaling` – Indicates whether the DB instance supports storage autoscaling.

For more information about storage, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

Turning Off Storage Autoscaling for a DB Instance

If you no longer need Amazon RDS to automatically increase the storage for an Amazon RDS DB instance, you can turn off storage autoscaling. After you do, you can still manually increase the amount of storage for your DB instance.

Console

To turn off storage autoscaling for a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the DB instance that you want to modify and choose **Modify**. The **Modify DB instance** page appears.
4. Clear the **Enable storage autoscaling** check box in the **Storage autoscaling** section. For more information, see [Modifying an Amazon RDS DB Instance \(p. 150\)](#).
5. When all the changes are as you want them, choose **Continue** and check the modifications.
6. On the confirmation page, review your changes. If they're correct, choose **Modify DB Instance** to save your changes. If they aren't correct, choose **Back** to edit your changes or **Cancel** to cancel your changes.

Changing the storage autoscaling limit occurs immediately. This setting ignores the **Apply immediately** setting.

AWS CLI

To turn off storage autoscaling for a DB instance, use the AWS CLI command `modify-db-instance` and the following parameter:

- `--max-allocated-storage` – Specify a value equal to the `--allocated-storage` setting to prevent further Amazon RDS storage autoscaling for the specified DB instance.

For more information about storage, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

Amazon RDS API

To turn off storage autoscaling for a DB instance, use the Amazon RDS API operation `ModifyDBInstance`. Set the following parameter:

- `MaxAllocatedStorage` – Specify a value equal to the `AllocatedStorage` setting to prevent further Amazon RDS storage autoscaling for the specified DB instance.

For more information about storage, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

Modifying SSD Storage Settings for Provisioned IOPS

You can modify the settings for a DB instance that uses Provisioned IOPS SSD storage by using the Amazon RDS console, AWS CLI, or Amazon RDS API. Specify the storage type, allocated storage, and the amount of Provisioned IOPS that you require. You can choose from a range between 1,000 IOPS and 100 GiB of storage up to 80,000 IOPS and 64 TiB (64,000 GiB) of storage. The range depends on your database engine and instance type.

Although you can reduce the amount of IOPS provisioned for your instance, you can't reduce the amount of General Purpose SSD or magnetic storage allocated.

In most cases, scaling storage doesn't require any outage and doesn't degrade performance of the server. After you modify the storage size for a DB instance, the status of the DB instance is **storage-optimization**. The DB instance is fully operational after a storage modification.

Note

You can't make further storage modifications until six (6) hours after storage optimization has completed on the instance.

Console

To change the Provisioned IOPS settings for a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.

Note

To filter the list of DB instances, for **Filter databases** enter a text string for Amazon RDS to use to filter the results. Only DB instances whose names contain the string appear.

3. Choose the DB instance with Provisioned IOPS that you want to modify.
4. Choose **Modify**.
5. On the **Modify DB Instance page**, choose Provisioned IOPS for **Storage type** and then provide a Provisioned IOPS value.

The screenshot shows the 'Modify DB Instance' page with the following settings:

- Storage type:** Provisioned IOPS (SSD)
- Allocated storage:** 16384 GiB (Minimum: 100 GiB, Maximum: 16384)
- Provisioned IOPS:** 80000

If the value you specify for either **Allocated storage** or **Provisioned IOPS** is outside the limits supported by the other parameter, a warning message is displayed. This message gives the range of values required for the other parameter.

6. Choose **Continue**.
7. To apply the changes to the DB instance immediately, choose **Apply immediately** in the **Scheduling of modifications** section. Or choose **Apply during the next scheduled maintenance window** to apply the changes during the next maintenance window.

An immediate outage occurs when the storage type changes. For more information about storage, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

8. Review the parameters to be changed, and choose **Modify DB instance** to complete the modification.

The new value for allocated storage or for Provisioned IOPS appears in the **Status** column.

AWS CLI

To change the Provisioned IOPS setting for a DB instance, use the AWS CLI command `modify-db-instance`. Set the following parameters:

- `--storage-type` – Set to `io1` for Provisioned IOPS.
- `--allocated-storage` – Amount of storage to be allocated for the DB instance, in gibibytes.
- `--iops` – The new amount of Provisioned IOPS for the DB instance, expressed in I/O operations per second.

- `--apply-immediately` – Use `--apply-immediately` to apply changes immediately. Use `--no-apply-immediately` (the default) to apply changes during the next maintenance window.

Amazon RDS API

To change the Provisioned IOPS settings for a DB instance, use the Amazon RDS API operation [ModifyDBInstance](#). Set the following parameters:

- `StorageType` – Set to `io1` for Provisioned IOPS.
- `AllocatedStorage` – Amount of storage to be allocated for the DB instance, in gibibytes.
- `Iops` – The new IOPS rate for the DB instance, expressed in I/O operations per second.
- `ApplyImmediately` – Set this option to `True` to apply changes immediately. Set this option to `False` (the default) to apply changes during the next maintenance window.

DB Instance Billing for Amazon RDS

Amazon RDS instances are billed based on the following components:

- DB instance hours (per hour) – Based on the DB instance class of the DB instance (for example, db.t2.small or db.m4.large). Pricing is listed on a per-hour basis, but bills are calculated down to the second and show times in decimal form. RDS usage is billed in one second increments, with a minimum of 10 minutes. For more information, see [Choosing the DB Instance Class \(p. 107\)](#).
- Storage (per GiB per month) – Storage capacity that you have provisioned to your DB instance. If you scale your provisioned storage capacity within the month, your bill is pro-rated. For more information, see [Amazon RDS DB Instance Storage \(p. 135\)](#).
- I/O requests (per 1 million requests per month) – Total number of storage I/O requests that you have made in a billing cycle, for Amazon RDS magnetic storage only.
- Provisioned IOPS (per IOPS per month) – Provisioned IOPS rate, regardless of IOPS consumed, for Amazon RDS Provisioned IOPS (SSD) storage only. Provisioned storage for EBS volumes are billed in one second increments, with a minimum of 10 minutes.
- Backup storage (per GiB per month) – *Backup storage* is the storage that is associated with automated database backups and any active database snapshots that you have taken. Increasing your backup retention period or taking additional database snapshots increases the backup storage consumed by your database. Per second billing doesn't apply to backup storage (metered in GB-month).

For more information, see [Backing Up and Restoring Amazon RDS DB Instances \(p. 253\)](#).

- Data transfer (per GB) – Data transfer in and out of your DB instance from or to the internet and other AWS Regions.

Amazon RDS provides the following purchasing options to enable you to optimize your costs based on your needs:

- **On-Demand Instances** – Pay by the hour for the DB instance hours that you use. Pricing is listed on a per-hour basis, but bills are calculated down to the second and show times in decimal form. RDS usage is now billed in one second increments, with a minimum of 10 minutes.
- **Reserved Instances** – Reserve a DB instance for a one-year or three-year term and get a significant discount compared to the on-demand DB instance pricing. With Reserved Instance usage, you can launch, delete, start, or stop multiple instances within an hour and get the Reserved Instance benefit for all of the instances.

For Amazon RDS pricing information, see the [Amazon RDS product page](#).

Topics

- [On-Demand DB Instances for Amazon RDS \(p. 240\)](#)
- [Reserved DB Instances for Amazon RDS \(p. 241\)](#)

On-Demand DB Instances for Amazon RDS

Amazon RDS on-demand DB instances are billed based on the class of the DB instance (for example, db.t2.small or db.m4.large). Partial DB instance hours consumed are billed as full hours. For Amazon RDS pricing information, see the [Amazon RDS product page](#).

Billing starts for a DB instance as soon as the DB instance is available. Pricing is listed on a per-hour basis, but bills are calculated down to the second and show times in decimal form. Amazon RDS usage is billed in one second increments, with a minimum of 10 minutes. In the case of billable configuration change, such as scaling compute or storage capacity, you're charged a 10-minute minimum. Billing continues until the DB instance terminates, which occurs when you delete the DB instance or if the DB instance fails.

If you no longer want to be charged for your DB instance, you must stop or delete it to avoid being billed for additional DB instance hours. For more information about the DB instance states for which you are billed, see [DB Instance Status \(p. 129\)](#).

Stopped DB Instances

While your DB instance is stopped, you're charged for provisioned storage, including Provisioned IOPS. You are also charged for backup storage, including storage for manual snapshots and automated backups within your specified retention window. You aren't charged for DB instance hours.

Multi-AZ DB Instances

If you specify that your DB instance should be a Multi-AZ deployment, you're billed according to the Multi-AZ pricing posted on the Amazon RDS pricing page.

Reserved DB Instances for Amazon RDS

Using reserved DB instances, you can reserve a DB instance for a one- or three-year term. Reserved DB instances provide you with a significant discount compared to on-demand DB instance pricing. Reserved DB instances are not physical instances, but rather a billing discount applied to the use of certain on-demand DB instances in your account. Discounts for reserved DB instances are tied to instance type and AWS Region.

The general process for working with reserved DB instances is: First get information about available reserved DB instance offerings, then purchase a reserved DB instance offering, and finally get information about your existing reserved DB instances.

Overview of Reserved DB Instances

When you purchase a reserved DB instance in Amazon RDS, you purchase a commitment to getting a discounted rate, on a specific DB instance type, for the duration of the reserved DB instance. To use an Amazon RDS reserved DB instance, you create a new DB instance just like you do for an on-demand instance. The new DB instance that you create must match the specifications of the reserved DB instance. If the specifications of the new DB instance match an existing reserved DB instance for your account, you are billed at the discounted rate offered for the reserved DB instance. Otherwise, the DB instance is billed at an on-demand rate.

For more information about reserved DB instances, including pricing, see [Amazon RDS Reserved Instances](#).

Offering Types

Reserved DB instances are available in three varieties—No Upfront, Partial Upfront, and All Upfront—that let you optimize your Amazon RDS costs based on your expected usage.

No Upfront

This option provides access to a reserved DB instance without requiring an upfront payment. Your No Upfront reserved DB instance bills a discounted hourly rate for every hour within the term, regardless of usage, and no upfront payment is required. This option is only available as a one-year reservation.

Partial Upfront

This option requires a part of the reserved DB instance to be paid upfront. The remaining hours in the term are billed at a discounted hourly rate, regardless of usage. This option is the replacement for the previous Heavy Utilization option.

All Upfront

Full payment is made at the start of the term, with no other costs incurred for the remainder of the term regardless of the number of hours used.

If you are using consolidated billing, all the accounts in the organization are treated as one account. This means that all accounts in the organization can receive the hourly cost benefit of reserved DB instances that are purchased by any other account. For more information about consolidated billing, see [Amazon RDS Reserved DB Instances](#) in the *AWS Billing and Cost Management User Guide*.

Size-Flexible Reserved DB Instances

When you purchase a reserved DB instance, one thing that you specify is the instance class, for example db.m4.large. For more information about instance classes, see [Choosing the DB Instance Class \(p. 107\)](#).

If you have a DB instance, and you need to scale it to larger capacity, your reserved DB instance is automatically applied to your scaled DB instance. That is, your reserved DB instances are automatically applied across all DB instance class sizes. Size-flexible reserved DB instances are available for DB instances with the same AWS Region and database engine. Size-flexible reserved DB instances can only scale in their instance class type. For example, a reserved DB instance for a db.m4.large can apply to a db.m4.xlarge, but not to a db.m5.large, because db.m4 and db.m5 are different instance class types. Reserved DB instance benefits also apply for both Multi-AZ and Single-AZ configurations.

Size-flexible reserved DB instances are available for the following Amazon RDS database engines:

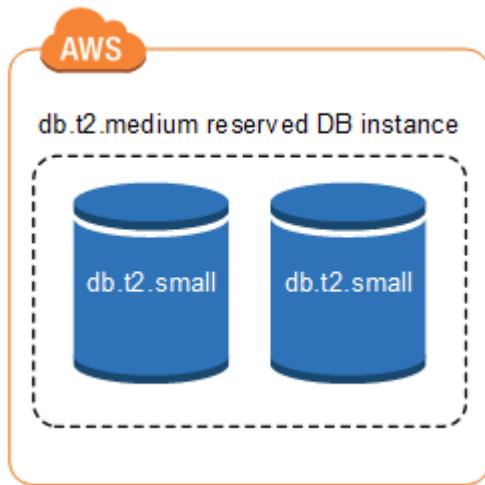
- MariaDB
- MySQL
- Oracle, Bring Your Own License
- PostgreSQL

For details about using size-flexible reserved instances with Aurora, see [Reserved DB Instances for Aurora](#).

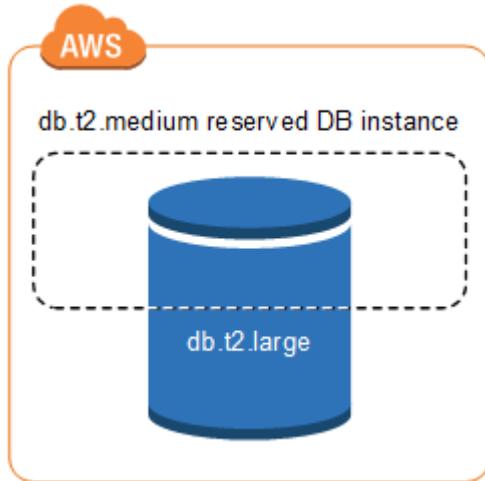
You can compare usage for different reserved DB instance sizes by using normalized units. For example, one unit of usage on two db.m3.large DB instances is equivalent to eight normalized units of usage on one db.m3.small. The following table shows the number of normalized units for each DB instance size.

Instance Size	Single-AZ Normalized Units	Multi-AZ Normalized Units
micro	0.5	1
small	1	2
medium	2	4
large	4	8
xlarge	8	16
2xlarge	16	32
4xlarge	32	64
8xlarge	64	128
10xlarge	80	160
16xlarge	128	256

For example, suppose that you purchase a db.t2.medium reserved DB instance, and you have two running db.t2.small DB instances in your account in the same AWS Region. In this case, the billing benefit is applied in full to both instances.



Alternatively, if you have one db.t2.large instance running in your account in the same AWS Region, the billing benefit is applied to 50 percent of the usage of the DB instance.



Reserved DB Instance Billing Example

The price for a reserved DB instance doesn't include regular costs associated with storage, backups, and I/O. The following example illustrates the total cost per month for a reserved DB instance:

- An Amazon RDS MySQL reserved Single-AZ db.r4.large DB instance class in US East (N. Virginia) with the No Upfront option at a cost of \$0.12 for the instance, or \$90 per month
- 400 GiB of General Purpose SSD (gp2) storage at a cost of 0.115 per GiB per month, or \$45.60 per month
- 600 GiB of backup storage at \$0.095, or \$19 per month (400 GiB free)

Add all of these options (\$90 + \$45.60 + \$19) with the reserved DB instance, and the total cost per month is \$154.60.

If you chose to use an on-demand DB instance instead of a reserved DB instance, an Amazon RDS MySQL Single-AZ db.r4.large DB instance class in US East (N. Virginia) costs \$0.1386 per hour, or \$101.18 per month. So, for an on-demand DB instance, add all of these options (\$101.18 + \$45.60 + \$19), and the total cost per month is \$165.78.

Note

The prices in this example are sample prices and might not match actual prices.
For Amazon RDS pricing information, see the [Amazon RDS product page](#).

Deleting a Reserved DB Instance

The terms for a reserved DB instance involve a one-year or three-year commitment. You can't cancel a reserved DB instance. However, you can delete a DB instance that is covered by a reserved DB instance discount. The process for deleting a DB instance that is covered by a reserved DB instance discount is the same as for any other DB instance.

Your upfront payment for a reserved DB instance reserves the resources for your use. Because these resources are reserved for you, you are billed for the resources regardless of whether you use them.

If you delete a DB instance that is covered by a reserved DB instance discount, you can launch another DB instance with compatible specifications. In this case, you continue to get the discounted rate during the reservation term (one or three years).

Console

You can use the AWS Management Console to work with reserved DB instances as shown in the following procedures.

To get pricing and information about available reserved DB instance offerings

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Reserved instances**.
3. Choose **Purchase Reserved DB Instance**.
4. For **Product description**, choose the DB engine and licensing type.
5. For **DB instance class**, choose the DB instance class.
6. For **Multi-AZ deployment**, choose whether you want a Multi-AZ deployment.
7. For **Term**, choose the length of time you want the DB instance reserved.
8. For **Offering type**, choose the offering type.

After you select the offering type, you can see the pricing information.

Important

Choose **Cancel** to avoid purchasing the reserved DB instance and incurring any charges.

After you have information about the available reserved DB instance offerings, you can use the information to purchase an offering as shown in the following procedure.

To purchase a reserved DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Reserved instances**.
3. Choose **Purchase Reserved DB Instance**.
4. For **Product description**, choose the DB engine and licensing type.
5. For **DB instance class**, choose the DB instance class.
6. For **Multi-AZ deployment**, choose whether you want a Multi-AZ deployment.
7. For **Term**, choose the length of time you want the DB instance reserved.
8. For **Offering type**, choose the offering type.

After you choose the offering type, you can see the pricing information, as shown following.

Purchase Reserved DB Instances

Choose from the options below, then enter the number of DB instances you wish to reserve with this order. When you are done, click the Continue button.

Options

Product description

aurora-mysql

DB instance class

db.r4.4xlarge — 16 vCPU, 122 GiB RAM

Multi AZ deployment

Multi-AZ deployment model is not applicable for this database engine and edition

- Yes
 No

Term

1 year

Offering type

All Upfront

Reserved Id (optional)

Optional tag to track your reservation

Number of DB instances

1

Pricing details

One-time payment (per instance)

Total one-time payment*

*Additional taxes may apply

Normalized units per hour [info](#)

32

Usage charges*

USD (hourly)

*Additional taxes may apply

This hourly rate is charged for every hour for each instance in the Reserved Instance term you purchase, regardless of instance usage

Charges for your usage will appear on your monthly bill.

[Cancel](#)

[Continue](#)

9. (Optional) You can assign your own identifier to the reserved DB instances that you purchase to help you track them. For **Reserved Id**, type an identifier for your reserved DB instance.
10. Choose **Continue**.

The **Purchase Reserved DB Instance** dialog box appears, with a summary of the reserved DB instance attributes that you've selected and the payment due, as shown following.

Purchase Reserved DB Instances

Summary of Purchase
You are about to purchase a Reserved DB Instance with the following information.

Region	US East (N. Virginia)
Product Description	aurora-mysql
DB Instance Class	db.r4.4xlarge
Offering Type	All Upfront
Multi AZ Deployment	No
Term	1 year
Reserved DB Instance	default
Quantity	1
Price Per Instance	[REDACTED]
Total Payment Due Now	[REDACTED]

⚠️ Purchasing this Reserved DB Instance will charge [REDACTED] to the payment method associated with this Amazon Web Services account. Are you sure you would like to proceed?

Cancel **Back** **Purchase**

11. On the confirmation page, review your reserved DB instance. If the information is correct, choose **Purchase** to purchase the reserved DB instance.

Alternatively, choose **Back** to edit your reserved DB instance.

After you have purchased reserved DB instances, you can get information about your reserved DB instances as shown in the following procedure.

To get information about reserved DB instances for your AWS account

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the **Navigation** pane, choose **Reserved instances**.

The reserved DB instances for your account appear. To see detailed information about a particular reserved DB instance, choose that instance in the list. You can then see detailed information about that instance in the detail pane at the bottom of the console.

AWS CLI

You can use the AWS CLI to work with reserved DB instances as shown in the following examples.

Example Get Available Reserved DB Instance Offerings

To get information about available reserved DB instance offerings, call the AWS CLI command `describe-reserved-db-instances-offerings`.

```
aws rds describe-reserved-db-instances-offerings
```

This call returns output similar to the following:

OFFERING	OfferingId	Class	Multi-AZ	Duration	Fixed
Price	Usage Price	Description	Offering Type		
OFFERING	438012d3-4052-4cc7-b2e3-8d3372e0e706	db.m1.large	y	1y	1820.00
USD	0.368	mysql	Partial Upfront		
OFFERING	649fd0c8-cf6d-47a0-bfa6-060f8e75e95f	db.m1.small	n	1y	227.50
USD	0.046	mysql	Partial Upfront		
OFFERING	123456cd-ab1c-47a0-bfa6-12345667232f	db.m1.small	n	1y	162.00
USD	0.00	mysql	All Upfront		
Recurring Charges:	Amount	Currency	Frequency		
Recurring Charges:	0.123	USD	Hourly		
OFFERING	123456cd-ab1c-37a0-bfa6-12345667232d	db.m1.large	y	1y	700.00
USD	0.00	mysql	All Upfront		
Recurring Charges:	Amount	Currency	Frequency		
Recurring Charges:	1.25	USD	Hourly		
OFFERING	123456cd-ab1c-17d0-bfa6-12345667234e	db.m1.xlarge	n	1y	4242.00
USD	2.42	mysql	No Upfront		

After you have information about the available reserved DB instance offerings, you can use the information to purchase an offering as shown in the following example.

Example Purchase a Reserved DB Instance

To purchase a reserved DB instance, use the AWS CLI command `purchase-reserved-db-instances-offering` with the following parameters:

- **--reserved-db-instances-offering-id** – the id of the offering that you want to purchase. See the preceding example to get the offering ID.
- **--reserved-db-instance-id** – you can assign your own identifier to the reserved DB instances that you purchase to help you track them.

The following example purchases the reserved DB instance offering with ID ***649fd0c8-cf6d-47a0-bfa6-060f8e75e95f***, and assigns the identifier of ***MyReservation***.

For Linux, OS X, or Unix:

```
aws rds purchase-reserved-db-instances-offering \
--reserved-db-instances-offering-id 649fd0c8-cf6d-47a0-bfa6-060f8e75e95f \
--reserved-db-instance-id MyReservation
```

For Windows:

```
aws rds purchase-reserved-db-instances-offering ^
--reserved-db-instances-offering-id 649fd0c8-cf6d-47a0-bfa6-060f8e75e95f ^
--reserved-db-instance-id MyReservation
```

The command returns output similar to the following:

RESERVATION	ReservationId	Class	Multi-AZ	Start Time	Duration
Fixed Price	Usage Price	Count	State	Description	Offering Type
RESERVATION	MyReservation	db.m1.small	y	2011-12-19T00:30:23.247Z	1y
455.00 USD	0.092 USD	1	payment-pending	mysql	Partial Upfront

After you have purchased reserved DB instances, you can get information about your reserved DB instances as shown in the following example.

Example Get Your Reserved DB Instances

To get information about reserved DB instances for your AWS account, call the AWS CLI command [describe-reserved-db-instances](#).

```
aws rds describe-reserved-db-instances
```

The command returns output similar to the following:

RESERVATION	ReservationId	Class	Multi-AZ	Start Time	Duration
Fixed Price	Usage Price	Count	State	Description	Offering Type
RESERVATION	MyReservation	db.m1.small	y	2011-12-09T23:37:44.720Z	1y
455.00 USD	0.092 USD	1	retired	mysql	Partial Upfront

RDS API

You can use the RDS API to work with reserved DB instances as shown in the following examples.

Example Get Available Reserved DB Instance Offerings

To get information about available reserved DB instance offerings, call the Amazon RDS API function [DescribeReservedDBInstancesOfferings](#).

```
https://rds.us-east-1.amazonaws.com/
?Action=DescribeReservedDBInstancesOfferings
&SignatureMethod=HmacSHA256
&SignatureVersion=4
```

```
&Version=2014-09-01
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140411/us-east-1/rds/aws4_request
&X-Amz-Date=20140411T203327Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=545f04acffeb4b80d2e778526b1c9da79d0b3097151c24f28e83e851d65422e2
```

This call returns output similar to the following:

```
<DescribeReservedDBInstancesOfferingsResponse xmlns="http://rds.amazonaws.com/doc/2014-10-31/">
  <DescribeReservedDBInstancesOfferingsResult>
    <ReservedDBInstancesOfferings>
      <ReservedDBInstancesOffering>
        <Duration>31536000</Duration>
        <OfferingType>Partial Upfront</OfferingType>
        <CurrencyCode>USD</CurrencyCode>
        <RecurringCharges/>
        <FixedPrice>1820.0</FixedPrice>
        <ProductDescription>mysql</ProductDescription>
        <UsagePrice>0.368</UsagePrice>
        <MultiAZ>true</MultiAZ>
        <ReservedDBInstancesOfferingId>438012d3-4052-4cc7-b2e3-8d3372e0e706</
      ReservedDBInstancesOfferingId>
        <DBInstanceClass>db.m1.large</DBInstanceClass>
      </ReservedDBInstancesOffering>
      <ReservedDBInstancesOffering>
        <Duration>31536000</Duration>
        <OfferingType>Partial Upfront</OfferingType>
        <CurrencyCode>USD</CurrencyCode>
        <RecurringCharges/>
        <FixedPrice>227.5</FixedPrice>
        <ProductDescription>mysql</ProductDescription>
        <UsagePrice>0.046</UsagePrice>
        <MultiAZ>false</MultiAZ>
        <ReservedDBInstancesOfferingId>649fd0c8-cf6d-47a0-bfa6-060f8e75e95f</
      ReservedDBInstancesOfferingId>
        <DBInstanceClass>db.m1.small</DBInstanceClass>
      </ReservedDBInstancesOffering>
    </ReservedDBInstancesOfferings>
  </DescribeReservedDBInstancesOfferingsResult>
  <ResponseMetadata>
    <RequestId>5e4ec40b-2978-11e1-9e6d-771388d6ed6b</RequestId>
  </ResponseMetadata>
</DescribeReservedDBInstancesOfferingsResponse>
```

After you have information about the available reserved DB instance offerings, you can use the information to purchase an offering as shown in the following example.

Example Purchase a Reserved DB Instance

To purchase a reserved DB instance, call the Amazon RDS API operation [PurchaseReservedDBInstancesOffering](#) with the following parameters:

- **--reserved-db-instances-offering-id** – the id of the offering that you want to purchase. See the preceding example to get the offering ID.
- **--reserved-db-instance-id** – you can assign your own identifier to the reserved DB instances that you purchase to help you track them.

The following example purchases the reserved DB instance offering with ID **649fd0c8-cf6d-47a0-bfa6-060f8e75e95f**, and assigns the identifier of **MyReservation**.

```
https://rds.us-east-1.amazonaws.com/
?Action=PurchaseReservedDBInstancesOffering
&ReservedDBInstanceId=MyReservation
&ReservedDBInstancesOfferingId=438012d3-4052-4cc7-b2e3-8d3372e0e706
&DBInstanceCount=10
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2014-09-01
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140415/us-east-1/rds/aws4_request
&X-Amz-Date=20140415T232655Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=c2ac761e8c8f54a8c0727f5a87ad0a766fbb0024510b9aa34ea6d1f7df52fb11
```

This call returns output similar to the following:

```
<PurchaseReservedDBInstancesOfferingResponse xmlns="http://rds.amazonaws.com/
doc/2014-10-31/">
  <PurchaseReservedDBInstancesOfferingResult>
    <ReservedDBInstance>
      <OfferingType>Partial Upfront</OfferingType>
      <CurrencyCode>USD</CurrencyCode>
      <RecurringCharges/>
      <ProductDescription>mysql</ProductDescription>
      <ReservedDBInstancesOfferingId>649fd0c8-cf6d-47a0-bfa6-060f8e75e95f</
      ReservedDBInstancesOfferingId>
      <MultiAZ>true</MultiAZ>
      <State>payment-pending</State>
      <ReservedDBInstanceId>MyReservation</ReservedDBInstanceId>
      <DBInstanceCount>10</DBInstanceCount>
      <StartTime>2011-12-18T23:24:56.577Z</StartTime>
      <Duration>31536000</Duration>
      <FixedPrice>123.0</FixedPrice>
      <UsagePrice>0.123</UsagePrice>
      <DBInstanceClass>db.m1.small</DBInstanceClass>
    </ReservedDBInstance>
  </PurchaseReservedDBInstancesOfferingResult>
  <ResponseMetadata>
    <RequestId>7f099901-29cf-11e1-bd06-6fe008f046c3</RequestId>
  </ResponseMetadata>
</PurchaseReservedDBInstancesOfferingResponse>
```

After you have purchased reserved DB instances, you can get information about your reserved DB instances as shown in the following example.

Example Get Your Reserved DB Instances

To get information about reserved DB instances for your AWS account, call the Amazon RDS API operation [DescribeReservedDBInstances](#).

```
https://rds.us-west-2.amazonaws.com/
?Action=DescribeReservedDBInstances
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2014-09-01
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140420/us-west-2/rds/aws4_request
&X-Amz-Date=20140420T162211Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=3312d17a4c43bcd209bc22a0778dd23e73f8434254abbd7ac53b89ade3dae88e
```

The API returns output similar to the following:

```
<DescribeReservedDBInstancesResponse xmlns="http://rds.amazonaws.com/doc/2014-10-31/">
<DescribeReservedDBInstancesResult>
  <ReservedDBInstances>
    <ReservedDBInstance>
      <OfferingType>Partial Upfront</OfferingType>
      <CurrencyCode>USD</CurrencyCode>
      <RecurringCharges/>
      <ProductDescription>mysql</ProductDescription>
      <ReservedDBInstancesOfferingId>649fd0c8-cf6d-47a0-bfa6-060f8e75e95f</
      ReservedDBInstancesOfferingId>
      <MultiAZ>false</MultiAZ>
      <State>payment-failed</State>
      <ReservedDBInstanceId>MyReservation</ReservedDBInstanceId>
      <DBInstanceCount>1</DBInstanceCount>
      <StartTime>2010-12-15T00:25:14.131Z</StartTime>
      <Duration>31536000</Duration>
      <FixedPrice>227.5</FixedPrice>
      <UsagePrice>0.046</UsagePrice>
      <DBInstanceClass>db.m1.small</DBInstanceClass>
    </ReservedDBInstance>
    <ReservedDBInstance>
      <OfferingType>Partial Upfront</OfferingType>
      <CurrencyCode>USD</CurrencyCode>
      <RecurringCharges/>
      <ProductDescription>mysql</ProductDescription>
      <ReservedDBInstancesOfferingId>649fd0c8-cf6d-47a0-bfa6-060f8e75e95f</
      ReservedDBInstancesOfferingId>
      <MultiAZ>false</MultiAZ>
      <State>payment-failed</State>
      <ReservedDBInstanceId>MyReservation</ReservedDBInstanceId>
      <DBInstanceCount>1</DBInstanceCount>
      <StartTime>2010-12-15T01:07:22.275Z</StartTime>
      <Duration>31536000</Duration>
      <FixedPrice>227.5</FixedPrice>
      <UsagePrice>0.046</UsagePrice>
      <DBInstanceClass>db.m1.small</DBInstanceClass>
    </ReservedDBInstance>
  </ReservedDBInstances>
</DescribeReservedDBInstancesResult>
<ResponseMetadata>
  <RequestId>23400d50-2978-11e1-9e6d-771388d6ed6b</RequestId>
</ResponseMetadata>
</DescribeReservedDBInstancesResponse>
```

Backing Up and Restoring Amazon RDS DB Instances

This section shows how to back up and restore a DB instance.

Topics

- [Working With Backups \(p. 254\)](#)
- [Creating a DB Snapshot \(p. 264\)](#)
- [Restoring from a DB Snapshot \(p. 266\)](#)
- [Copying a Snapshot \(p. 269\)](#)
- [Sharing a DB Snapshot \(p. 278\)](#)
- [Restoring a DB Instance to a Specified Time \(p. 285\)](#)
- [Deleting a Snapshot \(p. 287\)](#)
- [Tutorial: Restore a DB Instance from a DB Snapshot \(p. 289\)](#)

Working With Backups

Amazon RDS creates and saves automated backups of your DB instance. Amazon RDS creates a storage volume snapshot of your DB instance, backing up the entire DB instance and not just individual databases.

Amazon RDS creates automated backups of your DB instance during the backup window of your DB instance. Amazon RDS saves the automated backups of your DB instance according to the backup retention period that you specify. If necessary, you can recover your database to any point in time during the backup retention period.

Automated backups follow these rules:

- Your DB instance must be in the `AVAILABLE` state for automated backups to occur. Automated backups don't occur while your DB instance is in a state other than `AVAILABLE`, for example `STORAGE_FULL`.
- Automated backups and automated snapshots don't occur while a copy is executing in the same region for the same DB instance.

You can also back up your DB instance manually, by manually creating a DB snapshot. For more information about creating a DB snapshot, see [Creating a DB Snapshot \(p. 264\)](#).

The first snapshot of a DB instance contains the data for the full DB instance. Subsequent snapshots of the same DB instance are incremental, which means that only the data that has changed after your most recent snapshot is saved.

You can copy both automatic and manual DB snapshots, and share manual DB snapshots. For more information about copying a DB snapshot, see [Copying a Snapshot \(p. 269\)](#). For more information about sharing a DB snapshot, see [Sharing a DB Snapshot \(p. 278\)](#).

Note

You can also use AWS Backup to manage backups of Amazon RDS DB instances. Backups managed by AWS Backup are considered manual snapshots for the manual snapshot limit. For information about AWS Backup, see the [AWS Backup Developer Guide](#).

Backup Storage

Your Amazon RDS backup storage for each region is composed of the automated backups and manual DB snapshots for that region. Your backup storage is equivalent to the sum of the database storage for all instances in that region. Moving a DB snapshot to another region increases the backup storage in the destination region.

For more information about backup storage costs, see [Amazon RDS Pricing](#).

If you chose to retain automated backups when you delete a DB instance, the automated backups are saved for the full retention period. If you don't choose **Retain automated backups** when you delete a DB instance, all automated backups are deleted with the DB instance. After they are deleted, the automated backups can't be recovered. If you choose to have Amazon RDS create a final DB snapshot before it deletes your DB instance, you can use that to recover your DB instance. Or you can use a previously created manual snapshot. Manual snapshots are not deleted. You can have up to 100 manual snapshots per region.

Backup Window

Automated backups occur daily during the preferred backup window. If the backup requires more time than allotted to the backup window, the backup continues after the window ends, until it finishes. The backup window can't overlap with the weekly maintenance window for the DB instance.

During the automatic backup window, storage I/O might be suspended briefly while the backup process initializes (typically under a few seconds). You might experience elevated latencies for a few minutes during backups for Multi-AZ deployments. For MariaDB, MySQL, Oracle, and PostgreSQL, I/O activity is not suspended on your primary during backup for Multi-AZ deployments, because the backup is taken from the standby. For SQL Server, I/O activity is suspended briefly during backup for Multi-AZ deployments.

If you don't specify a preferred backup window when you create the DB instance, Amazon RDS assigns a default 30-minute backup window. This window is selected at random from an 8-hour block of time for each AWS Region. The following table lists the time blocks for each region from which the default backups windows are assigned.

Region	Time Block
US East (Ohio) Region	03:00–11:00 UTC
US East (N. Virginia) Region	03:00–11:00 UTC
US West (N. California) Region	06:00–14:00 UTC
US West (Oregon) Region	06:00–14:00 UTC
Asia Pacific (Hong Kong) Region	06:00–14:00 UTC
Asia Pacific (Mumbai) Region	16:30–00:30 UTC
Asia Pacific (Osaka-Local) Region	00:00–08:00 UTC
Asia Pacific (Seoul) Region	13:00–21:00 UTC
Asia Pacific (Singapore) Region	14:00–22:00 UTC
Asia Pacific (Sydney) Region	12:00–20:00 UTC
Asia Pacific (Tokyo) Region	13:00–21:00 UTC
Canada (Central) Region	06:29–14:29 UTC
China (Beijing) Region	06:00–14:00 UTC
China (Ningxia) Region	06:00–14:00 UTC
Europe (Frankfurt) Region	20:00–04:00 UTC
Europe (Ireland) Region	22:00–06:00 UTC
Europe (London) Region	06:00–14:00 UTC
Europe (Paris) Region	07:29–14:29 UTC
Europe (Stockholm) Region	23:00–07:00 UTC
Middle East (Bahrain) Region	06:00–14:00 UTC
South America (São Paulo) Region	23:00–07:00 UTC

Region	Time Block
AWS GovCloud (US-West)	03:00–11:00 UTC

Backup Retention Period

You can set the backup retention period when you create a DB instance. If you don't set the backup retention period, the default backup retention period is one day if you create the DB instance using the Amazon RDS API or the AWS CLI. The default backup retention period is seven days if you create the DB instance using the console. After you create a DB instance, you can modify the backup retention period. You can set the backup retention period to between 0 and 35 days. Setting the backup retention period to 0 disables automated backups. Manual snapshot limits (100 per region) do not apply to automated backups.

Important

An outage occurs if you change the backup retention period from 0 to a non-zero value or from a non-zero value to 0.

Disabling Automated Backups

You might want to temporarily disable automated backups in certain situations; for example, while loading large amounts of data.

Important

We highly discourage disabling automated backups because it disables point-in-time recovery. Disabling automatic backups for a DB instance deletes all existing automated backups for the instance. If you disable and then re-enable automated backups, you are only able to restore starting from the time you re-enabled automated backups.

In this example, you disable automated backups for a DB instance named *mydbinstance* by setting the backup retention parameter to 0.

Console

To disable automated backups immediately

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to modify.
3. Choose **Modify**. The **Modify DB Instance** page appears.
4. For **Backup Retention Period**, choose **0 days**.
5. Choose **Continue**.
6. Choose **Apply Immediately**.
7. On the confirmation page, choose **Modify DB Instance** to save your changes and disable automated backups.

AWS CLI

To disable automated backups immediately, use the `modify-db-instance` command and set the backup retention period to 0 with `--apply-immediately`.

Example

The following example immediately disabled automatic backups.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier mydbinstance \
--backup-retention-period 0 \
--apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier mydbinstance ^
--backup-retention-period 0 ^
--apply-immediately
```

To know when the modification is in effect, call `describe-db-instances` for the DB instance until the value for backup retention period is 0 and `mydbinstance` status is available.

```
aws rds describe-db-instances --db-instance-identifier mydbinstance
```

RDS API

To disable automated backups immediately, call the [ModifyDBInstance](#) operation with the following parameters:

- `DBInstanceIdentifier` = `mydbinstance`
- `BackupRetentionPeriod` = 0

Example

```
https://rds.amazonaws.com/
?Action=ModifyDBInstance
&DBInstanceIdentifier=mydbinstance
&BackupRetentionPeriod=0
&SignatureVersion=2
&SignatureMethod=HmacSHA256
&Timestamp=2009-10-14T17%3A48%3A21.746Z
&AWSAccessKeyId=<AWS Access Key ID>
&Signature=<Signature>
```

Enabling Automated Backups

If your DB instance doesn't have automated backups enabled, you can enable them at any time. You enable automated backups by setting the backup retention period to a positive non-zero value. When automated backups are enabled, your RDS instance and database is taken offline and a backup is immediately created.

In this example, you enable automated backups for a DB instance named `mydbinstance` by setting the backup retention period to a positive non-zero value (in this case, 3).

Console

To enable automated backups immediately

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.

2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to modify.
3. Choose **Modify**. The **Modify DB Instance** page appears.
4. For **Backup Retention Period**, choose a positive nonzero value, for example 3 days.
5. Choose **Continue**.
6. Choose **Apply Immediately**.
7. On the confirmation page, choose **Modify DB Instance** to save your changes and enable automated backups.

AWS CLI

To enable automated backups immediately, use the AWS CLI `modify-db-instance` command.

In this example, we enable automated backups by setting the backup retention period to three days.

Include the following parameters:

- `--db-instance-identifier`
- `--backup-retention-period`
- `--apply-immediately` or `--no-apply-immediately`

Example

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
  --db-instance-identifier mydbinstance \
  --backup-retention-period 3 \
  --apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
  --db-instance-identifier mydbinstance ^
  --backup-retention-period 3 ^
  --apply-immediately
```

RDS API

To enable automated backups immediately, use the RDS API `ModifyDBInstance` operation.

In this example, we enable automated backups by setting the backup retention period to three days.

Include the following parameters:

- `DBInstanceIdentifier`
- `BackupRetentionPeriod`
- `ApplyImmediately = true`

Example

```
https://rds.amazonaws.com/
?Action=ModifyDBInstance
&DBInstanceIdentifier=mydbinstance
```

```
&BackupRetentionPeriod=3
&ApplyImmediately=true
&SignatureVersion=2
&SignatureMethod=HmacSHA256
&Timestamp=2009-10-14T17%3A48%3A21.746Z
&AWSAccessKeyId=<AWS Access Key ID>
&Signature=<Signature>
```

Retaining Automated Backups

When you delete a DB instance, you can retain automated backups.

Retained automated backups contain system snapshots and transaction logs from a DB instance. They also include your DB instance properties like allocated storage and DB instance class, which are required to restore it to an active instance.

You can retain automated backups for RDS instances running MySQL, MariaDB, PostgreSQL, Oracle, and Microsoft SQL Server engines.

You can restore or remove retained automated backups using the AWS Management Console, RDS API, and AWS CLI.

Topics

- [Retention Period \(p. 259\)](#)
- [Restoration \(p. 259\)](#)
- [Retention Costs \(p. 260\)](#)
- [Limitations and Recommendations \(p. 260\)](#)
- [Deleting Retained Automated Backups \(p. 260\)](#)

Retention Period

The system snapshots and transaction logs in a retained automated backup expire the same way that they expire for the source DB instance. Because there are no new snapshots or logs created for this instance, the retained automated backups eventually expire completely. Effectively, they live as long their last system snapshot would have done, based on the settings for retention period the source instance had when you deleted it. Retained automated backups are removed by the system after their last system snapshot expires.

You can remove a retained automated backup in the same way that you can delete a DB instance.

You can remove retained automated backups using the console or the RDS API operation `DeleteDBInstanceAutomatedBackup`.

Final snapshots are independent of retained automated backups. We strongly suggest that you take a final snapshot even if you retain automated backups, because the retained automated backups eventually expire. The final snapshot doesn't expire.

Restoration

To view your retained automated backups, switch to the automated backups page. You can view individual snapshots associated with a retained automated backup on the database snapshots page in the console. Alternatively, you can describe individual snapshots associated with a retained automated backup. From there, you can restore a DB instance directly from one of those snapshots.

Restored DB instances are automatically associated with the default parameter and option groups. However, you can apply a custom parameter group and option group by specifying them during a restore.

In this example, you restore a DB instance to a point in time using the retained automated backup. First, you describe your retained automated backups, so you can see which of them to restore.

To describe your retained automated backups using the RDS API, call the [DescribeDBInstanceAutomatedBackups](#) action with one of the following parameters:

- `DBInstanceIdentifier`
- `DbiResourceId`

```
aws rds describe-db-instance-automated-backups --db-instance-  
identifier DBInstanceIdentifier  
OR  
aws rds describe-db-instance-automated-backups --dbi-resource-idDbiResourceId
```

Next, to restore your retained automated backup to a point in time, using the RDS API, call the [RestoreDBInstanceToPointInTime](#) action with the following parameters:

- `SourceDbiResourceId`
- `TargetDBInstanceIdentifier`

```
aws rds restore-db-instance-to-point-in-time --source-dbi-resource-id SourceDbiResourceId  
--target-db-instance-identifier TargetDBInstanceIdentifier --use-latest-restorable-time
```

Retention Costs

The cost of a retained automated backup is the cost of total storage of the system snapshots that are associated with it. There is no additional charge for transaction logs or instance metadata. All other pricing rules for backups apply to restorable instances.

For example, suppose that your total allocated storage of running instances is 100 GB. Suppose also that you have 50 GB of manual snapshots plus 75 GB of system snapshots associated with a retained automated backup. In this case, you are charged only for the additional 25 GB of backup storage, like this: $(50 \text{ GB} + 75 \text{ GB}) - 100 \text{ GB} = 25 \text{ GB}$.

Limitations and Recommendations

The following limitations apply to retained automated backups:

- The maximum number of retained automated backups in one AWS Region is 40. It's not included in the DB instances limit. You can have 40 running DB instances and an additional 40 retained automated backups at the same time.
- Retained automated backups don't contain information about parameters or option groups.
- You can restore a deleted instance to a point in time that is within the retention period at the time of delete.
- A retained automated backup can't be modified because it consists of system backups, transaction logs, and the DB instance properties that existed at the time you deleted the source instance.

Deleting Retained Automated Backups

You can delete retained automated backups when they are no longer needed.

Console

To delete a retained automated backup

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Automated backups**.
3. Choose **Retained**.

The screenshot shows the AWS RDS 'Automated backups' page. At the top, there are two tabs: 'Active' and 'Retained'. The 'Retained' tab is highlighted with a red oval. Below the tabs, the heading 'Retained backups (1)' is displayed. Underneath is a search bar labeled 'Filter retained backups'. A table follows, with columns 'DB Name', 'Earliest restorable time', and 'Latest restorable time'. The table contains one row for a MySQL instance named 'mysql56', with the earliest restorable time listed as 'Mon Nov 11 2019 23:25:34 GMT-0800' and the latest restorable time as 'Wed Nov 20 2019'.

DB Name	Earliest restorable time	Latest restorable time
mysql56	Mon Nov 11 2019 23:25:34 GMT-0800	Wed Nov 20 2019

4. Choose the retained automated backup that you want to delete.
5. For **Actions**, choose **Delete**.
6. On the confirmation page, enter **delete me** and choose **Delete**.

AWS CLI

You can delete a retained automated backup by using the AWS CLI command [delete-db-instance-automated-backup](#).

The following options are used to delete a retained automated backup:

- `--dbi-resource-id` – The resource identifier for the source DB instance.

You can find the resource identifier for the source DB instance of a retained automated backup by running the AWS CLI command [describe-db-instance-automated-backups](#).

Example

The following example deletes the retained automated backup with source DB instance resource identifier `db-123ABCEXAMPLE`.

For Linux, OS X, or Unix:

```
aws rds delete-db-instance-automated-backup \
--dbi-resource-id db-123ABCEXAMPLE
```

For Windows:

```
aws rds delete-db-instance-automated-backup ^
--dbi-resource-id db-123ABCEXAMPLE
```

RDS API

You can delete a retained automated backup by using the Amazon RDS API operation [DeleteDBInstanceAutomatedBackup](#).

The following parameters are used to delete a retained automated backup:

- `DbiResourceId` – The resource identifier for the source DB instance.

You can find the resource identifier for the source DB instance of a retained automated backup using the Amazon RDS API operation [DescribeDBInstanceAutomatedBackups](#).

Automated Backups with Unsupported MySQL Storage Engines

For the MySQL DB engine, automated backups are only supported for the InnoDB storage engine. Use of these features with other MySQL storage engines, including MyISAM, can lead to unreliable behavior while restoring from backups. Specifically, since storage engines like MyISAM don't support reliable crash recovery, your tables can be corrupted in the event of a crash. For this reason, we encourage you to use the InnoDB storage engine.

- To convert existing MyISAM tables to InnoDB tables, you can use the `ALTER TABLE` command, for example: `ALTER TABLE table_name ENGINE=innodb, ALGORITHM=COPY;`
- If you choose to use MyISAM, you can attempt to manually repair tables that become damaged after a crash by using the `REPAIR` command. For more information, see [REPAIR TABLE Syntax](#) in the MySQL documentation. However, as noted in the MySQL documentation, there is a good chance that you might not be able to recover all your data.
- If you want to take a snapshot of your MyISAM tables before restoring, follow these steps:

1. Stop all activity to your MyISAM tables (that is, close all sessions).

You can close all sessions by calling the `mysql.rds_kill` command for each process that is returned from the `SHOW FULL PROCESSLIST` command.

2. Lock and flush each of your MyISAM tables. For example, the following commands lock and flush two tables named `myisam_table1` and `myisam_table2`:

```
mysql> FLUSH TABLES myisam_table1, myisam_table2 WITH READ LOCK;
```

3. Create a snapshot of your DB instance. When the snapshot has completed, release the locks and resume activity on the MyISAM tables. You can release the locks on your tables using the following command:

```
mysql> UNLOCK TABLES;
```

These steps force MyISAM to flush data stored in memory to disk, which ensures a clean start when you restore from a DB snapshot. For more information on creating a DB snapshot, see [Creating a DB Snapshot \(p. 264\)](#).

Automated Backups with Unsupported MariaDB Storage Engines

For the MariaDB DB engine, automated backups are only supported with the InnoDB storage engine (version 10.2 and later) and XtraDB storage engine (versions 10.0 and 10.1). Use of these features with other MariaDB storage engines, including Aria, might lead to unreliable behavior while restoring from backups. Even though Aria is a crash-resistant alternative to MyISAM, your tables can still be corrupted in the event of a crash. For this reason, we encourage you to use the XtraDB storage engine.

- To convert existing Aria tables to InnoDB tables, you can use the `ALTER TABLE` command. For example: `ALTER TABLE table_name ENGINE=innodb, ALGORITHM=COPY;`
- To convert existing Aria tables to XtraDB tables, you can use the `ALTER TABLE` command. For example: `ALTER TABLE table_name ENGINE=xtradb, ALGORITHM=COPY;`
- If you choose to use Aria, you can attempt to manually repair tables that become damaged after a crash by using the `REPAIR TABLE` command. For more information, see <http://mariadb.com/kb/en/mariadb/repair-table/>.
- If you want to take a snapshot of your Aria tables before restoring, follow these steps:
 1. Stop all activity to your Aria tables (that is, close all sessions).
 2. Lock and flush each of your Aria tables.
 3. Create a snapshot of your DB instance. When the snapshot has completed, release the locks and resume activity on the Aria tables. These steps force Aria to flush data stored in memory to disk, thereby ensuring a clean start when you restore from a DB snapshot.

Creating a DB Snapshot

Amazon RDS creates a storage volume snapshot of your DB instance, backing up the entire DB instance and not just individual databases. Creating this DB snapshot on a Single-AZ DB instance results in a brief I/O suspension that can last from a few seconds to a few minutes, depending on the size and class of your DB instance. Multi-AZ DB instances are not affected by this I/O suspension since the backup is taken on the standby.

When you create a DB snapshot, you need to identify which DB instance you are going to back up, and then give your DB snapshot a name so you can restore from it later. The amount of time it takes to create a snapshot varies with the size your databases. Since the snapshot includes the entire storage volume, the size of files, such as temporary files, also affects the amount of time it takes to create the snapshot.

You can create a DB snapshot using the AWS Management Console, the AWS CLI, or the RDS API.

Console

To create a DB snapshot

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. In the list of DB instances, choose the DB instance for which you want to take a snapshot.
4. For **Actions**, choose **Take snapshot**.

The **Take DB Snapshot** window appears.

5. Type the name of the snapshot in the **Snapshot Name** box.

Take DB Snapshot

This feature is currently supported for InnoDB storage engine only. If you are using MyISAM, refer to details [here](#). 

Settings

To take a snapshot of this DB instance you must provide a name for the snapshot.

DB instance

The unique key that identifies a DB instance. This parameter isn't case-sensitive.

mydbinstance3

Snapshot name

The Identifier for the DB Snapshot.

Cancel

Take Snapshot

6. Choose **Take Snapshot**.

AWS CLI

When you create a DB snapshot using the AWS CLI, you need to identify which DB instance you are going to back up, and then give your DB snapshot a name so you can restore from it later. You can do this by using the AWS CLI [create-db-snapshot](#) command with the following parameters:

- `--db-instance-identifier`
- `--db-snapshot-identifier`

In this example, you create a DB snapshot called `mydbsnapshot` for a DB instance called `mydbinstance`.

Example

For Linux, OS X, or Unix:

```
aws rds create-db-snapshot \
--db-instance-identifier mydbinstance \
--db-snapshot-identifier mydbsnapshot
```

For Windows:

```
aws rds create-db-snapshot ^
--db-instance-identifier mydbinstance ^
--db-snapshot-identifier mydbsnapshot
```

RDS API

When you create a DB snapshot using the Amazon RDS API, you need to identify which DB instance you are going to back up, and then give your DB snapshot a name so you can restore from it later. You can do this by using the Amazon RDS API [CreateDBSnapshot](#) command with the following parameters:

- `DBInstanceIdentifier`
- `DBSnapshotIdentifier`

Restoring from a DB Snapshot

Amazon RDS creates a storage volume snapshot of your DB instance, backing up the entire DB instance and not just individual databases. You can create a DB instance by restoring from this DB snapshot. When you restore the DB instance, you provide the name of the DB snapshot to restore from, and then provide a name for the new DB instance that is created from the restore. You can't restore from a DB snapshot to an existing DB instance; a new DB instance is created when you restore.

You can restore a DB instance and use a different storage type than the source DB snapshot. In this case, the restoration process is slower because of the additional work required to migrate the data to the new storage type. If you restore to or from Magnetic (Standard) storage, the migration process is the slowest. That's because Magnetic storage doesn't have the IOPS capability of Provisioned IOPS or General Purpose (SSD) storage.

Note

You can't restore a DB instance from a DB snapshot that is both shared and encrypted. Instead, you can make a copy of the DB snapshot and restore the DB instance from the copy.

Parameter Group Considerations

We recommend that you retain the parameter group for any DB snapshots you create, so that you can associate your restored DB instance with the correct parameter group. You can specify the parameter group when you restore the DB instance.

Security Group Considerations

When you restore a DB instance, the default security group is associated with the restored instance by default.

Note

- If you're using the AWS CLI, you can specify a custom security group to associate with the instance by including the `--vpc-security-group-ids` option in the `restore-db-instance-from-db-snapshot` command.
- If you're using the Amazon RDS API, you can include the `VpcSecurityGroupIds.VpcSecurityGroupId.N` parameter in the `RestoreDBInstanceFromDBSnapshot` action.
- The Amazon RDS console has no option for associating a custom security group while restoring.

As soon as the restore is complete and your new DB instance is available, you can associate any custom security groups used by the snapshot you restored from. You must apply these changes by modifying the DB instance with the RDS console, the AWS CLI `modify-db-instance` command, or the `ModifyDBInstance` Amazon RDS API operation. For more information, see [Modifying an Amazon RDS DB Instance \(p. 150\)](#).

Option Group Considerations

When you restore a DB instance, the option group associated with the DB snapshot is associated with the restored DB instance after it is created. For example, if the DB snapshot you are restoring from uses Oracle Transparent Data Encryption, the restored DB instance will use the same option group.

When you assign an option group to a DB instance, the option group is also linked to the supported platform the DB instance is on, either VPC or EC2-Classic (non-VPC). If a DB instance is in a VPC, the option group associated with the DB instance is linked to that VPC. This means that you can't use the

option group assigned to a DB instance if you attempt to restore the instance into a different VPC or onto a different platform. If you restore a DB instance into a different VPC or onto a different platform, you must either assign the default option group to the instance, assign an option group that is linked to that VPC or platform, or create a new option group and assign it to the DB instance. For persistent or permanent options, when restoring a DB instance into a different VPC you must create a new option group that includes the persistent or permanent option.

Microsoft SQL Server Considerations

When you restore a Microsoft SQL Server DB snapshot to a new instance, you can always restore to the same edition as your snapshot. In some cases, you can also change the edition of the DB instance. The following are the limitations when you change editions:

- The DB snapshot must have enough storage allocated for the new edition.
- Only the following edition changes are supported:
 - From Standard Edition to Enterprise Edition
 - From Web Edition to Standard Edition or Enterprise Edition
 - From Express Edition to Web Edition, Standard Edition or Enterprise Edition

If you want to change from one edition to a new edition that is not supported by restoring a snapshot, you can try using the native backup and restore feature. SQL Server verifies whether or not your database is compatible with the new edition based on what SQL Server features you have enabled on the database. For more information, see [Importing and Exporting SQL Server Databases \(p. 657\)](#).

Oracle Considerations

If you use Oracle GoldenGate, always retain the parameter group with the `compatible` parameter. When you restore a DB instance from a DB snapshot, you must specify the parameter group that has a matching or greater `compatible` parameter value.

You can upgrade a DB snapshot while it is still a DB snapshot, before you restore it. For more information, see [Upgrading an Oracle DB Snapshot \(p. 984\)](#).

Restoring from a Snapshot

You can restore a DB instance from a DB snapshot using the AWS Management Console, the AWS CLI, or the RDS API.

Console

To restore a DB instance from a DB snapshot

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Snapshots**.
3. Choose the DB snapshot that you want to restore from.
4. For **Actions**, choose **Restore Snapshot**.
5. On the **Restore DB Instance** page, for **DB Instance Identifier**, enter the name for your restored DB instance.
6. Choose **Restore DB Instance**.
7. If you want to restore the functionality of the DB instance to that of the DB instance that the snapshot was created from, you must modify the DB instance to use the security group. The next

steps assume that your DB instance is in a VPC. If your DB instance is not in a VPC, use the EC2 Management Console to locate the security group you need for the DB instance.

- a. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
- b. In the navigation pane, choose **Security Groups**.
- c. Select the security group that you want to use for your DB instances. If necessary, add rules to link the security group to a security group for an EC2 instance. For more information, see [A DB Instance in a VPC Accessed by an EC2 Instance in the Same VPC \(p. 507\)](#).

AWS CLI

To restore a DB instance from a DB snapshot, use the AWS CLI command `restore-db-instance-from-db-snapshot`.

In this example, you restore from a previously created DB snapshot named *mydbsnapshot*. You restore to a new DB instance named *mynewdbinstance*.

Example

For Linux, OS X, or Unix:

```
aws rds restore-db-instance-from-db-snapshot \
  --db-instance-identifier mynewdbinstance \
  --db-snapshot-identifier mydbsnapshot
```

For Windows:

```
aws rds restore-db-instance-from-db-snapshot ^
  --db-instance-identifier mynewdbinstance ^
  --db-snapshot-identifier mydbsnapshot
```

This command returns output similar to the following:

```
DBINSTANCE mynewdbinstance db.m3.large MySQL      50          sa           creating  3  n
5.6.40 general-public-license
```

After the DB instance has been restored, you must add the DB instance to the security group used by the DB instance used to create the DB snapshot if you want the same functionality as that of the previous DB instance.

RDS API

To restore a DB instance from a DB snapshot, call the Amazon RDS API function `RestoreDBInstanceFromDBSnapshot` with the following parameters:

- `DBInstanceIdentifier`
- `DBSnapshotIdentifier`

Copying a Snapshot

With Amazon RDS, you can copy automated or manual DB snapshots. After you copy a snapshot, the copy is a manual snapshot.

You can copy a snapshot within the same AWS Region, you can copy a snapshot across AWS Regions, and you can copy shared snapshots.

Limitations

The following are some limitations when you copy snapshots:

- You can't copy a snapshot to or from the following AWS Regions: China (Beijing) or China (Ningxia).
- You can copy a snapshot between AWS GovCloud (US-East) and AWS GovCloud (US-West), but you can't copy a snapshot between these AWS GovCloud (US) regions and other AWS Regions.
- If you delete a source snapshot before the target snapshot becomes available, the snapshot copy may fail. Verify that the target snapshot has a status of **AVAILABLE** before you delete a source snapshot.
- You can have up to five snapshot copy requests in progress to a single destination region per account.
- Depending on the regions involved and the amount of data to be copied, a cross-region snapshot copy can take hours to complete. If there is a large number of cross-region snapshot copy requests from a given source AWS Region, Amazon RDS might put new cross-region copy requests from that source AWS Region into a queue until some in-progress copies complete. No progress information is displayed about copy requests while they are in the queue. Progress information is displayed when the copy starts.

Snapshot Retention

Amazon RDS deletes automated snapshots at the end of their retention period, when you disable automated snapshots for a DB instance, or when you delete a DB instance. If you want to keep an automated snapshot for a longer period, copy it to create a manual snapshot, which is retained until you delete it. Amazon RDS storage costs might apply to manual snapshots if they exceed your default storage space.

For more information about backup storage costs, see [Amazon RDS Pricing](#).

Copying Shared Snapshots

You can copy snapshots shared to you by other AWS accounts. If you are copying an encrypted snapshot that has been shared from another AWS account, you must have access to the KMS encryption key that was used to encrypt the snapshot.

You can copy a shared DB snapshot across regions, provided that the snapshot is unencrypted. However, if the shared DB snapshot is encrypted, you can only copy it in the same AWS Region.

Handling Encryption

You can copy a snapshot that has been encrypted using an AWS KMS encryption key. If you copy an encrypted snapshot, the copy of the snapshot must also be encrypted. If you copy an encrypted snapshot within the same AWS Region, you can encrypt the copy with the same KMS encryption key as the original snapshot, or you can specify a different KMS encryption key. If you copy an encrypted snapshot across regions, you can't use the same KMS encryption key for the copy as used for the source snapshot, because KMS keys are region-specific. Instead, you must specify a KMS key valid in the destination AWS Region. You can't use the default AWS KMS encryption key when copying snapshots across AWS Regions, even in the same AWS account.

You can also encrypt a copy of an unencrypted snapshot. This way, you can quickly add encryption to a previously unencrypted DB instance. That is, you can create a snapshot of your DB instance when you are ready to encrypt it, and then create a copy of that snapshot and specify a KMS encryption key to encrypt that snapshot copy. You can then restore an encrypted DB instance from the encrypted snapshot.

Copying Snapshots Across AWS Regions

When you copy a snapshot to an AWS Region that is different from the source snapshot's AWS Region, the first copy is a full snapshot copy, even if you copy an incremental snapshot. A full snapshot copy contains all of the data and metadata required to restore the DB instance. After the first snapshot copy, you can copy incremental snapshots of the same DB instance to the same destination region within the same AWS account.

An incremental snapshot contains only the data that has changed after the most recent snapshot of the same DB instance. Incremental snapshot copying is faster and results in lower storage costs than full snapshot copying. Incremental snapshot copying across AWS Regions is supported for both unencrypted and encrypted snapshots.

Important

For shared snapshots, copying incremental snapshots is not supported. For shared snapshots, all of the copies are full snapshots, even within the same region.

Depending on the AWS Regions involved and the amount of data to be copied, a cross-region snapshot copy can take hours to complete. In some cases, there might be a large number of cross-region snapshot copy requests from a given source AWS Region. In these cases, Amazon RDS might put new cross-region copy requests from that source AWS Region into a queue until some in-progress copies complete. No progress information is displayed about copy requests while they are in the queue. Progress information is displayed when the copy starts.

Note

When you copy a source snapshot that is a snapshot copy, the copy isn't incremental because the snapshot copy doesn't include the required metadata for incremental copies.

Option Group Considerations

Option groups are specific to the AWS Region that they are created in, and you can't use an option group from one AWS Region in another AWS Region.

When you copy a snapshot across regions, you can specify a new option group for the snapshot. We recommend that you prepare the new option group before you copy the snapshot. In the destination AWS Region, create an option group with the same settings as the original DB instance . If one already exists in the new AWS Region, you can use that one.

If you copy a snapshot and you don't specify a new option group for the snapshot, when you restore it the DB instance gets the default option group. To give the new DB instance the same options as the original, you must do the following:

1. In the destination AWS Region, create an option group with the same settings as the original DB instance . If one already exists in the new AWS Region, you can use that one.
2. After you restore the snapshot in the destination AWS Region, modify the new DB instance and add the new or existing option group from the previous step.

Parameter Group Considerations

When you copy a snapshot across regions, the copy doesn't include the parameter group used by the original DB instance . When you restore a snapshot to create a new DB instance , that DB instance gets

the default parameter group for the AWS Region it is created in. To give the new DB instance the same parameters as the original, you must do the following:

1. In the destination AWS Region, create a DB parameter group with the same settings as the original DB instance . If one already exists in the new AWS Region, you can use that one.
2. After you restore the snapshot in the destination AWS Region, modify the new DB instance and add the new or existing parameter group from the previous step.

Copying a DB Snapshot

Use the procedures in this topic to copy a DB snapshot. For an overview of copying a snapshot, see [Copying a Snapshot \(p. 269\)](#)

For each AWS account, you can copy up to five DB snapshots at a time from one AWS Region to another. If you copy a DB snapshot to another AWS Region, you create a manual DB snapshot that is retained in that AWS Region. Copying a DB snapshot out of the source AWS Region incurs Amazon RDS data transfer charges.

For more information about data transfer pricing, see [Amazon RDS Pricing](#).

After the DB snapshot copy has been created in the new AWS Region, the DB snapshot copy behaves the same as all other DB snapshots in that AWS Region.

You can copy a DB snapshot using the AWS Management Console, the AWS CLI, or the RDS API.

Console

This procedure copies an encrypted or unencrypted DB snapshot, in the same AWS Region or across regions, by using the AWS Management Console.

To copy a DB snapshot

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Snapshots**.
3. Select the DB snapshot that you want to copy.
4. For **Actions**, choose **Copy Snapshot**. The **Make Copy of DB Snapshot** page appears.

Make Copy of DB Snapshot?

Settings

Source DB Snapshot

DB Snapshot Identifier for the automated snapshot being copied.
mydbinstancesnapshot

Destination Region [Info](#)

US East (N. Virginia) ▾

New DB Snapshot Identifier

DB Snapshot Identifier for the new snapshot

Target Option Group (Optional) [Info](#)

No preference ▾

Copy Tags [Info](#)

i Please note that depending on the amount of data to be copied and the Region you choose, this operation could take several hours to complete and the display on the progress bar could be delayed until setup is complete.

Encryption

Encryption [Info](#)

Enable encryption [Learn more](#)

Select to encrypt the given instance. Master key ids and aliases appear in the list after they have been created using the Key Management Service(KMS) console.

Disable encryption

[Cancel](#)

Copy Snapshot

5. (Optional) To copy the DB snapshot to a different AWS Region, for **Destination Region**, choose the new AWS Region.

Note

The destination AWS Region must have the same database engine version available as the source AWS Region.

6. For **New DB Snapshot Identifier**, type the name of the DB snapshot copy.
7. (Optional) For **Target Option Group**, choose a new option group.

Specify this option if you are copying a snapshot from one AWS Region to another, and your DB instance uses a non-default option group.

If your source DB instance uses Transparent Data Encryption for Oracle or Microsoft SQL Server, you must specify this option when copying across regions. For more information, see [Option Group Considerations \(p. 270\)](#).

8. (Optional) Select **Copy Tags** to copy tags and values from the snapshot to the copy of the snapshot.
9. (Optional) For **Enable Encryption**, choose one of the following options:
 - Choose **Disable encryption** if the DB snapshot isn't encrypted and you don't want to encrypt the copy.
 - Choose **Enable encryption** if the DB snapshot isn't encrypted but you want to encrypt the copy. In this case, for **Master Key**, specify the KMS key identifier to use to encrypt the DB snapshot copy.
 - Choose **Enable encryption** if the DB snapshot is encrypted. In this case, you must encrypt the copy, so **Yes** is already selected. For **Master Key**, specify the KMS key identifier to use to encrypt the DB snapshot copy.
10. Choose **Copy Snapshot**.

AWS CLI

You can copy a DB snapshot by using the AWS CLI command `copy-db-snapshot`. If you are copying the snapshot to a new AWS Region, run the command in the new AWS Region.

The following options are used to copy a DB snapshot. Not all options are required for all scenarios. Use the descriptions and the examples that follow to determine which options to use.

- **--source-db-snapshot-identifier** – The identifier for the source DB snapshot.
 - If the source snapshot is in the same AWS Region as the copy, specify a valid DB snapshot identifier. For example, `rds:mysql-instance1-snapshot-20130805`.
 - If the source snapshot is in a different AWS Region than the copy, specify a valid DB snapshot ARN. For example, `arn:aws:rds:us-west-2:123456789012:snapshot:mysql-instance1-snapshot-20130805`.
 - If you are copying from a shared manual DB snapshot, this parameter must be the Amazon Resource Name (ARN) of the shared DB snapshot.
 - If you are copying an encrypted snapshot this parameter must be in the ARN format for the source AWS Region, and must match the `SourceDBSnapshotIdentifier` in the `PreSignedUrl` parameter.
- **--target-db-snapshot-identifier** – The identifier for the new copy of the encrypted DB snapshot.
- **--copy-tags** – Include the copy tags option to copy tags and values from the snapshot to the copy of the snapshot.
- **--option-group-name** – The option group to associate with the copy of the snapshot.

Specify this option if you are copying a snapshot from one AWS Region to another, and your DB instance uses a non-default option group.

If your source DB instance uses Transparent Data Encryption for Oracle or Microsoft SQL Server, you must specify this option when copying across regions. For more information, see [Option Group Considerations \(p. 270\)](#).

- **--kms-key-id** – The AWS KMS key ID for an encrypted DB snapshot. The KMS key ID is the Amazon Resource Name (ARN), KMS key identifier, or the KMS key alias for the KMS encryption key.
 - If you copy an encrypted DB snapshot from your AWS account, you can specify a value for this parameter to encrypt the copy with a new KMS encryption key. If you don't specify a value for this parameter, then the copy of the DB snapshot is encrypted with the same KMS key as the source DB snapshot.

- If you copy an encrypted DB snapshot that is shared from another AWS account, then you must specify a value for this parameter.
- If you specify this parameter when you copy an unencrypted snapshot, the copy is encrypted.
- If you copy an encrypted snapshot to a different AWS Region, then you must specify a KMS key for the destination AWS Region. KMS encryption keys are specific to the AWS Region that they are created in, and you cannot use encryption keys from one AWS Region in another AWS Region.
- `--source-region` – The ID of the AWS Region of the source DB snapshot. If you copy an encrypted snapshot to a different AWS Region, then you must specify this option.

Example From Unencrypted, To Same Region

The following code creates a copy of a snapshot, with the new name `mydbsnapshotcopy`, in the same AWS Region as the source snapshot. When the copy is made, all tags on the original snapshot are copied to the snapshot copy.

For Linux, OS X, or Unix:

```
aws rds copy-db-snapshot \  
  --source-db-snapshot-identifier mysql-instance1-snapshot-20130805 \  
  --target-db-snapshot-identifier mydbsnapshotcopy \  
  --copy-tags
```

For Windows:

```
aws rds copy-db-snapshot ^  
  --source-db-snapshot-identifier mysql-instance1-snapshot-20130805 ^  
  --target-db-snapshot-identifier mydbsnapshotcopy ^  
  --copy-tags
```

Example From Unencrypted, Across Regions

The following code creates a copy of a snapshot, with the new name `mydbsnapshotcopy`, in the AWS Region in which the command is run.

For Linux, OS X, or Unix:

```
aws rds copy-db-snapshot \  
  --source-db-snapshot-identifier arn:aws:rds:us-east-1:123456789012:snapshot:mysql-  
instance1-snapshot-20130805 \  
  --target-db-snapshot-identifier mydbsnapshotcopy
```

For Windows:

```
aws rds copy-db-snapshot ^  
  --source-db-snapshot-identifier arn:aws:rds:us-east-1:123456789012:snapshot:mysql-  
instance1-snapshot-20130805 ^  
  --target-db-snapshot-identifier mydbsnapshotcopy
```

Example From Encrypted, Across Regions

The following code example copies an encrypted DB snapshot from the us-west-2 region in the us-east-1 region. Run the command in the us-east-1 region.

For Linux, OS X, or Unix:

```
aws rds copy-db-snapshot \  
  --source-region us-west-2 \  
  --source-db-snapshot-identifier arn:aws:rds:us-west-2:123456789012:snapshot:mysql-  
instance1-snapshot-20130805 \  
  --target-region us-east-1 \  
  --target-db-snapshot-identifier mydbsnapshotcopy \  
  --kms-key-arn arn:aws:kms:us-east-1:123456789012:key/1234567890123456
```

```
--source-db-snapshot-identifier arn:aws:rds:us-west-2:123456789012:snapshot:mysql-
instance1-snapshot-20161115 \
--target-db-snapshot-identifier mydbsnapshotcopy \
--source-region us-west-2 \
--kms-key-id my-us-east-1-key \
--option-group-name custom-option-group-name
```

For Windows:

```
aws rds copy-db-snapshot ^
--source-db-snapshot-identifier arn:aws:rds:us-west-2:123456789012:snapshot:mysql-
instance1-snapshot-20161115 ^
--target-db-snapshot-identifier mydbsnapshotcopy ^
--source-region us-west-2 ^
--kms-key-id my-us-east-1-key ^
--option-group-name custom-option-group-name
```

RDS API

You can copy a DB snapshot by using the Amazon RDS API operation [CopyDBSnapshot](#). If you are copying the snapshot to a new AWS Region, perform the action in the new AWS Region.

The following parameters are used to copy a DB snapshot. Not all parameters are required for all scenarios. Use the descriptions and the examples that follow to determine which parameters to use.

- **SourceDBSnapshotIdentifier** – The identifier for the source DB snapshot.
 - If the source snapshot is in the same AWS Region as the copy, specify a valid DB snapshot identifier. For example, `rds:mysql-instance1-snapshot-20130805`.
 - If the source snapshot is in a different AWS Region than the copy, specify a valid DB snapshot ARN. For example, `arn:aws:rds:us-west-2:123456789012:snapshot:mysql-instance1-snapshot-20130805`.
 - If you are copying from a shared manual DB snapshot, this parameter must be the Amazon Resource Name (ARN) of the shared DB snapshot.
 - If you are copying an encrypted snapshot this parameter must be in the ARN format for the source AWS Region, and must match the `SourceDBSnapshotIdentifier` in the `PreSignedUrl` parameter.
- **TargetDBSnapshotIdentifier** – The identifier for the new copy of the encrypted DB snapshot.
- **CopyTags** – Set this parameter to `true` to copy tags and values from the snapshot to the copy of the snapshot. The default is `false`.
- **OptionGroupName** – The option group to associate with the copy of the snapshot.

Specify this parameter if you are copying a snapshot from one AWS Region to another, and your DB instance uses a non-default option group.

If your source DB instance uses Transparent Data Encryption for Oracle or Microsoft SQL Server, you must specify this parameter when copying across regions. For more information, see [Option Group Considerations \(p. 270\)](#).

- **KmsKeyId** – The AWS KMS key ID for an encrypted DB snapshot. The KMS key ID is the Amazon Resource Name (ARN), KMS key identifier, or the KMS key alias for the KMS encryption key.
 - If you copy an encrypted DB snapshot from your AWS account, you can specify a value for this parameter to encrypt the copy with a new KMS encryption key. If you don't specify a value for this parameter, then the copy of the DB snapshot is encrypted with the same KMS key as the source DB snapshot.
 - If you copy an encrypted DB snapshot that is shared from another AWS account, then you must specify a value for this parameter.

- If you specify this parameter when you copy an unencrypted snapshot, the copy is encrypted.
- If you copy an encrypted snapshot to a different AWS Region, then you must specify a KMS key for the destination AWS Region. KMS encryption keys are specific to the AWS Region that they are created in, and you cannot use encryption keys from one AWS Region in another AWS Region.
- **PreSignedUrl** – The URL that contains a Signature Version 4 signed request for the `CopyDBSnapshot` API operation in the source AWS Region that contains the source DB snapshot to copy.

You must specify this parameter when you copy an encrypted DB snapshot from another AWS Region by using the Amazon RDS API. You can specify the `sourceRegion` option instead of this parameter when you copy an encrypted DB snapshot from another AWS Region by using the AWS CLI.

The presigned URL must be a valid request for the `CopyDBSnapshot` API operation that can be executed in the source AWS Region that contains the encrypted DB snapshot to be copied. The presigned URL request must contain the following parameter values:

- `DestinationRegion` - The AWS Region that the encrypted DB snapshot will be copied to. This AWS Region is the same one where the `CopyDBSnapshot` action is called that contains this presigned URL.

For example, if you copy an encrypted DB snapshot from the us-west-2 region to the us-east-1 region, then you call the `CopyDBSnapshot` action in the us-east-1 region and provide a presigned URL that contains a call to the `CopyDBSnapshot` action in the us-west-2 region. For this example, the `DestinationRegion` in the presigned URL must be set to the us-east-1 region.

- `KmsKeyId` - The KMS key identifier for the key to use to encrypt the copy of the DB snapshot in the destination AWS Region. This is the same identifier for both the `CopyDBSnapshot` action that is called in the destination AWS Region, and the action contained in the presigned URL.
- `SourceDBSnapshotIdentifier` - The DB snapshot identifier for the encrypted snapshot to be copied. This identifier must be in the Amazon Resource Name (ARN) format for the source AWS Region. For example, if you are copying an encrypted DB snapshot from the us-west-2 region, then your `SourceDBSnapshotIdentifier` looks like the following example: `arn:aws:rds:us-west-2:123456789012:snapshot:mysql-instance1-snapshot-20161115`.

For more information on Signature Version 4 signed requests, see the following:

- [Authenticating Requests: Using Query Parameters \(AWS Signature Version 4\)](#) in the Amazon Simple Storage Service API Reference
- [Signature Version 4 Signing Process](#) in the AWS General Reference

Example From Unencrypted, To Same Region

The following code creates a copy of a snapshot, with the new name `mydbsnapshotcopy`, in the same AWS Region as the source snapshot. When the copy is made, all tags on the original snapshot are copied to the snapshot copy.

```
https://rds.us-west-1.amazonaws.com/
?Action=CopyDBSnapshot
&CopyTags=true
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&SourceDBSnapshotIdentifier=mysql-instance1-snapshot-20130805
&TargetDBSnapshotIdentifier=mydbsnapshotcopy
&Version=2013-09-09
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140429/us-west-1/rds/aws4_request
&X-Amz-Date=20140429T175351Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=9164337efa99caf850e874a1cb7ef62f3cea29d0b448b9e0e7c53b288ddffed2
```

Example From Unencrypted, Across Regions

The following code creates a copy of a snapshot, with the new name `mydbsnapshotcopy`, in the `us-west-1` region.

```
https://rds.us-west-1.amazonaws.com/
?Action=CopyDBSnapshot
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&SourceDBSnapshotIdentifier=arn%3Aaws%3Ards%3Aus-east-1%3A123456789012%3Asnapshot%3Amysql-
instance1-snapshot-20130805
&TargetDBSnapshotIdentifier=mydbsnapshotcopy
&Version=2013-09-09
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140429/us-west-1/rds/aws4_request
&X-Amz-Date=20140429T175351Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=9164337efa99caf850e874a1cb7ef62f3cea29d0b448b9e0e7c53b288ddffed2
```

Example From Encrypted, Across Regions

The following code creates a copy of a snapshot, with the new name `mydbsnapshotcopy`, in the `us-east-1` region.

```
https://rds.us-east-1.amazonaws.com/
?Action=CopyDBSnapshot
&KmsKeyId=my-us-east-1-key
&OptionGroupName=custom-option-group-name
&PreSignedUrl=https%253A%252F%252Frds.us-west-2.amazonaws.com%252F
%25253FAction%253DCopyDBSnapshot
%2526DestinationRegion%253Dus-east-1
%2526KmsKeyId%253Dmy-us-east-1-key
%2526SourceDBSnapshotIdentifier%253Darn%25253Aaws%25253Ards%25253Aus-
west-2%25253A123456789012%25253Asnapshot%25253Amysql-instance1-snapshot-20161115
%2526SignatureMethod%253DHmacSHA256
%2526SignatureVersion%253D4
%2526Version%253D2014-10-31
%2526X-Amz-Algorithm%253DAWS4-HMAC-SHA256
%2526X-Amz-Credential%253DAKIADQKE4SARGYLE%252F20161117%252Fus-west-2%252Frds
%252Faws4_request
%2526X-Amz-Date%253D20161117T215409Z
%2526X-Amz-Expires%253D3600
%2526X-Amz-SignedHeaders%253Dcontent-type%253Bhost%253Buser-agent%253Bx-amz-
content-sha256%253Bx-amz-date
%2526X-Amz-Signature
%253D255a0f17b4e717d3b67fad163c3ec26573b882c03a65523522cf890a67fcfa613
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&SourceDBSnapshotIdentifier=arn%3Aaws%3Ards%3Aus-west-2%3A123456789012%3Asnapshot
%3Amysql-instance1-snapshot-20161115
&TargetDBSnapshotIdentifier=mydbsnapshotcopy
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20161117/us-east-1/rds/aws4_request
&X-Amz-Date=20161117T221704Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=da4f2da66739d2e722c85fcfd225dc27bba7e2b8dbea8d8612434378e52adccf
```

Sharing a DB Snapshot

Using Amazon RDS, you can share a manual DB snapshot in the following ways:

- Sharing a manual DB snapshot, whether encrypted or unencrypted, enables authorized AWS accounts to copy the snapshot.
- Sharing an unencrypted manual DB snapshot enables authorized AWS accounts to directly restore a DB instance from the snapshot instead of taking a copy of it and restoring from that. However, you can't restore a DB instance from a DB snapshot that is both shared and encrypted. Instead, you can make a copy of the DB snapshot and restore the DB instance from the copy.

Note

To share an automated DB snapshot, create a manual DB snapshot by copying the automated snapshot, and then share that copy.

For more information on copying a snapshot, see [Copying a Snapshot \(p. 269\)](#). For more information on restoring a DB instance from a DB snapshot, see [Restoring from a DB Snapshot \(p. 266\)](#).

You can share a manual snapshot with up to 20 other AWS accounts. You can also share an unencrypted manual snapshot as public, which makes the snapshot available to all AWS accounts. Take care when sharing a snapshot as public so that none of your private information is included in any of your public snapshots.

The following limitations apply when sharing manual snapshots with other AWS accounts:

- When you restore a DB instance from a shared snapshot using the AWS Command Line Interface (AWS CLI) or Amazon RDS API, you must specify the Amazon Resource Name (ARN) of the shared snapshot as the snapshot identifier.
- You cannot share a DB snapshot that uses an option group with permanent or persistent options.

A *permanent option* cannot be removed from an option group. Option groups with persistent options cannot be removed from a DB instance once the option group has been assigned to the DB instance.

The following table lists permanent and persistent options and their related DB engines.

Option Name	Persistent	Permanent	DB Engine
TDE	Yes	No	Microsoft SQL Server Enterprise Edition
TDE	Yes	Yes	Oracle Enterprise Edition
Timezone	Yes	Yes	Oracle Enterprise Edition Oracle Standard Edition Oracle Standard Edition One Oracle Standard Edition Two

For Oracle DB instances, you can copy shared DB snapshots that have the options `Timezone` or `OLS` (or both). To do so, specify a target option group that includes these options when you copy the DB snapshot. The OLS option is permanent and persistent only for Oracle DB instances running Oracle version 12.2 or higher. For more information about these options, see [Oracle Time Zone \(p. 1071\)](#) and [Oracle Label Security \(p. 1040\)](#).

Sharing an Encrypted Snapshot

You can share DB snapshots that have been encrypted "at rest" using the AES-256 encryption algorithm, as described in [Encrypting Amazon RDS Resources \(p. 429\)](#). To do this, you must take the following steps:

1. Share the AWS Key Management Service (AWS KMS) encryption key that was used to encrypt the snapshot with any accounts that you want to be able to access the snapshot.

You can share AWS KMS encryption keys with another AWS account by adding the other account to the KMS key policy. For details on updating a key policy, see [Key Policies](#) in the *AWS KMS Developer Guide*. For an example of creating a key policy, see [Allowing Access to an AWS KMS Encryption Key \(p. 279\)](#) later in this topic.

2. Use the AWS Management Console, AWS CLI, or Amazon RDS API to share the encrypted snapshot with the other accounts.

These restrictions apply to sharing encrypted snapshots:

- You can't share encrypted snapshots as public.
- You can't share Oracle or Microsoft SQL Server snapshots that are encrypted using Transparent Data Encryption (TDE).
- You can't share a snapshot that has been encrypted using the default AWS KMS encryption key of the AWS account that shared the snapshot.

Allowing Access to an AWS KMS Encryption Key

For another AWS account to copy an encrypted DB snapshot shared from your account, the account that you share your snapshot with must have access to the KMS key that encrypted the snapshot. To allow another AWS account access to an AWS KMS key, update the key policy for the KMS key with the ARN of the AWS account that you are sharing to as a `Principal` in the KMS key policy, and then allow the `kms:CreateGrant` action.

After you have given an AWS account access to your KMS encryption key, to copy your encrypted snapshot, that AWS account must create an AWS Identity and Access Management (IAM) user if it doesn't already have one. In addition, that AWS account must also attach an IAM policy to that IAM user that allows the IAM user to copy an encrypted DB snapshot using your KMS key. The account must be an IAM user and cannot be a root AWS account identity due to KMS security restrictions.

In the following key policy example, user `111122223333` is the owner of the KMS encryption key, and user `444455556666` is the account that the key is being shared with. This updated key policy gives the AWS account access to the KMS key by including the ARN for the root AWS account identity for user `444455556666` as a `Principal` for the policy, and by allowing the `kms:CreateGrant` action.

```
{  
    "Id": "key-policy-1",  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Sid": "Allow use of the key",  
            "Effect": "Allow",  
            "Principal": {"AWS": [  
                "arn:aws:iam::111122223333:user/KeyUser",  
                "arn:aws:iam::444455556666:root"  
            ]},  
            "Action": [  
                "kms:CreateGrant",  
                "kms:Encrypt",  
                "kms:Decrypt",  
                "kms:ReEncrypt",  
                "kms:GenerateDataKey",  
                "kms:DescribeKey",  
                "kms:ListAliases",  
                "kms:ListGrants",  
                "kms:ListKeys",  
                "kms:PutGrant",  
                "kms:RevokeGrant"  
            ]  
        }  
    ]  
}
```

```

        "kms:Decrypt",
        "kms:ReEncrypt*",
        "kms:GenerateDataKey*",
        "kms:DescribeKey"
    ],
    "Resource": "*"
},
{
    "Sid": "Allow attachment of persistent resources",
    "Effect": "Allow",
    "Principal": {"AWS": [
        "arn:aws:iam::111122223333:user/KeyUser",
        "arn:aws:iam::444455556666:root"
    ]},
    "Action": [
        "kms>CreateGrant",
        "kms>ListGrants",
        "kms:RevokeGrant"
    ],
    "Resource": "*",
    "Condition": {"Bool": {"kms:GrantIsForAWSResource": true}}
}
]
}

```

Creating an IAM Policy to Enable Copying of the Encrypted Snapshot

Once the external AWS account has access to your KMS key, the owner of that AWS account can create a policy that allows an IAM user created for that account to copy an encrypted snapshot encrypted with that KMS key.

The following example shows a policy that can be attached to an IAM user for AWS account 444455556666 that enables the IAM user to copy a shared snapshot from AWS account 111122223333 that has been encrypted with the KMS key c989c1dd-a3f2-4a5d-8d96-e793d082ab26 in the us-west-2 region.

```

{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "AllowUseOfTheKey",
            "Effect": "Allow",
            "Action": [
                "kms:Encrypt",
                "kms:Decrypt",
                "kms:ReEncrypt*",
                "kms:GenerateDataKey*",
                "kms:DescribeKey",
                "kms>CreateGrant",
                "kms:RetireGrant"
            ],
            "Resource": ["arn:aws:kms:us-west-2:111122223333:key/c989c1dd-a3f2-4a5d-8d96-e793d082ab26"]
        },
        {
            "Sid": "AllowAttachmentOfPersistentResources",
            "Effect": "Allow",
            "Action": [
                "kms>CreateGrant",
                "kms>ListGrants",
                "kms:RevokeGrant"
            ],
        }
    ]
}

```

```
    "Resource": [ "arn:aws:kms:us-west-2:111122223333:key/c989c1dd-a3f2-4a5d-8d96-e793d082ab26" ],
      "Condition": {
        "Bool": {
          "kms:GrantIsForAWSResource": true
        }
      }
    ]
}
```

For details on updating a key policy, see [Key Policies in the AWS KMS Developer Guide](#).

Sharing a Snapshot

You can share a DB snapshot using the AWS Management Console, the AWS CLI, or the RDS API.

Console

Using the Amazon RDS console, you can share a manual DB snapshot with up to 20 AWS accounts. You can also use the console to stop sharing a manual snapshot with one or more accounts.

To share a manual DB snapshot by using the Amazon RDS console

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Snapshots**.
3. Select the manual snapshot that you want to share.
4. For **Actions**, choose **Share Snapshot**.
5. Choose one of the following options for **DB snapshot visibility**.
 - If the source is unencrypted, choose **Public** to permit all AWS accounts to restore a DB instance from your manual DB snapshot, or choose **Private** to permit only AWS accounts that you specify to restore a DB instance from your manual DB snapshot.

Warning

If you set **DB snapshot visibility** to **Public**, all AWS accounts can restore a DB instance from your manual DB snapshot and have access to your data. Do not share any manual DB snapshots that contain private information as **Public**.

- If the source is encrypted, **DB snapshot visibility** is set as **Private** because encrypted snapshots can't be shared as public.
6. For **AWS Account ID**, type the AWS account identifier for an account that you want to permit to restore a DB instance from your manual snapshot, and then choose **Add**. Repeat to include additional AWS account identifiers, up to 20 AWS accounts.

If you make an error when adding an AWS account identifier to the list of permitted accounts, you can delete it from the list by choosing **Delete** at the right of the incorrect AWS account identifier.

Snapshot permissions

Preferences

You are sharing an unencrypted DB snapshot. When you share an unencrypted DB snapshot, you give the other account permission to make a copy of the DB snapshot and to restore a database from your DB snapshot.

DB snapshot
testoracletags-snap

DB snapshot visibility

Private
 Public

AWS account ID

AWS account ID	Delete
Please add AWS account ID	

- After you have added identifiers for all of the AWS accounts that you want to permit to restore the manual snapshot, choose **Save** to save your changes.

To stop sharing a manual DB snapshot with an AWS account

- Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
- In the navigation pane, choose **Snapshots**.
- Select the manual snapshot that you want to stop sharing.
- Choose **Actions**, and then choose **Share Snapshot**.
- To remove permission for an AWS account, choose **Delete** for the AWS account identifier for that account from the list of authorized accounts.

Snapshot permissions

Preferences

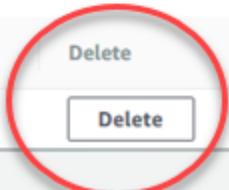
You are sharing an unencrypted DB snapshot. When you share an unencrypted DB snapshot, you give the other account permission to make a copy of the DB snapshot and to restore a database from your DB snapshot.

DB snapshot
testoracletags-snap

DB snapshot visibility
 Private
 Public

AWS account ID

AWS account ID



6. Choose **Save** to save your changes.

AWS CLI

To share a DB snapshot, use the `aws rds modify-db-snapshot-attribute` command. Use the `--values-to-add` parameter to add a list of the IDs for the AWS accounts that are authorized to restore the manual snapshot.

The following example permits two AWS account identifiers, 123451234512 and 123456789012, to restore the DB snapshot named `manual-snapshot1`, and removes the `all` attribute value to mark the snapshot as private.

```
aws rds modify-db-snapshot-attribute \
--db-snapshot-identifier manual-snapshot1 \
--attribute-name restore \
--values-to-add '[ "111122223333", "444455556666" ]'
```

To remove an AWS account identifier from the list, use the `--values-to-remove` parameter. The following example prevents AWS account ID 444455556666 from restoring the snapshot.

```
aws rds modify-db-snapshot-attribute \
--db-snapshot-identifier manual-snapshot1 \
--attribute-name restore \
--values-to-remove '[ "444455556666" ]'
```

RDS API

You can also share a manual DB snapshot with other AWS accounts by using the Amazon RDS API. To do so, call the [ModifyDBSnapshotAttribute](#) operation. Specify `restore` for `AttributeName`, and use the `ValuesToAdd` parameter to add a list of the IDs for the AWS accounts that are authorized to restore the manual snapshot.

To make a manual snapshot public and restorable by all AWS accounts, use the value `all`. However, take care not to add the `all` value for any manual snapshots that contain private information that you don't want to be available to all AWS accounts. Also, don't specify `all` for encrypted snapshots, because making such snapshots public isn't supported.

To remove sharing permission for an AWS account, use the [ModifyDBSnapshotAttribute](#) operation with `AttributeName` set to `restore` and the `ValuesToRemove` parameter. To mark a manual snapshot as private, remove the value `all` from the values list for the `restore` attribute.

To list all of the AWS accounts permitted to restore a snapshot, use the [DescribeDBSnapshotAttributes](#) API operation.

Restoring a DB Instance to a Specified Time

You can restore a DB instance to a specific point in time, creating a new DB instance. When you restore a DB instance to a point in time, the default DB security group is applied to the new DB instance. If you need custom DB security groups applied to your DB instance, you must apply them explicitly using the AWS Management Console, the AWS CLI `modify-db-instance` command, or the Amazon RDS API `ModifyDBInstance` operation after the DB instance is available.

RDS uploads transaction logs for DB instances to Amazon S3 every 5 minutes. To determine the latest restorable time for a DB instance, use the AWS CLI `describe-db-instances` command and look at the value returned in the `LatestRestorableTime` field for the DB instance. In the AWS Management Console, this property is visible as the **Latest restore time** for the DB instance. You can restore to any point in time during your backup retention period.

Several of the database engines used by Amazon RDS have special considerations when restoring from a point in time. When you restore an Oracle DB instance to a point in time, you can specify a different Oracle DB engine, license model, and DBName (SID) to be used by the new DB instance. When you restore a SQL Server DB instance to a point in time, each database within that instance is restored to a point in time within 1 second of each other database within the instance. Transactions that span multiple databases within the instance may be restored inconsistently. Also, for a SQL Server DB instance, the `OFFLINE`, `EMERGENCY`, and `SINGLE_USER` modes are not currently supported. Setting any database into one of these modes will cause the latest restorable time to stop moving ahead for the whole instance.

Some actions, such as changing the recovery model of a SQL Server database, can break the sequence of logs that are used for point-in-time recovery. In some cases, Amazon RDS can detect this issue and the latest restorable time is prevented from moving forward; in other cases, such as when a SQL Server database uses the `BULK_LOGGED` recovery model, the break in log sequence is not detected. It may not be possible to restore a SQL Server DB instance to a point in time if there is a break in the log sequence. For these reasons, Amazon RDS does not support changing the recovery model of SQL Server databases.

You can restore a DB instance to a point in time using the AWS Management Console, the AWS CLI, or the RDS API.

Console

To restore a DB instance to a specified time

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the DB instance that you want to restore.
4. For **Actions**, choose **Restore to point in time**.

The **Launch DB Instance** window appears.

5. Choose **Latest restorable time** to restore to the latest possible time, or choose **Custom** to choose a time.

If you chose **Custom**, enter the date and time that you want to restore the instance to.
6. For **DB instance identifier**, enter the name of the restored DB instance, and then complete the other options.
7. Choose **Launch DB Instance**.

AWS CLI

To restore a DB instance to a specified time, use the AWS CLI command [restore-db-instance-to-point-in-time](#) to create a new DB instance.

Example

For Linux, OS X, or Unix:

```
aws rds restore-db-instance-to-point-in-time \
--source-db-instance-identifier mysourcedbinstance \
--target-db-instance-identifier mytargetdbinstance \
--restore-time 2017-10-14T23:45:00.000Z
```

For Windows:

```
aws rds restore-db-instance-to-point-in-time ^
--source-db-instance-identifier mysourcedbinstance ^
--target-db-instance-identifier mytargetdbinstance ^
--restore-time 2017-10-14T23:45:00.000Z
```

RDS API

To restore a DB instance to a specified time, call the Amazon RDS API

[RestoreDBInstanceToPointInTime](#) operation with the following parameters:

- `SourceDBInstanceIdentifier`
- `TargetDBInstanceIdentifier`
- `RestoreTime`

Deleting a Snapshot

You can delete DB snapshots managed by Amazon RDS when you no longer need them.

Note

To delete backups managed by AWS Backup, use the AWS Backup console. For information about AWS Backup, see the [AWS Backup Developer Guide](#).

Deleting a DB Snapshot

You can delete a manual, shared, or public DB snapshot using the AWS Management Console, the AWS CLI, or the RDS API.

To delete a shared or public snapshot, you must sign in to the AWS account that owns the snapshot.

If you have automated DB snapshots that you want to delete without deleting the DB instance, change the backup retention period for the DB instance to 0. The automated snapshots are deleted when the change is applied. You can apply the change immediately if you don't want to wait until the next maintenance period. After the change is complete, you can then re-enable automatic backups by setting the backup retention period to a number greater than 0. For information about modifying a DB instance, see [Modifying an Amazon RDS DB Instance \(p. 150\)](#).

If you deleted a DB instance, you can delete its automated DB snapshots by removing the automated backups for the DB instance. For information about automated backups, see [Working With Backups \(p. 254\)](#).

Console

To delete a DB snapshot

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Snapshots**.
3. Choose the DB snapshot that you want to delete.
4. For **Actions**, choose **Delete Snapshot**.
5. Choose **Delete** on the confirmation page.

AWS CLI

You can delete a DB snapshot by using the AWS CLI command `delete-db-snapshot`.

The following options are used to delete a DB snapshot.

- `--db-snapshot-identifier` – The identifier for the DB snapshot.

Example

The following code deletes the `mydbsnapshot` DB snapshot.

For Linux, OS X, or Unix:

```
aws rds delete-db-snapshot \
--db-snapshot-identifier mydbsnapshot
```

For Windows:

```
aws rds delete-db-snapshot ^  
--db-snapshot-identifier mydbsnapshot
```

RDS API

You can delete a DB snapshot by using the Amazon RDS API operation [DeleteDBSnapshot](#).

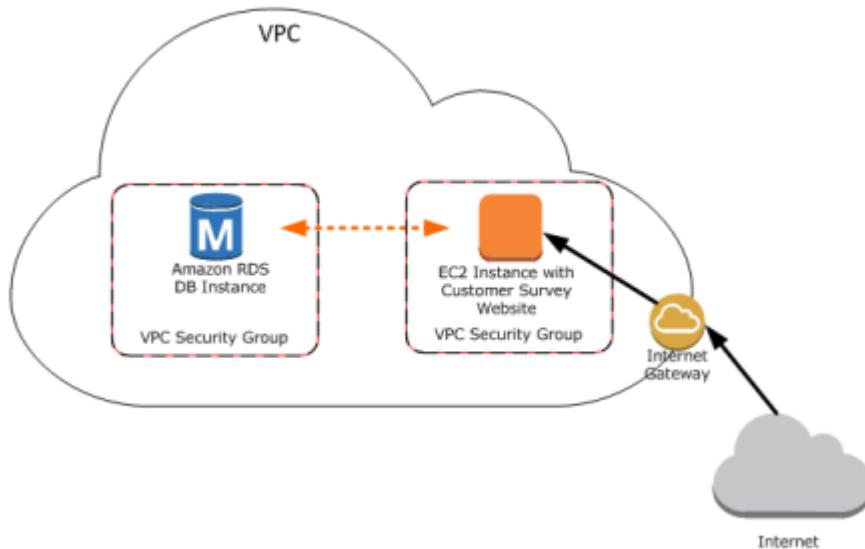
The following parameters are used to delete a DB snapshot.

- **DBSnapshotIdentifier** – The identifier for the DB snapshot.

Tutorial: Restore a DB Instance from a DB Snapshot

A common scenario when working with Amazon RDS is to have a DB instance that you work with occasionally but that you don't need full time. For example, you might have a quarterly customer survey that uses an Amazon Elastic Compute Cloud (Amazon EC2) instance to host a customer survey website and a DB instance that is used to store the survey results. One way to save money on such a scenario is to take a DB snapshot of the DB instance after the survey is completed, delete the DB instance, and then restore the DB instance when you need to conduct the survey again.

In the following illustration, you can see a possible scenario where an EC2 instance hosting a customer survey website is in the same Amazon Virtual Private Cloud (Amazon VPC) as a DB instance that retains the customer survey data. Note that each instance has its own security group; the EC2 instance security group allows access from the internet while the DB instance security group allows access only to and from the EC2 instance. When the survey is done, the EC2 instance can be stopped and the DB instance can be deleted after a final DB snapshot is created. When you need to conduct another survey, you can restart the EC2 instance and restore the DB instance from the DB snapshot.



For information about how to set up the needed VPC security groups for this scenario that allows the EC2 instance to connect with the DB instance, see [A DB Instance in a VPC Accessed by an EC2 Instance in the Same VPC \(p. 507\)](#).

You must create a DB snapshot before you can restore a DB instance from one. When you restore the DB instance, you provide the name of the DB snapshot to restore from, and then provide a name for the new DB instance that is created from the restore operation. You cannot restore from a DB snapshot to an existing DB instance; a new DB instance is created when you restore.

Prerequisites for Restoring a DB Instance from a DB Snapshot

Some settings on the restored DB instance are reset when the instance is restored, so you must retain the original resources to be able to restore the DB instance to its previous settings. For example, when you restore a DB instance from a DB snapshot, the default DB parameter and a default security group are

associated with the restored instance. That association means that the default security group does not allow access to the DB instance, and no custom parameter settings are available in the default parameter group. You need to retain the DB parameter group and security group associated with the DB instance that was used to create the DB snapshot.

The following are required before you can restore a DB instance from a DB snapshot:

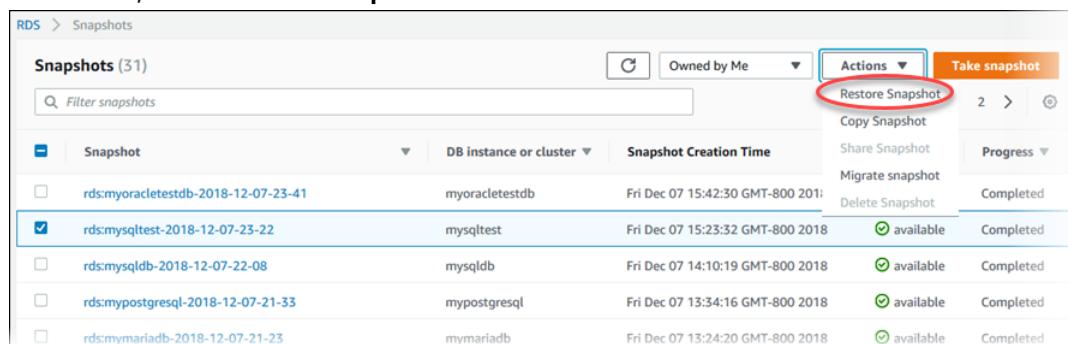
- You must have created a DB snapshot of a DB instance before you can restore a DB instance from that DB snapshot. For more information about creating a DB snapshot, see [Creating a DB Snapshot \(p. 264\)](#).
- You must retain the parameter group and security group associated with the DB instance you created the DB snapshot from.
- You need to determine the correct option group for the restored DB instance:
 - The option group associated with the DB snapshot that you restore from is associated with the restored DB instance once it is created. For example, if the DB snapshot you restore from uses Oracle Transparent Data Encryption (TDE), the restored DB instance uses the same option group, which had the TDE option.
 - You cannot use the option group associated with the original DB instance if you attempt to restore that instance into a different VPC or into a different platform. This restriction occurs because when an option group is assigned to a DB instance, it is also linked to the platform that the DB instance is on, either VPC or EC2-Classic (non-VPC). If a DB instance is in a VPC, the option group associated with the instance is linked to that VPC.
 - If you restore a DB instance into a different VPC or onto a different platform, you must either assign the default option group to the instance, assign an option group that is linked to that VPC or platform, or create a new option group and assign it to the DB instance. Note that with persistent or permanent options, such as Oracle TDE, you must create a new option group that includes the persistent or permanent option when restoring a DB instance into a different VPC. For more information about working with option groups, see [Working with Option Groups \(p. 196\)](#).

Restoring a DB Instance from a DB Snapshot

You can use the procedure following to restore from a snapshot in the AWS Management Console.

To restore a DB instance from a DB snapshot

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Snapshots**.
3. Choose the DB snapshot that you want to restore from.
4. For **Actions**, choose **Restore Snapshot**.



The **Restore DB Instance** page appears.

5. For **DB Instance Identifier** under **Settings**, enter the name that you want to use for the restored DB instance. If you are restoring from a DB instance that you deleted after you made the DB snapshot, you can use the name of that DB instance.
6. Choose **Restore DB Instance**.

Modifying a Restored DB Instance

As soon as the restore operation is complete, you should associate the custom security group used by the instance you restored from with any applicable custom DB parameter group that you might have. Only the default DB parameter and security groups are associated with the restored instance. If you want to restore the functionality of the DB instance to that of the DB instance that the snapshot was created from, you must modify the DB instance to use the security group and parameter group used by the previous DB instance.

You must apply any changes explicitly using the RDS console's **Modify** command, the `ModifyDBInstance` API, or the `aws rds modify-db-instance` command line tool, once the DB instance is available. We recommend that you retain parameter groups for any DB snapshots you have so that you can associate a restored instance with the correct parameter file.

You can modify other settings on the restored DB instance. For example, you can use a different storage type than the source DB snapshot. In this case the restoration process is slower because of the additional work required to migrate the data to the new storage type. In the case of restoring to or from Magnetic (Standard) storage, the migration process is the slowest, because Magnetic storage does not have the IOPS capability of Provisioned IOPS or General Purpose (SSD) storage.

The next steps assume that your DB instance is in a VPC. If your DB instance is not in a VPC, use the AWS Management Console to locate the DB security group you need for the DB instance.

To modify a restored DB instance to have the settings of the original DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the name of the DB instance created when you restored from the DB snapshot to display its details. Choose the **Connectivity** tab. The security group assigned to the DB instance might not allow access. If there are no inbound rules, no permissions exist that allow inbound access.

The screenshot shows the 'Connectivity' tab selected in the top navigation bar. Below it, the 'Security' section is highlighted with a red box. This section displays the following details:

Endpoint & port	Networking	Security
Endpoint restored-db-instance.cahj48jpj34l.us-west-2.rds.amazonaws.com	Availability zone us-west-2a	VPC security groups (active)
Port 3306	VPC	Public accessibility Yes
	Subnet group default	Certificate authority rds-ca-2015
	Subnets	Certificate authority d... Mar 5th, 2020

Security group (1)

Filter security group rules

Security group	Type
[Redacted]	No applicable security group roles

4. Choose **Modify**.
5. In the **Network & Security** section, choose the security group that you want to use for your DB instance. If you need to add rules to create a new security group to use with an EC2 instance, see [A DB Instance in a VPC Accessed by an EC2 Instance in the Same VPC \(p. 507\)](#) for more information.

You can also remove a security group by choosing the X associated with it.

Network & Security

Subnet group
Use this field to move the DB instance to a new subnet group in another VPC. [Learn more.](#)

default

Security group
List of DB security groups to associate with this DB instance.

Select security groups

default () () X

Certificate authority
Certificate authority for this DB instance

rds-ca-2015

Public accessibility [info](#)

Yes
EC2 instances and devices outside of the VPC hosting the DB instance will connect to the DB instances. You must also select one or more VPC security groups that specify which EC2 instances and devices can connect to the DB instance.

No
DB instance will not have a public IP address assigned. No EC2 instance or devices outside of the VPC will be able to connect.

6. Choose **Continue**, and then choose **Apply immediately**.
7. Choose **Modify DB Instance**.

After the instance status is available, choose the DB instance name to display its details. Choose the **Connectivity** tab, and confirm that the new security group has been applied, making the DB instance authorized for access.

Connectivity	Monitoring	Logs & events	Configuration	Maintenance & backups	Tags															
<h3>Connectivity</h3> <table border="1"> <tr> <td>Endpoint & port</td> <td>Networking</td> <td>Security</td> </tr> <tr> <td>Endpoint restored-db-instance.cahj48jpj34l.us-west-2.rds.amazonaws.com</td> <td>Availability zone us-west-2a</td> <td>VPC security groups rds-launch-wizard-9 (active)</td> </tr> <tr> <td>Port 3306</td> <td>VPC</td> <td>Public accessibility Yes</td> </tr> <tr> <td></td> <td>Subnet group default</td> <td>Certificate authority rds-ca-2015</td> </tr> <tr> <td></td> <td>Subnets</td> <td>Certificate authority date Mar 5th, 2020</td> </tr> </table>						Endpoint & port	Networking	Security	Endpoint restored-db-instance.cahj48jpj34l.us-west-2.rds.amazonaws.com	Availability zone us-west-2a	VPC security groups rds-launch-wizard-9 (active)	Port 3306	VPC	Public accessibility Yes		Subnet group default	Certificate authority rds-ca-2015		Subnets	Certificate authority date Mar 5th, 2020
Endpoint & port	Networking	Security																		
Endpoint restored-db-instance.cahj48jpj34l.us-west-2.rds.amazonaws.com	Availability zone us-west-2a	VPC security groups rds-launch-wizard-9 (active)																		
Port 3306	VPC	Public accessibility Yes																		
	Subnet group default	Certificate authority rds-ca-2015																		
	Subnets	Certificate authority date Mar 5th, 2020																		

Monitoring Amazon RDS

This section shows you how to monitor Amazon RDS.

Topics

- [Overview of Monitoring Amazon RDS \(p. 295\)](#)
- [Enhanced Monitoring \(p. 307\)](#)
- [Using Amazon RDS Performance Insights \(p. 319\)](#)
- [Using Amazon RDS Recommendations \(p. 365\)](#)
- [Using Amazon RDS Event Notification \(p. 370\)](#)
- [Viewing Amazon RDS Events \(p. 388\)](#)
- [Amazon RDS Database Log Files \(p. 390\)](#)
- [Logging Amazon RDS API Calls with AWS CloudTrail \(p. 423\)](#)

Overview of Monitoring Amazon RDS

Monitoring is an important part of maintaining the reliability, availability, and performance of Amazon RDS and your AWS solutions. You should collect monitoring data from all of the parts of your AWS solution so that you can more easily debug a multi-point failure if one occurs. Before you start monitoring Amazon RDS, we recommend that you create a monitoring plan that includes answers to the following questions:

- What are your monitoring goals?
- What resources will you monitor?
- How often will you monitor these resources?
- What monitoring tools will you use?
- Who will perform the monitoring tasks?
- Who should be notified when something goes wrong?

The next step is to establish a baseline for normal Amazon RDS performance in your environment, by measuring performance at various times and under different load conditions. As you monitor Amazon RDS, you should consider storing historical monitoring data. This stored data will give you a baseline to compare against with current performance data, identify normal performance patterns and performance anomalies, and devise methods to address issues.

For example, with Amazon RDS, you can monitor network throughput, I/O for read, write, and/or metadata operations, client connections, and burst credit balances for your DB instances. When performance falls outside your established baseline, you might need to change the instance class of your DB instance or the number of DB instances and Read Replicas that are available for clients in order to optimize your database availability for your workload.

In general, acceptable values for performance metrics depend on what your baseline looks like and what your application is doing. Investigate consistent or trending variances from your baseline. Advice about specific types of metrics follows:

- **High CPU or RAM consumption** – High values for CPU or RAM consumption might be appropriate, provided that they are in keeping with your goals for your application (like throughput or concurrency) and are expected.
- **Disk space consumption** – Investigate disk space consumption if space used is consistently at or above 85 percent of the total disk space. See if it is possible to delete data from the instance or archive data to a different system to free up space.
- **Network traffic** – For network traffic, talk with your system administrator to understand what expected throughput is for your domain network and Internet connection. Investigate network traffic if throughput is consistently lower than expected.
- **Database connections** – Consider constraining database connections if you see high numbers of user connections in conjunction with decreases in instance performance and response time. The best number of user connections for your DB instance will vary based on your instance class and the complexity of the operations being performed. You can determine the number of database connections by associating your DB instance with a parameter group where the `User_Connections` parameter is set to a value other than 0 (unlimited). You can either use an existing parameter group or create a new one. For more information, see [Working with DB Parameter Groups \(p. 211\)](#).
- **IOPS metrics** – The expected values for IOPS metrics depend on disk specification and server configuration, so use your baseline to know what is typical. Investigate if values are consistently different than your baseline. For best IOPS performance, make sure your typical working set will fit into memory to minimize read and write operations.

Monitoring Tools

AWS provides various tools that you can use to monitor Amazon RDS. You can configure some of these tools to do the monitoring for you, while some of the tools require manual intervention. We recommend that you automate monitoring tasks as much as possible.

Automated Monitoring Tools

You can use the following automated monitoring tools to watch Amazon RDS and report when something is wrong:

- **Amazon RDS Events** – Subscribe to Amazon RDS events to be notified when changes occur with a DB instance, DB snapshot, DB parameter group, or DB security group. For more information, see [Using Amazon RDS Event Notification \(p. 370\)](#).
- **Database log files** – View, download, or watch database log files using the Amazon RDS console or Amazon RDS API operations. You can also query some database log files that are loaded into database tables. For more information, see [Amazon RDS Database Log Files \(p. 390\)](#).
- **Amazon RDS Enhanced Monitoring** — Look at metrics in real time for the operating system. For more information, see [Enhanced Monitoring \(p. 307\)](#).

In addition, Amazon RDS integrates with Amazon CloudWatch for additional monitoring capabilities:

- **Amazon CloudWatch Metrics** – Amazon RDS automatically sends metrics to CloudWatch every minute for each active database. You are not charged additionally for Amazon RDS metrics in CloudWatch. For more information, see [the section called "Viewing DB Instance Metrics" \(p. 304\)](#).
- **Amazon CloudWatch Alarms** – You can watch a single Amazon RDS metric over a specific time period, and perform one or more actions based on the value of the metric relative to a threshold you set. For more information, see [Monitoring with Amazon CloudWatch \(p. 297\)](#)
- **Amazon CloudWatch Logs** – Most DB engines enable you to monitor, store, and access your database log files in CloudWatch Logs. For more information, see [Amazon CloudWatch Logs User Guide](#)

Manual Monitoring Tools

Another important part of monitoring Amazon RDS involves manually monitoring those items that the CloudWatch alarms don't cover. The Amazon RDS, CloudWatch, AWS Trusted Advisor and other AWS console dashboards provide an at-a-glance view of the state of your AWS environment. We recommend that you also check the log files on your DB instance.

- From the Amazon RDS console, you can monitor the following items for your resources:
 - The number of connections to a DB instance
 - The amount of read and write operations to a DB instance
 - The amount of storage that a DB instance is currently utilizing
 - The amount of memory and CPU being utilized for a DB instance
 - The amount of network traffic to and from a DB instance
- From the AWS Trusted Advisor dashboard, you can review the following cost optimization, security, fault tolerance, and performance improvement checks:
 - Amazon RDS Idle DB Instances
 - Amazon RDS Security Group Access Risk
 - Amazon RDS Backups
 - Amazon RDS Multi-AZ

For more information on these checks, see [Trusted Advisor Best Practices \(Checks\)](#).

- CloudWatch home page shows:
 - Current alarms and status
 - Graphs of alarms and resources
 - Service health status

In addition, you can use CloudWatch to do the following:

- Create [customized dashboards](#) to monitor the services you care about
- Graph metric data to troubleshoot issues and discover trends
- Search and browse all your AWS resource metrics
- Create and edit alarms to be notified of problems

Monitoring with Amazon CloudWatch

You can monitor DB instances using Amazon CloudWatch, which collects and processes raw data from Amazon RDS into readable, near real-time metrics. These statistics are recorded for a period of two weeks, so that you can access historical information and gain a better perspective on how your web application or service is performing. By default, Amazon RDS metric data is automatically sent to CloudWatch in 1-minute periods. For more information about CloudWatch, see [What Are Amazon CloudWatch, Amazon CloudWatch Events, and Amazon CloudWatch Logs?](#) in the *Amazon CloudWatch User Guide*.

Note

If you are using Amazon RDS Performance Insights, additional metrics are available. For more information, see [Performance Insights Metrics Published to Amazon CloudWatch \(p. 351\)](#).

Amazon RDS Metrics and Dimensions

When you use Amazon RDS resources, Amazon RDS sends metrics and dimensions to Amazon CloudWatch every minute. You can use the following procedures to view the metrics for Amazon RDS.

To view metrics using the Amazon CloudWatch console

Metrics are grouped first by the service namespace, and then by the various dimension combinations within each namespace.

1. Open the CloudWatch console at <https://console.aws.amazon.com/cloudwatch/>.
2. If necessary, change the AWS Region. From the navigation bar, choose the AWS Region where your AWS resources reside. For more information, see [Regions and Endpoints](#).
3. In the navigation pane, choose **Metrics**. Choose the **RDS** metric namespace.

The screenshot shows the AWS CloudWatch Metrics console. At the top, there are three tabs: "All metrics" (selected), "Graphed metrics", and "Graph options". Below the tabs, the navigation path is "All > RDS". A search bar contains the placeholder text "Search for any metric, dimension or resource id". The main area displays "168 Metrics" under three categories: "DbClusterIdentifier, EngineName" (3 Metrics), "Per-Database Metrics" (45 Metrics), and "By Database Class" (45 Metrics).

4. Choose a metric dimension, for example **By Database Class**.
5. To sort the metrics, use the column heading. To graph a metric, select the check box next to the metric. To filter by resource, choose the resource ID, and then choose **Add to search**. To filter by metric, choose the metric name, and then choose **Add to search**.

The screenshot shows the "By Database Class" filter interface. The top navigation path is "All > RDS > By Database Class". The left sidebar shows a checkbox for "DatabaseClass (45)". The main table lists metrics for different database classes. The first row, "db.r4.large BufferCacheHitRatio", has a checked checkbox. A context menu is open over this row, listing options: "Add to search", "Search for this only", "Remove from graph", "Graph this metric only", "Graph all search results", and "What is this?".

To view metrics using the AWS CLI

- At a command prompt, use the following command:

```
aws cloudwatch list-metrics --namespace AWS/RDS
```

Amazon RDS Metrics

The AWS/RDS namespace includes the following metrics.

Metric	Description
BinLogDiskUsage	<p>The amount of disk space occupied by binary logs on the master. Applies to MySQL read replicas.</p> <p>Units: Bytes</p>
BurstBalance	<p>The percent of General Purpose SSD (gp2) burst-bucket I/O credits available.</p> <p>Units: Percent</p>
CPUUtilization	<p>The percentage of CPU utilization.</p> <p>Units: Percent</p>
CPUCreditUsage	<p>(T2 instances) The number of CPU credits spent by the instance for CPU utilization. One CPU credit equals one vCPU running at 100 percent utilization for one minute or an equivalent combination of vCPUs, utilization, and time. For example, you might have one vCPU running at 50 percent utilization for two minutes or two vCPUs running at 25 percent utilization for two minutes.</p> <p>CPU credit metrics are available at a five-minute frequency only. If you specify a period greater than five minutes, use the <code>Sum</code> statistic instead of the <code>Average</code> statistic.</p> <p>Units: Credits (vCPU-minutes)</p>
CPUCreditBalance	<p>(T2 instances) The number of earned CPU credits that an instance has accrued since it was launched or started. For T2 Standard, the CPUCreditBalance also includes the number of launch credits that have been accrued.</p> <p>Credits are accrued in the credit balance after they are earned, and removed from the credit balance when they are spent. The credit balance has a maximum limit, determined by the instance size. After the limit is reached, any new credits that are earned are discarded. For T2 Standard, launch credits don't count towards the limit.</p> <p>The credits in the CPUCreditBalance are available for the instance to spend to burst beyond its baseline CPU utilization.</p> <p>When an instance is running, credits in the CPUCreditBalance don't expire. When the instance stops, the CPUCreditBalance does not persist, and all accrued credits are lost.</p> <p>CPU credit metrics are available at a five-minute frequency only.</p>

Metric	Description
	Units: Credits (vCPU-minutes)
DatabaseConnections	The number of database connections in use. Units: Count
DiskQueueDepth	The number of outstanding IOs (read/write requests) waiting to access the disk. Units: Count
FailedSQLServerAgentJobs	The number of failed SQL Server Agent jobs during the last minute. Unit: Count/Minute
FreeableMemory	The amount of available random access memory. For MariaDB, MySQL, Oracle, and PostgreSQL DB instances, this metric reports the value of the <code>MemAvailable</code> field of <code>/proc/meminfo</code> . Units: Bytes
FreeStorageSpace	The amount of available storage space. Units: Bytes
MaximumUsedTransactionIDs	The maximum transaction ID that has been used. Applies to PostgreSQL. Units: Count
NetworkReceiveThroughput	The incoming (Receive) network traffic on the DB instance, including both customer database traffic and Amazon RDS traffic used for monitoring and replication. Units: Bytes/Second
NetworkTransmitThroughput	The outgoing (Transmit) network traffic on the DB instance, including both customer database traffic and Amazon RDS traffic used for monitoring and replication. Units: Bytes/Second
OldestReplicationSlotLag	The lagging size of the replica lagging the most in terms of WAL data received. Applies to PostgreSQL. Units: Megabytes
ReadIOPS	The average number of disk read I/O operations per second. Units: Count/Second
ReadLatency	The average amount of time taken per disk I/O operation. Units: Seconds
ReadThroughput	The average number of bytes read from disk per second. Units: Bytes/Second

Metric	Description
ReplicaLag	The amount of time a Read Replica DB instance lags behind the source DB instance. Applies to MySQL, MariaDB, and PostgreSQL Read Replicas. Units: Seconds
ReplicationSlotDiskUsage	The disk space used by replication slot files. Applies to PostgreSQL. Units: Megabytes
SwapUsage	The amount of swap space used on the DB instance. This metric is not available for SQL Server. Units: Bytes
TransactionLogsDiskUsage	The disk space used by transaction logs. Applies to PostgreSQL. Units: Megabytes
TransactionLogsGeneration	The size of transaction logs generated per second. Applies to PostgreSQL. Units: Bytes/Second
WriteIOPS	The average number of disk write I/O operations per second. Units: Count/Second
WriteLatency	The average amount of time taken per disk I/O operation. Units: Seconds
WriteThroughput	The average number of bytes written to disk per second. Units: Bytes/Second

Amazon RDS Dimensions

Amazon RDS metrics data can be filtered by using any of the dimensions in the following table:

Dimension	Description
DBInstanceIdentifier	This dimension filters the data you request for a specific database instance.
DBClusterIdentifier	This dimension filters the data you request for a specific Amazon Aurora DB cluster.
DBClusterIdentifier, Role	This dimension filters the data you request for a specific Aurora DB cluster, aggregating the metric by instance role (WRITER/READER). For example, you can aggregate metrics for all READER instances that belong to a cluster.
DatabaseClass	This dimension filters the data you request for all instances in a database class. For example, you can aggregate metrics for all instances that belong to the database class db.m1.small

Dimension	Description
EngineName	This dimension filters the data you request for the identified engine name only. For example, you can aggregate metrics for all instances that have the engine name mysql.
SourceRegion	This dimension filters the data you request for the specified region only. For example, you can aggregate metrics for all instances in the region us-east-1.

Creating CloudWatch Alarms to Monitor Amazon RDS

You can create a CloudWatch alarm that sends an Amazon SNS message when the alarm changes state. An alarm watches a single metric over a time period you specify, and performs one or more actions based on the value of the metric relative to a given threshold over a number of time periods. The action is a notification sent to an Amazon SNS topic or Auto Scaling policy.

Alarms invoke actions for sustained state changes only. CloudWatch alarms will not invoke actions simply because they are in a particular state, the state must have changed and been maintained for a specified number of periods. The following procedures show how to create alarms for Amazon RDS.

To set alarms using the CloudWatch console

1. Sign in to the AWS Management Console and open the CloudWatch console at <https://console.aws.amazon.com/cloudwatch/>.
2. Choose **Alarms** and then choose **Create Alarm**. This launches the **Create Alarm Wizard**.
3. Choose **RDS Metrics** and scroll through the Amazon RDS metrics to locate the metric you want to place an alarm on. To display just the Amazon RDS metrics in this dialog box, search for the identifier of your resource. Choose the metric to create an alarm on and then choose **Next**.
4. Enter **Name**, **Description**, and **Whenever** values for the metric.
5. If you want CloudWatch to send you an email when the alarm state is reached, for **Whenever this alarm:**, choose **State is ALARM**. For **Send notification to:**, choose an existing SNS topic. If you choose **Create topic**, you can set the name and email addresses for a new email subscription list. This list is saved and appears in the field for future alarms.

Note

If you use **Create topic** to create a new Amazon SNS topic, the email addresses must be verified before they receive notifications. Emails are only sent when the alarm enters an alarm state. If this alarm state change happens before the email addresses are verified, they don't receive a notification.

6. At this point, the **Alarm Preview** area gives you a chance to preview the alarm you're about to create. Choose **Create Alarm**.

To set an alarm using the AWS CLI

- Call `put-metric-alarm`. For more information, see [AWS CLI Command Reference](#).

To set an alarm using the CloudWatch API

- Call `PutMetricAlarm`. For more information, see [Amazon CloudWatch API Reference](#)

Publishing Database Engine Logs to Amazon CloudWatch Logs

You can configure your Amazon RDS database engine to publish log data to a log group in Amazon CloudWatch Logs. With CloudWatch Logs, you can perform real-time analysis of the log data, and use CloudWatch to create alarms and view metrics. You can use CloudWatch Logs to store your log records in highly durable storage, which you can manage with the CloudWatch Logs Agent. For example, you can determine when to rotate log records from a host to the log service, so you can access the raw logs when you need to.

For engine-specific information, see the following topics:

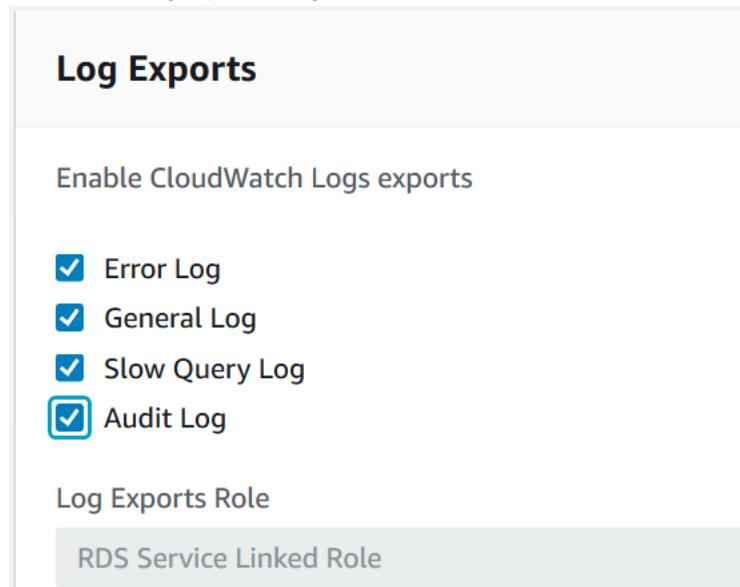
- the section called “[Publishing MariaDB Logs to Amazon CloudWatch Logs](#)” (p. 395)
- the section called “[Publishing MySQL Logs to CloudWatch Logs](#)” (p. 406)
- the section called “[Publishing Oracle Logs to Amazon CloudWatch Logs](#)” (p. 415)
- the section called “[Publishing PostgreSQL Logs to CloudWatch Logs](#)” (p. 420)
- the section called “[Publishing SQL Server Logs to Amazon CloudWatch Logs](#)” (p. 402)

Note

Before you enable log data publishing, you must have a service-linked role in AWS Identity and Access Management (IAM). For more information about service-linked roles, see [Using Service-Linked Roles for Amazon RDS](#) (p. 501).

Configuring CloudWatch Log Integration

To publish your database log files to CloudWatch Logs, choose which logs to publish. Make this choice in the **Advanced Settings** section when you create a new DB instance. You can also modify an existing DB instance to begin publishing.



After you have enabled publishing, Amazon RDS continuously streams all of the DB instance log records to a log group. For example, you have a log group `/aws/rds/instance/log` type for each type of log that you publish. This log group is in the same AWS Region as the database instance that generates the log.

After you have published log records, you can use CloudWatch Logs to search and filter the records. For more information about searching and filtering logs, see [Searching and Filtering Log Data](#).

Viewing DB Instance Metrics

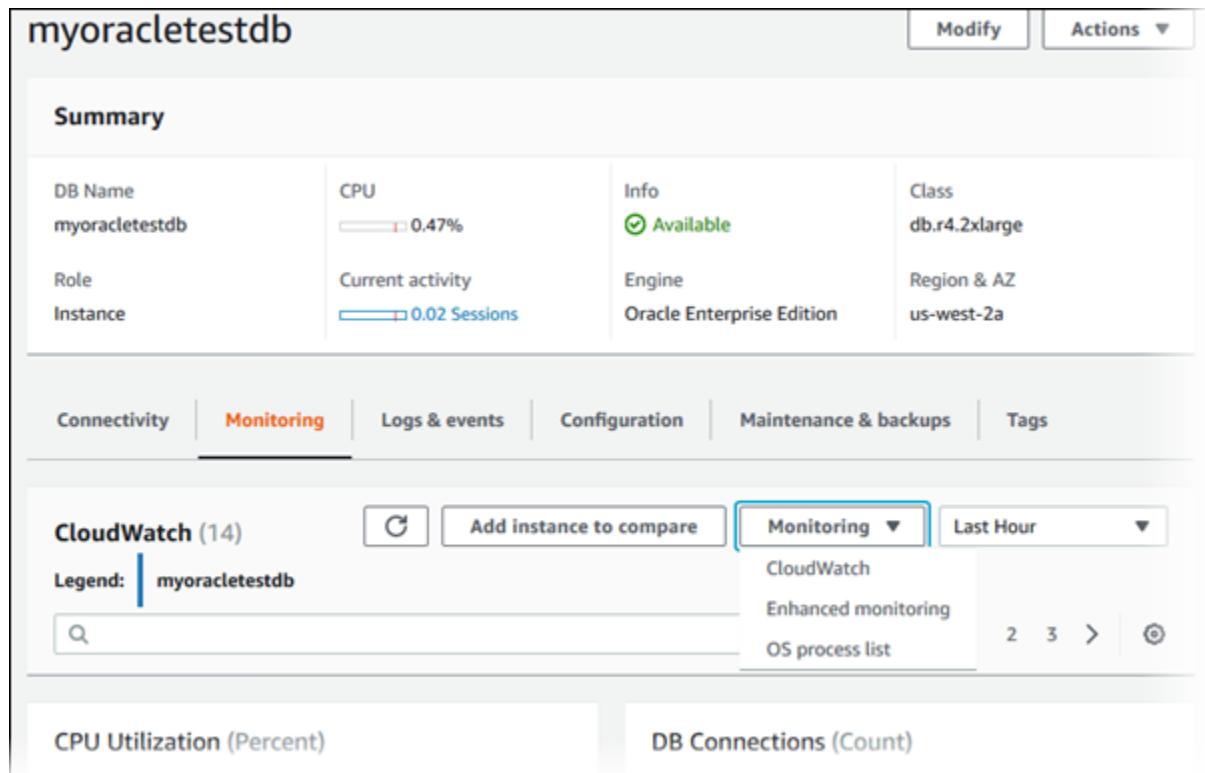
Amazon RDS provides metrics so that you can monitor the health of your DB instances. You can monitor both DB instance metrics and operating system (OS) metrics.

This section provides details on how you can view metrics for your DB instance using the RDS console and CloudWatch. For information on monitoring metrics for the operating system of your DB instance in real time using CloudWatch Logs, see [Enhanced Monitoring \(p. 307\)](#).

Viewing Metrics by Using the Console

To view DB and OS metrics for a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the name of the DB instance you need information about to show its details.
4. Choose the **Monitoring** tab.
5. For **Monitoring**, choose the option for how you want to view your metrics from these:
 - **CloudWatch** – Shows a summary of DB instance metrics available from Amazon CloudWatch. Each metric includes a graph showing the metric monitored over a specific time span.
 - **Enhanced monitoring** – Shows a summary of OS metrics available for a DB instance with Enhanced Monitoring enabled. Each metric includes a graph showing the metric monitored over a specific time span.
 - **OS Process list** – Shows details for each process running in the selected instance.



Tip

You can choose the time range of the metrics represented by the graphs with the time range list.

You can choose any graph to bring up a more detailed view. You can also apply metric-specific filters to the data.

Viewing DB Instance Metrics with the CLI or API

Amazon RDS integrates with CloudWatch metrics to provide a variety of DB instance metrics. You can view CloudWatch metrics using the RDS console, AWS CLI, or API.

For a complete list of Amazon RDS metrics, go to [Amazon RDS Dimensions and Metrics](#) in the *Amazon CloudWatch User Guide*.

Viewing DB Metrics by Using the CloudWatch CLI

Note

The following CLI example requires the CloudWatch command line tools. For more information on CloudWatch and to download the developer tools, see the [Amazon CloudWatch product page](#). The `StartTime` and `EndTime` values supplied in this example are for illustrative purposes. You must substitute appropriate start and end time values for your DB instance.

To view usage and performance statistics for a DB instance

- Use the CloudWatch command `mon-get-stats` with the following parameters.

```
PROMPT>mon-get-stats FreeStorageSpace --dimensions="DBInstanceIdentifier=mydbinstance"
--statistics= Average
```

```
--namespace="AWS/RDS" --start-time 2009-10-16T00:00:00 --end-time 2009-10-16T00:02:00
```

[Viewing DB Metrics by Using the CloudWatch API](#)

The `StartTime` and `EndTime` values supplied in this example are for illustrative purposes. You must substitute appropriate start and end time values for your DB instance.

To view usage and performance statistics for a DB instance

- Call the CloudWatch API `GetMetricStatistics` with the following parameters:
 - `Statistics.member.1 = Average`
 - `Namespace = AWS/RDS`
 - `StartTime = 2009-10-16T00:00:00`
 - `EndTime = 2009-10-16T00:02:00`
 - `Period = 60`
 - `MeasureName = FreeStorageSpace`

Enhanced Monitoring

Amazon RDS provides metrics in real time for the operating system (OS) that your DB instance runs on. You can view the metrics for your DB instance using the console, or consume the Enhanced Monitoring JSON output from Amazon CloudWatch Logs in a monitoring system of your choice.

By default, Enhanced Monitoring metrics are stored in the CloudWatch Logs for 30 days. To modify the amount of time the metrics are stored in the CloudWatch Logs, change the retention for the `RDSOSMetrics` log group in the CloudWatch console. For more information, see [Change Log Data Retention in CloudWatch Logs](#) in the *Amazon CloudWatch Logs User Guide*.

The cost for using Enhanced Monitoring depends on several factors:

- You are only charged for Enhanced Monitoring that exceeds the free tier provided by Amazon CloudWatch Logs.

For more information about pricing, see [Amazon CloudWatch Pricing](#).
- A smaller monitoring interval results in more frequent reporting of OS metrics and increases your monitoring cost.
- Usage costs for Enhanced Monitoring are applied for each DB instance that Enhanced Monitoring is enabled for. Monitoring a large number of DB instances is more expensive than monitoring only a few.
- DB instances that support a more compute-intensive workload have more OS process activity to report and higher costs for Enhanced Monitoring.

Enhanced Monitoring Availability

Enhanced Monitoring is available for the following database engines:

- MariaDB
- Microsoft SQL Server
- MySQL version 5.5 or later
- Oracle
- PostgreSQL

Enhanced Monitoring is available for all DB instance classes except for db.m1.small.

Differences Between CloudWatch and Enhanced Monitoring Metrics

CloudWatch gathers metrics about CPU utilization from the hypervisor for a DB instance, and Enhanced Monitoring gathers its metrics from an agent on the instance. As a result, you might find differences between the measurements, because the hypervisor layer performs a small amount of work. The differences can be greater if your DB instances use smaller instance classes, because then there are likely more virtual machines (VMs) that are managed by the hypervisor layer on a single physical instance. Enhanced Monitoring metrics are useful when you want to see how different processes or threads on a DB instance use the CPU.

Setting Up for and Enabling Enhanced Monitoring

To set up for and enable Enhanced Monitoring, take the steps listed following.

Before You Begin

Enhanced Monitoring requires permission to act on your behalf to send OS metric information to CloudWatch Logs. You grant Enhanced Monitoring the required permissions using an AWS Identity and Access Management (IAM) role.

The first time that you enable Enhanced Monitoring in the console, you can select the **Default** option for the **Monitoring Role** property to have RDS create the required IAM role. RDS then automatically creates a role named `rds-monitoring-role` for you, and uses it for the specified DB instance or Read Replica.

You can also create the required role before you enable Enhanced Monitoring, and then specify your new role's name when you enable Enhanced Monitoring. You must create this required role if you enable Enhanced Monitoring using the AWS CLI or the RDS API.

To create the appropriate IAM role to permit Amazon RDS to communicate with the Amazon CloudWatch Logs service on your behalf, take the following steps.

The user that enables Enhanced Monitoring must be granted the `PassRole` permission. For more information, see Example 2 in [Granting a User Permissions to Pass a Role to an AWS Service](#) in the *IAM User Guide*.

To create an IAM role for Amazon RDS Enhanced Monitoring

1. Open the [IAM Console](#) at <https://console.aws.amazon.com>.
2. In the navigation pane, choose **Roles**.
3. Choose **Create role**.
4. Choose the **AWS service** tab, and then choose **RDS** from the list of services.
5. Choose **RDS - Enhanced Monitoring**, and then choose **Next: Permissions**.
6. Ensure that the **Attached permissions policy** page shows `AmazonRDSEnhancedMonitoringRole`, and then choose **Next: Tags**.
7. On the **Add tags** page, choose **Next: Review**.
8. For **Role Name**, enter a name for your role, for example `emaccess`, and then choose **Create role**.

Enabling and Disabling Enhanced Monitoring

You can enable Enhanced Monitoring when you create a DB instance or Read Replica, or when you modify a DB instance. If you modify a DB instance to enable Enhanced Monitoring, you don't need to reboot your DB instance for the change to take effect.

You can enable Enhanced Monitoring in the RDS console when you do one of the following actions:

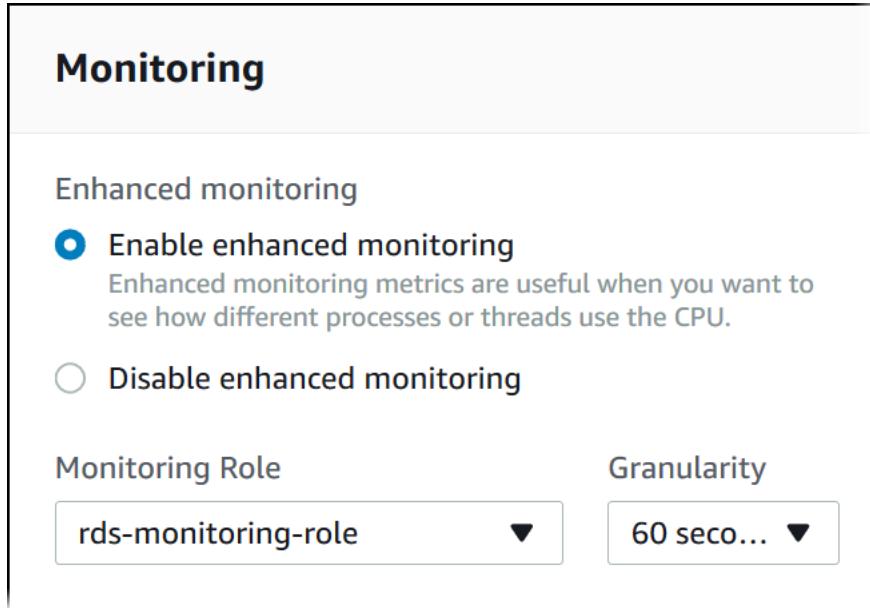
- **Create a DB Instance** – You can enable Enhanced Monitoring in the [Configure Advanced Settings](#) page.
- **Create a Read Replica** – You can enable Enhanced Monitoring in the [Configure Advanced Settings](#) page.
- **Modify a DB Instance** – You can enable Enhanced Monitoring in the [Modify DB Instance](#) page.

To enable Enhanced Monitoring by using the RDS console, scroll to the **Monitoring** section and do the following:

1. Choose **Enable enhanced monitoring** for your DB instance or Read Replica.
2. Set the **Monitoring Role** property to the IAM role that you created to permit Amazon RDS to communicate with Amazon CloudWatch Logs for you, or choose **Default** to have RDS create a role for you named `rds-monitoring-role`.

3. Set the **Granularity** property to the interval, in seconds, between points when metrics are collected for your DB instance or Read Replica. The **Granularity** property can be set to one of the following values: 1, 5, 10, 15, 30, or 60.

To disable Enhanced Monitoring, choose **Disable enhanced monitoring**.



Enabling Enhanced Monitoring doesn't require your DB instance to restart.

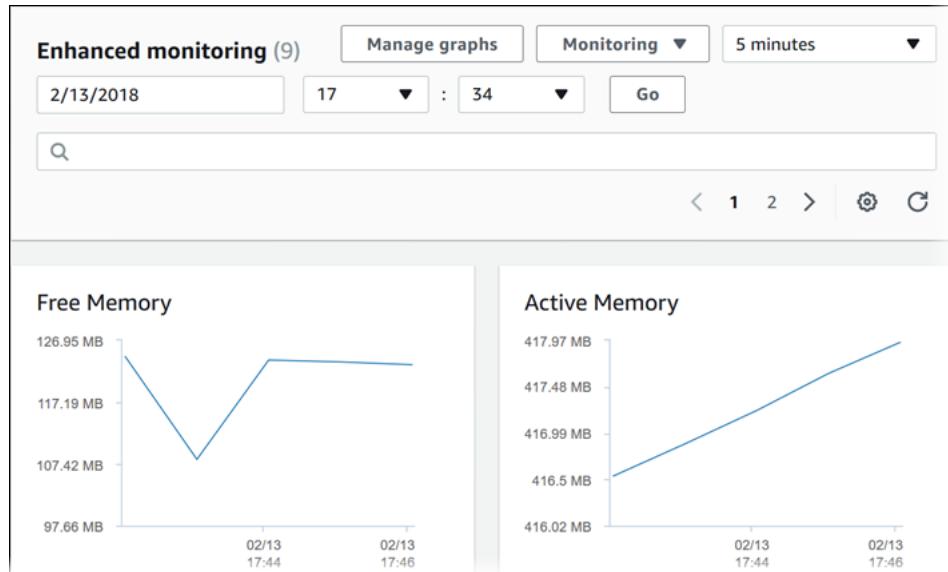
Note

The fastest that the RDS console refreshes is every 5 seconds. If you set the granularity to 1 second in the RDS console, you still see updated metrics only every 5 seconds. You can retrieve 1-second metric updates by using CloudWatch Logs.

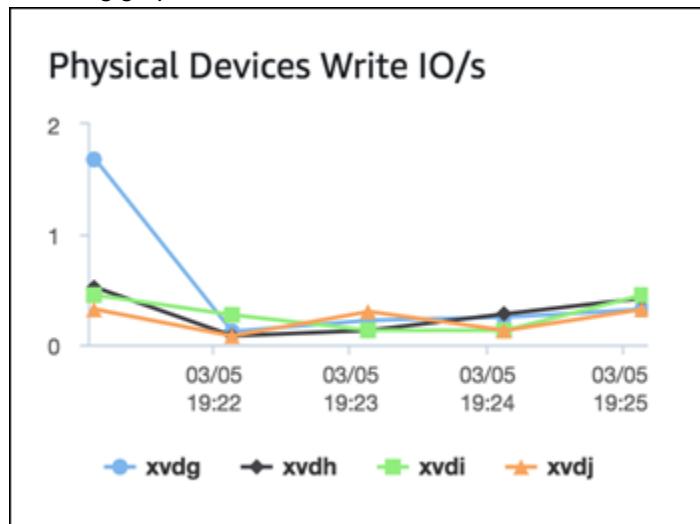
Viewing Enhanced Monitoring

You can view OS metrics reported by Enhanced Monitoring in the RDS console by choosing **Enhanced monitoring** for **Monitoring**.

The Enhanced Monitoring page is shown following.



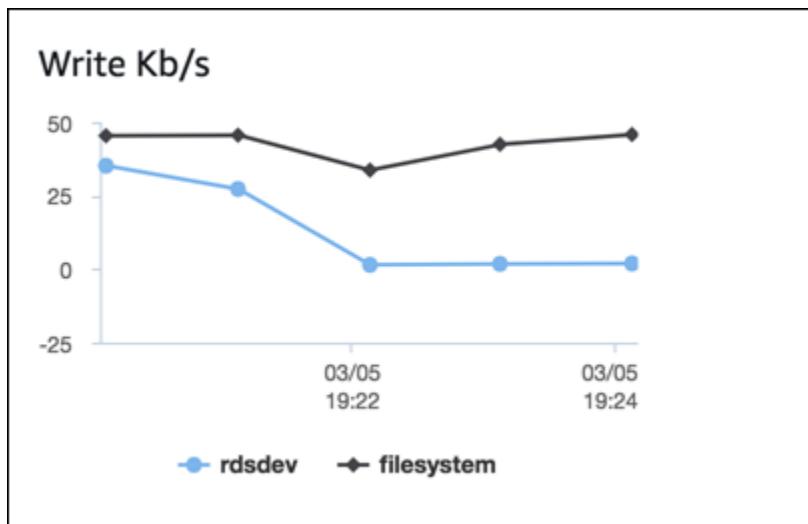
Some DB instances use more than one disk for the DB instance's data storage volume. On those DB instances, the **Physical Devices** graphs show metrics for each one of the disks. For example, the following graph shows metrics for four disks.



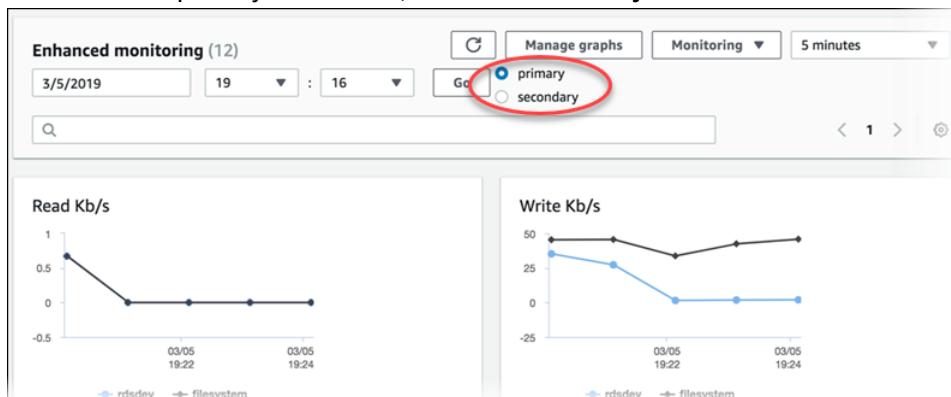
Note

Currently, **Physical Devices** graphs are not available for Microsoft SQL Server DB instances.

When you are viewing aggregated **Disk I/O** and **File system** graphs, the **rdsdev** device relates to the **/rdsdbdata** file system, where all database files and logs are stored. The **filesystem** device relates to the **/** file system (also known as root), where files related to the operating system are stored.



If the DB instance is a Multi-AZ deployment, you can view the OS metrics for the primary DB instance and its Multi-AZ standby replica. In the **Enhanced monitoring** view, choose **primary** to view the OS metrics for the primary DB instance, or choose **secondary** to view the OS metrics for the standby replica.



For more information about Multi-AZ deployments, see [High Availability \(Multi-AZ\) for Amazon RDS \(p. 144\)](#).

Note

Currently, viewing OS metrics for a Multi-AZ standby replica is not supported for MariaDB, Microsoft SQL Server, or PostgreSQL DB instances.

If you want to see details for the processes running on your DB instance, choose **OS process list** for **Monitoring**.

The **Process List** view is shown following.

Process List						
<input type="text"/> Filter process list						
NAME	VIRT	RES	CPU%	MEM%	VMLIMIT	
postgres [3181] ^t	283.55 MB	17.11 MB	0.02	1.72		
postgres:						
rdsadmin	384.7 MB	9.51 MB	0.02	0.95		
rdsadmin						
localhost(40156)						
idle [2953] ^t						

The Enhanced Monitoring metrics shown in the **Process list** view are organized as follows:

- **RDS child processes** – Shows a summary of the RDS processes that support the DB instance, for example mysqld for MySQL DB instances. Process threads appear nested beneath the parent process. Process threads show CPU utilization only as other metrics are the same for all threads for the process. The console displays a maximum of 100 processes and threads. The results are a combination of the top CPU consuming and memory consuming processes and threads. If there are more than 50 processes and more than 50 threads, the console displays the top 50 consumers in each category. This display helps you identify which processes are having the greatest impact on performance.
- **RDS processes** – Shows a summary of the resources used by the RDS management agent, diagnostics monitoring processes, and other AWS processes that are required to support RDS DB instances.
- **OS processes** – Shows a summary of the kernel and system processes, which generally have minimal impact on performance.

The items listed for each process are:

- **VIRT** – Displays the virtual size of the process.
- **RES** – Displays the actual physical memory being used by the process.
- **CPU%** – Displays the percentage of the CPU bandwidth consumed by the process.
- **MEM%** – Displays the percentage of the total memory consumed by the process.

The monitoring data that is shown in the RDS console is retrieved from Amazon CloudWatch Logs. You can also retrieve the metrics for a DB instance as a log stream from CloudWatch Logs. For more information, see [Viewing Enhanced Monitoring by Using CloudWatch Logs \(p. 313\)](#).

Enhanced Monitoring metrics are not returned during the following:

- A failover of the DB instance.
- Changing the instance class of the DB instance (scale compute).

Enhanced Monitoring metrics are returned during a reboot of a DB instance because only the database engine is rebooted. Metrics for the operating system are still reported.

Viewing Enhanced Monitoring by Using CloudWatch Logs

After you have enabled Enhanced Monitoring for your DB instance, you can view the metrics for your DB instance using CloudWatch Logs, with each log stream representing a single DB instance being monitored. The log stream identifier is the resource identifier (`DbiResourceId`) for the DB instance.

To view Enhanced Monitoring log data

1. Open the CloudWatch console at <https://console.aws.amazon.com/cloudwatch/>.
2. If necessary, choose the region that your DB instance is in. For more information, see [Regions and Endpoints](#) in the [Amazon Web Services General Reference](#).
3. Choose **Logs** in the navigation pane.
4. Choose **RDSOSMetrics** from the list of log groups.

In a Multi-AZ deployment, log files with `-secondary` appended to the name are for the Multi-AZ standby replica.

The screenshot shows the AWS CloudWatch Log Groups interface. The top navigation bar includes 'CloudWatch' > 'Log Groups' > 'Streams for RDSOSMetrics'. Below the navigation are three buttons: 'Search Log Group', 'Create Log Stream', and 'Delete Log Stream'. A 'Filter' dropdown is set to 'Log Stream Name Prefix'. The main area displays a table of log streams. The first row has a checkbox and the text 'Log Streams'. The second row contains a checkbox, the prefix 'db-SORJBHOPBSMWGRI5EJW3KMYOTU-', and the suffix '-secondary'. The third row contains a checkbox and the prefix 'db-SORJBHOPBSMWGRI5EJW3KMYOTU'. To the right of the table are two columns: 'Last Event Time' and the corresponding event times for each row. A red box highlights the '-secondary' entry in the list.

Log Streams	Last Event Time
db-SORJBHOPBSMWGRI5EJW3KMYOTU-secondary	2019-03-05 12:12 UTC-8
db-SORJBHOPBSMWGRI5EJW3KMYOTU	2019-03-05 12:07 UTC-8

5. Choose the log stream that you want to view from the list of log streams.

Available OS Metrics

The following tables list the OS metrics available using Amazon CloudWatch Logs.

Metrics for MariaDB, MySQL, Oracle, and PostgreSQL DB instances

Group	Metrics	Description
General	engine	The database engine for the DB instance.
	instanceID	The DB instance identifier.
	instanceResourceIdentifier	A region-unique, immutable identifier for the DB instance, also used as the log stream identifier.
	numVCpus	The number of virtual CPUs for the DB instance.
	timestamp	The time at which the metrics were taken.
	uptime	The amount of time that the DB instance has been active.
	version	The version of the OS metrics' stream JSON format.

Group	Metrics	Description
cpuUtilization	guest	The percentage of CPU in use by guest programs.
	idle	The percentage of CPU that is idle.
	irq	The percentage of CPU in use by software interrupts.
	nice	The percentage of CPU in use by programs running at lowest priority.
	steal	The percentage of CPU in use by other virtual machines.
	system	The percentage of CPU in use by the kernel.
	total	The total percentage of the CPU in use. This value includes the nice value.
	user	The percentage of CPU in use by user programs.
	wait	The percentage of CPU unused while waiting for I/O access.
diskIO	avgQueueLen	The number of requests waiting in the I/O device's queue.
	avgReqSz	The average request size, in kilobytes.
	await	The number of milliseconds required to respond to requests, including queue time and service time.
	device	The identifier of the disk device in use.
	readIOsPS	The number of read operations per second.
	readKb	The total number of kilobytes read.
	readKbPS	The number of kilobytes read per second.
	rrqmPS	The number of merged read requests queued per second.
	tps	The number of I/O transactions per second.
	util	The percentage of CPU time during which requests were issued.
	writeIOsPS	The number of write operations per second.
	writeKb	The total number of kilobytes written.
	writeKbPS	The number of kilobytes written per second.
	wrqmPS	The number of merged write requests queued per second.
fileSys	maxFiles	The maximum number of files that can be created for the file system.
	mountPoint	The path to the file system.
	name	The name of the file system.
	total	The total number of disk space available for the file system, in kilobytes.

Group	Metrics	Description
	used	The amount of disk space used by files in the file system, in kilobytes.
	usedFilePercent	The percentage of available files in use.
	usedFiles	The number of files in the file system.
	usedPercent	The percentage of the file-system disk space in use.
loadAverageMinutes	fifteen	The number of processes requesting CPU time over the last 15 minutes.
	five	The number of processes requesting CPU time over the last 5 minutes.
	one	The number of processes requesting CPU time over the last minute.
memory	active	The amount of assigned memory, in kilobytes.
	buffers	The amount of memory used for buffering I/O requests prior to writing to the storage device, in kilobytes.
	cached	The amount of memory used for caching file system-based I/O.
	dirty	The amount of memory pages in RAM that have been modified but not written to their related data block in storage, in kilobytes.
	free	The amount of unassigned memory, in kilobytes.
	hugePagesFree	The number of free huge pages. Huge pages are a feature of the Linux kernel.
	hugePagesRsvd	The number of committed huge pages.
	hugePagesSize	The size for each huge pages unit, in kilobytes.
	hugePagesSurp	The number of available surplus huge pages over the total.
	hugePagesTotal	The total number of huge pages for the system.
	inactive	The amount of least-frequently used memory pages, in kilobytes.
	mapped	The total amount of file-system contents that is memory mapped inside a process address space, in kilobytes.
	pageTables	The amount of memory used by page tables, in kilobytes.
	slab	The amount of reusable kernel data structures, in kilobytes.
	total	The total amount of memory, in kilobytes.
	writeback	The amount of dirty pages in RAM that are still being written to the backing storage, in kilobytes.
network	interface	The identifier for the network interface being used for the DB instance.

Group	Metrics	Description
	rx	The number of bytes received per second.
	tx	The number of bytes uploaded per second.
processList	cpuUsedPc	The percentage of CPU used by the process.
	id	The identifier of the process.
	memoryUsedPc	The amount of memory used by the process, in kilobytes.
	name	The name of the process.
	parentID	The process identifier for the parent process of the process.
	rss	The amount of RAM allocated to the process, in kilobytes.
	tgid	The thread group identifier, which is a number representing the process ID to which a thread belongs. This identifier is used to group threads from the same process.
swap	VIRT	The amount of virtual memory allocated to the process, in kilobytes.
	swap	The amount of swap memory available, in kilobytes.
	swap_in	The amount of memory, in kilobytes, swapped in from disk.
	swap_out	The amount of memory, in kilobytes, swapped out to disk.
	free	The amount of swap memory free, in kilobytes.
tasks	committed	The amount of swap memory, in kilobytes, used as cache memory.
	blocked	The number of tasks that are blocked.
	running	The number of tasks that are running.
	sleeping	The number of tasks that are sleeping.
	stopped	The number of tasks that are stopped.
	total	The total number of tasks.
	zombie	The number of child tasks that are inactive with an active parent task.

Metrics for Microsoft SQL Server DB instances

Group	Metrics	Description
General	engine	The database engine for the DB instance.
	instanceID	The DB instance identifier.
	instanceResourceIdentifier	A region-unique, immutable identifier for the DB instance, also used as the log stream identifier.

Group	Metrics	Description
	numVCpus	The number of virtual CPUs for the DB instance.
	timestamp	The time at which the metrics were taken.
	uptime	The amount of time that the DB instance has been active.
	version	The version of the OS metrics' stream JSON format.
cpuUtilization	idle	The percentage of CPU that is idle.
	kern	The percentage of CPU in use by the kernel.
	user	The percentage of CPU in use by user programs.
disks	name	The identifier for the disk.
	totalKb	The total space of the disk, in kilobytes.
	usedKb	The amount of space used on the disk, in kilobytes.
	usedPc	The percentage of space used on the disk.
	availKb	The space available on the disk, in kilobytes.
	availPc	The percentage of space available on the disk.
	rdCountPS	The number of read operations per second
	rdBytesPS	The number of bytes read per second.
	wrCountPS	The number of write operations per second.
	wrBytesPS	The amount of bytes written per second.
memory	commitTotKb	The amount of pagefile-backed virtual address space in use, that is, the current commit charge. This value is composed of main memory (RAM) and disk (pagefiles).
	commitLimitKb	The maximum possible value for the commitTotKb metric. This value is the sum of the current pagefile size plus the physical memory available for pageable contents—excluding RAM that is assigned to non-pageable areas.
	commitPeakKb	The largest value of the commitTotKb metric since the operating system was last started.
	kernTotKb	The sum of the memory in the paged and non-paged kernel pools, in kilobytes.
	kernPagedKb	The amount of memory in the paged kernel pool, in kilobytes.
	kernNonpagedKb	The amount of memory in the non-paged kernel pool, in kilobytes.
	pageSize	The size of a page, in bytes.
	physTotKb	The amount of physical memory, in kilobytes.
	physAvailKb	The amount of available physical memory, in kilobytes.

Group	Metrics	Description
	sqlServerTotKb	The amount of memory committed to Microsoft SQL Server, in kilobytes.
	sysCacheKb	The amount of system cache memory, in kilobytes.
network	interface	The identifier for the network interface being used for the DB instance.
	rdBytesPS	The number of bytes received per second.
	wrBytesPS	The number of bytes sent per second.
processList	cpuUsedPc	The percentage of CPU used by the process.
	memUsedPc	The percentage of total memory used by the process.
	name	The name of the process.
	pid	The identifier of the process. This value is not present for processes that are owned by Amazon RDS.
	ppid	The process identifier for the parent of this process. This value is only present for child processes.
	tid	The thread identifier. This value is only present for threads. The owning process can be identified by using the pid value.
	workingSetKb	The amount of memory in the private working set plus the amount of memory that is in use by the process and can be shared with other processes, in kilobytes.
	workingSetPrivKb	The amount of memory that is in use by a process, but can't be shared with other processes, in kilobytes.
	workingSetShareableKb	The amount of memory that is in use by a process and can be shared with other processes, in kilobytes.
	virtKb	The amount of virtual address space the process is using, in kilobytes. Use of virtual address space does not necessarily imply corresponding use of either disk or main memory pages.
system	handles	The number of handles that the system is using.
	processes	The number of processes running on the system.
	threads	The number of threads running on the system.

Using Amazon RDS Performance Insights

Amazon RDS Performance Insights monitors your Amazon RDS DB instance load so that you can analyze and troubleshoot your database performance. Amazon RDS Performance Insights is currently available for use with the following DB engines:

- Amazon Aurora with MySQL compatibility version 2.04.2 and higher 2.x versions (compatible with MySQL 5.7)
- Amazon Aurora with MySQL compatibility version 1.17.3 and higher 1.x versions (compatible with MySQL 5.6)
- Amazon Aurora with PostgreSQL compatibility
- Amazon RDS for MariaDB version 10.2.21 and higher 10.2 versions
- Amazon RDS for MySQL version 5.7.22 and higher 5.7 versions, and version 5.6.41 and higher 5.6 versions
- Amazon RDS for Microsoft SQL Server (all versions except SQL Server 2008)
- Amazon RDS for PostgreSQL version 10 and 11
- Amazon RDS for Oracle (all versions)

Note

Amazon RDS Performance Insights is not supported for MariaDB version 10.0, 10.1, or 10.3, or for MySQL version 5.5 or 8.0.

For Amazon RDS for MariaDB and MySQL, Performance Insights is not supported on the following DB instance classes: db.t2.micro, db.t2.small, db.t3.micro, and db.t3.small.

On Aurora MySQL, Performance Insights is not supported on db.t2 or db.t3 DB instance classes. Performance Insights is not supported for Aurora MySQL DB clusters enabled for parallel query.

Important

This guide describes using Amazon RDS Performance Insights with non-Aurora DB engines. For information about using Amazon RDS Performance Insights with Amazon Aurora, see the [Using Amazon RDS Performance Insights](#) in the [Amazon Aurora User Guide](#).

Performance Insights expands on existing Amazon RDS monitoring features to illustrate your database's performance and help you analyze any issues that affect it. With the Performance Insights dashboard, you can visualize the database load and filter the load by waits, SQL statements, hosts, or users.

Performance Insights is on by default in the console create wizard for all DB engines that work with Amazon RDS. If you have more than one database on a DB instance, performance data for all of the databases is aggregated for the DB instance.

The central metric for Performance Insights is `DB_Load`, which represents the average number of active sessions for the DB engine. The `DB_Load` metric is collected every second. An *active session* is a connection that has submitted work to the DB engine and is waiting for a response from it. For example, if you submit a SQL query to the DB engine, the database session is active while the DB engine is processing that query.

By combining `DB_Load` with wait event data, you can get a complete picture of the state for an active session. A *wait event* is a condition that causes the execution of a SQL statement to wait for a specific event to happen before it can continue. For example, SQL statement execution might wait until a locked resource is unlocked. Wait events vary by DB engine:

- For information about all MariaDB and MySQL wait events, see [Wait Event Summary Tables](#) in the MySQL documentation.
- For information about all PostgreSQL wait events, see [PostgreSQL Wait Events](#) in the PostgreSQL documentation.

- For information about all Oracle wait events, see [Descriptions of Wait Events](#) in the Oracle documentation.
- For information about all SQL Server wait events, see [Types of Waits](#) in the SQL Server documentation.

Note

For Oracle, background processes sometimes do work without an associated SQL statement. In these cases, Performance Insights reports the type of background process concatenated with a colon and the wait class associated with that background process. Types of background process include LGWR, ARC0, PMON, and so on.

For example, when the archiver is performing I/O, the Performance Insights report for it is similar to ARC1 : System I/O. Occasionally, the background process type is also missing, and Performance Insights only reports the wait class, for example :System I/O.

Session information is collected, aggregated, and displayed in the dashboard as the **Average Active Sessions** chart. The **Average Active Sessions** chart displays the **Max CPU** value as a line, so you can see if active sessions are exceeding it or not. The **Max CPU** value is determined by the number of **vCPU** (virtual CPU) cores for your DB instance.

If the load in the **Average Active Sessions** chart is often above the **Max CPU** line and the primary wait state is CPU, the system CPU is overloaded. In these cases, you might want to throttle connections to the instance, tune any SQL queries with a high CPU load, or consider a larger instance class. High and consistent instances of any wait state indicate that there might be bottlenecks or resource contention issues to resolve. This can be true even if the load doesn't cross the **Max CPU** line.

You can find an overview of Performance Insights in the following video.

[Using Performance Insights to Analyze Performance of Amazon Aurora PostgreSQL](#)

Topics

- [Enabling Performance Insights \(p. 320\)](#)
- [Access Control for Performance Insights \(p. 324\)](#)
- [Using the Performance Insights Dashboard \(p. 326\)](#)
- [Additional User Interface Features \(p. 337\)](#)
- [Performance Insights API \(p. 338\)](#)
- [Performance Insights Metrics Published to Amazon CloudWatch \(p. 351\)](#)
- [Performance Insights Counters \(p. 352\)](#)
- [Logging Performance Insights Calls by Using AWS CloudTrail \(p. 362\)](#)

Enabling Performance Insights

To use Performance Insights, you must enable it on your DB instance.

Console

You can use the console to enable Performance Insights when you create a new DB instance. You can also modify a DB instance to enable Performance Insights.

Topics

- [Enabling Performance Insights with the Console When Creating a DB Instance \(p. 321\)](#)
- [Enabling Performance Insights with the Console When Modifying a DB Instance \(p. 321\)](#)

Enabling Performance Insights with the Console When Creating a DB Instance

When you create a new DB instance, Performance Insights is enabled when you choose **Enable Performance Insights** in the **Performance Insights** section.

To create a DB instance, follow the instructions for your DB engine in [Creating an Amazon RDS DB Instance \(p. 148\)](#).

The following screenshot shows the **Performance Insights** section.

The screenshot shows the 'Performance Insights' configuration section. At the top, there is a radio button group for 'Enable Performance Insights' (selected) and 'Disable Performance Insights'. Below this is a 'Retention' dropdown menu with three options: '(Default) 7 days', '(Default) 7 days', and 'Long Term Retention (2 years)', with 'Long Term Retention (2 years)' highlighted. At the bottom, there is a table with two columns: 'Description' and 'Account'. The 'Description' row contains the text 'Default master key that protects my RDS database volumes when no other key is defined'. The 'Account' row contains the text 'This account([REDACTED])'. To the right of the table is a 'KMS key ID' column, which is currently empty.

You have the following options when you choose **Enable Performance Insights**:

- **Retention** – The amount of time to retain Performance Insights data. Choose either 7 days (the default) or 2 years.
- **Master key** – Specify your AWS Key Management Service (AWS KMS) key. Performance Insights encrypts all potentially sensitive data using your AWS KMS key. Data is encrypted in flight and at rest. For more information, see [Encrypting Amazon RDS Resources \(p. 429\)](#).

Enabling Performance Insights with the Console When Modifying a DB Instance

You can modify a DB instance to enable Performance Insights using the console.

To enable Performance Insights for a DB instance using the console

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. Choose **Databases**.
3. Choose the DB instance that you want to modify, and choose **Modify**.
4. In the **Performance Insights** section, choose **Enable Performance Insights**.

You have the following options when you choose **Enable Performance Insights**:

- **Retention** – The amount of time to retain Performance Insights data. Choose either 7 days (the default) or 2 years.

- **Master key** – Specify your AWS Key Management Service (AWS KMS) key. Performance Insights encrypts all potentially sensitive data using your AWS KMS key. Data is encrypted in flight and at rest. For more information, see [Encrypting Amazon RDS Resources \(p. 429\)](#).
5. Choose **Continue**.
 6. For **Scheduling of Modifications**, choose one of the following:
 - **Apply during the next scheduled maintenance window** – Wait to apply the **Performance Insights** modification until the next maintenance window.
 - **Apply immediately** – Apply the **Performance Insights** modification as soon as possible.
 7. Choose **Modify instance**.

AWS CLI

When you create a new DB instance using the [create-db-instance](#) AWS CLI command, Performance Insights is enabled when you specify `--enable-performance-insights`.

You can also specify the `--enable-performance-insights` value using the following AWS CLI commands:

- [create-db-instance-read-replica](#)
- [modify-db-instance](#)
- [restore-db-instance-from-s3](#)

The following procedure describes how to enable Performance Insights for a DB instance using the AWS CLI.

To enable Performance Insights for a DB instance using the AWS CLI

- Call the [modify-db-instance](#) AWS CLI command and supply the following values:
 - `--db-instance-identifier` – The name of the DB instance.
 - `--enable-performance-insights`

The following example enables Performance Insights for `sample-db-instance`.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
  --db-instance-identifier sample-db-instance \
  --enable-performance-insights
```

For Windows:

```
aws rds modify-db-instance ^
  --db-instance-identifier sample-db-instance ^
  --enable-performance-insights
```

When you enable Performance Insights, you can optionally specify the amount of time, in days, to retain Performance Insights data with the `--performance-insights-retention-period` option. Valid values are 7 (the default) or 731 (2 years).

The following example enables Performance Insights for `sample-db-instance` and specifies that Performance Insights data is retained for two years.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier sample-db-instance \
--enable-performance-insights \
--performance-insights-retention-period 731
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier sample-db-instance ^
--enable-performance-insights ^
--performance-insights-retention-period 731
```

RDS API

When you create a new DB instance using the [CreateDBInstance](#) operation Amazon RDS API operation, Performance Insights is enabled when you set `EnablePerformanceInsights` to True.

You can also specify the `EnablePerformanceInsights` value using the following API operations:

- [ModifyDBInstance](#)
- [CreateDBInstanceReadReplica](#)
- [RestoreDBInstanceFromS3](#)

When you enable Performance Insights, you can optionally specify the amount of time, in days, to retain Performance Insights data with the `PerformanceInsightsRetentionPeriod` parameter. Valid values are 7 (the default) or 731 (2 years).

Enabling the Performance Schema for Performance Insights on Amazon RDS for MariaDB or MySQL

For Amazon RDS for MariaDB or MySQL, Performance Insights provides more detailed information when the Performance Schema feature is enabled. The Performance Schema is enabled automatically when you create an Amazon RDS for MariaDB or MySQL DB instance with Performance Insights enabled. When you create the DB instance with Performance Insights enabled, the following subset of Performance Schema parameters is set to the specified values automatically:

Parameter Name	Parameter Value
<code>performance_schema</code>	1
<code>performance-schema-consumer-events-waits-current</code>	ON
<code>performance-schema-instrument</code>	<code>wait/=%=ON</code>

Parameter Name	Parameter Value
performance-schema-consumer-global-instrumentation	ON
performance-schema-consumer-thread-instrumentation	ON

Performance Schema is enabled automatically only if your parameter group doesn't have an explicitly set value for the `performance_schema` parameter. You can examine the `performance_schema` parameter, and if the value of source is `user`, then you set a value. If you want the Performance Schema parameters to be set automatically, then unset the value for the `performance_schema` parameter. You can view the source of a parameter value by viewing the parameter in the AWS Management Console or by running the AWS CLI [describe-db-parameters](#) command.

When you change the value of the `performance_schema` parameter, a DB instance reboot is required. If you're creating a new DB instance with Performance Insights enabled, the `performance_schema` parameter is set to 1 (enabled) by default.

Without the Performance Schema enabled, Performance Insights displays database load broken down by the list state of the MySQL process. With Performance Schema enabled, Performance Insights displays database load broken down by detailed wait events.

For more information, see [Using the Performance Insights Dashboard \(p. 326\)](#).

Access Control for Performance Insights

To access Performance Insights, you must have the appropriate permissions from AWS Identity and Access Management (IAM). There are two options available for granting access:

1. Attach the `AmazonRDSFullAccess` managed policy to an IAM user or role.
2. Create a custom IAM policy and attach it to an IAM user or role.

AmazonRDSFullAccess Managed Policy

`AmazonRDSFullAccess` is an AWS-managed policy that grants access to all of the Amazon RDS API operations. The policy also grants access to related services that are used by the Amazon RDS console—for example, event notifications using Amazon SNS.

In addition, `AmazonRDSFullAccess` contains all the permissions needed for using Performance Insights. If you attach this policy to an IAM user or role, the recipient can use Performance Insights, along with other console features.

Using a Custom IAM Policy

For users who don't have full access with the `AmazonRDSFullAccess` policy, you can grant access to Performance Insights by creating or modifying a user-managed IAM policy. When you attach the policy to an IAM user or role, the recipient can use Performance Insights.

To create a custom policy

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Policies**.
3. Choose **Create policy**.
4. On the **Create Policy** page, choose the JSON tab.

5. Copy and paste the following.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": "pi:*",
            "Resource": "arn:aws:pi:*::metrics/rds/*"
        }
    ]
}
```

6. Choose **Review policy**

Note

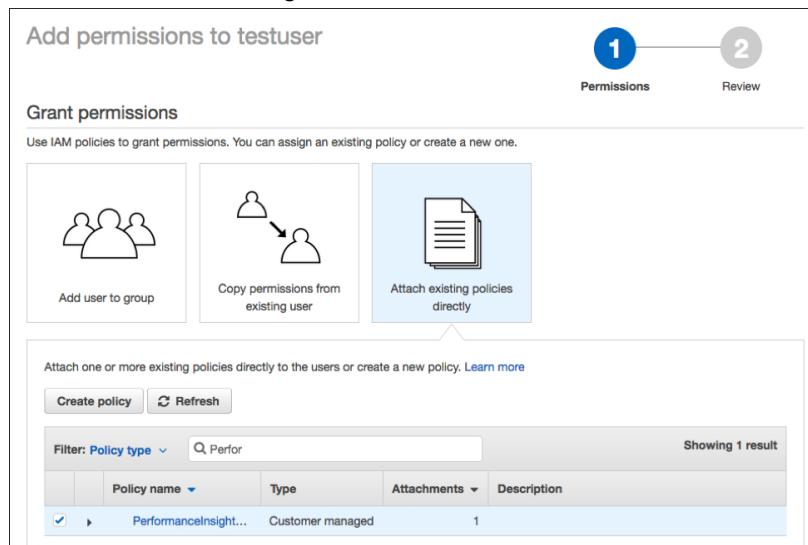
Currently, when you enter this policy, the **Visual editor** tab displays a warning that the **pi** resource is not recognized. You can ignore this warning.

7. Provide a name for the policy and optionally a description, and then choose **Create policy**.

You can now attach the policy to an IAM user or role. The following procedure assumes that you already have an IAM user available for this purpose.

To attach the policy to an IAM user

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.
 2. In the navigation pane, choose **Users**.
 3. Choose an existing user from the list.
- Important**
To use Performance Insights, the user must have access to Amazon RDS in addition to the custom policy. For example, the `AmazonRDSReadOnlyAccess` predefined policy provides read-only access to Amazon RDS. For more information, see [Managing Access Using Policies \(p. 447\)](#).
4. On the **Summary** page, choose **Add permissions**.
 5. Choose **Attach existing policies directly**. For **Search**, type the first few characters of your policy name, as shown following.



6. Choose your policy, and then choose **Next: Review**.

7. Choose **Add permissions**.

Using the Performance Insights Dashboard

The Performance Insights dashboard contains database performance information to help you analyze and troubleshoot performance issues. On the main dashboard page, you can view information about the database load. You can also drill into details for a particular wait state, SQL query, host, or user.

Topics

- [Opening the Performance Insights Dashboard \(p. 326\)](#)
- [Performance Insights Dashboard Components \(p. 327\)](#)
- [Analyzing Database Load Using the Performance Insights Dashboard \(p. 332\)](#)
- [Analyzing Statistics for Running Queries \(p. 333\)](#)
- [Viewing More SQL Text in the Performance Insights Dashboard \(p. 335\)](#)

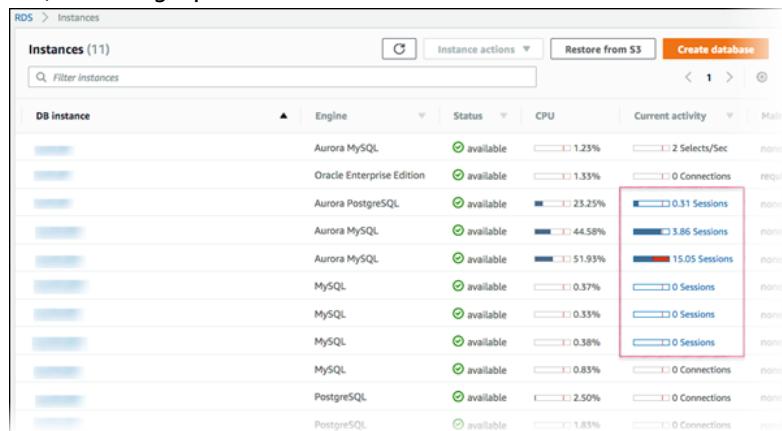
Opening the Performance Insights Dashboard

To see the Performance Insights dashboard, use the following procedure.

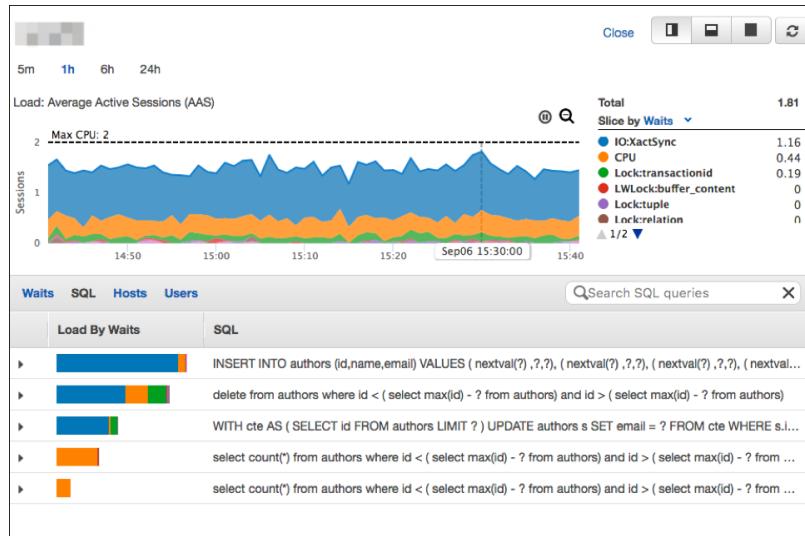
To view the Performance Insights dashboard in the AWS Management Console

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Performance Insights**.
3. Choose a DB instance. The Performance Insights dashboard is displayed for that DB instance.

For DB instances with Performance Insights enabled, you can also reach the dashboard by choosing the **Sessions** item in the list of DB instances. Under **Current activity**, the **Sessions** item shows the database load in average active sessions over the last five minutes. The bar graphically shows the load. When the bar is empty, the DB instance is idle. As the load increases, the bar fills with blue. When the load passes the number of virtual CPUs (vCPUs) on the DB instance class, the bar turns red, indicating a potential bottleneck.



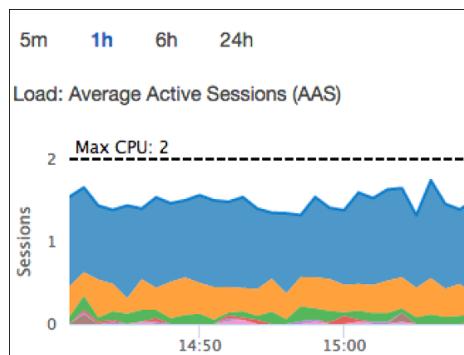
The following screenshot shows the dashboard for a DB instance.



By default, the Performance Insights dashboard shows data for the last 60 minutes. You can modify it to display data for the last 5 minutes, 60 minutes, 5 hours, 24 hours, or 1 week. You can also show all of the data available.

The Performance Insight dashboard automatically refreshes with new data. The refresh rate depends on the amount of data displayed:

- 5 minutes refreshes every 5 seconds.
- 1 hour and 5 hours both refresh every minute.
- 24 hours refreshes every 5 minutes.
- 1 week refreshes every hour.



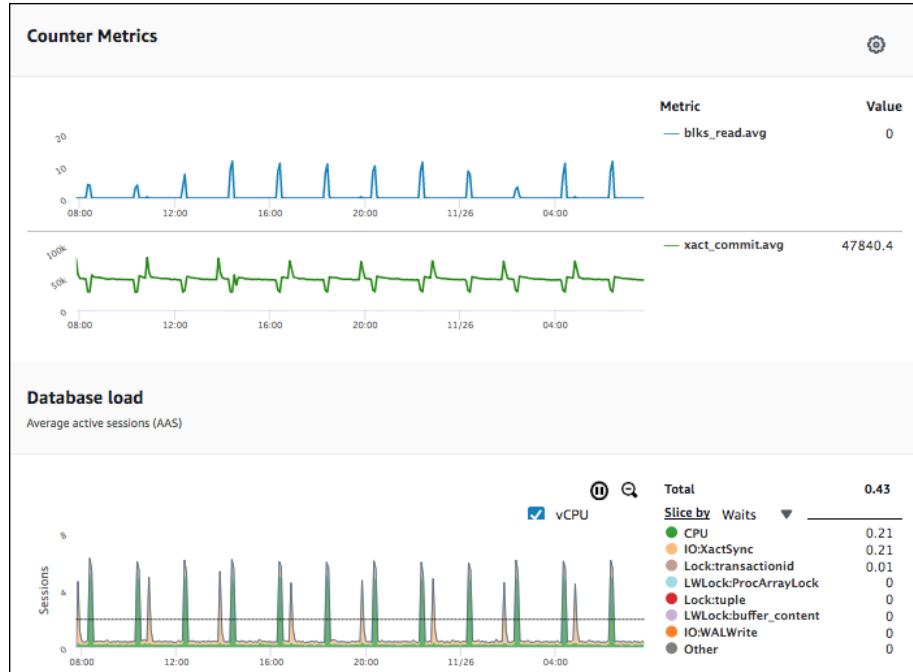
Performance Insights Dashboard Components

The dashboard is divided into three parts:

1. **Counter Metrics chart** – Shows data for specific performance counter metrics.
2. **Average Active Sessions chart** – Shows how the database load compares to DB instance capacity as represented by the **Max CPU** line.
3. **Top load items table** – Shows the top items contributing to database load.

Counter Metrics Chart

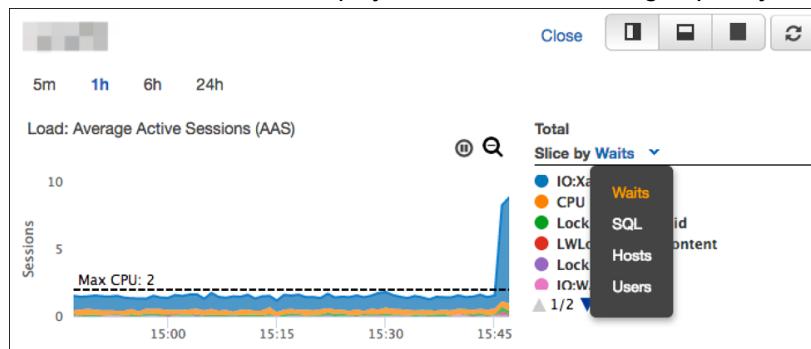
The **Counter Metrics** chart displays data for performance counters. The default metrics shown are `blk_s_read.avg` and `xact_commit.avg`. Choose which performance counters to display by selecting the gear icon in the upper-right corner of the chart.



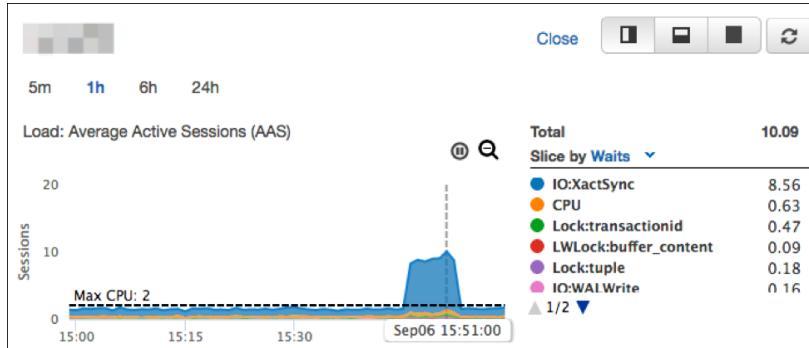
For more information, see [Performance Insights Counters \(p. 352\)](#).

Average Active Sessions Chart

The **Average Active Sessions** chart shows how the database load compares to DB instance capacity as represented by the **Max CPU** line. By default, load is shown as active sessions grouped by wait states. You can also choose instead to display load as active sessions grouped by SQL queries, hosts, or users.



To see details for any item for the selected time period in the legend, hover over that item on the **Average Active Sessions** chart.



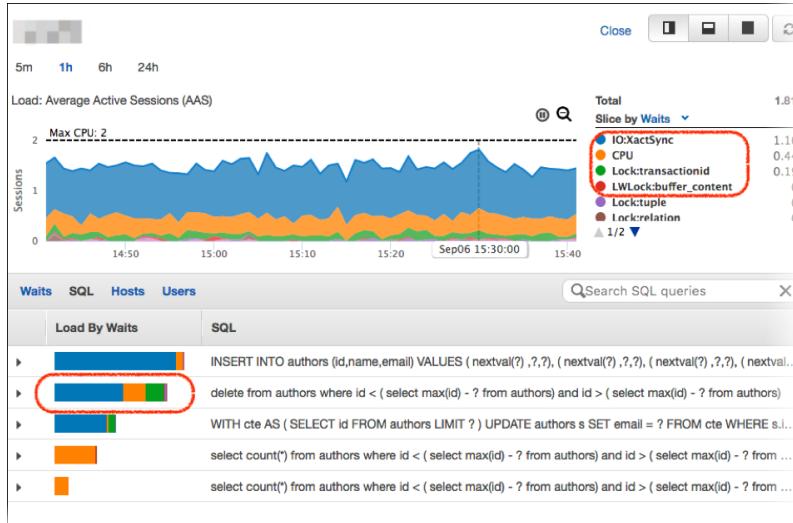
Top Load Items Table

The **Top Load Items** table shows the top items contributing to database load. By default, the top SQL queries that are contributing to the database load are shown. Queries are displayed as digests of multiple actual queries that are structurally similar but that possibly have different parameters. You can choose to display top wait states, hosts, or users instead.

	Waits	SQL	Hosts	Users	Search SQL queries	X
	Load By Waits	SQL				
▶		INSERT INTO authors (id,name,email) VALUES (nextval(?) ,?,?), (nextval(?) ,?,?), (nextval(?) ,?,?), (nextval(?) ,?,?)				
▶		delete from authors where id < (select max(id) - ? from authors) and id > (select max(id) - ? from authors)				
▶		WITH cte AS (SELECT id FROM authors LIMIT ?) UPDATE authors s SET email = ? FROM cte WHERE s.i...				
▶		select count(*) from authors where id < (select max(id) - ? from authors) and id > (select max(id) - ? from ...)				
▶		select count(*) from authors where id < (select max(id) - ? from authors) and id > (select max(id) - ? from ...)				

The percentage of the database load associated with each top load item is illustrated in the **DB Load by Waits** column. This column reflects the load for that item by whatever grouping is currently selected in the **Average Active Sessions** chart. Take the case where the **Average Active Sessions** chart is grouping by hosts and you are looking at SQL queries in the top load items table. In this case, the **DB Load by Waits** bar reflects the load that query represents on the related host. Here it's colored-coded to map to the representation of that host in the **Average Active Sessions** chart.

For another example, suppose that the **Average Active Sessions** chart is grouping by wait states and you are looking at SQL queries in the top load items table. In this case, the **DB Load by Waits** bar is sized, segmented, and color-coded to show how much of a given wait state that query is contributing to. It also shows what wait states are affecting that query.



In the **Top Load Items** table, you can view the following types of identifiers (IDs) that are associated with SQL statements:

- **SQL ID** – An ID that the database uses to uniquely identify a SQL statement.

For Oracle and SQL Server DB instances, you can use a SQL ID to find a specific SQL statement.

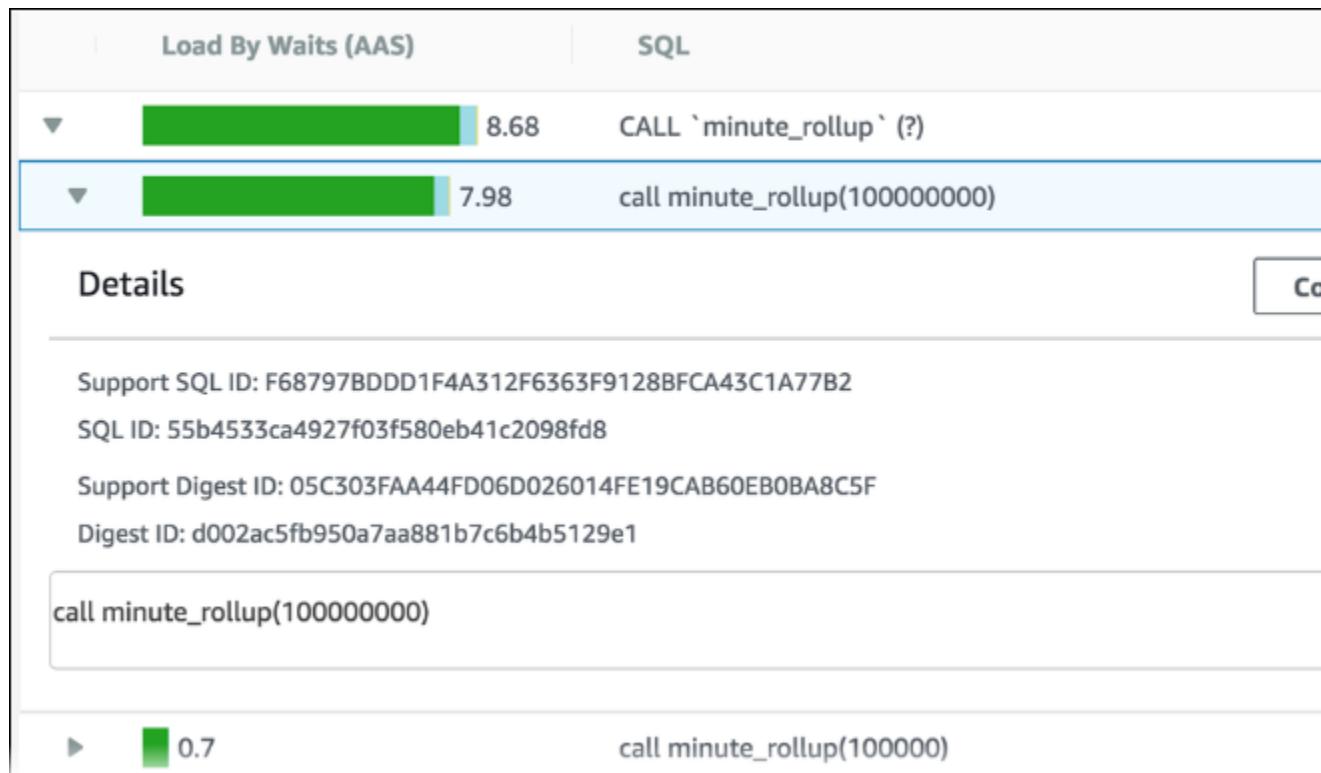
- **Support SQL ID** – A hash value of the SQL ID. This value is only for referencing a SQL ID when you are working with AWS Support. AWS Support doesn't have access to your actual SQL IDs and SQL text.
- **Digest ID** – An ID that the database uses to uniquely identify a SQL digest. A SQL digest can contain one or more SQL statements with literals removed and white space standardized. The literals are replaced with question marks (?).

For Amazon RDS for MariaDB, MySQL, and PostgreSQL DB instances, you can use a digest ID to find a specific SQL digest.

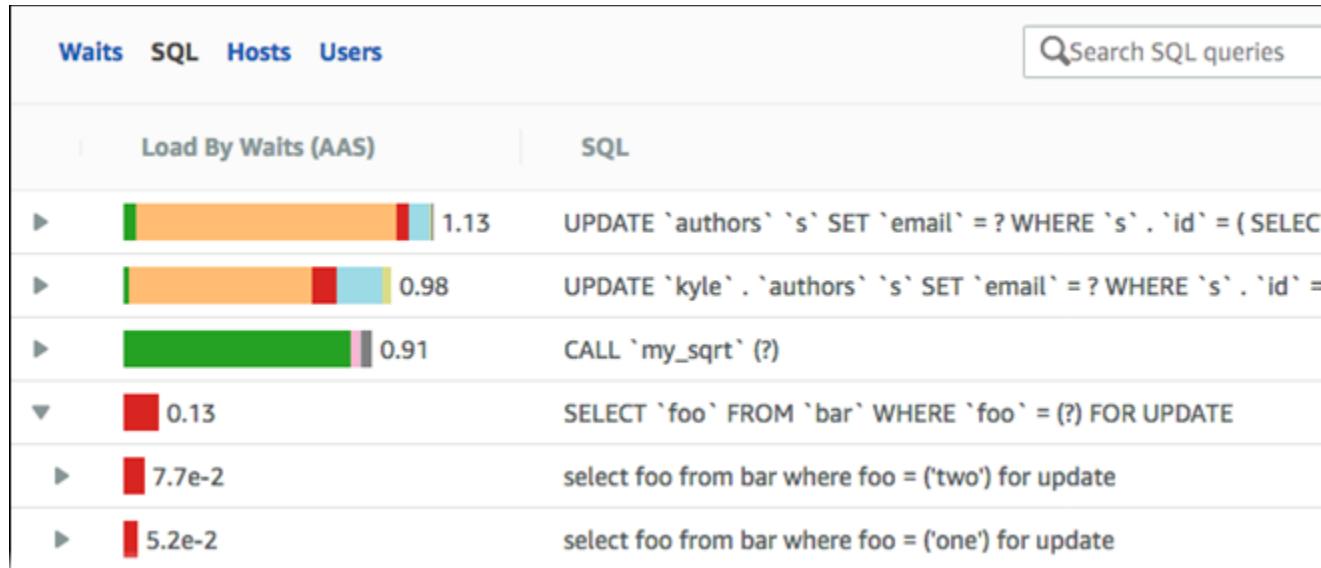
For Oracle and SQL Server DB instances, the digest ID is the same as the SQL ID. The top row in the **Top Load Items** table is the actual SQL statement, including the literals.

- **Support Digest ID** – A hash value of the digest ID. This value is only for referencing a digest ID when you are working with AWS Support. AWS Support doesn't have access to your actual digest IDs and SQL text.

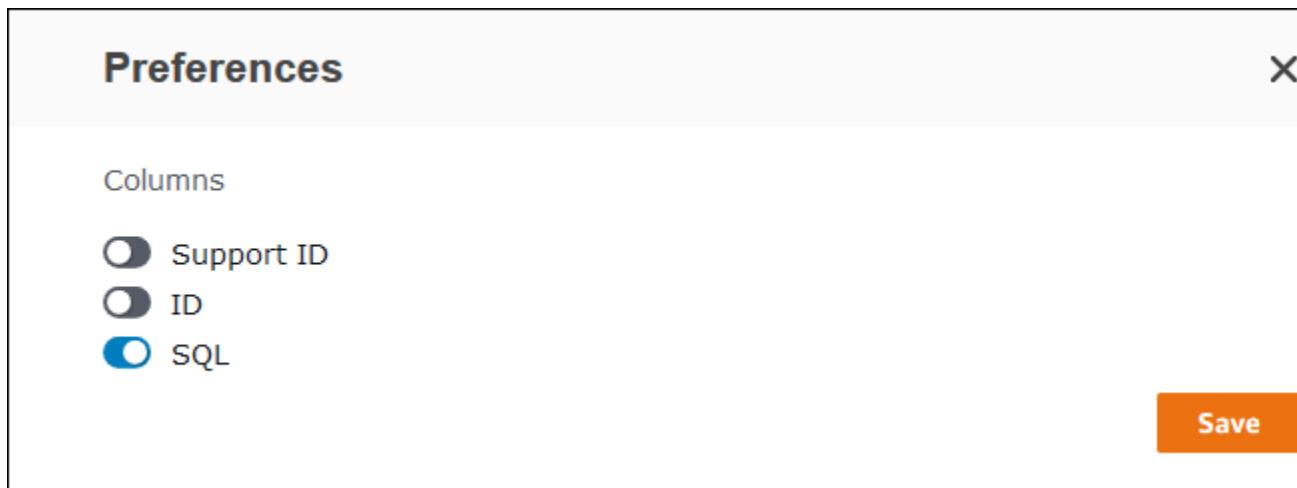
In the **Top Load Items** table, you can open a top statement to view its IDs. The following screenshot shows an open top statement.



You can control the IDs that the **Top Load Items** table shows by choosing the **Preferences** icon.



When you choose the **Preferences** icon, the **Preferences** window opens.



Enable the IDs that you want to have visible in the **Top Load Items** table, and choose **Save**.

Analyzing Database Load Using the Performance Insights Dashboard

If the **Average Active Sessions** chart shows a bottleneck, you can find out where the load is coming from. To do so, look at the top load items table below the **Average Active Sessions** chart. Choose a particular item, like a SQL query or a user, to drill down into that item and see details about it.

DB load grouped by waits and top SQL queries is the default Performance Insights dashboard view. This combination typically provides the most insight into performance issues. DB load grouped by waits shows if there are any resource or concurrency bottlenecks in the database. In this case, the **SQL** tab of the top load items table shows which queries are driving that load.

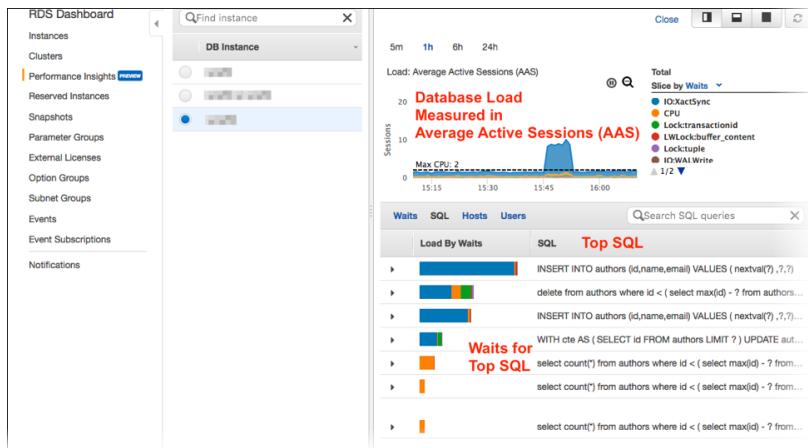
Your typical workflow for diagnosing performance issues is as follows:

1. Review the **Average Active Sessions** chart and see if there are any incidents of database load exceeding the **Max CPU** line.
2. If there is, look at the **Average Active Sessions** chart and identify which wait state or states are primarily responsible.
3. Identify the digest queries causing the load by seeing which of the queries the **SQL** tab on the top load items table are contributing most to those wait states. You can identify these by the **DB Load by Wait** column.
4. Choose one of these digest queries in the **SQL** tab to expand it and see the child queries that it is composed of.

For example, in the dashboard following, **IO:XactSync** waits are a frequent issue. **CPU** wait is less, but it still contributes to load.

The first four roll-up queries in the **SQL** tab of the top load items table correlate strongly to the first state. Thus, those are the ones to drill into and examine the child queries of. You do so to determine how they are contributing to the performance issue.

The last three roll-up queries are the major contributors to CPU. These are the queries to investigate if CPU load is an issue.



Analyzing Statistics for Running Queries

In Amazon RDS Performance Insights, you can find statistics on running queries in the **Top Load Items** section. To view these statistics, view top SQL. Performance Insights collects statistics only for the most common queries, and these usually match the top queries by load shown in the Performance Insights dashboard.

Topics

- [Statistics for Amazon RDS for Oracle \(p. 333\)](#)
- [Analyzing Oracle Metrics for SQL Statements That Are Running \(p. 334\)](#)

Statistics for Amazon RDS for Oracle

The following SQL statistics are available for Oracle DB instances.

Metric	Unit
db.sql.stats.executions_per_sec	Number of executions per second
db.sql.stats.elapsed_time_per_sec	Average active executions (AAE)
db.sql.stats.rows_processed_per_sec	Rows processed per second
db.sql.stats.buffer_gets_per_sec	Buffer gets per second
db.sql.stats.physical_read_requests_per_sec	Physical reads per second
db.sql.stats.physical_write_requests_per_sec	Physical writes per second
db.sql.stats.total_sharable_mem_per_sec	Total shareable memory per second (in bytes)
db.sql.stats.cpu_time_per_sec	CPU time per second (in ms)

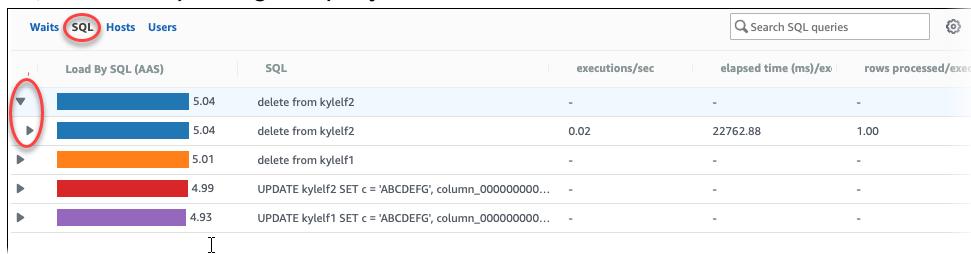
The following metrics provide per call statistics for a SQL statement.

Metric	Unit
db.sql.stats.elapsed_time_per_exec	Elapsed time per executions (in ms)

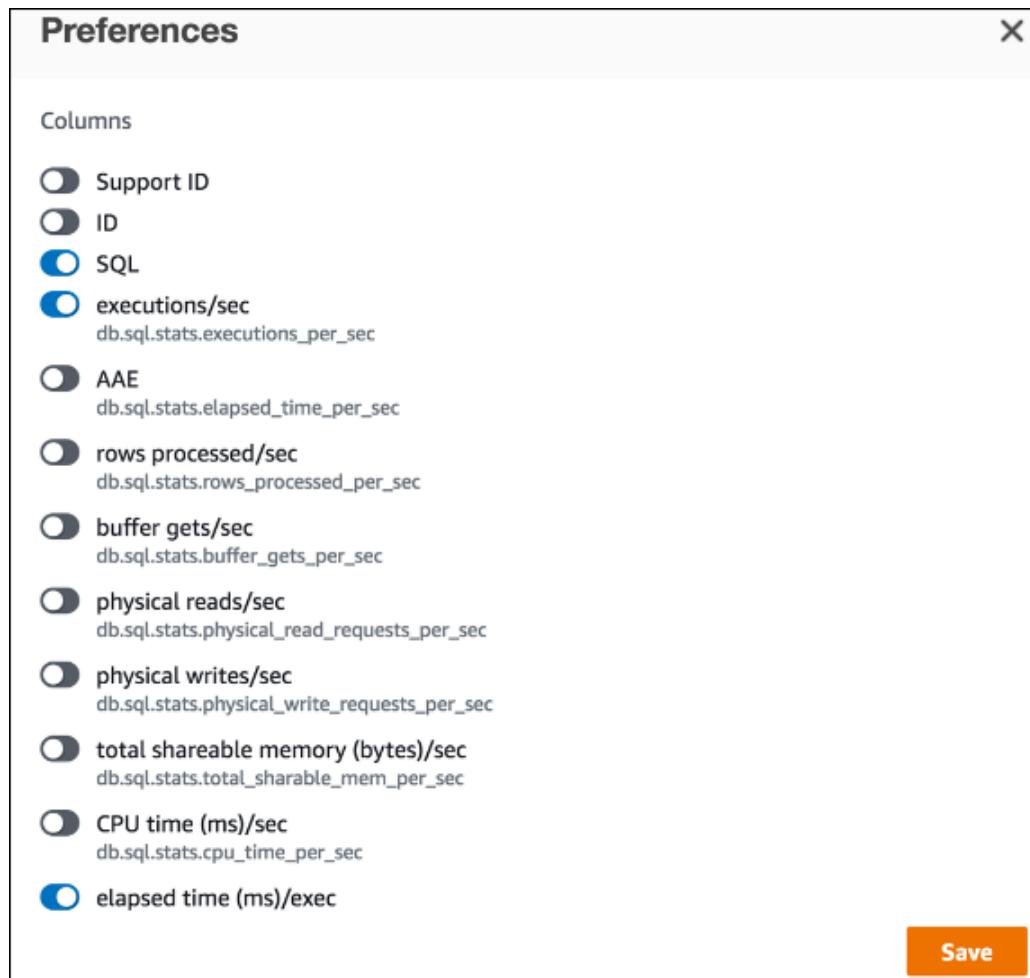
Metric	Unit
db.sql.stats.rows_processed_per_exec	Rows processed per execution
db.sql.stats.buffer_gets_per_exec	Buffer gets per execution
db.sql.stats.physical_read_requests_per_exec	Physical reads per execution
db.sql.stats.physical_write_requests_per_exec	Physical writes per execution
db.sql.stats.total_sharable_mem_per_exec	Total shareable memory per execution (in bytes)
db.sql.stats.cpu_time_per_exec	CPU time per execution (in ms)

Analyzing Oracle Metrics for SQL Statements That Are Running

Using the AWS Management Console, you can view the metrics for a running SQL query by choosing the **SQL** tab and expanding the query.



Choose which statistics to display by choosing the gear icon in the upper-right corner of the chart.



Viewing More SQL Text in the Performance Insights Dashboard

By default, each row in the **Top Load Items** table shows 500 bytes of SQL text for each SQL statement. When a SQL statement is larger than 500 bytes, you can view more of the SQL statement by opening the statement in the Performance Insights dashboard. The Performance Insights dashboard can display up to 4,096 bytes for a SQL statement. You can copy the displayed SQL statement. To view more than 4,096 bytes, choose **Download full SQL** to view the SQL text up to the DB engine limit.

The limit for SQL text depends on the DB engine. The following limits apply:

- Amazon RDS for Microsoft SQL Server – 4,096 characters
- Amazon RDS for MySQL and MariaDB – 1,024 bytes
- Amazon RDS for Oracle – 1,000 bytes
- Amazon RDS for PostgreSQL – Set by the `track_activity_query_size` DB instance parameter

For Amazon RDS for PostgreSQL DB instances, you can control the limit of the SQL text size by setting the `track_activity_query_size` DB instance parameter, up to 102,400 bytes. You can use the AWS Management Console to download SQL text up to the limit you set with this parameter. For more information, see [Setting the SQL Text Limit for Amazon RDS for PostgreSQL DB Instances \(p. 337\)](#).

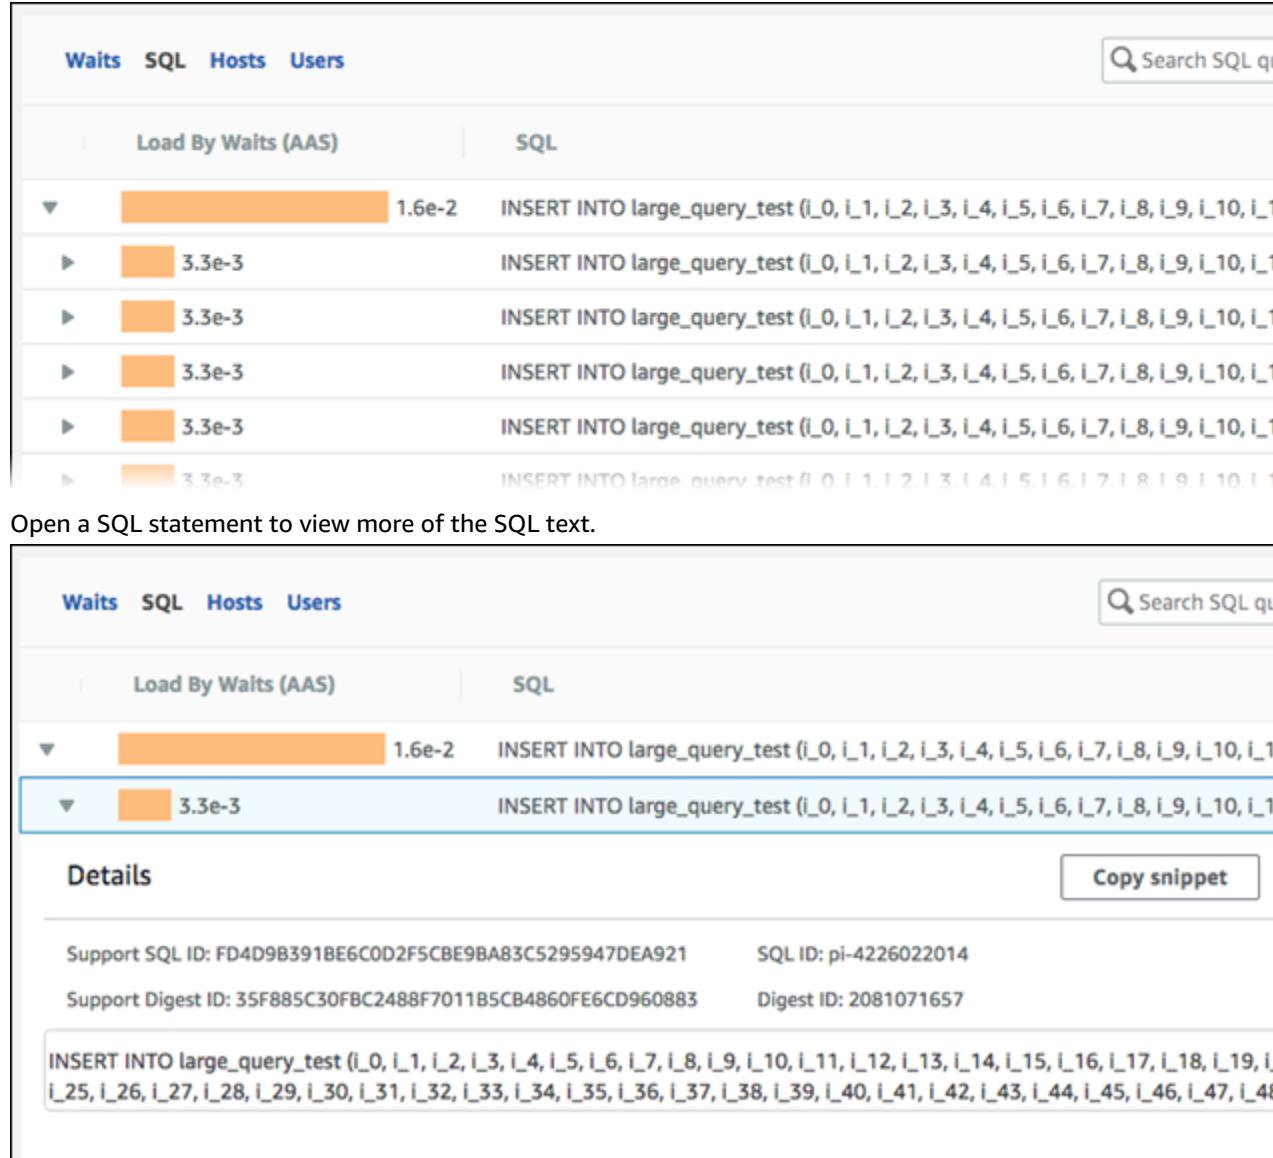
Important

Currently, you can only view and download more SQL text with the AWS Management Console. The AWS Performance Insights CLI and API can return a maximum of 500 bytes of text.

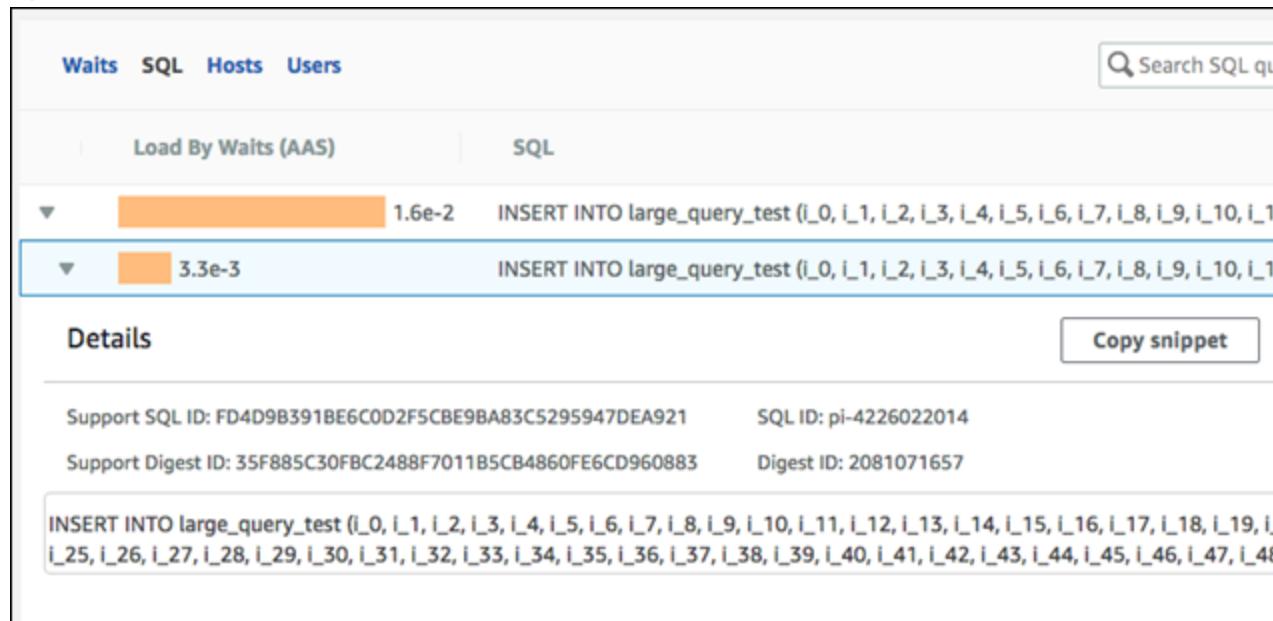
To view more SQL text in the Performance Insights dashboard

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Performance Insights**.
3. Choose a DB instance. The Performance Insights dashboard is displayed for that DB instance.

SQL statements with text larger than 500 bytes look similar to the following image.



4. Open a SQL statement to view more of the SQL text.



The Performance Insights dashboard can display up to 4,096 bytes for each SQL statement.

5. (Optional) Choose **Copy snippet** to copy the displayed SQL statement, or choose **Download full SQL** to download the SQL statement to view the SQL text up to the DB engine limit.

Note

To copy or download the SQL statement, disable pop-up blockers.

Setting the SQL Text Limit for Amazon RDS for PostgreSQL DB Instances

For Amazon RDS for PostgreSQL DB instances, you can control the limit for the SQL text that can be shown on the Performance Insights dashboard.

To do so, modify the `track_activity_query_size` DB instance parameter. The default setting for the `track_activity_query_size` parameter is 1,024 bytes.

You can increase the number of bytes to increase the SQL text size visible in the Performance Insights dashboard. The limit for the parameter is 10,240 bytes. For more information about the `track_activity_query_size` DB instance parameter, see [Run-time Statistics](#) in the PostgreSQL documentation.

To modify the parameter, change the parameter setting in the parameter group that is associated with the Amazon RDS for PostgreSQL DB instance.

If the Amazon RDS for PostgreSQL DB instance is using the default parameter group, complete the following steps:

1. Create a new DB instance parameter group for the appropriate DB engine and DB engine version.
2. Set the parameter in the new parameter group.
3. Associate the new parameter group with the DB instance.

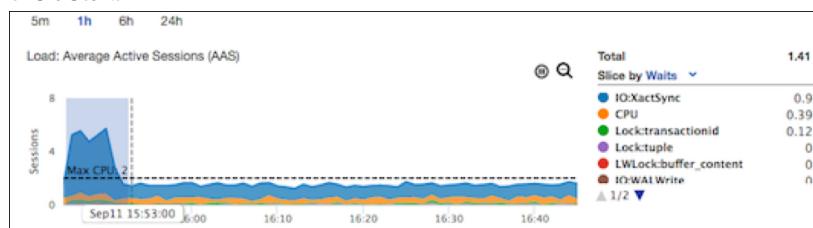
For information about setting a DB instance parameter, see [Modifying Parameters in a DB Parameter Group \(p. 213\)](#).

Additional User Interface Features

You can use other features of the Performance Insights user interface to help analyze performance data.

Click-and-Drag Zoom In

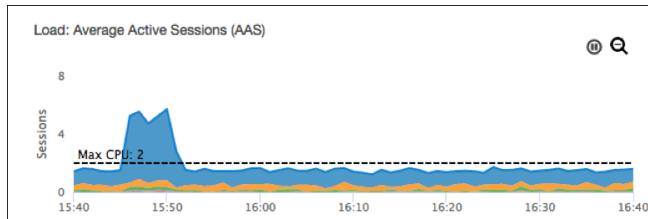
In the Performance Insights interface, you can choose a small portion of the load chart and zoom in on the detail.



To zoom in on a portion of the load chart, choose the start time and drag to the end of the time period you want. When you do this, the selected area is highlighted. When you release the mouse, the load chart zooms in on the selected region, and the **Top N** table is recalculated.

Pause and Zoom Out

In the upper-right corner of the load chart, you can find the **Pause** and **Zoom out** tools.



When you choose **Pause**, the load chart stops autorefreshing. When you choose **Pause** again, the chart resumes autorefreshing.

When you choose **Zoom out**, the load chart zooms out to the next largest time interval.

Performance Insights API

The Amazon RDS Performance Insights API provides visibility into the performance of your RDS instance, when Performance Insights is enabled for supported engine types. Amazon CloudWatch Logs provides the authoritative source for vended monitoring metrics for AWS services. Performance Insights offers a domain-specific view of database load measured as average active sessions and provided to API consumers as a two-dimensional time-series dataset. The time dimension of the data provides database load data for each time point in the queried time range. Each time point decomposes overall load in relation to the requested dimensions, such as `SQL`, `Wait-event`, `User`, or `Host`, measured at that time point.

Amazon RDS Performance Insights monitors your Amazon RDS DB instance so that you can analyze and troubleshoot database performance. One way to view Performance Insights data is in the AWS Management Console. Performance Insights also provides a public API so that you can query your own data. The API can be used to offload data into a database, add Performance Insights data to existing monitoring dashboards, or to build monitoring tools. To use the Performance Insights API, enable Performance Insights on one of your Amazon RDS DB instances. For information about enabling Performance Insights, see [Enabling Performance Insights \(p. 320\)](#).

The Performance Insights API provides the following operations.

Performance Insights Operation	AWS CLI Command	Description
<code>DescribeDimensionKeys</code>	<code>aws rds describe-dimension-keys</code>	Retrieves the top N dimension keys for a metric for a specific time period.
<code>GetResourceMetrics</code>	<code>aws rds pi get-resource-metrics</code>	Retrieves Performance Insights metrics for a set of data sources, over a time period. You can provide specific dimension groups and dimensions, and provide aggregation and filtering criteria for each group.

For more information about the Performance Insights API, see the [Amazon RDS Performance Insights API Reference](#).

AWS CLI for Performance Insights

You can view Performance Insights data using the AWS CLI. You can view help for the AWS CLI commands for Performance Insights by entering the following on the command line.

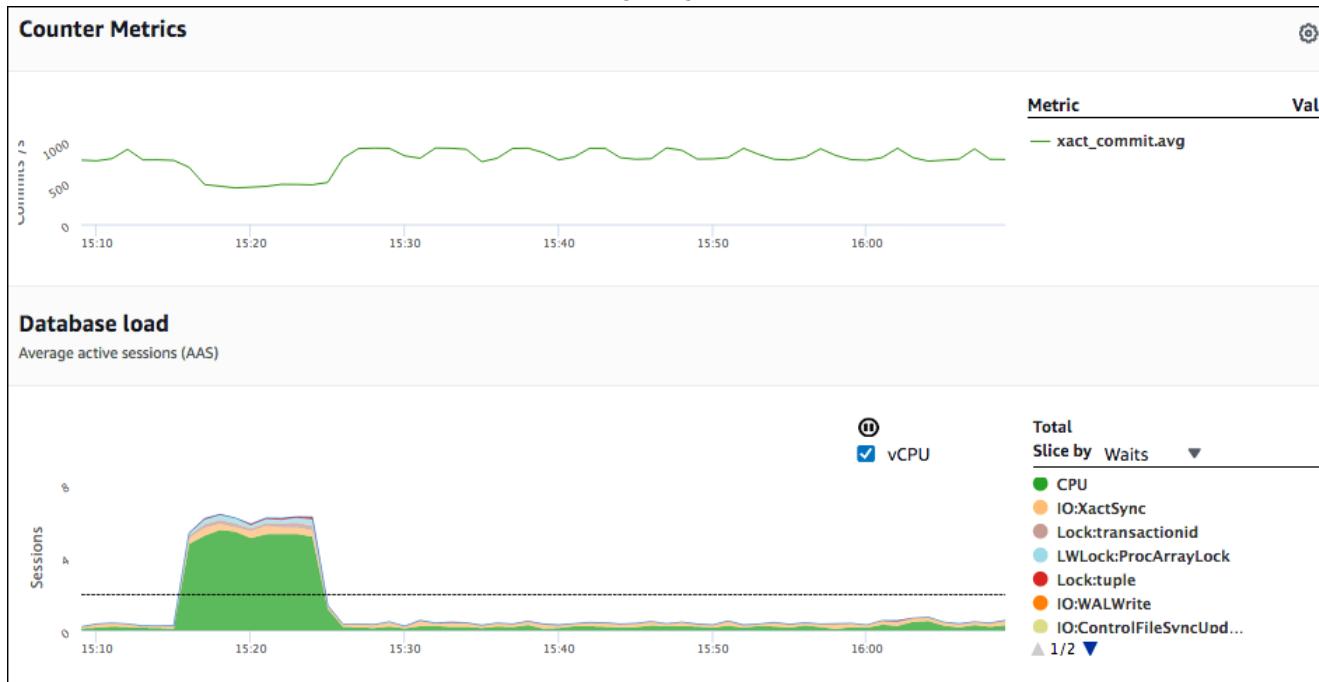
```
aws pi help
```

If you don't have the AWS CLI installed, see [Installing the AWS Command Line Interface](#) in the *AWS CLI User Guide* for information about installing it.

Retrieving Time-Series Metrics

The `GetResourceMetrics` operation retrieves one or more time-series metrics from the Performance Insights data. `GetResourceMetrics` requires a metric and time period, and returns a response with a list of data points.

For example, the AWS Management Console uses `GetResourceMetrics` in two places in the Performance Insights dashboard. `GetResourceMetrics` is used to populate the **Counter Metrics** chart and in the **Database Load** chart, as seen in the following image.



All the metrics returned by `GetResourceMetrics` are standard time-series metrics with one exception. The exception is `db.load`, which is the core metric in Performance Insights. This metric is displayed in the **Database Load** chart. The `db.load` metric is different from the other time-series metrics because you can break it into subcomponents called dimensions. In the previous image, `db.load` is broken down and grouped by the waits states that make up the `db.load`.

Note

`GetResourceMetrics` can also return the `db.sampleload` metric, but the `db.load` metric is appropriate in most cases.

For information about the counter metrics returned by `GetResourceMetrics`, see [Performance Insights Counters \(p. 352\)](#).

The following calculations are supported for the metrics:

- Average – The average value for the metric over a period of time. Append `.avg` to the metric name.

- Minimum – The minimum value for the metric over a period of time. Append `.min` to the metric name.
- Maximum – The maximum value for the metric over a period of time. Append `.max` to the metric name.
- Sum – The sum of the metric values over a period of time. Append `.sum` to the metric name.
- Sample count – The number of times the metric was collected over a period of time. Append `.sample_count` to the metric name.

For example, assume that a metric is collected for 300 seconds (5 minutes), and that the metric is collected one time each minute. The values for each minute are 1, 2, 3, 4, and 5. In this case, the following calculations are returned:

- Average – 3
- Minimum – 1
- Maximum – 5
- Sum – 15
- Sample count – 5

For information about using the `get-resource-metrics` AWS CLI command, see [get-resource-metrics](#).

For the `--metric-queries` option, specify one or more queries that you want to get results for. Each query consists of a mandatory `Metric` and optional `GroupBy` and `Filter` parameters. The following is an example of a `--metric-queries` option specification.

```
{  
    "Metric": "string",  
    "GroupBy": {  
        "Group": "string",  
        "Dimensions": ["string", ...],  
        "Limit": integer  
    },  
    "Filter": {"string": "string"  
              ...}  
}
```

AWS CLI Examples for Performance Insights

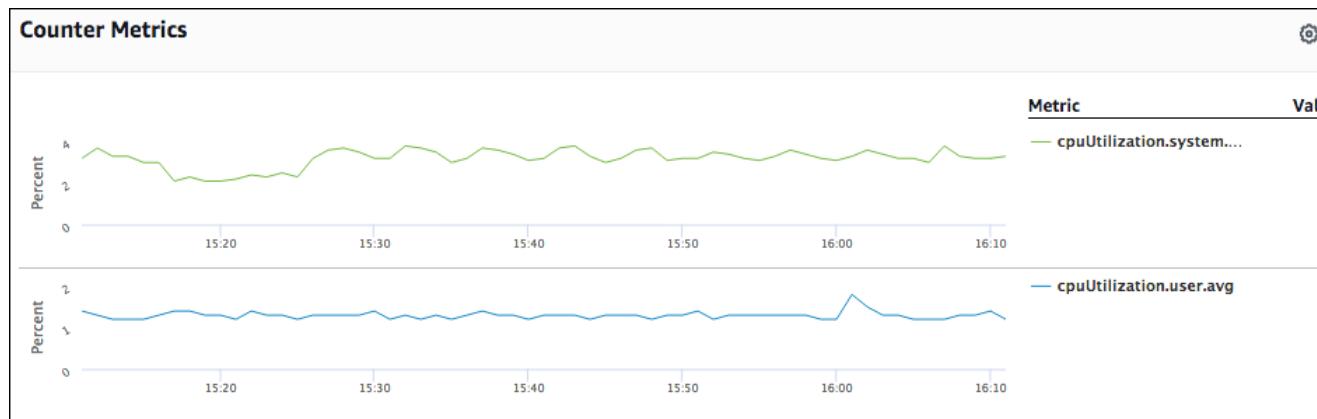
The following are several examples that use the AWS CLI for Performance Insights.

Topics

- [Retrieving Counter Metrics \(p. 340\)](#)
- [Retrieving the DB Load Average for Top Wait Events \(p. 343\)](#)
- [Retrieving the DB Load Average for Top SQL \(p. 345\)](#)
- [Retrieving the DB Load Average Filtered by SQL \(p. 348\)](#)

Retrieving Counter Metrics

The following screenshot shows two counter metrics charts in the AWS Management Console.



The following example shows how to gather the same data that the AWS Management Console uses to generate the two counter metric charts.

For Linux, OS X, or Unix:

```
aws pi get-resource-metrics \
--service-type RDS \
--identifier db-ID \
--start-time 2018-10-30T00:00:00Z \
--end-time 2018-10-30T01:00:00Z \
--period-in-seconds 60 \
--metric-queries '[{"Metric": "os.cpuUtilization.user.avg" },
 {"Metric": "os.cpuUtilization.idle.avg"}]
```

For Windows:

```
aws pi get-resource-metrics ^
--service-type RDS ^
--identifier db-ID ^
--start-time 2018-10-30T00:00:00Z ^
--end-time 2018-10-30T01:00:00Z ^
--period-in-seconds 60 ^
--metric-queries '[{"Metric": "os.cpuUtilization.user.avg" },
 {"Metric": "os.cpuUtilization.idle.avg"}]'
```

You can also make a command easier to read by specifying a file for the `--metrics-query` option. The following example uses a file called `query.json` for the option. The file has the following contents.

```
[ {
  "Metric": "os.cpuUtilization.user.avg"
},
{
  "Metric": "os.cpuUtilization.idle.avg"
}]
```

Run the following command to use the file.

For Linux, OS X, or Unix:

```
aws pi get-resource-metrics \
--service-type RDS \
--identifier db-ID \
--start-time 2018-10-30T00:00:00Z \
--end-time 2018-10-30T01:00:00Z \
--period-in-seconds 60 \
--metric-queries file://query.json
```

For Windows:

```
aws pi get-resource-metrics ^
--service-type RDS ^
--identifier db-ID ^
--start-time 2018-10-30T00:00:00Z ^
--end-time 2018-10-30T01:00:00Z ^
--period-in-seconds 60 ^
--metric-queries file://query.json
```

The preceding example specifies the following values for the options:

- **--service-type** – RDS for Amazon RDS
- **--identifier** – The resource ID for the DB instance
- **--start-time** and **--end-time** – The ISO 8601 DateTime values for the period to query, with multiple supported formats

It queries for a one-hour time range:

- **--period-in-seconds** – 60 for a per-minute query
- **--metric-queries** – An array of two queries, each just for one metric.

The metric name uses dots to classify the metric in a useful category, with the final element being a function. In the example, the function is `avg` for each query. As with Amazon CloudWatch, the supported functions are `min`, `max`, `total`, and `avg`.

The response looks similar to the following.

```
{
    "Identifier": "db-XXX",
    "AlignedStartTime": 1540857600.0,
    "AlignedEndTime": 1540861200.0,
    "MetricList": [
        { //A list of key/datapoints
            "Key": {
                "Metric": "os.cpuUtilization.user.avg" //Metric1
            },
            "DataPoints": [
                //Each list of datapoints has the same timestamps and same number of items
                {
                    "Timestamp": 1540857660.0, //Minute1
                    "Value": 4.0
                },
                ...
            ]
        }
    ]
}
```

```

        {
            "Timestamp": 1540857720.0, //Minute2
            "Value": 4.0
        },
        {
            "Timestamp": 1540857780.0, //Minute 3
            "Value": 10.0
        }
        //... 60 datapoints for the os.cpuUtilization.user.avg metric
    ]
},
{
    "Key": {
        "Metric": "os.cpuUtilization.idle.avg" //Metric2
    },
    "DataPoints": [
        {
            "Timestamp": 1540857660.0, //Minute1
            "Value": 12.0
        },
        {
            "Timestamp": 1540857720.0, //Minute2
            "Value": 13.5
        },
        //... 60 datapoints for the os.cpuUtilization.idle.avg metric
    ]
}
] //end of MetricList
} //end of response

```

The response has an `Identifier`, `AlignedStartTime` and `AlignedEndTime`. Because the `--period-in-seconds` value was 60, the start and end times have been aligned to the minute. If the `--period-in-seconds` was 3600, the start and end times would have been aligned to the hour.

The `MetricList` in the response has a number of entries, each with a `Key` and a `DataPoints` entry. Each `DataPoint` has a `Timestamp` and a `Value`. Each `DataPoints` list has 60 data points because the queries are for per-minute data over an hour, with `Timestamp1/Minute1`, `Timestamp2/Minute2`, and so on up to `Timestamp60/Minute60`.

Because the query is for two different counter metrics, there are two elements in the response `MetricList`.

Retrieving the DB Load Average for Top Wait Events

The following example is the same query that the AWS Management Console uses to generate a stacked area line graph. This example retrieves the `db.load.avg` for the last hour with load divided according to the top seven wait events. The command is the same as the command in [Retrieving Counter Metrics \(p. 340\)](#). However, the `query.json` file has the following contents.

```

[
    {
        "Metric": "db.load.avg",
        "GroupBy": { "Group": "db.wait_event", "Limit": 7 }
    }
]

```

Run the following command.

For Linux, OS X, or Unix:

```
aws pi get-resource-metrics \
--service-type RDS \
--identifier db-ID \
--start-time 2018-10-30T00:00:00Z \
--end-time 2018-10-30T01:00:00Z \
--period-in-seconds 60 \
--metric-queries file://query.json
```

For Windows:

```
aws pi get-resource-metrics ^
--service-type RDS ^
--identifier db-ID ^
--start-time 2018-10-30T00:00:00Z ^
--end-time 2018-10-30T01:00:00Z ^
--period-in-seconds 60 ^
--metric-queries file://query.json
```

The example specifies the metric of db.load.avg and a GroupBy of the top seven wait events. For details about valid values for this example, see [DimensionGroup](#) in the *Performance Insights API Reference*.

The response looks similar to the following.

```
{
    "Identifier": "db-XXX",
    "AlignedStartTime": 1540857600.0,
    "AlignedEndTime": 1540861200.0,
    "MetricList": [
        { //A list of key/datapoints
            "Key": {
                //A Metric with no dimensions. This is the total db.load.avg
                "Metric": "db.load.avg"
            },
            "DataPoints": [
                //Each list of datapoints has the same timestamps and same number of items
                {
                    "Timestamp": 1540857660.0, //Minute1
                    "Value": 0.5166666666666667
                },
                {
                    "Timestamp": 1540857720.0, //Minute2
                    "Value": 0.3833333333333336
                },
                {
                    "Timestamp": 1540857780.0, //Minute 3
                    "Value": 0.2666666666666666
                }
                //... 60 datapoints for the total db.load.avg key
            ]
        },
        {
            "Key": {
                //Another key. This is db.load.avg broken down by CPU
                "Metric": "db.load.avg",
                "Dimensions": {
                    "db.wait_event.name": "CPU",

```

```

        "db.wait_event.type": "CPU"
    }
},
"DataPoints": [
{
    "Timestamp": 1540857660.0, //Minute1
    "Value": 0.35
},
{
    "Timestamp": 1540857720.0, //Minute2
    "Value": 0.15
},
//... 60 datapoints for the CPU key
]
},
//... In total we have 8 key/datapoints entries, 1) total, 2-8) Top Wait Events
] //end of MetricList
} //end of response

```

In this response, there are eight entries in the `MetricList`. There is one entry for the total `db.load.avg`, and seven entries each for the `db.load.avg` divided according to one of the top seven wait events. Unlike in the first example, because there was a grouping dimension, there must be one key for each grouping of the metric. There can't be only one key for each metric, as in the basic counter metric use case.

Retrieving the DB Load Average for Top SQL

The following example groups `db.wait_events` by the top 10 SQL statements. There are two different groups for SQL statements:

- `db.sql` – The full SQL statement, such as `select * from customers where customer_id = 123`
- `db.sql_tokenized` – The tokenized SQL statement, such as `select * from customers where customer_id = ?`

When analyzing database performance, it can be useful to consider SQL statements that only differ by their parameters as one logic item. So, you can use `db.sql_tokenized` when querying. However, especially when you are interested in explain plans, sometimes it's more useful to examine full SQL statements with parameters, and query grouping by `db.sql`. There is a parent-child relationship between tokenized and full SQL, with multiple full SQL (children) grouped under the same tokenized SQL (parent).

The command in this example is the similar to the command in [Retrieving the DB Load Average for Top Wait Events \(p. 343\)](#). However, the `query.json` file has the following contents.

```

[
{
    "Metric": "db.load.avg",
    "GroupBy": { "Group": "db.sql_tokenized", "Limit": 10 }
}
]

```

The following example uses `db.sql_tokenized`.

For Linux, OS X, or Unix:

```
aws pi get-resource-metrics \
--service-type RDS \
--identifier db-ID \
--start-time 2018-10-29T00:00:00Z \
--end-time 2018-10-30T00:00:00Z \
--period-in-seconds 3600 \
--metric-queries file://query.json
```

For Windows:

```
aws pi get-resource-metrics ^
--service-type RDS ^
--identifier db-ID ^
--start-time 2018-10-29T00:00:00Z ^
--end-time 2018-10-30T00:00:00Z ^
--period-in-seconds 3600 ^
--metric-queries file://query.json
```

This example queries over 24 hours, with a one hour period-in-seconds.

The example specifies the metric of `db.load.avg` and a `GroupBy` of the top seven wait events. For details about valid values for this example, see [DimensionGroup](#) in the *Performance Insights API Reference*.

The response looks similar to the following.

```
{
    "AlignedStartTime": 1540771200.0,
    "AlignedEndTime": 1540857600.0,
    "Identifier": "db-XXX",

    "MetricList": [ //11 entries in the MetricList
        {
            "Key": { //First key is total
                "Metric": "db.load.avg"
            }
            "DataPoints": [ //Each DataPoints list has 24 per-hour Timestamps and a value
                {
                    "Value": 1.6964980544747081,
                    "Timestamp": 1540774800.0
                },
                //... 24 datapoints
            ]
        },
        {
            "Key": { //Next key is the top tokenized SQL
                "Dimensions": {
                    "db.sql_tokenized.statement": "INSERT INTO authors (id,name,email) VALUES\n( nextval(?) ,?,?)",
                    "db.sql_tokenized.db_id": "pi-2372568224",
                    "db.sql_tokenized.id": "AKIAIOSFODNN7EXAMPLE"
                },
                "Metric": "db.load.avg"
            },
            "DataPoints": [ //... 24 datapoints
            ]
        },
        // In total 11 entries, 10 Keys of top tokenized SQL, 1 total key
    ]
}
```

```
    ] //End of MetricList
} //End of response
```

This response has 11 entries in the `MetricList` (1 total, 10 top tokenized SQL), with each entry having 24 per-hour `DataPoints`.

For tokenized SQL, there are three entries in each dimensions list:

- `db.sql_tokenized.statement` – The tokenized SQL statement.
- `db.sql_tokenized.db_id` – Either the native database ID used to refer to the SQL, or a synthetic ID that Performance Insights generates for you if the native database ID isn't available. This example returns the `pi-2372568224` synthetic ID.
- `db.sql_tokenized.statement` – The ID of the query inside Performance Insights.

In the AWS Management Console, this ID is called the Support ID. It's named this because the ID is data that AWS Support can examine to help you troubleshoot an issue with your database. AWS takes the security and privacy of your data extremely seriously, and almost all data is stored encrypted with your AWS KMS key. Therefore, nobody inside AWS can look at this data. In the example preceding, both the `tokenized_statement` and the `tokenized.db_id` are stored encrypted. If you have an issue with your database, AWS Support can help you by referencing the Support ID.

When querying, it might be convenient to specify a `Group` in `GroupBy`. However, for finer-grained control over the data that's returned, specify the list of dimensions. For example, if all that is needed is the `db.sql_tokenized.statement`, then a `Dimensions` attribute can be added to the `query.json` file.

```
[  
  {  
    "Metric": "db.load.avg",  
    "GroupBy": {  
      "Group": "db.sql_tokenized",  
      "Dimensions": ["db.sql_tokenized.statement"],  
      "Limit": 10  
    }  
  }  
]
```

Retrieving the DB Load Average Filtered by SQL



The preceding image shows that a particular query is selected, and the top average active sessions stacked area line graph is scoped to that query. Although the query is still for the top seven overall wait events, the value of the response is filtered. The filter causes it to take into account only sessions that are a match for the particular filter.

The corresponding API query in this example is similar to the command in [Retrieving the DB Load Average for Top SQL \(p. 345\)](#). However, the `query.json` file has the following contents.

```
[  
  {  
    "Metric": "db.load.avg",  
    "GroupBy": { "Group": "db.wait_event", "Limit": 5 },  
    "Filter": { "db.sql_tokenized.id": "AKIAIOSFODNN7EXAMPLE" }  
  }  
]
```

For Linux, OS X, or Unix:

```
aws pi get-resource-metrics \  
  --service-type RDS \  
  --identifier db-ID \  
  --start-time 2018-10-30T00:00:00Z \  
  --end-time 2018-10-30T01:00:00Z \  
  --period-in-seconds 60 \  
  --metric-queries file://query.json
```

For Windows:

```
aws pi get-resource-metrics ^  
  --service-type RDS ^  
  --identifier db-ID ^
```

```
--start-time 2018-10-30T00:00:00Z ^
--end-time 2018-10-30T01:00:00Z ^
--period-in-seconds 60 ^
--metric-queries file://query.json
```

The response looks similar to the following.

```
{
  "Identifier": "db-XXX",
  "AlignedStartTime": 1556215200.0,
  "MetricList": [
    {
      "Key": {
        "Metric": "db.load.avg"
      },
      "DataPoints": [
        {
          "Timestamp": 1556218800.0,
          "Value": 1.4878117913832196
        },
        {
          "Timestamp": 1556222400.0,
          "Value": 1.192823803967328
        }
      ]
    },
    {
      "Key": {
        "Metric": "db.load.avg",
        "Dimensions": {
          "db.wait_event.type": "io",
          "db.wait_event.name": "wait/io/aurora_redo_log_flush"
        }
      },
      "DataPoints": [
        {
          "Timestamp": 1556218800.0,
          "Value": 1.1360544217687074
        },
        {
          "Timestamp": 1556222400.0,
          "Value": 1.058051341890315
        }
      ]
    },
    {
      "Key": {
        "Metric": "db.load.avg",
        "Dimensions": {
          "db.wait_event.type": "io",
          "db.wait_event.name": "wait/io/table/sql/handler"
        }
      },
      "DataPoints": [
        {
          "Timestamp": 1556218800.0,
          "Value": 0.16241496598639457
        },
        {
          "Timestamp": 1556222400.0,
          "Value": 0.05163360560093349
        }
      ]
    }
  ]
}
```

```

        },
        {
            "Key": {
                "Metric": "db.load.avg",
                "Dimensions": {
                    "db.wait_event.type": "synch",
                    "db.wait_event.name": "wait/synch/mutex/innodb/
aurora_lock_thread_slot_futex"
                }
            },
            "DataPoints": [
                {
                    "Timestamp": 1556218800.0,
                    "Value": 0.11479591836734694
                },
                {
                    "Timestamp": 1556222400.0,
                    "Value": 0.013127187864644107
                }
            ]
        },
        {
            "Key": {
                "Metric": "db.load.avg",
                "Dimensions": {
                    "db.wait_event.type": "CPU",
                    "db.wait_event.name": "CPU"
                }
            },
            "DataPoints": [
                {
                    "Timestamp": 1556218800.0,
                    "Value": 0.05215419501133787
                },
                {
                    "Timestamp": 1556222400.0,
                    "Value": 0.05805134189031505
                }
            ]
        },
        {
            "Key": {
                "Metric": "db.load.avg",
                "Dimensions": {
                    "db.wait_event.type": "synch",
                    "db.wait_event.name": "wait/synch/mutex/innodb/lock_wait_mutex"
                }
            },
            "DataPoints": [
                {
                    "Timestamp": 1556218800.0,
                    "Value": 0.017573696145124718
                },
                {
                    "Timestamp": 1556222400.0,
                    "Value": 0.002333722287047841
                }
            ]
        },
        ],
        "AlignedEndTime": 1556222400.0
    } //end of response
}

```

In this response, all values are filtered according to the contribution of tokenized SQL AKIAIOSFODNN7EXAMPLE specified in the query.json file. The keys also might follow a different order than a query without a filter, because it's the top five wait events that affected the filtered SQL.

Performance Insights Metrics Published to Amazon CloudWatch

Performance Insights automatically publishes metrics to Amazon CloudWatch. The same data can be queried from Performance Insights, but having the metrics in CloudWatch makes it easy to add CloudWatch alarms. It also makes it easy to add the metrics to existing CloudWatch Dashboards.

Metric	Description
DBLoad	The number of active sessions for the DB engine. Typically, you want the data for the average number of active sessions. In Performance Insights, this data is queried as db.load.avg.
DBLoadCPU	The number of active sessions where the wait event type is CPU. In Performance Insights, this data is queried as db.load.avg, filtered by the wait event type CPU.
DBLoadNonCPU	The number of active sessions where the wait event type is not CPU.

Note

These metrics are published to CloudWatch only if there is load on the DB instance.

You can examine these metrics using the CloudWatch console, the AWS CLI, or the CloudWatch API.

For example, you can get the statistics for the DBLoad metric by running the [get-metric-statistics](#) command.

```
aws cloudwatch get-metric-statistics --region us-west-2 --namespace AWS/RDS --metric-name DBLoad --period 60 --statistics Average --start-time 1532035185 --end-time 1532036185 --dimensions Name=DBInstanceIdentifier,Value=db-loadtest-0
```

This example generates output similar to the following.

```
{
  "Datapoints": [
    {
      "Timestamp": "2018-07-19T21:30:00Z",
      "Unit": "None",
      "Average": 2.1
    },
    {
      "Timestamp": "2018-07-19T21:34:00Z",
      "Unit": "None",
      "Average": 1.7
    },
    {
      "Timestamp": "2018-07-19T21:35:00Z",
      "Unit": "None",
      "Average": 1.7
    }
  ]
}
```

```
    "Unit": "None",
    "Average": 2.8
},
{
"Timestamp": "2018-07-19T21:31:00Z",
"Unit": "None",
"Average": 1.5
},
{
"Timestamp": "2018-07-19T21:32:00Z",
"Unit": "None",
"Average": 1.8
},
{
"Timestamp": "2018-07-19T21:29:00Z",
"Unit": "None",
"Average": 3.0
},
{
"Timestamp": "2018-07-19T21:33:00Z",
"Unit": "None",
"Average": 2.4
}
],
"Label": "DBLoad"
}
```

For more information about CloudWatch, see [What is Amazon CloudWatch?](#) in the *Amazon CloudWatch User Guide*.

Performance Insights Counters

With counter metrics, you can customize the Performance Insights dashboard to include up to 10 additional graphs. These graphs that show a selection of dozens of operating system and database performance metrics. This information can be correlated with database load to help identify and analyze performance problems.

Topics

- [Performance Insights Operating System Counters \(p. 352\)](#)
- [Performance Insights Counters for Amazon RDS for MariaDB and MySQL \(p. 354\)](#)
- [Performance Insights Counters for Amazon RDS for Microsoft SQL Server \(p. 357\)](#)
- [Performance Insights Counters for Amazon RDS for Oracle \(p. 359\)](#)
- [Performance Insights Counters for Amazon RDS for PostgreSQL \(p. 360\)](#)

Performance Insights Operating System Counters

The following operating system counters are available with Performance Insights for Aurora PostgreSQL. You can find definitions for these metrics in [Viewing Enhanced Monitoring by Using CloudWatch Logs \(p. 313\)](#).

Counter	Type	Metric
active	memory	os.memory.active
buffers	memory	os.memory.buffers
cached	memory	os.memory.cached

Counter	Type	Metric
dirty	memory	os.memory.dirty
free	memory	os.memory.free
hugePagesFree	memory	os.memory.hugePagesFree
hugePagesRsvd	memory	os.memory.hugePagesRsvd
hugePagesSize	memory	os.memory.hugePagesSize
hugePagesSurp	memory	os.memory.hugePagesSurp
hugePagesTotal	memory	os.memory.hugePagesTotal
inactive	memory	os.memory.inactive
mapped	memory	os.memory.mapped
pageTables	memory	os.memory.pageTables
slab	memory	os.memory.slab
total	memory	os.memory.total
writeback	memory	os.memory.writeback
guest	cpuUtilization	os.cpuUtilization.guest
idle	cpuUtilization	os.cpuUtilization.idle
irq	cpuUtilization	os.cpuUtilization.irq
nice	cpuUtilization	os.cpuUtilization.nice
steal	cpuUtilization	os.cpuUtilization.steal
system	cpuUtilization	os.cpuUtilization.system
total	cpuUtilization	os.cpuUtilization.total
user	cpuUtilization	os.cpuUtilization.user
wait	cpuUtilization	os.cpuUtilization.wait
avgQueueLen	diskIO	os.diskIO.avgQueueLen
avgReqSz	diskIO	os.diskIO.avgReqSz
await	diskIO	os.diskIO.await
readIOsPS	diskIO	os.diskIO.readIOsPS
readKb	diskIO	os.diskIO.readKb
readKbPS	diskIO	os.diskIO.readKbPS
rrqmPS	diskIO	os.diskIO.rrqmPS
tps	diskIO	os.diskIO.tps
util	diskIO	os.diskIO.util

Counter	Type	Metric
writelOsPS	diskIO	os.diskIO.writelOsPS
writeKb	diskIO	os.diskIO.writeKb
writeKbPS	diskIO	os.diskIO.writeKbPS
wrqmPS	diskIO	os.diskIO.wrqmPS
blocked	tasks	os.tasks.blocked
running	tasks	os.tasks.running
sleeping	tasks	os.tasks.sleeping
stopped	tasks	os.tasks.stopped
total	tasks	os.tasks.total
zombie	tasks	os.tasks.zombie
one	loadAverageMinute	os.loadAverageMinute.one
fifteen	loadAverageMinute	os.loadAverageMinute.fifteen
five	loadAverageMinute	os.loadAverageMinute.five
cached	swap	os.swap.cached
free	swap	os.swap.free
in	swap	os.swap.in
out	swap	os.swap.out
total	swap	os.swap.total
maxFiles	fileSys	os.fileSys.maxFiles
usedFiles	fileSys	os.fileSys.usedFiles
usedFilePercent	fileSys	os.fileSys.usedFilePercent
usedPercent	fileSys	os.fileSys.usedPercent
used	fileSys	os.fileSys.used
total	fileSys	os.fileSys.total
rx	network	os.network.rx
tx	network	os.network.tx
numVCPU	general	os.general.numVCPU

Performance Insights Counters for Amazon RDS for MariaDB and MySQL

The following database counters are available with Performance Insights for Amazon RDS for MariaDB and MySQL.

Topics

- [Native Counters for RDS MariaDB and RDS MySQL \(p. 355\)](#)
- [Non-Native Counters for Amazon RDS for MariaDB and MySQL \(p. 356\)](#)

Native Counters for RDS MariaDB and RDS MySQL

You can find definitions for these native metrics in [Server Status Variables](#) in the MySQL documentation.

Counter	Type	Unit	Metric
Com_analyze	SQL	Queries per second	db.SQL.Com_analyze
Com_optimize	SQL	Queries per second	db.SQL.Com_optimize
Com_select	SQL	Queries per second	db.SQL.Com_select
Innodb_rows_deleted	SQL	Rows per second	db.SQL.Innodb_rows_deleted
Innodb_rows_inserted	SQL	Rows per second	db.SQL.Innodb_rows_inserted
Innodb_rows_read	SQL	Rows per second	db.SQL.Innodb_rows_read
Innodb_rows_updated	SQL	Rows per second	db.SQL.Innodb_rows_updated
Select_full_join	SQL	Queries per second	db.SQL.Select_full_join
Select_full_range_join	SQL	Queries per second	db.SQL.Select_full_range_join
Select_range	SQL	Queries per second	db.SQL.Select_range
Select_range_check	SQL	Queries per second	db.SQL.Select_range_check
Select_scan	SQL	Queries per second	db.SQL.Select_scan
Slow_queries	SQL	Queries per second	db.SQL.Slow_queries
Sort_merge_passes	SQL	Queries per second	db.SQL.Sort_merge_passes
Sort_range	SQL	Queries per second	db.SQL.Sort_range
Sort_rows	SQL	Queries per second	db.SQL.Sort_rows
Sort_scan	SQL	Queries per second	db.SQL.Sort_scan
Questions	SQL	Queries per second	db.SQL.Questions
Innodb_row_lock_time	Locks	Milliseconds (average)	db.Locks.Innodb_row_lock_time
Table_locks_immediate	Locks	Requests per second	db.Locks.Table_locks_immediate
Table_locks_waited	Locks	Requests per second	db.Locks.Table_locks_waited
Aborted_clients	Users	Connections	db.Users.Aborted_clients
Aborted_connects	Users	Connections	db.Users.Aborted_connects
Threads_created	Users	Connections	db.Users.Threads_created
Threads_running	Users	Connections	db.Users.Threads_running
Innodb_data_writes	IO	Operations per second	db.IO.Innodb_data_writes

Counter	Type	Unit	Metric
Innodb_dblwr_writes	IO	Operations per second	db.IO.Innodb_dblwr_writes
Innodb_log_write_requests	IO	Operations per second	db.IO.Innodb_log_write_requests
Innodb_log_writes	IO	Operations per second	db.IO.Innodb_log_writes
Innodb_pages_written	IO	Pages per second	db.IO.Innodb_pages_written
Created_tmp_disk_tables	Temp	Tables per second	db.Temp.Created_tmp_disk_tables
Created_tmp_tables	Temp	Tables per second	db.Temp.Created_tmp_tables
Innodb_buffer_pool_pages_Cache	Cache	Pages	db.Cache.Innodb_buffer_pool_pages_Cache
Innodb_buffer_pool_pages_Cached	Cache	Pages	db.Cache.Innodb_buffer_pool_pages_Cached
Innodb_buffer_pool_read_Chests	Cache	Pages per second	db.Cache.Innodb_buffer_pool_read_Chests
Innodb_buffer_pool_readsCache	Cache	Pages per second	db.Cache.Innodb_buffer_pool_readsCache
Opened_tables	Cache	Tables	db.Cache.Opened_tables
Opened_table_definitions	Cache	Tables	db.Cache.Opened_table_definitions
Qcache_hits	Cache	Queries	db.Cache.Qcache_hits

Non-Native Counters for Amazon RDS for MariaDB and MySQL

Non-native counter metrics are counters defined by Amazon RDS. A non-native metric can be a metric that you get with a specific query. A non-native metric also can be a derived metric, where two or more native counters are used in calculations for ratios, hit rates, or latencies.

Counter	Type	Metric	Description	Definition
innodb_buffer_pool_hits		db.Cache.innodb_buffer_pool_hits	The number of reads that InnoDB could satisfy from the buffer pool.	innodb_buffer_pool_read_requests - innodb_buffer_pool_reads
innodb_buffer_pool_hit_rate		db.Cache.innodb_buffer_pool_hit_rate	The percentage of reads that InnoDB could satisfy from the buffer pool.	100 * innodb_buffer_pool_read_requests / (innodb_buffer_pool_read_requests + innodb_buffer_pool_reads)
innodb_buffer_pool_usage		db.Cache.innodb_buffer_pool_usage	The percentage of the InnoDB buffer pool that contains data (pages).	Innodb_buffer_pool_pages_data / Innodb_buffer_pool_pages_total * 100.0

Counter	Type	Metric	Description	Definition
			and Innodb_buffer_pool_pages_total in Server Status Variables in the MySQL documentation.	
query_cache_hits	db.Cache.query	MySQL_resultset cache (query cache) hit ratio.	Qcache_hits / (QCache_hits + Com_select) * 100	
innodb_datafile_writes_to_disk	db.IO.innodb_datafile_writes_to_disk	The total number of InnoDB data file writes to disk, excluding double write and redo logging write operations.	Innodb_data_writes - Innodb_log_writes - Innodb_dblwr_writes	
innodb_rows_changed	db.SQL.innodb_rows_changed	The total number of InnoDB row operations.	db.SQL.Innodb_rows_inserted + db.SQL.Innodb_rows_deleted + db.SQL.Innodb_rows_updated	
active_transactions	db.Transaction.active_transactions	The total number of active transactions.	SELECT COUNT(1) AS active_transactions FROM INFORMATION_SCHEMA.INNODB_TRX	
innodb_deadlocks	db.Locks.innodb_deadlocks	The total number of deadlocks.	SELECT COUNT AS innodb_deadlocks FROM INFORMATION_SCHEMA.INNODB_METRIC WHERE NAME='lock_deadlocks'	
innodb_lock_timeouts	db.Locks.innodb_lock_timeouts	The total number of deadlocks that timed out.	SELECT COUNT AS innodb_lock_timeouts FROM INFORMATION_SCHEMA.INNODB_METRIC WHERE NAME='lock_timeouts'	
innodb_row_lock_waits	db.Locks.innodb_row_lock_waits	The total number of row locks that resulted in a wait.	SELECT COUNT AS innodb_row_lock_waits FROM INFORMATION_SCHEMA.INNODB_METRIC WHERE NAME='lock_row_lock_waits'	

Performance Insights Counters for Amazon RDS for Microsoft SQL Server

The following database counters are available with Performance Insights for RDS for Microsoft SQL Server.

Native Counters for RDS for Microsoft SQL Server

You can find definitions for these native metrics in [Use SQL Server Objects](#) in the Microsoft SQL Server documentation.

Counter	Type	Unit	Metric
Forwarded Records	Access Methods	Records per second	db.Access Methods.Forwarded Records
Page Splits	Access Methods	Splits per second	db.Access Methods.Page Splits
Buffer cache hit ratio	Buffer Manager	Ratio	db.Buffer Manager.Buffer cache hit ratio
Page life expectancy	Buffer Manager	Expectancy in seconds	db.Buffer Manager.Page life expectancy
Page lookups	Buffer Manager	Lookups per second	db.Buffer Manager.Page lookups
Page reads	Buffer Manager	Reads per second	db.Buffer Manager.Page reads
Page writes	Buffer Manager	Writes per second	db.Buffer Manager.Page writes
Active Transactions	Databases	Transactions	db.Databases.Active Transactions (_Total)
Log Bytes Flushed	Databases	Bytes flushed per second	db.Databases.Log Bytes Flushed (_Total)
Log Flush Waits	Databases	Waits per second	db.Databases.Log Flush Waits (_Total)
Log Flushes	Databases	Flushes per second	db.Databases.Log Flushes (_Total)
Write Transactions	Databases	Transactions per second	db.Databases.Write Transactions (_Total)
Processes blocked	General Statistics	Processes blocked	db.General Statistics.Processes blocked
User Connections	General Statistics	Connections	db.General Statistics.User Connections
Latch Waits	Latches	Waits per second	db.Latches.Latch Waits
Number of Deadlocks	Locks	Deadlocks per second	db.Locks.Number of Deadlocks (_Total)

Counter	Type	Unit	Metric
Memory Grants Pending	Memory Manager	Memory grants	db.Memory Manager.Memory Grants Pending
Batch Requests	SQL Statistics	Requests per second	db.SQL Statistics.Batch Requests
SQL Compilations	SQL Statistics	Compilations per second	db.SQL Statistics.SQL Compilations
SQL Re-Compilations	SQL Statistics	Re-compilations per second	db.SQL Statistics.SQL Re-Compilations

Performance Insights Counters for Amazon RDS for Oracle

The following database counters are available with Performance Insights for RDS for Oracle.

Native Counters for RDS for Oracle

You can find definitions for these native metrics in [Statistics Descriptions](#) in the Oracle documentation.

Note

For the CPU used by this session counter metric, the unit has been transformed from the native centiseconds to active sessions to make the value easier to use. For example, CPU send in the DB Load chart represents the demand for CPU. The counter metric CPU used by this session represents the amount of CPU used by Oracle sessions. You can compare CPU send to the CPU used by this session counter metric. When demand for CPU is higher than CPU used, sessions are waiting for CPU time.

Counter	Type	Unit	Metric
CPU used by this session	User	Active sessions	db.User.CPU used by this session
SQL*Net roundtrips to/from client	User	Roundtrips per second	db.User.SQL*Net roundtrips to/from client
Bytes received via SQL*Net from client	User	Bytes per second	db.User.bytes received via SQL*Net from client
User commits	User	Commits per second	db.User.user commits
Logons cumulative	User	Logons per second	db.User.logons cumulative
User calls	User	Calls per second	db.User.user calls
Bytes sent via SQL*Net to client	User	Bytes per second	db.User.bytes sent via SQL*Net to client
User rollbacks	User	Rollbacks per second	db.User.user rollbacks
Redo size	Redo	Bytes per second	db.Redo.redo size
Parse count (total)	SQL	Parses per second	db.SQL.parse count (total)

Counter	Type	Unit	Metric
Parse count (hard)	SQL	Parses per second	db.SQL.parse count (hard)
Table scan rows gotten	SQL	Rows per second	db.SQL.table scan rows gotten
Sorts (memory)	SQL	Sorts per second	db.SQL.sorts (memory)
Sorts (disk)	SQL	Sorts per second	db.SQL.sorts (disk)
Sorts (rows)	SQL	Sorts per second	db.SQL.sorts (rows)
Physical read bytes	Cache	Bytes per second	db.Cache.physical read bytes
DB block gets	Cache	Blocks per second	db.Cache.db block gets
DBWR checkpoints	Cache	Checkpoints per minute	db.Cache.DBWR checkpoints
Physical reads	Cache	Reads per second	db.Cache.physical reads
Consistent gets from cache	Cache	Gets per second	db.Cache.consistent gets from cache
DB block gets from cache	Cache	Gets per second	db.Cache.db block gets from cache
Consistent gets	Cache	Gets per second	db.Cache.consistent gets

Performance Insights Counters for Amazon RDS for PostgreSQL

The following database counters are available with Performance Insights for Amazon RDS for PostgreSQL.

Topics

- [Native Counters for Amazon RDS for PostgreSQL \(p. 360\)](#)
- [Non-Native Counters for Amazon RDS for PostgreSQL \(p. 361\)](#)

Native Counters for Amazon RDS for PostgreSQL

You can find definitions for these native metrics in [Viewing Statistics](#) in the PostgreSQL documentation.

Counter	Type	Unit	Metric
blks_hit	Cache	Blocks per second	db.Cache.blks_hit
buffers_alloc	Cache	Blocks per second	db.Cache.buffers_alloc
buffers_checkpoint	Checkpoint	Blocks per second	db.Checkpoint.buffers_checkpoint
checkpoint_sync_time	Checkpoint	Milliseconds per checkpoint	db.Checkpoint.checkpoint_sync_time

Counter	Type	Unit	Metric
checkpoint_write_time	Checkpoint	Milliseconds per checkpoint	db.Checkpoint.checkpoint_write_time
checkpoints_req	Checkpoint	Checkpoints per minute	db.Checkpoint.checkpoints_req
checkpoints_timed	Checkpoint	Checkpoints per minute	db.Checkpoint.checkpoints_timed
maxwritten_clean	Checkpoint	Bgwriter clean stops per minute	db.Checkpoint.maxwritten_clean
deadlocks	Concurrency	Deadlocks per minute	db.Concurrency.deadlocks
blk_read_time	IO	Milliseconds	db.IO.blk_read_time
blks_read	IO	Blocks per second	db.IO.blks_read
buffers_backend	IO	Blocks per second	db.IO.buffers_backend
buffers_backend_fsync	IO	Blocks per second	db.IO.buffers_backend_fsync
buffers_clean	IO	Blocks per second	db.IO.buffers_clean
tup_deleted	SQL	Tuples per second	db.SQL.tup_deleted
tup_fetched	SQL	Tuples per second	db.SQL.tup_fetched
tup_inserted	SQL	Tuples per second	db.SQL.tup_inserted
tup_returned	SQL	Tuples per second	db.SQL.tup_returned
tup_updated	SQL	Tuples per second	db.SQL.tup_updated
temp_bytes	Temp	Bytes per second	db.Temp.temp_bytes
temp_files	Temp	Files per minute	db.Temp.temp_files
active_transactions	Transactions	Transactions	db.Transactions.active_transactions
blocked_transactions	Transactions	Transactions	db.Transactions.blocked_transactions
max_used_xact_ids	Transactions	Transactions	db.Transactions.max_used_xact_ids
xact_commit	Transactions	Commits per second	db.Transactions.xact_commit
xact_rollback	Transactions	Rollbacks per second	db.Transactions.xact_rollback
numbackends	User	Connections	db.User.numbackends
archived_count	WAL	Files per minute	db.WAL.archived_count
archive_failed_count	WAL	Files per minute	db.WAL.archive_failed_count

Non-Native Counters for Amazon RDS for PostgreSQL

Non-native counter metrics are counters defined by Amazon RDS. A non-native metric can be a metric that you get with a specific query. A non-native metric also can be a derived metric, where two or more native counters are used in calculations for ratios, hit rates, or latencies.

Counter	Type	Metric	Description	Definition
checkpoint_sync_time	Checkpoint	db.Checkpoint	The total time spent in <code>checkpoint_start_time</code> that has been spent in the portion of checkpoint processing where files are synchronized to disk.	<code>checkpoint_sync_time</code> / (<code>checkpoints_timed</code> + <code>checkpoints_req</code>)
checkpoint_write_time	Checkpoint	db.Checkpoint	The total time spent in <code>checkpoint_end_time</code> that has been spent in the portion of checkpoint processing where files are written to disk.	<code>checkpoint_write_time</code> / (<code>checkpoints_timed</code> + <code>checkpoints_req</code>)
read_latency	IO	db.IO.read_latency	The time spent reading data file blocks by backends in this instance.	<code>blk_read_time</code> / <code>blks_read</code>

Logging Performance Insights Calls by Using AWS CloudTrail

Performance Insights is integrated with AWS CloudTrail, a service that provides a record of actions taken by a user, role, or an AWS service in Performance Insights. CloudTrail captures all API calls for Performance Insights as events. This capture includes calls from the Amazon RDS console and from code calls to the Performance Insights API operations.

If you create a trail, you can enable continuous delivery of CloudTrail events to an Amazon S3 bucket, including events for Performance Insights. If you don't configure a trail, you can still view the most recent events in the CloudTrail console in [Event history](#). Using the data collected by CloudTrail, you can determine certain information. This information includes the request that was made to Performance Insights, the IP address from which the request was made, who made the request, when it was made, and additional details.

To learn more about CloudTrail, see the [AWS CloudTrail User Guide](#).

Working with Performance Insights Information in CloudTrail

CloudTrail is enabled on your AWS account when you create the account. When activity occurs in Performance Insights, that activity is recorded in a CloudTrail event along with other AWS service events in the CloudTrail console in [Event history](#). You can view, search, and download recent events in your AWS account. For more information, see [Viewing Events with CloudTrail Event History](#) in [AWS CloudTrail User Guide](#).

For an ongoing record of events in your AWS account, including events for Performance Insights, create a trail. A *trail* enables CloudTrail to deliver log files to an Amazon S3 bucket. By default, when you create a trail in the console, the trail applies to all AWS Regions. The trail logs events from all AWS Regions in the AWS partition and delivers the log files to the Amazon S3 bucket that you specify. Additionally, you can configure other AWS services to further analyze and act upon the event data collected in CloudTrail logs. For more information, see the following topics in [AWS CloudTrail User Guide](#):

- [Overview for Creating a Trail](#)
- [CloudTrail Supported Services and Integrations](#)
- [Configuring Amazon SNS Notifications for CloudTrail](#)

- [Receiving CloudTrail Log Files from Multiple Regions](#) and [Receiving CloudTrail Log Files from Multiple Accounts](#)

All Performance Insights operations are logged by CloudTrail and are documented in the [Performance Insights API Reference](#). For example, calls to the `DescribeDimensionKeys` and `GetResourceMetrics` operations generate entries in the CloudTrail log files.

Every event or log entry contains information about who generated the request. The identity information helps you determine the following:

- Whether the request was made with root or IAM user credentials.
- Whether the request was made with temporary security credentials for a role or federated user.
- Whether the request was made by another AWS service.

For more information, see the [CloudTrail userIdentity Element](#).

Understanding Performance Insights Log File Entries

A *trail* is a configuration that enables delivery of events as log files to an Amazon S3 bucket that you specify. CloudTrail log files contain one or more log entries. An *event* represents a single request from any source. Each event includes information about the requested operation, the date and time of the operation, request parameters, and so on. CloudTrail log files aren't an ordered stack trace of the public API calls, so they don't appear in any specific order.

The following example shows a CloudTrail log entry that demonstrates the `GetResourceMetrics` operation.

```
{  
    "eventVersion": "1.05",  
    "userIdentity": {  
        "type": "IAMUser",  
        "principalId": "AKIAIOSFODNN7EXAMPLE",  
        "arn": "arn:aws:iam::123456789012:user/johndoe",  
        "accountId": "123456789012",  
        "accessKeyId": "AKIAI44QH8DHBEXAMPLE",  
        "userName": "johndoe"  
    },  
    "eventTime": "2019-12-18T19:28:46Z",  
    "eventSource": "pi.amazonaws.com",  
    "eventName": "GetResourceMetrics",  
    "awsRegion": "us-east-1",  
    "sourceIPAddress": "72.21.198.67",  
    "userAgent": "aws-cli/1.16.240 Python/3.7.4 Darwin/18.7.0 botocore/1.12.230",  
    "requestParameters": {  
        "identifier": "db-YTDU5J5V66X7CXSCVDFD2V3SZM",  
        "metricQueries": [  
            {  
                "metric": "os.cpuUtilization.user.avg"  
            },  
            {  
                "metric": "os.cpuUtilization.idle.avg"  
            }  
        ],  
        "startTime": "Dec 18, 2019 5:28:46 PM",  
        "periodInSeconds": 60,  
        "endTime": "Dec 18, 2019 7:28:46 PM",  
        "serviceType": "RDS"  
    },  
    "responseElements": null,  
}
```

```
"requestID": "9ffbe15c-96b5-4fe6-bed9-9fccff1a0525",  
"eventID": "08908de0-2431-4e2e-ba7b-f5424f908433",  
"eventType": "AwsApiCall",  
"recipientAccountId": "123456789012"  
}
```

Using Amazon RDS Recommendations

Amazon RDS provides automated recommendations for database resources, such as DB instances, Read Replicas, and DB parameter groups. These recommendations provide best practice guidance by analyzing DB instance configuration, usage, and performance data.

You can find examples of these recommendations in the following table.

Type	Description	Recommendation	Additional Information
Engine version outdated	Your DB instance is not running the latest minor engine version.	We recommend that you upgrade to the latest version because it contains the latest security fixes and other improvements.	Upgrading a DB Instance Engine Version (p. 159)
Pending maintenance available	You have pending maintenance available on your DB instance.	We recommend that you perform the pending maintenance available on your DB instance. Updates to the operating system most often occur for security issues and should be done as soon as possible.	Maintaining a DB Instance (p. 152)
Automated backups disabled	Your DB instance has automated backups disabled.	We recommend that you enable automated backups on your DB instance. Automated backups enable point-in-time recovery of your DB instance. You receive backup storage up to the storage size of your DB instance at no additional charge.	Working With Backups (p. 254)
Magnetic volumes in use	Your DB instance is using magnetic storage.	Magnetic storage is not recommended for most DB instances. We recommend switching to General Purpose (SSD) storage or provisioned IOPS storage.	Amazon RDS DB Instance Storage (p. 135)
EC2-Classic platform in use	Your DB instance is using the legacy EC2-Classic platform.	We recommend moving your DB instance to the EC2-VPC platform for better network access control. Amazon VPC provides a virtual network that is logically isolated from other virtual networks in the AWS Cloud.	Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform (p. 506)
Enhanced Monitoring disabled	Your DB instance doesn't have Enhanced Monitoring enabled.	We recommend enabling Enhanced Monitoring. Enhanced Monitoring provides real-time operating system metrics for monitoring and troubleshooting.	Enhanced Monitoring (p. 307)
Encryption disabled	Your DB instance doesn't have encryption enabled.	We recommend enabling encryption. You can encrypt your existing Amazon RDS DB instances by restoring from an encrypted snapshot.	Encrypting Amazon RDS Resources (p. 429)
Previous generation	Your DB instance is running on a	Previous-generation DB instance classes have been replaced by DB	Choosing the DB Instance Class (p. 107)

Type	Description	Recommendation	Additional Information
DB instance class in use	previous-generation DB instance class.	instance classes with better price, better performance, or both. We recommend running your DB instance on a later generation DB instance class.	
Huge pages not used for an Oracle DB instance	The <code>use_large_pages</code> parameter is not set to <code>ONLY</code> in the DB parameter group used by your DB instance.	For increased database scalability, we recommend setting <code>use_large_pages</code> to <code>ONLY</code> in the DB parameter group used by your DB instance.	Using Huge Pages with an Oracle DB Instance (p. 935)
Nondefault custom memory parameters	Your DB parameter group sets memory parameters that diverge too much from the default values.	Settings that diverge too much from the default values can cause poor performance and errors. We recommend setting custom memory parameters to their default values in the DB parameter group used by the DB instance.	Working with DB Parameter Groups (p. 211)
Change buffering enabled for a MySQL DB instance	Your DB parameter group has change buffering enabled.	Change buffering allows a MySQL DB instance to defer some writes necessary to maintain secondary indexes. This configuration can improve performance slightly, but it can create a large delay in crash recovery. During crash recovery, the secondary index must be brought up to date. So, the benefits of change buffering are outweighed by the potentially very long crash recovery events. We recommend disabling change buffering.	Best practices for configuring parameters for Amazon RDS for MySQL, part 1: Parameters related to performance on the AWS Database Blog
Query cache enabled for a MySQL DB instance	Your DB parameter group has query cache parameter enabled.	The query cache can cause the DB instance to appear to stall when changes require the cache to be purged. Most workloads don't benefit from a query cache. The query cache was removed from MySQL version 8.0. We recommend that you disable the query cache parameter.	Best practices for configuring parameters for Amazon RDS for MySQL, part 1: Parameters related to performance on the AWS Database Blog
Logging to table	Your DB parameter group sets logging output to <code>TABLE</code> .	Setting logging output to <code>TABLE</code> uses more storage than setting this parameter to <code>FILE</code> . To avoid reaching the storage limit, we recommend setting the logging output parameter to <code>FILE</code> .	MySQL Database Log Files (p. 405)

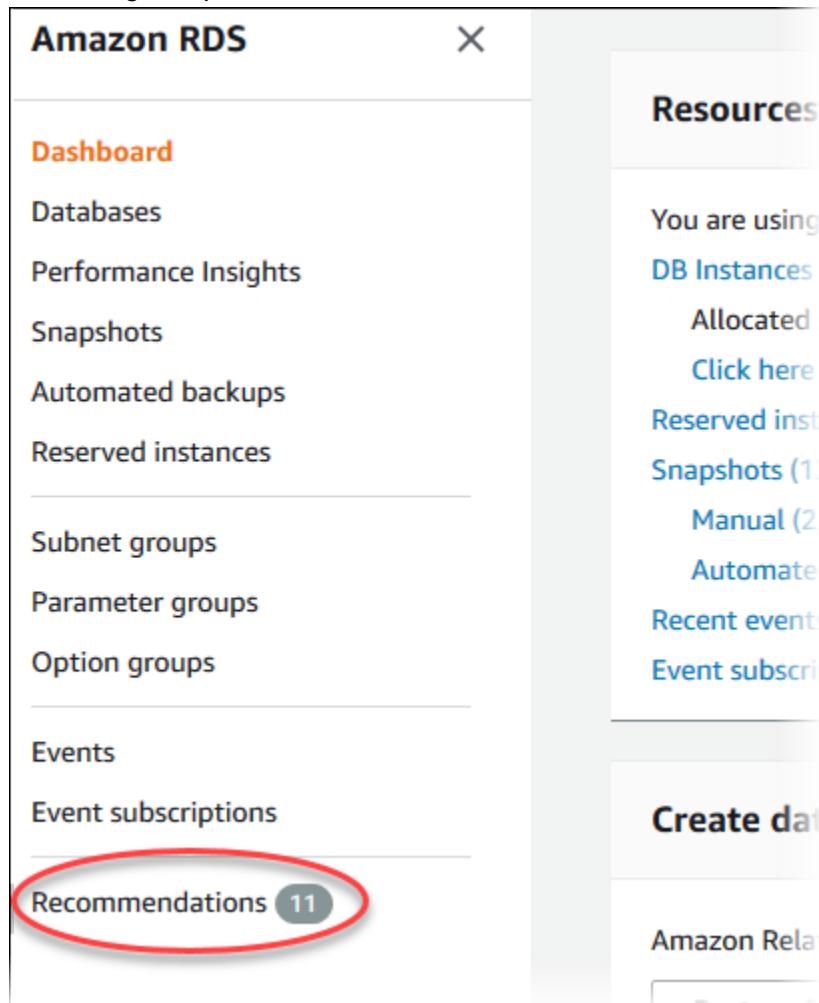
Amazon RDS generates recommendations for a resource when the resource is created or modified. Amazon RDS also periodically scans your resources and generates recommendations.

Responding to Amazon RDS Recommendations

You can find recommendations in the AWS Management Console. You can perform the recommended action immediately, schedule it for the next maintenance window, or dismiss it.

To respond to Amazon RDS recommendations

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Recommendations**.



The Recommendations page appears.

The screenshot shows a 'Recommendations' page with a header bar containing four tabs: Active (2), Dismissed (0), Scheduled (0), and Applied (1). The Active tab is selected. Below the tabs, there are two main sections: '▶ Engine version outdated for DB instances (1)' and '▶ Enhanced monitoring disabled (1)'. Each section contains a brief description and a 'Info' link.

3. On the **Recommendations** page, choose one of the following:

- **Active** – Shows the current recommendations that you can apply, dismiss, or schedule.
- **Dismissed** – Shows the recommendations that have been dismissed. When you choose **Dismissed**, you can apply these dismissed recommendations.
- **Scheduled** – Shows the recommendations that are scheduled but not yet applied. These recommendations will be applied in the next scheduled maintenance window.
- **Applied** – Shows the recommendations that are currently applied.

From any list of recommendations, you can open a section to view the recommendations in that section.

The screenshot shows the 'Engine version outdated for DB instances (1)' section from the previous screenshot. It includes a 'DB instances' table with a single row for 'database-1'. The table has columns for 'Resource' and 'Recommendation'. The 'Resource' column shows a checkbox and the name 'database-1'. The 'Recommendation' column contains a message about upgrading the MySQL version. There are also 'Dismiss', 'Schedule', and 'Apply now' buttons at the top right of the table. A 'Filter recommendations' search bar is located above the table.

To configure preferences for displaying recommendations in each section, choose the **Preferences** icon.

The screenshot shows the 'Recommendations' page in the Amazon RDS console. At the top, there are four tabs: 'Active (2)', 'Dismissed (0)', 'Scheduled (0)', and 'Applied (1)'. The 'Active (2)' tab is selected. Below the tabs, there are two sections: 'Engine version outdated for DB instances (1)' and 'Enhanced monitoring disabled (1)'. The first section has a note: 'DB instances that are not running the latest minor engine version.' The second section has a note: 'DB instances that don't have Enhanced Monitoring enabled.' In the 'DB instances' section, there is a table with a single row for 'database-1'. The table has columns for 'Resource' and 'Recommendation'. The 'Recommendation' column for 'database-1' contains the text: 'Your DB instance is running mysql version 5.6.34. We recommend that you upgrade to version 5.6.41 because it contains the latest security fixes and other improvements.' At the top right of the recommendations list, there are three buttons: 'Dismiss', 'Schedule', and 'Apply now'. The 'Apply now' button is circled in red.

From the **Preferences** window that appears, you can set display options. These options include the visible columns and the number of recommendations to display on the page.

4. Manage your active recommendations:

- a. Choose **Active** and open one or more sections to view the recommendations in them.
- b. Choose one or more recommendations and choose **Apply now** (to apply them immediately), **Schedule** (to apply them in next maintenance window), or **Dismiss**.

If the **Apply now** button appears for a recommendation but is unavailable (grayed out), the DB instance is not available. You can apply recommendations immediately only if the DB instance status is **available**. For example, you can't apply recommendations immediately to the DB instance if its status is **modifying**. In this case, wait for the DB instance to be available and then apply the recommendation.

If the **Active** button doesn't appear for a recommendation, you can't apply the recommendation using the **Recommendations** page. You can modify the DB instance to apply the recommendation manually.

For more information about modifying a DB instance, see [Modifying an Amazon RDS DB Instance \(p. 150\)](#).

Note

When you choose **Apply now**, a brief DB instance outage might result.

Using Amazon RDS Event Notification

Topics

- [Amazon RDS Event Categories and Event Messages \(p. 371\)](#)
- [Subscribing to Amazon RDS Event Notification \(p. 378\)](#)
- [Listing Your Amazon RDS Event Notification Subscriptions \(p. 380\)](#)
- [Modifying an Amazon RDS Event Notification Subscription \(p. 382\)](#)
- [Adding a Source Identifier to an Amazon RDS Event Notification Subscription \(p. 384\)](#)
- [Removing a Source Identifier from an Amazon RDS Event Notification Subscription \(p. 385\)](#)
- [Listing the Amazon RDS Event Notification Categories \(p. 386\)](#)
- [Deleting an Amazon RDS Event Notification Subscription \(p. 387\)](#)

Amazon RDS uses the Amazon Simple Notification Service (Amazon SNS) to provide notification when an Amazon RDS event occurs. These notifications can be in any notification form supported by Amazon SNS for an AWS Region, such as an email, a text message, or a call to an HTTP endpoint.

Amazon RDS groups these events into categories that you can subscribe to so that you can be notified when an event in that category occurs. You can subscribe to an event category for a DB instance, DB snapshot, DB parameter group, or DB security group. For example, if you subscribe to the Backup category for a given DB instance, you are notified whenever a backup-related event occurs that affects the DB instance. If you subscribe to a configuration change category for a DB security group, you are notified when the DB security group is changed. You also receive notification when an event notification subscription changes.

Event notifications are sent to the addresses that you provide when you create the subscription. You might want to create several different subscriptions, such as one subscription receiving all event notifications and another subscription that includes only critical events for your production DB instances. You can easily turn off notification without deleting a subscription by choosing **No** for **Enabled** in the Amazon RDS console or by setting the `Enabled` parameter to `false` using the AWS CLI or Amazon RDS API.

Important

Amazon RDS doesn't guarantee the order of events sent in an event stream. The event order is subject to change.

Note

Amazon RDS event notifications using SMS text messages are currently available for topic Amazon Resource Names (ARNs) and Amazon RDS resources in the US-East (Northern Virginia) Region. For more information on using text messages with SNS, see [Sending and Receiving SMS Notifications Using Amazon SNS](#) in the *Amazon Simple Notification Service Developer Guide*.

Amazon RDS uses the ARN of an Amazon SNS topic to identify each subscription. The Amazon RDS console creates the ARN for you when you create the subscription. If you use the CLI or API, you create the ARN by using the Amazon SNS console or the Amazon SNS API when you create a subscription.

Billing for Amazon RDS event notification is through the Amazon Simple Notification Service (Amazon SNS). Amazon SNS fees apply when using event notification. For more information on Amazon SNS billing, see [Amazon Simple Notification Service Pricing](#).

The process for subscribing to Amazon RDS event notification is as follows:

1. Create an Amazon RDS event notification subscription by using the Amazon RDS console, AWS CLI, or API.
2. Amazon RDS sends an approval email or SMS message to the addresses you submitted with your subscription. To confirm your subscription, choose the link in the notification you were sent.

3. When you have confirmed the subscription, the status of your subscription is updated in the Amazon RDS console's **My Event Subscriptions** section.
4. You then begin to receive event notifications.

Note

When Amazon SNS sends a notification to a subscribed HTTP or HTTPS endpoint, the POST message sent to the endpoint has a message body that contains a JSON document. For more information, see [Amazon SNS Message and JSON Formats](#) in the *Amazon Simple Notification Service Developer Guide*.

The following section lists all categories and events that you can be notified of. It also provides information about subscribing to and working with Amazon RDS event subscriptions.

Amazon RDS Event Categories and Event Messages

Amazon RDS generates a significant number of events in categories that you can subscribe to using the Amazon RDS Console, AWS CLI, or the API. Each category applies to a source type, which can be a DB instance, DB snapshot, DB security group, or DB parameter group.

The following table shows the event category and a list of events when a DB instance is the source type.

Category	Amazon RDS Event ID	Description
availability	RDS-EVENT-0006	The DB instance restarted.
availability	RDS-EVENT-0004	DB instance shutdown.
availability	RDS-EVENT-0022	An error has occurred while restarting MySQL or MariaDB.
backup	RDS-EVENT-0001	Backing up DB instance.
backup	RDS-EVENT-0002	Finished DB Instance backup.
configuration change	RDS-EVENT-0009	The DB instance has been added to a security group.
configuration change	RDS-EVENT-0024	The DB instance is being converted to a Multi-AZ DB instance.
configuration change	RDS-EVENT-0030	The DB instance is being converted to a Single-AZ DB instance.
configuration change	RDS-EVENT-0012	Applying modification to database instance class.
configuration change	RDS-EVENT-0018	The current storage settings for this DB instance are being changed.
configuration change	RDS-EVENT-0011	A parameter group for this DB instance has changed.
configuration change	RDS-EVENT-0092	A parameter group for this DB instance has finished updating.
configuration change	RDS-EVENT-0028	Automatic backups for this DB instance have been disabled.

Category	Amazon RDS Event ID	Description
configuration change	RDS-EVENT-0032	Automatic backups for this DB instance have been enabled.
configuration change	RDS-EVENT-0033	There are [count] users that match the master user name. Users not tied to a specific host have been reset.
configuration change	RDS-EVENT-0025	The DB instance has been converted to a Multi-AZ DB instance.
configuration change	RDS-EVENT-0029	The DB instance has been converted to a Single-AZ DB instance.
configuration change	RDS-EVENT-0014	The DB instance class for this DB instance has changed.
configuration change	RDS-EVENT-0017	The storage settings for this DB instance have changed.
configuration change	RDS-EVENT-0010	The DB instance has been removed from a security group.
configuration change	RDS-EVENT-0016	The master password for the DB instance has been reset.
configuration change	RDS-EVENT-0067	An attempt to reset the master password for the DB instance has failed.
configuration change	RDS-EVENT-0078	The Enhanced Monitoring configuration has been changed.
creation	RDS-EVENT-0005	DB instance created.
deletion	RDS-EVENT-0003	The DB instance has been deleted.
failover	RDS-EVENT-0034	Amazon RDS is not attempting a requested failover because a failover recently occurred on the DB instance.
failover	RDS-EVENT-0013	A Multi-AZ failover that resulted in the promotion of a standby instance has started.
failover	RDS-EVENT-0015	A Multi-AZ failover that resulted in the promotion of a standby instance is complete. It may take several minutes for the DNS to transfer to the new primary DB instance.
failover	RDS-EVENT-0065	The instance has recovered from a partial failover.
failover	RDS-EVENT-0049	A Multi-AZ failover has completed.
failover	RDS-EVENT-0050	A Multi-AZ activation has started after a successful instance recovery.
failover	RDS-EVENT-0051	A Multi-AZ activation is complete. Your database should be accessible now.

Category	Amazon RDS Event ID	Description
failure	RDS-EVENT-0031	The DB instance has failed due to an incompatible configuration or an underlying storage issue. Begin a point-in-time-restore for the DB instance.
failure	RDS-EVENT-0036	The DB instance is in an incompatible network. Some of the specified subnet IDs are invalid or do not exist.
failure	RDS-EVENT-0035	The DB instance has invalid parameters. For example, if the DB instance could not start because a memory-related parameter is set too high for this instance class, the customer action would be to modify the memory parameter and reboot the DB instance.
failure	RDS-EVENT-0058	Error while creating Statspack user account PERFSTAT. Please drop the account before adding the Statspack option.
failure	RDS-EVENT-0079	Enhanced Monitoring cannot be enabled without the enhanced monitoring IAM role. For information on creating the enhanced monitoring IAM role, see To create an IAM role for Amazon RDS Enhanced Monitoring (p. 308) .
failure	RDS-EVENT-0080	Enhanced Monitoring was disabled due to an error making the configuration change. It is likely that the enhanced monitoring IAM role is configured incorrectly. For information on creating the enhanced monitoring IAM role, see To create an IAM role for Amazon RDS Enhanced Monitoring (p. 308) .
failure	RDS-EVENT-0081	The IAM role that you use to access your Amazon S3 bucket for SQL Server native backup and restore is configured incorrectly. For more information, see Setting Up for Native Backup and Restore (p. 658) .
low storage	RDS-EVENT-0089	The DB instance has consumed more than 90% of its allocated storage. You can monitor the storage space for a DB instance using the Free Storage Space metric. For more information, see Viewing DB Instance Metrics (p. 304) .
low storage	RDS-EVENT-0007	The allocated storage for the DB instance has been exhausted. To resolve this issue, you should allocate additional storage for the DB instance. For more information, see the RDS FAQ . You can monitor the storage space for a DB instance using the Free Storage Space metric. For more information, see Viewing DB Instance Metrics (p. 304) .
maintenance	RDS-EVENT-0026	Offline maintenance of the DB instance is taking place. The DB instance is currently unavailable.
maintenance	RDS-EVENT-0027	Offline maintenance of the DB instance is complete. The DB instance is now available.
maintenance	RDS-EVENT-0047	Patching of the DB instance has completed.

Category	Amazon RDS Event ID	Description
maintenance	RDS-EVENT-0155	The DB instance has a DB engine minor version upgrade available.
notification	RDS-EVENT-0044	Operator-issued notification. For more information, see the event message.
notification	RDS-EVENT-0048	Patching of the DB instance has been delayed.
notification	RDS-EVENT-0054	The MySQL storage engine you are using is not InnoDB, which is the recommended MySQL storage engine for Amazon RDS. For information about MySQL storage engines, see Supported Storage Engines for MySQL on Amazon RDS .
notification	RDS-EVENT-0055	<p>The number of tables you have for your DB instance exceeds the recommended best practices for Amazon RDS. Please reduce the number of tables on your DB instance.</p> <p>For information about recommended best practices, see Amazon RDS Basic Operational Guidelines (p. 94).</p>
notification	RDS-EVENT-0056	<p>The number of databases you have for your DB instance exceeds the recommended best practices for Amazon RDS. Please reduce the number of databases on your DB instance.</p> <p>For information about recommended best practices, see Amazon RDS Basic Operational Guidelines (p. 94).</p>
notification	RDS-EVENT-0064	The TDE key has been rotated. For information about recommended best practices, see Amazon RDS Basic Operational Guidelines (p. 94) .
notification	RDS-EVENT-0084	You attempted to convert a DB instance to Multi-AZ, but it contains in-memory file groups that are not supported for Multi-AZ. For more information, see Multi-AZ Deployments for Microsoft SQL Server (p. 681) .
notification	RDS-EVENT-0087	The DB instance has been stopped.
notification	RDS-EVENT-0088	The DB instance has been started.
notification	RDS-EVENT-0154	The DB instance is being started due to it exceeding the maximum allowed time being stopped.
notification	RDS-EVENT-0157	<p>RDS can't modify the DB instance class because the target instance class can't support the number of databases that exist on the source DB instance. The error message appears as: "The instance has <i>N</i> databases, but after conversion it would only support <i>N</i>".</p> <p>For more information, see Limits for Microsoft SQL Server DB Instances (p. 600).</p>
notification	RDS-EVENT-0158	DB instance is in a state that can't be upgraded.

Category	Amazon RDS Event ID	Description
notification	RDS-EVENT-0159	DB snapshot export task failed.
notification	RDS-EVENT-0160	DB snapshot export task canceled.
notification	RDS-EVENT-0161	DB snapshot export task completed.
notification	RDS-EVENT-0162	DB cluster snapshot export task failed.
notification	RDS-EVENT-0163	DB cluster snapshot export task canceled.
notification	RDS-EVENT-0164	DB cluster snapshot export task completed.
read replica	RDS-EVENT-0045	An error has occurred in the read replication process. For more information, see the event message. For information on troubleshooting Read Replica errors, see Troubleshooting a MySQL Read Replica Problem (p. 856) .
read replica	RDS-EVENT-0046	The Read Replica has resumed replication. This message appears when you first create a Read Replica, or as a monitoring message confirming that replication is functioning properly. If this message follows an RDS-EVENT-0045 notification, then replication has resumed following an error or after replication was stopped.
read replica	RDS-EVENT-0057	Replication on the Read Replica was terminated.
read replica	RDS-EVENT-0062	Replication on the Read Replica was manually stopped.
read replica	RDS-EVENT-0063	Replication on the Read Replica was reset.
recovery	RDS-EVENT-0020	Recovery of the DB instance has started. Recovery time will vary with the amount of data to be recovered.
recovery	RDS-EVENT-0021	Recovery of the DB instance is complete.
recovery	RDS-EVENT-0023	A manual backup has been requested but Amazon RDS is currently in the process of creating a DB snapshot. Submit the request again after Amazon RDS has completed the DB snapshot.
recovery	RDS-EVENT-0052	Recovery of the Multi-AZ instance has started. Recovery time will vary with the amount of data to be recovered.
recovery	RDS-EVENT-0053	Recovery of the Multi-AZ instance is complete.
recovery	RDS-EVENT-0066	The SQL Server DB instance is re-establishing its mirror. Performance will be degraded until the mirror is reestablished. A database was found with non-FULL recovery model. The recovery model was changed back to FULL and mirroring recovery was started. (<dbname>: <recovery model found>[,...])"

Category	Amazon RDS Event ID	Description
restoration	RDS-EVENT-0008	The DB instance has been restored from a DB snapshot.
restoration	RDS-EVENT-0019	The DB instance has been restored from a point-in-time backup.

The following table shows the event category and a list of events when a DB parameter group is the source type.

Category	RDS Event ID	Description
configuration change	RDS-EVENT-0037	The parameter group was modified.

The following table shows the event category and a list of events when a DB security group is the source type.

Category	RDS Event ID	Description
configuration change	RDS-EVENT-0038	The security group has been modified.
failure	RDS-EVENT-0039	The Amazon EC2 security group owned by [user] does not exist; authorization for the security group has been revoked.

The following table shows the event category and a list of events when a DB snapshot is the source type.

Category	RDS Event ID	Description
creation	RDS-EVENT-0040	A manual DB snapshot is being created.
deletion	RDS-EVENT-0041	A DB snapshot has been deleted.
creation	RDS-EVENT-0042	A manual DB snapshot has been created.
restoration	RDS-EVENT-0043	A DB instance is being restored from a DB snapshot.
notification	RDS-EVENT-0059	Started the copy of the cross region DB snapshot [DB snapshot name] from source region [region name].
notification	RDS-EVENT-0060	Finished the copy of the cross region DB snapshot [DB snapshot name] from source region [region name] in [time] minutes.
notification	RDS-EVENT-0061	The copy of a cross region DB snapshot failed.
creation	RDS-EVENT-0090	An automated DB snapshot is being created.
creation	RDS-EVENT-0091	An automated DB snapshot has been created.

The following table shows the event category and a list of events when an Aurora DB cluster is the source type.

Category	RDS Event ID	Description
failover	RDS-EVENT-0069	A failover for the DB cluster has failed.
failover	RDS-EVENT-0070	A failover for the DB cluster has restarted.
failover	RDS-EVENT-0071	A failover for the DB cluster has finished.
failover	RDS-EVENT-0072	A failover for the DB cluster has begun within the same Availability Zone.
failover	RDS-EVENT-0073	A failover for the DB cluster has begun across Availability Zones.
failure	RDS-EVENT-0083	Aurora was unable to copy backup data from an Amazon S3 bucket. It is likely that the permissions for Aurora to access the Amazon S3 bucket are configured incorrectly. For more information, see Migrating Data from an External MySQL Database to an Amazon Aurora MySQL DB Cluster .
maintenance	RDS-EVENT-0156	The DB cluster has a DB engine minor version upgrade available.
notification	RDS-EVENT-0076	Migration to an Aurora DB cluster failed.
notification	RDS-EVENT-0077	An attempt to convert a table from the source database to InnoDB failed during the migration to an Aurora DB cluster.
notification	RDS-EVENT-0141	Scaling initiated for the Aurora Serverless DB cluster.
notification	RDS-EVENT-0142	Scaling completed for the Aurora Serverless DB cluster.
notification	RDS-EVENT-0143	Scaling failed for the Aurora Serverless DB cluster.
notification	RDS-EVENT-0144	Automatic pause initiated for the Aurora Serverless DB cluster.
notification	RDS-EVENT-0145	The Aurora Serverless DB cluster paused.
notification	RDS-EVENT-0146	Pause cancelled for the Aurora Serverless DB cluster.
notification	RDS-EVENT-0147	Resume initiated for the Aurora Serverless DB cluster.
notification	RDS-EVENT-0148	Resume completed for the Aurora Serverless DB cluster.
notification	RDS-EVENT-0149	Seamless scaling completed with the force option for the Aurora Serverless DB cluster. Connections might have been interrupted as required.
notification	RDS-EVENT-0150	The DB cluster stopped.
notification	RDS-EVENT-0151	The DB cluster started.
notification	RDS-EVENT-0152	The DB cluster stop failed.

Category	RDS Event ID	Description
notification	RDS-EVENT-0153	The DB cluster is being started due to it exceeding the maximum allowed time being stopped.

The following table shows the event category and a list of events when an Aurora DB cluster snapshot is the source type.

Category	RDS Event ID	Description
backup	RDS-EVENT-0074	Creation of a manual DB cluster snapshot has started.
backup	RDS-EVENT-0075	A manual DB cluster snapshot has been created.

Subscribing to Amazon RDS Event Notification

You can create an Amazon RDS event notification subscription so you can be notified when an event occurs for a given DB instance, DB snapshot, DB security group, or DB parameter group. The simplest way to create a subscription is with the RDS console. If you choose to create event notification subscriptions using the CLI or API, you must create an Amazon Simple Notification Service topic and subscribe to that topic with the Amazon SNS console or Amazon SNS API. You will also need to retain the Amazon Resource Name (ARN) of the topic because it is used when submitting CLI commands or API operations. For information on creating an SNS topic and subscribing to it, see [Getting Started with Amazon SNS](#) in the *Amazon Simple Notification Service Developer Guide*.

You can specify the type of source you want to be notified of and the Amazon RDS source that triggers the event. These are defined by the **SourceType** (type of source) and the **SourceIdentifier** (the Amazon RDS source generating the event). If you specify both the **SourceType** and **SourceIdentifier**, such as **SourceType** = `db-instance` and **SourceIdentifier** = `myDBInstance1`, you receive all the DB instance events for the specified source. If you specify a **SourceType** but don't specify a **SourceIdentifier**, you receive notice of the events for that source type for all your Amazon RDS sources. If you don't specify either the **SourceType** or the **SourceIdentifier**, you are notified of events generated from all Amazon RDS sources belonging to your customer account.

Note

Event notifications might take up to five minutes to be delivered.

Amazon RDS event notification is only available for unencrypted SNS topics. If you specify an encrypted SNS topic, event notifications aren't sent for the topic.

Console

To subscribe to RDS event notification

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In navigation pane, choose **Event subscriptions**.
3. In the **Event subscriptions** pane, choose **Create event subscription**.
4. In the **Create event subscription** dialog box, do the following:
 - a. For **Name**, enter a name for the event notification subscription.
 - b. For **Send notifications to**, choose an existing Amazon SNS ARN for an Amazon SNS topic, or choose **create topic** to enter the name of a topic and a list of recipients.
 - c. For **Source type**, choose a source type.

- d. Choose **Yes** to enable the subscription. If you want to create the subscription but to not have notifications sent yet, choose **No**.
- e. Depending on the source type you selected, choose the event categories and sources that you want to receive event notifications for.
- f. Choose **Create**.

The Amazon RDS console indicates that the subscription is being created.

Event subscriptions (2)				
	Name	Status	Source Type	Enabled
<input type="checkbox"/>	Configchangerdpgres	active	Instances	Yes
<input type="checkbox"/>	Test	creating	Instances	Yes

AWS CLI

To subscribe to RDS event notification, use the AWS CLI `create-event-subscription` command. Include the following required parameters:

- `--subscription-name`
- `--sns-topic-arn`

Example

For Linux, OS X, or Unix:

```
aws rds create-event-subscription \
--subscription-name myeventsubscription \
--sns-topic-arn arn:aws:sns:us-east-1:802#####:myawsuser-RDS \
--enabled
```

For Windows:

```
aws rds create-event-subscription ^
--subscription-name myeventsubscription ^
--sns-topic-arn arn:aws:sns:us-east-1:802#####:myawsuser-RDS ^
--enabled
```

API

To subscribe to Amazon RDS event notification, call the Amazon RDS API function `CreateEventSubscription`. Include the following required parameters:

- `SubscriptionName`
- `SnsTopicArn`

Listing Your Amazon RDS Event Notification Subscriptions

You can list your current Amazon RDS event notification subscriptions.

Console

To list your current Amazon RDS event notification subscriptions

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Event subscriptions**. The **Event subscriptions** pane shows all your event notification subscriptions.

Event subscriptions (2)		
	Name	Status
<input type="checkbox"/>	Configchangerdpgres	<input checked="" type="checkbox"/> active
<input type="checkbox"/>	Postgresnotification	<input checked="" type="checkbox"/> active

AWS CLI

To list your current Amazon RDS event notification subscriptions, use the AWS CLI [describe-event-subscriptions](#) command.

Example

The following example describes all event subscriptions.

```
aws rds describe-event-subscriptions
```

The following example describes the `myfirsteventsubscription`.

```
aws rds describe-event-subscriptions --subscription-name myfirsteventsubscription
```

API

To list your current Amazon RDS event notification subscriptions, call the Amazon RDS API [DescribeEventSubscriptions](#) action.

Example

The following code example lists up to 100 event subscriptions.

```
https://rds.us-east-1.amazonaws.com/?Action=DescribeEventSubscriptions
```

```
&MaxRecords=100
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140428/us-east-1/rds/aws4_request
&X-Amz-Date=20140428T161907Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=4208679fe967783a1a149c826199080a066085d5a88227a80c6c0cadb3e8c0d4
```

The following example describes the `myfirsteventsSubscription`.

```
https://rds.us-east-1.amazonaws.com/
?Action=DescribeEventSubscriptions
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&SubscriptionName=myfirsteventsSubscription
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140428/us-east-1/rds/aws4_request
&X-Amz-Date=20140428T161907Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=4208679fe967783a1a149c826199080a066085d5a88227a80c6c0cadb3e8c0d4
```

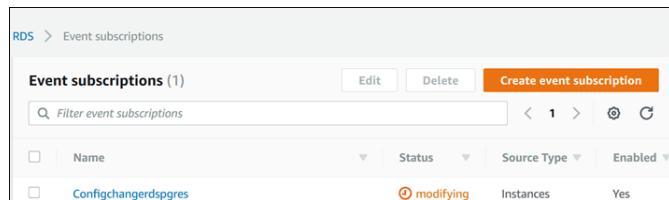
Modifying an Amazon RDS Event Notification Subscription

After you have created a subscription, you can change the subscription name, source identifier, categories, or topic ARN.

Console

To modify an Amazon RDS event notification subscription

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Event subscriptions**.
3. In the **Event subscriptions** pane, choose the subscription that you want to modify and choose **Edit**.
4. Make your changes to the subscription in either the **Target** or **Source** section.
5. Choose **Edit**. The Amazon RDS console indicates that the subscription is being modified.



AWS CLI

To modify an Amazon RDS event notification subscription, use the AWS CLI `modify-event-subscription` command. Include the following required parameter:

- `--subscription-name`

Example

The following code enables `myeventsSubscription`.

For Linux, OS X, or Unix:

```
aws rds modify-event-subscription \
--subscription-name myeventsSubscription \
--enabled
```

For Windows:

```
aws rds modify-event-subscription ^
--subscription-name myeventsSubscription ^
--enabled
```

API

To modify an Amazon RDS event, call the Amazon RDS API operation `ModifyEventSubscription`. Include the following required parameter:

- `SubscriptionName`

Adding a Source Identifier to an Amazon RDS Event Notification Subscription

You can add a source identifier (the Amazon RDS source generating the event) to an existing subscription.

Console

You can easily add or remove source identifiers using the Amazon RDS console by selecting or deselecting them when modifying a subscription. For more information, see [Modifying an Amazon RDS Event Notification Subscription \(p. 382\)](#).

AWS CLI

To add a source identifier to an Amazon RDS event notification subscription, use the AWS CLI [add-source-identifier-to-subscription](#) command. Include the following required parameters:

- `--subscription-name`
- `--source-identifier`

Example

The following example adds the source identifier `mysqldb` to the `myrdseventsSubscription` subscription.

For Linux, OS X, or Unix:

```
aws rds add-source-identifier-to-subscription \
  --subscription-name myrdseventsSubscription \
  --source-identifier mysqldb
```

For Windows:

```
aws rds add-source-identifier-to-subscription ^
  --subscription-name myrdseventsSubscription ^
  --source-identifier mysqldb
```

API

To add a source identifier to an Amazon RDS event notification subscription, call the Amazon RDS API [AddSourceIdentifierToSubscription](#). Include the following required parameters:

- `SubscriptionName`
- `SourceIdentifier`

Removing a Source Identifier from an Amazon RDS Event Notification Subscription

You can remove a source identifier (the Amazon RDS source generating the event) from a subscription if you no longer want to be notified of events for that source.

Console

You can easily add or remove source identifiers using the Amazon RDS console by selecting or deselecting them when modifying a subscription. For more information, see [Modifying an Amazon RDS Event Notification Subscription \(p. 382\)](#).

AWS CLI

To remove a source identifier from an Amazon RDS event notification subscription, use the AWS CLI `remove-source-identifier-from-subscription` command. Include the following required parameters:

- `--subscription-name`
- `--source-identifier`

Example

The following example removes the source identifier `mysqldb` from the `myrdseventsSubscription` subscription.

For Linux, OS X, or Unix:

```
aws rds remove-source-identifier-from-subscription \
--subscription-name myrdseventsSubscription \
--source-identifier mysqldb
```

For Windows:

```
aws rds remove-source-identifier-from-subscription ^
--subscription-name myrdseventsSubscription ^
--source-identifier mysqldb
```

API

To remove a source identifier from an Amazon RDS event notification subscription, use the Amazon RDS API `RemoveSourceIdentifierFromSubscription` command. Include the following required parameters:

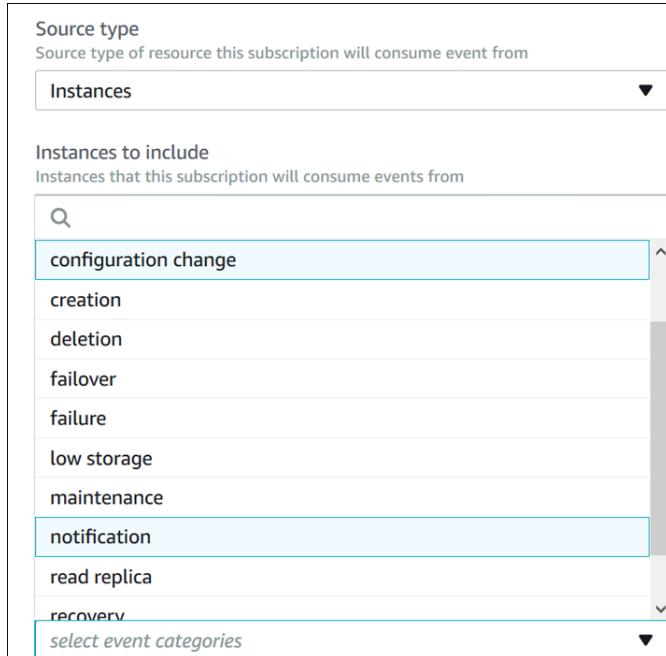
- `SubscriptionName`
- `SourceIdentifier`

Listing the Amazon RDS Event Notification Categories

All events for a resource type are grouped into categories. To view the list of categories available, use the following procedures.

Console

When you create or modify an event notification subscription, the event categories are displayed in the Amazon RDS console. For more information, see [Modifying an Amazon RDS Event Notification Subscription \(p. 382\)](#).



AWS CLI

To list the Amazon RDS event notification categories, use the AWS CLI `describe-event-categories` command. This command has no required parameters.

Example

```
aws rds describe-event-categories
```

API

To list the Amazon RDS event notification categories, use the Amazon RDS API `DescribeEventCategories` command. This command has no required parameters.

Deleting an Amazon RDS Event Notification Subscription

You can delete a subscription when you no longer need it. All subscribers to the topic will no longer receive event notifications specified by the subscription.

Console

To delete an Amazon RDS event notification subscription

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **DB Event Subscriptions**.
3. In the **My DB Event Subscriptions** pane, choose the subscription that you want to delete.
4. Choose **Delete**.
5. The Amazon RDS console indicates that the subscription is being deleted.

The screenshot shows the 'Event subscriptions' section of the AWS RDS console. It displays two entries:

Name	Status
Configchangerdpgres	active
Postgresnotification	active

AWS CLI

To delete an Amazon RDS event notification subscription, use the AWS CLI `delete-event-subscription` command. Include the following required parameter:

- `--subscription-name`

Example

The following example deletes the subscription `myrdssubscription`.

```
aws rds delete-event-subscription --subscription-name myrdssubscription
```

API

To delete an Amazon RDS event notification subscription, use the RDS API `DeleteEventSubscription` command. Include the following required parameter:

- `SubscriptionName`

Viewing Amazon RDS Events

Amazon RDS keeps a record of events that relate to your DB instances, DB snapshots, DB security groups, and DB parameter groups. This information includes the date and time of the event, the source name and source type of the event, and a message associated with the event.

You can retrieve events for your RDS resources through the AWS Management Console, which shows events from the past 24 hours. You can also retrieve events for your RDS resources by using the [describe-events](#) AWS CLI command, or the [DescribeEvents](#) RDS API operation. If you use the AWS CLI or the RDS API to view events, you can retrieve events for up to the past 14 days.

Console

To view all Amazon RDS instance events for the past 24 hours

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Events**. The available events appear in a list.
3. Use the **Filter** list to filter the events by type, and use the text box to the right of the **Filter** list to further filter your results. For example, the following screenshot shows a list of events filtered by the DB instance event type and containing the characters **1318**.

Identifier	Type
feb1318	DB cluster snapshots
feb1318	DB cluster snapshots

AWS CLI

To view all Amazon RDS instance events for the past 7 days

You can view all Amazon RDS instance events for the past 7 days by calling the [describe-events](#) AWS CLI command and setting the `--duration` parameter to 10080.

```
aws rds describe-events --duration 10080
```

API

To view all Amazon RDS instance events for the past 14 days

You can view all Amazon RDS instance events for the past 14 days by calling the [DescribeEvents](#) RDS API operation and setting the `Duration` parameter to 20160.

```
https://rds.us-west-2.amazonaws.com/  
?Action=DescribeEvents
```

```
&Duration=20160
&MaxRecords=100
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140421/us-west-2/rds/aws4_request
&X-Amz-Date=20140421T194733Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=8e313cabcd9766c56a2886b5b298fd944e0b7cfa248953c82705fdd0374f27
```

Amazon RDS Database Log Files

You can view, download, and watch database logs using the Amazon RDS console, the AWS Command Line Interface (AWS CLI), or the Amazon RDS API. Viewing, downloading, or watching transaction logs is not supported.

For engine-specific information, see the following:

- [MariaDB Database Log Files \(p. 394\)](#)
- [Microsoft SQL Server Database Log Files \(p. 402\)](#)
- [MySQL Database Log Files \(p. 405\)](#)
- [Oracle Database Log Files \(p. 413\)](#)
- [PostgreSQL Database Log Files \(p. 419\)](#)

Viewing and Listing Database Log Files

You can view database log files for your DB engine by using the Amazon RDS console. You can list what log files are available for download or monitoring by using the AWS CLI or Amazon RDS API.

Note

If you can't view the list of log files for an existing Oracle DB instance, reboot the instance to view the list.

Console

To view a database log file

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the name of the DB instance that has the log file that you want to view.
4. Choose the **Logs & events** tab.
5. Scroll down to the **Logs** section.
6. In the **Logs** section, choose the log that you want to view, and then choose **View**.

AWS CLI

To list the available database log files for a DB instance, use the AWS CLI `describe-db-log-files` command.

The following example returns a list of log files for a DB instance named `my-db-instance`.

Example

```
aws rds describe-db-log-files --db-instance-identifier my-db-instance
```

RDS API

To list the available database log files for a DB instance, use the Amazon RDS API `DescribeDBLogFiles` action.

Downloading a Database Log File

You can use the Amazon RDS console, AWS CLI or API to download a database log file.

Console

To download a database log file

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the name of the DB instance that has the log file that you want to view.
4. Choose the **Logs & events** tab.
5. Scroll down to the **Logs** section.
6. In the **Logs** section, choose the button next to the log that you want to download, and then choose **Download**.
7. Open the context (right-click) menu for the link provided, and then choose **Save Link As**. Enter the location where you want the log file to be saved, and then choose **Save**.



AWS CLI

To download a database log file, use the AWS CLI command `download-db-log-file-portion`. By default, this command downloads only the latest portion of a log file. However, you can download an entire file by specifying the parameter `--starting-token 0`.

The following example shows how to download the entire contents of a log file called `log/ERROR.4` and store it in a local file called `errorlog.txt`.

Example

For Linux, OS X, or Unix:

```
aws rds download-db-log-file-portion \
    --db-instance-identifier myexampledb \
    --starting-token 0 --output text \
    --log-file-name log/ERROR.4 > errorlog.txt
```

For Windows:

```
aws rds download-db-log-file-portion ^
    --db-instance-identifier myexampledb ^
    --starting-token 0 --output text ^
    --log-file-name log/ERROR.4 > errorlog.txt
```

RDS API

To download a database log file, use the Amazon RDS API `DownloadDBLogFilePortion` action.

Watching a Database Log File

You can monitor the contents of a log file by using the Amazon RDS console.

Console

To watch a database log file

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the name of the DB instance that has the log file that you want to view.
4. Choose the **Logs & events** tab.
5. In the **Logs** section, choose a log file, and then choose **Watch**.

Publishing Database Logs to Amazon CloudWatch Logs

In addition to viewing and downloading DB instance logs, you can publish logs to Amazon CloudWatch Logs. CloudWatch Logs lets you perform real-time analysis of the log data, store the data in highly durable storage, and manage the data with the CloudWatch Logs Agent. AWS retains log data published to CloudWatch Logs for an indefinite time period unless you specify a retention period. For more information, see [Change Log Data Retention in CloudWatch Logs](#).

For engine-specific information, see the following:

- the section called “Publishing MariaDB Logs to Amazon CloudWatch Logs” (p. 395)
- the section called “Publishing MySQL Logs to CloudWatch Logs” (p. 406)
- the section called “Publishing Oracle Logs to Amazon CloudWatch Logs” (p. 415)
- the section called “Publishing PostgreSQL Logs to CloudWatch Logs” (p. 420)
- the section called “Publishing SQL Server Logs to Amazon CloudWatch Logs” (p. 402)

Reading Log File Contents Using REST

Amazon RDS provides a REST endpoint that allows access to DB instance log files. This is useful if you need to write an application to stream Amazon RDS log file contents.

The syntax is:

```
GET /v13/downloadCompleteLogFile/DBInstanceIdentifier/LogFileName HTTP/1.1
Content-type: application/json
host: rds.region.amazonaws.com
```

The following parameters are required:

- *DBInstanceIdentifier*—the name of the DB instance that contains the log file you want to download.
- *LogFileName*—the name of the log file to be downloaded.

The response contains the contents of the requested log file, as a stream.

The following example downloads the log file named *log/ERROR.6* for the DB instance named *sample-sql* in the *us-west-2* region.

```
GET /v13/downloadCompleteLogFile/sample-sql/log/ERROR.6 HTTP/1.1
host: rds.us-west-2.amazonaws.com
X-Amz-Security-Token: AQoDYXdzEIH//////////wEa0AIXLhngC5zp9CyB1R6abwKrXHVR5efnAVN3XvR7IwqKYalFSn6UyJuEFTft9nObglx4QJ+GXV9cpACkETq=
X-Amz-Date: 20140903T233749Z
X-Amz-Algorithm: AWS4-HMAC-SHA256
X-Amz-Credential: AKIADQKE4SARGYLE/20140903/us-west-2/rds/aws4_request
X-Amz-SignedHeaders: host
X-Amz-Content-SHA256: e3b0c44298fc1c229afbf4c8996fb92427ae41e4649b934de495991b7852b855
X-Amz-Expires: 86400
X-Amz-Signature: 353a4f14b3f250142d9afc34f9f9948154d46ce7d4ec091d0cdabbcf8b40c558
```

If you specify a nonexistent DB instance, the response consists of the following error:

- **DBInstanceNotFound**—*DB Instance Identifier* does not refer to an existing DB instance. (HTTP status code: 404)

MariaDB Database Log Files

You can monitor the MariaDB error log, slow query log, and the general log. The MariaDB error log is generated by default; you can generate the slow query and general logs by setting parameters in your DB parameter group. Amazon RDS rotates all of the MariaDB log files; the intervals for each type are given following.

You can monitor the MariaDB logs directly through the Amazon RDS console, Amazon RDS API, Amazon RDS CLI, or AWS SDKs. You can also access MariaDB logs by directing the logs to a database table in the main database and querying that table. You can use the `mysqlbinlog` utility to download a binary log.

For more information about viewing, downloading, and watching file-based database logs, see [Amazon RDS Database Log Files \(p. 390\)](#).

Accessing MariaDB Error Logs

The MariaDB error log is written to the `<host-name>.err` file. You can view this file by using the Amazon RDS console or by retrieving the log using the Amazon RDS API, Amazon RDS CLI, or AWS SDKs. The `<host-name>.err` file is flushed every 5 minutes, and its contents are appended to `mysql-error-running.log`. The `mysql-error-running.log` file is then rotated every hour and the hourly files generated during the last 24 hours are retained. Each log file has the hour it was generated (in UTC) appended to its name. The log files also have a timestamp that helps you determine when the log entries were written.

MariaDB writes to the error log only on startup, shutdown, and when it encounters errors. A DB instance can go hours or days without new entries being written to the error log. If you see no recent entries, it's because the server did not encounter an error that resulted in a log entry.

Accessing the MariaDB Slow Query and General Logs

The MariaDB slow query log and the general log can be written to a file or a database table by setting parameters in your DB parameter group. For information about creating and modifying a DB parameter group, see [Working with DB Parameter Groups \(p. 211\)](#). You must set these parameters before you can view the slow query log or general log in the Amazon RDS console or by using the Amazon RDS API, AWS CLI, or AWS SDKs.

You can control MariaDB logging by using the parameters in this list:

- `slow_query_log`: To create the slow query log, set to 1. The default is 0.
- `general_log`: To create the general log, set to 1. The default is 0.
- `long_query_time`: To prevent fast-running queries from being logged in the slow query log, specify a value for the shortest query execution time to be logged, in seconds. The default is 10 seconds; the minimum is 0. If `log_output = FILE`, you can specify a floating point value that goes to microsecond resolution. If `log_output = TABLE`, you must specify an integer value with second resolution. Only queries whose execution time exceeds the `long_query_time` value are logged. For example, setting `long_query_time` to 0.1 prevents any query that runs for less than 100 milliseconds from being logged.
- `log_queries_not_using_indexes`: To log all queries that do not use an index to the slow query log, set this parameter to 1. The default is 0. Queries that do not use an index are logged even if their execution time is less than the value of the `long_query_time` parameter.
- `log_output option`: You can specify one of the following options for the `log_output` parameter:
 - **TABLE** (default)– Write general queries to the `mysql.general_log` table, and slow queries to the `mysql.slow_log` table.
 - **FILE**– Write both general and slow query logs to the file system. Log files are rotated hourly.

- **NONE**—Disable logging.

When logging is enabled, Amazon RDS rotates table logs or deletes log files at regular intervals. This measure is a precaution to reduce the possibility of a large log file either blocking database use or affecting performance. **FILE** and **TABLE** logging approach rotation and deletion as follows:

- When **FILE** logging is enabled, log files are examined every hour and log files older than 24 hours are deleted. In some cases, the remaining combined log file size after the deletion might exceed the threshold of 2 percent of a DB instance's allocated space. In these cases, the largest log files are deleted until the log file size no longer exceeds the threshold.
- When **TABLE** logging is enabled, in some cases log tables are rotated every 24 hours. This rotation occurs if the space used by the table logs is more than 20 percent of the allocated storage space or the size of all logs combined is greater than 10 GB. If the amount of space used for a DB instance is greater than 90 percent of the DB instance's allocated storage space, then the thresholds for log rotation are reduced. Log tables are then rotated if the space used by the table logs is more than 10 percent of the allocated storage space or the size of all logs combined is greater than 5 GB.

When log tables are rotated, the current log table is copied to a backup log table and the entries in the current log table are removed. If the backup log table already exists, then it is deleted before the current log table is copied to the backup. You can query the backup log table if needed. The backup log table for the `mysql.general_log` table is named `mysql.general_log_backup`. The backup log table for the `mysql.slow_log` table is named `mysql.slow_log_backup`.

You can rotate the `mysql.general_log` table by calling the `mysql.rds_rotate_general_log` procedure. You can rotate the `mysql.slow_log` table by calling the `mysql.rds_rotate_slow_log` procedure.

Table logs are rotated during a database version upgrade.

Amazon RDS records both **TABLE** and **FILE** log rotation in an Amazon RDS event and sends you a notification.

To work with the logs from the Amazon RDS console, Amazon RDS API, Amazon RDS CLI, or AWS SDKs, set the `log_output` parameter to **FILE**. Like the MariaDB error log, these log files are rotated hourly. The log files that were generated during the previous 24 hours are retained.

For more information about the slow query and general logs, go to the following topics in the MariaDB documentation:

- [Slow Query Log](#)
- [General Query Log](#)

Publishing MariaDB Logs to Amazon CloudWatch Logs

You can configure your Amazon RDS MariaDB DB instance to publish log data to a log group in Amazon CloudWatch Logs. With CloudWatch Logs, you can perform real-time analysis of the log data, and use CloudWatch to create alarms and view metrics. You can use CloudWatch Logs to store your log records in highly durable storage.

Amazon RDS publishes each MariaDB database log as a separate database stream in the log group. For example, if you configure the export function to include the slow query log, slow query data is stored in a slow query log stream in the `/aws/rds/instance/my_instance/slowquery` log group.

The error log is enabled by default. The following table summarizes the requirements for the other MariaDB logs.

Log	Requirement
Audit log	The DB instance must use a custom option group with the MARIADB_AUDIT_PLUGIN option.
General log	The DB instance must use a custom parameter group with the parameter setting general_log = 1 to enable the general log.
Slow query log	The DB instance must use a custom parameter group with the parameter setting slow_query_log = 1 to enable the slow query log.
Log output	The DB instance must use a custom parameter group with the parameter setting log_output = FILE to write logs to the file system and publish them to CloudWatch Logs.

Console

To publish MariaDB logs to CloudWatch Logs from the console

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to modify.
3. Choose **Modify**.
4. In the **Log exports** section, choose the logs that you want to start publishing to CloudWatch Logs.
5. Choose **Continue**, and then choose **Modify DB Instance** on the summary page.

AWS CLI

You can publish a MariaDB logs with the AWS CLI. You can call the `modify-db-instance` command with the following parameters:

- `--db-instance-identifier`
- `--cloudwatch-logs-export-configuration`

Note

A change to the `--cloudwatch-logs-export-configuration` option is always applied to the DB instance immediately. Therefore, the `--apply-immediately` and `--no-apply-immediately` options have no effect.

You can also publish MariaDB logs by calling the following AWS CLI commands:

- `create-db-instance`
- `restore-db-instance-from-db-snapshot`
- `restore-db-instance-from-s3`
- `restore-db-instance-to-point-in-time`

Run one of these AWS CLI commands with the following options:

- `--db-instance-identifier`

- --enable-cloudwatch-logs-exports
- --db-instance-class
- --engine

Other options might be required depending on the AWS CLI command you run.

Example

The following example modifies an existing MariaDB DB instance to publish log files to CloudWatch Logs. The --cloudwatch-logs-export-configuration value is a JSON object. The key for this object is EnableLogTypes, and its value is an array of strings with any combination of audit, error, general, and slowquery.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
  --db-instance-identifier mydbinstance \
  --cloudwatch-logs-export-configuration '{"EnableLogTypes": \
  ["audit","error","general","slowquery"]}'
```

For Windows:

```
aws rds modify-db-instance ^
  --db-instance-identifier mydbinstance ^
  --cloudwatch-logs-export-configuration '{"EnableLogTypes": \
  ["audit","error","general","slowquery"]}'
```

Example

The following command creates a MariaDB DB instance and publishes log files to CloudWatch Logs. The --enable-cloudwatch-logs-exports value is a JSON array of strings. The strings can be any combination of audit, error, general, and slowquery.

For Linux, OS X, or Unix:

```
aws rds create-db-instance \
  --db-instance-identifier mydbinstance \
  --enable-cloudwatch-logs-exports '[{"audit","error","general","slowquery"}]' \
  --db-instance-class db.m4.large \
  --engine mariadb
```

For Windows:

```
aws rds create-db-instance ^
  --db-instance-identifier mydbinstance ^
  --enable-cloudwatch-logs-exports '[{"audit","error","general","slowquery"}]' ^
  --db-instance-class db.m4.large ^
  --engine mariadb
```

RDS API

You can publish MariaDB logs with the RDS API. You can call the [ModifyDBInstance](#) operation with the following parameters:

- [DBInstanceIdentifier](#)
- [CloudwatchLogsExportConfiguration](#)

Note

A change to the `CloudwatchLogsExportConfiguration` parameter is always applied to the DB instance immediately. Therefore, the `ApplyImmediately` parameter has no effect.

You can also publish MariaDB logs by calling the following RDS API operations:

- [CreateDBInstance](#)
- [RestoreDBInstanceFromDBSnapshot](#)
- [RestoreDBInstanceFromS3](#)
- [RestoreDBInstanceToPointInTime](#)

Run one of these RDS API operations with the following parameters:

- `DBInstanceIdentifier`
- `EnableCloudwatchLogsExports`
- `Engine`
- `DBInstanceClass`

Other parameters might be required depending on the AWS CLI command you run.

Log File Size

The MariaDB slow query log, error log, and the general log file sizes are constrained to no more than 2 percent of the allocated storage space for a DB instance. To maintain this threshold, logs are automatically rotated every hour and log files older than 24 hours are removed. If the combined log file size exceeds the threshold after removing old log files, then the largest log files are deleted until the log file size no longer exceeds the threshold.

Managing Table-Based MariaDB Logs

You can direct the general and slow query logs to tables on the DB instance by creating a DB parameter group and setting the `log_output` server parameter to `TABLE`. General queries are then logged to the `mysql.general_log` table, and slow queries are logged to the `mysql.slow_log` table. You can query the tables to access the log information. Enabling this logging increases the amount of data written to the database, which can degrade performance.

Both the general log and the slow query logs are disabled by default. In order to enable logging to tables, you must also set the `general_log` and `slow_query_log` server parameters to 1.

Log tables keep growing until the respective logging activities are turned off by resetting the appropriate parameter to 0. A large amount of data often accumulates over time, which can use up a considerable percentage of your allocated storage space. Amazon RDS does not allow you to truncate the log tables, but you can move their contents. Rotating a table saves its contents to a backup table and then creates a new empty log table. You can manually rotate the log tables with the following command line procedures, where the command prompt is indicated by `PROMPT>`:

```
PROMPT> CALL mysql.rds_rotate_slow_log;
PROMPT> CALL mysql.rds_rotate_general_log;
```

To completely remove the old data and reclaim the disk space, call the appropriate procedure twice in succession.

Binary Logging Format

MariaDB on Amazon RDS supports the *row-based*, *statement-based*, and *mixed* binary logging formats. The default binary logging format is *mixed*. For details on the different MariaDB binary log formats, see [Binary Log Formats](#) in the MariaDB documentation.

If you plan to use replication, the binary logging format is important because it determines the record of data changes that is recorded in the source and sent to the replication targets. For information about the advantages and disadvantages of different binary logging formats for replication, see [Advantages and Disadvantages of Statement-Based and Row-Based Replication](#) in the MySQL documentation.

Important

Setting the binary logging format to row-based can result in very large binary log files. Large binary log files reduce the amount of storage available for a DB instance and can increase the amount of time to perform a restore operation of a DB instance.

Statement-based replication can cause inconsistencies between the source DB instance and a Read Replica. For more information, see [Unsafe Statements for Statement-based Replication](#) in the MariaDB documentation.

To set the MariaDB binary logging format

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Parameter groups**.
3. Choose the parameter group that is used by the DB instance that you want to modify.

You can't modify a default parameter group. If the DB instance is using a default parameter group, create a new parameter group and associate it with the DB instance.

For more information on DB parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#).

4. For **Parameter group actions**, choose **Edit**.
5. Set the `binlog_format` parameter to the binary logging format of your choice (**ROW**, **STATEMENT**, or **MIXED**).
6. Choose **Save changes** to save the updates to the DB parameter group.

Accessing MariaDB Binary Logs

You can use the `mysqlbinlog` utility to download binary logs in text format from MariaDB DB instances. The binary log is downloaded to your local computer. For more information about using the `mysqlbinlog` utility, go to [Using mysqlbinlog](#) in the MariaDB documentation.

To run the `mysqlbinlog` utility against an Amazon RDS instance, use the following options:

- Specify the `--read-from-remote-server` option.
- `--host`: Specify the DNS name from the endpoint of the instance.
- `--port`: Specify the port used by the instance.
- `--user`: Specify a MariaDB user that has been granted the replication slave permission.
- `--password`: Specify the password for the user, or omit a password value so the utility prompts you for a password.
- `--result-file`: Specify the local file that receives the output.
- Specify the names of one or more binary log files. To get a list of the available logs, use the SQL command `SHOW BINARY LOGS`.

For more information about mysqlbinlog options, go to [mysqlbinlog Options](#) in the MariaDB documentation.

The following is an example:

For Linux, OS X, or Unix:

```
mysqlbinlog \
--read-from-remote-server \
--host=mariadbinstance1.1234abcd.region.rds.amazonaws.com \
--port=3306 \
--user ReplUser \
--password <password> \
--result-file=/tmp/binlog.txt
```

For Windows:

```
mysqlbinlog ^
--read-from-remote-server ^
--host=mariadbinstance1.1234abcd.region.rds.amazonaws.com ^
--port=3306 ^
--user ReplUser ^
--password <password> ^
--result-file=/tmp/binlog.txt
```

Amazon RDS normally purges a binary log as soon as possible, but the binary log must still be available on the instance to be accessed by mysqlbinlog. To specify the number of hours for RDS to retain binary logs, use the `mysql.rds_set_configuration` stored procedure and specify a period with enough time for you to download the logs. After you set the retention period, monitor storage usage for the DB instance to ensure that the retained binary logs do not take up too much storage.

The following example sets the retention period to 1 day:

```
call mysql.rds_set_configuration('binlog retention hours', 24);
```

To display the current setting, use the `mysql.rds_show_configuration` stored procedure:

```
call mysql.rds_show_configuration;
```

Binary Log Annotation

In a MariaDB DB instance, you can use the `Annotate_rows` event to annotate a row event with a copy of the SQL query that caused the row event. This approach provides similar functionality to enabling the `binlog_rows_query_log_events` parameter on a DB instance on MySQL version 5.6 or later.

You can enable binary log annotations globally by creating a custom parameter group and setting the `binlog_annotation_row_events` parameter to 1. You can also enable annotations at the session level, by calling `SET SESSION binlog_annotation_row_events = 1`. Use the `replicate_annotation_row_events` to replicate binary log annotations to the slave instance if binary logging is enabled on it. No special privileges are required to use these settings.

The following is an example of a row-based transaction in MariaDB. The use of row-based logging is triggered by setting the transaction isolation level to read-committed.

```
CREATE DATABASE IF NOT EXISTS test;
USE test;
CREATE TABLE square(x INT PRIMARY KEY, y INT NOT NULL) ENGINE = InnoDB;
```

```
SET SESSION TRANSACTION ISOLATION LEVEL READ COMMITTED;
BEGIN
INSERT INTO square(x, y) VALUES(5, 5 * 5);
COMMIT;
```

Without annotations, the binary log entries for the transaction look like the following:

```
BEGIN
/*!*/
# at 1163
# at 1209
#150922 7:55:57 server id 1855786460 end_log_pos 1209      Table_map: `test`.`square`
mapped to number 76
#150922 7:55:57 server id 1855786460 end_log_pos 1247      Write_rows: table id 76
flags: STMT_END_F
### INSERT INTO `test`.`square`
### SET
### @1=5
### @2=25
# at 1247
#150922 7:56:01 server id 1855786460 end_log_pos 1274      Xid = 62
COMMIT/*!*/;
```

The following statement enables session-level annotations for this same transaction, and disables them after committing the transaction:

```
CREATE DATABASE IF NOT EXISTS test;
USE test;
CREATE TABLE square(x INT PRIMARY KEY, y INT NOT NULL) ENGINE = InnoDB;
SET SESSION TRANSACTION ISOLATION LEVEL READ COMMITTED;
SET SESSION binlog_annotation_row_events = 1;
BEGIN;
INSERT INTO square(x, y) VALUES(5, 5 * 5);
COMMIT;
SET SESSION binlog_annotation_row_events = 0;
```

With annotations, the binary log entries for the transaction look like the following:

```
BEGIN
/*!*/
# at 423
# at 483
# at 529
#150922 8:04:24 server id 1855786460 end_log_pos 483 Annotate_rows:
#Q> INSERT INTO square(x, y) VALUES(5, 5 * 5)
#150922 8:04:24 server id 1855786460 end_log_pos 529 Table_map: `test`.`square` mapped
to number 76
#150922 8:04:24 server id 1855786460 end_log_pos 567 Write_rows: table id 76 flags:
STMT_END_F
### INSERT INTO `test`.`square`
### SET
### @1=5
### @2=25
# at 567
#150922 8:04:26 server id 1855786460 end_log_pos 594 Xid = 88
COMMIT/*!*/;
```

Microsoft SQL Server Database Log Files

You can access Microsoft SQL Server error logs, agent logs, trace files, and dump files by using the Amazon RDS console, AWS CLI, or RDS API. For more information about viewing, downloading, and watching file-based database logs, see [Amazon RDS Database Log Files \(p. 390\)](#).

Retention Schedule

Log files are rotated each day and whenever your DB instance is restarted. The following is the retention schedule for Microsoft SQL Server logs on Amazon RDS.

Log Type	Retention Schedule
Error logs	A maximum of 30 error logs are retained. Amazon RDS may delete error logs older than 7 days.
Agent logs	A maximum of 10 agent logs are retained. Amazon RDS may delete agent logs older than 7 days.
Trace files	Trace files are retained according to the trace file retention period of your DB instance. The default trace file retention period is 7 days. To modify the trace file retention period for your DB instance, see Setting the Retention Period for Trace and Dump Files (p. 737) .
Dump files	Dump files are retained according to the dump file retention period of your DB instance. The default dump file retention period is 7 days. To modify the dump file retention period for your DB instance, see Setting the Retention Period for Trace and Dump Files (p. 737) .

Viewing the SQL Server Error Log by Using the `rds_read_error_log` Procedure

You can use the Amazon RDS stored procedure `rds_read_error_log` to view error logs and agent logs. For more information, see [Using the `rds_read_error_log` Procedure \(p. 736\)](#).

Publishing SQL Server Logs to Amazon CloudWatch Logs

With Amazon RDS for SQL Server, you can publish error and agent log events directly to Amazon CloudWatch Logs. Analyze the log data with CloudWatch Logs, then use CloudWatch to create alarms and view metrics.

With CloudWatch Logs, you can do the following:

- Store logs in highly durable storage space with a retention period that you define.
- Search and filter log data.
- Share log data between accounts.
- Export logs to Amazon S3.
- Stream data to Amazon Elasticsearch Service.
- Process log data in real time with Amazon Kinesis Data Streams.

Amazon RDS publishes each SQL Server database log as a separate database stream in the log group. For example, if you publish error logs, error data is stored in an error log stream in the `/aws/rds/instance/my_instance/error` log group.

Note

Publishing SQL Server logs to CloudWatch Logs isn't enabled by default. Publishing trace and dump files isn't supported. Publishing SQL Server logs to CloudWatch Logs is supported in all regions, except for Asia Pacific (Hong Kong).

Console

To publish SQL Server DB logs to CloudWatch Logs from the AWS Management Console

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to modify.
3. Choose **Modify**.
4. In the **Log exports** section, choose the logs that you want to start publishing to CloudWatch Logs.
You can choose **error**, **agent**, or both.
5. Choose **Continue**, and then choose **Modify DB Instance** on the summary page.

AWS CLI

To publish SQL Server logs, you can use the `modify-db-instance` command with the following parameters:

- `--db-instance-identifier`
- `--cloudwatch-logs-export-configuration`

Note

A change to the `--cloudwatch-logs-export-configuration` option is always applied to the DB instance immediately. Therefore, the `--apply-immediately` and `--no-apply-immediately` options have no effect.

You can also publish SQL Server logs using the following commands:

- `create-db-instance`
- `restore-db-instance-from-db-snapshot`
- `restore-db-instance-to-point-in-time`

Example

The following example creates an SQL Server DB instance with CloudWatch Logs publishing enabled. The `--enable-cloudwatch-logs-exports` value is a JSON array of strings that can include `error`, `agent`, or both.

For Linux, OS X, or Unix:

```
aws rds create-db-instance \
--db-instance-identifier mydbinstance \
--enable-cloudwatch-logs-exports '[{"error","agent"}]' \
--db-instance-class db.m4.large \
--engine sqlserver-se
```

For Windows:

```
aws rds create-db-instance ^
```

```
--db-instance-identifier mydbinstance ^
--enable-cloudwatch-logs-exports "[\"error\", \"agent\"]" ^
--db-instance-class db.m4.large ^
--engine sqlserver-se
```

Note

When using the Windows command prompt, you must escape double quotes ("") in JSON code by prefixing them with a backslash (\).

Example

The following example modifies an existing SQL Server DB instance to publish log files to CloudWatch Logs. The --cloudwatch-logs-export-configuration value is a JSON object. The key for this object is `EnableLogTypes`, and its value is an array of strings that can include `error`, `agent`, or both.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier mydbinstance \
--cloudwatch-logs-export-configuration '{"EnableLogTypes": ["error", "agent"]}'
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier mydbinstance ^
--cloudwatch-logs-export-configuration "{\"EnableLogTypes\": [\"error\", \"agent\"]}"
```

Note

When using the Windows command prompt, you must escape double quotes ("") in JSON code by prefixing them with a backslash (\).

Example

The following example modifies an existing SQL Server DB instance to disable publishing agent log files to CloudWatch Logs. The --cloudwatch-logs-export-configuration value is a JSON object. The key for this object is `DisableLogTypes`, and its value is an array of strings that can include `error`, `agent`, or both.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier mydbinstance \
--cloudwatch-logs-export-configuration '{"DisableLogTypes": ["agent"]}'
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier mydbinstance ^
--cloudwatch-logs-export-configuration "{\"DisableLogTypes\": [\"agent\"]}"
```

Note

When using the Windows command prompt, you must escape double quotes ("") in JSON code by prefixing them with a backslash (\).

MySQL Database Log Files

You can monitor the MySQL error log, slow query log, and the general log. The MySQL error log is generated by default; you can generate the slow query and general logs by setting parameters in your DB parameter group. Amazon RDS rotates all of the MySQL log files; the intervals for each type are given following.

You can monitor the MySQL logs directly through the Amazon RDS console, Amazon RDS API, AWS CLI, or AWS SDKs. You can also access MySQL logs by directing the logs to a database table in the main database and querying that table. You can use the `mysqlbinlog` utility to download a binary log.

For more information about viewing, downloading, and watching file-based database logs, see [Amazon RDS Database Log Files \(p. 390\)](#).

Accessing MySQL Error Logs

The MySQL error log is written to the `mysql-error.log` file. You can view `mysql-error.log` by using the Amazon RDS console or by retrieving the log using the Amazon RDS API, Amazon RDS CLI, or AWS SDKs. `mysql-error.log` is flushed every 5 minutes, and its contents are appended to `mysql-error-running.log`. The `mysql-error-running.log` file is then rotated every hour and the hourly files generated during the last 24 hours are retained. Note that the retention period is different between Amazon RDS and Aurora.

Each log file has the hour it was generated (in UTC) appended to its name. The log files also have a timestamp that helps you determine when the log entries were written.

MySQL writes to the error log only on startup, shutdown, and when it encounters errors. A DB instance can go hours or days without new entries being written to the error log. If you see no recent entries, it's because the server did not encounter an error that would result in a log entry.

Accessing the MySQL Slow Query and General Logs

The MySQL slow query log and the general log can be written to a file or a database table by setting parameters in your DB parameter group. For information about creating and modifying a DB parameter group, see [Working with DB Parameter Groups \(p. 211\)](#). You must set these parameters before you can view the slow query log or general log in the Amazon RDS console or by using the Amazon RDS API, Amazon RDS CLI, or AWS SDKs.

You can control MySQL logging by using the parameters in this list:

- `slow_query_log`: To create the slow query log, set to 1. The default is 0.
- `general_log`: To create the general log, set to 1. The default is 0.
- `long_query_time`: To prevent fast-running queries from being logged in the slow query log, specify a value for the shortest query execution time to be logged, in seconds. The default is 10 seconds; the minimum is 0. If `log_output = FILE`, you can specify a floating point value that goes to microsecond resolution. If `log_output = TABLE`, you must specify an integer value with second resolution. Only queries whose execution time exceeds the `long_query_time` value are logged. For example, setting `long_query_time` to 0.1 prevents any query that runs for less than 100 milliseconds from being logged.
- `log_queries_not_using_indexes`: To log all queries that do not use an index to the slow query log, set to 1. The default is 0. Queries that do not use an index are logged even if their execution time is less than the value of the `long_query_time` parameter.
- `log_output option`: You can specify one of the following options for the `log_output` parameter.
 - **TABLE** (default)– Write general queries to the `mysql.general_log` table, and slow queries to the `mysql.slow_log` table.

- **FILE**— Write both general and slow query logs to the file system. Log files are rotated hourly.
- **NONE**— Disable logging.

When logging is enabled, Amazon RDS rotates table logs or deletes log files at regular intervals. This measure is a precaution to reduce the possibility of a large log file either blocking database use or affecting performance. **FILE** and **TABLE** logging approach rotation and deletion as follows:

- When **FILE** logging is enabled, log files are examined every hour and log files older than 24 hours are deleted. In some cases, the remaining combined log file size after the deletion might exceed the threshold of 2 percent of a DB instance's allocated space. In these cases, the largest log files are deleted until the log file size no longer exceeds the threshold.
- When **TABLE** logging is enabled, in some cases log tables are rotated every 24 hours. This rotation occurs if the space used by the table logs is more than 20 percent of the allocated storage space or the size of all logs combined is greater than 10 GB. If the amount of space used for a DB instance is greater than 90 percent of the DB instance's allocated storage space, then the thresholds for log rotation are reduced. Log tables are then rotated if the space used by the table logs is more than 10 percent of the allocated storage space or the size of all logs combined is greater than 5 GB. You can subscribe to the `low_free_storage` event to be notified when log tables are rotated to free up space. For more information, see [Using Amazon RDS Event Notification \(p. 370\)](#).

When log tables are rotated, the current log table is copied to a backup log table and the entries in the current log table are removed. If the backup log table already exists, then it is deleted before the current log table is copied to the backup. You can query the backup log table if needed. The backup log table for the `mysql.general_log` table is named `mysql.general_log_backup`. The backup log table for the `mysql.slow_log` table is named `mysql.slow_log_backup`.

You can rotate the `mysql.general_log` table by calling the `mysql.rds_rotate_general_log` procedure. You can rotate the `mysql.slow_log` table by calling the `mysql.rds_rotate_slow_log` procedure.

Table logs are rotated during a database version upgrade.

To work with the logs from the Amazon RDS console, Amazon RDS API, Amazon RDS CLI, or AWS SDKs, set the `log_output` parameter to **FILE**. Like the MySQL error log, these log files are rotated hourly. The log files that were generated during the previous 24 hours are retained. Note that the retention period is different between Amazon RDS and Aurora.

For more information about the slow query and general logs, go to the following topics in the MySQL documentation:

- [The Slow Query Log](#)
- [The General Query Log](#)

Accessing the MySQL Audit Log

To access the audit log, the DB instance must use a custom option group with the `MARIADB_AUDIT_PLUGIN` option. For more information, see [MariaDB Audit Plugin Support \(p. 874\)](#).

Publishing MySQL Logs to CloudWatch Logs

You can configure your Amazon RDS MySQL DB instance to publish log data to a log group in Amazon CloudWatch Logs. With CloudWatch Logs, you can perform real-time analysis of the log data, and use CloudWatch to create alarms and view metrics. You can use CloudWatch Logs to store your log records in highly durable storage.

Amazon RDS publishes each MySQL database log as a separate database stream in the log group. For example, if you configure the export function to include the slow query log, slow query data is stored in a slow query log stream in the `/aws/rds/instance/my_instance/slowquery` log group.

The error log is enabled by default. The following table summarizes the requirements for the other MySQL logs.

Log	Requirement
Audit log	The DB instance must use a custom option group with the <code>MARIADB_AUDIT_PLUGIN</code> option.
General log	The DB instance must use a custom parameter group with the parameter setting <code>general_log = 1</code> to enable the general log.
Slow query log	The DB instance must use a custom parameter group with the parameter setting <code>slow_query_log = 1</code> to enable the slow query log.
Log output	The DB instance must use a custom parameter group with the parameter setting <code>log_output = FILE</code> to write logs to the file system and publish them to CloudWatch Logs.

Note

Publishing log files to CloudWatch Logs is only supported for MySQL versions 5.6, 5.7, and 8.0.

Console

To publish MySQL logs to CloudWatch Logs using the console

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to modify.
3. Choose **Modify**.
4. In the **Log exports** section, choose the logs that you want to start publishing to CloudWatch Logs.
5. Choose **Continue**, and then choose **Modify DB Instance** on the summary page.

AWS CLI

You can publish MySQL logs with the AWS CLI. You can call the `modify-db-instance` command with the following parameters:

- `--db-instance-identifier`
- `--cloudwatch-logs-export-configuration`

Note

A change to the `--cloudwatch-logs-export-configuration` option is always applied to the DB instance immediately. Therefore, the `--apply-immediately` and `--no-apply-immediately` options have no effect.

You can also publish MySQL logs by calling the following AWS CLI commands:

- `create-db-instance`
- `restore-db-instance-from-db-snapshot`
- `restore-db-instance-from-s3`
- `restore-db-instance-to-point-in-time`

Run one of these AWS CLI commands with the following options:

- `--db-instance-identifier`
- `--enable-cloudwatch-logs-exports`
- `--db-instance-class`
- `--engine`

Other options might be required depending on the AWS CLI command you run.

Example

The following example modifies an existing MySQL DB instance to publish log files to CloudWatch Logs. The `--cloudwatch-logs-export-configuration` value is a JSON object. The key for this object is `EnableLogTypes`, and its value is an array of strings with any combination of `audit`, `error`, `general`, and `slowquery`.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
    --db-instance-identifier mydbinstance \
    --cloudwatch-logs-export-configuration '{"EnableLogTypes": \
    ["audit","error","general","slowquery"]}'
```

For Windows:

```
aws rds modify-db-instance ^
    --db-instance-identifier mydbinstance ^
    --cloudwatch-logs-export-configuration '{"EnableLogTypes": \
    ["audit","error","general","slowquery"]}'
```

Example

The following example creates a MySQL DB instance and publishes log files to CloudWatch Logs. The `--enable-cloudwatch-logs-exports` value is a JSON array of strings. The strings can be any combination of `audit`, `error`, `general`, and `slowquery`.

For Linux, OS X, or Unix:

```
aws rds create-db-instance \
    --db-instance-identifier mydbinstance \
    --enable-cloudwatch-logs-exports '[{"audit","error","general","slowquery"}]' \
    --db-instance-class db.m4.large \
    --engine MySQL
```

For Windows:

```
aws rds create-db-instance ^
```

```
--db-instance-identifier mydbinstance ^
--enable-cloudwatch-logs-exports '[ "audit", "error", "general", "slowquery" ]' ^
--db-instance-class db.m4.large ^
--engine MySQL
```

RDS API

You can publish MySQL logs with the RDS API. You can call the [ModifyDBInstance](#) action with the following parameters:

- `DBInstanceIdentifier`
- `CloudwatchLogsExportConfiguration`

Note

A change to the `CloudwatchLogsExportConfiguration` parameter is always applied to the DB instance immediately. Therefore, the `ApplyImmediately` parameter has no effect.

You can also publish MySQL logs by calling the following RDS API operations:

- [CreateDBInstance](#)
- [RestoreDBInstanceFromDBSnapshot](#)
- [RestoreDBInstanceFromS3](#)
- [RestoreDBInstanceToPointInTime](#)

Run one of these RDS API operations with the following parameters:

- `DBInstanceIdentifier`
- `EnableCloudwatchLogsExports`
- `Engine`
- `DBInstanceClass`

Other parameters might be required depending on the AWS CLI command you run.

Log File Size

The MySQL slow query log, error log, and the general log file sizes are constrained to no more than 2 percent of the allocated storage space for a DB instance. To maintain this threshold, logs are automatically rotated every hour and log files older than 24 hours are removed. If the combined log file size exceeds the threshold after removing old log files, then the largest log files are deleted until the log file size no longer exceeds the threshold.

For MySQL, there is a size limit on BLOBS written to the redo log. To account for this limit, ensure that the `innodb_log_file_size` parameter for your MySQL DB instance is 10 times larger than the largest BLOB data size found in your tables, plus the length of other variable length fields (`VARCHAR`, `VARBINARY`, `TEXT`) in the same tables. For information on how to set parameter values, see [Working with DB Parameter Groups \(p. 211\)](#). For information on the redo log BLOB size limit, go to [Changes in MySQL 5.6.20](#).

Managing Table-Based MySQL Logs

You can direct the general and slow query logs to tables on the DB instance by creating a DB parameter group and setting the `log_output` server parameter to `TABLE`. General queries are then logged to the `mysql.general_log` table, and slow queries are logged to the `mysql.slow_log` table. You can query

the tables to access the log information. Enabling this logging increases the amount of data written to the database, which can degrade performance.

Both the general log and the slow query logs are disabled by default. In order to enable logging to tables, you must also set the `general_log` and `slow_query_log` server parameters to 1.

Log tables keep growing until the respective logging activities are turned off by resetting the appropriate parameter to 0. A large amount of data often accumulates over time, which can use up a considerable percentage of your allocated storage space. Amazon RDS does not allow you to truncate the log tables, but you can move their contents. Rotating a table saves its contents to a backup table and then creates a new empty log table. You can manually rotate the log tables with the following command line procedures, where the command prompt is indicated by `PROMPT>`:

```
PROMPT> CALL mysql.rds_rotate_slow_log;
PROMPT> CALL mysql.rds_rotate_general_log;
```

To completely remove the old data and reclaim the disk space, call the appropriate procedure twice in succession.

Binary Logging Format

MySQL on Amazon RDS supports the *row-based*, *statement-based*, and *mixed* binary logging formats for MySQL version 5.6 and later. The default binary logging format is mixed. For DB instances running MySQL versions 5.1 and 5.5, only mixed binary logging is supported. For details on the different MySQL binary log formats, see [Binary Logging Formats](#) in the MySQL documentation.

If you plan to use replication, the binary logging format is important because it determines the record of data changes that is recorded in the source and sent to the replication targets. For information about the advantages and disadvantages of different binary logging formats for replication, see [Advantages and Disadvantages of Statement-Based and Row-Based Replication](#) in the MySQL documentation.

Important

Setting the binary logging format to row-based can result in very large binary log files. Large binary log files reduce the amount of storage available for a DB instance and can increase the amount of time to perform a restore operation of a DB instance.

Statement-based replication can cause inconsistencies between the source DB instance and a Read Replica. For more information, see [Determination of Safe and Unsafe Statements in Binary Logging](#) in the MySQL documentation.

To set the MySQL binary logging format

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Parameter groups**.
3. Choose the parameter group used by the DB instance you want to modify.

You can't modify a default parameter group. If the DB instance is using a default parameter group, create a new parameter group and associate it with the DB instance.

For more information on DB parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#).

4. From **Parameter group actions**, choose **Edit**.
5. Set the `binlog_format` parameter to the binary logging format of your choice (**ROW**, **STATEMENT**, or **MIXED**).
6. Choose **Save changes** to save the updates to the DB parameter group.

Important

Changing the `default.mysql5.6`, `default.mysql5.7`, or `default.mysql8.0` DB parameter group affects all MySQL version DB instances that use that parameter group. If

If you want to specify different binary logging formats for different MySQL 5.6, 5.7, or 8.0 DB instances in an AWS Region, you need to create your own DB parameter group. This parameter group identifies the different logging format and assigns that DB parameter group to the intended DB instances.

Accessing MySQL Binary Logs

You can use the `mysqlbinlog` utility to download or stream binary logs from Amazon RDS instances running MySQL 5.6 or later. The binary log is downloaded to your local computer, where you can perform actions such as replaying the log using the `mysql` utility. For more information about using the `mysqlbinlog` utility, go to [Using mysqlbinlog to Back Up Binary Log Files](#).

To run the `mysqlbinlog` utility against an Amazon RDS instance, use the following options:

- Specify the `--read-from-remote-server` option.
- `--host`: Specify the DNS name from the endpoint of the instance.
- `--port`: Specify the port used by the instance.
- `--user`: Specify a MySQL user that has been granted the replication slave permission.
- `--password`: Specify the password for the user, or omit a password value so that the utility prompts you for a password.
- To have the file downloaded in binary format, specify the `--raw` option.
- `--result-file`: Specify the local file to receive the raw output.
- Specify the names of one or more binary log files. To get a list of the available logs, use the SQL command `SHOW BINARY LOGS`.
- To stream the binary log files, specify the `--stop-never` option.

For more information about `mysqlbinlog` options, go to [mysqlbinlog - Utility for Processing Binary Log Files](#).

For example, see the following.

For Linux, OS X, or Unix:

```
mysqlbinlog \
--read-from-remote-server \
--host=MySQL56Instance1.cg034hpkmmt.region.rds.amazonaws.com \
--port=3306 \
--user ReplUser \
--password \
--raw \
--result-file=/tmp/ \
binlog.00098
```

For Windows:

```
mysqlbinlog ^
--read-from-remote-server ^
--host=MySQL56Instance1.cg034hpkmmt.region.rds.amazonaws.com ^
--port=3306 ^
--user ReplUser ^
--password ^
--raw ^
--result-file=/tmp/ ^
binlog.00098
```

Amazon RDS normally purges a binary log as soon as possible, but the binary log must still be available on the instance to be accessed by mysqlbinlog. To specify the number of hours for RDS to retain binary logs, use the `mysql.rds_set_configuration` stored procedure and specify a period with enough time for you to download the logs. After you set the retention period, monitor storage usage for the DB instance to ensure that the retained binary logs don't take up too much storage.

Note

The `mysql.rds_set_configuration` stored procedure is only available for MySQL version 5.6 or later.

The following example sets the retention period to 1 day.

```
call mysql.rds_set_configuration('binlog retention hours', 24);
```

To display the current setting, use the `mysql.rds_show_configuration` stored procedure.

```
call mysql.rds_show_configuration;
```

Oracle Database Log Files

You can access Oracle alert logs, audit files, and trace files by using the Amazon RDS console or API. For more information about viewing, downloading, and watching file-based database logs, see [Amazon RDS Database Log Files \(p. 390\)](#).

The Oracle audit files provided are the standard Oracle auditing files. Amazon RDS supports the Oracle fine-grained auditing (FGA) feature. However, log access doesn't provide access to FGA events that are stored in the `SYS.FGA_LOG$` table and that are accessible through the `DBA_FGA_AUDIT_TRAIL` view.

The `DescribeDBLogFiles` API operation that lists the Oracle log files that are available for a DB instance ignores the `MaxRecords` parameter and returns up to 1,000 records.

Retention Schedule

The Oracle database engine might rotate logs files if they get very large. To retain audit or trace files, download them. Storing the files locally reduces your Amazon RDS storage costs and makes more space available for your data.

The following is the retention schedule for Oracle alert logs, audit files, and trace files on Amazon RDS.

Log Type	Retention Schedule
Alert logs	The text alert log is rotated daily with 30-day retention managed by Amazon RDS. The XML alert log is retained for at least seven days. You can access this log by using the <code>ALERTLOG</code> view.
Audit files	The default retention period for audit files is seven days. Amazon RDS might delete audit files older than seven days.
Trace files	The default retention period for trace files is seven days. Amazon RDS might delete trace files older than seven days.
Listener logs	The default retention period for the listener logs is seven days. Amazon RDS might delete listener logs older than seven days.

Note

Audit files and trace files share the same retention configuration.

Switching Online Log files

You can use the Amazon RDS procedure `rdsadmin.rdsadmin_util.switch_logfile` to switch online log files. For more information, see [Switching Online Log Files \(p. 1107\)](#).

Retrieving Archived Redo Logs

You can retain archived redo logs. For more information, see [Retaining Archived Redo Logs \(p. 1110\)](#).

Working with Oracle Trace Files

Following, you can find descriptions of Amazon RDS procedures to create, refresh, access, and delete trace files.

[Listing Files](#)

You can use either of two procedures to allow access to any file in the `background_dump_dest` path. The first procedure refreshes a view containing a listing of all files currently in `background_dump_dest`.

```
exec rdsadmin.manage_tracefiles.refresh_tracefile_listing;
```

After the view is refreshed, use the following view to access the results.

```
rdsadmin.tracefile_listing
```

An alternative to the previous process is to use `FROM table` to stream nontable data in a table-like format to list database directory contents.

```
SELECT * FROM table(rdsadmin.rds_file_util.listdir('BDUMP'));
```

The following query shows the text of a log file.

```
SELECT text FROM table(rdsadmin.rds_file_util.read_text_file('BDUMP','alert_xxx.log'));
```

Generating Trace Files and Tracing a Session

Because there are no restrictions on `alter session`, many standard methods to generate trace files in Oracle remain available to an Amazon RDS DB instance. The following procedures are provided for trace files that require greater access.

Oracle Method	Amazon RDS Method
<code>oradebug hanganalyze 3</code>	<code>exec rdsadmin.manage_tracefiles.hanganalyze;</code>
<code>oradebug dump systemstate 266</code>	<code>exec rdsadmin.manage_tracefiles.dump_systemstate;</code>

You can use many standard methods to trace individual sessions connected to an Oracle DB instance in Amazon RDS. To enable tracing for a session, you can run subprograms in PL/SQL packages supplied by Oracle, such as the `DBMS_SESSION` and `DBMS_MONITOR` packages. For more information, see [Enabling Tracing for a Session](#) in the Oracle documentation.

Retrieving Trace Files

You can retrieve any trace file in `background_dump_dest` using a standard SQL query on an Amazon RDS-managed external table. To use this method, you must execute the procedure to set the location for this table to the specific trace file.

For example, you can use the `rdsadmin.tracefile_listing` view mentioned preceding to list all of the trace files on the system. You can then set the `tracefile_table` view to point to the intended trace file using the following procedure.

```
exec
rdsadmin.manage_tracefiles.set_tracefile_table_location('CUST01_ora_3260_SYSTEMSTATE.trc');
```

The following example creates an external table in the current schema with the location set to the file provided. You can retrieve the contents into a local file using a SQL query.

```
# eg: send the contents of the tracefile to a local file:
```

```
sqlplus user/password@TNS alias << EOF > /tmp/tracefile.txt
select * from tracefile_table;
EOF
```

Purging Trace Files

Trace files can accumulate and consume disk space. Amazon RDS purges trace files by default and log files that are older than seven days. You can view and set the trace file retention period using the `show_configuration` procedure. You should run the command `SET SERVEROUTPUT ON` so that you can view the configuration results.

The following example shows the current trace file retention period, and then sets a new trace file retention period.

```
# Show the current tracefile retention
SQL> exec rdsadmin.rdsadmin_util.show_configuration;
NAME:tracefile retention
VALUE:10080
DESCRIPTION:tracefile expiration specifies the duration in minutes before tracefiles in
bdbump are automatically deleted.

# Set the tracefile retention to 24 hours:
SQL> exec rdsadmin.rdsadmin_util.set_configuration('tracefile retention',1440);

#show the new tracefile retention
SQL> exec rdsadmin.rdsadmin_util.show_configuration;
NAME:tracefile retention
VALUE:1440
DESCRIPTION:tracefile expiration specifies the duration in minutes before tracefiles in
bdbump are automatically deleted.
```

In addition to the periodic purge process, you can manually remove files from the `background_dump_dest`. The following example shows how to purge all files older than five minutes.

```
exec rdsadmin.manage_tracefiles.purge_tracefiles(5);
```

You can also purge all files that match a specific pattern (if you do, don't include the file extension, such as .trc). The following example shows how to purge all files that start with `SCHPOC1_ora_5935`.

```
exec rdsadmin.manage_tracefiles.purge_tracefiles('SCHPOC1_ora_5935');
```

Publishing Oracle Logs to Amazon CloudWatch Logs

You can configure your Amazon RDS Oracle DB instance to publish log data to a log group in Amazon CloudWatch Logs. With CloudWatch Logs, you can analyze the log data, and use CloudWatch to create alarms and view metrics. You can use CloudWatch Logs to store your log records in highly durable storage.

Amazon RDS publishes each Oracle database log as a separate database stream in the log group. For example, if you configure the export function to include the audit log, audit data is stored in an audit log stream in the `/aws/rds/instance/my_instance/audit` log group.

Console

To publish Oracle DB logs to CloudWatch Logs from the AWS Management Console

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.

2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to modify.
3. Choose **Modify**.
4. In the **Log exports** section, choose the logs that you want to start publishing to CloudWatch Logs.
5. Choose **Continue**, and then choose **Modify DB Instance** on the summary page.

AWS CLI

To publish Oracle logs, you can use the `modify-db-instance` command with the following parameters:

- `--db-instance-identifier`
- `--cloudwatch-logs-export-configuration`

Note

A change to the `--cloudwatch-logs-export-configuration` option is always applied to the DB instance immediately. Therefore, the `--apply-immediately` and `--no-apply-immediately` options have no effect.

You can also publish Oracle logs using the following commands:

- `create-db-instance`
- `restore-db-instance-from-db-snapshot`
- `restore-db-instance-from-s3`
- `restore-db-instance-to-point-in-time`

Example

The following example creates an Oracle DB instance with CloudWatch Logs publishing enabled. The `--enable-cloudwatch-logs-exports` value is a JSON array of strings. The strings can be any combination of `alert`, `audit`, `listener`, and `trace`.

For Linux, OS X, or Unix:

```
aws rds create-db-instance \
--db-instance-identifier mydbinstance \
--enable-cloudwatch-logs-exports '[{"trace","audit","alert","listener"}]' \
--db-instance-class db.m5.large \
--allocated-storage 20 \
--engine oracle-ee \
--engine-version 12.1.0.2.v18 \
--license-model bring-your-own-license \
--master-username myadmin \
--master-user-password mypassword
```

For Windows:

```
aws rds create-db-instance ^
--db-instance-identifier mydbinstance ^
--enable-cloudwatch-logs-exports trace alert audit listener ^
--db-instance-class db.m5.large ^
--allocated-storage 20 ^
--engine oracle-ee ^
--engine-version 12.1.0.2.v18 ^
```

```
--license-model bring-your-own-license ^
--master-username myadmin ^
--master-user-password mypassword
```

Example

The following example modifies an existing Oracle DB instance to publish log files to CloudWatch Logs. The --cloudwatch-logs-export-configuration value is a JSON object. The key for this object is `EnableLogTypes`, and its value is an array of strings with any combination of `alert`, `audit`, `listener`, and `trace`.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
  --db-instance-identifier mydbinstance \
  --cloudwatch-logs-export-configuration '{"EnableLogTypes": \
  ["trace", "alert", "audit", "listener"]}'
```

For Windows:

```
aws rds modify-db-instance ^
  --db-instance-identifier mydbinstance ^
  --cloudwatch-logs-export-configuration EnableLogTypes=\"trace\", \"alert\", \"audit\",
  \"listener\""
```

Example

The following example modifies an existing Oracle DB instance to disable publishing audit and listener log files to CloudWatch Logs. The --cloudwatch-logs-export-configuration value is a JSON object. The key for this object is `DisableLogTypes`, and its value is an array of strings with any combination of `audit`, `audit`, `listener`, and `trace`.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
  --db-instance-identifier mydbinstance \
  --cloudwatch-logs-export-configuration '{"DisableLogTypes": ["audit", "listener"]}'
```

For Windows:

```
aws rds modify-db-instance ^
  --db-instance-identifier mydbinstance ^
  --cloudwatch-logs-export-configuration DisableLogTypes=\"audit\", \"listener\""
```

RDS API

You can publish Oracle DB logs with the RDS API. You can call the [ModifyDBInstance](#) action with the following parameters:

- `DBInstanceIdentifier`
- `CloudwatchLogsExportConfiguration`

Note

A change to the `CloudwatchLogsExportConfiguration` parameter is always applied to the DB instance immediately. Therefore, the `ApplyImmediately` parameter has no effect.

You can also publish Oracle logs by calling the following RDS API operations:

- [CreateDBInstance](#)
- [RestoreDBInstanceFromDBSnapshot](#)
- [RestoreDBInstanceFromS3](#)
- [RestoreDBInstanceToPointInTime](#)

Run one of these RDS API operations with the following parameters:

- `DBInstanceIdentifier`
- `EnableCloudwatchLogsExports`
- `Engine`
- `DBInstanceClass`

Other parameters might be required depending on the RDS operation that you run.

Previous Methods for Accessing Alert Logs and Listener Logs

You can view the alert log using the Amazon RDS console. You can also use the following SQL statement to access the alert log.

```
select message_text from alertlog;
```

To access the listener log, use the following SQL statement.

```
select message_text from listenerlog;
```

Note

Oracle rotates the alert and listener logs when they exceed 10 MB, at which point they are unavailable from Amazon RDS views.

PostgreSQL Database Log Files

RDS PostgreSQL generates query and error logs. We write auto-vacuum information and rds_admin actions to the error log. PostgreSQL also logs connections, disconnections, and checkpoints to the error log. For more information, see the [Error Reporting and Logging](#) in the PostgreSQL documentation.

You can set the retention period for system logs using the `rds.log_retention_period` parameter in the DB parameter group associated with your DB instance. The unit for this parameter is minutes. For example, a setting of 1440 would retain logs for one day. The default value is 4320 (three days). The maximum value is 10080 (seven days). Note that your instance must have enough allocated storage to contain the retained log files.

You can enable query logging for your PostgreSQL DB instance by setting two parameters in the DB parameter group associated with your DB instance: `log_statement` and `log_min_duration_statement`. The `log_statement` parameter controls which SQL statements are logged. We recommend setting this parameter to `all` to log all statements when debugging issues in your DB instance. The default value is `none`. Alternatively, you can set this value to `ddl` to log all data definition language (DDL) statements (CREATE, ALTER, DROP, and so on) or to `mod` to log all DDL and data modification language (DML) statements (INSERT, UPDATE, DELETE, and so on).

The `log_min_duration_statement` parameter sets the limit in milliseconds of a statement to be logged. All SQL statements that run longer than the parameter setting are logged. This parameter is disabled and set to minus 1 (-1) by default. Enabling this parameter can help you find unoptimized queries.

If you are new to setting parameters in a DB parameter group and associating that parameter group with a DB instance, see [Working with DB Parameter Groups \(p. 211\)](#)

The following steps show how to set up query logging:

1. Set the `log_statement` parameter to `all`. The following example shows the information that is written to the `postgres.log` file:

```
2013-11-05 16:48:56 UTC::@[2952]:LOG: received SIGHUP, reloading configuration files
2013-11-05 16:48:56 UTC::@[2952]:LOG: parameter "log_statement" changed to "all"
```

Additional information is written to the `postgres.log` file when you execute a query. The following example shows the type of information written to the file after a query:

```
2013-11-05 16:41:07 UTC::@[2955]:LOG: checkpoint starting: time
2013-11-05 16:41:07 UTC::@[2955]:LOG: checkpoint complete: wrote 1 buffers (0.3%),
0 transaction log file(s) added, 0 removed, 1 recycled; write=0.000 s, sync=0.003 s,
total=0.012 s; sync files=1, longest=0.003 s, average=0.003 s
2013-11-05 16:45:14 UTC:[local]:master@postgres:[8839]:LOG: statement: SELECT d.datname
as "Name",
pg_catalog.pg_get_userbyid(d.datdba) as "Owner",
pg_catalog.pg_encoding_to_char(d.encoding) as "Encoding",
d.datcollate as "Collate",
d.datctype as "Ctype",
pg_catalog.array_to_string(d.datacl, E'\n') AS "Access privileges"
FROM pg_catalog.pg_database d
ORDER BY 1;
2013-11-05 16:45:
```

2. Set the `log_min_duration_statement` parameter. The following example shows the information that is written to the `postgres.log` file when the parameter is set to 1:

```
2013-11-05 16:48:56 UTC::@[2952]:LOG: received SIGHUP, reloading configuration files
```

```
2013-11-05 16:48:56 UTC:[@:[2952]:LOG: parameter "log_min_duration_statement" changed to "1"
```

Additional information is written to the `postgres.log` file when you execute a query that exceeds the duration parameter setting. The following example shows the type of information written to the file after a query:

```
2013-11-05 16:51:10 UTC:[local]:master@postgres:[9193]:LOG: statement: SELECT
c2.relname, i.indisprimary, i.indisunique, i.indisclustered, i.indisvalid,
pg_catalog.pg_get_indexdef(i.indexrelid, 0, true),
pg_catalog.pg_get_constraintdef(con.oid, true), contype, condeferrable, condeferred,
c2.reltablename
FROM pg_catalog.pg_class c, pg_catalog.pg_class c2, pg_catalog.pg_index i
LEFT JOIN pg_catalog.pg_constraint con ON (conrelid = i.indrelid AND conindid = i.indexrelid AND contype IN ('p','u','x'))
WHERE c.oid = '1255' AND c.oid = i.indrelid AND i.indexrelid = c2.oid
ORDER BY i.indisprimary DESC, i.indisunique DESC, c2.relname;
2013-11-05 16:51:10 UTC:[local]:master@postgres:[9193]:LOG: duration: 3.367 ms
2013-11-05 16:51:10 UTC:[local]:master@postgres:[9193]:LOG: statement: SELECT
c.oid::pg_catalog.regclass FROM pg_catalog.pg_class c, pg_catalog.pg_inherits i WHERE
c.oid=i.inhparent AND i.inhrelid = '1255' ORDER BY inhseqno;
2013-11-05 16:51:10 UTC:[local]:master@postgres:[9193]:LOG: duration: 1.002 ms
2013-11-05 16:51:10 UTC:[local]:master@postgres:[9193]:LOG: statement:
SELECT c.oid::pg_catalog.regclass FROM pg_catalog.pg_class c,
pg_catalog.pg_inherits i WHERE c.oid=i.inhrelid AND i.inhparent = '1255' ORDER BY
c.oid::pg_catalog.regclass::pg_catalog.text;
2013-11-05 16:51:18 UTC:[local]:master@postgres:[9193]:LOG: statement: select proname
from pg_proc;
2013-11-05 16:51:18 UTC:[local]:master@postgres:[9193]:LOG: duration: 3.469 ms
```

Publishing PostgreSQL Logs to CloudWatch Logs

You can configure your Amazon RDS for PostgreSQL DB instance to publish log data to a log group in Amazon CloudWatch Logs. With CloudWatch Logs, you can perform real-time analysis of the log data, and use CloudWatch to create alarms and view metrics. You can use CloudWatch Logs to store your log records in highly durable storage.

Note

Publishing log files to CloudWatch Logs is only supported for PostgreSQL versions 9.6.6 and above and 10.4 and above.

Following are the log types that can be published to CloudWatch Logs for Amazon RDS for PostgreSQL.

- Postgresql log
- Upgrade log

After you complete the configuration, Amazon RDS publishes the log events to log streams within a CloudWatch log group. For example, the PostgreSQL log data is stored within the log group `/aws/rds/instance/my_instance/postgresql`. To view your Amazon CloudWatch Logs, open <https://console.aws.amazon.com/cloudwatch/>.

Console

To publish PostgreSQL logs to CloudWatch Logs using the console

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Select the DB instance that you want to modify, and then choose **Modify**.

4. In the **Log exports** section, choose the logs you want to start publishing to CloudWatch Logs.

The **Log exports** section is only available for PostgreSQL versions that support publishing to CloudWatch Logs.

5. Choose **Continue**, and then choose **Modify DB Instance** on the summary page.

AWS CLI

You can publish PostgreSQL logs with the AWS CLI. You can call the [modify-db-instance](#) command with the following parameters:

- `--db-instance-identifier`
- `--cloudwatch-logs-export-configuration`

Note

A change to the `--cloudwatch-logs-export-configuration` option is always applied to the DB instance immediately. Therefore, the `--apply-immediately` and `--no-apply-immediately` options have no effect.

You can also publish PostgreSQL logs by calling the following AWS CLI commands:

- [create-db-instance](#)
- [restore-db-instance-from-db-snapshot](#)
- [restore-db-instance-to-point-in-time](#)

Run one of these AWS CLI commands with the following options:

- `--db-instance-identifier`
- `--enable-cloudwatch-logs-exports`
- `--db-instance-class`
- `--engine`

Other options might be required depending on the AWS CLI command you run.

Modify an instance to publish logs to CloudWatch Logs

Example

The following example modifies an existing PostgreSQL DB instance to publish log files to CloudWatch Logs. The `--cloudwatch-logs-export-configuration` value is a JSON object. The key for this object is `EnableLogTypes`, and its value is an array of strings with any combination of `postgresql` and `upgrade`.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier mydbinstance \
--cloudwatch-logs-export-configuration '{"EnableLogTypes": ["postgresql", "upgrade"]}'
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier mydbinstance ^
```

```
--cloudwatch-logs-export-configuration '{"EnableLogTypes":["postgresql","upgrade"]}'
```

Example

Create an instance to publish logs to CloudWatch Logs

The following example creates a PostgreSQL DB instance and publishes log files to CloudWatch Logs. The `--enable-cloudwatch-logs-exports` value is a JSON array of strings. The strings can be any combination of `postgresql` and `upgrade`.

For Linux, OS X, or Unix:

```
aws rds create-db-instance \
--db-instance-identifier mydbinstance \
--enable-cloudwatch-logs-exports '[ "postgresql", "upgrade" ]' \
--db-instance-class db.m4.large \
--engine postgres
```

For Windows:

```
aws rds create-db-instance ^
--db-instance-identifier mydbinstance ^
--enable-cloudwatch-logs-exports '[ "postgresql", "upgrade" ]' ^
--db-instance-class db.m4.large ^
--engine postgres
```

RDS API

You can publish PostgreSQL logs with the RDS API. You can call the [ModifyDBInstance](#) action with the following parameters:

- `DBInstanceIdentifier`
- `CloudwatchLogsExportConfiguration`

Note

A change to the `CloudwatchLogsExportConfiguration` parameter is always applied to the DB instance immediately. Therefore, the `ApplyImmediately` parameter has no effect.

You can also publish PostgreSQL logs by calling the following RDS API operations:

- [CreateDBInstance](#)
- [RestoreDBInstanceFromDBSnapshot](#)
- [RestoreDBInstanceToPointInTime](#)

Run one of these RDS API operations with the following parameters:

- `DBInstanceIdentifier`
- `EnableCloudwatchLogsExports`
- `Engine`
- `DBInstanceClass`

Other parameters might be required depending on the action you run.

Logging Amazon RDS API Calls with AWS CloudTrail

Amazon RDS is integrated with AWS CloudTrail, a service that provides a record of actions taken by a user, role, or an AWS service in Amazon RDS. CloudTrail captures all API calls for Amazon RDS as events, including calls from the Amazon RDS console and from code calls to the Amazon RDS APIs. If you create a trail, you can enable continuous delivery of CloudTrail events to an Amazon S3 bucket, including events for Amazon RDS. If you don't configure a trail, you can still view the most recent events in the CloudTrail console in **Event history**. Using the information collected by CloudTrail, you can determine the request that was made to Amazon RDS, the IP address from which the request was made, who made the request, when it was made, and additional details.

To learn more about CloudTrail, see the [AWS CloudTrail User Guide](#).

Amazon RDS Information in CloudTrail

CloudTrail is enabled on your AWS account when you create the account. When activity occurs in Amazon RDS, that activity is recorded in a CloudTrail event along with other AWS service events in **Event history**. You can view, search, and download recent events in your AWS account. For more information, see [Viewing Events with CloudTrail Event History](#).

For an ongoing record of events in your AWS account, including events for Amazon RDS, create a trail. A trail enables CloudTrail to deliver log files to an Amazon S3 bucket. By default, when you create a trail in the console, the trail applies to all regions. The trail logs events from all regions in the AWS partition and delivers the log files to the Amazon S3 bucket that you specify. Additionally, you can configure other AWS services to further analyze and act upon the event data collected in CloudTrail logs. For more information, see:

- [Overview for Creating a Trail](#)
- [CloudTrail Supported Services and Integrations](#)
- [Configuring Amazon SNS Notifications for CloudTrail](#)
- [Receiving CloudTrail Log Files from Multiple Regions](#) and [Receiving CloudTrail Log Files from Multiple Accounts](#)

All Amazon RDS actions are logged by CloudTrail and are documented in the [Amazon RDS API Reference](#). For example, calls to the `CreateDBInstance`, `ModifyDBInstance`, and `CreateDBParameterGroup` actions generate entries in the CloudTrail log files.

Every event or log entry contains information about who generated the request. The identity information helps you determine the following:

- Whether the request was made with root or IAM user credentials.
- Whether the request was made with temporary security credentials for a role or federated user.
- Whether the request was made by another AWS service.

For more information, see the [CloudTrail userIdentity Element](#).

Understanding Amazon RDS Log File Entries

A trail is a configuration that enables delivery of events as log files to an Amazon S3 bucket that you specify. CloudTrail log files contain one or more log entries. An event represents a single request from any source and includes information about the requested action, the date and time of the action, request

parameters, and so on. CloudTrail log files are not an ordered stack trace of the public API calls, so they do not appear in any specific order.

The following example shows a CloudTrail log entry that demonstrates the `CreateDBInstance` action.

```
{  
    "eventVersion": "1.04",  
    "userIdentity": {  
        "type": "IAMUser",  
        "principalId": "AKIAIOSFODNN7EXAMPLE",  
        "arn": "arn:aws:iam::123456789012:user/johndoe",  
        "accountId": "123456789012",  
        "accessKeyId": "AKIAI44QH8DHBEEXAMPLE",  
        "userName": "johndoe"  
    },  
    "eventTime": "2018-07-30T22:14:06Z",  
    "eventSource": "rds.amazonaws.com",  
    "eventName": "CreateDBInstance",  
    "awsRegion": "us-east-1",  
    "sourceIPAddress": "192.0.2.0",  
    "userAgent": "aws-cli/1.15.42 Python/3.6.1 Darwin/17.7.0 botocore/1.10.42",  
    "requestParameters": {  
        "enableCloudwatchLogsExports": [  
            "audit",  
            "error",  
            "general",  
            "slowquery"  
        ],  
        "dBInstanceIdentifier": "test-instance",  
        "engine": "mysql",  
        "masterUsername": "myawsuser",  
        "allocatedStorage": 20,  
        "dBInstanceClass": "db.m1.small",  
        "masterUserPassword": "*****"  
    },  
    "responseElements": {  
        "dBInstanceArn": "arn:aws:rds:us-east-1:123456789012:db:test-instance",  
        "storageEncrypted": false,  
        "preferredBackupWindow": "10:27-10:57",  
        "preferredMaintenanceWindow": "sat:05:47-sat:06:17",  
        "backupRetentionPeriod": 1,  
        "allocatedStorage": 20,  
        "storageType": "standard",  
        "engineVersion": "5.6.39",  
        "dBInstancePort": 0,  
        "optionGroupMemberships": [  
            {  
                "status": "in-sync",  
                "optionGroupName": "default:mysql-5-6"  
            }  
        ],  
        "dBParameterGroups": [  
            {  
                "dBParameterGroupName": "default.mysql5.6",  
                "parameterApplyStatus": "in-sync"  
            }  
        ],  
        "monitoringInterval": 0,  
        "dBInstanceClass": "db.m1.small",  
        "readReplicaDBInstanceIdentifiers": [],  
        "dBSubnetGroup": {  
            "dBSubnetGroupName": "default",  
            "dBSubnetGroupDescription": "default",  
            "subnets": [  
            ]  
        }  
    }  
}
```

```
{
    "subnetAvailabilityZone": {"name": "us-east-1b"},  

    "subnetIdentifier": "subnet-cbfff283",  

    "subnetStatus": "Active"
},  

{
    "subnetAvailabilityZone": {"name": "us-east-1e"},  

    "subnetIdentifier": "subnet-d7c825e8",  

    "subnetStatus": "Active"
},  

{
    "subnetAvailabilityZone": {"name": "us-east-1f"},  

    "subnetIdentifier": "subnet-6746046b",  

    "subnetStatus": "Active"
},  

{
    "subnetAvailabilityZone": {"name": "us-east-1c"},  

    "subnetIdentifier": "subnet-bac383e0",  

    "subnetStatus": "Active"
},  

{
    "subnetAvailabilityZone": {"name": "us-east-1d"},  

    "subnetIdentifier": "subnet-42599426",  

    "subnetStatus": "Active"
},  

{
    "subnetAvailabilityZone": {"name": "us-east-1a"},  

    "subnetIdentifier": "subnet-da327bf6",  

    "subnetStatus": "Active"
}
],
"vpcId": "vpc-136a4c6a",
"subnetGroupStatus": "Complete"
},
"masterUsername": "myawsuser",
"multiAZ": false,
"autoMinorVersionUpgrade": true,
"engine": "mysql",
"caCertificateIdentifier": "rds-ca-2015",
"dbiResourceId": "db-ETDZIIIXHEWY5N7GXVC4SH7H5IA",
"dBSecurityGroups": [],
"pendingModifiedValues": {
    "masterUserPassword": "*****",
    "pendingCloudwatchLogsExports": {
        "logTypesToEnable": [
            "audit",
            "error",
            "general",
            "slowquery"
        ]
    }
},
"dBInstanceStatus": "creating",
"publiclyAccessible": true,
"domainMemberships": [],
"copyTagsToSnapshot": false,
"dBInstanceIdentifier": "test-instance",
"licenseModel": "general-public-license",
"iamDatabaseAuthenticationEnabled": false,
"performanceInsightsEnabled": false,
"vpcSecurityGroups": [
    {
        "status": "active",
        "vpcSecurityGroupId": "sg-f839b688"
    }
]
}
```

```
},
"requestID": "daf2e3f5-96a3-4df7-a026-863f96db793e",
"eventID": "797163d3-5726-441d-80a7-6eeb7464acd4",
"eventType": "AwsApiCall",
"recipientAccountId": "123456789012"
}
```

Security in Amazon RDS

Cloud security at AWS is the highest priority. As an AWS customer, you benefit from a data center and network architecture that are built to meet the requirements of the most security-sensitive organizations.

Security is a shared responsibility between AWS and you. The [shared responsibility model](#) describes this as security of the cloud and security *in* the cloud:

- **Security of the cloud** – AWS is responsible for protecting the infrastructure that runs AWS services in the AWS Cloud. AWS also provides you with services that you can use securely. Third-party auditors regularly test and verify the effectiveness of our security as part of the [AWS compliance programs](#). To learn about the compliance programs that apply to Amazon RDS, see [AWS Services in Scope by Compliance Program](#).
- **Security in the cloud** – Your responsibility is determined by the AWS service that you use. You are also responsible for other factors including the sensitivity of your data, your organization's requirements, and applicable laws and regulations.

This documentation helps you understand how to apply the shared responsibility model when using Amazon RDS. The following topics show you how to configure Amazon RDS to meet your security and compliance objectives. You also learn how to use other AWS services that help you monitor and secure your Amazon RDS resources.

You can manage access to your Amazon RDS resources and your databases on a DB instance. The method you use to manage access depends on what type of task the user needs to perform with Amazon RDS:

- Run your DB instance in a virtual private cloud (VPC) based on the Amazon VPC service for the greatest possible network access control. For more information about creating a DB instance in a VPC, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).
- Use AWS Identity and Access Management (IAM) policies to assign permissions that determine who is allowed to manage Amazon RDS resources. For example, you can use IAM to determine who is allowed to create, describe, modify, and delete DB instances, tag resources, or modify security groups.
- Use security groups to control what IP addresses or Amazon EC2 instances can connect to your databases on a DB instance. When you first create a DB instance, its firewall prevents any database access except through rules specified by an associated security group.
- Use Secure Socket Layer (SSL) or Transport Layer Security (TLS) connections with DB instances running the MySQL, MariaDB, PostgreSQL, Oracle, or Microsoft SQL Server database engines. For more information on using SSL/TLS with a DB instance, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).
- Use Amazon RDS encryption to secure your DB instances and snapshots at rest. Amazon RDS encryption uses the industry standard AES-256 encryption algorithm to encrypt your data on the server that hosts your DB instance. For more information, see [Encrypting Amazon RDS Resources \(p. 429\)](#).
- Use network encryption and transparent data encryption with Oracle DB instances; for more information, see [Oracle Native Network Encryption \(p. 1049\)](#) and [Oracle Transparent Data Encryption \(p. 1074\)](#)
- Use the security features of your DB engine to control who can log in to the databases on a DB instance. These features work just as if the database was on your local network.

Note

You only have to configure security for your use cases. You don't have to configure security access for processes that Amazon RDS manages. These include creating backups, replicating data between a master and a Read Replica, and other processes.

For more information on managing access to Amazon RDS resources and your databases on a DB instance, see the following topics.

Topics

- [Data Protection in Amazon RDS \(p. 428\)](#)
- [Identity and Access Management in Amazon RDS \(p. 445\)](#)
- [Kerberos Authentication \(p. 480\)](#)
- [Logging and Monitoring in Amazon RDS \(p. 480\)](#)
- [Compliance Validation for Amazon RDS \(p. 483\)](#)
- [Resilience in Amazon RDS \(p. 484\)](#)
- [Infrastructure Security in Amazon RDS \(p. 485\)](#)
- [Security Best Practices for Amazon RDS \(p. 485\)](#)
- [Controlling Access with Security Groups \(p. 486\)](#)
- [Master User Account Privileges \(p. 499\)](#)
- [Using Service-Linked Roles for Amazon RDS \(p. 501\)](#)
- [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#)

Data Protection in Amazon RDS

Amazon RDS conforms to the AWS [shared responsibility model](#), which includes regulations and guidelines for data protection. AWS is responsible for protecting the global infrastructure that runs all the AWS services. AWS maintains control over data hosted on this infrastructure, including the security configuration controls for handling customer content and personal data. AWS customers and APN partners, acting either as data controllers or data processors, are responsible for any personal data that they put in the AWS Cloud.

For data protection, we recommend that you protect AWS account credentials and set up principals with AWS Identity and Access Management (IAM). Doing this means that each user is given only the permissions necessary to fulfill their job duties. We also recommend that you secure your data in the following ways:

- Use multi-factor authentication (MFA) with each account.
- Use SSL/TLS to communicate with AWS resources.
- Set up API and user activity logging with AWS CloudTrail.
- Use AWS encryption solutions, along with all default security controls within AWS services.
- Use advanced managed security services such as Amazon Macie, which assists in discovering and securing personal data that is stored in Amazon S3.

We strongly recommend that you never put sensitive identifying information, such as your customers' account numbers, into free-form fields such as a **Name** field. This recommendation includes when you work with Amazon RDS or other AWS services using the console, API, AWS CLI, or AWS SDKs. Any data that you enter into Amazon RDS or other services might get picked up for inclusion in diagnostic logs. When you provide a URL to an external server, don't include credentials information in the URL to validate your request to that server.

For more information about data protection, see the [AWS Shared Responsibility Model and GDPR blog post](#) on the [AWS Security Blog](#).

Topics

- [Protecting Data Using Encryption \(p. 429\)](#)

- [Internetwork Traffic Privacy \(p. 444\)](#)

Protecting Data Using Encryption

You can enable encryption for database resources. You can also encrypt connections to DB instances.

Topics

- [Encrypting Amazon RDS Resources \(p. 429\)](#)
- [Key Management \(p. 431\)](#)
- [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#)
- [Rotating Your SSL/TLS Certificate \(p. 433\)](#)

Encrypting Amazon RDS Resources

You can encrypt your Amazon RDS DB instances and snapshots at rest by enabling the encryption option for your Amazon RDS DB instances. Data that is encrypted at rest includes the underlying storage for DB instances, its automated backups, Read Replicas, and snapshots.

Amazon RDS encrypted DB instances use the industry standard AES-256 encryption algorithm to encrypt your data on the server that hosts your Amazon RDS DB instances. After your data is encrypted, Amazon RDS handles authentication of access and decryption of your data transparently with a minimal impact on performance. You don't need to modify your database client applications to use encryption.

Note

For encrypted and unencrypted DB instances, data that is in transit between the source and the Read Replicas is encrypted, even when replicating across AWS Regions.

Topics

- [Overview of Encrypting Amazon RDS Resources \(p. 429\)](#)
- [Enabling Amazon RDS Encryption for a DB Instance \(p. 430\)](#)
- [Availability of Amazon RDS Encryption \(p. 430\)](#)
- [Limitations of Amazon RDS Encrypted DB Instance \(p. 431\)](#)

Overview of Encrypting Amazon RDS Resources

Amazon RDS encrypted DB instances provide an additional layer of data protection by securing your data from unauthorized access to the underlying storage. You can use Amazon RDS encryption to increase data protection of your applications deployed in the cloud, and to fulfill compliance requirements for data-at-rest encryption.

Amazon RDS also supports encrypting an Oracle or SQL Server DB instance with Transparent Data Encryption (TDE). TDE can be used with encryption at rest, although using TDE and encryption at rest simultaneously might slightly affect the performance of your database. You must manage different keys for each encryption method. For more information on TDE, see [Oracle Transparent Data Encryption \(p. 1074\)](#) or [Support for Transparent Data Encryption in SQL Server \(p. 714\)](#).

To manage the keys used for encrypting and decrypting your Amazon RDS resources, you use the [AWS Key Management Service \(AWS KMS\)](#). AWS KMS combines secure, highly available hardware and software to provide a key management system scaled for the cloud. Using AWS KMS, you can create encryption keys and define the policies that control how these keys can be used. AWS KMS supports CloudTrail, so you can audit key usage to verify that keys are being used appropriately. You can use your AWS KMS keys with Amazon RDS and supported AWS services such as Amazon S3, Amazon EBS, and Amazon Redshift. For a list of services that support AWS KMS, see [Supported Services](#) in the [AWS Key Management Service Developer Guide](#).

For an Amazon RDS encrypted DB instance, all logs, backups, and snapshots are encrypted. A Read Replica of an Amazon RDS encrypted instance is also encrypted using the same key as the master instance when both are in the same AWS Region. If the master and Read Replica are in different AWS Regions, you encrypt using the encryption key for that AWS Region.

Enabling Amazon RDS Encryption for a DB Instance

To enable encryption for a new DB instance, choose **Enable encryption** on the Amazon RDS console. For information on creating a DB instance, see one of the following topics:

- [Creating a DB Instance Running the MySQL Database Engine \(p. 751\)](#)
- [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#)
- [Creating a DB Instance Running the Microsoft SQL Server Database Engine \(p. 616\)](#)
- [Creating a DB Instance Running the PostgreSQL Database Engine \(p. 1257\)](#)
- [Creating a DB Instance Running the MariaDB Database Engine \(p. 543\)](#)

If you use the `create-db-instance` AWS CLI command to create an encrypted DB instance, set the `--storage-encrypted` parameter to true. If you use the `CreateDBInstance` API operation, set the `StorageEncrypted` parameter to true.

When you create an encrypted DB instance, you can also supply the AWS KMS key identifier for your encryption key. If you don't specify an AWS KMS key identifier, then Amazon RDS uses your default encryption key for your new DB instance. AWS KMS creates your default encryption key for Amazon RDS for your AWS account. Your AWS account has a different default encryption key for each AWS Region.

Once you have created an encrypted DB instance, you can't change the type of encryption key used by that DB instance. Therefore, be sure to determine your encryption key requirements before you create your encrypted DB instance.

If you use the AWS CLI `create-db-instance` command to create an encrypted DB instance, set the `--kms-key-id` parameter to the Amazon Resource Name (ARN) for the AWS KMS encryption key for the DB instance. If you use the Amazon RDS API `CreateDBInstance` action, set the `KmsKeyId` parameter to the ARN for your AWS KMS key for the DB instance.

You can use the ARN of a key from another account to encrypt a DB instance. Or you might create a DB instance with the same AWS account that owns the AWS KMS encryption key used to encrypt that new DB instance. In this case, the AWS KMS key ID that you pass can be the AWS KMS key alias instead of the key's ARN.

Availability of Amazon RDS Encryption

Amazon RDS encryption is currently available for all database engines and storage types.

Amazon RDS encryption is available for most DB instance classes. The following table lists DB instance classes that *do not support* Amazon RDS encryption:

Instance Type	Instance Class
General Purpose (M1)	db.m1.small
	db.m1.medium
	db.m1.large
	db.m1.xlarge

Instance Type	Instance Class
Memory Optimized (M2)	db.m2.xlarge
	db.m2.2xlarge
	db.m2.4xlarge
Burst Capable (T2)	db.t2.micro

Note

Encryption at rest is not available for DB instances running SQL Server Express Edition.

Limitations of Amazon RDS Encrypted DB Instance

The following limitations exist for Amazon RDS encrypted DB instance:

- You can only enable encryption for an Amazon RDS DB instance when you create it, not after the DB instance is created.

However, because you can encrypt a copy of an unencrypted DB snapshot, you can effectively add encryption to an unencrypted DB instance. That is, you can create a snapshot of your DB instance, and then create an encrypted copy of that snapshot. You can then restore a DB instance from the encrypted snapshot, and thus you have an encrypted copy of your original DB instance. For more information, see [Copying a Snapshot \(p. 269\)](#).

- DB instances that are encrypted can't be modified to disable encryption.
- You can't have an encrypted Read Replica of an unencrypted DB instance or an unencrypted Read Replica of an encrypted DB instance.
- Encrypted Read Replicas must be encrypted with the same key as the source DB instance.
- You can't restore an unencrypted backup or snapshot to an encrypted DB instance.
- To copy an encrypted snapshot from one AWS Region to another, you must specify the KMS key identifier of the destination AWS Region. This is because KMS encryption keys are specific to the AWS Region that they are created in.

The source snapshot remains encrypted throughout the copy process. AWS Key Management Service uses envelope encryption to protect data during the copy process. For more information about envelope encryption, see [Envelope Encryption](#).

Key Management

You can manage keys used for Amazon RDS encrypted DB instances using the [AWS Key Management Service \(AWS KMS\)](#) in the IAM console. If you want full control over a key, then you must create a customer-managed key.

You can't delete, revoke, or rotate default keys provisioned by AWS KMS. You can't share a snapshot that has been encrypted using the default AWS KMS encryption key of the AWS account that shared the snapshot.

You can view audit logs of every action taken with a customer-managed key by using [AWS CloudTrail](#).

Important

When RDS encounters a DB instance encrypted by a key that RDS doesn't have access to, RDS puts the DB instance into a terminal state. In this state, the DB instance is no longer available and the current state of the database can't be recovered. To restore the DB instance, you must re-enable access to the encryption key for RDS, and then restore the DB instance from a backup.

Using SSL/TLS to Encrypt a Connection to a DB Instance

You can use Secure Socket Layer (SSL) or Transport Layer Security (TLS) from your application to encrypt a connection to a DB instance running MySQL, MariaDB, SQL Server, Oracle, or PostgreSQL. Each DB engine has its own process for implementing SSL/TLS. To learn how to implement SSL/TLS for your DB instance, use the link following that corresponds to your DB engine:

- [Using SSL with a MariaDB DB Instance \(p. 539\)](#)
- [Using SSL with a Microsoft SQL Server DB Instance \(p. 687\)](#)
- [Using SSL with a MySQL DB Instance \(p. 746\)](#)
- [Using SSL with an Oracle DB Instance \(p. 921\)](#)
- [Using SSL with a PostgreSQL DB Instance \(p. 1411\)](#)

Important

For information about rotating your certificate, see [Rotating Your SSL/TLS Certificate \(p. 433\)](#).

Note

All certificates are only available for download using SSL/TLS connections.

To get a root certificate that works for all AWS Regions, download from one of these locations:

- <https://s3.amazonaws.com/rds-downloads/rds-ca-2019-root.pem> (CA-2019)
- <https://s3.amazonaws.com/rds-downloads/rds-ca-2015-root.pem> (CA-2015)

This root certificate is a trusted root entity and should work in most cases but might fail if your application doesn't accept certificate chains. If your application doesn't accept certificate chains, download the AWS Region-specific certificate from the list of intermediate certificates found later in this section.

To get a certificate bundle that contains both the intermediate and root certificates, download from <https://s3.amazonaws.com/rds-downloads/rds-combined-ca-bundle.pem>.

If your application is on Microsoft Windows and requires a PKCS7 file, you can download the PKCS7 certificate bundle. This bundle contains both the intermediate and root certificates at <https://s3.amazonaws.com/rds-downloads/rds-combined-ca-bundle.p7b>.

Intermediate Certificates

You might need to use an intermediate certificate to connect to your AWS Region. For example, you must use an intermediate certificate to connect to the AWS GovCloud (US-West) Region using SSL/TLS. If you need an intermediate certificate for a particular AWS Region, download the certificate from the following table.

AWS Region	CA-2019	CA-2015
Asia Pacific (Hong Kong)	rds-ca-2019-ap-east-1.pem	No certificate available
Asia Pacific (Mumbai)	rds-ca-2019-ap-south-1.pem	rds-ca-2015-ap-south-1.pem
Asia Pacific (Tokyo)	rds-ca-2019-ap-northeast-1.pem	rds-ca-2015-ap-northeast-1.pem
Asia Pacific (Seoul)	rds-ca-2019-ap-northeast-2.pem	rds-ca-2015-ap-northeast-2.pem

AWS Region	CA-2019	CA-2015
Asia Pacific (Osaka-Local)	rds-ca-2019-ap-northeast-3.pem	rds-ca-2015-ap-northeast-3.pem
Asia Pacific (Singapore)	rds-ca-2019-ap-southeast-1.pem	rds-ca-2015-ap-southeast-1.pem
Asia Pacific (Sydney)	rds-ca-2019-ap-southeast-2.pem	rds-ca-2015-ap-southeast-2.pem
Canada (Central)	rds-ca-2019-ca-central-1.pem	rds-ca-2015-ca-central-1.pem
Europe (Frankfurt)	rds-ca-2019-eu-central-1.pem	rds-ca-2015-eu-central-1.pem
Europe (Ireland)	rds-ca-2019-eu-west-1.pem	rds-ca-2015-eu-west-1.pem
Europe (London)	rds-ca-2019-eu-west-2.pem	rds-ca-2015-eu-west-2.pem
Europe (Paris)	rds-ca-2019-eu-west-3.pem	rds-ca-2015-eu-west-3.pem
Europe (Stockholm)	rds-ca-2019-eu-north-1.pem	rds-ca-2015-eu-north-1.pem
Middle East (Bahrain)	rds-ca-2019-me-south-1.pem	No certificate available
South America (São Paulo)	rds-ca-2019-sa-east-1.pem	rds-ca-2015-sa-east-1.pem
US East (N. Virginia)	rds-ca-2019-us-east-1.pem	rds-ca-2015-us-east-1.pem
US East (Ohio)	rds-ca-2019-us-east-2.pem	rds-ca-2015-us-east-2.pem
US West (N. California)	rds-ca-2019-us-west-1.pem	rds-ca-2015-us-west-1.pem
US West (Oregon)	rds-ca-2019-us-west-2.pem	rds-ca-2015-us-west-2.pem

AWS GovCloud (US) Certificates

You can download a root certificate for the AWS GovCloud (US-West) Region at <https://s3-us-gov-west-1.amazonaws.com/rds-downloads/rds-GovCloud-Root-CA-2017.pem>.

To get a certificate bundle that contains both the intermediate and root certificates for the AWS GovCloud (US) Regions, download from <https://s3-us-gov-west-1.amazonaws.com/rds-downloads/rds-combined-ca-us-gov-bundle.pem>.

You can download the intermediate certificate for an AWS GovCloud (US) Region from the following list:

[AWS GovCloud \(US-East\) \(CA-2017\)](#)

[AWS GovCloud \(US-West\) \(CA-2017\)](#)

[AWS GovCloud \(US-West\) \(CA-2012\)](#)

Rotating Your SSL/TLS Certificate

Note

If your application connects to an RDS DB instance using Secure Socket Layer (SSL) or Transport Layer Security (TLS), you must take the following steps *before March 5, 2020*. Doing this means you can avoid interruption of connectivity between your applications and your RDS DB instances.

As of September 19, 2019, Amazon RDS has published new Certificate Authority (CA) certificates for connecting to your RDS DB instances using SSL/TLS. We provide these new CA certificates as an AWS security best practice. For information about the new certificates and the supported AWS Regions, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).

The current CA certificates expire on March 5, 2020. Therefore, we strongly recommend completing this change as soon as possible (and no later than February 5, 2020), to avoid disruption on the expiration date. If the change is not completed, your applications will fail to connect to your RDS DB instances using SSL/TLS after March 5, 2020.

We encourage you to test the steps listed following in a development or staging environment before taking them for your production environments.

Before you update your DB instances to use the new CA certificate, make sure that you update your clients or applications connecting to your RDS databases.

Any new RDS DB instances created after January 14, 2020 will use the new certificates by default. If you want to temporarily modify new DB instances manually to use the old (rds-ca-2015) certificates, you can do so using the AWS Management Console or the AWS CLI. Any DB instances created prior to January 14, 2020 use the rds-ca-2015 certificates until you update them to the rds-ca-2019 certificates.

Topics

- [Updating Your CA Certificate by Modifying Your DB Instance \(p. 434\)](#)
- [Updating Your CA Certificate by Applying DB Instance Maintenance \(p. 436\)](#)
- [Reverting an Update of a CA Certificate \(p. 442\)](#)
- [Sample Script for Importing Certificates Into Your Trust Store \(p. 443\)](#)

Updating Your CA Certificate by Modifying Your DB Instance

Complete the following steps to update your CA certificate by modifying your DB instance.

To update your CA certificate by modifying your DB instance

1. Download the new SSL/TLS certificate as described in [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).
2. Update your database applications to use the new SSL/TLS certificate.

The methods for updating applications for new SSL/TLS certificates depend on your specific applications. Work with your application developers to update the SSL/TLS certificates for your applications.

For information about checking for SSL/TLS connections and updating applications for each DB engine, see the following topics:

- [Updating Applications to Connect to MariaDB DB Instances Using New SSL/TLS Certificates \(p. 560\)](#)
- [Updating Applications to Connect to Microsoft SQL Server DB Instances Using New SSL/TLS Certificates \(p. 639\)](#)
- [Updating Applications to Connect to MySQL DB Instances Using New SSL/TLS Certificates \(p. 797\)](#)
- [Updating Applications to Connect to Oracle DB Instances Using New SSL/TLS Certificates \(p. 962\)](#)
- [Updating Applications to Connect to PostgreSQL DB Instances Using New SSL/TLS Certificates \(p. 1275\)](#)

For a sample script that updates a trust store for a Linux operating system, see [Sample Script for Importing Certificates Into Your Trust Store \(p. 443\)](#).

Note

The certificate bundle contains certificates for both the old and new CA, so you can upgrade your application safely and maintain connectivity during the transition period.

3. Modify the DB instance to change the CA from **rds-ca-2015** to **rds-ca-2019**.

Important

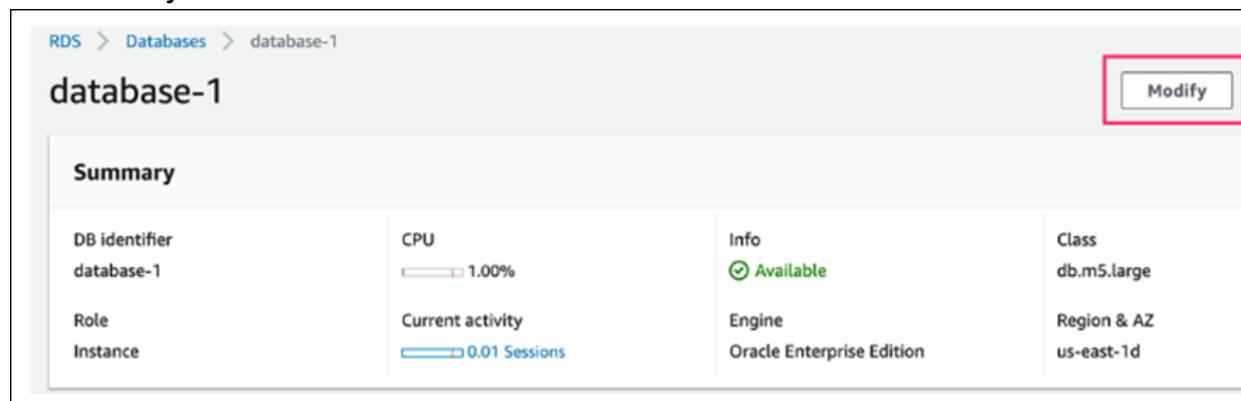
This operation reboots your DB instance. By default, this operation is scheduled to run during your next maintenance window. Or you can choose to run it immediately.

You can use the AWS Management Console or the AWS CLI to change the CA certificate from **rds-ca-2015** to **rds-ca-2019** for a DB instance.

Console

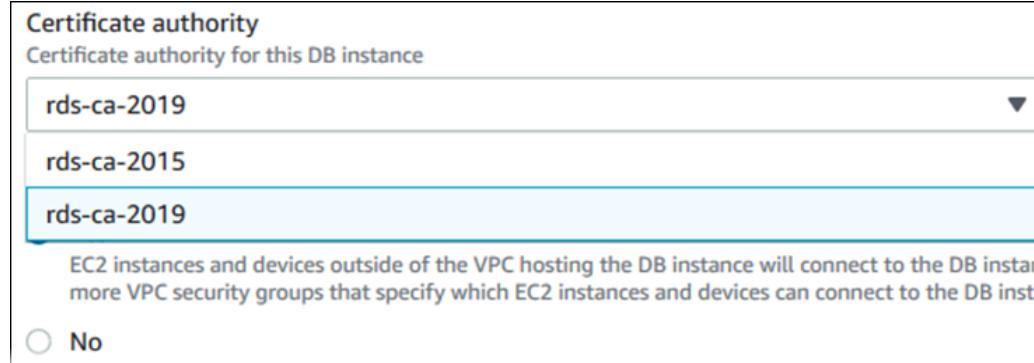
To change the CA from rds-ca-2015 to rds-ca-2019 for a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to modify.
3. Choose **Modify**.



The **Modify DB Instance** page appears.

4. In the **Network & Security** section, choose **rds-ca-2019**.



5. Choose **Continue** and check the summary of modifications.

6. To apply the changes immediately, choose **Apply immediately**.

Important

Choosing this option causes an outage.

7. On the confirmation page, review your changes. If they are correct, choose **Modify DB Instance** to save your changes.

Important

When you schedule this operation, make sure that you have updated your client-side trust store beforehand.

Or choose **Back** to edit your changes or **Cancel** to cancel your changes.

AWS CLI

To use the AWS CLI to change the CA from **rds-ca-2015** to **rds-ca-2019** for a DB instance, call the [modify-db-instance](#) command. Specify the DB instance identifier and the `--ca-certificate-identifier` option.

Important

When you schedule this operation, make sure that you have updated your client-side trust store beforehand.

Example

The following code modifies `mydbinstance` by setting the CA certificate to `rds-ca-2019`. The changes are applied during the next maintenance window by using `--no-apply-immediately`. Use `--apply-immediately` to apply the changes immediately.

Important

Using the `--apply-immediately` option causes an outage.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier mydbinstance \
--ca-certificate-identifier rds-ca-2019 \
--no-apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier mydbinstance ^
--ca-certificate-identifier rds-ca-2019 ^
--no-apply-immediately
```

Updating Your CA Certificate by Applying DB Instance Maintenance

Complete the following steps to update your CA certificate by applying DB instance maintenance.

To update your CA certificate by applying DB instance maintenance

1. Download the new SSL/TLS certificate as described in [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).
2. Update your database applications to use the new SSL/TLS certificate.

The methods for updating applications for new SSL/TLS certificates depend on your specific applications. Work with your application developers to update the SSL/TLS certificates for your applications.

For information about checking for SSL/TLS connections and updating applications for each DB engine, see the following topics:

- [Updating Applications to Connect to MariaDB DB Instances Using New SSL/TLS Certificates \(p. 560\)](#)
- [Updating Applications to Connect to Microsoft SQL Server DB Instances Using New SSL/TLS Certificates \(p. 639\)](#)
- [Updating Applications to Connect to MySQL DB Instances Using New SSL/TLS Certificates \(p. 797\)](#)
- [Updating Applications to Connect to Oracle DB Instances Using New SSL/TLS Certificates \(p. 962\)](#)
- [Updating Applications to Connect to PostgreSQL DB Instances Using New SSL/TLS Certificates \(p. 1275\)](#)

For a sample script that updates a trust store for a Linux operating system, see [Sample Script for Importing Certificates Into Your Trust Store \(p. 443\)](#).

Note

The certificate bundle contains certificates for both the old and new CA, so you can upgrade your application safely and maintain connectivity during the transition period.

3. Apply DB instance maintenance to change the CA from **rds-ca-2015** to **rds-ca-2019**.

Important

This operation reboots your DB instance. By default, this operation is scheduled to run during your next maintenance window. Or you can choose to run it immediately.

You can use the AWS Management Console to apply DB instance maintenance to change the CA certificate from **rds-ca-2015** to **rds-ca-2019** for a single DB instance or for multiple DB instances.

Topics

- [Updating Your CA Certificate by Applying Maintenance to a Single DB Instance \(p. 437\)](#)
- [Updating Your CA Certificate by Applying Maintenance to Multiple DB Instances \(p. 438\)](#)

[Updating Your CA Certificate by Applying Maintenance to a Single DB Instance](#)

Use the AWS Management Console to change the CA certificate for a single DB instance.

To change the CA from rds-ca-2015 to rds-ca-2019 for a single DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the name of the DB instance that you want to modify.
3. Choose **Maintenance & backups**.

If your DB instance is using the old CA certificate, the **Pending maintenance** section shows an action with the description **Rotation of CA certificate**. This pending maintenance action is scheduled by default for your maintenance window and before February 5, 2020. However, you can apply the rotation immediately by choosing the pending maintenance action and choosing **Apply now**.

Important

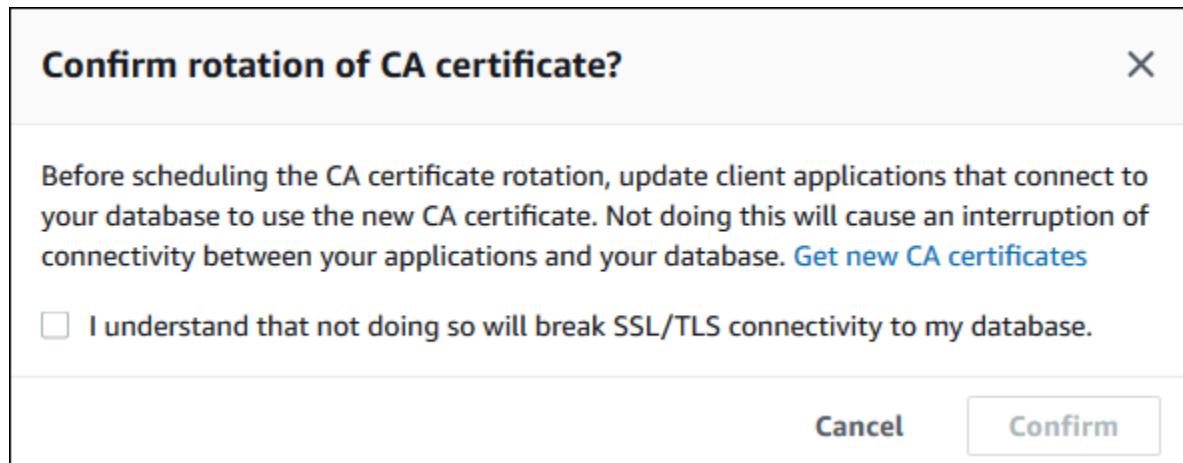
When your CA certificate is rotated, the operation reboots your DB instance.

The screenshot shows the 'Maintenance' section of the AWS RDS console. At the top, there are tabs for Connectivity & security, Monitoring, Logs & events, Configuration, Maintenance & backups (which is selected), and Tags. Below the tabs, there's a summary row with four items: Auto minor version upgrade (Enabled), Maintenance window (fri:09:09-fri:09:39 UTC (GMT)), Pending maintenance next window, and Pending modifications. The main area is titled 'Pending maintenance (1)' and lists a single item: 'Rotation of CA certificate' (Type: ca-certificate-rotation, Status: next window, Apply date: February 4th 2020, 1:54:04 pm UTC-8 (local)). There are buttons for 'Apply now' (circled in red) and 'Apply at next maintenance window'. A search bar and navigation controls are also present.

4. If you choose **Apply now** or **Apply at next maintenance window**, you are prompted to confirm the CA certificate rotation.

Important

Before scheduling the CA certificate rotation on your database, update any client applications that use SSL/TLS and the server certificate to connect. These updates are specific to your DB engine. To determine whether your applications use SSL/TLS and the server certificate to connect, see [Step 2: Update Your Database Applications to Use the New SSL/TLS Certificate \(p. 436\)](#). After you have updated these client applications, you can confirm the CA certificate rotation.



To continue, choose the check box, and then choose **Confirm**.

[Updating Your CA Certificate by Applying Maintenance to Multiple DB Instances](#)

Use the AWS Management Console to change the CA certificate for multiple DB instances.

To change the CA from rds-ca-2015 to rds-ca-2019 for multiple DB instances

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.

If you have at least one DB instance that is using the old CA certificate, the following banner appears at the top of the page.

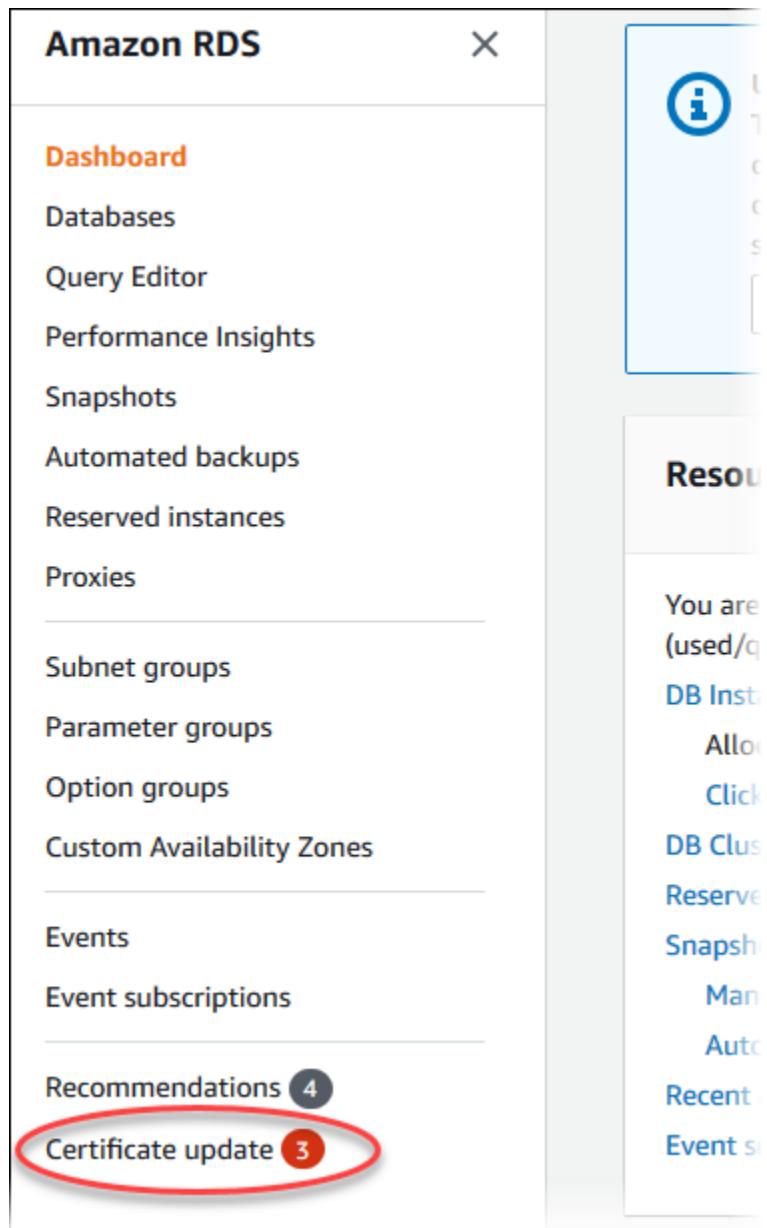


Update your Amazon RDS SSL/TLS certificates before March 5, 2020

To avoid interruption of your applications using RDS and Aurora databases, update the Certificate Authority (CA) certificates for these databases before March 5, 2020. We strongly recommend making your updates before February 5, 2020, to leave time for deployments, testing, and validation. New databases created after January 14, 2020, will default to using the new CA certificates. Make sure that you update your client applications with the new certificates first. Find the new CA certificates and info: [RDS](#) [Aurora](#)

[View pending maintenance actions](#)

In the navigation pane, there is also a **Certificate update** option that shows the total number of affected DB instances.



Either choose **View pending maintenance actions** in the banner, or choose **Certificate update** in the navigation pane.

The **Update your Amazon RDS SSL/TLS certificates** page appears.

DB identifier	DB cluster identifier	Status	Apply date
mydbinstancecf	-	Requires Update	-
mydbinstancecf2	-	Requires Update	-
oracledb	-	Requires Update	-

Note

This page only shows the DB instances for the current AWS Region. If you have DB instances in more than one AWS Region, check this page in each AWS Region to see all DB instances with old SSL/TLS certificates.

3. Choose the DB instance you want to update.

You can schedule the certificate rotation for your next maintenance window by choosing **Update at the next maintenance window**. Or you can apply the rotation immediately by choosing **Update now**.

Important

When your CA certificate is rotated, the operation reboots your DB instance.

4. If you choose **Update at the next maintenance window** or **Update now**, you are prompted to confirm the CA certificate rotation.

Important

Before scheduling the CA certificate rotation on your database, update any client applications that use SSL/TLS and the server certificate to connect. These updates are specific to your DB engine. To determine whether your applications use SSL/TLS and the server certificate to connect, see [Step 2: Update Your Database Applications to Use the New SSL/TLS Certificate \(p. 436\)](#). After you have updated these client applications, you can confirm the CA certificate rotation.

Confirm rotation of CA certificate?

Before scheduling the CA certificate rotation, update client applications that connect to your database to use the new CA certificate. Not doing this will cause an interruption of connectivity between your applications and your database. [Get new CA certificates](#)

I understand that not doing so will break SSL/TLS connectivity to my database.

Cancel

Confirm

To continue, choose the check box, and then choose **Confirm**.

5. Repeat steps 3 and 4 for each DB instance that you want to update.

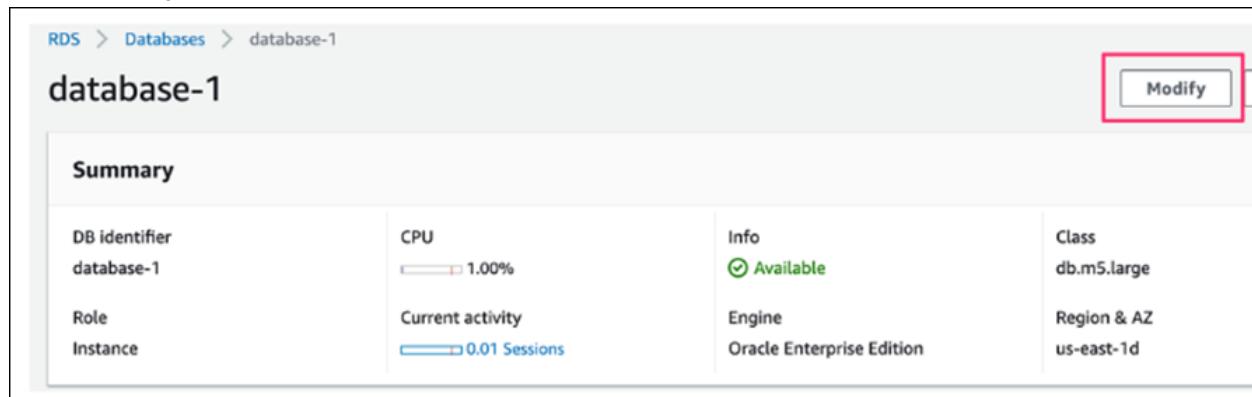
Reverting an Update of a CA Certificate

You can use the AWS Management Console or the AWS CLI to revert to a previous CA certificate for a DB instance.

Console

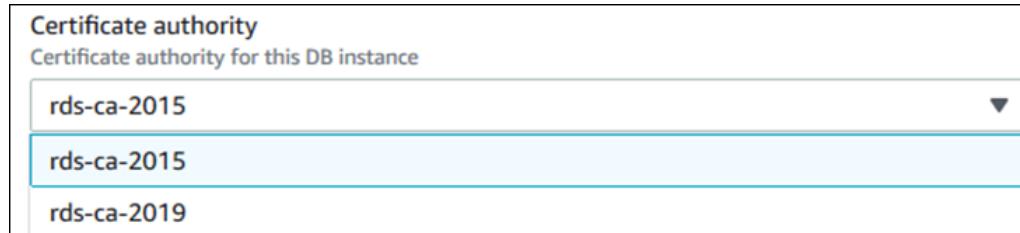
To revert to a previous CA certificate for a DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to modify.
3. Choose **Modify**.



The **Modify DB Instance** page appears.

4. In the **Network & Security** section, choose **rds-ca-2015**.



No

5. Choose **Continue** and check the summary of modifications.
6. To apply the changes immediately, choose **Apply immediately**.

Important

Choosing this option causes an outage.

7. On the confirmation page, review your changes. If they are correct, choose **Modify DB Instance** to save your changes.

Important

When you schedule this operation, make sure that you have updated your client-side trust store beforehand.

Or choose **Back** to edit your changes or **Cancel** to cancel your changes.

AWS CLI

To revert to a previous CA certificate for a DB instance, call the [modify-db-instance](#) command. Specify the DB instance identifier and the --ca-certificate-identifier option.

Important

When you schedule this operation, make sure that you have updated your client-side trust store beforehand.

Example

The following code example modifies `mydbinstance` by setting the CA certificate to `rds-ca-2015`. The changes are applied during the next maintenance window by using `--no-apply-immediately`. Use `--apply-immediately` to apply the changes immediately.

Important

Using the `--apply-immediately` option causes an outage.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier mydbinstance \
--ca-certificate-identifier rds-ca-2015 \
--no-apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier mydbinstance ^
--ca-certificate-identifier rds-ca-2015 ^
--no-apply-immediately
```

Sample Script for Importing Certificates Into Your Trust Store

The following is a sample shell script that imports the certificate bundle into a trust store on a Linux operating system.

```
mydir=/tmp/certs
truststore=${mydir}/rds-truststore.jks
storepassword=changeit

curl -sS "https://s3.amazonaws.com/rds-downloads/rds-combined-ca-bundle.pem" > ${mydir}/rds-combined-ca-bundle.pem
awk 'split_after == 1 {n++;split_after=0} /-----END CERTIFICATE-----/ {split_after=1}{print > "rds-ca-" n ".pem"}' < ${mydir}/rds-combined-ca-bundle.pem

for CERT in rds-ca-*; do
    alias=$(openssl x509 -noout -text -in $CERT | perl -ne 'next unless /Subject:/; s/.+CN=//; print')
    echo "Importing $alias"
    keytool -import -file ${CERT} -alias "${alias}" -storepass ${storepassword} -keystore ${truststore} -noprompt
    rm $CERT
done

rm ${mydir}/rds-combined-ca-bundle.pem

echo "Trust store content is: "

keytool -list -v -keystore "$truststore" -storepass ${storepassword} | grep Alias | cut -d " " -f3- | while read alias
```

```
do
    expiry=`keytool -list -v -keystore "$truststore" -storepass ${storepassword} -alias
"${alias}" | grep Valid | perl -ne 'if(/until: (.*)\n/) { print "$1\n"; }'`^
        echo " Certificate ${alias} expires in '$expiry'"
done
```

Internetwork Traffic Privacy

Connections are protected both between Amazon RDS and on-premises applications and between Amazon RDS and other AWS resources within the same AWS Region.

Traffic Between Service and On-Premises Clients and Applications

You have two connectivity options between your private network and AWS:

- An AWS Site-to-Site VPN connection. For more information, see [What is AWS Site-to-Site VPN?](#)
- An AWS Direct Connect connection. For more information, see [What is AWS Direct Connect?](#)

You get access to Amazon RDS through the network by using AWS-published API operations. Clients must support Transport Layer Security (TLS) 1.0. We recommend TLS 1.2. Clients must also support cipher suites with Perfect Forward Secrecy (PFS), such as Ephemeral Diffie-Hellman (DHE) or Elliptic Curve Diffie-Hellman Ephemeral (ECDHE). Most modern systems such as Java 7 and later support these modes. Additionally, you must sign requests using an access key ID and a secret access key that are associated with an IAM principal. Or you can use the [AWS Security Token Service \(STS\)](#) to generate temporary security credentials to sign requests.

Traffic Between AWS Resources in the Same Region

An Amazon Virtual Private Cloud (Amazon VPC) endpoint for Amazon RDS is a logical entity within a VPC that allows connectivity only to Amazon RDS. The Amazon VPC routes requests to Amazon RDS and routes responses back to the VPC. For more information, see [VPC Endpoints](#) in the *Amazon VPC User Guide*. For sample policies that you can use to control DB instance access from VPC endpoints, see [Creating and Using an IAM Policy for IAM Database Access \(p. 464\)](#).

Identity and Access Management in Amazon RDS

AWS Identity and Access Management (IAM) is an AWS service that helps an administrator securely control access to AWS resources. IAM administrators control who can be *authenticated* (signed in) and *authorized* (have permissions) to use Amazon RDS resources. IAM is an AWS service that you can use with no additional charge.

Topics

- [Audience \(p. 445\)](#)
- [Authenticating With Identities \(p. 445\)](#)
- [Managing Access Using Policies \(p. 447\)](#)
- [How Amazon RDS Works with IAM \(p. 449\)](#)
- [Amazon RDS Identity-Based Policy Examples \(p. 451\)](#)
- [IAM Database Authentication for MySQL and PostgreSQL \(p. 461\)](#)
- [Troubleshooting Amazon RDS Identity and Access \(p. 478\)](#)

Audience

How you use AWS Identity and Access Management (IAM) differs, depending on the work you do in Amazon RDS.

Service user – If you use the Amazon RDS service to do your job, then your administrator provides you with the credentials and permissions that you need. As you use more Amazon RDS features to do your work, you might need additional permissions. Understanding how access is managed can help you request the right permissions from your administrator. If you cannot access a feature in Amazon RDS, see [Troubleshooting Amazon RDS Identity and Access \(p. 478\)](#).

Service administrator – If you're in charge of Amazon RDS resources at your company, you probably have full access to Amazon RDS. It's your job to determine which Amazon RDS features and resources your employees should access. You must then submit requests to your IAM administrator to change the permissions of your service users. Review the information on this page to understand the basic concepts of IAM. To learn more about how your company can use IAM with Amazon RDS, see [How Amazon RDS Works with IAM \(p. 449\)](#).

IAM administrator – If you're an IAM administrator, you might want to learn details about how you can write policies to manage access to Amazon RDS. To view example Amazon RDS identity-based policies that you can use in IAM, see [Amazon RDS Identity-Based Policy Examples \(p. 451\)](#).

Authenticating With Identities

Authentication is how you sign in to AWS using your identity credentials. For more information about signing in using the AWS Management Console, see [The IAM Console and Sign-in Page](#) in the *IAM User Guide*.

You must be *authenticated* (signed in to AWS) as the AWS account root user, an IAM user, or by assuming an IAM role. You can also use your company's single sign-on authentication, or even sign in using Google or Facebook. In these cases, your administrator previously set up identity federation using IAM roles. When you access AWS using credentials from another company, you are assuming a role indirectly.

To sign in directly to the [AWS Management Console](#), use your password with your root user email or your IAM user name. You can access AWS programmatically using your root user or IAM user access keys. AWS provides SDK and command line tools to cryptographically sign your request using your credentials. If you don't use AWS tools, you must sign the request yourself. Do this using *Signature Version 4*, a protocol

for authenticating inbound API requests. For more information about authenticating requests, see [Signature Version 4 Signing Process](#) in the *AWS General Reference*.

Regardless of the authentication method that you use, you might also be required to provide additional security information. For example, AWS recommends that you use multi-factor authentication (MFA) to increase the security of your account. To learn more, see [Using Multi-Factor Authentication \(MFA\) in AWS](#) in the *IAM User Guide*.

AWS Account Root User

When you first create an AWS account, you begin with a single sign-in identity that has complete access to all AWS services and resources in the account. This identity is called the *AWS account root user* and is accessed by signing in with the email address and password that you used to create the account. We strongly recommend that you do not use the root user for your everyday tasks, even the administrative ones. Instead, adhere to the [best practice of using the root user only to create your first IAM user](#). Then securely lock away the root user credentials and use them to perform only a few account and service management tasks.

IAM Users and Groups

An *IAM user* is an identity within your AWS account that has specific permissions for a single person or application. An IAM user can have long-term credentials such as a user name and password or a set of access keys. To learn how to generate access keys, see [Managing Access Keys for IAM Users](#) in the *IAM User Guide*. When you generate access keys for an IAM user, make sure you view and securely save the key pair. You cannot recover the secret access key in the future. Instead, you must generate a new access key pair.

An *IAM group* is an identity that specifies a collection of IAM users. You can't sign in as a group. You can use groups to specify permissions for multiple users at a time. Groups make permissions easier to manage for large sets of users. For example, you could have a group named *IAMAdmins* and give that group permissions to administer IAM resources.

Users are different from roles. A user is uniquely associated with one person or application, but a role is intended to be assumable by anyone who needs it. Users have permanent long-term credentials, but roles provide temporary credentials. To learn more, see [When to Create an IAM User \(Instead of a Role\)](#) in the *IAM User Guide*.

You can authenticate to your DB instance using IAM database authentication.

IAM database authentication works with the following DB engines:

- Amazon RDS for MySQL
- Amazon RDS for PostgreSQL

For more information about authenticating to your DB instance using IAM, see [IAM Database Authentication for MySQL and PostgreSQL \(p. 461\)](#).

IAM Roles

An *IAM role* is an identity within your AWS account that has specific permissions. It is similar to an IAM user, but is not associated with a specific person. You can temporarily assume an IAM role in the AWS Management Console by [switching roles](#). You can assume a role by calling an AWS CLI or AWS API operation or by using a custom URL. For more information about methods for using roles, see [Using IAM Roles](#) in the *IAM User Guide*.

IAM roles with temporary credentials are useful in the following situations:

- **Temporary IAM user permissions** – An IAM user can assume an IAM role to temporarily take on different permissions for a specific task.
- **Federated user access** – Instead of creating an IAM user, you can use existing identities from AWS Directory Service, your enterprise user directory, or a web identity provider. These are known as *federated users*. AWS assigns a role to a federated user when access is requested through an [identity provider](#). For more information about federated users, see [Federated Users and Roles](#) in the *IAM User Guide*.
- **Cross-account access** – You can use an IAM role to allow someone (a trusted principal) in a different account to access resources in your account. Roles are the primary way to grant cross-account access. However, with some AWS services, you can attach a policy directly to a resource (instead of using a role as a proxy). To learn the difference between roles and resource-based policies for cross-account access, see [How IAM Roles Differ from Resource-based Policies](#) in the *IAM User Guide*.
- **AWS service access** – A service role is an IAM role that a service assumes to perform actions in your account on your behalf. When you set up some AWS service environments, you must define a role for the service to assume. This service role must include all the permissions that are required for the service to access the AWS resources that it needs. Service roles vary from service to service, but many allow you to choose your permissions as long as you meet the documented requirements for that service. Service roles provide access only within your account and cannot be used to grant access to services in other accounts. You can create, modify, and delete a service role from within IAM. For example, you can create a role that allows Amazon Redshift to access an Amazon S3 bucket on your behalf and then load data from that bucket into an Amazon Redshift cluster. For more information, see [Creating a Role to Delegate Permissions to an AWS Service](#) in the *IAM User Guide*.
- **Applications running on Amazon EC2** – You can use an IAM role to manage temporary credentials for applications that are running on an EC2 instance and making AWS CLI or AWS API requests. This is preferable to storing access keys within the EC2 instance. To assign an AWS role to an EC2 instance and make it available to all of its applications, you create an instance profile that is attached to the instance. An instance profile contains the role and enables programs that are running on the EC2 instance to get temporary credentials. For more information, see [Using an IAM Role to Grant Permissions to Applications Running on Amazon EC2 Instances](#) in the *IAM User Guide*.

To learn whether to use IAM roles, see [When to Create an IAM Role \(Instead of a User\)](#) in the *IAM User Guide*.

Managing Access Using Policies

You control access in AWS by creating policies and attaching them to IAM identities or AWS resources. A policy is an object in AWS that, when associated with an identity or resource, defines their permissions. AWS evaluates these policies when an entity (root user, IAM user, or IAM role) makes a request. Permissions in the policies determine whether the request is allowed or denied. Most policies are stored in AWS as JSON documents. For more information about the structure and contents of JSON policy documents, see [Overview of JSON Policies](#) in the *IAM User Guide*.

An IAM administrator can use policies to specify who has access to AWS resources, and what actions they can perform on those resources. Every IAM entity (user or role) starts with no permissions. In other words, by default, users can do nothing, not even change their own password. To give a user permission to do something, an administrator must attach a permissions policy to a user. Or the administrator can add the user to a group that has the intended permissions. When an administrator gives permissions to a group, all users in that group are granted those permissions.

IAM policies define permissions for an action regardless of the method that you use to perform the operation. For example, suppose that you have a policy that allows the `iam:GetRole` action. A user with that policy can get role information from the AWS Management Console, the AWS CLI, or the AWS API.

Identity-Based Policies

Identity-based policies are JSON permissions policy documents that you can attach to an identity, such as an IAM user, role, or group. These policies control what actions that identity can perform, on which resources, and under what conditions. To learn how to create an identity-based policy, see [Creating IAM Policies](#) in the *IAM User Guide*.

Identity-based policies can be further categorized as *inline policies* or *managed policies*. Inline policies are embedded directly into a single user, group, or role. Managed policies are standalone policies that you can attach to multiple users, groups, and roles in your AWS account. Managed policies include AWS managed policies and customer managed policies. To learn how to choose between a managed policy or an inline policy, see [Choosing Between Managed Policies and Inline Policies](#) in the *IAM User Guide*.

The following AWS managed policies, which you can attach to users in your account, are specific to Amazon RDS:

- **AmazonRDSReadOnlyAccess** – Grants read-only access to all Amazon RDS resources for the AWS account specified.
- **AmazonRDSFullAccess** – Grants full access to all Amazon RDS resources for the AWS account specified.

Other Policy Types

AWS supports additional, less-common policy types. These policy types can set the maximum permissions granted to you by the more common policy types.

- **Permissions boundaries** – A permissions boundary is an advanced feature in which you set the maximum permissions that an identity-based policy can grant to an IAM entity (IAM user or role). You can set a permissions boundary for an entity. The resulting permissions are the intersection of entity's identity-based policies and its permissions boundaries. Resource-based policies that specify the user or role in the Principal field are not limited by the permissions boundary. An explicit deny in any of these policies overrides the allow. For more information about permissions boundaries, see [Permissions Boundaries for IAM Entities](#) in the *IAM User Guide*.
- **Service control policies (SCPs)** – SCPs are JSON policies that specify the maximum permissions for an organization or organizational unit (OU) in AWS Organizations. AWS Organizations is a service for grouping and centrally managing multiple AWS accounts that your business owns. If you enable all features in an organization, then you can apply service control policies (SCPs) to any or all of your accounts. The SCP limits permissions for entities in member accounts, including each AWS account root user. For more information about Organizations and SCPs, see [How SCPs Work](#) in the *AWS Organizations User Guide*.
- **Session policies** – Session policies are advanced policies that you pass as a parameter when you programmatically create a temporary session for a role or federated user. The resulting session's permissions are the intersection of the user or role's identity-based policies and the session policies. Permissions can also come from a resource-based policy. An explicit deny in any of these policies overrides the allow. For more information, see [Session Policies](#) in the *IAM User Guide*.

Multiple Policy Types

When multiple types of policies apply to a request, the resulting permissions are more complicated to understand. To learn how AWS determines whether to allow a request when multiple policy types are involved, see [Policy Evaluation Logic](#) in the *IAM User Guide*.

For more information about identity and access management for Amazon RDS, continue to the following pages:

- [How Amazon RDS Works with IAM \(p. 449\)](#)
- [Troubleshooting Amazon RDS Identity and Access \(p. 478\)](#)

How Amazon RDS Works with IAM

Before you use IAM to manage access to Amazon RDS, you should understand what IAM features are available to use with Amazon RDS. To get a high-level view of how Amazon RDS and other AWS services work with IAM, see [AWS Services That Work with IAM](#) in the *IAM User Guide*.

Topics

- [Amazon RDS Identity-Based Policies \(p. 449\)](#)
- [Amazon RDS Resource-Based Policies \(p. 451\)](#)
- [Authorization Based on Amazon RDS Tags \(p. 451\)](#)
- [Amazon RDS IAM Roles \(p. 451\)](#)

Amazon RDS Identity-Based Policies

With IAM identity-based policies, you can specify allowed or denied actions and resources as well as the conditions under which actions are allowed or denied. Amazon RDS supports specific actions, resources, and condition keys. To learn about all of the elements that you use in a JSON policy, see [IAM JSON Policy Elements Reference](#) in the *IAM User Guide*.

Actions

The `Action` element of an IAM identity-based policy describes the specific action or actions that will be allowed or denied by the policy. Policy actions usually have the same name as the associated AWS API operation. The action is used in a policy to grant permissions to perform the associated operation.

Policy actions in Amazon RDS use the following prefix before the action: `rds::`. For example, to grant someone permission to describe DB instances with the Amazon RDS `DescribeDBInstances` API operation, you include the `rds:DescribeDBInstances` action in their policy. Policy statements must include either an `Action` or `NotAction` element. Amazon RDS defines its own set of actions that describe tasks that you can perform with this service.

To specify multiple actions in a single statement, separate them with commas as follows:

```
"Action": [  
    "rds:action1",  
    "rds:action2"]
```

You can specify multiple actions using wildcards (*). For example, to specify all actions that begin with the word `Describe`, include the following action:

```
"Action": "rds:Describe*"
```

To see a list of Amazon RDS actions, see [Actions Defined by Amazon RDS in the IAM User Guide](#).

Resources

The `Resource` element specifies the object or objects to which the action applies. Statements must include either a `Resource` or a `NotResource` element. You specify a resource using an ARN or using the wildcard (*) to indicate that the statement applies to all resources.

The DB instance resource has the following ARN:

```
arn:${Partition}:rds:${Region}:${Account}:{ResourceType}/${Resource}
```

For more information about the format of ARNs, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

For example, to specify the dbtest DB instance in your statement, use the following ARN:

```
"Resource": "arn:aws:rds:us-west-2:123456789012:db:dbtest"
```

To specify all DB instances that belong to a specific account, use the wildcard (*):

```
"Resource": "arn:aws:ec2:us-east-1:123456789012:db:/*"
```

Some RDS API operations, such as those for creating resources, cannot be performed on a specific resource. In those cases, you must use the wildcard (*).

```
"Resource": "*"
```

Many Amazon RDS API operations involve multiple resources. For example, `CreateDBInstance` creates a DB instance. You can specify that an IAM user must use a specific security group and parameter group when creating a DB instance. To specify multiple resources in a single statement, separate the ARNs with commas.

```
"Resource": [  
    "resource1",  
    "resource2"]
```

To see a list of Amazon RDS resource types and their ARNs, see [Resources Defined by Amazon RDS](#) in the *IAM User Guide*. To learn with which actions you can specify the ARN of each resource, see [Actions Defined by Amazon RDS](#).

Condition Keys

The Condition element (or Condition *block*) lets you specify conditions in which a statement is in effect. The Condition element is optional. You can build conditional expressions that use [condition operators](#), such as equals or less than, to match the condition in the policy with values in the request.

If you specify multiple Condition elements in a statement, or multiple keys in a single Condition element, AWS evaluates them using a logical AND operation. If you specify multiple values for a single condition key, AWS evaluates the condition using a logical OR operation. All of the conditions must be met before the statement's permissions are granted.

You can also use placeholder variables when you specify conditions. For example, you can grant an IAM user permission to access a resource only if it is tagged with their IAM user name. For more information, see [IAM Policy Elements: Variables and Tags](#) in the *IAM User Guide*.

Amazon RDS defines its own set of condition keys and also supports using some global condition keys. To see all AWS global condition keys, see [AWS Global Condition Context Keys](#) in the *IAM User Guide*.

All RDS API operations support the `aws:RequestedRegion` condition key.

To see a list of Amazon RDS condition keys, see [Condition Keys for Amazon RDS](#) in the *IAM User Guide*. To learn with which actions and resources you can use a condition key, see [Actions Defined by Amazon RDS](#).

Examples

To view examples of Amazon RDS identity-based policies, see [Amazon RDS Identity-Based Policy Examples \(p. 451\)](#).

Amazon RDS Resource-Based Policies

Amazon RDS does not support resource-based policies.

Authorization Based on Amazon RDS Tags

You can attach tags to Amazon RDS resources or pass tags in a request to Amazon RDS. To control access based on tags, you provide tag information in the `condition element` of a policy using the `rds:ResourceTag/key-name`, `aws:RequestTag/key-name`, or `aws:TagKeys` condition keys. For more information about tagging Amazon RDS resources, see [Specifying Conditions: Using Custom Tags \(p. 458\)](#).

To view an example identity-based policy for limiting access to a resource based on the tags on that resource, see [Grant Permission for Actions on a Resource with a Specific Tag with Two Different Values \(p. 456\)](#).

Amazon RDS IAM Roles

An [IAM role](#) is an entity within your AWS account that has specific permissions.

Using Temporary Credentials with Amazon RDS

You can use temporary credentials to sign in with federation, assume an IAM role, or to assume a cross-account role. You obtain temporary security credentials by calling AWS STS API operations such as [AssumeRole](#) or [GetFederationToken](#).

Amazon RDS supports using temporary credentials.

Service-Linked Roles

[Service-linked roles](#) allow AWS services to access resources in other services to complete an action on your behalf. Service-linked roles appear in your IAM account and are owned by the service. An IAM administrator can view but not edit the permissions for service-linked roles.

Amazon RDS supports service-linked roles. For details about creating or managing Amazon RDS service-linked roles, see [Using Service-Linked Roles for Amazon RDS \(p. 501\)](#).

Service Roles

This feature allows a service to assume a [service role](#) on your behalf. This role allows the service to access resources in other services to complete an action on your behalf. Service roles appear in your IAM account and are owned by the account. This means that an IAM administrator can change the permissions for this role. However, doing so might break the functionality of the service.

Amazon RDS supports service roles.

Amazon RDS Identity-Based Policy Examples

By default, IAM users and roles don't have permission to create or modify Amazon RDS resources. They also can't perform tasks using the AWS Management Console, AWS CLI, or AWS API. An IAM

administrator must create IAM policies that grant users and roles permission to perform specific API operations on the specified resources they need. The administrator must then attach those policies to the IAM users or groups that require those permissions.

To learn how to create an IAM identity-based policy using these example JSON policy documents, see [Creating Policies on the JSON Tab](#) in the *IAM User Guide*.

Topics

- [Policy Best Practices \(p. 452\)](#)
- [Using the Amazon RDS Console \(p. 452\)](#)
- [Allow Users to View Their Own Permissions \(p. 453\)](#)
- [Allow a User to Create DB Instances in an AWS account \(p. 453\)](#)
- [Permissions Required to Use the Console \(p. 455\)](#)
- [Allow a User to Perform Any Describe Action on Any RDS Resource \(p. 455\)](#)
- [Allow a User to Create a DB Instance That Uses the Specified DB Parameter and Security Groups \(p. 455\)](#)
- [Grant Permission for Actions on a Resource with a Specific Tag with Two Different Values \(p. 456\)](#)
- [Prevent a User from Deleting a DB Instance \(p. 456\)](#)
- [Example Policies: Using Condition Keys \(p. 456\)](#)
- [Specifying Conditions: Using Custom Tags \(p. 458\)](#)

Policy Best Practices

Identity-based policies are very powerful. They determine whether someone can create, access, or delete Amazon RDS resources in your account. These actions can incur costs for your AWS account. When you create or edit identity-based policies, follow these guidelines and recommendations:

- **Get Started Using AWS Managed Policies** – To start using Amazon RDS quickly, use AWS managed policies to give your employees the permissions they need. These policies are already available in your account and are maintained and updated by AWS. For more information, see [Get Started Using Permissions With AWS Managed Policies](#) in the *IAM User Guide*.
- **Grant Least Privilege** – When you create custom policies, grant only the permissions required to perform a task. Start with a minimum set of permissions and grant additional permissions as necessary. Doing so is more secure than starting with permissions that are too lenient and then trying to tighten them later. For more information, see [Grant Least Privilege](#) in the *IAM User Guide*.
- **Enable MFA for Sensitive Operations** – For extra security, require IAM users to use multi-factor authentication (MFA) to access sensitive resources or API operations. For more information, see [Using Multi-Factor Authentication \(MFA\) in AWS](#) in the *IAM User Guide*.
- **Use Policy Conditions for Extra Security** – To the extent that it's practical, define the conditions under which your identity-based policies allow access to a resource. For example, you can write conditions to specify a range of allowable IP addresses that a request must come from. You can also write conditions to allow requests only within a specified date or time range, or to require the use of SSL or MFA. For more information, see [IAM JSON Policy Elements: Condition](#) in the *IAM User Guide*.

Using the Amazon RDS Console

To access the Amazon RDS console, you must have a minimum set of permissions. These permissions must enable you to list and view details about the Amazon RDS resources in your AWS account. You can create an identity-based policy that is more restrictive than the minimum required permissions. However, if you do, the console doesn't function as intended for entities (IAM users or roles) with that policy.

To ensure that those entities can still use the Amazon RDS console, also attach the following AWS managed policy to the entities. For more information, see [Adding Permissions to a User](#) in the *IAM User Guide*.

AmazonRDSReadOnlyAccess

You don't need to allow minimum console permissions for users that are making calls only to the AWS CLI or the AWS API. Instead, allow access to only the actions that match the API operation that you're trying to perform.

Allow Users to View Their Own Permissions

This example shows how you might create a policy that allows IAM users to view the inline and managed policies that are attached to their user identity. This policy includes permissions to complete this action on the console or programmatically using the AWS CLI or AWS API.

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Sid": "ViewOwnUserInfo",  
            "Effect": "Allow",  
            "Action": [  
                "iam:GetUserPolicy",  
                "iam>ListGroupsForUser",  
                "iam>ListAttachedUserPolicies",  
                "iam>ListUserPolicies",  
                "iam GetUser"  
            ],  
            "Resource": [  
                "arn:aws:iam::*:user/${aws:username}"  
            ]  
        },  
        {  
            "Sid": "NavigateInConsole",  
            "Effect": "Allow",  
            "Action": [  
                "iam:GetGroupPolicy",  
                "iam:GetPolicyVersion",  
                "iam GetPolicy",  
                "iam>ListAttachedGroupPolicies",  
                "iam>ListGroupPolicies",  
                "iam>ListPolicyVersions",  
                "iam>ListPolicies",  
                "iam>ListUsers"  
            ],  
            "Resource": "*"  
        }  
    ]  
}
```

Allow a User to Create to DB Instances in an AWS account

The following is an example policy that allows the user with the ID 123456789012 to create DB instances for your AWS account. The policy requires that the name of the new DB instance begin with test. The new DB instance must also use the MySQL database engine and the db.t2.micro DB instance class. In addition, the new DB instance must use an option group and a DB parameter group that starts with default, and it must use the default subnet group.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "AllowCreateDBInstanceOnly",
            "Effect": "Allow",
            "Action": [
                "rds:CreateDBInstance"
            ],
            "Resource": [
                "arn:aws:rds::123456789012:db:test*",
                "arn:aws:rds::123456789012:og:default*",
                "arn:aws:rds::123456789012:pg:default*",
                "arn:aws:rds::123456789012:subgrp:default"
            ],
            "Condition": {
                "StringEquals": {
                    "rds:DatabaseEngine": "mysql",
                    "rds:DatabaseClass": "db.t2.micro"
                }
            }
        }
    ]
}
```

The policy includes a single statement that specifies the following permissions for the IAM user:

- The policy allows the IAM user to create a DB instance using the [CreateDBInstance](#) API operation (this also applies to the [create-db-instance](#) AWS CLI command and the AWS Management Console).
- The **Resource** element specifies that the user can perform actions on or with resources. You specify resources using an Amazon Resources Name (ARN). This ARN includes the name of the service that the resource belongs to (`rds`), the AWS Region (* indicates any region in this example), the user account number (123456789012 is the user ID in this example), and the type of resource. For more information about creating ARNs, see [Working with Amazon Resource Names \(ARNs\) in Amazon RDS \(p. 224\)](#).

The **Resource** element in the example specifies the following policy constraints on resources for the user:

- The DB instance identifier for the new DB instance must begin with `test` (for example, `testCustomerData1`, `test-region2-data`).
- The option group for the new DB instance must begin with `default`.
- The DB parameter group for the new DB instance must begin with `default`.
- The subnet group for the new DB instance must be the `default` subnet group.
- The **Condition** element specifies that the DB engine must be MySQL and the DB instance class must be db.t2.micro. The **Condition** element specifies the conditions when a policy should take effect. You can add additional permissions or restrictions by using the **Condition** element. For more information about specifying conditions, see [Condition Keys \(p. 450\)](#). This example specifies the `rds:DatabaseEngine` and `rds:DatabaseClass` conditions. For information about the valid condition values for `rds:DatabaseEngine`, see the list under the `Engine` parameter in [CreateDBInstance](#). For information about the valid condition values for `rds:DatabaseClass`, see [Supported DB Engines for All Available DB Instance Classes \(p. 112\)](#).

The policy doesn't specify the **Principal** element because in an identity-based policy you don't specify the principal who gets the permission. When you attach policy to a user, the user is the implicit principal. When you attach a permission policy to an IAM role, the principal identified in the role's trust policy gets the permissions.

To see a list of Amazon RDS actions, see [Actions Defined by Amazon RDS in the IAM User Guide](#).

Permissions Required to Use the Console

For a user to work with the console, that user must have a minimum set of permissions. These permissions allow the user to describe the Amazon RDS resources for their AWS account and to provide other related information, including Amazon EC2 security and network information.

If you create an IAM policy that is more restrictive than the minimum required permissions, the console doesn't function as intended for users with that IAM policy. To ensure that those users can still use the console, also attach the `AmazonRDSReadOnlyAccess` managed policy to the user, as described in [Managing Access Using Policies \(p. 447\)](#).

You don't need to allow minimum console permissions for users that are making calls only to the AWS CLI or the Amazon RDS API.

The following policy grants full access to all Amazon RDS resources for the root AWS account:

```
AmazonRDSFullAccess
```

Allow a User to Perform Any Describe Action on Any RDS Resource

The following permissions policy grants permissions to a user to run all of the actions that begin with `Describe`. These actions show information about an RDS resource, such as a DB instance. The wildcard character (*) in the `Resource` element indicates that the actions are allowed for all Amazon RDS resources owned by the account.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "AllowRDSDescribe",
            "Effect": "Allow",
            "Action": "rds:Describe*",
            "Resource": "*"
        }
    ]
}
```

Allow a User to Create a DB Instance That Uses the Specified DB Parameter and Security Groups

The following permissions policy grants permissions to allow a user to only create a DB instance that must use the `mysql-production` DB parameter group and the `db-production` DB security group.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "AllowMySQLProductionCreate",
            "Effect": "Allow",
            "Action": "rds>CreateDBInstance",
            "Resource": [
                "arn:aws:rds:us-west-2:123456789012:pg:mysql-production",

```

```
        "arn:aws:rds:us-west-2:123456789012:secgrp:db-production"
    }
}
}
```

Grant Permission for Actions on a Resource with a Specific Tag with Two Different Values

You can use conditions in your identity-based policy to control access to Amazon RDS resources based on tags. The following policy allows permission to perform the `ModifyDBInstance` and `CreateDBSnapshot` APIs on DB instances with either the `stage` tag set to `development` or `test`.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "AllowDevTestCreate",
      "Effect": "Allow",
      "Action": [
        "rds:ModifyDBInstance",
        "rds:CreateDBSnapshot"
      ],
      "Resource": "*",
      "Condition": {
        "StringEquals": {
          "rds:db-tag/stage": [
            "development",
            "test"
          ]
        }
      }
    }
  ]
}
```

Prevent a User from Deleting a DB Instance

The following permissions policy grants permissions to prevent a user from deleting a specific DB instance. For example, you might want to deny the ability to delete your production DB instances to any user that is not an administrator.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "DenyDelete1",
      "Effect": "Deny",
      "Action": "rds>DeleteDBInstance",
      "Resource": "arn:aws:rds:us-west-2:123456789012:db:my-mysql-instance"
    }
  ]
}
```

Example Policies: Using Condition Keys

Following are examples of how you can use condition keys in Amazon RDS IAM permissions policies.

Example 1: Grant Permission to Create a DB Instance that Uses a Specific DB Engine and Isn't MultiAZ

The following policy uses an RDS condition key and allows a user to create only DB instances that use the MySQL database engine and don't use MultiAZ. The Condition element indicates the requirement that the database engine is MySQL.

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Sid": "AllowMySQLCreate",  
            "Effect": "Allow",  
            "Action": "rds>CreateDBInstance",  
            "Resource": "*",  
            "Condition": {  
                "StringEquals": {  
                    "rds:DatabaseEngine": "mysql"  
                },  
                "Bool": {  
                    "rds:MultiAz": false  
                }  
            }  
        }  
    ]  
}
```

Example 2: Explicitly Deny Permission to Create DB Instances for Certain DB Instance Classes and Create DB Instances that Use Provisioned IOPS

The following policy explicitly denies permission to create DB instances that use the DB instance classes `r3.8xlarge` and `m4.10xlarge`, which are the largest and most expensive DB instance classes. This policy also prevents users from creating DB instances that use Provisioned IOPS, which incurs an additional cost.

Explicitly denying permission supersedes any other permissions granted. This ensures that identities do not accidentally get permission that you never want to grant.

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Sid": "DenyLargeCreate",  
            "Effect": "Deny",  
            "Action": "rds>CreateDBInstance",  
            "Resource": "*",  
            "Condition": {  
                "StringEquals": {  
                    "rds:DatabaseClass": [  
                        "db.r3.8xlarge",  
                        "db.m4.10xlarge"  
                    ]  
                }  
            }  
        },  
        {  
            "Sid": "DenyPIOPSCreate",  
            "Effect": "Deny",  
            "Action": "rds>CreateDBInstance",  
            "Resource": "*"  
        }  
    ]  
}
```

```

        "Condition": {
            "NumericNotEquals": {
                "rds:Piops": "0"
            }
        }
    ]
}

```

Example 3: Limit the Set of Tag Keys and Values That Can Be Used to Tag a Resource

The following policy uses an RDS condition key and allows the addition of a tag with the key `stage` to be added to a resource with the values `test`, `qa`, and `production`.

```

{
{
    "Version" : "2012-10-17",
    "Statement" : [
        {
            "Effect" : "Allow",
            "Action" : [ "rds:AddTagsToResource", "rds:RemoveTagsFromResource" ],
            "Resource" : "*",
            "Condition" : { "streq" : { "rds:req-tag/stage" : [ "test", "qa", "production" ] } }
        }
    ]
}
}

```

Specifying Conditions: Using Custom Tags

Amazon RDS supports specifying conditions in an IAM policy using custom tags.

For example, suppose that you add a tag named `environment` to your DB instances with values such as `beta`, `staging`, `production`, and so on. If you do, you can create a policy that restricts certain users to DB instances based on the `environment` tag value.

Note

Custom tag identifiers are case-sensitive.

The following table lists the RDS tag identifiers that you can use in a Condition element.

RDS Tag Identifier	Applies To
<code>db-tag</code>	DB instances, including Read Replicas
<code>snapshot-tag</code>	DB snapshots
<code>ri-tag</code>	Reserved DB instances
<code>secgrp-tag</code>	DB security groups
<code>og-tag</code>	DB option groups
<code>pg-tag</code>	DB parameter groups
<code>subgrp-tag</code>	DB subnet groups
<code>es-tag</code>	Event subscriptions

RDS Tag Identifier	Applies To
cluster-tag	DB clusters
cluster-pg-tag	DB cluster parameter groups
cluster-snapshot-tag	DB cluster snapshots

The syntax for a custom tag condition is as follows:

```
"Condition": {"StringEquals": {"rds:rds-tag-identifier/tag-name": ["value"]}}
```

For example, the following Condition element applies to DB instances with a tag named environment and a tag value of production.

```
"Condition": {"StringEquals": {"rds:db-tag/environment": ["production"]}}
```

For information about creating tags, see [Tagging Amazon RDS Resources \(p. 177\)](#).

Important

If you manage access to your RDS resources using tagging, we recommend that you secure access to the tags for your RDS resources. You can manage access to tags by creating policies for the AddTagsToResource and RemoveTagsFromResource actions. For example, the following policy denies users the ability to add or remove tags for all resources. You can then create policies to allow specific users to add or remove tags.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "DenyTagUpdates",
            "Effect": "Deny",
            "Action": [
                "rds:AddTagsToResource",
                "rds:RemoveTagsFromResource"
            ],
            "Resource": "*"
        }
    ]
}
```

To see a list of Amazon RDS actions, see [Actions Defined by Amazon RDS in the IAM User Guide](#).

Example Policies: Using Custom Tags

Following are examples of how you can use custom tags in Amazon RDS IAM permissions policies. For more information about adding tags to an Amazon RDS resource, see [Working with Amazon Resource Names \(ARNs\) in Amazon RDS \(p. 224\)](#).

Note

All examples use the us-west-2 region and contain fictitious account IDs.

Example 1: Grant Permission for Actions on a Resource with a Specific Tag with Two Different Values

The following policy allows permission to perform the `ModifyDBInstance` and `CreateDBSnapshot` APIs on DB instances with either the stage tag set to development or test.

```
{
    "Version": "2012-10-17",
```

```

"Statement": [
    {
        "Sid": "AllowDevTestCreate",
        "Effect": "Allow",
        "Action": [
            "rds:ModifyDBInstance",
            "rds>CreateDBSnapshot"
        ],
        "Resource": "*",
        "Condition": {
            "StringEquals": {
                "rds:db-tag/stage": [
                    "development",
                    "test"
                ]
            }
        }
    }
]
}

```

Example 2: Explicitly Deny Permission to Create a DB Instance that Uses Specified DB Parameter Groups

The following policy explicitly denies permission to create a DB instance that uses DB parameter groups with specific tag values. You might apply this policy if you require that a specific customer-created DB parameter group always be used when creating DB instances. Policies that use Deny are most often used to restrict access that was granted by a broader policy.

Explicitly denying permission supersedes any other permissions granted. This ensures that identities do not accidentally get permission that you never want to grant.

```

{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "DenyProductionCreate",
            "Effect": "Deny",
            "Action": "rds>CreateDBInstance",
            "Resource": "*",
            "Condition": {
                "StringEquals": {
                    "rds:pg-tag/usage": "prod"
                }
            }
        }
    ]
}

```

Example 3: Grant Permission for Actions on a DB Instance with an Instance Name that is Prefixed with a User Name

The following policy allows permission to call any API (except to `AddTagsToResource` or `RemoveTagsFromResource`) on a DB instance that has a DB instance name that is prefixed with the user's name and that has a tag called `stage` equal to `devo` or that has no tag called `stage`.

The Resource line in the policy identifies a resource by its Amazon Resource Name (ARN). For more information about using ARNs with Amazon RDS resources, see [Working with Amazon Resource Names \(ARNs\) in Amazon RDS \(p. 224\)](#).

```

{
    "Version": "2012-10-17",
}

```

```
"Statement": [
    {
        "Sid": "AllowFullDevAccessNoTags",
        "Effect": "Allow",
        "NotAction": [
            "rds:AddTagsToResource",
            "rds:RemoveTagsFromResource"
        ],
        "Resource": "arn:aws:rds:*:123456789012:db:${aws:username}*",
        "Condition": {
            "StringEqualsIfExists": {
                "rds:db-tag/stage": "devo"
            }
        }
    }
]
```

IAM Database Authentication for MySQL and PostgreSQL

You can authenticate to your DB instance using AWS Identity and Access Management (IAM) database authentication. IAM database authentication works with MySQL and PostgreSQL. With this authentication method, you don't need to use a password when you connect to a DB instance. Instead, you use an authentication token.

An *authentication token* is a unique string of characters that Amazon RDS generates on request. Authentication tokens are generated using AWS Signature Version 4. Each token has a lifetime of 15 minutes. You don't need to store user credentials in the database, because authentication is managed externally using IAM. You can also still use standard database authentication.

IAM database authentication provides the following benefits:

- Network traffic to and from the database is encrypted using Secure Sockets Layer (SSL).
- You can use IAM to centrally manage access to your database resources, instead of managing access individually on each DB instance.
- For applications running on Amazon EC2, you can use profile credentials specific to your EC2 instance to access your database instead of a password, for greater security.

Topics

- [Availability for IAM Database Authentication \(p. 461\)](#)
- [MySQL Limitations for IAM Database Authentication \(p. 462\)](#)
- [PostgreSQL Limitations for IAM Database Authentication \(p. 462\)](#)
- [Enabling and Disabling IAM Database Authentication \(p. 462\)](#)
- [Creating and Using an IAM Policy for IAM Database Access \(p. 464\)](#)
- [Creating a Database Account Using IAM Authentication \(p. 467\)](#)
- [Connecting to Your DB Instance Using IAM Authentication \(p. 468\)](#)

Availability for IAM Database Authentication

IAM database authentication is available for the following database engines and DB instance classes:

- MySQL 5.6, minor version 5.6.34 or higher.
- MySQL 5.7, minor version 5.7.16 or higher.

- MySQL 8.0, minor version 8.0.16 or higher.
- PostgreSQL versions 10.6 or higher, 9.6.11 or higher, and 9.5.15 or higher.

Note

IAM database authentication is not supported for MySQL 5.5.

MySQL Limitations for IAM Database Authentication

When using IAM database authentication with MySQL, you are limited to a maximum of 200 new connections per second. If you are using a db.t2.micro DB instance class, the limit is 10 connections per second.

The database engines that work with Amazon RDS don't impose any limits on authentication attempts per second. However, when you use IAM database authentication, your application must generate an authentication token. Your application then uses that token to connect to the DB instance. If you exceed the limit of maximum new connections per second, then the extra overhead of IAM database authentication can cause connection throttling. The extra overhead can cause even existing connections to drop. For information about the maximum total connections for MySQL, see [Maximum MySQL connections \(p. 769\)](#).

We recommend the following when using the MySQL engine:

- Use IAM database authentication as a mechanism for temporary, personal access to databases.
- Use IAM database authentication only for workloads that can be easily retried.
- Don't use IAM database authentication if your application requires more than 200 new connections per second.

PostgreSQL Limitations for IAM Database Authentication

When using IAM database authentication with PostgreSQL, note the following limitation:

- The maximum number of connections per second for your database instance may be limited depending on the instance type and your workload.

Enabling and Disabling IAM Database Authentication

By default, IAM database authentication is disabled on DB instances. You can enable IAM database authentication (or disable it again) using the AWS Management Console, AWS CLI, or the API.

IAM authentication for PostgreSQL DB instances requires that the SSL value be 1. You can't enable IAM authentication for a PostgreSQL DB instance if the SSL value is 0. You can't change the SSL value to 0 if IAM authentication is enabled for a PostgreSQL DB instance.

Console

To create a new DB instance with IAM authentication by using the console, see either [Creating a DB Instance Running the MySQL Database Engine \(p. 751\)](#) or [Creating a DB Instance Running the PostgreSQL Database Engine \(p. 1257\)](#).

Each creation workflow has a **Configure Advanced Settings** page, where you can enable IAM DB authentication. In that page's **Database Options** section, choose **Yes** for **Enable IAM DB Authentication**.

To enable or disable IAM authentication for an existing DB instance

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.

2. In the navigation pane, choose **Databases**.
3. Choose the DB instance that you want to modify.

Note
Make sure that all affected DB instances are compatible with IAM authentication. Check the compatibility requirements in [Availability for IAM Database Authentication \(p. 461\)](#). For an Aurora DB cluster, you can only enable IAM authentication if all DB instances in the cluster are compatible with IAM.
4. Choose **Modify**.
5. In the **Database options** section, for **IAM DB authentication** choose **Enable IAM DB authentication** or **Disable**, and then choose **Continue**.
6. To apply the changes immediately, choose **Apply immediately**.
7. Choose **Modify DB instance**.

To restore a DB instance

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Snapshots**.
3. Choose the snapshot that you want to restore, and then choose **Restore Snapshot** for **Actions**.
4. In the **Settings** section, enter an identifier for the DB instance for **DB Instance Identifier**.
5. In the **Database options** section, for **IAM DB authentication**, choose **Enable IAM DB authentication** or **Disable**.
6. Choose **Restore DB Instance**.

AWS CLI

To create a new DB instance with IAM authentication by using the AWS CLI, use the `create-db-instance` command. Specify the `--enable-iam-database-authentication` option, as shown in the following example.

```
aws rds create-db-instance \
--db-instance-identifier mydbinstance \
--db-instance-class db.m3.medium \
--engine MySQL \
--allocated-storage 20 \
--master-username masterawsuser \
--master-user-password masteruserpassword \
--enable-iam-database-authentication
```

To update an existing DB instance to have or not have IAM authentication, use the AWS CLI command `modify-db-instance`. Specify either the `--enable-iam-database-authentication` or `--no-enable-iam-database-authentication` option, as appropriate.

Note

Make sure that all affected DB instances are compatible with IAM authentication. Check the compatibility requirements in [Availability for IAM Database Authentication \(p. 461\)](#). For an Aurora DB cluster, you can only enable IAM authentication if all DB instances in the cluster are compatible with IAM.

By default, Amazon RDS performs the modification during the next maintenance window. If you want to override this and enable IAM DB authentication as soon as possible, use the `--apply-immediately` parameter.

The following example shows how to immediately enable IAM authentication for an existing DB instance.

```
aws rds modify-db-instance \
```

```
--db-instance-identifier mydbinstance \  
--apply-immediately \  
--enable-iam-database-authentication
```

If you are restoring a DB instance, use one of the following AWS CLI commands:

- [restore-db-instance-to-point-in-time](#)
- [restore-db-instance-from-db-snapshot](#)

The IAM database authentication setting defaults to that of the source snapshot. To change this setting, set the `--enable-iam-database-authentication` or `--no-enable-iam-database-authentication` option, as appropriate.

RDS API

To create a new DB instance with IAM authentication by using the API, use the API operation [CreateDBInstance](#). Set the `EnableIAMDatabaseAuthentication` parameter to `true`.

To update an existing DB instance to have IAM authentication, use the API operation [ModifyDBInstance](#). Set the `EnableIAMDatabaseAuthentication` parameter to `true` to enable IAM authentication, or `false` to disable it.

Note

Make sure that all affected DB instances are compatible with IAM authentication. Check the compatibility requirements in [Availability for IAM Database Authentication \(p. 461\)](#). For an Aurora DB cluster, you can only enable IAM authentication if all DB instances in the cluster are compatible with IAM.

If you are restoring a DB instance, use one of the following API operations:

- [RestoreDBInstanceToPointInTime](#)
- [RestoreDBInstanceFromDBSnapshot](#)

The IAM database authentication setting defaults to that of the source snapshot. To change this setting, set the `EnableIAMDatabaseAuthentication` parameter to `true` to enable IAM authentication, or `false` to disable it.

Creating and Using an IAM Policy for IAM Database Access

To allow an IAM user or role to connect to your DB instance, you must create an IAM policy. After that, you attach the policy to an IAM user or role.

Note

To learn more about IAM policies, see [Identity and Access Management in Amazon RDS \(p. 445\)](#).

The following example policy allows an IAM user to connect to a DB instance using IAM database authentication.

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "rds-db:connect"  
            ],  
            "Resource": [  
                "arn:aws:rds:region:account:dbinstance:iamuser"  
            ]  
        }  
    ]  
}
```

```
        "arn:aws:rds-db:us-east-2:1234567890:dbuser:db-ABCDEFGHIJKLM01234/db_user"
    }
}
```

Important

An IAM administrator user can access DB instances without explicit permissions in an IAM policy. The example in [Create an IAM User \(p. 5\)](#) creates an IAM administrator user. If you want to restrict administrator access to DB instances, you can create an IAM role with the appropriate, lesser privileged permissions and assign it to the administrator.

Note

Don't confuse the `rds-db:` prefix with other RDS API operation prefixes that begin with `rds:`. You use the `rds-db:` prefix and the `rds-db:connect` action only for IAM database authentication. They aren't valid in any other context.

Currently, the IAM console displays an error for policies with the `rds-db:connect` action. You can ignore this error.

The example policy includes a single statement with the following elements:

- **Effect** – Specify `Allow` to grant access to the DB instance. If you don't explicitly allow access, then access is denied by default.
- **Action** – Specify `rds-db:connect` to allow connections to the DB instance.
- **Resource** – Specify an Amazon Resource Name (ARN) that describes one database account in one DB instance. The ARN format is as follows.

```
arn:aws:rds-db:region:account-id:dbuser:DbiResourceId/db-user-name
```

In this format, replace the following:

- **region** is the AWS Region for the DB instance. In the example policy, the AWS Region is `us-east-2`.
- **account-id** is the AWS account number for the DB instance. In the example policy, the account number is `1234567890`.
- **DbiResourceId** is the identifier for the DB instance. This identifier is unique to an AWS Region and never changes. In the example policy, the identifier is `db-ABCDEFGHIJKLM01234`.

To find a DB instance resource ID in the AWS Management Console for Amazon RDS, choose the DB instance to see its details. Then choose the **Configuration** tab. The **Resource ID** is shown in the **Configuration** section.

Alternatively, you can use the AWS CLI command to list the identifiers and resource IDs for all of your DB instance in the current AWS Region, as shown following.

```
aws rds describe-db-instances --query "DBInstances[*].  
[DBInstanceIdentifier,DbiResourceId]"
```

- **db-user-name** is the name of the database account to associate with IAM authentication. In the example policy, the database account is `db_user`.

You can construct other ARNs to support various access patterns. The following policy allows access to two different database accounts in a DB instance.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "rds-db:connect"
            ],
            "Resource": [
                "arn:aws:rds-db:us-east-2:123456789012:dbuser:db-ABCDEFGHIJKLM01234/jane_doe",
                "arn:aws:rds-db:us-east-2:123456789012:dbuser:db-ABCDEFGHIJKLM01234/mary_roe"
            ]
        }
    ]
}
```

The following policy uses the "*" character to match all DB instances and database accounts for a particular AWS account and AWS Region.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "rds-db:connect"
            ],
            "Resource": [
                "arn:aws:rds-db:us-east-2:1234567890:dbuser:*/**"
            ]
        }
    ]
}
```

The following policy matches all of the DB instances for a particular AWS account and AWS Region. However, the policy only grants access to DB instances that have a jane_doe database account.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "rds-db:connect"
            ],
            "Resource": [
                "arn:aws:rds-db:us-east-2:123456789012:dbuser:*/jane_doe"
            ]
        }
    ]
}
```

The IAM user or role has access to only those databases that the database user does. For example, suppose that your DB instance has a database named *dev*, and another database named *test*. If the database user *jane_doe* has access only to *dev*, any IAM users or roles that access that DB instance with

the `jane_doe` user also have access only to `dev`. This access restriction is also true for other database objects, such as tables, views, and so on.

Attaching an IAM Policy to an IAM User or Role

After you create an IAM policy to allow database authentication, you need to attach the policy to an IAM user or role. For a tutorial on this topic, see [Create and Attach Your First Customer Managed Policy](#) in the [IAM User Guide](#).

As you work through the tutorial, you can use one of the policy examples shown in this section as a starting point and tailor it to your needs. At the end of the tutorial, you have an IAM user with an attached policy that can make use of the `rds-db:connect` action.

Note

You can map multiple IAM users or roles to the same database user account. For example, suppose that your IAM policy specified the following resource ARN.

```
arn:aws:rds-db:us-east-2:123456789012:dbuser:db-12ABC34DEFG5HIJ6KLMNOP78QR/jane_doe
```

If you attach the policy to IAM users *Jane*, *Bob*, and *Diego*, then each of those users can connect to the specified DB instance using the `jane_doe` database account.

Creating a Database Account Using IAM Authentication

With IAM database authentication, you don't need to assign database passwords to the user accounts you create. If you remove an IAM user that is mapped to a database account, you should also remove the database account with the `DROP USER` statement.

Using IAM Authentication with PostgreSQL

To use IAM authentication with PostgreSQL, connect to the DB instance, create database users, and then grant them the `rds_iam` role as shown in the following example.

```
CREATE USER db_userx WITH LOGIN;
GRANT rds_iam TO db_userx;
```

Using IAM Authentication with MySQL

With MySQL, authentication is handled by `AWSAuthenticationPlugin`—an AWS-provided plugin that works seamlessly with IAM to authenticate your IAM users. Connect to the DB instance and issue the `CREATE USER` statement, as shown in the following example.

```
CREATE USER jane_doe IDENTIFIED WITH AWSAuthenticationPlugin AS 'RDS';
```

The `IDENTIFIED WITH` clause allows MySQL to use the `AWSAuthenticationPlugin` to authenticate the database account (`jane_doe`). The `AS 'RDS'` clause refers to the authentication method, and the specified database account should have the same name as the IAM user or role. In this example, both the database account and the IAM user or role are named `jane_doe`.

Note

If you see the following message, it means that the AWS-provided plugin is not available for the current DB instance.

ERROR 1524 (HY000): Plugin 'AWSAuthenticationPlugin' is not loaded
To troubleshoot this error, verify that you are using a supported configuration and that
you have enabled IAM database authentication on your DB instance. For more information,
see [Availability for IAM Database Authentication \(p. 461\)](#) and [Enabling and Disabling IAM
Database Authentication \(p. 462\)](#).

After you create an account using `AWSAuthenticationPlugin`, you manage it in the same way as
other database accounts. For example, you can modify account privileges with `GRANT` and `REVOKE`
statements, or modify various account attributes with the `ALTER USER` statement.

Connecting to Your DB Instance Using IAM Authentication

With IAM database authentication, you use an authentication token when you connect to your DB
instance. An *authentication token* is a string of characters that you use instead of a password. After you
generate an authentication token, it's valid for 15 minutes before it expires. If you try to connect using
an expired token, the connection request is denied.

Every authentication token must be accompanied by a valid signature, using AWS signature version 4.
(For more information, see [Signature Version 4 Signing Process](#) in the *AWS General Reference*.) The AWS
CLI and the AWS SDK for Java can automatically sign each token you create.

You can use an authentication token when you connect to Amazon RDS from another AWS service, such
as AWS Lambda. By using a token, you can avoid placing a password in your code. Alternatively, you
can use the AWS SDK for Java to programmatically create and programmatically sign an authentication
token.

After you have a signed IAM authentication token, you can connect to an Amazon RDS DB instance.
Following, you can find out how to do this using either a command line tool or the AWS SDK for Java.

For more information, see [Use IAM authentication to connect with SQL Workbench/J to Amazon Aurora
MySQL or Amazon RDS for MySQL](#).

Topics

- [Connecting to Your DB Instance Using IAM Authentication from the Command Line: AWS CLI and mysql Client \(p. 468\)](#)
- [Connecting to Your DB Instance Using IAM Authentication from the Command Line: AWS CLI and psql Client \(p. 470\)](#)
- [Connecting to Your DB Instance Using IAM Authentication and the AWS SDK for Java \(p. 471\)](#)

Connecting to Your DB Instance Using IAM Authentication from the Command Line: AWS CLI and mysql Client

You can connect from the command line to an Amazon RDS DB instance with the AWS CLI and `mysql`
command line tool as described following.

Topics

- [Generating an IAM Authentication Token \(p. 468\)](#)
- [Connecting to a DB Instance \(p. 469\)](#)

Generating an IAM Authentication Token

The following example shows how to get a signed authentication token using the AWS CLI.

```
aws rds generate-db-auth-token \
--hostname rdsmysql.cdgmuiqadpid.us-west-2.rds.amazonaws.com \
```

```
--port 3306 \
--region us-west-2 \
--username jane_doe
```

In the example, the parameters are as follows:

- **--hostname** – The host name of the DB instance that you want to access.
- **--port** – The port number used for connecting to your DB instance.
- **--region** – The AWS Region where the DB instance is running.
- **--username** – The database account that you want to access.

The first several characters of the token look like the following.

```
rdsmysql.cdgmuqiadpid.us-west-2.rds.amazonaws.com:3306/?Action=connect&DBUser=jane_doe&X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Expires=900...
```

Connecting to a DB Instance

The general format for connecting is shown following.

```
mysql --host=hostName --port=portNumber --ssl-ca=[full path]rds-combined-ca-bundle.pem --enable-cleartext-plugin --user=userName --password=authToken
```

The parameters are as follows:

- **--host** – The host name of the DB instance that you want to access.
- **--port** – The port number used for connecting to your DB instance.
- **--ssl-ca** – The SSL certificate file that contains the public key. For more information, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).
- **--enable-cleartext-plugin** – A value that specifies that `AWSAuthenticationPlugin` must be used for this connection.
- **--user** – The database account that you want to access.
- **--password** – A signed IAM authentication token.

The authentication token consists of several hundred characters. It can be unwieldy on the command line. One way to work around this is to save the token to an environment variable, and then use that variable when you connect. The following example shows one way to perform this workaround.

```
RDSHOST="rdsmysql.cdgmuqiadpid.us-west-2.rds.amazonaws.com"
TOKEN=$(aws rds generate-db-auth-token --hostname $RDSHOST --port 3306 --region us-west-2
--username jane_doe )"

mysql --host=$RDSHOST --port=3306 --ssl-ca=/sample_dir/rds-combined-ca-bundle.pem --enable-
cleartext-plugin --user=jane_doe --password=$TOKEN
```

When you connect using `AWSAuthenticationPlugin`, the connection is secured using SSL. To verify this, type the following at the `mysql>` command prompt.

```
show status like 'Ssl%';
```

The following lines in the output show more details.

Variable_name	Value
...	...
Ssl_cipher	AES256-SHA
...	...
Ssl_version	TLSv1.1
...	...

Connecting to Your DB Instance Using IAM Authentication from the Command Line: AWS CLI and psql Client

You can connect from the command line to an Amazon RDS for PostgreSQL DB instance with the AWS CLI and psql command line tool as described following.

Topics

- [Generating an IAM Authentication Token \(p. 470\)](#)
- [Connecting to an Amazon RDS PostgreSQL Instance \(p. 470\)](#)

Generating an IAM Authentication Token

The authentication token consists of several hundred characters so it can be unwieldy on the command line. One way to work around this is to save the token to an environment variable, and then use that variable when you connect. The following example shows how to use the AWS CLI to get a signed authentication token using the `generate-db-auth-token` command, and store it in a `PGPASSWORD` environment variable.

```
export RDSSHOST="rdspostgres.cdgmuqiadpid.us-west-2.rds.amazonaws.com"
export PGPASSWORD="$(aws rds generate-db-auth-token --hostname $RDSSHOST --port 5432 --
region us-west-2 --username jane_doe )"
```

In the example, the parameters to the `generate-db-auth-token` command are as follows:

- `--hostname` – The host name of the DB instance that you want to access.
- `--port` – The port number used for connecting to your DB instance.
- `--region` – The AWS Region where the DB instance is running.
- `--username` – The database account that you want to access.

The first several characters of the generated token look like the following.

```
rdspostgres.cdgmuqiadpid.us-west-2.rds.amazonaws.com:5432/?
Action=connect&DBUser=jane_doe&X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Expires=900...
```

Connecting to an Amazon RDS PostgreSQL Instance

The general format for using psql to connect is shown following.

```
psql "host=hostName port=portNumber sslmode=verify-full sslrootcert=certificateFile  
dbname=DBName user=userName"
```

The parameters are as follows:

- **host** – The host name of the DB instance that you want to access.
- **port** – The port number used for connecting to your DB instance.
- **sslmode** – The SSL mode to use. When you use `sslmode=verify-full`, the SSL connection verifies the DB instance endpoint against the endpoint in the SSL certificate.
- **sslrootcert** – The SSL certificate file that contains the public key. For more information, see [Using SSL with a PostgreSQL DB Instance](#).
- **dbname** – The database that you want to access.
- **user** – The database account that you want to access.

The following example shows using the command to connect. The example uses the environment variables that were set when the token was generated in the previous section.

```
psql "host=$RDSHOST port=5432 sslmode=verify-full sslrootcert=/sample_dir/rds-combined-ca-bundle.pem dbname=DBName user=jane_doe"
```

Connecting to Your DB Instance Using IAM Authentication and the AWS SDK for Java

You can connect from the command line to an Amazon RDS MySQL or PostgreSQL DB instance with the AWS SDK for Java as described following.

Topics

- [Generating an IAM Authentication Token \(p. 471\)](#)
- [Manually Constructing an IAM Authentication Token \(p. 472\)](#)
- [Connecting to a DB Instance \(p. 475\)](#)

Generating an IAM Authentication Token

If you are writing programs using the AWS SDK for Java, you can get a signed authentication token using the `RdsIamAuthTokenGenerator` class. Using this class requires that you provide AWS credentials. To do this, you create an instance of the `DefaultAWSCredentialsProviderChain` class. `DefaultAWSCredentialsProviderChain` uses the first AWS access key and secret key that it finds in the [default credential provider chain](#). For more information about AWS access keys, see [Managing Access Keys for IAM Users](#).

After you create an instance of `RdsIamAuthTokenGenerator`, you can call the `getAuthToken` method to obtain a signed token. Provide the AWS Region, host name, port number, and user name. The following code example illustrates how to do this.

```
package com.amazonaws.codesamples;

import com.amazonaws.auth.DefaultAWSCredentialsProviderChain;
import com.amazonaws.services.rds.auth.GetIamAuthTokenRequest;
import com.amazonaws.services.rds.auth.RdsIamAuthTokenGenerator;

public class GenerateRDSToken {

    public static void main(String[] args) {

        String region = "us-west-2";
```

```
String hostname = "rdsmysql.cdgmuqiadpid.us-west-2.rds.amazonaws.com";
String port = "3306";
String username = "jane_doe";

System.out.println(generateAuthToken(region, hostname, port, username));

}

static String generateAuthToken(String region, String hostName, String port, String
username) {

RdsIamAuthTokenGenerator generator = RdsIamAuthTokenGenerator.builder()
.credentials(new DefaultAWSCredentialsProviderChain())
.region(region)
.build();

String authToken = generator.getAuthToken(
GetIamAuthTokenRequest.builder()
.hostname(hostName)
.port(Integer.parseInt(port))
.userName(username)
.build());

return authToken;
}

}
```

Manually Constructing an IAM Authentication Token

In Java, the easiest way to generate an authentication token is to use `RdsIamAuthTokenGenerator`. This class creates an authentication token for you, and then signs it using AWS signature version 4. For more information, see [Signature Version 4 Signing Process](#) in the *AWS General Reference*.

However, you can also construct and sign an authentication token manually, as shown in the following code example.

```
package com.amazonaws.codesamples;

import com.amazonaws.SdkClientException;
import com.amazonaws.auth.DefaultAWSCredentialsProviderChain;
import com.amazonaws.auth.SigningAlgorithm;
import com.amazonaws.util.BinaryUtils;
import org.apache.commons.lang3.StringUtils;

import javax.crypto.Mac;
import javax.crypto.spec.SecretKeySpec;
import java.nio.charset.Charset;
import java.security.MessageDigest;
import java.text.SimpleDateFormat;
import java.util.Date;
import java.util.SortedMap;
import java.util.TreeMap;

import static com.amazonaws.auth.internal.SignerConstants.AWS4_TERMINATOR;
import static com.amazonaws.util.StringUtils.UTF8;

public class CreateRDSAuthTokenManually {
    public static String httpMethod = "GET";
    public static String action = "connect";
    public static String canonicalURIParameter = "/";
    public static SortedMap<String, String> canonicalQueryParameters = new TreeMap();
    public static String payload = StringUtils.EMPTY;
    public static String signedHeader = "host";
    public static String algorithm = "AWS4-HMAC-SHA256";
```

```

public static String serviceName = "rds-db";
public static String requestWithoutSignature;

public static void main(String[] args) throws Exception {

    String region = "us-west-2";
    String instanceName = "rdsmysql.cdgmuqiadpid.us-west-2.rds.amazonaws.com";
    String port = "3306";
    String username = "jane_doe";

    Date now = new Date();
    String date = new SimpleDateFormat("yyyyMMdd").format(now);
    String dateTimeStamp = new SimpleDateFormat("yyyyMMdd'T'HHmmssZ").format(now);
    DefaultAWSCredentialsProviderChain creds = new
DefaultAWSCredentialsProviderChain();
    String awsAccessKey = creds.getCredentials().getAWSAccessKeyId();
    String awsSecretKey = creds.getCredentials().getAWSSecretKey();
    String expiryMinutes = "900";

    System.out.println("Step 1: Create a canonical request:");
    String canonicalString = createCanonicalString(username, awsAccessKey, date,
dateTimeStamp, region, expiryMinutes, instanceName, port);
    System.out.println(canonicalString);
    System.out.println();

    System.out.println("Step 2: Create a string to sign:");
    String stringToSign = createStringToSign(dateTimeStamp, canonicalString,
awsAccessKey, date, region);
    System.out.println(stringToSign);
    System.out.println();

    System.out.println("Step 3: Calculate the signature:");
    String signature = BinaryUtils.toHex(calculateSignature(stringToSign,
newSigningKey(awsSecretKey, date, region, serviceName)));
    System.out.println(signature);
    System.out.println();

    System.out.println("Step 4: Add the signing info to the request");
    System.out.println	appendSignature(signature));
    System.out.println();

}

//Step 1: Create a canonical request date should be in format YYYYMMDD and date
should be in format YYYYMMDDTHHMMSSZ
public static String createCanonicalString(String user, String accessKey, String date,
String date, String region, String expiryPeriod, String hostName, String port) throws
Exception {
    canonicalQueryParameters.put("Action", action);
    canonicalQueryParameters.put("DBUser", user);
    canonicalQueryParameters.put("X-Amz-Algorithm", "AWS4-HMAC-SHA256");
    canonicalQueryParameters.put("X-Amz-Credential", accessKey + "%2F" + date + "%2F" +
region + "%2F" + serviceName + "%2Faws4_request");
    canonicalQueryParameters.put("X-Amz-Date", date);
    canonicalQueryParameters.put("X-Amz-Expires", expiryPeriod);
    canonicalQueryParameters.put("X-Amz-SignedHeaders", signedHeader);
    String canonicalQueryString = "";
    while(!canonicalQueryParameters.isEmpty()) {
        String currentQueryParameter = canonicalQueryParameters.firstKey();
        String currentQueryParameterValue =
canonicalQueryParameters.remove(currentQueryParameter);
        canonicalQueryString = canonicalQueryString + currentQueryParameter + "=" +
currentQueryParameterValue;
        if (!currentQueryParameter.equals("X-Amz-SignedHeaders")) {
            canonicalQueryString += "&";
        }
    }
}

```

```

        }

        String canonicalHeaders = "host:" + hostName + ":" + port + '\n';
        requestWithoutSignature = hostName + ":" + port + "/" + canonicalQueryString;

        String hashedPayload = BinaryUtils.toHex(hash(payload));
        return httpMethod + '\n' + canonicalURIParameter + '\n' + canonicalQueryString +
        '\n' + canonicalHeaders + '\n' + signedHeader + '\n' + hashedPayload;

    }

    //Step 2: Create a string to sign using sig v4
    public static String createStringToSign(String dateTime, String canonicalRequest,
    String accessKey, String date, String region) throws Exception {
        String credentialScope = date + "/" + region + "/" + serviceName + "/aws4_request";
        return algorithm + '\n' + dateTime + '\n' + credentialScope + '\n' +
        BinaryUtils.toHex(hash(canonicalRequest));

    }

    //Step 3: Calculate signature
    /**
     * Step 3 of the AWS Signature version 4 calculation. It involves deriving
     * the signing key and computing the signature. Refer to
     * http://docs.aws.amazon.com/general/latest/gr/sigv4-calculate-signature.html
     */
    public static byte[] calculateSignature(String stringToSign,
                                           byte[] signingKey) {
        return sign(stringToSign.getBytes(Charset.forName("UTF-8")), signingKey,
                   SigningAlgorithm.HmacSHA256);
    }

    public static byte[] sign(byte[] data, byte[] key,
                           SigningAlgorithm algorithm) throwsSdkClientException {
        try {
            Mac mac = algorithm.getMac();
            mac.init(new SecretKeySpec(key, algorithm.toString()));
            return mac.doFinal(data);
        } catch (Exception e) {
            throw new SdkClientException(
                "Unable to calculate a request signature: "
                + e.getMessage(), e);
        }
    }

    public static byte[] newSigningKey(String secretKey,
                                      String dateStamp, String regionName, String serviceName)
    {
        byte[] kSecret = ("AWS4" + secretKey).getBytes(Charset.forName("UTF-8"));
        byte[] kDate = sign(dateStamp, kSecret, SigningAlgorithm.HmacSHA256);
        byte[] kRegion = sign(regionName, kDate, SigningAlgorithm.HmacSHA256);
        byte[] kService = sign(serviceName, kRegion,
                               SigningAlgorithm.HmacSHA256);
        return sign(AWS4_TERMINATOR, kService, SigningAlgorithm.HmacSHA256);
    }

    public static byte[] sign(String stringData, byte[] key,
                           SigningAlgorithm algorithm) throwsSdkClientException {
        try {
            byte[] data = stringData.getBytes(UTF8);
            return sign(data, key, algorithm);
        } catch (Exception e) {
            throw new SdkClientException(
                "Unable to calculate a request signature: "
                + e.getMessage(), e);
        }
    }
}

```

```

        }

    //Step 4: append the signature
    public static String appendSignature(String signature) {
        return requestWithoutSignature + "&X-Amz-Signature=" + signature;
    }

    public static byte[] hash(String s) throws Exception {
        try {
            MessageDigest md = MessageDigest.getInstance("SHA-256");
            md.update(s.getBytes(UTF8));
            return md.digest();
        } catch (Exception e) {
            throw new SdkClientException(
                "Unable to compute hash while signing request: "
                + e.getMessage(), e);
        }
    }
}

```

Connecting to a DB Instance

The following code example shows how to generate an authentication token, and then use it to connect to an instance running MySQL.

To run this code example, you need the [AWS SDK for Java](#), found on the AWS site. In addition, you need the following:

- MySQL Connector/J. This code example was tested with `mysql-connector-java-5.1.33-bin.jar`.
- An intermediate certificate for Amazon RDS that is specific to an AWS Region. (For more information, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).) At runtime, the class loader looks for the certificate in the same directory as this Java code example, so that the class loader can find it.
- Modify the values of the following variables as needed:
 - `RDS_INSTANCE_HOSTNAME` – The host name of the DB instance that you want to access.
 - `RDS_INSTANCE_PORT` – The port number used for connecting to your PostgreSQL DB instance.
 - `REGION_NAME` – The AWS Region where the DB instance is running.
 - `DB_USER` – The database account that you want to access.
 - `SSL_CERTIFICATE` – An SSL certificate for Amazon RDS that is specific to an AWS Region.

To download a certificate for your AWS Region, see [Intermediate Certificates \(p. 432\)](#). Place the SSL certificate in the same directory as this Java program file, so that the class loader can find the certificate at runtime.

This code example obtains AWS credentials from the [default credential provider chain](#).

```

package com.amazonaws.samples;

import com.amazonaws.services.rds.auth.RdsIamAuthTokenGenerator;
import com.amazonaws.services.rds.auth.GetIamAuthTokenRequest;
import com.amazonaws.auth.BasicAWSCredentials;
import com.amazonaws.auth.DefaultAWSCredentialsProviderChain;
import com.amazonaws.auth.AWSStaticCredentialsProvider;

import java.io.File;
import java.io.FileOutputStream;
import java.io.InputStream;
import java.security.KeyStore;

```

```

import java.security.cert.CertificateFactory;
import java.security.cert.X509Certificate;

import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.ResultSet;
import java.sql.Statement;
import java.util.Properties;

import java.net.URL;

public class IAMDatabaseAuthenticationTester {
    //AWS Credentials of the IAM user with policy enabling IAM Database Authenticated
    access to the db by the db user.
    private static final DefaultAWSCredentialsProviderChain creds = new
DefaultAWSCredentialsProviderChain();
    private static final String AWS_ACCESS_KEY =
creds.getCredentials().getAWSAccessKeyId();
    private static final String AWS_SECRET_KEY = creds.getCredentials().getAWSSecretKey();

    //Configuration parameters for the generation of the IAM Database Authentication token
    private static final String RDS_INSTANCE_HOSTNAME = "rdsmysql.cdgmuqiadpid.us-
west-2.rds.amazonaws.com";
    private static final int RDS_INSTANCE_PORT = 3306;
    private static final String REGION_NAME = "us-west-2";
    private static final String DB_USER = "jane_doe";
    private static final String JDBC_URL = "jdbc:mysql://" + RDS_INSTANCE_HOSTNAME + ":" +
RDS_INSTANCE_PORT;

    private static final String SSL_CERTIFICATE = "rds-ca-2015-us-west-2.pem";

    private static final String KEY_STORE_TYPE = "JKS";
    private static final String KEY_STORE_PROVIDER = "SUN";
    private static final String KEY_STORE_FILE_PREFIX = "sys-connect-via-ssl-test-cacerts";
    private static final String KEY_STORE_FILE_SUFFIX = ".jks";
    private static final String DEFAULT_KEY_STORE_PASSWORD = "changeit";

    public static void main(String[] args) throws Exception {
        //get the connection
        Connection connection = getDBConnectionUsingIam();

        //verify the connection is successful
        Statement stmt= connection.createStatement();
        ResultSet rs=stmt.executeQuery("SELECT 'Success!' FROM DUAL;");
        while (rs.next()) {
            String id = rs.getString(1);
            System.out.println(id); //Should print "Success!"
        }

        //close the connection
        stmt.close();
        connection.close();

        clearSslProperties();
    }

    /**
     * This method returns a connection to the db instance authenticated using IAM Database
     Authentication
     * @return
     * @throws Exception
     */
    private static Connection getDBConnectionUsingIam() throws Exception {
        setSslProperties();
        return DriverManager.getConnection(JDBC_URL, setMySqlConnectionProperties());
    }
}

```

```

    }

    /**
     * This method sets the mysql connection properties which includes the IAM Database
     * Authentication token
     * as the password. It also specifies that SSL verification is required.
     * @return
     */
    private static Properties setMySqlConnectionProperties() {
        Properties mysqlConnectionProperties = new Properties();
        mysqlConnectionProperties.setProperty("verifyServerCertificate","true");
        mysqlConnectionProperties.setProperty("useSSL", "true");
        mysqlConnectionProperties.setProperty("user",DB_USER);
        mysqlConnectionProperties.setProperty("password",generateAuthToken());
        return mysqlConnectionProperties;
    }

    /**
     * This method generates the IAM Auth Token.
     * An example IAM Auth Token would look like follows:
     * btusi123.cmz7kenwo2ye.rds.cn-north-1.amazonaws.com.cn:3306/?  

Action=connect&DBUser=iamtestuser&X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-  

Date=20171003T010726Z&X-Amz-SignedHeaders=host&X-Amz-Expires=899&X-Amz-  

Credential=AKIAPFXHGVDI5RNFO4AQ%2F20171003%2Fcn-north-1%2Frds-db%2Faws4_request&X-Amz-  

Signature=f9f45ef96c1f770cdad11a53e33ffa4c3730bc03fdee820cfdf1322eed15483b
     * @return
     */
    private static String generateAuthToken() {
        BasicAWSCredentials awsCredentials = new BasicAWSCredentials(AWS_ACCESS_KEY,
        AWS_SECRET_KEY);

        RdsIamAuthTokenGenerator generator = RdsIamAuthTokenGenerator.builder()
            .credentials(new
        AWSStaticCredentialsProvider(awsCredentials)).region(REGION_NAME).build();
        return generator.getAuthToken(GetIamAuthTokenRequest.builder()

        .hostname(RDS_INSTANCE_HOSTNAME).port(RDS_INSTANCE_PORT).userName(DB_USER).build());
    }

    /**
     * This method sets the SSL properties which specify the key store file, its type and
     * password:
     * @throws Exception
     */
    private static void setSslProperties() throws Exception {
        System.setProperty("javax.net.ssl.trustStore", createKeyStoreFile());
        System.setProperty("javax.net.ssl.trustStoreType", KEY_STORE_TYPE);
        System.setProperty("javax.net.ssl.trustStorePassword", DEFAULT_KEY_STORE_PASSWORD);
    }

    /**
     * This method returns the path of the Key Store File needed for the SSL verification
     * during the IAM Database Authentication to
     * the db instance.
     * @return
     * @throws Exception
     */
    private static String createKeyStoreFile() throws Exception {
        return createKeyStoreFile(createCertificate()).getPath();
    }

    /**
     * This method generates the SSL certificate
     * @return
     * @throws Exception
     */
}

```

```
private static X509Certificate createCertificate() throws Exception {
    CertificateFactory certFactory = CertificateFactory.getInstance("X.509");
    URL url = new File(SSL_CERTIFICATE).toURI().toURL();
    if (url == null) {
        throw new Exception();
    }
    try (InputStream certInputStream = url.openStream()) {
        return (X509Certificate) certFactory.generateCertificate(certInputStream);
    }
}

/**
 * This method creates the Key Store File
 * @param rootX509Certificate - the SSL certificate to be stored in the KeyStore
 * @return
 * @throws Exception
 */
private static File createKeyStoreFile(X509Certificate rootX509Certificate) throws
Exception {
    File keyStoreFile = File.createTempFile(KEY_STORE_FILE_PREFIX,
KEY_STORE_FILE_SUFFIX);
    try (FileOutputStream fos = new FileOutputStream(keyStoreFile.getPath())) {
        KeyStore ks = KeyStore.getInstance(KEY_STORE_TYPE, KEY_STORE_PROVIDER);
        ks.load(null);
        ks.setCertificateEntry("rootCaCertificate", rootX509Certificate);
        ks.store(fos, DEFAULT_KEY_STORE_PASSWORD.toCharArray());
    }
    return keyStoreFile;
}

/**
 * This method clears the SSL properties.
 * @throws Exception
 */
private static void clearSslProperties() throws Exception {
    System.clearProperty("javax.net.ssl.trustStore");
    System.clearProperty("javax.net.ssl.trustStoreType");
    System.clearProperty("javax.net.ssl.trustStorePassword");
}
```

Troubleshooting Amazon RDS Identity and Access

Use the following information to help you diagnose and fix common issues that you might encounter when working with Amazon RDS and IAM.

Topics

- [I Am Not Authorized to Perform an Action in Amazon RDS \(p. 478\)](#)
- [I Am Not Authorized to Perform iam:PassRole \(p. 479\)](#)
- [I Want to View My Access Keys \(p. 479\)](#)
- [I'm an Administrator and Want to Allow Others to Access Amazon RDS \(p. 479\)](#)
- [I Want to Allow People Outside of My AWS Account to Access My Amazon RDS Resources \(p. 480\)](#)

I Am Not Authorized to Perform an Action in Amazon RDS

If the AWS Management Console tells you that you're not authorized to perform an action, then you must contact your administrator for assistance. Your administrator is the person that provided you with your user name and password.

The following example error occurs when the `mateojackson` IAM user tries to use the console to view details about a `widget` but does not have `rds:GetWidget` permissions.

```
User: arn:aws:iam::123456789012:user/mateojackson is not authorized to perform:  
rds:GetWidget on resource: my-example-widget
```

In this case, Mateo asks his administrator to update his policies to allow him to access the `my-example-widget` resource using the `rds:GetWidget` action.

I Am Not Authorized to Perform iam:PassRole

If you receive an error that you're not authorized to perform the `iam:PassRole` action, then you must contact your administrator for assistance. Your administrator is the person that provided you with your user name and password. Ask that person to update your policies to allow you to pass a role to Amazon RDS.

Some AWS services allow you to pass an existing role to that service, instead of creating a new service role or service-linked role. To do this, you must have permissions to pass the role to the service.

The following example error occurs when an IAM user named `marymajor` tries to use the console to perform an action in Amazon RDS. However, the action requires the service to have permissions granted by a service role. Mary does not have permissions to pass the role to the service.

```
User: arn:aws:iam::123456789012:user/marymajor is not authorized to perform: iam:PassRole
```

In this case, Mary asks her administrator to update her policies to allow her to perform the `iam:PassRole` action.

I Want to View My Access Keys

After you create your IAM user access keys, you can view your access key ID at any time. However, you can't view your secret access key again. If you lose your secret key, you must create a new access key pair.

Access keys consist of two parts: an access key ID (for example, `AKIAIOSFODNN7EXAMPLE`) and a secret access key (for example, `wJalrXUtnFEMI/K7MDENG/bPxRfiCYEXAMPLEKEY`). Like a user name and password, you must use both the access key ID and secret access key together to authenticate your requests. Manage your access keys as securely as you do your user name and password.

Important

Do not provide your access keys to a third party, even to help [find your canonical user ID](#). By doing this, you might give someone permanent access to your account.

When you create an access key pair, you are prompted to save the access key ID and secret access key in a secure location. The secret access key is available only at the time you create it. If you lose your secret access key, you must add new access keys to your IAM user. You can have a maximum of two access keys. If you already have two, you must delete one key pair before creating a new one. To view instructions, see [Managing Access Keys](#) in the *IAM User Guide*.

I'm an Administrator and Want to Allow Others to Access Amazon RDS

To enable others to access Amazon RDS, you must create an IAM entity (user or role) for the person or application that needs access. They use the credentials for that entity to access AWS. You must then attach a policy to the entity that grants them the correct permissions in Amazon RDS.

To get started right away, see [Creating Your First IAM Delegated User and Group](#) in the *IAM User Guide*.

I Want to Allow People Outside of My AWS Account to Access My Amazon RDS Resources

You can create a role that users in other accounts or people outside of your organization can use to access your resources. You can specify who is trusted to assume the role. For services that support resource-based policies or access control lists (ACLs), you can use those policies to grant people access to your resources.

To learn more, consult the following:

- To learn whether Amazon RDS supports these features, see [How Amazon RDS Works with IAM \(p. 449\)](#).
- To learn how to provide access to your resources across AWS accounts that you own, see [Providing Access to an IAM User in Another AWS Account That You Own](#) in the *IAM User Guide*.
- To learn how to provide access to your resources to third-party AWS accounts, see [Providing Access to AWS Accounts Owned by Third Parties](#) in the *IAM User Guide*.
- To learn how to provide access through identity federation, see [Providing Access to Externally Authenticated Users \(Identity Federation\)](#) in the *IAM User Guide*.
- To learn the difference between using roles and resource-based policies for cross-account access, see [How IAM Roles Differ from Resource-based Policies](#) in the *IAM User Guide*.

Kerberos Authentication

Amazon RDS supports external authentication of database users using Kerberos and Microsoft Active Directory. Kerberos is a network authentication protocol that uses tickets and symmetric-key cryptography to eliminate the need to transmit passwords over the network. Kerberos has been built into Active Directory and is designed to authenticate users to network resources, such as databases.

Amazon RDS support for Kerberos and Active Directory provides the benefits of single sign-on and centralized authentication of database users. You can keep your user credentials in Active Directory. Active Directory provides a centralized place for storing and managing credentials for multiple DB instances.

You can enable your database users to authenticate against DB instances in two ways. They can use credentials stored either in AWS Directory Service for Microsoft Active Directory or in your on-premises Active Directory.

Currently, RDS supports Kerberos authentication for Microsoft SQL Server, Oracle, and PostgreSQL DB instances. For information about Kerberos authentication with a specific engine, see the following:

- [Using Windows Authentication with a Microsoft SQL Server DB Instance \(p. 691\)](#)
- [Using Kerberos Authentication with Amazon RDS for Oracle \(p. 1134\)](#)
- [Using Kerberos Authentication with Amazon RDS for PostgreSQL \(p. 1351\)](#)

Logging and Monitoring in Amazon RDS

Monitoring is an important part of maintaining the reliability, availability, and performance of Amazon RDS and your AWS solutions. You should collect monitoring data from all of the parts of your AWS solution so that you can more easily debug a multi-point failure if one occurs. AWS provides several tools for monitoring your Amazon RDS resources and responding to potential incidents:

Amazon CloudWatch Alarms

Using Amazon CloudWatch alarms, you watch a single metric over a time period that you specify. If the metric exceeds a given threshold, a notification is sent to an Amazon SNS topic or AWS Auto Scaling policy. CloudWatch alarms do not invoke actions because they are in a particular state. Rather the state must have changed and been maintained for a specified number of periods. For more information, see [Monitoring with Amazon CloudWatch \(p. 297\)](#).

AWS CloudTrail Logs

CloudTrail provides a record of actions taken by a user, role, or an AWS service in Amazon RDS. CloudTrail captures all API calls for Amazon RDS as events, including calls from the console and from code calls to Amazon RDS API operations. Using the information collected by CloudTrail, you can determine the request that was made to Amazon RDS, the IP address from which the request was made, who made the request, when it was made, and additional details. For more information, see [Logging Amazon RDS API Calls with AWS CloudTrail \(p. 423\)](#).

Enhanced Monitoring

Amazon RDS provides metrics in real time for the operating system (OS) that your DB instance runs on. You can view the metrics for your DB instance using the console, or consume the Enhanced Monitoring JSON output from Amazon CloudWatch Logs in a monitoring system of your choice. For more information, see [Enhanced Monitoring \(p. 307\)](#).

Amazon RDS Performance Insights

Performance Insights expands on existing Amazon RDS monitoring features to illustrate your database's performance and help you analyze any issues that affect it. With the Performance Insights dashboard, you can visualize the database load and filter the load by waits, SQL statements, hosts, or users. For more information, see [Using Amazon RDS Performance Insights \(p. 319\)](#).

Database Logs

You can view, download, and watch database logs using the AWS Management Console, AWS CLI, or RDS API. For more information, see [Amazon RDS Database Log Files \(p. 390\)](#).

Amazon RDS Recommendations

Amazon RDS provides automated recommendations for database resources. These recommendations provide best practice guidance by analyzing DB instance configuration, usage, and performance data. For more information, see [Using Amazon RDS Recommendations \(p. 365\)](#).

Amazon RDS Event Notification

Amazon RDS uses the Amazon Simple Notification Service (Amazon SNS) to provide notification when an Amazon RDS event occurs. These notifications can be in any notification form supported by Amazon SNS for an AWS Region, such as an email, a text message, or a call to an HTTP endpoint. For more information, see [Using Amazon RDS Event Notification \(p. 370\)](#).

AWS Trusted Advisor

Trusted Advisor draws upon best practices learned from serving hundreds of thousands of AWS customers. Trusted Advisor inspects your AWS environment and then makes recommendations when opportunities exist to save money, improve system availability and performance, or help close security gaps. All AWS customers have access to five Trusted Advisor checks. Customers with a Business or Enterprise support plan can view all Trusted Advisor checks.

Trusted Advisor has the following Amazon RDS-related checks:

- Amazon RDS Idle DB Instances
- Amazon RDS Security Group Access Risk
- Amazon RDS Backups
- Amazon RDS Multi-AZ

For more information on these checks, see [Trusted Advisor Best Practices \(Checks\)](#).

For more information about monitoring Amazon RDS, see [Monitoring Amazon RDS \(p. 294\)](#).

Compliance Validation for Amazon RDS

Third-party auditors assess the security and compliance of Amazon RDS as part of multiple AWS compliance programs. These include SOC, PCI, FedRAMP, HIPAA, and others.

For a list of AWS services in scope of specific compliance programs, see [AWS Services in Scope by Compliance Program](#). For general information, see [AWS Compliance Programs](#).

You can download third-party audit reports using AWS Artifact. For more information, see [Downloading Reports in AWS Artifact](#).

Your compliance responsibility when using Amazon RDS is determined by the sensitivity of your data, your organization's compliance objectives, and applicable laws and regulations. AWS provides the following resources to help with compliance:

- [Security and Compliance Quick Start Guides](#) – These deployment guides discuss architectural considerations and provide steps for deploying security- and compliance-focused baseline environments on AWS.
- [Architecting for HIPAA Security and Compliance Whitepaper](#) – This whitepaper describes how companies can use AWS to create HIPAA-compliant applications.
- [AWS Compliance Resources](#) – This collection of workbooks and guides that might apply to your industry and location.
- [AWS Config](#) – This AWS service assesses how well your resource configurations comply with internal practices, industry guidelines, and regulations.
- [AWS Security Hub](#) – This AWS service provides a comprehensive view of your security state within AWS that helps you check your compliance with security industry standards and best practices.

Resilience in Amazon RDS

The AWS global infrastructure is built around AWS Regions and Availability Zones. AWS Regions provide multiple physically separated and isolated Availability Zones, which are connected with low-latency, high-throughput, and highly redundant networking. With Availability Zones, you can design and operate applications and databases that automatically fail over between Availability Zones without interruption. Availability Zones are more highly available, fault tolerant, and scalable than traditional single or multiple data center infrastructures.

For more information about AWS Regions and Availability Zones, see [AWS Global Infrastructure](#).

In addition to the AWS global infrastructure, Amazon RDS offers features to help support your data resiliency and backup needs.

Backup and Restore

Amazon RDS creates and saves automated backups of your DB instance. Amazon RDS creates a storage volume snapshot of your DB instance, backing up the entire DB instance and not just individual databases.

Amazon RDS creates automated backups of your DB instance during the backup window of your DB instance. Amazon RDS saves the automated backups of your DB instance according to the backup retention period that you specify. If necessary, you can recover your database to any point in time during the backup retention period. You can also back up your DB instance manually, by manually creating a DB snapshot.

You can create a DB instance by restoring from this DB snapshot as a disaster recovery solution if the source DB instance fails.

For more information, see [Backing Up and Restoring Amazon RDS DB Instances \(p. 253\)](#).

Replication

Amazon RDS uses the MariaDB, MySQL, Oracle, and PostgreSQL DB engines' built-in replication functionality to create a special type of DB instance called a Read Replica from a source DB instance. Updates made to the source DB instance are asynchronously copied to the Read Replica. You can reduce the load on your source DB instance by routing read queries from your applications to the Read Replica. Using Read Replicas, you can elastically scale out beyond the capacity constraints of a single DB instance for read-heavy database workloads. You can promote a Read Replica to a standalone instance as a disaster recovery solution if the source DB instance fails. For some DB engines, Amazon RDS also supports other replication options.

For more information, see [Working with Read Replicas \(p. 182\)](#).

Failover

Amazon RDS provides high availability and failover support for DB instances using Multi-AZ deployments. Amazon RDS uses several different technologies to provide failover support. Multi-AZ deployments for Oracle, PostgreSQL, MySQL, and MariaDB DB instances use Amazon's failover technology. SQL Server DB instances use SQL Server Database Mirroring (DBM).

For more information, see [High Availability \(Multi-AZ\) for Amazon RDS \(p. 144\)](#).

Infrastructure Security in Amazon RDS

As a managed service, Amazon RDS is protected by the AWS global network security procedures that are described in the [Amazon Web Services: Overview of Security Processes](#) whitepaper.

You use AWS published API calls to access Amazon RDS through the network. Clients must support Transport Layer Security (TLS) 1.0. We recommend TLS 1.2 or later. Clients must also support cipher suites with perfect forward secrecy (PFS) such as Ephemeral Diffie-Hellman (DHE) or Elliptic Curve Ephemeral Diffie-Hellman (ECDHE). Most modern systems such as Java 7 and later support these modes.

Additionally, requests must be signed by using an access key ID and a secret access key that is associated with an IAM principal. Or you can use the [AWS Security Token Service](#) (AWS STS) to generate temporary security credentials to sign requests.

You can call these API operations from any network location. However, Amazon RDS also supports resource-based access policies, which can include restrictions based on the source IP address. In addition, you can use Amazon RDS policies to control access from specific Amazon VPC endpoints or specific VPCs. Effectively, this isolates network access to a given Amazon RDS resource from only the specific VPC within the AWS network.

In addition, Amazon RDS offers features to help support infrastructure security.

Security Groups

Security groups control the access that traffic has in and out of a DB instance. By default, network access is turned off to a DB instance. You can specify rules in a security group that allow access from an IP address range, port, or Amazon EC2 security group. After ingress rules are configured, the same rules apply to all DB instances that are associated with that security group.

For more information, see [Controlling Access with Security Groups \(p. 486\)](#).

Public Accessibility

When you launch a DB instance inside a virtual private cloud (VPC) based on the Amazon VPC service, you can turn on or off public accessibility for that instance. To designate whether the DB instance that you create has a DNS name that resolves to a public IP address, you use the *Public accessibility* parameter. By using this parameter, you can designate whether there is public access to the DB instance. You can modify a DB instance to turn on or off public accessibility by modifying the *Public accessibility* parameter.

For more information, see [Hiding a DB Instance in a VPC from the Internet \(p. 515\)](#).

Security Best Practices for Amazon RDS

Use AWS Identity and Access Management (IAM) accounts to control access to Amazon RDS API operations, especially operations that create, modify, or delete Amazon RDS resources. Such resources include DB instances, security groups, and parameter groups. Also use IAM to control actions that perform common administrative actions such as backing up and restoring DB instances.

- Assign an individual IAM account to each person who manages Amazon RDS resources. Don't use AWS root credentials to manage Amazon RDS resources; you should create an IAM user for everyone, including yourself.
- Grant each user the minimum set of permissions required to perform his or her duties.
- Use IAM groups to effectively manage permissions for multiple users.

- Rotate your IAM credentials regularly.
- Configure AWS Secrets Manager to automatically rotate the secrets for Amazon RDS. For more information, see [Rotating Your AWS Secrets Manager Secrets](#) in the *AWS Secrets Manager User Guide*.

For more information about IAM, see [AWS Identity and Access Management](#). For information on IAM best practices, see [IAM Best Practices](#).

Use the AWS Management Console, the AWS CLI, or the RDS API to change the password for your master user. If you use another tool, such as a SQL client, to change the master user password, it might result in privileges being revoked for the user unintentionally.

Controlling Access with Security Groups

Security groups control the access that traffic has in and out of a DB instance. Three types of security groups are used with Amazon RDS: DB security groups, VPC security groups, and Amazon EC2 security groups. In simple terms, these work as follows:

- A DB security group controls access to EC2-Classic DB instances that are not in a VPC.
- A VPC security group controls access to DB instances and EC2 instances inside a VPC.
- An EC2 security group controls access to an EC2 instance.

By default, network access is disabled for a DB instance. You can specify rules in a security group that allow access from an IP address range, port, or EC2 security group. Once ingress rules are configured, the same rules apply to all DB instances that are associated with that security group. You can specify up to 20 rules in a security group.

DB Security Groups

DB security groups are used with DB instances that are not in a VPC and on the EC2-Classic platform. Each DB security group rule enables a specific source to access a DB instance that is associated with that DB security group. The source can be a range of addresses (for example, 203.0.113.0/24), or an EC2 security group. When you specify an EC2 security group as the source, you allow incoming traffic from all EC2 instances that use that EC2 security group. DB security group rules apply to inbound traffic only; outbound traffic is not currently permitted for DB instances.

You don't need to specify a destination port number when you create DB security group rules. The port number defined for the DB instance is used as the destination port number for all rules defined for the DB security group. DB security groups can be created using the Amazon RDS API operations or the Amazon RDS page of the AWS Management Console.

For more information about working with DB security groups, see [Working with DB Security Groups \(EC2-Classic Platform\) \(p. 491\)](#).

VPC Security Groups

Each VPC security group rule enables a specific source to access a DB instance in a VPC that is associated with that VPC security group. The source can be a range of addresses (for example, 203.0.113.0/24), or another VPC security group. By specifying a VPC security group as the source, you allow incoming traffic from all instances (typically application servers) that use the source VPC security group. VPC security groups can have rules that govern both inbound and outbound traffic, though the outbound traffic rules typically do not apply to DB instances. Outbound traffic rules only apply if the DB instance acts as a client. For example, outbound traffic rules apply to an Oracle DB instance with outbound database

links. You must use the [Amazon EC2 API](#) or the **Security Group** option on the VPC Console to create VPC security groups.

When you create rules for your VPC security group that allow access to the instances in your VPC, you must specify a port for each range of addresses that the rule allows access for. For example, if you want to enable SSH access to instances in the VPC, then you create a rule allowing access to TCP port 22 for the specified range of addresses.

You can configure multiple VPC security groups that allow access to different ports for different instances in your VPC. For example, you can create a VPC security group that allows access to TCP port 80 for web servers in your VPC. You can then create another VPC security group that allows access to TCP port 3306 for RDS MySQL DB instances in your VPC.

For more information on VPC security groups, see [Security Groups](#) in the *Amazon Virtual Private Cloud User Guide*.

DB Security Groups vs. VPC Security Groups

The following table shows the key differences between DB security groups and VPC security groups.

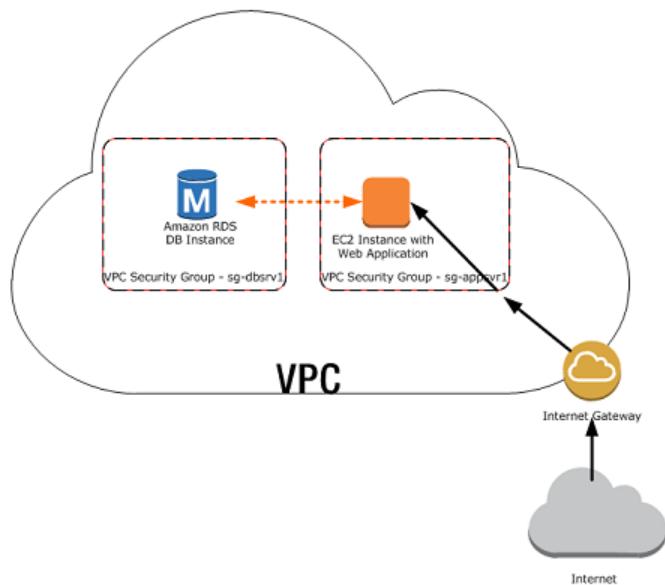
DB Security Group	VPC Security Group
Controls access to DB instances outside a VPC.	Controls access to DB instances in VPC.
Uses Amazon RDS API operations or the Amazon RDS page of the AWS Management Console to create and manage group and rules.	Uses Amazon EC2 API operations or the Amazon VPC page of the AWS Management Console to create and manage group and rules.
When you add a rule to a group, you don't need to specify port number or protocol.	When you add a rule to a group, specify the protocol as TCP. In addition, specify the same port number that you used to create the DB instances (or options) that you plan to add as members to the group.
Groups allow access from EC2 security groups in your AWS account or other accounts.	Groups allow access from other VPC security groups in your VPC only.

Security Group Scenario

A common use of a DB instance in a VPC is to share data with an application server running in an Amazon EC2 instance in the same VPC, which is accessed by a client application outside the VPC. For this scenario, you use the RDS and VPC pages on the AWS Management Console or the RDS and EC2 API operations to create the necessary instances and security groups:

1. Create a VPC security group (for example, sg-appsrv1) and define inbound rules that use the IP addresses of the client application as the source. This security group allows your client application to connect to EC2 instances in a VPC that uses this security group.
2. Create an EC2 instance for the application and add the EC2 instance to the VPC security group (sg-appsrv1) that you created in the previous step. The EC2 instance in the VPC shares the VPC security group with the DB instance.
3. Create a second VPC security group (for example, sg-dbsrv1) and create a new rule by specifying the VPC security group that you created in step 1 (sg-appsrv1) as the source.
4. Create a new DB instance and add the DB instance to the VPC security group (sg-dbsrv1) that you created in the previous step. When you create the DB instance, use the same port number as the one specified for the VPC security group (sg-dbsrv1) rule that you created in step 3.

The following diagram shows this scenario.



For more information about using a VPC, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).

Creating a VPC Security Group

You can create a VPC security group for a DB instance by using the VPC console. For information about creating a security group, see [Provide Access to Your DB Instance in Your VPC by Creating a Security Group \(p. 8\)](#) and [Security Groups](#) in the *Amazon Virtual Private Cloud User Guide*.

Associating a Security Group with a DB Instance

You can associate a security group with a DB instance by using **Modify** on the RDS console, the `ModifyDBInstance` Amazon RDS API, or the `modify-db-instance` AWS CLI command.

For information about modifying a DB instance, see [Modifying an Amazon RDS DB Instance \(p. 150\)](#). For security group considerations when you restore a DB instance from a DB snapshot, see [Security Group Considerations \(p. 266\)](#).

Deleting DB VPC Security Groups

DB VPC security groups are an RDS mechanism to synchronize security information with a VPC security group. However, this synchronization is no longer required, because RDS has been updated to use VPC security group information directly.

Note

DB VPC security groups are deprecated, and they are different from DB security groups, VPC security groups, and EC2 security groups.

We strongly recommend that you delete any DB VPC security groups that you currently use. If you don't delete your DB VPC security groups, you might encounter unintended behaviors with your DB instances, which can be as severe as losing access to a DB instance. The unintended behaviors might result from taking an action such as an update to a DB instance, a parameter group, or similar. Such updates cause RDS to resynchronize the DB VPC security group with the VPC security group. This resynchronization can

result in your security information being overwritten with incorrect and outdated security information. This in turn can have a severe impact on your ability to access to your DB instances.

How Can I Determine If I Have a DB VPC Security Group?

Because DB VPC security groups have been deprecated, they don't appear in the RDS console. However, you can call the [describe-db-security-groups](#) AWS CLI command or the [DescribeDBSecurityGroups](#) API operation to determine if you have any DB VPC security groups.

In this case, you can call the `describe-db-security-groups` AWS CLI command with JSON specified as the output format. If you do, you can identify DB VPC security groups by the VPC identifier on the second line of the output for the security group as shown in the following example.

```
{  
    "DBSecurityGroups": [  
        {  
            "VpcId": "vpc-abcd1234",  
            "DBSecurityGroupDescription": "default:vpc-abcd1234",  
            "IPRanges": [  
                {  
                    "Status": "authorized",  
                    "CIDRIP": "xxx.xxx.xxx.xxxx/n"  
                },  
                {  
                    "Status": "authorized",  
                    "CIDRIP": "xxx.xxx.xxx.xxxx/n "  
                }  
            ],  
            "OwnerId": "123456789012",  
            "EC2SecurityGroups": [],  
            "DBSecurityGroupName": "default:vpc-abcd1234"  
        }  
    ]  
}
```

If you run the `DescribeDBSecurityGroups` API operation, then you can identify DB VPC security groups using the `<VpcId>` response element as shown in the following example.

```
<DBSecurityGroup>  
    <EC2SecurityGroups/>  
    <DBSecurityGroupDescription>default:vpc-abcd1234</DBSecurityGroupDescription>  
    <IPRanges>  
        <IPRange>  
            <CIDRIP>xxx.xxx.xxx.xxxx/n</CIDRIP>  
            <Status>authorized</Status>  
        </IPRange>  
        <IPRange>  
            <CIDRIP>xxx.xxx.xxx.xxxx/n</CIDRIP>  
            <Status>authorized</Status>  
        </IPRange>  
    </IPRanges>  
    <VpcId>vpc-abcd1234</VpcId>  
    <OwnerId>123456789012</OwnerId>  
    <DBSecurityGroupName>default:vpc-abcd1234</DBSecurityGroupName>  
</DBSecurityGroup>
```

How Do I Delete a DB VPC Security Group?

Because DB VPC security groups don't appear in the RDS console, you must call the [delete-db-security-group](#) AWS CLI command or the [DeleteDBSecurityGroup](#) API operation to delete a DB VPC security group.

After you delete a DB VPC security group, your DB instances in your VPC continue to be secured by the VPC security group for that VPC. The DB VPC security group that was deleted was merely a copy of the VPC security group information.

Review Your AWS CloudFormation Templates

Older versions of AWS CloudFormation templates can contain instructions to create a DB VPC security group. Because DB VPC security groups are not yet fully deprecated, they can still be created. Make sure that any AWS CloudFormation templates that you use to provision a DB instance with security settings don't also create a DB VPC security group. Don't use AWS CloudFormation templates that create an RDS `DBSecurityGroup` with an `EC2VpcId` as shown in the following example.

```
{  
  "DbSecurityByEC2SecurityGroup" : {  
    "Type" : "AWS::RDS::DBSecurityGroup",  
    "Properties" : {  
      "GroupDescription" : "Ingress for Amazon EC2 security group",  
      "EC2VpcId" : "MyVPC",  
      "DBSecurityGroupIngress" : [ {  
        "EC2SecurityGroupId" : "sg-b0ff1111",  
        "EC2SecurityGroupOwnerId" : "111122223333"  
      }, {  
        "EC2SecurityGroupId" : "sg-ffd722222",  
        "EC2SecurityGroupOwnerId" : "111122223333"  
      } ]  
    }  
  }  
}
```

Instead, add security information for your DB instances in a VPC using VPC security groups, as shown in the following example.

```
{  
  "DBInstance" : {  
    "Type": "AWS::RDS::DBInstance",  
    "Properties": {  
      "DBName" : { "Ref" : "DBName" },  
      "Engine" : "MySQL",  
      "MultiAZ" : { "Ref": "MultiAZDatabase" },  
      "MasterUsername" : { "Ref" : "<master_username>" },  
      "DBInstanceClass" : { "Ref" : "DBCClass" },  
      "AllocatedStorage" : { "Ref" : "DBAllocatedStorage" },  
      "MasterUserPassword": { "Ref" : "<master_password>" },  
      "VPCSecurityGroups" : [ { "Fn::GetAtt": [ "VPCSecurityGroup", "GroupId" ] } ]  
    }  
  }  
}
```

Working with DB Security Groups (EC2-Classic Platform)

By default, network access is turned off to a DB instance. You can specify rules in a *security group* that allows access from an IP address range, port, or EC2 security group. Once ingress rules are configured, the same rules apply to all DB instances that are associated with that security group. You can specify up to 20 rules in a security group.

Amazon RDS supports two different kinds of security groups. The one you use depends on which Amazon RDS platform you are on:

- **VPC security groups** – for the EC2-VPC platform.
- **DB security groups** – for the EC2-Classic platform.

You are most likely on the EC2-VPC platform (and must use VPC security groups) if any of the following are true:

- If you are a new Amazon RDS customer.
- If you have never created a DB instance before.
- If you are creating a DB instance in an AWS Region you have not used before.

Otherwise, if you are on the EC2-Classic platform, you use DB security groups to manage access to your Amazon RDS DB instances. For more information about the differences between DB security groups and VPC security groups, see [Controlling Access with Security Groups \(p. 486\)](#).

Note

To determine which platform you are on, see [Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform \(p. 506\)](#).

If you are on the EC2-VPC platform, you must use VPC security groups instead of DB security groups. For more information about using a VPC, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).

Topics

- [Creating a DB Security Group \(p. 491\)](#)
- [Listing Available DB Security Groups \(p. 493\)](#)
- [Viewing a DB Security Group \(p. 493\)](#)
- [Associating a DB Security Group with a DB Instance \(p. 494\)](#)
- [Authorizing Network Access to a DB Security Group from an IP Range \(p. 495\)](#)
- [Authorizing Network Access to a DB Instance from an Amazon EC2 Instance \(p. 496\)](#)
- [Revoking Network Access to a DB Instance from an IP Range \(p. 498\)](#)

Creating a DB Security Group

To create a DB security group, you need to provide a name and a description.

Console

To create a DB security group

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. From the navigation pane, choose **Security Groups**.

Note

If you are on the EC2-VPC platform, the **Security Groups** option does not appear in the navigation pane. In this case, you must use VPC security groups instead of DB security groups. For more information about using a VPC, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).

3. Choose **Create DB Security Group**.
4. Type the name and description of the new DB security group in the **Name** and **Description** text boxes. The security group name can't contain spaces and can't start with a number.
5. Choose **Yes, Create**.

The DB security group is created.

A newly created DB security group doesn't provide access to a DB instance by default. You must specify a range of IP addresses or an Amazon EC2 security group that can have access to the DB instance. To specify IP addresses or an Amazon EC2 security group for a DB security group, see [Authorizing Network Access to a DB Security Group from an IP Range \(p. 495\)](#).

AWS CLI

To create a DB security group, use the AWS CLI command `create-db-security-group`.

Example

For Linux, OS X, or Unix:

```
aws rds create-db-security-group \
--db-security-group-name mydbsecuritygroup \
--db-security-group-description "My new security group"
```

For Windows:

```
aws rds create-db-security-group ^
--db-security-group-name mydbsecuritygroup ^
--db-security-group-description "My new security group"
```

A newly created DB security group doesn't provide access to a DB instance by default. You must specify a range of IP addresses or an Amazon EC2 security group that can have access to the DB instance. To specify IP addresses or an Amazon EC2 security group for a DB security group, see [Authorizing Network Access to a DB Security Group from an IP Range \(p. 495\)](#).

API

To create a DB security group, call the Amazon RDS function `CreateDBSecurityGroup` with the following parameters:

- `DBSecurityGroupName` = `mydbsecuritygroup`
- `Description` = `"My new security group"`

Example

```
https://rds.amazonaws.com/
?Action=CreateDBSecurityGroup
&DBSecurityGroupName=mydbsecuritygroup
&Description=My%20new%20db%20security%20group
&Version=2012-01-15
&SignatureVersion=2
```

```
&SignatureMethod=HmacSHA256
&Timestamp=2012-01-20T22%3A06%3A23.624Z
&AWSAccessKeyId=<AWS Access Key ID>
&Signature=<Signature>
```

A newly created DB security group doesn't provide access to a DB instance by default. You must specify a range of IP addresses or an Amazon EC2 security group that can have access to the DB instance. To specify IP addresses or an Amazon EC2 security group for a DB security group, see [Authorizing Network Access to a DB Security Group from an IP Range \(p. 495\)](#).

Listing Available DB Security Groups

You can list which DB security groups have been created for your AWS account.

Console

To list all available DB security groups for an AWS account

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. From the navigation pane, choose **Security Groups**.

The available DB security groups appear in the **DB Security Groups** list.

Note

If you are on the EC2-VPC platform, the **Security Groups** option does not appear in the navigation pane. In this case, you must use VPC security groups instead of DB security groups. For more information about using a VPC, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).

AWS CLI

To list all available DB security groups for an AWS account, Use the AWS CLI command `describe-db-security-groups` with no parameters.

Example

```
aws rds describe-db-security-groups
```

API

To list all available DB security groups for an AWS account, call `DescribeDBSecurityGroups` with no parameters.

Example

```
https://rds.amazonaws.com/
?Action=DescribeDBSecurityGroups
&MaxRecords=100
&Version=2009-10-16
&SignatureVersion=2
&SignatureMethod=HmacSHA256
&AWSAccessKeyId=<AWS Access Key ID>
&Signature=<Signature>
```

Viewing a DB Security Group

You can view detailed information about your DB security group to see what IP ranges have been authorized.

Console

To view properties of a specific DB security group

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. From the navigation pane, choose **Security Groups**.

Note

If you are on the EC2-VPC platform, the **Security Groups** option does not appear in the navigation pane. In this case, you must use VPC security groups instead of DB security groups. For more information about using a VPC, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).

3. Select the details icon for the DB security group you want to view. The detailed information for the DB security group is displayed.

AWS CLI

To view the properties of a specific DB security group use the AWS CLI `describe-db-security-groups`. Specify the DB security group you want to view.

Example

For Linux, OS X, or Unix:

```
aws rds describe-db-security-groups \
--db-security-group-name mydbsecuritygroup
```

For Windows:

```
aws rds describe-db-security-groups ^
--db-security-group-name mydbsecuritygroup
```

API

To view properties of a specific DB security group, call `DescribeDBSecurityGroups` with the following parameters:

- `DBSecurityGroupName=mydbsecuritygroup`

Example

```
https://rds.amazonaws.com/
?Action=DescribeDBSecurityGroups
&DBSecurityGroupName=mydbsecuritygroup
&Version=2009-10-16
&SignatureVersion=2
&SignatureMethod=HmacSHA256
&Timestamp=2009-10-16T22%3A23%3A07.107Z
&AWSAccessKeyId=<AWS Access Key ID>
&Signature=<Signature>
```

Associating a DB Security Group with a DB Instance

You can associate a DB security group with a DB instance using the RDS console's **Modify** option, the `ModifyDBInstance` Amazon RDS API, or the AWS CLI `modify-db-instance` command.

For information about modifying a DB instance, see [Modifying an Amazon RDS DB Instance \(p. 150\)](#).

Authorizing Network Access to a DB Security Group from an IP Range

By default, network access is turned off to a DB instance. If you want to access a DB instance that is not in a VPC, you must set access rules for a DB security group to allow access from specific EC2 security groups or CIDR IP ranges. You then must associate that DB instance with that DB security group. This process is called *ingress*. Once ingress is configured for a DB security group, the same ingress rules apply to all DB instances associated with that DB security group.

Warning

Talk with your network administrator if you are intending to access a DB instance behind a firewall to determine the IP addresses you should use.

In following example, you configure a DB security group with an ingress rule for a CIDR IP range.

Console

To configure a DB security group with an ingress rule for a CIDR IP range

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. From the navigation pane, choose **Security Groups**.

Note

If you are on the EC2-VPC platform, the **Security Groups** option does not appear in the navigation pane. In this case, you must use VPC security groups instead of DB security groups. For more information about using a VPC, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).

3. Select the details icon for the DB security group you want to authorize.
4. In the details page for your security group, select *CIDR/IP* from the **Connection Type** drop-down list, type the CIDR range for the ingress rule you want to add to this DB security group into the **CIDR** text box, and choose **Authorize**.

Tip

The AWS Management Console displays a CIDR IP based on your connection below the CIDR text field. If you are not accessing the DB instance from behind a firewall, you can use this CIDR IP.

5. The status of the ingress rule is **authorizing** until the new ingress rule has been applied to all DB instances that are associated with the DB security group that you modified. After the ingress rule has been successfully applied, the status changes to **authorized**.

AWS CLI

To configure a DB security group with an ingress rule for a CIDR IP range, use the AWS CLI command [authorize-db-security-group-ingress](#).

Example

For Linux, OS X, or Unix:

```
aws rds authorize-db-security-group-ingress \
--db-security-group-name mydbsecuritygroup \
--cidr 192.168.1.10/27
```

For Windows:

```
aws rds authorize-db-security-group-ingress ^
--db-security-group-name mydbsecuritygroup ^
--cidrIp 192.168.1.10/27
```

The command should produce output similar to the following.

```
SECURITYGROUP mydbsecuritygroup My new DBSecurityGroup
IP-RANGE 192.168.1.10/27 authorizing
```

API

To configure a DB security group with an ingress rule for a CIDR IP range, call the Amazon RDS API [AuthorizeDBSecurityGroupIngress](#) with the following parameters:

- DBSecurityGroupName = *mydbsecuritygroup*
- CIDRIP = *192.168.1.10/27*

Example

```
https://rds.amazonaws.com/
?Action=AuthorizeDBSecurityGroupIngress
&CIDRIP=192.168.1.10%2F27
&DBSecurityGroupName=mydbsecuritygroup
&Version=2009-10-16
&Action=AuthorizeDBSecurityGroupIngress
&SignatureVersion=2
&SignatureMethod=HmacSHA256
&Timestamp=2009-10-22T17%3A10%3A50.274Z
&AWSAccessKeyId=<AWS Access Key ID>
&Signature=<Signature>
```

Authorizing Network Access to a DB Instance from an Amazon EC2 Instance

If you want to access your DB instance from an Amazon EC2 instance, you must first determine if your EC2 instance and DB instance are in a VPC. If you are using a default VPC, you can assign the same EC2 or VPC security group that you used for your EC2 instance when you create or modify the DB instance that the EC2 instance accesses.

If your DB instance and EC2 instance are not in a VPC, you must configure the DB instance's security group with an ingress rule that allows traffic from the Amazon EC2 instance. You do this by adding the Amazon EC2 security group for the EC2 instance to the DB security group for the DB instance. In this example, you add an ingress rule to a DB security group for an Amazon EC2 security group.

Important

- Adding an ingress rule to a DB security group for an Amazon EC2 security group only grants access to your DB instances from Amazon EC2 instances associated with that Amazon EC2 security group.
- You can't authorize an Amazon EC2 security group that is in a different AWS Region than your DB instance. You can authorize an IP range, or specify an Amazon EC2 security group in the same AWS Region that refers to IP address in another AWS Region. If you specify an IP range, we recommend that you use the private IP address of your Amazon EC2 instance, which provides a more direct network route from your Amazon EC2 instance to your Amazon RDS DB instance, and doesn't incur network charges for data sent outside of the Amazon network.

Console

To add an EC2 security group to a DB security group

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. From the navigation pane, choose **Security Groups**.

Note

If you are on the EC2-VPC platform, the **Security Groups** option does not appear in the navigation pane. In this case, you must use VPC security groups instead of DB security groups. For more information about using a VPC, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).

3. Select the details icon for the DB security group you want to grant access.
4. In the details page for your security group, choose **EC2 Security Group** for **Connection Type**, and then select the Amazon EC2 security group you want to use. Then choose **Authorize**.
5. The status of the ingress rule is **authorizing** until the new ingress rule has been applied to all DB instances that are associated with the DB security group that you modified. After the ingress rule has been successfully applied, the status changes to **authorized**.

AWS CLI

To grant access to an Amazon EC2 security group, use the AWS CLI command [authorize-db-security-group-ingress](#).

Example

For Linux, OS X, or Unix:

```
aws rds authorize-db-security-group-ingress \
--db-security-group-name default \
--ec2-security-group-name myec2group \
--ec2-security-group-owner-id 987654321021
```

For Windows:

```
aws rds authorize-db-security-group-ingress ^
--db-security-group-name default ^
--ec2-security-group-name myec2group ^
--ec2-security-group-owner-id 987654321021
```

The command should produce output similar to the following:

```
SECGROUP  Name      Description
SECGROUP  default   default
          EC2-SECGROUP  myec2group  987654321021  authorizing
```

API

To authorize network access to an Amazon EC2 security group, call the Amazon RDS API function, https://docs.aws.amazon.com/AmazonRDS/latest/APIReference/API_AuthorizeDBSecurityGroupIngress.html with the following parameters:

- **EC2SecurityGroupName** = *myec2group*
- **EC2SecurityGroupOwnerId** = *987654321021*

Example

```
https://rds.amazonaws.com/  
    ?Action=AuthorizeDBSecurityGroupIngress  
    &EC2SecurityGroupOwnerId=987654321021  
    &EC2SecurityGroupName=myec2group  
    &Version=2009-10-16  
    &SignatureVersion=2  
    &SignatureMethod=HmacSHA256  
    &Timestamp=2009-10-22T17%3A10%3A50.274Z  
    &AWSAccessKeyId=<AWS Access Key ID>  
    &Signature=<Signature>
```

Revoking Network Access to a DB Instance from an IP Range

You can easily revoke network access from a CIDR IP range to DB instances belonging to a DB security group by revoking the associated CIDR IP ingress rule.

In this example, you revoke an ingress rule for a CIDR IP range on a DB security group.

Console

To revoke an ingress rule for a CIDR IP range on a DB Security Group

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. From the navigation pane, choose **Security Groups**.

Note

If you are on the EC2-VPC platform, the **Security Groups** option does not appear in the navigation pane. In this case, you must use VPC security groups instead of DB security groups. For more information about using a VPC, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).

3. Select the details icon for the DB security group that has the ingress rule you want to revoke.
4. In the details page for your security group, choose **Remove** next to the ingress rule you want to revoke.
5. The status of the ingress rule is **revoking** until the ingress rule has been removed from all DB instances that are associated with the DB security group that you modified. After the ingress rule has been successfully removed, the ingress rule is removed from the DB security group.

AWS CLI

To revoke an ingress rule for a CIDR IP range on a DB security group, use the AWS CLI command `revoke-db-security-group-ingress`.

Example

For Linux, OS X, or Unix:

```
aws rds revoke-db-security-group-ingress \  
    --db-security-group-name mydbsecuritygroup \  
    --cidrip 192.168.1.1/27
```

For Windows:

```
aws rds revoke-db-security-group-ingress ^  
    --db-security-group-name mydbsecuritygroup ^  
    --cidrip 192.168.1.1/27
```

The command should produce output similar to the following.

```
SECGROUP mydbsecuritygroup My new DBSecurityGroup
IP-RANGE 192.168.1.1/27 revoking
```

API

To revoke an ingress rule for a CIDR IP range on a DB security group, call the Amazon RDS API operation https://docs.aws.amazon.com/AmazonRDS/latest/APIReference/API_RevokeDBSecurityGroupIngress.html with the following parameters:

- DBSecurityGroupName = *mydbsecuritygroup*
- CIDRIP = *192.168.1.10/27*

Example

```
https://rds.amazonaws.com/
?Action=RevokeDBSecurityGroupIngress
&DBSecurityGroupName=mydbsecuritygroup
&CIDRIP=192.168.1.10%2F27
&Version=2009-10-16
&SignatureVersion=2&SignatureMethod=HmacSHA256
&Timestamp=2009-10-22T22%3A32%3A12.515Z
&AWSAccessKeyId=<AWS Access Key ID>
&Signature=<Signature>
```

Master User Account Privileges

When you create a new DB instance, the default master user that you use gets certain privileges for that DB instance. The following table shows the privileges and database roles the master user gets for each of the database engines.

Important

We strongly recommend that you do not use the master user directly in your applications. Instead, adhere to the best practice of using a database user created with the minimal privileges required for your application.

Note

If you accidentally delete the permissions for the master user, you can restore them by modifying the DB instance and setting a new master user password. For more information about modifying a DB instance, see [Modifying an Amazon RDS DB Instance \(p. 150\)](#).

Database Engine	System Privilege	Database Role
MySQL and MariaDB	SELECT, INSERT, UPDATE, DELETE, CREATE, DROP, RELOAD, PROCESS, REFERENCES, INDEX, ALTER, SHOW DATABASES, CREATE TEMPORARY TABLES, LOCK TABLES, EXECUTE, REPLICATION CLIENT, CREATE VIEW, SHOW VIEW, CREATE ROUTINE, ALTER ROUTINE, CREATE USER, EVENT, TRIGGER ON *.* WITH GRANT OPTION, REPLICATION SLAVE (only for Amazon RDS MySQL versions 5.6, 5.7 and 8.0, Amazon RDS MariaDB)	—

Database Engine	System Privilege	Database Role
PostgreSQL	CREATE ROLE, CREATE DB, PASSWORD VALID UNTIL INFINITY, CREATE EXTENSION, ALTER EXTENSION, DROP EXTENSION, CREATE TABLESPACE, ALTER < OBJECT> OWNER, CHECKPOINT, PG_CANCEL_BACKEND(), PG_TERMINATE_BACKEND(), SELECT PG_STAT_REPLICATION, EXECUTE PG_STAT_STATEMENTS_RESET(), OWN POSTGRES_FDW_HANDLER(), OWN POSTGRES_FDW_VALIDATOR(), OWN POSTGRES_FDW, EXECUTE PG_BUFFERCACHE_PAGES(), SELECT PG_BUFFERCACHE	RDS_SUPERUSER
Oracle	ALTER DATABASE LINK, ALTER PUBLIC DATABASE LINK, DROP ANY DIRECTORY, EXEMPT ACCESS POLICY, EXEMPT IDENTITY POLICY, GRANT ANY OBJECT PRIVILEGE, RESTRICTED SESSION, EXEMPT REDACTION POLICY	AQ_ADMINISTRATOR_ROLE, AQ_USER_ROLE, CONNECT, CTXAPP, DBA, EXECUTE_CATALOG_ROLE, RECOVERY_CATALOG_OWNER, RESOURCE, SELECT_CATALOG_ROLE
Microsoft SQL Server	ADMINISTER BULK OPERATIONS, ALTER ANY CONNECTION, ALTER ANY LINKED SERVER, ALTER ANY LOGIN, ALTER SERVER STATE, ALTER TRACE, CONNECT SQL, CREATE ANY DATABASE, VIEW ANY DATABASE, VIEW ANY DEFINITION, VIEW SERVER STATE, ALTER ANY SERVER ROLE, ALTER ANY USER, ALTER ON ROLE SQLAgentOperatorRole	DB_OWNER (Database Level Role) PROCESSADMIN (Server Level Role) SETUPADMIN(Server Level Role) SQLAgentUserRole(Server Level Role)

Using Service-Linked Roles for Amazon RDS

Amazon RDS uses AWS Identity and Access Management (IAM) [service-linked roles](#). A service-linked role is a unique type of IAM role that is linked directly to Amazon RDS. Service-linked roles are predefined by Amazon RDS and include all the permissions that the service requires to call other AWS services on your behalf.

A service-linked role makes using Amazon RDS easier because you don't have to manually add the necessary permissions. Amazon RDS defines the permissions of its service-linked roles, and unless defined otherwise, only Amazon RDS can assume its roles. The defined permissions include the trust policy and the permissions policy, and that permissions policy cannot be attached to any other IAM entity.

You can delete the roles only after first deleting their related resources. This protects your Amazon RDS resources because you can't inadvertently remove permission to access the resources.

For information about other services that support service-linked roles, see [AWS Services That Work with IAM](#) and look for the services that have **Yes** in the **Service-Linked Role** column. Choose a **Yes** with a link to view the service-linked role documentation for that service.

Service-Linked Role Permissions for Amazon RDS

Amazon RDS uses the service-linked role named **AWSServiceRoleForRDS** – to allow Amazon RDS to call AWS services on behalf of your DB instances.

The **AWSServiceRoleForRDS** service-linked role trusts the following services to assume the role:

- `rds.amazonaws.com`

The role permissions policy allows Amazon RDS to complete the following actions on the specified resources:

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "ec2:AuthorizeSecurityGroupIngress",  
                "ec2>CreateNetworkInterface",  
                "ec2>CreateSecurityGroup",  
                "ec2>DeleteNetworkInterface",  
                "ec2>DeleteSecurityGroup",  
                "ec2:DescribeAvailabilityZones",  
                "ec2:DescribeInternetGateways",  
                "ec2:DescribeSecurityGroups",  
                "ec2:DescribeSubnets",  
                "ec2:DescribeVpcAttribute",  
                "ec2:DescribeVpcs",  
                "ec2:ModifyNetworkInterfaceAttribute",  
                "ec2:ModifyVpcEndpoint",  
                "ec2:RevokeSecurityGroupIngress",  
                "ec2>CreateVpcEndpoint",  
                "ec2:DescribeVpcEndpoints",  
                "ec2>DeleteVpcEndpoints",  
                "ec2:AssignPrivateIpAddresses",  
                "ec2:UnassignPrivateIpAddresses"  
            ],  
            "Resource": "*"  
        }  
    ]  
}
```

```

},
{
    "Effect": "Allow",
    "Action": [
        "sns:Publish"
    ],
    "Resource": "*"
},
{
    "Effect": "Allow",
    "Action": [
        "logs>CreateLogGroup"
    ],
    "Resource": [
        "arn:aws:logs:***:log-group:/aws/rds/*",
        "arn:aws:logs:***:log-group:/aws/docdb/*",
        "arn:aws:logs:***:log-group:/aws/neptune/*"
    ]
},
{
    "Effect": "Allow",
    "Action": [
        "logs>CreateLogStream",
        "logs:PutLogEvents",
        "logs:DescribeLogStreams"
    ],
    "Resource": [
        "arn:aws:logs:***:log-group:/aws/rds/*:log-stream:*",
        "arn:aws:logs:***:log-group:/aws/docdb/*:log-stream:*",
        "arn:aws:logs:***:log-group:/aws/neptune/*:log-stream:*
    ]
},
{
    "Effect": "Allow",
    "Action": [
        "kinesis>CreateStream",
        "kinesis:PutRecord",
        "kinesis:PutRecords",
        "kinesis:DescribeStream",
        "kinesis:SplitShard",
        "kinesis:MergeShards",
        "kinesis>DeleteStream",
        "kinesis:UpdateShardCount"
    ],
    "Resource": [
        "arn:aws:kinesis:***:stream/aws-rds-das-*"
    ]
}
]
}

```

Note

You must configure permissions to allow an IAM entity (such as a user, group, or role) to create, edit, or delete a service-linked role. If you encounter the following error message:

Unable to create the resource. Verify that you have permission to create service linked role. Otherwise wait and try again later.

Make sure you have the following permissions enabled:

```
{
    "Action": "iam>CreateServiceLinkedRole",
    "Effect": "Allow",
    "Resource": "arn:aws:iam:***:role/aws-service-role/rds.amazonaws.com/
AWSServiceRoleForRDS",
```

```
"Condition": {  
    "StringLike": {  
        "iam:AWSServiceName": "rds.amazonaws.com"  
    }  
}
```

For more information, see [Service-Linked Role Permissions](#) in the *IAM User Guide*.

Creating a Service-Linked Role for Amazon RDS

You don't need to manually create a service-linked role. When you create a DB instance, Amazon RDS creates the service-linked role for you.

Important

If you were using the Amazon RDS service before December 1, 2017, when it began supporting service-linked roles, then Amazon RDS created the AWSRoleForRDS role in your account. To learn more, see [A New Role Appeared in My IAM Account](#).

If you delete this service-linked role, and then need to create it again, you can use the same process to recreate the role in your account. When you create a DB instance, Amazon RDS creates the service-linked role for you again.

Editing a Service-Linked Role for Amazon RDS

Amazon RDS does not allow you to edit the AWSRoleForRDS service-linked role. After you create a service-linked role, you cannot change the name of the role because various entities might reference the role. However, you can edit the description of the role using IAM. For more information, see [Editing a Service-Linked Role](#) in the *IAM User Guide*.

Deleting a Service-Linked Role for Amazon RDS

If you no longer need to use a feature or service that requires a service-linked role, we recommend that you delete that role. That way you don't have an unused entity that is not actively monitored or maintained. However, you must delete all of your DB instances before you can delete the service-linked role.

Cleaning Up a Service-Linked Role

Before you can use IAM to delete a service-linked role, you must first confirm that the role has no active sessions and remove any resources used by the role.

To check whether the service-linked role has an active session in the IAM console

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane of the IAM console, choose **Roles**. Then choose the name (not the check box) of the AWSRoleForRDS role.
3. On the **Summary** page for the chosen role, choose the **Access Advisor** tab.
4. On the **Access Advisor** tab, review recent activity for the service-linked role.

Note

If you are unsure whether Amazon RDS is using the AWSRoleForRDS role, you can try to delete the role. If the service is using the role, then the deletion fails and you can view the AWS Regions where the role is being used. If the role is being used, then you must wait

for the session to end before you can delete the role. You cannot revoke the session for a service-linked role.

If you want to remove the AWSServiceRoleForRDS role, you must first delete *all* of your DB instances .

Deleting All of Your Instances

Use one of these procedures to delete each of your instances.

To delete an instance (console)

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**.
3. Choose the instance that you want to delete.
4. For **Actions**, choose **Delete**.
5. If you are prompted for **Create final Snapshot?**, choose **Yes** or **No**.
6. If you chose **Yes** in the previous step, for **Final snapshot name** enter the name of your final snapshot.
7. Choose **Delete**.

To delete an instance (CLI)

See [delete-db-instance](#) in the *AWS CLI Command Reference*.

To delete an instance (API)

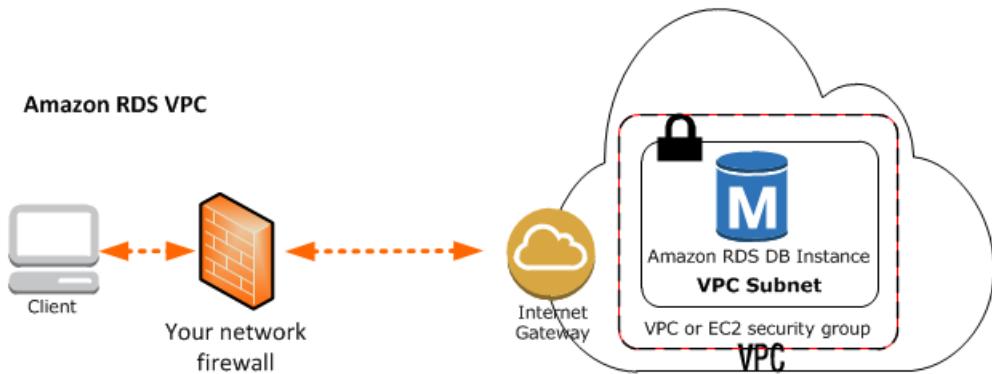
See [DeleteDBInstance](#) in the *Amazon RDS API Reference*.

You can use the IAM console, the IAM CLI, or the IAM API to delete the AWSServiceRoleForRDS service-linked role. For more information, see [Deleting a Service-Linked Role](#) in the *IAM User Guide*.

Amazon Virtual Private Cloud VPCs and Amazon RDS

There are two Amazon Elastic Compute Cloud (EC2) platforms that host Amazon RDS DB instances, *EC2-VPC* and *EC2-Classic*. Amazon Virtual Private Cloud (Amazon VPC) enables you to launch AWS resources, such as Amazon RDS DB instances, into a virtual private cloud (VPC).

When you use an Amazon VPC, you have control over your virtual networking environment: you can choose your own IP address range, create subnets, and configure routing and access control lists. The basic functionality of Amazon RDS is the same whether your DB instance is running in an Amazon VPC or not: Amazon RDS manages backups, software patching, automatic failure detection, and recovery. There is no additional cost to run your DB instance in an Amazon VPC.



Accounts that support only the *EC2-VPC* platform have a default VPC. All new DB instances are created in the default VPC unless you specify otherwise. If you are a new Amazon RDS customer, if you have never created a DB instance before, or if you are creating a DB instance in an AWS Region you have not used before, you are most likely on the *EC2-VPC* platform and have a default VPC.

Some legacy DB instances on the *EC2-Classic* platform are not in a VPC. The legacy *EC2-Classic* platform does not have a default VPC, but as is true for either platform, you can create your own VPC and specify that a DB instance be located in that VPC.

Topics

- [Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform \(p. 506\)](#)
- [Scenarios for Accessing a DB Instance in a VPC \(p. 507\)](#)
- [Working with a DB Instance in a VPC \(p. 514\)](#)
- [Updating the VPC for a DB Instance \(p. 519\)](#)
- [Tutorial: Create an Amazon VPC for Use with a DB Instance \(p. 521\)](#)

This documentation only discusses VPC functionality relevant to Amazon RDS DB instances. For more information about Amazon VPC, see [Amazon VPC Getting Started Guide](#) and [Amazon VPC User Guide](#). For information about using a network address translation (NAT) gateway, see [NAT Gateways](#) in the [Amazon Virtual Private Cloud User Guide](#).

Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform

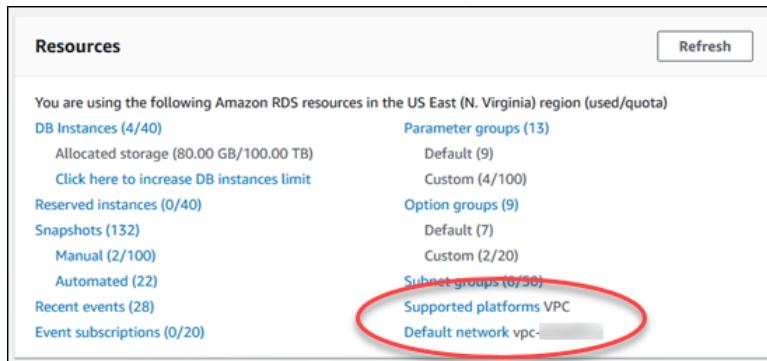
Your AWS account and the AWS Region you choose determines which of the two RDS platforms your DB instance is created on: *EC2-Classic* or *EC2-VPC*. The type of platform determines if you have a default VPC, and which type of security group you use to provide access to your DB instance. The legacy *EC2-Classic* platform is the original platform used by Amazon RDS; if you are on this platform and want to use a VPC, you must create the VPC using the Amazon VPC console or Amazon VPC API. Accounts that only support the *EC2-VPC* platform have a default VPC where all DB instances are created, and you must use either an EC2 or VPC security group to provide access to the DB instance.

Note

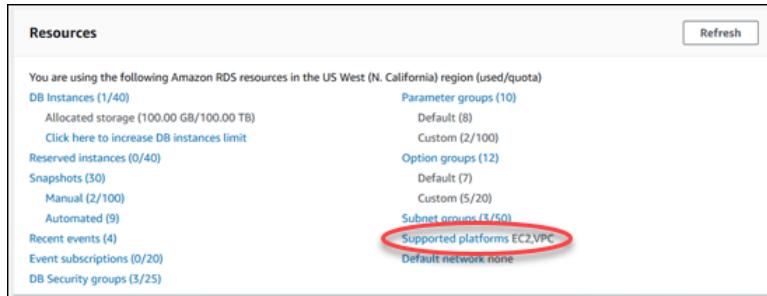
If you are a new Amazon RDS customer, if you have never created a DB instance before, or if you are creating a DB instance in an AWS Region you have not used before, in almost all cases you are on the *EC2-VPC* platform and have a default VPC.

You can tell which platform your AWS account in a given AWS Region is using by looking at the dashboard on the RDS console or EC2 console. If you are a new Amazon RDS customer, if you have never created a DB instance before, or if you are creating a DB instance in an AWS Region you have not used before, you might be redirected to the first-run console page and not see the home page following.

If **Supported Platforms** indicates VPC, as shown following, your AWS account in the current AWS Region uses the *EC2-VPC* platform, and uses a default VPC. The name of the default VPC is shown below the supported platform. To provide access to a DB instance created on the *EC2-VPC* platform, you must create a VPC security group. For information about creating a VPC security group, see [Tutorial: Create an Amazon VPC for Use with a DB Instance \(p. 521\)](#).



If **Supported Platforms** indicates EC2 , VPC, as shown following, your AWS account in the current AWS Region uses the *EC2-Classic* platform, and you do not have a default VPC. To provide access to a DB instance created on the *EC2-Classic* platform, you must create a DB security group. For information about creating a DB security group, see [Creating a DB Security Group \(p. 491\)](#).



Note

- You can create a VPC on the *EC2-Classic* platform, but one is not created for you by default as it is on accounts that support the *EC2-VPC* platform.
- If you are interested in moving an existing DB instance into a VPC, you can use the AWS Management Console to do it easily. For more information, see [Moving a DB Instance Not in a VPC into a VPC \(p. 520\)](#).

Scenarios for Accessing a DB Instance in a VPC

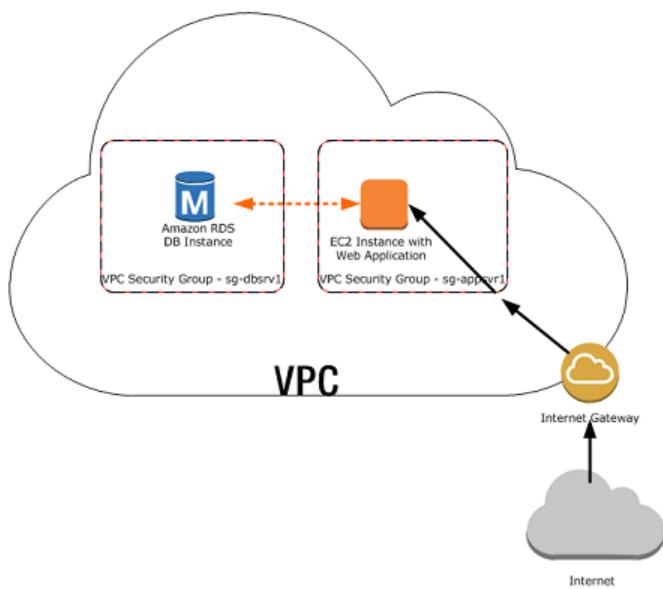
Amazon RDS supports the following scenarios for accessing a DB instance:

DB Instance	Accessed By
In a VPC	An EC2 Instance in the Same VPC (p. 507)
	An EC2 Instance in a Different VPC (p. 509)
	An EC2 Instance Not in a VPC (p. 509)
	A Client Application Through the Internet (p. 510)
Not in a VPC	An EC2 Instance in a VPC (p. 511)
	An EC2 Instance Not in a VPC (p. 512)
	A Client Application Through the Internet (p. 513)

A DB Instance in a VPC Accessed by an EC2 Instance in the Same VPC

A common use of a DB instance in a VPC is to share data with an application server that is running in an EC2 instance in the same VPC. This is the user scenario created if you use AWS Elastic Beanstalk to create an EC2 instance and a DB instance in the same VPC.

The following diagram shows this scenario.



The simplest way to manage access between EC2 instances and DB instances in the same VPC is to do the following:

- Create a VPC security group for your DB instances to be in. This security group can be used to restrict access to the DB instances. For example, you can create a custom rule for this security group that allows TCP access using the port you assigned to the DB instance when you created it and an IP address you use to access the DB instance for development or other purposes.
- Create a VPC security group for your EC2 instances (web servers and clients) to be in. This security group can, if needed, allow access to the EC2 instance from the Internet via the VPC's routing table. For example, you can set rules on this security group to allow TCP access to the EC2 instance over port 22.
- Create custom rules in the security group for your DB instances that allow connections from the security group you created for your EC2 instances. This would allow any member of the security group to access the DB instances.

For a tutorial that shows you how to create a VPC with both public and private subnets for this scenario, see [Tutorial: Create an Amazon VPC for Use with a DB Instance \(p. 521\)](#).

To create a rule in a VPC security group that allows connections from another security group, do the following:

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc>.
2. In the navigation pane, choose **Security Groups**.
3. Choose or create a security group for which you want to allow access to members of another security group. In the scenario preceding, this is the security group that you use for your DB instances. Choose the **Inbound Rules** tab, and then choose **Edit rule**.
4. On the **Edit inbound rules** page, choose **Add Rule**.
5. From **Type**, choose one of the **All ICMP** options. In the **Source** box, start typing the ID of the security group; this provides you with a list of security groups. Choose the security group with members that you want to have access to the resources protected by this security group. In the scenario preceding, this is the security group that you use for your EC2 instance.
6. Repeat the steps for the TCP protocol by creating a rule with **All TCP** as the **Type** and your security group in the **Source** box. If you intend to use the UDP protocol, create a rule with **All UDP** as the **Type** and your security group in the **Source** box.

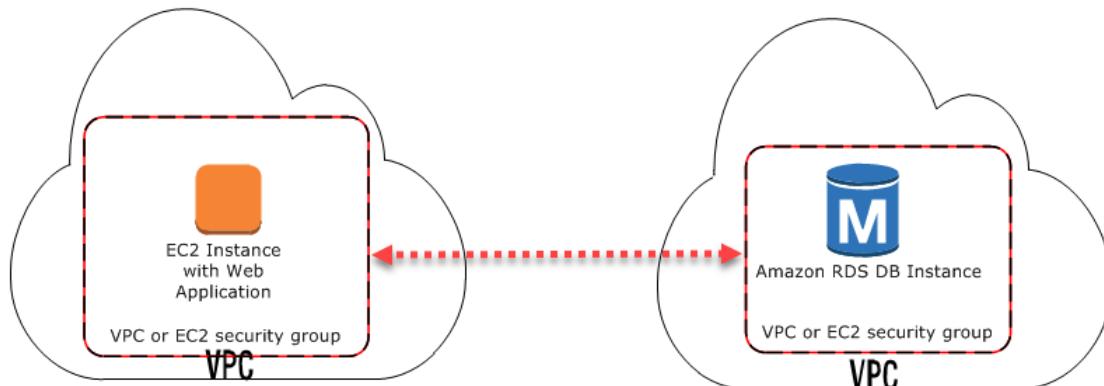
7. Create a custom TCP rule that permits access via the port you used when you created your DB instance, such as port 3306 for MySQL. Enter your security group or an IP address to use in the **Source** box.
8. Choose **Save** when you are done.



A DB Instance in a VPC Accessed by an EC2 Instance in a Different VPC

When your DB instance is in a different VPC from the EC2 instance you are using to access it, you can use VPC peering to access the DB instance.

The following diagram shows this scenario.

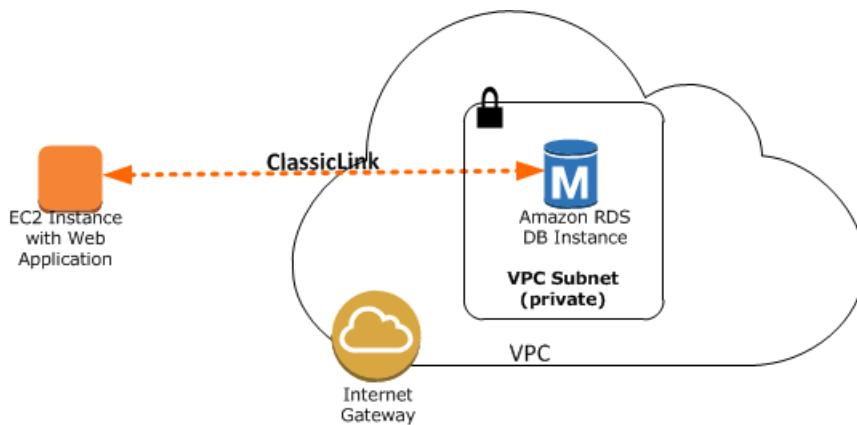


A VPC peering connection is a networking connection between two VPCs that enables you to route traffic between them using private IP addresses. Instances in either VPC can communicate with each other as if they are within the same network. You can create a VPC peering connection between your own VPCs, with a VPC in another AWS account, or with a VPC in a different AWS Region. To learn more about VPC peering, see [VPC Peering](#) in the *Amazon Virtual Private Cloud User Guide*.

A DB Instance in a VPC Accessed by an EC2 Instance Not in a VPC

You can communicate between an Amazon RDS DB instance that is in a VPC and an EC2 instance that is not in an Amazon VPC by using *ClassicLink*. When you use Classic Link, an application on the EC2 instance can connect to the DB instance by using the endpoint for the DB instance. ClassicLink is available at no charge.

The following diagram shows this scenario.



Using ClassicLink, you can connect an EC2 instance to a logically isolated database where you define the IP address range and control the access control lists (ACLs) to manage network traffic. You don't have to use public IP addresses or tunneling to communicate with the DB instance in the VPC. This arrangement provides you with higher throughput and lower latency connectivity for inter-instance communications.

To enable ClassicLink between a DB instance in a VPC and an EC2 instance not in a VPC

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc>.
2. In the navigation pane, choose **Your VPCs**.
3. Choose the VPC used by the DB instance.
4. In **Actions**, choose **Enable ClassicLink**. In the confirmation dialog box, choose **Yes, Enable**.
5. On the EC2 console, choose the EC2 instance you want to connect to the DB instance in the VPC.
6. In **Actions**, choose **ClassicLink**, and then choose **Link to VPC**.
7. On the **Link to VPC** page, choose the security group you want to use, and then choose **Link to VPC**.

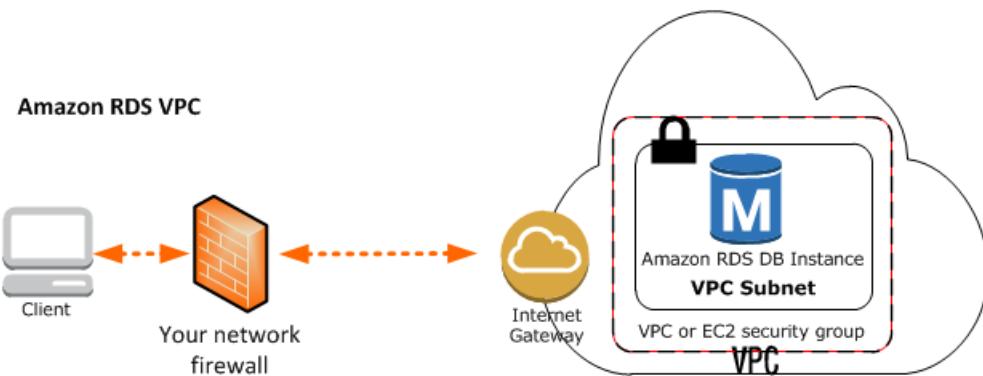
Note

The ClassicLink features are only visible in the consoles for accounts and regions that support EC2-Classic. For more information, see [ClassicLink](#) in the *Amazon EC2 User Guide for Linux Instances*.

A DB Instance in a VPC Accessed by a Client Application Through the Internet

To access a DB instance in a VPC from a client application through the internet, you configure a VPC with a single public subnet, and an Internet gateway to enable communication over the Internet.

The following diagram shows this scenario.



We recommend the following configuration:

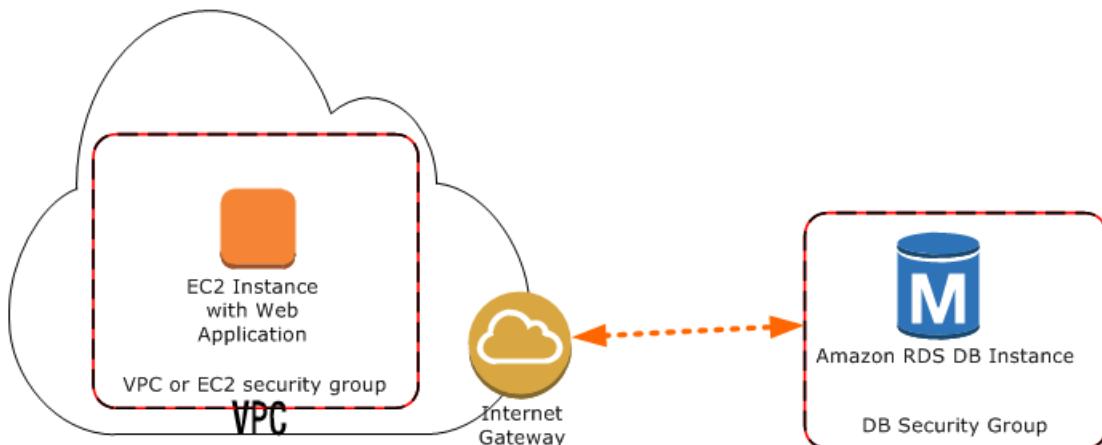
- A VPC of size /16 (for example CIDR: 10.0.0.0/16). This size provides 65,536 private IP addresses.
- A subnet of size /24 (for example CIDR: 10.0.0.0/24). This size provides 256 private IP addresses.
- An Amazon RDS DB instance that is associated with the VPC and the subnet. Amazon RDS assigns an IP address within the subnet to your DB instance.
- An Internet gateway which connects the VPC to the Internet and to other AWS products.
- A security group associated with the DB instance. The security group's inbound rules allow your client application to access to your DB instance.

For information about creating a DB instance in a VPC, see [Creating a DB Instance in a VPC \(p. 516\)](#).

A DB Instance Not in a VPC Accessed by an EC2 Instance in a VPC

In the case where you have an EC2 instance in a VPC and an RDS DB instance not in a VPC, you can connect them over the public Internet.

The following diagram shows this scenario.



Note

ClassicLink, as described in [A DB Instance in a VPC Accessed by an EC2 Instance Not in a VPC \(p. 509\)](#), is not available for this scenario.

To connect your DB instance and your EC2 instance over the public Internet, do the following:

- Ensure that the EC2 instance is in a public subnet in the VPC.
- Ensure that the RDS DB instance was marked as publicly accessible.
- A note about network ACLs here. A network ACL is like a firewall for your entire subnet. Therefore, all instances in that subnet are subject to network ACL rules. By default, network ACLs allow all traffic and you generally don't need to worry about them, unless you particularly want to add rules as an extra layer of security. A security group, on the other hand, is associated with individual instances, and you do need to worry about security group rules.
- Add the necessary ingress rules to the DB security group for the RDS DB instance.

An ingress rule specifies a network port and a CIDR/IP range. For example, you can add an ingress rule that allows port 3306 to connect to a MySQL RDS DB instance, and a CIDR/IP range of 203.0.113.25/32. For more information, see [Authorizing Network Access to a DB Security Group from an IP Range \(p. 495\)](#).

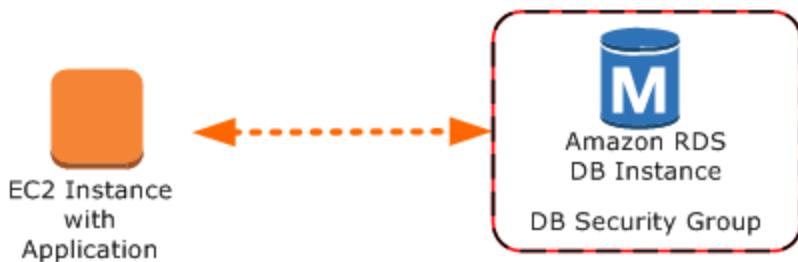
Note

If you are interested in moving an existing DB instance into a VPC, you can use the AWS Management Console to do it easily. For more information, see [Moving a DB Instance Not in a VPC into a VPC \(p. 520\)](#).

A DB Instance Not in a VPC Accessed by an EC2 Instance Not in a VPC

When neither your DB instance nor an application on an EC2 instance are in a VPC, you can access the DB instance by using its endpoint and port.

The following diagram shows this scenario.



You must create a security group for the DB instance that permits access from the port you specified when creating the DB instance. For example, you could use a connection string similar to this connection string used with *sqlplus* to access an Oracle DB instance:

```
PROMPT>sqlplus 'mydbusr@(DESCRIPTION=(ADDRESS=(PROTOCOL=TCP)(HOST=<endpoint>)(PORT=<port number>))(CONNECT_DATA=(SID=<database name>)))'
```

For more information, see the following documentation.

Database Engine	Relevant Documentation
MariaDB	Connecting to a DB Instance Running the MariaDB Database Engine (p. 555)
Microsoft SQL Server	Connecting to a DB Instance Running the Microsoft SQL Server Database Engine (p. 630)
MySQL	Connecting to a DB Instance Running the MySQL Database Engine (p. 764)
Oracle	Connecting to a DB Instance Running the Oracle Database Engine (p. 955)
PostgreSQL	Connecting to a DB Instance Running the PostgreSQL Database Engine (p. 1270)

Note

If you are interested in moving an existing DB instance into a VPC, you can use the AWS Management Console to do it easily. For more information, see [Moving a DB Instance Not in a VPC into a VPC \(p. 520\)](#).

A DB Instance Not in a VPC Accessed by a Client Application Through the Internet

New Amazon RDS customers can only create a DB instance in a VPC. However, you might need to connect to an existing Amazon RDS DB instance that is not in a VPC from a client application through the Internet.

The following diagram shows this scenario.



In this scenario, you must ensure that the DB security group for the RDS DB instance includes the necessary ingress rules for your client application to connect. An ingress rule specifies a network port and a CIDR/IP range. For example, you can add an ingress rule that allows port 3306 to connect to a MySQL RDS DB instance, and a CIDR/IP range of 203.0.113.25/32. For more information, see [Authorizing Network Access to a DB Security Group from an IP Range \(p. 495\)](#).

Warning

If you intend to access a DB instance behind a firewall, talk with your network administrator to determine the IP addresses you should use.

Note

If you are interested in moving an existing DB instance into a VPC, you can use the AWS Management Console to do it easily. For more information, see [Moving a DB Instance Not in a VPC into a VPC \(p. 520\)](#).

Working with a DB Instance in a VPC

Unless you are working with a legacy DB instance, your DB instance is in a virtual private cloud (VPC). A VPC is a virtual network that is logically isolated from other virtual networks in the AWS Cloud. Amazon VPC lets you launch AWS resources, such as an Amazon RDS DB instance or Amazon EC2 instance, into a VPC. The VPC can either be a default VPC that comes with your account or one that you create. All VPCs are associated with your AWS account.

Your default VPC has three subnets you can use to isolate resources inside the VPC. The default VPC also has an Internet Gateway that can be used to provide access to resources inside the VPC from outside the VPC.

For a list of scenarios involving Amazon RDS DB instances in a VPC and outside of a VPC, see [Scenarios for Accessing a DB Instance in a VPC \(p. 507\)](#).

For a tutorial that shows you how to create a VPC that you can use with a common Amazon RDS scenario, see [Tutorial: Create an Amazon VPC for Use with a DB Instance \(p. 521\)](#).

To learn how to work with DB instances inside a VPC, see the following:

Topics

- [Working with a DB Instance in a VPC \(p. 514\)](#)
- [Working with DB Subnet Groups \(p. 515\)](#)
- [Hiding a DB Instance in a VPC from the Internet \(p. 515\)](#)
- [Creating a DB Instance in a VPC \(p. 516\)](#)

Working with a DB Instance in a VPC

Here are some tips on working with a DB instance in a VPC:

- Your VPC must have at least two subnets. These subnets must be in two different Availability Zones in the AWS Region where you want to deploy your DB instance. A subnet is a segment of a VPC's IP address range that you can specify and that lets you group instances based on your security and operational needs.
- If you want your DB instance in the VPC to be publicly accessible, you must enable the VPC attributes *DNS hostnames* and *DNS resolution*.
- Your VPC must have a DB subnet group that you create (for more information, see the next section). You create a DB subnet group by specifying the subnets you created. Amazon RDS uses that DB subnet group and your preferred Availability Zone to choose a subnet and an IP address within that subnet to assign to your DB instance.
- Your VPC must have a VPC security group that allows access to the DB instance.
- The CIDR blocks in each of your subnets must be large enough to accommodate spare IP addresses for Amazon RDS to use during maintenance activities, including failover and compute scaling.
- A VPC can have an *instance tenancy* attribute of either *default* or *dedicated*. All default VPCs have the instance tenancy attribute set to default, and a default VPC can support any DB instance class.

If you choose to have your DB instance in a dedicated VPC where the instance tenancy attribute is set to dedicated, the DB instance class of your DB instance must be one of the approved Amazon EC2 dedicated instance types. For example, the m3.medium EC2 dedicated instance corresponds to the db.m3.medium DB instance class. For information about instance tenancy in a VPC, go to [Using EC2 Dedicated Instances in the Amazon Virtual Private Cloud User Guide](#).

For more information about the instance types that can be in a dedicated instance, see [Amazon EC2 Dedicated Instances](#) on the EC2 pricing page.

- When an option group is assigned to a DB instance, it is linked to the supported platform the DB instance is on, either VPC or EC2-Classic (non-VPC). Furthermore, if a DB instance is in a VPC, the option group associated with the DB instance is linked to that VPC. This linkage means that you cannot use the option group assigned to a DB instance if you attempt to restore the DB instance into a different VPC or onto a different platform.
- If you restore a DB instance into a different VPC or onto a different platform, you must either assign the default option group to the DB instance, assign an option group that is linked to that VPC or platform, or create a new option group and assign it to the DB instance. With persistent or permanent options, such as Oracle TDE, you must create a new option group that includes the persistent or permanent option when restoring a DB instance into a different VPC.

Working with DB Subnet Groups

Subnets are segments of a VPC's IP address range that you designate to group your resources based on security and operational needs. A DB subnet group is a collection of subnets (typically private) that you create in a VPC and that you then designate for your DB instances. A DB subnet group allows you to specify a particular VPC when creating DB instances using the CLI or API; if you use the console, you can just choose the VPC and subnets you want to use.

Each DB subnet group should have subnets in at least two Availability Zones in a given AWS Region. When creating a DB instance in a VPC, you must choose a DB subnet group. Amazon RDS uses that DB subnet group and your preferred Availability Zone to choose a subnet and an IP address within that subnet to associate with your DB instance. If the primary DB instance of a Multi-AZ deployment fails, Amazon RDS can promote the corresponding standby and subsequently create a new standby using an IP address of the subnet in one of the other Availability Zones.

When Amazon RDS creates a DB instance in a VPC, it assigns a network interface to your DB instance by using an IP address from your DB subnet group. However, we strongly recommend that you use the DNS name to connect to your DB instance because the underlying IP address changes during failover.

Note

For each DB instance that you run in a VPC, you should reserve at least one address in each subnet in the DB subnet group for use by Amazon RDS for recovery actions.

Hiding a DB Instance in a VPC from the Internet

One common Amazon RDS scenario is to have a VPC in which you have an EC2 instance with a public-facing web application and a DB instance with a database that is not publicly accessible. For example, you can create a VPC that has a public subnet and a private subnet. Amazon EC2 instances that function as web servers can be deployed in the public subnet, and the DB instances are deployed in the private subnet. In such a deployment, only the web servers have access to the DB instances. For an illustration of this scenario, see [A DB Instance in a VPC Accessed by an EC2 Instance in the Same VPC \(p. 507\)](#).

When you launch a DB instance inside a VPC, you can designate whether the DB instance you create has a DNS that resolves to a public IP address by using the *Public accessibility* parameter. This parameter lets you designate whether there is public access to the DB instance. Access to the DB instance is ultimately controlled by the security group it uses, and that public access is not permitted if the security group assigned to the DB instance does not permit it.

You can modify a DB instance to turn on or off public accessibility by modifying the *Public accessibility* parameter. This parameter is modified just like any other DB instance parameter. For more information, see the modifying section for your DB engine.

The following illustration shows the **Public accessibility** option in the **Network & Security** section.

The screenshot shows the 'Network & Security' configuration page. It includes fields for Subnet group (set to 'default'), Security group ('Select security groups'), and Certificate authority ('rds-ca-2015'). The 'Public accessibility info' section is highlighted with a red box. It contains two radio button options: 'Yes' (selected) and 'No'. The 'Yes' option describes public accessibility, while the 'No' option states that the DB instance will not have a public IP address. A warning message at the bottom of this section states: 'Setting your DB instance accessibility to private might result in the loss of connectivity to your database.' with a 'Learn more' link.

Creating a DB Instance in a VPC

The following procedures help you create a DB instance in a VPC. If your account has a default VPC, you can begin with step 3 because the VPC and DB subnet group have already been created for you. If your AWS account doesn't have a default VPC, or if you want to create an additional VPC, you can create a new VPC.

If you don't know if you have a default VPC, see [Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform \(p. 506\)](#).

Note

If you want your DB instance in the VPC to be publicly accessible, you must update the DNS information for the VPC by enabling the VPC attributes *DNS hostnames* and *DNS resolution*. For information about updating the DNS information for a VPC instance, see [Updating DNS Support for Your VPC](#).

Follow these steps to create a DB instance in a VPC:

- [Step 1: Create a VPC \(p. 516\)](#)
- [Step 2: Add Subnets to the VPC \(p. 516\)](#)
- [Step 3: Create a DB Subnet Group \(p. 517\)](#)
- [Step 4: Create a VPC Security Group \(p. 518\)](#)
- [Step 5: Create a DB Instance in the VPC \(p. 518\)](#)

Step 1: Create a VPC

If your AWS account does not have a default VPC or if you want to create an additional VPC, follow the instructions for creating a new VPC. See [Create a VPC with Private and Public Subnets \(p. 521\)](#), or see [Step 1: Create a VPC](#) in the Amazon VPC documentation.

Step 2: Add Subnets to the VPC

Once you have created a VPC, you need to create subnets in at least two Availability Zones. You use these subnets when you create a DB subnet group. If you have a default VPC, a subnet is automatically created for you in each Availability Zone in the AWS Region.

For instructions on how to create subnets in a VPC, see [Create a VPC with Private and Public Subnets \(p. 521\)](#).

Step 3: Create a DB Subnet Group

A DB subnet group is a collection of subnets (typically private) that you create for a VPC and that you then designate for your DB instances. A DB subnet group allows you to specify a particular VPC when you create DB instances using the CLI or API. If you use the console, you can just choose the VPC and subnets you want to use. Each DB subnet group must have at least one subnet in at least two Availability Zones in the AWS Region.

Note

For a DB instance to be publicly accessible, the subnets in the DB subnet group must have an Internet gateway. For more information about Internet gateways for subnets, go to [Internet Gateways](#) in the Amazon VPC documentation.

When you create a DB instance in a VPC, you must choose a DB subnet group. Amazon RDS then uses that DB subnet group and your preferred Availability Zone to choose a subnet and an IP address within that subnet. Amazon RDS creates and associates an Elastic Network Interface to your DB instance with that IP address. For Multi-AZ deployments, defining a subnet for two or more Availability Zones in an AWS Region allows Amazon RDS to create a new standby in another Availability Zone should the need arise. You need to do this even for Single-AZ deployments, just in case you want to convert them to Multi-AZ deployments at some point.

In this step, you create a DB subnet group and add the subnets you created for your VPC.

Console

To create a DB subnet group

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Subnet groups**.
3. Choose **Create DB Subnet Group**.
4. For **Name**, type the name of your DB subnet group.
5. For **Description**, type a description for your DB subnet group.
6. For **VPC**, choose the VPC that you created.
7. In the **Add subnets** section, choose **Add all the subnets related to this VPC**.

Create DB subnet group

To create a new Subnet Group give it a name, description, and select an existing VPC below. Once you select an existing VPC, you will be able to add subnets related to that VPC.

Subnet group details

Name	<input type="text" value="mydbsubnetgroup"/>
Description	<input type="text" value="My DB Subnet Group"/>
VPC	VPC Identifier corresponding to the subnets you want to use for the DB subnet group
	<input type="text" value="tutorial-vpc (vpc-971c12ee)"/>

Add subnets

Add subnet(s) to this subnet group. You may add subnets one at a time below or add all the subnets related to this VPC. You may make additions/edits after this group is created. A minimum of 2 subnets is required.

[Add all the subnets related to this VPC](#)

Availability zone	<input type="text" value="select an availability zone"/>
Subnet	<input type="text" value="select a subnet"/> Add subnet

Subnets in this subnet group (4)

Availability zone	Subnet ID	CIDR block	Action
us-east-1a	subnet-d8c8e7f4	10.0.2.0/24	Remove
us-east-1f	subnet-718fdc7d	10.0.3.0/24	Remove
us-east-1a	subnet-cbc8e7e7	10.0.1.0/24	Remove
us-east-1a	subnet-0ccde220	10.0.0.0/24	Remove

[Cancel](#) [Create](#)

8. Choose **Create**.

Your new DB subnet group appears in the DB subnet groups list on the RDS console. You can choose the DB subnet group to see details, including all of the subnets associated with the group, in the details pane at the bottom of the window.

Step 4: Create a VPC Security Group

Before you create your DB instance, you must create a VPC security group to associate with your DB instance. For instructions on how to create a security group for your DB instance, see [Create a VPC Security Group for a Private DB Instance \(p. 524\)](#), or see [Security Groups for Your VPC](#) in the Amazon VPC documentation.

Step 5: Create a DB Instance in the VPC

In this step, you create a DB instance and use the VPC name, the DB subnet group, and the VPC security group you created in the previous steps.

Note

If you want your DB instance in the VPC to be publicly accessible, you must enable the VPC attributes *DNS hostnames* and *DNS resolution*. For information on updating the DNS information for a VPC instance, see [Updating DNS Support for Your VPC](#).

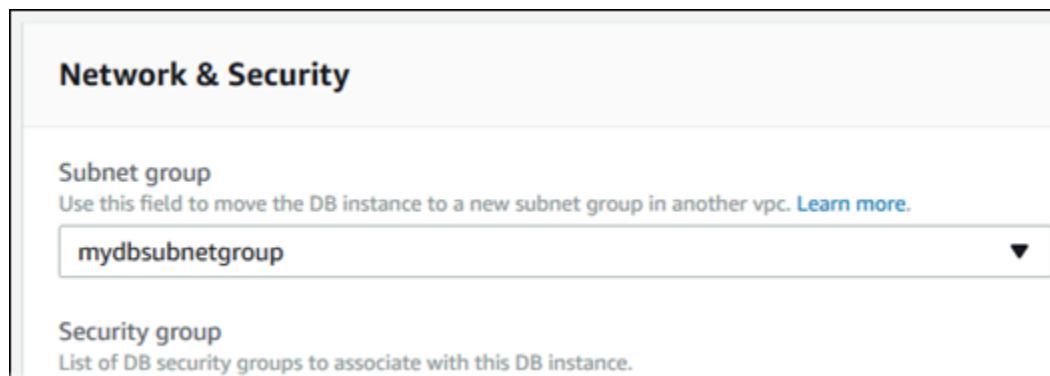
For details on how to create a DB instance for your DB engine, see the topic following that discusses your DB engine. For each engine, when prompted in the **Network & Security** section, enter the VPC name, the DB subnet group, and the VPC security group you created in the previous steps.

Database Engine	Relevant Documentation
MariaDB	Creating a DB Instance Running the MariaDB Database Engine (p. 543)
Microsoft SQL Server	Creating a DB Instance Running the Microsoft SQL Server Database Engine (p. 616)
MySQL	Creating a DB Instance Running the MySQL Database Engine (p. 751)
Oracle	Creating a DB Instance Running the Oracle Database Engine (p. 942)
PostgreSQL	Creating a DB Instance Running the PostgreSQL Database Engine (p. 1257)

Updating the VPC for a DB Instance

You can use the AWS Management Console to easily move your DB instance to a different VPC.

For details on how to modify a DB instance for your DB engine, see the topic in the table following that discusses your DB engine. In the **Network & Security** section of the modify page, shown following, for **Subnet group**, enter the new subnet group. The new subnet group must be a subnet group in a new VPC.



Database Engine	Relevant Documentation
MariaDB	Modifying a DB Instance Running the MariaDB Database Engine (p. 564)
Microsoft SQL Server	Modifying a DB Instance Running the Microsoft SQL Server Database Engine (p. 643)
MySQL	Modifying a DB Instance Running the MySQL Database Engine (p. 802)
Oracle	Modifying a DB Instance Running the Oracle Database Engine (p. 966)
PostgreSQL	Modifying a DB Instance Running the PostgreSQL Database Engine (p. 1279)

Moving a DB Instance Not in a VPC into a VPC

Some legacy DB instances on the EC2-Classic platform are not in a VPC. If your DB instance is not in a VPC, you can use the AWS Management Console to easily move your DB instance into a VPC. Before you can move a DB instance not in a VPC, into a VPC, you must create the VPC.

Follow these steps to create a VPC for your DB instance.

- [Step 1: Create a VPC \(p. 516\)](#)
- [Step 2: Add Subnets to the VPC \(p. 516\)](#)
- [Step 3: Create a DB Subnet Group \(p. 517\)](#)
- [Step 4: Create a VPC Security Group \(p. 518\)](#)

After you create the VPC, follow these steps to move your DB instance into the VPC.

- [Updating the VPC for a DB Instance \(p. 519\)](#)

The following are some limitations to moving your DB instance into the VPC.

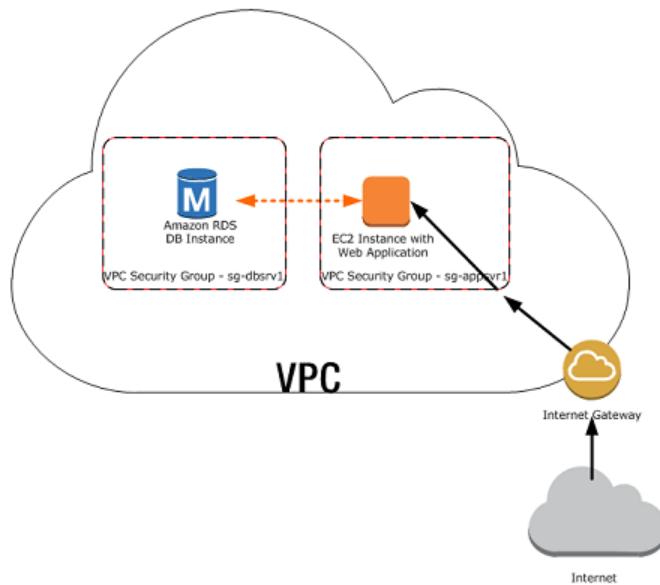
- Moving a Multi-AZ DB instance not in a VPC into a VPC is not currently supported.
- Moving a DB instance with Read Replicas not in a VPC into a VPC is not currently supported.

If you move your DB instance into a VPC, and you are using a custom option group with your DB instance, then you need to change the option group that is associated with your DB instance. Option groups are platform-specific, and moving to a VPC is a change in platform. To use a custom option group in this case, assign the default VPC option group to the DB instance, assign an option group that is used by other DB instances in the VPC you are moving to, or create a new option group and assign it to the DB instance. For more information, see [Working with Option Groups \(p. 196\)](#).

Tutorial: Create an Amazon VPC for Use with a DB Instance

A common scenario includes a DB instance in an Amazon VPC, that shares data with a web server that is running in the same VPC. In this tutorial you create the VPC for this scenario.

The following diagram shows this scenario. For information about other scenarios, see [Scenarios for Accessing a DB Instance in a VPC \(p. 507\)](#).



Because your DB instance only needs to be available to your web server, and not to the public Internet, you create a VPC with both public and private subnets. The web server is hosted in the public subnet, so that it can reach the public Internet. The DB instance is hosted in a private subnet. The web server is able to connect to the DB instance because it is hosted within the same VPC, but the DB instance is not available to the public Internet, providing greater security.

Create a VPC with Private and Public Subnets

Use the following procedure to create a VPC with both public and private subnets.

To create a VPC and subnets

1. Open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. In the top-right corner of the AWS Management Console, choose the region to create your VPC in. This example uses the US West (Oregon) region.
3. In the upper-left corner, choose **VPC Dashboard**. To begin creating a VPC, choose **Launch VPC Wizard**.
4. On the **Step 1: Select a VPC Configuration** page, choose **VPC with Public and Private Subnets**, and then choose **Select**.
5. On the **Step 2: VPC with Public and Private Subnets** page, set these values:
 - **IPv4 CIDR block:** 10.0.0.0/16
 - **IPv6 CIDR block:** No IPv6 CIDR Block
 - **VPC name:** tutorial-vpc

- **Public subnet's IPv4 CIDR:** 10.0.0.0/24
- **Availability Zone:** us-west-2a
- **Public subnet name:** Tutorial public
- **Private subnet's IPv4 CIDR:** 10.0.1.0/24
- **Availability Zone:** us-west-2a
- **Private subnet name:** Tutorial Private 1
- **Instance type:** t2.small

Important

If you don't see the **Instance type** box in the console, choose **Use a NAT instance instead**. This link is on the right.

Note

If the t2.small instance type is not listed, you can choose a different instance type.

- **Key pair name:** No key pair
- **Service endpoints:** Skip this field.
- **Enable DNS hostnames:** Yes
- **Hardware tenancy:** Default

6. When you're finished, choose **Create VPC**.

Create Additional Subnets

You must have either two private subnets or two public subnets available to create a DB subnet group for a DB instance to use in a VPC. Because the DB instance for this tutorial is private, add a second private subnet to the VPC.

To create an additional subnet

1. Open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. To add the second private subnet to your VPC, choose **VPC Dashboard**, choose **Subnets**, and then choose **Create subnet**.
3. On the **Create subnet** page, set these values:
 - **Name tag:** Tutorial private 2
 - **VPC:** Choose the VPC that you created in the previous step, for example: vpc-*identifier* (10.0.0.0/16) | tutorial-vpc
 - **Availability Zone:** us-west-2b

Note

Choose an Availability Zone that is different from the one that you chose for the first private subnet.

- **IPv4 CIDR block:** 10.0.2.0/24
4. When you're finished, choose **Create**. Next, choose **Close** on the confirmation page.
 5. To ensure that the second private subnet that you created uses the same route table as the first private subnet, choose **VPC Dashboard**, choose **Subnets**, and then choose the first private subnet that you created for the VPC, Tutorial private 1.
 6. Below the list of subnets, choose the **Route Table** tab, and note the value for **Route Table**—for example: rtb-98b613fd.
 7. In the list of subnets, deselect the first private subnet.
 8. In the list of subnets, choose the second private subnet Tutorial private 2, and choose the **Route Table** tab.

9. If the current route table is not the same as the route table for the first private subnet, choose **Edit route table association**. For **Route Table ID**, choose the route table that you noted earlier—for example: rtb-98b613fd. Next, to save your selection, choose **Save**.

Create a VPC Security Group for a Public Web Server

Next you create a security group for public access. To connect to public instances in your VPC, you add inbound rules to your VPC security group that allow traffic to connect from the internet.

To create a VPC security group

1. Open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. Choose **VPC Dashboard**, choose **Security Groups**, and then choose **Create security group**.
3. On the **Create security group** page, set these values:
 - **Security group name:** tutorial-securitygroup
 - **Description:** Tutorial Security Group
 - **VPC:** Choose the VPC that you created earlier, for example: vpc-*identifier* (10.0.0.0/16 | tutorial-vpc)
4. To create the security group, choose **Create**. Next, choose **Close** on the confirmation page.

Note the security group ID because you need it later in this tutorial.

To add inbound rules to the security group

1. Determine the IP address to use to connect to instances in your VPC. To determine your public IP address, you can use the service at <https://checkip.amazonaws.com>. An example of an IP address is 203.0.113.25/32.

If you are connecting through an Internet service provider (ISP) or from behind your firewall without a static IP address, you need to find out the range of IP addresses used by client computers.

Warning

If you use 0.0.0.0/0, you enable all IP addresses to access your public instances. This approach is acceptable for a short time in a test environment, but it's unsafe for production environments. In production, you'll authorize only a specific IP address or range of addresses to access your instances.

2. Open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
3. Choose **VPC Dashboard**, choose **Security Groups**, and then choose the tutorial-securitygroup security group that you created in the previous procedure.
4. Under the list of security groups, choose the **Inbound Rules** tab, and then choose **Edit rules**.
5. On the **Edit inbound rules** page, choose **Add Rule**.
6. Set the following values for your new inbound rule to allow Secure Shell (SSH) access to your EC2 instance. If you do this, you can connect to your EC2 instance to install the web server and other utilities, and to upload content for your web server.
 - **Type:** SSH
 - **Source:** The IP address or range from Step 1, for example: 203.0.113.25/32.
7. Choose **Add rule**.
8. Set the following values for your new inbound rule to allow HTTP access to your web server.
 - **Type:** HTTP
 - **Source:** 0.0.0.0/0.

9. To save your settings, choose **Save rules**. Next, choose **Close** on the confirmation page.

Create a VPC Security Group for a Private DB Instance

To keep your DB instance private, create a second security group for private access. To connect to private instances in your VPC, you add inbound rules to your VPC security group that allow traffic from your web server only.

To create a VPC security group

1. Open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. Choose **VPC Dashboard**, choose **Security Groups**, and then choose **Create security group**.
3. On the **Create security group** page, set these values:
 - **Security group name:** tutorial-db-securitygroup
 - **Description:** Tutorial DB Instance Security Group
 - **VPC:** Choose the VPC that you created earlier, for example: vpc-*identifier* (10.0.0.0/16) | tutorial-vpc
4. To create the security group, choose **Create**. Next, choose **Close** on the confirmation page.

To add inbound rules to the security group

1. Open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. Choose **VPC Dashboard**, choose **Security Groups**, and then choose the tutorial-db-securitygroup security group that you created in the previous procedure.
3. Under the list of security groups, choose the **Inbound Rules** tab, and then choose **Edit rules**.
4. On the **Edit inbound rules** page, choose **Add Rule**.
5. Set the following values for your new inbound rule to allow MySQL traffic on port 3306 from your EC2 instance. If you do this, you can connect from your web server to your DB instance to store and retrieve data from your web application to your database.
 - **Type:** MySQL/Aurora
 - **Source:** The identifier of the tutorial-securitygroup security group that you created previously in this tutorial, for example: sg-9edd5cfb.
6. To save your settings, choose **Save rules**. Next, choose **Close** on the confirmation page.

Create a DB Subnet Group

A DB subnet group is a collection of subnets that you create in a VPC and that you then designate for your DB instances. A DB subnet group allows you to specify a particular VPC when creating DB instances.

To create a DB subnet group

1. Open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Subnet groups**.
3. Choose **Create DB Subnet Group**.
4. On the **Create DB subnet group** page, set these values in **Subnet group details**:
 - **Name:** tutorial-db-subnet-group
 - **Description:** Tutorial DB Subnet Group
 - **VPC:** tutorial-vpc (vpc-*identifier*)

5. In the **Add subnets** section, choose **Add all the subnets related to this VPC**.
6. Choose **Create**.

Your new DB subnet group appears in the DB subnet groups list on the RDS console. You can click the DB subnet group to see details, including all of the subnets associated with the group, in the details pane at the bottom of the window.

Amazon RDS on AWS Outposts (Preview)

This is prerelease documentation for a service in preview release. It is subject to change.

By using AWS Outposts, you can work with the same AWS hardware infrastructure, services, APIs, and tools to build and run your applications on premises as in the cloud. AWS Outposts is ideal for workloads that need low latency access to on-premises applications or systems and for workloads that need to store and process data locally. For more information about AWS Outposts, see [AWS Outposts](#).

With Amazon RDS on AWS Outposts, you can create AWS-managed relational databases in your on-premises data centers. RDS on Outposts enables you to run RDS databases on AWS Outposts. You manage RDS on Outposts databases using the same AWS Management Console, AWS CLI, and RDS API that you use for Amazon RDS databases in the cloud.

Important

Currently, RDS on Outposts is in preview. Don't use RDS on Outposts for databases running production workloads. Also, we strongly recommend not putting sensitive data in databases that you use with the RDS on Outposts preview. You might encounter issues during the preview that cause incorrect results, corrupted data, or both. Over the duration of the preview, we might introduce breaking changes without advance notice.

RDS on Outposts supports automated backups of DB instances. Automated backups enable point-in-time recovery. You can also manually create DB snapshots. Network connectivity between your Outpost and your AWS Region is required to back up and restore DB instances. All DB snapshots and transaction logs from an Outpost are stored in the AWS Region. From the AWS Region, you can restore a DB instance from a DB snapshot to a different Outpost. For more information, see [Working With Backups \(p. 254\)](#).

RDS on Outposts supports automated maintenance and upgrades of DB instances. For more information, see [Maintaining a DB Instance \(p. 152\)](#).

RDS on Outposts uses encryption at rest for DB instances and DB snapshots using your AWS KMS key. For more information about encryption at rest, see [Encrypting Amazon RDS Resources \(p. 429\)](#).

If network connectivity between your Outpost and its AWS Region is lost, your databases continue to run. You can't create new databases or take new actions on existing databases. If an Outpost's capacity becomes unhealthy while you are disconnected, the database isn't automatically replaced. Additionally, automatic backups don't run if network connectivity is down during the backup window. If there is a hardware failure in disconnected mode, all reads and writes to your databases are disabled. They're disabled because the underlying Amazon EBS volumes stop accepting reads and writes to avoid data loss.

Note

If network connectivity is lost, we recommend that you restore network connectivity as soon as possible for the best RDS on Outposts experience.

Topics

- [Prerequisites for Amazon RDS on AWS Outposts \(p. 527\)](#)
- [Limitations for Amazon RDS on AWS Outposts \(p. 527\)](#)
- [Supported DB Instance Classes for Amazon RDS on AWS Outposts \(p. 527\)](#)
- [Creating DB Instances for Amazon RDS on AWS Outposts \(p. 528\)](#)

Prerequisites for Amazon RDS on AWS Outposts

The following are prerequisites for using Amazon RDS on AWS Outposts:

- You must have installed and configured AWS Outposts in your on-premises data center. For more information about AWS Outposts, see [AWS Outposts](#).
- You must have at least one subnet available for RDS on Outposts. You can use the same subnet for other workloads.
- You must have a reliable network connection between your Outpost and an AWS Region.

Limitations for Amazon RDS on AWS Outposts

The following are limitations for Amazon RDS on AWS Outposts:

- You can only create DB instances for RDS for MySQL and RDS for PostgreSQL DB instances. Currently, other DB engines aren't supported.
- RDS on Outposts doesn't support use cases that require all data to remain in your data center.

RDS on Outposts stores database backups and logs in your AWS Region and sends database metrics to Amazon CloudWatch.
- Some Amazon RDS features aren't supported.

Supported DB Instance Classes for Amazon RDS on AWS Outposts

Amazon RDS on AWS Outposts supports the following DB instance classes:

- General Purpose DB instance classes
 - db.m5.24xlarge
 - db.m5.12xlarge
 - db.m5.4xlarge
 - db.m5.2xlarge
 - db.m5.xlarge
 - db.m5.large
- Memory Optimized DB instance classes
 - db.r5.24xlarge
 - db.r5.12xlarge
 - db.r5.4xlarge
 - db.r5.2xlarge
 - db.r5.xlarge
 - db.r5.large

Only General Purpose SSD storage is supported for RDS on Outposts DB instances. For more information about DB instance classes, see [Choosing the DB Instance Class \(p. 107\)](#).

Creating DB Instances for Amazon RDS on AWS Outposts

Creating an Amazon RDS on AWS Outposts DB instance is similar to creating an Amazon RDS DB instance in the AWS Cloud. However, you must specify a DB subnet group that is associated with your Outpost.

An Amazon VPC can span all of the Availability Zones in an AWS Region. In AWS Outposts, Outposts are extensions of Availability Zones, and you can extend an Amazon VPC in an account to span multiple Availability Zones and associated Outpost locations. When you configure your Outpost, you associate a subnet group with it to extend your regional Amazon VPC environment to your on-premises facility. Outpost instances and related services appear as part of your regional Amazon VPC, similar to an Availability Zone with associated subnets.

Before you create an RDS on Outposts DB instance, you can create a DB subnet group that includes one subnet that is associated with your Outpost. When you create an RDS on Outposts DB instance, specify this DB subnet group.

For information about configuring AWS Outposts, see the [AWS Outposts User Guide](#).

Note

The RDS console is the only interface supported for the RDS on Outposts preview. The AWS CLI and RDS API aren't supported.

To create an RDS on Outposts DB instance using the console

1. Create a DB subnet group with one subnet that is associated with your Outpost.

Note

To create a DB subnet group for the AWS Cloud, you specify at least two subnets. However, for an Outpost DB subnet group, you can specify only one subnet.

- a. Sign in to the AWS Management Console for the preview, and open the Amazon RDS console.
- b. Choose **Subnet groups**, and then choose **Create DB Subnet Group**.

The **Create DB subnet group** page appears.

RDS > Subnet groups > Create DB subnet group

Create DB subnet group

To create a new subnet group, give it a name and a description, and choose an existing VPC. You will then be able to add subnets related to that VPC.

Subnet group details

Name
You won't be able to modify the name after your subnet group has been created.
 Must contain from 1 to 255 characters. Alphanumeric characters, spaces, hyphens, underscores, and periods are allowed.

Description

VPC
Choose a VPC identifier that corresponds to the subnets you want to use for your DB subnet group. You won't be able to choose a different VPC identifier after your subnet group has been created.

Add subnets

Add subnet(s) to this subnet group. You may add subnets one at a time below or add all the subnets related to this VPC. You may make additions/edits after this group is created. A minimum of 2 subnets is required.

Availability zone

- c. Set the following values for your new DB subnet group:
 - **Name** – The name of the DB subnet group
 - **Description** – A description for the DB subnet group
 - **VPC** – The VPC for which you're creating the DB subnet group
- d. For **Availability Zone**, choose the Availability Zone for your Outpost.
- e. For **Subnet**, choose the subnet for use by RDS on Outposts.
- f. Choose **Add subnet**.

Your DB subnet group must have only one subnet.

- g. Choose **Create** to create the DB subnet group.

2. Choose the Outpost for your DB instance.

The AWS Management Console detects available Outposts that you have configured. Choose the Outpost with the DB subnet group that you created previously.

RDS > Create database

Create database

Database location
Choose a location to meet your use case [Info](#)

Amazon Cloud
Use Amazon's cloud to store and provision a database instance with RDS.

On-premises
Create a DB instance on-premises using a AWS Outpost or VMware datacenter.

Choose an on-premises creation method [Info](#)

RDS on Outposts
Use an AWS Outpost to create a DB instance on-premises

RDS on VMware
Use a custom AZ to create a DB instance with an on-premises VMware datacenter.

Outposts Connectivity [Info](#)

Virtual Private Cloud (VPC) [Info](#)
VPC that defines the virtual networking environment for this DB instance.

outpost-vpc1 (vpc-412d2862) ▾

Only VPCs with an Outpost subnet are listed.

Subnet group [Info](#)
DB subnet group that defines the subnet and IP ranges of the VPC.

Choose existing
Choose existing subnet group

Create new
Create new subnet group

Select subnets for DB subnet group

Add DB subnet group ▾

subnet-23c23ada9a (Outpost01) X

Choose the following settings:

- **Database location – On-premises**
- **On-premises creation method – RDS on Outposts**
- **Outpost** – The Outpost that uses the Amazon VPC with the DB subnet group for the DB instance
- **Virtual Private Cloud (VPC)** – The Amazon VPC that contains the DB subnet group for the DB instance
- **Subnet group** – The DB subnet group for the DB instance

You can choose an existing DB subnet group that is associated with the Outpost. If you didn't create a DB subnet group, you can create a new DB subnet group for the Outpost. Only one subnet is allowed in the DB subnet group.

3. For the remaining sections, specify your DB instance settings.

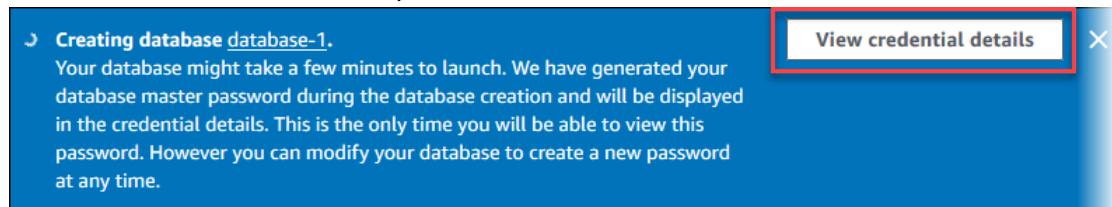
For information about each setting when creating a MySQL DB instance, see [Settings for MySQL DB Instances \(p. 758\)](#).

For information about each setting when creating a PostgreSQL DB instance, see [Settings for PostgreSQL DB Instances \(p. 1265\)](#).

4. Choose **Create database**.

If you chose to use an automatically generated password, the **View credential details** button appears on the **Databases** page.

To view the master user name and password for the DB instance, choose **View credential details**.



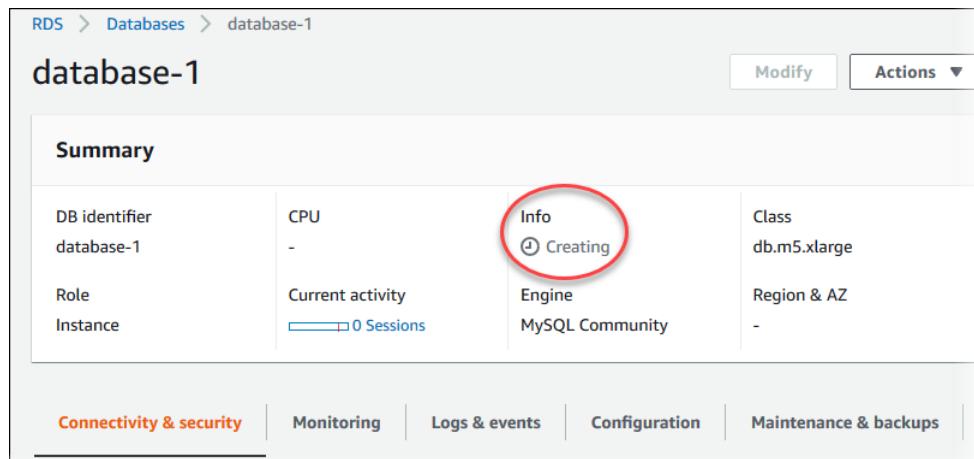
To connect to the DB instance as the master user, use the user name and password that appear.

Important

You can't view the master user password again. If you don't record it, you might have to change it. To change the master user password after the DB instance is available, modify the DB instance. For more information about modifying a DB instance, see [Modifying a DB Instance Running the MySQL Database Engine \(p. 802\)](#) or [Modifying a DB Instance Running the PostgreSQL Database Engine \(p. 1279\)](#).

5. For **Databases**, choose the name of the new DB instance.

On the RDS console, the details for the new DB instance appear. The DB instance has a status of **Creating** until the DB instance is created and ready for use. When the state changes to **Available**, you can connect to the DB instance. Depending on the DB instance class and storage allocated, it can take several minutes for the new DB instance to be available.



After the DB instance is available, you can manage it the same way that you manage RDS DB instances in the cloud.

MariaDB on Amazon RDS

Amazon RDS supports DB instances running several versions of MariaDB. You can use the following major versions:

- MariaDB 10.3
- MariaDB 10.2
- MariaDB 10.1
- MariaDB 10.0

For more information about minor version support, see [MariaDB on Amazon RDS Versions \(p. 534\)](#).

You first use the Amazon RDS management tools or interfaces to create an Amazon RDS MariaDB DB instance. You can then use the Amazon RDS tools to perform management actions for the DB instance, such as reconfiguring or resizing the DB instance, authorizing connections to the DB instance, creating and restoring from backups or snapshots, creating Multi-AZ secondaries, creating Read Replicas, and monitoring the performance of the DB instance. You use standard MariaDB utilities and applications to store and access the data in the DB instance.

MariaDB is available in all of the AWS Regions. For more information about AWS Regions, see [Regions and Availability Zones \(p. 132\)](#).

You can use Amazon RDS for MariaDB databases to build HIPAA-compliant applications. You can store healthcare-related information, including protected health information (PHI), under an executed Business Associate Agreement (BAA) with AWS. For more information, see [HIPAA Compliance](#). AWS Services in Scope have been fully assessed by a third-party auditor and result in a certification, attestation of compliance, or Authority to Operate (ATO). For more information, see [AWS Services in Scope by Compliance Program](#).

Before creating your first DB instance, you should complete the steps in the setting up section of this guide. For more information, see [Setting Up for Amazon RDS \(p. 5\)](#).

Common Management Tasks for MariaDB on Amazon RDS

The following are the common management tasks you perform with an Amazon RDS DB instance running MariaDB, with links to relevant documentation for each task.

Task Area	Relevant Documentation
Instance Classes, Storage, and PIOPS If you are creating a DB instance for production purposes, you should understand how instance classes, storage types, and Provisioned IOPS work in Amazon RDS.	Choosing the DB Instance Class (p. 107) Amazon RDS Storage Types (p. 135)
Multi-AZ Deployments Provide high availability with synchronous standby replication in a different Availability Zone, automatic failover, fault tolerance for DB instances using Multi-AZ deployments, and Read Replicas.	High Availability (Multi-AZ) for Amazon RDS (p. 144)

Task Area	Relevant Documentation
Amazon Virtual Private Cloud (VPC) If your AWS account has a default VPC, then your DB instance is automatically created inside the default VPC. If your account does not have a default VPC, and you want the DB instance in a VPC, you must create the VPC and subnet groups before you create the DB instance.	Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform (p. 506) Working with a DB Instance in a VPC (p. 514)
Security Groups By default, DB instances are created with a firewall that prevents access to them. You therefore must create a security group with the correct IP addresses and network configuration to access the DB instance. The security group you create depends on what Amazon EC2 platform your DB instance is on, and whether you access your DB instance from an Amazon EC2 instance. In general, if your DB instance is on the <i>EC2-Classic</i> platform, you will need to create a DB security group; if your DB instance is on the EC2-VPC platform, you will need to create a VPC security group.	Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform (p. 506) Controlling Access with Security Groups (p. 486)
Parameter Groups If your DB instance is going to require specific database parameters, you should create a parameter group before you create the DB instance.	Working with DB Parameter Groups (p. 211)
Importing and Exporting Data Establish procedures for importing or exporting data.	Importing Data into a MariaDB DB Instance (p. 586)
Replication You can offload read traffic from your primary MariaDB DB instance by creating Read Replicas.	Working with Read Replicas (p. 182)
Connecting to Your DB Instance Connect to your DB instance using a standard SQL client application.	Connecting to a DB Instance Running the MariaDB Database Engine (p. 555)
Backup and Restore When you create your DB instance, you can configure it to take automated backups. You can also back up and restore your databases manually by using full backup files (.bak files).	Working With Backups (p. 254)
Monitoring Monitor your RDS MariaDB DB instance by using Amazon CloudWatch RDS metrics, events, and Enhanced Monitoring. View log files for your RDS MariaDB DB instance.	Viewing DB Instance Metrics (p. 304) Viewing Amazon RDS Events (p. 388)
Log Files You can access the log files for your MariaDB DB instance.	Amazon RDS Database Log Files (p. 390) MariaDB Database Log Files (p. 394)

There are also advanced administrative tasks for working with DB instances running MariaDB. For more information, see the following documentation:

- [Parameters for MariaDB \(p. 590\)](#)
- [MariaDB on Amazon RDS SQL Reference \(p. 595\)](#)

MariaDB on Amazon RDS Versions

For MariaDB, version numbers are organized as version X.Y.Z. In Amazon RDS terminology, X.Y denotes the major version, and Z is the minor version number. For Amazon RDS implementations, a version change is considered major if the major version number changes, for example going from version 10.0 to 10.1. A version change is considered minor if only the minor version number changes, for example going from version 10.0.17 to 10.0.24.

Amazon RDS currently supports the following versions of MariaDB:

Major Version	Minor Version
MariaDB 10.3	<ul style="list-style-type: none">• 10.3.13 (supported in all AWS Regions)• 10.3.8 (supported in all AWS Regions)
MariaDB 10.2	<ul style="list-style-type: none">• 10.2.21 (supported in all AWS Regions)• 10.2.15 (supported in all AWS Regions)• 10.2.12 (supported in all AWS Regions)• 10.2.11 (supported in all AWS Regions)
MariaDB 10.1	<ul style="list-style-type: none">• 10.1.34 (supported in all AWS Regions)• 10.1.31 (supported in all AWS Regions)• 10.1.26 (supported in all AWS Regions)• 10.1.23 (supported in all AWS Regions)• 10.1.19 (supported in all AWS Regions)• 10.1.14 (supported in all AWS Regions except us-east-2)
MariaDB 10.0	<ul style="list-style-type: none">• 10.0.35 (supported in all AWS Regions)• 10.0.34 (supported in all AWS Regions)• 10.0.32 (supported in all AWS Regions)• 10.0.31 (supported in all AWS Regions)• 10.0.28 (supported in all AWS Regions)• 10.0.24 (supported in all AWS Regions)• 10.0.17 (supported in all AWS Regions except us-east-2, ca-central-1, eu-west-2)

You can specify any currently supported MariaDB version when creating a new DB instance. You can specify the major version (such as MariaDB 10.2), and any supported minor version for the specified major version. If no version is specified, Amazon RDS defaults to a supported version, typically the most recent version. If a major version is specified but a minor version is not, Amazon RDS defaults to a recent release of the major version you have specified. To see a list of supported versions, as well as defaults for newly created DB instances, use the [describe-db-engine-versions](#) AWS CLI command.

For information about the Amazon RDS deprecation policy for MariaDB, see [Amazon RDS FAQs](#).

Version and Feature Support on Amazon RDS

MariaDB 10.3 Support on Amazon RDS

Amazon RDS supports the following versions of MariaDB 10.3:

- 10.3.8 (supported in all AWS Regions)

Amazon RDS supports the following new features for your DB instances running MariaDB version 10.3 or later:

- **Oracle compatibility** – PL/SQL compatibility parser, sequences, INTERSECT and EXCEPT to complement UNION, new TYPE OF and ROW TYPE OF declarations, and invisible columns
- **Temporal data processing** – System versioned tables for querying of past and present states of the database
- **Flexibility** – User-defined aggregates, storage-independent column compression, and proxy protocol support to relay the client IP address to the server
- **Manageability** – Instant ADD COLUMN operations and fast-fail data definition language (DDL) operations

For a list of all MariaDB 10.3 features and their documentation, see [Changes & Improvements in MariaDB 10.3](#) and [Release Notes - MariaDB 10.3 Series](#) on the MariaDB website.

For a list of unsupported features, see [Features Not Supported \(p. 536\)](#).

MariaDB 10.2 Support on Amazon RDS

Amazon RDS supports the following versions of MariaDB 10.2:

- 10.2.15 (supported in all AWS Regions)
- 10.2.12 (supported in all AWS Regions)
- 10.2.11 (supported in all AWS Regions)

Amazon RDS supports the following new features for your DB instances running MariaDB version 10.2 or later:

- ALTER USER
- Common Table Expressions
- Compressing Events to Reduce Size of the Binary Log
- CREATE USER — new options for limiting resource usage and TLS/SSL
- EXECUTE IMMEDIATE
- Flashback
- InnoDB — now the default storage engine instead of XtraDB
- InnoDB — set the buffer pool size dynamically
- JSON Functions
- Window Functions
- WITH

For a list of all MariaDB 10.2 features and their documentation, see [Changes & Improvements in MariaDB 10.2](#) and [Release Notes - MariaDB 10.2 Series](#) on the MariaDB website.

For a list of unsupported features, see [Features Not Supported \(p. 536\)](#).

MariaDB 10.1 Support on Amazon RDS

Amazon RDS supports the following versions of MariaDB 10.1:

- 10.1.34 (supported in all AWS Regions)
- 10.1.31 (supported in all AWS Regions)
- 10.1.26 (supported in all AWS Regions)
- 10.1.23 (supported in all AWS Regions)
- 10.1.19 (supported in all AWS Regions)
- 10.1.14 (supported in all AWS Regions except us-east-2)

Amazon RDS supports the following new features for your DB instances running MariaDB version 10.1 or later:

- Optimistic in-order parallel replication
- Page Compression
- XtraDB data scrubbing and defragmentation

For a list of all MariaDB 10.1 features and their documentation, see [Changes & Improvements in MariaDB 10.1](#) and [Release Notes - MariaDB 10.1 Series](#) on the MariaDB website.

For a list of unsupported features, see [Features Not Supported \(p. 536\)](#).

MariaDB 10.0 Support on Amazon RDS

Amazon RDS supports the following versions of MariaDB 10.0:

- 10.0.35 (supported in all AWS Regions)
- 10.0.34 (supported in all AWS Regions)
- 10.0.32 (supported in all AWS Regions)
- 10.0.31 (supported in all AWS Regions)
- 10.0.28 (supported in all AWS Regions)
- 10.0.24 (supported in all AWS Regions)
- 10.0.17 (supported in all AWS Regions except us-east-2, ca-central-1, eu-west-2)

For a list of all MariaDB 10.0 features and their documentation, see [Changes & Improvements in MariaDB 10.0](#) and [Release Notes - MariaDB 10.0 Series](#) on the MariaDB website.

For a list of unsupported features, see [Features Not Supported \(p. 536\)](#).

Features Not Supported

The following MariaDB features are not supported on Amazon RDS:

- Authentication plugin – GSSAPI
- Authentication plugin – Unix Socket
- AWS Key Management encryption plugin

- Delayed replication
- Native MariaDB encryption at rest for XtraDB, InnoDB, and Aria.

You can enable encryption at rest for a MariaDB DB instance by following the instructions in [Encrypting Amazon RDS Resources \(p. 429\)](#).

- HandlerSocket
- JSON table type
- MariaDB ColumnStore
- MariaDB Galera Cluster
- Multisource replication
- MyRocks storage engine
- Password validation plugin, `simple_password_check`, and `cracklib_password_check`
- Replication filters
- Spider storage engine
- Sphinx storage engine
- TokuDB storage engine
- Storage engine-specific object attributes, as described in [Engine-defined New Table/Field/Index Attributes](#) in the MariaDB documentation
- Table and tablespace encryption

To deliver a managed service experience, Amazon RDS doesn't provide shell access to DB instances, and it restricts access to certain system procedures and tables that require advanced privileges. Amazon RDS supports access to databases on a DB instance using any standard SQL client application. Amazon RDS doesn't allow direct host access to a DB instance by using Telnet, Secure Shell (SSH), or Windows Remote Desktop Connection.

Supported Storage Engines for MariaDB on Amazon RDS

While MariaDB supports multiple storage engines with varying capabilities, not all of them are optimized for recovery and data durability. InnoDB (for version 10.2 and higher) and XtraDB (for version 10.0 and 10.1) are the recommended and supported storage engines for MariaDB DB instances on Amazon RDS. Amazon RDS features such as Point-In-Time Restore and snapshot restore require a recoverable storage engine and are supported only for the recommended storage engine for the MariaDB version. Amazon RDS also supports Aria, although using Aria might have a negative impact on recovery in the event of an instance failure. However, if you need to use spatial indexes to handle geographic data on MariaDB 10.1 or 10.0, you should use Aria because spatial indexes are not supported by XtraDB. On MariaDB 10.2 and higher, the InnoDB storage engine supports spatial indexes.

Other storage engines are not currently supported by Amazon RDS for MariaDB.

MariaDB Security on Amazon RDS

Security for Amazon RDS MariaDB DB instances is managed at three levels:

- AWS Identity and Access Management controls who can perform Amazon RDS management actions on DB instances. When you connect to AWS using IAM credentials, your IAM account must have IAM policies that grant the permissions required to perform Amazon RDS management operations. For more information, see [Identity and Access Management in Amazon RDS \(p. 445\)](#).

- When you create a DB instance, you use either a VPC security group or a DB security group to control which devices and Amazon EC2 instances can open connections to the endpoint and port of the DB instance. These connections can be made using Secure Socket Layer (SSL). In addition, firewall rules at your company can control whether devices running at your company can open connections to the DB instance.
- Once a connection has been opened to a MariaDB DB instance, authentication of the login and permissions are applied the same way as in a stand-alone instance of MariaDB. Commands such as `CREATE USER`, `RENAME USER`, `GRANT`, `REVOKE`, and `SET PASSWORD` work just as they do in stand-alone databases, as does directly modifying database schema tables.

When you create an Amazon RDS DB instance, the master user has the following default privileges:

- `alter`
- `alter routine`
- `create`
- `create routine`
- `create temporary tables`
- `create user`
- `create view`
- `delete`
- `drop`
- `event`
- `execute`
- `grant option`
- `index`
- `insert`
- `lock tables`
- `process`
- `references`
- `reload`

This privilege is limited on Amazon RDS MariaDB DB instances. It doesn't grant access to the `FLUSH LOGS` or `FLUSH TABLES WITH READ LOCK` operations.

- `replication client`
- `replication slave`
- `select`
- `show databases`
- `show view`
- `trigger`
- `update`

For more information about these privileges, see [User Account Management](#) in the MariaDB documentation.

Note

Although you can delete the master user on a DB instance, we don't recommend doing so. To recreate the master user, use the `ModifyDBInstance` API or the `modify-db-instance` AWS command line tool and specify a new master user password with the appropriate parameter. If the master user does not exist in the instance, the master user is created with the specified password.

To provide management services for each DB instance, the `rdsadmin` user is created when the DB instance is created. Attempting to drop, rename, change the password for, or change privileges for the `rdsadmin` account results in an error.

To allow management of the DB instance, the standard `kill` and `kill_query` commands have been restricted. The Amazon RDS commands `mysql.rds_kill`, `mysql.rds_kill_query`, and `mysql.rds_kill_query_id` are provided for use in MariaDB and also MySQL so that you can terminate user sessions or queries on DB instances.

Using SSL with a MariaDB DB Instance

Amazon RDS supports Secure Sockets Layer (SSL) connections with DB instances running the MariaDB database engine.

Amazon RDS creates an SSL certificate and installs the certificate on the DB instance when Amazon RDS provisions the instance. These certificates are signed by a certificate authority. The SSL certificate includes the DB instance endpoint as the Common Name (CN) for the SSL certificate to guard against spoofing attacks.

For information about downloading certificates, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).

MariaDB uses yaSSL for secure connections in the following versions:

- MariaDB version 10.1.26 and earlier 10.1 versions
- MariaDB version 10.0.32 and earlier 10.0 versions

MariaDB uses OpenSSL for secure connections in the following versions:

- MariaDB 10.3 versions
- MariaDB 10.2 versions
- MariaDB version 10.1.31 and later 10.1 versions
- MariaDB version 10.0.34 and later 10.0 versions

Amazon RDS for MariaDB supports Transport Layer Security (TLS) versions 1.0, 1.1, and 1.2. The following table shows the TLS support for MySQL versions.

MariaDB Version	TLS 1.0	TLS 1.1	TLS 1.2
MariaDB 10.3	Supported	Supported	Supported
MariaDB 10.2	Supported	Supported	Supported
MariaDB 10.1	Supported	Supported for 10.1.31 and later 10.1 versions	Supported for 10.1.31 and later 10.1 versions
MariaDB 10.0	Supported	Supported for 10.0.34 and later 10.0 versions	Supported for 10.0.34 and later 10.0 versions

To encrypt connections using the default `mysql` client, launch the `mysql` client using the `--ssl-ca` parameter to reference the public key, as shown in the examples following.

The following example shows how to launch the client using the `--ssl-ca` parameter for MariaDB 10.2 and later.

```
mysql -h myinstance.c9akciq32.rds-us-east-1.amazonaws.com  
--ssl-ca=[full path]rds-combined-ca-bundle.pem --ssl-mode=REQUIRED
```

The following example shows how to launch the client using the `--ssl-ca` parameter for MariaDB 10.1 and earlier.

```
mysql -h myinstance.c9akciq32.rds-us-east-1.amazonaws.com  
--ssl-ca=[full path]rds-combined-ca-bundle.pem --ssl-verify-server-cert
```

You can require SSL connections for specific users accounts. For example, you can use one of the following statements, depending on your MariaDB version, to require SSL connections on the user account `encrypted_user`.

For MariaDB 10.2 and later, use the following statement.

```
ALTER USER 'encrypted_user'@'%' REQUIRE SSL;
```

For MariaDB 10.1 and earlier, use the following statement.

```
GRANT USAGE ON *.* TO 'encrypted_user'@'%' REQUIRE SSL;
```

For more information on SSL connections with MariaDB, see [SSL Overview](#) in the MariaDB documentation.

Cache Warming

InnoDB (version 10.2 and later) and XtraDB (versions 10.0 and 10.1) cache warming can provide performance gains for your MariaDB DB instance by saving the current state of the buffer pool when the DB instance is shut down, and then reloading the buffer pool from the saved information when the DB instance starts up. This approach bypasses the need for the buffer pool to "warm up" from normal database use and instead preloads the buffer pool with the pages for known common queries. For more information on cache warming, see [Dumping and restoring the buffer pool](#) in the MariaDB documentation.

Cache warming is enabled by default on MariaDB 10.2 and higher DB instances. To enable it, set the `innodb_buffer_pool_dump_at_shutdown` and `innodb_buffer_pool_load_at_startup` parameters to 1 in the parameter group for your DB instance. Changing these parameter values in a parameter group affects all MariaDB DB instances that use that parameter group. To enable cache warming for specific MariaDB DB instances, you might need to create a new parameter group for those DB instances. For information on parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#).

Cache warming primarily provides a performance benefit for DB instances that use standard storage. If you use PIOPS storage, you don't commonly see a significant performance benefit.

Important

If your MariaDB DB instance doesn't shut down normally, such as during a failover, then the buffer pool state isn't saved to disk. In this case, MariaDB loads whatever buffer pool file is available when the DB instance is restarted. No harm is done, but the restored buffer pool might not reflect the most recent state of the buffer pool prior to the restart. To ensure that you have a recent state of the buffer pool available to warm the cache on startup, we recommend that you periodically dump the buffer pool "on demand." You can dump or load the buffer pool on demand.

You can create an event to dump the buffer pool automatically and at a regular interval. For example, the following statement creates an event named `periodic_buffer_pool_dump` that dumps the buffer pool every hour.

```
CREATE EVENT periodic_buffer_pool_dump
  ON SCHEDULE EVERY 1 HOUR
  DO CALL mysql.rds_innodb_buffer_pool_dump_now();
```

For more information, see [Events](#) in the MariaDB documentation.

Dumping and Loading the Buffer Pool on Demand

You can save and load the cache on demand using the following stored procedures:

- To dump the current state of the buffer pool to disk, call the [mysql.rds_innodb_buffer_pool_dump_now \(p. 907\)](#) stored procedure.
- To load the saved state of the buffer pool from disk, call the [mysql.rds_innodb_buffer_pool_load_now \(p. 907\)](#) stored procedure.
- To cancel a load operation in progress, call the [mysql.rds_innodb_buffer_pool_load_abort \(p. 907\)](#) stored procedure.

Database Parameters for MariaDB

By default, a MariaDB DB instance uses a DB parameter group that is specific to a MariaDB database. This parameter group contains some but not all of the parameters contained in the Amazon RDS DB parameter groups for the MySQL database engine. It also contains a number of new, MariaDB-specific parameters. For more information on the parameters available for the Amazon RDS MariaDB DB engine, see [Parameters for MariaDB \(p. 590\)](#).

Common DBA Tasks for MariaDB

Killing sessions or queries, skipping replication errors, working with InnoDB (version 10.2 and later) and XtraDB (versions 10.0 and 10.1) tablespaces to improve crash recovery times, and managing the global status history are common DBA tasks you might perform in a MariaDB DB instance. You can handle these tasks just as in an Amazon RDS MySQL DB instance, as described in [Common DBA Tasks for MySQL DB Instances \(p. 881\)](#). The crash recovery instructions there refer to the MySQL InnoDB engine, but they are applicable to a MariaDB instance running InnoDB or XtraDB as well.

Local Time Zone for MariaDB DB Instances

By default, the time zone for an RDS MariaDB DB instance is Universal Time Coordinated (UTC). You can set the time zone for your DB instance to the local time zone for your application instead.

To set the local time zone for a DB instance, set the `time_zone` parameter in the parameter group for your DB instance to one of the supported values listed later in this section. When you set the `time_zone` parameter for a parameter group, all DB instances and Read Replicas that are using that parameter group change to use the new local time zone. For information on setting parameters in a parameter group, see [Working with DB Parameter Groups \(p. 211\)](#).

After you set the local time zone, all new connections to the database reflect the change. If you have any open connections to your database when you change the local time zone, you won't see the local time zone update until after you close the connection and open a new connection.

You can set a different local time zone for a DB instance and one or more of its Read Replicas. To do this, use a different parameter group for the DB instance and the replica or replicas and set the `time_zone` parameter in each parameter group to a different local time zone.

If you are replicating across regions, then the replication master DB instance and the Read Replica use different parameter groups (parameter groups are unique to a region). To use the same local time zone for each instance, you must set the `time_zone` parameter in the instance's and Read Replica's parameter groups.

When you restore a DB instance from a DB snapshot, the local time zone is set to UTC. You can update the time zone to your local time zone after the restore is complete. If you restore a DB instance to a point in time, then the local time zone for the restored DB instance is the time zone setting from the parameter group of the restored DB instance.

You can set your local time zone to one of the following values.

Africa/Cairo	Asia/Bangkok	Australia/Darwin
Africa/Casablanca	Asia/Beirut	Australia/Hobart
Africa/Harare	Asia/Calcutta	Australia/Perth
Africa/Monrovia	Asia/Damascus	Australia/Sydney
Africa/Nairobi	Asia/Dhaka	Brazil/East
Africa/Tripoli	Asia/Irkutsk	Canada/Newfoundland
Africa/Windhoek	Asia/Jerusalem	Canada/Saskatchewan
America/Araguaina	Asia/Kabul	Europe/Amsterdam
America/Asuncion	Asia/Karachi	Europe/Athens
America/Bogota	Asia/Kathmandu	Europe/Dublin
America/Caracas	Asia/Krasnoyarsk	Europe/Helsinki
America/Chihuahua	Asia/Magadan	Europe/Istanbul
America/Cuiaba	Asia/Muscat	Europe/Kaliningrad
America/Denver	Asia/Novosibirsk	Europe/Moscow
America/Fortaleza	Asia/Riyadh	Europe/Paris
America/Guatemala	Asia/Seoul	Europe/Prague
America/Halifax	Asia/Shanghai	Europe/Sarajevo
America/Manaus	Asia/Singapore	Pacific/Auckland
America/Matamoros	Asia/Taipei	Pacific/Fiji
America/Monterrey	Asia/Tehran	Pacific/Guam
America/Montevideo	Asia/Tokyo	Pacific/Honolulu
America/Phoenix	Asia/Ulaanbaatar	Pacific/Samoa
America/Santiago	Asia/Vladivostok	US/Alaska

America/Tijuana	Asia/Yakutsk	US/Central
Asia/Amman	Asia/Yerevan	US/Eastern
Asia/Ashgabat	Atlantic/Azores	US/East-Indiana
Asia/Baghdad	Australia/Adelaide	US/Pacific
Asia/Baku	Australia/Brisbane	UTC

Creating a DB Instance Running the MariaDB Database Engine

The basic building block of Amazon RDS is the DB instance. The DB instance is where you create your MariaDB databases.

Important

Before you can create or connect to a DB instance, you must complete the tasks in [Setting Up for Amazon RDS \(p. 5\)](#).

For an example that walks you through the process of creating and connecting to a sample DB instance, see [Creating a MariaDB DB Instance and Connecting to a Database on a MariaDB DB Instance \(p. 11\)](#).

Note

If you are using the console, a new console interface is available for database creation. Choose either the **New Console** or the **Original Console** instructions based on the console that you are using. The **New Console** instructions are open by default.

New Console

You can create a DB instance running MariaDB with the AWS Management Console with **Easy Create** enabled or not enabled. With **Easy Create** enabled, you specify only the DB engine type, DB instance size, and DB instance identifier. **Easy Create** uses the default setting for other configuration options. With **Easy Create** not enabled, you specify more configuration options when you create a database, including ones for availability, security, backups, and maintenance.

Note

For this example, **Standard Create** is enabled, and **Easy Create** isn't enabled. For information about creating a MariaDB DB instance with **Easy Create** enabled, see [Creating a MariaDB Instance \(p. 11\)](#).

To create a MariaDB DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the AWS Management Console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
4. Choose **Create database**.
5. In **Choose a database creation method**, choose **Standard Create**.
6. In **Engine options**, choose **MariaDB**.

Create database

Choose a database creation method Info

Standard Create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy Create

Use recommended best-practice configuration options can be changed after the database is created.

Engine options

Engine type Info

Amazon Aurora



MySQL



MariaDB



PostgreSQL



Oracle

ORACLE®

Microsoft SQL



Version Info

MariaDB 10.2.21

7. In **Templates**, choose the template that matches your use case. If you choose **Production**, the following are also chosen in a later step:
 - **Multi-AZ failover option**
 - **Provisioned IOPS storage option**
 - **Enable deletion protection option**

We recommend these features for any production environment.

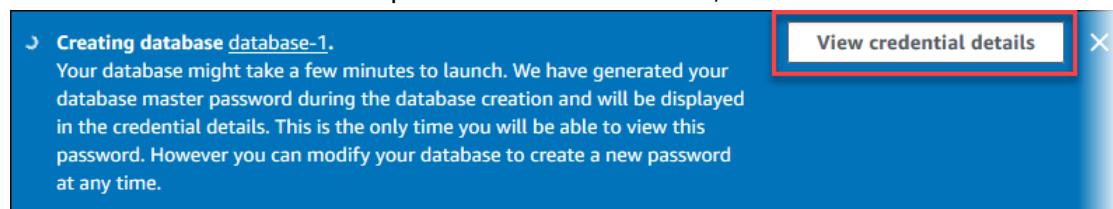
8. To enter your master password, do the following:
 - a. In the **Settings** section, open **Credential Settings**.
 - b. Clear the **Auto generate a password** check box.
 - c. (Optional) Change the **Master username** value and enter the same password in **Master password** and **Confirm password**.

By default, the new DB instance uses an automatically generated password for the master user.

9. For the remaining sections, specify your DB instance settings. For information about each setting, see [Settings for MariaDB DB Instances \(p. 551\)](#).
10. Choose **Create database**.

If you chose to use an automatically generated password, the **View credential details** button appears on the **Databases** page.

To view the master username and password for the DB instance, choose **View credential details**.



To connect to the DB instance as the master user, use the user name and password that appear.

Important

You can't view the master user password again. If you don't record it, you might have to change it. If you need to change the master user password after the DB instance is available, you can modify the DB instance to do so. For more information about modifying a DB instance, see [Modifying a DB Instance Running the MariaDB Database Engine \(p. 564\)](#).

11. In the **Databases** list, choose the name of the new MariaDB DB instance.

On the RDS console, the details for the new DB instance appear. The DB instance has a status of **creating** until the DB instance is created and ready for use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and storage allocated, it could take several minutes for the new instance to be available.

The screenshot shows the 'Summary' page for a database named 'database-1'. The 'Info' section displays the status as 'Creating'. Below the status, there are other details: 'DB identifier' is 'database-1', 'Role' is 'Instance', 'CPU' is listed under 'Current activity', and 'Engine' is 'MariaDB'. At the bottom, there are tabs for 'Connectivity & security', 'Monitoring', 'Logs & events', 'Configuration', and 'Maintenance & ...'. The 'Connectivity & security' tab is currently selected.

Original Console

To launch a MariaDB DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the top right corner of the AWS Management Console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
If the navigation pane is closed, choose the menu icon at the top left to open it.
4. Choose **Create database** to open the **Select engine** page.

Select engine

Engine options

Amazon Aurora
Amazon Aurora

MySQL


MariaDB


PostgreSQL


Oracle


Microsoft SQL Server


MariaDB
MariaDB Community Edition is a MySQL-compatible database with strong support from the open source community, and extra features and performance optimizations.

- Supports database size up to 16 TB.
- Instances offer up to 32 vCPUs and 244 GiB Memory.
- Supports automated backup and point-in-time recovery.
- Supports cross-region read replicas.
- Supports global transaction ID (GTID) and thread pooling.
- Developed and supported by the MariaDB open source community.

Only enable options eligible for RDS Free Usage Tier [info](#)

[Cancel](#) **Next**

5. Choose **MariaDB**, and then choose **Next**.
6. The **Choose use case** page asks if you are planning to use the DB instance you are creating for production. If you are, choose **Production - MariaDB**. If you choose **Production - MariaDB**, the following are also chosen in a later step:
 - **Multi-AZ failover** option
 - **Provisioned IOPS** storage option
 - **Enable deletion protection** option

We recommend these features for any production environment.

7. Choose **Next** to continue. The **Specify DB details** page appears.

On the **Specify DB details** page, specify your DB instance information. For information about each setting, see [Settings for MariaDB DB Instances \(p. 551\)](#).

Specify DB details

Instance specifications

Estimate your monthly costs for the DB Instance using the [AWS Simple Monthly Calculator](#) 

DB engine
MariaDB Community Edition

License model [Info](#)
▼

DB engine version [Info](#)
▼

DB instance class [Info](#)
▼

Multi-AZ deployment [Info](#)

Create replica in different zone
Creates a replica in a different Availability Zone (AZ) to provide data redundancy, eliminate I/O freezes, and minimize latency spikes during system backups.

No

Storage type [Info](#)
▼

8. Choose **Next** to continue.

On the **Configure advanced settings** page, provide additional information that Amazon RDS needs to launch the DB instance. For information about each setting, see [Settings for MySQL DB Instances \(p. 758\)](#).

9. Choose **Create database**.
10. On the final page, choose **View DB instance details**.

On the RDS console, the details for the new DB instance appear. The DB instance has a status of **creating** until the DB instance is created and ready for use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and storage allocated, it could take several minutes for the new instance to be available.

The screenshot shows the AWS RDS 'mariadb-instance1' summary page. It includes a 'Summary' section with engine details (MariaDB 10.1.26), DB instance class (db.t2.small), and status (creating). Below this is a 'CloudWatch (17)' section with a legend for 'mariadb-instance1' and a search bar.

AWS CLI

To create a MariaDB DB instance by using the AWS CLI, call the `create-db-instance` command with the parameters below. For information about each setting, see [Settings for MariaDB DB Instances \(p. 551\)](#).

- `--db-instance-identifier`
- `--db-instance-class`
- `--db-security-groups`
- `--db-subnet-group`
- `--engine`
- `--master-user-name`
- `--master-user-password`
- `--allocated-storage`
- `--backup-retention-period`

Note

If you require a specific minor version of MariaDB, include the `--engine-version` parameter.

Example

The following command creates a MariaDB instance named *mydbinstance*.

For Linux, OS X, or Unix:

```
aws rds create-db-instance \
--db-instance-identifier mydbinstance \
--db-instance-class db.m4.xlarge \
--engine mariadb \
--allocated-storage 20 \
--master-username masteruser \
--master-user-password masteruserpassword \
--backup-retention-period 3
```

For Windows:

```
aws rds create-db-instance ^
--db-instance-identifier mydbinstance ^
--db-instance-class db.m4.xlarge ^
--engine mariadb ^
--allocated-storage 20 ^
--master-username masteruser ^
--master-user-password masteruserpassword ^
--backup-retention-period 3
```

This command should produce output that begins with information that is similar to the following:

```
DBINSTANCE 20 True 3 rds-ca-2015 False arn:aws:rds:us-east-1:1234567890:db:mydbinstance
db.m4.xlarge mydbinstance creating 0 **** mariadb 10.1.26
```

RDS API

To create a MariaDB DB instance by using the Amazon RDS API, call the [CreateDBInstance](#) operation with the parameters below. For information about each setting, see [Settings for MariaDB DB Instances \(p. 551\)](#).

- AllocatedStorage
- BackupRetentionPeriod
- DBInstanceClass
- DBInstanceIdentifier
- DBSecurityGroups
- DBSubnetGroup
- Engine
- MasterUsername
- MasterUserPassword

Note

If you require a specific minor version of MariaDB, include the EngineVersion parameter.

Example

```
https://rds.us-west-2.amazonaws.com/
?Action=CreateDBInstance
&AllocatedStorage=20
&BackupRetentionPeriod=3
&DBInstanceClass=db.m4.xlarge
&DBInstanceIdentifier=mydbinstance
&DBName=mydatabase
&DBSecurityGroups.member.1=mysecuritygroup
&DBSubnetGroup=mydbsubnetgroup
&Engine=mariadb
&MasterUserPassword=masteruserpassword
&MasterUsername=masterawsuser
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140213/us-west-2/rds/aws4_request
&X-Amz-Date=20140213T162136Z
```

```
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=8052a76dfb18469393c5f0182cdab0ebc224a9c7c5c949155376c1c250fc7ec3
```

Settings for MariaDB DB Instances

The following table contains details about settings that you choose when you create a Maria DB instance.

Setting	Setting Description
Allocated storage	The amount of storage to allocate for your DB instance (in gigabytes). In some cases, allocating a higher amount of storage for your DB instance than the size of your database can improve I/O performance. For more information, see Amazon RDS DB Instance Storage (p. 135) .
Auto minor version upgrade	Enable auto minor version upgrade to enable your DB instance to receive preferred minor DB engine version upgrades automatically when they become available. Amazon RDS performs automatic minor version upgrades in the maintenance window.
Availability zone	The Availability Zone for your DB instance. Use the default value of No Preference unless you want to specify an Availability Zone. For more information, see Regions and Availability Zones (p. 132) .
Backup retention period	The number of days that you want automatic backups of your DB instance to be kept. For any nontrivial DB instance, set this value to 1 or greater. For more information, see Working With Backups (p. 254) .
Backup window	The time period during which Amazon RDS automatically takes a backup of your DB instance. Unless you have a specific time that you want to have your database backed up, use the default of No Preference . For more information, see Working With Backups (p. 254) .
Copy tags to snapshots	This option copies any DB instance tags to a DB snapshot when you create a snapshot. For more information, see Tagging Amazon RDS Resources (p. 177) .
Database name	The name for the database on your DB instance. The name must contain 1 to 64 alphanumeric characters. If you don't provide a name, Amazon RDS doesn't create a database on the DB instance you are creating. To create additional databases on your DB instance, connect to your DB instance and use the SQL command CREATE DATABASE. For more information, see Connecting to a DB Instance Running the MariaDB Database Engine (p. 555) .

Setting	Setting Description
Database port	<p>The port that you want to access the DB instance through. MariaDB installations default to port 3306. If you use a DB security group with your DB instance, this port value must be the same one that you provided when creating the DB security group.</p> <p>The firewalls at some companies block connections to the default MariaDB port. If your company firewall blocks the default port, enter another port for your DB instance.</p>
DB engine version	The version of MariaDB that you want to use.
DB instance class	<p>The configuration for your DB instance.</p> <p>If possible, choose an instance class large enough that a typical query working set can be held in memory. When working sets are held in memory, the system can avoid writing to disk, and this improves performance.</p> <p>For more information, see Choosing the DB Instance Class (p. 107).</p>
DB instance identifier	The name for your DB instance. Your DB instance identifier can contain up to 63 alphanumeric characters, and must be unique for your account in the AWS Region you chose. You can add some intelligence to the name, such as including the AWS Region you chose, for example <code>mariadb-instance1</code> .
DB parameter group	<p>A parameter group for your DB instance. You can choose the default parameter group or you can create a custom parameter group.</p> <p>For more information, see Working with DB Parameter Groups (p. 211).</p>
Deletion protection	<p>Enable deletion protection to prevent your DB instance from being deleted. If you create a production DB instance with the AWS Management Console, deletion protection is enabled by default.</p> <p>For more information, see Deleting a DB Instance (p. 174).</p>
Encryption	<p>Enable Encryption to enable encryption at rest for this DB instance.</p> <p>For more information, see Encrypting Amazon RDS Resources (p. 429).</p>
Enhanced Monitoring	<p>Enable enhanced monitoring to enable gathering metrics in real time for the operating system that your DB instance runs on.</p> <p>For more information, see Enhanced Monitoring (p. 307).</p>
License model	MariaDB has only one license model, general-public-license , the general license agreement for MariaDB.

Setting	Setting Description
Log exports	The types of MariaDB database log files to generate. For more information, see MariaDB Database Log Files (p. 394) .
Maintenance window	The 30-minute window in which pending modifications to your DB instance are applied. If the time period doesn't matter, choose No Preference . For more information, see The Amazon RDS Maintenance Window (p. 156) .
Master password	The password for your master user account. The password must contain from 8 to 41 printable ASCII characters (excluding /, ", a space, and @).
Master username	The name that you use as the master user name to log on to your DB instance. For more information, and a list of the default privileges for the master user, see MariaDB Security on Amazon RDS (p. 537) .
Multi-AZ deployment	Create replica in different zone to create a standby mirror of your DB instance in another Availability Zone for failover support. We recommend Multi-AZ for production workloads to maintain high availability. For development and testing, you can choose No . For more information, see High Availability (Multi-AZ) for Amazon RDS (p. 144) .
Option group	An option group for your DB instance. You can choose the default option group or you can create a custom option group. For more information, see Working with Option Groups (p. 196) .
Performance Insights	Enable Performance Insights to monitor your DB instance load so that you can analyze and troubleshoot your database performance. Choose a retention period to determine how much rolling data history to retain. The default of seven days is in the free tier. Long-term retention (two years) is priced per vCPU per month. Choose a master key to use to protect the key used to encrypt this database volume. Choose from the master keys in your account, or enter the key from a different account. For more information, see Using Amazon RDS Performance Insights (p. 319) .

Setting	Setting Description
Publicly accessible	<p>Yes to give the DB instance a public IP address, meaning that it's accessible outside the VPC. To be publicly accessible, the DB instance also has to be in a public subnet in the VPC.</p> <p>No to make the DB instance accessible only from inside the VPC.</p> <p>For more information, see Hiding a DB Instance in a VPC from the Internet (p. 515).</p>
Storage autoscaling	<p>Enable storage autoscaling to enable Amazon RDS to automatically increase storage when needed to avoid having your DB instance run out of storage space.</p> <p>Use Maximum storage threshold to set the upper limit for Amazon RDS to automatically increase storage for your DB instance. The default is 1,000 GiB.</p> <p>For more information, see Managing Capacity Automatically with Amazon RDS Storage Autoscaling (p. 232).</p>
Storage type	<p>The storage type for your DB instance.</p> <p>For more information, see Amazon RDS Storage Types (p. 135).</p>
Subnet group	<p>This setting depends on the platform that you are on. If you are a new customer to AWS, keep the subnet group set to default, which is the default DB subnet group that was created for your account.</p> <p>If you are creating a DB instance on the earlier E2-Classic platform, you might want your DB instance in a specific VPC. In this case, choose the DB subnet group that you created for that VPC.</p>
Virtual Private Cloud (VPC)	<p>This setting depends on the platform that you are on. If you are a new customer to AWS, choose the default VPC shown. If you are creating a DB instance on the earlier E2-Classic platform that doesn't use a VPC, choose Not in VPC.</p> <p>For more information, see Amazon Virtual Private Cloud VPCs and Amazon RDS (p. 505).</p>
VPC security groups	<p>If you are a new customer to AWS, Create new to create a new VPC security group. Otherwise, Choose existing, then choose from security groups that you previously created.</p> <p>When you choose Create new in the RDS console, a new security group is created. This new security group has an inbound rule that allows access to the DB instance from the IP address detected in your browser.</p> <p>For more information, see Working with DB Security Groups (EC2-Classic Platform) (p. 491).</p>

Connecting to a DB Instance Running the MariaDB Database Engine

After Amazon RDS provisions your DB instance, you can use any standard MariaDB client application or utility to connect to the instance. In the connection string, you specify the DNS address from the DB instance endpoint as the host parameter, and specify the port number from the DB instance endpoint as the port parameter.

You can use the AWS Management Console, the AWS CLI [describe-db-instances](#) command, or the Amazon RDS API [DescribeDBInstances](#) operation to list the details of an Amazon RDS DB instance, including its endpoint.

To find the endpoint for a MariaDB instance in the AWS Management Console

1. Open the RDS console and then choose **Databases** to display a list of your DB instances.
2. Choose the MariaDB DB instance name to display its details.
3. On the **Connectivity & security** tab, copy the endpoint. Also, note the port number. You need both the endpoint and the port number to connect to the DB instance.

The screenshot shows the 'Summary' page for a database instance named 'database-1'. The 'Connectivity & security' tab is active. The 'Endpoint' field, which contains the value 'database-1. [REDACTED].us-east-2.rds.amazonaws.com', is circled in red.

DB identifier	CPU
database-1	[Redacted]

Role	Curr
Instance	[Redacted]

Connectivity & security

Endpoint & port

Endpoint
database-1. [REDACTED].us-east-2.rds.amazonaws.com

Port
3306

If an endpoint value is `mariadb-instance1.123456789012.us-east-1.rds.amazonaws.com:3306`, then you specify the following values in a MariaDB connection string:

- For host or host name, specify `mariadb-instance1.123456789012.us-east-1.rds.amazonaws.com`
- For port, specify 3306

You can connect to an Amazon RDS MariaDB DB instance by using tools like the `mysql` command line utility. For more information on using the `mysql` utility, go to [mysql Command-line Client](#) in the MariaDB documentation. One GUI-based application you can use to connect is HeidiSQL. For more information, see the [Download HeidiSQL](#) page.

Two common causes of connection failures to a new DB instance are the following:

- The DB instance was created using a security group that doesn't authorize connections from the device or Amazon EC2 instance where the MariaDB application or utility is running. If the DB instance was created in an Amazon VPC, it must have a VPC security group that authorizes the connections. For more information, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).

You can add or edit an inbound rule in the security group. For **Source**, choose **My IP**. This allows access to the DB instance from the IP address detected in your browser.

If the DB instance was created outside of a VPC, it must have a DB security group that authorizes the connections.

- The DB instance was created using the default port of 3306, and your company has firewall rules blocking connections to that port from devices in your company network. To fix this failure, recreate the instance with a different port.

You can use SSL encryption on connections to an Amazon RDS MariaDB DB instance. For information, see [Using SSL with a MariaDB DB Instance \(p. 539\)](#).

Connecting from the mysql Utility

To connect to a DB instance using the mysql utility, type the following command at a command prompt on a client computer to connect to a database on a MariaDB DB instance. Substitute the DNS name (endpoint) for your DB instance for `<endpoint>`, the master user name you used for `<mymasteruser>`, and provide the master password you used when prompted for a password.

```
mysql -h <endpoint> -P 3306 -u <mymasteruser> -p
```

After you enter the password for the user, you see output similar to the following.

```
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 272
Server version: 5.5.5-10.0.17-MariaDB-log MariaDB Server

Copyright (c) 2000, 2015, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql >
```

Connecting with SSL

Amazon RDS creates an SSL certificate for your DB instance when the instance is created. If you enable SSL certificate verification, then the SSL certificate includes the DB instance endpoint as the Common Name (CN) for the SSL certificate to guard against spoofing attacks. To connect to your DB instance using SSL, follow these steps:

To connect to a DB instance with SSL using the mysql utility

1. Download a root certificate that works for all AWS Regions.

For information about downloading certificates, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).

2. Enter the following command at a command prompt to connect to a DB instance with SSL using the mysql utility. For the -h parameter, substitute the DNS name for your DB instance. For the --ssl-ca parameter, substitute the SSL certificate file name as appropriate.

```
mysql -h mariadb-instance1.123456789012.us-east-1.rds.amazonaws.com --ssl-ca=rds-ca-2015-root.pem -p
```

3. Include the --ssl-verify-server-cert parameter so that the SSL connection verifies the DB instance endpoint against the endpoint in the SSL certificate. For example:

```
mysql -h mariadb-instance1.123456789012.us-east-1.rds.amazonaws.com --ssl-ca=rds-ca-2015-root.pem --ssl-verify-server-cert -p
```

4. Enter the master user password when prompted.

You should see output similar to the following.

```
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 272
Server version: 5.5.5-10.0.17-MariaDB-log MariaDB Server

Copyright (c) 2000, 2015, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql >
```

Maximum MariaDB Connections

The maximum number of connections allowed to an Amazon RDS MariaDB DB instance is based on the amount of memory available for the DB instance class of the DB instance. A DB instance class with more memory available results in a larger number of connections available. For more information on DB instance classes, see [Choosing the DB Instance Class \(p. 107\)](#).

The connection limit for a DB instance is set by default to the maximum for the DB instance class for the DB instance. You can limit the number of concurrent connections to any value up to the maximum number of connections allowed using the `max_connections` parameter in the parameter group for the DB instance. For more information, see [Working with DB Parameter Groups \(p. 211\)](#).

You can retrieve the maximum number of connections allowed for an Amazon RDS MariaDB DB instance by executing the following query on your DB instance:

```
SELECT @@max_connections;
```

You can retrieve the number of active connections to an Amazon RDS MariaDB DB instance by executing the following query on your DB instance:

```
SHOW STATUS WHERE `variable_name` = 'Threads_connected';
```

Related Topics

- [Amazon RDS DB Instances \(p. 105\)](#)

- [Creating a DB Instance Running the MariaDB Database Engine \(p. 543\)](#)
- [Controlling Access with Security Groups \(p. 486\)](#)
- [Deleting a DB Instance \(p. 174\)](#)

Updating Applications to Connect to MariaDB DB Instances Using New SSL/TLS Certificates

As of September 19, 2019, Amazon RDS has published new Certificate Authority (CA) certificates for connecting to your RDS DB instances using Secure Socket Layer or Transport Layer Security (SSL/TLS). The previous CA certificates expire on March 5, 2020. Following, you can find information about updating your applications to use the new certificates. If your application connects to an RDS DB instance using SSL/TLS, you must take the following steps before **March 5, 2020**. Doing this means you can avoid interruption of connectivity between your applications and your RDS DB instances.

This topic can help you to determine whether your applications require certificate verification to connect to your DB instances.

Note

Some applications are configured to connect to MariaDB only if they can successfully verify the certificate on the server.

For such applications, you must update your client application trust stores to include the new CA certificates.

After you update your CA certificates in the client application trust stores, you can rotate the certificates on your DB instances. We strongly recommend testing these procedures in a development or staging environment before implementing them in your production environments.

For more information about certificate rotation, see [Rotating Your SSL/TLS Certificate \(p. 433\)](#). For information about using SSL/TLS with MariaDB DB instances, see [Using SSL with a MariaDB DB Instance \(p. 539\)](#).

Topics

- [Determining Whether a Client Requires Certificate Verification in Order to Connect \(p. 560\)](#)
- [Updating Your Application Trust Store \(p. 561\)](#)
- [Example Java Code for Establishing SSL Connections \(p. 562\)](#)

Determining Whether a Client Requires Certificate Verification in Order to Connect

You can check whether JDBC clients and MySQL clients require certificate verification to connect.

JDBC

The following example with MySQL Connector/J 8.0 shows one way to check an application's JDBC connection properties to determine whether successful connections require a valid certificate. For more information on all of the JDBC connection options for MySQL, see [Configuration Properties](#) in the MySQL documentation.

When using the MySQL Connector/J 8.0, an SSL connection requires verification against the server CA certificate if your connection properties have `sslMode` set to `VERIFY_CA` or `VERIFY_IDENTITY`, as in the following example.

```
Properties properties = new Properties();
properties.setProperty("sslMode", "VERIFY_IDENTITY");
properties.put("user", DB_USER);
properties.put("password", DB_PASSWORD);
```

MySQL

The following examples with the MySQL Client show two ways to check a script's MySQL connection to determine whether successful connections require a valid certificate. For more information on all of the connection options with the MySQL Client, see [Client-Side Configuration for Encrypted Connections](#) in the MySQL documentation.

When using the MySQL 5.7 or MySQL 8.0 Client, an SSL connection requires verification against the server CA certificate if for the `--ssl-mode` option you specify `VERIFY_CA` or `VERIFY_IDENTITY`, as in the following example.

```
mysql -h mysql-database.rds.amazonaws.com -uadmin -ppassword --ssl-ca=/tmp/ssl-cert.pem --ssl-mode=VERIFY_CA
```

When using the MySQL 5.6 Client, an SSL connection requires verification against the server CA certificate if you specify the `--ssl-verify-server-cert` option, as in the following example.

```
mysql -h mysql-database.rds.amazonaws.com -uadmin -ppassword --ssl-ca=/tmp/ssl-cert.pem --ssl-verify-server-cert
```

Updating Your Application Trust Store

For information about updating the trust store for MySQL applications, see [Using TLS/SSL with MariaDB Connector/J](#) in the MariaDB documentation.

Note

When you update the trust store, you can retain older certificates in addition to adding the new certificates.

Updating Your Application Trust Store for JDBC

You can update the trust store for applications that use JDBC for SSL/TLS connections.

To update the trust store for JDBC applications

1. Download the 2019 root certificate that works for all AWS Regions and put the file in the trust store directory.

For information about downloading the root certificate, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).

2. Convert the certificate to .der format using the following command.

```
openssl x509 -outform der -in rds-ca-2019-root.pem -out rds-ca-2019-root.der
```

Replace the file name with the one that you downloaded.

3. Import the certificate into the key store using the following command.

```
keytool -import -alias rds-root -keystore clientkeystore -file rds-ca-2019-root.der
```

4. Confirm that the key store was updated successfully.

```
keytool -list -v -keystore clientkeystore.jks
```

Enter the key store password when you are prompted for it.

Your output should contain the following.

```
rds-root, date, trustedCertEntry,  
Certificate fingerprint (SHA1):  
D4:0D:DB:29:E3:75:0D:FF:A6:71:C3:14:0B:BF:5F:47:8D:1C:80:96  
# This fingerprint should match the output from the below command  
openssl x509 -fingerprint -in rds-ca-2019-root.pem -noout
```

If you are using the MariaDB Connector/J JDBC driver in an application, set the following properties in the application.

```
System.setProperty("javax.net.ssl.trustStore", certs);  
System.setProperty("javax.net.ssl.trustStorePassword", "password");
```

When you start the application, set the following properties.

```
java -Djavax.net.ssl.trustStore=/path_to_truststore/MyTruststore.jks -  
Djavax.net.ssl.trustStorePassword=my_truststore_password com.companyName.MyApplication
```

Example Java Code for Establishing SSL Connections

The following code example shows how to set up the SSL connection using JDBC.

```
private static final String DB_USER = "admin";  
  
private static final String DB_USER = "user name";  
private static final String DB_PASSWORD = "password";  
// This key store has only the prod root ca.  
private static final String KEY_STORE_FILE_PATH = "file-path-to-keystore";  
private static final String KEY_STORE_PASS = "keystore-password";  
  
public static void main(String[] args) throws Exception {  
    Class.forName("org.mariadb.jdbc.Driver");  
  
    System.setProperty("javax.net.ssl.trustStore", KEY_STORE_FILE_PATH);
```

```
System.setProperty("javax.net.ssl.trustStorePassword", KEY_STORE_PASS);

Properties properties = new Properties();
properties.put("user", DB_USER);
properties.put("password", DB_PASSWORD);

Connection connection = DriverManager.getConnection("jdbc:mysql://ssl-mariadb-
public.cn162e2e7kwh.us-east-1.rds.amazonaws.com:3306?useSSL=true",properties);
Statement stmt=connection.createStatement();

ResultSet rs=stmt.executeQuery("SELECT 1 from dual");

return;
}
```

Important

After you have determined that your database connections use SSL/TLS and have updated your application trust store, you can update your database to use the rds-ca-2019 certificates. For instructions, see step 3 in [Updating Your CA Certificate by Modifying Your DB Instance \(p. 434\)](#).

Modifying a DB Instance Running the MariaDB Database Engine

You can change the settings of a DB instance to accomplish tasks such as adding additional storage or changing the DB instance class. This topic guides you through modifying an Amazon RDS MariaDB DB instance, and describes the settings for MariaDB instances.

We recommend that you test any changes on a test instance before modifying a production instance, so that you fully understand the impact of each change. This is especially important when upgrading database versions.

After you modify your DB instance settings, you can apply the changes immediately, or apply them during the next maintenance window for the DB instance. Some modifications cause an interruption by restarting the DB instance.

Console

To modify a MariaDB DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to modify.
3. Choose **Modify**. The **Modify DB instance** page appears.
4. Change any of the settings that you want. For information about each setting, see [Settings for MariaDB DB Instances \(p. 565\)](#).
5. When all the changes are as you want them, choose **Continue** and check the summary of modifications.
6. To apply the changes immediately, choose **Apply immediately**. Choosing this option can cause an outage in some cases. For more information, see [Using the Apply Immediately Parameter \(p. 150\)](#).
7. On the confirmation page, review your changes. If they are correct, choose **Modify DB Instance** to save your changes.

Alternatively, choose **Back** to edit your changes, or choose **Cancel** to cancel your changes.

AWS CLI

To modify a MariaDB DB instance by using the AWS CLI, call the `modify-db-instance` command. Specify the DB instance identifier, and the parameters for the settings that you want to modify. For information about each parameter, see [Settings for MariaDB DB Instances \(p. 565\)](#).

Example

The following code modifies `mydbinstance` by setting the backup retention period to 1 week (7 days). The code enables automatic minor version upgrades by using `--auto-minor-version-upgrade`. To disable automatic minor version upgrades, use `--no-auto-minor-version-upgrade`. The changes are applied during the next maintenance window by using `--no-apply-immediately`. Use `--apply-immediately` to apply the changes immediately. For more information, see [Using the Apply Immediately Parameter \(p. 150\)](#).

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
```

```
--db-instance-identifier mydbinstance \
--backup-retention-period 7 \
--auto-minor-version-upgrade \
--no-apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier mydbinstance ^
--backup-retention-period 7 ^
--auto-minor-version-upgrade ^
--no-apply-immediately
```

RDS API

To modify a MariaDB instance by using the Amazon RDS API, call the [ModifyDBInstance](#) operation. Specify the DB instance identifier, and the parameters for the settings that you want to modify. For information about each parameter, see [Settings for MariaDB DB Instances \(p. 565\)](#).

Example

The following code modifies *mydbinstance* by setting the backup retention period to 1 week (7 days) and enabling automatic minor version upgrades. These changes are applied during the next maintenance window.

```
https://rds.amazonaws.com/
?Action=ModifyDBInstance
&ApplyImmediately=false
&AutoMinorVersionUpgrade=true
&BackupRetentionPeriod=7
&DBInstanceIdentifier=mydbinstance
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20131016/us-west-1/rds/aws4_request
&X-Amz-Date=20131016T233051Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=087a8eb41cb1ab0fc9ec1575f23e73757ffc6a1e42d7d2b30b9cc0be988cff97
```

Settings for MariaDB DB Instances

The following table contains details about which settings you can modify, which settings you can't modify, when the changes can be applied, and whether the changes cause downtime for the DB instance.

Setting	Setting Description	When the Change Occurs	Downtime Notes
Allocated storage	<p>The storage, in gigabytes, that you want to allocate for your DB instance.</p> <p>You can't modify allocated storage if the DB instance status is storage-optimization or if the allocated storage for the DB instance has been modified in the last six hours.</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	No downtime. Performance might be degraded during the change.

Setting	Setting Description	When the Change Occurs	Downtime Notes
	The maximum storage allowed depends on the storage type. For more information, see Amazon RDS DB Instance Storage (p. 135) .		
Auto minor version upgrade	Enable auto minor version upgrade to enable your DB instance to receive preferred minor DB engine version upgrades automatically when they become available. Amazon RDS performs automatic minor version upgrades in the maintenance window.	–	–
Backup retention period	<p>The number of days that automatic backups are retained. To disable automatic backups, set the backup retention period to 0.</p> <p>For more information, see Working With Backups (p. 254).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false and you change the setting from a nonzero value to another nonzero value, the change is applied asynchronously, as soon as possible. Otherwise, the change occurs during the next maintenance window.</p>	An outage occurs if you change from 0 to a nonzero value, or from a nonzero value to 0.
Backup window	<p>The time range during which automated backups of your databases occur. The backup window is a start time in Universal Coordinated Time (UTC), and a duration in hours.</p> <p>For more information, see Working With Backups (p. 254).</p>	The change is applied asynchronously, as soon as possible.	–
Certificate authority	The certificate that you want to use.	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change.
Copy tags to snapshots	<p>If you have any DB instance tags, this option copies them when you create a DB snapshot.</p> <p>For more information, see Tagging Amazon RDS Resources (p. 177).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Database port	<p>The port that you want to use to access the database.</p> <p>The port value must not match any of the port values specified for options in the option group for the DB instance.</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	The DB instance is rebooted immediately.
DB engine version	<p>The version of the MariaDB database engine that you want to use. Before you upgrade your production DB instances, we recommend that you test the upgrade process on a test instance to verify its duration and to validate your applications.</p> <p>For more information, see Upgrading the MariaDB DB Engine (p. 573).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change.
DB instance class	<p>The DB instance class that you want to use.</p> <p>For more information, see Choosing the DB Instance Class (p. 107).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change.
DB instance identifier	<p>The DB instance identifier. This value is stored as a lowercase string.</p> <p>For more information about the effects of renaming a DB instance, see Renaming a DB Instance (p. 163).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change. The DB instance is rebooted.
DB parameter group	<p>The parameter group that you want associated with the DB instance.</p> <p>For more information, see Working with DB Parameter Groups (p. 211).</p>	The parameter group change occurs immediately.	<p>An outage doesn't occur during this change. When you change the parameter group, changes to some parameters are applied to the DB instance immediately without a reboot. Changes to other parameters are applied only after the DB instance is rebooted.</p> <p>For more information, see Rebooting a DB Instance (p. 166).</p>

Setting	Setting Description	When the Change Occurs	Downtime Notes
Deletion protection	<p>Enable deletion protection to prevent your DB instance from being deleted.</p> <p>For more information, see Deleting a DB Instance (p. 174).</p>	–	–
Enhanced Monitoring	<p>Enable enhanced monitoring to enable gathering metrics in real time for the operating system that your DB instance runs on.</p> <p>For more information, see Enhanced Monitoring (p. 307).</p>	–	–
Log exports	<p>The types of MariaDB database log files to publish to Amazon CloudWatch Logs.</p> <p>For more information, see MariaDB Database Log Files (p. 394).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–
Maintenance window	<p>The time range during which system maintenance occurs. System maintenance includes upgrades, if applicable. The maintenance window is a start time in Universal Coordinated Time (UTC), and a duration in hours.</p> <p>If you set the window to the current time, there must be at least 30 minutes between the current time and end of the window to ensure any pending changes are applied.</p> <p>For more information, see The Amazon RDS Maintenance Window (p. 156).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	If there are one or more pending actions that cause an outage, and the maintenance window is changed to include the current time, then those pending actions are applied immediately, and an outage occurs.
Multi-AZ deployment	<p>Yes to deploy your DB instance in multiple Availability Zones. Otherwise, No.</p> <p>For more information, see Regions and Availability Zones (p. 132).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	–
New master password	The password for your master user. The password must contain from 8 to 41 alphanumeric characters.	The change is applied asynchronously, as soon as possible. This setting ignores the Apply immediately setting.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Option group	<p>The option group that you want associated with the DB instance.</p> <p>For more information, see Working with Option Groups (p. 196).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	–
Performance Insights	<p>Enable Performance Insights to monitor your DB instance load so that you can analyze and troubleshoot your database performance.</p> <p>Choose a retention period to determine how much rolling data history to retain. The default of seven days is in the free tier. Long-term retention (two years) is priced per vCPU per month.</p> <p>You can't change the master key after the database is created.</p> <p>For more information, see Using Amazon RDS Performance Insights (p. 319).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–
Public accessibility	<p>Yes to give the DB instance a public IP address, meaning that it's accessible outside the VPC. To be publicly accessible, the DB instance also has to be in a public subnet in the VPC.</p> <p>No to make the DB instance accessible only from inside the VPC.</p> <p>For more information, see Hiding a DB Instance in a VPC from the Internet (p. 515).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–
Security group	<p>The security group that you want associated with the DB instance.</p> <p>For more information, see Working with DB Security Groups (EC2-Classic Platform) (p. 491).</p>	The change is applied asynchronously, as soon as possible. This setting ignores the Apply immediately setting.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Storage autoscaling	<p>Enable storage autoscaling to enable Amazon RDS to automatically increase storage when needed to avoid having your DB instance run out of storage space.</p> <p>Use Maximum storage threshold to set the upper limit for Amazon RDS to automatically increase storage for your DB instance. The default is 1,000 GiB.</p> <p>For more information, see Managing Capacity Automatically with Amazon RDS Storage Autoscaling (p. 232).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Storage type	<p>The storage type that you want to use.</p> <p>For more information, see Amazon RDS Storage Types (p. 135).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	<p>The following changes all result in a brief outage while the process starts. After that, you can use your database normally while the change takes place.</p> <ul style="list-style-type: none"> • From General Purpose (SSD) to Magnetic. • From General Purpose (SSD) to Provisioned IOPS (SSD). The outage only happens if the DB instance is single-AZ and you are using a custom parameter group. There is no outage for a multi-AZ DB instance. • From Magnetic to General Purpose (SSD). • From Magnetic to Provisioned IOPS (SSD). • From Provisioned IOPS (SSD) to Magnetic. • From Provisioned IOPS (SSD) to General Purpose (SSD). The outage only happens if the DB instance is single-AZ and you are using a custom parameter group. There is no outage for a multi-AZ DB instance.

Setting	Setting Description	When the Change Occurs	Downtime Notes
Subnet group	<p>The subnet group for the DB instance. You can use this setting to move your DB instance to a different VPC. If your DB instance isn't in a VPC, you can use this setting to move your DB instance into a VPC.</p> <p>For more information, see Moving a DB Instance Not in a VPC into a VPC (p. 520).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change. The DB instance is rebooted.

Upgrading the MariaDB DB Engine

When Amazon RDS supports a new version of a database engine, you can upgrade your DB instances to the new version. There are two kinds of upgrades: major version upgrades and minor version upgrades. You must modify the DB instance manually to perform a major version upgrade. Minor version upgrades occur automatically if you enable auto minor version upgrades on your DB instance. In all other cases, you must modify the DB instance manually to perform a minor version upgrade.

For more information about MariaDB supported versions and version management, see [MariaDB on Amazon RDS Versions \(p. 534\)](#).

Topics

- [Overview of Upgrading \(p. 573\)](#)
- [Upgrading a MariaDB DB Instance \(p. 574\)](#)

Overview of Upgrading

Major version upgrades can contain database changes that are not backward-compatible with existing applications. As a result, Amazon RDS doesn't apply major version upgrades automatically; you must manually modify your DB instance. You should thoroughly test any upgrade before applying it to your production instances.

Unless you specify otherwise, your DB instance will automatically be upgraded to new MariaDB minor versions as they are supported by Amazon RDS. This patching occurs during your scheduled maintenance window. You can modify a DB instance to turn off automatic minor version upgrades.

Amazon RDS takes two DB snapshots during the upgrade process. The first DB snapshot is of the DB instance before any upgrade changes have been made. If the upgrade doesn't work for your databases, you can restore this snapshot to create a DB instance running the old version. The second DB snapshot is taken when the upgrade completes.

Note

Amazon RDS only takes DB snapshots if you have set the backup retention period for your DB instance to a number greater than 0. To change your backup retention period, see [Modifying a DB Instance Running the MariaDB Database Engine \(p. 564\)](#).

After the upgrade is complete, you can't revert to the previous version of the database engine. If you want to return to the previous version, restore the first DB snapshot taken to create a new DB instance.

You control when to upgrade your DB instance to a new version supported by Amazon RDS. This level of control helps you maintain compatibility with specific database versions and test new versions with your application before deploying in production. When you are ready, you can perform version upgrades at the times that best fit your schedule.

If your DB instance is using read replication, you must upgrade all of the Read Replicas before upgrading the source instance.

If your DB instance is in a Multi-AZ deployment, both the primary and standby DB instances are upgraded. The primary and standby DB instances are upgraded at the same time and you will experience an outage until the upgrade is complete. The time for the outage varies based on your database engine, engine version, and the size of your DB instance.

If you are using a custom parameter group, and you perform a major version upgrade, you must specify either a default parameter group for the new DB engine version or create your own custom parameter group for the new DB engine version. Associating the new parameter group with the DB instance requires a customer-initiated database reboot after the upgrade completes. The instance's parameter group

status will show pending-reboot if the instance needs to be rebooted to apply the parameter group changes. An instance's parameter group status can be viewed in the AWS console or by using a "describe" call such as describe-db-instances.

Upgrading a MariaDB DB Instance

For information about manually or automatically upgrading a MariaDB DB instance, see [Upgrading a DB Instance Engine Version \(p. 159\)](#).

Migrating Data from a MySQL DB Snapshot to a MariaDB DB Instance

You can migrate an Amazon RDS MySQL DB snapshot to a new DB instance running MariaDB 10.1 using the AWS Management Console, AWS CLI, or Amazon RDS API. You must create the DB snapshot from an Amazon RDS DB instance running MySQL 5.6. To learn how to create an RDS MySQL DB snapshot, see [Creating a DB Snapshot \(p. 264\)](#).

After you migrate from MySQL to MariaDB, the MariaDB DB instance will be associated with the default DB parameter group and option group. After you restore the DB snapshot, you can associate a custom DB parameter group for the new DB instance. However, a MariaDB parameter group has a different set of configurable system variables. For information about the differences between MySQL and MariaDB system variables, see [System Variable Differences Between MariaDB 10.1 and MySQL 5.6](#). To learn about DB parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#). To learn about option groups, see [Working with Option Groups \(p. 196\)](#).

Incompatibilities Between MariaDB and MySQL

Incompatibilities between MySQL and MariaDB include the following:

- You can't migrate a DB snapshot created with MySQL 5.5 to MariaDB 10.1.
- You can't migrate a DB snapshot created with MySQL 5.7 to MariaDB.
- You can't migrate a DB snapshot created with MySQL 8.0 to MariaDB.
- You can't migrate an encrypted snapshot.
- If the source MySQL database uses a SHA256 password hash, you need to reset user passwords that are SHA256 hashed before you can connect to the MariaDB database. The following code shows how to reset a password that is SHA256 hashed:

```
SET old_passwords = 0;
UPDATE mysql.user SET plugin = 'mysql_native_password',
Password = PASSWORD('new_password')
WHERE (User, Host) = ('master_user_name', %);
FLUSH PRIVILEGES;
```

- If your RDS master user account uses the SHA-256 password hash, the password has to be reset using the RDS `modify-db-instance` AWS CLI command, `ModifyDBInstance` API operation, or the AWS Management Console. For information about modifying a MariaDB DB instance, see [Modifying a DB Instance Running the MariaDB Database Engine \(p. 564\)](#).
- MariaDB doesn't support the Memcached plugin; however, the data used by the Memcached plugin is stored as InnoDB tables. After you migrate a MySQL DB snapshot, you can access the data used by the Memcached plugin using SQL. For more information about the innodb_memcache database, see [InnoDB memcached Plugin Internals](#).

Console

To migrate a MySQL DB snapshot to a MariaDB DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.

2. In the navigation pane, choose **Snapshots**, and then select the MySQL DB snapshot you want to migrate.
3. For **Actions**, choose **Migrate Snapshot**. The **Migrate Database** page appears.
4. For **Migrate to DB Engine**, choose **mariadb**.
5. On the **Migrate Database** page, provide additional information that RDS needs to launch the MariaDB DB instance.
 - **DB Engine Version:** Choose the version of the MariaDB database engine that you want to use. For more information, see [Upgrading the MariaDB DB Engine \(p. 573\)](#).
 - **DB Instance Class:** Choose a DB instance class that has the required storage and capacity for your database, for example db.r3.large. For any production application that requires fast and consistent I/O performance, we recommend Provisioned IOPS storage. For more information, see [Provisioned IOPS SSD Storage \(p. 137\)](#). MariaDB 10.1 does not support previous-generation DB instance classes. For more information, see [Choosing the DB Instance Class \(p. 107\)](#).
 - **Multi-AZ Deployment:** Choose **Yes** to deploy your DB instance in multiple Availability Zones; otherwise, **No**. For more information, see [Regions and Availability Zones \(p. 132\)](#).
 - **DB Snapshot ID:** Type a name for the DB snapshot identifier.

The DB snapshot identifier has the following constraints:

- It must contain from 1 to 255 alphanumeric characters or hyphens.
- The character must be a letter.
- It cannot end with a hyphen or contain two consecutive hyphens.

If you are restoring from a shared manual DB snapshot, the DB snapshot identifier must be the Amazon Resource Name (ARN) of the shared DB snapshot.

- **DB Instance Identifier:** Type a name for the DB instance that is unique for your account in the AWS Region where the DB instance will reside. This identifier is used in the endpoint addresses for the instances in your DB instance.

The DB instance identifier has the following constraints:

- It must contain from 1 to 63 alphanumeric characters or hyphens.
- Its first character must be a letter.
- It cannot end with a hyphen or contain two consecutive hyphens.
- It must be unique for all DB instances for your AWS account, within an AWS Region.
- **Virtual Private Cloud (VPC):** If you have an existing VPC, then you can use that VPC with your MariaDB DB instance by selecting your VPC identifier, for example vpc-a464d1c1. For more information about VPC, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).

Otherwise, you can choose to have Amazon RDS create a VPC for you by selecting Create a new VPC.

You cannot create MariaDB instances in the EC2 Classic Network.

- **Subnet group:** If you have an existing subnet group, then you can use that subnet group with your MariaDB DB instance by selecting your subnet group identifier, for example gs-subnet-group1.

Otherwise, you can choose to have Amazon RDS create a subnet group for you by selecting Create a new subnet group.

- **Public accessibility:** Choose **No** to specify that instances in your DB instance can only be accessed by resources inside your VPC. Choose **Yes** to specify that instances in your DB instance can be accessed by resources on the public network. The default is **Yes**.
- **Availability zone:** Choose the **Availability Zone** to host the primary instance for your MariaDB DB instance. To have Amazon RDS choose an **Availability Zone** for you, choose **No Preference**.

- **Database Port:** Type the default port to be used when connecting to instances in the DB instance. The default is 3306.

You might be behind a corporate firewall that doesn't allow access to default ports such as the MySQL default port 3306. In this case, provide a port value that your corporate firewall allows.

- **Option Group:** Choose the option group that you want associated with the DB instance. For more information, see [Working with Option Groups \(p. 196\)](#).
- **Encryption:** Choose **Enable Encryption** for your new MariaDB DB instance to be encrypted "at rest." If you choose **Enable Encryption**, you will be required to choose an AWS KMS encryption key as the **Master Key** value.
- **Auto minor version upgrade:** Choose **Enable auto minor version upgrade** to enable your DB instance to receive preferred minor DB engine version upgrades automatically when they become available. Amazon RDS performs automatic minor version upgrades in the maintenance window. The **Auto minor version upgrade** option only applies to upgrades to MySQL minor engine versions for your MariaDB DB instance. It doesn't apply to regular patches applied to maintain system stability.

Migrate Database

Migrate this database to a new DB Engine by selecting your desired options for the migrated instance.

Instance specifications

Migrate to DB Engine
Name of the Database Engine
mariadb

DB Engine Version
Version Number of the Database Engine to be used for this instance
10.1.14 (default)

DB Instance Class
Contains the compute and memory capacity of the DB Instance.
db.m4.xlarge — 4 vCPU, 16 GiB RAM

Multi-AZ Deployment
Specifies if the DB Instance should have a standby deployed in another Availability Zone.

6. Choose **Migrate**.

AWS CLI

To migrate data from a MySQL DB snapshot to a MariaDB DB instance, use the AWS CLI [restore-db-instance-from-db-snapshot](#) command with the following parameters:

- **--db-instance-identifier** – Name of the DB instance to create from the DB snapshot.

- **--db-snapshot-identifier** – The identifier for the DB snapshot to restore from.
- **--engine** – The database engine to use for the new instance.

Example

For Linux, OS X, or Unix:

```
aws rds restore-db-instance-from-db-snapshot \
--db-instance-identifier newmariadbinstance \
--db-snapshot-identifier mysqlsnapshot \
--engine mariadb
```

For Windows:

```
aws rds restore-db-instance-from-db-snapshot \
--db-instance-identifier newmariadbinstance ^
--db-snapshot-identifier mysqlsnapshot ^
--engine mariadb
```

API

To migrate data from a MySQL DB snapshot to a MariaDB DB instance, call the Amazon RDS API operation [RestoreDBInstanceFromDBSnapshot](#).

Example

```
https://rds.us-west-2.amazonaws.com/
?Action=RestoreDBInstanceFromDBSnapshot
&DBInstanceIdentifier= newmariadbinstance
&DBSnapshotIdentifier= mysqlsnapshot
&Engine= mariadb
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2013-09-09
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140428/us-west-2/rds/aws4_request
&X-Amz-Date=20140428T232655Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=78ac761e8c8f54a8c0727f4e67ad0a766fbb0024510b9aa34ea6d1f7df52fe92
```

Working with MariaDB Replication

You usually use Read Replicas to configure replication between Amazon RDS DB instances. For general information about Read Replicas, see [Working with Read Replicas \(p. 182\)](#). For specific information about working with Read Replicas on Amazon RDS MariaDB, see [Working with MariaDB Read Replicas \(p. 579\)](#).

You can also configure replication based on binary log coordinates for a MariaDB DB instance. For MariaDB instances, you can also configure replication based on global transaction IDs (GTIDs), which provides better crash safety. For more information, see [Configuring GTID-Based Replication into an Amazon RDS MariaDB DB instance \(p. 582\)](#).

The following are other replication options available with Amazon RDS MariaDB:

- You can set up replication between an Amazon RDS MariaDB DB instance and a MySQL or MariaDB instance that is external to Amazon RDS. For information about configuring replication with an external source, see [Replication with a MySQL or MariaDB Instance Running External to Amazon RDS \(p. 862\)](#).
- You can configure replication to import databases from a MySQL or MariaDB instance that is external to Amazon RDS, or to export databases to such instances. For more information, see [Importing Data to an Amazon RDS MySQL or MariaDB DB Instance with Reduced Downtime \(p. 835\)](#) and [Exporting Data from a MySQL DB Instance by Using Replication \(p. 869\)](#).

For any of these replication options, you can use either row-based replication, statement-based, or mixed replication. Row-based replication only replicates the changed rows that result from a SQL statement. Statement-based replication replicates the entire SQL statement. Mixed replication uses statement-based replication when possible, but switches to row-based replication when SQL statements that are unsafe for statement-based replication are executed. In most cases, mixed replication is recommended. The binary log format of the DB instance determines whether replication is row-based, statement-based, or mixed. For information about setting the binary log format, see [Binary Logging Format \(p. 399\)](#).

Topics

- [Working with MariaDB Read Replicas \(p. 579\)](#)
- [Configuring GTID-Based Replication into an Amazon RDS MariaDB DB instance \(p. 582\)](#)

Working with MariaDB Read Replicas

This section contains specific information about working with Read Replicas on Amazon RDS MariaDB. For general information about Read Replicas and instructions for using them, see [Working with Read Replicas \(p. 182\)](#).

Topics

- [Read Replica Configuration with MariaDB \(p. 580\)](#)
- [Read Replica Updates with MariaDB \(p. 580\)](#)
- [Multi-AZ Read Replica Deployments with MariaDB \(p. 580\)](#)
- [Monitoring MariaDB Read Replicas \(p. 580\)](#)
- [Starting and Stopping Replication with MariaDB Read Replicas \(p. 581\)](#)
- [Deleting Read Replicas with MariaDB \(p. 581\)](#)
- [Troubleshooting a MariaDB Read Replica Problem \(p. 581\)](#)

Read Replica Configuration with MariaDB

Before a MariaDB DB instance can serve as a replication source, you must enable automatic backups on the source DB instance by setting the backup retention period to a value other than 0. This requirement also applies to a Read Replica that is the source DB instance for another Read Replica.

You can create up to five Read Replicas from one DB instance. For replication to operate effectively, each Read Replica should have as the same amount of compute and storage resources as the source DB instance. If you scale the source DB instance, you should also scale the Read Replicas.

If a Read Replica is running any version of MariaDB, you can specify it as the source DB instance for another Read Replica. For example, you can create ReadReplica1 from MyDBInstance, and then create ReadReplica2 from ReadReplica1. Updates made to MyDBInstance are replicated to ReadReplica1 and then replicated from ReadReplica1 to ReadReplica2. You can't have more than four instances involved in a replication chain. For example, you can create ReadReplica1 from MySourceDBInstance, and then create ReadReplica2 from ReadReplica1, and then create ReadReplica3 from ReadReplica2, but you can't create a ReadReplica4 from ReadReplica3.

If you promote a MariaDB Read Replica that is in turn replicating to other Read Replicas, those Read Replicas remain active. Consider an example where MyDBInstance1 replicates to MyDBInstance2, and MyDBInstance2 replicates to MyDBInstance3. If you promote MyDBInstance2, replication from MyDBInstance1 to MyDBInstance2 no longer occurs, but MyDBInstance2 still replicates to MyDBInstance3.

To enable automatic backups on a Read Replica for Amazon RDS MariaDB, first create the Read Replica, then modify the Read Replica to enable automatic backups.

You can run multiple concurrent Read Replica create or delete actions that reference the same source DB instance, as long as you stay within the limit of five Read Replicas for the source instance.

Read Replica Updates with MariaDB

Read Replicas are designed to support read queries, but you might need occasional updates. For example, you might need to add an index to speed the specific types of queries accessing the replica. You can enable updates by setting the `read_only` parameter to **0** in the DB parameter group for the Read Replica.

Multi-AZ Read Replica Deployments with MariaDB

You can create a Read Replica from either single-AZ or Multi-AZ DB instance deployments. You use Multi-AZ deployments to improve the durability and availability of critical data, but you can't use the Multi-AZ secondary to serve read-only queries. Instead, you can create Read Replicas from high-traffic Multi-AZ DB instances to offload read-only queries. If the source instance of a Multi-AZ deployment fails over to the secondary, any associated Read Replicas automatically switch to use the secondary (now primary) as their replication source. For more information, see [High Availability \(Multi-AZ\) for Amazon RDS \(p. 144\)](#).

You can create a Read Replica as a Multi-AZ DB instance. Amazon RDS creates a standby of your replica in another Availability Zone for failover support for the replica. Creating your Read Replica as a Multi-AZ DB instance is independent of whether the source database is a Multi-AZ DB instance.

Monitoring MariaDB Read Replicas

For MariaDB Read Replicas, you can monitor replication lag in Amazon CloudWatch by viewing the Amazon RDS `ReplicaLag` metric. The `ReplicaLag` metric reports the value of the `Seconds_Behind_Master` field of the `SHOW SLAVE STATUS` command.

Common causes for replication lag for MariaDB are the following:

- A network outage.

- Writing to tables with indexes on a Read Replica. If the `read_only` parameter is not set to 0 on the Read Replica, it can break replication.
- Using a nontransactional storage engine such as MyISAM. Replication is only supported for the InnoDB storage engine on MariaDB 10.2 and later and the XtraDB storage engine on MariaDB 10.1 and earlier.

When the `ReplicaLag` metric reaches 0, the replica has caught up to the source DB instance. If the `ReplicaLag` metric returns -1, then replication is currently not active. `ReplicaLag = -1` is equivalent to `Seconds_Behind_Master = NULL`.

Starting and Stopping Replication with MariaDB Read Replicas

You can stop and restart the replication process on an Amazon RDS DB instance by calling the system stored procedures [mysql.rds_stop_replication \(p. 903\)](#) and [mysql.rds_start_replication \(p. 901\)](#). You can do this when replicating between two Amazon RDS instances for long-running operations such as creating large indexes. You also need to stop and start replication when importing or exporting databases. For more information, see [Importing Data to an Amazon RDS MySQL or MariaDB DB Instance with Reduced Downtime \(p. 835\)](#) and [Exporting Data from a MySQL DB Instance by Using Replication \(p. 869\)](#).

If replication is stopped for more than 30 consecutive days, either manually or due to a replication error, Amazon RDS terminates replication between the master DB instance and all Read Replicas. It does so to prevent increased storage requirements on the master DB instance and long failover times. The Read Replica DB instance is still available. However, replication can't be resumed because the binary logs required by the Read Replica are deleted from the master DB instance after replication is terminated. You can create a new Read Replica for the master DB instance to reestablish replication.

Deleting Read Replicas with MariaDB

You must explicitly delete Read Replicas, using the same mechanisms for deleting a DB instance. If you delete the source DB instance without deleting the replicas, each replica is promoted to a standalone DB instance.

Troubleshooting a MariaDB Read Replica Problem

The replication technologies for MariaDB are asynchronous. Because they are asynchronous, occasional `BinLogDiskUsage` increases on the source DB instance and `ReplicaLag` on the Read Replica are to be expected. For example, a high volume of write operations to the source DB instance can occur in parallel. In contrast, write operations to the Read Replica are serialized using a single I/O thread, which can lead to a lag between the source instance and Read Replica. For more information about read-only replicas in the MariaDB documentation, go to [Replication Overview](#).

You can do several things to reduce the lag between updates to a source DB instance and the subsequent updates to the Read Replica, such as the following:

- Sizing a Read Replica to have a storage size and DB instance class comparable to the source DB instance.
- Ensuring that parameter settings in the DB parameter groups used by the source DB instance and the Read Replica are compatible. For more information and an example, see the discussion of the `max_allowed_packet` parameter later in this section.

Amazon RDS monitors the replication status of your Read Replicas and updates the `Replication State` field of the Read Replica instance to `Error` if replication stops for any reason. An example might be if DML queries run on your Read Replica conflict with the updates made on the source DB instance.

You can review the details of the associated error thrown by the MariaDB engine by viewing the `Replication Error` field. Events that indicate the status of the Read Replica are also generated,

including [RDS-EVENT-0045 \(p. 375\)](#), [RDS-EVENT-0046 \(p. 375\)](#), and [RDS-EVENT-0047 \(p. 373\)](#). For more information about events and subscribing to events, see [Using Amazon RDS Event Notification \(p. 370\)](#). If a MariaDB error message is returned, review the error in the [MariaDB error message documentation](#).

One common issue that can cause replication errors is when the value for the `max_allowed_packet` parameter for a Read Replica is less than the `max_allowed_packet` parameter for the source DB instance. The `max_allowed_packet` parameter is a custom parameter that you can set in a DB parameter group that is used to specify the maximum size of DML code that can be executed on the database. In some cases, the `max_allowed_packet` parameter value in the DB parameter group associated with a source DB instance is smaller than the `max_allowed_packet` parameter value in the DB parameter group associated with the source's Read Replica. In these cases, the replication process can throw an error (Packet bigger than '`max_allowed_packet`' bytes) and stop replication. You can fix the error by having the source and Read Replica use DB parameter groups with the same `max_allowed_packet` parameter values.

Other common situations that can cause replication errors include the following:

- Writing to tables on a Read Replica. If you are creating indexes on a Read Replica, you need to have the `read_only` parameter set to **0** to create the indexes. If you are writing to tables on the Read Replica, it might break replication.
- Using a non-transactional storage engine such as MyISAM. Read Replicas require a transactional storage engine. Replication is only supported for the InnoDB storage engine on MariaDB 10.2 and later and the XtraDB storage engine on MariaDB 10.1 and earlier.
- Using unsafe nondeterministic queries such as `SYSDATE()`. For more information, see [Determination of Safe and Unsafe Statements in Binary Logging](#).

If you decide that you can safely skip an error, you can follow the steps described in the section [Skipping the Current Replication Error \(p. 881\)](#). Otherwise, you can delete the Read Replica and create an instance using the same DB instance identifier so that the endpoint remains the same as that of your old Read Replica. If a replication error is fixed, the `Replication State` changes to *replicating*.

For MariaDB DB instances, in some cases Read Replicas can't be switched to the secondary if some binlog events aren't flushed during the failure. In these cases, you must manually delete and recreate the Read Replicas. You can reduce the chance of this happening by setting the following dynamic variable values: `sync_binlog=1`, `innodb_flush_log_at_trx_commit=1`, and `innodb_support_xa=1`. These settings might reduce performance, so test their impact before implementing the changes in a production environment.

Configuring GTID-Based Replication into an Amazon RDS MariaDB DB instance

You can set up GTID-based replication from an external MariaDB instance of version 10.0.24 or greater into an Amazon RDS MariaDB DB instance. Be sure to follow these guidelines when you set up an external replication master and a replica on Amazon RDS:

- Monitor failover events for the Amazon RDS MariaDB DB instance that is your replica. If a failover occurs, then the DB instance that is your replica might be recreated on a new host with a different network address. For information on how to monitor failover events, see [Using Amazon RDS Event Notification \(p. 370\)](#).
- Maintain the binlogs on your master instance until you have verified that they have been applied to the replica. This maintenance ensures that you can restore your master instance in the event of a failure.
- Turn on automated backups on your MariaDB DB instance on Amazon RDS. Turning on automated backups ensures that you can restore your replica to a particular point in time if you need to re-

synchronize your master and replica. For information on backups and Point-In-Time Restore, see [Backing Up and Restoring Amazon RDS DB Instances \(p. 253\)](#).

Note

The permissions required to start replication on an Amazon RDS MariaDB DB instance are restricted and not available to your Amazon RDS master user. Because of this, you must use the Amazon RDS [mysql.rds_set_external_master_gtid \(p. 595\)](#) and [mysql.rds_start_replication \(p. 901\)](#) commands to set up replication between your live database and your Amazon RDS MariaDB database.

To start replication between an external master instance and a MariaDB DB instance on Amazon RDS, use the following procedure.

To Start Replication

1. Make the source MariaDB instance read-only:

```
mysql> FLUSH TABLES WITH READ LOCK;
mysql> SET GLOBAL read_only = ON;
```

2. Get the current GTID of the external MariaDB instance. You can do this by using mysql or the query editor of your choice to run `SELECT @@gtid_current_pos;`.

The GTID is formatted as <domain-id>-<server-id>-<sequence-id>. A typical GTID looks something like **0-1234510749-1728**. For more information about GTIDs and their component parts, see [Global Transaction ID](#) in the MariaDB documentation.

3. Copy the database from the external MariaDB instance to the Amazon RDS MariaDB DB instance using mysqldump. For very large databases, you might want to use the procedure in [Importing Data to an Amazon RDS MySQL or MariaDB DB Instance with Reduced Downtime \(p. 835\)](#).

Note

Make sure there is not a space between the `-p` option and the entered password.

For Linux, OS X, or Unix:

```
mysqldump \
    --databases <database_name> \
    --single-transaction \
    --compress \
    --order-by-primary \
    -u <local_user> \
    -p<local_password> | mysql \
        --host=hostname \
        --port=3306 \
        -u <RDS_user_name> \
        -p <RDS_password>
```

For Windows:

```
mysqldump ^
    --databases <database_name> ^
    --single-transaction ^
    --compress ^
    --order-by-primary \
    -u <local_user> \
    -p<local_password> | mysql ^
        --host=hostname ^
        --port=3306 ^
        -u <RDS_user_name> ^
```

```
-p <RDS_password>
```

Use the `--host`, `--user (-u)`, `--port` and `-p` options in the `mysql` command to specify the host name, user name, port, and password to connect to your Amazon RDS MariaDB DB instance. The host name is the DNS name from the Amazon RDS MariaDB DB instance endpoint, for example `myinstance.123456789012.us-east-1.rds.amazonaws.com`. You can find the endpoint value in the instance details in the Amazon RDS Management Console.

4. Make the source MariaDB instance writeable again:

```
mysql> SET GLOBAL read_only = OFF;
mysql> UNLOCK TABLES;
```

5. In the Amazon RDS Management Console, add the IP address of the server that hosts the external MariaDB database to the VPC security group for the Amazon RDS MariaDB DB instance. For more information on modifying a VPC security group, go to [Security Groups for Your VPC](#) in the *Amazon Virtual Private Cloud User Guide*.

The IP address can change when the following conditions are met:

- You are using a public IP address for communication between the external master instance and the DB instance.
- The external master instance was stopped and restarted.

If these conditions are met, verify the IP address before adding it.

You might also need to configure your local network to permit connections from the IP address of your Amazon RDS MariaDB DB instance, so that it can communicate with your external MariaDB instance. To find the IP address of the Amazon RDS MariaDB DB instance, use the `host` command:

```
host <RDS_MariaDB_DB_host_name>
```

The host name is the DNS name from the Amazon RDS MariaDB DB instance endpoint.

6. Using the client of your choice, connect to the external MariaDB instance and create a MariaDB user to be used for replication. This account is used solely for replication and must be restricted to your domain to improve security. The following is an example:

```
CREATE USER 'repl_user'@'mydomain.com' IDENTIFIED BY '<password>';
```

7. For the external MariaDB instance, grant `REPLICATION CLIENT` and `REPLICATION SLAVE` privileges to your replication user. For example, to grant the `REPLICATION CLIENT` and `REPLICATION SLAVE` privileges on all databases for the '`repl_user`' user for your domain, issue the following command:

```
GRANT REPLICATION CLIENT, REPLICATION SLAVE ON *.*  
TO 'repl_user'@'mydomain.com'  
IDENTIFIED BY '<password>';
```

8. Make the Amazon RDS MariaDB DB instance the replica. Connect to the Amazon RDS MariaDB DB instance as the master user and identify the external MariaDB database as the replication master by using the [mysql.rds_set_external_master_gtid \(p. 595\)](#) command. Use the GTID that you determined in Step 2. The following is an example:

```
CALL mysql.rds_set_external_master_gtid ('mymasterserver.mydomain.com', 3306,  
'repl_user', '<password>', '<GTID>', 0);
```

9. On the Amazon RDS MariaDB DB instance, issue the [mysql.rds_start_replication \(p. 901\)](#) command to start replication:

```
CALL mysql.rds_start_replication;
```

Importing Data into a MariaDB DB Instance

Following, you can find information about methods to import your MariaDB data to an Amazon RDS DB instance running MariaDB.

To do an initial data import into a MariaDB DB instance, you can use the procedures documented in [Restoring a Backup into an Amazon RDS MySQL DB Instance \(p. 826\)](#), as follows:

- To move data from an Amazon RDS MySQL DB instance, a MariaDB or MySQL instance in Amazon Elastic Compute Cloud (Amazon EC2) in the same VPC as your Amazon RDS MariaDB DB instance, or a small on-premises instance of MariaDB or MySQL, you can use the procedure documented in [Importing Data from a MySQL or MariaDB DB to an Amazon RDS MySQL or MariaDB DB Instance \(p. 833\)](#).
- To move data from a large or production on-premises instance of MariaDB or MySQL, you can use the procedure documented in [Importing Data to an Amazon RDS MySQL or MariaDB DB Instance with Reduced Downtime \(p. 835\)](#).
- To move data from an instance of MariaDB or MySQL that is in EC2 in a different VPC than your Amazon RDS MariaDB DB instance, or to move data from any data source that can output delimited text files, you can use the procedure documented in [Importing Data From Any Source to a MySQL or MariaDB DB Instance \(p. 847\)](#).

You can also use AWS Database Migration Service (AWS DMS) to import data into an Amazon RDS DB instance. AWS DMS can migrate databases without downtime and, for many database engines, continue ongoing replication until you are ready to switch over to the target database. You can migrate to MariaDB from either the same database engine or a different database engine using AWS DMS. If you are migrating from a different database engine, you can use the AWS Schema Conversion Tool to migrate schema objects that are not migrated by AWS DMS. For more information about AWS DMS, see [What is AWS Database Migration Service](#).

You can configure replication into an Amazon RDS MariaDB DB instance using MariaDB global transaction identifiers (GTIDs) when the external instance is MariaDB version 10.0.24 or greater, or using binary log coordinates for MySQL instances or MariaDB instances on earlier versions than 10.0.24. Note that MariaDB GTIDs are implemented differently than MySQL GTIDs, which are not supported by Amazon RDS.

To configure replication into a MariaDB DB instance, you can use the following procedures:

- To configure replication into a MariaDB DB instance from an external MySQL instance or an external MariaDB instance running a version prior to 10.0.24, you can use the procedure documented in [Replication with a MySQL or MariaDB Instance Running External to Amazon RDS \(p. 862\)](#).
- To configure replication into a MariaDB DB instance from an external MariaDB instance running version 10.0.24 or greater, you can use the procedure documented in [Configuring GTID-Based Replication into an Amazon RDS MariaDB DB Instance \(p. 582\)](#).

Note

The mysql system database contains authentication and authorization information required to log into your DB instance and access your data. Dropping, altering, renaming, or truncating tables, data, or other contents of the mysql database in your DB instance can result in errors and might render the DB instance and your data inaccessible. If this occurs, the DB instance can be restored from a snapshot using the AWS CLI `restore-db-instance-from-db-snapshot` or recovered using `restore-db-instance-to-point-in-time` commands.

Options for MariaDB Database Engine

This appendix describes options, or additional features, that are available for Amazon RDS instances running the MariaDB DB engine. To enable these options, you add them to a custom option group, and then associate the option group with your DB instance. For more information about working with option groups, see [Working with Option Groups \(p. 196\)](#).

Amazon RDS supports the following options for MariaDB:

Option ID	Engine Versions
MARIADB_AUDIT_PLUGIN	MariaDB 10.0.24 and later

MariaDB Audit Plugin Support

Amazon RDS supports using the MariaDB Audit Plugin on MariaDB database instances. The MariaDB Audit Plugin records database activity such as users logging on to the database, queries run against the database, and more. The record of database activity is stored in a log file.

Audit Plugin Option Settings

Amazon RDS supports the following settings for the MariaDB Audit Plugin option.

Option Setting	Valid Values	Default Value	Description
SERVER_AUDIT_FILE	/rdsdbdata/log/audit/	/rdsdbdata/log/audit/	The location of the log file. The log file contains the record of the activity specified in SERVER_AUDIT_EVENTS. For more information, see Viewing and Listing Database Log Files (p. 390) and MariaDB Database Log Files (p. 394) .
SERVER_AUDIT_SIZE	1000000000	1000000	The size in bytes that when reached, causes the file to rotate. For more information, see Log File Size (p. 398) .
SERVER_AUDIT_ROTATIONS	100	9	The number of log rotations to save. For more information, see Log File Size (p. 398) and Downloading a Database Log File (p. 390) .
SERVER_AUDIT_EVENTS	CONNECT, QUERY, TABLE, QUERY_DDL, QUERY_DML, QUERY_DCL	CONNECT, QUERY	<p>The types of activity to record in the log. Installing the MariaDB Audit Plugin is itself logged.</p> <ul style="list-style-type: none">• CONNECT: Log successful and unsuccessful connections to the database, and disconnections from the database.• QUERY: Log the text of all queries run against the database.• TABLE: Log tables affected by queries when the queries are run against the database.• QUERY_DDL: Similar to the QUERY event, but returns only data definition language (DDL) queries (CREATE, ALTER, and so on).

Option Setting	Valid Values	Default Value	Description
			<ul style="list-style-type: none"> • <code>QUERY_DML</code>: Similar to the <code>QUERY</code> event, but returns only data manipulation language (DML) queries (<code>INSERT</code>, <code>UPDATE</code>, and so on, and also <code>SELECT</code>). • <code>QUERY_DCL</code>: Similar to the <code>QUERY</code> event, but returns only data control language (DCL) queries (<code>GRANT</code>, <code>REVOKE</code>, and so on).
<code>SERVER_AUDIT_INCL_USERS</code>	Multiple users comma-separated values	None	Include only activity from the specified users. By default, activity is recorded for all users. If a user is specified in both <code>SERVER_AUDIT_EXCL_USERS</code> and <code>SERVER_AUDIT_INCL_USERS</code> , then activity is recorded for the user.
<code>SERVER_AUDIT_EXCL_USERS</code>	Multiple users comma-separated values	None	Exclude activity from the specified users. By default, activity is recorded for all users. If a user is specified in both <code>SERVER_AUDIT_EXCL_USERS</code> and <code>SERVER_AUDIT_INCL_USERS</code> , then activity is recorded for the user.
			<p>The <code>rdsadmin</code> user queries the database every second to check the health of the database. Depending on your other settings, this activity can possibly cause the size of your log file to grow very large, very quickly. If you don't need to record this activity, add the <code>rdsadmin</code> user to the <code>SERVER_AUDIT_EXCL_USERS</code> list.</p> <p>Note CONNECT activity is always recorded for all users, even if the user is specified for this option setting.</p>
<code>SERVER_AUDIT_LOGGING</code>	ON	ON	Logging is active. The only valid value is ON. Amazon RDS does not support deactivating logging. If you want to deactivate logging, remove the MariaDB Audit Plugin. For more information, see Removing the MariaDB Audit Plugin (p. 589) .
<code>SERVER_AUDIT_QUERY_LENGTH</code>	1024	1024	The limit on the length of the query string in a record.

Adding the MariaDB Audit Plugin

The general process for adding the MariaDB Audit Plugin to a DB instance is the following:

1. Create a new option group, or copy or modify an existing option group.
2. Add the option to the option group.
3. Associate the option group with the DB instance.

After you add the MariaDB Audit Plugin, you don't need to restart your DB instance. As soon as the option group is active, auditing begins immediately.

To add the MariaDB Audit Plugin

1. Determine the option group you want to use. You can create a new option group or use an existing option group. If you want to use an existing option group, skip to the next step. Otherwise, create a custom DB option group. Choose **mariadb** for **Engine**, and choose **10.0** or later for **Major engine version**. For more information, see [Creating an Option Group \(p. 197\)](#).
2. Add the **MARIADB_AUDIT_PLUGIN** option to the option group, and configure the option settings. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#). For more information about each setting, see [Audit Plugin Option Settings \(p. 587\)](#).
3. Apply the option group to a new or existing DB instance.
 - For a new DB instance, you apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the MariaDB Database Engine \(p. 543\)](#).
 - For an existing DB instance, you apply the option group by modifying the instance and attaching the new option group. For more information, see [Modifying a DB Instance Running the MariaDB Database Engine \(p. 564\)](#).

Viewing and Downloading the MariaDB Audit Plugin Log

After you enable the MariaDB Audit Plugin, you access the results in the log files the same way you access any other text-based log files. The audit log files are located at `/rdsdbdata/log/audit/`. For information about viewing the log file in the console, see [Viewing and Listing Database Log Files \(p. 390\)](#). For information about downloading the log file, see [Downloading a Database Log File \(p. 390\)](#).

Modifying MariaDB Audit Plugin Settings

After you enable the MariaDB Audit Plugin, you can modify settings for the plugin. For more information about how to modify option settings, see [Modifying an Option Setting \(p. 204\)](#). For more information about each setting, see [Audit Plugin Option Settings \(p. 587\)](#).

Removing the MariaDB Audit Plugin

Amazon RDS doesn't support turning off logging in the MariaDB Audit Plugin. However, you can remove the plugin from a DB instance. When you remove the MariaDB Audit Plugin, the DB instance is restarted automatically to stop auditing.

To remove the MariaDB Audit Plugin from a DB instance, do one of the following:

- Remove the MariaDB Audit Plugin option from the option group it belongs to. This change affects all DB instances that use the option group. For more information, see [Removing an Option from an Option Group \(p. 207\)](#)
- Modify the DB instance and specify a different option group that doesn't include the plugin. This change affects a single DB instance. You can specify the default (empty) option group, or a different custom option group. For more information, see [Modifying a DB Instance Running the MariaDB Database Engine \(p. 564\)](#).

Parameters for MariaDB

By default, a MariaDB DB instance uses a DB parameter group that is specific to a MariaDB database. This parameter group contains some but not all of the parameters contained in the Amazon RDS DB parameter groups for the MySQL database engine. It also contains a number of new, MariaDB-specific parameters. The following MySQL parameters are not available in MariaDB-specific DB parameter groups:

- bind_address
- binlog_error_action
- binlog_gtid_simple_recovery
- binlog_max_flush_queue_time
- binlog_order_commits
- binlog_row_image
- binlog_rows_query_log_events
- binlogging_impossible_mode
- block_encryption_mode
- core_file
- default_tmp_storage_engine
- div_precision_increment
- end_markers_in_json
- enforce_gtid_consistency
- eq_range_index_dive_limit
- explicit_defaults_for_timestamp
- gtid_executed
- gtid-mode
- gtid_next
- gtid_owned
- gtid_purged
- log_bin_basename
- log_bin_index
- log_bin_use_v1_row_events
- log_slow_admin_statements
- log_slow_slave_statements
- log_throttle_queries_not_using_indexes
- master-info-repository
- optimizer_trace
- optimizer_trace_features
- optimizer_trace_limit
- optimizer_trace_max_mem_size
- optimizer_trace_offset
- relay_log_info_repository
- rpl_stop_slave_timeout
- slave_parallel_workers
- slave_pending_jobs_size_max
- slave_rows_search_algorithms

- storage_engine
- table_open_cache_instances
- timed_mutexes
- transaction_allow_batching
- validate_password
- validate_password_dictionary_file
- validate_password_length
- validate_password_mixed_case_count
- validate_password_number_count
- validate_password_policy
- validate_password_special_char_count

For more information on MySQL 5.6 parameters, go to the [MySQL 5.6 documentation](#).

The MariaDB-specific DB parameter groups also contain the following parameters that are applicable to MariaDB only. Acceptable ranges for the modifiable parameters are the same as specified in the MariaDB documentation except where noted. Amazon RDS MariaDB parameters are set to the default values of the storage engine you have selected.

- aria_block_size
- aria_checkpoint_interval
- aria_checkpoint_log_activity
- aria_force_start_after_recovery_failures
- aria_group_commit
- aria_group_commit_interval
- aria_log_dir_path
- aria_log_file_size
- aria_log_purge_type
- aria_max_sort_file_size
- aria_page_checksum
- aria_pagecache_age_threshold
- aria_pagecache_division_limit
- aria_recover

Amazon RDS MariaDB supports the values of NORMAL, OFF, and QUICK, but not FORCE or BACKUP.

- aria_repair_threads
- aria_sort_buffer_size
- aria_stats_method
- aria_sync_log_dir
- binlog_annotation_row_events
- binlog_commit_wait_count
- binlog_commit_wait_usec
- binlog_row_image (MariaDB version 10.1 and later)
- deadlock_search_depth_long
- deadlock_search_depth_short
- deadlock_timeout_long
- deadlock_timeout_short
- explicit_defaults_for_timestamp (MariaDB version 10.1 and later)

- `extra_max_connections`
- `extra_port`
- `feedback`
- `feedback_send_retry_wait`
- `feedback_send_timeout`
- `feedback_url`
- `feedback_user_info`
- `gtid_domain_id`
- `gtid_strict_mode`
- `histogram_size`
- `histogram_type`
- `innodb_adaptive_hash_index_partitions`
- `innodb_background_scrub_data_check_interval` (MariaDB version 10.1 and later)
- `innodb_background_scrub_data_compressed` (MariaDB version 10.1 and later)
- `innodb_background_scrub_data_interval` (MariaDB version 10.1 and later)
- `innodb_background_scrub_data_uncompressed` (MariaDB version 10.1 and later)
- `innodb_buf_dump_status_frequency` (MariaDB version 10.1 and later)
- `innodb_buffer_pool_populate`
- `innodb_cleaner_lsn_age_factor`
- `innodb_compression_algorithm` (MariaDB version 10.1 and later)
- `innodb_corrupt_table_action`
- `innodb_defragment` (MariaDB version 10.1 and later)
- `innodb_defragment_fill_factor` (MariaDB version 10.1 and later)
- `innodb_defragment_fill_factor_n_recs` (MariaDB version 10.1 and later)
- `innodb_defragment_frequency` (MariaDB version 10.1 and later)
- `innodb_defragment_n_pages` (MariaDB version 10.1 and later)
- `innodb_defragment_stats_accuracy` (MariaDB version 10.1 and later)
- `innodb_empty_free_List_algorithm`
- `innodb_fake_changes`
- `innodb_fatal_semaphore_wait_threshold` (MariaDB version 10.1 and later)
- `innodb_foreground_preflush`
- `innodb_idle_flush_pct` (MariaDB version 10.1 and later)
- `innodb_immediate_scrub_data_uncompressed` (MariaDB version 10.1 and later)
- `innodb_instrument_semaphores` (MariaDB version 10.1 and later)
- `innodb_locking_fake_changes`
- `innodb_log_arch_dir`
- `innodb_log_arch_expire_sec`
- `innodb_log_archive`
- `innodb_log_block_size`
- `innodb_log_checksum_algorithm`
- `innodb_max_bitmap_file_size`
- `innodb_max_changed_pages`
- `innodb_prefix_index_cluster_optimization` (MariaDB version 10.1 and later)
- `innodb_sched_priority_cleaner`
- `innodb_scrub_log` (MariaDB version 10.1 and later)
- `innodb_scrub_log_speed` (MariaDB version 10.1 and later)

- innodb_show_locks_held
- innodb_show_verbose_locks
- innodb_simulate_comp_failures
- innodb_stats_modified_counter
- innodb_stats_traditional
- innodb_use_atomic_writes
- innodb_use_fallocate
- innodb_use_global_flush_log_at_trx_commit
- innodb_use_stacktrace
- innodb_use_trim (MariaDB version 10.1 and later)
- join_buffer_space_limit
- join_cache_level
- key_cache_file_hash_size
- key_cache_segments
- max_digest_length (MariaDB version 10.1 and later)
- max_statement_time (MariaDB version 10.1 and later)
- mysql56_temporal_format (MariaDB version 10.1 and later)
- progress_report_time
- query_cache_strip_comments
- replicate_annotation_row_events
- replicate_do_db
- replicate_do_table
- replicate_events_marked_for_skip
- replicate_ignore_db
- replicate_ignore_table
- replicate_wild_ignore_table
- slave_domain_parallel_threads
- slave_parallel_max_queued
- slave_parallel_mode (MariaDB version 10.1 and later)
- slave_parallel_threads
- slave_run_triggers_for_rbr (MariaDB version 10.1 and later)
- sql_error_log_filename
- sql_error_log_rate
- sql_error_log_rotate
- sql_error_log_rotations
- sql_error_log_size_limit
- thread_handling
- thread_pool_idle_timeout
- thread_pool_max_threads
- thread_pool_min_threads
- thread_pool_oversubscribe
- thread_pool_size
- thread_pool_stall_limit
- transaction_write_set_extraction
- use_stat_tables
- userstat

For more information on MariaDB parameters, go to the [MariaDB documentation](#).

MariaDB on Amazon RDS SQL Reference

This appendix describes system stored procedures that are available for Amazon RDS instances running the MariaDB DB engine.

You can use all of the system stored procedures that are available for Amazon RDS MySQL DB instances for MariaDB DB instances also. These stored procedures are documented at [MySQL on Amazon RDS SQL Reference \(p. 889\)](#).

Additionally, the following system stored procedures are supported only for Amazon RDS DB instances running MariaDB:

- [mysql.rds_set_external_master_gtid \(p. 595\)](#)
- [mysql.rds_kill_query_id \(p. 597\)](#)

[mysql.rds_set_external_master_gtid](#)

Configures GTID-based replication from a MariaDB instance running external to Amazon RDS to an Amazon RDS MariaDB DB instance. This stored procedure is supported only where the external MariaDB instance is version 10.0.24 or greater. When setting up replication where one or both instances do not support MariaDB global transaction identifiers (GTIDs), use [mysql.rds_set_external_master \(p. 891\)](#).

Using GTIDs for replication provides crash-safety features not offered by binary log replication, so we recommend it in cases where the replicating instances support it.

Syntax

```
CALL mysql.rds_set_external_master_gtid(
    host_name
    , host_port
    , replication_user_name
    , replication_user_password
    , gtid
    , ssl_encryption
);
```

Parameters

host_name

String. The host name or IP address of the MariaDB instance running external to Amazon RDS that will become the replication master.

host_port

Integer. The port used by the MariaDB instance running external to Amazon RDS to be configured as the replication master. If your network configuration includes SSH port replication that converts the port number, specify the port number that is exposed by SSH.

replication_user_name

String. The ID of a user with REPLICATION SLAVE permissions in the MariaDB DB instance to be configured as the Read Replica.

replication_user_password

String. The password of the user ID specified in *replication_user_name*.

gtid

String. The global transaction ID on the master that replication should start from.

You can use `@@gtid_current_pos` to get the current GTID if the replication master has been locked while you are configuring replication, so the binary log doesn't change between the points when you get the GTID and when replication starts.

Otherwise, if you are using `mysqldump` version 10.0.13 or greater to populate the slave instance prior to starting replication, you can get the GTID position in the output by using the `--master-data` or `--dump-slave` options. If you are not using `mysqldump` version 10.0.13 or greater, you can run the `SHOW MASTER STATUS` or use those same `mysqldump` options to get the binary log file name and position, then convert them to a GTID by running `BINLOG_GTID_POS` on the external MariaDB instance:

```
SELECT BINLOG_GTID_POS('<binary log file name>', <binary log file position>);
```

For more information about the MariaDB implementation of GTIDs, go to [Global Transaction ID](#) in the MariaDB documentation.

ssl_encryption

Integer. This option is not currently implemented. The default is 0.

Usage Notes

The `mysql.rds_set_external_master_gtid` procedure must be run by the master user. It must be run on the MariaDB DB instance that you are configuring as the replication slave of a MariaDB instance running external to Amazon RDS. Before running `mysql.rds_set_external_master_gtid`, you must have configured the instance of MariaDB running external to Amazon RDS as a replication master. For more information, see [Importing Data into a MariaDB DB Instance \(p. 586\)](#).

Warning

Do not use `mysql.rds_set_external_master_gtid` to manage replication between two Amazon RDS DB instances. Use it only when replicating with a MariaDB instance running external to RDS. For information about managing replication between Amazon RDS DB instances, see [Working with Read Replicas \(p. 182\)](#).

After calling `mysql.rds_set_external_master_gtid` to configure an Amazon RDS DB instance as a Read Replica, you can call [mysql.rds_start_replication \(p. 901\)](#) on the replica to start the replication process. You can call [mysql.rds_reset_external_master \(p. 897\)](#) to remove the Read Replica configuration.

When `mysql.rds_set_external_master_gtid` is called, Amazon RDS records the time, user, and an action of "set master" in the `mysql.rds_history` and `mysql.rds_replication_status` tables.

Examples

When run on a MariaDB DB instance, the following example configures it as the replication slave of an instance of MariaDB running external to Amazon RDS.

```
call mysql.rds_set_external_master_gtid
('Sourcedb.some.com',3306,'ReplicationUser','SomePassWord','0-123-456',0);
```

Related Topics

- [mysql.rds_reset_external_master \(p. 897\)](#)

- [mysql.rds_start_replication \(p. 901\)](#)
- [mysql.rds_stop_replication \(p. 903\)](#)

mysql.rds_kill_query_id

Terminates a query running against the MariaDB server.

Syntax

```
CALL mysql.rds_kill_query_id(queryID);
```

Parameters

queryID

Integer. The identity of the query to be terminated.

Usage Notes

To terminate a query running against the MariaDB server, use the `mysql.rds_kill_query_id` procedure and pass in the ID of that query. To obtain the query ID, query the MariaDB [Information Schema PROCESSLIST Table](#), as shown following:

```
SELECT USER, HOST, COMMAND, TIME, STATE, INFO, QUERY_ID FROM
    INFORMATION_SCHEMA.PROCESSLIST WHERE USER = '<user name>';
```

The connection to the MariaDB server is retained.

Related Topics

- [mysql.rds_kill \(p. 910\)](#)
- [mysql.rds_kill_query \(p. 910\)](#)

Examples

The following example terminates a query with a query ID of 230040:

```
call mysql.rds_kill_query_id(230040);
```

Microsoft SQL Server on Amazon RDS

Amazon RDS supports DB instances running several versions and editions of Microsoft SQL Server. The most recent supported version of each major version is shown following. For the full list of supported versions, editions, and RDS engine versions, see [Version and Feature Support on Amazon RDS \(p. 605\)](#).

- SQL Server 2017 RTM CU13 14.00.3049.1, released per [KB4466404](#) on January 01, 2019.
- SQL Server 2016 SP2 CU8 13.00.5426.0, released per [KB4505830](#) on July 31, 2019.
- SQL Server 2014 SP2 CU10 12.00.5571.0, released per [KB4052725](#) on January 16, 2018.
- SQL Server 2012 SP4 GDR 11.00.7462.6, released per [KB4057116](#) on January 12, 2017.
- SQL Server 2008: It is no longer possible to provision new instances in any region. Amazon RDS is actively migrating existing instances off this version.

For information about licensing for SQL Server, see [Licensing Microsoft SQL Server on Amazon RDS \(p. 615\)](#). For information about SQL Server builds, see this Microsoft support article about [the latest SQL Server builds](#).

With Amazon RDS, you can create DB instances and DB snapshots, point-in-time restores, and automated or manual backups. DB instances running SQL Server can be used inside a VPC. You can also use SSL to connect to a DB instance running SQL Server, and you can use TDE to encrypt data at rest. Amazon RDS currently supports Multi-AZ deployments for SQL Server using SQL Server Database Mirroring (DBM) or Always On Availability Groups (AGs) as a high-availability, failover solution.

In order to deliver a managed service experience, Amazon RDS does not provide shell access to DB instances, and it restricts access to certain system procedures and tables that require advanced privileges. Amazon RDS supports access to databases on a DB instance using any standard SQL client application such as Microsoft SQL Server Management Studio. Amazon RDS does not allow direct host access to a DB instance via Telnet, Secure Shell (SSH), or Windows Remote Desktop Connection. When you create a DB instance, the master user is assigned to the *db_owner* role for all user databases on that instance, and has all database-level permissions except for those that are used for backups. Amazon RDS manages backups for you.

Before creating your first DB instance, you should complete the steps in the setting up section of this guide. For more information, see [Setting Up for Amazon RDS \(p. 5\)](#).

Common Management Tasks for Microsoft SQL Server on Amazon RDS

The following are the common management tasks you perform with an Amazon RDS SQL Server DB instance, with links to relevant documentation for each task.

Task Area	Relevant Documentation
Instance Classes, Storage, and PIOPS If you are creating a DB instance for production purposes, you should understand how instance classes, storage types, and Provisioned IOPS work in Amazon RDS.	DB Instance Class Support for Microsoft SQL Server (p. 602) Amazon RDS Storage Types (p. 135)

Task Area	Relevant Documentation
Multi-AZ Deployments A production DB instance should use Multi-AZ deployments. Multi-AZ deployments provide increased availability, data durability, and fault tolerance for DB instances. Multi-AZ deployments for SQL Server are implemented using SQL Server's native DBM or AGs technology.	High Availability (Multi-AZ) for Amazon RDS (p. 144) Multi-AZ Deployments Using Microsoft SQL Server Database Mirroring or Always On Availability Groups (p. 610)
Amazon Virtual Private Cloud (VPC) If your AWS account has a default VPC, then your DB instance is automatically created inside the default VPC. If your account does not have a default VPC, and you want the DB instance in a VPC, you must create the VPC and subnet groups before you create the DB instance.	Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform (p. 506) Working with a DB Instance in a VPC (p. 514)
Security Groups By default, DB instances are created with a firewall that prevents access to them. You therefore must create a security group with the correct IP addresses and network configuration to access the DB instance. The security group you create depends on what Amazon EC2 platform your DB instance is on, and whether you will access your DB instance from an Amazon EC2 instance. In general, if your DB instance is on the <i>EC2-Classic</i> platform, you will need to create a DB security group; if your DB instance is on the <i>EC2-VPC</i> platform, you will need to create a VPC security group.	Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform (p. 506) Controlling Access with Security Groups (p. 486)
Parameter Groups If your DB instance is going to require specific database parameters, you should create a parameter group before you create the DB instance.	Working with DB Parameter Groups (p. 211)
Option Groups If your DB instance is going to require specific database options, you should create an option group before you create the DB instance.	Options for the Microsoft SQL Server Database Engine (p. 712)
Connecting to Your DB Instance After creating a security group and associating it to a DB instance, you can connect to the DB instance using any standard SQL client application such as Microsoft SQL Server Management Studio.	Connecting to a DB Instance Running the Microsoft SQL Server Database Engine (p. 630)
Backup and Restore When you create your DB instance, you can configure it to take automated backups. You can also back up and restore your databases manually by using full backup files (.bak files).	Working With Backups (p. 254) Importing and Exporting SQL Server Databases (p. 657)

Task Area	Relevant Documentation
Monitoring You can monitor your SQL Server DB instance by using CloudWatch Amazon RDS metrics, events, and enhanced monitoring.	Viewing DB Instance Metrics (p. 304) Viewing Amazon RDS Events (p. 388)
Log Files You can access the log files for your SQL Server DB instance.	Amazon RDS Database Log Files (p. 390) Microsoft SQL Server Database Log Files (p. 402)

There are also advanced administrative tasks for working with SQL Server DB instances. For more information, see the following documentation:

- [Common DBA Tasks for Microsoft SQL Server \(p. 723\)](#).
- [Using Windows Authentication with a SQL Server DB Instance \(p. 691\)](#)
- [Accessing the tempdb Database \(p. 724\)](#)

Limits for Microsoft SQL Server DB Instances

The Amazon RDS implementation of Microsoft SQL Server on a DB instance has some limitations that you should be aware of:

- The maximum number of databases supported on a DB instance depends on the instance class type and the availability mode—Single-AZ, Multi-AZ Database Mirroring (DBM), or Multi-AZ Availability Groups (AGs). The Microsoft SQL Server system databases don't count toward this limit.

The following table shows the maximum number of supported databases for each instance class type and availability mode. Use this table to help you decide if you can move from one instance class type to another, or from one availability mode to another. If your source DB instance has more databases than the target instance class type or availability mode can support, modifying the DB instance fails. You can see the status of your request in the **Events** pane.

Instance Class Type	Single-AZ	Multi-AZ with DBM	Multi-AZ with Always On AGs
db.*.micro to db.*.medium	30	N/A	N/A
db.*.large	30	30	30
db.*.xlarge to db.*.16xlarge	100	50	75
db.*.24xlarge	100	50	100

* Represents the different instance class types.

For example, let's say that your DB instance runs on a db.*.16xlarge with Single-AZ and that it has 76 databases. You modify the DB instance to upgrade to using Multi-AZ Always On AGs. This upgrade

fails, because your DB instance contains more databases than your target configuration can support. If you upgrade your instance class type to db.*.24xlarge instead, the modification succeeds.

If the upgrade fails, you see events and messages similar to the following:

- Unable to modify database instance class. The instance has 76 databases, but after conversion it would only support 75.
- Unable to convert the DB instance to Multi-AZ: The instance has 76 databases, but after conversion it would only support 75.

If the point-in-time restore or snapshot restore fails, you see events and messages similar to the following:

- Database instance put into incompatible-restore. The instance has 76 databases, but after conversion it would only support 75.
- Some ports are reserved for Amazon RDS, and you can't use them when you create a DB instance.
- Client connections from IP addresses within the range 169.254.0.0/16 are not permitted. This is the Automatic Private IP Addressing Range (APIPA), which is used for local-link addressing.
- SQL Server Standard Edition will use only a subset of the available processors if the DB instance has more processors than the software limits (24 cores, 4 sockets, and 128GB RAM). Examples of this are the db.m5.24xlarge and db.r5.24xlarge instance classes.
- Amazon RDS for SQL Server doesn't support importing data into the msdb database.
- You can't rename databases on a DB instance in a SQL Server Multi-AZ deployment.
- The maximum storage size for SQL Server DB instances is the following:
 - General Purpose (SSD) storage – 16 TiB for all editions
 - Provisioned IOPS storage – 16 TiB for all editions
 - Magnetic storage – 1 TiB for all editions

If you have a scenario that requires a larger amount of storage, you can use sharding across multiple DB instances to get around the limit. This approach requires data-dependent routing logic in applications that connect to the sharded system. You can use an existing sharding framework, or you can write custom code to enable sharding. If you use an existing framework, the framework can't install any components on the same server as the DB instance.

- The minimum storage size for SQL Server DB instances is the following:
 - General Purpose (SSD) storage – 20 GiB for Enterprise, Standard, Web, and Express editions
 - Provisioned IOPS storage – 20 GiB for Enterprise and Standard editions, 100 GiB for Web and Express editions
 - Magnetic storage – 200 GiB for Enterprise and Standard editions, 20 GiB for Web and Express editions
- Amazon RDS doesn't support running these services on the same server as your Amazon RDS DB instance:
 - SQL Server Analysis Services
 - SQL Server Integration Services
 - SQL Server Reporting Services
 - Data Quality Services
 - Master Data Services

To use these features, we recommend that you install SQL Server on an Amazon EC2 instance, or use an on-premises SQL Server instance. In these cases, the EC2 or SQL Server instance acts as the Reporting, Analysis, Integration, or Master Data Services server for your SQL Server DB instance on Amazon RDS. You can install SQL Server on an Amazon EC2 instance with Amazon EBS storage, pursuant to Microsoft licensing policies.

- Because of limitations in Microsoft SQL Server, restoring to a point in time before successful execution of DROP DATABASE might not reflect the state of that database at that point in time. For example,

the dropped database is typically restored to its state up to 5 minutes before the DROP DATABASE command was issued. This type of restore means that you can't restore the transactions made during those few minutes on your dropped database. To work around this, you can reissue the DROP DATABASE command after the restore operation is completed. Dropping a database removes the transaction logs for that database.

DB Instance Class Support for Microsoft SQL Server

The computation and memory capacity of a DB instance is determined by its DB instance class. The DB instance class you need depends on your processing power and memory requirements. For more information, see [Choosing the DB Instance Class \(p. 107\)](#).

The following list of DB instance classes supported for Microsoft SQL Server is provided here for your convenience. For the most current list, see the RDS console: <https://console.aws.amazon.com/rds/>.

SQL Server Edition	2017 and 2016 Support Range	2014 and 2012 Support Range
Enterprise Edition	db.t3.xlarge-db.t3.2xlarge db.r3.xlarge-db.r3.8xlarge db.r4.xlarge-db.r4.16xlarge db.r5.xlarge-db.r5.24xlarge db.m4.xlarge-db.m4.16xlarge db.m5.xlarge-db.m5.24xlarge db.x1.16xlarge-db.x1.32xlarge db.x1e.xlarge-db.x1e.32xlarge	db.t3.xlarge-db.t3.2xlarge db.r3.xlarge-db.r3.8xlarge db.r4.xlarge-db.r4.8xlarge db.r5.xlarge-db.r5.24xlarge db.m4.xlarge-db.m4.10xlarge db.m5.xlarge-db.m5.24xlarge db.x1.16xlarge-db.x1.32xlarge db.x1e.xlarge-db.x1e.32xlarge
Standard Edition	db.t3.xlarge-db.t3.2xlarge db.r4.large-db.r4.16xlarge db.r5.large-db.r5.24xlarge db.m4.large-db.m4.16xlarge db.m5.large-db.m5.24xlarge db.x1.16xlarge-db.x1.32xlarge db.x1e.xlarge-db.x1e.32xlarge	db.t3.xlarge-db.t3.2xlarge db.r3.large-db.r3.8xlarge db.r4.large-db.r4.8xlarge db.r5.large-db.r5.24xlarge db.m3.medium-db.m3.2xlarge db.m4.large-db.m4.10xlarge db.m5.large-db.m5.24xlarge db.x1.16xlarge-db.x1.32xlarge db.x1e.xlarge-db.x1e.32xlarge
Web Edition	db.t2.small-db.t2.medium db.t3.small-db.t3.2xlarge	db.t2.small-db.t2.medium db.t3.small-db.t3.2xlarge

SQL Server Edition	2017 and 2016 Support Range	2014 and 2012 Support Range
	db.r4.large-db.r4.2xlarge db.r5.large-db.r5.4xlarge db.m4.large-db.m4.4xlarge db.m5.large-db.m5.4xlarge	db.r3.large-db.r3.2xlarge db.r4.large-db.r4.2xlarge db.r5.large-db.r5.4xlarge db.m3.medium-db.m3.2xlarge db.m4.large-db.m4.4xlarge db.m5.large-db.m5.4xlarge
Express Edition	db.t2.micro-db.t2.medium db.t3.small-db.t3.xlarge	db.t2.micro-db.t2.medium db.t3.small-db.t3.xlarge

Microsoft SQL Server Security

The Microsoft SQL Server database engine uses role-based security. The master user name you use when you create a DB instance is a SQL Server Authentication login that is a member of the `processadmin`, `public`, and `setupadmin` fixed server roles.

Any user who creates a database is assigned to the `db_owner` role for that database and has all database-level permissions except for those that are used for backups. Amazon RDS manages backups for you.

The following server-level roles are not currently available in Amazon RDS:

- `bulkadmin`
- `dbcreator`
- `diskadmin`
- `securityadmin`
- `serveradmin`
- `sysadmin`

The following server-level permissions are not available on SQL Server DB instances:

- `ALTER ANY CREDENTIAL`
- `ALTER ANY EVENT NOTIFICATION`
- `ALTER ANY EVENT SESSION`
- `ALTER RESOURCES`
- `ALTER SETTINGS` (you can use the DB parameter group API operations to modify parameters; for more information, see [Working with DB Parameter Groups \(p. 211\)](#))
- `AUTHENTICATE SERVER`
- `CONTROL_SERVER`
- `CREATE DDL EVENT NOTIFICATION`
- `CREATE ENDPOINT`
- `CREATE TRACE EVENT NOTIFICATION`

- EXTERNAL ACCESS ASSEMBLY
- SHUTDOWN (You can use the RDS reboot option instead)
- UNSAFE ASSEMBLY
- ALTER ANY AVAILABILITY GROUP (SQL Server 2012 only)
- CREATE ANY AVAILABILITY GROUP (SQL Server 2012 only)

Compliance Program Support for Microsoft SQL Server DB Instances

AWS Services in Scope have been fully assessed by a third-party auditor and result in a certification, attestation of compliance, or Authority to Operate (ATO). For more information, see [AWS Services in Scope by Compliance Program](#).

HIPAA Support for Microsoft SQL Server DB Instances

You can use Amazon RDS for Microsoft SQL Server databases to build HIPAA-compliant applications. You can store healthcare-related information, including protected health information (PHI), under an executed Business Associate Agreement (BAA) with AWS. For more information, see [HIPAA Compliance](#).

Amazon RDS for SQL Server supports HIPAA for the following versions and editions:

- SQL Server 2017 Enterprise, Standard, and Web Editions
- SQL Server 2016 Enterprise, Standard, and Web Editions
- SQL Server 2014 Enterprise, Standard, and Web Editions
- SQL Server 2012 Enterprise, Standard, and Web Editions

To enable HIPAA support on your DB instance, set up the following three components.

Component	Details
Auditing	To set up auditing, set the parameter <code>rds.sqlserver_audit</code> to the value <code>fedramp_hipaa</code> . If your DB instance is not already using a custom DB parameter group, you must create a custom parameter group and attach it to your DB instance before you can modify the <code>rds.sqlserver_audit</code> parameter. For more information, see Working with DB Parameter Groups (p. 211) .
Transport Encryption	To set up transport encryption, force all connections to your DB instance to use Secure Sockets Layer (SSL). For more information, see Forcing Connections to Your DB Instance to Use SSL (p. 687) .
Encryption at Rest	To set up encryption at rest, you have two options: <ol style="list-style-type: none">1. If you are running Enterprise Edition, you can choose to use Transparent Data Encryption (TDE) to achieve encryption at rest. For more information, see Support for Transparent Data Encryption in SQL Server (p. 714).2. You can set up encryption at rest by using AWS Key Management Service (AWS KMS) encryption keys. For more information, see Encrypting Amazon RDS Resources (p. 429).

SSL Support for Microsoft SQL Server DB Instances

You can use SSL to encrypt connections between your applications and your Amazon RDS DB instances running Microsoft SQL Server. You can also force all connections to your DB instance to use SSL. If you force connections to use SSL, it happens transparently to the client, and the client doesn't have to do any work to use SSL.

SSL is supported in all AWS Regions and for all supported SQL Server editions. For more information, see [Using SSL with a Microsoft SQL Server DB Instance \(p. 687\)](#).

Version and Feature Support on Amazon RDS

You can specify any currently supported Microsoft SQL Server version when creating a new DB instance. You can specify the Microsoft SQL Server major version (such as Microsoft SQL Server 14.00), and any supported minor version for the specified major version. If no version is specified, Amazon RDS defaults to a supported version, typically the most recent version. If a major version is specified but a minor version is not, Amazon RDS defaults to a recent release of the major version you have specified. To see a list of supported versions, as well as defaults for newly created DB instances, use the [describe-db-engine-versions](#) AWS CLI command.

Microsoft SQL Server 2017 Support on Amazon RDS

Amazon RDS supports the following versions of SQL Server 2017:

- SQL Server 2017 RTM (CU13) 14.00.3049.1, for all editions and all AWS Regions.

RDS API `EngineVersion` and CLI `engine-version`: 14.00.3049.1.v1

- SQL Server 2017 RTM CU13 14.00.3015.40, for all editions and all AWS Regions.

RDS API `EngineVersion` and CLI `engine-version`: 14.00.3015.40.v1

- Version 14.00.1000.169, RTM, for all editions, and all AWS Regions.

RDS API `EngineVersion` and CLI `engine-version`: 14.00.1000.169.v1

SQL Server 2017 includes many new features, such as the following:

- Adaptive query processing
- Automatic plan correction
- GraphDB
- Resumable index rebuilds

For the full list of SQL Server 2017 features, see [What's New in SQL Server 2017](#) in the Microsoft documentation.

For a list of unsupported features, see [Features Not Supported and Features with Limited Support \(p. 609\)](#).

Microsoft SQL Server 2016 Support on Amazon RDS

Amazon RDS supports the following versions of SQL Server 2016:

- SQL Server 2016, SP2 CU8 13.00.5426.0, for all editions and all AWS Regions
 - RDS API EngineVersion and CLI engine-version: 13.00.5426.0.v1
- SQL Server 2016, SP1 CU7 13.00.4466.4, for all editions and all AWS Regions
 - RDS API EngineVersion and CLI engine-version: 13.00.4466.4.v1
- Version 13.00.4451.0, SP1 CU5, for all editions, and all AWS Regions
 - RDS API EngineVersion and CLI engine-version: 13.00.4451.0.v1
- Version 13.00.4422.0, SP1 CU2, for all editions, and all AWS Regions
 - RDS API EngineVersion and CLI engine-version: 13.00.4422.0.v1
- Version 13.00.2164.0, RTM CU2, for all editions, and all AWS Regions
 - RDS API EngineVersion and CLI engine-version: 13.00.2164.0.v1

Amazon RDS supports the following features of SQL Server 2016:

- Always Encrypted
- JSON Support
- Operational Analytics
- Query Store
- Temporal Tables

For the full list of SQL Server 2016 features, see [What's New in SQL Server 2016](#) in the Microsoft documentation.

Microsoft SQL Server 2014 Support on Amazon RDS

Amazon RDS supports the following versions of SQL Server 2014:

- SQL Server 2014 SP2 CU10 12.00.5571.0, for all editions, and all AWS Regions.
 - RDS API EngineVersion and CLI engine-version: 12.00.5571.0.v1
- Version 12.00.5546.0, SP2 CU5, for all editions and all AWS Regions
 - RDS API EngineVersion and CLI engine-version: 12.00.5546.0.v1
- Version 12.00.5000.0, SP2, for all editions and all AWS Regions
 - RDS API EngineVersion and CLI engine-version: 12.00.5000.0.v1
- Version 12.00.4422.0, SP1 CU2, for all editions except Enterprise Edition, and all AWS Regions except Canada (Central), and Europe (London)
 - RDS API EngineVersion and CLI engine-version: 12.00.4422.0.v1

In addition to supported features of SQL Server 2012, Amazon RDS supports the new query optimizer available in SQL Server 2014, and also the delayed durability feature.

For a list of unsupported features, see [Features Not Supported and Features with Limited Support \(p. 609\)](#).

SQL Server 2014 supports all the parameters from SQL Server 2012 and uses the same default values. SQL Server 2014 includes one new parameter, backup checksum default. For more information, see

[How to enable the CHECKSUM option if backup utilities do not expose the option](#) in the Microsoft documentation.

Microsoft SQL Server 2012 Support on Amazon RDS

Amazon RDS supports the following versions of SQL Server 2012:

- SQL Server 2012 SP4 GDR 11.00.7462.6, for all editions, for all AWS Regions.
RDS API EngineVersion and CLI engine-version: 11.00.7462.6.v1
- Version 11.00.6594.0, SP3 CU8, for all editions and all AWS Regions
RDS API EngineVersion and CLI engine-version: 11.00.6594.0.v1
- Version 11.00.6020.0, SP3, for all editions and all AWS Regions
RDS API EngineVersion and CLI engine-version: 11.00.6020.0.v1
- Version 11.00.5058.0, SP2, for all editions, and all AWS Regions except US East (Ohio), Canada (Central), and Europe (London)
RDS API EngineVersion and CLI engine-version: 11.00.5058.0.v1

For more information about SQL Server 2012, see [Features Supported by the Editions of SQL Server 2012](#) in the Microsoft documentation.

In addition to supported features of SQL Server 2008 R2, Amazon RDS supports the following SQL Server 2012 features:

- Columnstore indexes (Enterprise Edition)
- Online Index Create, Rebuild and Drop for XML, varchar(max), nvarchar(max), and varbinary(max) data types (Enterprise Edition)
- Flexible Server Roles
- Service Broker is supported, Service Broker endpoints are not supported
- Partially Contained Databases
- Sequences
- Transparent Data Encryption (Enterprise Edition only)
- THROW statement
- New and enhanced spatial types
- UTF-16 Support
- ALTER ANY SERVER ROLE server-level permission

For a list of unsupported features, see [Features Not Supported and Features with Limited Support \(p. 609\)](#).

Some SQL Server parameters have changed in SQL Server 2012.

- The following parameters have been removed from SQL Server 2012: `awe enabled`, `precompute rank`, and `sql mail xps`. These parameters were not modifiable in SQL Server DB Instances and their removal should have no impact on your SQL Server use.
- A new `contained database authentication` parameter in SQL Server 2012 supports partially contained databases. When you enable this parameter and then create a partially contained database, an authorized user's user name and password is stored within the partially contained database instead of in the master database. For more information about partially contained databases, see [Contained Databases](#) in the Microsoft documentation.

Microsoft SQL Server 2008 R2 Deprecated on Amazon RDS

We are upgrading all existing instances that are still using SQL Server 2008 R2 to the latest minor version of SQL Server 2012. For more information, see [Microsoft SQL Server Engine Version Management in Amazon RDS \(p. 608\)](#).

For more information about SQL Server 2008 R2, see [Features Supported by the Editions of SQL Server 2008 R2](#) in the Microsoft documentation.

Microsoft SQL Server Engine Version Management in Amazon RDS

Amazon RDS includes flexible version management that enables you to control when and how your DB instance is patched or upgraded. This enables you to do the following for your DB engine:

- Maintain compatibility with database engine patch versions
- Test new patch versions to verify that they work with your application before you deploy them in production
- Plan and perform version upgrades to meet your service level agreements and timing requirements

Microsoft SQL Server Engine Patching in Amazon RDS

Amazon RDS periodically aggregates official Microsoft SQL Server database patches into a DB instance engine version that's specific to Amazon RDS. For more information about the Microsoft SQL Server patches in each engine version, see [Version and Feature Support on Amazon RDS](#).

Currently, you manually execute all engine upgrades on your DB instance. For more information, see [Upgrading the Microsoft SQL Server DB Engine \(p. 653\)](#).

Deprecation Schedule for Major Engine Versions of Microsoft SQL Server on Amazon RDS

The table following displays the planned schedule of deprecations for major engine versions of Microsoft SQL Server.

Date	Information
July 12, 2019	<p>The Amazon RDS team deprecated support for Microsoft SQL Server 2008 R2 in June 2019. Remaining instances of Microsoft SQL Server 2008 R2 are migrating to SQL Server 2012 (latest minor version available).</p> <p>To avoid an automatic upgrade from Microsoft SQL Server 2008 R2, you can upgrade at a time that is convenient to you. For more information, see Upgrading a DB Instance Engine Version (p. 159).</p>
April 25, 2019	Before the end of April 2019, you will no longer be able to create new Amazon RDS for SQL Server database instances using Microsoft SQL Server 2008R2.

Change Data Capture Support for Microsoft SQL Server DB Instances

Amazon RDS supports change data capture (CDC) for your DB instances running Microsoft SQL Server. CDC captures changes that are made to the data in your tables, and stores metadata about each change that you can access later. For more information, see [Change Data Capture](#) in the Microsoft documentation.

Amazon RDS supports CDC for the following SQL Server editions and versions:

- Microsoft SQL Server Enterprise Edition (All versions)
- Microsoft SQL Server Standard Edition:
 - 2017
 - 2016 version 13.00.4422.0 SP1 CU2 and later

To use CDC with your Amazon RDS DB instances, first enable or disable CDC at the database level by using RDS-provided stored procedures. After that, any user that has the `db_owner` role for that database can use the native Microsoft stored procedures to control CDC on that database. For more information, see [Using Change Data Capture \(p. 731\)](#).

You can use CDC and AWS Database Migration Service to enable ongoing replication from SQL Server DB instances.

Features Not Supported and Features with Limited Support

The following Microsoft SQL Server features are not supported on Amazon RDS:

- Stretch database
- Backing up to Microsoft Azure Blob Storage
- Buffer pool extension
- Data Quality Services
- Database Log Shipping
- Database Mail
- Distribution Transaction Coordinator (MSDTC)
- File tables
- FILESTREAM support
- Maintenance Plans
- Performance Data Collector
- Policy-Based Management
- PolyBase
- Machine Learning and R Services (requires OS access to install it)
- Replication
- Resource Governor
- Server-level triggers

- Service Broker endpoints
- T-SQL endpoints (all operations using CREATE ENDPOINT are unavailable)
- WCF Data Services

The following Microsoft SQL Server features have limited support on Amazon RDS:

- Distributed Queries / Linked Servers. For more information, see: [Implementing Linked Servers with Amazon RDS for Microsoft SQL Server](#).

Multi-AZ Deployments Using Microsoft SQL Server Database Mirroring or Always On Availability Groups

Amazon RDS supports Multi-AZ deployments for DB instances running Microsoft SQL Server by using SQL Server Database Mirroring (DBM) or Always On Availability Groups (AGs). Multi-AZ deployments provide increased availability, data durability, and fault tolerance for DB instances. In the event of planned database maintenance or unplanned service disruption, Amazon RDS automatically fails over to the up-to-date secondary replica so database operations can resume quickly without manual intervention. The primary and secondary instances use the same endpoint, whose physical network address transitions to the passive secondary replica as part of the failover process. You don't have to reconfigure your application when a failover occurs.

Amazon RDS manages failover by actively monitoring your Multi-AZ deployment and initiating a failover when a problem with your primary occurs. Failover doesn't occur unless the standby and primary are fully in sync. Amazon RDS actively maintains your Multi-AZ deployment by automatically repairing unhealthy DB instances and re-establishing synchronous replication. You don't have to manage anything. Amazon RDS handles the primary, the witness, and the standby instance for you. When you set up SQL Server Multi-AZ, RDS configures passive secondary instances for all of the databases on the instance.

For more information, see [Multi-AZ Deployments for Microsoft SQL Server \(p. 681\)](#).

Using Transparent Data Encryption to Encrypt Data at Rest

Amazon RDS supports Microsoft SQL Server Transparent Data Encryption (TDE), which transparently encrypts stored data. Amazon RDS uses option groups to enable and configure these features. For more information about the TDE option, see [Support for Transparent Data Encryption in SQL Server \(p. 714\)](#).

Local Time Zone for Microsoft SQL Server DB Instances

The time zone of an Amazon RDS DB instance running Microsoft SQL Server is set by default. The current default is Universal Coordinated Time (UTC). You can set the time zone of your DB instance to a local time zone instead, to match the time zone of your applications.

You set the time zone when you first create your DB instance. You can create your DB instance by using the [AWS Management Console](#), the Amazon RDS API [CreateDBInstance](#) action, or the AWS CLI [create-db-instance](#) command.

If your DB instance is part of a Multi-AZ deployment (using SQL Server DBM or AGs), then when you fail over, your time zone remains the local time zone that you set. For more information, see [Multi-AZ Deployments Using Microsoft SQL Server Database Mirroring or Always On Availability Groups \(p. 610\)](#).

When you request a point-in-time restore, you specify the time to restore to in UTC. During the restore process, the time is translated to the time zone of the DB instance. For more information, see [Restoring a DB Instance to a Specified Time \(p. 285\)](#).

The following are limitations to setting the local time zone on your DB instance:

- You can't modify the time zone of an existing SQL Server DB instance.
- You can't restore a snapshot from a DB instance in one time zone to a DB instance in a different time zone.
- We strongly recommend that you don't restore a backup file from one time zone to a different time zone. If you restore a backup file from one time zone to a different time zone, you must audit your queries and applications for the effects of the time zone change. For more information, see [Importing and Exporting SQL Server Databases \(p. 657\)](#).

Supported Time Zones

You can set your local time zone to one of the values listed in the following table.

Time Zone	Standard Time Offset	Description	Notes
Afghanistan Standard Time	(UTC+04:30)	Kabul	
Alaskan Standard Time	(UTC–09:00)	Alaska	
Arabian Standard Time	(UTC+04:00)	Abu Dhabi, Muscat	
Atlantic Standard Time	(UTC–04:00)	Atlantic Time (Canada)	
AUS Central Standard Time	(UTC+09:30)	Darwin	
AUS Eastern Standard Time	(UTC+10:00)	Canberra, Melbourne, Sydney	
Belarus Standard Time	(UTC+03:00)	Minsk	This time zone does not observe daylight savings time.
Canada Central Standard Time	(UTC–06:00)	Saskatchewan	
Cape Verde Standard Time	(UTC–01:00)	Cabo Verde Is.	
Cen. Australia Standard Time	(UTC+09:30)	Adelaide	
Central America Standard Time	(UTC–06:00)	Central America	
Central Asia Standard Time	(UTC+06:00)	Astana	

Time Zone	Standard Time Offset	Description	Notes
Central Brazilian Standard Time	(UTC-04:00)	Cuiaba	
Central Europe Standard Time	(UTC+01:00)	Belgrade, Bratislava, Budapest, Ljubljana, Prague	
Central European Standard Time	(UTC+01:00)	Sarajevo, Skopje, Warsaw, Zagreb	
Central Pacific Standard Time	(UTC+11:00)	Solomon Islands, New Caledonia	
Central Standard Time	(UTC-06:00)	Central Time (US and Canada)	
Central Standard Time (Mexico)	(UTC-06:00)	Guadalajara, Mexico City, Monterrey	
China Standard Time	(UTC+08:00)	Beijing, Chongqing, Hong Kong, Urumqi	
E. Africa Standard Time	(UTC+03:00)	Nairobi	This time zone does not observe daylight savings time.
E. Australia Standard Time	(UTC+10:00)	Brisbane	
E. Europe Standard Time	(UTC+02:00)	Chisinau	
E. South America Standard Time	(UTC-03:00)	Brasilia	
Eastern Standard Time	(UTC-05:00)	Eastern Time (US and Canada)	
Georgian Standard Time	(UTC+04:00)	Tbilisi	
GMT Standard Time	(UTC)	Dublin, Edinburgh, Lisbon, London	This time zone is not the same as Greenwich Mean Time. This time zone does observe daylight savings time.
Greenland Standard Time	(UTC-03:00)	Greenland	
Greenwich Standard Time	(UTC)	Monrovia, Reykjavik	This time zone does not observe daylight savings time.
GTB Standard Time	(UTC+02:00)	Athens, Bucharest	
Hawaiian Standard Time	(UTC-10:00)	Hawaii	
India Standard Time	(UTC+05:30)	Chennai, Kolkata, Mumbai, New Delhi	
Jordan Standard Time	(UTC+02:00)	Amman	

Time Zone	Standard Time Offset	Description	Notes
Korea Standard Time	(UTC+09:00)	Seoul	
Middle East Standard Time	(UTC+02:00)	Beirut	
Mountain Standard Time	(UTC−07:00)	Mountain Time (US and Canada)	
Mountain Standard Time (Mexico)	(UTC−07:00)	Chihuahua, La Paz, Mazatlan	
US Mountain Standard Time	(UTC−07:00)	Arizona	This time zone does not observe daylight savings time.
New Zealand Standard Time	(UTC+12:00)	Auckland, Wellington	
Newfoundland Standard Time	(UTC−03:30)	Newfoundland	
Pacific SA Standard Time	(UTC−03:00)	Santiago	
Pacific Standard Time	(UTC−08:00)	Pacific Time (US and Canada)	
Pacific Standard Time (Mexico)	(UTC−08:00)	Baja California	
Russian Standard Time	(UTC+03:00)	Moscow, St. Petersburg, Volgograd	This time zone does not observe daylight savings time.
SA Pacific Standard Time	(UTC−05:00)	Bogota, Lima, Quito, Rio Branco	This time zone does not observe daylight savings time.
SE Asia Standard Time	(UTC+07:00)	Bangkok, Hanoi, Jakarta	
Singapore Standard Time	(UTC+08:00)	Kuala Lumpur, Singapore	
Tokyo Standard Time	(UTC+09:00)	Osaka, Sapporo, Tokyo	
US Eastern Standard Time	(UTC−05:00)	Indiana (East)	
UTC	UTC	Coordinated Universal Time	This time zone does not observe daylight savings time.
UTC−02	(UTC−02:00)	Coordinated Universal Time−02	
UTC−08	(UTC−08:00)	Coordinated Universal Time−08	
UTC−09	(UTC−09:00)	Coordinated Universal Time−09	

Time Zone	Standard Time Offset	Description	Notes
UTC-11	(UTC-11:00)	Coordinated Universal Time-11	
UTC+12	(UTC+12:00)	Coordinated Universal Time+12	
W. Australia Standard Time	(UTC+08:00)	Perth	
W. Central Africa Standard Time	(UTC+01:00)	West Central Africa	
W. Europe Standard Time	(UTC+01:00)	Amsterdam, Berlin, Bern, Rome, Stockholm, Vienna	

Licensing Microsoft SQL Server on Amazon RDS

When you set up an Amazon RDS DB instance for Microsoft SQL Server, the software license is included.

This means that you don't need to purchase SQL Server licenses separately. AWS holds the license for the SQL Server database software. Amazon RDS pricing includes the software license, underlying hardware resources, and Amazon RDS management capabilities.

Amazon RDS supports the following Microsoft SQL Server editions:

- Enterprise
- Standard
- Web
- Express

Note

Licensing for SQL Server Web Edition supports only public and internet-accessible webpages, websites, web applications, and web services. This level of support is required for compliance with Microsoft's usage rights. For more information, see [AWS Service Terms](#).

Amazon RDS supports Multi-AZ deployments for DB instances running Microsoft SQL Server by using SQL Server Database Mirroring (DBM) or Always On Availability Groups (AGs). There are no additional licensing requirements for Multi-AZ deployments. For more information, see [Multi-AZ Deployments for Microsoft SQL Server \(p. 681\)](#).

Restoring License-Terminated DB Instances

Amazon RDS takes snapshots of license-terminated DB instances. If your instance is terminated for licensing issues, you can restore it from the snapshot to a new DB instance. New DB instances have a license included.

For more information, see [Restoring License-Terminated DB Instances \(p. 734\)](#).

Development and Test

Because of licensing requirements, we can't offer SQL Server Developer edition on Amazon RDS. You can use Express edition for many development, testing, and other nonproduction needs. However, if you need the full feature capabilities of an enterprise-level installation of SQL Server, you must use a dedicated host environment. You can download and install SQL Server Developer edition (and other MSDN products) on Amazon EC2. Dedicated infrastructure is not required for Developer edition. By using your own host, you also gain access to other programmability features that are not accessible on Amazon RDS. For more information on the difference between SQL Server editions, see [Editions and supported features of SQL Server 2017](#) in the Microsoft documentation.

Related Topics

- [Microsoft SQL Server on Amazon RDS \(p. 598\)](#)
- [Creating a DB Instance Running the Microsoft SQL Server Database Engine \(p. 616\)](#)

Creating a DB Instance Running the Microsoft SQL Server Database Engine

The basic building block of Amazon RDS is the DB instance. Your Amazon RDS DB instance is similar to your on-premises Microsoft SQL Server. After you create your SQL Server DB instance, you can add one or more custom databases to it.

Important

Before you can create or connect to a DB instance, you must complete the tasks in [Setting Up for Amazon RDS \(p. 5\)](#).

For an example that walks you through the process of creating and connecting to a sample DB instance, see [Creating a Microsoft SQL Server DB Instance and Connecting to a DB Instance \(p. 23\)](#).

Note

If you are using the console, a new console interface is available for database creation. Choose either the **New Console** or the **Original Console** instructions based on the console that you are using. The **New Console** instructions are open by default.

New Console

You can create a DB instance running Microsoft SQL Server with the AWS Management Console with **Easy Create** enabled or not enabled. With **Easy Create** enabled, you specify only the DB engine type, DB instance size, and DB instance identifier. **Easy Create** uses the default setting for other configuration options. With **Easy Create** not enabled, you specify more configuration options when you create a database, including ones for availability, security, backups, and maintenance.

Note

For this example, **Standard Create** is enabled, and **Easy Create** isn't enabled. For information about creating a Microsoft SQL Server DB instance with **Easy Create** enabled, see [Creating a Microsoft SQL Server DB Instance and Connecting to a DB Instance \(p. 23\)](#).

To create a SQL Server DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the Amazon RDS console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
4. Choose **Create database**.
5. In **Choose a database creation method**, select **Standard Create**.
6. In **Engine options**, choose **Microsoft SQL Server**.

Create database

Choose a database creation method Info

Standard Create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy Create

Use recommended best-practice configuration options can be changed after the database is created.

Engine options

Engine type Info

Amazon Aurora



MySQL



MariaDB



PostgreSQL



Oracle

ORACLE®

Microsoft SQL Server



Edition

SQL Server Express Edition

Affordable database management system that supports database sizes up to 10 GB.

SQL Server Web Edition

In accordance with Microsoft's licensing policies, it can only be used to support public and Internet-accessible webpages, websites, web applications, and web services.

SQL Server Standard Edition

Core data management and business intelligence capabilities for mission-critical applications and mixed workloads.

SQL Server Enterprise Edition

Comprehensive high-end capabilities for mission-critical applications with demanding database workloads and business intelligence requirements.

7. For **Edition**, choose the SQL Server DB engine edition that you want to use. The SQL Server editions that are available vary by AWS Region.
8. For **Version**, choose the SQL server version.
9. In **Templates**, choose the template that matches your use case. If you choose **Production**, the following are preselected in a later step:
 - **Multi-AZ failover option**
 - **Provisioned IOPS storage option**
 - **Enable deletion protection** option

We recommend these features for any production environment.

Note

Template choices vary by edition.

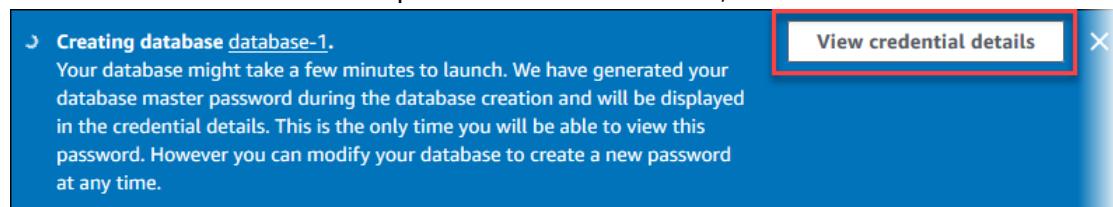
10. To enter your master password, do the following:
 - a. In the **Settings** section, open **Credential Settings**.
 - b. Clear the **Auto generate a password** check box.
 - c. (Optional) Change the **Master username** value and enter the same password in **Master password** and **Confirm password**.

By default, the new DB instance uses an automatically generated password for the master user.

11. For the remaining sections, specify your DB instance settings. For information about each setting, see [Settings for Microsoft SQL Server DB Instances \(p. 625\)](#).
12. Choose **Create database**.

If you chose to use an automatically generated password, the **View credential details** button appears on the **Databases** page.

To view the master user name and password for the DB instance, choose **View credential details**.



To connect to the DB instance as the master user, use the user name and password that appear.

Important

You can't view the master user password again. If you don't record it, you might have to change it. If you need to change the master user password after the DB instance is available, you can modify the DB instance to do so. For more information about modifying a DB instance, see [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#).

13. For **Databases**, choose the name of the new SQL Server DB instance.

On the RDS console, the details for the new DB instance appear. The DB instance has a status of **creating** until the DB instance is created and ready for use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and storage allocated, it can take several minutes for the new instance to be available.

The screenshot shows the 'Summary' page for a database named 'database-1'. The 'Info' section displays the status as 'Creating'. Below the summary, there are tabs for 'Connectivity & security', 'Monitoring', 'Logs & events', 'Configuration', and 'Maintenance & ...'. The 'Connectivity & security' tab is currently selected.

Original Console

To launch a SQL Server DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the top right corner of the Amazon RDS console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
If the navigation pane is closed, choose the menu icon at the top left to open it.
4. Choose **Create database** to open the **Select engine** page.
5. Choose the **Microsoft SQL Server** icon.

Select engine

Engine options

Amazon Aurora
Amazon Aurora

MySQL


MariaDB


PostgreSQL


Oracle
ORACLE

Microsoft SQL Server


Microsoft SQL Server

Edition

SQL Server Express Edition
Affordable database management system that supports database sizes up to 10 GiB.

SQL Server Web Edition
In accordance with Microsoft's licensing policies, it can only be used to support public and Internet-accessible webpages, websites, web applications, and web services.

SQL Server Standard Edition
Core data management and business intelligence capabilities for mission-critical applications and mixed workloads.

SQL Server Enterprise Edition
Comprehensive high-end capabilities for mission-critical applications with demanding database workloads and business intelligence requirements.

Aurora global database feature is now available.
This feature is now available in our new database creation flow.

[Try it now](#)

Only enable options eligible for RDS Free Usage Tier [Info](#)

[Cancel](#) **Next**

6. Choose the SQL Server DB engine edition that you want to use. The SQL Server editions that are available vary by AWS Region.
7. For some editions, the **Use Case** step asks if you are planning to use the DB instance you are creating for production. If you are, choose **Production**. If you choose **Production**, the following are all preselected in a later step:
 - Multi-AZ failover option

- **Provisioned IOPS** storage option
- **Enable deletion protection** option

We recommend these features for any production environment.

8. Choose **Next** to continue. The **Specify DB Details** page appears.

On the **Specify DB Details** page, specify your DB instance information. For information about each setting, see [Settings for Microsoft SQL Server DB Instances \(p. 625\)](#).

Specify DB details

Instance specifications

Estimate your monthly costs for the DB Instance using the [AWS Simple Monthly Calculator](#) 

DB engine

Microsoft SQL Server Express Edition

License model [Info](#)

license-included

DB engine version [Info](#)

SQL Server 2017 14.00.3035.2.v1



Free tier

The Amazon RDS Free Tier provides a single db.t2.micro instance as well as up to 20 GiB of storage, allowing new AWS customers to gain hands-on experience with Amazon RDS. Learn more about the RDS Free Tier and the instance restrictions [here](#).

Only enable options eligible for RDS Free Usage Tier [Info](#)

DB instance class [Info](#)

db.t2.small — 1 vCPU, 2 GiB RAM

Time zone (optional)

No preference

Storage type [Info](#)

9. Choose **Next** to continue. The **Configure Advanced Settings** page appears.

On the **Configure Advanced Settings** page, provide additional information that Amazon RDS needs to launch the DB instance. For information about each setting, see [Settings for Microsoft SQL Server DB Instances \(p. 625\)](#).

Configure advanced settings

Network & Security

Virtual Private Cloud (VPC) [Info](#)

VPC defines the virtual networking environment for this DB instance.

Default VPC (vpc-2aed394c)



Only VPCs with a corresponding DB subnet group are listed.

Subnet group [Info](#)

DB subnet group that defines which subnets and IP ranges the DB instance can use in the VPC you selected.

default



Public accessibility [Info](#)

Yes

EC2 instances and devices outside of the VPC hosting the DB instance will connect to the DB instances. You must also select one or more VPC security groups that specify which EC2 instances and devices can connect to the DB instance.

No

DB instance will not have a public IP address assigned. No EC2 instance or devices outside of the VPC will be able to connect.

Availability zone [Info](#)

No preference



VPC security groups

Security groups have rules authorizing connections from all the EC2 instances and devices that need to access the DB instance.

Create new VPC security group

Choose existing VPC security groups

10. Choose **Launch DB Instance**.

11. On the final page of the wizard, choose **Close**.

On the RDS console, the new DB instance appears in the list of DB instances. The DB instance has a status of **creating** until the DB instance is ready to use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and the amount of storage, it can take up to 20 minutes before the new instance is available.

The screenshot shows the 'Databases' section of the AWS RDS console. At the top left, it says 'RDS > Databases'. Below that is a table header with columns: DB Name, Role, Engine, and Region. The table lists five databases:

DB Name	Role	Engine	Region
mymariadb	Instance	MariaDB	us-east-1
myoracledb	Instance	Oracle Enterprise Edition	us-east-1
mypostgresql	Instance	PostgreSQL	us-east-1
mysqldb	Instance	SQL Server Express Edition	us-east-1
testauroramysql-cl	Regional	Aurora MySQL	us-east-1

AWS CLI

To create a Microsoft SQL Server DB instance by using the AWS CLI, call the [create-db-instance](#) command with the parameters below. For information about each setting, see [Settings for Microsoft SQL Server DB Instances \(p. 625\)](#).

- --db-instance-identifier
- --db-instance-class
- --db-security-groups
- --db-subnet-group
- --engine
- --master-user-name
- --master-user-password
- --allocated-storage
- --backup-retention-period

Example

For Linux, OS X, or Unix:

```
aws rds create-db-instance
  --engine sqlserver-se \
  --db-instance-identifier mymsftsqlserver \
  --allocated-storage 250 \
  --db-instance-class db.m1.large \
  --db-security-groups mydbsecuritygroup \
  --db-subnet-group mydbsubnetgroup \
  --master-user-name masterawsuser \
  --master-user-password masteruserpassword \
  --backup-retention-period 3
```

For Windows:

```
aws rds create-db-instance ^
```

```
--engine sqlserver-se ^
--db-instance-identifier mydbinstance ^
--allocated-storage 250 ^
--db-instance-class db.m1.large ^
--db-security-groups mydbsecuritygroup ^
--db-subnet-group mydbsubnetgroup ^
--master-user-name masterawsuser ^
--master-user-password masteruserpassword ^
--backup-retention-period 3
```

This command should produce output similar to the following:

```
DBINSTANCE mydbinstance db.m1.large sqlserver-se 250 sa creating 3 **** n
10.50.2789
SECGROUP default active
PARAMGRP default.sqlserver-se-10.5 in-sync
```

RDS API

To create a Microsoft SQL Server DB instance by using the Amazon RDS API, call the [CreateDBInstance](#) operation with the parameters below. For information about each setting, see [Settings for Microsoft SQL Server DB Instances \(p. 625\)](#).

- AllocatedStorage
- BackupRetentionPeriod
- DBInstanceClass
- DBInstanceIdentifier
- DBSecurityGroups
- DBSubnetGroup
- Engine
- MasterUsername
- MasterUserPassword

Example

```
https://rds.amazonaws.com/
?Action=CreateDBInstance
&AllocatedStorage=250
&BackupRetentionPeriod=3
&DBInstanceClass=db.m1.large
&DBInstanceIdentifier=mydbinstance
&DBSecurityGroups.member.1=mysecuritygroup
&DBSubnetGroup=mydbsubnetgroup
&Engine=sqlserver-se
&MasterUserPassword=masteruserpassword
&MasterUsername=masterawsuser
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140305/us-west-1/rds/aws4_request
&X-Amz-Date=20140305T185838Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=b441901545441d3c7a48f63b5b1522c5b2b37c137500c93c45e209d4b3a064a3
```

Settings for Microsoft SQL Server DB Instances

The following table contains details about settings that you choose when you create a SQL Server DB instance.

Setting	Setting Description
Allocated storage	The amount of storage to allocate for your DB instance (in gigabytes). In some cases, allocating a higher amount of storage for your DB instance than the size of your database can improve I/O performance. For more information, see Amazon RDS DB Instance Storage (p. 135) .
Auto minor version upgrade	Not supported on Amazon RDS for SQL Server
Availability zone	The Availability Zone for your DB instance. Use the default value of No Preference unless you want to specify an Availability Zone. For more information, see Regions and Availability Zones (p. 132) .
Backup retention period	The number of days that you want automatic backups of your DB instance to be retained. For any nontrivial DB instance, set this value to 1 or greater. For more information, see Working With Backups (p. 254) .
Backup window	The time period during which Amazon RDS automatically takes a backup of your DB instance. Unless you have a specific time that you want to have your database backed up, use the default of No Preference . For more information, see Working With Backups (p. 254) .
Collation	A server-level collation for your DB instance. For more information, see Server-Level Collation for Microsoft SQL Server (p. 728) .
Copy tags to snapshots	This option copies any DB instance tags to a DB snapshot when you create a snapshot. For more information, see Tagging Amazon RDS Resources (p. 177) .
Database port	The port that you want to access the DB instance through. SQL Server installations default to port 1433. If you use a DB security group with your DB instance, this port value must be the same one that you provided when creating the DB security group.
Database authentication	The database authentication option you want to use. Choose Password authentication to authenticate database users with database passwords only.

Setting	Setting Description
	<p>Choose Password and Kerberos authentication to authenticate database users with database passwords and Kerberos authentication through an AWS Managed Microsoft AD created with AWS Directory Service. Next, choose the directory or choose Create a new Directory. For more information, see Using Kerberos Authentication with Amazon RDS for Oracle (p. 1134).</p>
DB engine version	The version of Microsoft SQL Server that you want to use.
DB instance class	<p>The configuration for your DB instance. For example, a db.m1.small instance class has 1.7 GiB memory, 1 ECU (1 virtual core with 1 ECU), 64-bit platform, and moderate I/O capacity.</p> <p>If possible, choose an instance class large enough that a typical query working set can be held in memory. When working sets are held in memory the system can avoid writing to disk, and this improves performance.</p> <p>For more information, see Choosing the DB Instance Class (p. 107) and DB Instance Class Support for Microsoft SQL Server (p. 602).</p>
DB instance identifier	<p>The name for your DB instance. Name your DB instances in the same way that you name your on-premises servers. Your DB instance identifier can contain up to 63 alphanumeric characters, and must be unique for your account in the AWS Region you chose. You can add some intelligence to the name, such as including the AWS Region and DB engine you chose, for example sqlsv-instance1.</p>
DB parameter group	<p>A parameter group for your DB instance. You can choose the default parameter group or you can create a custom parameter group.</p> <p>For more information, see Working with DB Parameter Groups (p. 211).</p>
Deletion protection	<p>Enable deletion protection to prevent your DB instance from being deleted. If you create a production DB instance with the AWS Management Console, deletion protection is enabled by default.</p> <p>For more information, see Deleting a DB Instance (p. 174).</p>
Encryption	<p>Enable Encryption to enable encryption at rest for this DB instance.</p> <p>For more information, see Encrypting Amazon RDS Resources (p. 429).</p>
Enhanced Monitoring	<p>Enable enhanced monitoring to enable gathering metrics in real time for the operating system that your DB instance runs on.</p> <p>For more information, see Enhanced Monitoring (p. 307).</p>

Setting	Setting Description
License model	The license model that you want to use. Choose license-included to use the general license agreement for Microsoft SQL Server.
Maintenance window	<p>The 30-minute window in which pending modifications to your DB instance are applied. If the time period doesn't matter, choose No Preference.</p> <p>For more information, see The Amazon RDS Maintenance Window (p. 156).</p>
Master password	The password for your master user account. The password must contain from 8 to 128 printable ASCII characters (excluding /, ", a space, and @).
Master username	<p>The name that you use as the master user name to log on to your DB instance with all database privileges. The master user name is a SQL Server Authentication login that is a member of the processadmin, public, and setupadmin fixed server roles.</p> <p>For more information, see Microsoft SQL Server Security (p. 603).</p>
Multi-AZ deployment	<p>Create a standby instance to create a passive secondary replica of your DB instance in another Availability Zone for failover support. We recommend Multi-AZ for production workloads to maintain high availability.</p> <p>For development and testing, you can choose Do not create a standby instance.</p> <p>For more information, see Multi-AZ Deployments for Microsoft SQL Server (p. 681).</p>
Option group	<p>An option group for your DB instance. You can choose the default option group or you can create a custom option group.</p> <p>For more information, see Working with Option Groups (p. 196).</p>
Performance Insights	<p>Enable Performance Insights to monitor your DB instance load so that you can analyze and troubleshoot your database performance.</p> <p>Choose a retention period to determine how much rolling data history to keep. The default of seven days is in the free tier. Long-term retention (two years) is priced per vCPU per month.</p> <p>Choose a master key to use to protect the key used to encrypt this database volume. Choose from the master keys in your account, or enter the key from a different account.</p> <p>For more information, see Using Amazon RDS Performance Insights (p. 319).</p>

Setting	Setting Description
Publicly accessible	<p>Yes to give the DB instance a public IP address, meaning that it's accessible outside the VPC. To be publicly accessible, the DB instance also has to be in a public subnet in the VPC.</p> <p>No to make the DB instance accessible only from inside the VPC.</p> <p>For more information, see Hiding a DB Instance in a VPC from the Internet (p. 515).</p>
Storage autoscaling	<p>Enable storage autoscaling to enable Amazon RDS to automatically increase storage when needed to avoid having your DB instance run out of storage space.</p> <p>Use Maximum storage threshold to set the upper limit for Amazon RDS to automatically increase storage for your DB instance. The default is 1,000 GiB.</p> <p>For more information, see Managing Capacity Automatically with Amazon RDS Storage Autoscaling (p. 232).</p>
Storage type	<p>The storage type for your DB instance.</p> <p>For more information, see Amazon RDS Storage Types (p. 135).</p>
Subnet group	<p>This setting depends on the platform that you are on. If you are a new customer to AWS, choose default, which is the default DB subnet group that was created for your account.</p> <p>If you are creating a DB instance on the earlier E2-Classic platform, you might want your DB instance in a specific VPC. In this case, choose the DB subnet group that you created for that VPC.</p>
Time zone	<p>The time zone for your DB instance. If you don't choose a time zone, your DB instance uses the default time zone.</p> <p>For more information, see Local Time Zone for Microsoft SQL Server DB Instances (p. 610).</p>
Virtual Private Cloud (VPC)	<p>This setting depends on the platform that you are on. If you are a new customer to AWS, choose the default VPC shown. If you are creating a DB instance on the earlier E2-Classic platform that doesn't use a VPC, choose Not in VPC.</p> <p>For more information, see Amazon Virtual Private Cloud VPCs and Amazon RDS (p. 505).</p>

Setting	Setting Description
VPC security group	<p>If you are a new customer to AWS, Create new to create a new VPC security group. Otherwise, Choose existing, then choose from security groups that you previously created.</p> <p>When you choose Create new in the RDS console, a new security group is created. This new security group has an inbound rule that allows access to the DB instance from the IP address detected in your browser.</p> <p>For more information, see Working with DB Security Groups (EC2-Classic Platform) (p. 491).</p>

Connecting to a DB Instance Running the Microsoft SQL Server Database Engine

After Amazon RDS provisions your DB instance, you can use any standard SQL client application to connect to the DB instance. In this topic you connect to your DB instance by using either Microsoft SQL Server Management Studio (SSMS) or SQL Workbench/J.

For an example that walks you through the process of creating and connecting to a sample DB instance, see [Creating a Microsoft SQL Server DB Instance and Connecting to a DB Instance \(p. 23\)](#).

Connecting to Your DB Instance with Microsoft SQL Server Management Studio

In this procedure you connect to your sample DB instance by using Microsoft SQL Server Management Studio (SSMS). To download a stand-alone version of this utility, see [Download SQL Server Management Studio \(SSMS\)](#) in the Microsoft documentation.

To connect to a DB Instance using SSMS

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the Amazon RDS console, choose the AWS Region of your DB instance.
3. Find the DNS name and port number for your DB instance:
 - a. Open the RDS console and choose **Databases** to display a list of your DB instances.
 - b. Choose the SQL Server DB instance name to display its details.
 - c. On the **Connectivity & security** tab, copy the endpoint. Also, note the port number. You need both the endpoint and the port number to connect to the DB instance.

The screenshot shows the AWS RDS console for a database named "database-2". The "Summary" tab is selected, displaying basic information like DB identifier, Role, and Instance. Below it, the "Connectivity & security" tab is selected, showing the endpoint and port details.

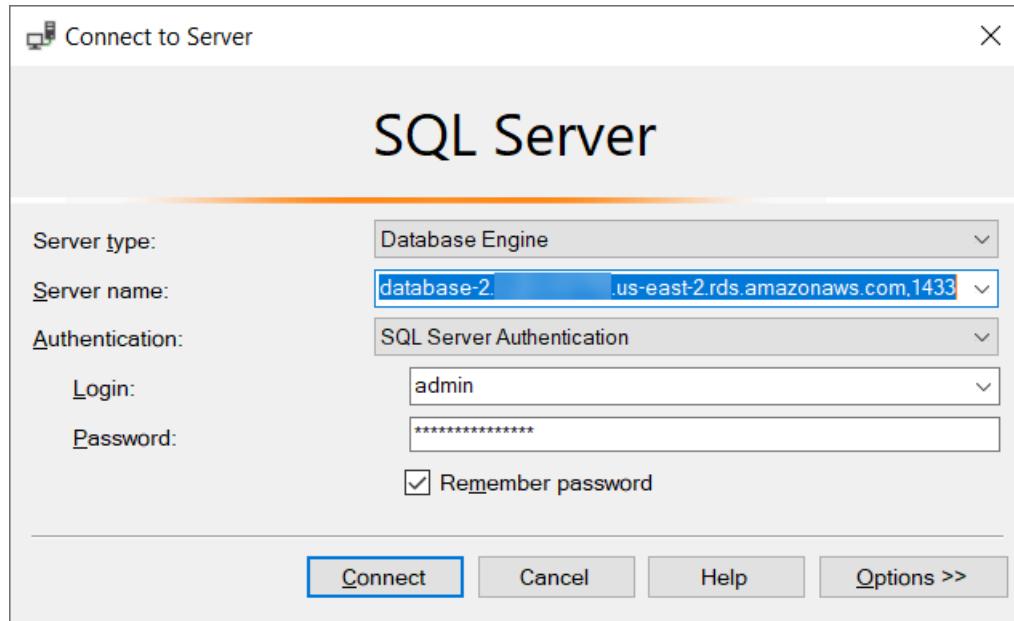
Summary	Connectivity & security
DB identifier database-2	CPU
Role	Current
Instance	

Connectivity & security

Endpoint	Port
database-2. [REDACTED].us-east-2.rds.amazonaws.com	1433

4. Start SQL Server Management Studio.

The **Connect to Server** dialog box appears.



5. Provide the information for your DB instance:
 - a. For **Server type**, choose **Database Engine**.
 - b. For **Server name**, enter the DNS name and port number of your DB Instance, separated by a comma.

Important

Change the colon between the DNS name and port number to a comma.

For example, your server name should look like the following.

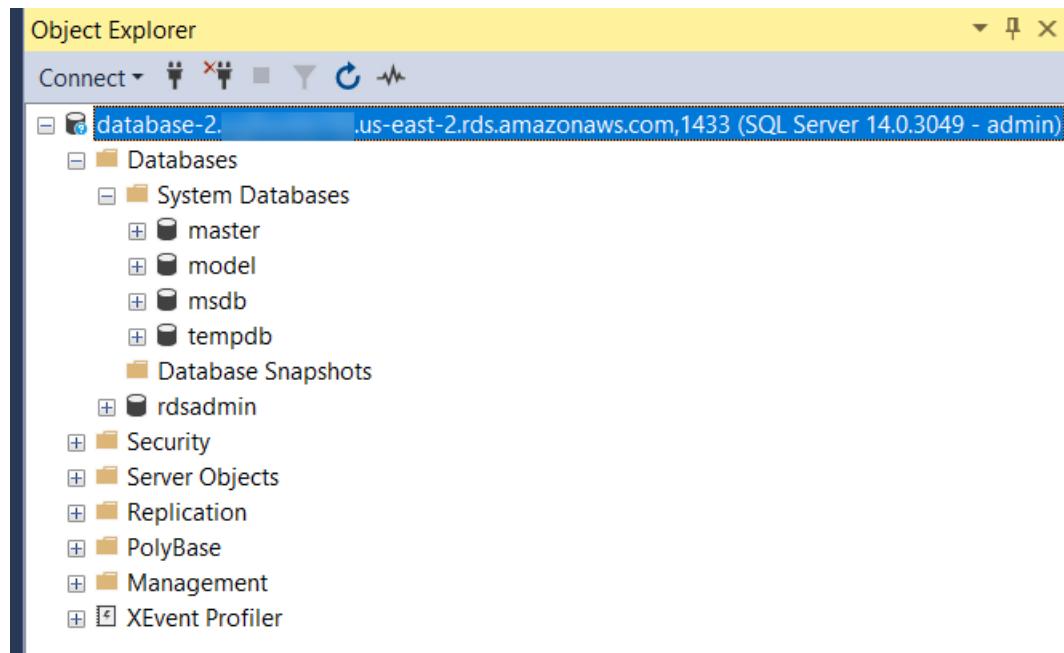
```
database-2.cg034hpkfake.us-east-1.rds.amazonaws.com,1433
```

- c. For **Authentication**, choose **SQL Server Authentication**.
- d. For **Login**, enter the master user name for your DB instance.
- e. For **Password**, enter the password for your DB instance.

6. Choose **Connect**.

After a few moments, SSMS connects to your DB instance. If you can't connect to your DB instance, see [Security Group Considerations \(p. 636\)](#) and [Troubleshooting the Connection to Your SQL Server DB Instance \(p. 637\)](#).

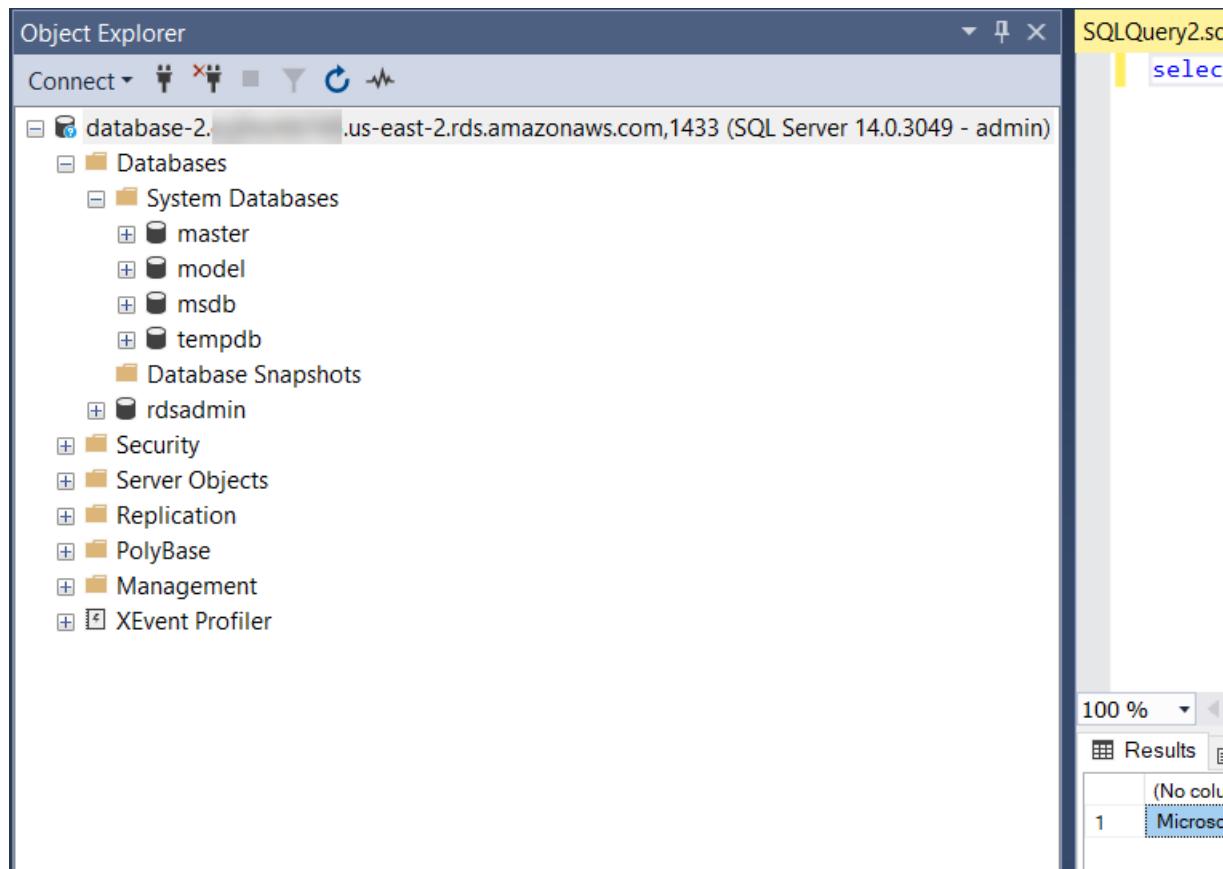
7. Your SQL Server DB instance comes with SQL Server's standard built-in system databases (master, model, msdb, and tempdb). To explore the system databases, do the following:
 - a. In SSMS, on the **View** menu, choose **Object Explorer**.
 - b. Expand your DB instance, expand **Databases**, and then expand **System Databases** as shown following.



8. Your SQL Server DB instance also comes with a database named `rdsadmin`. Amazon RDS uses this database to store the objects that it uses to manage your database. The `rdsadmin` database also includes stored procedures that you can run to perform advanced tasks. For more information, see [Common DBA Tasks for Microsoft SQL Server \(p. 723\)](#).
9. You can now start creating your own databases and running queries against your DB instance and databases as usual. To run a test query against your DB instance, do the following:
 - a. In SSMS, on the **File** menu point to **New** and then choose **Query with Current Connection**.
 - b. Enter the following SQL query.

```
select @@VERSION
```

- c. Run the query. SSMS returns the SQL Server version of your Amazon RDS DB instance.



Connecting to Your DB Instance with SQL Workbench/J

This example shows how to connect to a DB instance running the Microsoft SQL Server database engine by using the SQL Workbench/J database tool. To download SQL Workbench/J, see [SQL Workbench/J](#).

SQL Workbench/J uses JDBC to connect to your DB instance. You also need the JDBC driver for SQL Server. To download this driver, see [Microsoft JDBC Drivers 4.1 \(Preview\)](#) and [4.0 for SQL Server](#).

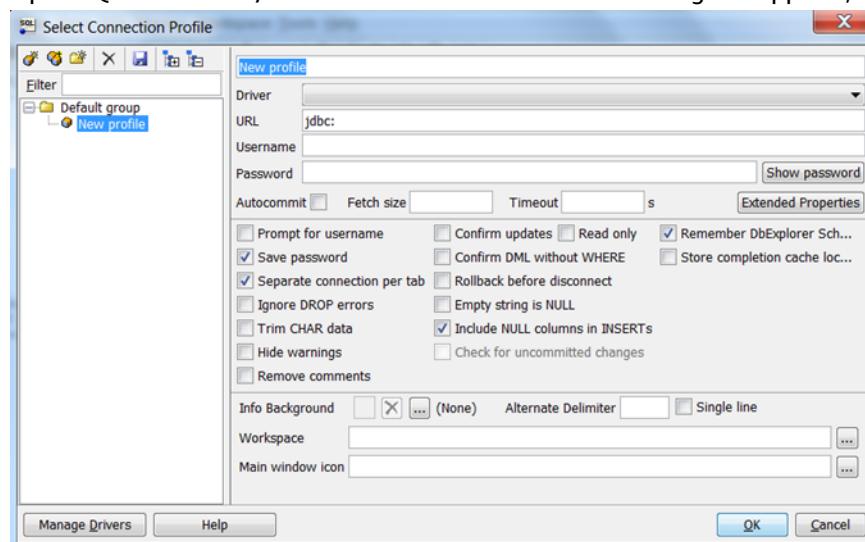
To connect to a DB instance using SQL Workbench

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the Amazon RDS console, choose the AWS Region of your DB instance.
3. Find the DNS name and port number for your DB instance:
 - a. Open the RDS console and then choose **Databases** to display a list of your DB instances.
 - b. Choose the name of your SQL Server DB instance to display its details.

The screenshot shows the 'Summary' section of the 'database-2' instance in the AWS RDS console. It includes fields for 'DB identifier' (database-2), 'Role' (Instance), and 'CPU' usage. Below this is the 'Connectivity & security' tab, which displays the 'Endpoint & port' information: 'Endpoint' is database-2.us-east-2.rds.amazonaws.com and 'Port' is 1433.

c. On the **Connectivity** tab, copy the endpoint. Also, note the port used by the DB instance.

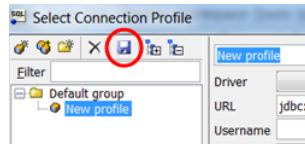
4. Open SQL Workbench/J. The **Select Connection Profile** dialog box appears, as shown following.



5. In the first box at the top of the dialog box, enter a name for the profile.
6. For **Driver**, choose **SQL JDBC 4.0**.
7. For **URL**, enter **jdbc:sqlserver://**, then enter the endpoint of your DB instance. For example, the URL value might be the following.

```
jdbc:sqlserver://sqlsvr-pdz.abcd12340.us-west-2.rds.amazonaws.com:1433
```

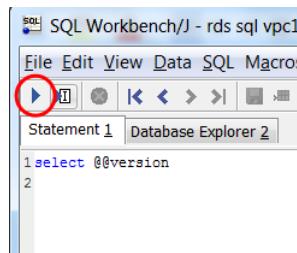
8. For **Username**, enter the master user name for the DB instance.
9. For **Password**, enter the password for the master user.
10. Choose the save icon in the dialog toolbar, as shown following.



11. Choose **OK**. After a few moments, SQL Workbench/J connects to your DB instance. If you can't connect to your DB instance, see [Security Group Considerations \(p. 636\)](#) and [Troubleshooting the Connection to Your SQL Server DB Instance \(p. 637\)](#).
12. In the query pane, enter the following SQL query.

```
select @@VERSION
```

13. Choose the execute icon in the toolbar, as shown following.



The query returns the version information for your DB instance, similar to the following.

```
Microsoft SQL Server 2012 - 11.0.2100.60 (X64)
```

Security Group Considerations

To connect to your DB instance, your DB instance must be associated with a security group that contains the IP addresses and network configuration that you use to access the DB instance. You may have associated your DB instance with an appropriate security group when you created your DB instance. If you assigned a default, non-configured security group when you created your DB instance, your DB instance firewall prevents connections.

If you need to create a new security group to enable access, the type of security group that you create will depend on what Amazon EC2 platform your DB instance is on. To determine your platform, see [Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform \(p. 506\)](#). In general, if your DB instance is on the *EC2-Classic* platform, you create a DB security group; if your DB instance is on the *VPC* platform, you create a VPC security group. For instructions on creating a new security group, see [Controlling Access with Security Groups \(p. 486\)](#).

After you have created the new security group, you modify your DB instance to associate it with the security group. For more information, see [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#).

You can enhance security by using SSL to encrypt connections to your DB instance. For more information, see [Using SSL with a Microsoft SQL Server DB Instance \(p. 687\)](#).

Troubleshooting the Connection to Your SQL Server DB Instance

The following are issues you might encounter when you attempt to connect to your SQL Server DB instance.

Issue	Troubleshooting Suggestions
Unable to connect to your DB instance	<p>For a newly created DB instance, the DB instance has a status of creating until the DB instance is ready to use. When the state changes to available, you can connect to the DB instance. Depending on the DB instance class and the amount of storage, it can take up to 20 minutes before the new instance is available.</p>
	<p>If you can't send or receive communications over the port that you specified when you created the DB instance, you can't connect to the DB instance. Check with your network administrator to verify that the port you specified for your DB instance allows inbound and outbound communication.</p>
	<p>The access rules enforced by your local firewall and the IP addresses that you authorized to access your DB instance in the security group for the DB instance might not match. The problem is most likely the inbound or outbound rules on your firewall.</p> <p>You can add or edit an inbound rule in the security group. For Source, choose My IP. This allows access to the DB instance from the IP address detected in your browser.</p> <p>For more information about security groups, see Controlling Access with Security Groups (p. 486). For a topic that walks you through the process of setting up rules for your security group, see Tutorial: Create an Amazon VPC for Use with a DB Instance (p. 521).</p> <p>Note Client connections from IP addresses within the range 169.254.0.0/16 are not permitted. This is the Automatic Private IP Addressing Range (APIPA), which is used for local-link addressing.</p>
Could not open a connection to SQL Server – Microsoft SQL Server, Error: 53	<p>Make sure specified the server name correctly. For Server name, enter the DNS name and port number of your sample DB Instance, separated by a comma.</p> <p>Important Change the colon between the DNS name and port number to a comma.</p> <p>For example, your server name should look like the following:</p>

Issue	Troubleshooting Suggestions
	sample-instance.cg034hpkmmjt.us-east-1.rds.amazonaws.com, 1433
No connection could be made because the target machine actively refused it – Microsoft SQL Server, Error: 10061	You were able to reach the DB instance but the connection was refused. This is usually caused by specifying the user name or password incorrectly. Verify the user name and password and then retry.

Updating Applications to Connect to Microsoft SQL Server DB Instances Using New SSL/TLS Certificates

As of September 19, 2019, Amazon RDS has published new Certificate Authority (CA) certificates for connecting to your RDS DB instances using Secure Socket Layer or Transport Layer Security (SSL/TLS). The previous CA certificates expire on March 5, 2020. Following, you can find information about updating your applications to use the new certificates. If your application connects to an RDS DB instance using SSL/TLS, you must take the following steps before **March 5, 2020**. Doing this means you can avoid interruption of connectivity between your applications and your RDS DB instances.

This topic can help you to determine whether any client applications use SSL/TLS to connect to your DB instances. If they do, you can further check whether those applications require certificate verification to connect.

Note

Some applications are configured to connect to SQL Server DB instances only if they can successfully verify the certificate on the server.

For such applications, you must update your client application trust stores to include the new CA certificates.

After you update your CA certificates in the client application trust stores, you can rotate the certificates on your DB instances. We strongly recommend testing these procedures in a development or staging environment before implementing them in your production environments.

For more information about certificate rotation, see [Rotating Your SSL/TLS Certificate \(p. 433\)](#). For information about using SSL/TLS with Microsoft SQL Server DB instances, see [Using SSL with a Microsoft SQL Server DB Instance \(p. 687\)](#).

Topics

- [Determining Whether Any Applications are Connecting to Your Microsoft SQL Server DB Instance Using SSL \(p. 639\)](#)
- [Determining Whether a Client Requires Certificate Verification in Order to Connect \(p. 640\)](#)
- [Updating Your Application Trust Store \(p. 641\)](#)

Determining Whether Any Applications are Connecting to Your Microsoft SQL Server DB Instance Using SSL

Check the DB instance configuration for the value of the `rds.force_ssl` parameter. By default, the `rds.force_ssl` parameter is set to 0 (off). If the `rds.force_ssl` parameter is set to 1 (on), clients are required to use SSL/TLS for connections. For more information about parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#). You can also find the setting for this parameter in the `sys.dm_server_registry` DMV.

Run the following query to get the current encryption option for all the open connections to a DB instance. The column `ENCRYPT_OPTION` returns `TRUE` if the connection is encrypted.

```
select SESSION_ID,
```

```
SELECT  
    ENCRYPT_OPTION,  
    NET_TRANSPORT,  
    AUTH_SCHEME  
FROM SYS.DM_EXEC_CONNECTIONS
```

This query shows only the current connections. It doesn't show whether applications that have connected and disconnected in the past have used SSL.

Determining Whether a Client Requires Certificate Verification in Order to Connect

You can check whether different types of clients require certificate verification to connect.

Note

If you use connectors other than the ones listed, see the specific connector's documentation for information about how it enforces encrypted connections. For more information, see [Connection modules for Microsoft SQL databases](#) in the Microsoft SQL Server documentation.

SQL Server Management Studio

Check whether encryption is enforced for SQL Server Management Studio connections:

1. Launch SQL Server Management Studio.
2. For **Connect to server**, enter the server information, login user name, and password.
3. Choose **Options**.
4. Check if **Encrypt connection** is selected in the connect page.

For more information about SQL Server Management Studio, see [Use SQL Server Management Studio](#).

sqlcmd

The following example with the sqlcmd client shows how to check a script's MySQL connection to determine whether successful connections require a valid certificate. For more information, see [Connecting with sqlcmd](#) in the Microsoft SQL Server documentation.

When using sqlcmd, an SSL connection requires verification against the server certificate if you use the -N command argument to encrypt connections, as in the following example.

```
$ sqlcmd -N -S dbinstance.rds.amazonaws.com -d ExampleDB
```

Note

If sqlcmd is invoked with the -C option, it trusts the server certificate, even if that doesn't match the client-side trust store.

ADO.NET

In the following example, the application connects using SSL, and the server certificate must be verified.

```
using SQLC = Microsoft.Data.SqlClient;
```

```
...
    static public void Main()
    {
        using (var connection = new SQLC.SqlConnection(
            "Server=tcp:dbinstance.rds.amazonaws.com;" +
            "Database=ExampleDB;User ID=LOGIN_NAME;" +
            "Password=YOUR_PASSWORD;" +
            "Encrypt=True;TrustServerCertificate=False;")
        )
        {
            connection.Open();
            ...
        }
    }
}
```

Java

In the following example, the application connects using SSL, and the server certificate must be verified.

```
String connectionUrl =
    "jdbc:sqlserver://dbinstance.rds.amazonaws.com;" +
    "databaseName=ExampleDB;integratedSecurity=true;" +
    "encrypt=true;trustServerCertificate=false";
```

To enable SSL encryption for clients that connect using JDBC, you might need to add the Amazon RDS certificate to the Java CA certificate store. For instructions, see [Configuring the client for encryption](#) in the Microsoft SQL Server documentation. You can also provide the trusted CA certificate file name directly by appending `trustStore=path-to-certificate-trust-store-file` to the connection string.

Note

If you use `TrustServerCertificate=true` (or its equivalent) in the connection string, the connection process skips the trust chain validation. In this case, the application connects even if the certificate can't be verified. Using `TrustServerCertificate=false` enforces certificate validation and is a best practice.

Updating Your Application Trust Store

You can update the trust store for applications that use Microsoft SQL Server. For instructions, see [Encrypting Specific Connections \(p. 688\)](#). Also, see [Configuring the client for encryption](#) in the Microsoft SQL Server documentation.

If you are using an operating system other than Microsoft Windows, see the software distribution documentation for SSL/TLS implementation for information about adding a new root CA certificate. For example, OpenSSL and GnuTLS are popular options. Use the implementation method to add trust to the RDS root CA certificate. Microsoft provides instructions for configuring certificates on some systems.

Note

When you update the trust store, you can retain older certificates in addition to adding the new certificates.

Updating Your Application Trust Store for JDBC

You can update the trust store for applications that use JDBC for SSL/TLS connections.

To update the trust store for JDBC applications

1. Download the 2019 root certificate that works for all AWS Regions and put the file in the trust store directory.

For information about downloading the root certificate, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).

2. Convert the certificate to .der format using the following command.

```
openssl x509 -outform der -in rds-ca-2019-root.pem -out rds-ca-2019-root.der
```

Replace the file name with the one that you downloaded.

3. Import the certificate into the key store using the following command.

```
keytool -import -alias rds-root -keystore clientkeystore -file rds-ca-2019-root.der
```

4. Confirm that the key store was updated successfully.

```
keytool -list -v -keystore clientkeystore.jks
```

Enter the key store password when you are prompted for it.

Your output should contain the following.

```
rds-root,date, trustedCertEntry,  
Certificate fingerprint (SHA1):  
D4:0D:DB:29:E3:75:0D:FF:A6:71:C3:14:0B:BF:5F:47:8D:1C:80:96  
# This fingerprint should match the output from the below command  
openssl x509 -fingerprint -in rds-ca-2019-root.pem -noout
```

Important

After you have determined that your database connections use SSL/TLS and have updated your application trust store, you can update your database to use the rds-ca-2019 certificates. For instructions, see step 3 in [Updating Your CA Certificate by Modifying Your DB Instance \(p. 434\)](#).

Modifying a DB Instance Running the Microsoft SQL Server Database Engine

You can change the settings of a DB instance to accomplish tasks such as changing the instance class or renaming the instance. This topic guides you through modifying an Amazon RDS DB instance running Microsoft SQL Server, and describes the settings for SQL Server DB instances.

We recommend that you test any changes on a test instance before modifying a production instance, so that you fully understand the impact of each change. This is especially important when upgrading database versions.

After you modify your DB instance settings, you can apply the changes immediately, or apply them during the next maintenance window for the DB instance. Some modifications cause an interruption by restarting the DB instance.

Console

To modify an SQL Server DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to modify.
3. Choose **Modify**. The **Modify DB Instance** page appears.
4. Change any of the settings that you want. For information about each setting, see [Settings for Microsoft SQL Server DB Instances \(p. 644\)](#).
5. When all the changes are as you want them, choose **Continue**.
6. To apply the changes immediately, choose **Apply Immediately**. Choosing this option can cause an outage in some cases. For more information, see [Using the Apply Immediately Parameter \(p. 150\)](#).
7. On the confirmation page, review your changes. If they are correct, choose **Modify DB Instance** to save your changes.

Alternatively, choose **Back** to edit your changes, or choose **Cancel** to cancel your changes.

AWS CLI

To modify a Microsoft SQL Server DB instance by using the AWS CLI, call the [modify-db-instance](#) command. Specify the DB instance identifier, and the parameters for the settings that you want to modify. For information about each parameter, see [Settings for Microsoft SQL Server DB Instances \(p. 644\)](#).

Example

The following code modifies `mydbinstance` by setting the backup retention period to 1 week (7 days). The code enables automatic minor version upgrades by using `--auto-minor-version-upgrade`. To disable automatic minor version upgrades, use `--no-auto-minor-version-upgrade`. The changes are applied during the next maintenance window by using `--no-apply-immediately`. Use `--apply-immediately` to apply the changes immediately. For more information, see [Using the Apply Immediately Parameter \(p. 150\)](#).

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier mydbinstance \
```

```
--backup-retention-period 7 \
--auto-minor-version-upgrade \
--no-apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier mydbinstance ^
--backup-retention-period 7 ^
--auto-minor-version-upgrade ^
--no-apply-immediately
```

RDS API

To modify a Microsoft SQL Server DB instance by using the Amazon RDS API, call the [ModifyDBInstance](#) operation. Specify the DB instance identifier, and the parameters for the settings that you want to modify. For information about each parameter, see [Settings for Microsoft SQL Server DB Instances \(p. 644\)](#).

Example

The following code modifies `mydbinstance` by setting the backup retention period to 1 week (7 days) and enabling automatic minor version upgrades. These changes are applied during the next maintenance window.

```
https://rds.amazonaws.com/
?Action=ModifyDBInstance
&ApplyImmediately=false
&AutoMinorVersionUpgrade=true
&BackupRetentionPeriod=7
&DBInstanceIdentifier=mydbinstance
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20131016/us-west-1/rds/aws4_request
&X-Amz-Date=20131016T233051Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=087a8eb41cb1ab0fc9ec1575f23e73757ffc6a1e42d7d2b30b9cc0be988cff97
```

Settings for Microsoft SQL Server DB Instances

The following table contains details about which settings you can modify, which settings you can't modify, when the changes can be applied, and whether the changes cause downtime for the DB instance.

Setting	Setting Description	When the Change Occurs	Downtime Notes
Allocated storage	<p>The storage, in gibibytes, that you want to allocate for your DB instance. You can only increase the allocated storage, you can't reduce the allocated storage. The maximum storage allowed is 16 TiB.</p> <p>Warning Once Amazon RDS begins to modify your DB instance</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	<p>A short outage of a few minutes may occur. After that, the DB instance is online but in the storage-optimization state. Performance might be degraded during storage</p>

Setting	Setting Description	When the Change Occurs	Downtime Notes
	<p>to increase the storage size or type, you can't submit another request to increase the storage size or type for 6 hours.</p> <p>You can't modify the storage of some older DB instances, and DB instances restored from older DB snapshots. The Allocated Storage option is disabled in the console if your DB instance isn't eligible. You can also check eligibility by using the AWS CLI command describe-valid-db-instance-modifications which returns the valid storage options for your DB instance.</p> <p>For more information, see Amazon RDS DB Instance Storage (p. 135).</p>		optimization. The storage optimization process is usually short, but can sometimes take up to and even beyond 24 hours.
Auto minor version upgrade	Not supported on Amazon RDS for SQL Server	–	–
Backup retention period	<p>The number of days that automatic backups are retained. To disable automatic backups, set the backup retention period to 0.</p> <p>For more information, see Working With Backups (p. 254).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false and you change the setting from a nonzero value to another nonzero value, the change is applied asynchronously, as soon as possible. Otherwise, the change occurs during the next maintenance window.</p>	An outage occurs if you change from 0 to a nonzero value, or from a nonzero value to 0.
Backup window	<p>The time range during which automated backups of your databases occur. The backup window is a start time in Universal Coordinated Time (UTC), and a duration in hours.</p> <p>For more information, see Working With Backups (p. 254).</p>	The change is applied asynchronously, as soon as possible.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Certificate authority	The certificate that you want to use.	If Apply Immediately is set to true, the change occurs immediately. If Apply Immediately is set to false, the change occurs during the next maintenance window.	An outage occurs during this change.
Copy tags to snapshots	If you have any DB instance tags, this option copies them when you create a DB snapshot. For more information, see Tagging Amazon RDS Resources (p. 177) .	The change occurs immediately. This setting ignores the Apply immediately setting.	–
Database authentication	The database authentication option you want to use. Choose Password authentication to authenticate database users with database passwords only. Choose Password and Kerberos authentication to authenticate database users with database passwords and Kerberos authentication through an AWS Managed Microsoft AD created with AWS Directory Service. Next, choose the directory or choose Create a new Directory . For more information, see Using Kerberos Authentication with Amazon RDS for Oracle (p. 1134) .	If Apply immediately is set to true, the change occurs immediately. If Apply immediately is set to false, the change occurs during the next maintenance window.	When you enable Password and Kerberos authentication , a brief outage occurs.
Database port	The port that you want to use to access the database. The port value must not match any of the port values specified for options in the option group for the DB instance.	The change occurs immediately. This setting ignores the Apply Immediately setting.	The DB instance is rebooted immediately.
DB engine version	The version of the SQL Server database engine that you want to use. Before you upgrade your production DB instances, we recommend that you test the upgrade process on a test instance to verify its duration and to validate your applications. For more information, see Upgrading the Microsoft SQL Server DB Engine (p. 653) .	If Apply Immediately is set to true, the change occurs immediately. If Apply Immediately is set to false, the change occurs during the next maintenance window.	An outage occurs during this change.

Setting	Setting Description	When the Change Occurs	Downtime Notes
DB instance class	<p>The DB instance class that you want to use.</p> <p>For more information, see Choosing the DB Instance Class (p. 107) and DB Instance Class Support for Microsoft SQL Server (p. 602).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change.
DB instance identifier	<p>The DB instance identifier.</p> <p>For more information about the effects of renaming a DB instance, see Renaming a DB Instance (p. 163).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change. The DB instance is rebooted.
DB parameter group	<p>The parameter group that you want associated with the DB instance.</p> <p>For more information, see Working with DB Parameter Groups (p. 211).</p>	The parameter group change occurs immediately.	<p>An outage doesn't occur during this change. When you change the parameter group, changes to some parameters are applied to the DB instance immediately without a reboot. Changes to other parameters are applied only after the DB instance is rebooted.</p> <p>For more information, see Rebooting a DB Instance (p. 166).</p>
Deletion protection	Select Enable deletion protection to prevent your DB instance from being deleted. For more information, see Deleting a DB Instance (p. 174) .	–	–
Domain	<p>The Active Directory Domain to move the instance to. Specify none to remove the instance from its current domain. The domain must exist prior to this operation.</p> <p>For more information, see Using Windows Authentication with a Microsoft SQL Server DB Instance (p. 691).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	<p>A brief outage occurs during this change. For single-AZ DB instances, the outage is approximately 5-10 minutes. For multi-AZ DB instances, the outage is approximately 1 minute.</p>

Setting	Setting Description	When the Change Occurs	Downtime Notes
Domain IAM role name	<p>The name of the IAM role to use when accessing the Active Directory Service.</p> <p>For more information, see Using Windows Authentication with a Microsoft SQL Server DB Instance (p. 691).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	A brief outage occurs during this change.
Enhanced monitoring	<p>Enable enhanced monitoring to enable gathering metrics in real time for the operating system that your DB instance runs on.</p> <p>For more information, see Enhanced Monitoring (p. 307).</p>	–	–
License model	<p>Choose license-included to use the general license agreement for Microsoft SQL Server.</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change.
Maintenance window	<p>The time range during which system maintenance occurs. System maintenance includes upgrades, if applicable. The maintenance window is a start time in Universal Coordinated Time (UTC), and a duration in hours.</p> <p>If you set the window to the current time, there must be at least 30 minutes between the current time and end of the window to ensure any pending changes are applied.</p> <p>For more information, see The Amazon RDS Maintenance Window (p. 156).</p>	<p>The change occurs immediately. This setting ignores the Apply Immediately setting.</p>	<p>If there are one or more pending actions that cause an outage, and the maintenance window is changed to include the current time, then those pending actions are applied immediately, and an outage occurs.</p>

Setting	Setting Description	When the Change Occurs	Downtime Notes
Multi-AZ deployment	<p>Yes to deploy your DB instance in multiple Availability Zones. Otherwise, No.</p> <p>If your DB instance is running Database Mirroring (DBM) – not Always On Availability Groups (AGs) – with SQL Server 2014, 2016, or 2017 Enterprise Edition, and has in-memory optimization enabled, disable in-memory optimization before you add Multi-AZ. If it is running AGs, it doesn't require this step.</p> <p>For more information, see Multi-AZ Deployments for Microsoft SQL Server (p. 681).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	–
New master password	<p>The password for your master user. The password must contain from 8 to 128 printable ASCII characters (excluding /, ", a space, and @). By resetting the master password, you also reset permissions for the DB instance.</p> <p>For more information, see Resetting the DB Instance Owner Role Password (p. 1422).</p>	The change is applied asynchronously, as soon as possible. This setting ignores the Apply Immediately setting.	–
Option group	<p>The option group that you want associated with the DB instance.</p> <p>For more information, see Working with Option Groups (p. 196).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Performance Insights	<p>Select Enable Performance Insights to monitor your DB instance load so that you can analyze and troubleshoot your database performance.</p> <p>Choose a Retention period to determine how much rolling data history to retain. The default of seven days is in the free tier; long-term retention (two years) is priced per vCPU per month.</p> <p>You cannot change the Master key after the database is created.</p> <p>For more information, see Using Amazon RDS Performance Insights (p. 319).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–
Public accessibility	<p>Yes to give the DB instance a public IP address, meaning that it's accessible outside the VPC. To be publicly accessible, the DB instance also has to be in a public subnet in the VPC.</p> <p>No to make the DB instance accessible only from inside the VPC.</p> <p>For more information, see Hiding a DB Instance in a VPC from the Internet (p. 515).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–
Security group	<p>The security group that you want associated with the DB instance.</p> <p>For more information, see Working with DB Security Groups (EC2-Classic Platform) (p. 491).</p>	The change is applied asynchronously, as soon as possible. This setting ignores the Apply Immediately setting.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Storage autoscaling	<p>Select Enable storage autoscaling to enable Amazon RDS to automatically increase storage when needed to avoid having your DB instance run out of storage space.</p> <p>Use Maximum storage threshold to set the upper limit for Amazon RDS to automatically increase storage for your DB instance. The default is 1000 GiB.</p> <p>For more information, see Managing Capacity Automatically with Amazon RDS Storage Autoscaling (p. 232).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–
Storage type	<p>The storage type that you want to use.</p> <p>You can't change from or to magnetic storage.</p> <p>For more information, see Amazon RDS Storage Types (p. 135).</p> <p>Warning Once Amazon RDS begins to modify your DB instance to change the storage size or type, you can't submit another request to change the storage size or type for 6 hours.</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	<p>The following changes all result in a brief outage while the process starts. After that, you can use your database normally while the change takes place.</p> <ul style="list-style-type: none"> From General Purpose (SSD) to Provisioned IOPS (SSD). From Provisioned IOPS (SSD) to General Purpose (SSD).
Subnet group	<p>The subnet group for the DB instance. You can use this setting to move your DB instance to a different VPC. If your DB instance is not in a VPC, you can use this setting to move your DB instance into a VPC.</p> <p>For more information, see Moving a DB Instance Not in a VPC into a VPC (p. 520).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change. The DB instance is rebooted.

Related Topics

- [Rebooting a DB Instance \(p. 166\)](#)
- [Connecting to a DB Instance Running the Microsoft SQL Server Database Engine \(p. 630\)](#)
- [Upgrading the Microsoft SQL Server DB Engine \(p. 653\)](#)
- [Deleting a DB Instance \(p. 174\)](#)

Upgrading the Microsoft SQL Server DB Engine

When Amazon RDS supports a new version of a database engine, you can upgrade your DB instances to the new version. There are two kinds of upgrades for SQL Server DB instances: major version upgrades and minor version upgrades.

Major version upgrades can contain database changes that are not backward-compatible with existing applications. As a result, you must manually perform major version upgrades of your DB instances. You can initiate a major version upgrade by modifying your DB instance. However, before you perform a major version upgrade, we recommend that you test the upgrade by following the steps described in [Testing an Upgrade \(p. 655\)](#).

In contrast, *minor version upgrades* include only changes that are backward-compatible with existing applications. You can initiate a minor version upgrade manually by modifying your DB instance.

Amazon RDS on SQL Server doesn't support automatic minor version upgrades. You can confirm this by using the `describe-db-engine-versions` AWS CLI command. For example:

```
aws rds describe-db-engine-versions --engine sqlserver-se --engine-version 14.00.3049.1.v1
```

In this example, the CLI command returns the following response that indicates that the upgrade will *not* be automatic, even if **Auto minor version upgrade** had been enabled:

```
...
"ValidUpgradeTarget": [
  {
    "Engine": "sqlserver-se",
    "EngineVersion": "14.00.3192.2.v1",
    "Description": "SQL Server 2017 14.00.3192.2.v1",
    "AutoUpgrade": false,
    "IsMajorVersionUpgrade": false
  }
]
```

For more information about performing upgrades, see [Upgrading a SQL Server DB Instance \(p. 656\)](#). For information about what SQL Server versions are available on Amazon RDS, see [Microsoft SQL Server on Amazon RDS \(p. 598\)](#).

Topics

- [Overview of Upgrading \(p. 653\)](#)
- [Major Version Upgrades \(p. 654\)](#)
- [Multi-AZ and In-Memory Optimization Considerations \(p. 655\)](#)
- [Option and Parameter Group Considerations \(p. 655\)](#)
- [Testing an Upgrade \(p. 655\)](#)
- [Upgrading a SQL Server DB Instance \(p. 656\)](#)
- [Upgrading Deprecated DB Instances Before Support Ends \(p. 656\)](#)

Overview of Upgrading

Amazon RDS takes two DB snapshots during the upgrade process. The first DB snapshot is of the DB instance before any upgrade changes have been made. If the upgrade doesn't work for your databases,

you can restore this snapshot to create a DB instance running the old version. The second DB snapshot is taken after the upgrade completes.

Note

Amazon RDS only takes DB snapshots if you have set the backup retention period for your DB instance to a number greater than 0. To change your backup retention period, see [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#).

After an upgrade is complete, you can't revert to the previous version of the database engine. If you want to return to the previous version, restore the DB snapshot that was taken before the upgrade to create a new DB instance.

During a minor or major version upgrade of SQL Server, the **Free Storage Space** and **Disk Queue Depth** metrics will display -1. After the upgrade is complete, both metrics will return to normal.

Major Version Upgrades

Amazon RDS currently supports the following major version upgrades to a Microsoft SQL Server DB instance.

You can upgrade your existing DB instance to SQL Server 2017 from any version except SQL Server 2008. To upgrade from SQL Server 2008, first upgrade to one of the other versions.

Current Version	Supported Upgrade Versions
SQL Server 2016	SQL Server 2017
SQL Server 2014	SQL Server 2017 SQL Server 2016
SQL Server 2012	SQL Server 2017 SQL Server 2016 SQL Server 2014
SQL Server 2008 R2 (Deprecated)	SQL Server 2016 SQL Server 2014 SQL Server 2012

Database Compatibility Level

You can use Microsoft SQL Server database compatibility levels to adjust some database behaviors to mimic previous versions of SQL Server. For more information, see [Compatibility Level](#) in the Microsoft documentation.

When you upgrade your DB instance, all existing databases remain at their original compatibility level. For example, if you upgrade from SQL Server 2012 to SQL Server 2014, all existing databases have a compatibility level of 110. Any new database created after the upgrade have compatibility level 120.

You can change the compatibility level of a database by using the ALTER DATABASE command. For example, to change a database named `customeracct` to be compatible with SQL Server 2014, issue the following command:

```
ALTER DATABASE customeracct SET COMPATIBILITY_LEVEL = 120
```

Multi-AZ and In-Memory Optimization Considerations

Amazon RDS supports Multi-AZ deployments for DB instances running Microsoft SQL Server by using SQL Server Database Mirroring (DBM) or Always On Availability Groups (AGs). For more information, see [Multi-AZ Deployments for Microsoft SQL Server \(p. 681\)](#).

If your DB instance is in a Multi-AZ deployment, both the primary and standby instances are upgraded. Amazon RDS does rolling upgrades. You have an outage only for the duration of a failover.

SQL Server 2014/2016/2017 Enterprise Edition supports in-memory optimization.

Option and Parameter Group Considerations

Option Group Considerations

If your DB instance uses a custom option group, in some cases Amazon RDS can't automatically assign your DB instance a new option group. For example, when you upgrade to a new major version. In that case, you must specify a new option group when you upgrade. We recommend that you create a new option group, and add the same options to it as your existing custom option group.

For more information, see [Creating an Option Group \(p. 197\)](#) or [Making a Copy of an Option Group \(p. 199\)](#).

Parameter Group Considerations

If your DB instance uses a custom parameter group, in some cases Amazon RDS can't automatically assign your DB instance a new parameter group. For example, when you upgrade to a new major version. In that case, you must specify a new parameter group when you upgrade. We recommend that you create a new parameter group, and configure the parameters as in your existing custom parameter group.

For more information, see [Creating a DB Parameter Group \(p. 212\)](#) or [Copying a DB Parameter Group \(p. 217\)](#).

Testing an Upgrade

Before you perform a major version upgrade on your DB instance, you should thoroughly test your database, and all applications that access the database, for compatibility with the new version. We recommend that you use the following procedure.

To test a major version upgrade

1. Review the upgrade documentation for the new version of the database engine to see if there are compatibility issues that might affect your database or applications:
 - [Upgrade to SQL Server 2016 or 2017](#)
 - [Upgrade to SQL Server 2014](#)
 - [Upgrade to SQL Server 2012](#)
2. If your DB instance uses a custom option group, create a new option group compatible with the new version you are upgrading to. For more information, see [Option Group Considerations \(p. 655\)](#).

3. If your DB instance uses a custom parameter group, create a new parameter group compatible with the new version you are upgrading to. For more information, see [Parameter Group Considerations \(p. 655\)](#).
4. Create a DB snapshot of the DB instance to be upgraded. For more information, see [Creating a DB Snapshot \(p. 264\)](#).
5. Restore the DB snapshot to create a new test DB instance. For more information, see [Restoring from a DB Snapshot \(p. 266\)](#).
6. Modify this new test DB instance to upgrade it to the new version, by using one of the following methods:
 - [Console \(p. 159\)](#)
 - [AWS CLI \(p. 160\)](#)
 - [RDS API \(p. 160\)](#)
7. Evaluate the storage used by the upgraded instance to determine if the upgrade requires additional storage.
8. Run as many of your quality assurance tests against the upgraded DB instance as needed to ensure that your database and application work correctly with the new version. Implement any new tests needed to evaluate the impact of any compatibility issues you identified in step 1. Test all stored procedures and functions. Direct test versions of your applications to the upgraded DB instance.
9. If all tests pass, then perform the upgrade on your production DB instance. We recommend that you do not allow write operations to the DB instance until you confirm that everything is working correctly.

Upgrading a SQL Server DB Instance

For information about manually or automatically upgrading a SQL Server DB instance, see the following:

- [Upgrading a DB Instance Engine Version \(p. 159\)](#)
- [Best practices for upgrading SQL Server 2008 R2 to SQL Server 2016 on Amazon RDS for SQL Server](#)

Important

If you have any snapshots that are encrypted using KMS, we recommend that you initiate an upgrade before support ends.

Upgrading Deprecated DB Instances Before Support Ends

After a major version is deprecated, you can't install it on new DB instances. RDS will try to automatically upgrade all existing DB instances.

If you need to restore a deprecated DB instance, you can do a point-in-time restore (PITR) or restore a snapshot. Doing this gives you temporary access a DB instance that uses the version that is being deprecated. However, after a major version is fully deprecated, these DB instances will also be automatically upgraded to a supported version.

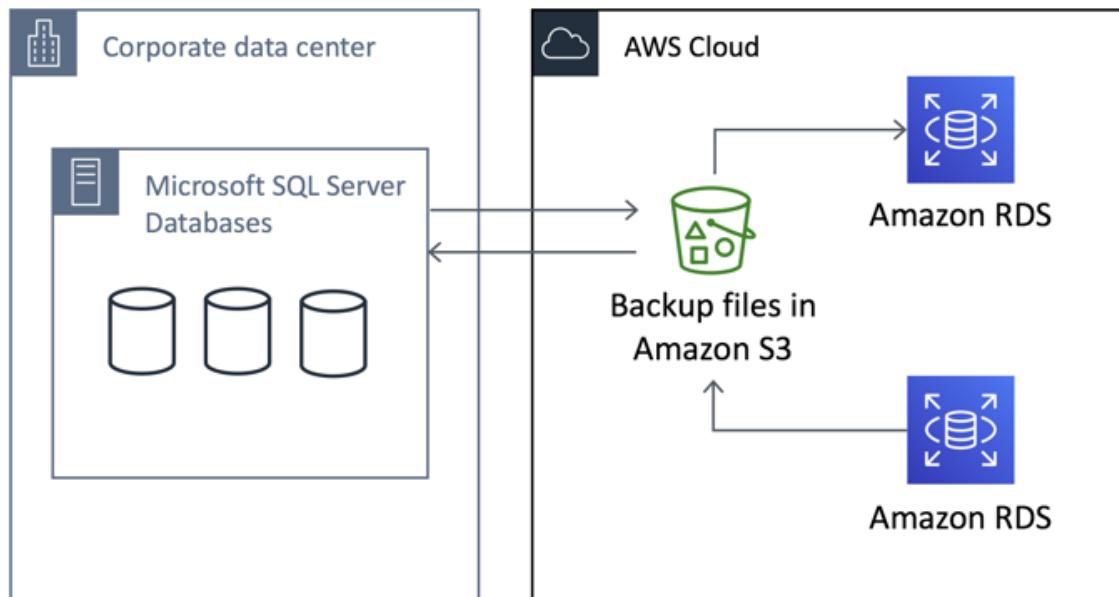
Importing and Exporting SQL Server Databases

Amazon RDS supports native backup and restore for Microsoft SQL Server databases using full backup files (.bak files). When you use RDS, you access files stored in Amazon S3 rather than using the local file system on the database server.

For example, you can create a full backup from your local server, store it on S3, and then restore it onto an existing Amazon RDS DB instance. You can also make backups from RDS, store them on S3, and then restore them wherever you want.

Native backup and restore is available in all AWS Regions, and for both Single-AZ and Multi-AZ DB instances. Native backup and restore is available for all editions of Microsoft SQL Server supported on Amazon RDS.

The following diagram shows the supported scenarios.



Using native .bak files to back up and restore databases is usually the fastest way to back up and restore databases. There are many additional advantages to using native backup and restore. For example, you can do the following:

- Migrate databases to or from Amazon RDS.
- Move databases between RDS SQL Server DB instances.
- Migrate data, schemas, stored procedures, triggers, and other database code inside .bak files.
- Backup and restore single databases, instead of entire DB instances.
- Create copies of databases for development, testing, training, and demonstrations.
- Store and transfer backup files with Amazon S3, for an added layer of protection for disaster recovery.

Limitations and Recommendations

The following are some limitations to using native backup and restore:

- You can't back up to, or restore from, an Amazon S3 bucket in a different AWS Region from your Amazon RDS DB instance.
- We strongly recommend that you don't restore backups from one time zone to a different time zone. If you restore backups from one time zone to a different time zone, you must audit your queries and applications for the effects of the time zone change.
- Native backups of databases larger than 1 TB aren't supported.
- You can't restore from more than 10 backup files at the same time.
- You can run up to two backup or restore tasks at the same time.
- You can't perform native log backups from SQL Server on Amazon RDS.
- RDS supports native restores of databases up to 16 TB. Native restores of databases on SQL Server Express are limited by the MSSQL edition to 10 GB or less.
- You can't do a native backup during the maintenance window, or any time Amazon RDS is in the process of taking a snapshot of the database.
- On Multi-AZ DB instances, you can only natively restore databases that are backed up in the full recovery model.
- Restoring from differential backups on Multi-AZ instances isn't supported.
- Calling the RDS procedures for native backup and restore within a transaction isn't supported.
- Use a symmetric AWS KMS customer master key (CMK) to encrypt your backups. Amazon RDS doesn't support asymmetric CMKs. For more information, see [Using Symmetric and Asymmetric Keys](#) in the *AWS Key Management Service Developer Guide*.
- Native backup files are encrypted with the specified AWS KMS key using the "Encryption-Only" crypto mode. When you are restoring encrypted backup files, be aware that they were encrypted with the "Encryption-Only" crypto mode.
- You can't restore a database that contains a FILESTREAM file group.

If your database can be offline while the backup file is created, copied, and restored, we recommend that you use native backup and restore to migrate it to RDS. If your on-premises database can't be offline, we recommend that you use the AWS Database Migration Service to migrate your database to Amazon RDS. For more information, see [What Is AWS Database Migration Service?](#)

Native backup and restore isn't intended to replace the data recovery capabilities of the cross-region snapshot copy feature. We recommend that you use snapshot copy to copy your database snapshot to another AWS Region for cross-region disaster recovery in Amazon RDS. For more information, see [Copying a Snapshot \(p. 269\)](#).

Setting Up for Native Backup and Restore

To set up for native backup and restore, you need three components:

1. An Amazon S3 bucket to store your backup files.

You must have an S3 bucket to use for your backup files and then upload backups you want to migrate to RDS. If you already have an Amazon S3 bucket, you can use that. If you don't, you can [create a bucket](#). Alternatively, you can choose to have a new bucket created for you when you add the `SOLSERVER_BACKUP_RESTORE` option by using the AWS Management Console.

For information on using S3, see the *Amazon Simple Storage Service Getting Started Guide* for a simple introduction. For more depth, see the *Amazon Simple Storage Service Console User Guide*.

2. An AWS Identity and Access Management (IAM) role to access the bucket.

If you already have an IAM role, you can use that. If you don't have an IAM role, you can create a new one manually. Alternatively, you can choose to have a new IAM role created for you when you add the `SOLSERVER_BACKUP_RESTORE` option by using the AWS Management Console.

If you want to create a new IAM role manually, take the approach discussed in the next section. Also, if you want to attach trust and permissions policies to an existing IAM role, take the approach discussed in the next section.

3. The `SQLSERVER_BACKUP_RESTORE` option added to an option group on your DB instance.

To enable native backup and restore on your DB instance, you add the `SQLSERVER_BACKUP_RESTORE` option to an option group on your DB instance. For more information and instructions, see [Support for Native Backup and Restore in SQL Server \(p. 712\)](#).

Manually Creating an IAM Role for Native Backup and Restore

If you want to manually create a new IAM role to use with native backup and restore, you can do so. In this case, you create a role to delegate permissions from the Amazon RDS service to your Amazon S3 bucket. When you create an IAM role, you attach trust and permissions policies. The trust policy allows RDS to assume this role. The permission policy defines the actions this Role can do. For more information about creating the role, see [Creating a Role to Delegate Permissions to an AWS Service](#).

For the native backup and restore feature, use trust and permissions policies similar to the examples in this section. In following example, we use the service principle name `rds.amazonaws.com` as an alias for all service accounts. In the other examples, we specify an Amazon Resource Name (ARN) to identify another account, user, or role that we're granting access to in the trust policy.

Example Trust Policy for Native Backup and Restore

```
{  
    "Version": "2012-10-17",  
    "Statement":  
    [ {  
        "Effect": "Allow",  
        "Principal": {"Service": "rds.amazonaws.com"},  
        "Action": "sts:AssumeRole"  
    }]  
}
```

The following example uses an ARN to specify a resource. For more information on using ARNs, see [Amazon Resource Names \(ARNs\)](#).

Example Permissions Policy for Native Backup and Restore Without Encryption Support

```
{  
    "Version": "2012-10-17",  
    "Statement":  
    [ {  
        "Effect": "Allow",  
        "Action":  
        [  
            "s3>ListBucket",  
            "s3:GetBucketLocation"  
        ],  
        "Resource": "arn:aws:s3:::bucket_name"  
    },  
    {  
        "Effect": "Allow",  
        "Action":  
        [  
            "s3GetObject",  
            "s3PutObject",  
            "s3DeleteObject"  
        ]  
    }]  
}
```

```
        "s3>ListMultipartUploadParts",
        "s3AbortMultipartUpload"
    ],
    "Resource": "arn:aws:s3:::bucket_name/*"
}
]
```

Example Permissions Policy for Native Backup and Restore with Encryption Support

If you want to encrypt your backup files, include an encryption key in your permissions policy. For more information about encryption keys, see [Getting Started](#) in the *AWS Key Management Service Developer Guide*.

Note

You must use a symmetric AWS KMS customer master key (CMK) to encrypt your backups. Amazon RDS doesn't support asymmetric CMKs. For more information, see [Using Symmetric and Asymmetric Keys](#) in the *AWS Key Management Service Developer Guide*.

```
{
    "Version": "2012-10-17",
    "Statement":
    [
        {
            "Effect": "Allow",
            "Action":
            [
                "kms:DescribeKey",
                "kms:GenerateDataKey",
                "kms:Encrypt",
                "kms:Decrypt"
            ],
            "Resource": "arn:aws:kms:region:account-id:key/key-id"
        },
        {
            "Effect": "Allow",
            "Action":
            [
                "s3>ListBucket",
                "s3:GetBucketLocation"
            ],
            "Resource": "arn:aws:s3:::bucket_name"
        },
        {
            "Effect": "Allow",
            "Action":
            [
                "s3GetObject",
                "s3PutObject",
                "s3ListMultipartUploadParts",
                "s3AbortMultipartUpload"
            ],
            "Resource": "arn:aws:s3:::bucket_name/*"
        }
    ]
}
```

Using Native Backup and Restore

After you have enabled and configured native backup and restore, you can start using it. First, you connect to your Microsoft SQL Server database, and then you call an Amazon RDS stored procedure to

do the work. For instructions on connecting to your database, see [Connecting to a DB Instance Running the Microsoft SQL Server Database Engine \(p. 630\)](#).

Some of the stored procedures require that you provide an Amazon Resource Name (ARN) to your Amazon S3 bucket and file. The format for your ARN is `arn:aws:s3:::bucket_name/file_name.extension`. Amazon S3 doesn't require an account number or AWS Region in ARNs.

If you also provide an optional AWS KMS encryption key, the format for the ARN of the key is `arn:aws:kms:region:account-id:key/key-id`. For more information, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#). You must use a symmetric AWS KMS customer master key (CMK) to encrypt your backups. Amazon RDS doesn't support asymmetric CMKs. For more information, see [Using Symmetric and Asymmetric Keys in the AWS Key Management Service Developer Guide](#).

For instructions on how to call each stored procedure, see the following topics:

- [Backing Up a Database \(p. 661\)](#)
- [Restoring a Database \(p. 662\)](#)
- [Restoring a Log \(p. 664\)](#)
- [Finishing a Database Restore \(p. 666\)](#)
- [Working with Partially Restored Databases \(p. 666\)](#)
- [Canceling a Task \(p. 667\)](#)
- [Tracking the Status of Tasks \(p. 667\)](#)

Backing Up a Database

To back up your database, call the `rds_backup_database` stored procedure.

Note

You can't back up a database during the maintenance window, or while Amazon RDS is taking a snapshot.

Usage

```
exec msdb.dbo.rds_backup_database
@source_db_name='database_name',
@s3_arn_to_backup_to='arn:aws:s3:::bucket_name/file_name.extension',
[@kms_master_key_arn='arn:aws:kms:region:account-id:key/key-id'],
[@overwrite_s3_backup_file=0/1],
[@type='DIFFERENTIAL/FULL'];
```

The following parameters are required:

- `@source_db_name` – The name of the database to back up.
- `@s3_arn_to_backup_to` – The ARN indicating the Amazon S3 bucket to use for the backup, plus the name of the backup file.

The file can have any extension, but `.bak` is usually used.

The following parameters are optional:

- `@kms_master_key_arn` – The ARN for the symmetric KMS customer master key (CMK) to use to encrypt the item. Amazon RDS doesn't support asymmetric CMKs.

If you don't specify a KMS key identifier, then RDS uses the default encryption key. Your AWS account has a different default encryption key for each AWS Region. For more information, see [Encrypting Amazon RDS Resources](#).

- `@overwrite_s3_backup_file` – A value that indicates whether to overwrite an existing backup file.
 - 0 – Doesn't overwrite an existing file. Setting this to 0 returns an error if the file already exists. This value is the default.
 - 1 – Overwrites an existing file that has the specified name, even if it isn't a backup file.
- `@type` – The type of backup.
 - `DIFFERENTIAL` – Makes a differential backup.
 - `FULL` – Makes a full backup. This value is the default.

A differential backup is based on the last full backup. For differential backups to work, you can't take a snapshot between the last full backup and the differential backup. If you want a differential backup, but a snapshot exists, then do another full backup before proceeding with the differential backup.

Examples

Example of Differential Backup

```
exec msdb.dbo.rds_backup_database
@source_db_name='mydatabase',
@s3_arn_to_backup_to='arn:aws:s3:::aws_example_bucket/backup1.bak',
@overwrite_s3_backup_file=1,
@type='DIFFERENTIAL';
```

Example of Full Backup with Encryption

```
exec msdb.dbo.rds_backup_database
@source_db_name='mydatabase',
@s3_arn_to_backup_to='arn:aws:s3:::aws_example_bucket/backup1.bak',
@kms_master_key_arn='arn:aws:kms:us-east-1:123456789012:key/AKIAIOSFODNN7EXAMPLE',
@overwrite_s3_backup_file=1,
@type='FULL';
```

You can look for the last full backup or snapshot using the following example SQL.

```
select top 1
database_name
, backup_start_date
, backup_finish_date
from msdb.dbo.backupset
where database_name='mydatabase'
and type = 'D'
order by backup_start_date desc;
```

Restoring a Database

To restore your database, call the `rds_restore_database` stored procedure. Amazon RDS creates an initial snapshot of the database after the restore task is complete and the database is open.

Usage

```
exec msdb.dbo.rds_restore_database
@restore_db_name='database_name',
```

```
@s3_arn_to_restore_from='arn:aws:s3:::bucket_name/file_name.extension',  
@with_norecovery=0/1,  
[@kms_master_key_arn='arn:aws:kms:region:account-id:key/key-id'],  
[@type='DIFFERENTIAL/FULL'];
```

The following parameters are required:

- `@restore_db_name` – The name of the database to restore.
- `@s3_arn_to_restore_from` – The ARN indicating the Amazon S3 prefix and names of the backup files used to restore the database.
 - For a single-file backup, provide the entire file name.
 - For a multiple-file backup, provide the prefix that the files have in common, then suffix that with an asterisk (*).

If `@s3_arn_to_restore_from` is empty, the following error message is returned: S3 ARN prefix cannot be empty.

The following parameter is required for differential restores, but optional for full restores:

- `@with_norecovery` – The recovery clause to use for the restore operation.
 - Set it to 0 to restore with RECOVERY. In this case, the database is online after the restore.
 - Set it to 1 to restore with NORECOVERY. In this case, the database remains in the RESTORING state after restore task completion. With this approach, you can do later differential restores.

For DIFFERENTIAL restores, specify 0 or 1. For FULL restores, this value defaults to 0.

The following parameters are optional:

- `@kms_master_key_arn` – If you encrypted the backup file, the KMS key to use to decrypt the file.
- `@type` – The type of restore. Valid types are DIFFERENTIAL and FULL. The default value is FULL.

Note

For differential restores, either the database must be in the RESTORING state or a task must already exist that restores with NORECOVERY.

You can't restore later differential backups while the database is online.

You can't submit a restore task for a database that already has a pending restore task with RECOVERY.

Full restores with NORECOVERY and differential restores aren't supported on Multi-AZ instances.

Examples

Example of Single-File Restore

```
exec msdb.dbo.rds_restore_database  
@restore_db_name='mydatabase',  
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/backup1.bak';
```

Example of Multiple-File Restore

To avoid errors when restoring multiple files, make sure that all the backup files have the same prefix, and that no other files use that prefix.

```
exec msdb.dbo.rds_restore_database  
@restore_db_name='mydatabase',
```

```
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/backup*';
```

Example of Full Database Restore with RECOVERY

The following three examples perform the same task, full restore with RECOVERY.

```
exec msdb.dbo.rds_restore_database
@restore_db_name='mydatabase',
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/backup1.bak';
```

```
exec msdb.dbo.rds_restore_database
@restore_db_name='mydatabase',
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/backup1.bak',
[@type='DIFFERENTIAL/FULL'];
```

```
exec msdb.dbo.rds_restore_database
@restore_db_name='mydatabase',
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/backup1.bak',
@type='FULL',
@with_norecovery=0;
```

Example of Full Database Restore with Encryption

```
exec msdb.dbo.rds_restore_database
@restore_db_name='mydatabase',
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/backup1.bak',
@kms_master_key_arn='arn:aws:kms:us-east-1:123456789012:key/AKIAIOSFODNN7EXAMPLE';
```

Example of Full Database Restore with NORECOVERY

```
exec msdb.dbo.rds_restore_database
@restore_db_name='mydatabase',
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/backup1.bak',
@type='FULL',
@with_norecovery=1;
```

Example of Differential Restore with NORECOVERY

```
exec msdb.dbo.rds_restore_database
@restore_db_name='mydatabase',
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/backup1.bak',
@type='DIFFERENTIAL',
@with_norecovery=1;
```

Example of Differential Restore with RECOVERY

```
exec msdb.dbo.rds_restore_database
@restore_db_name='mydatabase',
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/backup1.bak',
@type='DIFFERENTIAL',
@with_norecovery=0;
```

Restoring a Log

To restore your log, call the `rds_restore_log` stored procedure.

Usage

```
exec msdb.dbo.rds_restore_log
@restore_db_name='database_name',
@s3_arn_to_restore_from='arn:aws:s3:::bucket_name/log_file_name.extension',
[@kms_master_key_arn='arn:aws:kms:region:account-id:key/key-id'],
[@with_norecovery=0/1],
[@stopat='datetime'];
```

The following parameters are required:

- `@restore_db_name` – The name of the database whose log to restore.
- `@s3_arn_to_restore_from` – The ARN indicating the Amazon S3 prefix and name of the log file used to restore the log. The file can have any extension, but `.trn` is usually used.

If `@s3_arn_to_restore_from` is empty, the following error message is returned: S3 ARN prefix cannot be empty.

The following parameters are optional:

- `@kms_master_key_arn` – If you encrypted the log, the KMS key to use to decrypt the log.
- `@with_norecovery` – The recovery clause to use for the restore operation. This value defaults to 1.
 - Set it to 0 to restore with RECOVERY. In this case, the database is online after the restore. You can't restore further log backups while the database is online.
 - Set it to 1 to restore with NORECOVERY. In this case, the database remains in the RESTORING state after restore task completion. With this approach, you can do later log restores.
- `@stopat` – A value that specifies that the database is restored to its state at the date and time specified (in datetime format). Only transaction log records written before the specified date and time are applied to the database.

If this parameter isn't specified (it is NULL), the complete log is restored.

Note

For log restores, either the database must be in a state of restoring or a task must already exist that restores with NORECOVERY.

You can't restore log backups while the database is online.

You can't submit a log restore task on a database that already has a pending restore task with RECOVERY.

Log restores aren't supported on Multi-AZ instances.

Examples

Example of Log Restore

```
exec msdb.dbo.rds_restore_log
@restore_db_name='mydatabase',
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/mylog.trn';
```

Example of Log Restore with Encryption

```
exec msdb.dbo.rds_restore_log
@restore_db_name='mydatabase',
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/mylog.trn',
[@kms_master_key_arn='arn:aws:kms:us-east-1:123456789012:key/AKIAIOSFODNN7EXAMPLE';
```

Example of Log Restore with NORECOVERY

The following two examples perform the same task, log restore with NORECOVERY.

```
exec msdb.dbo.rds_restore_log
@restore_db_name='mydatabase',
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/mylog.trn',
@with_norecovery=1;
```

```
exec msdb.dbo.rds_restore_log
@restore_db_name='mydatabase',
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/mylog.trn';
```

Example of Log Restore with RECOVERY

```
exec msdb.dbo.rds_restore_log
@restore_db_name='mydatabase',
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/mylog.trn',
@with_norecovery=0;
```

Example of Log Restore with STOPAT Clause

```
exec msdb.dbo.rds_restore_log
@restore_db_name='mydatabase',
@s3_arn_to_restore_from='arn:aws:s3:::aws_example_bucket/mylog.trn',
@with_norecovery=0,
@stopat='2019-12-01 03:57:09';
```

Finishing a Database Restore

If the last restore task on the database was performed using `@with_norecovery=1`, the database is now in the RESTORING state. Open this database for normal operation by using the `rds_finish_restore` stored procedure.

Usage

```
exec msdb.dbo.rds_finish_restore @db_name='database_name';
```

Note

To use this approach, the database must be in the RESTORING state without any pending restore tasks.

The `rds_finish_restore` procedure isn't supported on Multi-AZ instances.

To finish restoring the database, use the master login. Or use the user login that most recently restored the database or log with NORECOVERY.

Working with Partially Restored Databases

Dropping a Partially Restored Database

To drop a partially restored database (left in the RESTORING state), use the `rds_drop_database` stored procedure.

```
exec msdb.dbo.rds_drop_database @db_name='database_name';
```

Note

You can't submit a DROP database request for a database that already has a pending restore or finish restore task.

To drop the database, use the master login. Or use the user login that most recently restored the database or log with NORECOVERY.

Snapshot Restore and Point-in-Time Recovery Behavior for Partially Restored Databases

Partially restored databases in the source instance (left in the RESTORING state) are dropped from the target instance during snapshot restore and point-in-time recovery.

Cancelling a Task

To cancel a backup or restore task, call the `rds_cancel_task` stored procedure.

Note

You can't cancel a FINISH_RESTORE task.

Usage

```
exec msdb.dbo.rds_cancel_task @task_id=ID_number;
```

The following parameter is required:

- `@task_id` – The ID of the task to cancel. You can get the task ID by calling `rds_task_status`.

Tracking the Status of Tasks

To track the status of your backup and restore tasks, call the `rds_task_status` stored procedure. If you don't provide any parameters, the stored procedure returns the status of all tasks. The status for tasks is updated approximately every two minutes. Task history is retained for 36 days.

Usage

```
exec msdb.dbo.rds_task_status
[@db_name='database_name'],
[@task_id=ID_number];
```

The following parameters are optional:

- `@db_name` – The name of the database to show the task status for.
- `@task_id` – The ID of the task to show the task status for.

Examples

Example of Listing the Status for a Specific Task

```
exec msdb.dbo.rds_task_status @task_id=5;
```

Example of Listing the Status for a Specific Database and Task

```
exec msdb.dbo.rds_task_status
```

```
@db_name='my_database',  
@task_id=5;
```

Example of Listing All Tasks and Their Statuses on a Specific Database

```
exec msdb.dbo.rds_task_status @db_name='my_database';
```

Example of Listing All Tasks and Their Statuses on the Current Instance

```
exec msdb.dbo.rds_task_status;
```

Response

The `rds_task_status` stored procedure returns the following columns.

Column	Description
<code>task_id</code>	The ID of the task.
<code>task_type</code>	Task type depending on the input parameters, as follows: <ul style="list-style-type: none">• For backup tasks:<ul style="list-style-type: none">• <code>BACKUP_DB</code> – Full database backup• <code>BACKUP_DB_DIFFERENTIAL</code> – Differential database backup• For restore tasks:<ul style="list-style-type: none">• <code>RESTORE_DB</code> – Full database restore with <code>RECOVERY</code>• <code>RESTORE_DB_NORECOVERY</code> – Full database restore with <code>NORECOVERY</code>• <code>RESTORE_DB_DIFFERENTIAL</code> – Differential database restore with <code>RECOVERY</code>• <code>RESTORE_DB_DIFFERENTIAL_NORECOVERY</code> – Differential database restore with <code>NORECOVERY</code>• <code>RESTORE_DB_LOG</code> – Log restore with <code>RECOVERY</code>• <code>RESTORE_DB_LOG_NORECOVERY</code> – Log restore with <code>NORECOVERY</code>• For tasks that finish a restore:<ul style="list-style-type: none">• <code>FINISH_RESTORE</code> – Finish restore and open database
	Amazon RDS creates an initial snapshot of the database after it is open on completion of the following restore tasks: <ul style="list-style-type: none">• <code>RESTORE_DB</code>• <code>RESTORE_DB_DIFFERENTIAL</code>• <code>RESTORE_DB_LOG</code>• <code>FINISH_RESTORE</code>
<code>database_name</code>	The name of the database that the task is associated with.
<code>% complete</code>	The progress of the task as a percent value.
<code>duration (mins)</code>	The amount of time spent on the task, in minutes.
<code>lifecycle</code>	The status of the task. The possible statuses are the following:

Column	Description
	<ul style="list-style-type: none"> • CREATED – As soon as you call <code>rds_backup_database</code> or <code>rds_restore_database</code>, a task is created and the status is set to CREATED. • IN_PROGRESS – After a backup or restore task starts, the status is set to IN_PROGRESS. It can take up to 5 minutes for the status to change from CREATED to IN_PROGRESS. • SUCCESS – After a backup or restore task completes, the status is set to SUCCESS. • ERROR – If a backup or restore task fails, the status is set to ERROR. For more information about the error, see the <code>task_info</code> column. • CANCEL_REQUESTED – As soon as you call <code>rds_cancel_task</code>, the status of the task is set to CANCEL_REQUESTED. • CANCELLED – After a task is successfully canceled, the status of the task is set to CANCELLED.
<code>task_info</code>	<p>Additional information about the task.</p> <p>If an error occurs while backing up or restoring a database, this column contains information about the error. For a list of possible errors, and mitigation strategies, see Troubleshooting (p. 670).</p>
<code>last_updated</code>	The date and time that the task status was last updated. The status is updated after every 5 percent of progress.
<code>created_at</code>	The date and time that the task was created.
<code>S3_object_arn</code>	The ARN indicating the Amazon S3 prefix and the name of the file that is being backed up or restored.
<code>overwrite_s3_backup_file</code>	The value of the <code>@overwrite_s3_backup_file</code> parameter specified when calling a backup task. For more information, see Backing Up a Database (p. 661) .
<code>KMS_master_key_arn</code>	The ARN for the KMS customer master key used for encryption (for backup) and decryption (for restore).
<code>filepath</code>	Not applicable to native backup and restore tasks.
<code>overwrite_file</code>	Not applicable to native backup and restore tasks.

Compressing Backup Files

To save space in your Amazon S3 bucket, you can compress your backup files. For more information about compressing backup files, see [Backup Compression](#) in the Microsoft documentation.

Compressing your backup files is supported for the following database editions:

- Microsoft SQL Server Enterprise Edition
- Microsoft SQL Server Standard Edition

To turn on compression for your backup files, run the following code:

```
exec rdsadmin..rds_set_configuration 'S3 backup compression', 'true';
```

To turn off compression for your backup files, run the following code:

```
exec rdsadmin..rds_set_configuration 'S3 backup compression', 'false';
```

Troubleshooting

The following are issues you might encounter when you use native backup and restore.

Issue	Troubleshooting Suggestions
Access Denied	<p>The backup or restore process is unable to access the backup file. This is usually caused by issues like the following:</p> <ul style="list-style-type: none"> Referencing the incorrect bucket. Referencing the bucket using an incorrect format. Referencing a file name without using the ARN. Incorrect permissions on the bucket file. For example, if it is created by a different account that is trying to access it now, add the correct permissions. An IAM policy that is incorrect or incomplete. Your IAM role must include all the necessary elements, including, for example, the correct version. These are highlighted in Importing and Exporting SQL Server Databases (p. 657).
BACKUP DATABASE WITH COMPRESSION isn't supported on <edition_name> Edition	<p>Compressing your backup files is only supported for Microsoft SQL Server Enterprise Edition and Standard Edition.</p> <p>For more information, see Compressing Backup Files (p. 669).</p>
Key <ARN> does not exist	<p>You attempted to restore an encrypted backup, but didn't provide a valid encryption key. Check your encryption key and retry.</p> <p>For more information, see Restoring a Database (p. 662).</p>
Please reissue task with correct type and overwrite property	<p>If you attempt to back up your database and provide the name of a file that already exists, but set the overwrite property to false, the save operation fails. To fix this error, either provide the name of a file that doesn't already exist, or set the overwrite property to true.</p> <p>For more information, see Backing Up a Database (p. 661).</p> <p>It's also possible that you intended to restore your database, but called the <code>rds_backup_database</code> stored procedure accidentally. In that case, call the <code>rds_restore_database</code> stored procedure instead.</p> <p>For more information, see Restoring a Database (p. 662).</p> <p>If you intended to restore your database and called the <code>rds_restore_database</code> stored procedure, make sure that you provided the name of a valid backup file.</p> <p>For more information, see Using Native Backup and Restore (p. 660).</p>
Please specify a bucket that is in the same region as RDS instance	<p>You can't back up to, or restore from, an Amazon S3 bucket in a different AWS Region than your Amazon RDS DB instance. You can use Amazon S3 replication to copy the backup file to the correct AWS Region.</p> <p>For more information, see Cross-Region Replication in the Amazon S3 documentation.</p>

Issue	Troubleshooting Suggestions
The specified bucket does not exist	<p>Verify that you have provided the correct ARN for your bucket and file, in the correct format.</p> <p>For more information, see Using Native Backup and Restore (p. 660).</p>
User <ARN> is not authorized to perform <kms action> on resource <ARN>	<p>You requested an encrypted operation, but didn't provide correct AWS KMS permissions. Verify that you have the correct permissions, or add them.</p> <p>For more information, see Setting Up for Native Backup and Restore (p. 658).</p>
The Restore task is unable to restore from more than 10 backup file(s). Please reduce the number of files matched and try again.	Reduce the number of files you're trying to restore from. You can make each individual file larger if necessary.

Importing and Exporting SQL Server Data Using Other Methods

Following, you can find information about using snapshots to import your Microsoft SQL Server data to Amazon RDS. You can also find information about using snapshots to export your data from an RDS DB instance running SQL Server.

If your scenario supports it, it's easier to move data in and out of Amazon RDS by using the native backup and restore functionality. For more information, see [Importing and Exporting SQL Server Databases \(p. 657\)](#).

Note

Amazon RDS for Microsoft SQL Server does not support importing data into the `msdb` database.

Importing Data into SQL Server on Amazon RDS by Using a Snapshot

To import data into a SQL Server DB instance by using a snapshot

1. Create a DB instance. For more information, see [Creating a DB Instance Running the Microsoft SQL Server Database Engine \(p. 616\)](#).
2. Stop applications from accessing the destination DB instance.

If you prevent access to your DB instance while you are importing data, data transfer is faster. Additionally, you don't need to worry about conflicts while data is being loaded if other applications cannot write to the DB instance at the same time. If something goes wrong and you have to roll back to an earlier database snapshot, the only changes that you lose are the imported data. You can import this data again after you resolve the issue.

For information about controlling access to your DB instance, see [Working with DB Security Groups \(EC2-Classic Platform\) \(p. 491\)](#).

3. Create a snapshot of the target database.

If the target database is already populated with data, we recommend that you take a snapshot of the database before you import the data. If something goes wrong with the data import or you want to discard the changes, you can restore the database to its previous state by using the snapshot. For information about database snapshots, see [Creating a DB Snapshot \(p. 264\)](#).

Note

When you take a database snapshot, I/O operations to the database are suspended for a moment (milliseconds) while the backup is in progress.

4. Disable automated backups on the target database.

Disabling automated backups on the target DB instance improves performance while you are importing your data because Amazon RDS doesn't log transactions when automatic backups are disabled. However, there are some things to consider. Automated backups are required to perform a point-in-time recovery. Thus, you can't restore the database to a specific point in time while you are importing data. Additionally, any automated backups that were created on the DB instance are erased unless you choose to retain them.

Choosing to retain the automated backups can help protect you against accidental deletion of data. Amazon RDS also saves the database instance properties along with each automated backup to make it easy to recover. Using this option lets you can restore a deleted database instance to a specified point in time within the backup retention period even after deleting it. Automated backups are automatically deleted at the end of the specified backup window, just as they are for an active database instance.

You can also use previous snapshots to recover the database, and any snapshots that you have taken remain available. For information about automated backups, see [Working With Backups \(p. 254\)](#).

5. Disable foreign key constraints, if applicable.

If you need to disable foreign key constraints, you can do so with the following script.

```
--Disable foreign keys on all tables
DECLARE @table_name SYSNAME;
DECLARE @cmd NVARCHAR(MAX);
DECLARE table_cursor CURSOR FOR SELECT name FROM sys.tables;

OPEN table_cursor;
FETCH NEXT FROM table_cursor INTO @table_name;

WHILE @@FETCH_STATUS = 0 BEGIN
    SELECT @cmd = 'ALTER TABLE '+QUOTENAME(@table_name)+' NOCHECK CONSTRAINT ALL';
    EXEC (@cmd);
    FETCH NEXT FROM table_cursor INTO @table_name;
END

CLOSE table_cursor;
DEALLOCATE table_cursor;

GO
```

6. Drop indexes, if applicable.
7. Disable triggers, if applicable.

If you need to disable triggers, you can do so with the following script.

```
--Disable triggers on all tables
DECLARE @enable BIT = 0;
DECLARE @trigger SYSNAME;
DECLARE @table SYSNAME;
DECLARE @cmd NVARCHAR(MAX);
DECLARE trigger_cursor CURSOR FOR SELECT trigger_object.name trigger_name,
    table_object.name table_name
FROM sysobjects trigger_object
JOIN sysobjects table_object ON trigger_object.parent_obj = table_object.id
WHERE trigger_object.type = 'TR';

OPEN trigger_cursor;
FETCH NEXT FROM trigger_cursor INTO @trigger, @table;

WHILE @@FETCH_STATUS = 0 BEGIN
    IF @enable = 1
        SET @cmd = 'ENABLE ';
    ELSE
        SET @cmd = 'DISABLE ';

    SET @cmd = @cmd + ' TRIGGER dbo.'+QUOTENAME(@trigger)+'' ON
dbo.'+QUOTENAME(@table)+'' ;
    EXEC (@cmd);
    FETCH NEXT FROM trigger_cursor INTO @trigger, @table;
END

CLOSE trigger_cursor;
DEALLOCATE trigger_cursor;

GO
```

8. Query the source SQL Server instance for any logins that you want to import to the destination DB instance.

SQL Server stores logins and passwords in the `master` database. Because Amazon RDS doesn't grant access to the `master` database, you cannot directly import logins and passwords into your destination DB instance. Instead, you must query the `master` database on the source SQL Server instance to generate a data definition language (DDL) file. This file should include all logins and passwords that you want to add to the destination DB instance. This file also should include role memberships and permissions that you want to transfer.

For information about querying the `master` database, see [How to Transfer the Logins and the Passwords Between Instances of SQL Server 2005 and SQL Server 2008](#) in the Microsoft Knowledge Base.

The output of the script is another script that you can run on the destination DB instance. The script in the Knowledge Base article has the following code:

```
p.type IN
```

Every place `p.type` appears, use the following code instead:

```
p.type = 'S'
```

9. Import the data using the method in [Import the Data \(p. 675\)](#).
10. Grant applications access to the target DB instance.

When your data import is complete, you can grant access to the DB instance to those applications that you blocked during the import. For information about controlling access to your DB instance, see [Working with DB Security Groups \(EC2-Classic Platform\) \(p. 491\)](#).

11. Enable automated backups on the target DB instance.

For information about automated backups, see [Working With Backups \(p. 254\)](#).

12. Enable foreign key constraints.

If you disabled foreign key constraints earlier, you can now enable them with the following script.

```
--Enable foreign keys on all tables
DECLARE @table_name SYSNAME;
DECLARE @cmd NVARCHAR(MAX);
DECLARE table_cursor CURSOR FOR SELECT name FROM sys.tables;

OPEN table_cursor;
FETCH NEXT FROM table_cursor INTO @table_name;

WHILE @@FETCH_STATUS = 0 BEGIN
    SELECT @cmd = 'ALTER TABLE '+QUOTENAME(@table_name)+' CHECK CONSTRAINT ALL';
    EXEC (@cmd);
    FETCH NEXT FROM table_cursor INTO @table_name;
END

CLOSE table_cursor;
DEALLOCATE table_cursor;
```

13. Enable indexes, if applicable.
14. Enable triggers, if applicable.

If you disabled triggers earlier, you can now enable them with the following script.

```
--Enable triggers on all tables
DECLARE @enable BIT = 1;
DECLARE @trigger SYSNAME;
DECLARE @table SYSNAME;
DECLARE @cmd NVARCHAR(MAX);
DECLARE trigger_cursor CURSOR FOR SELECT trigger_object.name trigger_name,
    table_object.name table_name
FROM sysobjects trigger_object
JOIN sysobjects table_object ON trigger_object.parent_obj = table_object.id
WHERE trigger_object.type = 'TR';

OPEN trigger_cursor;
FETCH NEXT FROM trigger_cursor INTO @trigger, @table;

WHILE @@FETCH_STATUS = 0 BEGIN
    IF @enable = 1
        SET @cmd = 'ENABLE ';
    ELSE
        SET @cmd = 'DISABLE ';

    SET @cmd = @cmd + ' TRIGGER dbo.'+QUOTENAME(@trigger)+' ON
dbo.'+QUOTENAME(@table)+';
    EXEC (@cmd);
    FETCH NEXT FROM trigger_cursor INTO @trigger, @table;
END

CLOSE trigger_cursor;
DEALLOCATE trigger_cursor;
```

Import the Data

Microsoft SQL Server Management Studio is a graphical SQL Server client that is included in all Microsoft SQL Server editions except the Express Edition. SQL Server Management Studio Express is available from Microsoft as a free download. To find this download, see [the Microsoft website](#).

Note

SQL Server Management Studio is available only as a Windows-based application.

SQL Server Management Studio includes the following tools, which are useful in importing data to a SQL Server DB instance:

- Generate and Publish Scripts wizard
- Import and Export wizard
- Bulk copy

Generate and Publish Scripts Wizard

The Generate and Publish Scripts wizard creates a script that contains the schema of a database, the data itself, or both. You can generate a script for a database in your local SQL Server deployment. You can then run the script to transfer the information that it contains to an Amazon RDS DB instance.

Note

For databases of 1 GiB or larger, it's more efficient to script only the database schema. You then use the Import and Export wizard or the bulk copy feature of SQL Server to transfer the data.

For detailed information about the Generate and Publish Scripts wizard, see the [Microsoft SQL Server documentation](#).

In the wizard, pay particular attention to the advanced options on the **Set Scripting Options** page to ensure that everything you want your script to include is selected. For example, by default, database triggers are not included in the script.

When the script is generated and saved, you can use SQL Server Management Studio to connect to your DB instance and then run the script.

Import and Export Wizard

The Import and Export wizard creates a special Integration Services package, which you can use to copy data from your local SQL Server database to the destination DB instance. The wizard can filter which tables and even which tuples within a table are copied to the destination DB instance.

Note

The Import and Export wizard works well for large datasets, but it might not be the fastest way to remotely export data from your local deployment. For an even faster way, consider the SQL Server bulk copy feature.

For detailed information about the Import and Export wizard, see the [Microsoft SQL Server documentation](#).

In the wizard, on the **Choose a Destination** page, do the following:

- For **Server Name**, type the name of the endpoint for your DB instance.
- For the server authentication mode, choose **Use SQL Server Authentication**.
- For **User name** and **Password**, type the credentials for the master user that you created for the DB instance.

Bulk Copy

The SQL Server bulk copy feature is an efficient means of copying data from a source database to your DB instance. Bulk copy writes the data that you specify to a data file, such as an ASCII file. You can then run bulk copy again to write the contents of the file to the destination DB instance.

This section uses the **bcp** utility, which is included with all editions of SQL Server. For detailed information about bulk import and export operations, see [the Microsoft SQL Server documentation](#).

Note

Before you use bulk copy, you must first import your database schema to the destination DB instance. The Generate and Publish Scripts wizard, described earlier in this topic, is an excellent tool for this purpose.

The following command connects to the local SQL Server instance. It generates a tab-delimited file of a specified table in the C:\ root directory of your existing SQL Server deployment. The table is specified by its fully qualified name, and the text file has the same name as the table that is being copied.

```
bcp dbname.schema_name.table_name out C:\table_name.txt -n -S localhost -U username -P password -b 10000
```

The preceding code includes the following options:

- **-n** specifies that the bulk copy uses the native data types of the data to be copied.
- **-S** specifies the SQL Server instance that the *bcp* utility connects to.
- **-U** specifies the user name of the account to log in to the SQL Server instance.
- **-P** specifies the password for the user specified by **-U**.

- **-b** specifies the number of rows per batch of imported data.

Note

There might be other parameters that are important to your import situation. For example, you might need the **-E** parameter that pertains to identity values. For more information; see the full description of the command line syntax for the **bcp** utility in the [Microsoft SQL Server documentation](#).

For example, suppose that a database named **store** that uses the default schema, **dbo**, contains a table named **customers**. The user account **admin**, with the password **insecure**, copies 10,000 rows of the **customers** table to a file named **customers.txt**.

```
bcp store.dbo.customers out C:\customers.txt -n -S localhost -U admin -P insecure -b 10000
```

After you generate the data file, you can upload the data to your DB instance by using a similar command. Beforehand, create the database and schema on the target DB instance. Then use the **in** argument to specify an input file instead of **out** to specify an output file. Instead of using **localhost** to specify the local SQL Server instance, specify the endpoint of your DB instance. If you use a port other than 1433, specify that too. The user name and password are the master user and password for your DB instance. The syntax is as follows.

```
bcp dbname.schema_name.table_name in C:\table_name.txt -n -S endpoint,port -  
U master_user_name -P master_user_password -b 10000
```

To continue the previous example, suppose that the master user name is **admin**, and the password is **insecure**. The endpoint for the DB instance is **rds.ckz2kqd4qsn1.us-east-1.rds.amazonaws.com**, and you use port 4080. The command is as follows.

```
bcp store.dbo.customers in C:\customers.txt -n -S rds.ckz2kqd4qsn1.us-east-1.rds.amazonaws.com,4080 -U admin -P insecure -b 10000
```

Exporting Data from SQL Server on Amazon RDS

You can choose one of the following options to export data from an RDS SQL Server DB instance:

- **Native database backup using a full backup file (.bak)** – Using .bak files to backup databases is heavily optimized, and is usually the fastest way to export data. For more information, see [Importing and Exporting SQL Server Databases \(p. 657\)](#).
- **SQL Server Import and Export wizard** – For more information, see [SQL Server Import and Export Wizard \(p. 677\)](#).
- **SQL Server Generate and Publish Scripts wizard and bcp utility** – For more information, see [SQL Server Generate and Publish Scripts Wizard and bcp Utility \(p. 679\)](#).

SQL Server Import and Export Wizard

You can use the SQL Server Import and Export wizard to copy one or more tables, views, or queries from your RDS SQL Server DB instance to another data store. This choice is best if the target data store is not SQL Server. For more information, see [SQL Server Import and Export Wizard](#) in the SQL Server documentation.

The SQL Server Import and Export wizard is available as part of Microsoft SQL Server Management Studio. This graphical SQL Server client is included in all Microsoft SQL Server editions except the Express Edition. SQL Server Management Studio is available only as a Windows-based application.

SQL Server Management Studio Express is available from Microsoft as a free download. To find this download, see [the Microsoft website](#).

To use the SQL Server Import and Export wizard to export data

1. In SQL Server Management Studio, connect to your RDS SQL Server DB instance. For details on how to do this, see [Connecting to a DB Instance Running the Microsoft SQL Server Database Engine \(p. 630\)](#).
2. In **Object Explorer**, expand **Databases**, open the context (right-click) menu for the source database, choose **Tasks**, and then choose **Export Data**. The wizard appears.
3. On the **Choose a Data Source** page, do the following:
 - a. For **Data source**, choose **SQL Server Native Client 11.0**.
 - b. Verify that the **Server name** box shows the endpoint of your RDS SQL Server DB instance.
 - c. Select **Use SQL Server Authentication**. For **User name** and **Password**, type the master user name and password of your Amazon RDS SQL DB.
 - d. Verify that the **Database** box shows the database from which you want to export data.
 - e. Choose **Next**.
4. On the **Choose a Destination** page, do the following:
 - a. For **Destination**, choose **SQL Server Native Client 11.0**.

Note

Other target data sources are available. These include .NET Framework data providers, OLE DB providers, SQL Server Native Client providers, ADO.NET providers, Microsoft Office Excel, Microsoft Office Access, and the Flat File source. If you choose to target one of these data sources, skip the remainder of step 4. For details on the connection information to provide next, see [Choose a Destination](#) in the SQL Server documentation.

- b. For **Server name**, type the server name of the target SQL Server DB instance.
- c. Choose the appropriate authentication type. Type a user name and password if necessary.
- d. For **Database**, choose the name of the target database, or choose **New** to create a new database to contain the exported data.

If you choose **New**, see [Create Database](#) in the SQL Server documentation for details on the database information to provide.

- e. Choose **Next**.
5. On the **Table Copy or Query** page, choose **Copy data from one or more tables or views** or **Write a query to specify the data to transfer**. Choose **Next**.
6. If you chose **Write a query to specify the data to transfer**, you see the **Provide a Source Query** page. Type or paste in a SQL query, and then choose **Parse** to verify it. Once the query validates, choose **Next**.
7. On the **Select Source Tables and Views** page, do the following:
 - a. Select the tables and views that you want to export, or verify that the query you provided is selected.
 - b. Choose **Edit Mappings** and specify database and column mapping information. For more information, see [Column Mappings](#) in the SQL Server documentation.
 - c. (Optional) To see a preview of data to be exported, select the table, view, or query, and then choose **Preview**.
 - d. Choose **Next**.
8. On the **Run Package** page, verify that **Run immediately** is selected. Choose **Next**.
9. On the **Complete the Wizard** page, verify that the data export details are as you expect. Choose **Finish**.

10. On the **The execution was successful** page, choose **Close**.

SQL Server Generate and Publish Scripts Wizard and bcp Utility

You can use the SQL Server Generate and Publish Scripts wizard to create scripts for an entire database or just selected objects. You can run these scripts on a target SQL Server DB instance to recreate the scripted objects. You can then use the bcp utility to bulk export the data for the selected objects to the target DB instance. This choice is best if you want to move a whole database (including objects other than tables) or large quantities of data between two SQL Server DB instances. For a full description of the bcp command-line syntax, see [bcp Utility](#) in the Microsoft SQL Server documentation.

The SQL Server Generate and Publish Scripts wizard is available as part of Microsoft SQL Server Management Studio. This graphical SQL Server client is included in all Microsoft SQL Server editions except the Express Edition. SQL Server Management Studio is available only as a Windows-based application. SQL Server Management Studio Express is available from Microsoft as a [free download](#).

To use the SQL Server Generate and Publish Scripts wizard and the bcp utility to export data

1. In SQL Server Management Studio, connect to your RDS SQL Server DB instance. For details on how to do this, see [Connecting to a DB Instance Running the Microsoft SQL Server Database Engine \(p. 630\)](#).
2. In **Object Explorer**, expand the **Databases** node and select the database you want to script.
3. Follow the instructions in [Generate and Publish Scripts Wizard](#) in the SQL Server documentation to create a script file.
4. In SQL Server Management Studio, connect to your target SQL Server DB instance.
5. With the target SQL Server DB instance selected in **Object Explorer**, choose **Open** on the **File** menu, choose **File**, and then open the script file.
6. If you have scripted the entire database, review the CREATE DATABASE statement in the script. Make sure that the database is being created in the location and with the parameters that you want. For more information, see [CREATE DATABASE](#) in the SQL Server documentation.
7. If you are creating database users in the script, check to see if server logins exist on the target DB instance for those users. If not, create logins for those users; the scripted commands to create the database users fail otherwise. For more information, see [Create a Login](#) in the SQL Server documentation.
8. Choose **!Execute** on the SQL Editor menu to execute the script file and create the database objects. When the script finishes, verify that all database objects exist as expected.
9. Use the bcp utility to export data from the RDS SQL Server DB instance into files. Open a command prompt and type the following command.

```
bcp database_name.schema_name.table_name out data_file -n -S aws_rds_sql_endpoint -U
username -P password
```

The preceding code includes the following options:

- *table_name* is the name of one of the tables that you've recreated in the target database and now want to populate with data.
- *data_file* is the full path and name of the data file to be created.
- *-n* specifies that the bulk copy uses the native data types of the data to be copied.
- *-S* specifies the SQL Server DB instance to export from.
- *-U* specifies the user name to use when connecting to the SQL Server DB instance.
- *-P* specifies the password for the user specified by *-U*.

The following shows an example command.

```
bcp world.dbo.city out C:\Users\JohnDoe\city.dat -n -S sql-jdoe.1234abcd.us-west-2.rds.amazonaws.com,1433 -U JohnDoe -P ClearTextPassword
```

Repeat this step until you have data files for all of the tables you want to export.

10. Prepare your target DB instance for bulk import of data by following the instructions at [Basic Guidelines for Bulk Importing Data](#) in the SQL Server documentation.
11. Decide on a bulk import method to use after considering performance and other concerns discussed in [About Bulk Import and Bulk Export Operations](#) in the SQL Server documentation.
12. Bulk import the data from the data files that you created using the bcp utility. To do so, follow the instructions at either [Import and Export Bulk Data by Using the bcp Utility](#) or [Import Bulk Data by Using BULK INSERT or OPENROWSET\(BULK...\)](#) in the SQL Server documentation, depending on what you decided in step 11.

Multi-AZ Deployments for Microsoft SQL Server

Multi-AZ deployments provide increased availability, data durability, and fault tolerance for DB instances. In the event of planned database maintenance or unplanned service disruption, Amazon RDS automatically fails over to the up-to-date secondary DB instance. This functionality lets database operations resume quickly without manual intervention. The primary and standby instances use the same endpoint, whose physical network address transitions to the secondary replica as part of the failover process. You don't have to reconfigure your application when a failover occurs.

Amazon RDS supports Multi-AZ deployments for Microsoft SQL Server by using either SQL Server Database Mirroring (DBM) or Always On Availability Groups (AGs). Amazon RDS monitors and maintains the health of your Multi-AZ deployment. If problems occur, RDS automatically repairs unhealthy DB instances, reestablishes synchronization, and initiates failovers. Failover only occurs if the standby and primary are fully in sync. You don't have to manage anything.

When you set up SQL Server Multi-AZ, RDS automatically configures all databases on the instance to use DBM or AGs. Amazon RDS handles the primary, the witness, and the secondary DB instance for you. Because configuration is automatic, RDS selects DBM or Always On AGs based on the version of SQL Server that you deploy.

Amazon RDS supports Multi-AZ with Always On AGs for the following SQL Server versions and editions:

- SQL Server 2017: Enterprise Edition 14.00.3049.1 or later
- SQL Server 2016: Enterprise Edition 13.00.5216.0 or later

Amazon RDS supports Multi-AZ with DBM for the following SQL Server versions and editions, except for the versions of Enterprise Edition noted previously:

- SQL Server 2017: Standard and Enterprise Editions
- SQL Server 2016: Standard and Enterprise Editions
- SQL Server 2014: Standard and Enterprise Editions
- SQL Server 2012: Standard and Enterprise Editions

Amazon RDS supports Multi-AZ for SQL Server in all AWS Regions, with the following exceptions:

- US West (N. California): Neither DBM nor Always On AGs are supported here
- Asia Pacific (Osaka-Local): Neither DBM nor Always On AGs are supported here
- Asia Pacific (Sydney): Supported for [DB instances in VPCs](#)
- Asia Pacific (Tokyo): Supported for [DB instances in VPCs](#)
- South America (São Paulo): Supported on all [DB instance classes](#) except m1 and m2

Adding Multi-AZ to a Microsoft SQL Server DB Instance

When you create a new SQL Server DB instance using the AWS Management Console, you can add Multi-AZ with Database Mirroring (DBM) or Always On AGs. You do so by choosing **Yes (Mirroring / Always On)** from the **Multi-AZ Deployment** list on the **Specify DB Details** page. For more information, see [Creating a DB Instance Running the Microsoft SQL Server Database Engine \(p. 616\)](#).

When you modify an existing SQL Server DB instance using the AWS Management Console, you can add Multi-AZ with DBM or AGs by choosing **Yes (Mirroring / Always On)** from the **Multi-AZ Deployment**

list on the **Modify DB Instance** page. For more information, see [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#).

Microsoft SQL Server Multi-AZ Deployment Notes and Recommendations

The following are some restrictions when working with Multi-AZ deployments for Microsoft SQL Server DB instances:

- Cross-region Multi-AZ is not currently supported.
- You can't configure the secondary to accept database read activity.
- Multi-AZ with Always On Availability Groups (AGs) supports in-memory optimization.
- You can't rename a database on a SQL Server DB instance that is in a SQL Server Multi-AZ deployment. If you need to rename a database on such an instance, first turn off Multi-AZ for the DB instance, then rename the database. Finally, turn Multi-AZ back on for the DB instance.
- You can only restore Multi-AZ DB instances that are backed up using the full recovery model.

The following are some notes about working with Multi-AZ deployments for Microsoft SQL Server DB instances:

- Amazon RDS exposes the Always On AGs [availability group listener endpoint](#). The endpoint is visible in the console, and is returned by the `DescribeDBInstances` API as an entry in the endpoints field.
- Amazon RDS supports [availability group multisubnet failovers](#).
- To use SQL Server Multi-AZ with a SQL Server DB instance in a VPC, first create a DB subnet group that has subnets in at least two distinct Availability Zones. Then assign the DB subnet group to the primary replica of the SQL Server DB instance.
- When a DB instance is modified to be a Multi-AZ deployment, during the modification it has a status of [modifying](#). Amazon RDS creates the standby, and makes a backup of the primary DB instance. After the process is complete, the status of the primary DB instance becomes [available](#).
- Multi-AZ deployments maintain all databases on the same node. If a database on the primary host fails over, all your SQL Server databases fail over as one atomic unit to your standby host. Amazon RDS provisions a new healthy host, and replaces the unhealthy host.
- Multi-AZ with DBS or AGs supports a single standby replica.
- Users, logins, and permissions are automatically replicated for you on the secondary. You don't need to recreate them. User-defined server roles (a SQL Server 2012 feature) are only replicated in Multi-AZ instances for AGs instances.
- If you have SQL Server Agent jobs, recreate them on the secondary. You do so because these jobs are stored in the msdb database, and you can't replicate this database by using Database Mirroring (DBM) or Always On Availability Groups (AGs). Create the jobs first in the original primary, then fail over, and create the same jobs in the new primary.
- You might observe elevated latencies compared to a standard DB instance deployment (in a single Availability Zone) because of the synchronous data replication.
- Failover times are affected by the time it takes to complete the recovery process. Large transactions increase the failover time.

The following are some recommendations for working with Multi-AZ deployments for Microsoft SQL Server DB instances:

- For databases used in production or preproduction, we recommend the following options:
 - Multi-AZ deployments for high availability

- "Provisioned IOPS" for fast, consistent performance
- "Memory optimized" rather than "General purpose"
- You can't select the Availability Zone (AZ) for the secondary instance, so when you deploy application hosts, take this into account. Your database might fail over to another AZ, and the application hosts might not be in the same AZ as the database. For this reason, we recommend that you balance your application hosts across all AZs in the given AWS Region.
- For best performance, don't enable Database Mirroring or Always On AGs during a large data load operation. If you want your data load to be as fast as possible, finish loading data before you convert your DB instance to a Multi-AZ deployment.
- Applications that access the SQL Server databases should have exception handling that catches connection errors. The following code sample shows a try/catch block that catches a communication error.

```

for (int iRetryCount = 0; (iRetryCount < RetryMaxAttempts && keepInserting); iRetryCount++)
{
    using (SqlConnection connection = new SqlConnection(DatabaseConnString))
    {
        using (SqlCommand command = connection.CreateCommand())
        {
            command.CommandText = "INSERT INTO SOME_TABLE VALUES ('SomeValue');";

            try
            {
                connection.Open();

                while (keepInserting)
                {
                    command.ExecuteNonQuery();
                    intervalCount++;
                }
                connection.Close();
            }

            catch (Exception ex)
            {
                Logger(ex.Message);
            }
        }
    }

    if (iRetryCount < RetryMaxAttempts && keepInserting)
    {
        Thread.Sleep(RetryIntervalPeriodInSeconds * 1000);
    }
}

```

- Don't use the `Set Partner Off` command when working with Multi-AZ instances. For example, don't do the following.

```
--Don't do this
ALTER DATABASE db1 SET PARTNER off
```

- Don't set the recovery mode to `simple`. For example, don't do the following.

```
--Don't do this
ALTER DATABASE db1 SET RECOVERY simple
```

- Don't use the `DEFAULT_DATABASE` parameter when creating new logins on Multi-AZ DB instances, because these settings can't be applied to the standby mirror. For example, don't do the following.

```
--Don't do this
CREATE LOGIN [test_dba] WITH PASSWORD=foo, DEFAULT_DATABASE=[db2]
```

Also, don't do the following.

```
--Don't do this
ALTER LOGIN [test_dba] SET DEFAULT_DATABASE=[db3]
```

Determining the Location of the Secondary

You can determine the location of the secondary replica by using the AWS Management Console. You need to know the location of the secondary if you are setting up your primary DB instance in a VPC.

Connectivity & security	Monitoring	Logs & events	Configuration	Maintenance & backups	Tags
Instance					
Configuration			Instance class		Storage
DB instance id			Instance class		Encryption
database-1			db.m4.large		Enabled
Engine version			vCPU		KMS key
14.00.3192.2.v1			2		aws/rds 
DB name			RAM		Storage type
-			8 GB		General Purpose (
License model			Availability		IOPS
License Included			Master username		-
Collation			admin		Storage
SQL_Latin1_General_CI_AS			IAM db authentication		20 GiB
Option groups			Not Enabled		Storage autoscaling
default:sqlserver-se-14-00			Multi AZ		Enabled
ARN			Yes (Mirroring)		Maximum storage
arn:aws:rds:us-west-2: XXXXXXXXXX :db:database-1					1000 GiB
Resource id			Secondary Zone		
db- XXXXXXXXXX			us-west-2c		

You can also view the Availability Zone of the secondary using the AWS CLI command `describe-db-instances` or RDS API operation `DescribeDBInstances`. The output shows the secondary AZ where the standby mirror is located.

Migrating from Database Mirroring to Always On Availability Groups

In version 14.00.3049.1 of Microsoft SQL Server Enterprise edition, Always On Availability Groups (AGs) is enabled by default.

To migrate from Database Mirroring (DBM) to AGs, first check your version. If you are using a DB instance with a version prior to Enterprise Edition 13.00.5216.0, modify the instance to patch it to 13.00.5216.0 or later. If you are using a DB instance with a version prior to Enterprise Edition 14.00.3049.1, modify the instance to patch it to 14.00.3049.1 or later.

If you want to upgrade a mirrored DB instance to use AGs, run the upgrade first, modify the instance to remove Multi-AZ, and then modify it again to add Multi-AZ. This converts your instance to use Always On AGs.

Additional Features for Microsoft SQL Server on Amazon RDS

In the following sections, you can find information about augmenting Amazon RDS instances running the Microsoft SQL Server DB engine by adding SSL, Windows Authentication, and S3 integration.

Topics

- [Using SSL with a Microsoft SQL Server DB Instance \(p. 687\)](#)
- [Using Windows Authentication with a Microsoft SQL Server DB Instance \(p. 691\)](#)
- [Integrating an Amazon RDS for SQL Server DB Instance with Amazon S3 \(p. 700\)](#)

Using SSL with a Microsoft SQL Server DB Instance

You can use Secure Sockets Layer (SSL) to encrypt connections between your client applications and your Amazon RDS DB instances running Microsoft SQL Server. SSL support is available in all AWS regions for all supported SQL Server editions.

When you create a SQL Server DB instance, Amazon RDS creates an SSL certificate for it. The SSL certificate includes the DB instance endpoint as the Common Name (CN) for the SSL certificate to guard against spoofing attacks.

There are 2 ways to use SSL to connect to your SQL Server DB instance:

- Force SSL for all connections — this happens transparently to the client, and the client doesn't have to do any work to use SSL.
- Encrypt specific connections — this sets up an SSL connection from a specific client computer, and you must do work on the client to encrypt connections.

For information about Transport Layer Security (TLS) support for SQL Server, see [TLS 1.2 support for Microsoft SQL Server](#).

Forcing Connections to Your DB Instance to Use SSL

You can force all connections to your DB instance to use SSL. If you force connections to use SSL, it happens transparently to the client, and the client doesn't have to do any work to use SSL.

If you want to force SSL, use the `rds.force_ssl` parameter. By default, the `rds.force_ssl` parameter is set to 0 (off). Set the `rds.force_ssl` parameter to 1 (on) to force connections to use SSL. The `rds.force_ssl` parameter is static, so after you change the value, you must reboot your DB instance for the change to take effect.

To force all connections to your DB instance to use SSL

1. Determine the parameter group that is attached to your DB instance:
 - a. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
 - b. In the top right corner of the Amazon RDS console, choose the AWS Region of your DB instance.
 - c. In the navigation pane, choose **Databases**, and then choose the name of your DB instance to show its details.
 - d. Choose the **Configuration** tab. Find the **Parameter group** in the section.
2. If necessary, create a new parameter group. If your DB instance uses the default parameter group, you must create a new parameter group. If your DB instance uses a nondefault parameter group, you can choose to edit the existing parameter group or to create a new parameter group. If you edit an existing parameter group, the change affects all DB instances that use that parameter group.

To create a new parameter group, follow the instructions in [Creating a DB Parameter Group \(p. 212\)](#).

3. Edit your new or existing parameter group to set the `rds.force_ssl` parameter to `true`. To edit the parameter group, follow the instructions in [Modifying Parameters in a DB Parameter Group \(p. 213\)](#).
4. If you created a new parameter group, modify your DB instance to attach the new parameter group. Modify the **DB Parameter Group** setting of the DB instance. For more information, see [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#).
5. Reboot your DB instance. For more information, see [Rebooting a DB Instance \(p. 166\)](#).

Encrypting Specific Connections

You can force all connections to your DB instance to use SSL, or you can encrypt connections from specific client computers only. To use SSL from a specific client, you must obtain certificates for the client computer, import certificates on the client computer, and then encrypt the connections from the client computer.

Note

All SQL Server instances created after August 5, 2014, use the DB instance endpoint in the Common Name (CN) field of the SSL certificate. Prior to August 5, 2014, SSL certificate verification was not available for VPC-based SQL Server instances. If you have a VPC-based SQL Server DB instance that was created before August 5, 2014, and you want to use SSL certificate verification and ensure that the instance endpoint is included as the CN for the SSL certificate for that DB instance, then rename the instance. When you rename a DB instance, a new certificate is deployed and the instance is rebooted to enable the new certificate.

Obtaining Certificates for Client Computers

To encrypt connections from a client computer to an Amazon RDS DB instance running Microsoft SQL Server, you need a certificate on your client computer.

To obtain that certificate, download the certificate to your client computer. You can download a root certificate that works for all regions. You can also download a certificate bundle that contains both the old and new root certificate. In addition, you can download region-specific intermediate certificates. For more information about downloading certificates, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).

After you have downloaded the appropriate certificate, import the certificate into your Microsoft Windows operating system by following the procedure in the section following.

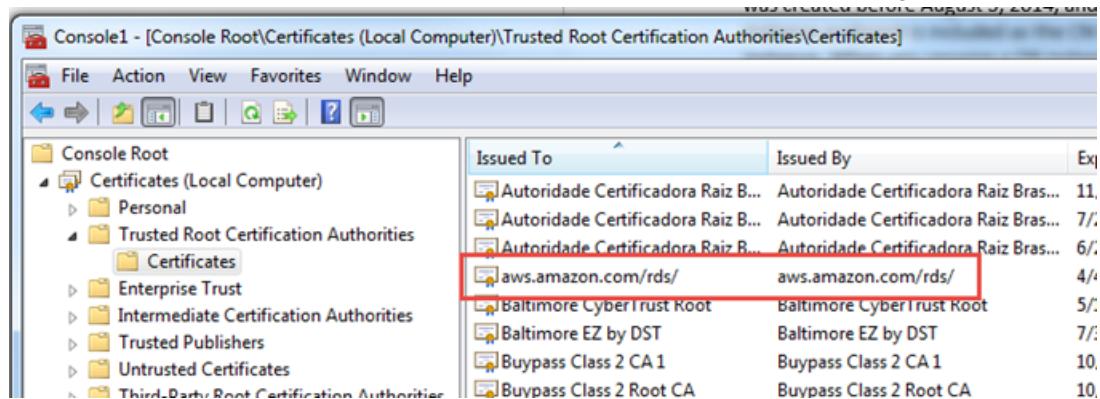
Importing Certificates on Client Computers

You can use the following procedure to import your certificate into the Microsoft Windows operating system on your client computer.

To import the certificate into your Windows operating system:

1. On the **Start** menu, type **Run** in the search box and press **Enter**.
2. In the **Open** box, type **MMC** and then choose **OK**.
3. In the MMC console, on the **File** menu, choose **Add/Remove Snap-in**.
4. In the **Add or Remove Snap-ins** dialog box, for **Available snap-ins**, select **Certificates**, and then choose **Add**.
5. In the MMC console, on the **File** menu, choose **Add/Remove Snap-in**.
6. In the **Certificates snap-in** dialog box, choose **Computer account**, and then choose **Next**.
7. In the **Select computer** dialog box, choose **Finish**.
8. In the **Add or Remove Snap-ins** dialog box, choose **OK**.
9. In the MMC console, expand **Certificates**, open the context (right-click) menu for **Trusted Root Certification Authorities**, choose **All Tasks**, and then choose **Import**.
10. On the first page of the Certificate Import Wizard, choose **Next**.
11. On the second page of the Certificate Import Wizard, choose **Browse**. In the browse window, change the file type to **All files (*.*)** because .pem is not a standard certificate extension. Locate the .pem file that you downloaded previously.
12. Choose **Open** to select the certificate file, and then choose **Next**.
13. On the third page of the Certificate Import Wizard, choose **Next**.

14. On the fourth page of the Certificate Import Wizard, choose **Finish**. A dialog box appears indicating that the import was successful.
15. In the MMC console, expand **Certificates**, expand **Trusted Root Certification Authorities**, and then choose **Certificates**. Locate the certificate to confirm it exists, as shown following.



16. Restart your computer.

Encrypting Connections to an Amazon RDS DB Instance Running Microsoft SQL Server

After you have imported a certificate into your client computer, you can encrypt connections from the client computer to an Amazon RDS DB instance running Microsoft SQL Server.

For SQL Server Management Studio, use the following procedure. For more information about SQL Server Management Studio, see [Use SQL Server Management Studio](#).

To encrypt connections from SQL Server Management Studio

1. Launch SQL Server Management Studio.
2. For **Connect to server**, type the server information, login user name, and password.
3. Choose **Options**.
4. Select **Encrypt connection**.
5. Choose **Connect**.
6. Confirm that your connection is encrypted by running the following query. Verify that the query returns true for `encrypt_option`.

```
select ENCRYPT_OPTION from SYS.DM_EXEC_CONNECTIONS where SESSION_ID = @@SPID
```

For any other SQL client, use the following procedure.

To encrypt connections from other SQL clients

1. Append `encrypt=true` to your connection string. This string might be available as an option, or as a property on the connection page in GUI tools.

Note

To enable SSL encryption for clients that connect using JDBC, you might need to add the Amazon RDS SQL certificate to the Java CA certificate (cacerts) store. You can do this by using the [keytool](#) utility.

2. Confirm that your connection is encrypted by running the following query. Verify that the query returns true for `encrypt_option`.

```
select ENCRYPT_OPTION from SYS.DM_EXEC_CONNECTIONS where SESSION_ID = @@SPID
```

Using Windows Authentication with a Microsoft SQL Server DB Instance

You can use Windows Authentication to authenticate users when they connect to your Amazon RDS DB instance running Microsoft SQL Server. The DB instance works with AWS Directory Service for Microsoft Active Directory, also called AWS Managed Microsoft AD, to enable Windows Authentication. When users authenticate with a SQL Server DB instance joined to the trusting domain, authentication requests are forwarded to the domain directory that you create with AWS Directory Service.

Amazon RDS supports Windows Authentication for SQL Server in all AWS Regions except the following:

- US West (N. California)
- Asia Pacific (Mumbai)
- South America (São Paulo)
- AWS GovCloud (US-East)
- AWS GovCloud (US-West)

Amazon RDS uses Mixed Mode for Windows Authentication. This approach means that the *master user* (the name and password used to create your SQL Server DB instance) uses SQL Authentication. Because the master user account is a privileged credential, you should restrict access to this account.

To get Windows Authentication using an on-premises or self-hosted Microsoft Active Directory, you need to create a forest trust. The trust can be one-way or two-way. For more information on setting up forest trusts using AWS Directory Service, see [When to Create a Trust Relationship](#).

To set up Windows authentication for a SQL Server DB instance, do the following steps (explained in greater detail in this section):

1. Use the AWS Directory Service for Microsoft Active Directory, also called AWS Managed Microsoft AD, either from the AWS console or AWS Directory Service API to create an AWS Managed Microsoft AD directory.
2. If you use the AWS CLI or Amazon RDS API to create your SQL Server DB instance, you need to create an AWS Identity and Access Management (IAM) role that uses the managed IAM policy `AmazonRDSDirectoryServiceAccess`. The role allows Amazon RDS to make calls to your directory. If you use the AWS console to create your SQL Server DB instance, AWS creates the IAM role for you.

For the role to allow access, the AWS Security Token Service (AWS STS) endpoint must be activated in the AWS Region for your AWS account. AWS STS endpoints are active by default in all AWS Regions, and you can use them without any further actions. For more information, see [Managing AWS STS in an AWS Region](#).

3. Create and configure users and groups in the AWS Managed Microsoft AD directory using the Microsoft Active Directory tools. For more information about creating users and groups in your Active Directory, see [Manage Users and Groups in AWS Managed Microsoft AD](#).
4. Use Amazon RDS to create a new SQL Server DB instance either from the AWS console, AWS CLI, or Amazon RDS API. In the create request, you provide the domain identifier ("d-*" identifier) that was generated when you created your directory and the name of the role you created. You can also modify an existing SQL Server DB instance to use Windows Authentication by setting the domain and IAM role parameters for the DB instance, and locating the DB instance in the same VPC as the domain directory.
5. Use the Amazon RDS *master user* credentials to connect to the SQL Server DB instance as you would any other DB instance. Because the DB instance is joined to the AWS Managed Microsoft AD domain, you can provision SQL Server logins and users from the Active Directory users and groups in their domain (known as SQL Server "Windows" logins). Database permissions are managed through standard SQL Server permissions granted and revoked to these windows logins.

Creating the Endpoint for Kerberos Authentication

Kerberos-based authentication requires that the endpoint be the customer-specified host name, a period, and then the fully qualified domain name (FQDN). For example, the following is an example of an endpoint you would use with Kerberos-based authentication. In this example, the SQL Server DB instance host name is ad-test and the domain name is corp-ad.company.com:

```
ad-test.corp-ad.company.com
```

If you want to check to make sure your connection is using Kerberos, you can run the following query:

```
SELECT net_transport, auth_scheme
  FROM sys.dm_exec_connections
 WHERE session_id = @@SPID;
```

Setting Up Windows Authentication for SQL Server DB Instances

You use AWS Directory Service for Microsoft Active Directory, also called AWS Managed Microsoft AD, to set up Windows Authentication for a SQL Server DB instance. To set up Windows Authentication, you take the following steps:

Step 1: Create a Directory Using the AWS Directory Service for Microsoft Active Directory

AWS Directory Service creates a fully managed, Microsoft Active Directory in the AWS cloud. When you create an AWS Managed Microsoft AD directory, AWS Directory Service creates two domain controllers and DNS servers on your behalf. The directory servers are created in two subnets in two different Availability Zones within a VPC. This redundancy helps ensure that your directory remains accessible even if a failure occurs.

When you create an AWS Managed Microsoft AD directory, AWS Directory Service performs the following tasks on your behalf:

- Sets up a Microsoft Active Directory within the VPC.
- Creates a directory administrator account with the user name Admin and the specified password. You use this account to manage your directory.

Note

Be sure to save this password. AWS Directory Service does not store this password and it cannot be retrieved or reset.

- Creates a security group for the directory controllers.

When you launch an AWS Directory Service for Microsoft Active Directory, AWS creates an Organizational Unit (OU) that contains all your directory's objects. This OU, which has the NetBIOS name that you typed when you created your directory, is located in the domain root. The domain root is owned and managed by AWS.

The *admin* account that was created with your AWS Managed Microsoft AD directory has permissions for the most common administrative activities for your OU:

- Create update, or delete users, groups, and computers
- Add resources to your domain such as file or print servers, and then assign permissions for those resources to users and groups in your OU
- Create additional OUs and containers

- Delegate authority
- Create and link group policies
- Restore deleted objects from the Active Directory Recycle Bin
- Run AD and DNS Windows PowerShell modules on the Active Directory Web Service

The *admin* account also has rights to perform the following domain-wide activities:

- Manage DNS configurations (Add, remove, or update records, zones, and forwarders)
- View DNS event logs
- View security event logs

To create a directory with AWS Managed Microsoft AD

1. In the [AWS Directory Service console](#) navigation pane, select **Directories** and choose **Set up directory**.
2. Choose **AWS Managed Microsoft AD**. This is the only option currently supported for use with Amazon RDS.
3. Choose **Next**.
4. On the **Enter directory information** page, provide the following information:

Edition

Choose the edition that meets your requirements.

Directory DNS name

The fully qualified name for the directory, such as corp.example.com.

Directory NetBIOS name

An optional short name for the directory, such as CORP.

Directory description

An optional description for the directory.

Admin password

The password for the directory administrator. The directory creation process creates an administrator account with the user name Admin and this password.

The directory administrator password and cannot include the word "admin." The password is case-sensitive and must be between 8 and 64 characters in length, inclusive. It must also contain at least one character from three of the following four categories:

- Lowercase letters (a-z)
- Uppercase letters (A-Z)
- Numbers (0-9)
- Non-alphanumeric characters (~!@#\$%^&*_+=`|\{};:"<>,.?/)

Confirm password

Retype the administrator password.

5. Choose **Next**.
6. On the **Choose VPC and subnets** page, provide the following information:

VPC

Select the VPC for the directory. The SQL Server DB instance must be created in this same VPC.

Subnets

Select the subnets for the directory servers. The two subnets must be in different Availability Zones.

7. Choose **Next**.
8. Review the directory information. If changes are needed, choose **Previous**. When the information is correct, choose **Create directory**.

Review & create

Review	
Directory type	VPC
Microsoft AD	vpc-8b6b78e9 (172.31.0.0/16)
Directory DNS name	Subnets
corp.example.com	subnet-75128d10 (172.31.64.0/20, us-east-1a) subnet-f51665dd (172.31.32.0/20, us-east-1b)
Directory NetBIOS name	
CORP	
Directory description	
My directory	

Pricing	
Edition	Free trial eligible Learn more
Standard	30-day limited trial
~USD 86.4000/mo (USD 0.1200/hr)*	
* Includes two domain controllers, USD 43.2000/mo for each additional domain controller.	

Cancel [Previous](#) **Create directory**

It takes several minutes for the directory to be created. When it has been successfully created, the **Status** value changes to **Active**.

To see information about your directory, choose the directory ID in the directory listing. Make a note of the **Directory ID**. You will need this value when you create or modify your SQL Server DB instance.

The screenshot shows the 'Directory details' page for a Microsoft AD directory. The 'Directory ID' field, which contains the value 'd-9067004760', is highlighted with a red circle. Other visible fields include 'Directory type' (Microsoft AD), 'Edition' (Standard), 'VPC' (vpc-8b6b78e9), 'Status' (Creating), 'Subnets' (subnet-f51665dd, subnet-75128d10), 'Last updated' (Friday, July 26, 2019), 'Availability zones' (us-east-1a, us-east-1b), 'Launch time' (Friday, July 26, 2019), 'DNS address' (corp.example.com), 'Directory NetBIOS name' (CORP), and a 'Description' section.

Step 2: Create the IAM role for Use by Amazon RDS

If you use the AWS console to create your SQL Server DB instance, you can skip this step. If you used the AWS CLI or Amazon RDS API to create your SQL Server DB instance, you must create an IAM role that uses the `AmazonRDSDirectoryServiceAccess` managed IAM policy. This role allows Amazon RDS to make calls to the AWS Directory Service for you.

If you are using a custom policy for joining a domain, rather than using the AWS-managed `AmazonRDSDirectoryServiceAccess` policy, you must allow the `"ds:GetAuthorizedApplicationDetails"` action. This is effective starting July 2019, due to a change in AWS Directory Service's API.

The following IAM policy, `AmazonRDSDirectoryServiceAccess`, provides access to AWS Directory Service:

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Action": [  
        "ds:DescribeDirectories",  
        "ds:AuthorizeApplication",  
        "ds:UnauthorizeApplication",  
        "ds:GetAuthorizedApplicationDetails"  
      ],  
      "Effect": "Allow",  
      "Resource": "*"  
    }  
  ]}
```

```
    ]  
}
```

Create an IAM role using this policy. For more information about creating IAM roles, see [Creating Customer Managed Policies](#).

Step 3: Create and Configure Users and Groups

You can create users and groups with the Active Directory Users and Computers tool, which is part of the Active Directory Domain Services and Active Directory Lightweight Directory Services tools. Users represent individual people or entities that have access to your directory. Groups are very useful for giving or denying privileges to groups of users, rather than having to apply those privileges to each individual user.

To create users and groups in an AWS Directory Service directory, you must be connected to a Windows EC2 instance that is a member of the AWS Directory Service directory, and be logged in as a user that has privileges to create users and groups. For more information, see [Add Users and Groups \(Simple AD and AWS Managed Microsoft AD\)](#).

Step 4: Create or Modify a SQL Server DB Instance

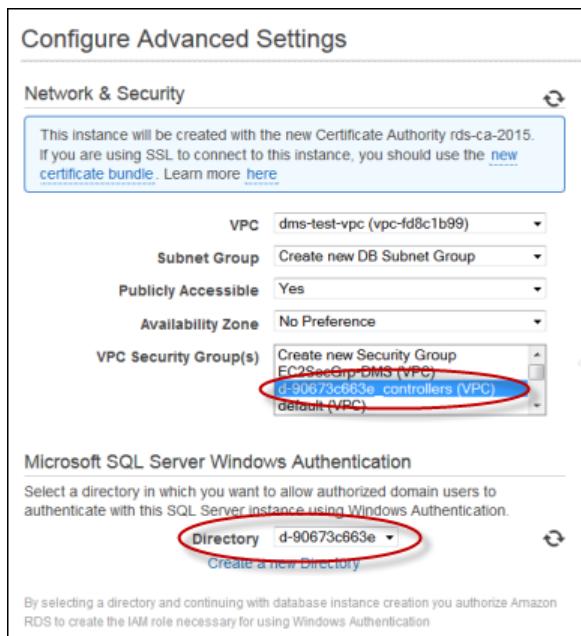
Next, you create or modify a Microsoft SQL Server DB instance for use with the directory. You can do this in one of the following ways:

- Create a new SQL Server DB instance.
- Modify an existing SQL Server DB instance.
- Restore a SQL Server DB instance from a DB snapshot.
- Restore a SQL Server DB instance from a point-in-time restore (PITR).

Windows Authentication is only supported for SQL Server DB instances in a VPC, and the DB instance must be in the same VPC as the directory.

Several parameters are required for the DB instance to be able to use the domain directory you created:

- For the **Directory** parameter, you must enter the domain identifier ("d-*" identifier) generated when you created the directory.
- Use the same VPC that was used when you created the directory.
- Use a security group that allows egress within the VPC so the DB instance can communicate with the directory.



Step 5: Create Windows Authentication SQL Server Logins

Use the Amazon RDS *master user* credentials to connect to the SQL Server DB instance as you would any other DB instance. Because the DB instance is joined to the AWS Managed Microsoft AD domain, you can provision SQL Server logins and users from the Active Directory users and groups in your domain. Database permissions are managed through standard SQL Server permissions granted and revoked to these windows logins.

To allow an Active Directory user to authenticate with SQL Server, a SQL Server Windows login must exist for the user or a group that the user is a member of. Fine-grained access control is handled through granting and revoking permissions on these SQL Server logins. If a user does not have a corresponding SQL Server login and is not a member of a group with a corresponding SQL Server login, that user cannot access the SQL Server DB instance.

The ALTER ANY LOGIN permission is required to create an Active Directory SQL Server login. If you have not yet created any logins with this permission, connect as the DB instance's *master user* using SQL Server Authentication. Run the following data definition language (DDL) command to create a SQL Server login for an Active Directory user or group:

```
CREATE LOGIN [<user or group>] FROM WINDOWS WITH DEFAULT_DATABASE = [master],  
DEFAULT_LANGUAGE = [us_english];
```

Users or groups must be specified using the pre-Windows 2000 login name in the format *domainName\login_name*. You cannot use a User Principle Name (UPN) in the format *login_name@DomainName*. For more information about CREATE LOGIN, go to <https://msdn.microsoft.com/en-us/library/ms189751.aspx> in the Microsoft Developer Network documentation.

Users (both humans and applications) from your domain can now connect to the RDS SQL Server instance from a domain joined client machine using Windows authentication.

Managing a DB Instance in a Domain

You can use the AWS console, AWS CLI, or the Amazon RDS API to manage your DB instance and its relationship with your domain, such as moving the DB instance into, out of, or between domains.

For example, using the Amazon RDS API, you can do the following:

- To re-attempt a domain join for a failed membership, use the `ModifyDBInstance` API operation and specify the current membership's directory ID.
- To update the IAM role name for membership, use the `ModifyDBInstance` API operation and specify the current membership's directory ID and the new IAM role.
- To remove a DB instance from a domain, use the `ModifyDBInstance` API operation and specify `none` as the domain parameter.
- To move a DB instance from one domain to another, use the `ModifyDBInstance` API operation and specify the domain identifier of the new domain as the domain parameter.
- To list membership for each DB instance, use the `DescribeDBInstances` API operation.

Understanding Domain Membership

After you create or modify your DB instance, the instance becomes a member of the domain. The AWS console indicates the status of the domain membership for the DB instance. The status of the DB instance can be one of the following:

- **joined** – The instance is a member of the domain.
- **joining** – The instance is in the process of becoming a member of the domain.
- **pending-join** – The instance membership is pending .
- **pending-maintenance-join** – AWS will attempt to make the instance a member of the domain during the next scheduled maintenance window.
- **pending-removal** – The removal of the instance from the domain is pending.
- **pending-maintenance-removal** – AWS will attempt to remove the instance from the domain during the next scheduled maintenance window.
- **failed** – A configuration problem has prevented the instance from joining the domain. Check and fix your configuration before re-issuing the instance modify command.
- **removing** – The instance is being removed from the domain.

A request to become a member of a domain can fail because of a network connectivity issue or an incorrect IAM role. If you create a DB instance or modify an existing instance and the attempt to become a member of a domain fails, you should re-issue the modify command or modify the newly created instance to join the domain.

Connecting to SQL Server with Windows Authentication

To connect to SQL Server with Windows Authentication, you must be logged into a domain-joined computer as a domain user. After launching SQL Server Management Studio, choose **Windows Authentication** as the authentication type, as shown following.



Restoring a SQL Server DB Instance and then Adding It to a Domain

You can restore a DB snapshot or do a point-in-time restore (PITR) for a SQL Server DB instance and then add it to a domain. Once the DB instance is restored, modify the instance using the process explained in [Step 4: Create or Modify a SQL Server DB Instance \(p. 696\)](#) to add the DB instance to a domain.

Related Topics

- [Microsoft SQL Server on Amazon RDS \(p. 598\)](#)
- [Security in Amazon RDS \(p. 427\)](#)

Integrating an Amazon RDS for SQL Server DB Instance with Amazon S3

You can transfer files between a DB instance running Amazon RDS for SQL Server and an Amazon S3 bucket. By doing this, you can use Amazon S3 with SQL Server features such as BULK INSERT. For example, you can download .csv, .xml, .txt, and other files from Amazon S3 to the DB instance host and import the data from D:\S3\ into the database. All files are stored in D:\S3\ on the DB instance.

The following limitations apply:

- The S3 integration feature is not supported on Multi-AZ instances.
- The DB instance and the S3 bucket must be in the same AWS Region.
- S3 integration tasks are run sequentially and share the same queue as native backup and restore tasks. At maximum, you can have only two tasks in progress at any time in this queue.
- You must re-enable the S3 integration feature on restored instances. It isn't propagated from the source instance to the restored instance. Files in D:\S3 will be cleaned up on a restored instance.
- Downloading to the DB instance is limited to 100 files. In other words, there can't be more than 100 files in D:\S3\.
- Only files without file extensions or with the following file extensions are supported for download: .xml, .txt, .fmt, .lst, .dat, .csv, .info, .bcp, and .tbl.
- The S3 bucket must have the same owner as the related AWS Identity and Access Management (IAM) role. Also, the bucket can't be open to the public.
- File size for uploads is limited to 50 GB per file.

Topics

- [Prerequisites for Integrating RDS SQL Server with S3 \(p. 700\)](#)
- [Enabling RDS SQL Server Integration with S3 \(p. 704\)](#)
- [Transferring Files Between RDS SQL Server and an S3 Bucket \(p. 705\)](#)
- [Disabling RDS SQL Server Integration with S3 \(p. 710\)](#)

Prerequisites for Integrating RDS SQL Server with S3

Before you begin, find or create the S3 bucket that you want to use. Also, add permissions so that the RDS DB instance can access the S3 bucket. To configure this access, you create both an IAM policy and an IAM role.

Console

To create an IAM policy for access to Amazon S3

1. In the [IAM Management Console](#), choose **Policies** in the navigation pane.
2. Create a new policy, and use the **Visual editor** tab for the following steps.
3. For **Service**, enter **s3** and then choose the **S3** service.
4. For **Actions**, choose the following to grant the access that your DB instance requires:
 - **ListAllMyBuckets** – required
 - **ListBucket** – required
 - **GetBucketACL** – required
 - **GetBucketLocation** – required
 - **GetObject** – required for downloading files from S3 to D:\S3\

- `PutObject` – required for uploading files from `D:\S3\` to S3
 - `ListMultipartUploadParts` – required for uploading files from `D:\S3\` to S3
 - `AbortMultipartUpload` – required for uploading files from `D:\S3\` to S3
5. For **Resources**, the options that display depend on which actions you choose in the previous step. You might see options for **bucket**, **object**, or both. For each of these, add the appropriate Amazon Resource Name (ARN).
- For **bucket**, add the ARN for the bucket that you want to use. For example, if your bucket is named `example-bucket`, set the ARN to `arn:aws:s3:::example-bucket`.
- For **object**, enter the ARN for the bucket and then choose one of the following:
- To grant access to all files in the specified bucket, choose **Any** for both **Bucket name** and **Object name**.
 - To grant access to specific files or folders in the bucket, provide ARNs for the specific buckets and objects that you want SQL Server to access.
6. Follow the instructions in the console until you finish creating the policy.

The preceding is an abbreviated guide to setting up a policy. For more detailed instructions on creating IAM policies, see [Creating IAM Policies](#) in the *IAM User Guide*.

To create an IAM role that uses the IAM policy from the previous procedure

1. In the [IAM Management Console](#), choose **Roles** in the navigation pane.
2. Create a new IAM role, and choose the following options as they appear in the console:
 - **AWS service**
 - **RDS**
 - **RDS – Add Role to Database**

Then choose **Next:Permissions** at the bottom.

3. For **Attach permissions policies**, enter the name of the IAM policy that you previously created. Then choose the policy from the list.
4. Follow the instructions in the console until you finish creating the role.

The preceding is an abbreviated guide to setting up a role. If you want more detailed instructions on creating roles, see [IAM Roles](#) in the *IAM User Guide*.

AWS CLI

To grant Amazon RDS access to an Amazon S3 bucket

1. Create an IAM policy that grants Amazon RDS access to an S3 bucket.

Include the appropriate actions to grant the access your DB instance requires:

- `ListAllMyBuckets` – required
- `ListBucket` – required
- `GetBucketACL` – required
- `GetBucketLocation` – required
- `GetObject` – required for downloading files from S3 to `D:\S3\`
- `PutObject` – required for uploading files from `D:\S3\` to S3

- `ListMultipartUploadParts` – required for uploading files from D:\S3\ to S3
- `AbortMultipartUpload` – required for uploading files from D:\S3\ to S3

The following AWS CLI command creates an IAM policy named `rds-s3-integration-policy` with these options. It grants access to a bucket named `your-s3-bucket-arn`.

Example

For Linux, OS X, or Unix:

```
aws iam create-policy \
--policy-name rds-s3-integration-policy \
--policy-document '{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": "s3>ListAllMyBuckets",
            "Resource": "*"
        },
        {
            "Effect": "Allow",
            "Action": [
                "s3>ListBucket",
                "s3:GetBucketACL",
                "s3:GetBucketLocation"
            ],
            "Resource": "arn:aws:s3:::bucket_name"
        },
        {
            "Effect": "Allow",
            "Action": [
                "s3.GetObject",
                "s3:PutObject",
                "s3>ListMultipartUploadParts",
                "s3:AbortMultipartUpload"
            ],
            "Resource": "arn:aws:s3:::bucket_name/key_prefix/*"
        }
    ]
}'
```

For Windows:

Make sure to change the line endings to the ones supported by your interface (^ instead of \). Also, in Windows, you must escape all double quotes with a \. To avoid the need to escape the quotes in the JSON, you can save it to a file instead and pass that in as a parameter.

First, create the `policy.json` file with the following permission policy:

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": "s3>ListAllMyBuckets",
            "Resource": "*"
        },
        {
            "Effect": "Allow",
            "Action": [
```

```

        "s3>ListBucket",
        "s3:GetBucketACL",
        "s3:GetBucketLocation"
    ],
    "Resource": "arn:aws:s3:::bucket_name"
},
{
    "Effect": "Allow",
    "Action": [
        "s3GetObject",
        "s3PutObject",
        "s3ListMultipartUploadParts",
        "s3AbortMultipartUpload"
    ],
    "Resource": "arn:aws:s3:::bucket_name/key_prefix/*"
}
]
}

```

Then use the following command to create the policy:

```
aws iam create-policy ^
--policy-name rds-s3-integration-policy ^
--policy-document file://policy_file_path
```

2. After the policy is created, note the Amazon Resource Name (ARN) of the policy. You need the ARN for a later step.
3. Create an IAM role that Amazon RDS can assume on your behalf to access your S3 buckets.

The following AWS CLI command creates the *rds-s3-integration-role* for this purpose.

Example

For Linux, OS X, or Unix:

```
aws iam create-role \
--role-name rds-s3-integration-role \
--assume-role-policy-document '{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Principal": {
                "Service": "rds.amazonaws.com"
            },
            "Action": "sts:AssumeRole"
        }
    ]
}'
```

For Windows:

Make sure to change the line endings to the ones supported by your interface (^ instead of \). Also, in Windows, you must escape all double quotes with a \. To avoid the need to escape the quotes in the JSON, you can save it to a file instead and pass that in as a parameter.

First, create the *assume_role_policy.json* file with the following policy:

```
{
    "Version": "2012-10-17",
    "Statement": [
```

```
{  
    "Effect": "Allow",  
    "Principal": {  
        "Service": [  
            "rds.amazonaws.com"  
        ]  
    },  
    "Action": "sts:AssumeRole"  
}  
]  
}
```

Then use the following command to create the IAM role:

```
aws iam create-role ^  
--role-name rds-s3-integration-role ^  
--assume-role-policy-document file://assume_role_policy_file_path
```

For more information, see [Creating a Role to Delegate Permissions to an IAM User](#) in the *IAM User Guide*.

4. After the IAM role is created, note the ARN of the role. You need the ARN for a later step.
5. Attach the IAM policy that you created to the IAM role that you created.

The following AWS CLI command attaches the policy to the role named *rds-s3-integration-role*.

Example

For Linux, OS X, or Unix:

```
aws iam attach-role-policy \  
--policy-arn your-policy-arn \  
--role-name rds-s3-integration-role
```

For Windows:

```
aws iam attach-role-policy ^  
--policy-arn your-policy-arn ^  
--role-name rds-s3-integration-role
```

Replace *your-policy-arn* with the policy ARN that you noted in a previous step.

Enabling RDS SQL Server Integration with S3

In the following section, you can find how to enable Amazon S3 integration with Amazon RDS for SQL Server. To work with S3 integration, your DB instance must be associated with the IAM role that you previously created before you use the `S3_INTEGRATION` feature-name parameter.

Note

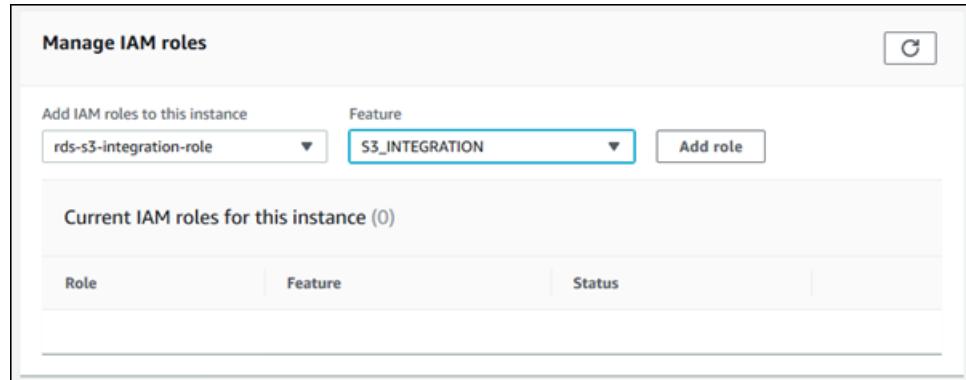
To add an IAM role to a DB instance, the status of the DB instance must be **available**.

Console

To associate your IAM role with your DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.

2. Choose the RDS SQL Server DB instance name to display its details.
3. On the **Connectivity & security** tab, in the **Manage IAM roles** section, choose the IAM role to add for **Add IAM roles to this instance**.
4. For **Feature**, choose **S3_INTEGRATION**.



5. Choose **Add role**.

AWS CLI

To add the IAM role to the RDS SQL Server DB instance

- The following AWS CLI command adds your IAM role to an RDS SQL Server DB instance named *mydbinstance*.

Example

For Linux, OS X, or Unix:

```
aws rds add-role-to-db-instance \
    --db-instance-identifier mydbinstance \
    --feature-name S3_INTEGRATION \
    --role-arn your-role-arn
```

For Windows:

```
aws rds add-role-to-db-instance ^
    --db-instance-identifier mydbinstance ^
    --feature-name S3_INTEGRATION ^
    --role-arn your-role-arn
```

Replace *your-role-arn* with the role ARN that you noted in a previous step. S3_INTEGRATION must be specified for the --feature-name option.

Transferring Files Between RDS SQL Server and an S3 Bucket

You can use Amazon RDS stored procedures to download and upload files between S3 and your RDS DB instance. You can also use Amazon RDS stored procedures to list and delete files on the RDS instance.

The files that you download from and upload to Amazon S3 are stored in the D:\S3 folder. This is the only folder that you can use to access your files. You can organize your files into subfolders, which are created for you when you include the destination folder during download.

Some of the stored procedures require that you provide an Amazon Resource Name (ARN) to your Amazon S3 bucket and file. The format for your ARN is `arn:aws:s3:::bucket_name/file_name`. Amazon S3 doesn't require an account number or AWS Region in ARNs.

S3 integration tasks run sequentially and share the same queue as native backup and restore tasks. At maximum, you can have only two tasks in progress at any time in this queue. It can take up to five minutes for the task to begin processing.

Downloading Files from an Amazon S3 Bucket to an SQL Server DB Instance

To download files from an S3 bucket to an RDS SQL Server DB instance, use the Amazon RDS stored procedure `msdb.dbo.rds_download_from_s3` with the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>@s3_arn_of_file</code>	NVARCHAR	–	Required	The S3 ARN of the file to download, for example: <code>arn:aws:s3:::bucket_name/mydata.csv</code>
<code>@rds_file_path</code>	NVARCHAR	–	Optional	The file path for the RDS instance. If not specified, the file path is <code>D:\S3\<filename in s3></code> . RDS supports absolute paths and relative paths. If you want to create a subfolder, include it in the file path.
<code>@overwrite_file</code>	INT	0	Optional	Overwrite the existing file: 0 = Don't overwrite 1 = Overwrite

You can download files without a file extension and files with the following file extensions: `.xml`, `.txt`, `.fmt`, `.lst`, `.dat`, `.csv`, `.info`, `.bcp`, and `.tbl`.

The following example shows the stored procedure to download files from S3.

```
exec msdb.dbo.rds_download_from_s3
    @s3_arn_of_file='arn:aws:s3:::bucket_name/bulk_data.csv',
    @rds_file_path='D:\S3\seed_data\data.csv',
    @overwrite_file=1;
```

The example `rds_download_from_s3` operation creates a folder named `seed_data` in `D:\S3\`, if the folder doesn't exist yet. Then the example downloads the source file `bulk_data.csv` from S3 to a new file named `data.csv` on the DB instance. If the file previously existed, it's overwritten because the `@overwrite_file` parameter is set to 1.

Uploading Files from an SQL Server DB Instance to an Amazon S3 Bucket

To upload files from an RDS SQL Server DB instance to an S3 bucket, use the Amazon RDS stored procedure `msdb.dbo.rds_upload_to_s3` with the following parameters.

Parameter Name	Data Type	Default	Required	Description
@s3_arn_of_file	NVARCHAR	–	Required	The S3 ARN of the file to be created in S3, for example: <code>arn:aws:s3:::bucket_name/mydata.csv</code>
@rds_file_path	NVARCHAR	–	Required	The file path of the file to upload to S3. Absolute and relative paths are supported.
@overwrite_file	INT	–	Optional	Overwrite the existing file: 0 = Don't overwrite 1 = Overwrite

The following example uploads the file named `data.csv` from the specified location in `D:\S3\seed_data\` to a file `new_data.csv` in the S3 bucket specified by the ARN.

```
exec msdb.dbo.rds_upload_to_s3
    @rds_file_path='D:\S3\seed_data\data.csv',
    @s3_arn_of_file='arn:aws:s3:::bucket_name/new_data.csv',
    @overwrite_file=1;
```

If the file previously existed in S3, it's overwritten because the `@overwrite_file` parameter is set to 1.

Listing Files on the RDS DB Instance

To list the files available on the DB instance, use both a stored procedure and a function. First, run the following stored procedure to gather file details from the files in `D:\S3\`.

```
exec msdb.dbo.rds_gather_file_details;
```

The stored procedure returns the ID of the task. Like other tasks, this stored procedure runs asynchronously. As soon as the status of the task is `SUCCESS`, you can use the task ID in the `rds_fn_list_file_details` function to list the existing files and directories in `D:\S3\`, as shown following.

```
SELECT * FROM msdb.dbo.rds_fn_list_file_details(TASK_ID);
```

The `rds_fn_list_file_details` function returns a table with the following columns.

Output Parameter	Description
<code>filepath</code>	Absolute path of the file (for example, <code>D:\S3\mydata.csv</code>)
<code>size_in_bytes</code>	File size (in bytes)
<code>last_modified_utc</code>	Last modification date and time in UTC format
<code>is_directory</code>	Option that indicates whether the item is a directory (<code>true/false</code>)

Deleting Files on the RDS Instance

To delete the files available on the DB instance, use the Amazon RDS stored procedure `msdb.dbo.rds_delete_from_filesystem` with the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>@rds_file_path</code>	NVARCHAR	–	Required	The file path of the file to delete. Absolute and relative paths are supported.
<code>@force_delete</code>	INT	0	Optional	To delete a directory, this flag must be included and set to 1. 1 = delete a directory This parameter is ignored if you are deleting a file.

To delete a directory, the `@rds_file_path` must end with a backslash (\) and `@force_delete` must be set to 1.

The following example deletes the file D:\S3\delete_me.txt.

```
exec msdb.dbo.rds_delete_from_filesystem
    @rds_file_path='D:\S3\delete_me.txt';
```

The following example deletes the directory D:\S3\example_folder\.

```
exec msdb.dbo.rds_delete_from_filesystem
    @rds_file_path='D:\S3\example_folder\' ,
    @force_delete=1;
```

Monitoring the Status of a File Transfer Task

To track the status of your S3 integration task, call the `rds_fn_task_status` function. It takes two parameters. The first parameter should always be `NULL` because it doesn't apply to S3 integration. The second parameter accepts a task ID. Set the second parameter to 0 to get results for all tasks.

To see a list of all tasks, set the first parameter to `NULL` and the second parameter to 0, as shown in the following example.

```
SELECT * FROM msdb.dbo.rds_fn_task_status(null,0);
```

To get a specific task, set the first parameter to `NULL` and the second parameter to the task ID, as shown in the following example.

```
SELECT * FROM msdb.dbo.rds_fn_task_status(null,42);
```

The `rds_fn_task_status` function returns the following information.

Output Parameter	Description
task_id	The ID of the task
task_type	For S3 integration, tasks can have the following task types: <ul style="list-style-type: none"> • DOWNLOAD_FROM_S3 • UPLOAD_TO_S3 • LIST_FILES_ON_DISK • DELETE_FILES_ON_DISK
database_name	Not applicable to S3 integration tasks.
% complete	The progress of the task as a percentage.
duration (mins)	The amount of time spent on the task, in minutes.
lifecycle	The status of the task. Possible statuses are the following: <ul style="list-style-type: none"> • CREATED – After you call one of the S3 integration stored procedures, a task is created and the status is set to CREATED. • IN_PROGRESS – After a task starts, the status is set to IN_PROGRESS. It can take up to five minutes for the status to change from CREATED to IN_PROGRESS. • SUCCESS – After a task completes, the status is set to SUCCESS. • ERROR – If a task fails, the status is set to ERROR. For more information about the error, see the task_info column. • CANCEL_REQUESTED – After you call rds_cancel_task, the status of the task is set to CANCEL_REQUESTED. • CANCELLED – After a task is successfully canceled, the status of the task is set to CANCELLED.
task_info	Additional information about the task. If an error occurs during processing, this column contains information about the error.
last_updated	The date and time that the task status was last updated.
created_at	The date and time that the task was created.
S3_object_arn	The ARN of the S3 object downloaded from or uploaded to.
overwrite_S3_backup_file	Not applicable to S3 integration tasks.
KMS_master_key_arn	Not applicable to S3 integration tasks.
filepath	The file path on the RDS DB instance.

Output Parameter	Description
overwrite_file	An option that indicates if an existing file is overwritten.
task_metadata	Not applicable to S3 integration tasks.

Canceling a Task

To cancel S3 integration tasks, use the `msdb.dbo.rds_cancel_task` stored procedure with the `task_id` parameter. Delete and list tasks that are in progress can't be cancelled. The following example shows a request to cancel a task.

```
exec msdb.dbo.rds_cancel_task @task_id = 1234;
```

To get an overview of all tasks and their task IDs, use the `rds_fn_task_status` function as described in [Monitoring the Status of a File Transfer Task \(p. 708\)](#).

Disabling RDS SQL Server Integration with S3

Following, you can find how to disable Amazon S3 integration with Amazon RDS for SQL Server. Files in `D:\S3\` will not be deleted when disabling S3 integration.

Note

To remove an IAM role from a DB instance, the status of the DB instance must be available.

Console

To disassociate your IAM role from your DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. Choose the RDS SQL Server DB instance name to display its details.
3. On the **Connectivity & security** tab, in the **Manage IAM roles** section, choose the IAM role to remove.
4. Choose **Delete**.

AWS CLI

To remove the IAM role from the RDS SQL Server DB instance

- The following AWS CLI command removes the IAM role from a RDS SQL Server DB instance named `mydbinstance`.

Example

For Linux, OS X, or Unix:

```
aws rds remove-role-from-db-instance \
--db-instance-identifier mydbinstance \
--feature-name S3_INTEGRATION \
--role-arn your-role-arn
```

For Windows:

```
aws rds remove-role-from-db-instance ^
--db-instance-identifier mydbinstance ^
--feature-name S3_INTEGRATION ^
--role-arn your-role-arn
```

Replace *your-role-arn* with the appropriate IAM role ARN for the --feature-name option.

Options for the Microsoft SQL Server Database Engine

In this section, you can find descriptions for options that are available for Amazon RDS instances running the Microsoft SQL Server DB engine. To enable these options, you add them to an option group, and then associate the option group with your DB instance. For more information, see [Working with Option Groups \(p. 196\)](#).

If you are looking for optional features that aren't added through RDS option groups—such as SSL, Windows Authentication, and S3 integration, see [Additional Features for Microsoft SQL Server on Amazon RDS \(p. 686\)](#).

Amazon RDS supports the following options for Microsoft SQL Server DB instances.

Option	Option ID	Engine Editions
Native Backup and Restore (p. 712)	SOLSERVER_BACKUP_RESTORE	SQL Server Enterprise Edition SQL Server Standard Edition SQL Server Web Edition SQL Server Express Edition
Transparent Data Encryption (p. 714)	TRANSPARENT_DATA_ENCRYPTION	SQL Server Enterprise Edition
SQL Server Audit (p. 717)	SOLSERVER_AUDIT	In RDS, starting with SQL Server 2012, all editions of SQL Server support server level audits, and Enterprise edition also supports database level audits. Starting with SQL Server SQL Server 2016 (13.x) SP1, all editions support both server and database level audits. For more information, see SQL Server Audit (Database Engine)

Support for Native Backup and Restore in SQL Server

By using native backup and restore for SQL Server databases, you can create a full backup of your on-premises database and store the backup files on Amazon S3. You can then restore to an existing Amazon RDS DB instance running SQL Server. You can also back up an RDS SQL Server database, store it on Amazon S3, and restore it in other locations. In addition, you can restore the backup to an on-premises

server, or a different Amazon RDS DB instance running SQL Server. For more information, see [Importing and Exporting SQL Server Databases \(p. 657\)](#).

Amazon RDS supports native backup and restore for Microsoft SQL Server databases by using full backup files (.bak files).

Adding the Native Backup and Restore Option

The general process for adding the native backup and restore option to a DB instance is the following:

1. Create a new option group, or copy or modify an existing option group.
2. Add the **SQLSERVER_BACKUP_RESTORE** option to the option group.
3. Associate an IAM role with the option. The IAM role must have access to an S3 bucket to store the database backups.

That is, it must have as its option setting a valid Amazon Resource Name (ARN) in the format `arn:aws:iam::account-id:role/role-name`. For more information, see [AWS Identity and Access Management \(IAM\)](#).

4. Associate the option group with the DB instance.

After you add the native backup and restore option, you don't need to restart your DB instance. As soon as the option group is active, you can begin backing up and restoring immediately.

To add the native backup and restore option

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Option groups**.
3. Create a new option group or use an existing option group. For information on how to create a custom DB option group, see [Creating an Option Group \(p. 197\)](#).

To use an existing option group, skip to the next step.

4. Add the **SQLSERVER_BACKUP_RESTORE** option to the option group. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#).
5. Do one of the following:
 - To use an existing IAM role and Amazon S3 settings, choose an existing IAM role for **IAM Role**. If you use an existing IAM role, RDS uses the Amazon S3 settings configured for this role.
 - To create a new role and configure new Amazon S3 settings, do the following:
 1. For **IAM Role**, choose **Create a New Role**.
 2. For **Select S3 Bucket**, either create an S3 bucket or use an existing one. To create a new bucket, choose **Create a New S3 Bucket**. To use an existing bucket, choose it from the list.
 3. For **S3 folder path prefix (optional)**, specify a prefix to use for the files stored in your Amazon S3 bucket.

This prefix can include a file path but doesn't have to. If you provide a prefix, RDS attaches that prefix to all backup files. RDS then uses the prefix during a restore to identify related files and ignore irrelevant files. For example, you might use the S3 bucket for purposes besides holding backup files. In this case, you can use the prefix to have RDS perform native backup and restore only on a particular folder and its subfolders.

If you leave the prefix blank, then RDS doesn't use a prefix to identify backup files or files to restore. As a result, during a multiple-file restore, RDS attempts to restore every file in every folder of the S3 bucket.

4. For **Enable Encryption**, choose **Yes** to encrypt the backup file. Choose **No** to leave the backup file unencrypted.

If you choose **Yes**, choose an encryption key for **Master Key**. For more information about encryption keys, see [Getting Started](#) in the AWS Key Management Service (AWS KMS) documentation.

6. Choose **Add option**.
7. Apply the option group to a new or existing DB instance:
 - For a new DB instance, apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the Microsoft SQL Server Database Engine \(p. 616\)](#).
 - For an existing DB instance, apply the option group by modifying the instance and attaching the new option group. For more information, see [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#).

Modifying Native Backup and Restore Option Settings

After you enable the native backup and restore option, you can modify the settings for the option. For more information about how to modify option settings, see [Modifying an Option Setting \(p. 204\)](#).

Removing the Native Backup and Restore Option

You can turn off native backup and restore by removing the option from your DB instance. After you remove the native backup and restore option, you don't need to restart your DB instance.

To remove the native backup and restore option from a DB instance, do one of the following:

- Remove the option from the option group it belongs to. This change affects all DB instances that use the option group. For more information, see [Removing an Option from an Option Group \(p. 207\)](#).
- Modify the DB instance and specify a different option group that doesn't include the native backup and restore option. This change affects a single DB instance. You can specify the default (empty) option group, or a different custom option group. For more information, see [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#).

Support for Transparent Data Encryption in SQL Server

Amazon RDS supports using Transparent Data Encryption (TDE) to encrypt stored data on your DB instances running Microsoft SQL Server. TDE automatically encrypts data before it is written to storage, and automatically decrypts data when the data is read from storage.

Amazon RDS supports TDE for the following SQL Server versions and editions:

- SQL Server 2017 Enterprise Edition
- SQL Server 2016 Enterprise Edition
- SQL Server 2014 Enterprise Edition
- SQL Server 2012 Enterprise Edition

To enable transparent data encryption for an RDS SQL Server DB instance, specify the TDE option in an RDS option group that is associated with that DB instance.

Transparent data encryption for SQL Server provides encryption key management by using a two-tier key architecture. A certificate, which is generated from the database master key, is used to protect the data encryption keys. The database encryption key performs the actual encryption and decryption of data on the user database. Amazon RDS backs up and manages the database master key and the TDE certificate. To comply with several security standards, Amazon RDS is working to implement automatic periodic master key rotation.

Transparent data encryption is used in scenarios where you need to encrypt sensitive data. For example, you might want to provide data files and backups to a third party, or address security-related regulatory compliance issues. You can't encrypt the system databases for SQL Server, such as the Model or Master databases.

A detailed discussion of transparent data encryption is beyond the scope of this guide, but you should understand the security strengths and weaknesses of each encryption algorithm and key. For information about transparent data encryption for SQL Server, see [Transparent Data Encryption \(TDE\)](#) on the Microsoft website.

You should determine if your DB instance is already associated with an option group that has the TDE option. To view the option group that a DB instance is associated with, you can use the RDS console, the [describe-db-instance](#) AWS CLI command, or the API operation [DescribeDBInstances](#).

The process for enabling transparent data encryption on a SQL Server DB instance is as follows:

1. If the DB instance isn't associated with an option group that has TDE enabled, you have two choices. You can create an option group and add the TDE option, or you can modify the associated option group to add it. For information about creating or modifying an option group, see [Working with Option Groups \(p. 196\)](#). For information about adding an option to an option group, see [Adding an Option to an Option Group \(p. 200\)](#).
2. Associate the DB instance with the option group with the TDE option. For information about associating a DB instance with an option group, see [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#).

When the TDE option is added to an option group, Amazon RDS generates a certificate that is used in the encryption process. You can then use the certificate to run SQL statements that encrypt data in a database on the DB instance. The following example uses the RDS-created certificate called RDSTDECertificateName to encrypt a database called customerDatabase.

```
----- Enabling TDE -----  
  
-- Find a RDSTDECertificate to use  
USE [master]  
GO  
SELECT name FROM sys.certificates WHERE name LIKE 'RDSTDECertificate%'  
GO  
  
USE [customerDatabase]  
GO  
-- Create DEK using one of the certificates from the previous step  
CREATE DATABASE ENCRYPTION KEY  
WITH ALGORITHM = AES_128  
ENCRYPTION BY SERVER CERTIFICATE [RDSTDECertificateName]  
GO  
  
-- Enable encryption on the database  
ALTER DATABASE [customerDatabase]  
SET ENCRYPTION ON  
GO  
  
-- Verify that the database is encrypted  
USE [master]
```

```
GO
SELECT name FROM sys.databases WHERE is_encrypted = 1
GO
SELECT db_name(database_id) as DatabaseName, * FROM sys.dm_database_encryption_keys
GO
```

The time that it takes to encrypt a SQL Server database using TDE depends on several factors. These include the size of the DB instance, whether PIOPS is enabled for the instance, the amount of data, and other factors.

The TDE option is a persistent option that you can't remove from an option group unless all DB instances and backups are disassociated from the option group. After you add the TDE option to an option group, the option group can only be associated with DB instances that use TDE. For more information about persistent options in an option group, see [Option Groups Overview \(p. 196\)](#).

Because the TDE option is a persistent option, you can have a conflict between the option group and an associated DB instance. You can have a conflict between the option group and an associated DB instance in the following situations:

- The current option group has the TDE option, and you replace it with an option group that does not have the TDE option.
- You restore from a DB snapshot to a new DB instance that does not have an option group that contains the TDE option. For more information about this scenario, see [Option Group Considerations \(p. 270\)](#).

To disable TDE for a DB instance, first ensure that there are no encrypted objects left on the DB instance by either unencrypting the objects or by dropping them. If any encrypted objects exist on the DB instance, you can't disable TDE for the DB instance. When you use the console to remove the TDE option from an option group, the console indicates that it is processing. In addition, an error event is created if the option group is associated with an encrypted DB instance or DB snapshot.

The following example removes the TDE encryption from a database called `customerDatabase`.

```
----- Removing TDE -----
USE [customerDatabase]
GO

-- Disable encryption on the database
ALTER DATABASE [customerDatabase]
SET ENCRYPTION OFF
GO

-- Wait until the encryption state of the database becomes 1. The state is 5 (Decryption in
-- progress) for a while
SELECT db_name(database_id) as DatabaseName, * FROM sys.dm_database_encryption_keys
GO

-- Drop the DEK used for encryption
DROP DATABASE ENCRYPTION KEY
GO

-- Alter to SIMPLE Recovery mode so that your encrypted log gets truncated
USE [master]
GO
ALTER DATABASE [customerDatabase] SET RECOVERY SIMPLE
GO
```

When all objects are unencrypted, you can modify the DB instance to be associated with an option group without the TDE option or you can remove the TDE option from the option group.

SQL Server Performance Considerations

The performance of a SQL Server DB instance can be impacted by using transparent data encryption.

Performance for unencrypted databases can also be degraded if the databases are on a DB instance that has at least one encrypted database. As a result, we recommend that you keep encrypted and unencrypted databases on separate DB instances.

Because of the nature of encryption, the database size and the size of the transaction log is larger than for an unencrypted database. You could run over your allocation of free backup space. The nature of TDE causes an unavoidable performance hit. If you need high performance and TDE, measure the impact and make sure that it meets your needs. There is less of an impact on performance if you use Provisioned IOPS and at least an M3.Large DB instance class.

SQL Server Audit

In Amazon RDS, you can audit Microsoft SQL Server databases by using the built-in SQL Server auditing mechanism. You can create audits and audit specifications in the same way that you create them for on-premises database servers.

RDS uploads the completed audit logs to your S3 bucket, using the IAM role that you provide. If you enable retention, RDS keeps your audit logs on your DB instance for the configured period of time.

For more information, see [SQL Server Audit \(Database Engine\)](#) in the Microsoft SQL Server documentation.

Note

SQL Server Audit is supported in all regions, except for Asia Pacific (Hong Kong).

Topics

- [Support for SQL Server Audit \(p. 717\)](#)
- [Adding SQL Server Audit to the DB Instance Options \(p. 718\)](#)
- [Using SQL Server Audit \(p. 719\)](#)
- [Viewing Audit Logs \(p. 720\)](#)
- [Using SQL Server Audit with Multi-AZ Instances \(p. 720\)](#)
- [Configuring an S3 Bucket \(p. 720\)](#)
- [Manually Creating an IAM Role for SQL Server Audit \(p. 721\)](#)

Support for SQL Server Audit

In Amazon RDS, starting with SQL Server 2012, all editions of SQL Server support server-level audits, and the Enterprise edition also supports database-level audits. Starting with SQL Server 2016 (13.x) SP1, all editions support both server-level and database-level audits. For more information, see [SQL Server Audit \(Database Engine\)](#) in the SQL Server documentation.

RDS supports configuring the following option settings for SQL Server Audit.

Option Setting	Valid Values	Description
IAM_ROLE_ARN	A valid Amazon Resource Name (ARN) in the format <code>arn:aws:iam::account-id:role/role-name</code> .	The ARN of the IAM role that grants access to the S3 bucket where you want to store your audit logs. For more information, see AWS Identity and Access Management (IAM) .

Option Setting	Valid Values	Description
S3_BUCKET_ARN	A valid ARN in the format <code>arn:aws:s3:::bucket-name</code> or <code>arn:aws:s3:::bucket-name/key-prefix</code>	The ARN for the S3 bucket where you want to store your audit logs.
ENABLE_COMPRESSION	true or false	Controls audit log compression. By default, compression is enabled (set to true).
RETENTION_TIME	0 to 840	The retention time (in hours) that SQL Server audit records are kept on your RDS instance. By default, retention is disabled.

Adding SQL Server Audit to the DB Instance Options

Enabling SQL Server Audit requires two steps: enabling the option on the DB instance, and enabling the feature inside SQL Server. The process for adding the SQL Server Audit option to a DB instance is as follows:

1. Create a new option group, or copy or modify an existing option group.
2. Add and configure all required options.
3. Associate the option group with the DB instance.

After you add the SQL Server Audit option, you don't need to restart your DB instance. As soon as the option group is active, you can create audits and store audit logs in your S3 bucket.

To add and configure SQL Server Audit on a DB instance's option group

1. Choose one of the following:
 - Use an existing option group.
 - Create a custom DB option group and use that option group. For more information, see [Creating an Option Group \(p. 197\)](#).
2. Add the **SQLSERVER_AUDIT** option to the option group, and configure the option settings. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#).
 - For **IAM role**, if you already have an IAM role with the required policies, you can choose that role. To create a new IAM role, choose [Create a New Role](#). For information about the required policies, see [Manually Creating an IAM Role for SQL Server Audit \(p. 721\)](#).
 - For **Select S3 destination**, if you already have an S3 bucket that you want to use, choose it. To create an S3 bucket, choose [Create a New S3 Bucket](#).
 - For **Enable Compression**, leave this option checked to compress audit files. Compression is enabled by default. To disable compression, clear **Enable Compression**.
 - For **Audit log retention**, to keep audit records on the DB instance, choose this option. Specify a retention time in hours. The maximum retention time is 35 days.
3. Apply the option group to a new or existing DB instance. Choose one of the following:
 - If you are creating a new DB instance, apply the option group when you launch the instance.
 - On an existing DB instance, apply the option group by modifying the instance and then attaching the new option group. For more information, see [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#).

Modifying the SQL Server Audit Option

After you enable the SQL Server Audit option, you can modify the settings. For information about how to modify option settings, see [Modifying an Option Setting \(p. 204\)](#).

Removing SQL Server Audit from the DB Instance Options

You can turn off the SQL Server Audit feature by disabling audits and then deleting the option.

To remove auditing

1. Disable all of the audit settings inside SQL Server. To learn where audits are running, query the SQL Server security catalog views. For more information, see [Security Catalog Views](#) in the Microsoft SQL Server documentation.
2. Delete the SQL Server Audit option from the DB instance. Choose one of the following:
 - Delete the SQL Server Audit option from the option group that the DB instance uses. This change affects all DB instances that use the same option group. For more information, see [Removing an Option from an Option Group \(p. 207\)](#).
 - Modify the DB instance, and then choose an option group without the SQL Server Audit option. This change affects only the DB instance that you modify. You can specify the default (empty) option group, or a different custom option group. For more information, see [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#).
3. After you delete the SQL Server Audit option from the DB instance, you don't need to restart the instance. Remove unneeded audit files from your S3 bucket.

Using SQL Server Audit

You can control server audits, server audit specifications, and database audit specifications the same way that you control them for on-premises database servers.

Creating Audits

You create server audits in the same way that you create them for on-premises database servers. For information about how to create server audits, see [CREATE SERVER AUDIT](#) in the Microsoft SQL Server documentation.

To avoid errors, adhere to the following limitations:

- Don't exceed the maximum number of supported server audits per instance of 50.
- Instruct SQL Server to write data to a binary file.
- Don't use RDS_ as a prefix in the server audit name.
- For FILEPATH, specify D:\rdsdbdata\SQLAudit.
- For MAXSIZE, specify a size between 2 MB and 50 MB.
- Don't configure MAX_ROLLOVER_FILES or MAX_FILES.
- Don't configure SQL Server to shut down the DB instance if it fails to write the audit record.

Creating Audit Specifications

You create server audit specifications and database audit specifications the same way that you create them for on-premises database servers. For information about creating audit specifications, see [CREATE SERVER AUDIT SPECIFICATION](#) and [CREATE DATABASE AUDIT SPECIFICATION](#) in the Microsoft SQL Server documentation.

To avoid errors, don't use RDS_ as a prefix in the name of the database audit specification or server audit specification.

Viewing Audit Logs

Your audit logs are stored in D:\rdsdbdata\SQLAudit.

After SQL Server finishes writing to an audit log file—when the file reaches its size limit—Amazon RDS uploads the file to your S3 bucket. If retention is enabled, Amazon RDS moves the file into the retention folder: D:\rdsdbdata\SQLAudit\transmitted.

For information about configuring retention, see [Adding SQL Server Audit to the DB Instance Options \(p. 718\)](#).

Audit records are kept on the DB instance until the audit log file is uploaded. You can view the audit records by running the following command.

```
SELECT *
FROM msdb.dbo.rds_fn_get_audit_file
      ('D:\rdsdbdata\SQLAudit\*.sqlaudit'
      , default
      , default )
```

You can use the same command to view audit records in your retention folder by changing the filter to D:\rdsdbdata\SQLAudit\transmitted*.sqlaudit.

```
SELECT *
FROM msdb.dbo.rds_fn_get_audit_file
      ('D:\rdsdbdata\SQLAudit\transmitted\*.sqlaudit'
      , default
      , default )
```

Using SQL Server Audit with Multi-AZ Instances

For Multi-AZ instances, the process for sending audit log files to Amazon S3 is similar to the process for Single-AZ instances. However, there are some important differences:

- Database audit specification objects are replicated to all nodes.
- Server audits and server audit specifications aren't replicated to secondary nodes. Instead, you have to create or modify them manually.

To capture server audits or a server audit specification from both nodes:

1. Create a server audit or a server audit specification on the primary node.
2. Fail over to the secondary node and create a server audit or a server audit specification with the same name and GUID on the secondary node. Use the AUDIT_GUID parameter to specify the GUID.

Configuring an S3 Bucket

The audit log files are automatically uploaded from the DB instance to your S3 bucket. The following restrictions apply to the S3 bucket that you use as a target for audit files:

- It must be in the same AWS Region as the DB instance.
- It must not be open to the public.
- The bucket owner must also be the IAM role owner.

The target key that is used to store the data follows this naming schema: `bucket-name/key-prefix/instance-name/audit-name/node_file-name.ext`

Note

You set both the bucket name and the key prefix values with the `(S3_BUCKET_ARN` option setting.

The schema is composed of the following elements:

- **bucket-name** – The name of your S3 bucket.
- **key-prefix** – The custom key prefix you want to use for audit logs.
- **instance-name** – The name of your Amazon RDS instance.
- **audit-name** – The name of the audit.
- **node** – The identifier of the node that is the source of the audit logs (`node1` or `node2`). There is one node for a Single-AZ instance and two replication nodes for a Multi-AZ instance. These are not primary and secondary nodes, because the roles of primary and secondary change over time. Instead, the node identifier is a simple label.
 - **node1** – The first replication node (Single-AZ has one node only).
 - **node2** – The second replication node (Multi-AZ has two nodes).
- **file-name** – The target file name. The file name is taken as-is from SQL Server.
- **ext** – The extension of the file (`zip` or `sqlaudit`):
 - **zip** – If compression is enabled (default).
 - **sqlaudit** – If compression is disabled.

Manually Creating an IAM Role for SQL Server Audit

Typically, when you create a new option, the AWS Management Console creates the IAM role and the IAM trust policy for you. However, you can manually create a new IAM role to use with SQL Server Audits, so that you can customize it with any additional requirements you might have. To do this, you create an IAM role and delegate permissions so that the Amazon RDS service can use your Amazon S3 bucket. When you create this IAM role, you attach trust and permissions policies. The trust policy allows Amazon RDS to assume this role. The permission policy defines the actions that this role can do. For more information, see [Creating a Role to Delegate Permissions to an AWS Service](#) in the *AWS Identity and Access Management User Guide*.

You can use the examples in this section to create the trust and permissions policies you need.

The following example shows a trust policy for SQL Server Audit. The policy uses the *service principal* `rds.amazonaws.com` to allow RDS to write to the S3 bucket. A *service principal* is an identifier that is used to grant permissions to a service. Anytime you allow access to `rds.amazonaws.com` in this way, you are allowing RDS to perform an action on your behalf. For more information about service principals, see [AWS JSON Policy Elements: Principal](#).

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Principal": {  
                "Service": "rds.amazonaws.com"  
            },  
            "Action": "sts:AssumeRole"  
        }  
    ]  
}
```

In the following example of a permissions policy for SQL Server Audit, we specify an Amazon Resource Name (ARN) for the Amazon S3 bucket. You can use ARNs to identify a specific account, user, or role that you want grant access to. For more information about using ARNs, see [Amazon Resource Names \(ARNs\)](#).

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": "s3>ListAllMyBuckets",  
            "Resource": "*"  
        },  
        {  
            "Effect": "Allow",  
            "Action": [  
                "s3>ListBucket",  
                "s3>GetBucketACL",  
                "s3>GetBucketLocation"  
            ],  
            "Resource": "arn:aws:s3:::bucket_name"  
        },  
        {  
            "Effect": "Allow",  
            "Action": [  
                "s3>PutObject",  
                "s3>ListMultipartUploadParts",  
                "s3>AbortMultipartUpload"  
            ],  
            "Resource": "arn:aws:s3:::bucket_name/key_prefix/*"  
        }  
    ]  
}
```

Common DBA Tasks for Microsoft SQL Server

This section describes the Amazon RDS-specific implementations of some common DBA tasks for DB instances that are running the Microsoft SQL Server database engine. In order to deliver a managed service experience, Amazon RDS does not provide shell access to DB instances, and it restricts access to certain system procedures and tables that require advanced privileges.

Note

When working with a SQL Server DB instance, you can run scripts to modify a newly created database, but you cannot modify the [model] database, the database used as the model for new databases.

Topics

- [Accessing the tempdb Database on Microsoft SQL Server DB Instances on Amazon RDS \(p. 724\)](#)
- [Analyzing Your Database Workload on an Amazon RDS DB Instance with SQL Server Tuning Advisor \(p. 726\)](#)
- [Collations and Character Sets for Microsoft SQL Server \(p. 728\)](#)
- [Determining a Recovery Model for Your Microsoft SQL Server Database \(p. 730\)](#)
- [Dropping a Microsoft SQL Server Database \(p. 731\)](#)
- [Using Change Data Capture \(p. 731\)](#)
- [Renaming a Microsoft SQL Server Database in a Multi-AZ Deployment \(p. 733\)](#)
- [Resetting the db_owner Role Password \(p. 734\)](#)
- [Restoring License-Terminated DB Instances \(p. 734\)](#)
- [Transitioning a Microsoft SQL Server Database from OFFLINE to ONLINE \(p. 735\)](#)
- [Using SQL Server Agent \(p. 735\)](#)
- [Working with Microsoft SQL Server Logs \(p. 736\)](#)
- [Working with Trace and Dump Files \(p. 737\)](#)

Accessing the tempdb Database on Microsoft SQL Server DB Instances on Amazon RDS

You can access the tempdb database on your Microsoft SQL Server DB instances on Amazon RDS. You can run code on tempdb by using Transact-SQL through Microsoft SQL Server Management Studio (SSMS), or any other standard SQL client application. For more information about connecting to your DB instance, see [Connecting to a DB Instance Running the Microsoft SQL Server Database Engine \(p. 630\)](#).

The master user for your DB instance is granted CONTROL access to tempdb so that this user can modify the tempdb database options. The master user isn't the database owner of the tempdb database. If necessary, the master user can grant CONTROL access to other users so that they can also modify the tempdb database options.

Note

You can't run Database Console Commands (DBCC) on the tempdb database.

Modifying tempdb Database Options

You can modify the database options on the tempdb database on your Amazon RDS DB instances. For more information about which options can be modified, see [tempdb Database](#) in the Microsoft documentation.

Database options such as the maximum file size options are persistent after you restart your DB instance. You can modify the database options to optimize performance when importing data, and to prevent running out of storage.

Optimizing Performance when Importing Data

To optimize performance when importing large amounts of data into your DB instance, set the SIZE and FILEGROWTH properties of the tempdb database to large numbers. For more information about how to optimize tempdb, see [Optimizing tempdb Performance](#) in the Microsoft documentation.

The following example demonstrates setting the size to 100 GB and file growth to 10 percent.

```
alter database[tempdb] modify file (NAME = N'templog', SIZE=100GB, FILEGROWTH = 10%)
```

Preventing Storage Problems

To prevent the tempdb database from using all available disk space, set the MAXSIZE property. The following example demonstrates setting the property to 2048 MB.

```
alter database [tempdb] modify file (NAME = N'templog', MAXSIZE = 2048MB)
```

Shrinking the tempdb Database

There are two ways to shrink the tempdb database on your Amazon RDS DB instance. You can use the `rds_shrink_tempdbfile` procedure, or you can set the SIZE property,

Using the `rds_shrink_tempdbfile` Procedure

You can use the Amazon RDS procedure `msdb.dbo.rds_shrink_tempdbfile` to shrink the tempdb database. You can only call `rds_shrink_tempdbfile` if you have CONTROL access to tempdb. When you call `rds_shrink_tempdbfile`, there is no downtime for your DB instance.

The `rds_shrink_tempdbfile` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>@temp_filename</code>	SYSNAME	—	required	The logical name of the file to shrink.
<code>@target_size</code>	int	null	optional	The new size for the file, in megabytes.

The following example gets the names of the files for the tempdb database.

```
use tempdb;
GO

select name, * from sys.sysfiles;
GO
```

The following example shrinks a tempdb database file named `test_file`, and requests a new size of 10 megabytes:

```
exec msdb.dbo.rds_shrink_tempdbfile @temp_filename = N'test_file', @target_size = 10;
```

Setting the SIZE Property

You can also shrink the tempdb database by setting the `SIZE` property and then restarting your DB instance. For more information about restarting your DB instance, see [Rebooting a DB Instance \(p. 166\)](#).

The following example demonstrates setting the `SIZE` property to 1024 MB.

```
alter database [tempdb] modify file (NAME = N'templog', SIZE = 1024MB)
```

Considerations for Multi-AZ Deployments

If your Amazon RDS DB instance is in a Multi-AZ Deployment for Microsoft SQL Server with Database Mirroring (DBM) or Always On Availability Groups (AGs), there are some things to consider.

The tempdb database can't be replicated. No data that you store on your primary instance is replicated to your secondary instance.

If you modify any database options on the tempdb database, you can capture those changes on the secondary by using one of the following methods:

- First modify your DB instance and turn Multi-AZ off, then modify tempdb, and finally turn Multi-AZ back on. This method doesn't involve any downtime.

For more information, see [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#).

- First modify tempdb in the original primary instance, then fail over manually, and finally modify tempdb in the new primary instance. This method involves downtime.

For more information, see [Rebooting a DB Instance \(p. 166\)](#).

Analyzing Your Database Workload on an Amazon RDS DB Instance with SQL Server Tuning Advisor

The Database Engine Tuning Advisor is a client application provided by Microsoft that analyzes database workload and recommends an optimal set of indexes for your Microsoft SQL Server databases based on the kinds of queries you run. Like SQL Server Management Studio, you run Tuning Advisor from a client computer that connects to your Amazon RDS DB instance that is running SQL Server. The client computer can be a local computer that you run on premises within your own network or it can be an Amazon EC2 Windows instance that is running in the same region as your Amazon RDS DB instance.

This section shows how to capture a workload for Tuning Advisor to analyze. This is the preferred process for capturing a workload because Amazon RDS restricts host access to the SQL Server instance. The full documentation on Tuning Advisor can be found on [MSDN](#).

To use Tuning Advisor, you must provide what is called a workload to the advisor. A workload is a set of Transact-SQL statements that execute against a database or databases that you want to tune. Database Engine Tuning Advisor uses trace files, trace tables, Transact-SQL scripts, or XML files as workload input when tuning databases. When working with Amazon RDS, a workload can be a file on a client computer or a database table on an Amazon RDS SQL Server DB accessible to your client computer. The file or the table must contain queries against the databases you want to tune in a format suitable for replay.

For Tuning Advisor to be most effective, a workload should be as realistic as possible. You can generate a workload file or table by performing a trace against your DB instance. While a trace is running, you can either simulate a load on your DB instance or run your applications with a normal load.

There are two types of traces: client-side and server-side. A client-side trace is easier to set up and you can watch trace events being captured in real-time in SQL Server Profiler. A server-side trace is more complex to set up and requires some Transact-SQL scripting. In addition, because the trace is written to a file on the Amazon RDS DB instance, storage space is consumed by the trace. It is important to track of how much storage space a running server-side trace uses because the DB instance could enter a storage-full state and would no longer be available if it runs out of storage space.

For a client-side trace, when a sufficient amount of trace data has been captured in the SQL Server Profiler, you can then generate the workload file by saving the trace to either a file on your local computer or in a database table on a DB instance that is available to your client computer. The main disadvantage of using a client-side trace is that the trace may not capture all queries when under heavy loads. This could weaken the effectiveness of the analysis performed by the Database Engine Tuning Advisor. If you need to run a trace under heavy loads and you want to ensure that it captures every query during a trace session, you should use a server-side trace.

For a server-side trace, you must get the trace files on the DB instance into a suitable workload file or you can save the trace to a table on the DB instance after the trace completes. You can use the SQL Server Profiler to save the trace to a file on your local computer or have the Tuning Advisor read from the trace table on the DB instance.

Running a Client-Side Trace on a SQL Server DB Instance

To run a client-side trace on a SQL Server DB instance

1. Start SQL Server Profiler. It is installed in the Performance Tools folder of your SQL Server instance folder. You must load or define a trace definition template to start a client-side trace.
2. In the SQL Server Profiler File menu, choose **New Trace**. In the **Connect to Server** dialog box, enter the DB instance endpoint, port, master user name, and password of the database you would like to run a trace on.
3. In the **Trace Properties** dialog box, enter a trace name and choose a trace definition template. A default template, `TSQL_Replay`, ships with the application. You can edit this template to define your

trace. Edit events and event information under the **Events Selection** tab of the **Trace Properties** dialog box. For more information about trace definition templates and using the SQL Server Profiler to specify a client-side trace see the documentation in [MSDN](#).

4. Start the client-side trace and watch SQL queries in real-time as they execute against your DB instance.
5. Select **Stop Trace** from the **File** menu when you have completed the trace. Save the results as a file or as a trace table on your DB instance.

Running a Server-Side Trace on a SQL Server DB Instance

Writing scripts to create a server-side trace can be complex and is beyond the scope of this document. This section contains sample scripts that you can use as examples. As with a client-side trace, the goal is to create a workload file or trace table that you can open using the Database Engine Tuning Advisor.

The following is an abridged example script that starts a server-side trace and captures details to a workload file. The trace initially saves to the file RDSTrace.trc in the D:\RDSDBDATA\Log directory and rolls-over every 100 MB so subsequent trace files are named RDSTrace_1.trc, RDSTrace_2.trc, etc.

```
DECLARE @file_name NVARCHAR(245) = 'D:\RDSDBDATA\Log\RDSTrace';
DECLARE @max_file_size BIGINT = 100;
DECLARE @on BIT = 1
DECLARE @rc INT
DECLARE @traceid INT

EXEC @rc = sp_trace_create @traceid OUTPUT, 2, @file_name, @max_file_size
IF (@rc = 0) BEGIN
    EXEC sp_trace_setevent @traceid, 10, 1, @on
    EXEC sp_trace_setevent @traceid, 10, 2, @on
    EXEC sp_trace_setevent @traceid, 10, 3, @on
    .
    .
    EXEC sp_trace_setfilter @traceid, 10, 0, 7, N'SQL Profiler'
    EXEC sp_trace_setstatus @traceid, 1
END
```

The following example is a script that stops a trace. Note that a trace created by the previous script continues to run until you explicitly stop the trace or the process runs out of disk space.

```
DECLARE @traceid INT
SELECT @traceid = traceid FROM ::fn_trace_getinfo(default)
WHERE property = 5 AND value = 1 AND traceid <> 1

IF @traceid IS NOT NULL BEGIN
    EXEC sp_trace_setstatus @traceid, 0
    EXEC sp_trace_setstatus @traceid, 2
END
```

You can save server-side trace results to a database table and use the database table as the workload for the Tuning Advisor by using the fn_trace_gettable function. The following commands load the results of all files named RDSTrace.trc in the D:\rdsdbdata\Log directory, including all rollover files like RDSTrace_1.trc, into a table named RDSTrace in the current database.

```
SELECT * INTO RDSTrace
FROM fn_trace_gettable('D:\rdsdbdata\Log\RDSTrace.trc', default);
```

To save a specific rollover file to a table, for example the RDSTrace_1.trc file, specify the name of the rollover file and substitute 1 instead of default as the last parameter to fn_trace_gettable.

```
SELECT * INTO RDSTrace_1
FROM fn_trace_gettable('D:\rdsdbdata\Log\RDSTrace_1.trc', 1);
```

Running Tuning Advisor with a Trace

Once you create a trace, either as a local file or as a database table, you can then run Tuning Advisor against your DB instance. Microsoft includes documentation on using the Database Engine Tuning Advisor in [MSDN](#). Using Tuning Advisor with Amazon RDS is the same process as when working with a standalone, remote SQL Server instance. You can either use the Tuning Advisor UI on your client machine or use the dta.exe utility from the command line. In both cases, you must connect to the Amazon RDS DB instance using the endpoint for the DB instance and provide your master user name and master user password when using Tuning Advisor.

The following code example demonstrates using the dta.exe command line utility against an Amazon RDS DB instance with an endpoint of **dta.cnazcmklse1.us-east-1.rds.amazonaws.com**. The example includes the master user name **admin** and the master user password **test**, the example database to tune is named **RDSDTA** and the input workload is a trace file on the local machine named **C:\RDSTrace.trc**. The example command line code also specifies a trace session named **RDSTrace1** and specifies output files to the local machine named **RDSTrace.sql** for the SQL output script, **RDSTrace.txt** for a result file, and **RDSTrace.xml** for an XML file of the analysis. There is also an error table specified on the RDSDTA database named **RDSTraceErrors**.

```
dta -S dta.cnazcmklse1.us-east-1.rds.amazonaws.com -U admin -P test -D RDSDTA -if C:
\RDSTrace.trc -s RDSTrace1 -of C:\ RDSTrace.sql -or C:\ RDSTrace.txt -ox C:\ RDSTrace.xml -
e RDSDTA.dbo.RDSTraceErrors
```

Here is the same example command line code except the input workload is a table on the remote Amazon RDS instance named **RDSTrace** which is on the **RDSDTA** database.

```
dta -S dta.cnazcmklse1.us-east-1.rds.amazonaws.com -U admin -P test -D RDSDTA -it
RDSDTA.dbo.RDSTrace -s RDSTrace1 -of C:\ RDSTrace.sql -or C:\ RDSTrace.txt -ox C:\ RDSTrace.xml -
e RDSDTA.dbo.RDSTraceErrors
```

A full list of dta utility command-line parameters can be found in [MSDN](#).

Collations and Character Sets for Microsoft SQL Server

SQL Server supports collations at multiple levels. You set the default server collation when you create the DB instance. You can override the collation in the database, table, or column level.

Topics

- [Server-Level Collation for Microsoft SQL Server \(p. 728\)](#)
- [Database-Level Collation for Microsoft SQL Server \(p. 730\)](#)

Server-Level Collation for Microsoft SQL Server

When you create a Microsoft SQL Server DB instance, you can set the server collation that you want to use. If you don't choose a different collation, the server-level collation defaults to **SQL_Latin1_General_CI_AS**. The server collation is applied by default to all databases and database objects.

Currently, Amazon RDS supports the following server collations:

Collation	Description
Chinese_PRC_CI_AS	Chinese-PRC, case-insensitive, accent-sensitive, kanatype-insensitive, width-insensitive
Chinese_Taiwan_Stroke_CI_AS	Chinese-Taiwan-Stroke, case-insensitive, accent-sensitive, kanatype-insensitive, width-insensitive
French_CI_AS	French, case-insensitive, accent-sensitive, kanatype-insensitive, width-insensitive
Hebrew_BIN	Hebrew, binary sort
Japanese_CI_AS	Japanese, case-insensitive, accent-sensitive, kanatype-insensitive, width-insensitive
Korean_Wansung_CI_AS	Korean-Wansung, case-insensitive, accent-sensitive, kanatype-insensitive, width-insensitive
Latin1_General_100_BIN	Latin1-General-100, binary sort
Latin1_General_100_BIN2	Latin1-General-100, binary code point comparison sort
Latin1_General_BIN	Latin1-General, binary sort
Latin1_General_CI_AI	Latin1-General, case-insensitive, accent-insensitive, kanatype-insensitive, width-insensitive
Latin1_General_CI_AS	Latin1-General, case-insensitive, accent-sensitive, kanatype-insensitive, width-insensitive
Latin1_General_CS_AS	Latin1-General, case-sensitive, accent-sensitive, kanatype-insensitive, width-insensitive
Modern_Spanish_CI_AS	Modern-Spanish, case-insensitive, accent-sensitive, kanatype-insensitive, width-insensitive
SQL_Latin1_General_CP1_CI_AI	Latin1-General, case-insensitive, accent-insensitive, kanatype-insensitive, width-insensitive for Unicode Data, SQL Server Sort Order 54 on Code Page 1252 for non-Unicode Data
SQL_Latin1_General_CP1_CI_AS (default)	Latin1-General, case-insensitive, accent-sensitive, kanatype-insensitive, width-insensitive for Unicode Data, SQL Server Sort Order 52 on Code Page 1252 for non-Unicode Data
SQL_Latin1_General_CP1_CS_AS	Latin1-General, case-sensitive, accent-sensitive, kanatype-insensitive, width-insensitive for Unicode Data, SQL Server Sort Order 51 on Code Page 1252 for non-Unicode Data
SQL_Latin1_General_CP437_CI_AI	Latin1-General, case-insensitive, accent-insensitive, kanatype-insensitive, width-insensitive for Unicode Data, SQL Server Sort Order 34 on Code Page 437 for non-Unicode Data

Collation	Description
SQL_Latin1_General_CP850_BIN2	Latin1-General, binary code point comparison sort for Unicode Data, SQL Server Sort Order 40 on Code Page 850 for non-Unicode Data
SQL_Latin1_General_CP850_CI_AS	Latin1-General, case-insensitive, accent-sensitive, kanatype-insensitive, width-insensitive for Unicode Data, SQL Server Sort Order 42 on Code Page 850 for non-Unicode Data

To choose the collation:

- If you're using the Amazon RDS console, when creating a new DB instance choose **Additional configuration**, then choose the collation from the **Collation** menu under **Database options**. For more information, see [Creating a DB Instance Running the Microsoft SQL Server Database Engine \(p. 616\)](#).
- If you're using the AWS CLI, use the --character-set-name option with the `create-db-instance` command. For more information, see [create-db-instance](#).
- If you're using the Amazon RDS API, use the `CharacterSetName` parameter with the `CreateDBInstance` operation. For more information, see [CreateDBInstance](#).

Database-Level Collation for Microsoft SQL Server

You can change the default collation at the database, table, or column level by overriding the collation when creating a new database or database object. For example, if your default server collation is SQL_Latin1_General_CI_AS, you can change it to Mohawk_100_CI_AS for Mohawk collation support. Even arguments in a query can be type-cast to use a different collation if necessary.

For example, the following query would change the default collation for the AccountName column to Mohawk_100_CI_AS

```
CREATE TABLE [dbo].[Account]
(
    [AccountID] [nvarchar](10) NOT NULL,
    [AccountName] [nvarchar](100) COLLATE Mohawk_100_CI_AS NOT NULL
) ON [PRIMARY];
```

The Microsoft SQL Server DB engine supports Unicode by the built-in NCHAR, NVARCHAR, and NTEXT data types. For example, if you need CJK support, use these Unicode data types for character storage and override the default server collation when creating your databases and tables. Here are several links from Microsoft covering collation and Unicode support for SQL Server:

- [Working with Collations](#)
- [Collation and International Terminology](#)
- [Using SQL Server Collations](#)
- [International Considerations for Databases and Database Engine Applications](#)

Determining a Recovery Model for Your Microsoft SQL Server Database

In Amazon RDS, the recovery model, retention period, and database status are linked.

It's important to understand the consequences before making a change to one of these settings. Each setting can affect the others. For example:

- If you change a database's recovery model to SIMPLE while backup retention is enabled, Amazon RDS resets the recovery model back to FULL within five minutes. This also results in Amazon RDS taking a snapshot of the DB instance.
- If you set the backup retention to "0" days, Amazon RDS sets the recovery mode to SIMPLE.
- If you change a database's recovery model from SIMPLE to any other option while backup retention is set to "0" days, Amazon RDS resets the recovery model back to SIMPLE.

Important

Never change the recovery model on Multi-AZ instances, even if it seems you can do so—for example, by using ALTER DATABASE. Backup retention, and therefore FULL recovery mode, is required for Multi-AZ. If you alter the recovery model, RDS immediately changes it back to FULL. This automatic reset forces RDS to completely rebuild the mirror. During this rebuild, the availability of the database is degraded for about 30–90 minutes until the mirror is ready for failover. The DB instance also experiences performance degradation in the same way it does during a conversion from Single-AZ to Multi-AZ. How long performance is degraded depends on the database storage size—the bigger the stored database, the longer the degradation.

Dropping a Microsoft SQL Server Database

You can drop a database on an Amazon RDS DB instance running Microsoft SQL Server in a Single-AZ or Multi-AZ deployment. To drop the database, use the following command:

```
--replace your-database-name with the name of the database you want to drop
EXECUTE msdb.dbo.rds_drop_database N'your-database-name'
```

Note

Use straight single quotes in the command. Smart quotes will cause an error.

After you use this procedure to drop the database, Amazon RDS drops all existing connections to the database and removes the database's backup history.

Using Change Data Capture

Amazon RDS supports change data capture (CDC) for your DB instances running Microsoft SQL Server. CDC captures changes that are made to the data in your tables. It stores metadata about each change, which you can access later. For more information about how CDC works, see [Change Data Capture](#) in the Microsoft documentation.

Before you use CDC with your Amazon RDS DB instances, enable it in the database by running msdb.dbo.rds_cdc_enable_db. You must have master user privileges to enable CDC in the Amazon RDS DB instance. After CDC is enabled, any user who is db_owner of that database can enable or disable CDC on tables in that database.

Important

During restores, CDC will be disabled. All of the related metadata is automatically removed from the database. This applies to snapshot restores, point-in-time restores, and SQL Server Native restores from S3. After performing one of these types of restores, you can re-enable CDC and re-specify tables to track.

```
--Enable CDC for RDS DB Instance
exec msdb.dbo.rds_cdc_enable_db '<database name>'
```

To disable CDC, msdb.dbo.rds_cdc_disable_db run .

```
--Disable CDC for RDS DB Instance
exec msdb.dbo.rds_cdc_disable_db '<database name>'
```

Topics

- [Tracking Tables with Change Data Capture \(p. 732\)](#)
- [Change Data Capture Jobs \(p. 732\)](#)
- [Change Data Capture for Multi-AZ Instances \(p. 733\)](#)

Tracking Tables with Change Data Capture

After CDC is enabled on the database, you can start tracking specific tables. You can choose the tables to track by running [sys.sp_cdc_enable_table](#).

```
--Begin tracking a table
exec sys.sp_cdc_enable_table
    @source_schema      = N'<source_schema>',
    , @source_name       = N'<source_name>',
    , @role_name         = N'<role name>'

--The following parameters are optional:

--, @capture_instance      = '<capture_instance>'
--, @supports_net_changes  = '<supports_net_changes>'
--, @index_name             = '<index_name>'
--, @captured_column_list   = '<captured_column_list>'
--, @filegroup_name         = '<filegroup_name>'
--, @allow_partition_switch = '<allow_partition_switch>'
;
```

To view the CDC configuration for your tables, run [sys.sp_cdc_help_change_data_capture](#).

```
--View CDC configuration
exec sys.sp_cdc_help_change_data_capture

--The following parameters are optional and must be used together.
--  '<schema name>', '<table name>'
;
```

For more information on CDC tables, functions, and stored procedures in SQL Server documentation, see the following:

- [Change Data Capture Stored Procedures \(Transact-SQL\)](#)
- [Change Data Capture Functions \(Transact-SQL\)](#)
- [Change Data Capture Tables \(Transact-SQL\)](#)

Change Data Capture Jobs

When you enable CDC, SQL Server creates the CDC jobs. Database owners (`db_owner`) can view, create, modify, and delete the CDC jobs. However, the RDS system account owns them. Therefore, the jobs aren't visible from native views, procedures, or in SQL Server Management Studio.

To control behavior of CDC in a database, use native SQL Server procedures such as [sp_cdc_enable_table](#) and [sp_cdc_start_job](#). To change CDC job parameters, like `maxtrans` and `maxscans`, you can use [sp_cdc_change_jobs](#).

To get more information regarding the CDC jobs, you can query the following dynamic management views:

- sys.dm_cdc_errors
- sys.dm_cdc_log_scan_sessions
- sysjobs
- sysjobhistory

Change Data Capture for Multi-AZ Instances

If you use CDC on a Multi-AZ instance, make sure the mirror's CDC job configuration matches the one on the principal. CDC jobs are mapped to the `database_id`. If the database IDs on the secondary are different from the principal, then the jobs won't be associated with the correct database. To try to prevent errors after failover, RDS drops and recreates the jobs on the new principal. The recreated jobs use the parameters that the principal recorded before failover.

Although this process runs quickly, it's still possible that the CDC jobs might run before RDS can correct them. Here are three ways to force parameters to be consistent between primary and secondary replicas:

- Use the same job parameters for all the databases that have CDC enabled.
- Before you change the CDC job configuration, convert the Multi-AZ instance to Single-AZ.
- Manually transfer the parameters whenever you change them on the principal.

To view and define the CDC parameters that are used to recreate the CDC jobs after a failover, use `rds_show_configuration` and `rds_set_configuration`.

The following example returns the value set for `cdc_capture_maxtrans`. For any parameter that is set to `RDS_DEFAULT`, RDS automatically configures the value.

```
-- Show configuration for each parameter on either primary and secondary replicas.  
exec rdsadmin.dbo.rds_show_configuration 'cdc_capture_maxtrans'
```

To set the configuration on the secondary, run `rdsadmin.dbo.rds_set_configuration`. This procedure sets the parameter values for all of the databases on the secondary server. These settings are used only after a failover. The following example sets the `maxtrans` for all CDC capture jobs to `1000`:

```
--To set values on secondary. These are used after failover.  
exec rdsadmin..rds_set_configuration 'cdc_capture_maxtrans' , 1000
```

To set the CDC job parameters on the principal, use `sys.sp_cdc_change_jobs` instead.

Renaming a Microsoft SQL Server Database in a Multi-AZ Deployment

To rename a Microsoft SQL Server database instance that uses Multi-AZ, use the following procedure:

1. First, turn off Multi-AZ for the DB instance.
2. Rename the database by running `rdsadmin.dbo.rds_modify_db_name`.
3. Then, turn on Multi-AZ Mirroring or Always On Availability Groups for the DB instance, to return it to its original state.

For more information, see [Adding Multi-AZ to a Microsoft SQL Server DB Instance \(p. 681\)](#).

Note

If your instance doesn't use Multi-AZ, you don't need to change any settings before or after running `rdsadmin.dbo.rds_modify_db_name`.

Example: In the following example, the `rdsadmin.dbo.rds_modify_db_name` stored procedure renames a database from **MOO** to **ZAR**. This is similar to running the statement `DDL ALTER DATABASE [MOO] MODIFY NAME = [ZAR]`.

```
EXEC rdsadmin.dbo.rds_modify_db_name N'MOO', N'ZAR'  
GO
```

Resetting the db_owner Role Password

If you lock yourself out of the `db_owner` role on your Microsoft SQL Server database, you can reset the `db_owner` role password by modifying the DB instance master password. By changing the DB instance master password, you can regain access to the DB instance, access databases using the modified password for the `db_owner`, and restore privileges for the `db_owner` role that may have been accidentally revoked. You can change the DB instance password by using the Amazon RDS console, the AWS CLI command [modify-db-instance](#), or by using the [ModifyDBInstance](#) operation. For more information about modifying a SQL Server DB instance, see [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#).

Restoring License-Terminated DB Instances

Microsoft has requested that some Amazon RDS customers who did not report their Microsoft License Mobility information terminate their DB instance. Amazon RDS takes snapshots of these DB instances, and you can restore from the snapshot to a new DB instance that has the License Included model.

You can restore from a snapshot of Standard Edition to either Standard Edition or Enterprise Edition.

You can restore from a snapshot of Enterprise Edition to either Standard Edition or Enterprise Edition.

To restore from a SQL Server snapshot after Amazon RDS has created a final snapshot of your instance:

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Snapshots**.
3. Choose the snapshot of your SQL Server DB instance. Amazon RDS creates a final snapshot of your DB instance. The name of the terminated instance snapshot is in the format '<name of instance>-final-snapshot'. For example, if your DB instance name is `mytest.cdxgahs1ksma.us-east-1.rds.com`, the final snapshot is called `mytest-final-snapshot` and is located in the same AWS Region as the original DB instance.
4. For **Actions**, choose **Restore Snapshot**.

The **Restore DB Instance** window appears.

5. For **License Model**, choose **license-included**.
6. Choose the SQL Server DB engine that you want to use.
7. For **DB Instance Identifier**, enter the name for the restored DB instance.
8. Choose **Restore DB Instance**.

For more information about restoring from a snapshot, see [Restoring from a DB Snapshot \(p. 266\)](#).

Transitioning a Microsoft SQL Server Database from OFFLINE to ONLINE

You can transition your Microsoft SQL Server database on an Amazon RDS DB instance from OFFLINE to ONLINE.

SQL Server method	Amazon RDS method
ALTER DATABASE <i>name</i> SET ONLINE;	EXEC rdsadmin.dbo.rds_set_database_online <i>name</i>

Using SQL Server Agent

With Amazon RDS, you can use SQL Server Agent on a DB instance running Microsoft SQL Server Standard, Web Edition, or Enterprise Edition. SQL Server Agent is a Microsoft Windows service that executes scheduled administrative tasks, which are called jobs. You can use SQL Server Agent to run T-SQL jobs to rebuild indexes, run corruption checks, and aggregate data in a SQL Server DB instance.

SQL Server Agent can run a job on a schedule, in response to a specific event, or on demand. For more information, see [SQL Server Agent](#) in the SQL Server documentation. You should avoid scheduling jobs to run during the maintenance and backup windows for your DB instance because these maintenance and backup processes that are launched by AWS could interrupt the job or cause it to be cancelled. Because Amazon RDS backs up your DB instance, you do not use SQL Server Agent to create backups.

To view the history of an individual SQL Server Agent job in the SQL Server Management Studio, you open Object Explorer, right-click the job, and then click **View History**.

Because SQL Server Agent is running on a managed host in a DB instance, there are some actions that are not supported. Running replication jobs and running command-line scripts by using ActiveX, Windows command shell, or Windows PowerShell are not supported. In addition, you cannot manually start, stop, or restart SQL Server Agent because its operation is managed by the host. Email notifications through SQL Server Agent are not available from a DB instance.

When you create a SQL Server DB instance, the master user name is enrolled in the SQLAgentUserRole role. To allow an additional login/user to use SQL Server Agent, you must log in as the master user and do the following.

1. Create another server-level login by using the `CREATE LOGIN` command.
2. Create a user in msdb using `CREATE USER` command, and then link this user to the login that you created in the previous step.
3. Add the user to the SQLAgentUserRole using the `sp_addrolemember` system stored procedure.

For example, suppose your master user name is `myawsmaster` and you want to give access to SQL Server Agent to a user named `theirname` with a password `theirpassword`. You would log in using the master user name and run the following commands.

```
--Initially set context to master database
USE [master];
GO
--Create a server-level login named theirname with password theirpassword
CREATE LOGIN [theirname] WITH PASSWORD = 'theirpassword';
GO
--Set context to msdb database
USE [msdb];
```

```
GO
--Create a database user named theirname and link it to server-level login theirname
CREATE USER [theirname] FOR LOGIN [theirname];
GO
--Added database user theirname in msdb to SQLAgentUserRole in msdb
EXEC sp_addrolemember [SQLAgentUserRole], [theirname];
```

To delete a SQL Server Agent job, run the following T-SQL statement.

```
EXEC msdb..sp_delete_job @job_name = '<job-name>';
```

Note

Don't use the UI in SQL Server Management Console (SSMS) to delete a SQL Server Agent job. If you do, you get an error message similar to the following:

```
The EXECUTE permission was denied on the object 'xp_regrid', database
'mssqlsystemresource', schema 'sys'.
```

This error occurs because, as a managed service, RDS is restricted from running procedures that access the Windows registry. When you use SSMS to delete the job, it tries to run a process that RDS isn't authorized to do.

Working with Microsoft SQL Server Logs

You can use the Amazon RDS console to view, watch, and download SQL Server Agent logs and Microsoft SQL Server error logs.

Watching Log Files

If you view a log in the Amazon RDS console, you can see its contents as they exist at that moment. Watching a log in the console opens it in a dynamic state so that you can see updates to it in near real time.

Only the latest log is active for watching. For example, suppose you have the logs shown following:

Name	Last Written	Size	view	watch	download
log/ERROR	January 14, 2015 at 5:17:35 AM UTC-8	6.1 kB	view	watch	download
log/ERROR.1	January 13, 2015 at 3:59:00 PM UTC-8	53.3 kB	view	watch	download
log/ERROR.2	January 12, 2015 at 3:59:00 PM UTC-8	5.9 kB	view	watch	download
log/ERROR.3	January 11, 2015 at 3:59:00 PM UTC-8	5.9 kB	view	watch	download
log/ERROR.4	January 10, 2015 at 3:59:00 PM UTC-8	5.9 kB	view	watch	download

Only log/ERROR, as the most recent log, is being actively updated. You can choose to watch others, but they are static and will not update.

Archiving Log Files

The Amazon RDS console shows logs for the past week through the current day. You can download and archive logs to keep them for reference past that time. One way to archive logs is to load them into an Amazon S3 bucket. For instructions on how to set up an Amazon S3 bucket and upload a file, see [Amazon S3 Basics](#) in the *Amazon Simple Storage Service Getting Started Guide* and click **Get Started**.

Using the rds_read_error_log Procedure

To view Microsoft SQL server error and agent logs, use the Amazon RDS stored procedure `rds_read_error_log` with the following parameters:

- **@index** – the version of the log to retrieve. The default value is 0, which retrieves the current error log. Specify 1 to retrieve the previous log, specify 2 to retrieve the one before that, and so on.
- **@type** – the type of log to retrieve. Specify 1 to retrieve an error log. Specify 2 to retrieve an agent log.

Example

The following example requests the current error log.

```
EXEC rdsadmin.dbo.rds_read_error_log @index = 0, @type = 1;
```

Working with Trace and Dump Files

This section describes working with trace files and dump files for your Amazon RDS DB instances running Microsoft SQL Server.

Generating a Trace SQL Query

```
declare @rc int
declare @TraceID int
declare @maxfilesize bigint

set @maxfilesize = 5

exec @rc = sp_trace_create @TraceID output, 0, N'D:\rdsdbdata\log\rdstest', @maxfilesize,
NULL
```

Viewing an Open Trace

```
select * from ::fn_trace_getinfo(default)
```

Viewing Trace Contents

```
select * from ::fn_trace_gettable('D:\rdsdbdata\log\rdstest.trc', default)
```

Setting the Retention Period for Trace and Dump Files

Trace and dump files can accumulate and consume disk space. By default, Amazon RDS purges trace and dump files that are older than seven days.

To view the current trace and dump file retention period, use the `rds_show_configuration` procedure, as shown in the following example.

```
exec rdsadmin..rds_show_configuration;
```

To modify the retention period for trace files, use the `rds_set_configuration` procedure and set the `tracefile retention` in minutes. The following example sets the trace file retention period to 24 hours.

```
exec rdsadmin..rds_set_configuration 'tracefile retention', 1440;
```

To modify the retention period for dump files, use the `rds_set_configuration` procedure and set the `dumpfile retention` in minutes. The following example sets the dump file retention period to 3 days.

```
exec rdsadmin..rds_set_configuration 'dumpfile retention', 4320;
```

For security reasons, you cannot delete a specific trace or dump file on a SQL Server DB instance. To delete all unused trace or dump files, set the retention period for the files to 0.

MySQL on Amazon RDS

Amazon RDS supports DB instances running several versions of MySQL. You can use the following major versions:

- MySQL 8.0
- MySQL 5.7
- MySQL 5.6
- MySQL 5.5

For more information about minor version support, see [MySQL on Amazon RDS Versions \(p. 742\)](#).

You first use the Amazon RDS management tools or interfaces to create an Amazon RDS MySQL DB instance. You can then resize the DB instance, authorize connections to the DB instance, create and restore from backups or snapshots, create Multi-AZ secondaries, create Read Replicas, and monitor the performance of the DB instance. You use standard MySQL utilities and applications to store and access the data in the DB instance.

Amazon RDS for MySQL is compliant with many industry standards. For example, you can use Amazon RDS for MySQL databases to build HIPAA-compliant applications and to store healthcare related information, including protected health information (PHI) under an executed Business Associate Agreement (BAA) with AWS. Amazon RDS for MySQL also meets Federal Risk and Authorization Management Program (FedRAMP) security requirements and has received a FedRAMP Joint Authorization Board (JAB) Provisional Authority to Operate (P-ATO) at the FedRAMP HIGH Baseline within the AWS GovCloud (US-West) Region. For more information on supported compliance standards, see [AWS Cloud Compliance](#).

For information about the features in each version of MySQL, see [The Main Features of MySQL](#) in the MySQL documentation.

Common Management Tasks for MySQL on Amazon RDS

The following are the common management tasks you perform with an Amazon RDS MySQL DB instance, with links to relevant documentation for each task.

Task Area	Relevant Documentation
Understanding Amazon RDS Understand key Amazon RDS components, including DB instances, regions, Availability Zones, security groups, parameter groups, and option groups.	What Is Amazon Relational Database Service (Amazon RDS)? (p. 1)
Setting up Amazon RDS for first time use Set up Amazon RDS so that you can create MySQL DB instances in Amazon Web Services (AWS).	Setting Up for Amazon RDS (p. 5)

Task Area	Relevant Documentation
Understanding Amazon RDS DB instances Create virtual MySQL server instances that run in AWS. Because DB instances are the building blocks of Amazon RDS, we recommend that you understand their principles.	Amazon RDS DB Instances (p. 105)
Creating a DB instance for production Create a DB instance for production purposes. Creating an instance includes choosing a DB instance class with appropriate processing power and memory capacity and choosing a storage type that supports the way you expect to use your database.	Choosing the DB Instance Class (p. 107) Amazon RDS Storage Types (p. 135) Creating a DB Instance Running the MySQL Database Engine (p. 751)
Managing security for your DB instance By default, DB instances are created with a firewall that prevents access to them. You must create a security group with the correct IP addresses and network configuration to access the DB instance. You can also use AWS Identity and Access Management (IAM) policies to assign permissions that determine who is allowed to manage RDS resources.	Security in Amazon RDS (p. 427) Managing Access Using Policies (p. 447) Controlling Access with Security Groups (p. 486) Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform (p. 506)
Connecting to your DB instance Connect to your DB instance using a standard SQL client application such as the MySQL command line utility or MySQL Workbench.	Connecting to a DB Instance Running the MySQL Database Engine (p. 764)
Configuring high availability for a production DB instance Provide high availability with synchronous standby replication in a different Availability Zone, automatic failover, fault tolerance for DB instances using Multi-AZ deployments, and Read Replicas.	High Availability (Multi-AZ) for Amazon RDS (p. 144)
Configuring a DB instance in an Amazon Virtual Private Cloud Configure a virtual private cloud (VPC) in the Amazon VPC service. An Amazon VPC is a virtual network logically isolated from other virtual networks in AWS.	Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform (p. 506) Working with a DB Instance in a VPC (p. 514)

Task Area	Relevant Documentation
Configuring specific MySQL database parameters and features Configure specific MySQL database parameters with a parameter group that can be associated with many DB instances. You can also configure specific MySQL database features with an option group that can be associated with many DB instances.	Working with DB Parameter Groups (p. 211) Working with Option Groups (p. 196) Options for MySQL DB Instances (p. 873)
Modifying a DB instance running the MySQL database engine Change the settings of a DB instance to accomplish tasks such as adding additional storage or changing the DB instance class.	Modifying a DB Instance Running the MySQL Database Engine (p. 802) Modifying an Amazon RDS DB Instance (p. 150)
Configuring database backup and restore Configure your DB instance to take automated backups. You can also back up and restore your databases manually by using full backup files.	Working With Backups (p. 254) Backing Up and Restoring Amazon RDS DB Instances (p. 253)
Importing and exporting data Import data from other RDS MySQL DB instances, MySQL instances running external to Amazon RDS, and other types of data sources, and export data to MySQL instances running external to Amazon RDS.	Restoring a Backup into an Amazon RDS MySQL DB Instance (p. 826)
Monitoring a MySQL DB instance Monitor your RDS MySQL DB instance by using Amazon CloudWatch RDS metrics, events, and Enhanced Monitoring. View log files for your RDS MySQL DB instance.	Monitoring Amazon RDS (p. 294) Viewing DB Instance Metrics (p. 304) Viewing Amazon RDS Events (p. 388) Amazon RDS Database Log Files (p. 390) MySQL Database Log Files (p. 405)
Replicating your data Create a MySQL Read Replica, in the same AWS Region or a different one. You can use Read Replicas for load balancing, disaster recovery, and processing read-heavy database workloads, such as for analysis and reporting.	Working with Read Replicas (p. 182) Replication with a MySQL or MariaDB Instance Running External to Amazon RDS (p. 862)

There are also several sections with useful information about working with Amazon RDS MySQL DB instances:

- [Common DBA Tasks for MySQL DB Instances \(p. 881\)](#)
- [Options for MySQL DB Instances \(p. 873\)](#)
- [MySQL on Amazon RDS SQL Reference \(p. 889\)](#)

MySQL on Amazon RDS Versions

For MySQL, version numbers are organized as version = X.Y.Z. In Amazon RDS terminology, X.Y denotes the major version, and Z is the minor version number. For Amazon RDS implementations, a version change is considered major if the major version number changes—for example, going from version 5.6 to 5.7. A version change is considered minor if only the minor version number changes—for example, going from version 5.7.16 to 5.7.21.

Amazon RDS currently supports the following versions of MySQL:

Major Version	Minor Version
MySQL 8.0	<ul style="list-style-type: none">8.0.168.0.158.0.138.0.11
MySQL 5.7	<ul style="list-style-type: none">5.7.265.7.255.7.245.7.235.7.225.7.215.7.195.7.175.7.16
MySQL 5.6	<ul style="list-style-type: none">5.6.445.6.435.6.415.6.405.6.395.6.375.6.355.6.34
MySQL 5.5	<ul style="list-style-type: none">5.5.615.5.595.5.575.5.545.5.535.5.46

You can specify any currently supported MySQL version when creating a new DB instance. You can specify the major version (such as MySQL 5.7), and any supported minor version for the specified major version. If no version is specified, Amazon RDS defaults to a supported version, typically the most recent version. If a major version is specified but a minor version is not, Amazon RDS defaults to a recent release of the major version you have specified. To see a list of supported versions, as well as defaults for newly created DB instances, use the `describe-db-engine-versions` AWS CLI command.

With Amazon RDS, you control when to upgrade your MySQL instance to a new version supported by Amazon RDS. You can maintain compatibility with specific MySQL versions, test new versions with your application before deploying in production, and perform version upgrades at times that best fit your schedule.

Unless you specify otherwise, your DB instance will automatically be upgraded to new MySQL minor versions as they are supported by Amazon RDS. This patching occurs during your scheduled maintenance window. You can modify a DB instance to turn off automatic minor version upgrades.

If you opt out of automatically scheduled upgrades, you can manually upgrade to a supported minor version release by following the same procedure as you would for a major version update. For information, see [Upgrading a DB Instance Engine Version \(p. 159\)](#).

Amazon RDS currently supports the major version upgrades from MySQL version 5.5 to version 5.6, from MySQL version 5.6 to version 5.7, and from MySQL version 5.7 to version 8.0. Because major version upgrades involve some compatibility risk, they do not occur automatically; you must make a request to modify the DB instance. You should thoroughly test any upgrade before upgrading your production instances. For information about upgrading a MySQL DB instance, see [Upgrading the MySQL DB Engine \(p. 812\)](#).

You can test a DB instance against a new version before upgrading by creating a DB snapshot of your existing DB instance, restoring from the DB snapshot to create a new DB instance, and then initiating a version upgrade for the new DB instance. You can then experiment safely on the upgraded clone of your DB instance before deciding whether or not to upgrade your original DB instance.

For information about the Amazon RDS deprecation policy for MySQL, see [Amazon RDS FAQs](#).

MySQL Features Not Supported By Amazon RDS

Amazon RDS doesn't currently support the following MySQL features:

- Authentication Plugin
- Error Logging to the System Log
- Group Replication Plugin
- InnoDB Tablespace Encryption
- MariaDB Audit Plugin (not supported for Amazon RDS MySQL version 8.0 only)

The MariaDB Audit Plugin is supported for Amazon RDS MySQL version 5.5, 5.6, and 5.7.

- Password Strength Plugin
- Persisted system variables
- Replication filters
- Semisynchronous replication
- Transportable tablespace
- X Plugin

Note

Global transaction IDs are supported for MySQL 5.7.23 and later MySQL 5.7 versions. Global transaction IDs are not supported for Amazon RDS MySQL 5.5, 5.6, or 8.0.

IAM database authentication is supported for MySQL 5.6 and 5.7. IAM database authentication is not supported for MySQL 5.5 or 8.0.

Amazon RDS Performance Insights is supported for MySQL 5.6 and 5.7. Amazon RDS Performance Insights is not supported for MySQL 5.5 or 8.0.

To deliver a managed service experience, Amazon RDS doesn't provide shell access to DB instances. It also restricts access to certain system procedures and tables that require advanced privileges. Amazon RDS supports access to databases on a DB instance using any standard SQL client application. Amazon RDS doesn't allow direct host access to a DB instance by using Telnet, Secure Shell (SSH), or Windows Remote Desktop Connection. When you create a DB instance, you are assigned to the *db_owner* role for all databases on that instance, and you have all database-level permissions except for those used for backups. Amazon RDS manages backups for you.

Supported Storage Engines for MySQL on Amazon RDS

While MySQL supports multiple storage engines with varying capabilities, not all of them are optimized for recovery and data durability. Amazon RDS fully supports the InnoDB storage engine for MySQL DB instances. Amazon RDS features such as Point-In-Time restore and snapshot restore require a recoverable storage engine and are supported for the InnoDB storage engine only. You must be running an instance of MySQL 5.6 or later to use the InnoDB memcached interface. For more information, see [MySQL memcached Support \(p. 877\)](#).

The Federated Storage Engine is currently not supported by Amazon RDS for MySQL.

For user-created schemas, the MyISAM storage engine does not support reliable recovery and can result in lost or corrupt data when MySQL is restarted after a recovery, preventing Point-In-Time restore or snapshot restore from working as intended. However, if you still choose to use MyISAM with Amazon RDS, snapshots can be helpful under some conditions.

Note

System tables in the `mysql` schema can be in MyISAM storage.

If you want to convert existing MyISAM tables to InnoDB tables, you can use the `ALTER TABLE` command (for example, `alter table TABLE_NAME engine=innodb;`). Bear in mind that MyISAM and InnoDB have different strengths and weaknesses, so you should fully evaluate the impact of making this switch on your applications before doing so.

MySQL 5.1 is no longer supported in Amazon RDS. However, you can restore existing MySQL 5.1 snapshots. When you restore a MySQL 5.1 snapshot, the instance is automatically upgraded to MySQL 5.5.

MySQL Security on Amazon RDS

Security for Amazon RDS MySQL DB instances is managed at three levels:

- AWS Identity and Access Management controls who can perform Amazon RDS management actions on DB instances. When you connect to AWS using IAM credentials, your IAM account must have IAM policies that grant the permissions required to perform Amazon RDS management operations. For more information, see [Identity and Access Management in Amazon RDS \(p. 445\)](#).
- When you create a DB instance, you use either a VPC security group or a DB security group to control which devices and Amazon EC2 instances can open connections to the endpoint and port of the DB instance. These connections can be made using Secure Sockets Layer (SSL). In addition, firewall rules at your company can control whether devices running at your company can open connections to the DB instance.
- To authenticate login and permissions for a MySQL DB instance, you can take either of the following approaches, or a combination of them.

You can take the same approach as with a stand-alone instance of MySQL. Commands such as `CREATE USER`, `RENAME USER`, `GRANT`, `REVOKE`, and `SET PASSWORD` work just as they do in on-premises databases, as does directly modifying database schema tables. For information, see [MySQL User Account Management](#) in the MySQL documentation.

You can also use IAM database authentication. With IAM database authentication, you authenticate to your DB instance by using an IAM user or IAM role and an authentication token. An *authentication token* is a unique value that is generated using the Signature Version 4 signing process. By using IAM database authentication, you can use the same credentials to control access to your AWS resources and your databases. For more information, see [IAM Database Authentication for MySQL and PostgreSQL \(p. 461\)](#).

When you create an Amazon RDS DB instance, the master user has the following default privileges:

- `alter`
- `alter routine`
- `create`
- `create routine`
- `create temporary tables`
- `create user`
- `create view`
- `delete`
- `drop`
- `event`
- `execute`
- `grant option`
- `index`
- `insert`
- `lock tables`
- `process`
- `references`
- `replication client`
- `replication slave` (MySQL 5.6 and later)
- `select`
- `show databases`
- `show view`
- `trigger`
- `update`

Note

Although it is possible to delete the master user on the DB instance, it is not recommended. To recreate the master user, use the [ModifyDBInstance](#) RDS API operation or the [modify-db-instance](#) AWS CLI command and specify a new master user password with the appropriate parameter. If the master user does not exist in the instance, the master user is created with the specified password.

To provide management services for each DB instance, the `rdsadmin` user is created when the DB instance is created. Attempting to drop, rename, change the password, or change privileges for the `rdsadmin` account will result in an error.

To allow management of the DB instance, the standard `kill` and `kill_query` commands have been restricted. The Amazon RDS commands `rds_kill` and `rds_kill_query` are provided to allow you to terminate user sessions or queries on DB instances.

Using the Password Validation Plugin

MySQL provides the `validate_password` plugin for improved security. The plugin enforces password policies using parameters in the DB parameter group for your MySQL DB instance. The plugin is supported for DB instances running MySQL version 5.6, 5.7, and 8.0. For more information about the `validate_password` plugin, see [The Password Validation Plugin](#) in the MySQL documentation.

To enable the `validate_password` plugin for a MySQL DB instance

1. Connect to your MySQL DB instance and run the following command.

```
INSTALL PLUGIN validate_password SONAME 'validate_password.so';
```

2. Configure the parameters for the plugin in the DB parameter group used by the DB instance.

For more information about the parameters, see [Password Validation Plugin Options and Variables](#) in the MySQL documentation.

For more information about modifying DB instance parameters, see [Modifying Parameters in a DB Parameter Group \(p. 213\)](#).

After installing and enabling the `password_validate` plugin, reset existing passwords to comply with your new validation policies.

Amazon RDS doesn't validate passwords. The MySQL DB instance performs password validation. If you set a user password with the AWS Management Console, the `modify-db-instance` AWS CLI command, or the `ModifyDBInstance` RDS API operation, the change can succeed even if the new password doesn't satisfy your password policies. However, a new password is set in the MySQL DB instance only if it satisfies the password policies. In this case, Amazon RDS records the following event.

```
"RDS-EVENT-0067" - An attempt to reset the master password for the DB instance has failed.
```

For more information about Amazon RDS events, see [Using Amazon RDS Event Notification \(p. 370\)](#).

Using SSL with a MySQL DB Instance

Amazon RDS supports Secure Sockets Layer (SSL) connections with DB instances running the MySQL database engine.

Amazon RDS creates an SSL certificate and installs the certificate on the DB instance when Amazon RDS provisions the instance. These certificates are signed by a certificate authority. The SSL certificate includes the DB instance endpoint as the Common Name (CN) for the SSL certificate to guard against spoofing attacks.

An SSL certificate created by Amazon RDS is the trusted root entity and should work in most cases but might fail if your application does not accept certificate chains. If your application does not accept

certificate chains, you might need to use an intermediate certificate to connect to your region. For example, you must use an intermediate certificate to connect to the AWS GovCloud (US-West) Region using SSL.

For information about downloading certificates, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#). For more information about using SSL with MySQL, see [Updating Applications to Connect to MySQL DB Instances Using New SSL/TLS Certificates \(p. 797\)](#).

MySQL uses yaSSL for secure connections in the following versions:

- MySQL version 5.7.19 and earlier 5.7 versions
- MySQL version 5.6.37 and earlier 5.6 versions
- MySQL version 5.5.57 and earlier 5.5 versions

MySQL uses OpenSSL for secure connections in the following versions:

- MySQL version 8.0
- MySQL version 5.7.21 and later 5.7 versions
- MySQL version 5.6.39 and later 5.6 versions
- MySQL version 5.5.59 and later 5.5 versions

Amazon RDS for MySQL supports Transport Layer Security (TLS) versions 1.0, 1.1, and 1.2. The following table shows the TLS support for MySQL versions.

MySQL Version	TLS 1.0	TLS 1.1	TLS 1.2
MySQL 8.0	Supported	Supported	Supported
MySQL 5.7	Supported	Supported	Supported for MySQL 5.7.21 and later
MySQL 5.6	Supported	Not supported	Not supported
MySQL 5.5	Supported	Not supported	Not supported

To encrypt connections using the default mysql client, launch the mysql client using the --ssl-ca parameter to reference the public key, as shown in the examples following.

The following example shows how to launch the client using the --ssl-ca parameter for MySQL 5.7 and later.

```
mysql -h myinstance.c9akciq32.rds-us-east-1.amazonaws.com  
--ssl-ca=[full path]rds-combined-ca-bundle.pem --ssl-mode=VERIFY_IDENTITY
```

The following example shows how to launch the client using the --ssl-ca parameter for MySQL 5.6 and earlier.

```
mysql -h myinstance.c9akciq32.rds-us-east-1.amazonaws.com  
--ssl-ca=[full path]rds-combined-ca-bundle.pem --ssl-verify-server-cert
```

You can require SSL connections for specific user accounts. For example, you can use one of the following statements, depending on your MySQL version, to require SSL connections on the user account encrypted_user.

For MySQL 5.7 and later, use the following statement.

```
ALTER USER 'encrypted_user'@'%' REQUIRE SSL;
```

For MySQL 5.6 and earlier, use the following statement.

```
GRANT USAGE ON *.* TO 'encrypted_user'@'%' REQUIRE SSL;
```

For more information on SSL connections with MySQL, see the [Using Encrypted Connections](#) in the MySQL documentation.

Using memcached and Other Options with MySQL

Most Amazon RDS DB engines support option groups that allow you to select additional features for your DB instance. DB instances on MySQL version 5.6 and later support the `memcached` option, a simple, key-based cache. For more information about `memcached` and other options, see [Options for MySQL DB Instances \(p. 873\)](#). For more information about working with option groups, see [Working with Option Groups \(p. 196\)](#).

InnoDB Cache Warming

InnoDB cache warming can provide performance gains for your MySQL DB instance by saving the current state of the buffer pool when the DB instance is shut down, and then reloading the buffer pool from the saved information when the DB instance starts up. This bypasses the need for the buffer pool to "warm up" from normal database use and instead preloads the buffer pool with the pages for known common queries. The file that stores the saved buffer pool information only stores metadata for the pages that are in the buffer pool, and not the pages themselves. As a result, the file does not require much storage space. The file size is about 0.2 percent of the cache size. For example, for a 64 GiB cache, the cache warming file size is 128 MiB. For more information on InnoDB cache warming, see [Saving and Restoring the Buffer Pool State](#) in the MySQL documentation.

MySQL on Amazon RDS supports InnoDB cache warming for MySQL version 5.6 and later. To enable InnoDB cache warming, set the `innodb_buffer_pool_dump_at_shutdown` and `innodb_buffer_pool_load_at_startup` parameters to 1 in the parameter group for your DB instance. Changing these parameter values in a parameter group will affect all MySQL DB instances that use that parameter group. To enable InnoDB cache warming for specific MySQL DB instances, you might need to create a new parameter group for those instances. For information on parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#).

InnoDB cache warming primarily provides a performance benefit for DB instances that use standard storage. If you use PIOPS storage, you do not commonly see a significant performance benefit.

Important

If your MySQL DB instance does not shut down normally, such as during a failover, then the buffer pool state will not be saved to disk. In this case, MySQL loads whatever buffer pool file is available when the DB instance is restarted. No harm is done, but the restored buffer pool might not reflect the most recent state of the buffer pool prior to the restart. To ensure that you have a recent state of the buffer pool available to warm the InnoDB cache on startup, we recommend that you periodically dump the buffer pool "on demand." You can dump or load the buffer pool on demand if your DB instance is running MySQL version 5.6.19 or later.

You can create an event to dump the buffer pool automatically and on a regular interval. For example, the following statement creates an event named `periodic_buffer_pool_dump` that dumps the buffer pool every hour.

```
CREATE EVENT periodic_buffer_pool_dump
ON SCHEDULE EVERY 1 HOUR
DO CALL mysql.rds_innodb_buffer_pool_dump_now();
```

For more information on MySQL events, see [Event Syntax](#) in the MySQL documentation.

Dumping and Loading the Buffer Pool on Demand

For MySQL version 5.6.19 and later, you can save and load the InnoDB cache "on demand."

- To dump the current state of the buffer pool to disk, call the [mysql.rds_innodb_buffer_pool_dump_now \(p. 907\)](#) stored procedure.
- To load the saved state of the buffer pool from disk, call the [mysql.rds_innodb_buffer_pool_load_now \(p. 907\)](#) stored procedure.
- To cancel a load operation in progress, call the [mysql.rds_innodb_buffer_pool_load_abort \(p. 907\)](#) stored procedure.

Local Time Zone for MySQL DB Instances

By default, the time zone for an RDS MySQL DB instance is Universal Time Coordinated (UTC). You can set the time zone for your DB instance to the local time zone for your application instead.

To set the local time zone for a DB instance, set the `time_zone` parameter in the parameter group for your DB instance to one of the supported values listed later in this section. When you set the `time_zone` parameter for a parameter group, all DB instances and Read Replicas that are using that parameter group change to use the new local time zone. For information on setting parameters in a parameter group, see [Working with DB Parameter Groups \(p. 211\)](#).

After you set the local time zone, all new connections to the database reflect the change. If you have any open connections to your database when you change the local time zone, you won't see the local time zone update until after you close the connection and open a new connection.

You can set a different local time zone for a DB instance and one or more of its Read Replicas. To do this, use a different parameter group for the DB instance and the replica or replicas and set the `time_zone` parameter in each parameter group to a different local time zone.

If you are replicating across regions, then the replication master DB instance and the Read Replica use different parameter groups (parameter groups are unique to a region). To use the same local time zone for each instance, you must set the `time_zone` parameter in the instance's and Read Replica's parameter groups.

When you restore a DB instance from a DB snapshot, the local time zone is set to UTC. You can update the time zone to your local time zone after the restore is complete. If you restore a DB instance to a point in time, then the local time zone for the restored DB instance is the time zone setting from the parameter group of the restored DB instance.

You can set your local time zone to one of the following values.

Africa/Cairo	Asia/Bangkok	Australia/Darwin
Africa/Casablanca	Asia/Beirut	Australia/Hobart

Africa/Harare	Asia/Calcutta	Australia/Perth
Africa/Monrovia	Asia/Damascus	Australia/Sydney
Africa/Nairobi	Asia/Dhaka	Brazil/East
Africa/Tripoli	Asia/Irkutsk	Canada/Newfoundland
Africa/Windhoek	Asia/Jerusalem	Canada/Saskatchewan
America/Araguaina	Asia/Kabul	Europe/Amsterdam
America/Asuncion	Asia/Karachi	Europe/Athens
America/Bogota	Asia/Kathmandu	Europe/Dublin
America/Caracas	Asia/Krasnoyarsk	Europe/Helsinki
America/Chihuahua	Asia/Magadan	Europe/Istanbul
America/Cuiaba	Asia/Muscat	Europe/Kaliningrad
America/Denver	Asia/Novosibirsk	Europe/Moscow
America/Fortaleza	Asia/Riyadh	Europe/Paris
America/Guatemala	Asia/Seoul	Europe/Prague
America/Halifax	Asia/Shanghai	Europe/Sarajevo
America/Manaus	Asia/Singapore	Pacific/Auckland
America/Matamoros	Asia/Taipei	Pacific/Fiji
America/Monterrey	Asia/Tehran	Pacific/Guam
America/Montevideo	Asia/Tokyo	Pacific/Honolulu
America/Phoenix	Asia/Ulaanbaatar	Pacific/Samoa
America/Santiago	Asia/Vladivostok	US/Alaska
America/Tijuana	Asia/Yakutsk	US/Central
Asia/Amman	Asia/Yerevan	US/Eastern
Asia/Ashgabat	Atlantic/Azores	US/East-Indiana
Asia/Baghdad	Australia/Adelaide	US/Pacific
Asia/Baku	Australia/Brisbane	UTC

Known Issues and Limitations for MySQL on Amazon RDS

There are some known issues and limitations for working with MySQL on Amazon RDS. For more information, see [Known Issues and Limitations for MySQL on Amazon RDS \(p. 885\)](#).

Creating a DB Instance Running the MySQL Database Engine

The basic building block of Amazon RDS is the DB instance. The DB instance is where you create your MySQL databases.

Important

Before you can create or connect to a DB instance, you must complete the tasks in [Setting Up for Amazon RDS \(p. 5\)](#).

For an example that walks you through the process of creating and connecting to a sample DB instance, see [Creating a MySQL DB Instance and Connecting to a Database on a MySQL DB Instance \(p. 33\)](#).

Note

If you are using the console, a new console interface is available for database creation. Choose either the **New Console** or the **Original Console** instructions based on the console that you are using. The **New Console** instructions are open by default.

New Console

You can create a DB instance running MySQL with the AWS Management Console with **Easy Create** enabled or not enabled. With **Easy Create** enabled, you specify only the DB engine type, DB instance size, and DB instance identifier. **Easy Create** uses the default setting for other configuration options. With **Easy Create** not enabled, you specify more configuration options when you create a database, including ones for availability, security, backups, and maintenance.

Note

For this example, **Standard Create** is enabled, and **Easy Create** isn't enabled. For information about creating a MySQL DB instance with **Easy Create** enabled, see [Creating a MySQL DB Instance and Connecting to a Database on a MySQL DB Instance \(p. 33\)](#).

To create a MySQL DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the AWS Management Console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
4. Choose **Create database**.
5. In **Choose a database creation method**, choose **Standard Create**.
6. In **Engine options**, choose **MySQL**.

Create database

Choose a database creation method Info

Standard Create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy Create

Use recommended best-practice configurations. Some configuration options can be changed after the database is created.

Engine options

Engine type Info

Amazon Aurora



MySQL



MariaDB

PostgreSQL



Oracle

ORACLE®

Microsoft SQL Server



Edition

MySQL Community

Version Info

MySQL 5.7.20

7. In **Templates**, choose the template that matches your use case. If you choose **Production**, the following are also chosen in a later step:

- Multi-AZ failover option

- **Provisioned IOPS** storage option
- **Enable deletion protection** option

We recommend these features for any production environment.

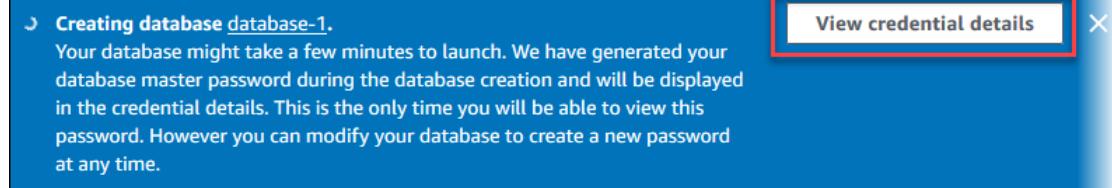
8. To enter your master password, do the following:
 - a. In the **Settings** section, open **Credential Settings**.
 - b. Clear the **Auto generate a password** check box.
 - c. (Optional) Change the **Master username** value and enter the same password in **Master password** and **Confirm password**.

By default, the new DB instance uses an automatically generated password for the master user.

9. For the remaining sections, specify your DB instance settings. For information about each setting, see [Settings for MySQL DB Instances \(p. 758\)](#).
10. Choose **Create database**.

If you chose to use an automatically generated password, the **View credential details** button appears on the **Databases** page.

To view the master user name and password for the DB instance, choose **View credential details**.



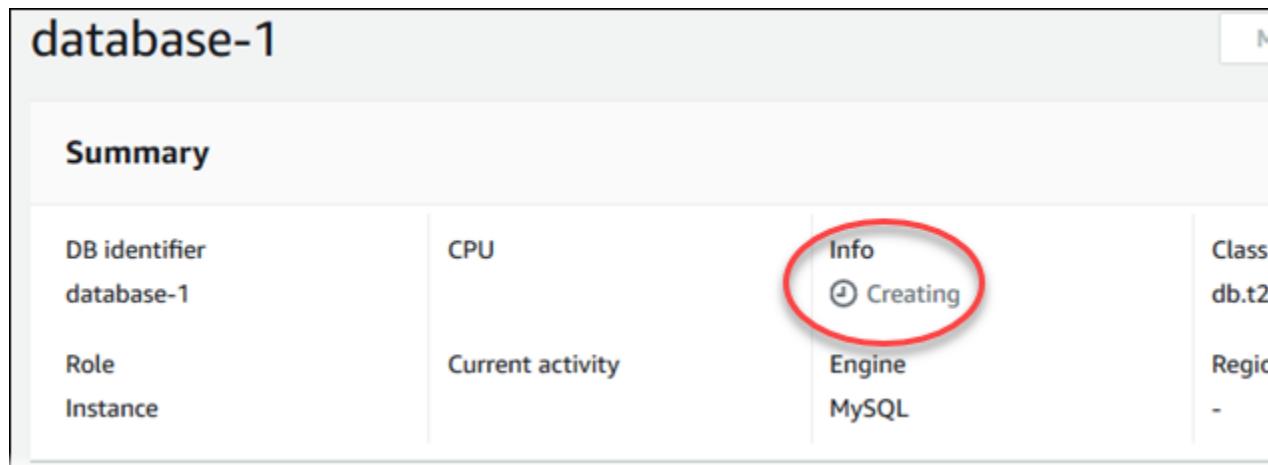
To connect to the DB instance as the master user, use the user name and password that appear.

Important

You can't view the master user password again. If you don't record it, you might have to change it. If you need to change the master user password after the DB instance is available, you can modify the DB instance to do so. For more information about modifying a DB instance, see [Modifying a DB Instance Running the MySQL Database Engine \(p. 802\)](#).

11. For **Databases**, choose the name of the new MySQL DB instance.

On the RDS console, the details for the new DB instance appear. The DB instance has a status of **creating** until the DB instance is created and ready for use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and storage allocated, it can take several minutes for the new instance to be available.



Original Console

To launch a MySQL DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the top right corner of the AWS Management Console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
If the navigation pane is closed, choose the menu icon at the top left to open it.
4. Choose **Create database** to open the **Select engine** page.

Select engine

Engine options

- Amazon Aurora
Amazon Aurora
- MySQL

- MariaDB

- PostgreSQL

- Oracle

- Microsoft SQL Server


MySQL

MySQL is the most popular open source database in the world. MySQL on RDS offers the rich features of the MySQL community edition with the flexibility to easily scale compute resources or storage capacity for your database.

- Supports database size up to 16 TB.
- Instances offer up to 32 vCPUs and 244 GiB Memory.
- Supports automated backup and point-in-time recovery.
- Supports cross-region read replicas.

Only enable options eligible for RDS Free Usage Tier [info](#)

[Cancel](#) **Next**

5. In the **Select engine** window, choose **MySQL**, and then choose **Next**.
6. The **Choose use case** page asks if you are planning to use the DB instance you are creating for production. If you are, choose **Production - MySQL**. If you choose **Production - MySQL**, the following are also chosen in a later step:
 - **Multi-AZ failover option**
 - **Provisioned IOPS storage option**
 - **Enable deletion protection option**

We recommend these features for any production environment.

7. Choose **Next** to continue. The **Specify DB details** page appears.

On the **Specify DB details** page, specify your DB instance information. For information about each setting, see [Settings for MySQL DB Instances \(p. 758\)](#).

Specify DB details

Instance specifications

Estimate your monthly costs for the DB Instance using the [AWS Simple Monthly Calculator](#)

DB engine

MySQL Community Edition

License model [Info](#)

general-public-license

DB engine version [Info](#)

MySQL 5.6.40



Known Issues/Limitations

Review the [Known Issues/Limitations](#) to learn about potential compatibility issues with specific database versions.

DB instance class [Info](#)

db.r4.xlarge — 4 vCPU, 30.5 GiB RAM

Multi-AZ deployment [Info](#)

Create replica in different zone

Creates a replica in a different Availability Zone (AZ) to provide data redundancy, eliminate I/O freezes, and minimize latency spikes during system backups.

No

Storage type [Info](#)

Provisioned IOPS (SSD)

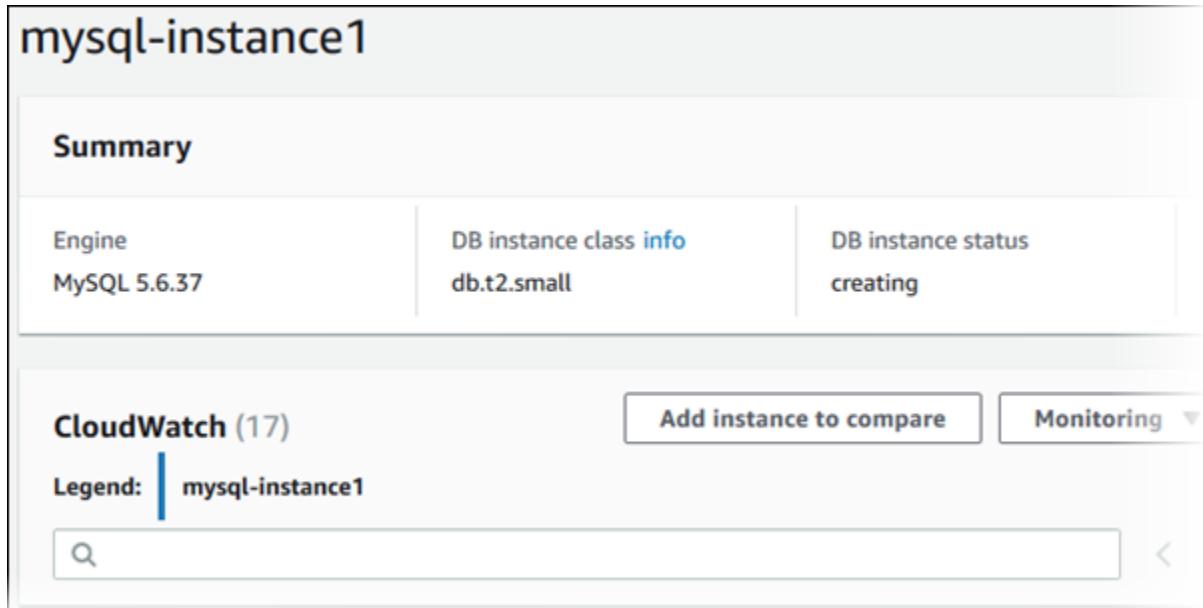
- Choose **Next** to continue. The **Configure advanced settings** page appears.

On the **Configure advanced settings** page, provide additional information that Amazon RDS needs to launch the DB instance. For information about each setting, see [Settings for MySQL DB Instances \(p. 758\)](#).

- Choose **Create database**.

- On the final page, choose **View DB instance details**.

On the RDS console, the details for the new DB instance appear. The DB instance has a status of **creating** until the DB instance is created and ready for use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and storage allocated, it could take several minutes for the new instance to be available.



AWS CLI

To create a MySQL DB instance by using the AWS CLI, call the [create-db-instance](#) command with the parameters below. For information about each setting, see [Settings for MySQL DB Instances \(p. 758\)](#).

- `--db-instance-identifier`
- `--db-instance-class`
- `--db-security-groups`
- `--db-subnet-group`
- `--engine`
- `--master-user-name`
- `--master-user-password`
- `--allocated-storage`
- `--backup-retention-period`

Example

The following example creates a MySQL DB instance named mydbinstance.

For Linux, OS X, or Unix:

```
aws rds create-db-instance \
--db-instance-identifier mydbinstance \
--db-instance-class db.m1.small \
--engine MySQL \
--allocated-storage 20 \
--master-username masterawsuser \
--master-user-password masteruserpassword \
--backup-retention-period 3
```

For Windows:

```
aws rds create-db-instance ^
```

```
--db-instance-identifier mydbinstance ^
--db-instance-class db.m3.medium ^
--engine MySQL ^
--allocated-storage 20 ^
--master-username masterawsuser ^
--master-user-password masteruserpassword ^
--backup-retention-period 3
```

This command should produce output similar to the following:

```
DBINSTANCE mydbinstance db.m3.medium mysql 20 sa creating 3 **** n 5.6.40
SECGROUP default active
PARAMGRP default.mysql5.6 in-sync
```

RDS API

To create a MySQL DB instance by using the Amazon RDS API, call the [CreateDBInstance](#) operation with the parameters below. For information about each setting, see [Settings for MySQL DB Instances \(p. 758\)](#).

- AllocatedStorage
- BackupRetentionPeriod
- DBInstanceClass
- DBInstanceIdentifier
- DBSecurityGroups
- DBSubnetGroup
- Engine
- MasterUsername
- MasterUserPassword

Example

The following example creates a MySQL DB instance named mydbinstance.

```
https://rds.us-west-2.amazonaws.com/
?Action=CreateDBInstance
&AllocatedStorage=20
&BackupRetentionPeriod=3
&DBInstanceClass=db.m3.medium
&DBInstanceIdentifier=mydbinstance
&DBName=mydatabase
&DBSecurityGroups.member.1=mymysecuritygroup
&DBSubnetGroup=mymydbsubnetgroup
&Engine=mysql
&MasterUserPassword=masteruserpassword
&MasterUsername=masterawsuser
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140213/us-west-2/rds/aws4_request
&X-Amz-Date=20140213T162136Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=8052a76dfb18469393c5f0182cdab0ebc224a9c7c5c949155376c1c250fc7ec3
```

Settings for MySQL DB Instances

For details about settings that you use when you create a MySQL DB instance, see the following table.

Setting	Setting Description
Allocated storage	The amount of storage to allocate for your DB instance (in gigabytes). In some cases, allocating a higher amount of storage for your DB instance than the size of your database can improve I/O performance. For more information, see Amazon RDS DB Instance Storage (p. 135) .
Auto minor version upgrade	Enable auto minor version upgrade to enable your DB instance to receive preferred minor DB engine version upgrades automatically when they become available. Amazon RDS performs automatic minor version upgrades in the maintenance window.
Availability zone	The Availability Zone for your DB instance. Use the default value of No Preference unless you want to specify an Availability Zone. For more information, see Regions and Availability Zones (p. 132) .
Backup retention period	The number of days that you want automatic backups of your DB instance to be kept. For any nontrivial DB instance, set this value to 1 or greater. For more information, see Working With Backups (p. 254) .
Backup window	The time period during which Amazon RDS automatically takes a backup of your DB instance. Unless you have a specific time that you want to have your database backed up, use the default of No Preference . For more information, see Working With Backups (p. 254) .
Copy tags to snapshots	This option copies any DB instance tags to a DB snapshot when you create a snapshot. For more information, see Tagging Amazon RDS Resources (p. 177) .
Database authentication	The database authentication option you want to use. Choose Password authentication to authenticate database users with database passwords only. Choose Password and IAM DB authentication to authenticate database users with database passwords and user credentials through IAM users and roles. For more information, see IAM Database Authentication for MySQL and PostgreSQL (p. 461) .
Database name	The name for the database on your DB instance. The name must contain 1 to 64 alphanumeric characters. If you don't provide a name, Amazon RDS doesn't create a database on the DB instance that you are creating.

Setting	Setting Description
	To create additional databases on your DB instance, connect to your DB instance and use the SQL command <code>CREATE DATABASE</code> . For more information, see Connecting to a DB Instance Running the MySQL Database Engine (p. 764) .
Database port	The port that you want to access the DB instance through. MySQL installations default to port 3306. If you use a DB security group with your DB instance, this port value must be the same one that you provided when creating the DB security group. The firewalls at some companies block connections to the default MySQL port. If your company firewall blocks the default port, enter another port for your DB instance.
DB engine version	The version of MySQL that you want to use.
DB instance class	The configuration for your DB instance. For example, a db.m1.small instance class has 1.7 GiB memory, 1 ECU (1 virtual core with 1 ECU), 64-bit platform, and moderate I/O capacity. If possible, choose an instance class large enough that a typical query working set can be held in memory. When working sets are held in memory, the system can avoid writing to disk, and this improves performance. For more information, see Choosing the DB Instance Class (p. 107) .
DB instance identifier	The name for your DB instance. Your DB instance identifier can contain up to 63 alphanumeric characters, and must be unique for your account in the AWS Region you chose. You can add some intelligence to the name, such as including the AWS Region you chose, for example mysql-instance1 .
DB parameter group	A parameter group for your DB instance. You can choose the default parameter group or you can create a custom parameter group. For more information, see Working with DB Parameter Groups (p. 211) .
Deletion protection	Enable deletion protection to prevent your DB instance from being deleted. If you create a production DB instance with the AWS Management Console, deletion protection is enabled by default. For more information, see Deleting a DB Instance (p. 174) .
Encryption	Enable Encryption to enable encryption at rest for this DB instance. For more information, see Encrypting Amazon RDS Resources (p. 429) .

Setting	Setting Description
Enhanced Monitoring	<p>Enable enhanced monitoring to enable gathering metrics in real time for the operating system that your DB instance runs on.</p> <p>For more information, see Enhanced Monitoring (p. 307).</p>
License model	MySQL has only one license model, general-public-license , the general license agreement for MySQL.
Log exports	The types of MySQL database log files to generate. For more information, see MySQL Database Log Files (p. 405) .
Maintenance window	<p>The 30-minute window in which pending modifications to your DB instance are applied. If the time period doesn't matter, choose No Preference.</p> <p>For more information, see The Amazon RDS Maintenance Window (p. 156).</p>
Master password	The password for your master user account. The password must contain from 8 to 16 printable ASCII characters (excluding /", a space, and @).
Master username	<p>The name that you use as the master user name to log on to your DB instance.</p> <p>For more information, and a list of the default privileges for the master user, see MySQL Security on Amazon RDS (p. 744).</p>
Multi-AZ deployment	<p>Create a standby instance to create a passive secondary replica of your DB instance in another Availability Zone for failover support. We recommend Multi-AZ for production workloads to maintain high availability.</p> <p>For development and testing, you can choose Do not create a standby instance.</p> <p>For more information, see High Availability (Multi-AZ) for Amazon RDS (p. 144).</p>
Option group	<p>An option group for your DB instance. You can choose the default option group or you can create a custom option group.</p> <p>For more information, see Working with Option Groups (p. 196).</p>

Setting	Setting Description
Performance Insights	<p>Enable Performance Insights to monitor your DB instance load so that you can analyze and troubleshoot your database performance.</p> <p>Choose a retention period to determine how much rolling data history to keep. The default of seven days is in the free tier. Long-term retention (two years) is priced per vCPU per month.</p> <p>Choose a master key to use to protect the key used to encrypt this database volume. Choose from the master keys in your account, or enter the key from a different account.</p> <p>For more information, see Using Amazon RDS Performance Insights (p. 319).</p>
Public accessibility	<p>Yes to give the DB instance a public IP address, meaning that it's accessible outside the VPC. To be publicly accessible, the DB instance also has to be in a public subnet in the VPC.</p> <p>No to make the DB instance accessible only from inside the VPC.</p> <p>For more information, see Hiding a DB Instance in a VPC from the Internet (p. 515).</p>
Storage autoscaling	<p>Enable storage autoscaling to enable Amazon RDS to automatically increase storage when needed to avoid having your DB instance run out of storage space.</p> <p>Use Maximum storage threshold to set the upper limit for Amazon RDS to automatically increase storage for your DB instance. The default is 1,000 GiB.</p> <p>For more information, see Managing Capacity Automatically with Amazon RDS Storage Autoscaling (p. 232).</p>
Storage type	<p>The storage type for your DB instance.</p> <p>For more information, see Amazon RDS Storage Types (p. 135).</p>
Subnet group	<p>This setting depends on the platform that you are on. If you are a new customer to AWS, keep the subnet group set to default, which is the default DB subnet group that was created for your account.</p> <p>If you are creating a DB instance on the earlier E2-Classic platform, you might want your DB instance in a specific VPC. In this case, choose the DB subnet group that you created for that VPC.</p>

Setting	Setting Description
Virtual Private Cloud (VPC)	This setting depends on the platform that you are on. If you are a new customer to AWS, choose the default VPC shown. If you are creating a DB instance on the earlier E2-Classic platform that doesn't use a VPC, choose Not in VPC . For more information, see Amazon Virtual Private Cloud VPCs and Amazon RDS (p. 505) .
VPC security groups	If you are a new customer to AWS, Create new to create a new VPC security group. Otherwise, Choose existing , then choose from security groups that you previously created. When you choose Create new in the RDS console, a new security group is created. This new security group has an inbound rule that allows access to the DB instance from the IP address detected in your browser. For more information, see Working with DB Security Groups (EC2-Classic Platform) (p. 491) .

Connecting to a DB Instance Running the MySQL Database Engine

Before you can connect to a DB instance running the MySQL database engine, you must create a DB instance. For information, see [Creating a DB Instance Running the MySQL Database Engine \(p. 751\)](#). After Amazon RDS provisions your DB instance, you can use any standard MySQL client application or utility to connect to the instance. In the connection string, you specify the DNS address from the DB instance endpoint as the host parameter, and specify the port number from the DB instance endpoint as the port parameter.

To authenticate to your RDS DB instance, you can use one of the authentication methods for MySQL and IAM database authentication.

- To learn how to authenticate to MySQL using one of the authentication methods for MySQL, see [Authentication Method](#) in the MySQL documentation.
- To learn how to authenticate to MySQL using IAM database authentication, see [IAM Database Authentication for MySQL and PostgreSQL \(p. 461\)](#).

You can use the AWS Management Console, the AWS CLI `describe-db-instances` command, or the Amazon RDS API `DescribeDBInstances` operation to list the details of an Amazon RDS DB instance, including its endpoint.

To find the endpoint for a MySQL DB instance in the AWS Management Console:

1. Open the RDS console and then choose **Databases** to display a list of your DB instances.
2. Choose the MySQL DB instance name to display its details.
3. On the **Connectivity & security** tab, copy the endpoint. Also, note the port number. You need both the endpoint and the port number to connect to the DB instance.

The screenshot shows the 'Summary' page for a database instance named 'database-1'. The 'Connectivity & security' tab is active. The 'Endpoint' field is highlighted with a red oval, containing the value 'database-1.' followed by a blurred IP address and '.us-east-2.rds.amazonaws.com'. The 'Port' is listed as 3306.

If an endpoint value is `mysql-instance1.123456789012.us-east-1.rds.amazonaws.com` and the port value is 3306, then you would specify the following values in a MySQL connection string:

- For host or host name, specify `mysql-instance1.123456789012.us-east-1.rds.amazonaws.com`
- For port, specify 3306

You can connect to an Amazon RDS MySQL DB instance by using tools like the MySQL command line utility. For more information on using the MySQL client, go to [mysql - The MySQL Command Line Tool](#) in the MySQL documentation. One GUI-based application you can use to connect is MySQL Workbench. For more information, go to the [Download MySQL Workbench](#) page. For information about installing MySQL (including the MySQL client), see [Installing and Upgrading MySQL](#).

Two common causes of connection failures to a new DB instance are:

- The DB instance was created using a security group that doesn't authorize connections from the device or Amazon EC2 instance where the MySQL application or utility is running. If the DB instance was created in a VPC, it must have a VPC security group that authorizes the connections. For more information, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).

You can add or edit an inbound rule in the security group. For **Source**, choose **My IP**. This allows access to the DB instance from the IP address detected in your browser.

If the DB instance was created outside of a VPC, it must have a DB security group that authorizes the connections.

- The DB instance was created using the default port of 3306, and your company has firewall rules blocking connections to that port from devices in your company network. To fix this failure, recreate the instance with a different port.

You can use SSL encryption on connections to an Amazon RDS MySQL DB instance. For information, see [Using SSL with a MySQL DB Instance \(p. 746\)](#). If you are using IAM database authentication, you must use an SSL connection. For information, see [IAM Database Authentication for MySQL and PostgreSQL \(p. 461\)](#).

For information on connecting to a MariaDB DB instance, see [Connecting to a DB Instance Running the MariaDB Database Engine \(p. 555\)](#).

Connecting from the MySQL Client

To connect to a DB instance using the MySQL client, type the following command at a command prompt to connect to a DB instance using the MySQL client. For the -h parameter, substitute the DNS name (endpoint) for your DB instance. For the -P parameter, substitute the port for your DB instance. Enter the master user password when prompted.

```
mysql -h mysql-instance1.123456789012.us-east-1.rds.amazonaws.com -P 3306 -u mymasteruser -p
```

After you enter the password for the user, you will see output similar to the following.

```
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 350
Server version: 5.6.40-log MySQL Community Server (GPL)

Type 'help;' or '\h' for help. Type '\c' to clear the buffer.

mysql>
```

Connecting with SSL

Amazon RDS creates an SSL certificate for your DB instance when the instance is created. If you enable SSL certificate verification, then the SSL certificate includes the DB instance endpoint as the Common Name (CN) for the SSL certificate to guard against spoofing attacks. To connect to your DB instance using SSL, you can use native password authentication or IAM database authentication. To connect to your DB instance using IAM database authentication, see [IAM Database Authentication for MySQL and PostgreSQL \(p. 461\)](#). To connect to your DB instance using native password authentication, you can follow these steps:

To connect to a DB instance with SSL using the MySQL client

- Download a root certificate that works for all AWS Regions.

For information about downloading certificates, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).

2. Enter the following command at a command prompt to connect to a DB instance with SSL using the MySQL client. For the -h parameter, substitute the DNS name for your DB instance. For the --ssl-ca parameter, substitute the SSL certificate file name as appropriate.

```
mysql -h mysql-instance1.123456789012.us-east-1.rds.amazonaws.com --ssl-ca=rds-ca-2015-root.pem -p
```

3. You can require that the SSL connection verifies the DB instance endpoint against the endpoint in the SSL certificate.

For MySQL 5.7 and later:

```
mysql -h mysql-instance1.123456789012.us-east-1.rds.amazonaws.com --ssl-ca=rds-ca-2015-root.pem --ssl-mode=VERIFY_IDENTITY -p
```

For MySQL 5.6 and earlier:

```
mysql -h mysql-instance1.123456789012.us-east-1.rds.amazonaws.com --ssl-ca=rds-ca-2015-root.pem --ssl-verify-server-cert -p
```

4. Enter the master user password when prompted.

You will see output similar to the following.

```
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 350
Server version: 5.6.40-log MySQL Community Server (GPL)

Type 'help;' or '\h' for help. Type '\c' to clear the buffer.

mysql>
```

Connecting from MySQL Workbench

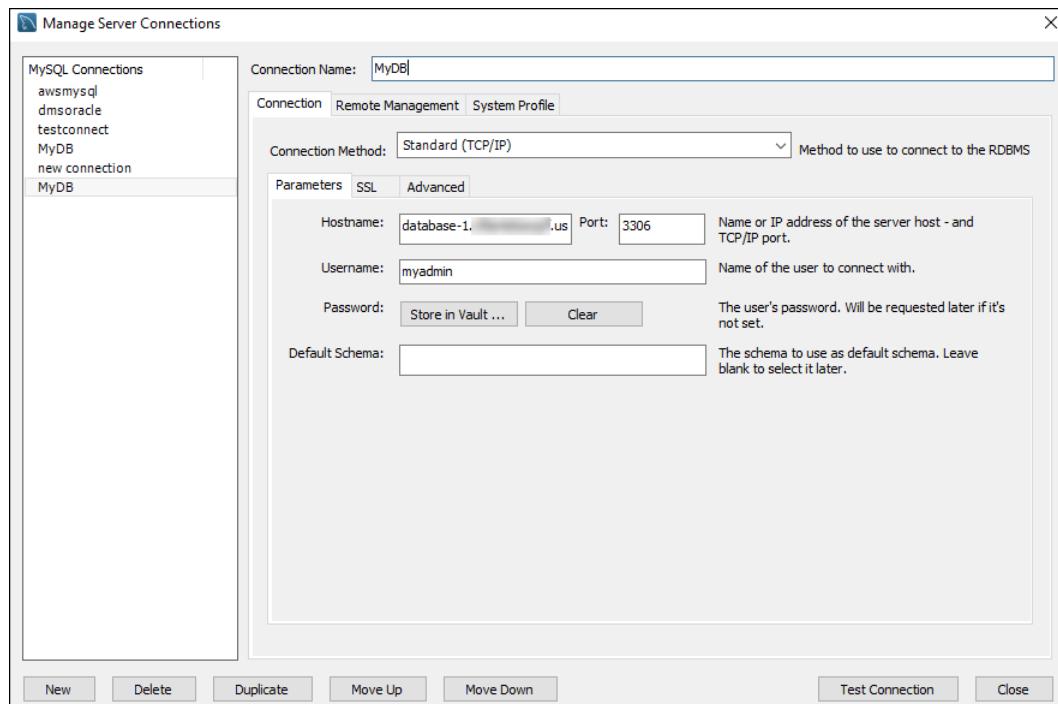
To connect from MySQL Workbench

1. Download and install MySQL Workbench at [Download MySQL Workbench](#).
2. Open MySQL Workbench.



3. From **Database**, choose **Manage Connections**.
4. In the **Manage Server Connections** window, choose **New**.
5. In the **Connect to Database** window, enter the following information:
 - **Stored Connection** – Enter a name for the connection, such as **MyDB**.
 - **Hostname** – Enter the DB instance endpoint.
 - **Port** – Enter the port used by the DB instance.
 - **Username** – Enter the username of a valid database user, such as the master user.
 - **Password** – Optionally, choose **Store in Vault** and then enter and save the password for the user.

The window looks similar to the following:



6. Optionally, choose **Test Connection** to confirm that the connection to the DB instance is successful.
7. Choose **Close**.
8. From **Database**, choose **Connect to Database**.
9. From **Stored Connection**, choose your connection.
10. Choose **OK**.

For information about using MySQL Workbench, see the [MySQL Workbench documentation](#).

Maximum MySQL connections

The maximum number of connections allowed to an Amazon RDS MySQL DB instance is based on the amount of memory available for the DB instance class of the DB instance. A DB instance class with more memory available will result in a larger amount of connections available. For more information on DB instance classes, see [Choosing the DB Instance Class \(p. 107\)](#).

The connection limit for a DB instance is set by default to the maximum for the DB instance class for the DB instance. You can limit the number of concurrent connections to any value up to the maximum number of connections allowed using the `max_connections` parameter in the parameter group for the DB instance. For more information, see [Working with DB Parameter Groups \(p. 211\)](#).

You can retrieve the maximum number of connections allowed for an Amazon RDS MySQL DB instance by executing the following query on your DB instance:

```
SELECT @@max_connections;
```

You can retrieve the number of active connections to an Amazon RDS MySQL DB instance by executing the following query on your DB instance:

```
SHOW STATUS WHERE `variable_name` = 'Threads_connected';
```

Managing Connections with Amazon RDS Proxy (Preview)

This is preview documentation for Amazon RDS Proxy. It is subject to change.

Amazon RDS Proxy allows applications to pool and share database connections to improve scalability. It makes applications more resilient to database failures by automatically connecting to a standby DB instance while preserving application connections. RDS Proxy also allows you to enforce AWS IAM (Identity and Access Management) authentication to databases, and securely store credentials in Secrets Manager. RDS Proxy is fully compatible with MySQL and can be enabled for most applications with no code change.

You can handle unpredictable surges in database traffic, which otherwise could cause issues due to oversubscribing connections or creating new connections at a fast rate. RDS Proxy establishes a database connection pool and reuses connections in this pool without the memory and CPU overhead of opening a new database connection each time. To protect the database against oversubscription, you can control the number of database connections that are created. RDS Proxy sequences or throttles application connections that can't be served immediately from the pool of connections. Although latencies might increase, your application can continue to scale without abruptly failing or overwhelming the database. If connection requests exceed the limits you specify, RDS Proxy rejects application connections (sheds load) while maintaining predictable performance for the load that can be served with available capacity.

You can reduce the overhead to process credentials and establish a secure connection for each new connection. RDS Proxy can handle some of that work on behalf of the database.

Important

Currently, Amazon RDS Proxy is in preview. Don't use RDS Proxy for production workloads. Also, we strongly recommend not putting sensitive data in databases that you use with RDS Proxy preview. You might encounter issues during the preview that cause incorrect results, corrupted data, or both. Over the duration of the preview, we might introduce breaking changes without advance notice. These might include upgrades and changes to our APIs.

Topics

- [RDS Proxy Concepts and Terminology \(p. 770\)](#)
- [Planning for and Setting Up RDS Proxy \(p. 773\)](#)
- [Connecting to a Database through RDS Proxy \(p. 785\)](#)
- [Managing RDS Proxy \(p. 786\)](#)
- [Monitoring RDS Proxy \(p. 789\)](#)
- [Limitations for RDS Proxy \(p. 791\)](#)
- [Command-Line Examples for RDS Proxy \(p. 792\)](#)
- [Troubleshooting for RDS Proxy \(p. 793\)](#)

RDS Proxy Concepts and Terminology

You can simplify connection management for your Amazon RDS DB instances and Aurora DB clusters by using RDS Proxy.

RDS Proxy handles the network traffic between the client application and the database. It does so in an active way by understanding the database protocol and adjusting its behavior based on the SQL operations from your application and the result sets from the database. RDS Proxy reduces the memory and CPU overhead for connection management on your database. The database needs less memory and

CPU resources when applications open many simultaneous connections. It also doesn't require logic in your applications to close and reopen connections that stay idle for a long time. Similarly, it requires less application logic to reestablish connections in case of a database problem.

The infrastructure for RDS Proxy is independent of the database servers for the RDS DB instances and Aurora DB clusters. This separation helps lower overhead on your database servers, so that they can devote their resources to serving database workloads.

Topics

- [Overview of RDS Proxy Terminology \(p. 771\)](#)
- [Connection Pooling \(p. 772\)](#)
- [RDS Proxy Security \(p. 772\)](#)
- [Failover \(p. 772\)](#)
- [Transactions \(p. 773\)](#)

Overview of RDS Proxy Terminology

RDS Proxy handles the infrastructure to perform connection pooling and the other features described in this documentation. You see the proxies represented in the AWS Management Console for RDS on the [Proxies](#) page.

Each proxy handles connections to a single RDS DB instance or Aurora DB cluster. For RDS multi-AZ DB instances and Aurora provisioned clusters, the proxy determines the current writer instance.

The connections that a proxy keeps open and available for your database application to use form the *connection pool*.

By default, RDS Proxy can reuse a connection after each transaction in your session. This transaction-level reuse is called *multiplexing*. When the RDS Proxy can't be sure that it's safe to run different transactions from a session using different connections, it keeps the session on the same connection until the session ends. This fallback behavior is called *pinning*.

A proxy has an endpoint. You connect to this endpoint when you work with an RDS DB instance or Aurora DB cluster, instead of the read-write endpoint that connects directly to the instance or cluster. The special-purpose endpoints for an Aurora cluster remain available for you to use. For example, you can still connect to the cluster endpoint for read-write connections without connection pooling. You can still connect to the reader endpoint for load-balanced read-only connections. You can still connect to the instance endpoints for diagnosis and troubleshooting of specific DB instances within an Aurora cluster. If you are using other AWS services such as AWS Lambda to connect to RDS databases, you change their connection settings to use the proxy endpoint. For example, you specify the proxy endpoint to allow Lambda functions to access your database while taking advantage of RDS Proxy functionality.

Each proxy contains a target group. This *target group* embodies the RDS DB instance or Aurora DB cluster that the proxy can connect to. For an Aurora cluster, by default the target group is associated with all the DB instances in that cluster. That way, the proxy can connect to whichever Aurora DB instance is promoted to be the writer instance in the cluster. The RDS DB instance associated with a proxy, or the Aurora DB cluster and its instances, are referred to as the *targets* of that proxy. For convenience, when you create a proxy through the AWS Management Console, RDS Proxy also creates the corresponding target group and registers the associated targets automatically.

An *engine family* is a related set of database engines that use the same DB protocol. You choose the engine family for each proxy that you create. During the preview, RDS Proxy supports one engine family, consisting of RDS MySQL 5.6 and 5.7, and Aurora MySQL. Within the MySQL engine family, currently RDS Proxy supports Aurora provisioned clusters, Aurora parallel query clusters, and Aurora Global Databases. For a global database, you can create a proxy for the primary AWS Region, but not for the read-only secondary AWS Regions. RDS Proxy currently doesn't support Aurora Serverless clusters and Aurora multi-master clusters.

Connection Pooling

Each proxy performs connection pooling for the writer instance of its associated RDS or Aurora database. *Connection pooling* is an optimization that reduces the overhead associated with opening and closing connections, and keeping many connections open simultaneously. This overhead includes memory needed to handle each new connection. It also involves CPU overhead to close each connection and open a new one, such as TLS handshaking, authentication, negotiating capabilities, and so on. Connection pooling simplifies your application logic. You don't need to write application code to minimize the number of simultaneous open connections.

Each proxy also performs connection multiplexing, also known as connection reuse. With *multiplexing*, RDS Proxy performs all the operations for a transaction using one underlying database connection, then can use a different connection for the next transaction. You can open many simultaneous connections to the proxy, and the proxy keeps a smaller number of connections open to the DB instance or cluster. Doing so further minimizes the memory overhead for connections on the database server. This technique also reduces the chance of "too many connections" errors.

RDS Proxy Security

RDS Proxy uses the existing RDS security mechanisms such as Transport Layer Security (TLS) and AWS Identity and Access Management (IAM). For general information about those security features, see [Security in Amazon RDS \(p. 427\)](#). If you aren't familiar with how RDS and Aurora work with authentication, authorization, and other areas of security, consult those resources first.

RDS Proxy can act as an additional layer of security between client applications and the underlying database. For example, you can connect to the proxy using TLS 1.2, even if the underlying DB instance only supports TLS 1.0 or 1.1. You can connect to the proxy using an IAM role, even if the proxy connects to the database using the native user/password authentication method. By using this technique, you can enforce strong authentication requirements for database applications without a costly migration effort for the DB instances themselves.

You store the database credentials used by RDS Proxy in AWS Secrets Manager. Each database user for the RDS DB instance or Aurora DB cluster accessed by a proxy must have a corresponding secret in Secrets Manager. You can also set up IAM authentication for users of RDS Proxy. By doing so, you can enforce IAM authentication for database access even if the databases still use native password authentication. These security features are a preferable alternative to embedding database credentials in your application code.

Failover

Failover is a high-availability feature that replaces a database instance with another one when the original instance becomes unavailable. A failover might happen because of a problem with a database instance. It could also be part of normal maintenance procedures, such as while performing a database upgrade. Failover applies to RDS DB instances in a multi-AZ configuration, and Aurora DB clusters with one or more reader instances in addition to the writer instance.

Without RDS Proxy, a failover involves a brief outage. During the outage, you can't perform write operations on that database. Any existing database connections are disrupted and your application must reopen them. The database becomes available for new connections and write operations when a read-only DB instance is promoted to take the place of the one that's unavailable.

Connecting through a proxy makes your application more resilient to database failovers. When the original DB instance becomes unavailable, RDS Proxy connects to the standby database without dropping idle application connections. Doing so helps to speed up and simplify the failover process. The result is faster failover that's less disruptive to your application than a typical reboot or database problem.

During DB failovers, RDS Proxy continues to accept connections at the same IP address and automatically directs connections to the new master. Clients connecting through RDS Proxy are not susceptible to

DNS propagation delays on failover, local DNS caching, connection timeouts, uncertainty about which DB instance is the current writer, or waiting for a query response from a former writer that became unavailable without closing connections.

For applications that maintain their own connection pool, going through RDS Proxy means that most connections stay alive during failovers or other disruptions. Only connections that are in the middle of a transaction or SQL statement are cancelled. RDS Proxy immediately accepts new connections. When the database writer is unavailable, RDS Proxy queues up incoming requests.

For applications that don't maintain their own connection pools, RDS Proxy offers faster connection rates and more open connections. It offloads the expensive overhead of frequent reconnects from the database. It does so by reusing database connections maintained in the RDS Proxy connection pool. This is particularly important for TLS connections, where setup costs are significant.

Transactions

All the statements within a single transaction always use the same underlying database connection. The connection becomes available for use by a different session when the transaction ends. Using the transaction as the unit of granularity has the following consequences:

- Connection reuse can happen after each individual statement when the `autocommit` setting is enabled.
- Conversely, when the `autocommit` setting is disabled, the first statement you issue in a session begins a new transaction. Thus, if you enter a sequence of `SELECT`, `INSERT`, `UPDATE`, and other data manipulation language (DML) statements, connection reuse doesn't happen until you issue a `COMMIT`, `ROLLBACK`, or otherwise end the transaction.
- Entering a data definition language (DDL) statement causes the transaction to end after that statement completes.

RDS Proxy detects when a transaction ends through the network protocol used by the database client application. Transaction detection doesn't rely on keywords such as `COMMIT` or `ROLLBACK` appearing in the text of the SQL statement.

If RDS Proxy detects a database request that makes it impractical to move your session to a different connection, it turns off multiplexing for that connection the remainder of your session. The same rule applies if RDS Proxy can't be certain that multiplexing is practical for the session. This operation is called pinning. For ways to detect and minimize pinning, see [Pinning \(p. 787\)](#).

Planning for and Setting Up RDS Proxy

This is preview documentation for Amazon RDS Proxy. It is subject to change.

In the following sections, you can find how to set up RDS Proxy. You can also find how to set the related security options that control who can access each proxy and how each proxy connects to DB instances.

Topics

- [Identifying DB Instances, Clusters, and Applications to Use with RDS Proxy \(p. 774\)](#)
- [Setting Up Network Prerequisites \(p. 774\)](#)
- [Setting Up Database Credentials in AWS Secrets Manager \(p. 774\)](#)
- [Setting Up AWS Identity and Access Management \(IAM\) Policies \(p. 775\)](#)
- [Creating an RDS Proxy \(p. 777\)](#)
- [Viewing RDS Proxy \(p. 779\)](#)

- [Deleting an RDS Proxy \(p. 781\)](#)
- [Modifying an RDS Proxy \(p. 781\)](#)
- [Adding a New Database User \(p. 785\)](#)
- [Changing the Password for a Database User \(p. 785\)](#)

Identifying DB Instances, Clusters, and Applications to Use with RDS Proxy

You can determine which of your DB instances, clusters, and applications might benefit the most from using RDS Proxy. To do so, consider these factors:

- Any DB instance or cluster that ever encounters "too many connections" errors is a good candidate for associating with a proxy. The proxy enables applications to open many client connections, while the proxy manages a smaller number of long-lived connections to the DB instance or cluster.
- For DB instances or clusters that use smaller AWS instance classes, such as T2 or T3, using a proxy can help avoid out-of-memory conditions and reduce the CPU overhead for establishing connections. These conditions can occur when dealing with large numbers of connections.
- You can monitor certain Amazon CloudWatch metrics to determine whether a DB instance or cluster is approaching its limits for number of connections or the memory associated with connection management. You can also monitor certain CloudWatch metrics to determine whether a DB instance or cluster is handling many short-lived connections. Opening and closing such connections can impose performance overhead on your database. For information about the metrics to monitor, see [Monitoring RDS Proxy \(p. 789\)](#).
- AWS Lambda functions can also be good candidates for using with a proxy. These functions make frequent short database connections that benefit from connection pooling offered by RDS Proxy. You can take advantage of any IAM authentication you already have for Lambda functions, instead of managing database credentials in your Lambda application code.
- Applications that use languages and frameworks such as PHP and Ruby on Rails are typically good candidates for using with a proxy. Such applications typically open and close large numbers of database connections, and don't have built-in connection pooling mechanisms.
- Applications that keep a large number of connections open for long periods are typically good candidates for using with a proxy. Applications in industries such as software as a service (SaaS) or ecommerce often minimize the latency for database requests by leaving connections open.
- You might not have adopted IAM authentication and Secrets Manager due to the complexity of setting up such authentication for all DB instances and clusters. If so, you can leave the existing authentication methods in place and delegate the authentication to a proxy. The proxy can enforce the authentication policies for client connections for particular applications. You can take advantage of any IAM authentication you already have for Lambda functions, instead of managing database credentials in your Lambda application code.

Setting Up Network Prerequisites

Using RDS Proxy requires you to have a set of networking resources in place, such as a virtual private cloud (VPC), two or more subnets, EC2 instance within the same VPC, and Internet gateway. If you've successfully connected to any RDS DB instances or Aurora DB clusters, you already have the required network resources.

Setting Up Database Credentials in AWS Secrets Manager

For each proxy that you create, you first use the Secrets Manager service to store sets of user name and password credentials. You create a separate Secrets Manager secret for each database user account that the proxy connects to on the RDS DB instance or Aurora DB cluster.

In Secrets Manager, you create these secrets with the `username` and `password` fields. Doing so allows the proxy to connect to the corresponding database users on whichever RDS DB instances or Aurora

DB clusters that you associate with the proxy. To do this, you can use the setting **Credentials for other database**, **Credentials for RDS database**, or **Other type of secrets**.

The proxy ignores other fields such as host and port if they're present in the secret. Those details are automatically supplied by the proxy. If you choose **Credentials for other database**, that option prompts you for the user name and password, but not the other connection details, which you specify in the settings for the proxy itself. If you choose **Other type of secrets**, you create the secret with keys named `username` and `password`.

Don't choose any rotation period for the secrets.

Because the secrets aren't tied to a specific database server, you can reuse a secret across multiple proxies if you use the same credentials across multiple database servers. For example, you might use the same credentials across a group of development and test servers.

If a password associated with a secret is incorrect, you can update the associated secret in Secrets Manager. For example, you might update the secret if the password for the database account changes. Until you update the secret, you can't connect through the proxy to that database account. You can still connect to other accounts where the passwords match the credentials stored in the secrets.

When you create a proxy through the AWS CLI or RDS API, you specify the Amazon Resource Names (ARNs) of the corresponding secrets for all the DB user accounts that the proxy can access. In the AWS Management Console, you choose the secrets by their descriptive names.

For instructions about creating secrets in Secrets Manager, see the [Creating a Secret](#) page in the Secrets Manager documentation. Use one of the following techniques. For instructions about creating secrets in Secrets Manager, see [Creating a Secret](#) in the *AWS Secrets Manager User Guide*. Use one of the following techniques:

- Use [Secrets Manager](#) in the console.
- To use the CLI to create a Secrets Manager secret for use with RDS Proxy, use a command such as the following.

```
aws secretsmanager create-secret \
--name "secret_name"
--description "secret_description"
--region region_name
--secret-string "{\"username\":\"db_user\",\"password\":\"db_user_password\"}"
```

For example, the following commands create Secrets Manager secrets for two database users, one named `admin` and the other named `app-user`

```
aws secretsmanager create-secret \
--name $ADMIN_SECRET --description 'db admin user' \
--secret-string '{"username":"admin","password":"choose_your_own_password"}'

aws secretsmanager create-secret \
--name $PROXY_SECRET --description 'application user' \
--secret-string '{"username":"app-user","password":"choose_your_own_password"}'
```

Setting Up AWS Identity and Access Management (IAM) Policies

After you create the secrets in Secrets Manager, you create an IAM policy that can access those secrets.

- Follow the [Create Role](#) process. Include the **Add Role to Database** step.

- For the new role, perform the **Add inline policy** step. Paste the following JSON, substituting the ARN for the secret that you created.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "VisualEditor0",
            "Effect": "Allow",
            "Action": [
                "secretsmanager:GetRandomPassword",
                "secretsmanager>CreateSecret",
                "secretsmanager>ListSecrets"
            ],
            "Resource": "*"
        },
        {
            "Sid": "VisualEditor1",
            "Effect": "Allow",
            "Action": "secretsmanager:*",
            "Resource": [
                "your_secret_ARN"
            ]
        }
    ]
}
```

- Optionally, edit the trust policy for this IAM policy. Paste the following JSON.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Principal": {
                "Service": "rds.amazonaws.com"
            },
            "Action": "sts:AssumeRole"
        },
        {
            "Sid": "",
            "Effect": "Allow",
            "Principal": {
                "Service": "rds.amazonaws.com"
            },
            "Action": "sts:AssumeRole"
        }
    ]
}
```

The following commands perform the same operation through the AWS CLI.

```
PREFIX=choose_an_identifier

aws iam create-role --role-name choose_role_name \
--assume-role-policy-document "{\"Version\":\"2012-10-17\",\"Statement\":[{\"Effect\":\"Allow\",\"Principal\":{\"Service\":[\\"rds.amazonaws.com\\\"]},\"Action\":\"sts:AssumeRole\"}]}"
aws iam put-role-policy --role-name same_role_name_as_previous \
```

```
--policy-name $PREFIX-secret-reader-policy --policy-document
'{"Version":"2012-10-17","Statement":[{"Sid":"getsecretvalue","Effect":"Allow","Action":
["secretsmanager:GetSecretValue","kms:Decrypt"],"Resource": "*"}]}'
```

```
aws kms create-key --description "$PREFIX-test-key" --policy "{\"Id\":\"$PREFIX-kms-policy\",
\"Version\":\"2012-10-17\", \"Statement\":[{\\"Sid\\":\\"Enable IAM User Permissions\\",
\"Effect\\\":\\\"Allow\\\", \\"Principal\\\":{\\\"AWS\\\":\\\"arn:aws:iam::account_id:root\\\"}, \\"Action\\\":
\\\"kms:*\\\", \\"Resource\\\":\\\"*\\\"}, {\\"Sid\\\":\\\"Allow access for Key Administrators\\\", \\"Effect
\\\":\\\"Allow\\\", \\"Principal\\\":{\\\"AWS\\\":[$USER_ARN], \\\"arn:aws:iam::account_id:role/
Admin\\\"}}, \\"Action\\\": [\\\"kms>Create\\\", \\"kms:Describe\\\", \\"kms:Enable*\\\", \\"kms>List*
\\\", \\"kms:Put*\\\", \\"kms:Update*\\\", \\"kms:Revoke*\\\", \\"kms:Disable*\\\", \\"kms:Get*\\\",
\\\"kms>Delete*\\\", \\"kms:TagResource\\\", \\"kms:UntagResource\\\", \\"kms:ScheduleKeyDeletion
\\\", \\"kms:CancelKeyDeletion\\\"], \\"Resource\\\":\\\"*\\\"}, {\\"Sid\\\":\\\"Allow use of the key\\\",
\\\"Effect\\\":\\\"Allow\\\", \\"Principal\\\":{\\\"AWS\\\":[$ROLE_ARN]}, \\"Action\\\": [\\\"kms:Decrypt\\\",
\\\"kms:DescribeKey\\\"], \\"Resource\\\":\\\"*\\\"}]}"
```

Creating an RDS Proxy

To manage connections for a specified set of DB instances, you can create a proxy. You can associate a proxy with an RDS MySQL DB instance or a Aurora MySQL DB cluster.

AWS Management Console

To create a proxy

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Proxies**.
3. Choose **Create proxy**.
4. Choose all the settings for your proxy. For the public preview, some choices are restricted or required. Consider the following guidance:

Proxy configuration:

- **Proxy identifier.** Specify a name of your choosing, unique within your AWS account ID and current AWS Region.
- **Engine compatibility.** Currently, **MySQL** is the only choice.
- **Require Transport Layer Security.** Choose this setting if you want the proxy to enforce TLS / SSL for all client connections. When you use an encrypted or unencrypted connection to a proxy, the proxy uses the same encryption setting when it makes a connection to the underlying database.
- **Idle client connection timeout.** Choose a time period that a client connection can be idle before the proxy can close it. A client connection is considered idle when the application doesn't submit a new request within the specified time after the previous request completed. The underlying database connection stays open and is returned to the connection pool. Thus, it's available to be reused for new client connections.

Target group configuration:

- **Database.** Choose one RDS DB instance or Aurora DB cluster to access through this proxy. The list only includes DB instances and clusters with compatible database engines, engine versions, and other settings. If the list is empty, create a new DB instance or cluster that's compatible with RDS Proxy. Then try creating the proxy again.
- **Connection pool maximum connections.** Specify a value between 1 and 100. This setting represents the percentage of the `max_connections` value that RDS Proxy can use for its connections. If you only intend to use one proxy with this DB instance or cluster, you can set it to 100. For details about how RDS Proxy uses this setting, see [Connection Limits and Timeouts \(p. 786\)](#).

- **Session pinning filters.** This is an advanced setting, for troubleshooting performance issues with particular applications. Currently, the only choice is `EXCLUDE_VARIABLE_SETS`. Only choose a filter if your application isn't reusing connections due to certain kinds of SQL statements, and if you can verify that reusing connections with those SQL statements doesn't affect application correctness.
- **Connection borrow timeout.** If you expect the proxy to sometimes use up all available connections, you can specify how long the proxy waits for a connection to become available before returning a timeout error. This setting only applies when the proxy has the maximum number of connections open and all connections are already in use.

Connectivity:

- **Secrets Manager ARNs.** Choose at least one Secrets Manager secret associated with the RDS DB instance or Aurora DB cluster that you intend to access with this proxy.
- **IAM role.** Choose an IAM role that has permission to access the Secrets Manager secrets you chose earlier. You can also choose for the AWS Management Console to create a new IAM role for you and use that.
- **IAM Authentication.** Choose whether to require or disallow IAM authentication for connections to your proxy.
- **Subnets.** This field is prepopulated with all the subnets associated with your VPC. Remove any subnets that you don't need for this proxy. You must leave at least two subnets.

Additional connectivity configuration:

- **VPC security group.** Choose an existing VPC security group. You can also choose for the AWS Management Console to create a new security group for you and use that.

Advanced configuration:

- **Enable enhanced logging.** You can enable this setting to troubleshoot proxy compatibility or performance issues. When this setting is enabled, RDS Proxy includes detailed information about SQL statements in its logs. This information helps you to debug issues involving SQL behavior or the performance and scalability of the proxy connections. The debug information includes the text of SQL statements that you submit through the proxy. Thus, only enable this setting when needed for debugging, and only when you have security measures in place to safeguard any sensitive information that appears in the logs.
5. Select **RDS Proxy is now available in Preview** to acknowledge the terms and conditions for the preview.
 6. Choose **Create Proxy**.

AWS CLI

To create a DB instance, use the AWS CLI command `create-db-proxy`.

Example

For Linux, OS X, or Unix:

```
aws rds create-db-proxy \
--db-proxy-name proxy_name \
--role-arn iam_role \
--engine-family { MySQL } \
--vpc-subnet-ids space_separated_list \
```

```
[--vpc-security-group-ids space_separated_list] \
[--auth ProxyAuthenticationConfig_JSON_string] \
[--require-tls | --no-require-tls] \
[--idle-client-timeout value] \
[--debug-logging | --no-debug-logging] \
[--tags comma_separated_list]
```

For Windows:

```
aws rds create-db-proxy ^
--db-proxy-name proxy_name ^
--role-arn iam_role ^
--engine-family { MYSQL } ^
--vpc-subnet-ids space_separated_list ^
[--vpc-security-group-ids space_separated_list] ^
[--auth ProxyAuthenticationConfig_JSON_string] ^
[--require-tls | --no-require-tls] ^
[--idle-client-timeout value] ^
[--debug-logging | --no-debug-logging] ^
[--tags comma_separated_list]
```

To create the required information and associations for the proxy, you also use these commands:

```
aws rds create-db-proxy-target-group
--name target_group_name
--db-proxy-name proxy_name
--connection-pool-defaults ConnectionPoolConfiguration
--tags comma_separated_list
```

```
aws rds register-db-proxy-targets
--db-proxy-name value
[--target-group-name target_group_name]
[--db-instance-identifiers space_separated_list] # rds db instances, or
[--db-cluster-identifiers cluster_id]           # rds db cluster (all instances), or
[--db-cluster-endpoint endpoint_name]          # rds db cluster endpoint (all
instances)
```

RDS API

To create an RDS Proxy, you call the Amazon RDS API function [CreateDBProxy](#). You pass a parameter with the [AuthConfig](#) data structure.

You also create the associated target group by calling the function [CreateDBProxyTargetGroup](#). Then you associate an RDS DB instance or Aurora DB cluster with the target group by calling the function [RegisterDBProxyTargets](#).

Viewing RDS Proxy

After you create one or more RDS proxies, you can view them all to examine their configuration details and choose which ones to modify, delete, and so on.

In particular, you need the endpoint for the proxy to use in the connection string for any database applications that use the proxy.

AWS Management Console

To view your RDS Proxy

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the AWS Management Console, choose the AWS Region in which you created the RDS Proxy.
3. In the navigation pane, choose **Proxies**.
4. Choose the name of an RDS proxy to display its details.
5. On the details page, the **Target groups** section shows how the proxy is associated with a specific RDS DB instance or Aurora DB cluster. You can follow the link to the **default** target group page to see more details about the association between the proxy and the database. This page is where you see settings that you specified when creating the proxy, such as maximum connection percentage, connection borrow timeout, engine compatibility, and session pinning filters.

CLI

To view your RDS Proxy using the CLI, use the `describe-db-proxies` command. By default, it displays all proxies owned by your AWS account. You can specify `--db-proxy-name` parameter to see details for a single proxy.

```
aws rds describe-db-proxies
[--db-proxy-name proxy_name]
```

To view the other information associated with the proxy, use these commands:

```
aws rds describe-db-proxy-target-groups \
--db-proxy-name proxy_name

aws rds describe-db-proxy-targets \
--db-proxy-name proxy_name
```

Use the following sequence of commands to see more detail about the things that are associated with the proxy:

- Run `describe-db-proxies` to get a list of proxies.
- Run `describe-db-proxy-target-groups --db-proxy-name` using the name of the proxy as the parameter. This output shows connection parameters such as the maximum percentage of connections that the proxy can use.
- Run `describe-db-proxy-targets` to see the details of the RDS DB instance or Aurora DB cluster associated with the returned target group.

RDS API

To view your proxies using the RDS API, use the [DescribeDBProxies](#) operation. It returns values of the [DBProxy](#) data type.

You can use the proxy identifiers from the return value to see details of the connection settings for the proxy by using the [DescribeDBProxyTargetGroups](#) operation. It returns values of the [DBProxyTargetGroup](#) data type.

You can see the RDS instance or Aurora DB cluster associated with the proxy by using the [DescribeDBProxyTargets](#) operation. It returns values of the [DBProxyTarget](#) data type.

Deleting an RDS Proxy

You can delete a proxy if you no longer need it. You might delete a proxy because the application that was using it is no longer relevant. Or you might delete a proxy if you take the DB instance or cluster associated with it out of service.

AWS Management Console

To delete a proxy

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Proxies**.
3. Choose the proxy to delete from the list.
4. Choose **Delete Proxy**.

AWS CLI

To delete a DB proxy, use the AWS CLI command `delete-db-proxy`.

Example

```
aws rds delete-db-proxy --name proxy_name
```

```
aws rds deregister-db-proxy-targets
  --db-proxy-name proxy_name
  [--target-group-name target_group_name]
  [--target-ids comma_separated_list]      # or
  [--db-instance-identifiers instance_id]    # or
  [--db-cluster-identifiers cluster_id]
```

RDS API

To delete a DB proxy, call the Amazon RDS API function `DeleteDBProxy`. To delete related items and associations, you also call the functions `DeleteDBProxyTargetGroup` and `DeregisterDBProxyTargets`.

Modifying an RDS Proxy

You can change certain settings associated with a proxy after you create the proxy. You do so by modifying the proxy itself, its associated target group, or both. Each proxy has an associated target group.

AWS Management Console

To modify the settings for a proxy

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Proxies**.
3. In the list of proxies, choose the proxy whose settings you want to modify or go to its details page.
4. For **Actions**, choose **Modify**.

5. Enter or choose the properties to modify.
 - You can rename the proxy by entering a new identifier.
 - You can turn the requirement for Transport layer Security (TLS) on or off.
 - You can enter a time period for the idle connection timeout.
 - You can add or remove Secrets Manager secrets. These secrets correspond to database user names and passwords.
 - You can change the IAM role used to retrieve the secrets from Secrets Manager.
 - You can require or disallow IAM authentication for connections to the proxy.
 - You can add or remove VPC subnets for the proxy to use.
 - You can add or remove VPC security groups for the proxy to use.
 - You can enable or disable enhanced logging.
6. Choose **Modify**.

If you didn't find the settings you intended to change, continue with the following procedure to update the target group for the proxy. Each proxy has one associated target group, which is created automatically along with the proxy.

To modify the settings for a proxy target group

The target group associated with a proxy controls the settings related to the physical database connections. You can only modify the target group from the proxy details page, not from the list on the **Proxies** page.

1. From the **Proxies** page, go to the details page for a proxy.
2. For **Target groups**, choose the **default** link. Currently, all proxies have a single target group named **default**.
3. On the details page for the **default** target group, choose **Modify**.
4. Certain properties, such as the target group identifier and the database engine, are fixed. You can't change these properties.

Choose new settings for the properties you can modify:

- You can choose a different RDS DB instance or Aurora cluster.
- You can adjust what percentage of the maximum available connections the proxy can use.
- You can select a session pinning filter. This setting can help reduce performance issues due to insufficient transaction-level reuse for connections. Using this setting requires understanding of application behavior and the circumstances under which RDS Proxy pins a session to a database connection.
- You can adjust the connection borrow timeout interval. This setting determines how long the proxy waits for a connection to become available before returning a timeout error, when the proxy already has the maximum number of connections being used.

5. Choose **Modify target group**.

AWS CLI

To modify a proxy using the AWS CLI, you use the commands [modify-db-proxy](#), [modify-db-proxy-target-group](#), [deregister-db-proxy-targets](#), and [register-db-proxy-targets](#)

With the `modify-db-proxy` command, you can change properties such as the set of Secrets Manager secrets used by the proxy, whether TLS is required, the idle client timeout, whether to log additional information from SQL statements for debugging, the IAM role used to retrieve Secrets Manager secrets, and the security groups used by the proxy.

The following example shows how to rename an existing proxy.

```
aws rds modify-db-proxy --db-proxy-name the-proxy --new-db-proxy-name the-new-name
```

With the `modify-db-proxy-target-group` command, you can modify connection-related settings or rename the target group. Currently, all proxies have a single target group named `default`. When working with this target group, you specify the name of the proxy and `default` for the name of the target group.

The following example shows how to first check the `MaxConnectionsPercent` setting for a proxy and then change it, using the target group.

```
$ aws rds describe-db-proxy-target-groups --db-proxy-name the-proxy
{
    "TargetGroups": [
        {
            "Status": "available",
            "UpdatedDate": "2019-11-30T16:49:30.342Z",
            "ConnectionPoolConfig": {
                "MaxIdleConnectionsPercent": 50,
                "ConnectionBorrowTimeout": 120,
                "MaxConnectionsPercent": 100,
                "SessionPinningFilters": []
            },
            "TargetGroupName": "default",
            "CreatedDate": "2019-11-30T16:49:27.940Z",
            "DBProxyName": "the-proxy",
            "IsDefault": true
        }
    ]
}

$ aws rds modify-db-proxy-target-group --db-proxy-name the-proxy --target-group-name default --connection-pool-config '{ "MaxIdleConnectionsPercent": 75 }'
{
    "DBProxyTargetGroup": {
        "Status": "available",
        "UpdatedDate": "2019-12-02T04:09:50.420Z",
        "ConnectionPoolConfig": {
            "MaxIdleConnectionsPercent": 75,
            "ConnectionBorrowTimeout": 120,
            "MaxConnectionsPercent": 100,
            "SessionPinningFilters": []
        },
        "TargetGroupName": "default",
        "CreatedDate": "2019-11-30T16:49:27.940Z",
        "DBProxyName": "the-proxy",
        "IsDefault": true
    }
}
```

With the `deregister-db-proxy-targets` and `register-db-proxy-targets` commands, you change which RDS DB instance or Aurora DB cluster the proxy is associated with through its target group. Currently, each proxy can connect to one RDS DB instance or Aurora DB cluster. The target group tracks the connection details for all the RDS DB instances in a multi-AZ configuration, or all the DB instances in an Aurora cluster.

The following example starts with a proxy that is associated with an Aurora cluster named `cluster-56-2019-11-14-1399`. The example shows how to change the proxy so that it can connect

to a different cluster named `provisioned-cluster`. When you work with an RDS DB instance, you specify the `--db-instance-identifier` option. When you work with an Aurora DB cluster, you specify the `--db-cluster-identifier` option instead.

```
$ aws rds describe-db-proxy-targets --db-proxy-name the-proxy
{
    "Targets": [
        {
            "Endpoint": "instance-9814.demo.us-east-1.rds.amazonaws.com",
            "Type": "RDS_INSTANCE",
            "Port": 3306,
            "RdsResourceId": "instance-9814"
        },
        {
            "Endpoint": "instance-8898.demo.us-east-1.rds.amazonaws.com",
            "Type": "RDS_INSTANCE",
            "Port": 3306,
            "RdsResourceId": "instance-8898"
        },
        {
            "Endpoint": "instance-1018.demo.us-east-1.rds.amazonaws.com",
            "Type": "RDS_INSTANCE",
            "Port": 3306,
            "RdsResourceId": "instance-1018"
        },
        {
            "Type": "TRACKED_CLUSTER",
            "Port": 0,
            "RdsResourceId": "cluster-56-2019-11-14-1399"
        },
        {
            "Endpoint": "instance-4330.demo.us-east-1.rds.amazonaws.com",
            "Type": "RDS_INSTANCE",
            "Port": 3306,
            "RdsResourceId": "instance-4330"
        }
    ]
}

$ aws rds deregister-db-proxy-targets --db-proxy-name the-proxy --db-cluster-identifier
cluster-56-2019-11-14-1399

$ aws rds describe-db-proxy-targets --db-proxy-name the-proxy
{
    "Targets": []
}

$ aws rds register-db-proxy-targets --db-proxy-name the-proxy --db-cluster-identifier
provisioned-cluster
{
    "DBProxyTargets": [
        {
            "Type": "TRACKED_CLUSTER",
            "Port": 0,
            "RdsResourceId": "provisioned-cluster"
        },
        {
            "Endpoint": "gkldje.demo.us-east-1.rds.amazonaws.com",
            "Type": "RDS_INSTANCE",
            "Port": 3306,
            "RdsResourceId": "gkldje"
        },
        {
            "Endpoint": "provisioned-1.demo.us-east-1.rds.amazonaws.com",
            "Type": "RDS_INSTANCE",
            "Port": 3306,
            "RdsResourceId": "provisioned-1"
        }
    ]
}
```

```
        "Type": "RDS_INSTANCE",
        "Port": 3306,
        "RdsResourceId": "provisioned-1"
    ]
}
```

RDS API

To modify a proxy using the RDS API, you use the operations [ModifyDBProxy](#), [ModifyDBProxyTargetGroup](#), [DeregisterDBProxyTargets](#), and [RegisterDBProxyTargets](#) operations.

With [ModifyDBProxy](#), you can change properties such as the set of Secrets Manager secrets used by the proxy, whether TLS is required, the idle client timeout, whether to log additional information from SQL statements for debugging, the IAM role used to retrieve Secrets Manager secrets, and the security groups used by the proxy.

With [ModifyDBProxyTargetGroup](#), you can modify connection-related settings or rename the target group. Currently, all proxies have a single target group named default. When working with this target group, you specify the name of the proxy and default for the name of the target group.

With [DeregisterDBProxyTargets](#) and [RegisterDBProxyTargets](#), you change which RDS DB instance or Aurora DB cluster the proxy is associated with through its target group. Currently, each proxy can connect to one RDS DB instance or Aurora DB cluster. The target group tracks the connection details for all the RDS DB instances in a multi-AZ configuration, or all the DB instances in an Aurora cluster.

Adding a New Database User

In some cases, you might add a new database user to an RDS DB instance or Aurora cluster that's associated with a proxy. If so, add or repurpose a Secrets Manager secret to store the credentials for that user. To do this, choose one of the following options:

- Create a new Secrets Manager secret, using the procedure described in [Setting Up Database Credentials in AWS Secrets Manager \(p. 774\)](#).
- If the new user takes the place of an existing one, update the credentials stored in the proxy's Secrets Manager secret for the existing user.

Changing the Password for a Database User

In some cases, you might change the password for a database user in an RDS DB instance or Aurora cluster that's associated with a proxy. If so, update the corresponding Secrets Manager secret with the new password.

Connecting to a Database through RDS Proxy

This is preview documentation for Amazon RDS Proxy. It is subject to change.

You connect to an RDS DB instance or Aurora DB cluster through a proxy in generally the same way as you connect directly to the database. The main difference is that you specify the proxy endpoint instead of the instance or cluster endpoint.

For an Aurora DB cluster, all proxy connections have read-write capability and use the writer instance. If you use the reader endpoint for read-only connections, you continue using the reader endpoint the same way.

Connecting to a Proxy Using Native Authentication

Use the following general procedure to connect to a proxy using native authentication:

- Find the proxy endpoint. In the AWS Management Console, you can find the endpoint on the details page for the corresponding proxy. With the AWS CLI, you can use the `describe-db-proxies` command. The following example shows how:

```
# Add --output text to get output as a simple tab-separated list.
$ aws rds describe-db-proxies --query '*[*].{DBProxyName:DBProxyName,Endpoint:Endpoint}'
[
  [
    {
      "Endpoint": "the-proxy.proxy-demo.us-east-1.rds.amazonaws.com",
      "DBProxyName": "the-proxy"
    },
    {
      "Endpoint": "the-proxy-other-secret.proxy-demo.us-east-1.rds.amazonaws.com",
      "DBProxyName": "the-proxy-other-secret"
    },
    {
      "Endpoint": "the-proxy-rds-secret.proxy-demo.us-east-1.rds.amazonaws.com",
      "DBProxyName": "the-proxy-rds-secret"
    },
    {
      "Endpoint": "the-proxy-t3.proxy-demo.us-east-1.rds.amazonaws.com",
      "DBProxyName": "the-proxy-t3"
    }
  ]
]
```

- Specify that endpoint as the host parameter in the connection string for your client application. For example, specify the proxy endpoint as the parameter for the `mysql -h` option.
- Supply the same database user name and password as you normally do.

Connecting to a Proxy Using IAM Authentication

Follow the same general procedure as for connecting to an RDS DB instance or Aurora cluster using IAM authentication. Instead of the instance, cluster, or reader endpoint, specify the proxy endpoint. For details of that procedure, see [Connecting to Your DB Instance from the Command Line: AWS CLI and mysql Client](#).

Managing RDS Proxy

The following explain aspects of RDS Proxy operation and configuration that help your application to make the most efficient use of database connections and achieve maximum connection reuse. The more you can take advantage of connection reuse, the more CPU and memory overhead you can save. This in turn reduces latency for your application and enables the database to devote more of its resources to processing application requests.

Connection Limits and Timeouts

RDS Proxy makes use of the `max_connections` setting for your RDS DB instance or Aurora DB cluster. This setting represents the overall upper limit on the connections that the proxy can open at any one time. In Aurora clusters and RDS multi-AZ configurations, the `max_connections` value that the proxy uses is the one for the Aurora primary instance or the RDS writer instance.

The proxy setting for maximum connections represents a percentage of the `max_connections` value. If you have multiple applications all using the same database, you can effectively divide their connection

quotas by using a proxy for each application with a specific percentage of `max_connections`. If you do so, ensure that the percentages add up to 100 or less for all proxies associated with the same database.

RDS Proxy periodically disconnects idle connections and returns them to the connection pool. You can adjust this timeout interval. Doing so helps your applications to deal with stale resources, especially if the application mistakenly leaves a connection open while holding important database resources.

Connection Pooling

As described in [Connection Pooling \(p. 772\)](#), connection pooling is a crucial RDS Proxy feature. Following, you can learn how to make the most efficient use of connection pooling and transaction-level connection reuse (multiplexing).

Because the connection pool is managed by RDS Proxy, you can monitor it and adjust connection limits and timeout intervals without changing your application code.

For each proxy, you can specify an upper limit on the number of connections used by the connection pool. You specify the limit as a percentage. This percentage applies to the maximum connections configured in the database. The exact number varies depending on the DB instance size and configuration settings. For example, suppose that you configured RDS Proxy to use 75% of the maximum connections for the database. For MySQL, the maximum value is defined by the `max_connections` configuration parameter. In this case, the other 25% of maximum connections would remain available to assign to other proxies or for connections that don't go through a proxy. If the database doesn't have many simultaneous connections, or some connections stay idle for long periods, the proxy might keep less than 75% of the maximum connections open at a particular time.

The overall number of connections available for the connection pool changes as you update the `max_connections` configuration setting that applies to an RDS DB instance or a Aurora MySQL cluster.

The proxy doesn't reserve all of these connections in advance. Thus, you can specify a relatively large percentage, and those connections are only opened when the proxy becomes busy enough to need them.

You can choose how long to wait for a connection to become available for use by your application. This setting is represented by the **Connection borrow timeout** option when you create a proxy. This setting specifies how long to wait for a connection to become available in the connection pool before returning a timeout error. It applies when the number of connections is at the maximum, and so no connections are available in the connection pool. It also applies if no writer instance is available because a failover operation is in process. This setting enables you to limit the wait period that is most suitable for your application without having to change the query timeout in your application code.

Pinning

Multiplexing is more efficient when database requests don't rely on state information from previous requests. In that case, RDS Proxy can reuse a connection at the conclusion of each transaction. Examples of such state information include most variables and configuration parameters that you can change through `SET` or `SELECT` statements. SQL transactions on a client connection can multiplex between underlying database connections by default.

Your connections to the proxy can enter a state known as *pinning*. When a connection is pinned, each later transaction uses the same underlying database connection until the session ends. RDS Proxy automatically pins a client connection to a specific DB connection when it detects a session state change that isn't appropriate for other sessions. Pinning reduces the effectiveness of connection reuse. If all or almost all of your connections experience pinning, you might modify your application code or workload to reduce the conditions that cause the pinning.

For example, if your application changes a session variable or configuration parameter, later statements can rely on the new variable or parameter to be in effect. Thus, when RDS Proxy processes requests to change session variables or configuration settings, it pins that session to the DB connection. That way, the session state remains in effect for all later transactions in the same session. This rule doesn't apply to all parameters you can set. RDS Proxy tracks changes to the character set, collation, time zone,

autocommit, current database, SQL mode, and session_track_schema settings. RDS Proxy doesn't pin the session when you modify them. In that case, RDS Proxy only reuses the connection for other sessions that have the same values for those settings.

Performance tuning for RDS Proxy involves trying to maximize transaction-level connection reuse (multiplexing) by minimizing pinning. You can do so by using the following techniques:

- Avoid unnecessary database requests that could cause pinning.
- Set variables and configuration settings consistently across all connections. That way, later sessions are more likely to reuse connections that have those particular settings.
- Apply a session pinning filter to the proxy. You can exempt certain kinds of operations from pinning the session if you know that doing so doesn't affect the correct operation of your application.
- Monitor the CloudWatch metric `DatabaseConnectionsCurrentlySessionPinned` to see how frequently pinning occurs. For information about this and other CloudWatch metrics, see [Monitoring RDS Proxy \(p. 789\)](#).
- If you use `SET` statements to perform identical initialization for each client connection, you can do so while still preserving transaction-level multiplexing. In this case, you move the statements that set up the initial session state into the initialization query used by a proxy. This property is a string containing one or more SQL statements, separated by semicolons. For example, you can define an initialization query for a proxy that sets certain configuration parameters. Then, RDS Proxy applies those settings whenever it sets up a new connection for that proxy. You can remove the corresponding `SET` statements from your application code, so that they don't interfere with transaction-level multiplexing. You can work with the initialization query through the AWS CLI and RDS API. Currently, you can't set this property through the AWS Management Console.

The proxy pins the session to the current connection in the following situations where multiplexing might cause unexpected behavior:

- Any statement with a text size greater than 4 KB causes the proxy to pin the session.
- Prepared statements cause the proxy to pin the session. This rule applies whether the prepared statement uses SQL text or the binary protocol.
- An explicit `LOCK TABLE`, `LOCK TABLES`, or `FLUSH TABLES WITH READ LOCK` statement causes the proxy to pin the session.
- Setting a user variable or a system variable (with some exceptions) causes the proxy to pin the session. If this situation reduces your connection reuse too much, you can choose for `SET` operations not to cause pinning.
- Creating a temporary table causes the proxy to pin the session. That way, the contents of the temporary table are preserved throughout the session regardless of transaction boundaries.
- Calling the `ROW_COUNT()`, `FOUND_ROWS()`, or `LAST_INSERT_ID()` functions sometimes causes pinning and sometimes not. The behavior might be different between Aurora MySQL versions that are compatible with MySQL 5.6 and MySQL 5.7.

Calling stored procedures and stored functions doesn't cause pinning. RDS Proxy doesn't detect any session state changes resulting from such calls. Therefore, ensure your application doesn't change session state inside stored routines and rely on that session state to persist across transactions. For example, if a stored procedure creates a temporary table that is intended to persist across transactions, that application currently isn't compatible with RDS Proxy.

If you have expert knowledge about your application behavior, you can choose to skip the pinning behavior for certain application statements. You do so by choosing the **Session pinning filters** option when creating the proxy. Currently, you can opt out of session pinning for the following kinds of application statements:

- Setting session variables and configuration settings.

To see metrics about how often pinning occurs for a proxy, see [Monitoring RDS Proxy \(p. 789\)](#).

Monitoring RDS Proxy

This is preview documentation for Amazon RDS Proxy. It is subject to change.

You can monitor RDS Proxy using Amazon CloudWatch. CloudWatch collects and processes raw data from the proxies into readable, near real-time metrics. To find these metrics in the CloudWatch console, choose **Metrics**, then choose **RDS**, and choose **Per-Proxy Metrics**.

Note

RDS publishes these metrics for each underlying EC2 instance associated with the proxy. A single proxy might be served by more than one EC2 instance. Use CloudWatch statistics to aggregate the values for a proxy across all the associated instances.

All RDS Proxy metrics are in the group `proxy`, with a dimension of `ProxyName`.

Metric	Valid Period	Description
<code>ClientConnectionsReceived</code>	1 minute and above	The number of client connection requests received. The most useful statistic for this metric is Sum.
<code>ClientConnectionsSetupSucceeded</code>	1 minute and above	The number of client connections successfully established with any authentication mechanism with or without TLS. The most useful statistic for this metric is Sum.
<code>ClientConnectionsSetupFailedAuth</code>	1 minute and above	The number of client connection attempts which failed due to bad authentication or TLS misconfiguration. The most useful statistic for this metric is Sum.
<code>ClientConnectionsClosed</code>	1 minute and above	The number of client connections closed. The most useful statistic for this metric is Sum.
<code>ClientConnections</code>	1 minute	The current number of client connections. This is reported every minute. The most useful statistic for this metric is Sum.
<code>QueryRequests</code>	1 minute and above	The number of queries received. Note: Multi-

Metric	Valid Period	Description
		statement query is counted as one query. The most useful statistic for this metric is Sum.
DatabaseConnectionRequests	1 minute and above	The number of requests to create a database connection. The most useful statistic for this metric is Sum.
DatabaseConnectionsSetupSucceeded	1 minute and above	The number of database connections successfully established with or without TLS. The most useful statistic for this metric is Sum.
DatabaseConnectionsSetupFailed	1 minute and above	The number of database connection requests which failed. The most useful statistic for this metric is Sum.
MaxDatabaseConnectionsAllowed	1 minute	The maximum number of database connections allowed. This is reported every minute. The most useful statistic for this metric is Sum.
DatabaseConnections	1 minute	The current number of database connections. This is reported every minute. The most useful statistic for this metric is Sum.
DatabaseConnectionsCurrentlyBorrowed	1 minute	The current number of database connections in the borrow state. This is reported every minute. The most useful statistic for this metric is Sum.

Metric	Valid Period	Description
DatabaseConnectionsCurrentlySessionPinned	1 minute	The current number of database connections currently pinned due to session state-altering operations in client requests. This is reported every minute. The most useful statistic for this metric is Sum.
DatabaseConnectionsCurrentlyInTransaction	1 minute	The current number of database connections in transaction. This is reported every minute. The most useful statistic for this metric is Sum.

Limitations for RDS Proxy

This is preview documentation for Amazon RDS Proxy. It is subject to change.

The following limitations apply to RDS Proxy during the public preview:

- The public preview is available only in these AWS Regions: US East (N. Virginia), US East (Ohio), US West (Oregon), Asia Pacific (Tokyo), and Europe (Ireland).
- You can have up to 20 proxies for each AWS account ID.
- In an Aurora cluster, all of the connections in the connection pool are handled by the Aurora primary instance. To perform load balancing for read-intensive workloads, you still use the reader endpoint directly for the Aurora cluster.
- Currently, RDS Proxy is only available for the MySQL engine family. This engine family includes RDS MySQL 5.6 and 5.7, and Aurora versions 1 and 2. Proxy support for PostgreSQL databases isn't available in the public preview.
- You can't use RDS Proxy with RDS MySQL 8.0.
- You can't use RDS Proxy with Aurora Serverless clusters.
- You can't use RDS Proxy with Aurora multi-master clusters.
- You can use RDS Proxy with Amazon RDS MySQL and Aurora MySQL. You can't use it with self-managed MySQL databases in EC2 instances.
- Currently, all proxies listen on port 3306.
- Your RDS Proxy must be in the same VPC as the database. Although the database can be publicly accessible, the proxy can't be.
- Currently, proxies don't track any changes to the set of DB instances within an Aurora DB cluster. Those changes include operations such as host replacements, instance renames, port changes, scaling instances up or down, or adding or removing DB instances.
- Not all logic is implemented to pin sessions to database connections based on SQL statements and functions. For the most current pinning behavior, see [Pinning \(p. 787\)](#).

Command-Line Examples for RDS Proxy

This is preview documentation for Amazon RDS Proxy. It is subject to change.

To see how combinations of `mysql` connection commands and SQL statements interact with RDS Proxy, look at the following examples.

Example Preserving Connections Across a Failover

This example demonstrates how open connections continue working during a failover, for example when you reboot a database or it becomes unavailable due to a problem. This example uses a proxy named `the-proxy` and an Aurora DB cluster with DB instances `instance-8898` and `instance-9814`. When the `failover-db-cluster` command is run from the Linux command line, the writer instance that the proxy is connected to changes to a different DB instance. You can see that the DB instance associated with the proxy changes while the connection remains open.

```
$ mysql -h the-proxy.proxy-demo.us-east-1.rds.amazonaws.com -u admin_user -p
Enter password:
...
mysql> select @@aurora_server_id;
+-----+
| @@aurora_server_id |
+-----+
| instance-9814     |
+-----+
1 row in set (0.01 sec)

mysql>
[1]+  Stopped                  mysql -h the-proxy.proxy-demo.us-east-1.rds.amazonaws.com -
u admin_user -p
$ # Initially, instance-9814 is the writer.
$ aws rds failover-db-cluster --db-cluster-id cluster-56-2019-11-14-1399
JSON output
$ # After a short time, the console shows that the failover operation is complete.
$ # Now instance-8898 is the writer.
$ fg
mysql -h the-proxy.proxy-demo.us.us-east-1.rds.amazonaws.com -u admin_user -p

mysql> select @@aurora_server_id;
+-----+
| @@aurora_server_id |
+-----+
| instance-8898      |
+-----+
1 row in set (0.01 sec)

mysql>
[1]+  Stopped                  mysql -h the-proxy.proxy-demo.us-east-1.rds.amazonaws.com -
u admin_user -p
$ aws rds failover-db-cluster --db-cluster-id cluster-56-2019-11-14-1399
JSON output
$ # After a short time, the console shows that the failover operation is complete.
$ # Now instance-9814 is the writer again.
$ fg
mysql -h the-proxy.proxy-demo.us-east-1.rds.amazonaws.com -u admin_user -p

mysql> select @@aurora_server_id;
+-----+
| @@aurora_server_id |
+-----+
```

```
+-----+
| instance-9814      |
+-----+
1 row in set (0.01 sec)
+-----+-----+
| Variable_name | Value      |
+-----+-----+
| hostname      | ip-10-1-3-178 |
+-----+-----+
1 row in set (0.02 sec)
```

Example Adjusting the max_connections Setting for an Aurora DB Cluster

This example demonstrates how you can adjust the `max_connections` setting for an Aurora MySQL DB cluster. To do so, you create your own DB cluster parameter group based on the default parameter settings for clusters that are compatible with MySQL 5.6 or 5.7. You specify a value for the `max_connections` setting, overriding the formula that sets the default value. You associate the DB cluster parameter group with your DB cluster.

```
export REGION=us-east-1
export CLUSTER_PARAM_GROUP=rds-proxy-mysql-56-max-connections-demo
export CLUSTER_NAME=rds-proxy-mysql-56

aws rds create-db-parameter-group --region $REGION \
--db-parameter-group-family aurora5.6 \
--db-parameter-group-name $CLUSTER_PARAM_GROUP \
--description "Aurora MySQL 5.6 cluster parameter group for RDS Proxy demo."

aws rds modify-db-cluster --region $REGION \
--db-cluster-identifier $CLUSTER_NAME \
--db-cluster-parameter-group-name $CLUSTER_PARAM_GROUP

echo "New cluster param group is assigned to cluster:"
aws rds describe-db-clusters --region $REGION \
--db-cluster-identifier $CLUSTER_NAME \
--query '*[*].{DBClusterParameterGroup:DBClusterParameterGroup}'

echo "Current value for max_connections:"
aws rds describe-db-cluster-parameters --region $REGION \
--db-cluster-parameter-group-name $CLUSTER_PARAM_GROUP \
--query '*[*].{ParameterName:ParameterName,ParameterValue:ParameterValue}' \
--output text | grep "max_connections"

echo -n "Enter number for max_connections setting: "
read answer

aws rds modify-db-cluster-parameter-group --region $REGION --db-cluster-parameter-group-name $CLUSTER_PARAM_GROUP \
--parameters "ParameterName=max_connections,ParameterValue=$answer,ApplyMethod=immediate"

echo "Updated value for max_connections:"
aws rds describe-db-cluster-parameters --region $REGION \
--db-cluster-parameter-group-name $CLUSTER_PARAM_GROUP \
--query '*[*].{ParameterName:ParameterName,ParameterValue:ParameterValue}' \
--output text | grep "max_connections"
```

Troubleshooting for RDS Proxy

This is preview documentation for Amazon RDS Proxy. It is subject to change.

Common Issues and Solutions

You can refer to the following list of symptoms to see the possible causes and solutions to problems you might encounter using RDS Proxy. During the public preview, error handling code isn't complete, so you might experience cryptic error messages from underlying services.

You might encounter the following issues while creating a new proxy.

Error or Symptom	Causes or Workarounds
403: The security token included in the request is invalid	<ul style="list-style-type: none"> Select an existing IAM role instead of choosing to create a new one.

You might encounter the following issues while connecting to a proxy.

Error or Symptom	Causes or Workarounds
ERROR 2003 (HY000): Can't connect to MySQL server on ' <i>proxy_endpoint</i> ' (110)	<p>The proxy endpoint exists, but couldn't establish a connection to the RDS DB instance or Aurora DB cluster. This error typically occurs after a timeout period. Some possible reasons include the following:</p> <ul style="list-style-type: none"> You might have specified a nonexistent database user name. You might have entered an incorrect database user password. You might not have a Secrets Manager secret with the credentials of the specified database user. The credentials in the Secrets Manager secret might not match the credentials of the corresponding database user. You might not have permission to retrieve the corresponding Secrets Manager secret.
ERROR 2005 (HY000): Unknown MySQL server host ' <i>proxy_endpoint</i> ' (0)	<p>The proxy endpoint doesn't exist. This error typically occurs immediately. Some possible reasons include the following:</p> <ul style="list-style-type: none"> You might have deleted the proxy. You might have copied or typed the proxy endpoint incorrectly. You might have specified an incorrect port. Currently, all proxies listen on port 3306.
ERROR 2013 (HY000): Lost connection to MySQL server at 'reading initial communication packet', system error: 2	<p>The proxy couldn't log into the database using the credentials from the Secrets Manager secrets. Some possible reasons include the following:</p> <ul style="list-style-type: none"> The Secrets Manager secrets associated with the proxy don't include one with the user name you're connecting as. You can add a secret with correct credentials to your proxy if so. The Secrets Manager secret for the user name you're connecting as has the wrong password. You can update the secret with the incorrect credentials if so. The Secrets Manager secret for some other user, not the one you're connecting as, has an incorrect password. You can remove the secret with the incorrect credentials if so.

CloudWatch Logs

You can find logs of RDS Proxy activity under CloudWatch in the AWS Management Console. Each proxy has an entry in the **Log groups** page.

Important

These logs are intended for human consumption for troubleshooting purposes and not for programmatic access. The format and content of the logs is subject to change (especially during preview).

Verifying Connectivity to a Proxy

You can use the following commands to verify that all components of the connection mechanism can communicate with the other components.

Examine the proxy itself. Also examine the associated target groups and the targets specified in the target groups. Check that the details of the targets match the RDS DB instance or Aurora DB cluster that you intend to associate with the proxy.

```
aws rds describe-db-proxies --db-proxy-name $DB_PROXY_NAME
aws rds describe-db-proxy-target-groups --db-proxy-name $DB_PROXY_NAME
aws rds describe-db-proxy-targets --db-proxy-name $DB_PROXY_NAME
```

If the following nc command reports success, you can access the proxy endpoint from the EC2 instance or other system where you're logged in. This command reports failure if you're not in the same VPC as the proxy and the associated database. You might be able to log directly in to the database without being in the same VPC. However, you can't log into the proxy unless you're in the same VPC.

```
nc -zx proxy_endpoint 3306
```

You can use the following commands to make sure that your EC2 instance has the required properties. In particular, the VPC for the EC2 instance must be the same as the VPC for the RDS DB instance or Aurora DB cluster that the proxy connects to.

```
aws ec2 describe-instances --instance-ids your_ec2_instance_id
```

Examine the Secrets Manager secrets used for the proxy.

```
aws secretsmanager list-secrets
aws secretsmanager get-secret-value --secret-id your_secret_id
```

Make sure that the `SecretString` field displayed by `get-secret-value` is encoded as a JSON string that includes `username` and `password` fields. The following example shows the format of the `SecretString` field.

```
{
    "ARN": "some_arn",
    "Name": "some_name",
    "VersionId": "some_version_id",
    "SecretString": "{\"username\":\"some_username\",\"password\":\"some_password\"}",
    "VersionStages": [ "some_stage" ],
    "CreatedDate": some_timestamp
```

}

Examine the RDS DB instance or Aurora DB cluster that you associated with the proxy. For an Aurora DB cluster, also examine each of the DB instances in the cluster. Make sure that their properties match what's expected for the proxy, such as the port being 3306.

```
aws rds describe-db-clusters --db-cluster-id $DB_CLUSTER_ID
aws rds describe-db-instances --db-instance-id $DB_INSTANCE_ID1
aws rds describe-db-instances --db-instance-id $DB_INSTANCE_ID2
```

Updating Applications to Connect to MySQL DB Instances Using New SSL/TLS Certificates

As of September 19, 2019, Amazon RDS has published new Certificate Authority (CA) certificates for connecting to your RDS DB instances using Secure Socket Layer or Transport Layer Security (SSL/TLS). The previous CA certificates expire on March 5, 2020. Following, you can find information about updating your applications to use the new certificates. If your application connects to an RDS DB instance using SSL/TLS, you must take the following steps before **March 5, 2020**. Doing this means you can avoid interruption of connectivity between your applications and your RDS DB instances.

This topic can help you to determine whether any client applications use SSL/TLS to connect to your DB instances. If they do, you can further check whether those applications require certificate verification to connect.

Note

Some applications are configured to connect to MySQL DB instances only if they can successfully verify the certificate on the server.

For such applications, you must update your client application trust stores to include the new CA certificates.

After you update your CA certificates in the client application trust stores, you can rotate the certificates on your DB instances. We strongly recommend testing these procedures in a development or staging environment before implementing them in your production environments.

For more information about certificate rotation, see [Rotating Your SSL/TLS Certificate \(p. 433\)](#).

For information about using SSL/TLS with MySQL DB instances, see [Using SSL with a MySQL DB Instance \(p. 746\)](#).

Topics

- [Determining Whether any Applications are Connecting to Your MySQL DB Instance Using SSL \(p. 797\)](#)
- [Determining Whether a Client Requires Certificate Verification to Connect \(p. 798\)](#)
- [Updating Your Application Trust Store \(p. 799\)](#)
- [Example Java Code for Establishing SSL Connections \(p. 800\)](#)

Determining Whether any Applications are Connecting to Your MySQL DB Instance Using SSL

If you are using Amazon RDS for MySQL version 5.7 or 8.0 and the Performance Schema is enabled, run the following query to check if connections are using SSL/TLS. For information about enabling the Performance Schema, see [Performance Schema Quick Start](#) in the MySQL documentation.

```
mysql> SELECT id, user, host, connection_type
    FROM performance_schema.threads pst
    INNER JOIN information_schema.processlist isp
    ON pst.processlist_id = isp.id;
```

In this sample output, you can see both your own session (`admin`) and an application logged in as `webapp1` are using SSL.

```
+---+-----+-----+-----+
| id | user           | host          | connection_type |
+---+-----+-----+-----+
| 8  | admin          | 10.0.4.249:42590 | SSL/TLS        |
| 4  | event_scheduler | localhost      | NULL          |
| 10 | webapp1        | 159.28.1.1:42189 | SSL/TLS        |
+---+-----+-----+-----+
3 rows in set (0.00 sec)
```

If you are using Amazon RDS for MySQL versions 5.5 or 5.6, then you can't determine from the server side whether applications are connecting with or without SSL. For those versions, you can determine whether SSL is used by examining the application's connection method. In the following section, you can find more information on examining the client connection configuration.

Determining Whether a Client Requires Certificate Verification to Connect

You can check whether JDBC clients and MySQL clients require certificate verification to connect.

JDBC

The following example with MySQL Connector/J 8.0 shows one way to check an application's JDBC connection properties to determine whether successful connections require a valid certificate. For more information on all of the JDBC connection options for MySQL, see [Configuration Properties](#) in the MySQL documentation.

When using the MySQL Connector/J 8.0, an SSL connection requires verification against the server CA certificate if your connection properties have `sslMode` set to `VERIFY_CA` or `VERIFY_IDENTITY`, as in the following example.

```
Properties properties = new Properties();
properties.setProperty("sslMode", "VERIFY_IDENTITY");
properties.put("user", DB_USER);
properties.put("password", DB_PASSWORD);
```

MySQL

The following examples with the MySQL Client show two ways to check a script's MySQL connection to determine whether successful connections require a valid certificate. For more information on all of the connection options with the MySQL Client, see [Client-Side Configuration for Encrypted Connections](#) in the MySQL documentation.

When using the MySQL 5.7 or MySQL 8.0 Client, an SSL connection requires verification against the server CA certificate if for the `--ssl-mode` option you specify `VERIFY_CA` or `VERIFY_IDENTITY`, as in the following example.

```
mysql -h mysql-database.rds.amazonaws.com -uadmin -ppassword --ssl-ca=/tmp/ssl-cert.pem --
ssl-mode=VERIFY_CA
```

When using the MySQL 5.6 Client, an SSL connection requires verification against the server CA certificate if you specify the `--ssl-verify-server-cert` option, as in the following example.

```
mysql -h mysql-database.rds.amazonaws.com -uadmin -ppassword --ssl-ca=/tmp/ssl-cert.pem --ssl-verify-server-cert
```

Updating Your Application Trust Store

For information about updating the trust store for MySQL applications, see [Installing SSL Certificates](#) in the MySQL documentation.

Note

When you update the trust store, you can retain older certificates in addition to adding the new certificates.

Updating Your Application Trust Store for JDBC

You can update the trust store for applications that use JDBC for SSL/TLS connections.

To update the trust store for JDBC applications

1. Download the 2019 root certificate that works for all AWS Regions and put the file in the trust store directory.

For information about downloading the root certificate, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).

2. Convert the certificate to .der format using the following command.

```
openssl x509 -outform der -in rds-ca-2019-root.pem -out rds-ca-2019-root.der
```

Replace the file name with the one that you downloaded.

3. Import the certificate into the key store using the following command.

```
keytool -import -alias rds-root -keystore clientkeystore -file rds-ca-2019-root.der
```

4. Confirm that the key store was updated successfully.

```
keytool -list -v -keystore clientkeystore.jks
```

Enter the key store password when you are prompted for it.

Your output should contain the following.

```
rds-root,date, trustedCertEntry,  
Certificate fingerprint (SHA1):  
D4:0D:DB:29:E3:75:0D:FF:A6:71:C3:14:0B:BF:5F:47:8D:1C:80:96  
# This fingerprint should match the output from the below command  
openssl x509 -fingerprint -in rds-ca-2019-root.pem -noout
```

If you are using the mysql JDBC driver in an application, set the following properties in the application.

```
System.setProperty("javax.net.ssl.trustStore", certs);
System.setProperty("javax.net.ssl.trustStorePassword", "password");
```

When you start the application, set the following properties.

```
java -Djavax.net.ssl.trustStore=/path_to_truststore/MyTruststore.jks -
Djavax.net.ssl.trustStorePassword=my_truststore_password com.companyName.MyApplication
```

Example Java Code for Establishing SSL Connections

The following code example shows how to set up the SSL connection that validates the server certificate using JDBC.

```
public class MySQLSSLTest {

    private static final String DB_USER = "user name";
    private static final String DB_PASSWORD = "password";
    // This key store has only the prod root ca.
    private static final String KEY_STORE_FILE_PATH = "file-path-to-keystore";
    private static final String KEY_STORE_PASS = "keystore-password";

    public static void test(String[] args) throws Exception {
        Class.forName("com.mysql.jdbc.Driver");

        System.setProperty("javax.net.ssl.trustStore", KEY_STORE_FILE_PATH);
        System.setProperty("javax.net.ssl.trustStorePassword", KEY_STORE_PASS);

        Properties properties = new Properties();
        properties.setProperty("sslMode", "VERIFY_IDENTITY");
        properties.put("user", DB_USER);
        properties.put("password", DB_PASSWORD);

        Connection connection = DriverManager.getConnection("jdbc:mysql://jagdeeps-ssl-
test.cni62e2e7kwh.us-east-1.rds.amazonaws.com:3306", properties);
        Statement stmt=connection.createStatement();

        ResultSet rs=stmt.executeQuery("SELECT 1 from dual");

        return;
    }
}
```

Important

After you have determined that your database connections use SSL/TLS and have updated your application trust store, you can update your database to use the rds-ca-2019 certificates. For instructions, see step 3 in [Updating Your CA Certificate by Modifying Your DB Instance \(p. 434\)](#).

Modifying a DB Instance Running the MySQL Database Engine

You can change the settings of a DB instance to accomplish tasks such as adding additional storage or changing the DB instance class. This topic guides you through modifying an Amazon RDS MySQL DB instance, and describes the settings for MySQL instances.

We recommend that you test any changes on a test instance before modifying a production instance, so that you fully understand the impact of each change. This is especially important when upgrading database versions.

After you modify your DB instance settings, you can apply the changes immediately, or apply them during the next maintenance window for the DB instance. Some modifications cause an interruption by restarting the DB instance.

Note

When you modify a DB instance, Amazon RDS will reboot the instance if both of the following are true:

- You change the DB instance class.
- You specify a custom parameter group.

Console

To modify a MySQL DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to modify.
3. Choose **Modify**. The **Modify DB Instance** page appears.
4. Change any of the settings that you want. For information about each setting, see [Settings for MySQL DB Instances \(p. 803\)](#).
5. When all the changes are as you want them, choose **Continue** and check the summary of modifications.
6. To apply the changes immediately, choose **Apply immediately**. Choosing this option can cause an outage in some cases. For more information, see [Using the Apply Immediately Parameter \(p. 150\)](#).
7. On the confirmation page, review your changes. If they are correct, choose **Modify DB Instance** to save your changes.

Alternatively, choose **Back** to edit your changes, or choose **Cancel** to cancel your changes.

AWS CLI

To modify a MySQL DB instance by using the AWS CLI, call the `modify-db-instance` command. Specify the DB instance identifier, and the parameters for the settings that you want to modify. For information about each parameter, see [Settings for MySQL DB Instances \(p. 803\)](#).

Example

The following code modifies `mydbinstance` by setting the backup retention period to 1 week (7 days). The code enables automatic minor version upgrades by using `--auto-minor-version-upgrade`. To disable automatic minor version upgrades, use `--no-auto-minor-version-upgrade`. The

changes are applied during the next maintenance window by using `--no-apply-immediately`. Use `--apply-immediately` to apply the changes immediately. For more information, see [Using the Apply Immediately Parameter \(p. 150\)](#).

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
    --db-instance-identifier mydbinstance \
    --backup-retention-period 7 \
    --auto-minor-version-upgrade \
    --no-apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
    --db-instance-identifier mydbinstance ^
    --backup-retention-period 7 ^
    --auto-minor-version-upgrade ^
    --no-apply-immediately
```

RDS API

To modify a MySQL instance by using the Amazon RDS API, call the [ModifyDBInstance](#) operation. Specify the DB instance identifier, and the parameters for the settings that you want to modify. For information about each parameter, see [Settings for MySQL DB Instances \(p. 803\)](#).

Example

The following code modifies `mydbinstance` by setting the backup retention period to 1 week (7 days) and enabling automatic minor version upgrades. These changes are applied during the next maintenance window.

```
https://rds.amazonaws.com/
?Action=ModifyDBInstance
&ApplyImmediately=false
&AutoMinorVersionUpgrade=true
&BackupRetentionPeriod=7
&DBInstanceIdentifier=mydbinstance
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20131016/us-west-1/rds/aws4_request
&X-Amz-Date=20131016T233051Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=087a8eb41cb1ab0fc9ec1575f23e73757ffc6a1e42d7d2b30b9cc0be988cff97
```

Settings for MySQL DB Instances

The following table contains details about which settings you can modify, which settings you can't modify, when the changes can be applied, and whether the changes cause downtime for the DB instance.

Setting	Setting Description	When the Change Occurs	Downtime Notes
Allocated storage	The storage, in gigabytes, that you want to allocate for your DB instance.	If Apply immediately is set to true, the change occurs immediately.	No downtime. Performance might be

Setting	Setting Description	When the Change Occurs	Downtime Notes
	For more information, see Amazon RDS DB Instance Storage (p. 135) .	If Apply immediately is set to false, the change occurs during the next maintenance window.	degraded during the change.
Auto minor version upgrade	Enable auto minor version upgrade to enable your DB instance to receive preferred minor DB engine version upgrades automatically when they become available. Amazon RDS performs automatic minor version upgrades in the maintenance window.	–	–
Backup retention period	The number of days that automatic backups are retained. To disable automatic backups, set the backup retention period to 0. For more information, see Working With Backups (p. 254) .	If Apply immediately is set to true, the change occurs immediately. If Apply immediately is set to false and you change the setting from a nonzero value to another nonzero value, the change is applied asynchronously, as soon as possible. Otherwise, the change occurs during the next maintenance window.	An outage occurs if you change from 0 to a nonzero value, or from a nonzero value to 0.
Backup window	The time range during which automated backups of your databases occur. The backup window is a start time in Universal Coordinated Time (UTC), and a duration in hours. For more information, see Working With Backups (p. 254) .	The change is applied asynchronously, as soon as possible.	–
Certificate authority	The certificate that you want to use.	If Apply Immediately is set to true, the change occurs immediately. If Apply Immediately is set to false, the change occurs during the next maintenance window.	An outage occurs during this change.
Copy tags to snapshots	If you have any DB instance tags, this option copies them when you create a DB snapshot. For more information, see Tagging Amazon RDS Resources (p. 177) .	The change occurs immediately. This setting ignores the Apply immediately setting.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Database authentication	<p>The database authentication option you want to use.</p> <p>Choose Password authentication to authenticate database users with database passwords only.</p> <p>Choose Password and IAM DB authentication to authenticate database users with database passwords and user credentials through IAM users and roles. For more information, see IAM Database Authentication for MySQL and PostgreSQL (p. 461).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	–
Database port	<p>The port that you want to use to access the database.</p> <p>The port value must not match any of the port values specified for options in the option group for the DB instance.</p>	<p>The change occurs immediately. This setting ignores the Apply immediately setting.</p>	The DB instance is rebooted immediately.
DB engine version	<p>The version of the MySQL database engine that you want to use. Before you upgrade your production DB instances, we recommend that you test the upgrade process on a test instance to verify its duration and to validate your applications.</p> <p>For more information, see Upgrading the MySQL DB Engine (p. 812).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change.
DB instance class	<p>The DB instance class that you want to use.</p> <p>For more information, see Choosing the DB Instance Class (p. 107).</p> <p>MySQL DB instances created after April 23, 2014, can change to a different DB instance class by modifying the DB instance just as with any other modification. MySQL DB instances running MySQL version 5.5 and created before April 23, 2014, must first upgrade to MySQL version 5.6. For more information, see Upgrading the MySQL DB Engine (p. 812).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change.

Setting	Setting Description	When the Change Occurs	Downtime Notes
DB instance identifier	<p>The DB instance identifier. This value is stored as a lowercase string.</p> <p>For more information about the effects of renaming a DB instance, see Renaming a DB Instance (p. 163).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change. The DB instance is rebooted.
DB parameter group	<p>The parameter group that you want associated with the DB instance.</p> <p>For more information, see Working with DB Parameter Groups (p. 211).</p>	The parameter group change occurs immediately.	<p>An outage doesn't occur during this change. When you change the parameter group, changes to some parameters are applied to the DB instance immediately without a reboot. Changes to other parameters are applied only after the DB instance is rebooted.</p> <p>For more information, see Rebooting a DB Instance (p. 166).</p>
Deletion protection	<p>Enable deletion protection to prevent your DB instance from being deleted.</p> <p>For more information, see Deleting a DB Instance (p. 174).</p>	–	–
Enhanced Monitoring	<p>Enable Enhanced Monitoring to enable gathering metrics in real time for the operating system that your DB instance runs on.</p> <p>For more information, see Enhanced Monitoring (p. 307).</p>	–	–
Log exports	<p>The types of MySQL database log files to publish to Amazon CloudWatch Logs.</p> <p>For more information, see MySQL Database Log Files (p. 405).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Maintenance window	<p>The time range during which system maintenance occurs. System maintenance includes upgrades, if applicable. The maintenance window is a start time in Universal Coordinated Time (UTC), and a duration in hours.</p> <p>If you set the window to the current time, there must be at least 30 minutes between the current time and end of the window to ensure any pending changes are applied.</p> <p>For more information, see The Amazon RDS Maintenance Window (p. 156).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	If there are one or more pending actions that cause an outage, and the maintenance window is changed to include the current time, then those pending actions are applied immediately, and an outage occurs.
Multi-AZ deployment	<p>Yes to deploy your DB instance in multiple Availability Zones. Otherwise, No.</p> <p>For more information, see Regions and Availability Zones (p. 132).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	–
New master password	The password for your master user. The password must contain from 8 to 41 alphanumeric characters.	The change is applied asynchronously, as soon as possible. This setting ignores the Apply immediately setting.	–
Option group	<p>The option group that you want associated with the DB instance.</p> <p>For more information, see Working with Option Groups (p. 196).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Performance Insights	<p>Enable Performance Insights to monitor your DB instance load so that you can analyze and troubleshoot your database performance.</p> <p>Choose a retention period to determine how much rolling data history to keep. The default of seven days is in the free tier. Long-term retention (two years) is priced per vCPU per month.</p> <p>You can't change the master key after the database is created.</p> <p>For more information, see Using Amazon RDS Performance Insights (p. 319).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–
Public accessibility	<p>Yes to give the DB instance a public IP address, meaning that it's accessible outside the VPC. To be publicly accessible, the DB instance also has to be in a public subnet in the VPC.</p> <p>No to make the DB instance accessible only from inside the VPC.</p> <p>For more information, see Hiding a DB Instance in a VPC from the Internet (p. 515).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–
Security group	<p>The security group that you want associated with the DB instance.</p> <p>For more information, see Working with DB Security Groups (EC2-Classic Platform) (p. 491).</p>	The change is applied asynchronously, as soon as possible. This setting ignores the Apply immediately setting.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Storage autoscaling	<p>Enable storage autoscaling to enable Amazon RDS to automatically increase storage when needed to avoid having your DB instance run out of storage space.</p> <p>Use Maximum storage threshold to set the upper limit for Amazon RDS to automatically increase storage for your DB instance. The default is 1,000 GiB.</p> <p>For more information, see Managing Capacity Automatically with Amazon RDS Storage Autoscaling (p. 232).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Storage type	<p>The storage type that you want to use.</p> <p>For more information, see Amazon RDS Storage Types (p. 135).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	<p>The following changes all result in a brief outage while the process starts. After that, you can use your database normally while the change takes place.</p> <ul style="list-style-type: none"> • From General Purpose (SSD) to Magnetic. • From General Purpose (SSD) to Provisioned IOPS (SSD). The outage only happens if the DB instance is single-AZ and you are using a custom parameter group. There is no outage for a multi-AZ DB instance. • From Magnetic to General Purpose (SSD). • From Magnetic to Provisioned IOPS (SSD). • From Provisioned IOPS (SSD) to Magnetic. • From Provisioned IOPS (SSD) to General Purpose (SSD). The outage only happens if the DB instance is single-AZ and you are using a custom parameter group. There is no outage for a multi-AZ DB instance.

Setting	Setting Description	When the Change Occurs	Downtime Notes
Subnet group	<p>The subnet group for the DB instance. You can use this setting to move your DB instance to a different VPC. If your DB instance isn't in a VPC, you can use this setting to move your DB instance into a VPC.</p> <p>For more information, see Moving a DB Instance Not in a VPC into a VPC (p. 520).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change. The DB instance is rebooted.

Upgrading the MySQL DB Engine

When Amazon RDS supports a new version of a database engine, you can upgrade your DB instances to the new version. There are two kinds of upgrades: major version upgrades and minor version upgrades. In general, a major engine version upgrade can introduce changes that are not compatible with existing applications. In contrast, a minor version upgrade includes only changes that are backward-compatible with existing applications.

To perform a major version upgrade, modify the DB instance manually. Minor version upgrades occur automatically if you enable auto minor version upgrades on your DB instance. In all other cases, modify the DB instance manually to perform a minor version upgrade.

Topics

- [Overview of Upgrading \(p. 812\)](#)
- [Major Version Upgrades for MySQL \(p. 812\)](#)
- [Testing an Upgrade \(p. 815\)](#)
- [Upgrading a MySQL DB Instance \(p. 816\)](#)
- [Upgrading a MySQL Database with Reduced Downtime \(p. 816\)](#)

Overview of Upgrading

Amazon RDS takes two DB snapshots during the upgrade process. The first DB snapshot is of the DB instance before any upgrade changes have been made. If the upgrade doesn't work for your databases, you can restore this snapshot to create a DB instance running the old version. The second DB snapshot is taken when the upgrade completes.

Note

Amazon RDS only takes DB snapshots if you have set the backup retention period for your DB instance to a number greater than 0. To change your backup retention period, see [Modifying a DB Instance Running the MySQL Database Engine \(p. 802\)](#).

After the upgrade is complete, you can't revert to the previous version of the database engine. If you want to return to the previous version, restore the first DB snapshot taken to create a new DB instance.

You control when to upgrade your DB instance to a new version supported by Amazon RDS. This level of control helps you maintain compatibility with specific database versions and test new versions with your application before deploying in production. When you are ready, you can perform version upgrades at the times that best fit your schedule.

If your DB instance is using read replication, upgrade all of the Read Replicas before upgrading the source instance.

If your DB instance is in a Multi-AZ deployment, both the primary and standby DB instances are upgraded. The primary and standby DB instances are upgraded at the same time and you experience an outage until the upgrade is complete. The time for the outage varies based on the size of your DB instance.

Major Version Upgrades for MySQL

Amazon RDS supports the following in-place upgrades for major versions of the MySQL database engine:

- MySQL 5.5 to MySQL 5.6
- MySQL 5.6 to MySQL 5.7
- MySQL 5.7 to MySQL 8.0

Note

You can only create MySQL version 5.7 and 8.0 DB instances with latest-generation and current-generation DB instance classes, in addition to the db.m3 previous-generation DB instance class. In some cases, you want to upgrade a MySQL version 5.6 DB instance running on a previous-generation DB instance class (other than db.m3) to a MySQL version 5.7 DB instance. In these cases, first modify the DB instance to use a latest-generation or current-generation DB instance class. After you do this, you can then modify the DB instance to use the MySQL version 5.7 database engine. For information on Amazon RDS DB instance classes, see [Choosing the DB Instance Class \(p. 107\)](#).

Topics

- [Overview of MySQL Major Version Upgrades \(p. 813\)](#)
- [Upgrades to MySQL Version 5.7 Might Be Slow \(p. 813\)](#)
- [Prechecks for Upgrades from MySQL 5.7 to 8.0 \(p. 814\)](#)

Overview of MySQL Major Version Upgrades

Major version upgrades can contain database changes that are not backward-compatible with existing applications. As a result, Amazon RDS doesn't apply major version upgrades automatically; you must manually modify your DB instance. We recommend that you thoroughly test any upgrade before applying it to your production instances.

To perform a major version upgrade for a MySQL version 5.5 DB instance on Amazon RDS to MySQL version 5.6 or later, first perform any available OS updates. After OS updates are complete, upgrade to each major version: 5.5 to 5.6, then 5.6 to 5.7, and then 5.7 to 8.0. MySQL DB instances created before April 24, 2014, show an available OS update until the update has been applied. For more information on OS updates, see [Applying Updates for a DB Instance \(p. 154\)](#).

During a major version upgrade of MySQL, Amazon RDS runs the MySQL binary `mysql_upgrade` to upgrade tables, if necessary. Also, Amazon RDS empties the `slow_log` and `general_log` tables during a major version upgrade. To preserve log information, save the log contents before the major version upgrade.

MySQL major version upgrades typically complete in about 10 minutes. Some upgrades might take longer because of the DB instance class size or because the instance doesn't follow certain operational guidelines in [Best Practices for Amazon RDS \(p. 94\)](#). If you upgrade a DB instance from the Amazon RDS console, the status of the DB instance indicates when the upgrade is complete. If you upgrade using the AWS Command Line Interface (AWS CLI), use the `describe-db-instances` command and check the `Status` value.

If you use a custom parameter group, for the new DB engine version you either specify a default parameter group or create your own custom parameter group. Associating the new parameter group with the DB instance requires a customer-initiated database reboot after the upgrade completes. The DB instance's parameter group status shows `pending-reboot` if the DB instance needs to be rebooted to apply the parameter group changes. You can view a DB instance's parameter group status in the console or by using a call such as `describe-db-instances`.

Upgrades to MySQL Version 5.7 Might Be Slow

MySQL version 5.6.4 introduced a new date and time format for the `datetime`, `time`, and `timestamp` columns that allows fractional components in date and time values. When upgrading a DB instance to MySQL version 5.7, MySQL forces the conversion of all date and time column types to the new format.

Because this conversion rebuilds your tables, it might take a considerable amount of time to complete the DB instance upgrade. The forced conversion occurs for any DB instances that are running a version

before MySQL version 5.6.4. It also occurs for any DB instances that were upgraded from a version before MySQL version 5.6.4 to a version other than 5.7.

If your DB instance runs a version before MySQL version 5.6.4, or was upgraded from a version before 5.6.4, we recommend an extra step. In these cases, we recommend that you convert the `datetime`, `time`, and `timestamp` columns in your database before upgrading your DB instance to MySQL version 5.7. This conversion can significantly reduce the amount of time required to upgrade the DB instance to MySQL version 5.7. To upgrade your date and time columns to the new format, issue the `ALTER TABLE <table_name> FORCE;` command for each table that contains date or time columns. Because altering a table locks the table as read-only, we recommend that you perform this update during a maintenance window.

To find all tables in your database that have `datetime`, `time`, or `timestamp` columns and create an `ALTER TABLE <table_name> FORCE;` command for each table, use the following query.

```
SELECT DISTINCT CONCAT('ALTER TABLE `',
    REPLACE(is_tables.TABLE_SCHEMA, '`', '``'), `.``,
    REPLACE(is_tables.TABLE_NAME, '`', '``'), ` ` FORCE;`)
FROM information_schema.TABLES is_tables
INNER JOIN information_schema.COLUMNS col ON col.TABLE_SCHEMA =
is_tables.TABLE_SCHEMA
    AND col.TABLE_NAME = is_tables.TABLE_NAME
LEFT OUTER JOIN information_schema.INNODB_SYS_TABLES systables ON
    SUBSTRING_INDEX(systables.NAME, '#', 1) =
CONCAT(is_tables.TABLE_SCHEMA,'/',is_tables.TABLE_NAME)
    LEFT OUTER JOIN information_schema.INNODB_SYS_COLUMNS syscolumns ON
        syscolumns.TABLE_ID = systables.TABLE_ID AND syscolumns.NAME = col.COLUMN_NAME
WHERE col.COLUMN_TYPE IN ('time','timestamp','datetime')
    AND is_tables.TABLE_TYPE = 'BASE TABLE'
    AND is_tables.TABLE_SCHEMA NOT IN ('mysql','information_schema','performance_schema')
    AND (is_tables.ENGINE = 'InnoDB' AND syscolumns.MTYPE = 6);
```

Prechecks for Upgrades from MySQL 5.7 to 8.0

MySQL 8.0 includes a number of incompatibilities with MySQL 5.7. These incompatibilities can cause problems during an upgrade from MySQL 5.7 to MySQL 8.0. So, some preparation might be required on your database for the upgrade to be successful. The following is a general list of these incompatibilities:

- There must be no tables that use obsolete data types or functions.
- There must be no orphan *.frm files.
- Triggers must not have a missing or empty definer or an invalid creation context.
- There must be no partitioned table that uses a storage engine that does not have native partitioning support.
- There must be no keyword or reserved word violations. Some keywords might be reserved in MySQL 8.0 that were not reserved previously.

For more information, see [Keywords and Reserved Words](#) in the MySQL documentation.

- There must be no tables in the MySQL 5.7 mysql system database that have the same name as a table used by the MySQL 8.0 data dictionary.
- There must be no obsolete SQL modes defined in your `sql_mode` system variable setting.
- There must be no tables or stored procedures with individual `ENUM` or `SET` column elements that exceed 255 characters or 1020 bytes in length.
- Before upgrading to MySQL 8.0.13 or higher, there must be no table partitions that reside in shared InnoDB tablespaces.
- There must be no queries and stored program definitions from MySQL 8.0.12 or lower that use `ASC` or `DESC` qualifiers for `GROUP BY` clauses.

- Your MySQL 5.7 installation must not use features that are not supported in MySQL 8.0.
For more information, see [Features Removed in MySQL 8.0](#) in the MySQL documentation.
- There must be no foreign key constraint names longer than 64 characters.
- For improved Unicode support, consider converting objects that use the `utf8mb3` charset to use the `utf8mb4` charset. The `utf8mb3` character set is deprecated. Also, consider using `utf8mb4` for character set references instead of `utf8`, because currently `utf8` is an alias for the `utf8mb3` charset.
For more information, see [The utf8mb3 Character Set \(3-Byte UTF-8 Unicode Encoding\)](#) in the MySQL documentation.

When you start an upgrade from MySQL 5.7 to 8.0, Amazon RDS runs prechecks automatically to detect these incompatibilities. For information about upgrading to MySQL 8.0, see [Upgrading MySQL](#) in the MySQL documentation.

These prechecks are mandatory. You can't choose to skip them. The prechecks provide the following benefits:

- They enable you to avoid unplanned downtime during the upgrade.
- If there are incompatibilities, Amazon RDS prevents the upgrade and provides a log for you to learn about them. You can then use the log to prepare your database for the upgrade to MySQL 8.0 by eliminating the incompatibilities. For detailed information about removing incompatibilities, see [Preparing Your Installation for Upgrade](#) in the MySQL documentation and [Upgrading to MySQL 8.0? Here is what you need to know...](#) on the MySQL Server Blog.

The prechecks include some that are included with MySQL and some that were created specifically by the Amazon RDS team. For information about the prechecks provided by MySQL, see [Upgrade Checker Utility](#).

The prechecks run before the DB instance is stopped for the upgrade, meaning that they don't cause any downtime when they run. If the prechecks find an incompatibility, Amazon RDS automatically cancels the upgrade before the DB instance is stopped. Amazon RDS also generates an event for the incompatibility. For more information about Amazon RDS events, see [Using Amazon RDS Event Notification \(p. 370\)](#).

Amazon RDS records detailed information about each incompatibility in the log file `PrePatchCompatibility.log`. In most cases, the log entry includes a link to the MySQL documentation for correcting the incompatibility. For more information about viewing log files, see [Viewing and Listing Database Log Files \(p. 390\)](#).

Due to the nature of the prechecks, they analyze the objects in your database. This analysis results in resource consumption and increases the time for the upgrade to complete.

Note

Amazon RDS runs prechecks only for an upgrade from MySQL 5.7 to MySQL 8.0. They aren't run for upgrades to releases lower than MySQL 8.0. For example, prechecks aren't run for an upgrade from MySQL 5.6 to MySQL 5.7.

Testing an Upgrade

Before you perform a major version upgrade on your DB instance, thoroughly test your database for compatibility with the new version. In addition, thoroughly test all applications that access the database for compatibility with the new version. We recommend that you use the following procedure.

To test a major version upgrade

1. Review the upgrade documentation for the new version of the database engine to see if there are compatibility issues that might affect your database or applications:

- [Changes in MySQL 5.6](#)
 - [Changes in MySQL 5.7](#)
 - [Changes in MySQL 8.0](#)
2. If your DB instance is a member of a custom DB parameter group, create a new DB parameter group with your existing settings that is compatible with the new major version. Specify the new DB parameter group when you upgrade your test instance, so your upgrade testing ensures that it works correctly. For more information about creating a DB parameter group, see [Working with DB Parameter Groups \(p. 211\)](#).
 3. Create a DB snapshot of the DB instance to be upgraded. For more information, see [Creating a DB Snapshot \(p. 264\)](#).
 4. Restore the DB snapshot to create a new test DB instance. For more information, see [Restoring from a DB Snapshot \(p. 266\)](#).
 5. Modify this new test DB instance to upgrade it to the new version, using one of the methods detailed following. If you created a new parameter group in step 2, specify that parameter group.
 6. Evaluate the storage used by the upgraded instance to determine if the upgrade requires additional storage.
 7. Run as many of your quality assurance tests against the upgraded DB instance as needed to ensure that your database and application work correctly with the new version. Implement any new tests needed to evaluate the impact of any compatibility issues that you identified in step 1. Test all stored procedures and functions. Direct test versions of your applications to the upgraded DB instance.
 8. If all tests pass, then perform the upgrade on your production DB instance. We recommend that you don't allow write operations to the DB instance until you confirm that everything is working correctly.

Upgrading a MySQL DB Instance

For information about manually or automatically upgrading a MySQL DB instance, see [Upgrading a DB Instance Engine Version \(p. 159\)](#).

Upgrading a MySQL Database with Reduced Downtime

If your MySQL DB instance is currently in use with a production application, you can use the following procedure to upgrade the database version for your DB instance. This procedure can reduce the amount of downtime for your application.

The following procedure shows an example of upgrading from MySQL version 5.5 to MySQL version 5.6. You can use the same general steps for upgrades to other major versions.

To upgrade an MySQL database while a DB instance is in use

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. Create a Read Replica of your MySQL 5.5 DB instance. This process creates an upgradable copy of your database.
 - a. On the console, choose **Databases**, and then choose the DB instance that you want to upgrade.
 - b. For **Actions**, choose **Create read replica**.
 - c. Provide a value for **DB instance identifier** for your Read Replica and ensure that the **DB instance class** and other settings match your MySQL 5.5 DB instance.

- d. Choose **Create read replica**.
3. When the Read Replica has been created and **Status** shows **available**, upgrade the Read Replica to MySQL 5.6:
 - a. On the console, choose **Databases**, and then choose the Read Replica that you just created.
 - b. Choose **Modify**.
 - c. For **DB engine version**, choose the MySQL 5.6 version to upgrade to, and then choose **Continue**.
 - d. For **Scheduling of Modifications**, choose **Apply immediately**.
 - e. Choose **Modify DB instance** to start the upgrade.
4. When the upgrade is complete and **Status** shows **available**, verify that the upgraded Read Replica is up-to-date with the master MySQL 5.5 DB instance. You can do this by connecting to the Read Replica and issuing the `SHOW SLAVE STATUS` command. If the `Seconds_Behind_Master` field is 0, then replication is up-to-date.
5. Make your MySQL 5.6 Read Replica a master DB instance.

Important

When you promote your MySQL 5.6 Read Replica to a standalone, single-AZ DB instance, it no longer is a replication replica to your MySQL 5.5 DB instance. We recommend that you promote your MySQL 5.6 Read Replica during a maintenance window when your source MySQL 5.5 DB instance is in read-only mode and all write operations are suspended. When the promotion is completed, you can direct your write operations to the upgraded MySQL 5.6 DB instance to ensure that no write operations are lost.

In addition, we recommend that before promoting your MySQL 5.6 Read Replica you perform all necessary data definition language (DDL) operations on your MySQL 5.6 Read Replica. An example is creating indexes. This approach avoids negative effects on the performance of the MySQL 5.6 Read Replica after it has been promoted. To promote a Read Replica, use the following procedure.

- a. On the console, choose **Databases**, and then choose the Read Replica that you just upgraded.
 - b. For **Actions**, choose **Promote read replica**.
 - c. Choose **Yes** to enable automated backups for the Read Replica instance. For more information, see [Working With Backups \(p. 254\)](#).
- Choose **Continue**.
- d. Choose **Promote Read Replica**.
6. You now have an upgraded version of your MySQL database. At this point, you can direct your applications to the new MySQL 5.6 DB instance, add Read Replicas, set up Multi-AZ support, and so on.

Upgrading a MySQL DB Snapshot

With Amazon RDS, you can create a storage volume DB snapshot of your MySQL DB instance. When you create a DB snapshot, the snapshot is based on the engine version used by your Amazon RDS instance. In addition to upgrading the DB engine version of your DB instance, you can also upgrade the engine version for your DB snapshots. For example, you can upgrade DB snapshots created from the MySQL 5.1 engine to DB snapshots for the MySQL 5.5 engine. After restoring a DB snapshot upgraded to a new engine version, you should test that the upgrade was successful. To learn how to test a major version upgrade, see [Testing an Upgrade \(p. 815\)](#). To learn how to restore a DB snapshot, see [Restoring from a DB Snapshot \(p. 266\)](#).

Amazon RDS supports upgrading a MySQL DB snapshot from MySQL 5.1 to MySQL 5.5.

Upgrading a MySQL DB Snapshot

You can upgrade manual DB snapshots, which can be encrypted or not encrypted, from MySQL 5.1 to MySQL 5.5 within the same AWS Region. You can't upgrade automated DB snapshots that are created during the automated backup process.

Console

To upgrade a DB snapshot

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Snapshots**.
3. For **Actions**, choose **Modify Snapshot**. The **Modify DB Snapshot** page appears.
4. Choose **Modify Snapshot** to upgrade the snapshot. During the upgrade process, all snapshot actions are disabled. Also, the DB snapshot status changes from **available** to **upgrading**, and then changes to **active** upon completion. If the DB snapshot can't be upgraded because of snapshot corruption issues, the status changes to **unavailable**. You can't recover the snapshot from this state.

AWS CLI

To upgrade a DB snapshot to a new database engine version, use the AWS CLI `modify-db-snapshot` command.

Parameters

- `--db-snapshot-identifier` – The identifier of the DB snapshot to upgrade. The identifier must be a unique Amazon Resource Name (ARN). For more information, see [Working with Amazon Resource Names \(ARNs\) in Amazon RDS \(p. 224\)](#).
- `--engine-version` – The engine version to upgrade the DB snapshot to.

Example

For Linux, OS X, or Unix:

```
aws rds modify-db-snapshot \
--db-snapshot-identifier <mydbsnapshot> \
--engine-version <new_version>
```

For Windows:

```
aws rds modify-db-snapshot ^
--db-snapshot-identifier <mydbsnapshot> ^
--engine-version <new_version>
```

RDS API

To upgrade a DB snapshot to a new database engine version, call the Amazon RDS API [ModifyDBSnapshot](#) operation.

- **DBSnapshotIdentifier** – The identifier of the DB snapshot to upgrade. The identifier must be a unique Amazon Resource Name (ARN). For more information, see [Working with Amazon Resource Names \(ARNs\) in Amazon RDS \(p. 224\)](#).
- **EngineVersion** – The engine version to upgrade the DB snapshot to.

Example

```
https://rds.us-west-2.amazonaws.com/
?Action=ModifyDBSnapshot
&DBSnapshotIdentifier=mydbsnapshot
&EngineVersion=newversion
&SignatureVersion=4
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20161222/us-west-1/rds/aws4_request
&X-Amz-Date=20161222T233051Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=8052a76dfb18469393c5f0182cdab0ebc224a9c7c5c949155376c1c250fc7ec3
```

Importing Data into a MySQL DB Instance

You can use several different techniques to import data into an Amazon RDS for MySQL DB instance. The best approach depends on the source of the data, the amount of data, and whether the import is done one time or is ongoing. If you are migrating an application along with the data, also consider the amount of downtime that you are willing to experience.

Overview

Find techniques to import data into an Amazon RDS for MySQL DB instance in the following table.

Source	Amount of Data	One Time or Ongoing	Application Downtime	Technique	More Information
Existing MySQL database on premises or on Amazon EC2	Any	One time	Some	Create a backup of your on-premises database, store it on Amazon S3, and then restore the backup file to a new Amazon RDS DB instance running MySQL.	Restoring a Backup into an Amazon RDS MySQL DB Instance (p. 826)
Any existing database	Any	One time or ongoing	Minimal	Use AWS Database Migration Service to migrate the database with minimal downtime and, for many database DB engines, continue ongoing replication.	What is AWS Database Migration Service in the AWS Database Migration Service User Guide
Existing Amazon RDS MySQL DB instance	Any	One time	Minimal	Create a Read Replica, and then promote the Read Replica.	Working with Read Replicas (p. 182)
Existing MySQL or MariaDB database	Small	One time	Some	Copy the data directly to your Amazon RDS MySQL DB instance using a command-line utility.	Importing Data from a MySQL or MariaDB DB to an Amazon RDS MySQL or MariaDB

Source	Amount of Data	One Time or Ongoing	Application Downtime	Technique	More Information
					DB Instance (p. 833)
Data not stored in an existing database	Medium	One time	Some	Create flat files and import them using the <code>mysqlimport</code> utility.	Importing Data From Any Source to a MySQL or MariaDB DB Instance (p. 847)
Existing MySQL or MariaDB database on premises or on Amazon EC2	Any	Ongoing	Minimal	Configure replication with an existing MySQL database as the replication source.	Replication with a MySQL or MariaDB Instance Running External to Amazon RDS (p. 862) or Importing Data to an Amazon RDS MySQL or MariaDB DB Instance with Reduced Downtime (p. 835)

Note

The 'mysql' system database contains authentication and authorization information required to log in to your DB instance and access your data. Dropping, altering, renaming, or truncating tables, data, or other contents of the 'mysql' database in your DB instance can result in error and might render the DB instance and your data inaccessible. If this occurs, you can restore the DB instance from a snapshot using the AWS CLI `restore-db-instance-from-db-snapshot` command. You can recover the DB instance using the AWS CLI `restore-db-instance-to-point-in-time` command.

Importing Data Considerations

Following, you can find additional technical information related to loading data into MySQL. This information is intended for advanced users who are familiar with the MySQL server architecture. All comments related to LOAD DATA LOCAL INFILE also apply to `mysqlimport`.

Binary Log

Data loads incur a performance penalty and require additional free disk space (up to four times more) when binary logging is enabled versus loading the same data with binary logging turned off. The severity of the performance penalty and the amount of free disk space required is directly proportional to the size of the transactions used to load the data.

Transaction Size

Transaction size plays an important role in MySQL data loads. It has a major influence on resource consumption, disk space utilization, resume process, time to recover, and input format (flat files or SQL). This section describes how transaction size affects binary logging and makes the case for disabling binary logging during large data loads. As noted earlier, binary logging is enabled and disabled by setting the Amazon RDS automated backup retention period. Non-zero values enable binary logging, and zero disables it. We also describe the impact of large transactions on InnoDB and why it's important to keep transaction sizes small.

Small Transactions

For small transactions, binary logging doubles the number of disk writes required to load the data. This effect can severely degrade performance for other database sessions and increase the time required to load the data. The degradation experienced depends in part upon the upload rate, other database activity taking place during the load, and the capacity of your Amazon RDS DB instance.

The binary logs also consume disk space roughly equal to the amount of data loaded until they are backed up and removed. Fortunately, Amazon RDS minimizes this by backing up and removing binary logs on a frequent basis.

Large Transactions

Large transactions incur a 3X penalty for IOPS and disk consumption with binary logging enabled. This is due to the binary log cache spilling to disk, consuming disk space and incurring additional IO for each write. The cache cannot be written to the binlog until the transaction commits or rolls back, so it consumes disk space in proportion to the amount of data loaded. When the transaction commits, the cache must be copied to the binlog, creating a third copy of the data on disk.

Because of this, there must be at least three times as much free disk space available to load the data compared to loading with binary logging disabled. For example, 10 GiB of data loaded as a single transaction consumes at least 30 GiB disk space during the load. It consumes 10 GiB for the table + 10 GiB for the binary log cache + 10 GiB for the binary log itself. The cache file remains on disk until the session that created it terminates or the session fills its binary log cache again during another transaction. The binary log must remain on disk until backed up, so it might be some time before the extra 20 GiB is freed.

If the data was loaded using LOAD DATA LOCAL INFILE, yet another copy of the data is created if the database has to be recovered from a backup made before the load. During recovery, MySQL extracts the data from the binary log into a flat file. MySQL then executes LOAD DATA LOCAL INFILE, just as in the original transaction. However, this time the input file is local to the database server. Continuing with the example preceding, recovery fails unless there is at least 40 GiB free disk space available.

Disable Binary Logging

Whenever possible, disable binary logging during large data loads to avoid the resource overhead and addition disk space requirements. In Amazon RDS, disabling binary logging is as simple as setting the backup retention period to zero. If you do this, we recommend that you take a DB snapshot of the database instance immediately before the load. By doing this, you can quickly and easily undo changes made during loading if you need to.

After the load, set the backup retention period back to an appropriate (no zero) value.

You can't set the backup retention period to zero if the DB instance is a source DB instance for Read Replicas.

InnoDB

The information in this section provides a strong argument for keeping transaction sizes small when using InnoDB.

Undo

InnoDB generates undo to support features such as transaction rollback and MVCC. Undo is stored in the InnoDB system tablespace (usually `ibdata1`) and is retained until removed by the purge thread. The purge thread cannot advance beyond the undo of the oldest active transaction, so it is effectively blocked until the transaction commits or completes a rollback. If the database is processing other transactions during the load, their undo also accumulates in the system tablespace and cannot be removed even if they commit and no other transaction needs the undo for MVCC. In this situation, all transactions (including read-only transactions) that access any of the rows changed by any transaction (not just the load transaction) slow down. The slowdown occurs because transactions scan through undo that could have been purged if not for the long-running load transaction.

Undo is stored in the system tablespace, and the system tablespace never shrinks in size. Thus, large data load transactions can cause the system tablespace to become quite large, consuming disk space that you can't reclaim without recreating the database from scratch.

Rollback

InnoDB is optimized for commits. Rolling back a large transaction can take a very, very long time. In some cases, it might be faster to perform a point-in-time recovery or restore a DB snapshot.

Input Data Format

MySQL can accept incoming data in one of two forms: flat files and SQL. This section points out some key advantages and disadvantages of each.

Flat Files

Loading flat files with `LOAD DATA LOCAL INFILE` can be the fastest and least costly method of loading data as long as transactions are kept relatively small. Compared to loading the same data with SQL, flat files usually require less network traffic, lowering transmission costs and load much faster due to the reduced overhead in the database.

One Big Transaction

`LOAD DATA LOCAL INFILE` loads the entire flat file as one transaction. This isn't necessarily a bad thing. If the size of the individual files can be kept small, this has a number of advantages:

- Resume capability – Keeping track of which files have been loaded is easy. If a problem arises during the load, you can pick up where you left off with little effort. Some data might have to be retransmitted to Amazon RDS, but with small files, the amount retransmitted is minimal.

- Load data in parallel – If you've got IOPS and network bandwidth to spare with a single file load, loading in parallel might save time.
- Throttle the load rate – Data load having a negative impact on other processes? Throttle the load by increasing the interval between files.

Be Careful

The advantages of LOAD DATA LOCAL INFILE diminish rapidly as transaction size increases. If breaking up a large set of data into smaller ones isn't an option, SQL might be the better choice.

SQL

SQL has one main advantage over flat files: it's easy to keep transaction sizes small. However, SQL can take significantly longer to load than flat files and it can be difficult to determine where to resume the load after a failure. For example, mysqldump files are not restartable. If a failure occurs while loading a mysqldump file, the file requires modification or replacement before the load can resume. The alternative is to restore to the point in time before the load and replay the file after the cause of the failure has been corrected.

Take Checkpoints Using Amazon RDS Snapshots

If you have a load that's going to take several hours or even days, loading without binary logging isn't a very attractive prospect unless you can take periodic checkpoints. This is where the Amazon RDS DB snapshot feature comes in very handy. A DB snapshot creates a point-in-time consistent copy of your database instance which can be used to restore the database to that point in time after a crash or other mishap.

To create a checkpoint, simply take a DB snapshot. Any previous DB snapshots taken for checkpoints can be removed without affecting durability or restore time.

Snapshots are fast too, so frequent checkpointing doesn't add significantly to load time.

Decreasing Load Time

Here are some additional tips to reduce load times:

- Create all secondary indexes before loading. This is counter-intuitive for those familiar with other databases. Adding or modifying a secondary index causes MySQL to create a new table with the index changes, copy the data from the existing table to the new table, and drop the original table.
- Load data in PK order. This is particularly helpful for InnoDB tables, where load times can be reduced by 75–80 percent and data file size cut in half.
- Disable foreign key constraints `foreign_key_checks=0`. For flat files loaded with LOAD DATA LOCAL INFILE, this is required in many cases. For any load, disabling FK checks provides significant performance gains. Just be sure to enable the constraints and verify the data after the load.
- Load in parallel unless already near a resource limit. Use partitioned tables when appropriate.
- Use multi-value inserts when loading with SQL to minimize statement execution overhead. When using mysqldump, this is done automatically.
- Reduce InnoDB log IO `innodb_flush_log_at_trx_commit=0`
- If you are loading data into a DB instance that does not have Read Replicas, set the `sync_binlog` parameter to 0 while loading data. When data loading is complete, set the `sync_binlog` parameter back to 1.
- Load data before converting the DB instance to a Multi-AZ deployment. However, if the DB instance already uses a Multi-AZ deployment, switching to a Single-AZ deployment for data loading is not recommended, because doing so only provides marginal improvements.

Note

Using `innodb_flush_log_at_trx_commit=0` causes InnoDB to flush its logs every second instead of at each commit. This provides a significant speed advantage, but can lead to data loss during a crash. Use with caution.

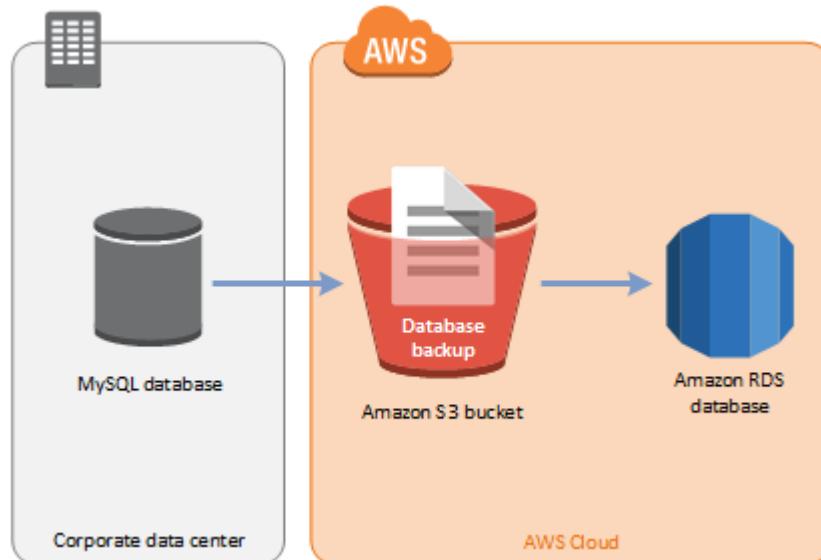
Topics

- [Restoring a Backup into an Amazon RDS MySQL DB Instance \(p. 826\)](#)
- [Importing Data from a MySQL or MariaDB DB to an Amazon RDS MySQL or MariaDB DB Instance \(p. 833\)](#)
- [Importing Data to an Amazon RDS MySQL or MariaDB DB Instance with Reduced Downtime \(p. 835\)](#)
- [Importing Data From Any Source to a MySQL or MariaDB DB Instance \(p. 847\)](#)

Restoring a Backup into an Amazon RDS MySQL DB Instance

Amazon RDS supports importing MySQL databases by using backup files. You can create a backup of your on-premises database, store it on Amazon S3, and then restore the backup file onto a new Amazon RDS DB instance running MySQL.

You can find the supported scenario in the following diagram.



Importing backup files from Amazon S3 is supported for MySQL version 5.6 and 5.7. Importing backup files from Amazon S3 is available in all AWS Regions.

We recommend that you import your database to Amazon RDS by using backup files if your database can be offline while the backup file is created, copied, and restored. If your on-premises database can't be offline, you can use binlog replication to update your database after you have migrated to Amazon RDS through Amazon S3 as explained in this topic. For more information, see [Replication with a MySQL or MariaDB Instance Running External to Amazon RDS \(p. 862\)](#). You can also use the AWS Database Migration Service to migrate your database to Amazon RDS. For more information, see [What Is AWS Database Migration Service?](#)

Limitations and Recommendations for Importing Backup Files from Amazon S3 to Amazon RDS

The following are some limitations and recommendations for importing backup files from Amazon S3:

- You can only import your data to a new DB instance, not an existing DB instance.
- You must use Percona XtraBackup to create the backup of your on-premises database.
- You can't migrate from a source database that has tables defined outside of the default MySQL data directory.
- You can't import a MySQL 5.5 or 8.0 database.
- You can't import an on-premises MySQL 5.6 database to an Amazon RDS MySQL 5.7 or 8.0 database. You can upgrade your DB instance after you complete the import.
- You can't restore databases larger than the maximum database size supported by Amazon RDS for MySQL. For more information about storage limits, see [General Purpose SSD Storage \(p. 135\)](#) and [Provisioned IOPS SSD Storage \(p. 137\)](#).

- You can't restore from an encrypted source database, but you can restore to an encrypted Amazon RDS DB instance.
- You can't restore from an encrypted backup in the Amazon S3 bucket.
- You can't restore from an Amazon S3 bucket in a different AWS Region than your Amazon RDS DB instance.
- Importing from Amazon S3 is not supported on the db.t2.micro DB instance class. However, you can restore to a different DB instance class, and then change the instance class later. For more information about instance classes, see [Hardware Specifications for All Available DB Instance Classes \(p. 109\)](#).
- Amazon S3 limits the size of a file uploaded to an Amazon S3 bucket to 5 TB. If a backup file exceeds 5 TB, then you must split the backup file into smaller files.
- Amazon RDS limits the number of files uploaded to an Amazon S3 bucket to 1 million. If the backup data for your database, including all full and incremental backups, exceeds 1 million files, use a tarball (.tar.gz) file to store full and incremental backup files in the Amazon S3 bucket.
- User accounts are not imported automatically. Save your user accounts from your source database and add them to your new DB instance later.
- Functions are not imported automatically. Save your functions from your source database and add them to your new DB instance later.
- Stored procedures are not imported automatically. Save your stored procedures from your source database and add them to your new DB instance later.
- Time zone information is not imported automatically. Record the time zone information for your source database, and set the time zone of your new DB instance later. For more information, see [Local Time Zone for MySQL DB Instances \(p. 749\)](#).
- Backward migration is not supported for both major versions and minor versions. For example, you can't migrate from version 5.7 to version 5.6, and you can't migrate from version 5.6.39 to version 5.6.37.

Overview of Setting Up to Import Backup Files from Amazon S3 to Amazon RDS

These are the components you need to set up to import backup files from Amazon S3 to Amazon RDS:

- An Amazon S3 bucket to store your backup files.
- A backup of your on-premises database created by Percona XtraBackup.
- An AWS Identity and Access Management (IAM) role to allow Amazon RDS to access the bucket.

If you already have an Amazon S3 bucket, you can use that. If you don't have an Amazon S3 bucket, you can create a new one. If you want to create a new bucket, see [Creating a Bucket](#).

Use the Percona XtraBackup tool to create your backup. For more information, see [Creating Your Database Backup \(p. 827\)](#).

If you already have an IAM role, you can use that. If you don't have an IAM role, you can create a new one manually. Alternatively, you can choose to have a new IAM role created for you in your account by the wizard when you restore the database by using the AWS Management Console. If you want to create a new IAM role manually, or attach trust and permissions policies to an existing IAM role, see [Creating an IAM Role Manually \(p. 829\)](#). If you want to have a new IAM role created for you, follow the procedure in [Console \(p. 830\)](#).

Creating Your Database Backup

Use the Percona XtraBackup software to create your backup. You can install Percona XtraBackup from [Download Percona XtraBackup](#).

Note

For MySQL 5.7 migration, you must use Percona XtraBackup 2.4. For earlier MySQL versions, use Percona XtraBackup 2.3 or 2.4.

You can create a full backup of your MySQL database files using Percona XtraBackup. Alternatively, if you already use Percona XtraBackup to back up your MySQL database files, you can upload your existing full and incremental backup directories and files.

For more information about backing up your database with Percona XtraBackup, see [Percona XtraBackup - Documentation](#) and [The xtrabackup Binary](#) on the Percona website.

Creating a Full Backup With Percona XtraBackup

To create a full backup of your MySQL database files that can be restored from Amazon S3, use the Percona XtraBackup utility (`xtrabackup`) to back up your database.

For example, the following command creates a backup of a MySQL database and stores the files in the folder `/on-premises/s3-restore/backup`.

```
xtrabackup --backup --user=<myuser> --password=<password> --target-dir=</on-premises/s3-restore/backup>
```

If you want to compress your backup into a single file (which can be split later, if needed), you can save your backup in one of the following formats:

- Gzip (.gz)
- tar (.tar)
- Percona xbstream (.xbstream)

The following command creates a backup of your MySQL database split into multiple Gzip files.

```
xtrabackup --backup --user=<myuser> --password=<password> --stream=tar \  
--target-dir=</on-premises/s3-restore/backup> | gzip - | split -d --bytes=500MB \  
- </on-premises/s3-restore/backup/backup>.tar.gz
```

The following command creates a backup of your MySQL database split into multiple tar files.

```
xtrabackup --backup --user=<myuser> --password=<password> --stream=tar \  
--target-dir=</on-premises/s3-restore/backup> | split -d --bytes=500MB \  
- </on-premises/s3-restore/backup/backup>.tar
```

The following command creates a backup of your MySQL database split into multiple xbstream files.

```
xtrabackup --backup --user=<myuser> --password=<password> --stream=xbstream \  
--target-dir=</on-premises/s3-restore/backup> | split -d --bytes=500MB \  
- </on-premises/s3-restore/backup/backup>.xbstream
```

Using Incremental Backups With Percona XtraBackup

If you already use Percona XtraBackup to perform full and incremental backups of your MySQL database files, you don't need to create a full backup and upload the backup files to Amazon S3. Instead, you can save a significant amount of time by copying your existing backup directories and files to your Amazon S3 bucket. For more information about creating incremental backups using Percona XtraBackup, see [Incremental Backup](#).

When copying your existing full and incremental backup files to an Amazon S3 bucket, you must recursively copy the contents of the base directory. Those contents include the full backup and also all

incremental backup directories and files. This copy must preserve the directory structure in the Amazon S3 bucket. Amazon RDS iterates through all files and directories. Amazon RDS uses the `xtrabackup-checkpoints` file that is included with each incremental backup to identify the base directory, and to order incremental backups by log sequence number (LSN) range.

Backup Considerations for Percona XtraBackup

Amazon RDS consumes your backup files based on the file name. Name your backup files with the appropriate file extension based on the file format—for example, `.xbstream` for files stored using the Percona xbstream format.

Amazon RDS consumes your backup files in alphabetical order and also in natural number order. Use the `split` option when you issue the `xtrabackup` command to ensure that your backup files are written and named in the proper order.

Amazon RDS doesn't support partial backups created using Percona XtraBackup. You can't use the following options to create a partial backup when you back up the source files for your database: `--tables`, `--tables-exclude`, `--tables-file`, `--databases`, `--databases-exclude`, or `--databases-file`.

Amazon RDS supports incremental backups created using Percona XtraBackup. For more information about creating incremental backups using Percona XtraBackup, see [Incremental Backup](#).

Creating an IAM Role Manually

If you don't have an IAM role, you can create a new one manually. Alternatively, you can choose to have a new IAM role created for you by the wizard when you restore the database by using the AWS Management Console. If you want to have a new IAM role created for you, follow the procedure in [Console \(p. 830\)](#).

To manually create a new IAM role for importing your database from Amazon S3, create a role to delegate permissions from Amazon RDS to your Amazon S3 bucket. When you create an IAM role, you attach trust and permissions policies. To import your backup files from Amazon S3, use trust and permissions policies similar to the examples following. For more information about creating the role, see [Creating a Role to Delegate Permissions to an AWS Service](#).

Alternatively, you can choose to have a new IAM role created for you by the wizard when you restore the database by using the AWS Management Console. If you want to have a new IAM role created for you, follow the procedure in [Console \(p. 830\)](#)

The trust and permissions policies require that you provide an Amazon Resource Name (ARN). For more information about ARN formatting, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

Example Trust Policy for Importing from Amazon S3

```
{  
    "Version": "2012-10-17",  
    "Statement":  
    [ {  
        "Effect": "Allow",  
        "Principal": {"Service": "rds.amazonaws.com"},  
        "Action": "sts:AssumeRole"  
    }]  
}
```

Example Permissions Policy for Importing from Amazon S3 — IAM User Permissions

```
{  
    "Version": "2012-10-17",
```

```
"Statement":  
[  
    {  
        "Sid": "AllowS3AccessRole",  
        "Effect": "Allow",  
        "Action": "iam:PassRole",  
        "Resource": "arn:aws:iam::IAM User ID:role/S3Access"  
    }  
]
```

Example Permissions Policy for Importing from Amazon S3 — Role Permissions

```
{  
    "Version": "2012-10-17",  
    "Statement":  
    [  
        {  
            "Effect": "Allow",  
            "Action":  
                [  
                    "s3>ListBucket",  
                    "s3:GetBucketLocation"  
                ],  
            "Resource": "arn:aws:s3:::bucket_name"  
        },  
        {  
            "Effect": "Allow",  
            "Action":  
                [  
                    "s3GetObject"  
                ],  
            "Resource": "arn:aws:s3:::bucket_name/prefix*"  
        }  
    ]  
}
```

Note

If you include a file name prefix, include the asterisk (*) after the prefix. If you don't want to specify a prefix, specify only an asterisk.

Importing Data From Amazon S3 to a New MySQL DB Instance

You can import data from Amazon S3 to a new MySQL DB instance using the AWS Management Console, AWS CLI, or RDS API.

Console

To import data from Amazon S3 to a new MySQL DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the top right corner of the Amazon RDS console, choose the AWS Region in which to create your DB instance. Choose the same AWS Region as the Amazon S3 bucket that contains your database backup.
3. In the navigation pane, choose **Databases**.
4. Choose **Restore from S3** to launch the wizard.

The wizard opens on the **Select engine** page.

5. On the **Select engine** page, choose **MySQL**, and then choose **Next**.

The **Specify source backup details** page appears.

Specify source backup details

Source database specifications

Source engine: mysql

Source engine version: 5.6

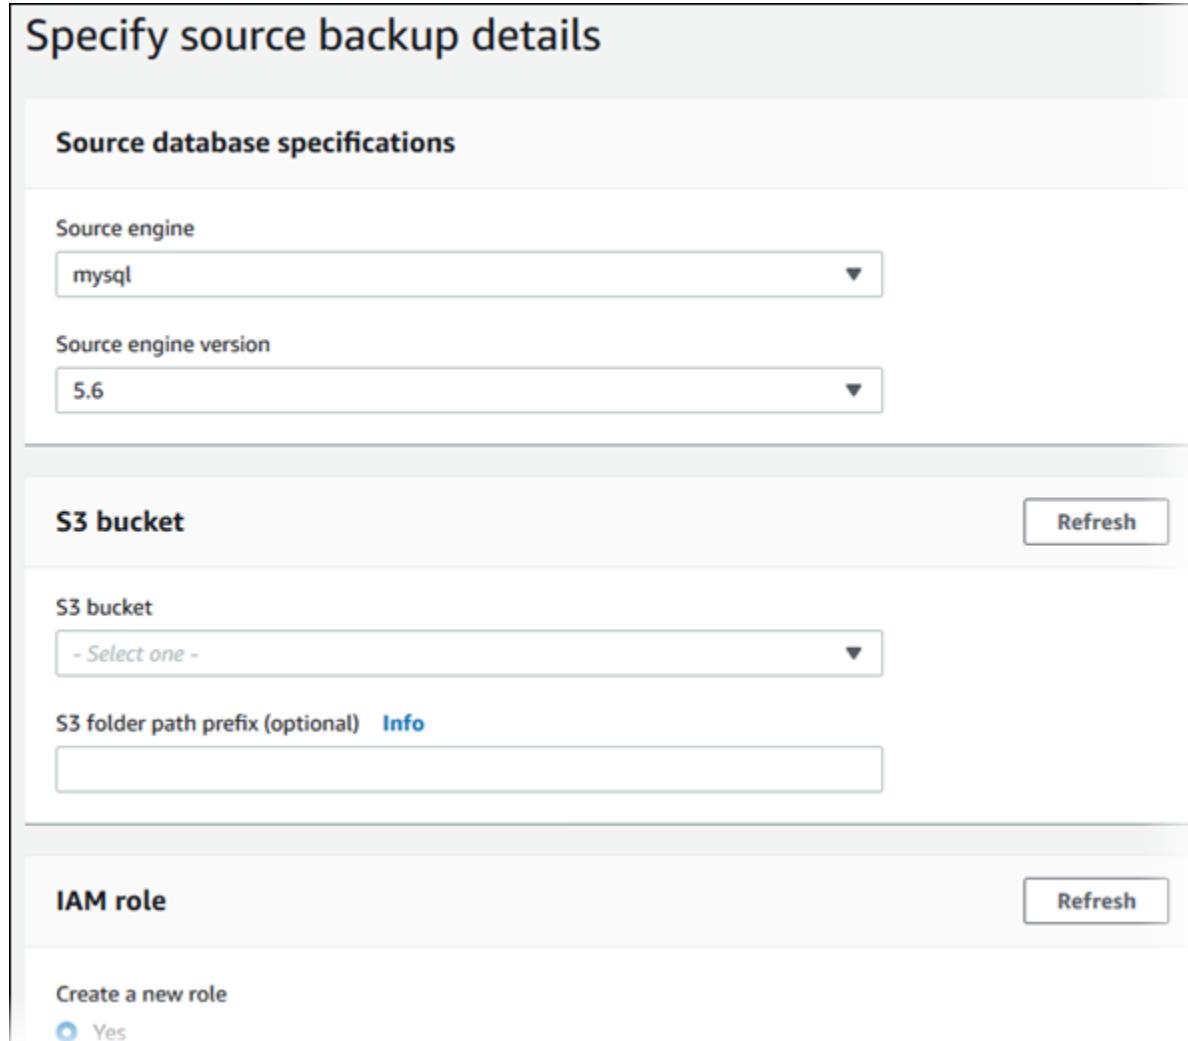
S3 bucket Refresh

S3 bucket: - Select one -

S3 folder path prefix (optional) [Info](#)

IAM role Refresh

Create a new role Yes



6. On the **Specify source backup details** page, specify your backup information.
 - a. For **Source engine**, choose **mysql**.
 - b. For **Source engine version**, choose the MySQL version of your source database.
 - c. For **S3 bucket**, choose your Amazon S3 bucket.
 - d. (Optional) For **S3 folder path prefix**, enter a file path prefix for the files stored in your Amazon S3 bucket. If you don't specify a prefix, then RDS creates your DB instance using all of the files and folders in the root folder of the S3 bucket. If you do specify a prefix, then RDS creates your DB instance using the files and folders in the S3 bucket where the path for the file begins with the specified prefix. For example, suppose that you store your backup files on S3 in a subfolder named backups, and you have multiple sets of backup files, each in its own directory (gzip_backup1, gzip_backup2, and so on). In this case, you specify a prefix of backups/gzip_backup1 to restore from the files in the gzip_backup1 folder.
 - e. For **Create a new role**, choose **Yes** to create a new IAM role in your account, or choose **No** to select an existing IAM role.
 - f. For **IAM role**, select an existing IAM role, or for **IAM role name**, specify the name for a new IAM role. You can choose to have a new IAM role created for you by choosing **Yes** for **Create a new role**.

7. Choose **Next** to continue. The **Specify DB details** page appears.

On the **Specify DB details** page, specify your DB instance information. For information about each setting, see [Settings for MySQL DB Instances \(p. 758\)](#).

Note

Be sure to allocate enough memory for your new DB instance so that the restore can succeed. You can also allocate additional memory for future growth.

8. Choose **Next** to continue. The **Configure advanced settings** page appears.

Provide additional information that Amazon RDS needs to launch the DB instance. For information about each setting, see [Settings for MySQL DB Instances \(p. 758\)](#).

9. Choose **Create database**.

AWS CLI

To import data from Amazon S3 to a new MySQL DB instance by using the AWS CLI, call the `restore-db-instance-from-s3` command with the parameters following. For information about each setting, see [Settings for MySQL DB Instances \(p. 758\)](#).

Note

Be sure to allocate enough memory for your new DB instance so that the restore can succeed. You can also allocate additional memory for future growth.

- `--allocated-storage`
- `--db-instance-identifier`
- `--db-instance-class`
- `--engine`
- `--master-user-name`
- `--master-user-password`
- `--s3-bucket-name`
- `--s3-ingestion-role-arn`
- `--s3-prefix`
- `--source-engine`
- `--source-engine-version`

Example

For Linux, OS X, or Unix:

```
aws rds restore-db-instance-from-s3 \
--allocated-storage 250 \
--db-instance-identifier myidentifier \
--db-instance-class db.m4.large \
--engine mysql \
--master-user-name masterawsuser \
--master-user-password masteruserpassword \
--s3-bucket-name mybucket \
--s3-ingestion-role-arn arn:aws:iam::account-number:role/rolename \
--s3-prefix bucketprefix \
--source-engine mysql \
--source-engine-version 5.6.40
```

For Windows:

```
aws rds restore-db-instance-from-s3 ^
```

```
--allocated-storage 250 ^
--db-instance-identifier myidentifier ^
--db-instance-class db.m4.large ^
--engine mysql ^
--master-user-name masterawsuser ^
--master-user-password masteruserpassword ^
--s3-bucket-name mybucket ^
--s3-ingestion-role-arn arn:aws:iam::account-number:role/rolename ^
--s3-prefix bucketprefix ^
--source-engine mysql ^
--source-engine-version 5.6.40
```

RDS API

To import data from Amazon S3 to a new MySQL DB instance by using the Amazon RDS API, call the [RestoreDBInstanceFromS3](#) operation.

Importing Data from a MySQL or MariaDB DB to an Amazon RDS MySQL or MariaDB DB Instance

If your scenario supports it, it is easier to move data in and out of Amazon RDS by using backup files and Amazon S3. For more information, see [Restoring a Backup into an Amazon RDS MySQL DB Instance \(p. 826\)](#).

You can also import data from an existing MySQL or MariaDB database to an Amazon RDS MySQL or MariaDB DB instance. You do so by copying the database with [mysqldump](#) and piping it directly into the Amazon RDS MySQL or MariaDB DB instance. The [mysqldump](#) command-line utility is commonly used to make backups and transfer data from one MySQL or MariaDB server to another. It is included with MySQL and MariaDB client software.

A typical [mysqldump](#) command to move data from an external database to an Amazon RDS DB instance looks similar to the following:

```
mysqldump -u <local_user> \
    --databases <database_name> \
    --single-transaction \
    --compress \
    --order-by-primary \
    -p<local_password> | mysql -u <RDS_user> \
        --port=<port_number> \
        --host=<host_name> \
        -p<RDS_password>
```

Important

Make sure not to leave a space between the `-p` option and the entered password.

Note

- Exclude the following schemas from the dump file: `sys`, `performance_schema`, and `information_schema`. The [mysqldump](#) utility excludes these schemas by default.
- If you need to migrate users and privileges, consider using a tool that generates the data control language (DCL) for recreating them, such as the [pt-show-grants](#) utility.

The parameters used are as follows:

- `-u <local_user>` – Use to specify a user name. In the first usage of this parameter, you specify the name of a user account on the local MySQL or MariaDB database identified by the `--databases` parameter.

- `--databases <database_name>` – Use to specify the name of the database on the local MySQL or MariaDB instance that you want to import into Amazon RDS.
- `--single-transaction` – Use to ensure that all of the data loaded from the local database is consistent with a single point in time. If there are other processes changing the data while `mysqldump` is reading it, using this option helps maintain data integrity.
- `--compress` – Use to reduce network bandwidth consumption by compressing the data from the local database before sending it to Amazon RDS.
- `--order-by-primary` – Use to reduce load time by sorting each table's data by its primary key.
- `-p<local_password>` – Use to specify a password. In the first usage of this parameter, you specify the password for the user account identified by the first `-u` parameter.
- `-u <RDS_user>` – Use to specify a user name. In the second usage of this parameter, you specify the name of a user account on the default database for the Amazon RDS MySQL or MariaDB DB instance identified by the `--host` parameter.
- `--port <port_number>` – Use to specify the port for your Amazon RDS MySQL or MariaDB DB instance. By default, this is 3306 unless you changed the value when creating the instance.
- `--host <host_name>` – Use to specify the DNS name from the Amazon RDS DB instance endpoint, for example, `myinstance.123456789012.us-east-1.rds.amazonaws.com`. You can find the endpoint value in the instance details in the Amazon RDS Management Console.
- `-p<RDS_password>` – Use to specify a password. In the second usage of this parameter, you specify the password for the user account identified by the second `-u` parameter.

You must create any stored procedures, triggers, functions, or events manually in your Amazon RDS database. If you have any of these objects in the database that you are copying, then exclude them when you run `mysqldump` by including the following parameters with your `mysqldump` command: `--routines=0 --triggers=0 --events=0`.

The following example copies the `world` sample database on the local host to an Amazon RDS MySQL DB instance.

For Linux, OS X, or Unix:

```
sudo mysqldump -u localuser \
    --databases world \
    --single-transaction \
    --compress \
    --order-by-primary \
    -plocalpassword | mysql -u rdsuser \
        --port=3306 \
        --host=myinstance.123456789012.us-east-1.rds.amazonaws.com \
        -prdspassword
```

For Windows, the following command needs to be run in a command prompt that has been opened by right-clicking **Command Prompt** on the Windows programs menu and choosing **Run as administrator**:

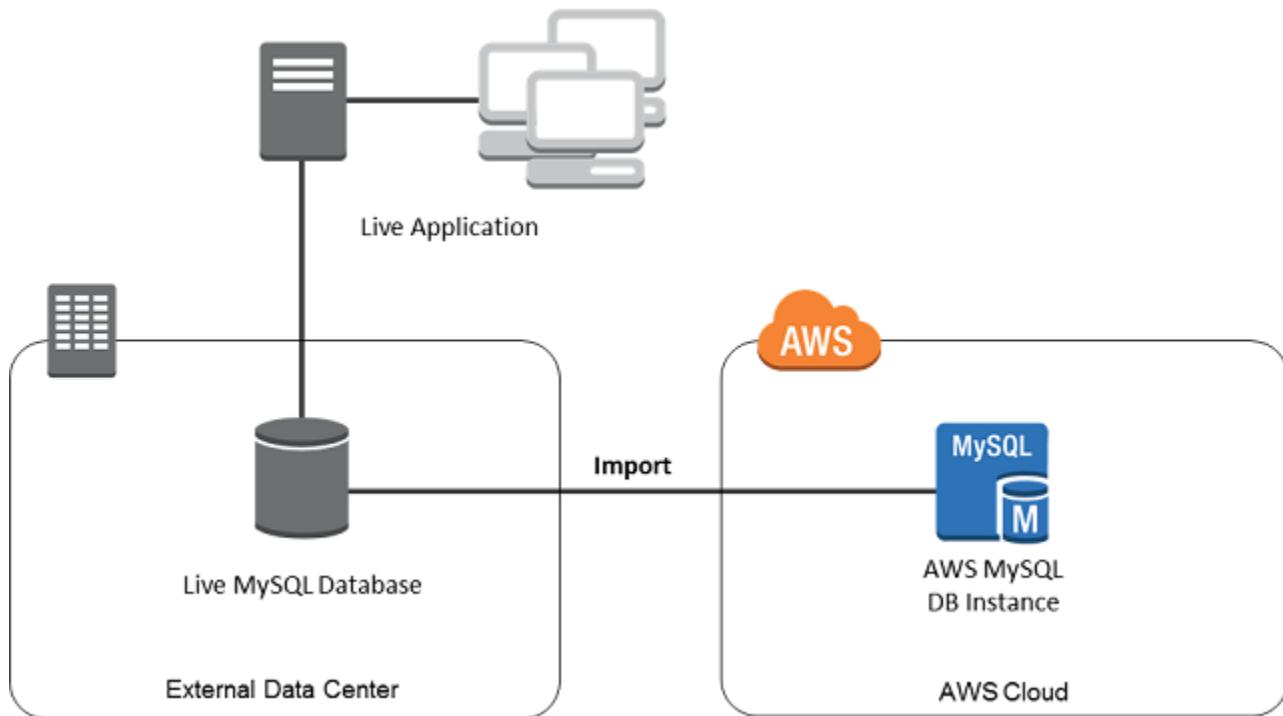
```
mysqldump -u localuser ^
    --databases world ^
    --single-transaction ^
    --compress ^
    --order-by-primary ^
    -plocalpassword | mysql -u rdsuser ^
        --port=3306 ^
        --host=myinstance.123456789012.us-east-1.rds.amazonaws.com ^
        -prdspassword
```

Importing Data to an Amazon RDS MySQL or MariaDB DB Instance with Reduced Downtime

If your scenario supports it, it is easier to move data in and out of Amazon RDS by using backup files and Amazon S3. For more information, see [Restoring a Backup into an Amazon RDS MySQL DB Instance \(p. 826\)](#).

In some cases, you might need to import data from an external MySQL or MariaDB database that supports a live application to an Amazon RDS MySQL or MariaDB DB instance. In these cases, you can use the following procedure to minimize the impact on application availability. This procedure can also help if you are working with a very large database. Here, the procedure helps because you can reduce the cost of the import by reducing the amount of data that is passed across the network to AWS.

In this procedure, you transfer a copy of your database data to an Amazon EC2 instance and import the data into a new Amazon RDS DB instance. You then use replication to bring the Amazon RDS DB instance up-to-date with your live external instance, before redirecting your application to the Amazon RDS DB instance. You configure MariaDB replication based on global transaction identifiers (GTIDs) if the external instance is MariaDB 10.0.2 or greater and the target instance is Amazon RDS MariaDB; otherwise, you configure replication based on binary log coordinates. We recommend GTID-based replication if your external database supports it due to its enhanced crash-safety features. For more information, see [Global Transaction ID](#) in the MariaDB documentation.

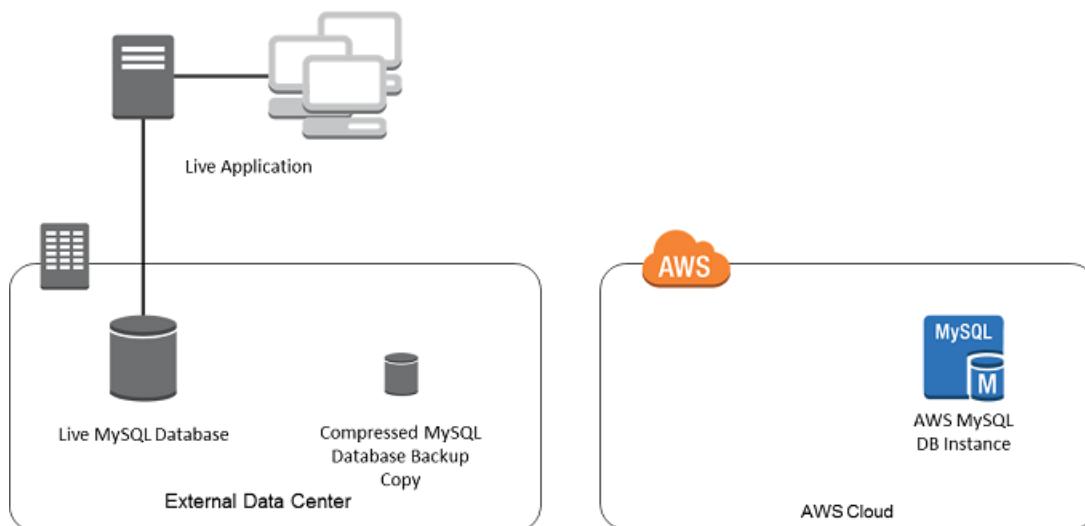


Note

We don't recommend that you use this procedure with source MySQL databases from MySQL versions earlier than version 5.1, due to potential replication issues. For more information, see [Replication Compatibility Between MySQL Versions](#) in the MySQL documentation.

Create a Copy of Your Existing Database

The first step in the process of migrating a large amount of data to an Amazon RDS MySQL or MariaDB DB instance with minimal downtime is to create a copy of the source data.



You can use the `mysqldump` utility to create a database backup in either SQL or delimited-text format. You should do a test run with each format in a nonproduction environment to see which method minimizes the amount of time that `mysqldump` runs.

You should also weigh `mysqldump` performance against the benefit offered by using the delimited-text format for loading. A backup using delimited-text format creates a tab-separated text file for each table being dumped. You can load these files in parallel using the `LOAD DATA LOCAL INFILE` command to reduce the amount of time required to import your database. For more information about choosing a `mysqldump` format and then loading the data, see [Using mysqldump For Backups](#) in the MySQL documentation.

Before you start the backup operation, you must set the replication options on the MySQL or MariaDB database that you are copying to Amazon RDS. The replication options include enabling binary logging and setting a unique server ID. Setting these options causes your server to start logging database transactions and prepares it to be a replication master later in this process.

Note

- Your database needs to be stopped to set the replication options and be in read-only mode while the backup copy is created, so you need to schedule a maintenance window for these operations.
- Exclude the following schemas from the dump file: `sys`, `performance_schema`, and `information_schema`. The `mysqldump` utility excludes these schemas by default.
- If you need to migrate users and privileges, consider using a tool that generates the data control language (DCL) for recreating them, such as the [pt-show-grants](#) utility.

To Set Replication Options

1. Edit the `my.cnf` file (this file is usually under `/etc`).

```
sudo vi /etc/my.cnf
```

Add the `log_bin` and `server_id` options to the `[mysqld]` section. The `log_bin` option provides a file name identifier for binary log files. The `server_id` option provides a unique identifier for the server in master-replica relationships.

The following example shows the updated `[mysqld]` section of a `my.cnf` file:

```
[mysqld]
log-bin=mysql-bin
server-id=1
```

For more information, see [Setting the Replication Master Configuration](#) in the MySQL documentation.

2. Restart the mysql service.

```
sudo service mysqld restart
```

To Create a Backup Copy of Your Existing Database

1. Create a backup of your data using the mysqldump utility, specifying either SQL or delimited-text format.

You must specify `--master-data=2` in order to create a backup file that can be used to start replication between servers. For more information, see the [mysqldump](#) documentation.

To improve performance and ensure data integrity, use the `--order-by-primary` and `--single-transaction` options of mysqldump.

To avoid including the MySQL system database in the backup, do not use the `--all-databases` option with mysqldump. For more information, see [Creating a Dump Snapshot Using mysqldump](#) in the MySQL documentation.

Use `chmod` if necessary to make sure that the directory where the backup file is being created is writeable.

Important

On Windows, run the command window as an administrator.

- To produce SQL output, use the following command.

For Linux, OS X, or Unix:

```
sudo mysqldump \
--databases <database_name> \
--master-data=2 \
--single-transaction \
--order-by-primary \
-r backup.sql \
-u <local_user> \
-p <password>
```

For Windows:

```
mysqldump ^
--databases <database_name> ^
--master-data=2 ^
--single-transaction ^
--order-by-primary ^
-r backup.sql ^
-u <local_user> ^
-p <password>
```

- To produce delimited-text output, use the following command.

For Linux, OS X, or Unix:

```
sudo mysqldump \
--tab=<target_directory> \
--fields-terminated-by ',' \
--fields-enclosed-by '"' \
--lines-terminated-by 0x0d0a \
<database_name> \
--master-data=2 \
--single-transaction \
--order-by-primary \
-p <password>
```

For Windows:

```
mysqldump ^
--tab=<target_directory> ^
--fields-terminated-by ',' ^
--fields-enclosed-by '"' ^
--lines-terminated-by 0x0d0a ^
<database_name> ^
--master-data=2 ^
--single-transaction ^
--order-by-primary ^
-p <password>
```

Note

You must create any stored procedures, triggers, functions, or events manually in your Amazon RDS database. If you have any of these objects in the database that you are copying, exclude them when you run `mysqldump` by including the following arguments with your `mysqldump` command: `--routines=0 --triggers=0 --events=0`.

When using the delimited-text format, a CHANGE MASTER TO comment is returned when you run `mysqldump`. This comment contains the master log file name and position. If the external instance is other than MariaDB version 10.0.2 or greater, note the values for `MASTER_LOG_FILE` and `MASTER_LOG_POS`; you need these values when setting up replication.

```
-- Position to start replication or point-in-time recovery from
--
-- CHANGE MASTER TO MASTER_LOG_FILE='mysql-bin-changelog.000031', MASTER_LOG_POS=107;
```

If you are using SQL format, you can get the master log file name and position in step 4 of the procedure at [Replicate Between Your External Database and New Amazon RDS DB Instance \(p. 843\)](#). If the external instance is MariaDB version 10.0.2 or greater, you can get the GTID in the next step.

2. If the external instance you are using is MariaDB version 10.0.2 or greater, you use GTID-based replication. Run `SHOW MASTER STATUS` on the external MariaDB instance to get the binary log file name and position, then convert them to a GTID by running `BINLOG_GTID_POS` on the external MariaDB instance.

```
SELECT BINLOG_GTID_POS('<binary log file name>', <binary log file position>);
```

Note the GTID returned; you need it to configure replication.

3. Compress the copied data to reduce the amount of network resources needed to copy your data to the Amazon RDS DB instance. Take note of the size of the backup file; you need this information when determining how large an Amazon EC2 instance to create. When you are done, compress the backup file using GZIP or your preferred compression utility.
 - To compress SQL output, use the following command.

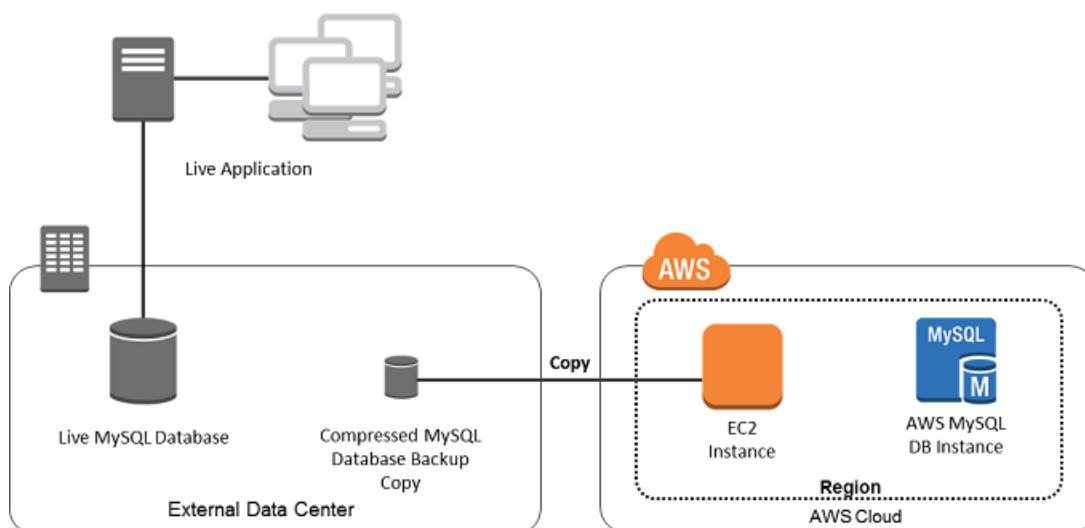
```
gzip backup.sql
```

- To compress delimited-text output, use the following command.

```
tar -zcvf backup.tar.gz <target_directory>
```

Create an Amazon EC2 Instance and Copy the Compressed Database

Copying your compressed database backup file to an Amazon EC2 instance takes fewer network resources than doing a direct copy of uncompressed data between database instances. After your data is in Amazon EC2, you can copy it from there directly to your Amazon RDS MySQL or MariaDB DB instance. For you to save on the cost of network resources, your Amazon EC2 instance must be in the same AWS Region as your Amazon RDS DB instance. Having the Amazon EC2 instance in the same AWS Region as your Amazon RDS DB instance also reduces network latency during the import.



To Create an Amazon EC2 Instance and Copy Your Data

1. In the AWS Region where you plan to create the RDS DB instance to run your MySQL database engine, create a VPC, a VPC security group, and a VPC subnet. Ensure that the inbound rules for your VPC security group allow the IP addresses required for your application to connect to AWS. This can be a range of IP addresses (for example, 203.0.113.0/24), or another VPC security group. You can use the [Amazon VPC Management Console](#) to create and manage VPCs, subnets, and security groups. For more information, see [Getting Started with Amazon VPC](#) in the *Amazon Virtual Private Cloud Getting Started Guide*.

Note

Older AWS accounts can also launch instances in Amazon EC2-Classic mode. In this case, make sure that the inbound rules in the DB security group for your Amazon RDS instance allow access for your EC2-Classic instance using the Amazon EC2 private IP address. For more information, see [Working with DB Security Groups \(EC2-Classic Platform\) \(p. 491\)](#).

2. Open the [Amazon EC2 Management Console](#) and select the AWS Region to contain both your Amazon EC2 instance and your Amazon RDS DB instance. Launch an Amazon EC2 instance using the VPC, subnet, and security group that you created in Step 1. Ensure that you select an instance type with enough storage for your database backup file when it is uncompressed. For details on Amazon EC2

instances, see [Getting Started with Amazon EC2 Linux Instances](#) in the *Amazon Elastic Compute Cloud User Guide for Linux*.

3. To connect to your Amazon RDS DB instance from your Amazon EC2 instance, you need to edit your VPC security group, and add an inbound rule specifying the private IP address of your EC2 instance. You can find the private IP address on the **Details** tab of the **Instance** pane in the EC2 console window. To edit the VPC security group and add an inbound rule, choose **Security Groups** in the EC2 console navigation pane, choose your security group, and then add an inbound rule for MySQL/Aurora specifying the private IP address of your EC2 instance. To learn how to add an inbound rule to a VPC security group, see [Adding and Removing Rules](#).
4. Copy your compressed database backup file from your local system to your Amazon EC2 instance. Use `chmod` if necessary to make sure you have write permission for the target directory of the Amazon EC2 instance. You can use `scp` or an SSH client to copy the file. The following is an example:

```
$ scp -r -i <key pair>.pem backup.sql.gz ec2-user@<EC2 DNS>:</target_directory>/  
backup.sql.gz
```

Important

Be sure to copy sensitive data using a secure network transfer protocol.

5. Connect to your Amazon EC2 instance and install the latest updates and the MySQL client tools using the following commands:

```
sudo yum update -y  
sudo yum install mysql -y
```

For more information, see [Connect to Your Instance](#) in the *Amazon Elastic Compute Cloud User Guide for Linux*.

6. While connected to your Amazon EC2 instance, decompress your database backup file. For example:
 - To decompress SQL output, use the following command:

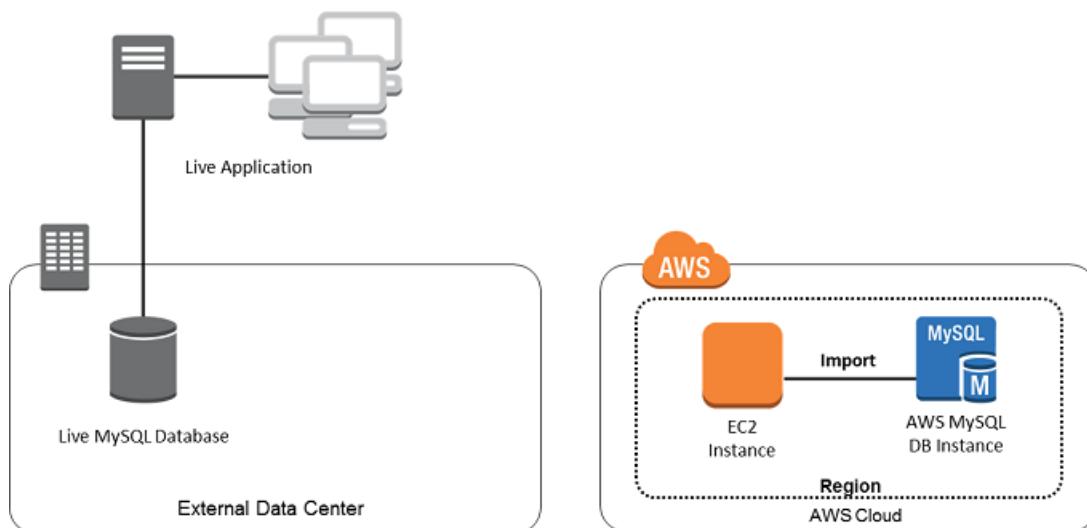
```
gzip backup.sql.gz -d
```

- To decompress delimited-text output, use the following command:

```
tar xzvf backup.tar.gz
```

Create an Amazon RDS MySQL or MariaDB DB instance and Import Data from Your Amazon EC2 Instance

By creating an Amazon RDS MySQL or MariaDB DB instance in the same AWS Region as your Amazon EC2 instance, you can import the database backup file from EC2 faster than over the internet.



To Create an Amazon RDS MySQL or MariaDB DB Instance and Import Your Data

1. Determine which DB instance class and what amount of storage space is required to support the expected workload for this Amazon RDS DB instance. As part of this process, decide what is sufficient space and processing capacity for your data load procedures, and also what is required to handle the production workload. You can estimate this based on the size and resources of the source MySQL or MariaDB database. For more information, see [Choosing the DB Instance Class \(p. 107\)](#).
2. Determine if Amazon RDS provisioned input/output operations per second (IOPS) is required to support the workloads. Provisioned IOPS storage delivers fast throughput for online transaction processing (OLTP) workloads, which are I/O intensive. For more information, see [Provisioned IOPS SSD Storage \(p. 137\)](#).
3. Open the [Amazon RDS console](#). In the upper-right corner, choose the AWS Region that contains your Amazon EC2 instance.
4. In the navigation pane, choose **Databases**.
5. Choose **Create database**, and then go through the steps to choose options for your DB instance:
 - a. Make sure that **Standard Create** is chosen.
 - b. In the **Engine options** section, choose **MySQL** or **MariaDB**, as appropriate.
 - c. For **Version**, choose the version that is compatible with your source MySQL instance, as follows:
 - If your source instance is MySQL 5.1.x, the Amazon RDS DB instance must be MySQL 5.5.x.
 - If your source instance is MySQL 5.5.x, the Amazon RDS DB instance must be MySQL 5.5.x or greater.
 - If your source instance is MySQL 5.6.x, the Amazon RDS DB instance must be MySQL 5.6.x or MariaDB.
 - If your source instance is MySQL 5.7.x, the Amazon RDS DB instance must be MySQL 5.7.x, 5.6.x, or MariaDB.
 - If your source instance is MySQL 8.0.x, the Amazon RDS DB instance must be MySQL 8.0.x.
 - If your source instance is MariaDB 5.1, 5.2, or 5.3, the Amazon RDS DB instance must be MySQL 5.1.x.
 - If your source instance is MariaDB 5.5 or greater, the Amazon RDS DB instance must be MariaDB.
 - d. In the **Templates** section, choose **Dev/Test** to skip configuring Multi-AZ deployment and provisioned IOPS storage.

- e. In the **Settings** section, specify the requested **DB instance identifier** and user information.
- f. In the **DB instance class** and **Storage** sections, specify the DB instance class and allocated storage size that you want.
- g. In the **Availability & durability** section, choose **Do not create a standby instance for Multi-AZ deployment**.
- h. In the **Connectivity** section, choose the same virtual private cloud (VPC) and VPC security group as for your Amazon EC2 instance. This approach ensures that your Amazon EC2 instance and your Amazon RDS instance are visible to each other over the network. Set **Publicly accessible** to **Yes**. To set up replication with your source database as described later, your DB instance must be publicly accessible.

Use the default values for the other settings in this section.

In the **Backup** section, set the backup retention period to **0 days**.

Use the default values for the other settings in this section.

- i. Open the **Additional configuration** section, and specify an **Initial database name**.

Set the **Backup retention period** to **0 days**

Use the default values for the other settings in this section.

- j. Choose **Create database**.

Your new DB instance appears in the **Databases** list with the status **Creating**. Wait for the **Status** of your new DB instance to show as **Available**.

Don't configure multiple Availability Zones, backup retention, or Read Replicas until after you have imported the database backup. When that import is done, you can set Multi-AZ and backup retention the way you want them for the production instance. For a detailed walkthrough of creating an Amazon RDS MySQL DB instance, see [Creating a DB Instance Running the MySQL Database Engine \(p. 751\)](#). For a detailed walkthrough of creating an Amazon RDS MariaDB DB instance, see [Creating a DB Instance Running the MariaDB Database Engine \(p. 543\)](#).

6. Review the default configuration options for the Amazon RDS DB instance. In the RDS console navigation pane, choose **Parameter groups**, and then choose the magnifying glass icon next to the **default.mysqlx.x** or **default.mariadb.x** parameter group. If this parameter group doesn't have the configuration options that you want, find a different one that does or create a new parameter group. For more information on creating a parameter group, see [Working with DB Parameter Groups \(p. 211\)](#).

If you decide to use a different parameter group than the default, associate it with your Amazon RDS DB instance. For more information, see [Modifying a DB Instance Running the MySQL Database Engine \(p. 802\)](#) or [Modifying a DB Instance Running the MariaDB Database Engine \(p. 564\)](#).

7. Connect to the new Amazon RDS DB instance as the master user, and create the users required to support the administrators, applications, and services that need to access the instance. The host name for the Amazon RDS DB instance is the **Endpoint** value for this instance without including the port number, for example `mysampledbs.claxc2oy9ak1.us-west-2.rds.amazonaws.com`. You can find the endpoint value in the instance details in the Amazon RDS Management Console.
8. Connect to your Amazon EC2 instance. For more information, see [Connect to Your Instance](#) in the [Amazon Elastic Compute Cloud User Guide for Linux](#).
9. Connect to your Amazon RDS DB instance as a remote host from your Amazon EC2 instance using the `mysql` command. The following is an example.

```
mysql -h <host_name> -P 3306 -u <db_master_user> -p
```

The host name is the DNS name from the Amazon RDS DB instance endpoint.

10. At the mysql prompt, run the source command and pass it the name of your database dump file to load the data into the Amazon RDS DB instance:

- For SQL format, use the following command.

```
mysql> source backup.sql;
```

- For delimited-text format, first create the database (if it isn't the default database you created when setting up the Amazon RDS DB instance).

```
mysql> create database <database_name>;  
$ mysql> use <database_name>;
```

Then create the tables.

```
mysql> source <table1>.sql  
$ mysql> source <table2>.sql  
etc...
```

Then import the data.

```
mysql> LOAD DATA LOCAL INFILE 'table1.txt' INTO TABLE table1 FIELDS TERMINATED BY ','  
ENCLOSED BY '\"' LINES TERMINATED BY '0x0d0a';  
$ mysql> LOAD DATA LOCAL INFILE 'table2.txt' INTO TABLE table2 FIELDS TERMINATED BY  
'','' ENCLOSED BY '\"' LINES TERMINATED BY '0x0d0a';  
etc...
```

To improve performance, you can perform these operations in parallel from multiple connections so that all of your tables get created and then loaded at the same time.

Note

If you used any data-formatting options with mysqldump when you initially dumped the table, you must use the same options with mysqlimport or LOAD DATA LOCAL INFILE to ensure proper interpretation of the data file contents.

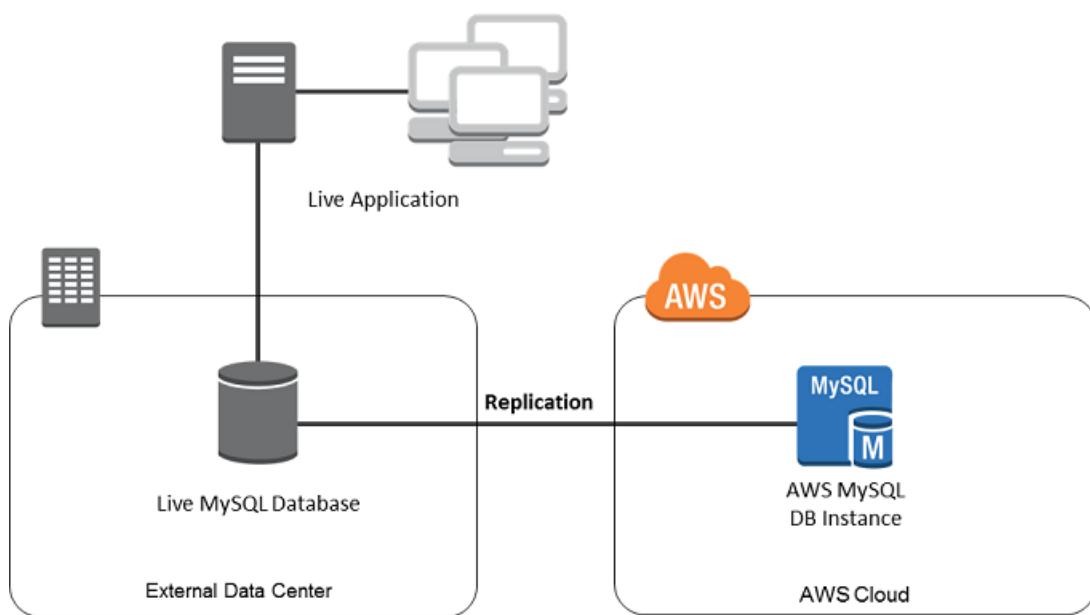
11. Run a simple SELECT query against one or two of the tables in the imported database to verify that the import was successful.

Note

If you no longer need the Amazon EC2 instance used in this procedure, terminate the EC2 instance to reduce your AWS resource usage. To terminate an EC2 instance, see [Terminating an Instance](#).

Replicate Between Your External Database and New Amazon RDS DB Instance

Your source database was likely updated during the time that it took to copy and transfer the data to the Amazon RDS MySQL or MariaDB DB instance. That being the case, you can use replication to bring the copied database up-to-date with the source database.



Note

The permissions required to start replication on an Amazon RDS DB instance are restricted and not available to your Amazon RDS master user. Because of this, you must use either the Amazon RDS [mysql.rds_set_external_master \(p. 891\)](#) command or the [mysql.rds_set_external_master_gtid \(p. 595\)](#) command to configure replication, and the [mysql.rds_start_replication \(p. 901\)](#) command to start replication between your live database and your Amazon RDS database.

To Start Replication

Earlier, you enabled binary logging and set a unique server ID for your source database. Now you can set up your Amazon RDS DB instance as a replica with your live database as the replication master.

1. In the Amazon RDS Management Console, add the IP address of the server that hosts the source database to the VPC security group for the Amazon RDS DB instance. For more information on modifying a VPC security group, see [Security Groups for Your VPC](#) in the *Amazon Virtual Private Cloud User Guide*.

You might also need to configure your local network to permit connections from the IP address of your Amazon RDS DB instance, so that it can communicate with your source instance. To find the IP address of the Amazon RDS DB instance, use the host command.

```
host <RDS_MySQL_DB_host_name>
```

The host name is the DNS name from the Amazon RDS DB instance endpoint, for example `myinstance.123456789012.us-east-1.rds.amazonaws.com`. You can find the endpoint value in the instance details in the Amazon RDS Management Console.

2. Using the client of your choice, connect to the source instance and create a user to be used for replication. This account is used solely for replication and must be restricted to your domain to improve security. The following is an example.

```
CREATE USER 'repl_user'@'mydomain.com' IDENTIFIED BY '<password>;'
```

3. For the source instance, grant REPLICATION CLIENT and REPLICATION SLAVE privileges to your replication user. For example, to grant the REPLICATION CLIENT and REPLICATION SLAVE privileges on all databases for the 'repl_user' user for your domain, issue the following command.

```
GRANT REPLICATION CLIENT, REPLICATION SLAVE ON *.* TO 'repl_user'@'mydomain.com'  
IDENTIFIED BY '<password>';
```

4. If you used SQL format to create your backup file and the external instance is not MariaDB 10.0.2 or greater, look at the contents of that file.

```
cat backup.sql
```

The file includes a CHANGE MASTER TO comment that contains the master log file name and position. This comment is included in the backup file when you use the --master-data option with mysqldump. Note the values for MASTER_LOG_FILE and MASTER_LOG_POS.

```
--  
-- Position to start replication or point-in-time recovery from  
--  
-- CHANGE MASTER TO MASTER_LOG_FILE='mysql-bin-changelog.000031', MASTER_LOG_POS=107;
```

If you used delimited text format to create your backup file and the external instance is not MariaDB 10.0.2 or greater, you should already have binary log coordinates from step 1 of the procedure at [To Create a Backup Copy of Your Existing Database \(p. 837\)](#).

If the external instance is MariaDB 10.0.2 or greater, you should already have the GTID from which to start replication from step 2 of the procedure at [To Create a Backup Copy of Your Existing Database \(p. 837\)](#).

5. Make the Amazon RDS DB instance the replica. If the external instance is not MariaDB 10.0.2 or greater, connect to the Amazon RDS DB instance as the master user and identify the source database as the replication master by using the [mysql.rds_set_external_master \(p. 891\)](#) command. Use the master log file name and master log position that you determined in the previous step if you have a SQL format backup file. Alternatively, use the name and position that you determined when creating the backup files if you used delimited-text format. The following is an example.

```
CALL mysql.rds_set_external_master ('mymasterserver.mydomain.com', 3306,  
'repl_user', '<password>', 'mysql-bin-changelog.000031', 107, 0);
```

If the external instance is MariaDB 10.0.2 or greater, connect to the Amazon RDS DB instance as the master user and identify the source database as the replication master by using the [mysql.rds_set_external_master_gtid \(p. 595\)](#) command. Use the GTID that you determined in step 2 of the procedure at [To Create a Backup Copy of Your Existing Database \(p. 837\)](#). The following is an example.

```
CALL mysql.rds_set_external_master_gtid ('<master_server_ip_address>', 3306,  
'ReplicationUser', '<password>', '<GTID>', 0);
```

The `master_server_ip_address` is the IP address of master MySQL instance. An EC2 private DNS address is currently not supported.

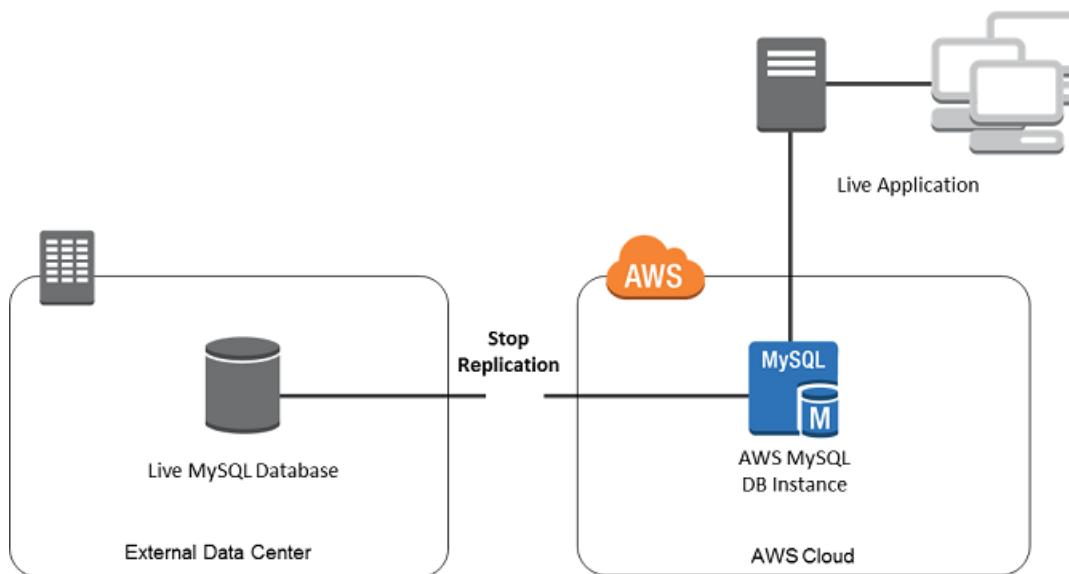
6. On the Amazon RDS DB instance, issue the [mysql.rds_start_replication \(p. 901\)](#) command to start replication.

```
CALL mysql.rds_start_replication;
```

7. On the Amazon RDS DB instance, run the [SHOW SLAVE STATUS](#) command to determine when the replica is up-to-date with the replication master. The results of the SHOW SLAVE STATUS command include the Seconds_Behind_Master field. When the Seconds_Behind_Master field returns 0, then the replica is up-to-date with the master.
8. After the Amazon RDS DB instance is up-to-date, enable automated backups so you can restore that database if needed. You can enable or modify automated backups for your Amazon RDS DB instance using the [Amazon RDS Management Console](#). For more information, see [Working With Backups \(p. 254\)](#).

Redirect Your Live Application to Your Amazon RDS Instance

After the Amazon RDS MySQL or MariaDB DB instance is up-to-date with the replication master, you can now update your live application to use the Amazon RDS instance.



To Redirect Your Live Application to Your Amazon RDS MySQL or MariaDB DB Instance and Stop Replication

1. To add the VPC security group for the Amazon RDS DB instance, add the IP address of the server that hosts the application. For more information on modifying a VPC security group, see [Security Groups for Your VPC](#) in the *Amazon Virtual Private Cloud User Guide*.
2. Verify that the Seconds_Behind_Master field in the [SHOW SLAVE STATUS](#) command results is 0, which indicates that the replica is up-to-date with the replication master.

```
SHOW SLAVE STATUS;
```

3. Close all connections to the source when their transactions complete.
4. Update your application to use the Amazon RDS DB instance. This update typically involves changing the connection settings to identify the host name and port of the Amazon RDS DB instance, the user account and password to connect with, and the database to use.
5. Stop replication for the Amazon RDS instance using the [mysql.rds_stop_replication \(p. 903\)](#) command.

```
CALL mysql.rds_stop_replication;
```

6. Run the [mysql.rds_reset_external_master \(p. 897\)](#) command on your Amazon RDS DB instance to reset the replication configuration so this instance is no longer identified as a replica.

```
CALL mysql.rds_reset_external_master;
```

7. Enable additional Amazon RDS features such as Multi-AZ support and Read Replicas. For more information, see [High Availability \(Multi-AZ\) for Amazon RDS \(p. 144\)](#) and [Working with Read Replicas \(p. 182\)](#).

Note

If you no longer need the Amazon RDS instance used in this procedure, you should delete the RDS instance to reduce your Amazon AWS resource usage. To delete an RDS instance, see [Deleting a DB Instance \(p. 174\)](#).

Importing Data From Any Source to a MySQL or MariaDB DB Instance

If you have more than 1 GiB of data to load, or if your data is coming from somewhere other than a MySQL or MariaDB database, we recommend creating flat files and loading them with `mysqldump`. `mysqldump` is another command line utility bundled with the MySQL and MariaDB client software whose purpose is to load flat files into MySQL or MariaDB. For information about `mysqldump`, see [mysqldump - A Data Import Program](#) in the MySQL documentation.

We also recommend creating DB snapshots of the target Amazon RDS DB instance before and after the data load. Amazon RDS DB snapshots are complete backups of your DB instance that can be used to restore your DB instance to a known state. When you initiate a DB snapshot, I/O operations to your database instance are momentarily suspended while your database is backed up.

Creating a DB snapshot immediately before the load lets you restore the database to its state before the load, if you need to. A DB snapshot taken immediately after the load protects you from having to load the data again in case of a mishap and can also be used to seed new database instances.

The following list shows the steps to take. Each step is discussed in more detail below.

1. Create flat files containing the data to be loaded.
2. Stop any applications accessing the target DB instance.
3. Create a DB snapshot.
4. Consider disabling Amazon RDS automated backups.
5. Load the data using `mysqldump`.
6. Enable automated backups again.

Step 1: Create Flat Files Containing the Data to be Loaded

Use a common format, such as CSV (Comma-Separated Values), to store the data to be loaded. Each table must have its own file; data for multiple tables cannot be combined in the same file. Give each file the same name as the table it corresponds to. The file extension can be anything you like. For example, if the table name is "sales", the file name could be "sales.csv" or "sales.txt", but not "sales_01.csv".

Whenever possible, order the data by the primary key of the table being loaded. This drastically improves load times and minimizes disk storage requirements.

The speed and efficiency of this procedure is dependent upon keeping the size of the files small. If the uncompressed size of any individual file is larger than 1 GiB, split it into multiple files and load each one separately.

On Unix-like systems (including Linux), use the 'split' command. For example, the following command splits the sales.csv file into multiple files of less than 1 GiB, splitting only at line breaks (-C 1024m). The new files are named sales.part_00, sales.part_01, and so on.

```
split -C 1024m -d sales.csv sales.part_
```

Similar utilities are available on other operating systems.

Step 2: Stop Any Applications Accessing the Target DB Instance

Before starting a large load, stop all application activity accessing the target DB instance that you plan to load to. We recommend this particularly if other sessions will be modifying the tables being loaded or tables they reference. Doing this reduces the risk of constraint violations occurring during the load and improves load performance. It also makes it possible to restore the database instance to the point just before the load without losing changes made by processes not involved in the load.

Of course, this might not be possible or practical. If you are unable to stop applications from accessing the DB instance before the load, take steps to ensure the availability and integrity of your data. The specific steps required vary greatly depending upon specific use cases and site requirements.

Step 3: Create a DB Snapshot

If you plan to load data into a new DB instance that contains no data, you can skip this step. Otherwise, creating a DB snapshot of your DB instance allows you to restore the DB instance to the point just before the load, if it becomes necessary. As previously mentioned, when you initiate a DB snapshot, I/O operations to your database instance are suspended for a few minutes while the database is backed up.

In the example below, we use the AWS CLI `create-db-snapshot` command to create a DB snapshot of our AcmeRDS instance and give the DB snapshot the identifier "preload".

For Linux, OS X, or Unix:

```
aws rds create-db-snapshot \
--db-instance-identifier AcmeRDS \
--db-snapshot-identifier preload
```

For Windows:

```
aws rds create-db-snapshot ^
--db-instance-identifier AcmeRDS ^
--db-snapshot-identifier preload
```

You can also use the restore from DB snapshot functionality in order to create test database instances for dry runs or to "undo" changes made during the load.

Keep in mind that restoring a database from a DB snapshot creates a new DB instance that, like all DB instances, has a unique identifier and endpoint. If you need to restore the database instance without changing the endpoint, you must first delete the DB instance so that the endpoint can be reused.

For example, to create a DB instance for dry runs or other testing, you would give the DB instance its own identifier. In the example, "AcmeRDS-2" is the identifier and we would connect to the database instance using the endpoint associated with AcmeRDS-2.

For Linux, OS X, or Unix:

```
aws rds restore-db-instance-from-db-snapshot \  
  --db-instance-identifier AcmeRDS-2 \  
  --db-snapshot-identifier preload
```

For Windows:

```
aws rds restore-db-instance-from-db-snapshot ^  
  --db-instance-identifier AcmeRDS-2 ^  
  --db-snapshot-identifier preload
```

To reuse the existing endpoint, we must first delete the database instance and then give the restored database the same identifier.

For Linux, OS X, or Unix:

```
aws rds delete-db-instance \  
  --db-instance-identifier AcmeRDS \  
  --final-db-snapshot-identifier AcmeRDS-Final  
  
aws rds restore-db-instance-from-db-snapshot \  
  --db-instance-identifier AcmeRDS \  
  --db-snapshot-identifier preload
```

For Windows:

```
aws rds delete-db-instance ^  
  --db-instance-identifier AcmeRDS ^  
  --final-db-snapshot-identifier AcmeRDS-Final  
  
aws rds restore-db-instance-from-db-snapshot ^  
  --db-instance-identifier AcmeRDS ^  
  --db-snapshot-identifier preload
```

The example takes a final DB snapshot of the database instance before deleting it. This is optional, but recommended.

Step 4: Consider Disabling Amazon RDS Automated Backups

Warning

Do not disable automated backups if you need the ability to perform point-in-time recovery.

Disabling automated backups erases all existing backups, so point-in-time recovery is not possible after automated backups have been disabled. Disabling automated backups is a performance optimization and is not required for data loads. DB snapshots are not affected by disabling automated backups. All existing DB snapshots are still available for restore.

Disabling automated backups reduces load time by about 25 percent and reduce the amount of storage space required during the load. If you plan to load data into a new DB instance that contains no data, disabling backups is an easy way to speed up the load and avoid using the additional storage needed for backups. However, if you plan to load into a DB instance that already contains data, weigh the benefits of disabling backups against the impact of losing the ability to perform point-in-time-recovery.

DB instances have automated backups enabled by default (with a one day retention period). In order to disable automated backups, you must set the backup retention period to zero. After the load, you can re-enable backups by setting the backup retention period to a non-zero value. In order to enable or disable

backups, Amazon RDS must shut the DB instance down and restart it in order to turn MySQL or MariaDB logging on or off.

Use the AWS CLI `modify-db-instance` command to set the backup retention to zero and apply the change immediately. Setting the retention period to zero requires a DB instance restart, so wait until the restart has completed before proceeding.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier AcmeRDS \
--apply-immediately \
--backup-retention-period 0
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier AcmeRDS ^
--apply-immediately ^
--backup-retention-period 0
```

You can check the status of your DB instance with the AWS CLI `describe-db-instances` command. The example displays the status of the AcmeRDS database instance and includes the `--headers` option to show column headings.

For Linux, OS X, or Unix:

```
aws rds describe-db-instances \
--db-instance-identifier AcmeRDS \
--headers
```

For Windows:

```
aws rds describe-db-instances ^
--db-instance-identifier AcmeRDS ^
--headers
```

When the Status column shows that the database is available, you're ready to proceed.

Step 5: Load the Data

Use the `mysqlimport` utility to load the flat files into Amazon RDS. In the example we tell `mysqlimport` to load all of the files named "sales" with an extension starting with "part_". This is a convenient way to load all of the files created in the "split" example. Use the `--compress` option to minimize network traffic. The `--fields-terminated-by=','` option is used for CSV files and the `--local` option specifies that the incoming data is located on the client. Without the `--local` option, the Amazon RDS DB instance looks for the data on the database host, so always specify the `--local` option.

For Linux, OS X, or Unix:

```
mysqlimport --local \
--compress \
--user=username \
--password \
--host=hostname \
--fields-terminated-by=',' Acme sales.part_*
```

For Windows:

```
mysqlimport --local ^
--compress ^
--user=username ^
--password ^
--host=hostname ^
--fields-terminated-by=',' Acme sales.part_*
```

For very large data loads, take additional DB snapshots periodically between loading files and note which files have been loaded. If a problem occurs, you can easily resume from the point of the last DB snapshot, avoiding lengthy reloads.

Step 6: Enable Amazon RDS Automated Backups

After the load is finished, re-enable Amazon RDS automated backups by setting the backup retention period back to its pre-load value. As noted earlier, Amazon RDS restarts the DB instance, so be prepared for a brief outage.

In the example, we use the AWS CLI modify-db-instance command to enable automated backups for the AcmeRDS DB instance and set the retention period to 1 day.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier AcmeRDS \
--backup-retention-period 1 \
--apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier AcmeRDS ^
--backup-retention-period 1 ^
--apply-immediately
```

Working with MySQL Replication

You usually use Read Replicas to configure replication between Amazon RDS DB instances. For general information about Read Replicas, see [Working with Read Replicas \(p. 182\)](#). For specific information about working with Read Replicas on Amazon RDS MySQL, see [Working with MySQL Read Replicas \(p. 852\)](#).

You can use global transaction identifiers (GTIDs) for replication with Amazon RDS MySQL. For more information, see [Using GTID-Based Replication for Amazon RDS MySQL \(p. 858\)](#).

You can also set up replication between an Amazon RDS MySQL DB instance and a MySQL or MariaDB instance that is external to Amazon RDS. For information about configuring replication with an external source, see [Replication with a MySQL or MariaDB Instance Running External to Amazon RDS \(p. 862\)](#).

For any of these replication options, you can use either row-based replication, statement-based, or mixed replication. Row-based replication only replicates the changed rows that result from a SQL statement. Statement-based replication replicates the entire SQL statement. Mixed replication uses statement-based replication when possible, but switches to row-based replication when SQL statements that are unsafe for statement-based replication are executed. In most cases, mixed replication is recommended. The binary log format of the DB instance determines whether replication is row-based, statement-based, or mixed. For information about setting the binary log format, see [Binary Logging Format \(p. 410\)](#).

Note

You can configure replication to import databases from a MySQL or MariaDB instance that is external to Amazon RDS, or to export databases to such instances. For more information, see [Importing Data to an Amazon RDS MySQL or MariaDB DB Instance with Reduced Downtime \(p. 835\)](#) and [Exporting Data from a MySQL DB Instance by Using Replication \(p. 869\)](#).

Topics

- [Working with MySQL Read Replicas \(p. 852\)](#)
- [Using GTID-Based Replication for Amazon RDS MySQL \(p. 858\)](#)
- [Replication with a MySQL or MariaDB Instance Running External to Amazon RDS \(p. 862\)](#)

Working with MySQL Read Replicas

This section contains specific information about working with Read Replicas on Amazon RDS MySQL. For general information about Read Replicas and instructions for using them, see [Working with Read Replicas \(p. 182\)](#).

Topics

- [Read Replica Configuration with MySQL \(p. 852\)](#)
- [Configuring Delayed Replication with MySQL \(p. 853\)](#)
- [Read Replica Updates with MySQL \(p. 855\)](#)
- [Multi-AZ Read Replica Deployments with MySQL \(p. 855\)](#)
- [Monitoring MySQL Read Replicas \(p. 856\)](#)
- [Starting and Stopping Replication with MySQL Read Replicas \(p. 856\)](#)
- [Deleting Read Replicas with MySQL \(p. 856\)](#)
- [Troubleshooting a MySQL Read Replica Problem \(p. 856\)](#)

Read Replica Configuration with MySQL

Before a MySQL DB instance can serve as a replication source, you must enable automatic backups on the source DB instance by setting the backup retention period to a value other than 0. This requirement

also applies to a Read Replica that is the source DB instance for another Read Replica. Automatic backups are supported only for Read Replicas running any version of MySQL 5.6 and later. You can configure replication based on binary log coordinates for a MySQL DB instance.

On Amazon RDS MySQL version 5.7.23 and later MySQL 5.7 versions, you can configure replication using global transaction identifiers (GTIDs). For more information, see [Using GTID-Based Replication for Amazon RDS MySQL \(p. 858\)](#).

You can create up to five Read Replicas from one DB instance. For replication to operate effectively, each Read Replica should have as the same amount of compute and storage resources as the source DB instance. If you scale the source DB instance, you should also scale the Read Replicas.

If a Read Replica is running any version of MySQL 5.6 and later, you can specify it as the source DB instance for another Read Replica. For example, you can create ReadReplica1 from MyDBInstance, and then create ReadReplica2 from ReadReplica1. Updates made to MyDBInstance are replicated to ReadReplica1 and then replicated from ReadReplica1 to ReadReplica2. You can't have more than four instances involved in a replication chain. For example, you can create ReadReplica1 from MySourceDBInstance, and then create ReadReplica2 from ReadReplica1, and then create ReadReplica3 from ReadReplica2, but you can't create a ReadReplica4 from ReadReplica3.

If you promote a MySQL Read Replica that is in turn replicating to other Read Replicas, those Read Replicas remain active. Consider an example where MyDBInstance1 replicates to MyDBInstance2, and MyDBInstance2 replicates to MyDBInstance3. If you promote MyDBInstance2, replication from MyDBInstance1 to MyDBInstance2 no longer occurs, but MyDBInstance2 still replicates to MyDBInstance3.

To enable automatic backups on a Read Replica for Amazon RDS MySQL version 5.6 and later, first create the Read Replica, then modify the Read Replica to enable automatic backups.

You can run multiple concurrent Read Replica create or delete actions that reference the same source DB instance, as long as you stay within the limit of five Read Replicas for the source instance.

Preparing MySQL DB Instances That Use MyISAM

If your MySQL DB instance uses a nontransactional engine such as MyISAM, you need to perform the following steps to successfully set up your Read Replica. These steps are required to make sure that the Read Replica has a consistent copy of your data. These steps are not required if all of your tables use a transactional engine such as InnoDB.

1. Stop all data manipulation language (DML) and data definition language (DDL) operations on non-transactional tables in the source DB instance and wait for them to complete. SELECT statements can continue running.
2. Flush and lock the tables in the source DB instance.
3. Create the Read Replica using one of the methods in the following sections.
4. Check the progress of the Read Replica creation using, for example, the `DescribeDBInstances` API operation. Once the Read Replica is available, unlock the tables of the source DB instance and resume normal database operations.

Configuring Delayed Replication with MySQL

You can use delayed replication as a strategy for disaster recovery. With delayed replication, you specify the minimum amount of time, in seconds, to delay replication from the source to the Read Replica. In the event of a disaster, such as a table deleted unintentionally, you complete the following steps to recover from the disaster quickly:

- Stop replication to the Read Replica before the change that caused the disaster is sent to it.

Use the [mysql.rds_stop_replication \(p. 903\)](#) stored procedure to stop replication.

- Start replication and specify that replication stops automatically at a log file location.

You specify a location just before the disaster using the [mysql.rds_start_replication_until \(p. 901\)](#) stored procedure.

- Promote the Read Replica to be the new source DB instance by using the instructions in [Promoting a Read Replica to Be a Standalone DB Instance \(p. 186\)](#).

Note

- On Amazon RDS MySQL 5.7, delayed replication is supported for MySQL 5.7.22 and later. On Amazon RDS MySQL 5.6, delayed replication is supported for MySQL 5.6.40 and later. Delayed replication is not supported on Amazon RDS MySQL 8.0.
- You must use stored procedures to configure delayed replication. You can't configure delayed replication with the AWS Management Console, the AWS CLI, or the Amazon RDS API.
- On Amazon RDS MySQL 5.7.23 and later MySQL 5.7 versions, you can use GTID-based replication in a delayed replication configuration. If you use GTID-based replication, use the [mysql.rds_start_replication_until_gtid \(p. 902\)](#) stored procedure instead of the [mysql.rds_start_replication_until \(p. 901\)](#) stored procedure. For more information about GTID-based replication, see [Using GTID-Based Replication for Amazon RDS MySQL \(p. 858\)](#).

Topics

- [Configuring Delayed Replication During Read Replica Creation \(p. 854\)](#)
- [Modifying Delayed Replication for an Existing Read Replica \(p. 854\)](#)
- [Setting a Location to Stop Replication to a Read Replica \(p. 855\)](#)

Configuring Delayed Replication During Read Replica Creation

To configure delayed replication for any future Read Replica created from a DB instance, run the [mysql.rds_set_configuration \(p. 908\)](#) stored procedure with the `target_delay` parameter.

To configure delayed replication during Read Replica creation

- Using a MySQL client, connect to the MySQL DB instance that will be the source for Read Replicas as the master user.
- Run the [mysql.rds_set_configuration \(p. 908\)](#) stored procedure with the `target_delay` parameter.

For example, run the following stored procedure to specify that replication is delayed by at least one hour (3600 seconds) for any Read Replica created from the current DB instance.

```
call mysql.rds_set_configuration('target_delay', 3600);
```

Note

After running this stored procedure, any Read Replica you create using the AWS CLI or Amazon RDS API is configured with replication delayed by the specified number of seconds.

Modifying Delayed Replication for an Existing Read Replica

To modify delayed replication for an existing Read Replica, run the [mysql.rds_set_source_delay \(p. 900\)](#) stored procedure.

To modify delayed replication for an existing Read Replica

1. Using a MySQL client, connect to the Read Replica as the master user.
2. Use the [mysql.rds_stop_replication \(p. 903\)](#) stored procedure to stop replication.
3. Run the [mysql.rds_set_source_delay \(p. 900\)](#) stored procedure.

For example, run the following stored procedure to specify that replication to the Read Replica is delayed by at least one hour (3600 seconds).

```
call mysql.rds_set_source_delay(3600);
```

4. Use the [mysql.rds_start_replication \(p. 901\)](#) stored procedure to start replication.

Setting a Location to Stop Replication to a Read Replica

After stopping replication to the Read Replica, you can start replication and then stop it at a specified binary log file location using the [mysql.rds_start_replication_until \(p. 901\)](#) stored procedure.

To start replication to a Read Replica and stop replication at a specific location

1. Using a MySQL client, connect to the source MySQL DB instance as the master user.
2. Run the [mysql.rds_start_replication_until \(p. 901\)](#) stored procedure.

The following example initiates replication and replicates changes until it reaches location 120 in the `mysql-bin-changelog.000777` binary log file. In a disaster recovery scenario, assume that location 120 is just before the disaster.

```
call mysql.rds_start_replication_until(
    'mysql-bin-changelog.000777',
    120);
```

Replication stops automatically when the stop point is reached. The following RDS event is generated:
Replication has been stopped since the replica reached the stop point specified by the `rds_start_replication_until` stored procedure.

After replication is stopped, in a disaster recovery scenario, you can [Promoting a Read Replica to Be a Standalone DB Instance \(p. 186\)](#) promote the Read Replica to be the new source DB instance. For information about promoting the Read Replica, see [Promoting a Read Replica to Be a Standalone DB Instance \(p. 186\)](#).

Read Replica Updates with MySQL

Read Replicas are designed to support read queries, but you might need occasional updates. For example, you might need to add an index to optimize the specific types of queries accessing the replica. You can enable updates by setting the `read_only` parameter to 0 in the DB parameter group for the Read Replica. Be careful when disabling read-only on a Read Replica because it can cause problems if the Read Replica becomes incompatible with the source DB instance. Change the value of the `read_only` parameter back to 1 as soon as possible.

Multi-AZ Read Replica Deployments with MySQL

You can create a Read Replica from either single-AZ or Multi-AZ DB instance deployments. You use Multi-AZ deployments to improve the durability and availability of critical data, but you can't use the Multi-AZ secondary to serve read-only queries. Instead, you can create Read Replicas from high-traffic Multi-AZ DB instances to offload read-only queries. If the source instance of a Multi-AZ deployment fails over to

the secondary, any associated Read Replicas automatically switch to use the secondary (now primary) as their replication source. For more information, see [High Availability \(Multi-AZ\) for Amazon RDS \(p. 144\)](#).

You can create a Read Replica as a Multi-AZ DB instance. Amazon RDS creates a standby of your replica in another Availability Zone for failover support for the replica. Creating your Read Replica as a Multi-AZ DB instance is independent of whether the source database is a Multi-AZ DB instance.

Note

To create a Read Replica as a Multi-AZ DB instance, the DB instance must be MySQL 5.6 or later.

Monitoring MySQL Read Replicas

For MySQL Read Replicas, you can monitor replication lag in Amazon CloudWatch by viewing the Amazon RDS `ReplicaLag` metric. The `ReplicaLag` metric reports the value of the `Seconds_Behind_Master` field of the `SHOW SLAVE STATUS` command.

Common causes for replication lag for MySQL are the following:

- A network outage.
- Writing to tables that have different indexes on a Read Replica. If the `read_only` parameter is set to 0 on the Read Replica, replication can break if the Read Replica becomes incompatible with the source DB instance. After you've performed maintenance tasks on the Read Replica, we recommend that you set the `read_only` parameter back to 1.
- Using a nontransactional storage engine such as MyISAM. Replication is only supported for the InnoDB storage engine on MySQL.

When the `ReplicaLag` metric reaches 0, the replica has caught up to the source DB instance. If the `ReplicaLag` metric returns -1, then replication is currently not active. `ReplicaLag = -1` is equivalent to `Seconds_Behind_Master = NULL`.

Starting and Stopping Replication with MySQL Read Replicas

You can stop and restart the replication process on an Amazon RDS DB instance by calling the system stored procedures [mysql.rds_stop_replication \(p. 903\)](#) and [mysql.rds_start_replication \(p. 901\)](#). You can do this when replicating between two Amazon RDS instances for long-running operations such as creating large indexes. You also need to stop and start replication when importing or exporting databases. For more information, see [Importing Data to an Amazon RDS MySQL or MariaDB DB Instance with Reduced Downtime \(p. 835\)](#) and [Exporting Data from a MySQL DB Instance by Using Replication \(p. 869\)](#).

If replication is stopped for more than 30 consecutive days, either manually or due to a replication error, Amazon RDS terminates replication between the source DB instance and all Read Replicas. It does so to prevent increased storage requirements on the source DB instance and long failover times. The Read Replica DB instance is still available. However, replication can't be resumed because the binary logs required by the Read Replica are deleted from the source DB instance after replication is terminated. You can create a new Read Replica for the source DB instance to reestablish replication.

Deleting Read Replicas with MySQL

You must explicitly delete Read Replicas, using the same mechanisms for deleting a DB instance. If you delete the source DB instance without deleting the replicas, each replica is promoted to a standalone DB instance.

Troubleshooting a MySQL Read Replica Problem

For MySQL DB instances, in some cases Read Replicas present replication errors or data inconsistencies between the Read Replica and its source DB instance. This problem occurs when some binlog

events or InnoDB redo logs aren't flushed during a failure of the Read Replica or the source DB instance. In these cases, you must manually delete and recreate the Read Replicas. You can reduce the chance of this happening by setting the following dynamic variable values: `sync_binlog=1`, `innodb_flush_log_at_trx_commit=1`, and `innodb_support_xa=1`. These settings might reduce performance, so test their impact before implementing the changes in a production environment. For MySQL 5.5, `sync_binlog` defaults to 0, but in MySQL 5.6 and later, problems are less likely to occur because these parameters are all set to the recommended values by default.

The replication technologies for MySQL are asynchronous. Because they are asynchronous, occasional `BinLogDiskUsage` increases on the source DB instance and `ReplicaLag` on the Read Replica are to be expected. For example, a high volume of write operations to the source DB instance can occur in parallel. In contrast, write operations to the Read Replica are serialized using a single I/O thread, which can lead to a lag between the source instance and Read Replica. For more information about read-only replicas in the MySQL documentation, see [Replication Implementation Details](#).

You can do several things to reduce the lag between updates to a source DB instance and the subsequent updates to the Read Replica, such as the following:

- Sizing a Read Replica to have a storage size and DB instance class comparable to the source DB instance.
- Ensuring that parameter settings in the DB parameter groups used by the source DB instance and the Read Replica are compatible. For more information and an example, see the discussion of the `max_allowed_packet` parameter later in this section.

Amazon RDS monitors the replication status of your Read Replicas and updates the `Replication State` field of the Read Replica instance to `Error` if replication stops for any reason. An example might be if DML queries run on your Read Replica conflict with the updates made on the source DB instance.

You can review the details of the associated error thrown by the MySQL engine by viewing the `Replication Error` field. Events that indicate the status of the Read Replica are also generated, including [RDS-EVENT-0045 \(p. 375\)](#), [RDS-EVENT-0046 \(p. 375\)](#), and [RDS-EVENT-0047 \(p. 373\)](#). For more information about events and subscribing to events, see [Using Amazon RDS Event Notification \(p. 370\)](#). If a MySQL error message is returned, review the error number in the [MySQL error message documentation](#).

One common issue that can cause replication errors is when the value for the `max_allowed_packet` parameter for a Read Replica is less than the `max_allowed_packet` parameter for the source DB instance. The `max_allowed_packet` parameter is a custom parameter that you can set in a DB parameter group that is used to specify the maximum size of DML code that can be executed on the database. In some cases, the `max_allowed_packet` parameter value in the DB parameter group associated with a source DB instance is smaller than the `max_allowed_packet` parameter value in the DB parameter group associated with the source's Read Replica. In these cases, the replication process can throw an error (Packet bigger than 'max_allowed_packet' bytes) and stop replication. You can fix the error by having the source and Read Replica use DB parameter groups with the same `max_allowed_packet` parameter values.

Other common situations that can cause replication errors include the following:

- Writing to tables on a Read Replica. If you are creating indexes on a Read Replica that are different from the indexes on the source DB instance, you must have the `read_only` parameter set to 0 to create the indexes. If you are writing to tables on the Read Replica, it might break replication if the read replica becomes incompatible with the source DB instance. After you've performed maintenance tasks on the Read Replica, we recommend that you set the `read_only` parameter back to 1.
- Using a non-transactional storage engine such as MyISAM. Read replicas require a transactional storage engine. Replication is only supported for the InnoDB storage engine on MySQL.
- Using unsafe nondeterministic queries such as `SYSDATE()`. For more information, see [Determination of Safe and Unsafe Statements in Binary Logging](#).

If you decide that you can safely skip an error, you can follow the steps described in the section [Skipping the Current Replication Error \(p. 881\)](#). Otherwise, you can delete the Read Replica and create an instance using the same DB instance identifier so that the endpoint remains the same as that of your old Read Replica. If a replication error is fixed, the `Replication State` changes to *replicating*.

Using GTID-Based Replication for Amazon RDS MySQL

Following, you can learn how to use global transaction identifiers (GTIDs) with binary log (binlog) replication among Amazon RDS MySQL DB instances.

If you use binlog replication and aren't familiar with GTID-based replication with MySQL, see [Replication with Global Transaction Identifiers](#) in the MySQL documentation for background.

Note

GTID-based replication is supported for RDS MySQL version 5.7.23 and later MySQL 5.7 versions. All RDS MySQL DB instances in a replication configuration must meet this requirement. GTID-based replication isn't supported for RDS MySQL 5.5, 5.6, or 8.0.

Topics

- [Overview of Global Transaction Identifiers \(GTIDs\) \(p. 858\)](#)
- [Parameters for GTID-Based Replication \(p. 858\)](#)
- [Configuring GTID-Based Replication for New Read Replicas \(p. 859\)](#)
- [Configuring GTID-Based Replication for Existing Read Replicas \(p. 860\)](#)
- [Disabling GTID-Based Replication for an RDS MySQL DB Instance with Read Replicas \(p. 861\)](#)

Note

For information about configuring GTID-based replication with an external database, see [Replication with a MySQL or MariaDB Instance Running External to Amazon RDS \(p. 862\)](#).

Overview of Global Transaction Identifiers (GTIDs)

Global transaction identifiers (GTIDs) are unique identifiers generated for committed MySQL transactions. You can use GTIDs to make binlog replication simpler and easier to troubleshoot.

MySQL uses two different types of transactions for binlog replication:

- *GTID transactions* – Transactions that are identified by a GTID.
- *Anonymous transactions* – Transactions that don't have a GTID assigned.

In a replication configuration, GTIDs are unique across all DB instances. GTIDs simplify replication configuration because when you use them, you don't have to refer to log file positions. GTIDs also make it easier to track replicated transactions and determine whether masters and replicas are consistent.

You can use GTID-based replication to replicate data with Amazon RDS MySQL Read Replicas or with an external MySQL database. For RDS MySQL Read Replicas, you can configure GTID-based replication when you are creating new Read Replicas, or you can convert existing Read Replicas to use GTID-based replication.

You can also use GTID-based replication in a delayed replication configuration with RDS MySQL . For more information, see [Configuring Delayed Replication with MySQL \(p. 853\)](#).

Parameters for GTID-Based Replication

Use the following parameters to configure GTID-based replication.

Parameter	Valid Values	Description
gtid_mode	OFF, OFF_PERMISSIVE, ON_PERMISSIVE, ON	<p>OFF specifies that new transactions are anonymous transactions (that is, don't have GTIDs), and a transaction must be anonymous to be replicated.</p> <p>OFF_PERMISSIVE specifies that new transactions are anonymous transactions, but all transactions can be replicated.</p> <p>ON_PERMISSIVE specifies that new transactions are GTID transactions, but all transactions can be replicated.</p> <p>ON specifies that new transactions are GTID transactions, and a transaction must be a GTID transaction to be replicated.</p>
enforce_gtid_consistency	OFF, ON, WARN	<p>OFF allows transactions to violate GTID consistency.</p> <p>ON prevents transactions from violating GTID consistency.</p> <p>WARN allows transactions to violate GTID consistency but generates a warning when a violation occurs.</p>

For GTID-based replication, use these settings for the parameter group for your DB instance or Read Replica:

- ON and ON_PERMISSIVE apply only to outgoing replication from an RDS DB instance or Aurora MySQL cluster. Both of these values cause your RDS DB instance or Aurora DB cluster to use GTIDs for transactions that are replicated to an external database. ON requires that the external database also use GTID-based replication. ON_PERMISSIVE makes GTID-based replication optional on the external database.
- OFF_PERMISSIVE, if set, means that your RDS DB instances or Aurora DB cluster can accept incoming replication from an external database. It can do this whether the external database uses GTID-based replication or not.
- OFF, if set, means that your RDS DB instances or Aurora DB cluster only accept incoming replication from external databases that don't use GTID-based replication.

For more information about parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#).

Note

In the AWS Management Console, the gtid_mode parameter appears as gtid-mode.

Configuring GTID-Based Replication for New Read Replicas

When GTID-based replication is enabled for an RDS MySQL DB instance, GTID-based replication is configured automatically for Read Replicas of the DB instance.

To enable GTID-based replication for new Read Replicas

1. Make sure that the parameter group associated with the DB instance has the following parameter settings:

- `gtid_mode` – ON or ON_PERMISSIVE
- `enforce_gtid_consistency` – ON

For more information about setting configuration parameters using parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#).

2. If you changed the parameter group of the DB instance, reboot the DB instance. For more information on how to do so, see [Rebooting a DB Instance \(p. 166\)](#).
3. Create one or more Read Replicas of the DB instance. For more information on how to do so, see [Creating a Read Replica \(p. 184\)](#).

Amazon RDS attempts to establish GTID-based replication between the MySQL DB instance and the Read Replicas using the `MASTER_AUTO_POSITION`. If the attempt fails, Amazon RDS uses log file positions for replication with the Read Replicas. For more information about the `MASTER_AUTO_POSITION`, see [GTID Auto-Positioning](#) in the MySQL documentation.

Configuring GTID-Based Replication for Existing Read Replicas

For an existing RDS MySQL DB instance with Read Replicas that doesn't use GTID-based replication, you can configure GTID-based replication between the DB instance and the Read Replicas.

To enable GTID-based replication for existing Read Replicas

1. If the DB instance or any Read Replica is using RDS MySQL version 5.7.22 or lower, upgrade the DB instance or Read Replica. Upgrade to RDS MySQL version 5.7.23 or a later MySQL 5.7 version.

For more information, see [Upgrading the MySQL DB Engine \(p. 812\)](#).

2. (Optional) Reset the GTID parameters and test the behavior of the DB instance and Read Replicas:
 - a. Make sure that the parameter group associated with the DB instance and each Read Replica has the `enforce_gtid_consistency` parameter set to `WARN`.
For more information about setting configuration parameters using parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#).
 - b. If you changed the parameter group of the DB instance, reboot the DB instance. If you changed the parameter group for a Read Replica, reboot the Read Replica.
For more information, see [Rebooting a DB Instance \(p. 166\)](#).
 - c. Run your DB instance and Read Replicas with your normal workload and monitor the log files.

If you see warnings about GTID-incompatible transactions, adjust your application so that it only uses GTID-compatible features. Make sure that the DB instance is not generating any warnings about GTID-incompatible transactions before proceeding to the next step.

3. Reset the GTID parameters for GTID-based replication that allows anonymous transactions until the Read Replicas have processed all of them.
 - a. Make sure that the parameter group associated with the DB instance and each Read Replica has the following parameter settings:
 - `gtid_mode` – ON_PERMISSIVE
 - `enforce_gtid_consistency` – ON
 - b. If you changed the parameter group of the DB instance, reboot the DB instance. If you changed the parameter group for a Read Replica, reboot the Read Replica.
4. Wait for all of your anonymous transactions to be replicated. To check that these are replicated, do the following:

- a. Run the following statement on your primary DB instance.

```
SHOW MASTER STATUS;
```

Note the values in the `File` and `Position` columns.

- b. On each Read Replica, use the file and position information from its master in the previous step to run the following query.

```
SELECT MASTER_POS_WAIT(file, position);
```

For example, if the file name is `mysql-bin-changelog.000031` and the position is 107, run the following statement.

```
SELECT MASTER_POS_WAIT(mysql-bin-changelog.000031, 107);
```

If the Read Replica is past the specified position, the query returns immediately. Otherwise, the function waits. Proceed to the next step when the query returns for all Read Replicas.

5. Reset the GTID parameters for GTID-based replication only.

- a. Make sure that the parameter group associated with the DB instance and each Read Replica has the following parameter settings:

- `gtid_mode` – ON
- `enforce_gtid_consistency` – ON

- b. Reboot the DB instance and each Read Replica.

6. On each Read Replica, run the following procedure.

```
CALL mysql.rds_set_master_auto_position(1);
```

Disabling GTID-Based Replication for an RDS MySQL DB Instance with Read Replicas

You can disable GTID-based replication for an RDS MySQL DB instance with Read Replicas.

To disable GTID-based replication for an RDS MySQL DB instance with Read Replicas

1. On each Read Replica, run the following procedure.

```
CALL mysql.rds_set_master_auto_position(0);
```

2. Reset the `gtid_mode` to `ON_PERMISSIVE`.

- a. Make sure that the parameter group associated with the RDS MySQL DB instance and each Read Replica has `gtid_mode` set to `ON_PERMISSIVE`.

For more information about setting configuration parameters using parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#).

- b. Reboot the RDS MySQL DB instance and each Read Replica. For more information about rebooting, see [Rebooting a DB Instance \(p. 166\)](#).
3. Reset the `gtid_mode` to `OFF_PERMISSIVE`:
 - a. Make sure that the parameter group associated with the RDS MySQL DB instance and each Read Replica has `gtid_mode` set to `OFF_PERMISSIVE`.
 - b. Reboot the RDS MySQL DB instance and each Read Replica.
4. Wait for all of the GTID transactions to be applied on all of the Read Replicas. To check that these are applied, do the following:

Wait for all of the GTID transactions to be applied on the Aurora primary instance. To check that these are applied, do the following:

 - a. On the RDS MySQL DB instance, run the `SHOW MASTER STATUS` command.

Your output should be similar to the following.

File	Position
mysql-bin-changelog.000031	107

Note the file and position in your output.

- b. On each Read Replica, use the file and position information from its master in the previous step to run the following query.

```
SELECT MASTER_POS_WAIT(file, position);
```

For example, if the file name is `mysql-bin-changelog.000031` and the position is `107`, run the following statement.

```
SELECT MASTER_POS_WAIT(mysql-bin-changelog.000031, 107);
```

If the Read Replica is past the specified position, the query returns immediately. Otherwise, the function waits. When the query returns for all Read Replicas, go to the next step.

5. Reset the GTID parameters to disable GTID-based replication:
 - a. Make sure that the parameter group associated with the RDS MySQL DB instance and each Read Replica has the following parameter settings:
 - `gtid_mode` – OFF
 - `enforce_gtid_consistency` – OFF
 - b. Reboot the RDS MySQL DB instance and each Read Replica.

Replication with a MySQL or MariaDB Instance Running External to Amazon RDS

You can set up replication between an Amazon RDS MySQL or MariaDB DB instance and a MySQL or MariaDB instance that is external to Amazon RDS.

Topics

- [Before You Begin \(p. 863\)](#)
- [Configuring Binary Log File Position Replication with an External Master Instance \(p. 863\)](#)
- [Configuring GTID-Based Replication with an External Master Instance \(p. 866\)](#)

Before You Begin

You can configure replication using the binary log file position of replicated transactions. On Amazon RDS MySQL 5.7.23 and later MySQL 5.7 versions, you can also configure replication using global transaction identifiers (GTIDs).

The permissions required to start replication on an Amazon RDS DB instance are restricted and not available to your Amazon RDS master user. Because of this, you must use the Amazon RDS [mysql.rds_set_external_master \(p. 891\)](#) and [mysql.rds_start_replication \(p. 901\)](#) commands to set up replication between your live database and your Amazon RDS database.

To set the binary logging format for a MySQL or MariaDB database, update the `binlog_format` parameter. If your DB instance uses the default DB instance parameter group, create a new DB parameter group to modify `binlog_format` settings. We recommend that you use the default setting for `binlog_format`, which is `MIXED`. However, you can also set `binlog_format` to `ROW` or `STATEMENT` if you need a specific binlog format. Reboot your DB instance for the change to take effect.

For information about setting the `binlog_format` parameter, see [Binary Logging Format \(p. 410\)](#). For information about the implications of different MySQL replication types, see [Advantages and Disadvantages of Statement-Based and Row-Based Replication](#) in the MySQL documentation.

Note

Use the procedure in this topic to configure replication in all cases except when the external instance is MariaDB version 10.0.2 or greater and the Amazon RDS instance is MariaDB. In that case, use the procedure at [Configuring GTID-Based Replication into an Amazon RDS MariaDB DB instance \(p. 582\)](#) to set up GTID-based replication.

Configuring Binary Log File Position Replication with an External Master Instance

Follow these guidelines when you set up an external replication master and a replica on Amazon RDS:

- Monitor failover events for the Amazon RDS DB instance that is your replica. If a failover occurs, then the DB instance that is your replica might be recreated on a new host with a different network address. For information on how to monitor failover events, see [Using Amazon RDS Event Notification \(p. 370\)](#).
- Maintain the binary logs (binlogs) on your master instance until you have verified that they have been applied to the replica. This maintenance makes sure that you can restore your master instance in the event of a failure.
- Turn on automated backups on your Amazon RDS DB instance. Turning on automated backups makes sure that you can restore your replica to a particular point in time if you need to re-synchronize your master and replica. For information on backups and point-in-time restore, see [Backing Up and Restoring Amazon RDS DB Instances \(p. 253\)](#).

To configure binary log file replication with an external master instance

1. Make the source MySQL or MariaDB instance read-only.

```
mysql> FLUSH TABLES WITH READ LOCK;
mysql> SET GLOBAL read_only = ON;
```

- Run the `SHOW MASTER STATUS` command on the source MySQL or MariaDB instance to determine the binlog location.

You receive output similar to the following example.

```
File          Position
-----
mysql-bin-changelog.000031      107
-----
```

- Copy the database from the external instance to the Amazon RDS DB instance using `mysqldump`. For very large databases, you might want to use the procedure in [Importing Data to an Amazon RDS MySQL or MariaDB DB Instance with Reduced Downtime \(p. 835\)](#).

For Linux, OS X, or Unix:

```
mysqldump --databases <database_name> \
    --single-transaction \
    --compress \
    --order-by-primary \
    -u <local_user> \
    -p<local_password> | mysql \
        --host=hostname \
        --port=3306 \
        -u <RDS_user_name> \
        -p<RDS_password>
```

For Windows:

```
mysqldump --databases <database_name> ^
    --single-transaction ^
    --compress ^
    --order-by-primary ^
    -u <local_user> ^
    -p<local_password> | mysql ^
        --host=hostname ^
        --port=3306 ^
        -u <RDS_user_name> ^
        -p<RDS_password>
```

Note

Make sure that there isn't a space between the `-p` option and the entered password.

To specify the host name, user name, port, and password to connect to your Amazon RDS DB instance, use the `--host`, `--user` (`-u`), `--port` and `-p` options in the `mysql` command. The host name is the Domain Name Service (DNS) name from the Amazon RDS DB instance endpoint, for example, `myinstance.123456789012.us-east-1.rds.amazonaws.com`. You can find the endpoint value in the instance details in the AWS Management Console.

- Make the source MySQL or MariaDB instance writeable again.

```
mysql> SET GLOBAL read_only = OFF;
mysql> UNLOCK TABLES;
```

For more information on making backups for use with replication, see [Backing Up a Master or Slave by Making It Read Only](#) in the MySQL documentation.

- In the AWS Management Console, add the IP address of the server that hosts the external database to the VPC security group for the Amazon RDS DB instance. For more information on modifying a VPC security group, see [Security Groups for Your VPC](#) in the *Amazon Virtual Private Cloud User Guide*.

The IP address can change when the following conditions are met:

- You are using a public IP address for communication between the external master instance and the DB instance.
- The external master instance was stopped and restarted.

If these conditions are met, verify the IP address before adding it.

You might also need to configure your local network to permit connections from the IP address of your Amazon RDS DB instance, so that it can communicate with your external MySQL or MariaDB instance. To find the IP address of the Amazon RDS DB instance, use the host command.

```
host <RDS_MySQL_DB_host_name>
```

The host name is the DNS name from the Amazon RDS DB instance endpoint.

6. Using the client of your choice, connect to the external instance and create a user to use for replication. Use this account solely for replication. and restrict it to your domain to improve security. The following is an example.

```
CREATE USER 'repl_user'@'mydomain.com' IDENTIFIED BY '<password>;'
```

7. For the external instance, grant REPLICATION CLIENT and REPLICATION SLAVE privileges to your replication user. For example, to grant the REPLICATION CLIENT and REPLICATION SLAVE privileges on all databases for the 'repl_user' user for your domain, issue the following command.

```
GRANT REPLICATION CLIENT, REPLICATION SLAVE ON *.* TO 'repl_user'@'mydomain.com'  
IDENTIFIED BY '<password>;'
```

8. Make the Amazon RDS DB instance the replica. To do so, connect to the Amazon RDS DB instance as the master user and identify the external MySQL or MariaDB database as the replication master by using the [mysql.rds_set_external_master \(p. 891\)](#) command. Use the master log file name and master log position that you determined in step 2. The following is an example.

```
CALL mysql.rds_set_external_master ('mymasterserver.mydomain.com', 3306, 'repl_user',  
'<password>', 'mysql-bin-changelog.000031', 107, 0);
```

Note

On Amazon RDS MySQL, you can choose to use delayed replication by running the [mysql.rds_set_external_master_with_delay \(p. 893\)](#) stored procedure instead.

One reason to use delayed replication is to enable disaster recovery with the [mysql.rds_start_replication_until \(p. 901\)](#) stored procedure. Currently, delayed replication is not supported on Amazon RDS MariaDB.

9. On the Amazon RDS DB instance, issue the [mysql.rds_start_replication \(p. 901\)](#) command to start replication:

```
CALL mysql.rds_start_replication;
```

Configuring GTID-Based Replication with an External Master Instance

When you set up an external replication master and a replica on Amazon RDS, monitor failover events for the Amazon RDS DB instance that is your replica. If a failover occurs, then the DB instance that is your replica might be recreated on a new host with a different network address. For information on how to monitor failover events, see [Using Amazon RDS Event Notification \(p. 370\)](#).

Important

GTID-based replication is only supported on Amazon RDS MySQL version 5.7.23 and later MySQL 5.7 versions. GTID-based replication is not supported for Amazon RDS MySQL 5.5, 5.6, or 8.0.

To configure GTID-based replication with an external master instance

1. Prepare for GTID-based replication:

- a. Make sure that the external MySQL or MariaDB database has GTID-based replication enabled. To do so, make sure that the external database has the following parameters set to the specified values:

```
gtid_mode = ON  
enforce_gtid_consistency = ON
```

For more information, see [Replication with Global Transaction Identifiers](#) in the MySQL documentation or [Global Transaction ID](#) in the MariaDB documentation.

- b. Make sure that the parameter group associated with the DB instance has the following parameter settings:

- gtid_mode = ON, ON_PERMISSIVE, or OFF_PERMISSIVE
- enforce_gtid_consistency = ON

For more information about parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#).

- c. If you changed the parameter group of the DB instance, reboot the DB instance. For more information, see [Rebooting a DB Instance \(p. 166\)](#).

2. Make the source MySQL or MariaDB instance read-only.

```
mysql> FLUSH TABLES WITH READ LOCK;  
mysql> SET GLOBAL read_only = ON;
```

3. Copy the database from the external instance to the Amazon RDS DB instance using mysqldump. For very large databases, you might want to use the procedure in [Importing Data to an Amazon RDS MySQL or MariaDB DB Instance with Reduced Downtime \(p. 835\)](#).

For Linux, OS X, or Unix:

```
mysqldump --databases <database_name> \  
--single-transaction \  
--compress \  
--order-by-primary \  
-u <local_user> \  
-p<local_password> | mysql \  
--host=hostname \  
--port=3306 \  
-u <RDS_user_name> \  
-p<RDS_password>
```

For Windows:

```
mysqldump --databases <database_name> ^
    --single-transaction ^
    --compress ^
    --order-by-primary ^
    -u <local_user> ^
    -p<local_password> | mysql ^
        --host=hostname ^
        --port=3306 ^
        -u <RDS_user_name> ^
        -p<RDS_password>
```

Note

Make sure that there is not a space between the `-p` option and the entered password.

To specify the host name, user name, port, and password to connect to your Amazon RDS DB instance, use the `--host`, `--user` (`-u`), `--port` and `-p` options in the `mysql` command. The host name is the DNS name from the Amazon RDS DB instance endpoint, for example, `myinstance.123456789012.us-east-1.rds.amazonaws.com`. You can find the endpoint value in the instance details in the AWS Management Console.

4. Make the source MySQL or MariaDB instance writeable again.

```
mysql> SET GLOBAL read_only = OFF;
mysql> UNLOCK TABLES;
```

For more information on making backups for use with replication, see [Backing Up a Master or Slave by Making It Read Only](#) in the MySQL documentation.

5. In the AWS Management Console, add the IP address of the server that hosts the external database to the VPC security group for the Amazon RDS DB instance. For more information on modifying a VPC security group, see [Security Groups for Your VPC](#) in the *Amazon Virtual Private Cloud User Guide*.

The IP address can change when the following conditions are met:

- You are using a public IP address for communication between the external master instance and the DB instance.
- The external master instance was stopped and restarted.

If these conditions are met, verify the IP address before adding it.

You might also need to configure your local network to permit connections from the IP address of your Amazon RDS DB instance, so that it can communicate with your external MySQL or MariaDB instance. To find the IP address of the Amazon RDS DB instance, use the `host` command.

```
host <RDS_SQL_DB_host_name>
```

The host name is the DNS name from the Amazon RDS DB instance endpoint.

6. Using the client of your choice, connect to the external instance and create a user to use for replication. Use this account solely for replication. and restrict it to your domain to improve security. The following is an example.

```
CREATE USER 'repl_user'@'mydomain.com' IDENTIFIED BY '<password>';
```

7. For the external instance, grant REPLICATION CLIENT and REPLICATION SLAVE privileges to your replication user. For example, to grant the REPLICATION CLIENT and REPLICATION SLAVE privileges on all databases for the 'repl_user' user for your domain, issue the following command.

```
GRANT REPLICATION CLIENT, REPLICATION SLAVE ON *.* TO 'repl_user'@'mydomain.com'  
IDENTIFIED BY '<password>';
```

8. Make the Amazon RDS DB instance the replica. To do so, connect to the Amazon RDS DB instance as the master user and identify the external MySQL or MariaDB database as the replication master by using the [mysql.rds_set_external_master_with_auto_position \(p. 895\)](#) command. The following is an example.

```
CALL mysql.rds_set_external_master_with_auto_position ('mymasterserver.mydomain.com',  
3306, 'repl_user', '<password>', 0, 0);
```

Note

On Amazon RDS MySQL, you can choose to use delayed replication by running the [mysql.rds_set_external_master_with_delay \(p. 893\)](#) stored procedure instead.

One reason to use delayed replication is to enable disaster recovery with the [mysql.rds_start_replication_until_gtid \(p. 902\)](#) stored procedure. Currently, delayed replication is not supported on Amazon RDS MariaDB.

9. On the Amazon RDS DB instance, issue the [mysql.rds_start_replication \(p. 901\)](#) command to start replication.

```
CALL mysql.rds_start_replication;
```

Exporting Data from a MySQL DB Instance by Using Replication

You can use replication to export data from a MySQL 5.6 or later DB instance to a MySQL instance running external to Amazon RDS. The MySQL instance external to Amazon RDS can be running either on-premises in your data center, or on an Amazon EC2 instance. The MySQL DB instance must be running version 5.6.13 or later. The MySQL instance external to Amazon RDS must be running the same version as the Amazon RDS instance, or a later version.

Replication to an instance of MySQL running external to Amazon RDS is only supported during the time it takes to export a database from a MySQL DB instance. The replication should be terminated when the data has been exported and applications can start accessing the external instance.

The following list shows the steps to take. Each step is discussed in more detail in later sections.

1. Prepare an instance of MySQL running external to Amazon RDS.
2. Configure the MySQL DB instance to be the replication source.
3. Use `mysqldump` to transfer the database from the Amazon RDS instance to the instance external to Amazon RDS.
4. Start replication to the instance running external to Amazon RDS.
5. After the export completes, stop replication.

Prepare an Instance of MySQL External to Amazon RDS

Install an instance of MySQL external to Amazon RDS.

Connect to the instance as the master user, and create the users required to support the administrators, applications, and services that access the instance.

Follow the directions in the MySQL documentation to prepare the instance of MySQL running external to Amazon RDS as a replica. For more information, see [Setting the Replication Slave Configuration](#).

Configure an egress rule for the external instance to operate as a Read Replica during the export. The egress rule will allow the MySQL Read Replica to connect to the MySQL DB instance during replication. Specify an egress rule that allows TCP connections to the port and IP address of the source Amazon RDS MySQL DB instance.

If the Read Replica is running in an Amazon EC2 instance in an Amazon VPC, specify the egress rules in a VPC security group. If the Read Replica is running in an Amazon EC2 instance that is not in a VPC, specify the egress rule in an Amazon EC2 security group. If the Read Replica is installed on-premises, specify the egress rule in a firewall.

If the Read Replica is running in a VPC, configure VPC ACL rules in addition to the security group egress rule. For more information about Amazon VPC network ACLs, see [Network ACLs](#).

- ACL ingress rule allowing TCP traffic to ports 1024-65535 from the IP address of the source MySQL DB instance.
- ACL egress rule: allowing outbound TCP traffic to the port and IP address of the source MySQL DB instance.

We recommend setting the `max_allowed_packet` parameter to the maximum size to avoid replication errors.

Prepare the Replication Source

Prepare the MySQL DB instance as the replication source.

Ensure your client computer has enough disk space available to save the binary logs while setting up replication.

Create a replication account by following the directions in [Creating a User For Replication](#).

Configure ingress rules on the system running the replication source MySQL DB instance that will allow the external MySQL Read Replica to connect during replication. Specify an ingress rule that allows TCP connections to the port used by the Amazon RDS instance from the IP address of the MySQL Read Replica running external to Amazon RDS.

If the Amazon RDS instance is running in a VPC, specify the ingress rules in a VPC security group. If the Amazon RDS instance is not running in an in a VPC, specify the ingress rules in a database security group.

If the Amazon RDS instance is running in a VPC, configure VPC ACL rules in addition to the security group ingress rule. For more information about Amazon VPC network ACLs, see [Network ACLs](#).

- ACL ingress rule: allow TCP connections to the port used by the Amazon RDS instance from the IP address of the external MySQL Read Replica.
- ACL egress rule: allow TCP connections from ports 1024-65535 to the IP address of the external MySQL Read Replica.

Ensure that the backup retention period is set long enough that no binary logs are purged during the export. If any of the logs are purged before the export is complete, you must restart replication from the beginning. For more information about setting the backup retention period, see [Working With Backups \(p. 254\)](#).

Use the `mysql.rds_set_configuration` stored procedure to set the binary log retention period long enough that the binary logs are not purged during the export. For more information, see [Accessing MySQL Binary Logs \(p. 411\)](#).

To further ensure that the binary logs of the source instance are not purged, create an Amazon RDS Read Replica from the source instance. For more information, see [Creating a Read Replica \(p. 184\)](#). After the Amazon RDS Read Replica has been created, call the `mysql.rds_stop_replication` stored procedure to stop the replication process. The source instance will no longer purge its binary log files, so they will be available for the replication process.

We recommend setting both the `max_allowed_packet` parameter and the `slave_max_allowed_packet` parameter to the maximum size to avoid replication errors. The maximum size for both parameters is 1 GB. For information about setting parameters, see [Modifying Parameters in a DB Parameter Group \(p. 213\)](#).

Copy the Database

Run the MySQL `SHOW SLAVE STATUS` statement on the RDS read replica, and note the values for the following:

- `master_host`
- `master_port`
- `master_log_file`
- `exec_master_log_pos`

Use the `mysqldump` utility to create a snapshot, which copies the data from Amazon RDS to your local client computer. Then run another utility to load the data into the MySQL instance running external to RDS. Ensure your client computer has enough space to hold the `mysqldump` files from the databases to be replicated. This process can take several hours for very large databases. Follow the directions in [Creating a Dump Snapshot Using mysqldump](#).

The following example shows how to run `mysqldump` on a client, and then pipe the dump into the `mysql` client utility, which loads the data into the external MySQL instance.

For Linux, OS X, or Unix:

```
mysqldump -h RDS instance endpoint \
-u user \
-p password \
--port=3306 \
--single-transaction \
--routines \
--triggers \
--databases database database2 \
--compress \
--port 3306
```

For Windows:

```
mysqldump -h RDS instance endpoint ^
-u user ^
-p password ^
--port=3306 ^
--single-transaction ^
--routines ^
--triggers ^
--databases database database2 ^
--compress ^
--port 3306
```

The following example shows how to run `mysqldump` on a client and write the dump to a file.

For Linux, OS X, or Unix:

```
mysqldump -h RDS instance endpoint \
-u user \
-p password \
--port=3306 \
--single-transaction \
--routines \
--triggers \
--databases database database2 > path/rds-dump.sql
```

For Windows:

```
mysqldump -h RDS instance endpoint ^
-u user ^
-p password ^
--port=3306 ^
--single-transaction ^
--routines ^
--triggers ^
--databases database database2 > path\rds-dump.sql
```

Complete the Export

After you have loaded the `mysqldump` files to create the databases on the MySQL instance running external to Amazon RDS, start replication from the source MySQL DB instance to export all source changes that have occurred after you stopped replication from the Amazon RDS Read Replica.

Use the MySQL `CHANGE MASTER` statement to configure the external MySQL instance. Specify the ID and password of the user granted `REPLICATION SLAVE` permissions. Specify the `master_host`, `master_port`, `relay_master_log_file` and `exec_master_log_pos` values you got from the `Mysql SHOW SLAVE STATUS` statement you ran on the RDS Read Replica. For more information, see [Setting the Master Configuration on the Slave](#).

Use the MySQL `START SLAVE` command to initiate replication from the source MySQL DB instance and the MySQL replica.

Run the MySQL `SHOW SLAVE STATUS` command on the Amazon RDS instance to verify that it is operating as a Read Replica. For more information about interpreting the results, see [SHOW SLAVE STATUS Syntax](#).

After replication on the MySQL instance has caught up with the Amazon RDS source, use the MySQL `STOP SLAVE` command to terminate replication from the source MySQL DB instance.

On the Amazon RDS Read Replica, call the `mysql.rds_start_replication` stored procedure. This will allow Amazon RDS to start purging the binary log files from the source MySQL DB instance.

Related Topics

- [Restoring a Backup into an Amazon RDS MySQL DB Instance \(p. 826\)](#)
- [Backing Up and Restoring Amazon RDS DB Instances \(p. 253\)](#)

Options for MySQL DB Instances

This appendix describes options, or additional features, that are available for Amazon RDS instances running the MySQL DB engine. To enable these options, you can add them to a custom option group, and then associate the option group with your DB instance. For more information about working with option groups, see [Working with Option Groups \(p. 196\)](#).

Amazon RDS supports the following options for MySQL:

Option	Option ID	Engine Versions
MariaDB Audit Plugin Support (p. 874)	MARIADB_AUDIT_PLUGIN	All MySQL 5.6 versions
		MySQL 5.7.16 and later 5.7 versions
MySQL memcached Support (p. 877)	MEMCACHED	All MySQL 5.6, 5.7, and 8.0 versions

MariaDB Audit Plugin Support

Amazon RDS supports using the MariaDB Audit Plugin on MySQL database instances. The MariaDB Audit Plugin records database activity such as users logging on to the database, queries run against the database, and more. The record of database activity is stored in a log file.

Note

Currently, the MariaDB Audit Plugin is only supported for the following Amazon RDS MySQL versions:

- All 5.6 versions
 - MySQL 5.7.16 and later 5.7 versions

Audit Plugin Option Settings

Amazon RDS supports the following settings for the MariaDB Audit Plugin option.

Option Setting	Valid Values	Default Value	Description
SERVER_AUDIT	/rdsdbdata/ log/audit/	/rdsdbdata/ log/audit/	The location of the log file. The log file contains the record of the activity specified in SERVER_AUDIT_EVENTS. For more information, see Viewing and Listing Database Log Files (p. 390) and MySQL Database Log Files (p. 405) .
SERVER_AUDIT_SIZE	100000000	1000000	The size in bytes that when reached, causes the file to rotate. For more information, see Log File Size (p. 409) .
SERVER_AUDIT_ROTATIONS	9		The number of log rotations to save. For more information, see Log File Size (p. 409) and Downloading a Database Log File (p. 390) .
SERVER_AUDIT_EVENTS	CONNECT, QUERY, QUERY_DDL, QUERY_DML, QUERY_DCL	CONNECT, QUERY	<p>The types of activity to record in the log. Installing the MariaDB Audit Plugin is itself logged.</p> <ul style="list-style-type: none"> • CONNECT: Log successful and unsuccessful connections to the database, and disconnections from the database. • QUERY: Log the text of all queries run against the database. • QUERY_DDL: Similar to the QUERY event, but returns only data definition language (DDL) queries (CREATE, ALTER, and so on). • QUERY_DML: Similar to the QUERY event, but returns only data manipulation language (DML) queries (INSERT, UPDATE, and so on, and also SELECT). • QUERY_DCL: Similar to the QUERY event, but returns only data control language (DCL) queries (GRANT, REVOKE, and so on). <p>For MySQL, TABLE is not supported.</p>

Option Setting	Valid Values	Default Value	Description
SERVER_AUDIT_USERS	Multiple comma-separated values	None	<p>Include only activity from the specified users. By default, activity is recorded for all users. If a user is specified in both SERVER_AUDIT_EXCL_USERS and SERVER_AUDIT_INCL_USERS, then activity is recorded for the user.</p>
SERVER_AUDIT_EXCL_USERS	Multiple comma-separated values	None	<p>Exclude activity from the specified users. By default, activity is recorded for all users. If a user is specified in both SERVER_AUDIT_EXCL_USERS and SERVER_AUDIT_INCL_USERS, then activity is recorded for the user.</p> <p>The <code>rdsadmin</code> user queries the database every second to check the health of the database. Depending on your other settings, this activity can possibly cause the size of your log file to grow very large, very quickly. If you don't need to record this activity, add the <code>rdsadmin</code> user to the SERVER_AUDIT_EXCL_USERS list.</p> <p>Note CONNECT activity is always recorded for all users, even if the user is specified for this option setting.</p>
SERVER_AUDIT_LOGGING	ON	ON	<p>Logging is active. The only valid value is ON. Amazon RDS does not support deactivating logging. If you want to deactivate logging, remove the MariaDB Audit Plugin. For more information, see Removing the MariaDB Audit Plugin (p. 876).</p>

Adding the MariaDB Audit Plugin

The general process for adding the MariaDB Audit Plugin to a DB instance is the following:

- Create a new option group, or copy or modify an existing option group
- Add the option to the option group
- Associate the option group with the DB instance

After you add the MariaDB Audit Plugin, you don't need to restart your DB instance. As soon as the option group is active, auditing begins immediately.

Important

Adding the MariaDB Audit Plugin to a DB instance might cause an outage. We recommend adding the MariaDB Audit Plugin during a maintenance window or during a time of low database workload.

To add the MariaDB Audit Plugin

1. Determine the option group you want to use. You can create a new option group or use an existing option group. If you want to use an existing option group, skip to the next step. Otherwise, create a custom DB option group. Choose `mysql` for **Engine**, and choose **5.6** or **5.7** for **Major engine version**. For more information, see [Creating an Option Group \(p. 197\)](#).

2. Add the **MARIADB_AUDIT_PLUGIN** option to the option group, and configure the option settings. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#). For more information about each setting, see [Audit Plugin Option Settings \(p. 874\)](#).
3. Apply the option group to a new or existing DB instance.
 - For a new DB instance, you apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the MySQL Database Engine \(p. 751\)](#).
 - For an existing DB instance, you apply the option group by modifying the instance and attaching the new option group. For more information, see [Modifying a DB Instance Running the MySQL Database Engine \(p. 802\)](#).

Viewing and Downloading the MariaDB Audit Plugin Log

After you enable the MariaDB Audit Plugin, you access the results in the log files the same way you access any other text-based log files. The audit log files are located at `/rdsdbdata/log/audit/`. For information about viewing the log file in the console, see [Viewing and Listing Database Log Files \(p. 390\)](#). For information about downloading the log file, see [Downloading a Database Log File \(p. 390\)](#).

Modifying MariaDB Audit Plugin Settings

After you enable the MariaDB Audit Plugin, you can modify the settings. For more information about how to modify option settings, see [Modifying an Option Setting \(p. 204\)](#). For more information about each setting, see [Audit Plugin Option Settings \(p. 874\)](#).

Removing the MariaDB Audit Plugin

Amazon RDS doesn't support turning off logging in the MariaDB Audit Plugin. However, you can remove the plugin from a DB instance. When you remove the MariaDB Audit Plugin, the DB instance is restarted automatically to stop auditing.

To remove the MariaDB Audit Plugin from a DB instance, do one of the following:

- Remove the MariaDB Audit Plugin option from the option group it belongs to. This change affects all DB instances that use the option group. For more information, see [Removing an Option from an Option Group \(p. 207\)](#)
- Modify the DB instance and specify a different option group that doesn't include the plugin. This change affects a single DB instance. You can specify the default (empty) option group, or a different custom option group. For more information, see [Modifying a DB Instance Running the MySQL Database Engine \(p. 802\)](#).

MySQL memcached Support

Amazon RDS supports using the memcached interface to InnoDB tables that was introduced in MySQL 5.6. The memcached API enables applications to use InnoDB tables in a manner similar to NoSQL key-value data stores.

The memcached interface is a simple, key-based cache. Applications use memcached to insert, manipulate, and retrieve key-value data pairs from the cache. MySQL 5.6 introduced a plugin that implements a daemon service that exposes data from InnoDB tables through the memcached protocol. For more information about the MySQL memcached plugin, see [InnoDB Integration with memcached](#).

To enable memcached support for an Amazon RDS MySQL 5.6 or later instance

1. Determine the security group to use for controlling access to the memcached interface. If the set of applications already using the SQL interface are the same set that will access the memcached interface, you can use the existing VPC or DB security group used by the SQL interface. If a different set of applications will access the memcached interface, define a new VPC or DB security group. For more information about managing security groups, see [Controlling Access with Security Groups \(p. 486\)](#)
2. Create a custom DB option group, selecting MySQL as the engine type and a 5.6 or later version. For more information about creating an option group, see [Creating an Option Group \(p. 197\)](#).
3. Add the MEMCACHED option to the option group. Specify the port that the memcached interface will use, and the security group to use in controlling access to the interface. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#).
4. Modify the option settings to configure the memcached parameters, if necessary. For more information about how to modify option settings, see [Modifying an Option Setting \(p. 204\)](#).
5. Apply the option group to an instance. Amazon RDS enables memcached support for that instance when the option group is applied:
 - You enable memcached support for a new instance by specifying the custom option group when you launch the instance. For more information about launching a MySQL instance, see [Creating a DB Instance Running the MySQL Database Engine \(p. 751\)](#).
 - You enable memcached support for an existing instance by specifying the custom option group when you modify the instance. For more information about modifying a MySQL instance, see [Modifying a DB Instance Running the MySQL Database Engine \(p. 802\)](#).
6. Specify which columns in your MySQL tables can be accessed through the memcached interface. The memcached plug-in creates a catalog table named `containers` in a dedicated database named `innodb_memcache`. You insert a row into the `containers` table to map an InnoDB table for access through memcached. You specify a column in the InnoDB table that is used to store the memcached key values, and one or more columns that are used to store the data values associated with the key. You also specify a name that a memcached application uses to refer to that set of columns. For details on inserting rows in the `containers` table, see [Internals of the InnoDB memcached Plugin](#). For an example of mapping an InnoDB table and accessing it through memcached, see [Specifying the Table and Column Mappings for an InnoDB + memcached Application](#).
7. If the applications accessing the memcached interface are on different computers or EC2 instances than the applications using the SQL interface, add the connection information for those computers to the VPC or DB security group associated with the MySQL instance. For more information about managing security groups, see [Controlling Access with Security Groups \(p. 486\)](#).

You turn off the memcached support for an instance by modifying the instance and specifying the default option group for your MySQL version. For more information about modifying a MySQL instance, see [Modifying a DB Instance Running the MySQL Database Engine \(p. 802\)](#).

MySQL memcached Security Considerations

The memcached protocol does not support user authentication. For more information about MySQL memcached security considerations, see [memcached Deployment](#) and [Using memcached as a MySQL Caching Layer](#).

You can take the following actions to help increase the security of the memcached interface:

- Specify a different port than the default of 11211 when adding the **MEMCACHED** option to the option group.
- Ensure that you associate the memcached interface with either a VPC or DB security group that limits access to known, trusted client addresses or EC2 instances. For more information about managing security groups, see [Controlling Access with Security Groups \(p. 486\)](#).

MySQL memcached Connection Information

To access the memcached interface, an application must specify both the DNS name of the Amazon RDS instance and the memcached port number. For example, if an instance has a DNS name of `my-cache-instance.cg034hpkmmt.region.rds.amazonaws.com` and the memcached interface is using port 11212, the connection information specified in PHP would be:

```
<?php  
  
$cache = new Memcache;  
$cache->connect('my-cache-instance.cg034hpkmmt.region.rds.amazonaws.com', 11212);  
?>
```

To find the DNS name and memcached port of an Amazon RDS MySQL instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the top right corner of the AWS Management Console, select the region that contains the DB instance.
3. In the navigation pane, choose **Databases**.
4. Choose the MySQL DB instance name to display its details.
5. In the **Connect** section, note the value of the **Endpoint** field. The DNS name is the same as the endpoint. Also, note that the port in the **Connect** section is not used to access the memcached interface.
6. In the **Details** section, note the name listed in the **Option Group** field.
7. In the navigation pane, choose **Option groups**.
8. Choose the name of the option group used by the MySQL DB instance to show the option group details. In the **Options** section, note the value of the **Port** setting for the **MEMCACHED** option.

MySQL memcached Option Settings

Amazon RDS exposes the MySQL memcached parameters as option settings in the Amazon RDS **MEMCACHED** option.

MySQL memcached Parameters

- **DAEMON_MEMCACHED_R_BATCH_SIZE** – an integer that specifies how many memcached read operations (get) to perform before doing a COMMIT to start a new transaction. The allowed values are 1 to 4294967295; the default is 1. The option does not take effect until the instance is restarted.

- **DAEMON_MEMCACHED_W_BATCH_SIZE** – an integer that specifies how many memcached write operations, such as add, set, or incr, to perform before doing a COMMIT to start a new transaction. The allowed values are 1 to 4294967295; the default is 1. The option does not take effect until the instance is restarted.
- **INNODB_API_BK_COMMIT_INTERVAL** – an integer that specifies how often to auto-commit idle connections that use the InnoDB memcached interface. The allowed values are 1 to 1073741824; the default is 5. The option takes effect immediately, without requiring that you restart the instance.
- **INNODB_API_DISABLE_ROWLOCK** – a Boolean that disables (1 (true)) or enables (0 (false)) the use of row locks when using the InnoDB memcached interface. The default is 0 (false). The option does not take effect until the instance is restarted.
- **INNODB_API_ENABLE_MDL** – a Boolean that when set to 0 (false) locks the table used by the InnoDB memcached plugin, so that it cannot be dropped or altered by DDL through the SQL interface. The default is 0 (false). The option does not take effect until the instance is restarted.
- **INNODB_API_TRX_LEVEL** – an integer that specifies the transaction isolation level for queries processed by the memcached interface. The allowed values are 0 to 3. The default is 0. The option does not take effect until the instance is restarted.

Amazon RDS configures these MySQL memcached parameters, and they cannot be modified: DAEMON_MEMCACHED_LIB_NAME, DAEMON_MEMCACHED_LIB_PATH, and INNODB_API_ENABLE_BINLOG. The parameters that MySQL administrators set by using `daemon_memcached_options` are available as individual MEMCACHED option settings in Amazon RDS.

MySQL `daemon_memcached_options` Parameters

- **BINDING_PROTOCOL** – a string that specifies the binding protocol to use. The allowed values are auto, ascii, or binary. The default is auto, which means the server automatically negotiates the protocol with the client. The option does not take effect until the instance is restarted.
- **BACKLOG_QUEUE_LIMIT** – an integer that specifies how many network connections can be waiting to be processed by memcached. Increasing this limit may reduce errors received by a client that is not able to connect to the memcached instance, but does not improve the performance of the server. The allowed values are 1 to 2048; the default is 1024. The option does not take effect until the instance is restarted.
- **CAS_DISABLED** – a Boolean that enables (1 (true)) or disables (0 (false)) the use of compare and swap (CAS), which reduces the per-item size by 8 bytes. The default is 0 (false). The option does not take effect until the instance is restarted.
- **CHUNK_SIZE** – an integer that specifies the minimum chunk size, in bytes, to allocate for the smallest item's key, value, and flags. The allowed values are 1 to 48. The default is 48 and you can significantly improve memory efficiency with a lower value. The option does not take effect until the instance is restarted.
- **CHUNK_SIZE_GROWTH_FACTOR** – a float that controls the size of new chunks. The size of a new chunk is the size of the previous chunk times CHUNK_SIZE_GROWTH_FACTOR. The allowed values are 1 to 2; the default is 1.25. The option does not take effect until the instance is restarted.
- **ERROR_ON_MEMORY_EXHAUSTED** – a Boolean that when set to 1 (true) specifies that memcached will return an error rather than evicting items when there is no more memory to store items. If set to 0 (false), memcached will evict items if there is no more memory. The default is 0 (false). The option does not take effect until the instance is restarted.
- **MAX_SIMULTANEOUS_CONNECTIONS** – an integer that specifies the maximum number of concurrent connections. Setting this value to anything under 10 prevents MySQL from starting. The allowed values are 10 to 1024; the default is 1024. The option does not take effect until the instance is restarted.
- **VERBOSITY** – a string that specifies the level of information logged in the MySQL error log by the memcached service. The default is v. The option does not take effect until the instance is restarted. The allowed values are:

- **v** – Logs errors and warnings while executing the main event loop.
- **vv** – In addition to the information logged by v, also logs each client command and the response.
- **vvv** – In addition to the information logged by vv, also logs internal state transitions.

Amazon RDS configures these MySQL DAEMON_MEMCACHED_OPTIONS parameters, they cannot be modified: DAEMON_PROCESS, LARGE_MEMORY_PAGES, MAXIMUM_CORE_FILE_LIMIT, MAX_ITEM_SIZE, LOCK_DOWN_PAGE_MEMORY, MASK, IDFILE, REQUESTS_PER_EVENT, SOCKET, and USER.

Common DBA Tasks for MySQL DB Instances

This section describes the Amazon RDS-specific implementations of some common DBA tasks for DB instances running the MySQL database engine. In order to deliver a managed service experience, Amazon RDS does not provide shell access to DB instances, and it restricts access to certain system procedures and tables that require advanced privileges.

For information about working with MySQL log files on Amazon RDS, see [MySQL Database Log Files \(p. 405\)](#)

Topics

- [Killing a Session or Query \(p. 881\)](#)
- [Skipping the Current Replication Error \(p. 881\)](#)
- [Working with InnoDB Tablespaces to Improve Crash Recovery Times \(p. 882\)](#)
- [Managing the Global Status History \(p. 883\)](#)

Killing a Session or Query

You can terminate user sessions or queries on DB instances by using the `rds_kill` and `rds_kill_query` commands. First connect to your MySQL database instance, then issue the appropriate command as shown following. For more information, see [Connecting to a DB Instance Running the MySQL Database Engine \(p. 764\)](#).

```
CALL mysql.rds_kill(thread-ID)
CALL mysql.rds_kill_query(thread-ID)
```

For example, to kill the session that is running on thread 99, you would type the following:

```
CALL mysql.rds_kill(99);
```

To kill the query that is running on thread 99, you would type the following:

```
CALL mysql.rds_kill_query(99);
```

Skipping the Current Replication Error

Amazon RDS provides a mechanism for you to skip an error on your Read Replicas if the error is causing your Read Replica to hang and the error doesn't affect the integrity of your data. First connect to your MySQL database instance, then issue the appropriate commands as shown following. For more information, see [Connecting to a DB Instance Running the MySQL Database Engine \(p. 764\)](#).

Note

You should first verify that the error can be safely skipped. In a MySQL utility, connect to the Read Replica and run the following MySQL command:

```
SHOW SLAVE STATUS\G
```

For information about the values returned, go to [SHOW SLAVE STATUS Syntax](#) in the MySQL documentation.

To skip the error, you can issue the following command:

```
CALL mysql.rds_skip_repl_error;
```

This command has no effect if you run it on the source DB instance, or on a Read Replica that has not encountered a replication error.

For more information, such as the versions of MySQL that support `mysql.rds_skip_repl_error`, see [mysql.rds_skip_repl_error \(p. 904\)](#).

Important

If you attempt to call `mysql.rds_skip_repl_error` and encounter the following error: `ERROR 1305 (42000): PROCEDURE mysql.rds_skip_repl_error does not exist`, then upgrade your MySQL DB instance to the latest minor version or one of the minimum minor versions listed in [mysql.rds_skip_repl_error \(p. 904\)](#).

Working with InnoDB Tablespaces to Improve Crash Recovery Times

Every table in MySQL consists of a table definition, data, and indexes. The MySQL storage engine InnoDB stores table data and indexes in a *tablespace*. InnoDB creates a global shared tablespace that contains a data dictionary and other relevant metadata, and it can contain table data and indexes. InnoDB can also create separate tablespaces for each table and partition. These separate tablespaces are stored in files with a .ibd extension and the header of each tablespace contains a number that uniquely identifies it.

Amazon RDS provides a parameter in a MySQL parameter group called `innodb_file_per_table`. This parameter controls whether InnoDB adds new table data and indexes to the shared tablespace (by setting the parameter value to 0) or to individual tablespaces (by setting the parameter value to 1). Amazon RDS sets the default value for `innodb_file_per_table` parameter to 1, which allows you to drop individual InnoDB tables and reclaim storage used by those tables for the DB instance. In most use cases, setting the `innodb_file_per_table` parameter to 1 is the recommended setting.

You should set the `innodb_file_per_table` parameter to 0 when you have a large number of tables, such as over 1000 tables when you use standard (magnetic) or general purpose SSD storage or over 10,000 tables when you use Provisioned IOPS storage. When you set this parameter to 0, individual tablespaces are not created and this can improve the time it takes for database crash recovery.

MySQL processes each metadata file, which includes tablespaces, during the crash recovery cycle. The time it takes MySQL to process the metadata information in the shared tablespace is negligible compared to the time it takes to process thousands of tablespace files when there are multiple tablespaces. Because the tablespace number is stored within the header of each file, the aggregate time to read all the tablespace files can take up to several hours. For example, a million InnoDB tablespaces on standard storage can take from five to eight hours to process during a crash recovery cycle. In some cases, InnoDB can determine that it needs additional cleanup after a crash recovery cycle so it will begin another crash recovery cycle, which will extend the recovery time. Keep in mind that a crash recovery cycle also entails rolling-back transactions, fixing broken pages, and other operations in addition to the processing of tablespace information.

Since the `innodb_file_per_table` parameter resides in a parameter group, you can change the parameter value by editing the parameter group used by your DB instance without having to reboot the DB instance. After the setting is changed, for example, from 1 (create individual tables) to 0 (use shared tablespace), new InnoDB tables will be added to the shared tablespace while existing tables continue to have individual tablespaces. To move an InnoDB table to the shared tablespace, you must use the `ALTER TABLE` command.

Migrating Multiple Tablespaces to the Shared Tablespace

You can move an InnoDB table's metadata from its own tablespace to the shared tablespace, which will rebuild the table metadata according to the `innodb_file_per_table` parameter setting. First connect

to your MySQL database instance, then issue the appropriate commands as shown following. For more information, see [Connecting to a DB Instance Running the MySQL Database Engine \(p. 764\)](#).

```
ALTER TABLE table_name ENGINE = InnoDB, ALGORITHM=COPY;
```

For example, the following query returns an ALTER TABLE statement for every InnoDB table that is not in the shared tablespace.

```
SELECT CONCAT('ALTER TABLE `',
REPLACE(LEFT(NAME , INSTR((NAME), '/') - 1), ``, ``), ``.``,
REPLACE(SUBSTR(NAME FROM INSTR(NAME, '/') + 1), ``, ``), `` ENGINE=InnoDB,
ALGORITHM=COPY;') AS Query
FROM INFORMATION_SCHEMA.INNODB_SYS_TABLES
WHERE SPACE <> 0 AND LEFT(NAME, INSTR((NAME), '/') - 1) NOT IN ('mysql','');
```

Note

This query is supported on MySQL 5.6 and later.

Rebuilding a MySQL table to move the table's metadata to the shared tablespace requires additional storage space temporarily to rebuild the table, so the DB instance must have storage space available. During rebuilding, the table is locked and inaccessible to queries. For small tables or tables not frequently accessed, this may not be an issue; for large tables or tables frequently accessed in a heavily concurrent environment, you can rebuild tables on a Read Replica.

You can create a Read Replica and migrate table metadata to the shared tablespace on the Read Replica. While the ALTER TABLE statement blocks access on the Read Replica, the source DB instance is not affected. The source DB instance will continue to generate its binary logs while the Read Replica lags during the table rebuilding process. Because the rebuilding requires additional storage space and the replay log file can become large, you should create a Read Replica with storage allocated that is larger than the source DB instance.

The following steps should be followed to create a Read Replica and rebuild InnoDB tables to use the shared tablespace:

1. Ensure that backup retention is enabled on the source DB instance so that binary logging is enabled
2. Use the AWS Console or AWS CLI to create a Read Replica for the source DB instance. Since the creation of a Read Replica involves many of the same processes as crash recovery, the creation process may take some time if there are a large number of InnoDB tablespaces. Allocate more storage space on the Read Replica than is currently used on the source DB instance.
3. When the Read Replica has been created, create a parameter group with the parameter settings `read_only = 0` and `innodb_file_per_table = 0`, and then associate the parameter group with the Read Replica.
4. Issue `ALTER TABLE <name> ENGINE = InnoDB` against all tables you want migrated on the replica.
5. When all of your `ALTER TABLE` statements have completed on the Read Replica, verify that the Read Replica is connected to the source DB instance and that the two instances are in-sync.
6. When ready, use the AWS Console or AWS CLI to promote the Read Replica to be the master instance. Make sure that the parameter group used for the new master has the `innodb_file_per_table` parameter set to 0. Change the name of the new master, and point any applications to the new master instance.

Managing the Global Status History

MySQL maintains many status variables that provide information about its operation. Their value can help you detect locking or memory issues on a DB instance . The values of these status variables are

cumulative since last time the DB instance was started. You can reset most status variables to 0 by using the `FLUSH STATUS` command.

To allow for monitoring of these values over time, Amazon RDS provides a set of procedures that will snapshot the values of these status variables over time and write them to a table, along with any changes since the last snapshot. This infrastructure, called Global Status History (GoSH), is installed on all MySQL DB instances starting with versions 5.5.23. GoSH is disabled by default.

To enable GoSH, you first enable the event scheduler from a DB parameter group by setting the parameter `event_scheduler` to ON. For information about creating and modifying a DB parameter group, see [Working with DB Parameter Groups \(p. 211\)](#).

You can then use the procedures in the following table to enable and configure GoSH. First connect to your MySQL database instance, then issue the appropriate commands as shown following. For more information, see [Connecting to a DB Instance Running the MySQL Database Engine \(p. 764\)](#). For each procedure, type the following:

```
CALL procedure-name;
```

Where *procedure-name* is one of the procedures in the table.

Procedure	Description
<code>rds_enable_gsh_collector</code>	Enables GoSH to take default snapshots at intervals specified by <code>rds_set_gsh_collector</code> .
<code>rds_set_gsh_collector</code>	Specifies the interval, in minutes, between snapshots. Default value is 5.
<code>rds_disable_gsh_collector</code>	Disables snapshots.
<code>rds_collect_global_status_history</code>	Takes a snapshot on demand.
<code>rds_enable_gsh_rotation</code>	Enables rotation of the contents of the <code>mysql.rds_global_status_history</code> table to <code>mysql.rds_global_status_history_old</code> at intervals specified by <code>rds_set_gsh_rotation</code> .
<code>rds_set_gsh_rotation</code>	Specifies the interval, in days, between table rotations. Default value is 7.
<code>rds_disable_gsh_rotation</code>	Disables table rotation.
<code>rds_rotate_global_status_history</code>	Rotates the contents of the <code>mysql.rds_global_status_history</code> table to <code>mysql.rds_global_status_history_old</code> on demand.

When GoSH is running, you can query the tables that it writes to. For example, to query the hit ratio of the Innodb buffer pool, you would issue the following query:

```
select a.collection_end, a.collection_start, (( a.variable_Delta-b.variable_delta)/
a.variable_delta)*100 as "HitRatio"
  from mysql.rds_global_status_history as a join mysql.rds_global_status_history as b on
a.collection_end = b.collection_end
    where a.variable_name = 'Innodb_buffer_pool_read_requests' and b.variable_name =
'Innodb_buffer_pool_reads'
```

Known Issues and Limitations for MySQL on Amazon RDS

Known issues and limitations for working with MySQL on Amazon RDS are as follows.

Topics

- [Inconsistent InnoDB Buffer Pool Size \(p. 885\)](#)
- [Index Merge Optimization Returns Wrong Results \(p. 885\)](#)
- [Log File Size \(p. 886\)](#)
- [MySQL Parameter Exceptions for Amazon RDS DB Instances \(p. 886\)](#)
- [MySQL File Size Limits \(p. 887\)](#)
- [MySQL Keyring Plugin Not Supported \(p. 888\)](#)

Inconsistent InnoDB Buffer Pool Size

For MySQL 5.7, there is currently a bug in the way that the InnoDB buffer pool size is managed. MySQL 5.7 might adjust the value of the `innodb_buffer_pool_size` parameter to a large value that can result in the InnoDB buffer pool growing too large and using up too much memory. This effect can cause the MySQL database engine to stop running or can prevent the MySQL database engine from starting. This issue is more common for DB instance classes that have less memory available.

To resolve this issue, set the value of the `innodb_buffer_pool_size` parameter to a multiple of the product of the `innodb_buffer_pool_instances` parameter value and the `innodb_buffer_pool_chunk_size` parameter value. For example, you might set the `innodb_buffer_pool_size` parameter value to a multiple of eight times the product of the `innodb_buffer_pool_instances` and `innodb_buffer_pool_chunk_size` parameter values, as shown in the following example.

```
innodb_buffer_pool_chunk_size = 536870912
innodb_buffer_pool_instances = 4
innodb_buffer_pool_size = (536870912 * 4) * 8 = 17179869184
```

For details on this MySQL 5.7 bug, go to <https://bugs.mysql.com/bug.php?id=79379> in the MySQL documentation.

Index Merge Optimization Returns Wrong Results

Queries that use index merge optimization might return wrong results due to a bug in the MySQL query optimizer that was introduced in MySQL 5.5.37. When you issue a query against a table with multiple indexes the optimizer scans ranges of rows based on the multiple indexes, but does not merge the results together correctly. For more information on the query optimizer bug, go to <http://bugs.mysql.com/bug.php?id=72745> and <http://bugs.mysql.com/bug.php?id=68194> in the MySQL bug database.

For example, consider a query on a table with two indexes where the search arguments reference the indexed columns.

```
SELECT * FROM table1
WHERE indexed_col1 = 'value1' AND indexed_col2 = 'value2';
```

In this case, the search engine will search both indexes. However, due to the bug, the merged results are incorrect.

To resolve this issue, you can do one of the following:

- Set the `optimizer_switch` parameter to `index_merge=off` in the DB parameter group for your MySQL DB instance. For information on setting DB parameter group parameters, see [Working with DB Parameter Groups \(p. 211\)](#).
- Upgrade your MySQL DB instance to MySQL version 5.6, 5.7, or 8.0. For more information, see [Upgrading a MySQL DB Snapshot \(p. 818\)](#).
- If you cannot upgrade your instance or change the `optimizer_switch` parameter, you can work around the bug by explicitly identifying an index for the query, for example:

```
SELECT * FROM table1
USE INDEX covering_index
WHERE indexed_col1 = 'value1' AND indexed_col2 = 'value2';
```

For more information, go to [Index Merge Optimization](#).

Log File Size

For MySQL, there is a size limit on BLOBS written to the redo log. To account for this limit, ensure that the `innodb_log_file_size` parameter for your MySQL DB instance is 10 times larger than the largest BLOB data size found in your tables, plus the length of other variable length fields (`VARCHAR`, `VARBINARY`, `TEXT`) in the same tables. For information on how to set parameter values, see [Working with DB Parameter Groups \(p. 211\)](#). For information on the redo log BLOB size limit, go to [Changes in MySQL 5.6.20](#).

MySQL Parameter Exceptions for Amazon RDS DB Instances

Some MySQL parameters require special considerations when used with an Amazon RDS DB instance.

[lower_case_table_names](#)

Because Amazon RDS uses a case-sensitive file system, setting the value of the `lower_case_table_names` server parameter to 2 ("names stored as given but compared in lowercase") is not supported. The following are the supported values for Amazon RDS for MySQL DB instances:

- 0 ("names stored as given and comparisons are case-sensitive") is supported for all Amazon RDS for MySQL versions.
- 1 ("names stored in lowercase and comparisons are not case-sensitive") is supported for Amazon RDS for MySQL versions 5.5, 5.6, and 5.7.

Note

For MySQL 8.0 DB instances, setting `lower_case_table_names` to 1 isn't supported.

The `lower_case_table_names` parameter should be set as part of a custom DB parameter group before creating a DB instance. You should avoid changing the `lower_case_table_names` parameter for existing database instances because doing so could cause inconsistencies with point-in-time recovery backups and Read Replica DB instances.

Read Replicas should always use the same `lower_case_table_names` parameter value as the master DB instance.

long_query_time

You can set the `long_query_time` parameter to a floating point value which allows you to log slow queries to the MySQL slow query log with microsecond resolution. You can set a value such as 0.1 seconds, which would be 100 milliseconds, to help when debugging slow transactions that take less than one second.

MySQL File Size Limits

For Amazon RDS MySQL DB instances, the maximum provisioned storage limit constrains the size of a table to a maximum size of 16 TB when using InnoDB file-per-table tablespaces. This limit also constrains the system tablespace to a maximum size of 16 TB. InnoDB file-per-table tablespaces (with tables each in their own tablespace) is set by default for Amazon RDS MySQL DB instances.

Note

Some existing DB instances have a lower limit. For example, MySQL DB instances created prior to April 2014 have a file and table size limit of 2 TB. This 2 TB file size limit also applies to DB instances or Read Replicas created from DB snapshots taken prior to April 2014, regardless of when the DB instance was created.

There are advantages and disadvantages to using InnoDB file-per-table tablespaces, depending on your application. To determine the best approach for your application, go to [InnoDB File-Per-Table Mode](#) in the MySQL documentation.

We don't recommend allowing tables to grow to the maximum file size. In general, a better practice is to partition data into smaller tables, which can improve performance and recovery times.

One option that you can use for breaking a large table up into smaller tables is partitioning. Partitioning distributes portions of your large table into separate files based on rules that you specify. For example, if you store transactions by date, you can create partitioning rules that distribute older transactions into separate files using partitioning. Then periodically, you can archive the historical transaction data that doesn't need to be readily available to your application. For more information, go to <https://dev.mysql.com/doc/refman/5.6/en/partitioning.html> in the MySQL documentation.

To determine the file size of a table

- Use the following SQL command to determine if any of your tables are too large and are candidates for partitioning.

```
SELECT TABLE_SCHEMA, TABLE_NAME,
round(((DATA_LENGTH + INDEX_LENGTH) / 1024 / 1024), 2) As "Approximate size (MB)"
FROM information_schema.TABLES
WHERE TABLE_SCHEMA NOT IN ('mysql', 'information_schema', 'performance_schema');
```

To enable InnoDB file-per-table tablespaces

- To enable InnoDB file-per-table tablespaces, set the `innodb_file_per_table` parameter to 1 in the parameter group for the DB instance.

To disable InnoDB file-per-table tablespaces

- To disable InnoDB file-per-table tablespaces, set the `innodb_file_per_table` parameter to 0 in the parameter group for the DB instance.

For information on updating a parameter group, see [Working with DB Parameter Groups \(p. 211\)](#).

When you have enabled or disabled InnoDB file-per-table tablespaces, you can issue an `ALTER TABLE` command to move a table from the global tablespace to its own tablespace, or from its own tablespace to the global tablespace as shown in the following example:

```
ALTER TABLE table_name ENGINE=InnoDB;
```

MySQL Keyring Plugin Not Supported

Currently, Amazon RDS for MySQL does not support the MySQL `keyring_aws` Amazon Web Services Keyring Plugin.

MySQL on Amazon RDS SQL Reference

This appendix describes system stored procedures that are available for Amazon RDS instances running the MySQL DB engine.

Overview

The following system stored procedures are supported for Amazon RDS DB instances running MySQL.

Replication

- [mysql.rds_set_master_auto_position \(p. 890\)](#)
- [mysql.rds_set_external_master \(p. 891\)](#)
- [mysql.rds_set_external_master_with_delay \(p. 893\)](#)
- [mysql.rds_set_external_master_with_auto_position \(p. 895\)](#)
- [mysql.rds_reset_external_master \(p. 897\)](#)
- [mysql.rds_import_binlog_ssl_material \(p. 898\)](#)
- [mysql.rds_remove_binlog_ssl_material \(p. 899\)](#)
- [mysql.rds_set_source_delay \(p. 900\)](#)
- [mysql.rds_start_replication \(p. 901\)](#)
- [mysql.rds_start_replication_until \(p. 901\)](#)
- [mysql.rds_start_replication_until_gtid \(p. 902\)](#)
- [mysql.rds_stop_replication \(p. 903\)](#)
- [mysql.rds_skip_transaction_with_gtid \(p. 904\)](#)
- [mysql.rds_skip_repl_error \(p. 904\)](#)
- [mysql.rds_next_master_log \(p. 905\)](#)

InnoDB cache warming

- [mysql.rds_innodb_buffer_pool_dump_now \(p. 907\)](#)
- [mysql.rds_innodb_buffer_pool_load_now \(p. 907\)](#)
- [mysql.rds_innodb_buffer_pool_load_abort \(p. 907\)](#)

Managing additional configuration (for example, binlog file retention)

- [mysql.rds_set_configuration \(p. 908\)](#)
- [mysql.rds_show_configuration \(p. 909\)](#)

Terminating a session or query

- [mysql.rds_kill \(p. 910\)](#)
- [mysql.rds_kill_query \(p. 910\)](#)

Logging

- [mysql.rds_rotate_general_log \(p. 911\)](#)
- [mysql.rds_rotate_slow_log \(p. 911\)](#)

Managing the global status history

- [mysql.rds_enable_gsh_collector \(p. 911\)](#)
- [mysql.rds_set_gsh_collector \(p. 912\)](#)
- [mysql.rds_disable_gsh_collector \(p. 912\)](#)
- [mysql.rds_collect_global_status_history \(p. 912\)](#)
- [mysql.rds_enable_gsh_rotation \(p. 912\)](#)
- [mysql.rds_set_gsh_rotation \(p. 913\)](#)
- [mysql.rds_disable_gsh_rotation \(p. 913\)](#)
- [mysql.rds_rotate_global_status_history \(p. 913\)](#)

SQL Reference Conventions

Following, you can find explanations for the conventions that are used to describe the syntax of the system stored procedures and tables described in the SQL reference section.

Character	Description
UPPERCASE	Words in uppercase are keywords.
[]	Square brackets indicate optional arguments.
{ }	Braces indicate that you are required to choose one of the arguments inside the braces.
	Pipes separate arguments that you can choose.
<i>italics</i>	Words in italics indicate placeholders. You must insert the appropriate value in place of the word in italics.
...	An ellipsis indicates that you can repeat the preceding element.
'	Words in single quotes indicate that you must type the quotes.

mysql.rds_set_master_auto_position

Sets the replication mode to be based on either binary log file positions or on global transaction identifiers (GTIDs).

Syntax

```
CALL mysql.rds_set_master_auto_position (
    auto_position_mode
);
```

Parameters

auto_position_mode

A value that indicates whether to use log file position replication or GTID-based replication:

- 0 – Use the replication method based on binary log file position. The default is 0.
- 1 – Use the GTID-based replication method.

Usage Notes

The master user must run the `mysql.rds_set_master_auto_position` procedure.

For Amazon RDS MySQL 5.7, this procedure is supported for MySQL 5.7.23 and later MySQL 5.7 versions. This procedure is not supported for Amazon RDS MySQL 5.5, 5.6, or 8.0.

mysql.rds_set_external_master

Configures a MySQL DB instance to be a Read Replica of an instance of MySQL running external to Amazon RDS.

Note

You can use the [mysql.rds_set_external_master_with_delay \(p. 893\)](#) stored procedure to configure an external master and delayed replication.

Syntax

```
CALL mysql.rds_set_external_master (
    host_name
    , host_port
    , replication_user_name
    , replication_user_password
    , mysql_binary_log_file_name
    , mysql_binary_log_file_location
    , ssl_encryption
);
```

Parameters

host_name

The host name or IP address of the MySQL instance running external to Amazon RDS to become the replication master.

host_port

The port used by the MySQL instance running external to Amazon RDS to be configured as the replication master. If your network configuration includes Secure Shell (SSH) port replication that converts the port number, specify the port number that is exposed by SSH.

replication_user_name

The ID of a user with REPLICATION CLIENT and REPLICATION SLAVE permissions on the MySQL instance running external to Amazon RDS. We recommend that you provide an account that is used solely for replication with the external instance.

replication_user_password

The password of the user ID specified in `replication_user_name`.

mysql_binary_log_file_name

The name of the binary log on the replication master that contains the replication information.

mysql_binary_log_file_location

The location in the `mysql_binary_log_file_name` binary log at which replication starts reading the replication information.

ssl_encryption

A value that specifies whether Secure Socket Layer (SSL) encryption is used on the replication connection. 1 specifies to use SSL encryption, 0 specifies to not use encryption. The default is 0.

Note

This parameter currently is only implemented for Amazon Aurora with MySQL compatibility. On MySQL DB instances, only the default is allowed.

Usage Notes

The master user must run the `mysql.rds_set_external_master` procedure. This procedure must be run on the MySQL DB instance to be configured as the Read Replica of a MySQL instance running external to Amazon RDS.

Before you run `mysql.rds_set_external_master`, you must configure the instance of MySQL running external to Amazon RDS to be a replication master. To connect to the MySQL instance running external to Amazon RDS, you must specify `replication_user_name` and `replication_user_password` values that indicate a replication user that has `REPLICATION CLIENT` and `REPLICATION SLAVE` permissions on the external instance of MySQL.

To configure an external instance of MySQL as a replication master

1. Using the MySQL client of your choice, connect to the external instance of MySQL and create a user account to be used for replication. The following is an example.

```
CREATE USER 'repl_user'@'mydomain.com' IDENTIFIED BY 'password'
```

2. On the external instance of MySQL, grant `REPLICATION CLIENT` and `REPLICATION SLAVE` privileges to your replication user. The following example grants `REPLICATION CLIENT` and `REPLICATION SLAVE` privileges on all databases for the 'repl_user' user for your domain.

```
GRANT REPLICATION CLIENT, REPLICATION SLAVE ON *.* TO 'repl_user'@'mydomain.com'  
IDENTIFIED BY 'password'
```

To use encrypted replication, configure the master to use SSL connections. Also, import the certificate authority certificate, client certificate, and client key into the DB instance or DB cluster using the [mysql.rds_import_binlog_ssl_material \(p. 898\)](#) procedure.

Note

We recommend that you use Read Replicas to manage replication between two Amazon RDS DB instances when possible. When you do so, we recommend that you use only this and other replication-related stored procedures. These practices enable more complex replication topologies between Amazon RDS DB instances. We offer these stored procedures primarily to enable replication with MySQL instances running external to Amazon RDS. For information about managing replication between Amazon RDS DB instances, see [Working with Read Replicas \(p. 182\)](#).

After calling `mysql.rds_set_external_master` to configure an Amazon RDS DB instance as a Read Replica, you can call [mysql.rds_start_replication \(p. 901\)](#) on the Read Replica to start the replication process. You can call [mysql.rds_reset_external_master \(p. 897\)](#) to remove the Read Replica configuration.

When `mysql.rds_set_external_master` is called, Amazon RDS records the time, user, and an action of "set master" in the `mysql.rds_history` and `mysql.rds_replication_status` tables.

Examples

When run on a MySQL DB instance, the following example configures the DB instance to be a Read Replica of an instance of MySQL running external to Amazon RDS.

```
call mysql.rds_set_external_master(
  'Externaldb.some.com',
  3306,
  'repl_user',
  'password',
  'mysql-bin-changelog.0777',
  120,
  0);
```

mysql.rds_set_external_master_with_delay

Configures an Amazon RDS MySQL DB instance to be a Read Replica of an instance of MySQL running external to Amazon RDS and configures delayed replication.

Syntax

```
CALL mysql.rds_set_external_master_with_delay (
  host_name
, host_port
, replication_user_name
, replication_user_password
, mysql_binary_log_file_name
, mysql_binary_log_file_location
, ssl_encryption
, delay
);
```

Parameters

host_name

The host name or IP address of the MySQL instance running external to Amazon RDS that will become the replication master.

host_port

The port used by the MySQL instance running external to Amazon RDS to be configured as the replication master. If your network configuration includes SSH port replication that converts the port number, specify the port number that is exposed by SSH.

replication_user_name

The ID of a user with REPLICATION CLIENT and REPLICATION SLAVE permissions on the MySQL instance running external to Amazon RDS. We recommend that you provide an account that is used solely for replication with the external instance.

replication_user_password

The password of the user ID specified in *replication_user_name*.

mysql_binary_log_file_name

The name of the binary log on the replication master contains the replication information.

mysql_binary_log_file_location

The location in the *mysql_binary_log_file_name* binary log at which replication will start reading the replication information.

ssl_encryption

This option is not currently implemented. The default is 0.

delay

The minimum number of seconds to delay replication from the master.

The limit for this parameter is one day (86400 seconds).

Usage Notes

The master user must run the `mysql.rds_set_external_master_with_delay` procedure. This procedure must be run on the MySQL DB instance to be configured as the Read Replica of a MySQL instance running external to Amazon RDS.

Before you run `mysql.rds_set_external_master_with_delay`, you must configure the instance of MySQL running external to Amazon RDS to be a replication master. To connect to the MySQL instance running external to Amazon RDS, you must specify values for `replication_user_name` and `replication_user_password`. These values must indicate a replication user that has `REPLICATION CLIENT` and `REPLICATION SLAVE` permissions on the external instance of MySQL.

To configure an external instance of MySQL as a replication master

1. Using the MySQL client of your choice, connect to the external instance of MySQL and create a user account to be used for replication. The following is an example.

```
CREATE USER 'repl_user'@'mydomain.com' IDENTIFIED BY 'SomePassW0rd'
```

2. On the external instance of MySQL, grant `REPLICATION CLIENT` and `REPLICATION SLAVE` privileges to your replication user. The following example grants `REPLICATION CLIENT` and `REPLICATION SLAVE` privileges on all databases for the '`repl_user`' user for your domain.

```
GRANT REPLICATION CLIENT, REPLICATION SLAVE ON *.* TO 'repl_user'@'mydomain.com'  
IDENTIFIED BY 'SomePassW0rd'
```

For more information, see [Replication with a MySQL or MariaDB Instance Running External to Amazon RDS \(p. 862\)](#).

Note

We recommend that you use Read Replicas to manage replication between two Amazon RDS DB instances when possible. When you do so, we recommend that you use only this and other replication-related stored procedures. These practices enable more complex replication topologies between Amazon RDS DB instances. We offer these stored procedures primarily to enable replication with MySQL instances running external to Amazon RDS. For information about managing replication between Amazon RDS DB instances, see [Working with Read Replicas \(p. 182\)](#).

After calling `mysql.rds_set_external_master_with_delay` to configure an Amazon RDS DB instance as a Read Replica, you can call [mysql.rds_start_replication \(p. 901\)](#) on the Read Replica to start the replication process. You can call [mysql.rds_reset_external_master \(p. 897\)](#) to remove the Read Replica configuration.

When you call `mysql.rds_set_external_master_with_delay`, Amazon RDS records the time, the user, and an action of "set master" in the `mysql.rds_history` and `mysql.rds_replication_status` tables.

For disaster recovery, you can use this procedure with the [mysql.rds_start_replication_until \(p. 901\)](#) or [mysql.rds_start_replication_until_gtid \(p. 902\)](#) stored procedure. To roll forward changes to a delayed Read Replica to the time just before a disaster, you can run

the `mysql.rds_set_external_master_with_delay` procedure. After the `mysql.rds_start_replication_until` procedure stops replication, you can promote the Read Replica to be the new master DB instance by using the instructions in [Promoting a Read Replica to Be a Standalone DB Instance \(p. 186\)](#).

To use the `mysql.rds_start_replication_until_gtid` procedure, GTID-based replication must be enabled. To skip a specific GTID-based transaction that is known to cause disaster, you can use the [mysql.rds_skip_transaction_with_gtid \(p. 904\)](#) stored procedure. For more information about working with GTID-based replication, see [Using GTID-Based Replication for Amazon RDS MySQL \(p. 858\)](#).

The `mysql.rds_set_external_master_with_delay` procedure is available in these versions of Amazon RDS MySQL:

- MySQL 5.6.40 and later 5.6 versions
- MySQL 5.7.22 and later 5.7 versions

Examples

When run on a MySQL DB instance, the following example configures the DB instance to be a Read Replica of an instance of MySQL running external to Amazon RDS. It sets the minimum replication delay to one hour (3,600 seconds) on the MySQL DB instance. A change from the MySQL master running external to Amazon RDS is not applied on the MySQL DB instance Read Replica for at least one hour.

```
call mysql.rds_set_external_master_with_delay(
    'Externaldb.some.com',
    3306,
    'repl_user',
    'SomePassW0rd',
    'mysql-bin-changelog.000777',
    120,
    0,
    3600);
```

mysql.rds_set_external_master_with_auto_position

Configures an Amazon RDS MySQL DB instance to be a Read Replica of an instance of MySQL running external to Amazon RDS. This procedure also configures delayed replication and replication based on global transaction identifiers (GTIDs).

Syntax

```
CALL mysql.rds_set_external_master_with_auto_position (
    host_name
    , host_port
    , replication_user_name
    , replication_user_password
    , ssl_encryption
    , delay
);
```

Parameters

host_name

The host name or IP address of the MySQL instance running external to Amazon RDS to become the replication master.

host_port

The port used by the MySQL instance running external to Amazon RDS to be configured as the replication master. If your network configuration includes Secure Shell (SSH) port replication that converts the port number, specify the port number that is exposed by SSH.

replication_user_name

The ID of a user with REPLICATION CLIENT and REPLICATION SLAVE permissions on the MySQL instance running external to Amazon RDS. We recommend that you provide an account that is used solely for replication with the external instance.

replication_user_password

The password of the user ID specified in *replication_user_name*.

ssl_encryption

This option is not currently implemented. The default is 0.

delay

The minimum number of seconds to delay replication from the master.

The limit for this parameter is one day (86,400 seconds).

Usage Notes

The master user must run the `mysql.rds_set_external_master_with_auto_position` procedure. This procedure must be run on the MySQL DB instance to be configured as the Read Replica of a MySQL instance running external to Amazon RDS.

For Amazon RDS MySQL 5.7, this procedure is supported for MySQL 5.7.23 and later MySQL 5.7 versions. This procedure is not supported for Amazon RDS MySQL 5.5, 5.6, or 8.0.

Before you run `mysql.rds_set_external_master_with_auto_position`, you must configure the instance of MySQL running external to Amazon RDS to be a replication master. To connect to the MySQL instance running external to Amazon RDS, you must specify values for *replication_user_name* and *replication_user_password*. These values must indicate a replication user that has REPLICATION CLIENT and REPLICATION SLAVE permissions on the external instance of MySQL.

To configure an external instance of MySQL as a replication master

1. Using the MySQL client of your choice, connect to the external instance of MySQL and create a user account to be used for replication. The following is an example.

```
CREATE USER 'repl_user'@'mydomain.com' IDENTIFIED BY 'SomePassW0rd'
```

2. On the external instance of MySQL, grant REPLICATION CLIENT and REPLICATION SLAVE privileges to your replication user. The following example grants REPLICATION CLIENT and REPLICATION SLAVE privileges on all databases for the 'repl_user' user for your domain.

```
GRANT REPLICATION CLIENT, REPLICATION SLAVE ON *.* TO 'repl_user'@'mydomain.com'  
IDENTIFIED BY 'SomePassW0rd'
```

For more information, see [Replication with a MySQL or MariaDB Instance Running External to Amazon RDS \(p. 862\)](#).

Note

We recommend that you use Read Replicas to manage replication between two Amazon RDS DB instances when possible. When you do so, we recommend that you use only this and

other replication-related stored procedures. These practices enable more complex replication topologies between Amazon RDS DB instances. We offer these stored procedures primarily to enable replication with MySQL instances running external to Amazon RDS. For information about managing replication between Amazon RDS DB instances, see [Working with Read Replicas \(p. 182\)](#).

After calling `mysql.rds_set_external_master_with_auto_position` to configure an Amazon RDS DB instance as a Read Replica, you can call [mysql.rds_start_replication \(p. 901\)](#) on the Read Replica to start the replication process. You can call [mysql.rds_reset_external_master \(p. 897\)](#) to remove the Read Replica configuration.

When you call `mysql.rds_set_external_master_with_auto_position`, Amazon RDS records the time, the user, and an action of "set master" in the `mysql.rds_history` and `mysql.rds_replication_status` tables.

For disaster recovery, you can use this procedure with the [mysql.rds_start_replication_until \(p. 901\)](#) or [mysql.rds_start_replication_until_gtid \(p. 902\)](#) stored procedure. To roll forward changes to a delayed Read Replica to the time just before a disaster, you can run the `mysql.rds_set_external_master_with_auto_position` procedure. After the `mysql.rds_start_replication_until_gtid` procedure stops replication, you can promote the Read Replica to be the new master DB instance by using the instructions in [Promoting a Read Replica to Be a Standalone DB Instance \(p. 186\)](#).

To use the `mysql.rds_start_replication_until_gtid` procedure, GTID-based replication must be enabled. To skip a specific GTID-based transaction that is known to cause disaster, you can use the [mysql.rds_skip_transaction_with_gtid \(p. 904\)](#) stored procedure. For more information about working with GTID-based replication, see [Using GTID-Based Replication for Amazon RDS MySQL \(p. 858\)](#).

Examples

When run on a MySQL DB instance, the following example configures the DB instance to be a Read Replica of an instance of MySQL running external to Amazon RDS. It sets the minimum replication delay to one hour (3,600 seconds) on the MySQL DB instance. A change from the MySQL master running external to Amazon RDS is not applied on the MySQL DB instance Read Replica for at least one hour.

```
call mysql.rds_set_external_master_with_auto_position(
  'Externaldb.some.com',
  3306,
  'repl_user',
  'SomePassW0rd',
  0,
  3600);
```

mysql.rds_reset_external_master

Reconfigures a MySQL DB instance to no longer be a Read Replica of an instance of MySQL running external to Amazon RDS.

Syntax

```
CALL mysql.rds_reset_external_master;
```

Usage Notes

The master user must run the `mysql.rds_reset_external_master` procedure. This procedure must be run on the MySQL DB instance to be removed as a Read Replica of a MySQL instance running external to Amazon RDS.

Note

We recommend that you use Read Replicas to manage replication between two Amazon RDS DB instances when possible. When you do so, we recommend that you use only this and other replication-related stored procedures. These practices enable more complex replication topologies between Amazon RDS DB instances. We offer these stored procedures primarily to enable replication with MySQL instances running external to Amazon RDS. For information about managing replication between Amazon RDS DB instances, see [Working with Read Replicas \(p. 182\)](#).

For more information about using replication to import data from an instance of MySQL running external to Amazon RDS, see [Restoring a Backup into an Amazon RDS MySQL DB Instance \(p. 826\)](#).

mysql.rds_import_binlog_ssl_material

Imports the certificate authority certificate, client certificate, and client key into an Aurora MySQL DB cluster. The information is required for SSL communication and encrypted replication.

Note

Currently, this procedure is only supported for Aurora MySQL version 5.6.

Syntax

```
CALL mysql.rds_import_binlog_ssl_material (
    ssl_material
);
```

Parameters

ssl_material

JSON payload that contains the contents of the following .pem format files for a MySQL client:

- "ssl_ca": "*Certificate authority certificate*"
- "ssl_cert": "*Client certificate*"
- "ssl_key": "*Client key*"

Usage Notes

Prepare for encrypted replication before you run this procedure:

- If you don't have SSL enabled on the external MySQL master database and don't have a client key and client certificate prepared, enable SSL on the MySQL database server and generate the required client key and client certificate.
- If SSL is enabled on the external master, supply a client key and certificate for the Aurora MySQL DB cluster. If you don't have these, generate a new key and certificate for the Aurora MySQL DB cluster. To sign the client certificate, you must have the certificate authority key you used to configure SSL on the external MySQL master database.

For more information, see [Creating SSL Certificates and Keys Using openssl](#) in the MySQL documentation.

Important

After you prepare for encrypted replication, use an SSL connection to run this procedure. The client key must not be transferred across an insecure connection.

This procedure imports SSL information from an external MySQL database into an Aurora MySQL DB cluster. The SSL information is in .pem format files that contain the SSL information for the Aurora MySQL DB cluster. During encrypted replication, the Aurora MySQL DB cluster acts a client to the MySQL database server. The certificates and keys for the Aurora MySQL client are in files in .pem format.

You can copy the information from these files into the `ssl_material` parameter in the correct JSON payload. To support encrypted replication, import this SSL information into the Aurora MySQL DB cluster.

The JSON payload must be in the following format.

```
'{"ssl_ca":"-----BEGIN CERTIFICATE-----  
ssl_ca_pem_body_code  
-----END CERTIFICATE-----\n","ssl_cert":"-----BEGIN CERTIFICATE-----  
ssl_cert_pem_body_code  
-----END CERTIFICATE-----\n","ssl_key":"-----BEGIN RSA PRIVATE KEY-----  
ssl_key_pem_body_code  
-----END RSA PRIVATE KEY-----\n"}'
```

Examples

The following example imports SSL information into an Aurora MySQL DB cluster. In .pem format files, the body code typically is longer than the body code shown in the example.

```
call mysql.rds_import_binlog_ssl_material(  
'{"ssl_ca":"-----BEGIN CERTIFICATE-----  
AAAAAB3NzaC1yc2EAAAQABAAQClKsfkNkuSevGj3eYhCe53pcjqP3maAhDFcvBS706V  
hz2ItxCih+PnDSUaw+WNQn/mZphTk/a/gU8jEzoOWbkM4yxyb/wB96xbiFveSFJuOp/d6RJhJOI0iBXr  
lsLnB1tntckij7FbtxJMLvvwJryDUilBMTjYtwB+QhYXUMOzce5Pjz5/i8SeJtjnV3iAoG/cQk+0Fzz  
qaeJAAHco+CY/5WrUBkrHmFJr6HcxkvJdWPkyQS3xqC0+FmUzofz221CBt5IMucxxPpkX4rWi+z7wB3Rb  
BQoQzd8v7yeb7Oz1PnWOyN0qFU0XA246RA8QFYiCNYwi3f05p6KLxEXAMPLE  
-----END CERTIFICATE-----\n","ssl_cert":"-----BEGIN CERTIFICATE-----  
AAAAAB3NzaC1yc2EAAAQABAAQClKsfkNkuSevGj3eYhCe53pcjqP3maAhDFcvBS706V  
hz2ItxCih+PnDSUaw+WNQn/mZphTk/a/gU8jEzoOWbkM4yxyb/wB96xbiFveSFJuOp/d6RJhJOI0iBXr  
lsLnB1tntckij7FbtxJMLvvwJryDUilBMTjYtwB+QhYXUMOzce5Pjz5/i8SeJtjnV3iAoG/cQk+0Fzz  
qaeJAAHco+CY/5WrUBkrHmFJr6HcxkvJdWPkyQS3xqC0+FmUzofz221CBt5IMucxxPpkX4rWi+z7wB3Rb  
BQoQzd8v7yeb7Oz1PnWOyN0qFU0XA246RA8QFYiCNYwi3f05p6KLxEXAMPLE  
-----END CERTIFICATE-----\n","ssl_key":"-----BEGIN RSA PRIVATE KEY-----  
AAAAAB3NzaC1yc2EAAAQABAAQClKsfkNkuSevGj3eYhCe53pcjqP3maAhDFcvBS706V  
hz2ItxCih+PnDSUaw+WNQn/mZphTk/a/gU8jEzoOWbkM4yxyb/wB96xbiFveSFJuOp/d6RJhJOI0iBXr  
lsLnB1tntckij7FbtxJMLvvwJryDUilBMTjYtwB+QhYXUMOzce5Pjz5/i8SeJtjnV3iAoG/cQk+0Fzz  
qaeJAAHco+CY/5WrUBkrHmFJr6HcxkvJdWPkyQS3xqC0+FmUzofz221CBt5IMucxxPpkX4rWi+z7wB3Rb  
BQoQzd8v7yeb7Oz1PnWOyN0qFU0XA246RA8QFYiCNYwi3f05p6KLxEXAMPLE  
-----END RSA PRIVATE KEY-----\n"}');
```

Note

For information about using Amazon Aurora, see the [Amazon Aurora User Guide](#).

mysql.rds_remove_binlog_ssl_material

Removes the certificate authority certificate, client certificate, and client key for SSL communication and encrypted replication. This information is imported by using [mysql.rds_import_binlog_ssl_material \(p. 898\)](#).

Note

Currently, this procedure is only supported for Aurora MySQL version 5.6.

Syntax

```
CALL mysql.rds_remove_binlog_ssl_material;
```

mysql.rds_set_source_delay

Sets the minimum number of seconds to delay replication from the master to the current Read Replica. Use this procedure when you are connected to a Read Replica to delay replication from its master.

Syntax

```
CALL mysql.rds_set_source_delay(  
    delay  
) ;
```

Parameters

delay

The minimum number of seconds to delay replication from the master.

The limit for this parameter is one day (86400 seconds).

Usage Notes

The master user must run the `mysql.rds_set_source_delay` procedure.

For disaster recovery, you can use this procedure with the [mysql.rds_start_replication_until \(p. 901\)](#) stored procedure or the [mysql.rds_start_replication_until_gtid \(p. 902\)](#) stored procedure. To roll forward changes to a delayed Read Replica to the time just before a disaster, you can run the `mysql.rds_set_source_delay` procedure. After the `mysql.rds_start_replication_until` or `mysql.rds_start_replication_until_gtid` procedure stops replication, you can promote the Read Replica to be the new master DB instance by using the instructions in [Promoting a Read Replica to Be a Standalone DB Instance \(p. 186\)](#).

To use the `mysql.rds_start_replication_until_gtid` procedure, GTID-based replication must be enabled. To skip a specific GTID-based transaction that is known to cause disaster, you can use the [mysql.rds_skip_transaction_with_gtid \(p. 904\)](#) stored procedure. For more information on GTID-based replication, see [Using GTID-Based Replication for Amazon RDS MySQL \(p. 858\)](#).

The `mysql.rds_set_source_delay` procedure is available in these versions of Amazon RDS MySQL:

- MySQL 5.6.40 and later 5.6 versions
- MySQL 5.7.22 and later 5.7 versions

Examples

To delay replication from the master to the current Read Replica for at least one hour (3,600 seconds), you can call `mysql.rds_set_source_delay` with the following parameter:

```
CALL mysql.rds_set_source_delay(3600);
```

mysql.rds_start_replication

Initiates replication from a MySQL DB instance.

Note

You can use the [mysql.rds_start_replication_until \(p. 901\)](#) or [mysql.rds_start_replication_until_gtid \(p. 902\)](#) stored procedure to initiate replication from an Amazon RDS MySQL DB instance and stop replication at the specified binary log file location.

Syntax

```
CALL mysql.rds_start_replication;
```

Usage Notes

The master user must run the `mysql.rds_start_replication` procedure.

If you are configuring replication to import data from an instance of MySQL running external to Amazon RDS, you call `mysql.rds_start_replication` on the Read Replica to start the replication process after you have called [mysql.rds_set_external_master \(p. 891\)](#) to build the replication configuration. For more information, see [Restoring a Backup into an Amazon RDS MySQL DB Instance \(p. 826\)](#).

If you are configuring replication to export data to an instance of MySQL external to Amazon RDS, you call `mysql.rds_start_replication` and `mysql.rds_stop_replication` on the Read Replica to control some replication actions, such as purging binary logs. For more information, see [Exporting Data from a MySQL DB Instance by Using Replication \(p. 869\)](#).

You can also call `mysql.rds_start_replication` on the Read Replica to restart any replication process that you previously stopped by calling [mysql.rds_stop_replication \(p. 903\)](#). For more information, see [Working with Read Replicas \(p. 182\)](#).

mysql.rds_start_replication_until

Initiates replication from an Amazon RDS MySQL DB instance and stops replication at the specified binary log file location.

Syntax

```
CALL mysql.rds_start_replication_until (
  replication_log_file
  , replication_stop_point
);
```

Parameters

replication_log_file

The name of the binary log on the replication master contains the replication information.

replication_stop_point

The location in the `replication_log_file` binary log at which replication will stop.

Usage Notes

The master user must run the `mysql.rds_start_replication_until` procedure.

You can use this procedure with delayed replication for disaster recovery. If you have delayed replication configured, you can use this procedure to roll forward changes to a delayed Read Replica to the time just before a disaster. After this procedure stops replication, you can promote the Read Replica to be the new master DB instance by using the instructions in [Promoting a Read Replica to Be a Standalone DB Instance \(p. 186\)](#).

You can configure delayed replication using the following stored procedures:

- [mysql.rds_set_configuration \(p. 908\)](#)
- [mysql.rds_set_external_master_with_delay \(p. 893\)](#)
- [mysql.rds_set_source_delay \(p. 900\)](#)

The file name specified for the `replication_log_file` parameter must match the master binlog file name.

When the `replication_stop_point` parameter specifies a stop location that is in the past, replication is stopped immediately.

The `mysql.rds_start_replication_until` procedure is available in these versions of Amazon RDS MySQL:

- MySQL 5.6.40 and later 5.6 versions
- MySQL 5.7.22 and later 5.7 versions

Examples

The following example initiates replication and replicates changes until it reaches location 120 in the `mysql-bin-changelog.000777` binary log file.

```
call mysql.rds_start_replication_until(
  'mysql-bin-changelog.000777',
  120);
```

mysql.rds_start_replication_until_gtid

Initiates replication from an Amazon RDS MySQL DB instance and stops replication immediately after the specified global transaction identifier (GTID).

Syntax

```
CALL mysql.rds_start_replication_until_gtid (
  gtid
);
```

Parameters

gtid

The GTID after which replication is to stop.

Usage Notes

The master user must run the `mysql.rds_start_replication_until_gtid` procedure.

For Amazon RDS MySQL 5.7, this procedure is supported for MySQL 5.7.23 and later MySQL 5.7 versions. This procedure is not supported for Amazon RDS MySQL 5.5, 5.6, or 8.0.

You can use this procedure with delayed replication for disaster recovery. If you have delayed replication configured, you can use this procedure to roll forward changes to a delayed Read Replica to the time just before a disaster. After this procedure stops replication, you can promote the Read Replica to be the new master DB instance by using the instructions in [Promoting a Read Replica to Be a Standalone DB Instance \(p. 186\)](#).

You can configure delayed replication using the following stored procedures:

- [mysql.rds_set_configuration \(p. 908\)](#)
- [mysql.rds_set_external_master_with_auto_position \(p. 895\)](#)
- [mysql.rds_set_source_delay \(p. 900\)](#)

When the `gtid` parameter specifies a transaction that has already been executed by the replica, replication is stopped immediately.

Examples

The following example initiates replication and replicates changes until it reaches GTID `3E11FA47-71CA-11E1-9E33-C80AA9429562:23`.

```
call mysql.rds_start_replication_until_gtid(
  '3E11FA47-71CA-11E1-9E33-C80AA9429562:23');
```

mysql.rds_stop_replication

Terminates replication from a MySQL DB instance.

Syntax

```
CALL mysql.rds_stop_replication;
```

Usage Notes

The master user must run the `mysql.rds_stop_replication` procedure.

If you are configuring replication to import data from an instance of MySQL running external to Amazon RDS, you call `mysql.rds_stop_replication` on the Read Replica to stop the replication process after the import has completed. For more information, see [Restoring a Backup into an Amazon RDS MySQL DB Instance \(p. 826\)](#).

If you are configuring replication to export data to an instance of MySQL external to Amazon RDS, you call `mysql.rds_start_replication` and `mysql.rds_stop_replication` on the Read Replica to control some replication actions, such as purging binary logs. For more information, see [Exporting Data from a MySQL DB Instance by Using Replication \(p. 869\)](#).

You can also use `mysql.rds_stop_replication` to stop replication between two Amazon RDS DB instances. You typically stop replication to perform a long running operation on the Read Replica, such as creating a large index on the Read Replica. You can restart any replication process that you stopped by calling [mysql.rds_start_replication \(p. 901\)](#) on the Read Replica. For more information, see [Working with Read Replicas \(p. 182\)](#).

mysql.rds_skip_transaction_with_gtid

Skips replication of a transaction with the specified global transaction identifier (GTID) on a MySQL DB instance.

You can use this procedure for disaster recovery when a specific GTID transaction is known to cause a problem. Use this stored procedure to skip the problematic transaction. Examples of problematic transactions include transactions that disable replication, delete important data, or cause the DB instance to become unavailable.

Syntax

```
CALL mysql.rds_skip_transaction_with_gtid (
    gtid_to_skip
);
```

Parameters

gtid_to_skip

The GTID of the replication transaction to skip.

Usage Notes

The master user must run the `mysql.rds_skip_transaction_with_gtid` procedure.

For Amazon RDS MySQL 5.7, this procedure is supported for MySQL 5.7.23 and later MySQL 5.7 versions. This procedure is not supported for Amazon RDS MySQL 5.5, 5.6, or 8.0.

mysql.rds_skip_repl_error

Skips and deletes a replication error on a MySQL DB instance.

Syntax

```
CALL mysql.rds_skip_repl_error;
```

Usage Notes

The master user must run the `mysql.rds_skip_repl_error` procedure.

To determine if there are errors, run the MySQL `show slave status\G` command. If a replication error isn't critical, you can run `mysql.rds_skip_repl_error` to skip the error. If there are multiple errors, `mysql.rds_skip_repl_error` deletes the first error, then warns that others are present. You can then use `show slave status\G` to determine the correct course of action for the next error. For information about the values returned, see [SHOW SLAVE STATUS Syntax](#) in the MySQL documentation.

For more information about addressing replication errors with Amazon RDS, see [Troubleshooting a MySQL Read Replica Problem \(p. 856\)](#).

Important

If you try to call `mysql.rds_skip_repl_error`, you might encounter the following error:
`ERROR 1305 (42000): PROCEDURE mysql.rds_skip_repl_error does not exist.`

If you do, upgrade your MySQL DB instance to the latest minor version or one of the minimum minor versions listed in this topic.

Slave Down or Disabled Error

When you call the `mysql.rds_skip_repl_error` command, you might receive the following error message: `Slave is down or disabled`.

This error message appears because replication has stopped and could not be restarted.

If you need to skip a large number of errors, the replication lag can increase beyond the default retention period for binary log (binlog) files. In this case, you might encounter a fatal error due to binlog files being purged before they have been replayed on the Read Replica. This purge causes replication to stop, and you can no longer call the `mysql.rds_skip_repl_error` command to skip replication errors.

You can mitigate this issue by increasing the number of hours that binlog files are retained on your replication master. After you have increased the binlog retention time, you can restart replication and call the `mysql.rds_skip_repl_error` command as needed.

To set the binlog retention time, use the [mysql.rds_set_configuration \(p. 908\)](#) procedure and specify a configuration parameter of '`'binlog retention hours'`' along with the number of hours to retain binlog files on the DB cluster. The following example sets the retention period for binlog files to 48 hours.

```
CALL mysql.rds_set_configuration('binlog retention hours', 48);
```

mysql.rds_next_master_log

Changes the replication master log position to the start of the next binary log on the master. Use this procedure only if you are receiving replication I/O error 1236 on a Read Replica.

Syntax

```
CALL mysql.rds_next_master_log(  
curr_master_log  
);
```

Parameters

curr_master_log

The index of the current master log file. For example, if the current file is named `mysql-bin-changelog.012345`, then the index is 12345. To determine the current master log file name, run the `SHOW SLAVE STATUS` command and view the `Master_Log_File` field.

Usage Notes

The master user must run the `mysql.rds_next_master_log` procedure.

Warning

Call `mysql.rds_next_master_log` only if replication fails after a failover of a Multi-AZ DB instance that is the replication source, and the `Last_IO_Errno` field of `SHOW SLAVE STATUS` reports I/O error 1236.

Calling `mysql.rds_next_master_log` may result in data loss in the Read Replica if transactions in the source instance were not written to the binary log on disk before the failover

event occurred. You can reduce the chance of this happening by configuring the source instance parameters sync_binlog = 1 and innodb_support_xa = 1, although this may reduce performance. For more information, see [Working with Read Replicas \(p. 182\)](#).

Examples

Assume replication fails on an Amazon RDS Read Replica. Running `SHOW SLAVE STATUS\G` on the Read Replica returns the following result:

```
***** 1. row *****
Slave_IO_State:
    Master_Host: myhost.XXXXXXXXXXXXXXX.rr-rrrr-1.rds.amazonaws.com
    Master_User: MasterUser
    Master_Port: 3306
    Connect_Retry: 10
    Master_Log_File: mysql-bin-changelog.012345
    Read_Master_Log_Pos: 1219393
        Relay_Log_File: relaylog.012340
        Relay_Log_Pos: 30223388
    Relay_Master_Log_File: mysql-bin-changelog.012345
        Slave_IO_Running: No
        Slave_SQL_Running: Yes
            Replicate_Do_DB:
            Replicate_Ignore_DB:
            Replicate_Do_Table:
            Replicate_Ignore_Table:
            Replicate_Wild_Do_Table:
            Replicate_Wild_Ignore_Table:
                Last_Error:
                Skip_Counter: 0
            Exec_Master_Log_Pos: 30223232
                Relay_Log_Space: 5248928866
                Until_Condition: None
                Until_Log_File:
                Until_Log_Pos: 0
            Master_SSL_Allowed: No
            Master_SSL_CA_File:
            Master_SSL_CA_Path:
                Master_SSL_Cert:
                Master_SSL_Cipher:
                Master_SSL_Key:
            Seconds_Behind_Master: NULL
Master_SSL_Verify_Server_Cert: No
    Last_IO_Errno: 1236
    Last_IO_Error: Got fatal error 1236 from master when reading data from
binary log: 'Client requested master to start replication from impossible position; the
first event 'mysql-bin-changelog.013406' at 1219393, the last event read from '/rdsdbdata/
log/binlog/mysql-bin-changelog.012345' at 4, the last byte read from '/rdsdbdata/log/
binlog/mysql-bin-changelog.012345' at 4.'
    Last_SQL_Errno: 0
    Last_SQL_Error:
Replicate_Ignore_Server_Ids:
    Master_Server_Id: 67285976
```

The `Last_IO_Errno` field shows that the instance is receiving I/O error 1236. The `Master_Log_File` field shows that the file name is `mysql-bin-changelog.012345`, which means that the log file index is 12345. To resolve the error, you can call `mysql.rds_next_master_log` with the following parameter:

```
CALL mysql.rds_next_master_log(12345);
```

mysql.rds_innodb_buffer_pool_dump_now

Dumps the current state of the buffer pool to disk. For more information, see [InnoDB Cache Warming \(p. 748\)](#).

Syntax

```
CALL mysql.rds_innodb_buffer_pool_dump_now();
```

Usage Notes

The master user must run the `mysql.rds_innodb_buffer_pool_dump_now` procedure.

The `mysql.rds_innodb_buffer_pool_dump_now` procedure is available in these versions of Amazon RDS MySQL:

- MySQL 5.6
- MySQL 5.7
- MySQL 8.0

mysql.rds_innodb_buffer_pool_load_now

Loads the saved state of the buffer pool from disk. For more information, see [InnoDB Cache Warming \(p. 748\)](#).

Syntax

```
CALL mysql.rds_innodb_buffer_pool_load_now();
```

Usage Notes

The master user must run the `mysql.rds_innodb_buffer_pool_load_now` procedure.

The `mysql.rds_innodb_buffer_pool_load_now` procedure is available in these versions of Amazon RDS MySQL:

- MySQL 5.6
- MySQL 5.7
- MySQL 8.0

mysql.rds_innodb_buffer_pool_load_abort

Cancels a load of the saved buffer pool state while in progress. For more information, see [InnoDB Cache Warming \(p. 748\)](#).

Syntax

```
CALL mysql.rds_innodb_buffer_pool_load_abort();
```

Usage Notes

The master user must run the `mysql.rds_innodb_buffer_pool_load_abort` procedure.

The `mysql.rds_innodb_buffer_pool_load_abort` procedure is available in these versions of Amazon RDS MySQL:

- MySQL 5.6
- MySQL 5.7
- MySQL 8.0

mysql.rds_set_configuration

Specifies the number of hours to retain binary logs or the number of seconds to delay replication.

Syntax

```
CALL mysql.rds_set_configuration(name,value);
```

Parameters

name

The name of the configuration parameter to set.

value

The value of the configuration parameter.

Usage Notes

The `mysql.rds_set_configuration` stored procedure is available in these versions of Amazon RDS MySQL:

- MySQL 5.6
- MySQL 5.7
- MySQL 8.0

The `mysql.rds_set_configuration` procedure supports the following configuration parameters:

- [binlog retention hours \(p. 908\)](#)
- [target delay \(p. 909\)](#)

binlog retention hours

The `binlog retention hours` parameter is used to specify the number of hours to retain binary log files. Amazon RDS normally purges a binary log as soon as possible, but the binary log might still be required for replication with a MySQL database external to Amazon RDS. The default value of `binlog retention hours` is `NULL` (do not retain binary logs).

To specify the number of hours for Amazon RDS to retain binary logs on a DB instance, use the `mysql.rds_set_configuration` stored procedure and specify a period with enough time for replication to occur, as shown in the following example.

```
call mysql.rds_set_configuration('binlog retention hours', 24);
```

For MySQL DB instances, the maximum binlog retention hours value is 168 (7 days).

After you set the retention period, monitor storage usage for the DB instance to make sure that the retained binary logs don't take up too much storage.

target delay

Use the target_delay parameter to specify the number of seconds to delay replication from the master to the Read Replica. The specified delay applies to new replicas created from the current DB instance. Amazon RDS normally replicates changes as soon as possible, but some environments might want to delay replication. For example, when replication is delayed, you can roll forward a delayed Read Replica to the time just before a disaster. If a table is dropped accidentally, you can use delayed replication to recover it quickly. The default value of target_delay is 0 (don't delay replication).

For disaster recovery, you can use this configuration parameter with the [mysql.rds_start_replication_until \(p. 901\)](#) stored procedure or the [mysql.rds_start_replication_until_gtid \(p. 902\)](#) stored procedure. To roll forward changes to a delayed Read Replica to the time just before a disaster, you can run the mysql.rds_set_configuration procedure with this parameter set. After the mysql.rds_start_replication_until or mysql.rds_start_replication_until_gtid procedure stops replication, you can promote the Read Replica to be the new master DB instance by using the instructions in [Promoting a Read Replica to Be a Standalone DB Instance \(p. 186\)](#).

To use the mysql.rds_start_replication_until_gtid procedure, GTID-based replication must be enabled. To skip a specific GTID-based transaction that is known to cause disaster, you can use the [mysql.rds_skip_transaction_with_gtid \(p. 904\)](#) stored procedure. For more information about working with GTID-based replication, see [Using GTID-Based Replication for Amazon RDS MySQL \(p. 858\)](#).

To specify the number of seconds for Amazon RDS to delay replication to a Read Replica, use the mysql.rds_set_configuration stored procedure and specify the number of seconds to delay replication. The following example specifies that replication is delayed by at least one hour (3,600 seconds).

```
call mysql.rds_set_configuration('target delay', 3600);
```

The limit for the target_delay parameter is one day (86400 seconds).

Note

The target_delay parameter is only supported for Amazon RDS MySQL.

The target_delay parameter is not supported for Amazon RDS MySQL version 8.0.

mysql.rds_show_configuration

The number of hours that binary logs are retained.

Syntax

```
CALL mysql.rds_show_configuration;
```

Usage Notes

To verify the number of hours that Amazon RDS retains binary logs, use the mysql.rds_show_configuration stored procedure.

The mysql.rds_show_configuration procedure is available in these versions of Amazon RDS MySQL:

- MySQL 5.6

- MySQL 5.7
- MySQL 8.0

Examples

The following example displays the retention period:

```
call mysql.rds_show_configuration;
      name          value      description
      binlog retention hours    24      binlog retention hours specifies the
duration in hours before binary logs are automatically deleted.
```

mysql.rds_kill

Terminates a connection to the MySQL server.

Syntax

```
CALL mysql.rds_kill(processID);
```

Parameters

processID

The identity of the connection thread to be terminated.

Usage Notes

Each connection to the MySQL server runs in a separate thread. To terminate a connection, use the `mysql.rds_kill` procedure and pass in the thread ID of that connection. To obtain the thread ID, use the MySQL `SHOW PROCESSLIST` command.

Examples

The following example terminates a connection with a thread ID of 4243:

```
call mysql.rds_kill(4243);
```

mysql.rds_kill_query

Terminates a query running against the MySQL server.

Syntax

```
CALL mysql.rds_kill_query(queryID);
```

Parameters

queryID

The identity of the query to be terminated.

Usage Notes

To terminate a query running against the MySQL server, use the `mysql_rds_kill_query` procedure and pass in the ID of that query. To obtain the query ID, use the MySQL [INFORMATION_SCHEMA PROCESSLIST](#) command. The connection to the MySQL server is retained.

Examples

The following example terminates a query with a thread ID of 230040:

```
call mysql.rds_kill_query(230040);
```

mysql.rds_rotate_general_log

Rotates the `mysql.general_log` table to a backup table. For more information, see [MySQL Database Log Files \(p. 405\)](#).

Syntax

```
CALL mysql.rds_rotate_general_log;
```

Usage Notes

You can rotate the `mysql.general_log` table to a backup table by calling the `mysql.rds_rotate_general_log` procedure. When log tables are rotated, the current log table is copied to a backup log table and the entries in the current log table are removed. If a backup log table already exists, then it is deleted before the current log table is copied to the backup. You can query the backup log table if needed. The backup log table for the `mysql.general_log` table is named `mysql.general_log_backup`.

mysql.rds_rotate_slow_log

Rotates the `mysql.slow_log` table to a backup table. For more information, see [MySQL Database Log Files \(p. 405\)](#).

Syntax

```
CALL mysql.rds_rotate_slow_log;
```

Usage Notes

You can rotate the `mysql.slow_log` table to a backup table by calling the `mysql.rds_rotate_slow_log` procedure. When log tables are rotated, the current log table is copied to a backup log table and the entries in the current log table are removed. If a backup log table already exists, then it is deleted before the current log table is copied to the backup.

You can query the backup log table if needed. The backup log table for the `mysql.slow_log` table is named `mysql.slow_log_backup`.

mysql.rds_enable_gsh_collector

Enables the Global Status History (GoSH) to take default snapshots at intervals specified by `rds_set_gsh_collector`. For more information, see [Managing the Global Status History \(p. 883\)](#).

Syntax

```
CALL mysql.rds_enable_gsh_collector;
```

mysql.rds_set_gsh_collector

Specifies the interval, in minutes, between snapshots taken by the Global Status History (GoSH). Default value is For more information, see [Managing the Global Status History \(p. 883\)](#).

Syntax

```
CALL mysql.rds_set_gsh_collector(intervalPeriod);
```

Parameters

intervalPeriod

The interval, in minutes, between snapshots. Default value is

mysql.rds_disable_gsh_collector

Disables snapshots taken by the Global Status History (GoSH). For more information, see [Managing the Global Status History \(p. 883\)](#).

Syntax

```
CALL mysql.rds_disable_gsh_collector;
```

mysql.rds_collect_global_status_history

Takes a snapshot on demand for the Global Status History (GoSH). For more information, see [Managing the Global Status History \(p. 883\)](#).

Syntax

```
CALL mysql.rds_collect_global_status_history;
```

mysql.rds_enable_gsh_rotation

Enables rotation of the contents of the `mysql.global_status_history` table to `mysql.global_status_history_old` at intervals specified by `rds_set_gsh_rotation`. For more information, see [Managing the Global Status History \(p. 883\)](#).

Syntax

```
CALL mysql.rds_enable_gsh_rotation;
```

mysql.rds_set_gsh_rotation

Specifies the interval, in days, between rotations of the `mysql.global_status_history` table. Default value is 7. For more information, see [Managing the Global Status History \(p. 883\)](#).

Syntax

```
CALL mysql.rds_set_gsh_rotation(intervalPeriod);
```

Parameters

intervalPeriod

The interval, in days, between table rotations. Default value is 7.

mysql.rds_disable_gsh_rotation

Disables rotation of the `mysql.global_status_history` table. For more information, see [Managing the Global Status History \(p. 883\)](#).

Syntax

```
CALL mysql.rds_disable_gsh_rotation;
```

mysql.rds_rotate_global_status_history

Rotates the contents of the `mysql.global_status_history` table to `mysql.global_status_history_old` on demand. For more information, see [Managing the Global Status History \(p. 883\)](#).

Syntax

```
CALL mysql.rds_rotate_global_status_history;
```

Oracle on Amazon RDS

Amazon RDS supports DB instances running several versions and editions of Oracle Database. You can use the following versions and editions:

- Oracle 19c, Version 19.0.0.0
- Oracle 18c, Version 18.0.0.0
- Oracle 12c, Version 12.2.0.1
- Oracle 12c, Version 12.1.0.2
- Oracle 11g, Version 11.2.0.4

Amazon RDS also currently supports the following versions and editions that are on deprecation paths, because Oracle no longer provides patches for them:

- Oracle 12c, Version 12.1.0.1 ([Deprecation of Oracle 12.1.0.1 \(p. 935\)](#))
- Oracle 11g, Version 11.2.0.3 ([Deprecation of Oracle 11.2.0.3 \(p. 935\)](#))
- Oracle 11g, Version 11.2.0.2 ([Deprecation of Oracle 11.2.0.2 \(p. 934\)](#))

You can create DB instances and DB snapshots, point-in-time restores and automated or manual backups. DB instances running Oracle can be used inside a VPC. You can also enable various options to add additional features to your Oracle DB instance. Amazon RDS supports Multi-AZ deployments for Oracle as a high-availability, failover solution.

To deliver a managed service experience, Amazon RDS doesn't provide shell access to DB instances, and it restricts access to certain system procedures and tables that require advanced privileges. Amazon RDS supports access to databases on a DB instance using any standard SQL client application such as Oracle SQL Plus. Amazon RDS doesn't enable direct host access to a DB instance by using Telnet or Secure Shell (SSH).

When you create a DB instance, the master account that you use to create the instance gets DBA user privileges (with some limitations). Use this account for any administrative tasks such as creating additional user accounts in the database. The SYS user, SYSTEM user, and other administrative accounts can't be used.

Before creating a DB instance, complete the steps in the [Setting Up for Amazon RDS \(p. 5\)](#) section of this guide.

Common Management Tasks for Oracle on Amazon RDS

The following are the common management tasks you perform with an Amazon RDS Oracle DB instance, with links to relevant documentation for each task.

Task Area	Relevant Documentation
Instance Classes, Storage, and PIOPS If you are creating a DB instance for production purposes, you should understand how instance classes, storage types, and Provisioned IOPS work in Amazon RDS.	DB Instance Class Support for Oracle (p. 917) Amazon RDS Storage Types (p. 135)

Task Area	Relevant Documentation
Multi-AZ Deployments A production DB instance should use Multi-AZ deployments. Multi-AZ deployments provide increased availability, data durability, and fault tolerance for DB instances.	High Availability (Multi-AZ) for Amazon RDS (p. 144)
Amazon Virtual Private Cloud (VPC) If your AWS account has a default VPC, then your DB instance is automatically created inside the default VPC. If your account doesn't have a default VPC, and you want the DB instance in a VPC, create the VPC and subnet groups before you create the DB instance.	Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform (p. 506) Working with a DB Instance in a VPC (p. 514)
Security Groups By default, DB instances are created with a firewall that prevents access to them. You therefore must create a security group with the correct IP addresses and network configuration to access the DB instance. The security group you create depends on what Amazon EC2 platform your DB instance is on, and whether you will access your DB instance from an Amazon EC2 instance. In general, if your DB instance is on the <i>EC2-Classic</i> platform, you will need to create a DB security group; if your DB instance is on the <i>EC2-VPC</i> platform, you will need to create a VPC security group.	Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform (p. 506) Controlling Access with Security Groups (p. 486)
Parameter Groups If your DB instance is going to require specific database parameters, you should create a parameter group before you create the DB instance.	Working with DB Parameter Groups (p. 211)
Option Groups If your DB instance is going to require specific database options, you should create an option group before you create the DB instance.	Options for Oracle DB Instances (p. 1004)
Connecting to Your DB Instance After creating a security group and associating it to a DB instance, you can connect to the DB instance using any standard SQL client application such as Oracle SQL Plus.	Connecting to a DB Instance Running the Oracle Database Engine (p. 955)
Backup and Restore You can configure your DB instance to take automated backups, or take manual snapshots, and then restore instances from the backups or snapshots.	Backing Up and Restoring Amazon RDS DB Instances (p. 253)
Monitoring You can monitor an Oracle DB instance by using CloudWatch Amazon RDS metrics, events, and enhanced monitoring.	Viewing DB Instance Metrics (p. 304) Viewing Amazon RDS Events (p. 388)

Task Area	Relevant Documentation
Log Files You can access the log files for your Oracle DB instance.	Amazon RDS Database Log Files (p. 390)

There are also advanced tasks and optional features for working with Oracle DB instances. For more information, see the following documentation:

- For information on common DBA tasks for Oracle on Amazon RDS, see [Common DBA Tasks for Oracle DB Instances \(p. 1080\)](#).
- For information on Oracle GoldenGate support, see [Using Oracle GoldenGate with Amazon RDS \(p. 1153\)](#).
- For information on Siebel Customer Relationship Management (CRM) support, see [Installing a Siebel Database on Oracle on Amazon RDS \(p. 1170\)](#).

Oracle Licensing

There are two licensing options available for Amazon RDS for Oracle: License Included and Bring Your Own License (BYOL). After you create an Oracle DB instance on Amazon RDS, you can change the licensing model by using the [AWS Management Console](#), the Amazon RDS API [ModifyDBInstance](#) operation, or the AWS CLI [modify-db-instance](#) command.

License Included

In the License Included model, you don't need to purchase Oracle licenses separately. AWS holds the license for the Oracle database software. In this model, if you have an AWS Support account with case support, you contact AWS Support for both Amazon RDS and Oracle Database service requests.

The License Included model is supported on Amazon RDS for the following Oracle database editions:

- Oracle Database Standard Edition One (SE1)
- Oracle Database Standard Edition Two (SE2)

Bring Your Own License (BYOL)

In the Bring Your Own License model, you can use your existing Oracle Database licenses to run Oracle deployments on Amazon RDS. You must have the appropriate Oracle Database license (with Software Update License and Support) for the DB instance class and Oracle Database edition you wish to run. You must also follow Oracle's policies for licensing Oracle Database software in the cloud computing environment. For more information on Oracle's licensing policy for Amazon EC2, see [Licensing Oracle Software in the Cloud Computing Environment](#).

In this model, you continue to use your active Oracle support account, and you contact Oracle directly for Oracle Database service requests. If you have an AWS Support account with case support, you can contact AWS Support for Amazon RDS issues. Amazon Web Services and Oracle have a multi-vendor support process for cases which require assistance from both organizations.

The Bring Your Own License model is supported on Amazon RDS for the following Oracle database editions:

- Oracle Database Enterprise Edition (EE)
- Oracle Database Standard Edition (SE)

- Oracle Database Standard Edition One (SE1)
- Oracle Database Standard Edition Two (SE2)

Licensing Oracle Multi-AZ Deployments

Amazon RDS supports Multi-AZ deployments for Oracle as a high-availability, failover solution. We recommend Multi-AZ for production workloads. For more information, see [High Availability \(Multi-AZ\) for Amazon RDS \(p. 144\)](#).

If you use the Bring Your Own License model, you must have a license for both the primary DB instance and the standby DB instance in a Multi-AZ deployment.

Migrating Between Oracle Editions

For the BYOL model, you can migrate from any Standard Edition (SE, SE1, or SE2) to Enterprise Edition (EE), assuming you have an unused Oracle license appropriate for the edition and class of DB instance you plan to run. You can't migrate from Enterprise Edition to other editions.

To change the edition and retain your data

1. Create a snapshot of the DB instance.

For more information, see [Creating a DB Snapshot \(p. 264\)](#).

2. Restore the snapshot to a new DB instance, and select the Oracle database edition you want to use.

For more information, see [Restoring from a DB Snapshot \(p. 266\)](#).

3. (Optional) Delete the old DB instance, unless you want to keep it running and have the appropriate Oracle Database licenses for it.

For more information, see [Deleting a DB Instance \(p. 174\)](#).

DB Instance Class Support for Oracle

The computation and memory capacity of a DB instance is determined by its DB instance class. The DB instance class you need depends on your processing power and memory requirements. For more information, see [Choosing the DB Instance Class \(p. 107\)](#).

The following are the DB instance classes supported for Oracle.

Oracle Edition	19c, 18c, and 12c Version 12.2.0.1 Support	12c Version 12.1.0.2 Support	11g Version 11.2.0.4 Support
Enterprise Edition (EE)	db.m5.large–db.m5.24xlarge	db.m5.large–db.m5.24xlarge	db.m5.large–db.m5.24xlarge
Bring Your Own License (BYOL)	db.m4.large–db.m4.16xlarge	db.m4.large–db.m4.16xlarge	db.m4.large–db.m4.16xlarge
	db.z1d.large–db.z1d.12xlarge	db.z1d.large–db.z1d.12xlarge	db.z1d.large–db.z1d.12xlarge
	db.x1e.xlarge–db.x1e.32xlarge	db.x1e.xlarge–db.x1e.32xlarge	db.x1e.xlarge–db.x1e.32xlarge

Oracle Edition	19c, 18c, and 12c Version 12.2.0.1 Support	12c Version 12.1.0.2 Support	11g Version 11.2.0.4 Support
	db.x1.16xlarge– db.x1.32xlarge db.r5.large–db.r5.24xlarge db.r4.large–db.r4.16xlarge db.t3.small–db.t3.2xlarge	db.x1.16xlarge– db.x1.32xlarge db.r5.large–db.r5.24xlarge db.r4.large–db.r4.16xlarge db.t3.micro–db.t3.2xlarge	db.x1.16xlarge– db.x1.32xlarge db.r5.large–db.r5.24xlarge db.r4.large–db.r4.16xlarge db.t3.micro–db.t3.2xlarge
Standard Edition 2 (SE2)	db.m5.large–db.m5.4xlarge db.m4.large–db.m4.4xlarge	db.m5.large–db.m5.4xlarge db.m4.large–db.m4.4xlarge	—
Bring Your Own License (BYOL)	db.z1d.large– db.z1d.3xlarge db.x1e.xlarge– db.x1e.4xlarge db.r5.large–db.r5.4xlarge db.r4.large–db.r4.4xlarge db.t3.small–db.t3.2xlarge	db.z1d.large– db.z1d.3xlarge db.x1e.xlarge– db.x1e.4xlarge db.r5.large–db.r5.4xlarge db.r4.large–db.r4.4xlarge db.t3.micro–db.t3.2xlarge	
Standard Edition 2 (SE2)	db.m5.large–db.m5.4xlarge db.m4.large–db.m4.4xlarge	db.m5.large–db.m5.4xlarge db.m4.large–db.m4.4xlarge	—
License Included	db.r5.large–db.r5.4xlarge db.r4.large–db.r4.4xlarge db.t3.small–db.t3.2xlarge	db.r5.large–db.r5.4xlarge db.r4.large–db.r4.4xlarge db.t3.micro–db.t3.2xlarge	
Standard Edition 1 (SE1) Bring Your Own License (BYOL)	—	—	db.m5.large–db.m5.4xlarge db.m4.large–db.m4.4xlarge db.z1d.large– db.z1d.3xlarge db.x1e.xlarge– db.x1e.4xlarge db.r5.large–db.r5.4xlarge db.r4.large–db.r4.4xlarge db.t3.micro–db.t3.2xlarge
Standard Edition 1 (SE1) License Included	—	—	db.m5.large–db.m5.4xlarge db.m4.large–db.m4.4xlarge db.r5.large–db.r5.4xlarge db.t3.micro–db.t3.2xlarge

Oracle Edition	19c, 18c, and 12c Version 12.2.0.1 Support	12c Version 12.1.0.2 Support	11g Version 11.2.0.4 Support
Standard Edition (SE)	—	—	db.m5.large–db.m5.8xlarge db.m4.large–db.m4.4xlarge db.z1d.large–db.z1d.6xlarge db.x1e.xlarge–db.x1e.8xlarge db.r5.large–db.r5.8xlarge db.r4.large–db.r4.8xlarge db.t3.micro–db.t3.2xlarge
Bring Your Own License (BYOL)			

Note

We encourage all bring-your-own-license customers to consult their licensing agreement to assess the impact of Amazon RDS for Oracle deprecations. For more information on the compute capacity of DB Instance classes supported by Amazon RDS for Oracle, see [Choosing the DB Instance Class \(p. 107\)](#) and [Configuring the Processor for a DB Instance Class \(p. 117\)](#).

Deprecation of db.t2 DB Instance Classes for Oracle

The db.t2 DB instance classes are on a deprecation path for Amazon RDS for Oracle. The db.t2 DB instance classes will be replaced by the better performing db.t3 DB instance classes that are generally available at a lower cost. Starting on September 15, 2019, Amazon RDS for Oracle will disable all new creates of the db.t2 DB instance classes. Starting on January 15, 2020, Amazon RDS for Oracle will automatically scale db.t2 DB instances to comparable db.t3 DB instance classes.

If you have DB instances that use db.t2 DB instance classes, Amazon RDS will modify each one automatically to use a comparable DB instance class that is not deprecated. You can change the DB instance class for a DB instance yourself by modifying the DB instance. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

If you have DB snapshots of DB instances that were using db.t2 DB instance classes, you can choose a DB instance class that is not deprecated when you restore the DB snapshots. For more information, see [Restoring from a DB Snapshot \(p. 266\)](#).

Note

The db.t3 DB instance classes have hyper-threading enabled by default. When DB instances running the db.t2 DB instance class are migrated, the number of vCPUs is set automatically to the default number of the comparable db.t3 DB instance class. To learn more about the vCPU management features available on Amazon RDS for Oracle, and the default settings for each db.t3 DB instance class, see [Configuring the Processor for a DB Instance Class \(p. 117\)](#).

Deprecated db.m3 and db.r3 DB Instance Classes for Oracle

The db.m3 and db.r3 DB instance classes are deprecated for Amazon RDS for Oracle. These DB instance classes have been replaced by better performing DB instance classes that are generally available at a lower cost. Starting on September 30, 2019, Amazon RDS for Oracle will automatically scale DB instances to DB instance classes that are not deprecated.

If you have DB instances that use db.m3 and db.r3 DB instance classes, Amazon RDS will modify each one automatically to use a comparable DB instance class that is not deprecated. You can change the DB instance class for a DB instance yourself by modifying the DB instance. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

If you have DB snapshots of DB instances that were using db.m3 or db.r3 DB instance classes, you can choose a DB instance class that is not deprecated when you restore the DB snapshots. For more information, see [Restoring from a DB Snapshot \(p. 266\)](#).

Deprecated db.m1 and db.m2 DB Instance Classes for Oracle

The db.m1 and db.m2 DB instance classes are deprecated for Amazon RDS for Oracle. These DB instance classes have been replaced by better performing DB instance classes that are generally available at a lower cost. Starting on September 12, 2018, Amazon RDS for Oracle will automatically scale DB instances to DB instance classes that are not deprecated.

If you have DB instances that use db.m1 and db.m2 DB instance classes, Amazon RDS will modify each one automatically to use a comparable DB instance class that is not deprecated. You can change the DB instance class for a DB instance yourself by modifying the DB instance. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

If you have DB snapshots of DB instances that were using db.m1 or db.m2 DB instance classes, you can choose a DB instance class that is not deprecated when you restore the DB snapshots. For more information, see [Restoring from a DB Snapshot \(p. 266\)](#).

Oracle Security

The Oracle database engine uses role-based security. A role is a collection of privileges that can be granted to or revoked from a user. A predefined role, named *DBA*, normally allows all administrative privileges on an Oracle database engine. The following privileges are not available for the DBA role on an Amazon RDS DB instance using the Oracle engine:

- Alter database
- Alter system
- Create any directory
- Drop any directory
- Grant any privilege
- Grant any role

When you create a DB instance, the master account that you use to create the instance gets DBA user privileges (with some limitations). Use this account for any administrative tasks such as creating additional user accounts in the database. The *SYS* user, *SYSTEM* user, and other administrative accounts can't be used.

Amazon RDS Oracle supports SSL/TLS encrypted connections and also the Oracle Native Network Encryption (NNE) option to encrypt connections between your application and your Oracle DB instance. For more information about using SSL with Oracle on Amazon RDS, see [Using SSL with an Oracle DB Instance \(p. 921\)](#). For more information about the Oracle Native Network Encryption option, see [Oracle Native Network Encryption \(p. 1049\)](#).

Using SSL with an Oracle DB Instance

Secure Sockets Layer (SSL) is an industry standard protocol used for securing network connections between client and server. After SSL version 3.0, the name was changed to Transport Layer Security (TLS), but it is still often referred to as SSL and we refer to the protocol as SSL. Amazon RDS supports SSL encryption for Oracle DB instances. Using SSL, you can encrypt a connection between your application client and your Oracle DB instance. SSL support is available in all AWS regions for Oracle.

You enable SSL encryption for an Oracle DB instance by adding the Oracle SSL option to the option group associated with the DB instance. Amazon RDS uses a second port, as required by Oracle, for SSL connections which allows both clear text and SSL-encrypted communication to occur at the same time between a DB instance and an Oracle client. For example, you can use the port with clear text communication to communicate with other resources inside a VPC while using the port with SSL-encrypted communication to communicate with resources outside the VPC.

For more information, see [Oracle Secure Sockets Layer \(p. 1051\)](#).

Note

You can't use both SSL and Oracle native network encryption (NNE) on the same DB instance. Before you can use SSL encryption, you must disable any other connection encryption.

Oracle 19c with Amazon RDS

Amazon RDS supports Oracle version 19c, which includes Oracle Enterprise Edition and Oracle Standard Edition Two.

Oracle 19c version 19.0.0.0 includes many new features and updates from the previous version. In this section, you can find the features and changes important to using Oracle 19c version 19.0.0.0 on Amazon RDS. For a complete list of the changes, see the [Oracle Database 19c](#) documentation. For a complete list of features supported by each Oracle 19c edition, see [Permitted Features, Options, and Management Packs by Oracle Database Offering](#) in the Oracle documentation.

Amazon RDS Parameter Changes for Oracle 19c Version 19.0.0.0

Oracle 19c version 19.0.0.0 includes several new parameters and parameters with new ranges and new default values.

The following table shows the new Amazon RDS parameters for Oracle 19c version 19.0.0.0.

Name	Values	Modifi:	Description
<code>lob_signature_enable</code>	TRUE, FALSE (default)	Y	Enables or disables the LOB locator signature feature.
<code>max_datapump_parallel_per_job</code>	1 to 1024, or AUTO	Y	Specifies the maximum number of parallel processes allowed for each Oracle Data Pump job.

The `compatible` parameter has a new maximum value for Oracle 19c version 19.0.0.0 on Amazon RDS. The following table shows the new default value.

Parameter Name	Oracle 19c Version 19.0.0.0 Maximum Value	Oracle 18c Version 18.0.0.0 Maximum Value
compatible	19.0.0	18.0.0

The following parameters were removed in Oracle 19c Version 19.0.0.0:

- `exafusion_enabled`
- `max_connections`
- `o7_dictionary_access`

Oracle 18c with Amazon RDS

Amazon RDS supports Oracle version 18c, which includes Oracle Enterprise Edition and Oracle Standard Edition Two.

Oracle 18c version 18.0.0.0 includes many new features and updates from the previous version. In this section, you can find the features and changes important to using Oracle 18c version 18.0.0.0 on Amazon RDS. For a complete list of the changes, see the [Oracle Database 18c](#). For a complete list of features supported by each Oracle 18c edition, see [Permitted Features, Options, and Management Packs by Oracle Database Offering](#) in the Oracle documentation.

Amazon RDS Parameter Changes for Oracle 18c Version 18.0.0.0

Oracle 18c version 18.0.0.0 includes several new parameters and parameters with new ranges and new default values.

The following table shows the new Amazon RDS parameters for Oracle 18c version 18.0.0.0.

Name	Values	Modifi	Description
adg_account_info_tracking	LOCAL (default), GLOBAL	N	Controls login attempts of users on Active Data Guard Standby databases. It extends the control of user account security information.
inmemory_optimized_arithmetic	ENABLE, DISABLE (default)	Y	Encodes the NUMBER data type in in-memory tables compressed with QUERY LOW as a fixed-width native integer scaled by a common exponent.
optimizer_ignore_hints	TRUE, FALSE (default)	Y	Specifies whether embedded hints are ignored.
optimizer_ignore_parallel_hints	TRUE, FALSE (default)	Y	Specifies that embedded parallel hints are ignored.
parallel_min_degree	1 (default) and higher or CPU	Y	Controls the minimum degree of parallelism computed by automatic degree of parallelism.

The `compatible` parameter has a new default value for Oracle 18c version 18.0.0.0 on Amazon RDS. The following table shows the new default value.

Parameter Name	Oracle 18c Version 18.0.0.0 Default Value	Oracle 12c Version 12.2.0.1 Default Value
<code>compatible</code>	18.0.0	12.2.0

The following parameters were removed in Oracle 18c Version 18.0.0.0:

- `standby_archive_dest`
- `utl_file_dir`

Oracle 12c with Amazon RDS

Amazon RDS supports Oracle version 12c, which includes Oracle Enterprise Edition and Oracle Standard Edition Two. Oracle version 12c includes two major versions:

- [Oracle 12c Version 12.2.0.1 with Amazon RDS \(p. 923\)](#)
- [Oracle 12c Version 12.1.0.2 with Amazon RDS \(p. 926\)](#)

Oracle 12c Version 12.2.0.1 with Amazon RDS

Oracle 12c version 12.2.0.1 includes many new features and updates from the previous version. In this section, you can find the features and changes important to using Oracle 12c version 12.2.0.1 on Amazon RDS. For a complete list of the changes, see the [Oracle 12c version 12.2 documentation](#). For a complete list of features supported by each Oracle 12c edition, see [Permitted Features, Options, and Management Packs by Oracle Database Offering](#) in the Oracle documentation.

Amazon RDS Parameter Changes for Oracle 12c Version 12.2.0.1

Oracle 12c version 12.2.0.1 includes 20 new parameters in addition to several parameters with new ranges and new default values.

The following table shows the new Amazon RDS parameters for Oracle 12c version 12.2.0.1.

Name	Values	Modifi	Description
<code>allow_global_dblinks</code>	TRUE, FALSE (default)	Y	Specifies whether LDAP lookup for database links is allowed for the database.
<code>approx_for_aggregation</code>	TRUE, FALSE (default)	Y	Replaces exact query processing for aggregation queries with approximate query processing.
<code>approx_for_count_distinct</code>	TRUE, FALSE (default)	Y	Automatically replaces COUNT (DISTINCT <i>expr</i>) queries with APPROX_COUNT_DISTINCT queries.
<code>approx_for_percentile</code>	NONE (default), PERCENTILE_CONT, PERCENTILE_CONT	Y	Converts exact percentile functions to their approximate percentile function counterparts.

Name	Values	Modifiable	Description
	DETERMINISTIC, PERCENTILE_DISC, PERCENTILE_DISC DETERMINISTIC, ALL, ALL DETERMINISTIC		
cursor_invalidation	DEFERRED, IMMEDIATE (default)	Y	Controls whether deferred cursor invalidation or immediate cursor invalidation is used for DDL statements by default.
data_guard_sync_latency	0 (default) to the number of seconds specified by the NET_TIMEOUT attribute for the LOG_ARCHIVE_DEST_n parameter	Y	Controls how many seconds the Log Writer (LGWR) process waits beyond the response of the first in a series of Oracle Data Guard SYNC redo transport mode connections.
data_transfer_cache_size	0 – 512M, rounded up to the next granule size	Y	Sets the size of the data transfer cache (in bytes) used to receive data blocks (typically from a primary database in an Oracle Data Guard environment) for consumption by an instance during execution of an RMAN RECOVER ... NONLOGGED BLOCK command.
inmemory_adg_enabled	TRUE (default), FALSE	Y	Indicates whether in-memory for Active Data Guard is enabled in addition to the in-memory cache size.
inmemory_expressions_usage	STATIC_ONLY, DYNAMIC_ONLY,ENABLE (default), DISABLE	Y	Controls which In-Memory Expressions (IM expressions) are populated into the In-Memory Column Store (IM column store) and are available for queries.
inmemory_virtual_columns	ENABLE, MANUAL (default), DISABLE	Y	Controls which In-Memory Expressions (IM expressions) are populated into the In-Memory Column Store (IM column store) and are available for queries.
instance_abort_delay_time	0 (default) and higher	Y	Specifies how much time to delay an internal initiated instance abort (in seconds), such as when a fatal process dies or an unrecoverable instance error occurs.
instance_mode	READ-WRITE (default), READ-ONLY, READ-MOSTLY	N	Indicates whether the instance is read-write, read-only, or read-mostly.
long_module_action	TRUE (default), FALSE	Y	Enables the use of longer lengths for modules and actions.

Name	Values	Modifiable	Description
max_idle_time	0 (default) to the maximum integer. The value of 0 indicates that there is no limit.	Y	Specifies the maximum number of minutes that a session can be idle. After that point, the session is automatically terminated.
optimizer_adaptive_plans	TRUE (default), FALSE	Y	Controls adaptive plans. Adaptive plans are execution plans built with alternative choices that are decided at run time based on statistics collected as the query executes.
optimizer_adaptive_statistics	TRUE, FALSE (default)	Y	Controls adaptive statistics. Some query shapes are too complex to rely on base table statistics alone, so the optimizer augments these statistics with adaptive statistics.
outbound_dblink_protocols	ALL (default), NONE, TCP, TCPS, IPC	N	Specifies the network protocols allowed for communicating for outbound database links in the database.
resource_manage_goldengate	TRUE, FALSE (default)	Y	Determines whether Oracle GoldenGate apply processes in the database are resource managed.
standby_db_preserve_states	NONE (default), SESSION, ALL	N	Controls whether user sessions and other internal states of the instance are retained when a readable physical standby database is converted to a primary database.
uniform_log_timestamp_format	TRUE (default), FALSE	Y	Specifies that a uniform timestamp format be used in Oracle Database trace (.trc) files and log files (such as the alert log).

The `compatible` parameter has a new default value for Oracle 12c version 12.2.0.1 on Amazon RDS. The following table shows the new default value.

Parameter Name	Oracle 12c Version 12.2.0.1 Default Value	Oracle 12c Version 12.1.0.2 Default Value
compatible	12.2.0	12.0.0

The `optimizer_features_enable` parameter has a new value range for Oracle 12c version 12.2.0.1 on Amazon RDS. For the old and new value ranges, see the following table.

Parameter Name	12c Version 12.2.0.1 Range	12c Version 12.1.0.2 Range
optimizer_features_enable	8.0.0 to 12.2.0.1	8.0.0 to 12.1.0.2

The following parameters were removed in Oracle 12c Version 12.2.0.1:

- global_context_pool_size
- max_enabled_roles
- optimizer_adaptive_features
- parallel_automatic_tuning
- parallel_degree_level
- use_indirect_data_buffers

The following parameter is not supported in Oracle 12c Version 12.2.0.1:

- sec_case_sensitive_logon

Amazon RDS Security Changes for Oracle 12c Version 12.2.0.1

In Oracle 12c version 12.2.0.1, direct grant of the privilege `ADMINISTER DATABASE TRIGGER` is required for the owners of database-level triggers. During a major version upgrade to Oracle 12c version 12.2.0.1, Amazon RDS grants this privilege to any user that owns a trigger so that the trigger owner has the required privileges. For more information, see the My Oracle Support document [2275535.1](#).

Oracle 12c Version 12.1.0.2 with Amazon RDS

Oracle 12c version 12.1.0.2 brings over 500 new features and updates from the previous version. In this section, you can find the features and changes important to using Oracle 12c version 12.1.0.2 on Amazon RDS. For a complete list of the changes, see the [Oracle 12c version 12.1 documentation](#). For a complete list of features supported by each Oracle 12c edition, see [Permitted Features, Options, and Management Packs by Oracle Database Edition](#) in the Oracle documentation.

Oracle 12c version 12.1.0.2 includes 16 new parameters that impact your Amazon RDS DB instance, and also 18 new system privileges, several no longer supported packages, and several new option group settings. For more information on these changes, see the following sections.

Amazon RDS Parameter Changes for Oracle 12c Version 12.1.0.2

Oracle 12c version 12.1.0.2 includes 16 new parameters in addition to several parameters with new ranges and new default values.

The following table shows the new Amazon RDS parameters for Oracle 12c version 12.1.0.2.

Name	Values	Modifi	Description
<code>connection_brokers</code>	<code>CONNECTION_BROKERSN = broker_description[,...]</code>		Specifies connection broker types, the number of connection brokers of each type, and the maximum number of connections per broker.
<code>db_index_compression_inheritance</code>	<code>TABLESPACE, TABL, ALL, NONE</code>	Y	Displays the options that are set for table or tablespace level compression inheritance.
<code>db_big_table_cache_percent_target</code>	<code>0-90</code>	Y	Specifies the cache section target size for automatic big table caching, as a percentage of the buffer cache.

Name	Values	Modifiable	Description
heat_map	ON, OFF	Y	Enables the database to track read and write access of all segments and modification of database blocks that are due to data manipulation language (DML) and data definition language (DDL) statements.
inmemory_clause_default	INMEMORY, NO INMEMORY	Y	INMEMORY_CLAUSE_DEFAULT enables you to specify a default In-Memory Column Store (IM column store) clause for new tables and materialized views.
inmemory_clause_default_memcomp	NO MEMCOMPRESS, MEMCOMPRESS FOR DML, MEMCOMPRESS FOR QUERY, MEMCOMPRESS FOR QUERY LOW, MEMCOMPRESS FOR QUERY HIGH, MEMCOMPRESS FOR CAPACITY, MEMCOMPRESS FOR CAPACITY LOW, MEMCOMPRESS FOR CAPACITY HIGH	Y	See INMEMORY_CLAUSE_DEFAULT.
inmemory_clause_default_priority	PRIORITY LOW, PRIORITY MEDIUM, PRIORITY HIGH, PRIORITY CRITICAL, PRIORITY NONE	Y	See INMEMORY_CLAUSE_DEFAULT.
inmemory_force	DEFAULT, OFF	Y	INMEMORY_FORCE allows you to specify whether tables and materialized view that are specified as INMEMORY are populated into the In-Memory Column Store (IM column store) or not.
inmemory_max_populate_servers	Null	N	INMEMORY_MAX_POPULATE_SERVERS specifies the maximum number of background populate servers to use for In-Memory Column Store (IM column store) population, so that these servers don't overload the rest of the system.
inmemory_query	ENABLE (default), DISABLE	Y	INMEMORY_QUERY is used to enable or disable in-memory queries for the entire database at the session or system level.

Name	Values	Modifiable	Description
inmemory_size	0, 104857600-274877906944	Y	INMEMORY_SIZE sets the size of the In-Memory Column Store (IM column store) on a database instance.
inmemory_trickle_repopulate_servers_percent	0 to 50 percent	Y	INMEMORY_TRICKLE_REPOPULATE_SERVERS_PERCENT limits the maximum number of background populate servers used for In-Memory Column Store (IM column store) repopulation. This limit is applied because trickle repopulation is designed to use only a small percentage of the populate servers.
max_string_size	STANDARD (default), EXTENDED	N	Controls the maximum size of VARCHAR2, NVARCHAR2, and RAW. For more information, see Using Extended Data Types (p. 940) .
optimizer_adaptive_features	TRUE (default), FALSE	Y	Enables or disables all of the adaptive optimizer features.
optimizer_adaptive_reporting_only	TRUE, FALSE (default)	Y	Controls reporting-only mode for adaptive optimizations.
pdb_file_name_convert	There is no default value.	N	Maps names of existing files to new file names.
pga_aggregate_limit	1-max of memory	Y	Specifies a limit on the aggregate PGA memory consumed by the instance.
processor_group_name	There is no default value.	N	Instructs the database instance to run itself within the specified operating system processor group.
spatial_vector_acceleration	TRUE, FALSE	N	Enables or disables the spatial vector acceleration, part of spatial option.
temp_undo_enabled	TRUE, FALSE (default)	Y	Determines whether transactions within a particular session can have a temporary undo log.
threaded_execution	TRUE, FALSE	N	Enables the multithreaded Oracle model, but prevents OS authentication.
unified_audit_sga_queue_size	1 MB - 30 MB	Y	Specifies the size of the system global area (SGA) queue for unified auditing.
use_dedicated_broker	TRUE, FALSE	N	Determines how dedicated servers are spawned.

Several parameter have new value ranges for Oracle 12c version 12.1.0.2 on Amazon RDS. For the old and new value ranges, see the following table.

Parameter Name	12c Version 12.1.0.2 Range	11g Range
audit_trail	os db [, extended] xml [, extended]	os db [, extended] xml [, extended] true false
compatible	For DB instances upgraded from Oracle 11g, automatically set to 12.0.0 on Amazon RDS unless a lower value is explicitly provided during the upgrade (as low as 11.2.0) For new Oracle 12c version 12.1.0.2 DB instances, starts with 12.0.0 on Amazon RDS	Starts with 11.2.0 on Amazon RDS
db_securefile	PERMITTED PREFERRED ALWAYS IGNORE FORCE	PERMITTED ALWAYS IGNORE FORCE
db_writer_processes	1-100	1-36
optimizer_features_e8.0.0 to 12.1.0.2		8.0.0 to 11.2.0.4
parallel_degree_polid	MANUAL,LIMITED,AUTO,ADAPTIVE	MANUAL,LIMITED,AUTO
parallel_min_server	0 to parallel_max_servers	CPU_COUNT * PARALLEL_THREADS_PER_CPU * 2 to parallel_max_servers

One parameter has a new default value for Oracle 12c on Amazon RDS. The following table shows the new default value.

Parameter Name	Oracle 12c Default Value	Oracle 11g Default Value
job_queue_processes	50	1000

Amazon RDS System Privileges for Oracle 12c Version 12.1.0.2

Several new system privileges have been granted to the system account for Oracle 12c version 12.1.0.2. These new system privileges include the following:

- ALTER ANY CUBE BUILD PROCESS
- ALTER ANY MEASURE FOLDER
- ALTER ANY SQL TRANSLATION PROFILE
- CREATE ANY SQL TRANSLATION PROFILE
- CREATE SQL TRANSLATION PROFILE
- DROP ANY SQL TRANSLATION PROFILE
- EM EXPRESS CONNECT
- EXEMPT DDL REDACTION POLICY
- EXEMPT DML REDACTION POLICY
- EXEMPT REDACTION POLICY
- LOGMINING
- REDEFINE ANY TABLE

- SELECT ANY CUBE BUILD PROCESS
- SELECT ANY MEASURE FOLDER
- USE ANY SQL TRANSLATION PROFILE

Amazon RDS Options for Oracle 12c Version 12.1.0.2

Several Oracle options changed between Oracle 11g and Oracle 12c version 12.1.0.2, though most of the options remain the same between the two versions. The Oracle 12c version 12.1.0.2 changes include the following:

- Oracle Enterprise Manager Database Express 12c replaced Oracle Enterprise Manager 11g Database Control. For more information, see [Oracle Enterprise Manager Database Express \(p. 1025\)](#).
- The option XMLDB is installed by default in Oracle 12c version 12.1.0.2. You no longer need to install this option yourself.

Amazon RDS PL/SQL Packages for Oracle 12c Version 12.1.0.2

Oracle 12c version 12.1.0.2 includes a number of new built-in PL/SQL packages. The packages included with Amazon RDS Oracle 12c version 12.1.0.2 include the following.

Package Name	Description
CTX_ANL	The CTX_ANL package is used with AUTO_LEXER and provides procedures for adding and dropping a custom dictionary from the lexer.
DBMS_APP_CONT	The DBMS_APP_CONT package provides an interface to determine if the in-flight transaction on a now unavailable session committed or not, and if the last call on that session completed or not.
DBMS_AUTO_REPORT	The DBMS_AUTO_REPORT package provides an interface to view SQL Monitoring and Real-time Automatic Database Diagnostic Monitor (ADDM) data that has been captured into Automatic Workload Repository (AWR).
DBMS_GOLDENGATE_AUTH	The DBMS_GOLDENGATE_AUTH package provides subprograms for granting privileges to and revoking privileges from GoldenGate administrators.
DBMS_HEAT_MAP	The DBMS_HEAT_MAP package provides an interface to externalize heatmaps at various levels of storage including block, extent, segment, object and tablespace.
DBMS_JLM	The DBMS_JLM package provides an interface for implementing Information Lifecycle Management (ILM) strategies using Automatic Data Optimization (ADO) policies.
DBMS_ILM_ADMIN	The DBMS_ILM_ADMIN package provides an interface to customize Automatic Data Optimization (ADO) policy execution.
DBMS_PART	The DBMS_PART package provides an interface for maintenance and management operations on partitioned objects.
DBMS_PRIVILEGE_CAPTURE	The DBMS_PRIVILEGE_CAPTURE package provides an interface to database privilege analysis.

Package Name	Description
DBMS_QOPATCH	The DBMS_QOPATCH package provides an interface to view the installed database patches.
DBMS_REDACT	The DBMS_REDACT package provides an interface to Oracle Data Redaction, which enables you to mask (redact) data that is returned from queries issued by low-privileged users or an application.
DBMS_SPD	The DBMS_SPD package provides subprograms for managing SQL plan directives (SPD).
DBMS_SQL_TRANSLATOR	The DBMS_SQL_TRANSLATOR package provides an interface for creating, configuring, and using SQL translation profiles.
DBMS_SQL_MONITOR	The DBMS_SQL_MONITOR package provides information about real-time SQL Monitoring and real-time Database Operation Monitoring.
DBMS_SYNC_REFRESH	The DBMS_SYNC_REFRESH package provides an interface to perform a synchronous refresh of materialized views.
DBMS_TSDP_MANAGE	The DBMS_TSDP_MANAGE package provides an interface to import and manage sensitive columns and sensitive column types in the database, and is used in conjunction with the DBMS_TSDP_PROTECT package with regard to transparent sensitive data protection (TSDP) policies. DBMS_TSDP_MANAGE is available with the Enterprise Edition only.
DBMS_TSDP_PROTECT	The DBMS_TSDP_PROTECT package provides an interface to configure transparent sensitive data protection (TSDP) policies in conjunction with the DBMS_TSDP_MANAGE package. DBMS_TSDP_PROTECT is available with the Enterprise Edition only.
DBMS_XDB_CONFIG	The DBMS_XDB_CONFIG package provides an interface for configuring Oracle XML DB and its repository.
DBMS_XDB_CONSTANTS	The DBMS_XDB_CONSTANTS package provides an interface to commonly used constants. Oracle recommends using constants instead of dynamic strings to avoid typographical errors.
DBMS_XDB_REPOS	The DBMS_XDB_REPOS package provides an interface to operate on the Oracle XML database Repository.
DBMS_XMLSCHEMA_ANNOTATE	The DBMS_XMLSCHEMA_ANNOTATE package provides an interface to manage and configure the structured storage model, mainly through the use of pre-registration schema annotations.
DBMS_XMLSTORAGE_MANAGE	The DBMS_XMLSTORAGE_MANAGE package provides an interface to manage and modify XML storage after schema registration has been completed.
DBMS_XSTREAM_ADMIN	The DBMS_XSTREAM_ADMIN package provides interfaces for streaming database changes between an Oracle database and other systems. XStream enables applications to stream out or stream in database changes.
DBMS_XSTREAM_AUTH	The DBMS_XSTREAM_AUTH package provides subprograms for granting privileges to and revoking privileges from XStream administrators.

Package Name	Description
UTL_CALL_STACK	The UTL_CALL_STACK package provides an interface to provide information about currently executing subprograms.

Oracle 12c Version 12.1.0.2 Packages Not Supported

Several Oracle 11g PL/SQL packages are not supported in Oracle 12c version 12.1.0.2. These packages include the following:

- DBMS_AUTO_TASK_IMMEDIATE
- DBMS_CDC_PUBLISH
- DBMS_CDC_SUBSCRIBE
- DBMS_EXPFIL
- DBMS_OBFUSCATION_TOOLKIT
- DBMS_RLMGR
- SDO_NET_MEM

Oracle Database Feature Support

Oracle Database supports a wide variety of features and capabilities, most of which are supported in Amazon RDS Oracle. Some features might have limited support or restricted privileges. Some features are only available in Enterprise Edition, and some require additional licenses. For more information about Oracle Database features for specific Oracle Database versions, see the *Oracle Database Licensing Information User Manual* for the version you're using.

Note

These lists are not exhaustive.

Amazon RDS Oracle supports the following Oracle Database features:

- Advanced Compression
- Application Express (APEX)

For more information, see [Oracle Application Express \(p. 1016\)](#).

- Automatic Memory Management
- Automatic Undo Management
- Automatic Workload Repository (AWR)

For more information, see [Working with Automatic Workload Repository \(AWR\) \(p. 1099\)](#).

- Active Data Guard with Maximum Performance within the same AWS Region

For more information, see [Working with Oracle Read Replicas for Amazon RDS \(p. 998\)](#).

- Data Redaction
- Database In-Memory (version 12.1 and later)
- Distributed Queries and Transactions
- Edition-Based Redefinition

For more information, see [Setting the Default Edition for a DB Instance \(p. 1100\)](#).

- Enterprise Manager Database Control (11g) and EM Express (12c)

For more information, see [Oracle Enterprise Manager \(p. 1024\)](#).

- Fine-Grained Auditing
- Flashback Table, Flashback Query, Flashback Transaction Query
- Import/export (legacy and Data Pump) and SQL*Loader

For more information, see [Importing Data into Oracle on Amazon RDS \(p. 987\)](#).

- Java Virtual Machine (JVM)

For more information, see [Oracle Java Virtual Machine \(p. 1037\)](#).

- Label Security (version 12.1 and later)

For more information, see [Oracle Label Security \(p. 1040\)](#).

- Locator

For more information, see [Oracle Locator \(p. 1043\)](#).

- Materialized Views
- Multimedia

For more information, see [Oracle Multimedia \(p. 1046\)](#).

- Network encryption

For more information, see [Oracle Native Network Encryption \(p. 1049\)](#) and [Oracle Secure Sockets Layer \(p. 1051\)](#).

- Partitioning
- Spatial and Graph

For more information, see [Oracle Spatial \(p. 1060\)](#).

- Star Query Optimization
- Streams and Advanced Queuing
- Summary Management – Materialized View Query Rewrite
- Text (File and URL data store types are not supported)
- Total Recall
- Transparent Data Encryption (TDE)

For more information, see [Oracle Transparent Data Encryption \(p. 1074\)](#).

- Unified Auditing, Mixed Mode (version 12.2 and later)

For more information, see [Mixed Mode Auditing](#) in the Oracle documentation.

- XML DB (without the XML DB Protocol Server)

For more information, see [Oracle XML DB \(p. 1078\)](#).

- Virtual Private Database

Amazon RDS Oracle doesn't support the following Oracle Database features:

- Automatic Storage Management (ASM)
- Database Vault
- Flashback Database
- Multitenant
- Oracle Enterprise Manager Cloud Control Management Repository
- Real Application Clusters (Oracle RAC)

- Real Application Testing
- Unified Auditing, Pure Mode

Oracle Database Parameter Support

In Amazon RDS, you manage parameters using parameter groups. For more information, see [Working with DB Parameter Groups \(p. 211\)](#). To view the supported parameters for a specific Oracle edition and version, you can run the AWS CLI `describe-engine-default-parameters` command.

For example, to view the supported parameters for Oracle Enterprise Edition version 12.2, run the following command:

```
aws rds describe-engine-default-parameters --db-parameter-group-family oracle-ee-12.2
```

Oracle Engine Version Management

DB Engine Version Management is a feature of Amazon RDS that enables you to control when and how the database engine software running your DB instances is patched and upgraded. This feature gives you the flexibility to maintain compatibility with database engine patch versions, test new patch versions to ensure they work effectively with your application before deploying in production, and perform version upgrades on your own terms and timelines.

Note

Amazon RDS periodically aggregates official Oracle database patches using an Amazon RDS-specific DB Engine version. To see a list of which Oracle patches are contained in an Amazon RDS Oracle-specific engine version, go to [Oracle Database Engine Release Notes \(p. 1173\)](#).

Currently, you perform all Oracle database upgrades manually. For more information about upgrading an Oracle DB instance, see [Upgrading the Oracle DB Engine \(p. 979\)](#).

Deprecation of Oracle 11.2.0.2

In 2017, Amazon RDS is deprecating support for Oracle version 11.2.0.2. Oracle is no longer providing patches for this version. Therefore, to provide the best experience for AWS customers, we are deprecating this version.

There are no longer any production DB instances running Oracle version 11.2.0.2. You might still have a snapshot of an 11.2.0.2 DB instance.

Amazon RDS is deprecating support for Oracle version 11.2.0.2 according to the following schedule.

Date	Information
August 4, 2016	You can no longer create DB instances that use Oracle version 11.2.0.2.
December 18, 2019	Any 11.2.0.2 snapshots are upgraded to 11.2.0.4. You can upgrade your snapshots yourself prior to this date. For more information, see Upgrading an Oracle DB Snapshot (p. 984) .

Deprecation of Oracle 11.2.0.3

In 2017, Amazon RDS is deprecating support for Oracle version 11.2.0.3. Oracle is no longer providing patches for this version. Therefore, to provide the best experience for AWS customers, we are deprecating this version.

There are no longer any production DB instances running Oracle version 11.2.0.3. You might still have a snapshot of an 11.2.0.3 DB instance.

Amazon RDS is deprecating support for Oracle version 11.2.0.3 according to the following schedule.

Date	Information
August 4, 2016	You can no longer create DB instances that use Oracle version 11.2.0.3.
December 18, 2019	Any 11.2.0.3 snapshots are upgraded to 11.2.0.4. You can upgrade your snapshots yourself prior to this date. For more information, see Upgrading an Oracle DB Snapshot (p. 984) .

Deprecation of Oracle 12.1.0.1

In 2017, Amazon RDS is deprecating support for Oracle version 12.1.0.1. Oracle is no longer providing patches for this version. Therefore, to provide the best experience for AWS customers, we are deprecating this version.

There are no longer any production DB instances running Oracle version 12.1.0.1. You might still have a snapshot of a 12.1.0.1 DB instance.

Amazon RDS will deprecate support for Oracle version 12.1.0.1 according to the following schedule.

Date	Information
February 15, 2017	You can no longer create DB instances that use Oracle version 12.1.0.1.
December 18, 2019	Any 12.1.0.1 snapshots are upgraded to 12.1.0.2. You can upgrade your snapshots yourself prior to this date. For more information, see Upgrading an Oracle DB Snapshot (p. 984) .

Using Huge Pages with an Oracle DB Instance

Amazon RDS for Oracle supports Linux kernel huge pages for increased database scalability. The use of huge pages results in smaller page tables and less CPU time spent on memory management, increasing the performance of large database instances. For more information, see [Overview of HugePages](#) in the Oracle documentation.

You can use huge pages with the following versions and editions of Oracle:

- 19.0.0.0, all editions

- 18.0.0.0, all editions
- 12.2.0.1, all editions
- 12.1.0.2, all editions
- 11.2.0.4, all editions

The `use_large_pages` parameter controls whether huge pages are enabled for a DB instance. The possible settings for this parameter are `ONLY`, `FALSE`, and `{DBInstanceClassHugePagesDefault}`. The `use_large_pages` parameter is set to `{DBInstanceClassHugePagesDefault}` in the default DB parameter group for Oracle.

To control whether huge pages are enabled for a DB instance automatically, you can use the `DBInstanceClassHugePagesDefault` formula variable in parameter groups. The value is determined as follows:

- For the DB instance classes mentioned in the table below, `DBInstanceClassHugePagesDefault` always evaluates to `FALSE` by default, and `use_large_pages` evaluates to `FALSE`. You can enable huge pages manually for these DB instance classes if the DB instance class has at least 14 GiB of memory.
- For DB instance classes not mentioned in the table below, if the DB instance class has less than 14 GiB of memory, `DBInstanceClassHugePagesDefault` always evaluates to `FALSE`, and `use_large_pages` evaluates to `FALSE`.
- For DB instance classes not mentioned in the table below, if the instance class has at least 14 GiB of memory and less than 100 GiB of memory, `DBInstanceClassHugePagesDefault` evaluates to `TRUE` by default, and `use_large_pages` evaluates to `ONLY`. You can disable huge pages manually by setting `use_large_pages` to `FALSE`.
- For DB instance classes not mentioned in the table below, if the instance class has at least 100 GiB of memory, `DBInstanceClassHugePagesDefault` always evaluates to `TRUE`, `use_large_pages` evaluates to `ONLY`, and huge pages can't be disabled.

Huge pages are not enabled by default for the following DB instance classes.

DB Instance Class Family	DB Instance Classes with Huge Pages Not Enabled by Default
db.m5	db.m5.large
db.m4	db.m4.large, db.m4.xlarge, db.m4.2xlarge, db.m4.4xlarge, db.m4.10xlarge
db.m3	db.m3.medium, db.m3.large, db.m3.xlarge, db.m3.2xlarge
db.r3	db.r3.large, db.r3.xlarge, db.r3.2xlarge, db.r3.4xlarge, db.r3.8xlarge
db.t3	db.t3.micro, db.t3.small, db.t3.medium, db.t3.large
db.t2	db.t2.micro, db.t2.small, db.t2.medium, db.t2.large

For more information about DB instance classes, see [Hardware Specifications for All Available DB Instance Classes \(p. 109\)](#).

To enable huge pages for new or existing DB instances manually, set the `use_large_pages` parameter to `ONLY`. You can't use huge pages with Oracle Automatic Memory Management (AMM). If you set the parameter `use_large_pages` to `ONLY`, then you must also set both `memory_target` and `memory_max_target` to 0. For more information about setting DB parameters for your DB instance, see [Working with DB Parameter Groups \(p. 211\)](#).

You can also set the `sga_target`, `sga_max_size`, and `pga_aggregate_target` parameters. When you set system global area (SGA) and program global area (PGA) memory parameters, add the values together. Subtract this total from your available instance memory (`DBInstanceClassMemory`) to determine the free memory beyond the huge pages allocation. You must leave free memory of at least 2 GiB, or 10 percent of the total available instance memory, whichever is smaller.

After you configure your parameters, you must reboot your DB instance for the changes to take effect. For more information, see [Rebooting a DB Instance \(p. 166\)](#).

The following is a sample parameter configuration for huge pages that enables huge pages manually. You should set the values to meet your needs.

```
memory_target      = 0
memory_max_target = 0
pga_aggregate_target = {DBInstanceClassMemory*1/8}
sga_target        = {DBInstanceClassMemory*3/4}
sga_max_size     = {DBInstanceClassMemory*3/4}
use_large_pages   = ONLY
```

Assume the following parameters values are set in a parameter group.

```
memory_target      = IF({DBInstanceClassHugePagesDefault}, 0,
{DBInstanceClassMemory*3/4})
memory_max_target = IF({DBInstanceClassHugePagesDefault}, 0,
{DBInstanceClassMemory*3/4})
pga_aggregate_target = IF({DBInstanceClassHugePagesDefault},
{DBInstanceClassMemory*1/8}, 0)
sga_target        = IF({DBInstanceClassHugePagesDefault},
{DBInstanceClassMemory*3/4}, 0)
sga_max_size     = IF({DBInstanceClassHugePagesDefault},
{DBInstanceClassMemory*3/4}, 0)
use_large_pages   = {DBInstanceClassHugePagesDefault}
```

The parameter group is used by a db.r4 DB instance class with less than 100 GiB of memory and a db.r3 instance with more than 100 GiB memory. With these parameter settings and `use_large_pages` set to `{DBInstanceClassHugePagesDefault}`, huge pages are enabled on the db.r4 instance, but disabled on the db.r3 instance.

Consider another example with following parameters values set in a parameter group.

```
memory_target      = IF({DBInstanceClassHugePagesDefault}, 0,
{DBInstanceClassMemory*3/4})
memory_max_target = IF({DBInstanceClassHugePagesDefault}, 0,
{DBInstanceClassMemory*3/4})
pga_aggregate_target = IF({DBInstanceClassHugePagesDefault},
{DBInstanceClassMemory*1/8}, 0)
sga_target        = IF({DBInstanceClassHugePagesDefault},
{DBInstanceClassMemory*3/4}, 0)
sga_max_size     = IF({DBInstanceClassHugePagesDefault},
{DBInstanceClassMemory*3/4}, 0)
use_large_pages   = FALSE
```

The parameter group is used by a db.r4 DB instance class and a db.r5 DB instance class, both with less than 100 GiB of memory. The parameter group is also used by a db.r3 instance with more than 100

GiB memory. With these parameter settings, huge pages are disabled on the db.r4 instance, the db.r5 instance, and the db.r3 instance.

Note

If this parameter group is used by a db.r4 DB instance class or db.r5 DB instance class with at least 100 GiB of memory, the `FALSE` setting for `use_large_pages` is overridden and set to `ONLY`. In this case, a customer notification regarding the override is sent.

After huge pages are active on your DB instance, you can view huge pages information by enabling enhanced monitoring. For more information, see [Enhanced Monitoring \(p. 307\)](#).

Using utl_http, utl_tcp, and utl_smtp with an Oracle DB Instance

Amazon RDS supports outbound network access on your DB instances running Oracle. You can use `utl_http`, `utl_tcp`, and `utl_smtp` to connect from your DB instance to the network.

Note the following about working with outbound network access:

- To use `utl_http` on DB instances running Oracle 11g, you must install the XMLDB option. For more information, see [Oracle XML DB \(p. 1078\)](#).
- Outbound network access with `utl_http`, `utl_tcp`, and `utl_smtp` is supported only for Oracle DB instances in a VPC. To determine whether or not your DB instance is in a VPC, see [Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform \(p. 506\)](#). To move a DB instance not in a VPC into a VPC, see [Moving a DB Instance Not in a VPC into a VPC \(p. 520\)](#).
- To use SMTP with the UTL_MAIL option, see [Oracle UTL_MAIL \(p. 1076\)](#).
- The Domain Name Server (DNS) name of the remote host can be any of the following:
 - Publicly resolvable.
 - The endpoint of an Amazon RDS DB instance.
 - Resolvable through a custom DNS server. For more information, see [Setting Up a Custom DNS Server \(p. 1094\)](#).
 - The private DNS name of an Amazon EC2 instance in the same VPC or a peered VPC. In this case, make sure that the name is resolvable through a custom DNS server. Alternatively, to use the DNS provided by Amazon, you can enable the `enableDnsSupport` attribute in the VPC settings and enable DNS resolution support for the VPC peering connection. For more information, see [DNS Support in Your VPC](#) and [Modifying Your VPC Peering Connection](#).

To connect securely to remote SSL/TLS resources, you can create and upload customized Oracle wallets. By using the Amazon S3 integration with Amazon RDS for Oracle feature, you can download a wallet from Amazon S3 into Oracle DB instances. For information about Amazon S3 integration for Oracle, see [Amazon S3 Integration \(p. 1005\)](#).

The following example creates a wallet for accessing <https://status.aws.amazon.com/robots.txt> over `utl_http`.

1. Obtain the certificate for `Amazon Root CA 1` from the [Amazon Trust Services Repository](#).
2. Create a new wallet and add the following certificate:

```
orapki wallet create -wallet . -auto_login_only
orapki wallet add -wallet . -trusted_cert -cert AmazonRootCA1.pem.pem -
auto_login_only orapki wallet display -wallet .
```

3. Upload the wallet to your Amazon S3 bucket.
4. Complete the prerequisites for Amazon S3 integration with Oracle, and add the S3_INTEGRATION option to your Oracle DB instance. Ensure that the IAM role for the option has access to the Amazon S3 bucket you are using.

For more information, see [Amazon S3 Integration \(p. 1005\)](#).

5. Connect to the DB instance, and create a directory in the database to hold the wallet. The following example creates a directory called SSL_WALLET:

```
exec rdsadmin.rdsadmin_util.create_directory('SSL_WALLET');
```

6. Download the wallet from your Amazon S3 bucket to the Oracle DB instance.

The following example downloads a wallet to the DB instance directory SSL_WALLET:

```
SELECT rdsadmin.rdsadmin_s3_tasks.download_from_s3(  
    p_bucket_name    => 'bucket_name',  
    p_s3_prefix      => 'wallet_name',  
    p_directory_name => 'SSL_WALLET')  
AS TASK_ID FROM DUAL;
```

Replace `bucket_name` with the name of the bucket you are using, and replace `wallet_name` with the name of the wallet.

7. Set this wallet for utl_http transactions by running the following procedure:

```
DECLARE  
    l_wallet_path all_directories.directory_path%type;  
BEGIN  
    select directory_path into l_wallet_path from all_directories  
    where upper(directory_name)='SSL_WALLET';  
  
    utl_http.set_wallet('file://' || l_wallet_path);  
END;
```

8. Access the URL from above over SSL/TLS.

```
SELECT utl_http.request('https://status.aws.amazon.com/robots.txt') AS ROBOTS_TXT FROM  
DUAL;  
  
ROBOTS_TXT  
-----  
User-agent: *  
Allow: /
```

Note

The specific certificates that are required for your wallet vary by service. For AWS services, the certificates can typically be found in the [Amazon Trust Services Repository](#).

You can use a similar setup to send emails through UTL_SMTP over SSL/TLS (including [Amazon Simple Email Service](#)).

You can establish database links between Oracle DB instances over an SSL/TLS endpoint if the Oracle SSL option is configured for each instance. No further configuration is required. For more information, see [Oracle Secure Sockets Layer \(p. 1051\)](#).

Using OEM, APEX, TDE, and Other Options

Most Amazon RDS DB engines support option groups that allow you to select additional features for your DB instance. Oracle DB instances support several options, including Oracle Enterprise Manager (OEM), Transparent Data Encryption (TDE), Application Express (APEX), and Native Network Encryption. For a complete list of supported Oracle options, see [Options for Oracle DB Instances \(p. 1004\)](#). For more information about working with option groups, see [Working with Option Groups \(p. 196\)](#).

Using Extended Data Types

Amazon RDS Oracle version 12c supports extended data types. With extended data types, the maximum size is 32,767 bytes for the VARCHAR2, NVARCHAR2, and RAW data types. To use extended data types, set the `MAX_STRING_SIZE` parameter to `EXTENDED`. For more information, see [Extended Data Types](#) in the Oracle documentation.

If you don't want to use extended data types, keep the `MAX_STRING_SIZE` parameter set to `STANDARD` (the default). When this parameter is set to `STANDARD`, the size limits are 4,000 bytes for the VARCHAR2 and NVARCHAR2 data types, and 2,000 bytes for the RAW data type.

You can enable extended data types on a new or existing DB instance. For new DB instances, DB instance creation time is typically longer when you enable extended data types. For existing DB instances, the DB instance is unavailable during the conversion process.

The following are considerations for a DB instance with extended data types enabled:

- When you enable extended data types for a DB instance, you can't change the DB instance back to use the standard size for data types. After a DB instance is converted to use extended data types, if you set the `MAX_STRING_SIZE` parameter back to `STANDARD` it results in the `incompatible-parameters` status.
- When you restore a DB instance that uses extended data types, you must specify a parameter group with the `MAX_STRING_SIZE` parameter set to `EXTENDED`. During restore, if you specify the default parameter group or any other parameter group with `MAX_STRING_SIZE` set to `STANDARD` it results in the `incompatible-parameters` status.
- We recommend that you don't enable extended data types for Oracle DB instances running on the t2.micro DB instance class.

When the DB instance status is `incompatible-parameters` because of the `MAX_STRING_SIZE` setting, the DB instance remains unavailable until you set the `MAX_STRING_SIZE` parameter to `EXTENDED` and reboot the DB instance.

Enabling Extended Data Types for a New DB Instance

To enable extended data types for a new DB instance

1. Set the `MAX_STRING_SIZE` parameter to `EXTENDED` in a parameter group.

To set the parameter, you can either create a new parameter group or modify an existing parameter group.

For more information, see [Working with DB Parameter Groups \(p. 211\)](#).

2. Create a new Amazon RDS Oracle DB instance, and associate the parameter group with `MAX_STRING_SIZE` set to `EXTENDED` with the DB instance.

For more information, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).

Enabling Extended Data Types for an Existing DB Instance

When you modify a DB instance to enable extended data types, the data in the database is converted to use the extended sizes. The DB instance is unavailable during the conversion. The amount of time it takes to convert the data depends on the DB instance class used by the DB instance and the size of the database.

Note

After you enable extended data types, you can't perform a point-in-time restore to a time during the conversion. You can restore to the time immediately before the conversion or after the conversion.

To enable extended data types for an existing DB instance

1. Take a snapshot of the database.

If there are invalid objects in the database, Amazon RDS tries to recompile them. The conversion to extended data types can fail if Amazon RDS can't recompile an invalid object. The snapshot enables you to restore the database if there is a problem with the conversion. Always check for invalid objects before conversion and fix or drop those invalid objects. For production databases, we recommend testing the conversion process on a copy of your DB instance first.

For more information, see [Creating a DB Snapshot \(p. 264\)](#).

2. Set the `MAX_STRING_SIZE` parameter to `EXTENDED` in a parameter group.

To set the parameter, you can either create a new parameter group or modify an existing parameter group.

For more information, see [Working with DB Parameter Groups \(p. 211\)](#).

3. Modify the DB instance to associate it with the parameter group with `MAX_STRING_SIZE` set to `EXTENDED`.

For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

4. Reboot the DB instance for the parameter change to take effect.

For more information, see [Rebooting a DB Instance \(p. 166\)](#).

Public Synonyms

You can create public synonyms referencing objects in your own schemas. Don't create or modify public synonyms for Oracle-maintained schemas, including `SYS`, `SYSTEM`, and `RDSADMIN`. Doing so might result in invalidation of core database components and affect the availability of your DB instance.

Creating a DB Instance Running the Oracle Database Engine

The basic building block of Amazon RDS is the DB instance. This is the environment in which you run your Oracle databases.

Important

Before you can create or connect to a DB instance, you must complete the tasks in [Setting Up for Amazon RDS \(p. 5\)](#).

For an example that walks you through the process of creating and connecting to a sample DB instance, see [Creating an Oracle DB Instance and Connecting to a Database on an Oracle DB Instance \(p. 45\)](#).

Note

If you are using the console, a new console interface is available for database creation. Choose either the **New Console** or the **Original Console** instructions based on the console that you are using. The **New Console** instructions are open by default.

New Console

You can create a DB instance running Oracle with the AWS Management Console with **Easy Create** enabled or disabled. With **Easy Create** enabled, you specify only the DB engine type, DB instance size, and DB instance identifier. **Easy Create** uses the default setting for other configuration options. With **Easy Create** not enabled, you specify more configuration options when you create a database, including ones for availability, security, backups, and maintenance.

Note

For this example, **Standard Create** is enabled, and **Easy Create** isn't enabled. For information about creating an Oracle DB instance with **Easy Create** enabled, see [Creating an Oracle DB Instance and Connecting to a Database on an Oracle DB Instance \(p. 45\)](#).

To create an Oracle DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the AWS Management Console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
4. Choose **Create database**.
5. In **Choose a database creation method**, choose **Standard Create**.

The Oracle editions that are available vary by AWS Region.

Create database

Choose a database creation method [Info](#)

Standard Create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy Create

Use recommended best-practice configuration options can be changed after the database is created.

Engine options

Engine type [Info](#)

Amazon Aurora



MySQL



MariaDB



PostgreSQL



Oracle

ORACLE®

Microsoft SQL Server



Edition

Oracle Enterprise Edition

Efficient, reliable, and secure database management system that delivers comprehensive high-end capabilities for mission-critical applications and demanding database workloads.

Version [Info](#)

6. In **Engine options**, choose **Oracle**.

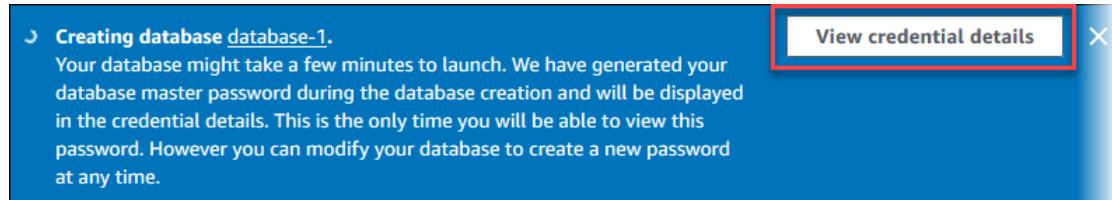
7. For **Edition**, choose the Oracle DB engine edition that you want to use. The Oracle editions that are available vary by AWS Region.
8. For **Version**, choose the Oracle version.
9. The next step asks if you are planning to use the DB instance that you are creating for production. If you are, choose **Production**. When you choose **Production**, the following are also chosen in a later step:
 - **Multi-AZ deployment** failover option
 - **Provisioned IOPS** storage option
 - **Enable deletion protection** option
10. To enter your master password, do the following:
 - a. In the **Settings** section, open **Credential Settings**.
 - b. Clear the **Auto generate a password** check box.
 - c. (Optional) Change the **Master username** value and enter the same password in **Master password** and **Confirm password**.

By default, the new DB instance uses an automatically generated password for the master user.

11. For the remaining sections, specify your DB instance settings. For information about each setting, see [Settings for Oracle DB Instances \(p. 949\)](#).
12. Choose **Create database**.

If you chose to use an automatically generated password, the **View credential details** button appears on the **Databases** page.

To view the master user name and password for the DB instance, choose **View credential details**.



To connect to the DB instance as the master user, use the user name and password that appear.

Important

You can't view the master user password again. If you don't record it, you might have to change it. If you need to change the master user password after the DB instance is available, you can modify the DB instance to do so. For more information about modifying a DB instance, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

13. For **Databases**, choose the name of the new Oracle DB instance.

On the RDS console, the details for the new DB instance appear. The DB instance has a status of **creating** until the DB instance is created and ready for use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and storage allocated, it can take several minutes for the new instance to be available.

The screenshot shows the 'Summary' tab for a database named 'database-1'. The 'Info' section displays the status as 'Creating'. Other details include the DB identifier 'database-1', Role 'Instance', CPU usage, Current activity (0 Sessions), Engine 'Oracle Enterprise Edition', and Class 'db.m5'. Below the summary, there are tabs for 'Connectivity & security' (highlighted in orange), 'Monitoring', 'Logs & events', 'Configuration', and 'Maintenance & b'.

Original Console

To launch an Oracle DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the top right corner of the AWS Management Console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
If the navigation pane is closed, choose the menu icon at the top left to open it.
4. Choose **Create database** to open the **Select engine** page.

The Oracle editions that are available vary by AWS Region.

Select engine

Engine options

Amazon Aurora

Amazon
Aurora

MySQL



MariaDB



PostgreSQL



Oracle

ORACLE®

Microsoft SQL Server



Oracle

Edition

Oracle Enterprise Edition

Efficient, reliable, and secure database management system that delivers comprehensive high-end capabilities for mission-critical applications and demanding database workloads.

Oracle Standard Edition

Affordable and full-featured database management system supporting up to 32 vCPUs.

Oracle Standard Edition One

Affordable and full-featured database management system supporting up to 16 vCPUs.

Oracle Standard Edition Two

Affordable and full-featured database management system supporting up to 16 vCPUs. Oracle Database Standard Edition Two is a replacement for Standard Edition and Standard Edition One.

Only enable options eligible for RDS Free Usage Tier [info](#)

Cancel

Next

5. In the **Select engine** window, choose **Select** for the Oracle DB engine you want to use and then choose **Next**.
6. The next step asks if you are planning to use the DB instance you are creating for production. If you are, choose **Production**. When you choose **Production**, the following are also chosen in a later step:
 - **Multi-AZ deployment** failover option
 - **Provisioned IOPS** storage option
 - **Enable deletion protection** option
7. Choose **Next** to continue. The **Specify DB details** page appears.

On the **Specify DB details** page, specify your DB instance information. For information about each setting, see [Settings for Oracle DB Instances \(p. 949\)](#).

Specify DB details

Instance specifications

Estimate your monthly costs for the DB Instance using the [AWS Simple Monthly Calculator](#)

DB engine
Oracle Database Enterprise Edition

License model [Info](#)

DB engine version [Info](#)

DB instance class [Info](#)

Multi-AZ deployment [Info](#)
 Create replica in different zone
Creates a replica in a different Availability Zone (AZ) to provide data redundancy, eliminate I/O freezes, and minimize latency spikes during system backups.
 No

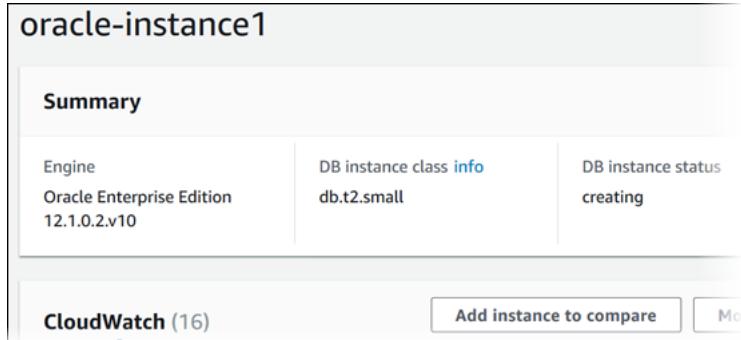
Storage type [Info](#)

8. Choose **Next** to continue. The **Configure advanced settings** page appears.

On the **Configure advanced settings** page, provide additional information that RDS needs to launch the DB instance. For information about each setting, see [Settings for Oracle DB Instances \(p. 949\)](#).

9. Choose **Create database**.
10. On the final page, choose **View DB instance details**.

On the RDS console, the details for the new DB instance appear. The DB instance has a status of **creating** until the DB instance is created and ready for use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and storage allocated, it could take several minutes for the new instance to be available.



AWS CLI

To create an Oracle DB instance by using the AWS CLI, call the [create-db-instance](#) command with the parameters below. For information about each setting, see [Settings for Oracle DB Instances \(p. 949\)](#).

- `--db-instance-identifier`
- `--db-instance-class`
- `--db-security-groups`
- `--db-subnet-group`
- `--engine`
- `--master-user-name`
- `--master-user-password`
- `--allocated-storage`
- `--backup-retention-period`

Example

The following command will launch the example DB instance.

For Linux, OS X, or Unix:

```
aws rds create-db-instance \
  --engine oracle-se1 \
  --db-instance-identifier mydbinstance \
  --allocated-storage 20 \
  --db-instance-class db.m1.small \
  --db-security-groups mydbsecuritygroup \
  --db-subnet-group mydbsubnetgroup \
  --master-username masterawsuser \
  --master-user-password masteruserpassword \
  --backup-retention-period 3
```

For Windows:

```
aws rds create-db-instance ^
  --engine oracle-se1 ^
  --db-instance-identifier mydbinstance ^
  --allocated-storage 20 ^
  --db-instance-class db.m1.small ^
  --db-security-groups mydbsecuritygroup ^
  --db-subnet-group mydbsubnetgroup ^
```

```
--master-username masterawsuser ^
--master-user-password masteruserpassword ^
--backup-retention-period 3
```

This command should produce output similar to the following:

```
DBINSTANCE mydbinstance db.m1.small oracle-se1 20 sa creating 3 **** n
11.2.0.4.v1
SECGROUP default active
PARAMGRP default.oracle-se1-11.2 in-sync
```

RDS API

To create an Oracle DB instance by using the Amazon RDS API, call the [CreateDBInstance](#) operation with the parameters below. For information about each setting, see [Settings for Oracle DB Instances \(p. 949\)](#).

- AllocatedStorage
- BackupRetentionPeriod
- DBInstanceClass
- DBInstanceIdentifier
- DBSecurityGroups
- DBSubnetGroup
- Engine
- MasterUsername
- MasterUserPassword

Example

```
https://rds.amazonaws.com/
?Action=CreateDBInstance
&AllocatedStorage=250
&BackupRetentionPeriod=3
&DBInstanceClass=db.m1.large
&DBInstanceIdentifier=mydbinstance
&DBSecurityGroups.member.1=mysecuritygroup
&DBSubnetGroup=mydbsubnetgroup
&Engine=oracle-se1
&MasterUserPassword=masteruserpassword
&MasterUsername=masterawsuser
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20140305/us-west-1/rds/aws4_request
&X-Amz-Date=20140305T185838Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=b441901545441d3c7a48f63b5b1522c5b2b37c137500c93c45e209d4b3a064a3
```

Settings for Oracle DB Instances

The following table contains details about settings that you choose when you create an Oracle DB instance.

Setting	Setting Description
Allocated storage	<p>The amount of storage to allocate your DB instance (in gigabytes). In some cases, allocating a higher amount of storage for your DB instance than the size of your database can improve I/O performance.</p> <p>For more information, see Amazon RDS DB Instance Storage (p. 135).</p>
Auto minor version upgrade	<p>Enable auto minor version upgrade to enable your DB instance to receive preferred minor DB engine version upgrades automatically when they become available. Amazon RDS performs automatic minor version upgrades in the maintenance window.</p> <p>Some options, such as Oracle Locator, Oracle Multimedia, and Oracle Spatial, require that you enable automatic minor version upgrades. Upgrades for DB instances that use these options are installed during your scheduled maintenance window, and an outage occurs during the upgrade. You can't disable automatic minor version upgrades at the same time as you modify the option group to remove such an option.</p>
Availability zone	<p>The Availability Zone for your DB instance. Use the default of No Preference unless you need to specify a particular Availability Zone.</p> <p>For more information, see Regions and Availability Zones (p. 132).</p>
Backup retention period	<p>The number of days that you want automatic backups of your DB instance to be kept. For any nontrivial instance, set this value to 1 or greater.</p> <p>For more information, see Working With Backups (p. 254).</p>
Backup window	<p>The time period during which Amazon RDS automatically takes a backup of your DB instance. Unless you have a specific time that you want to have your database backed up, use the default of No Preference.</p> <p>For more information, see Working With Backups (p. 254).</p>
Character set	<p>The character set for your DB instance. The default value of AL32UTF8 is for the Unicode 5.0 UTF-8 Universal character set. You can't change the character set after the DB instance is created.</p> <p>For more information, see Oracle Character Sets Supported in Amazon RDS (p. 1001).</p>
Copy tags to snapshots	<p>This option copies any DB instance tags to a DB snapshot when you create a snapshot.</p> <p>For more information, see Tagging Amazon RDS Resources (p. 177).</p>
Database authentication	<p>The database authentication option you want to use.</p>

Setting	Setting Description
	<p>Choose Password authentication to authenticate database users with database passwords only.</p> <p>Choose Password and Kerberos authentication to authenticate database users with database passwords and Kerberos authentication through an AWS Managed Microsoft AD created with AWS Directory Service. Next, choose the directory or choose Create a new Directory. For more information, see Using Kerberos Authentication with Amazon RDS for Oracle (p. 1134).</p>
Database name	The name for the database on your DB instance. The name must begin with a letter and contain up to eight alphanumeric characters. You can't specify the string NULL, or any other reserved word, for the database name. If you don't provide a name, Amazon RDS creates a database named ORCL.
Database port	The port that you want to access the DB instance through. Oracle installations default to port 1521.
DB engine version	The version of Oracle that you want to use.
DB instance class	<p>The DB instance class that you want to use.</p> <p>For more information, see Choosing the DB Instance Class (p. 107) and DB Instance Class Support for Oracle (p. 917).</p>
DB instance identifier	The name for your DB instance. The name must be unique for your account and AWS Region. You can add some intelligence to the name, such as including the AWS Region and DB engine you chose, for example <code>oracle-instance1</code> .
DB parameter group	<p>A parameter group for your DB instance. You can choose the default parameter group or you can create a custom parameter group.</p> <p>For more information, see Working with DB Parameter Groups (p. 211) and Modifying Oracle sqlnet.ora Parameters (p. 976).</p>
Deletion protection	<p>Enable deletion protection to prevent your DB instance from being deleted. If you create a production DB instance with the AWS Management Console, deletion protection is enabled by default.</p> <p>For more information, see Deleting a DB Instance (p. 174).</p>
Encryption	<p>Enable Encryption to enable encryption at rest for this DB instance.</p> <p>For more information, see Encrypting Amazon RDS Resources (p. 429).</p>

Setting	Setting Description
Enhanced Monitoring	<p>Enable enhanced monitoring to enable gathering metrics in real time for the operating system that your DB instance runs on.</p> <p>For more information, see Enhanced Monitoring (p. 307).</p>
License model	<p>The license model that you want to use. Choose license-included to use the general license agreement for Oracle. Choose bring-your-own-license to use your existing Oracle license.</p> <p>For more information, see Oracle Licensing (p. 916).</p>
Maintenance window	<p>The 30-minute window in which pending modifications to your DB instance are applied. If the time period doesn't matter, choose No Preference.</p> <p>For more information, see The Amazon RDS Maintenance Window (p. 156).</p>
Master password	<p>The password for your master user account. The password must contain from 8 to 30 printable ASCII characters (excluding /, ", and @).</p>
Master username	<p>The name that you use as the master user name to log on to your DB instance with all database privileges. This user account is used to log in to the DB instance and is granted DBA privileges.</p> <p>For more information, see Oracle Security (p. 920).</p>
Multi-AZ deployment	<p>Create a standby instance to create a passive secondary replica of your DB instance in another Availability Zone for failover support. We recommend Multi-AZ for production workloads to maintain high availability.</p> <p>For development and testing, you can choose Do not create a standby instance.</p> <p>For more information, see Regions and Availability Zones (p. 132).</p>
Option group	<p>An option group for your DB instance. You can choose the default option group or you can create a custom option group.</p> <p>For more information, see Working with Option Groups (p. 196).</p>

Setting	Setting Description
Performance Insights	<p>Enable Performance Insights to monitor your DB instance load so that you can analyze and troubleshoot your database performance.</p> <p>Choose a retention period to determine how much rolling data history to retain. The default of seven days is in the free tier. Long-term retention (two years) is priced per vCPU per month.</p> <p>Choose a master key to protect the key used to encrypt this database volume. Choose from the master keys in your account, or enter the key from a different account.</p> <p>For more information, see Using Amazon RDS Performance Insights (p. 319).</p>
Publicly accessible	<p>Yes to give the DB instance a public IP address, meaning that it's accessible outside the VPC. To be publicly accessible, the DB instance also has to be in a public subnet in the VPC.</p> <p>No to make the DB instance accessible only from inside the VPC.</p> <p>For more information, see Hiding a DB Instance in a VPC from the Internet (p. 515).</p>
Storage autoscaling	<p>Enable storage autoscaling to enable Amazon RDS to automatically increase storage when needed to avoid having your DB instance run out of storage space.</p> <p>Use Maximum storage threshold to set the upper limit for Amazon RDS to automatically increase storage for your DB instance. The default is 1,000 GiB.</p> <p>For more information, see Managing Capacity Automatically with Amazon RDS Storage Autoscaling (p. 232).</p>
Storage type	<p>The storage type for your DB instance.</p> <p>For more information, see Amazon RDS Storage Types (p. 135).</p>
Subnet group	<p>This setting depends on the platform that you are on. If you are a new customer to AWS, keep the subnet group set to default, which is the default DB subnet group that was created for your account.</p> <p>If you are creating a DB instance on the earlier E2-Classic platform, you might want your DB instance in a specific VPC. In this case, choose the DB subnet group that you created for that VPC.</p>

Setting	Setting Description
Virtual Private Cloud (VPC)	This setting depends on the platform that you are on. If you are a new customer to AWS, choose the default VPC shown. If you are creating a DB instance on the earlier EC2-Classic platform that doesn't use a VPC, choose Not in VPC . For more information, see Amazon Virtual Private Cloud VPCs and Amazon RDS (p. 505) .
VPC security groups	If you are a new customer to AWS, Create new to create a new VPC security group. Otherwise, Choose existing , then choose from security groups that you previously created. When you choose Create new in the RDS console, a new security group is created. This new security group has an inbound rule that allows access to the DB instance from the IP address detected in your browser. For more information, see Working with DB Security Groups (EC2-Classic Platform) (p. 491) .

Connecting to a DB Instance Running the Oracle Database Engine

After Amazon RDS provisions your Oracle DB instance, you can use any standard SQL client application to connect to the DB instance. In this topic, you connect to a DB instance that is running the Oracle database engine by using Oracle SQL Developer or SQL*Plus.

For an example that walks you through the process of creating and connecting to a sample DB instance, see [Creating an Oracle DB Instance and Connecting to a Database on an Oracle DB Instance \(p. 45\)](#).

Finding the Endpoint of Your DB Instance

Each Amazon RDS DB instance has an endpoint, and each endpoint has the DNS name and port number for the DB instance. To connect to your DB instance using a SQL client application, you need the DNS name and port number for your DB instance.

You can find the endpoint for a DB instance using the Amazon RDS console or the AWS CLI.

Note

If you are using Kerberos authentication, see [Connecting to Oracle with Kerberos Authentication \(p. 1145\)](#).

Console

To find the endpoint using the console

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the console, choose the AWS Region of your DB instance.
3. Find the DNS name and port number for your DB Instance.
 - a. Choose **Databases** to display a list of your DB instances.
 - b. Choose the Oracle DB instance name to display the instance details.
 - c. On the **Connectivity & security** tab, copy the endpoint. Also, note the port number. You need both the endpoint and the port number to connect to the DB instance.

The screenshot shows the AWS RDS console for an instance named 'database-1'. The 'Summary' section displays the DB identifier as 'database-1' and the instance role as 'Instance'. Below the summary, there are tabs for 'Connectivity & security', 'Monitoring', 'Logs & events', and 'Config'. The 'Connectivity & security' tab is selected. Under this tab, the 'Endpoint & port' section is visible. It shows the endpoint as 'database-1.us-west-2.rds.amazonaws.com' and the port as '1521'. The 'Endpoint' part of this section is highlighted with a red oval.

AWS CLI

To find the endpoint of an Oracle DB instance by using the AWS CLI, call the [describe-db-instances](#) command.

Example To find the endpoint using the AWS CLI

```
aws rds describe-db-instances
```

Search for `Endpoint` in the output to find the DNS name and port number for your DB instance. The `Address` line in the output contains the DNS name. The following is an example of the JSON endpoint output.

```
"Endpoint": {  
    "HostedZoneId": "Z1PVIF0B656C1W",  
    "Port": 3306,  
    "Address": "myinstance.123456789012.us-west-2.rds.amazonaws.com"  
},
```

Note

The output might contain information for multiple DB instances.

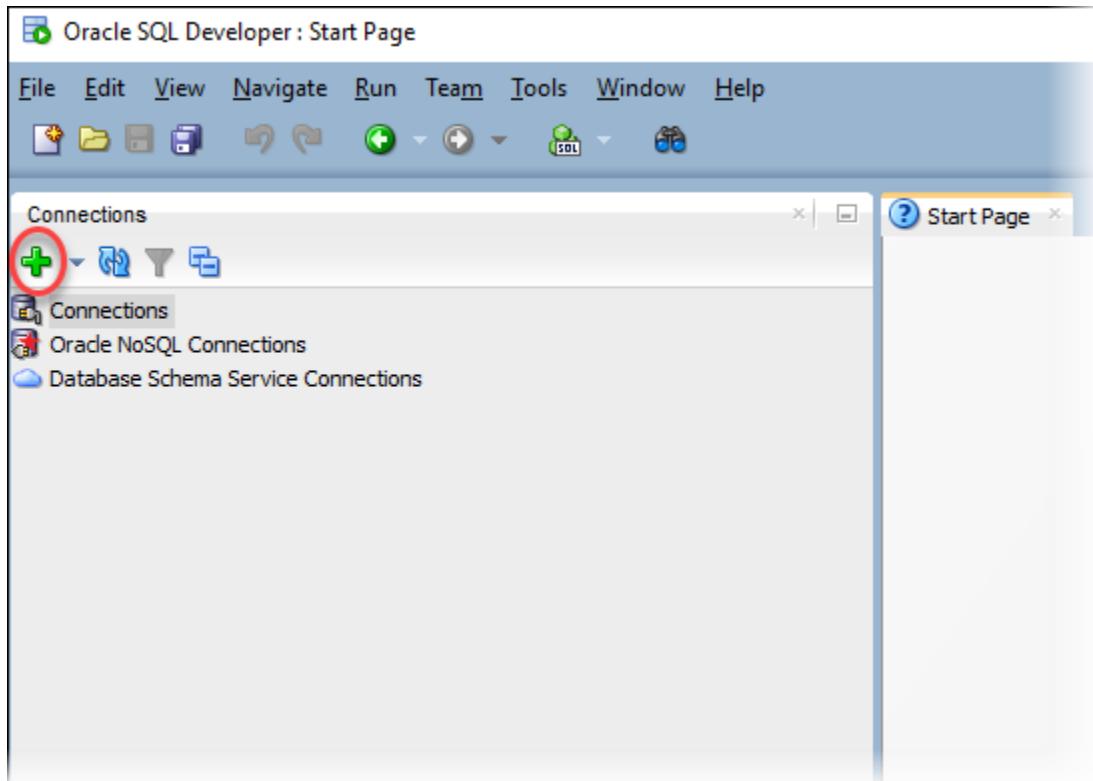
Connecting to Your DB Instance Using Oracle SQL Developer

In this procedure, you connect to your DB instance by using Oracle SQL Developer. To download a standalone version of this utility, see the [Oracle SQL Developer Downloads page](#).

To connect to your DB instance, you need its DNS name and port number. For information about finding the DNS name and port number for a DB instance, see [Finding the Endpoint of Your DB Instance \(p. 955\)](#).

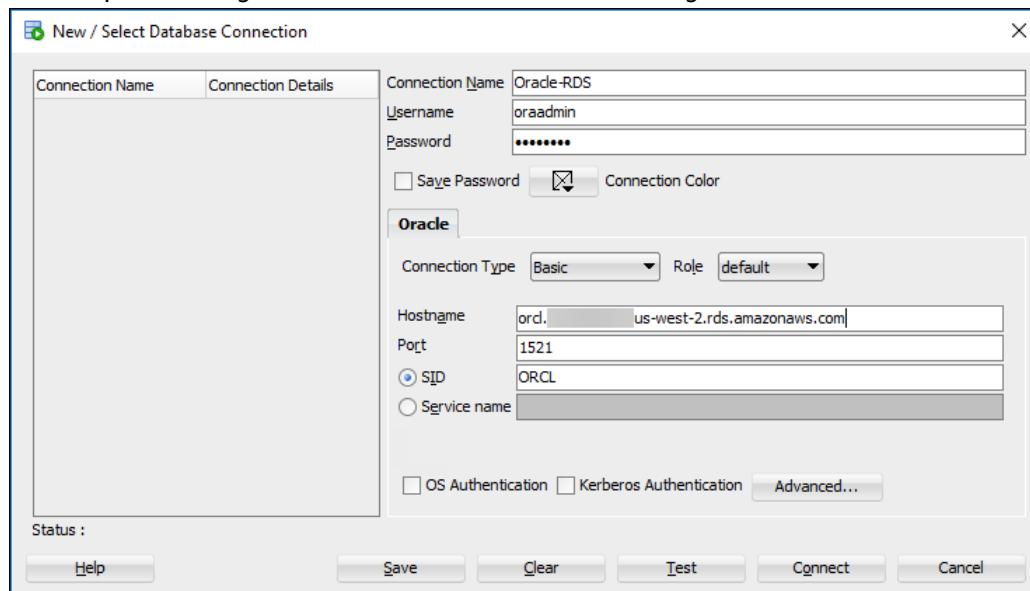
To connect to a DB instance using SQL Developer

1. Start Oracle SQL Developer.
2. On the **Connections** tab, choose the **add (+)** icon.



3. In the **New/Select Database Connection** dialog box, provide the information for your DB instance:
 - For **Connection Name**, enter a name that describes the connection, such as Oracle-RDS.
 - For **Username**, enter the name of the database administrator for the DB instance.
 - For **Password**, enter the password for the database administrator.
 - For **Hostname**, enter the DNS name of the DB instance.
 - For **Port**, enter the port number.
 - For **SID**, enter the Oracle database SID.

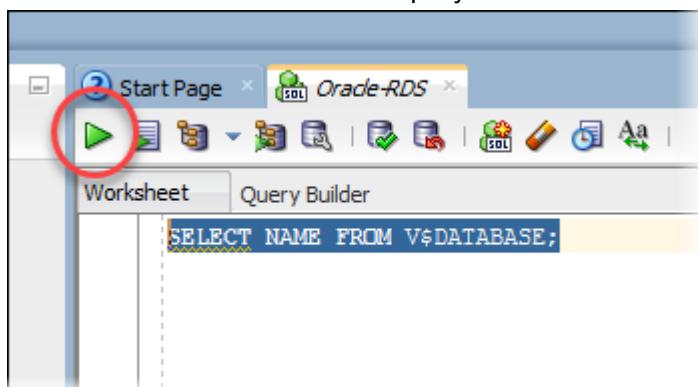
The completed dialog box should look similar to the following.



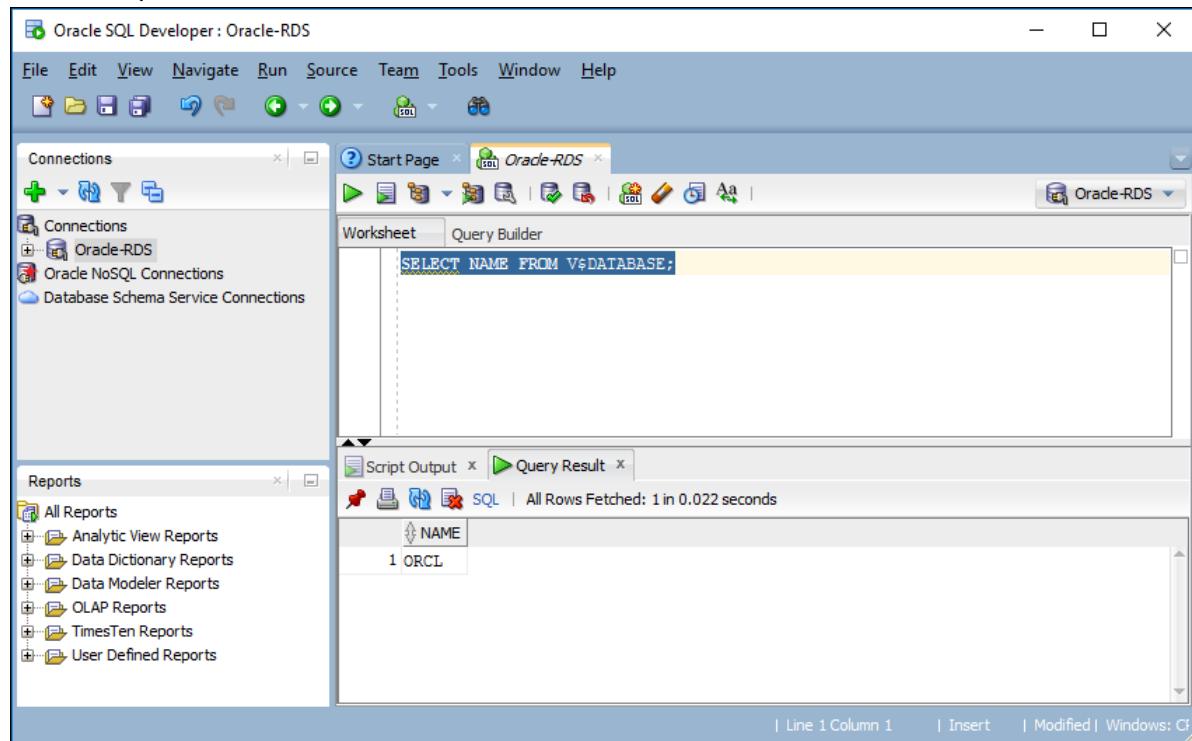
4. Choose **Connect**.
5. You can now start creating your own databases and running queries against your DB instance and databases as usual. To run a test query against your DB instance, do the following:
 - a. In the **Worksheet** tab for your connection, enter the following SQL query.

```
SELECT NAME FROM V$DATABASE;
```

- b. Choose the **execute** icon to run the query.



SQL Developer returns the database name.



Connecting to Your DB Instance Using SQL*Plus

You can use a utility like SQL*Plus to connect to an Amazon RDS DB instance running Oracle. To download a standalone version of SQL*Plus, see [SQL*Plus User's Guide and Reference](#).

To connect to your DB instance, you need its DNS name and port number. For information about finding the DNS name and port number for a DB instance, see [Finding the Endpoint of Your DB Instance \(p. 955\)](#).

Example To connect to an Oracle DB instance using SQL*Plus

In the following examples, substitute the user name of your DB instance administrator. Also, substitute the DNS name for your DB instance, and then include the port number and the Oracle SID. The SID value is the name of the DB instance's database that you specified when you created the DB instance, and not the name of the DB instance.

For Linux, OS X, or Unix:

```
sqlplus 'user_name@(DESCRIPTION=(ADDRESS=(PROTOCOL=TCP)(HOST=dns_name)(PORT=port))(CONNECT_DATA=(SID=database_name)))'
```

For Windows:

```
sqlplus user_name@(DESCRIPTION=(ADDRESS=(PROTOCOL=TCP)(HOST=dns_name)(PORT=port))(CONNECT_DATA=(SID=database_name)))
```

You should see output similar to the following.

```
SQL*Plus: Release 12.1.0.2.0 Production on Mon Aug 21 09:42:20 2017
```

After you enter the password for the user, the SQL prompt appears.

```
SQL>
```

Note

The shorter format connection string (Easy connect or EZCONNECT), such as `sqlplus USER/PASSWORD@LONGER-THAN-63-CHARS-RDS-ENDPOINT-HERE:1521/DATABASE_IDENTIFIER`, might encounter a maximum character limit and should not be used to connect.

Security Group Considerations

For you to connect to your DB instance, it must be associated with a security group that contains the IP addresses and network configuration that you use to access the DB instance. You might have associated your DB instance with an appropriate security group when you created it. If you assigned a default, nonconfigured security group when you created the DB instance, the DB instance firewall prevents connections.

If you need to create a new security group to enable access, the type of security group that you create depends on which Amazon EC2 platform your DB instance is on. To determine your platform, see [Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform \(p. 506\)](#). In general, if your DB instance is on the *EC2-Classic* platform, you create a DB security group; if your DB instance is on the *VPC* platform, you create a VPC security group. For information about creating a new security group, see [Controlling Access with Security Groups \(p. 486\)](#).

After you create the new security group, you modify your DB instance to associate it with the security group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

You can enhance security by using SSL to encrypt connections to your DB instance. For more information, see [Oracle Secure Sockets Layer \(p. 1051\)](#).

Dedicated and Shared Server Processes

Server processes handle user connections to an Oracle DB instance. By default, the Oracle DB instance uses dedicated server processes. With dedicated server processes, each server process services only one user process. You can optionally configure shared server processes. With shared server processes, each server process can service multiple user processes.

You might consider using shared server processes when a high number of user sessions are using too much memory on the server. You might also consider shared server processes when sessions connect and disconnect very often, resulting in performance issues. There are also disadvantages to using shared server processes. For example, they can strain CPU resources, and they are more complicated to configure and administer.

For more information about dedicated and shared server processes, see [About Dedicated and Shared Server Processes](#) in the Oracle documentation. For more information about configuring shared server

processes on an Amazon RDS Oracle DB instance, see [How do I configure Amazon RDS for Oracle Database to work with shared servers?](#) in the Knowledge Center.

Troubleshooting the Connection to Your Oracle DB Instance

The following are issues you might encounter when you try to connect to your Oracle DB instance.

Issue	Troubleshooting Suggestions
Unable to connect to your DB instance.	For a newly created DB instance, the DB instance has a status of creating until it is ready to use. When the state changes to available , you can connect to the DB instance. Depending on the DB instance class and the amount of storage, it can take up to 20 minutes before the new DB instance is available.
Unable to connect to your DB instance.	If you can't send or receive communications over the port that you specified when you created the DB instance, you can't connect to the DB instance. Check with your network administrator to verify that the port you specified for your DB instance allows inbound and outbound communication.
Unable to connect to your DB instance.	<p>The access rules enforced by your local firewall and the IP addresses you authorized to access your DB instance in the security group for the DB instance might not match. The problem is most likely the inbound or outbound rules on your firewall.</p> <p>You can add or edit an inbound rule in the security group. For Source, choose My IP. This allows access to the DB instance from the IP address detected in your browser. For more information, see Amazon Virtual Private Cloud VPCs and Amazon RDS (p. 505).</p> <p>For more information about security groups, see Controlling Access with Security Groups (p. 486).</p> <p>To walk through the process of setting up rules for your security group, see Tutorial: Create an Amazon VPC for Use with a DB Instance (p. 521).</p>
Connect failed because target host or object does not exist – Oracle, Error: ORA-12545	<p>Make sure that you specified the server name and port number correctly. For Server name, enter the DNS name from the console.</p> <p>For information about finding the DNS name and port number for a DB instance, see Finding the Endpoint of Your DB Instance (p. 955).</p>
Invalid username/password; logon denied – Oracle, Error: ORA-01017	You were able to reach the DB instance, but the connection was refused. This is usually caused by providing an incorrect user name or password. Verify the user name and password, and then retry.

Updating Applications to Connect to Oracle DB Instances Using New SSL/TLS Certificates

As of September 19, 2019, Amazon RDS has published new Certificate Authority (CA) certificates for connecting to your RDS DB instances using Secure Socket Layer or Transport Layer Security (SSL/TLS). The previous CA certificates expire on March 5, 2020. Following, you can find information about updating your applications to use the new certificates. If your application connects to an RDS DB instance using SSL/TLS, you must take the following steps before **March 5, 2020**. Doing this means you can avoid interruption of connectivity between your applications and your RDS DB instances.

This topic can help you to determine whether any client applications use SSL/TLS to connect to your DB instances.

Note

For client applications that use SSL/TLS to connect to your DB instances, you must update your client application trust stores to include the new CA certificates.

After you update your CA certificates in the client application trust stores, you can rotate the certificates on your DB instances. We strongly recommend testing these procedures in a development or staging environment before implementing them in your production environments.

For more information about certificate rotation, see [Rotating Your SSL/TLS Certificate \(p. 433\)](#). For information about using SSL/TLS with Oracle DB instances, see [Oracle Secure Sockets Layer \(p. 1051\)](#).

Topics

- [Determining Whether Any Applications Are Connecting to Your Oracle DB Instance Using SSL \(p. 962\)](#)
- [Updating Your Application Trust Store \(p. 963\)](#)
- [Example Java Code for Establishing SSL Connections \(p. 964\)](#)

Determining Whether Any Applications Are Connecting to Your Oracle DB Instance Using SSL

If your Oracle DB instance uses an option group with the `SSL` option added, you might be using SSL. Check this by following the instructions in [Listing the Options and Option Settings for an Option Group \(p. 203\)](#). For information about the `SSL` option, see [Oracle Secure Sockets Layer \(p. 1051\)](#).

Check the listener log to determine whether there are SSL connections. The following is sample output in a listener log.

```
date time * (CONNECT_DATA=(CID=(PROGRAM=program)
(HOST=host)(USER=user))(SID=sid)) *
(ADDRESS=(PROTOCOL=tcp)(HOST=host)(PORT=port)) * establish * ORCL * 0
```

When `PROTOCOL` has the value `tcp` for an entry, it shows an SSL connection. However, when `HOST` is `127.0.0.1`, you can ignore the entry. Connections from `127.0.0.1` are a local management agent on the DB instance. These connections aren't external SSL connections. Therefore, you have applications connecting using SSL if you see listener log entries where `PROTOCOL` is `tcp` and `HOST` is not `127.0.0.1`.

To check the listener log, you can publish the log to Amazon CloudWatch Logs. For more information, see [Publishing Oracle Logs to Amazon CloudWatch Logs \(p. 415\)](#).

Updating Your Application Trust Store

You can update the trust store for applications that use SQL*Plus or JDBC for SSL/TLS connections.

Updating Your Application Trust Store for SQL*Plus

You can update the trust store for applications that use SQL*Plus for SSL/TLS connections.

Note

When you update the trust store, you can retain older certificates in addition to adding the new certificates.

To update the trust store for SQL*Plus applications

1. Download the 2019 root certificate that works for all AWS Regions and put the file in the `ssl_wallet` directory.

For information about downloading the root certificate, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).

2. Run the following command to update the Oracle wallet.

```
prompt>orapki wallet add -wallet $ORACLE_HOME/ssl_wallet -trusted_cert -cert  
$ORACLE_HOME/ssl_wallet/rds-ca-2019-root.pem -auto_login_only
```

Replace the file name with the one that you downloaded.

3. Run the following command to confirm that the wallet was updated successfully.

```
prompt>orapki wallet display -wallet $ORACLE_HOME/ssl_wallet
```

Your output should contain the following.

```
Trusted Certificates:  
Subject: CN=Amazon RDS Root 2019 CA,OU=Amazon RDS,O=Amazon Web Services\,  
Inc.,L=Seattle,ST=Washington,C=US
```

Updating Your Application Trust Store for JDBC

You can update the trust store for applications that use JDBC for SSL/TLS connections.

To update the trust store for JDBC applications

1. Download the 2019 root certificate that works for all AWS Regions and put the file in the `ssl_wallet` directory.

For information about downloading the root certificate, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).

2. Convert the certificate to .der format using the following command.

```
openssl x509 -outform der -in rds-ca-2019-root.pem -out rds-ca-2019-root.der
```

Replace the file name with the one that you downloaded.

3. Import the certificate into the key store using the following command.

```
keytool -import -alias rds-root -keystore clientkeystore -file rds-ca-2019-root.der
```

4. Confirm that the key store was updated successfully.

```
keytool -list -v -keystore clientkeystore.jks
```

Enter the key store password when you are prompted for it.

Your output should contain the following.

```
rds-root, date, trustedCertEntry,  
Certificate fingerprint (SHA1):  
D4:0D:DB:29:E3:75:0D:FF:A6:71:C3:14:0B:BF:5F:47:8D:1C:80:96  
# This fingerprint should match the output from the below command  
openssl x509 -fingerprint -in rds-ca-2019-root.pem -noout
```

Example Java Code for Establishing SSL Connections

The following code example shows how to set up the SSL connection using JDBC.

```
import java.sql.Connection;  
import java.sql.DriverManager;  
import java.sql.SQLException;  
import java.util.Properties;  
  
public class OracleSslConnectionTest {  
    private static final String DB_SERVER_NAME = "<dns-name-provided-by-amazon-rds>";  
    private static final Integer SSL_PORT = "<ssl-option-port-configured-in-option-group>";  
    private static final String DB_SID = "<oracle-sid>";  
    private static final String DB_USER = "<user name>";  
    private static final String DB_PASSWORD = "<password>";  
    // This key store has only the prod root ca.  
    private static final String KEY_STORE_FILE_PATH = "<file-path-to-keystore>";  
    private static final String KEY_STORE_PASS = "<keystore-password>";  
  
    public static void main(String[] args) throws SQLException {  
        final Properties properties = new Properties();  
        final String connectionString = String.format(  
            "jdbc:oracle:thin:@(DESCRIPTION=(ADDRESS=(PROTOCOL=TCPS)(HOST=%s)(PORT=%d))  
(CONNECT_DATA=(SID=%s)))",  
            DB_SERVER_NAME, SSL_PORT, DB_SID);  
        properties.put("user", DB_USER);  
        properties.put("password", DB_PASSWORD);  
        properties.put("oracle.jdbc.J2EE13Compliant", "true");
```

```
        properties.put("javax.net.ssl.trustStore", KEY_STORE_FILE_PATH);
        properties.put("javax.net.ssl.trustStoreType", "JKS");
        properties.put("javax.net.ssl.trustStorePassword", KEY_STORE_PASS);
        final Connection connection = DriverManager.getConnection(connectionString,
properties);
        // If no exception, that means handshake has passed, and an SSL connection can be
opened
    }
}
```

Important

After you have determined that your database connections use SSL/TLS and have updated your application trust store, you can update your database to use the rds-ca-2019 certificates. For instructions, see step 3 in [Updating Your CA Certificate by Modifying Your DB Instance \(p. 434\)](#).

Modifying a DB Instance Running the Oracle Database Engine

You can change the settings of a DB instance to accomplish tasks such as adding additional storage or changing the DB instance class. In this topic, you learn how to modify an Amazon RDS Oracle DB instance, and about the settings for Oracle instances. We recommend that you test any changes on a test instance before modifying a production instance, so that you fully understand the impact of each change. This practice is especially important when upgrading database versions.

After you modify your DB instance settings, you can apply the changes immediately, or apply them during the next maintenance window for the DB instance. Some modifications cause an interruption by restarting the DB instance.

In addition to modifying Oracle instances as described directly following, you can also change settings for sqlnet.ora parameters for an Oracle DB instance as described in [Modifying Oracle sqlnet.ora Parameters \(p. 976\)](#), at the end of this topic.

Console

To modify an Oracle DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to modify.
3. Choose **Modify**. The **Modify DB Instance** page appears.
4. Change any of the settings that you want. For information about each setting, see [Settings for Oracle DB Instances \(p. 967\)](#).
5. When all the changes are as you want them, choose **Continue** and check the summary of modifications.
6. To apply the changes immediately, choose **Apply immediately**. Choosing this option can cause an outage in some cases. For more information, see [Using the Apply Immediately Parameter \(p. 150\)](#).
7. On the confirmation page, review your changes. If they are correct, choose **Modify DB Instance** to save your changes.

Alternatively, choose **Back** to edit your changes, or choose **Cancel** to cancel your changes.

AWS CLI

To modify an Oracle DB instance by using the AWS CLI, call the `modify-db-instance` command. Specify the DB instance identifier, and the parameters for the settings that you want to modify. For information about each parameter, see [Settings for Oracle DB Instances \(p. 967\)](#).

Example

The following code modifies `mydbinstance` by setting the backup retention period to 1 week (7 days). The code enables automatic minor version upgrades by using `--auto-minor-version-upgrade`. To disable automatic minor version upgrades, use `--no-auto-minor-version-upgrade`. The changes are applied during the next maintenance window by using `--no-apply-immediately`. Use `--apply-immediately` to apply the changes immediately. For more information, see [Using the Apply Immediately Parameter \(p. 150\)](#).

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier mydbinstance \
--backup-retention-period 7 \
--auto-minor-version-upgrade \
--no-apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier mydbinstance ^
--backup-retention-period 7 ^
--auto-minor-version-upgrade ^
--no-apply-immediately
```

RDS API

To modify an Oracle DB instance by using the Amazon RDS API, call the [ModifyDBInstance](#) operation. Specify the DB instance identifier, and the parameters for the settings that you want to modify. For information about each parameter, see [Settings for Oracle DB Instances \(p. 967\)](#).

Example

The following code modifies *mydbinstance* by setting the backup retention period to 1 week (7 days) and enabling automatic minor version upgrades. These changes are applied during the next maintenance window.

```
https://rds.amazonaws.com/
?Action=ModifyDBInstance
&ApplyImmediately=false
&AutoMinorVersionUpgrade=true
&BackupRetentionPeriod=7
&DBInstanceIdentifier=mydbinstance
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20131016/us-west-1/rds/aws4_request
&X-Amz-Date=20131016T233051Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=087a8eb41cb1ab0fc9ec1575f23e73757ffc6a1e42d7d2b30b9cc0be988cff97
```

Settings for Oracle DB Instances

The following table contains details about which settings you can modify, which settings you can't modify, when the changes can be applied, and whether the changes cause downtime for the DB instance.

Setting	Setting Description	When the Change Occurs	Downtime Notes
Allocated storage	<p>The storage, in gigabytes, that you want to allocate for your DB instance. You can only increase the allocated storage, you can't reduce the allocated storage.</p> <p>You can't modify allocated storage if the DB instance status is storage-locked.</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	No downtime. Performance might be degraded during the change.

Setting	Setting Description	When the Change Occurs	Downtime Notes
	<p>optimization or if the allocated storage for the DB instance has been modified in the last six hours.</p> <p>The maximum storage allowed depends on the storage type. For more information, see Amazon RDS DB Instance Storage (p. 135).</p>		
Auto minor version upgrade	<p>Enable auto minor version upgrade to enable your DB instance to receive preferred minor DB engine version upgrades automatically when they become available. Amazon RDS performs automatic minor version upgrades in the maintenance window.</p>	–	–
Backup retention period	<p>The number of days that automatic backups are retained. To disable automatic backups, set the backup retention period to 0.</p> <p>For more information, see Working With Backups (p. 254).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false and you change the setting from a nonzero value to another nonzero value, the change is applied asynchronously, as soon as possible. Otherwise, the change occurs during the next maintenance window.</p>	An outage occurs if you change from 0 to a nonzero value, or from a nonzero value to 0.
Backup window	<p>The time range during which automated backups of your databases occur. The backup window is a start time in Universal Coordinated Time (UTC), and a duration in hours.</p> <p>For more information, see Working With Backups (p. 254).</p>	The change is applied asynchronously, as soon as possible.	–
Certificate authority	The certificate that you want to use.	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change.

Setting	Setting Description	When the Change Occurs	Downtime Notes
Copy tags to snapshots	If you have any DB instance tags, this option copies them when you create a DB snapshot. For more information, see Tagging Amazon RDS Resources (p. 177) .	The change occurs immediately. This setting ignores the Apply immediately setting.	–
Database authentication	The database authentication option you want to use. Choose Password authentication to authenticate database users with database passwords only. Choose Password and Kerberos authentication to authenticate database users with database passwords and Kerberos authentication through an AWS Managed Microsoft AD created with AWS Directory Service. Next, choose the directory or choose Create a new Directory . For more information, see Using Kerberos Authentication with Amazon RDS for Oracle (p. 1134) .	If Apply immediately is set to true, the change occurs immediately. If Apply immediately is set to false, the change occurs during the next maintenance window.	When you enable Password and Kerberos authentication , a brief outage occurs.
Database port	The port that you want to use to access the database. The port value must not match any of the port values specified for options in the option group for the DB instance.	The change occurs immediately. This setting ignores the Apply immediately setting.	The DB instance is rebooted immediately.
DB engine version	The version of the Oracle database engine that you want to use. Before you upgrade your production DB instances, we recommend that you test the upgrade process on a test instance to verify its duration and to validate your applications. We do not recommend upgrading micro DB instances because they have limited CPU resources and the upgrade process may take hours to complete. An alternative to upgrading micro DB instances with small storage (10-20 GiB) is to copy your data using Data Pump, where we also recommend testing before migrating your production instances. For more information, see Upgrading the Oracle DB Engine (p. 979) .	If Apply immediately is set to true, the change occurs immediately. If Apply immediately is set to false, the change occurs during the next maintenance window.	An outage occurs during this change.

Setting	Setting Description	When the Change Occurs	Downtime Notes
DB instance class	<p>The DB instance class that you want to use.</p> <p>For more information, see Choosing the DB Instance Class (p. 107) and DB Instance Class Support for Oracle (p. 917).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change.
DB instance identifier	<p>The DB instance identifier. This value is stored as a lowercase string.</p> <p>For more information about the effects of renaming a DB instance, see Renaming a DB Instance (p. 163).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change. The DB instance is rebooted.
DB parameter group	<p>The parameter group that you want associated with the DB instance.</p> <p>For more information, see Working with DB Parameter Groups (p. 211) and Modifying Oracle sqlnet.ora Parameters (p. 976).</p>	The parameter group change occurs immediately.	<p>An outage doesn't occur during this change. When you change the parameter group, changes to some parameters are applied to the DB instance immediately without a reboot. Changes to other parameters are applied only after the DB instance is rebooted.</p> <p>For more information, see Rebooting a DB Instance (p. 166).</p>
Deletion protection	<p>Enable deletion protection to prevent your DB instance from being deleted.</p> <p>For more information, see Deleting a DB Instance (p. 174).</p>	–	–
Enhanced Monitoring	<p>Enable enhanced monitoring to enable gathering metrics in real time for the operating system that your DB instance runs on.</p> <p>For more information, see Enhanced Monitoring (p. 307).</p>	–	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
License model	<p>license-included to use the general license agreement for Oracle. bring-your-own-license to use your existing Oracle license.</p> <p>For more information, see Oracle Licensing (p. 916).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change.
Log exports	<p>The types of Oracle database log files to publish to Amazon CloudWatch Logs.</p> <p>For more information, see Oracle Database Log Files (p. 413).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–
Maintenance Window	<p>The time range during which system maintenance occurs. System maintenance includes upgrades, if applicable. The maintenance window is a start time in Universal Coordinated Time (UTC), and a duration in hours.</p> <p>If you set the window to the current time, there must be at least 30 minutes between the current time and end of the window to ensure any pending changes are applied.</p> <p>For more information, see The Amazon RDS Maintenance Window (p. 156).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	If there are one or more pending actions that cause an outage, and the maintenance window is changed to include the current time, then those pending actions are applied immediately, and an outage occurs.
Multi-AZ deployment	<p>Yes to deploy your DB instance in multiple Availability Zones. Otherwise, No.</p> <p>For more information, see Regions and Availability Zones (p. 132).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	–
New master password	The password for your master user. The password must contain from 8 to 30 alphanumeric characters.	The change is applied asynchronously, as soon as possible. This setting ignores the Apply immediately setting.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Option group	<p>The option group that you want associated with the DB instance.</p> <p>For more information, see Working with Option Groups (p. 196).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	<p>When you add the APEX options to an existing DB instance, a brief outage occurs while your DB instance is automatically restarted.</p> <p>When you add the OEM option to an existing DB instance, the change can cause a brief (sub-second) period during which new connections are rejected. Existing connections are not interrupted.</p>
Performance Insights	<p>Enable Performance Insights to monitor your DB instance load so that you can analyze and troubleshoot your database performance.</p> <p>Choose a retention period to determine how much rolling data history to keep. The default of seven days is in the free tier. Long-term retention (two years) is priced per vCPU per month.</p> <p>You can't change the master key after the database is created.</p> <p>For more information, see Using Amazon RDS Performance Insights (p. 319).</p>	<p>The change occurs immediately. This setting ignores the Apply immediately setting.</p>	–
Public accessibility	<p>Yes to give the DB instance a public IP address, meaning that it's accessible outside the VPC. To be publicly accessible, the DB instance also has to be in a public subnet in the VPC.</p> <p>No to make the DB instance accessible only from inside the VPC.</p> <p>For more information, see Hiding a DB Instance in a VPC from the Internet (p. 515).</p>	<p>The change occurs immediately. This setting ignores the Apply immediately setting.</p>	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Security group	<p>The security group that you want associated with the DB instance.</p> <p>For more information, see Working with DB Security Groups (EC2-Classic Platform) (p. 491).</p>	The change is applied asynchronously, as soon as possible. This setting ignores the Apply immediately setting.	–
Storage autoscaling	<p>Enable storage autoscaling to enable Amazon RDS to automatically increase storage when needed to avoid having your DB instance run out of storage space.</p> <p>Use Maximum storage threshold to set the upper limit for Amazon RDS to automatically increase storage for your DB instance. The default is 1,000 GiB.</p> <p>For more information, see Managing Capacity Automatically with Amazon RDS Storage Autoscaling (p. 232).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Storage type	<p>The storage type that you want to use.</p> <p>For more information, see Amazon RDS Storage Types (p. 135).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	<p>The following changes all result in a brief outage while the process starts. After that, you can use your database normally while the change takes place.</p> <ul style="list-style-type: none"> • From General Purpose (SSD) to Magnetic. • From General Purpose (SSD) to Provisioned IOPS (SSD). The outage only happens if the DB instance is single-AZ and you are using a custom parameter group. There is no outage for a multi-AZ DB instance. • From Magnetic to General Purpose (SSD). • From Magnetic to Provisioned IOPS (SSD). • From Provisioned IOPS (SSD) to Magnetic. • From Provisioned IOPS (SSD) to General Purpose (SSD). The outage only happens if the DB instance is single-AZ and you are using a custom parameter group. There is no outage for a multi-AZ DB instance.

Setting	Setting Description	When the Change Occurs	Downtime Notes
Subnet group	<p>The subnet group for the DB instance. You can use this setting to move your DB instance to a different VPC. If your DB instance isn't in a VPC, you can use this setting to move your DB instance into a VPC.</p> <p>For more information, see Moving a DB Instance Not in a VPC into a VPC (p. 520).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change. The DB instance is rebooted.

Modifying Oracle sqlnet.ora Parameters

The sqlnet.ora file includes parameters that configure Oracle Net features on Oracle database servers and clients. Using the parameters in the sqlnet.ora file, you can modify properties for connections in and out of the database.

For more information about why you might set sqlnet.ora parameters, see [Configuring Profile Parameters](#) in the Oracle documentation.

Setting sqlnet.ora Parameters

Amazon RDS Oracle parameter groups include a subset of sqlnet.ora parameters. You set them in the same way that you set other Oracle parameters. The `sqlnetora.` prefix identifies which parameters are sqlnet.ora parameters. For example, in an Oracle parameter group in Amazon RDS, the `default_sdu_size` sqlnet.ora parameter is `sqlnetora.default_sdu_size`.

For information about managing parameter groups and setting parameter values, see [Working with DB Parameter Groups \(p. 211\)](#).

Supported sqlnet.ora Parameters

Amazon RDS supports the following sqlnet.ora parameters. Changes to dynamic sqlnet.ora parameters take effect immediately.

Parameter	Valid Values	Static/ Dynamic	Description
<code>sqlnetora.default_sdu_size</code>	Oracle 11g: 512 to 65535 Oracle 12c: – 512 to 2097152	Dynamic	The session data unit (SDU) size, in bytes. The SDU is the amount of data that is put in a buffer and sent across the network at one time.
<code>sqlnetora.diag_ora_file_enabled</code>	ON or OFF	Dynamic	A value that enables or disables Automatic Diagnostic Repository (ADR) tracing. ON specifies that ADR file tracing is used. OFF specifies that non-ADR file tracing is used.
<code>sqlnetora.recv_192f_size</code>	268435456	Dynamic	The buffer space limit for receive operations of sessions, supported by the TCP/IP, TCP/IP with SSL, and SDP protocols.
<code>sqlnetora.send_192f_size</code>	268435456	Dynamic	The buffer space limit for send operations of sessions, supported by the TCP/IP, TCP/IP with SSL, and SDP protocols.
<code>sqlnetora.sqlnet_80_allowed_1_dynamiversion</code>		Dynamic	Minimum authentication protocol version allowed for clients, and servers acting as clients, to establish a connection to Oracle DB instances.
<code>sqlnetora.sqlnet_12_allowed_1_dynamiversion</code>	12, 12a	Dynamic	Minimum authentication protocol version allowed to establish a connection to Oracle DB instances.
<code>sqlnetora.sqlnet_max_expire_time</code>		Dynamic	Time interval, in minutes, to send a check to verify that client-server connections are active.

Parameter	Valid Values	Static/ Dynamic	Description
sqlnet.ora.sqlnet.inbound_connect_timeout	7200	Dynamic	Time, in seconds, for a client to connect with the database server and provide the necessary authentication information.
sqlnet.ora.sqlnet.outbound_connect_timeout	7200	Dynamic	Time, in seconds, for a client to establish an Oracle Net connection to the DB instance.
sqlnet.ora.sqlnet.reto_time	7200	Dynamic	Time, in seconds, for a database server to wait for client data after establishing a connection.
sqlnet.ora.sqlnet.send_time	7200	Dynamic	Time, in seconds, for a database server to complete a send operation to clients after establishing a connection.
sqlnet.ora.tcp_connect_timeout	7200	Dynamic	Time, in seconds, for a client to establish a TCP connection to the database server.
sqlnet.ora.tracing_level_server	16, OFF, USER, ADMIN, SUPPORT	Dynamic	For non-ADR tracing, turns server tracing on at a specified level or turns it off.

The default value for each supported sqlnet.ora parameter is the Oracle default for the release. For information about default values for Oracle 12c, see [Parameters for the sqlnet.ora File](#) in the 12c Oracle documentation. For information about default values for Oracle 11g, see [Parameters for the sqlnet.ora File](#) in the 11g Oracle documentation.

Viewing sqlnet.ora Parameters

You can view sqlnet.ora parameters and their settings using the AWS Management Console, the AWS CLI, or a SQL client.

Viewing sqlnet.ora Parameters Using the Console

For information about viewing parameters in a parameter group, see [Working with DB Parameter Groups \(p. 211\)](#).

In Oracle parameter groups, the `sqlnet.ora.` prefix identifies which parameters are sqlnet.ora parameters.

Viewing sqlnet.ora Parameters Using the AWS CLI

To view the sqlnet.ora parameters that were configured in an Oracle parameter group, use the AWS CLI `describe-db-parameters` command.

To view all of the sqlnet.ora parameters for an Oracle DB instance, call the AWS CLI `download-db-log-file-portion` command. Specify the DB instance identifier, the log file name, and the type of output.

Example

The following code lists all of the sqlnet.ora parameters for `mydbinstance`.

For Linux, OS X, or Unix:

```
aws rds download-db-log-file-portion \
--db-instance-identifier mydbinstance \
--log-file-name trace/sqlnet-parameters \
--output text
```

For Windows:

```
aws rds download-db-log-file-portion ^
--db-instance-identifier mydbinstance ^
--log-file-name trace/sqlnet-parameters ^
--output text
```

Viewing sqlnet.ora Parameters Using a SQL Client

After you connect to the Oracle DB instance in a SQL client, the following query lists the sqlnet.ora parameters.

```
SELECT * FROM TABLE
(rdsadmin.rds_file_util.read_text_file(
  p_directory => 'BDUMP',
  p_filename  => 'sqlnet-parameters'));
```

For information about connecting to an Oracle DB instance in a SQL client, see [Connecting to a DB Instance Running the Oracle Database Engine \(p. 955\)](#).

Upgrading the Oracle DB Engine

When Amazon RDS supports a new version of Oracle, you can upgrade your DB instances to the new version. Amazon RDS supports the following upgrades to an Oracle DB instance:

- Major Version Upgrades
- Minor Version Upgrades

In general, a major engine version upgrade can introduce changes that are not compatible with existing applications. In contrast, a minor version upgrade includes only changes that are backward-compatible with existing applications.

You must modify the DB instance manually to perform a major version upgrade. Minor version upgrades occur automatically if you enable auto minor version upgrades on your DB instance. In all other cases, you must modify the DB instance manually to perform a minor version upgrade.

An outage occurs while the upgrade takes place. The time for the outage varies based on your engine version and the size of your DB instance.

For information about what Oracle versions are available on Amazon RDS, see [Oracle Database Engine Release Notes \(p. 1173\)](#).

Topics

- [Overview of Upgrading \(p. 979\)](#)
- [Major Version Upgrades \(p. 980\)](#)
- [Oracle Minor Version Upgrades \(p. 981\)](#)
- [Oracle SE2 Upgrade Paths \(p. 981\)](#)
- [Option and Parameter Group Considerations \(p. 982\)](#)
- [Testing an Upgrade \(p. 982\)](#)
- [Upgrading an Oracle DB Instance \(p. 983\)](#)

Overview of Upgrading

Amazon RDS takes two DB snapshots during the upgrade process. The first DB snapshot is of the DB instance before any upgrade changes have been made. If the upgrade doesn't work for your databases, you can restore this snapshot to create a DB instance running the old version. The second DB snapshot is taken after the upgrade completes.

Note

Amazon RDS only takes DB snapshots if you have set the backup retention period for your DB instance to a number greater than 0. To change your backup retention period, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

After an upgrade is complete, you can't revert to the previous version of the database engine. If you want to return to the previous version, restore the DB snapshot that was taken before the upgrade to create a new DB instance.

If your DB instance is in a Multi-AZ deployment, both the primary and standby replicas are upgraded. If no operating system updates are required, the primary and standby DB instances are upgraded at the same time, and you experience an outage until the upgrade is complete.

If your DB instance is in a Multi-AZ deployment, and operating system updates are required, the operating system updates are applied when you request the database upgrade. In this case, the operating system is updated on the standby DB instance, and the standby DB instance is upgraded.

After that upgrade completes, the primary DB instance fails over to the standby DB instance, and the operating system is updated on the new standby DB instance (the former primary DB instance), and that database is upgraded.

Major Version Upgrades

Amazon RDS supports the following major version upgrades.

Current Version	Upgrade Supported
18.0.0.0	19.0.0.0
12.2.0.1	19.0.0.0
	18.0.0.0
12.1.0.2	19.0.0.0
	18.0.0.0
	12.2.0.1
11.2.0.4	19.0.0.0
	18.0.0.0
	12.2.0.1
	12.1.0.2.v5 and higher 12.1 versions

To perform a major version upgrade, modify the DB instance manually. Major version upgrades don't occur automatically.

In some cases, your current Oracle DB instance might be running on a DB instance class that isn't supported for the version to which you are upgrading. In such a case, you must migrate the DB instance to a supported DB instance class before you upgrade. For more information about the supported DB instance classes for each version and edition of Amazon RDS Oracle, see [Choosing the DB Instance Class \(p. 107\)](#).

Before you perform a major version upgrade, Oracle recommends that you gather optimizer statistics on the DB instance that you are upgrading. Gathering optimizer statistics can reduce DB instance downtime during the upgrade. To gather optimizer statistics, connect to the DB instance as the master user, and run the `DBMS_STATS.GATHER_DICTIONARY_STATS` procedure, as in the following example.

```
EXEC DBMS_STATS.GATHER_DICTIONARY_STATS;
```

For more information, see [Gathering Optimizer Statistics to Decrease Oracle Database Downtime](#) in the Oracle documentation.

Note

Major version upgrades aren't supported for deprecated Oracle versions, such as Oracle version 11.2.0.3 and 11.2.0.2.

Major version downgrades aren't supported.

A major version upgrade from 11g to 12c must upgrade to an Oracle Patch Set Update (PSU) that was released in the same month or later.

For example, a major version upgrade from Oracle version 11.2.0.4.v14 to Oracle version 12.1.0.2.v11 is supported. However, a major version upgrade from Oracle version 11.2.0.4.v14

to Oracle version 12.1.0.2.v9 isn't supported. This is because Oracle version 11.2.0.4.v14 was released in October 2017, and Oracle version 12.1.0.2.v9 was released in July 2017. For information about the release date for each Oracle PSU, see [Oracle Database Engine Release Notes \(p. 1173\)](#).

Oracle Minor Version Upgrades

A minor version upgrade applies an Oracle Database Patch Set Update (PSU) or Release Update (RU) in a major version.

An Amazon RDS for Oracle DB instance is scheduled to be upgraded automatically during its next maintenance window when it meets the following conditions:

- The DB instance has the **Auto minor version upgrade** option enabled.
- The DB instance is not running the latest minor DB engine version.

The DB instance is upgraded to the latest quarterly PSU or RU four to six weeks after it is made available by Amazon RDS for Oracle. For more information about PSUs and RUs, see [Oracle Database Engine Release Notes \(p. 1173\)](#).

The following minor version upgrades aren't supported.

Current Version	Upgrade Not Supported
12.1.0.2.v6	12.1.0.2.v7
12.1.0.2.v5	12.1.0.2.v7
12.1.0.2.v5	12.1.0.2.v6

Note

Minor version downgrades aren't supported.

Oracle SE2 Upgrade Paths

The following table shows supported upgrade paths to Standard Edition Two (SE2). For more information about the License Included and Bring Your Own License (BYOL) models, see [Oracle Licensing \(p. 916\)](#).

Your Existing Configuration	Supported SE2 Configuration
12.2.0.1 SE2, BYOL	12.2.0.1 SE2, BYOL or License Included
12.1.0.2 SE2, BYOL	12.2.0.1 SE2, BYOL or License Included 12.1.0.2 SE2, BYOL or License Included
11.2.0.4 SE1, BYOL or License Included	12.2.0.1 SE2, BYOL or License Included
11.2.0.4 SE, BYOL	12.1.0.2 SE2, BYOL or License Included

To upgrade from your existing configuration to a supported SE2 configuration, use a supported upgrade path. For more information, see [Major Version Upgrades \(p. 980\)](#).

Option and Parameter Group Considerations

Option Group Considerations

If your DB instance uses a custom option group, in some cases Amazon RDS can't automatically assign your DB instance a new option group. For example, this occurs when you upgrade to a new major version. In those cases, you must specify a new option group when you upgrade. We recommend that you create a new option group, and add the same options to it as in your existing custom option group.

For more information, see [Creating an Option Group \(p. 197\)](#) or [Making a Copy of an Option Group \(p. 199\)](#).

If your DB instance uses a custom option group that contains the APEX option, in some cases you can reduce the time it takes to upgrade your DB instance by upgrading your version of APEX at the same time as your DB instance. For more information, see [Upgrading the APEX Version \(p. 1023\)](#).

Parameter Group Considerations

If your DB instance uses a custom parameter group, in some cases Amazon RDS can't automatically assign your DB instance a new parameter group. For example, this occurs when you upgrade to a new major version. In those cases, you must specify a new parameter group when you upgrade. We recommend that you create a new parameter group, and configure the parameters as in your existing custom parameter group.

For more information, see [Creating a DB Parameter Group \(p. 212\)](#) or [Copying a DB Parameter Group \(p. 217\)](#).

Testing an Upgrade

Before you perform a major version upgrade on your DB instance, you should thoroughly test your database and all applications that access the database for compatibility with the new version. We recommend that you use the following procedure.

To test a major version upgrade

1. Review the Oracle upgrade documentation for the new version of the database engine to see if there are compatibility issues that might affect your database or applications. For more information, see [Database Upgrade Guide](#) in the Oracle documentation.
2. If your DB instance uses a custom option group, create a new option group compatible with the new version you are upgrading to. For more information, see [Option Group Considerations \(p. 982\)](#).
3. If your DB instance uses a custom parameter group, create a new parameter group compatible with the new version you are upgrading to. For more information, see [Parameter Group Considerations \(p. 982\)](#).
4. Create a DB snapshot of the DB instance to be upgraded. For more information, see [Creating a DB Snapshot \(p. 264\)](#).
5. Restore the DB snapshot to create a new test DB instance. For more information, see [Restoring from a DB Snapshot \(p. 266\)](#).
6. Modify this new test DB instance to upgrade it to the new version, by using one of the following methods:
 - [Console \(p. 159\)](#)
 - [AWS CLI \(p. 160\)](#)

- [RDS API \(p. 160\)](#)
7. Perform testing:
- Run as many of your quality assurance tests against the upgraded DB instance as needed to ensure that your database and application work correctly with the new version.
 - Implement any new tests needed to evaluate the impact of any compatibility issues that you identified in step 1.
 - Test all stored procedures, functions, and triggers.
 - Direct test versions of your applications to the upgraded DB instance. Verify that the applications work correctly with the new version.
 - Evaluate the storage used by the upgraded instance to determine if the upgrade requires additional storage. You might need to choose a larger instance class to support the new version in production. For more information, see [Choosing the DB Instance Class \(p. 107\)](#).
8. If all tests pass, then perform the upgrade on your production DB instance. We recommend that you don't allow write operations to the DB instance until you confirm that everything is working correctly.

Upgrading an Oracle DB Instance

For information about manually or automatically upgrading an Oracle DB instance, see [Upgrading a DB Instance Engine Version \(p. 159\)](#).

Upgrading an Oracle DB Snapshot

If you have existing manual DB snapshots, you might want to upgrade a snapshot to a later version of the Oracle database engine.

When Oracle stops providing patches for a version, and therefore Amazon RDS deprecates the version, you can upgrade your snapshots that correspond to the deprecated version. For more information, see [Oracle Engine Version Management \(p. 934\)](#).

The following snapshot upgrades are currently supported.

Current Snapshot Version	Supported Snapshot Upgrade
12.1.0.1	12.1.0.2.v8
11.2.0.3	11.2.0.4.v11
11.2.0.2	11.2.0.4.v12

Amazon RDS supports upgrading snapshots in all AWS Regions.

Console

To upgrade an Oracle DB snapshot

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Snapshots**, and then select the DB snapshot that you want to upgrade.
3. For **Actions**, choose **Modify Snapshot**. The **Modify DB Snapshot** page appears.
4. For **DB engine version**, choose the version to upgrade the snapshot to.
5. (Optional) For **Option group**, choose the option group for the upgraded DB snapshot. The same option group considerations apply when upgrading a DB snapshot as when upgrading a DB instance. For more information, see [Option Group Considerations \(p. 982\)](#).
6. Choose **Modify Snapshot** to save your changes.

Alternatively, choose **Cancel** to cancel your changes.

AWS CLI

To upgrade an Oracle DB snapshot by using the AWS CLI, call the `modify-db-snapshot` command with the following parameters:

- `--db-snapshot-identifier` – The name of the DB snapshot.
- `--engine-version` – The version to upgrade the snapshot to.

You might also need to include the following parameter. The same option group considerations apply when upgrading a DB snapshot as when upgrading a DB instance. For more information, see [Option Group Considerations \(p. 982\)](#).

- `--option-group-name` – The option group for the upgraded DB snapshot.

Example

The following example upgrades a DB snapshot.

For Linux, OS X, or Unix:

```
aws rds modify-db-snapshot \
    --db-snapshot-identifier <mydbsnapshot> \
    --engine-version <11.2.0.4.v12> \
    --option-group-name <default:oracle-se1-11-2>
```

For Windows:

```
aws rds modify-db-snapshot ^
    --db-snapshot-identifier <mydbsnapshot> ^
    --engine-version <11.2.0.4.v12> ^
    --option-group-name <default:oracle-se1-11-2>
```

RDS API

To upgrade an Oracle DB snapshot by using the Amazon RDS API, call the [ModifyDBSnapshot](#) operation with the following parameters:

- `DBSnapshotIdentifier` – The name of the DB snapshot.
- `EngineVersion` – The version to upgrade the snapshot to.

You might also need to include the following parameter. The same option group considerations apply when upgrading a DB snapshot as when upgrading a DB instance. For more information, see [Option Group Considerations \(p. 982\)](#).

- `OptionGroupName` – The option group for the upgraded DB snapshot.

Example

The following example upgrades a DB snapshot.

```
https://rds.amazonaws.com/
    ?Action=ModifyDBSnapshot
    &DBSnapshotIdentifier=mydbsnapshot
    &EngineVersion=11.2.0.4.v12
    &OptionGroupName=default:oracle-se1-11-2
    &SignatureMethod=HmacSHA256
    &SignatureVersion=4
    &Version=2014-10-31
    &X-Amz-Algorithm=AWS4-HMAC-SHA256
    &X-Amz-Credential=AKIADQKE4SARGYLE/20131016/us-west-1/rds/aws4_request
    &X-Amz-Date=20131016T233051Z
    &X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
    &X-Amz-Signature=087a8eb41cb1ab5f99e81575f23e73757fffc6a1e42d7d2b30b9cc0be988cff97
```

Related Topics

- [Oracle Database Engine Release Notes \(p. 1173\)](#)
- [Upgrading the Oracle DB Engine \(p. 979\)](#)
- [Applying Updates for a DB Instance \(p. 154\)](#)

Importing Data into Oracle on Amazon RDS

How you import data into an Amazon RDS DB instance depends on the amount of data you have and the number and variety of database objects in your database. For example, you can use Oracle SQL Developer to import a simple, 20 MB database. You can use Oracle Data Pump to import complex databases, or databases that are several hundred megabytes or several terabytes in size.

You can also use AWS Database Migration Service (AWS DMS) to import data into an Amazon RDS DB instance. AWS DMS can migrate databases without downtime and, for many database engines, continue ongoing replication until you are ready to switch over to the target database. You can migrate to Oracle from either the same database engine or a different database engine using AWS DMS. If you are migrating from a different database engine, you can use the AWS Schema Conversion Tool to migrate schema objects that are not migrated by AWS DMS. For more information about AWS DMS, see [What is AWS Database Migration Service](#).

Before you use any of these migration techniques, we recommend the best practice of taking a backup of your database. After you import the data, you can back up your Amazon RDS DB instances by creating snapshots. Later, you can restore the database from the snapshots. For more information, see [Backing Up and Restoring Amazon RDS DB Instances \(p. 253\)](#).

Note

You can also import data into Oracle using files from Amazon S3. For example, you can download Data Pump files from Amazon S3 to the DB instance host. For more information, see [Amazon S3 Integration \(p. 1005\)](#).

Importing Using Oracle SQL Developer

For small databases, you can use Oracle SQL Developer, a graphical Java tool distributed without cost by Oracle. You can install this tool on your desktop computer (Windows, Linux, or Mac) or on one of your servers. Oracle SQL Developer provides options for migrating data between two Oracle databases, or for migrating data from other databases, such as MySQL, to Oracle. Oracle SQL Developer is best suited for migrating small databases. We recommend that you read the Oracle SQL Developer product documentation before you begin migrating your data.

After you install SQL Developer, you can use it to connect to your source and target databases. Use the **Database Copy** command on the Tools menu to copy your data to your Amazon RDS instance.

To download Oracle SQL Developer, go to <http://www.oracle.com/technetwork/developer-tools/sql-developer>.

Oracle also has documentation on how to migrate from other databases, including MySQL and SQL Server. For more information, see <http://www.oracle.com/technetwork/database/migration> in the Oracle documentation.

Importing Using Oracle Data Pump

Oracle Data Pump is a long-term replacement for the Oracle Export/Import utilities. Oracle Data Pump is the preferred way to move large amounts of data from an Oracle installation to an Amazon RDS DB instance. You can use Oracle Data Pump for several scenarios:

- Import data from an Oracle database (either on-premises or Amazon EC2 instance) to an Amazon RDS for Oracle DB instance.
- Import data from an RDS Oracle DB instance to an Oracle database (either on-premises or Amazon EC2 instance).
- Import data between RDS Oracle DB instances (for example, to migrate data from EC2-Classic to VPC).

To download Oracle Data Pump utilities, see [Oracle Database Software Downloads](#) on the Oracle Technology Network website.

For compatibility considerations when migrating between versions of Oracle Database, see [the Oracle documentation](#).

When you import data with Oracle Data Pump, you must transfer the dump file that contains the data from the source database to the target database. You can transfer the dump file using an Amazon S3 bucket or by using a database link between the two databases.

The following are best practices for using Oracle Data Pump to import data into an Amazon RDS for Oracle DB instance:

- Perform imports in `schema` or `table` mode to import specific schemas and objects.
- Limit the schemas you import to those required by your application.
- Do not import in `full` mode.

Because Amazon RDS for Oracle does not allow access to `SYS` or `SYSDBA` administrative users, importing in `full` mode, or importing schemas for Oracle-maintained components, might damage the Oracle data dictionary and affect the stability of your database.

- When loading large amounts of data, transfer the dump file to the target Amazon RDS for Oracle DB instance, take a DB snapshot of your instance, and then test the import to verify that it succeeds. If database components are invalidated, you can delete the DB instance and re-create it from the DB snapshot. The restored DB instance includes any dump files staged on the DB instance when you took the DB snapshot.
- Do not import dump files that were created using the Oracle Data Pump export parameters `TRANSPORT_TABLESPACES`, `TRANSPORTABLE`, or `TRANSPORT_FULL_CHECK`. Amazon RDS for Oracle DB instances do not support importing these dump files.

Topics

- [Importing Data with Oracle Data Pump and an Amazon S3 Bucket \(p. 988\)](#)
- [Importing Data with Oracle Data Pump and a Database Link \(p. 991\)](#)

Importing Data with Oracle Data Pump and an Amazon S3 Bucket

The following import process uses Oracle Data Pump and an Amazon S3 bucket. The process exports data on the source database using the Oracle `DBMS_DATAPUMP` package and puts the dump file in an Amazon S3 bucket. It then downloads the dump file from the Amazon S3 bucket to the `DATA_PUMP_DIR` directory on the target Amazon RDS Oracle DB instance. The final step imports the data from the copied dump file into the Amazon RDS Oracle DB instance using the `DBMS_DATAPUMP` package.

The process has the following requirements:

- You must have an Amazon S3 bucket available for file transfers, and the Amazon S3 bucket must be in the same AWS Region as the DB instance. For instructions, see [Create a Bucket](#) in the *Amazon Simple Storage Service Getting Started Guide*.
- You must prepare the Amazon S3 bucket for Amazon RDS integration by following the instructions in [Prerequisites for Amazon RDS Oracle Integration with Amazon S3 \(p. 1005\)](#).
- You must ensure that you have enough storage space to store the dump file on the source instance and the target DB instance.

Note

This process imports a dump file into the DATA_PUMP_DIR directory, a preconfigured directory on all Oracle DB instances. This directory is located on the same storage volume as your data files. When you import the dump file, the existing Oracle data files use more space. Thus, you should make sure that your DB instance can accommodate that additional use of space. The imported dump file is not automatically deleted or purged from the DATA_PUMP_DIR directory. To remove the imported dump file, use [UTL_FILE.FREMOVE](#), found on the Oracle website.

The import process using Oracle Data Pump and an Amazon S3 bucket has the following steps.

Topics

- [Step 1: Grant Privileges to the User on the Amazon RDS Target Instance \(p. 989\)](#)
- [Step 2: Use DBMS_DATAPUMP to Create a Dump File \(p. 989\)](#)
- [Step 3: Upload the Dump File to Your Amazon S3 Bucket \(p. 990\)](#)
- [Step 4: Copy the Exported Dump File from the Amazon S3 Bucket to the Target DB Instance \(p. 990\)](#)
- [Step 5: Use DBMS_DATAPUMP to Import the Data File on the Target DB Instance \(p. 991\)](#)
- [Step 6: Clean Up \(p. 991\)](#)

Step 1: Grant Privileges to the User on the Amazon RDS Target Instance

To grant privileges to the user on the RDS target instance, take the following steps:

1. Use SQL Plus or Oracle SQL Developer to connect to the Amazon RDS target Oracle DB instance into which the data will be imported. Connect as the Amazon RDS master user. For information about connecting to the DB instance, see [Connecting to a DB Instance Running the Oracle Database Engine \(p. 955\)](#).
2. Create the required tablespaces before you import the data. For more information, see [Creating and Sizing Tablespaces \(p. 1096\)](#).
3. If the user account into which the data is imported doesn't exist, create the user account and grant the necessary permissions and roles. If you plan to import data into multiple user schemas, create each user account and grant the necessary privileges and roles to it.

For example, the following commands create a new user and grant the necessary permissions and roles to import the data into the user's schema.

```
create user schema_1 identified by <password>;
grant create session, resource to schema_1;
alter user schema_1 quota 100M on users;
```

This example grants the new user the CREATE SESSION privilege and the RESOURCE role. Additional privileges and roles might be required depending on the database objects that you import.

Note

Replace *schema_1* with the name of your schema in this step and in the following steps.

Step 2: Use DBMS_DATAPUMP to Create a Dump File

Use SQL Plus or Oracle SQL Developer to connect to the source Oracle instance with an administrative user. If the source database is an Amazon RDS Oracle DB instance, connect with the Amazon RDS master user. Next, use the Oracle Data Pump utility to create a dump file.

The following script creates a dump file named *sample.dmp* in the DATA_PUMP_DIR directory.

```
DECLARE
hdnl NUMBER;
BEGIN
hdnl := DBMS_DATAPUMP.OPEN( operation => 'EXPORT', job_mode => 'SCHEMA', job_name=>null);
DBMS_DATAPUMP.ADD_FILE( handle => hdnl, filename => 'sample.dmp', directory =>
'DATA_PUMP_DIR', filetype => dbms_datapump.ku$_file_type_dump_file);
DBMS_DATAPUMP.ADD_FILE( handle => hdnl, filename => 'exp.log', directory =>
'DATA_PUMP_DIR', filetype => dbms_datapump.ku$_file_type_log_file);
DBMS_DATAPUMP.METADATA_FILTER(hdnl,'SCHEMA_EXPR','IN (''SCHEMA_1'')');
DBMS_DATAPUMP.START_JOB(hdnl);
END;
/
```

Note

Data Pump jobs are started asynchronously. For information about monitoring a Data Pump job, see [Monitoring Job Status](#) in the Oracle documentation.

Step 3: Upload the Dump File to Your Amazon S3 Bucket

Upload the dump file to the Amazon S3 bucket.

Use the Amazon RDS procedure `rdsadmin.rdsadmin_s3_tasks.upload_to_s3` to copy the dump file to the Amazon S3 bucket. The following example uploads all of the files from the `DATA_PUMP_DIR` directory to an Amazon S3 bucket named `mys3bucket`.

```
SELECT rdsadmin.rdsadmin_s3_tasks.upload_to_s3(
    p_bucket_name => 'mys3bucket',
    p_directory_name => 'DATA_PUMP_DIR')
AS TASK_ID FROM DUAL;
```

The SELECT statement returns the ID of the task in a VARCHAR2 data type.

For more information, see [Uploading Files from an Oracle DB Instance to an Amazon S3 Bucket \(p. 1011\)](#).

Step 4: Copy the Exported Dump File from the Amazon S3 Bucket to the Target DB Instance

Use SQL Plus or Oracle SQL Developer to connect to the Amazon RDS target Oracle DB instance. Next, use the Amazon RDS procedure `rdsadmin.rdsadmin_s3_tasks.download_from_s3` to copy the dump file from the Amazon S3 bucket to the target DB instance. The following example downloads all of the files from an Amazon S3 bucket named `mys3bucket` to the `DATA_PUMP_DIR` directory.

```
SELECT rdsadmin.rdsadmin_s3_tasks.download_from_s3(
    p_bucket_name => 'mys3bucket',
    p_directory_name => 'DATA_PUMP_DIR')
AS TASK_ID FROM DUAL;
```

The SELECT statement returns the ID of the task in a VARCHAR2 data type.

For more information, see [Downloading Files from an Amazon S3 Bucket to an Oracle DB Instance \(p. 1013\)](#).

Step 5: Use DBMS_DATAPUMP to Import the Data File on the Target DB Instance

Use Oracle Data Pump to import the schema in the DB instance. Additional options such as METADATA_REMAP might be required.

Connect to the DB instance with the Amazon RDS master user account to perform the import.

```
DECLARE
hdnl NUMBER;
BEGIN
hdnl := DBMS_DATAPUMP.OPEN( operation => 'IMPORT', job_mode => 'SCHEMA', job_name=>null);
DBMS_DATAPUMP.ADD_FILE( handle => hdnl, filename => 'sample_copied.dmp', directory =>
'DATA_PUMP_DIR', filetype => dbms_datapump.ku$_file_type_dump_file);
DBMS_DATAPUMP.METADATA_FILTER(hdnl,'SCHEMA_EXPR','IN ('''SCHEMA_1'''');
DBMS_DATAPUMP.START_JOB(hdnl);
END;
/
```

You can verify the data import by viewing the user's tables on the DB instance. For example, the following query returns the number of tables for **SCHEMA_1**.

```
select count(*) from dba_tables where owner='SCHEMA_1';
```

Step 6: Clean Up

After the data has been imported, you can delete the files that you don't want to keep. You can list the files in the DATA_PUMP_DIR using the following command.

```
select * from table(RDSADMIN.RDS_FILE_UTIL.LISTDIR('DATA_PUMP_DIR')) order by mtime;
```

To delete files in the DATA_PUMP_DIR that you no longer require, use the following command.

```
exec utl_file.fremove('DATA_PUMP_DIR','<file name>');
```

For example, the following command deletes the file named "sample_copied.dmp".

```
exec utl_file.fremove('DATA_PUMP_DIR','sample_copied.dmp');
```

Importing Data with Oracle Data Pump and a Database Link

The following import process uses Oracle Data Pump and the Oracle [DBMS_FILE_TRANSFER](#) package. The process connects to a source Oracle instance, which can be an on-premises or Amazon EC2 instance, or an Amazon RDS for Oracle DB instance. The process then exports data using the [DBMS_DATAPUMP](#) package. Next, it uses the [DBMS_FILE_TRANSFER.PUT_FILE](#) method to copy the dump file from the Oracle instance to the DATA_PUMP_DIR directory on the target Amazon RDS Oracle DB instance that is connected using a database link. The final step imports the data from the copied dump file into the Amazon RDS Oracle DB instance using the [DBMS_DATAPUMP](#) package.

The process has the following requirements:

- You must have execute privileges on the DBMS_FILE_TRANSFER and DBMS_DATAPUMP packages.
- You must have write privileges to the DATA_PUMP_DIR directory on the source DB instance.
- You must ensure that you have enough storage space to store the dump file on the source instance and the target DB instance.

Note

This process imports a dump file into the DATA_PUMP_DIR directory, a preconfigured directory on all Oracle DB instances. This directory is located on the same storage volume as your data files. When you import the dump file, the existing Oracle data files use more space. Thus, you should make sure that your DB instance can accommodate that additional use of space. The imported dump file is not automatically deleted or purged from the DATA_PUMP_DIR directory. To remove the imported dump file, use [UTL_FILE.FREMOVE](#), found on the Oracle website.

The import process using Oracle Data Pump and the DBMS_FILE_TRANSFER package has the following steps.

Topics

- [Step 1: Grant Privileges to the User on the Amazon RDS Target Instance \(p. 989\)](#)
- [Step 2: Grant Privileges to the User on the Source Database \(p. 992\)](#)
- [Step 3: Use DBMS_DATAPUMP to Create a Dump File \(p. 989\)](#)
- [Step 4: Create a Database Link to the Target DB Instance \(p. 990\)](#)
- [Step 5: Use DBMS_FILE_TRANSFER to Copy the Exported Dump File to the Target DB Instance \(p. 990\)](#)
- [Step 6: Use DBMS_DATAPUMP to Import the Data File to the Target DB Instance \(p. 991\)](#)
- [Step 7: Clean Up \(p. 991\)](#)

Step 1: Grant Privileges to the User on the Amazon RDS Target Instance

To grant privileges to the user on the RDS target instance, take the following steps:

1. Use SQL Plus or Oracle SQL Developer to connect to the Amazon RDS target Oracle DB instance into which the data will be imported. Connect as the Amazon RDS master user. For information about connecting to the DB instance, see [Connecting to a DB Instance Running the Oracle Database Engine \(p. 955\)](#).
2. Create the required tablespaces before you import the data. For more information, see [Creating and Sizing Tablespaces \(p. 1096\)](#).
3. If the user account into which the data is imported doesn't exist, create the user account and grant the necessary permissions and roles. If you plan to import data into multiple user schemas, create each user account and grant the necessary privileges and roles to it.

For example, the following commands create a new user and grant the necessary permissions and roles to import the data into the user's schema.

```
create user schema_1 identified by <password>;
grant create session, resource to schema_1;
alter user schema_1 quota 100M on users;
```

This example grants the new user the CREATE SESSION privilege and the RESOURCE role. Additional privileges and roles might be required depending on the database objects that you import.

Note

Replace `schema_1` with the name of your schema in this step and in the following steps.

Step 2: Grant Privileges to the User on the Source Database

Use SQL Plus or Oracle SQL Developer to connect to the Oracle instance that contains the data to be imported. If necessary, create a user account and grant the necessary permissions.

Note

If the source database is an Amazon RDS instance, you can skip this step. You use your Amazon RDS master user account to perform the export.

The following commands create a new user and grant the necessary permissions.

```
create user export_user identified by <password>;
grant create session, create table, create database link to export_user;
alter user export_user quota 100M on users;
grant read, write on directory data_pump_dir to export_user;
grant select_catalog_role to export_user;
grant execute on dbms_datapump to export_user;
grant execute on dbms_file_transfer to export_user;
```

Step 3: Use DBMS_DATAPUMP to Create a Dump File

Use SQL Plus or Oracle SQL Developer to connect to the source Oracle instance with an administrative user or with the user you created in step 2. If the source database is an Amazon RDS Oracle DB instance, connect with the Amazon RDS master user. Next, use the Oracle Data Pump utility to create a dump file.

The following script creates a dump file named *sample.dmp* in the DATA_PUMP_DIR directory.

```
DECLARE
hdnl NUMBER;
BEGIN
hdnl := DBMS_DATAPUMP.OPEN( operation => 'EXPORT', job_mode => 'SCHEMA', job_name=>null);
DBMS_DATAPUMP.ADD_FILE( handle => hdnl, filename => 'sample.dmp', directory =>
'DATA_PUMP_DIR', filetype => dbms_datapump.ku$_file_type_dump_file);
DBMS_DATAPUMP.ADD_FILE( handle => hdnl, filename => 'exp.log', directory =>
'DATA_PUMP_DIR', filetype => dbms_datapump.ku$_file_type_log_file);
DBMS_DATAPUMP.METADATA_FILTER(hdnl,'SCHEMA_EXPR','IN (''SCHEMA_1'')');
DBMS_DATAPUMP.START_JOB(hdnl);
END;
/
```

Note

Data Pump jobs are started asynchronously. For information about monitoring a Data Pump job, see [Monitoring Job Status](#) in the Oracle documentation.

Step 4: Create a Database Link to the Target DB Instance

Create a database link between your source instance and your target DB instance. Your local Oracle instance must have network connectivity to the DB instance in order to create a database link and to transfer your export dump file.

Perform this step connected with the same user account as the previous step.

If you are creating a database link between two DB instances inside the same VPC or peered VPCs, the two DB instances should have a valid route between them. The security group of each DB instance must allow ingress to and egress from the other DB instance. The security group inbound and outbound rules can refer to security groups from the same VPC or a peered VPC. For more information, see [Adjusting Database Links for Use with DB Instances in a VPC \(p. 1099\)](#).

The following command creates a database link named *to_rds* that connects to the Amazon RDS master user at the target DB instance.

```
create database link to_rds connect to <master_user_account> identified by <password>
using '(DESCRIPTION=(ADDRESS=(PROTOCOL=TCP)(HOST=<dns or ip address of remote db>)
(PORT=<listener port>))(CONNECT_DATA=(SID=<remote SID>)))';
```

Step 5: Use DBMS_FILE_TRANSFER to Copy the Exported Dump File to the Target DB Instance

Use DBMS_FILE_TRANSFER to copy the dump file from the source database instance to the target DB instance. The following script copies a dump file named sample.dmp from the source instance to a target database link named *to_rds* (created in the previous step).

```
BEGIN
DBMS_FILE_TRANSFER.PUT_FILE(
source_directory_object      => 'DATA_PUMP_DIR',
source_file_name             => 'sample.dmp',
destination_directory_object => 'DATA_PUMP_DIR',
destination_file_name        => 'sample_copied.dmp',
destination_database         => 'to_rds'
);
END;
/
```

Step 6: Use DBMS_DATAPUMP to Import the Data File to the Target DB Instance

Use Oracle Data Pump to import the schema in the DB instance. Additional options such as METADATA_REMAP might be required.

Connect to the DB instance with the Amazon RDS master user account to perform the import.

```
DECLARE
hdnl NUMBER;
BEGIN
hdnl := DBMS_DATAPUMP.OPEN( operation => 'IMPORT', job_mode => 'SCHEMA', job_name=>null);
DBMS_DATAPUMP.ADD_FILE( handle => hdnl, filename => 'sample_copied.dmp', directory =>
'DATA_PUMP_DIR', filetype => dbms_datapump.ku$_file_type_dump_file);
DBMS_DATAPUMP.METADATA_FILTER(hdnl, 'SCHEMA_EXPR', 'IN (''SCHEMA_1'')');
DBMS_DATAPUMP.START_JOB(hdnl);
END;
/
```

You can verify the data import by viewing the user's tables on the DB instance. For example, the following query returns the number of tables for *schema_1*.

```
select count(*) from dba_tables where owner='SCHEMA_1';
```

Step 7: Clean Up

After the data has been imported, you can delete the files that you don't want to keep. You can list the files in the DATA_PUMP_DIR using the following command.

```
select * from table(RDSADMIN.RDS_FILE_UTIL.LISTDIR('DATA_PUMP_DIR')) order by mtme;
```

To delete files in the DATA_PUMP_DIR that you no longer require, use the following command.

```
exec utl_file.fremove('DATA_PUMP_DIR','<file name>');
```

For example, the following command deletes the file named "sample_copied.dmp".

```
exec utl_file.fremove('DATA_PUMP_DIR','sample_copied.dmp');
```

Oracle Export/Import Utilities

The Oracle Export/Import utilities are best suited for migrations where the data size is small and data types such as binary float and double are not required. The import process creates the schema objects so you do not need to run a script to create them beforehand, making this process well suited for databases with small tables. The following example demonstrates how these utilities can be used to export and import specific tables.

To download Oracle export and import utilities, go to <http://www.oracle.com/technetwork/database/enterprise-edition/downloads/index.html>.

Export the tables from the source database using the command below. Substitute username/password as appropriate.

```
exp cust_dba@ORCL FILE=exp_file.dmp TABLES=(tab1,tab2,tab3) LOG=exp_file.log
```

The export process creates a binary dump file that contains both the schema and data for the specified tables. Now this schema and data can be imported into a target database using the command:

```
imp cust_dba@targetdb FROMUSER=cust_schema TOUSER=cust_schema \
TABLES=(tab1,tab2,tab3) FILE=exp_file.dmp LOG=imp_file.log
```

There are other variations of the Export and Import commands that might be better suited to your needs. See Oracle's documentation for full details.

Oracle SQL*Loader

Oracle SQL*Loader is well suited for large databases that have a limited number of objects in them. Since the process involved in exporting from a source database and loading to a target database is very specific to the schema, the following example creates the sample schema objects, exports from a source, and then loads it into a target database.

To download Oracle SQL*Loader, go to <http://www.oracle.com/technetwork/database/enterprise-edition/downloads/index.html>.

1. Create a sample source table using the command below.

```
create table customer_0 tablespace users as select rownum id, o.* from
all_objects o, all_objects x where rownum <= 1000000;
```

2. On the target Amazon RDS instance, create a destination table that is used to load the data.

```
create table customer_1 tablespace users as select 0 as id, owner,
object_name, created from all_objects where 1=2;
```

3. The data is exported from the source database to a flat file with delimiters. This example uses SQL*Plus for this purpose. For your data, you will likely need to generate a script that does the export for all the objects in the database.

```
alter session set nls_date_format = 'YYYY/MM/DD HH24:MI:SS'; set linesize 800
HEADING OFF FEEDBACK OFF array 5000 pagesize 0 spool customer_0.out SET
```

```
MARKUP HTML PREFORMAT ON SET COLSEP ',' SELECT id, owner, object_name,  
created FROM customer_0; spool off
```

4. You need to create a control file to describe the data. Again, depending on your data, you will need to build a script that does this step.

```
cat << EOF > sqlldr_1.ctl  
load data  
infile customer_0.out  
into table customer_1  
APPEND  
fields terminated by "," optionally enclosed by ''''  
(  
id POSITION(01:10) INTEGER EXTERNAL,  
owner POSITION(12:41) CHAR,  
object_name POSITION(43:72) CHAR,  
created POSITION(74:92) date "YYYY/MM/DD HH24:MI:SS"  
)
```

If needed, copy the files generated by the preceding code to a staging area, such as an Amazon EC2 instance.

5. Finally, import the data using SQL*Loader with the appropriate username and password for the target database.

```
sqlldr cust_dba@targetdb control=sqlldr_1.ctl BINDSIZE=10485760 READSIZE=10485760  
ROWS=1000
```

Oracle Materialized Views

You can also make use of Oracle materialized view replication to migrate large datasets efficiently. Replication allows you to keep the target tables in sync with the source on an ongoing basis, so the actual cutover to Amazon RDS can be done later, if needed. The replication is set up using a database link from the Amazon RDS instance to the source database.

One requirement for materialized views is to allow access from the target database to the source database. In the following example, access rules were enabled on the source database to allow the Amazon RDS target database to connect to the source over SQLNet.

1. Create a user account on both source and Amazon RDS target instances that can authenticate with the same password.

```
create user dblink_user identified by <password>  
default tablespace users  
temporary tablespace temp;  
  
grant create session to dblink_user;  
  
grant select any table to dblink_user;  
  
grant select any dictionary to dblink_user;
```

2. Create a database link from the Amazon RDS target instance to the source instance using the newly created dblink_user.

```
create database link remote_site
```

```
connect to dblink_user identified by <password>
using '(description=(address=(protocol=tcp) (host=<myhost>)
(port=<listener port>)) (connect_data=(sid=<sourcedb sid>)))';
```

3. Test the link:

```
select * from v$instance@remote_site;
```

4. Create a sample table with primary key and materialized view log on the source instance.

```
create table customer_0 tablespace users as select rownum id, o.* from
all_objects o, all_objects x where rownum <= 1000000;

alter table customer_0 add constraint pk_customer_0 primary key (id) using index;

create materialized view log on customer_0;
```

5. On the target Amazon RDS instance, create a materialized view.

```
CREATE MATERIALIZED VIEW customer_0 BUILD IMMEDIATE REFRESH FAST AS
SELECT * FROM cust_dba.customer_0@remote_site;
```

6. On the target Amazon RDS instance, refresh the materialized view.

```
exec DBMS_MV.REFRESH('CUSTOMER_0', 'f');
```

7. Drop the materialized view and include the PRESERVE TABLE clause to retain the materialized view container table and its contents.

```
DROP MATERIALIZED VIEW customer_0 PRESERVE TABLE;
```

The retained table has the same name as the dropped materialized view.

Working with Oracle Read Replicas for Amazon RDS

You usually use Read Replicas to configure replication between Amazon RDS DB instances. For general information about Read Replicas, see [Working with Read Replicas \(p. 182\)](#).

In this section, you can find specific information about working with Read Replicas on Amazon RDS for Oracle.

Topics

- [Configuring Read Replicas for Oracle \(p. 998\)](#)
- [Read Replica Limitations with Oracle \(p. 999\)](#)
- [Troubleshooting an Oracle Read Replica Problem \(p. 1000\)](#)

Configuring Read Replicas for Oracle

Oracle Read Replicas use Oracle Data Guard to replicate database changes from the source DB instance to its Read Replicas. For information about Oracle Data Guard, see [Oracle Data Guard Concepts and Administration](#) in the Oracle documentation.

Before a DB instance can serve as a source DB instance, you must enable automatic backups on the source DB instance. To do so, you set the backup retention period to a value other than 0. We recommend that you enable force logging mode on the DB instance. To enable force logging mode, connect to the DB instance and run the following procedure.

```
exec rdsadmin.rdsadmin_util.force_logging(p_enable => true);
```

For more information about this procedure, see [Setting Force Logging \(p. 1106\)](#).

If you plan changes to your logging configuration, we recommend that you complete the changes before making a DB instance the source for one or more Read Replicas. You can modify the logging configuration for a DB instance by using the Amazon RDS procedures `rdsadmin.rdsadmin_util.add_logfile` and `rdsadmin.rdsadmin_util.drop_logfile`. For more information, see [Adding Online Redo Logs \(p. 1107\)](#) and [Dropping Online Redo Logs \(p. 1107\)](#). As a best practice, we recommend that you not modify the logging configuration after Read Replicas are created. Modifications can cause the online redo logging configuration to get out of sync with the standby logging configuration.

Creating an Oracle Read Replica doesn't require an outage for the master DB instance. Amazon RDS sets the necessary parameters and permissions for the source DB instance and the Read Replica without any service interruption. A snapshot is taken of the source DB instance, and this snapshot becomes the Read Replica. No outage occurs when you delete a Read Replica.

You can create up to five Read Replicas from one source DB instance. Before you create a Read Replica, make sure that the setting of the `max_string_size` parameter is same on the source DB instance and the Read Replica. You can do this by associating them with the same parameter group. If you have different parameters groups for the source and the Read Replica, you can do this by setting `max_string_size` to the same value.

The Oracle DB engine version of the source DB instance and all of its Read Replicas must be the same. Amazon RDS upgrades the Read Replicas immediately after upgrading the source DB instance, regardless of a Read Replica's maintenance window. For major version upgrades of cross-region Read Replicas,

Amazon RDS automatically generates an option group for the target version, copies all options and option settings from the original option group to the new option group, and associates the upgraded cross-region Read Replica with the new option group. For more information about upgrading the DB engine version, see [Upgrading the Oracle DB Engine \(p. 979\)](#).

For a Read Replica to receive and apply changes from the source, it should have sufficient compute and storage resources. If a Read Replica reaches compute, network, or storage resource capacity, the Read Replica stops receiving or applying changes from its source. You can modify the storage and CPU resources of a Read Replica independently from its source and other Read Replicas.

Read Replica Limitations with Oracle

The following are limitations for Oracle Read Replicas:

- You must have an Active Data Guard license.
- Oracle Read Replicas are only available on the Oracle Enterprise Edition (EE) engine.
- Oracle Read Replicas are available for Oracle version 12.1.0.2.v10 and higher 12.1 versions, for all 12.2 versions, for all 18 versions, and for all 19 versions.
- Oracle Read Replicas are only available for DB instances on the EC2-VPC platform.
- Oracle Read Replicas are only available for DB instances running on DB instance classes with two or more vCPUs. A source DB instance can't use the db.t3.micro instance class.
- Amazon RDS for Oracle doesn't intervene to mitigate high replica lag between a source DB instance and its Read Replicas. Ensure that the source DB instance and its Read Replicas are sized properly, in terms of compute and storage, to suit their operational load.
- An Oracle Read Replica that is in the same AWS Region as its source DB instance must belong to the same option group as the source DB instance. Modifications to the source option group or source option group membership propagate to Read Replicas. These changes are applied to the Read Replicas immediately after they are applied to the source DB instance, regardless of the Read Replica's maintenance window.

For more information about option groups, see [Working with Option Groups \(p. 196\)](#).

- If a DB instance is a source for one or more cross-region Read Replicas, it retains its redo logs until they are applied on all cross-region Read Replicas. The redo logs might result in increased storage consumption.
- When you create an Oracle cross-region Read Replica, Amazon RDS creates a dedicated option group for it.

You can't remove an Oracle cross-region Read Replica from its dedicated option group. No other DB instances can use the dedicated option group for an Oracle cross-region Read Replica.

Only the following nonreplicated options can be added to or removed from a dedicated option group:

- NATIVE_NETWORK_ENCRYPTION
- OEM
- OEM_AGENT
- SSL

To add other options to an Oracle cross-region Read Replica, add them to the source DB instance's option group. The option is also installed on all of the source DB instance's Read Replicas. For licensed options, you must ensure that there are sufficient licenses for the Read Replicas.

When you promote an Oracle cross-region Read Replica, the promoted Read Replica behaves the same as other Oracle DB instances, including the management of its options. You can promote a Read Replica explicitly or implicitly by deleting its source DB instance.

For more information about option groups, see [Working with Option Groups \(p. 196\)](#).

Troubleshooting an Oracle Read Replica Problem

You can monitor replication lag in Amazon CloudWatch by viewing the Amazon RDS `ReplicaLag` metric. For information about replication lag time, see [Monitoring Read Replication \(p. 194\)](#).

If replication lag is too long, you can query the following views for information about the lag:

- `V$ARCHIVED_LOG` – Shows which commits have been applied to the Read Replica.
- `V$DATAGUARD_STATS` – Shows a detailed breakdown of the components that make up the `replicaLag` metric.
- `V$DATAGUARD_STATUS` – Shows the log output from Oracle's internal replication processes.

Oracle Character Sets Supported in Amazon RDS

The following table lists the Oracle database character sets that are supported in Amazon RDS. You can use a value from this table with the `--character-set-name` parameter of the AWS CLI [create-db-instance](#) command or with the `CharacterSetName` parameter of the Amazon RDS API [CreateDBInstance](#) operation.

Setting the `NLS_LANG` environment parameter in your client's environment is the simplest way to specify locale behavior for Oracle. This parameter sets the language and territory used by the client application and the database server. It also indicates the client's character set, which corresponds to the character set for data entered or displayed by a client application. Amazon RDS lets you set the character set when you create a DB instance. For more information on the `NLS_LANG` and character sets, see [What is a Character set or Code Page? in the Oracle documentation](#).

Value	Description
AL32UTF8	Unicode 5.0 UTF-8 Universal character set (default)
AR8ISO8859P6	ISO 8859-6 Latin/Arabic
AR8MSWIN1256	Microsoft Windows Code Page 1256 8-bit Latin/Arabic
BLT8ISO8859P13	ISO 8859-13 Baltic
BLT8MSWIN1257	Microsoft Windows Code Page 1257 8-bit Baltic
CL8ISO8859P5	ISO 8859-5 Latin/Cyrillic
CL8MSWIN1251	Microsoft Windows Code Page 1251 8-bit Latin/Cyrillic
EE8ISO8859P2	ISO 8859-2 East European
EL8ISO8859P7	ISO 8859-7 Latin/Greek
EE8MSWIN1250	Microsoft Windows Code Page 1250 8-bit East European
EL8MSWIN1253	Microsoft Windows Code Page 1253 8-bit Latin/Greek
IW8ISO8859P8	ISO 8859-8 Latin/Hebrew
IW8MSWIN1255	Microsoft Windows Code Page 1255 8-bit Latin/Hebrew
JA16EUC	EUC 24-bit Japanese
JA16EUCTILDE	Same as JA16EUC except for mapping of wave dash and tilde to and from Unicode
JA16SJIS	Shift-JIS 16-bit Japanese
JA16SJISTILDE	Same as JA16SJIS except for mapping of wave dash and tilde to and from Unicode
KO16MSWIN949	Microsoft Windows Code Page 949 Korean

Value	Description
NE8ISO8859P10	ISO 8859-10 North European
NEE8ISO8859P4	ISO 8859-4 North and Northeast European
TH8TISASCII	Thai Industrial Standard 620-2533-ASCII 8-bit
TR8MSWIN1254	Microsoft Windows Code Page 1254 8-bit Turkish
US7ASCII	ASCII 7-bit American
UTF8	Unicode 3.0 UTF-8 Universal character set, CESU-8 compliant
VN8MSWIN1258	Microsoft Windows Code Page 1258 8-bit Vietnamese
WE8ISO8859P1	Western European 8-bit ISO 8859 Part 1
WE8ISO8859P15	ISO 8859-15 West European
WE8ISO8859P9	ISO 8859-9 West European and Turkish
WE8MSWIN1252	Microsoft Windows Code Page 1252 8-bit West European
ZHS16GBK	GBK 16-bit Simplified Chinese
ZHT16HKSCS	Microsoft Windows Code Page 950 with Hong Kong Supplementary Character Set HKSCS-2001. Character set conversion is based on Unicode 3.0.
ZHT16MSWIN950	Microsoft Windows Code Page 950 Traditional Chinese
ZHT32EUC	EUC 32-bit Traditional Chinese

You can also set the following National Language Support (NLS) initialization parameters at the instance level for an Oracle DB instance in Amazon RDS:

- NLS_DATE_FORMAT
- NLS_LENGTH_SEMANTICS
- NLS_NCHAR_CONV_EXCP
- NLS_TIME_FORMAT
- NLS_TIME_TZ_FORMAT
- NLS_TIMESTAMP_FORMAT
- NLS_TIMESTAMP_TZ_FORMAT

For information about modifying instance parameters, see [Working with DB Parameter Groups \(p. 211\)](#).

You can set other NLS initialization parameters in your SQL client. For example, the following statement sets the NLS_LANGUAGE initialization parameter to GERMAN in a SQL client that is connected to an Oracle DB instance:

```
ALTER SESSION SET NLS_LANGUAGE=GERMAN;
```

For information about connecting to an Oracle DB instance with a SQL client, see [Connecting to a DB Instance Running the Oracle Database Engine \(p. 955\)](#).

Options for Oracle DB Instances

Following, you can find a description of options, or additional features, that are available for Amazon RDS instances running the Oracle DB engine. To enable these options, you add them to an option group, and then associate the option group with your DB instance. For more information, see [Working with Option Groups \(p. 196\)](#).

Some options require additional memory to run on your DB instance. For example, Oracle Enterprise Manager Database Control uses about 300 MB of RAM. If you enable this option for a small DB instance, you might encounter performance problems due to memory constraints. You can adjust the Oracle parameters so that the database requires less RAM. Alternatively, you can scale up to a larger DB instance.

Amazon RDS supports the following options for Oracle DB instances.

Option	Option ID
Amazon S3 Integration (p. 1005)	S3_INTEGRATION
Oracle Application Express (p. 1016)	APEX
	APEX-DEV
Oracle Enterprise Manager (p. 1024)	OEM
	OEM_AGENT
Oracle Java Virtual Machine (p. 1037)	JVM
Oracle Label Security (p. 1040)	OLS
Oracle Locator (p. 1043)	LOCATOR
Oracle Multimedia (p. 1046)	MULTIMEDIA
Oracle Native Network Encryption (p. 1049)	NATIVE_NETWORK_ENCRYPTION
Oracle Secure Sockets Layer (p. 1051)	SSL
Oracle Spatial (p. 1060)	SPATIAL
Oracle SQLT (p. 1063)	SQLT
Oracle Statspack (p. 1068)	STATSPACK
Oracle Time Zone (p. 1071)	Timezone
Oracle Transparent Data Encryption (p. 1074)	TDE
Oracle UTL_MAIL (p. 1076)	UTL_MAIL
Oracle XML DB (p. 1078)	XMLDB

Amazon S3 Integration

You can transfer files between an Amazon RDS for Oracle DB instance and an Amazon S3 bucket. You can use Amazon S3 integration with Oracle features such as Data Pump. For example, you can download Data Pump files from Amazon S3 to the DB instance host.

Note

The DB instance and the Amazon S3 bucket must be in the same AWS Region.

Topics

- [Prerequisites for Amazon RDS Oracle Integration with Amazon S3 \(p. 1005\)](#)
- [Adding the Amazon S3 Integration Option \(p. 1010\)](#)
- [Transferring Files Between Amazon RDS for Oracle and an Amazon S3 Bucket \(p. 1011\)](#)
- [Removing the Amazon S3 Integration Option \(p. 1015\)](#)

Prerequisites for Amazon RDS Oracle Integration with Amazon S3

To work with Amazon RDS for Oracle integration with Amazon S3, the Amazon RDS DB instance must have access to an Amazon S3 bucket. For this, you create an AWS Identity and Access Management (IAM) policy and an IAM role. The Amazon VPC used by your DB instance doesn't need to provide access to the Amazon S3 endpoints.

Note

To add a role to a DB instance, the status of the DB instance must be available.

Console

To create an IAM policy to allow Amazon RDS access to an Amazon S3 bucket

1. Open the [IAM Management Console](#).
2. In the navigation pane, choose **Policies**.
3. Choose **Create policy**.
4. On the **Visual editor** tab, choose **Choose a service**, and then choose **S3**.
5. For **Actions**, choose **Expand all**, and then choose the bucket permissions and object permissions needed for the IAM policy.

Include the appropriate actions in the policy based on the type of access required:

- **GetObject** – Required to transfer files from an Amazon S3 bucket to Amazon RDS.
- **ListBucket** – Required to transfer files from an Amazon S3 bucket to Amazon RDS.
- **PutObject** – Required to transfer files from Amazon RDS to an Amazon S3 bucket.

Object permissions are permissions for object operations in Amazon S3, and need to be granted for objects in a bucket, not the bucket itself. For more information about permissions for object operations in Amazon S3, see [Permissions for Object Operations](#).

6. Choose **Resources**, and choose **Add ARN for bucket**.
7. In the **Add ARN(s)** dialog box, provide the details about your resource, and choose **Add**.

Specify the Amazon S3 bucket to allow access to. For instance, if you want to allow Amazon RDS to access the Amazon S3 bucket named `example-bucket`, then set the Amazon Resource Name (ARN) value to `arn:aws:s3:::example-bucket`.

8. If the **object** resource is listed, choose **Add ARN for object**.
9. In the **Add ARN(s)** dialog box, provide the details about your resource.

For the Amazon S3 bucket, specify the Amazon S3 bucket to allow access to. For the object, you can choose **Any** to grant permissions to any object in the bucket.

Note

You can set **Amazon Resource Name (ARN)** to a more specific ARN value to allow Amazon RDS to access only specific files or folders in an Amazon S3 bucket. For more information about how to define an access policy for Amazon S3, see [Managing Access Permissions to Your Amazon S3 Resources](#).

10. (Optional) Choose **Add additional permissions** to add another Amazon S3 bucket to the policy, and repeat the previous steps for the bucket.

Note

You can repeat this to add corresponding bucket permission statements to your policy for each Amazon S3 bucket that you want Amazon RDS to access. Optionally, you can also grant access to all buckets and objects in Amazon S3.

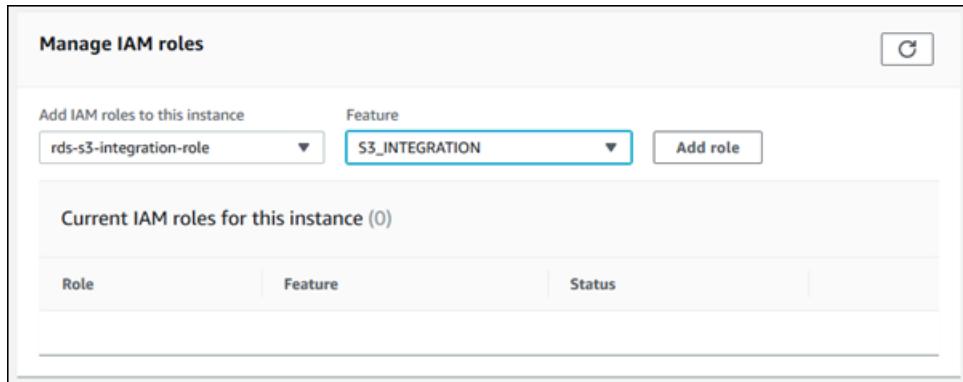
11. Choose **Review policy**.
12. For **Name**, enter a name for your IAM policy, for example `rds-s3-integration-policy`. You use this name when you create an IAM role to associate with your DB instance. You can also add an optional **Description** value.
13. Choose **Create policy**.

To create an IAM role to allow Amazon RDS access to an Amazon S3 bucket

1. In the navigation pane, choose **Roles**.
2. Choose **Create role**.
3. For **AWS service**, choose **RDS**.
4. For **Select your use case**, choose **RDS – Add Role to Database**.
5. Choose **Next: Permissions**.
6. For **Search under Attach permissions policies**, enter the name of the IAM policy you created, and choose the policy when it appears in the list.
7. Choose **Next: Tags** and then **Next: Review**.
8. Set **Role name** to a name for your IAM role, for example `rds-s3-integration-role`. You can also add an optional **Role description** value.
9. Choose **Create Role**.

To associate your IAM role with your DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. Choose the Oracle DB instance name to display its details.
3. On the **Connectivity & security** tab, in the **Manage IAM roles** section, choose the role to add under **Add IAM roles to this instance**.
4. For **Feature**, choose **S3_INTEGRATION**.



5. Choose **Add role**.

AWS CLI

To grant Amazon RDS access to an Amazon S3 bucket

1. Create an AWS Identity and Access Management (IAM) policy that grants Amazon RDS access to an Amazon S3 bucket.

Include the appropriate actions in the policy based on the type of access required:

- `GetObject` – Required to transfer files from an Amazon S3 bucket to Amazon RDS.
- `ListBucket` – Required to transfer files from an Amazon S3 bucket to Amazon RDS.
- `PutObject` – Required to transfer files from Amazon RDS to an Amazon S3 bucket.

The following AWS CLI command creates an IAM policy named `rds-s3-integration-policy` with these options. It grants access to a bucket named `your-s3-bucket-arn`.

Example

For Linux, OS X, or Unix:

```
aws iam create-policy \
--policy-name rds-s3-integration-policy \
--policy-document '{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "s3integration",
            "Action": [
                "s3:GetObject",
                "s3>ListBucket",
                "s3:PutObject"
            ],
            "Effect": "Allow",
            "Resource": [
                "arn:aws:s3:::your-s3-bucket-arn",
                "arn:aws:s3:::your-s3-bucket-arn/*"
            ]
        }
    ]
}'
```

For Windows:

```
aws iam create-policy ^
--policy-name rds-s3-integration-policy ^
--policy-document '{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "s3integration",
            "Action": [
                "s3:GetObject",
                "s3>ListBucket",
                "s3:PutObject"
            ],
            "Effect": "Allow",
            "Resource": [
                "arn:aws:s3:::your-s3-bucket-arn",
                "arn:aws:s3:::your-s3-bucket-arn/*"
            ]
        }
    ]
}'
```

2. After the policy is created, note the Amazon Resource Name (ARN) of the policy. You need the ARN for a subsequent step.
3. Create an IAM role that Amazon RDS can assume on your behalf to access your Amazon S3 buckets.

The following AWS CLI command creates the [rds-s3-integration-role](#) for this purpose.

Example

For Linux, OS X, or Unix:

```
aws iam create-role \
--role-name rds-s3-integration-role \
--assume-role-policy-document '{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Principal": {
                "Service": "rds.amazonaws.com"
            },
            "Action": "sts:AssumeRole"
        }
    ]
}'
```

For Windows:

```
aws iam create-role ^
--role-name rds-s3-integration-role ^
--assume-role-policy-document '{
    "Version": "2012-10-17",
    "Statement": [
        {
```

```
        "Effect": "Allow",
        "Principal": {
            "Service": "rds.amazonaws.com"
        },
        "Action": "sts:AssumeRole"
    }
}'
```

For more information, see [Creating a Role to Delegate Permissions to an IAM User](#) in the *IAM User Guide*.

4. After the role is created, note the ARN of the role. You need the ARN for a subsequent step.
5. Attach the policy you created to the role you created.

The following AWS CLI command attaches the policy to the role named `rds-s3-integration-role`.

Example

For Linux, OS X, or Unix:

```
aws iam attach-role-policy \
--policy-arn your-policy-arn \
--role-name rds-s3-integration-role
```

For Windows:

```
aws iam attach-role-policy ^
--policy-arn your-policy-arn ^
--role-name rds-s3-integration-role
```

Replace `your-policy-arn` with the policy ARN that you noted in a previous step.

6. Add the role to the Oracle DB instance.

The following AWS CLI command adds the role to an Oracle DB instance named `mydbinstance`.

Example

For Linux, OS X, or Unix:

```
aws rds add-role-to-db-instance \
--db-instance-identifier mydbinstance \
--feature-name S3_INTEGRATION \
--role-arn your-role-arn
```

For Windows:

```
aws rds add-role-to-db-instance ^
--db-instance-identifier mydbinstance ^
--feature-name S3_INTEGRATION ^
```

```
--role-arn your-role-arn
```

Replace *your-role-arn* with the role ARN that you noted in a previous step. S3_INTEGRATION must be specified for the --feature-name option.

Adding the Amazon S3 Integration Option

To use Amazon RDS for Oracle Integration with Amazon S3, your Amazon RDS Oracle DB instance must be associated with an option group that includes the S3_INTEGRATION option.

Console

To configure an option group for Amazon S3 integration

1. Create a new option group or identify an existing option group to which you can add the S3_INTEGRATION option.

For information about creating an option group, see [Creating an Option Group \(p. 197\)](#).

2. Add the S3_INTEGRATION option to the option group.

For information about adding an option to an option group, see [Adding an Option to an Option Group \(p. 200\)](#).

3. Create a new Oracle DB instance and associate the option group with it, or modify an Oracle DB instance to associate the option group with it.

For information about creating an Oracle DB instance, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).

For information about modifying an Oracle DB instance, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

AWS CLI

To configure an option group for Amazon S3 integration

1. Create a new option group or identify an existing option group to which you can add the S3_INTEGRATION option.

For information about creating an option group, see [Creating an Option Group \(p. 197\)](#).

2. Add the S3_INTEGRATION option to the option group.

For example, the following AWS CLI command adds the S3_INTEGRATION option to an option group named **myoptiongroup**.

Example

For Linux, OS X, or Unix:

```
aws rds add-option-to-option-group \
--option-group-name myoptiongroup \
--options OptionName=S3_INTEGRATION,OptionVersion=1.0
```

For Windows:

```
aws rds add-option-to-option-group ^
--option-group-name myoptiongroup ^
--options OptionName=S3_INTEGRATION,OptionVersion=1.0
```

3. Create a new Oracle DB instance and associate the option group with it, or modify an Oracle DB instance to associate the option group with it.

For information about creating an Oracle DB instance, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).

For information about modifying an Oracle DB instance, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Transferring Files Between Amazon RDS for Oracle and an Amazon S3 Bucket

You can use Amazon RDS procedures to upload files from an Oracle DB instance to an Amazon S3 bucket. You can also use Amazon RDS procedures to download files from an Amazon S3 bucket to an Oracle DB instance.

Note

These procedures upload or download the files in a single directory. You can't include subdirectories in these operations. The upload and download operations are processed serially.

Topics

- [Uploading Files from an Oracle DB Instance to an Amazon S3 Bucket \(p. 1011\)](#)
- [Downloading Files from an Amazon S3 Bucket to an Oracle DB Instance \(p. 1013\)](#)
- [Monitoring the Status of a File Transfer \(p. 1015\)](#)

Uploading Files from an Oracle DB Instance to an Amazon S3 Bucket

To upload files from an Oracle DB instance to an Amazon S3 bucket, use the Amazon RDS procedure `rdsadmin.rdsadmin_s3_tasks.upload_to_s3`. For example, you can upload Oracle Recovery Manager (RMAN) backup files. For more information about performing RMAN backups, see [Common DBA Recovery Manager \(RMAN\) Tasks for Oracle DB Instances \(p. 1112\)](#).

The `rdsadmin.rdsadmin_s3_tasks.upload_to_s3` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>p_bucket_name</code>	VARCHAR2	–	required	The name of the Amazon S3 bucket to upload files to.
<code>p_directory_name</code>	VARCHAR2	–	required	The name of the Oracle directory object to upload files from. The directory can be any user-created directory object or the Data Pump directory, such as <code>DATA_PUMP_DIR</code> .

Parameter Name	Data Type	Default	Required	Description
				Note You can only upload files from the specified directory. You can't upload files in subdirectories in the specified directory.
p_s3_prefix	VARCHAR2	–	required	An Amazon S3 file name prefix that files are uploaded to. An empty prefix uploads all files to the top level in the specified Amazon S3 bucket and doesn't add a prefix to the file names. For example, if the prefix is <code>folder_1/oradb</code> , files are uploaded to <code>folder_1</code> . In this case, the <code>oradb</code> prefix is added to each file.
p_prefix	VARCHAR2	–	required	A file name prefix that file names must match to be uploaded. An empty prefix uploads all files in the specified directory.

The following example uploads all of the files in the `DATA_PUMP_DIR` directory to the Amazon S3 bucket named `mys3bucket`.

```

SELECT rdsadmin.rdsadmin_s3_tasks.upload_to_s3(
    p_bucket_name => 'mys3bucket',
    p_prefix      => '',
    p_s3_prefix   => '',
    p_directory_name => 'DATA_PUMP_DIR')
AS TASK_ID FROM DUAL;
  
```

The following example uploads all of the files with the prefix `db` in the `DATA_PUMP_DIR` directory to the Amazon S3 bucket named `mys3bucket`.

```

SELECT rdsadmin.rdsadmin_s3_tasks.upload_to_s3(
    p_bucket_name => 'mys3bucket',
    p_prefix      => 'db',
    p_s3_prefix   => '',
    p_directory_name => 'DATA_PUMP_DIR')
AS TASK_ID FROM DUAL;
  
```

The following example uploads all of the files in the `DATA_PUMP_DIR` directory to the Amazon S3 bucket named `mys3bucket`. The files are uploaded to a `dbfiles` folder.

```
SELECT rdsadmin.rdsadmin_s3_tasks.upload_to_s3(
    p_bucket_name      => 'mys3bucket',
    p_prefix           => '',
    p_s3_prefix        => 'dbfiles/',
    p_directory_name   => 'DATA_PUMP_DIR')
AS TASK_ID FROM DUAL;
```

The following example uploads all of the files in the `DATA_PUMP_DIR` directory to the Amazon S3 bucket named `mys3bucket`. The files are uploaded to a `dbfiles` folder and `ora` is added to the beginning of each file name.

```
SELECT rdsadmin.rdsadmin_s3_tasks.upload_to_s3(
    p_bucket_name      => 'mys3bucket',
    p_prefix           => '',
    p_s3_prefix        => 'dbfiles/ora',
    p_directory_name   => 'DATA_PUMP_DIR')
AS TASK_ID FROM DUAL;
```

In each example, the `SELECT` statement returns the ID of the task in a `VARCHAR2` data type.

Downloading Files from an Amazon S3 Bucket to an Oracle DB Instance

To download files from an Amazon S3 bucket to an Oracle DB instance, use the Amazon RDS procedure `rdsadmin.rdsadmin_s3_tasks.download_from_s3`. The `rdsadmin.rdsadmin_s3_tasks.download_from_s3` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>p_bucket_name</code>	<code>VARCHAR2</code>	–	required	The name of the Amazon S3 bucket to download files from.
<code>p_directory_name</code>	<code>VARCHAR2</code>	–	required	The name of the Oracle directory object to download files to. The directory can be any user-created directory object or the Data Pump directory, such as <code>DATA_PUMP_DIR</code> .
<code>p_s3_prefix</code>	<code>VARCHAR2</code>	“”	optional	A file name prefix that file names must match to be downloaded. An empty prefix downloads all of the top level files in the specified Amazon S3 bucket, but not the files in folders in the bucket. The procedure downloads Amazon S3 objects only

Parameter Name	Data Type	Default	Required	Description
				<p>from the first level folder that matches the prefix. Nested directory structures matching the specified prefix are not downloaded.</p> <p>For example, suppose that an Amazon S3 bucket has the folder structure <code>folder_1/folder_2/folder_3</code>. Suppose also that you specify the <code>'folder_1/folder_2/'</code> prefix. In this case, only the files in <code>folder_2</code> are downloaded, not the files in <code>folder_1</code> or <code>folder_3</code>.</p> <p>If, instead, you specify the <code>'folder_1/folder_2'</code> prefix, all files in <code>folder_1</code> that match the <code>'folder_2'</code> prefix are downloaded, and no files in <code>folder_2</code> are downloaded.</p>

The following example downloads all of the files in the Amazon S3 bucket named `mys3bucket` to the `DATA_PUMP_DIR` directory.

```
SELECT rdsadmin.rdsadmin_s3_tasks.download_from_s3(
    p_bucket_name => 'mys3bucket',
    p_directory_name => 'DATA_PUMP_DIR')
AS TASK_ID FROM DUAL;
```

The following example downloads all of the files with the prefix `db` in the Amazon S3 bucket named `mys3bucket` to the `DATA_PUMP_DIR` directory.

```
SELECT rdsadmin.rdsadmin_s3_tasks.download_from_s3(
    p_bucket_name => 'mys3bucket',
    p_s3_prefix => 'db',
    p_directory_name => 'DATA_PUMP_DIR')
AS TASK_ID FROM DUAL;
```

In each example, the SELECT statement returns the ID of the task in a `VARCHAR2` data type.

The following example downloads all of the files in the folder `myfolder/` in the Amazon S3 bucket named `mys3bucket` to the `DATA_PUMP_DIR` directory. Use the prefix parameter setting to specify the Amazon S3 folder.

```
SELECT rdsadmin.rdsadmin_s3_tasks.download_from_s3(
```

```
p_bucket_name      => 'mys3bucket',
p_s3_prefix        => 'myfolder/',
p_directory_name  => 'DATA_PUMP_DIR')
AS TASK_ID FROM DUAL;
```

In each example, the `SELECT` statement returns the ID of the task in a `VARCHAR2` data type.

Note

You can use the `UTL_FILE.FREMOVE` Oracle procedure to remove files from a directory. For more information, see [FREMOVE Procedure](#) in the Oracle documentation.

Monitoring the Status of a File Transfer

File transfer tasks publish Amazon RDS events when they start and when they complete. For information about viewing events, see [Viewing Amazon RDS Events \(p. 388\)](#).

You can view the status of an ongoing task in a bdump file. The bdump files are located in the `/rdsdbdata/log/trace` directory. Each bdump file name is in the following format.

```
dbtask-task-id.log
```

Replace `task-id` with the ID of the task that you want to monitor.

You can use the `rdsadmin.rds_file_util.read_text_file` stored procedure to view the contents of bdump files. For example, the following query returns the contents of the `dbtask-1546988886389-2444.log` bdump file.

```
SELECT text FROM
  table(rdsadmin.rds_file_util.read_text_file('BDUMP','dbtask-1546988886389-2444.log'));
```

Removing the Amazon S3 Integration Option

You can remove Amazon S3 integration option from a DB instance.

To remove the Amazon S3 integration option from a DB instance, do one of the following:

- To remove the Amazon S3 integration option from multiple DB instances, remove the `S3_INTEGRATION` option from the option group to which the DB instances belong. This change affects all DB instances that use the option group. For more information, see [Removing an Option from an Option Group \(p. 207\)](#).
- To remove the Amazon S3 integration option from a single DB instance, modify the DB instance and specify a different option group that doesn't include the `S3_INTEGRATION` option. You can specify the default (empty) option group or a different custom option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Oracle Application Express

Amazon RDS supports Oracle Application Express (APEX) through the use of the `APEX` and `APEX-DEV` options. Oracle APEX can be deployed as a run-time environment or as a full development environment for web-based applications. Using Oracle APEX, developers can build applications entirely within the web browser. For more information, see [Oracle Application Express](#) in the Oracle documentation.

Oracle APEX consists of two main components:

- A *repository* that stores the metadata for APEX applications and components. The repository consists of tables, indexes, and other objects that are installed in your Amazon RDS DB instance.
- A *listener* that manages HTTP communications with Oracle APEX clients. The listener accepts incoming connections from web browsers, forwards them to the Amazon RDS DB instance for processing, and then sends results from the repository back to the browsers. The APEX Listener was renamed Oracle Rest Data Services (ORDS) in Oracle 12c.

When you add the Amazon RDS APEX options to your DB instance, Amazon RDS installs the Oracle APEX repository only. Install the Oracle APEX Listener on a separate host, such as an Amazon EC2 instance, an on-premises server at your company, or your desktop computer.

The APEX option uses storage on the DB instance class for your DB instance. Following are the supported versions and approximate storage requirements for Oracle APEX.

APEX Version	Storage Requirements	Oracle 19c	Oracle 18c	Oracle 12c version 12.2	Oracle 12c version 12.1	Oracle 11g
Oracle APEX version 19.1.v1	148 MiB	Supported	Supported	Supported	Supported	Supported
Oracle APEX version 18.2.v1	146 MiB	Not supported	Supported	Supported	Supported	Supported
Oracle APEX version 18.1.v1	145 MiB	Not supported	Supported	Supported	Supported	Supported
Oracle APEX version 5.1.4.v1	220 MiB	Not supported	Supported	Supported	Supported	Supported
Oracle APEX version 5.1.2.v1	150 MiB	Not supported	Not supported	Not supported	Supported	Supported
Oracle APEX version 5.0.4.v1	140 MiB	Not supported	Not supported	Not supported	Supported	Supported
Oracle APEX version 4.2.6.v1	160 MiB	Not supported	Not supported	Not supported	Supported	Supported
Oracle APEX version 4.1.1.v1	130 MiB	Not supported	Not supported	Not supported	Not supported	Supported

Note

Oracle APEX 5 for Oracle 11g isn't supported when the DB instance class used by the DB instance has only one vCPU. For information about DB instance classes, see [Choosing the DB Instance Class \(p. 107\)](#).

Prerequisites for Oracle APEX and APEX Listener

The following are prerequisites for using Oracle APEX and APEX Listener:

- Have SQL*Plus on your DB instance to perform administrative tasks.
- Have the following software installed on the host computer that acts as the Oracle APEX Listener:
 - The Java Runtime Environment (JRE).
 - Oracle Net Services, to enable the Oracle APEX Listener to connect to your Amazon RDS instance.

Adding the Amazon RDS APEX Options

The general process for adding the Amazon RDS APEX options to a DB instance is the following:

1. Create a new option group, or copy or modify an existing option group.
2. Add the options to the option group.
3. Associate the option group with the DB instance.

When you add the Amazon RDS APEX options, a brief outage occurs while your DB instance is automatically restarted.

To add the APEX options to a DB instance

1. Determine the option group you want to use. You can create a new option group or use an existing option group. If you want to use an existing option group, skip to the next step. Otherwise, create a custom DB option group with the following settings:
 - a. For **Engine**, choose the Oracle edition that you want to use. The APEX options are supported on all editions.
 - b. For **Major engine version**, choose the version of your DB instance.

For more information, see [Creating an Option Group \(p. 197\)](#).

2. Add the options to the option group. If you want to deploy only the Oracle APEX run-time environment, add only the **APEX** option. If you want to deploy the full development environment, add both the **APEX** and **APEX-DEV** options.
 - For Oracle 12c, add the **APEX** and **APEX-DEV** options.
 - For Oracle 11g, first add the **XMLDB** option as a prerequisite, then add the **APEX** and **APEX-DEV** options.

For **Version**, choose the version of **APEX** that you want to use. If you don't choose a version, version 4.1.1.v1 is the default for 11g, and version 4.2.6.v1 is the default for 12c.

Important

If you add the APEX options to an existing option group that is already attached to one or more DB instances, a brief outage occurs. During this outage, all the DB instances are automatically restarted.

For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#).

3. Apply the option group to a new or existing DB instance:

- For a new DB instance, you apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).
- For an existing DB instance, you apply the option group by modifying the instance and attaching the new option group. When you add the APEX options to an existing DB instance, a brief outage occurs while your DB instance is automatically restarted. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Unlocking the Public User Account

After the Amazon RDS APEX options are installed, you must change the password for the APEX public user account, and then unlock the account. You can do this by using the Oracle SQL*Plus command line utility. Connect to your DB instance as the master user, and issue the following commands. Replace `new_password` with a password of your choice.

```
alter user APEX_PUBLIC_USER identified by new_password;
alter user APEX_PUBLIC_USER account unlock;
```

Configuring RESTful Services for Oracle APEX

To configure RESTful services in APEX (not needed for APEX 4.1.1.V1), use SQL*Plus to connect to your DB instance as the master user. After you do this, run the `rdsadmin.rdsadmin_run_apex_rest_config` stored procedure. When you run the stored procedure, you provide passwords for the following users:

- `APEX_LISTENER`
- `APEX_REST_PUBLIC_USER`

The stored procedure runs the `apex_rest_config.sql` script, which creates new database accounts for these users.

Note

Configuration isn't required for Oracle APEX version 4.1.1.v1. For this Oracle APEX version only, you don't need to run the stored procedure.

The following command runs the stored procedure.

```
exec rdsadmin.rdsadmin_run_apex_rest_config('apex_listener_password',
'apex_rest_public_user_password');
```

Installing and Configuring a Listener for Use with Oracle APEX

You are now ready to install and configure a listener for use with Oracle APEX. You can use one of these products for this purpose:

- For APEX version 5.0 and later, use Oracle Rest Data Services (ORDS) version 19.1 and higher.
- For APEX version 4.1.1, use Oracle APEX Listener version 1.1.4.
- Oracle HTTP Server and `mod_plsql`.

For more information about these listener types, see [About Choosing a Web Listener](#) in the Oracle documentation.

Note

Amazon RDS doesn't support the Oracle XML DB HTTP server with the embedded PL/SQL gateway; you can't use this as a listener for APEX. In general, Oracle recommends against using the embedded PL/SQL gateway for applications that run on the internet.

Install the listener on a separate host such as an Amazon EC2 instance, an on-premises server at your company, or your desktop computer.

The following procedures show you how to install and configure either ORDS or the APEX Listener. We assume that the name of your host is `myapexhost.example.com`, and that your host is running Linux.

To set up for listener installation

1. Log in to `myapexhost.example.com` as `root`.
2. Create a nonprivileged OS user to own the listener installation. The following command creates a new user named `apexuser`.

```
useradd -d /home/apexuser apexuser
```

The following command assigns a password to the new user.

```
passwd apexuser;
```

3. Log in to `myapexhost.example.com` as `apexuser`, and download the APEX installation file from Oracle to your `/home/apexuser` directory:
 - <http://www.oracle.com/technetwork/developer-tools/apex/downloads/index.html>
 - [Oracle Application Express Prior Release Archives](#)
4. Unzip the file in the `/home/apexuser` directory.

```
unzip apex_<version>.zip
```

After you unzip the file, there is an `apex` directory in the `/home/apexuser` directory.

5. While you are still logged into `myapexhost.example.com` as `apexuser`, download the APEX Listener file from Oracle to your `/home/apexuser` directory: <http://www.oracle.com/technetwork/developer-tools/apex-listener/downloads/index.html>

While you are still in the directory from the previous procedure, install and run the listener program that you chose, as described in one of the procedures following:

- [To install and configure ORDS for use with Oracle APEX \(p. 1019\)](#)
- [To install APEX Listener for use with Oracle APEX \(p. 1021\)](#)

To install and configure ORDS for use with Oracle APEX

1. Create a new directory based on ORDS and open the listener file.

Run the following code:

```
mkdir /home/apexuser/ORDS
cd /home/apexuser/ORDS
unzip ../ords.<version>.zip
```

2. Download the file `ords.version.number.zip` from [Oracle REST Data Services](#).
3. Unzip the file in the `/home/apexuser/ORDS` directory.
4. Grant the master user the required privileges to install ORDS.

After the Amazon RDS APEX option is installed, give the master user the required privileges to install the ORDS schema. You can do this by connecting to the database and running the following commands. Replace `master_user` with the name of your master user.

```
exec rdsadmin.rdsadmin_util.grant_sys_object('DBA_OBJECTS', 'master_user', 'SELECT', true);
exec rdsadmin.rdsadmin_util.grant_sys_object('DBA_ROLE_PRIVS', 'master_user', 'SELECT', true);
exec rdsadmin.rdsadmin_util.grant_sys_object('DBA_TAB_COLUMNS', 'master_user', 'SELECT', true);
exec rdsadmin.rdsadmin_util.grant_sys_object('USER_CONS_COLUMNS', 'master_user', 'SELECT', true);
exec rdsadmin.rdsadmin_util.grant_sys_object('USER_CONSTRAINTS', 'master_user', 'SELECT', true);
exec rdsadmin.rdsadmin_util.grant_sys_object('USER_OBJECTS', 'master_user', 'SELECT', true);
exec rdsadmin.rdsadmin_util.grant_sys_object('USER_PROCEDURES', 'master_user', 'SELECT', true);
exec rdsadmin.rdsadmin_util.grant_sys_object('USER_TAB_COLUMNS', 'master_user', 'SELECT', true);
exec rdsadmin.rdsadmin_util.grant_sys_object('USER_TABLES', 'master_user', 'SELECT', true);
exec rdsadmin.rdsadmin_util.grant_sys_object('USER_VIEWS', 'master_user', 'SELECT', true);
exec rdsadmin.rdsadmin_util.grant_sys_object('WPIUTL', 'master_user', 'EXECUTE', true);
exec rdsadmin.rdsadmin_util.grant_sys_object('DBMS_SESSION', 'master_user', 'EXECUTE', true);
exec rdsadmin.rdsadmin_util.grant_sys_object('DBMSUTILITY', 'master_user', 'EXECUTE', true);
```

Note

These commands apply to ORDS version 19.1 and 19.2.

5. Install the ORDS schema using the downloaded `ords.war` file.

```
java -jar ords.war setup
```

The program prompts you for the following information. The default values are in brackets.

- The location to store configuration data
Enter `/home/apexuser/ORDS`.
- The name of the database server [localhost]: `DB_instance_endpoint`

Choose the default or enter the correct value.

- The database listen port [1521]: `DB_instance_port`

Choose the default or enter the correct value.

- Database service name or database SID [1]

Choose 1 to specify the database service name, or choose 2 to specify the database SID.

- Database SID [xe]

Choose the default or enter the correct value.

- Verify or install Oracle REST Data Services schema or skip this step [1]

Choose 1.

- Enter the database password for ORDS_PUBLIC_USER, and then confirm the password.
- Log in with administrator privileges to verify the Oracle REST Data Services schema.

Enter the administrator user name: *master_user*

Enter the database password for *master_user*: *master_user_password*

Confirm the password: *master_user_password*

- Enter the default tablespace for ORDS_METADATA [SYSAUX].

Enter the temporary tablespace for ORDS_METADATA [TEMP].

Enter the default tablespace for ORDS_PUBLIC_USER [USERS].

Enter the temporary tablespace for ORDS_PUBLIC_USER [TEMP].

- Enter 1 if you want to use PL/SQL Gateway or 2 to skip this step.

If you're using Oracle Application Express or migrating from mod_plsql, you must enter 1 [1].

Choose the default.

- PL/SQL Gateway database user name [APEX_PUBLIC_USER]

Choose the default.

- Database password for APEX_PUBLIC_USER

Enter the password.

- Specify passwords for Application Express RESTful Services database users (APEX_LISTENER, APEX_REST_PUBLIC_USER) or skip this step [1]

Choose 2 for APEX 4.1.1.V1; choose 1 for all other APEX versions.

- [Not needed for APEX 4.1.1.v1] Database password for APEX_LISTENER

Enter the password (if required).

- [Not needed for APEX 4.1.1.v1] Database password for APEX_REST_PUBLIC_USER

Enter the password (if required).

After you install ORDS, configure it by following the instructions in [To configure your listener \(p. 1022\)](#).

To install APEX Listener for use with Oracle APEX

1. Create a new directory based on the Apex Listener and open the listener file.

Run the following code:

```
language="bash">>mkdir /home/apexuser/apexlistener
cd /home/apexuser/apexlistener
unzip ..../apex_listener.<version>.zip
```

2. Run the following code.

```
java -Dapex.home=../apex -Dapex.images=/home/apexuser/apex/images -Dapex.erase -jar ../apex.war
1021
```

3. Enter information for the program prompts following:
 - The APEX Listener Administrator user name. The default is *adminlistener*.
 - A password for the APEX Listener Administrator.
 - The APEX Listener Manager user name. The default is *managerlistener*.
 - A password for the APEX Listener Administrator.

The program prints a URL that you need to complete the configuration, as follows.

```
INFO: Please complete configuration at: http://localhost:8080/apex/listenerConfigure
Database is not yet configured
```

4. Leave the APEX Listener running so that you can use Oracle Application Express. When you have finished this configuration procedure, you can run the listener in the background.
5. From your web browser, go to the URL provided by the APEX Listener program. The Oracle Application Express Listener administration window appears. Enter the following information:
 - **Username** – `APEX_PUBLIC_USER`
 - **Password** – the password for `APEX_PUBLIC_USER`. This password is the one that you specified earlier when you configured the APEX repository. For more information, see [Unlocking the Public User Account \(p. 1018\)](#).
 - **Connection Type** – Basic
 - **Hostname** – the endpoint of your Amazon RDS DB instance, such as `mydb.f9rbfa893tft.us-east-1.rds.amazonaws.com`.
 - **Port** – 1521
 - **SID** – the name of the database on your Amazon RDS DB instance, such as `mydb`.
6. Choose **Apply**. The APEX administration window appears.

After you install the APEX Listener, configure it as described following.

To configure your listener

1. Set a password for the APEX admin user. To do this, use SQL*Plus to connect to your DB instance as the master user, and then run the following commands.

```
EXEC rdsadmin.rdsadmin_util.grant_apex_admin_role;
grant APEX_ADMINISTRATOR_ROLE to master;
@/home/apexuser/apex/apxchpwd.sql
```

Replace `master` with your master user name. When the `apxchpwd.sql` script prompts you, enter a new admin password.

2. For ORDS, start the APEX Listener. Run the following code.

```
java -jar ords.war
```

The first time you start the APEX Listener, you are prompted to provide the location of the APEX Static resources. This images folder is located in the `/apex/images` directory in the installation directory for APEX.

3. Return to the APEX administration window in your browser and choose **Administration**. Next, choose **Application Express Internal Administration**. When you are prompted for credentials, enter the following information:

- **User name** – admin
- **Password** – the password you set using the `apxchpwd.sql` script

Choose [Login](#), and then set a new password for the `admin` user.

Your listener is now ready for use.

Upgrading the APEX Version

Important

Back up your DB instance before you upgrade APEX. For more information, see [Creating a DB Snapshot \(p. 264\)](#) and [Testing an Upgrade \(p. 982\)](#).

To upgrade APEX with your DB instance, do the following:

- Create a new option group for the upgraded version of your DB instance.
- Add the upgraded versions of APEX and APEX-DEV to the new option group. Be sure to include any other options that your DB instance uses. For more information, see [Option Group Considerations \(p. 982\)](#).
- When you upgrade your DB instance, specify the new option group for your upgraded DB instance.

After you upgrade your version of APEX, the APEX schema for the previous version might still exist in your database. If you don't need it anymore, you can drop the old APEX schema from your database after you upgrade.

If you upgrade the APEX version and RESTful services were not configured in the previous APEX version, we recommend that you configure RESTful services. For more information, see [Configuring RESTful Services for Oracle APEX \(p. 1018\)](#).

In some cases when you plan to do a major version upgrade of your DB instance, you might find that you're using an APEX version that isn't compatible with your target database version. In these cases, you can upgrade your version of APEX before you upgrade your DB instance. Upgrading APEX first can reduce the amount of time that it takes to upgrade your DB instance.

Note

After upgrading APEX, install and configure a listener for use with the upgraded version. For instructions, see [Installing and Configuring a Listener for Use with Oracle APEX \(p. 1018\)](#).

Removing the APEX Option

You can remove the Amazon RDS APEX options from a DB instance. To remove the APEX options from a DB instance, do one of the following:

- To remove the APEX options from multiple DB instances, remove the APEX options from the option group they belong to. This change affects all DB instances that use the option group. When you remove the APEX options from an option group that is attached to multiple DB instances, a brief outage occurs while all the DB instances are restarted.

For more information, see [Removing an Option from an Option Group \(p. 207\)](#).

- To remove the APEX options from a single DB instance, modify the DB instance and specify a different option group that doesn't include the APEX options. You can specify the default (empty) option group, or a different custom option group. When you remove the APEX options, a brief outage occurs while your DB instance is automatically restarted.

For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

When you remove the APEX options from a DB instance, the APEX schema is removed from your database.

Oracle Enterprise Manager

Amazon RDS supports Oracle Enterprise Manager (OEM). OEM is the Oracle product line for integrated management of enterprise information technology.

Amazon RDS supports OEM through the following options.

Option	Option ID	Support For
OEM Database (p. 1025)	OEM	OEM Database Express 12c
		OEM 11g Database Control
OEM Management Agent (p. 1028)	OEM_AGENT	OEM Cloud Control for 13c
		OEM Cloud Control for 12c

Note

You can use OEM Database or OEM Management Agent, but not both.

Oracle Enterprise Manager Database Express

Amazon RDS supports Oracle Enterprise Manager (OEM) Database Express through the use of the OEM option. Amazon RDS supports the following versions of OEM database:

- Oracle Enterprise Manager Database Express for Oracle 18c and Oracle 12c
- Oracle Enterprise Manager 11g Database Control for Oracle 11g

OEM Database Express and Database Control are similar tools that have a web-based interface for Oracle database administration. For more information about these tools, see [Accessing Enterprise Manager Database Express 18c](#), [Accessing Enterprise Manager Database Express 12c](#), and [Accessing Enterprise Manager 11g Database Control](#) in the Oracle documentation.

The following are some limitations to using OEM Database:

- OEM Database is not supported on the following DB instance classes: db.t2.micro, db.t2.small, db.t3.micro, db.t3.small.

For more information about DB instance classes, see [DB Instance Class Support for Oracle \(p. 917\)](#).

- OEM 11g Database Control is not compatible with the following time zones: America/Argentina/Buenos_Aires, America/Matamoros, America/Monterrey, America/Toronto, Asia/Ashgabat, Asia/Dhaka, Asia/Kathmandu, Asia/Kolkata, Asia/Ulaanbaatar, Atlantic/Cape_Verde, Australia/Eucla, Pacific/Kiritimati.

For more information about time zone support, see [Oracle Time Zone \(p. 1071\)](#).

OEM Database Option Settings

Amazon RDS supports the following settings for the OEM option.

Option Setting	Valid Values	Description
Port	An integer value	The port on the DB instance that listens for OEM Database. The default for OEM Database Express is 5500. The default for OEM 11g Database Control is 1158.
Security Groups	—	A security group that has access to Port .

Adding the OEM Database Option

The general process for adding the OEM option to a DB instance is the following:

1. Create a new option group, or copy or modify an existing option group.
2. Add the option to the option group.
3. Associate the option group with the DB instance.

When you add the OEM option for an Oracle 19c, Oracle 18c, or Oracle 12c DB instance, a brief outage occurs while your DB instance is automatically restarted. So, after you add the OEM option, you don't need to restart your DB instance. As soon as the option group is active, the OEM Database is active.

When you add the OEM option for an Oracle 11g DB instance, there is not outage, and you don't need to restart your DB instance.

To add the OEM option to a DB instance

1. Determine the option group you want to use. You can create a new option group or use an existing option group. If you want to use an existing option group, skip to the next step. Otherwise, create a custom DB option group with the following settings:
 - a. For **Engine** choose the oracle edition for your DB instance.
 - b. For **Major engine version** choose the version of your DB instance.

For more information, see [Creating an Option Group \(p. 197\)](#).

2. Add the OEM option to the option group, and configure the option settings. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#). For more information about each setting, see [OEM Database Option Settings \(p. 1025\)](#).

Note

If you add the OEM option to an existing option group that is already attached to one or more Oracle 19c, Oracle 18c, or Oracle 12c DB instances, a brief outage occurs while all the DB instances are automatically restarted.

3. Apply the option group to a new or existing DB instance:

- For a new DB instance, you apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).
- For an existing DB instance, you apply the option group by modifying the instance and attaching the new option group. When you add the OEM option for an Oracle 19c, Oracle 18c, or Oracle 12c DB instance, a brief outage occurs while your DB instance is automatically restarted. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Using OEM Database

After you enable the OEM option, you can begin using the OEM Database tool from your web browser.

You can access either OEM Database Control or OEM Database Express from your web browser. For example, if the endpoint for your Amazon RDS DB instance is `mydb.f9rbfa893tft.us-east-1.rds.amazonaws.com`, and your OEM port is 1158, then the URL to access the OEM Database Control the following.

```
https://mydb.f9rbfa893tft.us-east-1.rds.amazonaws.com:1158/em
```

When you access either tool from you web browser, a login window appears that prompts you for a user name and password. Type the master user name and master password for your DB instance. You are now ready to manage your Oracle databases.

Modifying OEM Database Settings

After you enable OEM Database, you can modify the Security Groups setting for the option.

You can't modify the OEM port number after you have associated the option group with a DB instance. To change the OEM port number for a DB instance, do the following:

1. Create a new option group.
2. Add the OEM option with the new port number to the new option group.
3. Remove the existing option group from the DB instance.

4. Add the new option group to the DB instance.

For more information about how to modify option settings, see [Modifying an Option Setting \(p. 204\)](#). For more information about each setting, see [OEM Database Option Settings \(p. 1025\)](#).

Removing the OEM Database Option

You can remove the OEM option from a DB instance. When you remove the OEM option for an Oracle 19c, Oracle 18c, or Oracle 12c DB instance, a brief outage occurs while your DB instance is automatically restarted. So, after you remove the OEM option, you don't need to restart your DB instance. When you remove the OEM option for an Oracle 11g DB instance, there is not outage, and you don't need to restart your DB instance.

To remove the OEM option from a DB instance, do one of the following:

- Remove the OEM option from the option group it belongs to. This change affects all DB instances that use the option group. For more information, see [Removing an Option from an Option Group \(p. 207\)](#)
- Modify the DB instance and specify a different option group that doesn't include the OEM option. This change affects a single DB instance. You can specify the default (empty) option group, or a different custom option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Oracle Management Agent for Enterprise Manager Cloud Control

Amazon RDS supports Oracle Enterprise Manager (OEM) Management Agent through the use of the OEM_AGENT option. Amazon RDS supports Management Agent for the following versions of OEM:

- Oracle Enterprise Manager Cloud Control for 13c
- Oracle Enterprise Manager Cloud Control for 12c

Management Agent is a software component that monitors targets running on hosts and communicates that information to the middle-tier Oracle Management Service (OMS). For more information, see [Overview of Oracle Enterprise Manager Cloud Control 12c](#) and [Overview of Oracle Enterprise Manager Cloud Control 13c](#) in the Oracle documentation.

Following are the supported Oracle versions for each Management Agent version.

Management Agent Version	Oracle 19c	Oracle 18c	Oracle 12c version 12.2	Oracle 12c version 12.1	Oracle 11g
13.3.0.0.v2	Supported	Supported	Supported	Supported	Supported
13.3.0.0.v1	Supported	Supported	Supported	Supported	Supported
13.2.0.0.v3	Supported	Supported	Supported	Supported	Supported
13.2.0.0.v2	Supported	Supported	Supported	Supported	Supported
13.2.0.0.v1	Supported	Supported	Supported	Supported	Supported
13.1.0.0.v1	Supported	Supported	Supported	Supported	Supported
12.1.0.5.v1	Not supported	Supported	Supported	Supported	Supported
12.1.0.4.v1	Not supported	Supported	Supported	Supported	Supported

The following are some limitations to using Management Agent:

- Administrative tasks such as job execution and database patching, that require host credentials, are not supported.
- Host metrics and the process list are not guaranteed to reflect the actual system state.
- Autodiscovery is not supported. You must manually add database targets.
- OMS module availability depends on your database edition. For example, the database performance diagnosis and tuning module is only available for Oracle Database Enterprise Edition.
- Management Agent consumes additional memory and computing resources. If you experience performance problems after enabling the OEM_AGENT option, we recommend that you scale up to a larger DB instance class. For more information, see [Choosing the DB Instance Class \(p. 107\)](#) and [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).
- Because operating system access on the alert log isn't granted to the user running the OEM_AGENT on the Amazon RDS host, it isn't possible to collect metrics for DB Alert Log and DB Alert Log Error Status in OEM.

Prerequisites for Management Agent

The following are prerequisites for using Management Agent:

- An Amazon RDS DB instance running Oracle version 19.0.0.0, 18.0.0.0, 12.2.0.1, 12.1.0.2, or 11.2.0.4.
- At least 3.3 GiB of storage space for OEM 13c Release 3.
- At least 3.3 GiB of storage space for OEM 13c Release 2.
- At least 3 GiB of storage space for OEM 13c Release 1.
- At least 2 GiB of storage space for OEM 12c.
- For Oracle version 19.0.0.0, the minimum `AGENT_VERSION` is 13.1.0.0.v1.
- An Oracle Management Service (OMS), configured to connect to your Amazon RDS DB instance.

For an Amazon RDS DB instance running Oracle version 18.0.0.0 or higher, meet the following requirements:

- For OMS 13c2, apply the Enterprise Manager 13.2 Master Bundle Patch List, which includes plugins 13.2.1, 13.2.2, 13.2.3, 13.2.4 (Oracle Doc ID 2219797.1).
- For OMS 13c2, apply the OMS PSU System Patch 28970534.
- For OMS 13c2, apply the OMS-Side Plugin System 13.2.2.0.190131 Patch 29201709.

For an Amazon RDS DB instance running Oracle version 12.2.0.1 or lower, meet the following requirements:

- For OMS 13c Release 2 with Oracle patch 25163555 applied, use OEM Agent 13.2.0.0.v2 or later.
 - Use OMSPatcher to apply the patch.
- For unpatched OMS 13c Release 2, use OEM Agent 13.2.0.0.v1.

Use OMSPatcher to apply patches.

- In most cases, you need to configure your VPC to allow connections from OMS to your DB instance. If you are not familiar with Amazon Virtual Private Cloud (Amazon VPC), we recommend that you complete the steps in [Tutorial: Create an Amazon VPC for Use with a DB Instance \(p. 521\)](#) before continuing.
- If you are using Management Agent versions `OEM_AGENT 13.2.0.0.v3` and `13.3.0.0.v2`, and you want to use TCPS connectivity, follow the instructions in [Configuring Third Party CA Certificates for Communication With Target Databases](#) in the Oracle documentation. Also, update the JDK on your OMS by following the instructions in the Oracle document with the Oracle Doc ID 2241358.1. Doing so ensures that OMS supports all the cipher suites that the database supports.

Note

TCPS connectivity between the Management Agent and the DB instance is only supported for Management Agent versions `OEM_AGENT 13.2.0.0.v3` and `13.3.0.0.v2`.

Additional configuration is required to allow your OMS host and your Amazon RDS DB instance to communicate. You must also do the following:

- To connect from the Management Agent to your OMS, if your OMS is behind a firewall, you must add the IP addresses of your DB instances to your OMS.

Make sure the firewall for the OMS allows traffic from both the DB listener port (default 1521) and the OEM Agent port (default 3872), originating from the IP address of the DB instance.

- To connect from your OMS to the Management Agent, if your OMS has a publicly resolvable host name, you must add the OMS address to a security group. Your security group must have inbound rules that allow access to the DB listener port and the Management Agent port. For an example of creating a security and adding inbound rules, see [Tutorial: Create an Amazon VPC for Use with a DB Instance \(p. 521\)](#).
- To connect from your OMS to the Management Agent, if your OMS doesn't have a publicly resolvable host name, use one of the following:

- If your OMS is hosted on an Amazon Elastic Compute Cloud (Amazon EC2) instance in a private VPC, you can set up VPC peering to connect from OMS to Management Agent. For more information, see [A DB Instance in a VPC Accessed by an EC2 Instance in a Different VPC \(p. 509\)](#).
- If your OMS is hosted on-premises, you can set up a VPN connection to allow access from OMS to Management Agent. For more information, see [A DB Instance in a VPC Accessed by a Client Application Through the Internet \(p. 510\)](#) or [VPN Connections](#).

Option Settings for Management Agent

Amazon RDS supports the following settings for the Management Agent option.

Note

All of the settings are required.

Option Setting	Valid Values	Description
Version (AGENT_VERSION)	13.3.0.0.v2 13.3.0.0.v1 13.2.0.0.v3 13.2.0.0.v2 13.2.0.0.v1 13.1.0.0.v1 12.1.0.5.v1 12.1.0.4.v1	The version of the Management Agent software. The AWS CLI option name is OptionVersion. Note In the AWS GovCloud (US-West) Region, 12.1 and 13.1 versions aren't available.
Port (AGENT_PORT)	An integer value	The port on the DB instance that listens for the OMS host. The default is 3872. Your OMS host must belong to a security group that has access to this port. The AWS CLI option name is Port.
Security Groups	Existing security groups	A security group that has access to Port . Your OMS host must belong to this security group. The AWS CLI option name is VpcSecurityGroupMemberships or DBSecurityGroupMemberships.
OMS_HOST	A string value, for example <i>my.example.oms</i>	The publicly accessible host name or IP address of the OMS. The AWS CLI option name is OMS_HOST.
OMS_PORT	An integer value	The HTTPS upload port on the OMS Host that listens for the Management Agent.

Option Setting	Valid Values	Description
		<p>To determine the HTTPS upload port, connect to the OMS host, and run the following command (which requires the SYSMAN password): <code>emctl status oms - details</code></p> <p>The AWS CLI option name is <code>OMS_PORT</code>.</p>
<code>AGENT_REGISTRATION_PASSWORD</code>	string value	<p>The password that the Management Agent uses to authenticate itself with the OMS. We recommend that you create a persistent password in your OMS before enabling the OEM_AGENT option. With a persistent password you can share a single Management Agent option group among multiple Amazon RDS databases.</p> <p>The AWS CLI option name is <code>AGENT_REGISTRATION_PASSWORD</code>.</p>

Adding the Management Agent Option

The general process for adding the Management Agent option to a DB instance is the following:

1. Create a new option group, or copy or modify an existing option group.
2. Add the option to the option group.
3. Associate the option group with the DB instance.

If you encounter errors, you can check [My Oracle Support](#) documents for information about resolving specific problems.

After you add the Management Agent option, you don't need to restart your DB instance. As soon as the option group is active, the OEM Agent is active.

If your OMS host is using an untrusted third-party certificate, Amazon RDS returns the following error.

```
You successfully installed the OEM_AGENT option. Your OMS host is using an untrusted third
party certificate.
Configure your OMS host with the trusted certificates from your third party.
```

If this error is returned, the Management Agent option isn't enabled until the problem is corrected. For information about correcting the problem, see the My Oracle Support document [2202569.1](#).

Console

To add the Management Agent option to a DB instance

1. Determine the option group you want to use. You can create a new option group or use an existing option group. If you want to use an existing option group, skip to the next step. Otherwise, create a custom DB option group with the following settings:
 - a. For **Engine** choose the oracle edition for your DB instance.
 - b. For **Major engine version** choose the version of your DB instance.

For more information, see [Creating an Option Group \(p. 197\)](#).

2. Add the **OEM_AGENT** option to the option group, and configure the option settings. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#). For more information about each setting, see [Option Settings for Management Agent \(p. 1030\)](#).
3. Apply the option group to a new or existing DB instance:
 - For a new DB instance, you apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).
 - For an existing DB instance, you apply the option group by modifying the instance and attaching the new option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

AWS CLI

The following example uses the AWS CLI `add-option-to-option-group` command to add the **OEM_AGENT** option to an option group called `myoptiongroup`.

For Linux, OS X, or Unix:

```
aws rds add-option-to-option-group \
--option-group-name "myoptiongroup" \
--options
{OptionName=OEM_AGENT,OptionVersion=13.1.0.0.v1,Port=3872,VpcSecurityGroupMemberships=sg-1234567890,Opt
{Name=OMS_PORT,Value=4903},{Name=AGENT_REGISTRATION_PASSWORD,Value=password}}] \
--apply-immediately
```

For Windows:

```
aws rds add-option-to-option-group ^
--option-group-name "myoptiongroup" ^
--options
{OptionName=OEM_AGENT,OptionVersion=13.1.0.0.v1,Port=3872,VpcSecurityGroupMemberships=sg-1234567890,Opt
{Name=OMS_PORT,Value=4903},{Name=AGENT_REGISTRATION_PASSWORD,Value=password}}] ^
--apply-immediately
```

Using the Management Agent

After you enable the Management Agent option, take the following steps to begin using it.

To use the Management Agent

1. Unlock and reset the DBSNMP account credential. Do this by running the following code on your target database on your DB instance and using your master user account.

```
ALTER USER dbsnmp IDENTIFIED BY new_password ACCOUNT UNLOCK;
```

2. Add your targets to the OMS console manually:
 - a. In your OMS console, choose **Setup, Add Target, Add Targets Manually**.
 - b. Choose **Add Targets Declaratively by Specifying Target Monitoring Properties**.
 - c. For **Target Type**, choose **Database Instance**.
 - d. For **Monitoring Agent**, choose the agent with the identifier that is the same as your RDS DB instance identifier.
 - e. Choose **Add Manually**.
 - f. Enter the endpoint for the Amazon RDS DB instance, or choose it from the host name list. Make sure that the specified host name matches the endpoint of the Amazon RDS DB instance.

For information about finding the endpoint for your Amazon RDS DB instance, see [Finding the Endpoint of Your DB Instance \(p. 955\)](#).
 - g. Specify the following database properties:
 - For **Target name**, enter a name.
 - For **Database system name**, enter a name.
 - For **Monitor username**, enter **dbsnmp**.
 - For **Monitor password**, enter the password from step 1.
 - For **Role**, enter **normal**.
 - For **Oracle home path**, enter **/oracle**.
 - For **Listener Machine name**, the agent identifier already appears.
 - For **Port**, enter the database port. The RDS default port is 1521.
 - For **Database name**, enter the name of your database.
 - h. Choose **Test Connection**.
 - i. Choose **Next**. The target database appears in your list of monitored resources.

Modifying Management Agent Settings

After you enable the Management Agent, you can modify settings for the option. For more information about how to modify option settings, see [Modifying an Option Setting \(p. 204\)](#). For more information about each setting, see [Option Settings for Management Agent \(p. 1030\)](#).

Performing Database Tasks with the Management Agent

You can use Amazon RDS procedures to run certain EMCTL commands on the Management Agent. By running these procedures, you can do the tasks listed following.

Tasks

- [Getting the Management Agent's Status \(p. 1034\)](#)
- [Restarting the Management Agent \(p. 1034\)](#)
- [Listing the Targets Monitored by the Management Agent \(p. 1034\)](#)
- [Clearing the Management Agent's State \(p. 1034\)](#)
- [Having the Management Agent Upload Its OMS \(p. 1034\)](#)
- [Pinging the OMS \(p. 1035\)](#)
- [Viewing the Status of an Ongoing Task \(p. 1035\)](#)

Getting the Management Agent's Status

To get the Management Agent's status, run the Amazon RDS procedure `rdsadmin.rdsadmin_oem_agent_tasks.get_status_oem_agent`. This procedure is equivalent to the `emctl status agent` command.

The following procedure gets the Management Agent's status,

```
SELECT rdsadmin.rdsadmin_oem_agent_tasks.get_status_oem_agent() as TASK_ID from DUAL;
```

Restarting the Management Agent

To restart the Management Agent, run the Amazon RDS procedure `rdsadmin.rdsadmin_oem_agent_tasks.get_status_oem_agent`. This procedure is equivalent to running the `emctl stop agent` and `emctl start agent` commands.

The following procedure restarts the Management Agent.

```
SELECT rdsadmin.rdsadmin_oem_agent_tasks.restart_oem_agent() as TASK_ID from DUAL;
```

Listing the Targets Monitored by the Management Agent

To list the targets monitored by the Management Agent, run the Amazon RDS procedure `rdsadmin.rdsadmin_oem_agent_tasks.list_targets_oem_agent`. This procedure is equivalent to running the `emctl config agent listtargets` command.

The following procedure lists the targets monitored by the Management Agent.

```
SELECT rdsadmin.rdsadmin_oem_agent_tasks.list_targets_oem_agent() as TASK_ID from DUAL;
```

Clearing the Management Agent's State

To clear the Management Agent's state, run the Amazon RDS procedure `rdsadmin.rdsadmin_oem_agent_tasks.clearstate_oem_agent`. This procedure is equivalent to running the `emctl clearstate agent` command.

The following procedure clears the Management Agent's state.

```
SELECT rdsadmin.rdsadmin_oem_agent_tasks.clearstate_oem_agent() as TASK_ID from DUAL;
```

Having the Management Agent Upload Its OMS

To have the Management Agent upload the Oracle Management Server (OMS) associated with it, run the Amazon RDS procedure `rdsadmin.rdsadmin_oem_agent_tasks.upload_oem_agent`. This procedure is equivalent to running the `emctl upload agent` command.

If you run the following procedure, the Management Agent uploads its associated OMS.

```
SELECT rdsadmin.rdsadmin_oem_agent_tasks.upload_oem_agent() as TASK_ID from DUAL;
```

Pinging the OMS

To ping the Management Agent's OMS, run the Amazon RDS procedure `rdsadmin.rdsadmin_oem_agent_tasks.ping_oms_oem_agent`. This procedure is equivalent to running the `emctl pingOMS` command.

The following procedure pings the Management Agent's OMS.

```
SELECT rdsadmin.rdsadmin_oem_agent_tasks.ping_oms_oem_agent() as TASK_ID from DUAL;
```

Viewing the Status of an Ongoing Task

You can view the status of an ongoing task in a bdump file. The bdump files are located in the `/rdsdbdata/log/trace` directory. Each bdump file name is in the following format.

```
dbtask-task-id.log
```

When you want to monitor a task, replace `task-id` with the ID of the task that you want to monitor.

To view the contents of bdump files, run the Amazon RDS procedure `rdsadmin.rds_file_util.read_text_file`. The following query returns the contents of the `dbtask-1546988886389-2444.log` bdump file.

```
SELECT text FROM
  table(rdsadmin.rds_file_util.read_text_file('BDUMP','dbtask-1546988886389-2444.log'));
```

For more information about the Amazon RDS procedure `rdsadmin.rds_file_util.read_text_file`, see [Reading Files in a DB Instance Directory \(p. 1132\)](#).

Removing the Management Agent Option

You can remove the OEM Agent from a DB instance. After you remove the OEM Agent, you don't need to restart your DB instance.

To remove the OEM Agent from a DB instance, do one of the following:

- Remove the OEM Agent option from the option group it belongs to. This change affects all DB instances that use the option group. For more information, see [Removing an Option from an Option Group \(p. 207\)](#).
- Modify the DB instance and specify a different option group that doesn't include the OEM Agent option. This change affects a single DB instance. You can specify the default (empty) option group, or a different custom option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Oracle Java Virtual Machine

Amazon RDS supports Oracle Java Virtual Machine (JVM) through the use of the `JVM` option. Oracle Java provides a SQL schema and functions that facilitate Oracle Java features in an Oracle database. For more information, see [Introduction to Java in Oracle Database](#) in the Oracle documentation.

You can use Oracle JVM with the following Oracle Database versions:

- Oracle 19c, 19.0.0.0, all versions
- Oracle 18c, 18.0.0.0, all versions
- Oracle 12c, 12.2.0.1, all versions
- Oracle 12c, 12.1.0.2.v13 or later
- Oracle 11g, 11.2.0.4.v17 or later

Java implementation in Amazon RDS has a limited set of permissions. The master user is granted the `RDS_JAVA_ADMIN` role, which grants a subset of the privileges granted by the `JAVA_ADMIN` role. To list the privileges granted to the `RDS_JAVA_ADMIN` role, run the following query on your DB instance:

```
SELECT * FROM dba_java_policy
WHERE grantee IN ('RDS_JAVA_ADMIN', 'PUBLIC')
AND enabled = 'ENABLED'
ORDER BY type_name, name, grantee;
```

Prerequisites for Oracle JVM

The following are prerequisites for using Oracle Java:

- Your DB instance must be inside a virtual private cloud (VPC). For more information, see [Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform \(p. 506\)](#).
- Your DB instance must be of a large enough class. Oracle Java isn't supported for the `db.t2.micro`, `db.t2.small`, `db.t3.micro`, or `db.t3.small` DB instance classes. For more information, see [Choosing the DB Instance Class \(p. 107\)](#).
- Your DB instance must have **Auto Minor Version Upgrade** enabled. This option enables your DB instance to receive minor DB engine version upgrades automatically when they become available. Amazon RDS uses this option to update your DB instance to the latest Oracle Patch Set Update (PSU) or Release Update (RU). For more information, see [Settings for Oracle DB Instances \(p. 967\)](#).
- If your DB instance is running on major version 11.2, you must install the `XMLDB` option. For more information, see [Oracle XML DB \(p. 1078\)](#).

Best Practices for Oracle JVM

The following are best practices for using Oracle Java:

- For maximum security, use the `JVM` option with Secure Sockets Layer (SSL). For more information, see [Oracle Secure Sockets Layer \(p. 1051\)](#).
- Configure your DB instance to restrict network access. For more information, see [Scenarios for Accessing a DB Instance in a VPC \(p. 507\)](#) and [Working with a DB Instance in a VPC \(p. 514\)](#).

Adding the Oracle JVM Option

The following is the general process for adding the `JVM` option to a DB instance:

1. Create a new option group, or copy or modify an existing option group.
2. Add the option to the option group.
3. Associate the option group with the DB instance.

There is a brief outage while the JVM option is added. After you add the option, you don't need to restart your DB instance. As soon as the option group is active, Oracle Java is available.

Note

During this outage, password verification functions are disabled briefly. You can also expect to see events related to password verification functions during the outage. Password verification functions are enabled again before the Oracle DB instance is available.

To add the JVM option to a DB instance

1. Determine the option group that you want to use. You can create a new option group or use an existing option group. If you want to use an existing option group, skip to the next step. Otherwise, create a custom DB option group with the following settings:
 - For **Engine**, choose the DB engine used by the DB instance (**oracle-ee**, **oracle-se**, **oracle-se1**, or **oracle-se2**).
 - For **Major engine version**, choose the version of your DB instance.

For more information, see [Creating an Option Group \(p. 197\)](#).

2. Add the **JVM** option to the option group. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#).
3. Apply the option group to a new or existing DB instance:
 - For a new DB instance, apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).
 - For an existing DB instance, apply the option group by modifying the instance and attaching the new option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).
4. Grant the required permissions to users.

The Amazon RDS master user has the permissions to use the JVM option by default. If other users require these permissions, connect to the DB instance as the master user in a SQL client and grant the permissions to the users.

The following example grants the permissions to use the JVM option to the `test_proc` user.

```
create user test_proc identified by password;
CALL dbms_java.grant_permission('TEST_PROC',
  'oracle.aurora.security.JServerPermission', 'LoadClassInPackage.*', '');
```

After the user is granted the permissions, the following query should return output.

```
select * from dba_java_policy where grantee='TEST_PROC';
```

Note

The Oracle user name is case-sensitive, and it usually has all uppercase characters.

Removing the Oracle JVM Option

You can remove the `JVM` option from a DB instance. There is a brief outage while the option is removed. After you remove the `JVM` option, you don't need to restart your DB instance.

Warning

Removing the `JVM` option can result in data loss if the DB instance is using data types that were enabled as part of the option. Back up your data before proceeding. For more information, see [Backing Up and Restoring Amazon RDS DB Instances \(p. 253\)](#).

To remove the `JVM` option from a DB instance, do one of the following:

- Remove the `JVM` option from the option group it belongs to. This change affects all DB instances that use the option group. For more information, see [Removing an Option from an Option Group \(p. 207\)](#).
- Modify the DB instance and specify a different option group that doesn't include the `JVM` option. This change affects a single DB instance. You can specify the default (empty) option group, or a different custom option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Oracle Label Security

Amazon RDS supports Oracle Label Security for Oracle Enterprise Edition, version 12c, through the use of the OLS option.

Most database security controls access at the object level. Oracle Label Security provides fine-grained control of access to individual table rows. For example, you can use Label Security to enforce regulatory compliance with a policy-based administration model. You can use Label Security policies to control access to sensitive data, and restrict access to only users with the appropriate clearance level. For more information, see [Introduction to Oracle Label Security](#) in the Oracle documentation.

Important

For Oracle 19c, Oracle 18c, and Oracle 12c version 12.2 on Amazon RDS, Oracle Label Security is a permanent and persistent option. You can't remove Oracle Label Security from an Oracle version 19c, 18c, or 12.2 DB instance.

Prerequisites for Oracle Label Security

The following are prerequisites for using Oracle Label Security:

- Your DB instance must use the Bring Your Own License model. For more information, see [Oracle Licensing \(p. 916\)](#).
- You must have a valid license for Oracle Enterprise Edition with Software Update License and Support.
- Your Oracle license must include the Label Security option.

Adding the Oracle Label Security Option

The general process for adding the Oracle Label Security option to a DB instance is the following:

1. Create a new option group, or copy or modify an existing option group.
2. Add the option to the option group.
3. Associate the option group with the DB instance.

After you add the Label Security option, as soon as the option group is active, Label Security is active.

To add the Label Security option to a DB instance

1. Determine the option group you want to use. You can create a new option group or use an existing option group. If you want to use an existing option group, skip to the next step. Otherwise, create a custom DB option group with the following settings:
 - a. For **Engine**, choose **oracle-ee**.
 - b. For **Major engine version**, choose the version of your DB instance.

For more information, see [Creating an Option Group \(p. 197\)](#).

2. Add the **OLS** option to the option group. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#).

Important

If you add Label Security to an existing option group that is already attached to one or more DB instances, all the DB instances are restarted.

3. Apply the option group to a new or existing DB instance:

- For a new DB instance, you apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).

- For an existing DB instance, you apply the option group by modifying the instance and attaching the new option group. When you add the Label Security option to an existing DB instance, a brief outage occurs while your DB instance is automatically restarted. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Using Oracle Label Security

To use Oracle Label Security, you create policies that control access to specific rows in your tables. For more information, see [Creating an Oracle Label Security Policy](#) in the Oracle documentation.

When you work with Label Security, you perform all actions as the LBAC_DBA role. The master user for your DB instance is granted the LBAC_DBA role. You can grant the LBAC_DBA role to other users so that they can administer Label Security policies.

For Amazon RDS for Oracle 19c, 18c, and 12.2 DB instances, you must grant access to the OLS_ENFORCEMENT package to any new users who require access to Oracle Label Security. To grant access to the OLS_ENFORCEMENT package, connect to the DB instance as the master user and run the following SQL statement:

```
GRANT ALL ON LBACSYS.OLS_ENFORCEMENT TO username;
```

You can configure Label Security through the Oracle Enterprise Manager (OEM) Cloud Control. Amazon RDS supports the OEM Cloud Control through the Management Agent option. For more information, see [Oracle Management Agent for Enterprise Manager Cloud Control \(p. 1028\)](#).

Removing the Oracle Label Security Option

You can remove Oracle Label Security from a DB instance.

To remove Label Security from a DB instance, do one of the following:

- To remove Label Security from multiple DB instances, remove the Label Security option from the option group they belong to. This change affects all DB instances that use the option group. When you remove Label Security from an option group that is attached to multiple DB instances, all the DB instances are restarted. For more information, see [Removing an Option from an Option Group \(p. 207\)](#).
- To remove Label Security from a single DB instance, modify the DB instance and specify a different option group that doesn't include the Label Security option. You can specify the default (empty) option group, or a different custom option group. When you remove the Label Security option, a brief outage occurs while your DB instance is automatically restarted. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Troubleshooting

The following are issues you might encounter when you use Oracle Label Security.

Issue	Troubleshooting Suggestions
When you try to create a policy, you see an error message similar to the following: <code>insufficient authorization for the SYSDBA package.</code>	A known issue with Oracle's Label Security feature prevents users with usernames of 16 or 24 characters from running Label Security commands. You can create a new user with a different number of characters,

Issue	Troubleshooting Suggestions
	grant LBAC_DBA to the new user, log in as the new user, and run the OLS commands as the new user. For additional information, please contact Oracle support.

Related Topics

- [Working with Option Groups \(p. 196\)](#)
- [Options for Oracle DB Instances \(p. 1004\)](#)

Oracle Locator

Amazon RDS supports Oracle Locator through the use of the `LOCATOR` option. Oracle Locator provides capabilities that are typically required to support internet and wireless service-based applications and partner-based GIS solutions. Oracle Locator is a limited subset of Oracle Spatial. For more information, see [Oracle Locator](#) in the Oracle documentation.

Important

If you use Oracle Locator, Amazon RDS automatically updates your DB instance to the latest Oracle PSU if there are security vulnerabilities with a Common Vulnerability Scoring System (CVSS) score of 9+ or other announced security vulnerabilities.

Amazon RDS supports Oracle Locator for the following editions and versions of Oracle:

- Oracle Standard Edition (SE2) or Enterprise Edition, version 19.0.0.0, all versions
- Oracle Standard Edition (SE2) or Enterprise Edition, version 18.0.0.0, all versions
- Oracle Standard Edition (SE2) or Enterprise Edition, version 12.2.0.1, all versions
- Oracle Standard Edition (SE2) or Enterprise Edition, version 12.1.0.2.v13 or later
- Oracle Standard Edition (SE, SE1) or Enterprise Edition, version 11.2.0.4.v17 or later

Prerequisites for Oracle Locator

The following are prerequisites for using Oracle Locator:

- Your DB instance must be inside a virtual private cloud (VPC). For more information, see [Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform \(p. 506\)](#).
- Your DB instance must be of sufficient class. Oracle Locator is not supported for the db.m1.small, db.t2.micro, db.t2.small, db.t3.micro, or db.t3.small DB instance classes. For more information, see [DB Instance Class Support for Oracle \(p. 917\)](#).
- Your DB instance must have **Auto Minor Version Upgrade** enabled. This option enables your DB instance to receive minor DB engine version upgrades automatically when they become available. Amazon RDS uses this option to update your DB instance to the latest Oracle Patch Set Update (PSU) or Release Update (RU). For more information, see [Settings for Oracle DB Instances \(p. 967\)](#).
- If your DB instance is running on major version 11.2, you must install the `XMLDB` option. For more information, see [Oracle XML DB \(p. 1078\)](#).

Best Practices for Oracle Locator

The following are best practices for using Oracle Locator:

- For maximum security, use the `LOCATOR` option with Secure Sockets Layer (SSL). For more information, see [Oracle Secure Sockets Layer \(p. 1051\)](#).
- Configure your DB instance to restrict access to your DB instance. For more information, see [Scenarios for Accessing a DB Instance in a VPC \(p. 507\)](#) and [Working with a DB Instance in a VPC \(p. 514\)](#).

Adding the Oracle Locator Option

The following is the general process for adding the `LOCATOR` option to a DB instance:

1. Create a new option group, or copy or modify an existing option group.
2. Add the option to the option group.
3. Associate the option group with the DB instance.

There is a brief outage while the `LOCATOR` option is added. After you add the option, you don't need to restart your DB instance. As soon as the option group is active, Oracle Locator is available.

Note

During this outage, password verification functions are disabled briefly. You can also expect to see events related to password verification functions during the outage. Password verification functions are enabled again before the Oracle DB instance is available.

To add the `LOCATOR` option to a DB instance

1. Determine the option group that you want to use. You can create a new option group or use an existing option group. If you want to use an existing option group, skip to the next step. Otherwise, create a custom DB option group with the following settings:
 - a. For **Engine**, choose the oracle edition for your DB instance.
 - b. For **Major engine version**, choose the version of your DB instance.

For more information, see [Creating an Option Group \(p. 197\)](#).
2. Add the `LOCATOR` option to the option group. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#).
3. Apply the option group to a new or existing DB instance:
 - For a new DB instance, you apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).
 - For an existing DB instance, you apply the option group by modifying the instance and attaching the new option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Using Oracle Locator

After you enable the Oracle Locator option, you can begin using it. You should only use Oracle Locator features. Don't use any Oracle Spatial features unless you have a license for Oracle Spatial.

For a list of features that are supported for Oracle Locator, see [Features Included with Locator](#) in the Oracle documentation.

For a list of features that are not supported for Oracle Locator, see [Features Not Included with Locator](#) in the Oracle documentation.

Removing the Oracle Locator Option

You can remove the `LOCATOR` option from a DB instance. There is a brief outage while the option is removed. After you remove the `LOCATOR` option, you don't need to restart your DB instance.

Warning

Removing the `LOCATOR` option can result in data loss if the DB instance is using data types that were enabled as part of the option. Back up your data before proceeding. For more information, see [Backing Up and Restoring Amazon RDS DB Instances \(p. 253\)](#).

To remove the `LOCATOR` option from a DB instance, do one of the following:

- Remove the `LOCATOR` option from the option group it belongs to. This change affects all DB instances that use the option group. For more information, see [Removing an Option from an Option Group \(p. 207\)](#).
- Modify the DB instance and specify a different option group that doesn't include the `LOCATOR` option. This change affects a single DB instance. You can specify the default (empty) option group or a

different custom option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Related Topics

- [Oracle Spatial \(p. 1060\)](#)
- [Options for Oracle DB Instances \(p. 1004\)](#)
- [Working with Option Groups \(p. 196\)](#)

Oracle Multimedia

Amazon RDS supports Oracle Multimedia through the use of the `MULTIMEDIA` option. You can use Oracle Multimedia to store, manage, and retrieve images, audio, video, and other heterogeneous media data. For more information, see [Oracle Multimedia](#) in the Oracle documentation.

Important

If you use Oracle Multimedia, Amazon RDS automatically updates your DB instance to the latest Oracle PSU if there are security vulnerabilities with a Common Vulnerability Scoring System (CVSS) score of 9+ or other announced security vulnerabilities.

Amazon RDS supports Oracle Multimedia for the following editions and versions of Oracle:

- Oracle Standard Edition (SE2) or Oracle Enterprise Edition, version 18.0.0.0, all versions
- Oracle Standard Edition (SE2) or Oracle Enterprise Edition, version 12.2.0.1, all versions
- Oracle Standard Edition (SE2) or Oracle Enterprise Edition, version 12.1.0.2.v13 or later
- Oracle Standard Edition (SE and SE1) or Oracle Enterprise Edition, version 11.2.0.4.v17 or later

Note

Oracle desupported Oracle Multimedia in Oracle Database 19c. So, Oracle Multimedia isn't supported for Oracle 19c DB instances. For more information, see [Desupport of Oracle Multimedia](#) in the Oracle documentation.

Prerequisites for Oracle Multimedia

The following are prerequisites for using Oracle Multimedia:

- Your DB instance must be inside a virtual private cloud (VPC). For more information, see [Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform \(p. 506\)](#).
- Your DB instance must be of sufficient class. Oracle Multimedia is not supported for the db.m1.small, db.t2.micro, or db.t2.small DB instance classes. For more information, see [DB Instance Class Support for Oracle \(p. 917\)](#).
- Your DB instance must have **Auto Minor Version Upgrade** enabled. This option enables your DB instance to receive minor DB engine version upgrades automatically when they become available. Amazon RDS uses this option to update your DB instance to the latest Oracle Patch Set Update (PSU) or Release Update (RU). For more information, see [Settings for Oracle DB Instances \(p. 967\)](#).
- If your DB instance is running on major version 11.2, you must install the `XMLDB` option. For more information, see [Oracle XML DB \(p. 1078\)](#).

Best Practices for Oracle Multimedia

The following are best practices for using Oracle Multimedia:

- For maximum security, use the `MULTIMEDIA` option with Secure Sockets Layer (SSL). For more information, see [Oracle Secure Sockets Layer \(p. 1051\)](#).
- Configure your DB instance to restrict access to your DB instance. For more information, see [Scenarios for Accessing a DB Instance in a VPC \(p. 507\)](#) and [Working with a DB Instance in a VPC \(p. 514\)](#).

Adding the Oracle Multimedia Option

The following is the general process for adding the `MULTIMEDIA` option to a DB instance:

1. Create a new option group, or copy or modify an existing option group.

2. Add the option to the option group.
3. Associate the option group with the DB instance.

There is a brief outage while the **MULTIMEDIA** option is added. After you add the option, you don't need to restart your DB instance. As soon as the option group is active, Oracle Multimedia is available.

Note

During this outage, password verification functions are disabled briefly. You can also expect to see events related to password verification functions during the outage. Password verification functions are enabled again before the Oracle DB instance is available.

To add the **MULTIMEDIA** option to a DB instance

1. Determine the option group that you want to use. You can create a new option group or use an existing option group. If you want to use an existing option group, skip to the next step. Otherwise, create a custom DB option group with the following settings:
 - a. For **Engine**, choose **oracle-ee**.
 - b. For **Major engine version**, choose the version of your DB instance.

For more information, see [Creating an Option Group \(p. 197\)](#).
2. Add the **MULTIMEDIA** option to the option group. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#).
3. Apply the option group to a new or existing DB instance:
 - For a new DB instance, you apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).
 - For an existing DB instance, you apply the option group by modifying the instance and attaching the new option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Removing the Oracle Multimedia Option

You can remove the **MULTIMEDIA** option from a DB instance. There is a brief outage while the option is removed. After you remove the **MULTIMEDIA** option, you don't need to restart your DB instance.

Warning

Removing the **MULTIMEDIA** option can result in data loss if the DB instance is using data types that were enabled as part of the option. Back up your data before proceeding. For more information, see [Backing Up and Restoring Amazon RDS DB Instances \(p. 253\)](#).

To remove the **MULTIMEDIA** option from a DB instance, do one of the following:

- Remove the **MULTIMEDIA** option from the option group it belongs to. This change affects all DB instances that use the option group. For more information, see [Removing an Option from an Option Group \(p. 207\)](#).
- Modify the DB instance and specify a different option group that doesn't include the **MULTIMEDIA** option. This change affects a single DB instance. You can specify the default (empty) option group or a different custom option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Related Topics

- [Working with Option Groups \(p. 196\)](#)

- [Options for Oracle DB Instances \(p. 1004\)](#)

Oracle Native Network Encryption

Amazon RDS supports Oracle native network encryption (NNE). With native network encryption, you can encrypt data as it moves to and from a DB instance. Amazon RDS supports NNE for all editions of Oracle.

A detailed discussion of Oracle native network encryption is beyond the scope of this guide, but you should understand the strengths and weaknesses of each algorithm and key before you decide on a solution for your deployment. For information about the algorithms and keys that are available through Oracle native network encryption, see [Configuring Network Data Encryption](#) in the Oracle documentation. For more information about AWS security, see the [AWS Security Center](#).

Note

You can use Native Network Encryption or Secure Sockets Layer, but not both. For more information, see [Oracle Secure Sockets Layer \(p. 1051\)](#).

NNE Option Settings

Amazon RDS supports the following settings for the NNE option.

Option Setting	Valid Values	Default Value	Description
<code>SQLNET.ENCRYPTION_SERVER</code>	Rejected, Requested, Required	Requested	The encryption behavior when a client, or a server acting as a client, connects to the DB instance. Requested indicates that the DB instance does not require traffic from the client to be encrypted.
<code>SQLNET.CRYPTO_CHECKSUM_SERVER</code>	Rejected, Requested, Required	Requested	The data integrity behavior when a client, or a server acting as a client, connects to the DB instance. Requested indicates that the DB instance does not require the client to perform a checksum.
<code>SQLNET.ENCRYPTION_TYPES_SERVER</code>	AES256, AES192, AES128, 3DES168, RC4_128, AES128, 3DES112, RC4_56, DES, RC4_40, DES40	AES256, AES192, AES128, 3DES168, RC4_128, AES128, 3DES112, RC4_56, DES, RC4_40, DES40	A list of encryption algorithms used by the DB instance. The DB instance will use each algorithm, in order, to attempt to decrypt the client input until an algorithm succeeds or until the end of the list is reached. Amazon RDS uses the following default list from Oracle. You can change the order or limit the algorithms that the DB instance will accept. <ol style="list-style-type: none"> 1. <code>RC4_256</code>: RSA RC4 (256-bit key size) 2. <code>AES256</code>: AES (256-bit key size) 3. <code>AES192</code>: AES (192-bit key size) 4. <code>3DES168</code>: 3-key Triple-DES (112-bit effective key size)

Option Setting	Valid Values	Default Value	Description
			5. RC4_128: RSA RC4 (128-bit key size) 6. AES128: AES (128-bit key size) 7. 3DES112: 2-key Triple-DES (80-bit effective key size) 8. RC4_56: RSA RC4 (56-bit key size) 9. DES: Standard DES (56-bit key size) 10. RC4_40: RSA RC4 (40-bit key size) 11. DES40: DES40 (40-bit key size)
SQLNET.CRYPTO_CHECKSUM_TYPES_SERVER	SHA256, SHA384, SHA512, SHA1, MD5	MD5	The checksum algorithm.

Adding the NNE Option

The general process for adding the NNE option to a DB instance is the following:

1. Create a new option group, or copy or modify an existing option group.
2. Add the option to the option group.
3. Associate the option group with the DB instance.

After you add the NNE option, as soon as the option group is active, NNE is active.

To add the NNE option to a DB instance

1. For **Engine**, choose the Oracle edition that you want to use. NNE is supported on all editions.
2. For **Major engine version**, choose the version of your DB instance.

For more information, see [Creating an Option Group \(p. 197\)](#).

3. Add the **NNE** option to the option group. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#).

Note

After you add the NNE option, you don't need to restart your DB instances. As soon as the option group is active, NNE is active.

4. Apply the option group to a new or existing DB instance:
 - For a new DB instance, you apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).
 - For an existing DB instance, you apply the option group by modifying the instance and attaching the new option group. After you add the NNE option, you don't need to restart your DB instance. As soon as the option group is active, NNE is active. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Using NNE

With Oracle native network encryption, you can also specify network encryption on the client side. On the client (the computer used to connect to the DB instance), you can use the sqlnet.ora file to specify the following client settings: SQLNET.CRYPTO_CHECKSUM_CLIENT, SQLNET.CRYPTO_CHECKSUM_TYPES_CLIENT, SQLNET.ENCRYPTION_CLIENT, and SQLNET.ENCRYPTION_TYPES_CLIENT. For information, see [Configuring Network Data Encryption and Integrity for Oracle Servers and Clients](#) in the Oracle documentation.

Sometimes, the DB instance will reject a connection request from an application, for example, if there is a mismatch between the encryption algorithms on the client and on the server.

To test Oracle native network encryption , add the following lines to the sqlnet.ora file on the client:

```
DIAG_ADR_ENABLED=off
TRACE_DIRECTORY_CLIENT=/tmp
TRACE_FILE_CLIENT=nettrace
TRACE_LEVEL_CLIENT=16
```

These lines generate a trace file on the client called /tmp/nettrace* when the connection is attempted. The trace file contains information on the connection. For more information about connection-related issues when you are using Oracle Native Network Encryption, see [About Negotiating Encryption and Integrity](#) in the Oracle documentation.

Modifying NNE Settings

After you enable NNE, you can modify settings for the option. For more information about how to modify option settings, see [Modifying an Option Setting \(p. 204\)](#). For more information about each setting, see [NNE Option Settings \(p. 1049\)](#).

Removing the NNE Option

You can remove NNE from a DB instance.

To remove NNE from a DB instance, do one of the following:

- To remove NNE from multiple DB instances, remove the NNE option from the option group they belong to. This change affects all DB instances that use the option group. After you remove the NNE option, you don't need to restart your DB instances. For more information, see [Removing an Option from an Option Group \(p. 207\)](#).
- To remove NNE from a single DB instance, modify the DB instance and specify a different option group that doesn't include the NNE option. You can specify the default (empty) option group, or a different custom option group. After you remove the NNE option, you don't need to restart your DB instance. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Related Topics

- [Working with Option Groups \(p. 196\)](#)
- [Options for Oracle DB Instances \(p. 1004\)](#)

Oracle Secure Sockets Layer

You enable Secure Sockets Layer (SSL) encryption for an Oracle DB instance by adding the Oracle SSL option to the option group associated with an Oracle DB instance. You specify the port you want to communicate over using SSL. You must configure SQL*Plus as shown in this following section.

You enable SSL encryption for an Oracle DB instance by adding the Oracle SSL option to the option group associated with the DB instance. Amazon RDS uses a second port, as required by Oracle, for SSL connections. This approach allows both clear text and SSL-encrypted communication to occur at the same time between a DB instance and SQL*Plus. For example, you can use the port with clear text communication to communicate with other resources inside a VPC while using the port with SSL-encrypted communication to communicate with resources outside the VPC.

Note

You can use Secure Sockets Layer or Native Network Encryption, but not both. For more information, see [Oracle Native Network Encryption \(p. 1049\)](#).

You can use SSL encryption with the following Oracle database versions and editions:

- 19.0.0.0: All versions, all editions including Standard Edition Two
- 18.0.0.0: All versions, all editions including Standard Edition Two
- 12.2.0.1: All versions, all editions including Standard Edition Two
- 12.1.0.2: All versions, all editions including Standard Edition Two
- 11.2.0.4: All versions, Enterprise Edition
- 11.2.0.4: Version 6 and later, Standard Edition, Standard Edition One, Enterprise Edition

Note

You cannot use both SSL and Oracle native network encryption (NNE) on the same instance. If you use SSL encryption, you must disable any other connection encryption.

TLS Versions for the Oracle SSL Option

Amazon RDS for Oracle supports Transport Layer Security (TLS) versions 1.0 and 1.2. To use the Oracle SSL option, use the `SQLNET.SSL_VERSION` option setting. The following values are allowed for this option setting:

- "1.0" – Clients can connect to the DB instance using TLS 1.0 only.
- "1.2" – Clients can connect to the DB instance using TLS 1.2 only.
- "1.2 or 1.0" – Clients can connect to the DB instance using either TLS 1.2 or 1.0.

To use the Oracle SSL option, the `SQLNET.SSL_VERSION` option setting is also required:

- For existing Oracle SSL options, `SQLNET.SSL_VERSION` is set to "1.0" automatically. You can change the setting if necessary.
- When you add a new Oracle SSL option, you must set `SQLNET.SSL_VERSION` explicitly to a valid value.

The following table shows the TLS option settings that are supported for different Oracle engine versions and editions.

Oracle Engine Version	<code>SQLNET.SSL_VERSION = "1.0"</code>	<code>SQLNET.SSL_VERSION = "1.2"</code>	<code>SQLNET.SSL_VERSION = "1.2 or 1.0"</code>
18.0.0.0 (All editions)	Supported	Supported	Supported
12.2.0.1 (All editions)	Supported	Supported	Supported
12.1.0.2 (All editions)	Supported	Supported	Supported

Oracle Engine Version	<code>SQLNET.SSL_VERSION = "1.0"</code>	<code>SQLNET.SSL_VERSION = "1.2"</code>	<code>SQLNET.SSL_VERSION = "1.2 or 1.0"</code>
11.2.0.4 (Oracle EE)	Supported	Supported for 11.2.0.4.v8 and later	Supported for 11.2.0.4.v8 and later
11.2.0.4 (Oracle SE1)	Supported	Not supported	Not supported
11.2.0.4 (Oracle SE)	Supported	Not supported	Not supported

Cipher Suites for the Oracle SSL Option

Amazon RDS for Oracle supports multiple SSL cipher suites. By default, the Oracle SSL option is configured to use the `SSL_RSA_WITH_AES_256_CBC_SHA` cipher suite. To specify a different cipher suite to use over SSL connections, use the `SQLNET.CIPHER_SUITE` option setting. Following are the allowed values for this option setting:

- `"SSL_RSA_WITH_AES_256_CBC_SHA"` – The default setting, which is compatible with TLS 1.0 and TLS 1.2
- `"SSL_RSA_WITH_AES_256_CBC_SHA256"` – Only compatible with TLS 1.2
- `"SSL_RSA_WITH_AES_256_GCM_SHA384"` – Only compatible with TLS 1.2

For existing Oracle SSL options, `SQLNET.CIPHER_SUITE` is set to `"SSL_RSA_WITH_AES_256_CBC_SHA"` automatically. You can change the setting if necessary.

The following table shows the cipher suite option settings that are supported for different Oracle engine versions and editions.

Oracle Engine Version	<code>SQLNET.CIPHER_SUITE = "SSL_RSA_WITH_AES_256_CBC_SHA"</code>	<code>SQLNET.CIPHER_SUITE = "SSL_RSA_WITH_AES_256_CBC_SHA256"</code>	<code>SQLNET.CIPHER_SUITE = "SSL_RSA_WITH_AES_256_GCM_SHA384"</code>
19.0.0.0 (All editions)	Supported	Supported	Supported
18.0.0.0 (All editions)	Supported	Supported	Supported
12.2.0.1 (All editions)	Supported	Supported	Supported
12.1.0.2 (All editions)	Supported	Supported	Supported
11.2.0.4 (Oracle EE)	Supported	Not supported	Not supported
11.2.0.4 (Oracle SE1)	Supported	Not supported	Not supported
11.2.0.4 (Oracle SE)	Supported	Not supported	Not supported

Adding the SSL Option

To use SSL, your Amazon RDS Oracle DB instance must be associated with an option group that includes the `SSL` option.

Console

To add the SSL option to an option group

1. Create a new option group or identify an existing option group to which you can add the `SSL` option.

For information about creating an option group, see [Creating an Option Group \(p. 197\)](#).

2. Add the `SSL` option to the option group.

For information about adding an option to an option group, see [Adding an Option to an Option Group \(p. 200\)](#).

3. Create a new Oracle DB instance and associate the option group with it, or modify an Oracle DB instance to associate the option group with it.

For information about creating an Oracle DB instance, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).

For information about modifying an Oracle DB instance, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

AWS CLI

To add the `SSL` option to an option group

1. Create a new option group or identify an existing option group to which you can add the `SSL` option.

For information about creating an option group, see [Creating an Option Group \(p. 197\)](#).

2. Add the `SSL` option to the option group.

Specify the following option settings:

- `Port` – The SSL port number
- `VpcSecurityGroupMemberships` – The VPC security group for which the option is enabled
- `SQLNET.SSL_VERSION` – The TLS version that client can use to connect to the DB instance

For example, the following AWS CLI command adds the `SSL` option to an option group named `ora-option-group`.

Example

For Linux, OS X, or Unix:

```
aws rds add-option-to-option-group --option-group-name ora-option-group \
--options
'OptionName=SSL,Port=2484,VpcSecurityGroupMemberships="sg-68184619",OptionSettings=[{Name=SQLNET.SSL_VERSI...]
```

For Windows:

```
aws rds add-option-to-option-group --option-group-name ora-option-group ^
--options
'OptionName=SSL,Port=2484,VpcSecurityGroupMemberships="sg-68184619",OptionSettings=[{Name=SQLNET.SSL_VERSI...]
```

3. Create a new Oracle DB instance and associate the option group with it, or modify an Oracle DB instance to associate the option group with it.

For information about creating an Oracle DB instance, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).

For information about modifying an Oracle DB instance, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Configuring SQL*Plus to Use SSL with an Oracle DB Instance

You must configure SQL*Plus before connecting to an Oracle DB instance that uses the Oracle SSL option.

Note

To allow access to the DB instance from the appropriate clients, ensure that your security groups are configured correctly. For more information, see [Controlling Access with Security Groups \(p. 486\)](#). Also, these instructions are for SQL*Plus and other clients that directly use an Oracle home. For JDBC connections, see [Setting Up an SSL Connection Over JDBC \(p. 1057\)](#).

To configure SQL*Plus to use SSL to connect to an Oracle DB instance

1. Set the ORACLE_HOME environment variable to the location of your Oracle home directory.

The path to your Oracle home directory depends on your installation. The following example sets the ORACLE_HOME environment variable.

```
prompt>export ORACLE_HOME=/home/user/app/user/product/12.1.0/dbhome_1
```

For information about setting Oracle environment variables, see [SQL*Plus Environment Variables](#) in the Oracle documentation, and also see the Oracle installation guide for your operating system.

2. Append \$ORACLE_HOME/lib to the LD_LIBRARY_PATH environment variable.

The following is an example that sets the LD_LIBRARY_PATH environment variable.

```
prompt>export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$ORACLE_HOME/lib
```

3. Create a directory for the Oracle wallet at \$ORACLE_HOME/ssl_wallet.

The following is an example that creates the Oracle wallet directory.

```
prompt>mkdir $ORACLE_HOME/ssl_wallet
```

4. Download the root certificate that works for all AWS Regions and put the file in the ssl_wallet directory.

For information about downloading the root certificate, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).

5. In the \$ORACLE_HOME/network/admin directory, modify or create the tnsnames.ora file and include the following entry.

```
<net_service_name>= (DESCRIPTION = (ADDRESS_LIST = (ADDRESS = (PROTOCOL = TCPS)
(HOST = <endpoint>) (PORT = <ssl port number>)))(CONNECT_DATA = (SID = <database
name>))
```

```
(SECURITY = (SSL_SERVER_CERT_DN =
"C=US,ST=Washington,L=Seattle,O=Amazon.com,OU=RDS,CN=<endpoint>"))
```

- In the same directory, modify or create the sqlnet.ora file and include the following parameters.

Note

To communicate with entities over a TLS secured connection, Oracle requires a wallet with the necessary certificates for authentication. You can use Oracle's ORAPKI utility to create and maintain Oracle wallets, as shown in step 7. For more information, see [Setting Up Oracle Wallet Using ORAPKI](#) in the Oracle documentation.

```
WALLET_LOCATION = (SOURCE = (METHOD = FILE) (METHOD_DATA = (DIRECTORY = $ORACLE_HOME/
ssl_wallet)))
SSL_CLIENT_AUTHENTICATION = FALSE
SSL_VERSION = 1.0
SSL_CIPHER_SUITES = (SSL_RSA_WITH_AES_256_CBC_SHA)
SSL_SERVER_DN_MATCH = ON
```

Note

You can set SSL_VERSION to a higher value if your DB instance supports it.

- Run the following commands to create the Oracle wallet.

```
prompt>orapki wallet create -wallet $ORACLE_HOME/ssl_wallet -auto_login_only
prompt>orapki wallet add -wallet $ORACLE_HOME/ssl_wallet -trusted_cert -cert
$ORACLE_HOME/ssl_wallet/rds-ca-2019-root.pem -auto_login_only
```

Replace the file name with the one you downloaded.

Connecting to an Oracle DB Instance Using SSL

After you configure SQL*Plus to use SSL as described previously, you can connect to the Oracle DB instance with the SSL option. Optionally, you can first export the TNS_ADMIN value that points to the directory that contains the tnsnames.ora and sqlnet.ora files. Doing so ensures that SQL*Plus can find these files consistently. The following example exports the TNS_ADMIN value.

```
export TNS_ADMIN = ${ORACLE_HOME}/network/admin
```

Connect to the DB instance. For example, you can connect using SQL*Plus and a [`<net_service_name>`](#) in a tnsnames.ora file.

```
sqlplus <mydbuser>@<net_service_name>
```

You can also connect to the DB instance using SQL*Plus without using a tnsnames.ora file by using the following command.

```
sqlplus '<mydbuser>@(DESCRIPTION = (ADDRESS = (PROTOCOL = TCPS)(HOST = <endpoint>) (PORT = <ssl port number>))(CONNECT_DATA = (SID = <database name>)))'
```

You can also connect to the Oracle DB instance without using SSL. For example, the following command connects to the DB instance through the clear text port without SSL encryption.

```
sqlplus '<mydbuser>@(DESCRIPTION = (ADDRESS = (PROTOCOL = TCP)(HOST = <endpoint>) (PORT = <port number>))(CONNECT_DATA = (SID = <database name>)))'
```

If you want to close Transmission Control Protocol (TCP) port access, create a security group with no IP address ingresses and add it to the instance. This addition closes connections over the TCP port, while still allowing connections over the SSL port that are specified from IP addresses within the range permitted by the SSL option security group.

Setting Up an SSL Connection Over JDBC

To use an SSL connection over JDBC, you must create a keystore, trust the Amazon RDS root CA certificate, and use the code snippet specified following.

To create the keystore in JKS format, use the following command. For more information about creating the keystore, see the [Oracle documentation](#).

```
keytool -keystore clientkeystore -genkey -alias client
```

Next, take the following steps to trust the Amazon RDS root CA certificate.

To trust the Amazon RDS root CA certificate

1. Download the root certificate that works for all AWS Regions and put the file in the ssl_wallet directory.

For information about downloading the root certificate, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).

2. Convert the certificate to .der format using the following command.

```
openssl x509 -outform der -in rds-ca-2019-root.pem -out rds-ca-2019-root.der
```

Replace the file name with the one you downloaded.

3. Import the certificate into the keystore using the following command.

```
keytool -import -alias rds-root -keystore clientkeystore -file rds-ca-2019-root.der
```

4. Confirm that the key store was created successfully.

```
keytool -list -v -keystore clientkeystore.jks
```

Enter the keystore password when you are prompted for it.

The following code example shows how to set up the SSL connection using JDBC.

```
import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.SQLException;
import java.util.Properties;

public class OracleSslConnectionTest {
    private static final String DB_SERVER_NAME = "<dns-name-provided-by-amazon-rds>";
    private static final Integer SSL_PORT = "<ssl-option-port-configured-in-option-group>";
    private static final String DB_SID = "<oracle-sid>";
    private static final String DB_USER = "<user name>";
    private static final String DB_PASSWORD = "<password>";
    // This key store has only the prod root ca.
    private static final String KEY_STORE_FILE_PATH = "<file-path-to-keystore>";
    private static final String KEY_STORE_PASS = "<keystore-password>";

    public static void main(String[] args) throws SQLException {
        final Properties properties = new Properties();
        final String connectionString = String.format(
            "jdbc:oracle:thin:@(DESCRIPTION=(ADDRESS=(PROTOCOL=TCPS)(HOST=%s)(PORT=%d))" +
            "(CONNECT_DATA=(SID=%s)))",
            DB_SERVER_NAME, SSL_PORT, DB_SID);
        properties.put("user", DB_USER);
        properties.put("password", DB_PASSWORD);
        properties.put("oracle.jdbc.J2EE13Compliant", "true");
        properties.put("javax.net.ssl.trustStore", KEY_STORE_FILE_PATH);
        properties.put("javax.net.ssl.trustStoreType", "JKS");
        properties.put("javax.net.ssl.trustStorePassword", KEY_STORE_PASS);
        final Connection connection = DriverManager.getConnection(connectionString,
        properties);
        // If no exception, that means handshake has passed, and an SSL connection can be
        opened
    }
}
```

Enforcing a DN Match with an SSL Connection

You can use the Oracle parameter `SSL_SERVER_DN_MATCH` to enforce that the distinguished name (DN) for the database server matches its service name. If you enforce the match verifications, then SSL ensures that the certificate is from the server. If you don't enforce the match verification, then SSL performs the check but allows the connection, regardless if there is a match. If you do not enforce the match, you allow the server to potentially fake its identify.

To enforce DN matching, add the DN match property and use the connection string specified below.

Add the property to the client connection to enforce DN matching.

```
properties.put("oracle.net.ssl_server_dn_match", "TRUE");
```

Use the following connection string to enforce DN matching when using SSL.

```
final String connectionString = String.format(
    "jdbc:oracle:thin:@(DESCRIPTION=(ADDRESS=(PROTOCOL=TCPS)(HOST=%s)(PORT=%d))" +
    "(CONNECT_DATA=(SID=%s))" +
    "(SECURITY = (SSL_SERVER_CERT_DN =
\"C=US,ST=Washington,L=Seattle,O=Amazon.com,OU=RDS,CN=%s\")))",
    DB_SERVER_NAME, SSL_PORT, DB_SID, DB_SERVER_NAME);
```

Oracle Spatial

Amazon RDS supports Oracle Spatial through the use of the `SPATIAL` option. Oracle Spatial provides a SQL schema and functions that facilitate the storage, retrieval, update, and query of collections of spatial data in an Oracle database. For more information, see [Spatial Concepts](#) in the Oracle documentation.

Important

If you use Oracle Spatial, Amazon RDS automatically updates your DB instance to the latest Oracle PSU if there are security vulnerabilities with a Common Vulnerability Scoring System (CVSS) score of 9+ or other announced security vulnerabilities.

Amazon RDS supports Oracle Spatial for the following editions and versions of Oracle:

- Oracle Enterprise Edition, version 19.0.0.0, all versions
- Oracle Enterprise Edition, version 18.0.0.0, all versions
- Oracle Enterprise Edition, version 12.2.0.1, all versions
- Oracle Enterprise Edition, version 12.1.0.2.v13 or later
- Oracle Enterprise Edition, version 11.2.0.4.v17 or later

Prerequisites for Oracle Spatial

The following are prerequisites for using Oracle Spatial:

- Your DB instance must be inside a virtual private cloud (VPC). For more information, see [Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform \(p. 506\)](#).
- Your DB instance must be of sufficient class. Oracle Spatial is not supported for the db.m1.small, db.t2.micro, or db.t2.small DB instance classes. For more information, see [DB Instance Class Support for Oracle \(p. 917\)](#).
- Your DB instance must have **Auto Minor Version Upgrade** enabled. This option enables your DB instance to receive minor DB engine version upgrades automatically when they become available. Amazon RDS uses this option to update your DB instance to the latest Oracle Patch Set Update (PSU) or Release Update (RU). For more information, see [Settings for Oracle DB Instances \(p. 967\)](#).
- If your DB instance is running on major version 11.2, you must install the `XMLDB` option. For more information, see [Oracle XML DB \(p. 1078\)](#).
- An Oracle Spatial license from Oracle. For more information, see [Oracle Spatial and Graph](#) in the Oracle documentation.

Best Practices for Oracle Spatial

The following are best practices for using Oracle Spatial:

- For maximum security, use the `SPATIAL` option with Secure Sockets Layer (SSL). For more information, see [Oracle Secure Sockets Layer \(p. 1051\)](#).
- Configure your DB instance to restrict access to your DB instance. For more information, see [Scenarios for Accessing a DB Instance in a VPC \(p. 507\)](#) and [Working with a DB Instance in a VPC \(p. 514\)](#).

Adding the Oracle Spatial Option

The following is the general process for adding the `SPATIAL` option to a DB instance:

1. Create a new option group, or copy or modify an existing option group.

2. Add the option to the option group.
3. Associate the option group with the DB instance.

There is a brief outage while the **SPATIAL** option is added. After you add the option, you don't need to restart your DB instance. As soon as the option group is active, Oracle Spatial is available.

Note

During this outage, password verification functions are disabled briefly. You can also expect to see events related to password verification functions during the outage. Password verification functions are enabled again before the Oracle DB instance is available.

To add the **SPATIAL** option to a DB instance

1. Determine the option group that you want to use. You can create a new option group or use an existing option group. If you want to use an existing option group, skip to the next step. Otherwise, create a custom DB option group with the following settings:
 - a. For **Engine**, choose **oracle-ee**.
 - b. For **Major engine version**, choose the version of your DB instance.

For more information, see [Creating an Option Group \(p. 197\)](#).
2. Add the **SPATIAL** option to the option group. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#).
3. Apply the option group to a new or existing DB instance:
 - For a new DB instance, you apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).
 - For an existing DB instance, you apply the option group by modifying the instance and attaching the new option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Removing the Oracle Spatial Option

You can remove the **SPATIAL** option from a DB instance. There is a brief outage while the option is removed. After you remove the **SPATIAL** option, you don't need to restart your DB instance.

Warning

Removing the **SPATIAL** option can result in data loss if the DB instance is using data types that were enabled as part of the option. Back up your data before proceeding. For more information, see [Backing Up and Restoring Amazon RDS DB Instances \(p. 253\)](#).

To remove the **SPATIAL** option from a DB instance, do one of the following:

- Remove the **SPATIAL** option from the option group it belongs to. This change affects all DB instances that use the option group. For more information, see [Removing an Option from an Option Group \(p. 207\)](#).
- Modify the DB instance and specify a different option group that doesn't include the **SPATIAL** option. This change affects a single DB instance. You can specify the default (empty) option group, or a different custom option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Related Topics

- [Oracle Locator \(p. 1043\)](#)

- [Options for Oracle DB Instances \(p. 1004\)](#)
- [Working with Option Groups \(p. 196\)](#)

Oracle SQLT

Amazon RDS supports Oracle SQLTXPLAIN (SQLT) through the use of the SQLT option.

The Oracle `EXPLAIN PLAN` statement can determine the execution plan of a SQL statement. It can verify whether the Oracle optimizer chooses a certain execution plan, such as a nested loops join. It also helps you understand the optimizer's decisions, such as why it chose a nested loops join over a hash join. So `EXPLAIN PLAN` helps you understand the statement's performance.

SQLT is an Oracle utility that produces a report. The report includes object statistics, object metadata, optimizer-related initialization parameters, and other information that a database administrator can use to tune a SQL statement for optimal performance. SQLT produces an HTML report with hyperlinks to all of the sections in the report.

Unlike Automatic Workload Repository or Statspack reports, SQLT works on individual SQL statements. SQLT is a collection of SQL, PL/SQL, and SQL*Plus files that collect, store, and display performance data.

Following are the supported Oracle versions for each SQLT version.

SQLT Version	Oracle 19c	Oracle 18c	Oracle 12c version 12.2	Oracle 12c version 12.1	Oracle 11g
12.2.180725	Supported	Supported	Supported	Supported	Supported
12.2.180331	Not supported	Supported	Supported	Supported	Supported
12.1.160429	Not supported	Not supported	Supported	Supported	Supported

To download SQLT and access instructions for using it:

- Log in to your My Oracle Support account, and open the following documents:
 - To download SQLT: [Document 215187.1](#)
 - For SQLT usage instructions: [Document 1614107.1](#)
 - For frequently asked questions about SQLT: [Document 1454160.1](#)
 - For information about reading SQLT output: [Document 1456176.1](#)
 - For interpreting the Main report: [Document 1922234.1](#)

You can use SQLT with any edition of the following Oracle Database versions:

- Oracle 19c, 19.0.0.0
- Oracle 18c, 18.0.0.0
- Oracle 12c, 12.2.0.1
- Oracle 12c, 12.1.0.2
- Oracle 11g, 11.2.0.4

Amazon RDS does not support the following SQLT methods:

- `XPORE`
- `XHUME`

Prerequisites for SQLT

The following are prerequisites for using SQLT:

- You must remove users and roles that are required by SQLT, if they exist.

The SQLT option creates the following users and roles on a DB instance:

- `SQLTXPLAIN` user
- `SQLTXADMIN` user
- `SQLT_USER_ROLE` role

If your DB instance has any of these users or roles, log in to the DB instance using a SQL client, and drop them using the following statements:

```
DROP USER SQLTXPLAIN CASCADE;
DROP USER SQLTXADMIN CASCADE;
DROP ROLE SQLT_USER_ROLE CASCADE;
```

- You must remove tablespaces that are required by SQLT, if they exist.

The SQLT option creates the following tablespaces on a DB instance:

- `RDS_SQLT_TS`
- `RDS_TEMP_SQLT_TS`

If your DB instance has these tablespaces, log in to the DB instance using a SQL client, and drop them.

SQLT Option Settings

SQLT can work with licensed features that are provided by the Oracle Tuning Pack and the Oracle Diagnostics Pack. The Oracle Tuning Pack includes the SQL Tuning Advisor, and the Oracle Diagnostics Pack includes the Automatic Workload Repository. The SQLT settings enable or disable access to these features from SQLT.

Amazon RDS supports the following settings for the SQLT option.

Option Setting	Valid Values	Default Value	Description
<code>LICENSE_PACK</code> , T, D, N	T		<p>The Oracle Management Packs that you want to access with SQLT. Enter one of the following values:</p> <ul style="list-style-type: none"> • T indicates that you have a license for the Oracle Tuning Pack and the Oracle Diagnostics Pack, and you want to access the SQL Tuning Advisor and Automatic Workload Repository from SQLT. • D indicates that you have a license for the Oracle Diagnostics Pack, and you want to access the Automatic Workload Repository from SQLT. • N indicates that you don't have a license for the Oracle Tuning Pack and the Oracle Diagnostics Pack, or that you have a license for one or both of them, but you don't want SQLT to access them.

Option Setting	Valid Values	Default Value	Description
			<p>Note Amazon RDS does not provide licenses for these Oracle Management Packs. If you indicate that you want to use a pack that is not included in your DB instance, you can use SQLT with the DB instance. However, SQLT can't access the pack, and the SQLT report doesn't include the data for the pack. For example, if you specify <code>T</code>, but the DB instance doesn't include the Oracle Tuning Pack, SQLT works on the DB instance, but the report it generates doesn't contain data related to the Oracle Tuning Pack.</p>
VERSION	2016-04-29 .v1, 2018-03-31 .v1, 2018-07-25 .v1	2016-04-29 .v1	<p>The version of SQLT that you want to install.</p> <p>Note For Oracle version 19.0.0.0, the only supported version is 2018-07-25 .v1. This version is also the default for Oracle version 19.0.0.0.</p>

Adding the SQLT Option

The following is the general process for adding the SQLT option to a DB instance:

1. Create a new option group, or copy or modify an existing option group.
2. Add the SQLT option to the option group.
3. Associate the option group with the DB instance.

After you add the SQLT option, as soon as the option group is active, SQLT is active.

To add the SQLT option to a DB instance

1. Determine the option group that you want to use. You can create a new option group or use an existing option group. If you want to use an existing option group, skip to the next step. Otherwise, create a custom DB option group with the following settings:
 - a. For **Engine**, choose the Oracle edition that you want to use. The SQLT option is supported on all editions.
 - b. For **Major engine version**, choose the version of your DB instance.

For more information, see [Creating an Option Group \(p. 197\)](#).
2. Add the **SQLT** option to the option group. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#).
3. Apply the option group to a new or existing DB instance:
 - For a new DB instance, you apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).
 - For an existing DB instance, you apply the option group by modifying the instance and attaching the new option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).
4. (Optional) Verify the SQLT installation on each DB instance with the SQLT option.

- a. Use a SQL client to connect to the DB instance as the master user.

For information about connecting to an Oracle DB instance using a SQL client, see [Connecting to a DB Instance Running the Oracle Database Engine \(p. 955\)](#).

- b. Run the following query:

```
SELECT sqltxplain.sqlt$a.get_param('tool_version') sqlt_version FROM DUAL;
```

The query returns the current version of the SQLT option on Amazon RDS. 12.1.160429 is an example of a version of SQLT that is available on Amazon RDS.

5. Change the passwords of the users that are created by the SQLT option.

- a. Use a SQL client to connect to the DB instance as the master user.
- b. Run the following SQL statement to change the password for the SQLTXADMIN user:

```
ALTER USER SQLTXADMIN IDENTIFIED BY new_password ACCOUNT UNLOCK;
```

- c. Run the following SQL statement to change the password for the SQLTXPLAIN user:

```
ALTER USER SQLTXPLAIN IDENTIFIED BY new_password ACCOUNT UNLOCK;
```

Note

Upgrading SQLT requires uninstalling an older version of SQLT and then installing the new version. So, all SQLT metadata can be lost when you upgrade SQLT. A major version upgrade of a database also uninstalls and re-installs SQLT. An example of a major version upgrade is an upgrade from Oracle 11g to Oracle 12c.

Using SQLT

SQLT works with the Oracle SQL*Plus utility.

To use SQLT

1. Download the SQLT .zip file from [Document 215187.1](#) on the My Oracle Support site.

Note

You can't download SQLT 12.1.160429 from the My Oracle Support site. Oracle has deprecated this older version.

2. Unzip the SQLT .zip file.
3. From a command prompt, change to the `sqlt/run` directory on your file system.
4. From the command prompt, open SQL*Plus, and connect to the DB instance as the master user.

For information about connecting to a DB instance using SQL*Plus, see [Connecting to a DB Instance Running the Oracle Database Engine \(p. 955\)](#).

5. Get the SQL ID of a SQL statement:

```
SELECT SQL_ID FROM V$SQL WHERE SQL_TEXT='sql_statement';
```

Your output is similar to the following:

```
SQL_ID
-----
chvsmttqjzjkn
```

6. Analyze a SQL statement with SQLT:

```
START sqlxtract.sql sql_id sqlxplain_user_password
```

For example, for the SQL ID *chvsmttqjzjkn*, enter the following:

```
START sqlxtract.sql chvsmttqjzjkn sqlxplain_user_password
```

SQLT generates the HTML report and related resources as a .zip file in the directory from which the SQLT command was run.

7. (Optional) To enable application users to diagnose SQL statements with SQLT, grant *SQLT_USER_ROLE* to each application user with the following statement:

```
GRANT ROLE SQLT_USER_ROLE TO application_user_name;
```

Note

Oracle does not recommend running SQLT with the *SYS* user or with users that have the *DBA* role. It is a best practice to run SQLT diagnostics using the application user's account, by granting *SQLT_USER_ROLE* to the application user.

Upgrading the SQLT Option

With Amazon RDS for Oracle, you can upgrade the SQLT option from your existing version to a higher version. To upgrade the SQLT option, complete steps 1–3 in [Using SQLT \(p. 1066\)](#) for the new version of SQLT. Also, if you granted privileges for the previous version of SQLT in step 7 of that section, grant the privileges again for the new SQLT version.

Upgrading the SQLT option results in the loss of the older SQLT version's metadata. The older SQLT version's schema and related objects are dropped, and the newer version of SQLT is installed. For more information about the changes in the latest SQLT version, see [Document 1614201.1](#) on the My Oracle Support site.

Note

Version downgrades are not supported.

Modifying SQLT Settings

After you enable SQLT, you can modify the `LICENSE_PACK` and `VERSION` settings for the option.

For more information about how to modify option settings, see [Modifying an Option Setting \(p. 204\)](#). For more information about each setting, see [SQLT Option Settings \(p. 1064\)](#).

Removing the SQLT Option

You can remove SQLT from a DB instance.

To remove SQLT from a DB instance, do one of the following:

- To remove SQLT from multiple DB instances, remove the SQLT option from the option group to which the DB instances belong. This change affects all DB instances that use the option group. For more information, see [Removing an Option from an Option Group \(p. 207\)](#).
- To remove SQLT from a single DB instance, modify the DB instance and specify a different option group that doesn't include the SQLT option. You can specify the default (empty) option group or a different custom option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Related Topics

- [Working with Option Groups \(p. 196\)](#)
- [Options for Oracle DB Instances \(p. 1004\)](#)

Oracle Statspack

The Oracle Statspack option installs and enables the Oracle Statspack performance statistics feature. Oracle Statspack is a collection of SQL, PL/SQL, and SQL*Plus scripts that collect, store, and display performance data. For information about using Oracle Statspack, see [Oracle Statspack](#) in the Oracle documentation.

Note

Oracle Statspack is no longer supported by Oracle and has been replaced by the more advanced Automatic Workload Repository (AWR). AWR is available only for Oracle Enterprise Edition customers who have purchased the Diagnostics Pack. Oracle Statspack can be used with any Oracle DB engine on Amazon RDS.

The following steps show you how to work with Oracle Statspack on Amazon RDS:

1. If you have an existing DB instance that has the `PERFSTAT` account already created and you want to use Oracle Statspack with it, you must drop the `PERFSTAT` account before adding the Statspack option to the option group associated with your DB instance. If you attempt to add the Statspack option to an option group associated with a DB instance that already has the `PERFSTAT` account created, you get an error and the RDS event RDS-Event-0058 is generated.

If you have already installed Statspack, and the `PERFSTAT` account is associated with Statspack, then skip this step, and do not drop the `PERFSTAT` user.

You can drop the `PERFSTAT` account by running the following command:

```
DROP USER perfstat CASCADE;
```

2. Add the Statspack option to an option group and then associate that option group with your DB instance. Amazon RDS installs the Statspack scripts on the DB instance and then sets up the PERFSTAT user account, the account you use to run the Statspack scripts. If you have installed Statspack, skip this step.
3. After Amazon RDS has installed Statspack on your DB instance, you must log in to the DB instance using your master user name and master password. You must then reset the PERFSTAT password from the randomly generated value Amazon RDS created when Statspack was installed. After you have reset the PERFSTAT password, you can log in using the PERFSTAT user account and run the Statspack scripts.

Use the following command to reset the password:

```
ALTER USER perfstat IDENTIFIED BY <new_password> ACCOUNT UNLOCK;
```

4. After you have logged on using the PERFSTAT account, you can either manually create a Statspack snapshot or create a job that will take a Statspack snapshot after a given time interval. For example, the following job creates a Statspack snapshot every hour:

```
variable jn number;
execute dbms_job.submit(:jn, 'statspack.snap;', sysdate, 'trunc(SYSDATE+1/24, ''HH24'')");
commit;
```

5. Once you have created at least two Statspack snapshots, you can view them using the following query:

```
select snap_id, snap_time from stats$snapshot order by 1;
```

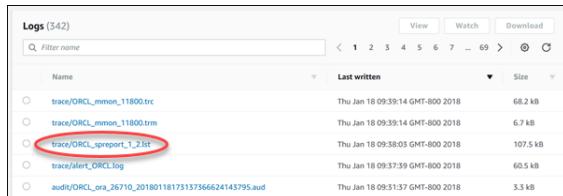
6. To create a Statspack report, you choose two snapshots to analyze and run the following Amazon RDS command:

```
exec RDSADMIN.RDS_RUN_SPREPORT(<begin snap>,<end snap>);
```

For example, the following Amazon RDS command would create a report based on the interval between Statspack snapshots 1 and 2:

```
exec RDSADMIN.RDS_RUN_SPREPORT(1,2);
```

The file name of the Statspack report that is generated includes the number of the two Statspack snapshots used. For example, a report file created using Statspack snapshots 1 and 2 would be named ORCL_spreport_1_2.lst. You can download the Statspack report by selecting the report in the **Log** section of the DB instance details on the RDS console and choosing **Download** or you can use the trace file procedures explained in [Working with Oracle Trace Files \(p. 413\)](#).



Name	Last written	Size
trace/ORCL_mmon_11800.trc	Thu Jan 18 09:39:14 GMT-800 2018	68.2 kB
trace/ORCL_mmon_11800.trm	Thu Jan 18 09:39:14 GMT-800 2018	6.7 kB
trace/ORCL_spreport_1_2.lst	Thu Jan 18 09:38:03 GMT-800 2018	107.5 kB
trace/alert_ORCL.log	Thu Jan 18 09:37:39 GMT-800 2018	60.5 kB
audit/ORCL ora_26710_20180118173157366624143795.aud	Thu Jan 18 09:51:37 GMT-800 2018	3.5 kB

If an error occurs when producing the report, an error file is created using the same naming conventions but with an extension of .err. For example, if an error occurred while creating a report using Statspack snapshots 1 and 7, the report file would be named ORCL_spreport_1_7.err. You can download the error report by selecting the report in the Log section of the RDS console and choosing **Download** or use the trace file procedures explained in [Working with Oracle Trace Files \(p. 413\)](#).

Oracle Statspack does some basic checking before running the report, so you could also see error messages displayed at the command prompt. For example, if you attempt to generate a report based on an invalid range, such as the beginning Statspack snapshot value is larger than the ending Statspack snapshot value, the error message is displayed at the command prompt and no error file is created.

```
exec RDSADMIN.RDS_RUN_SPREPORT(2,1);
*
ERROR at line 1:
ORA-20000: Invalid snapshot IDs. Find valid ones in perfstat.stats$snapshot.
```

If you use an invalid number for one of the Statspack snapshots, the error message will also be displayed at the command prompt. For example, if you have 20 Statspack snapshots but request that a report be run using Statspack snapshots 1 and 50, the command prompt will display an error.

```
exec RDSADMIN.RDS_RUN_SPREPORT(1,50);
*
ERROR at line 1:
ORA-20000: Could not find both snapshot IDs
```

For more information about how to use Oracle Statspack, including information on adjusting the amount of data captured by adjusting the snapshot level, go to the Oracle [Statspack documentation page](#).

To remove Oracle Statspack files, use the following command:

```
execute statspack.purge(<begin snap>, <end snap>);
```

Oracle Time Zone

You can use the time zone option to change the system time zone used by your Oracle DB instance. For example, you might change the time zone of a DB instance to be compatible with an on-premises environment, or a legacy application. The time zone option changes the time zone at the host level. Changing the time zone impacts all date columns and values, including SYSDATE and SYSTIMESTAMP.

The time zone option differs from the `rdsadmin_util.alter_db_time_zone` command. The `alter_db_time_zone` command changes the time zone only for certain data types. The time zone option changes the time zone for all date columns and values. For more information about `alter_db_time_zone`, see [Setting the Database Time Zone \(p. 1097\)](#).

Note

You can't share a DB snapshot that uses this option. For more information about sharing DB snapshots, see [Sharing a DB Snapshot \(p. 278\)](#).

Prerequisites for Time Zone

The time zone option is a permanent and persistent option. You can't remove the option from an option group after you add it. You can't remove the option group from a DB instance after you add it. You can't modify the time zone setting of the option to a different time zone.

We strongly urge you to take a DB snapshot of your DB instance before adding the time zone option to a DB instance. By using a snapshot you can recover the DB instance if you set the time zone option incorrectly. For more information, see [Creating a DB Snapshot \(p. 264\)](#).

We strongly urge you to test the time zone option on a test DB instance before you add it to a production DB instance. Adding the time zone option can cause problems with tables that use system date to add dates or times. You should analyze your data and applications to determine the impact of changing the time zone.

Time Zone Option Settings

Amazon RDS supports the following settings for the time zone option.

Option Setting	Valid Values	Description
<code>TIME_ZONE</code>	One of the available time zones. For the full list, see Available Time Zones (p. 1073) .	The new time zone for your DB instance.

Adding the Time Zone Option

The general process for adding the time zone option to a DB instance is the following:

1. Create a new option group, or copy or modify an existing option group.
2. Add the option to the option group.
3. Associate the option group with the DB instance.

When you add the time zone option, a brief outage occurs while your DB instance is automatically restarted.

Console

To add the time zone option to a DB instance

1. Determine the option group you want to use. You can create a new option group or use an existing option group. If you want to use an existing option group, skip to the next step. Otherwise, create a custom DB option group with the following settings:
 - a. For **Engine** choose the oracle edition for your DB instance.
 - b. For **Major engine version** choose the version of your DB instance.

For more information, see [Creating an Option Group \(p. 197\)](#).

2. Add the **Timezone** option to the option group, and configure the option settings.

Important

If you add the time zone option to an existing option group that is already attached to one or more DB instances, a brief outage occurs while all the DB instances are automatically restarted.

For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#). For more information about each setting, see [Time Zone Option Settings \(p. 1071\)](#).

3. Apply the option group to a new or existing DB instance:

- For a new DB instance, you apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).
- For an existing DB instance, you apply the option group by modifying the instance and attaching the new option group. When you add the time zone option to an existing DB instance, a brief outage occurs while your DB instance is automatically restarted. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

AWS CLI

The following example uses the AWS CLI `add-option-to-option-group` command to add the `Timezone` option and the `TIME_ZONE` option setting to an option group called `myoptiongroup`. The time zone is set to `Africa/Cairo`.

For Linux, OS X, or Unix:

```
aws rds add-option-to-option-group \
--option-group-name "myoptiongroup" \
--options "OptionName=Timezone,OptionSettings=[{Name=TIME_ZONE,Value=Africa/Cairo}]" \
--apply-immediately
```

For Windows:

```
aws rds add-option-to-option-group ^
--option-group-name "myoptiongroup" ^
--options "OptionName=Timezone,OptionSettings=[{Name=TIME_ZONE,Value=Africa/Cairo}]" ^
--apply-immediately
```

Modifying Time Zone Settings

The time zone option is a permanent and persistent option. You can't remove the option from an option group after you add it. You can't remove the option group from a DB instance after you add it. You can't

modify the time zone setting of the option to a different time zone. If you set the time zone incorrectly, restore a snapshot of your DB instance from before you added the time zone option.

Removing the Time Zone Option

The time zone option is a permanent and persistent option. You can't remove the option from an option group after you add it. You can't remove the option group from a DB instance after you add it. To remove the time zone option, restore a snapshot of your DB instance from before you added the time zone option.

Available Time Zones

The following values can be used for the time zone option.

Zone	Time Zone
Africa	Africa/Cairo, Africa/Casablanca, Africa/Harare, Africa/Lagos, Africa/Luanda, Africa/Monrovia, Africa/Nairobi, Africa/Tripoli, Africa/Windhoek
America	America/Araguaina, America/Argentina/Buenos_Aires, America/Asuncion, America/Bogota, America/Caracas, America/Chicago, America/Chihuahua, America/Cuiaba, America/Denver, America/Detroit, America/Fortaleza, America/Godthab, America/Guatemala, America/Halifax, America/Lima, America/Los_Angeles, America/Manaus, America/Matamoros, America/Mexico_City, America/Monterrey, America/Montevideo, America/New_York, America/Phoenix, America/Santiago, America/Sao_Paulo, America/Tijuana, America/Toronto
Asia	Asia/Amman, Asia/Ashgabat, Asia/Baghdad, Asia/Baku, Asia/Bangkok, Asia/Beirut, Asia/Calcutta, Asia/Damascus, Asia/Dhaka, Asia/Hong_Kong, Asia/Irkutsk, Asia/Jakarta, Asia/Jerusalem, Asia/Kabul, Asia/Karachi, Asia/Kathmandu, Asia/Kolkata, Asia/Krasnoyarsk, Asia/Magadan, Asia/Manila, Asia/Muscat, Asia/Novosibirsk, Asia/Rangoon, Asia/Riyadh, Asia/Seoul, Asia/Shanghai, Asia/Singapore, Asia/Taipei, Asia/Tehran, Asia/Tokyo, Asia/Ulaanbaatar, Asia/Vladivostok, Asia/Yakutsk, Asia/Yerevan
Atlantic	Atlantic/Azores, Atlantic/Cape_Verde
Australia	Australia/Adelaide, Australia/Brisbane, Australia/Darwin, Australia/Eucla, Australia/Hobart, Australia/Lord_Howe, Australia/Perth, Australia/Sydney
Brazil	Brazil/DeNoronha, Brazil/East
Canada	Canada/Newfoundland, Canada/Saskatchewan
Etc	Etc/GMT-3
Europe	Europe/Amsterdam, Europe/Athens, Europe/Berlin, Europe/Dublin, Europe/Helsinki, Europe/Kaliningrad, Europe/London, Europe/Madrid, Europe/Moscow, Europe/Paris, Europe/Prague, Europe/Rome, Europe/Sarajevo
Pacific	Pacific/Apia, Pacific/Auckland, Pacific/Chatham, Pacific/Fiji, Pacific/Guam, Pacific/Honolulu, Pacific/Kiritimati, Pacific/Marquesas, Pacific/Samoa, Pacific/Tongatapu, Pacific/Wake
US	US/Alaska, US/Central, US/East-Indiana, US/Eastern, US/Pacific
UTC	UTC

Related Topics

- [Working with Option Groups \(p. 196\)](#)
- [Options for Oracle DB Instances \(p. 1004\)](#)

Oracle Transparent Data Encryption

Amazon RDS supports Oracle Transparent Data Encryption (TDE), a feature of the Oracle Advanced Security option available in Oracle Enterprise Edition. This feature automatically encrypts data before it is written to storage and automatically decrypts data when the data is read from storage.

Oracle Transparent Data Encryption is used in scenarios where you need to encrypt sensitive data in case data files and backups are obtained by a third party or when you need to address security-related regulatory compliance issues.

The TDE option is a permanent option that can't be removed from an option group. You can't disable TDE from a DB instance once that instance is associated with an option group with the Oracle TDE option. You can change the option group of a DB instance that is using the TDE option, but the option group associated with the DB instance must include the TDE option. You can also modify an option group that includes the TDE option by adding or removing other options.

A detailed explanation about Oracle Transparent Data Encryption is beyond the scope of this guide. For information about using Oracle Transparent Data Encryption, see [Securing Stored Data Using Transparent Data Encryption](#). For more information about Oracle Advanced Security, see [Oracle Advanced Security](#) in the Oracle documentation. For more information on AWS security, see the [AWS Security Center](#).

Note

You can't share a DB snapshot that uses this option. For more information about sharing DB snapshots, see [Sharing a DB Snapshot \(p. 278\)](#).

TDE Encryption Modes

Oracle Transparent Data Encryption supports two encryption modes: TDE tablespace encryption and TDE column encryption. TDE tablespace encryption is used to encrypt entire application tables. TDE column encryption is used to encrypt individual data elements that contain sensitive data. You can also apply a hybrid encryption solution that uses both TDE tablespace and column encryption.

Note

Amazon RDS manages the Oracle Wallet and TDE master key for the DB instance. You do not need to set the encryption key using the command `ALTER SYSTEM set encryption key`.

For information about TDE best practices, see [Oracle Advanced Security Transparent Data Encryption Best Practices](#).

Once the option is enabled, you can check the status of the Oracle Wallet by using the following command:

```
SELECT * FROM v$encryption_wallet;
```

To create an encrypted tablespace, use the following command:

```
CREATE TABLESPACE encrypt_ts ENCRYPTION DEFAULT STORAGE (ENCRYPT);
```

To specify the encryption algorithm, use the following command:

```
CREATE TABLESPACE encrypt_ts ENCRYPTION USING 'AES256' DEFAULT STORAGE (ENCRYPT);
```

Note that the previous commands for encrypting a tablespace are the same as the commands you would use with an Oracle installation not on Amazon RDS, and the ALTER TABLE syntax to encrypt a column is also the same as the commands you would use for an Oracle installation not on Amazon RDS.

You should determine if your DB instance is associated with an option group that has the **TDE** option. To view the option group that a DB instance is associated with, you can use the RDS console, the [describe-db-instance](#) AWS CLI command, or the API operation [DescribeDBInstances](#).

To comply with several security standards, Amazon RDS is working to implement automatic periodic master key rotation.

Adding the TDE Option

The process for using Oracle Transparent Data Encryption (TDE) with Amazon RDS is as follows:

1. If the DB instance is not associated with an option group that has the **TDE** option enabled, you must either create an option group and add the **TDE** option or modify the associated option group to add the **TDE** option. For information about creating or modifying an option group, see [Working with Option Groups \(p. 196\)](#). For information about adding an option to an option group, see [Adding an Option to an Option Group \(p. 200\)](#).
2. Associate the DB instance with the option group with the **TDE** option. For information about associating a DB instance with an option group, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Removing the TDE Option

If you no longer want to use the TDE option with a DB instance, you must decrypt all your data on the DB instance, copy the data to a new DB instance that is not associated with an option group with TDE enabled, and then delete the original instance. You can rename the new instance to be the same name as the previous DB instance if you prefer.

Using TDE with Data Pump

You can use Oracle Data Pump to import or export encrypted dump files. Amazon RDS supports the password encryption mode (ENCRYPTION_MODE=PASSWORD) for Oracle Data Pump. Amazon RDS does not support transparent encryption mode (ENCRYPTION_MODE=TRANSPARENT) for Oracle Data Pump. For more information about using Oracle Data Pump with Amazon RDS, see [Importing Using Oracle Data Pump \(p. 987\)](#).

Related Topics

- [Working with Option Groups \(p. 196\)](#)
- [Options for Oracle DB Instances \(p. 1004\)](#)

Oracle UTL_MAIL

Amazon RDS supports Oracle UTL_MAIL through the use of the UTL_MAIL option and SMTP servers. You can send email directly from your database by using the UTL_MAIL package. Amazon RDS supports UTL_MAIL for the following versions of Oracle:

- Oracle version 19.0.0.0, all versions
- Oracle version 18.0.0.0, all versions
- Oracle version 12.2.0.1, all versions
- Oracle version 12.1.0.2.v5 and later
- Oracle version 11.2.0.4.v9 and later

The following are some limitations to using UTL_MAIL:

- UTL_MAIL does not support Transport Layer Security (TLS) and therefore emails are not encrypted.

To connect securely to remote SSL/TLS resources by creating and uploading custom Oracle wallets, follow the instructions in [Using utl_http, utl_tcp, and utl_smtp with an Oracle DB Instance \(p. 938\)](#).

The specific certificates that are required for your wallet vary by service. For AWS services, these can typically be found in the [Amazon Trust Services Repository](#).

- UTL_MAIL does not support authentication with SMTP servers.
- You can only send a single attachment in an email.
- You can't send attachments larger than 32 K.
- You can only use ASCII and Extended Binary Coded Decimal Interchange Code (EBCDIC) character encodings.
- SMTP port (25) is throttled based on the elastic network interface owner's policies.

When you enable UTL_MAIL, only the master user for your DB instance is granted the execute privilege. If necessary, the master user can grant the execute privilege to other users so that they can use UTL_MAIL.

Important

We recommend that you enable Oracle's built-in auditing feature to track the use of UTL_MAIL procedures.

Prerequisites for Oracle UTL_MAIL

The following are prerequisites for using Oracle UTL_MAIL:

- One or more SMTP servers, and the corresponding IP addresses or public or private Domain Name Server (DNS) names. For more information about private DNS names resolved through a custom DNS server, see [Setting Up a Custom DNS Server \(p. 1094\)](#).
- For Oracle versions prior to 12c, your DB instance must also use the XML DB option. For more information, see [Oracle XML DB \(p. 1078\)](#).

Adding the Oracle UTL_MAIL Option

The general process for adding the Oracle UTL_MAIL option to a DB instance is the following:

1. Create a new option group, or copy or modify an existing option group.
2. Add the option to the option group.

3. Associate the option group with the DB instance.

After you add the UTL_MAIL option, as soon as the option group is active, UTL_MAIL is active.

To add the UTL_MAIL option to a DB instance

1. Determine the option group you want to use. You can create a new option group or use an existing option group. If you want to use an existing option group, skip to the next step. Otherwise, create a custom DB option group with the following settings:

- a. For **Engine**, choose the edition of Oracle you want to use.
- b. For **Major engine version**, choose the version of your DB instance.

For more information, see [Creating an Option Group \(p. 197\)](#).

2. Add the **UTL_MAIL** option to the option group. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#).
3. Apply the option group to a new or existing DB instance:

- For a new DB instance, you apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).
- For an existing DB instance, you apply the option group by modifying the instance and attaching the new option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Using Oracle UTL_MAIL

After you enable the UTL_MAIL option, you must configure the SMTP server before you can begin using it.

You configure the SMTP server by setting the **SMTP_OUT_SERVER** parameter to a valid IP address or public DNS name. For the **SMTP_OUT_SERVER** parameter, you can specify a comma-separated list of the addresses of multiple servers. If the first server is unavailable, UTL_MAIL tries the next server, and so on.

You can set the default **SMTP_OUT_SERVER** for a DB instance by using a [DB parameter group](#). You can set the **SMTP_OUT_SERVER** parameter for a session by running the following code on your database on your DB instance.

```
ALTER SESSION SET smtp_out_server = mailserver.domain.com:25;
```

After the UTL_MAIL option is enabled, and your **SMTP_OUT_SERVER** is configured, you can send mail by using the **SEND** procedure. For more information, see [UTL_MAIL](#) in the Oracle documentation.

Removing the Oracle UTL_MAIL Option

You can remove Oracle UTL_MAIL from a DB instance.

To remove UTL_MAIL from a DB instance, do one of the following:

- To remove UTL_MAIL from multiple DB instances, remove the UTL_MAIL option from the option group they belong to. This change affects all DB instances that use the option group. For more information, see [Removing an Option from an Option Group \(p. 207\)](#).
- To remove UTL_MAIL from a single DB instance, modify the DB instance and specify a different option group that doesn't include the UTL_MAIL option. You can specify the default (empty) option group,

or a different custom option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Troubleshooting

The following are issues you might encounter when you use UTL_MAIL with Amazon RDS.

- Throttling. SMTP port (25) is throttled based on the elastic network interface owner's policies. If you can successfully send email by using UTL_MAIL, and you see the error ORA-29278: `SMTP transient error: 421 Service not available`, you are possibly being throttled. If you experience throttling with email delivery, we recommend that you implement a backoff algorithm. For more information about backoff algorithms, see [Error Retries and Exponential Backoff in AWS](#) and [How to handle a "Throttling – Maximum sending rate exceeded" error](#).

You can request that this throttle be removed. For more information, see [How do I remove the throttle on port 25 from my EC2 instance?](#).

Related Topics

- [Working with Option Groups \(p. 196\)](#)
- [Options for Oracle DB Instances \(p. 1004\)](#)

Oracle XML DB

Oracle XML DB adds native XML support to your DB instance. With XML DB, you can store and retrieve structured or unstructured XML, in addition to relational data.

XML DB is pre-installed on Oracle version 12c and later. Amazon RDS supports Oracle XML DB for version 11g through the use of the XMLDB option. After you apply the XMLDB option to your DB instance, you have full access to the Oracle XML DB repository; no post-installation tasks are required.

Note

The Amazon RDS XMLDB option does not provide support for the Oracle XML DB Protocol Server.

Adding the Oracle XML DB Option

The general process for adding the Oracle XML DB option to a DB instance is the following:

1. Create a new option group, or copy or modify an existing option group.
2. Add the option to the option group.
3. Associate the option group with the DB instance.

After you add the XML DB option, as soon as the option group is active, XML DB is active.

To add the XML DB option to a DB instance

1. Determine the option group you want to use. You can create a new option group or use an existing option group. If you want to use an existing option group, skip to the next step. Otherwise, create a custom DB option group with the following settings:
 - a. For **Engine**, choose the edition of Oracle you want to use.
 - b. For **Major engine version**, choose **11.2**.

- For more information, see [Creating an Option Group \(p. 197\)](#).
2. Add the **XMLDB** option to the option group. For more information about adding options, see [Adding an Option to an Option Group \(p. 200\)](#).
 3. Apply the option group to a new or existing DB instance:
 - For a new DB instance, you apply the option group when you launch the instance. For more information, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).
 - For an existing DB instance, you apply the option group by modifying the instance and attaching the new option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Removing the Oracle XML DB Option

You can remove the XML DB option from a DB instance running version 11g.

To remove the XML DB option from a DB instance running version 11g, do one of the following:

- To remove the XMLDB option from multiple DB instances, remove the XMLDB option from the option group they belong to. This change affects all DB instances that use the option group. For more information, see [Removing an Option from an Option Group \(p. 207\)](#).
- To remove the XMLDB option from a single DB instance, modify the DB instance and specify a different option group that doesn't include the XMLDB option. You can specify the default (empty) option group, or a different custom option group. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Common DBA Tasks for Oracle DB Instances

Following, you can find a description for the Amazon RDS-specific implementations of some common DBA tasks for DB instances running the Oracle database engine. To deliver a managed service experience, Amazon RDS doesn't provide shell access to DB instances, and restricts access to certain system procedures and tables that require advanced privileges.

The following are common DBA tasks for DB instances running Oracle:

- [System Tasks \(p. 1085\)](#)

Disconnecting a Session (p. 1085)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.disconnect</code> Oracle method: <code>alter system disconnect session</code>
Killing a Session (p. 1086)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.kill</code> Oracle method: <code>alter system kill session</code>
Cancelled a SQL Statement in a Session (p. 1086)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.cancel</code> Oracle method: <code>alter system cancel sql</code>
Enabling and Disabling Restricted Sessions (p. 1087)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.restricted_s</code> Oracle method: <code>alter system enable restricted session</code>
Flushing the Shared Pool (p. 1088)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.flush_shared</code> Oracle method: <code>alter system flush shared_pool</code>
Flushing the Buffer Cache (p. 1088)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.flush_buffer</code> Oracle method: <code>alter system flush buffer_cache</code>
Granting SELECT or EXECUTE Privileges to SYS Objects (p. 1088)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.grant_sys_o</code> Oracle method: <code>grant</code>
Revoking SELECT or EXECUTE Privileges on SYS Objects (p. 1089)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.revoke_sys_o</code> Oracle method: <code>revoke</code>
Granting Privileges to Non-Master Users (p. 1090)	Amazon RDS method: <code>grant</code>

	Oracle method: <code>grant</code>
Creating Custom Functions to Verify Passwords (p. 1091)	Amazon RDS method: <code>rdsadmin.rdsadmin_password_verify.c</code>
Setting Up a Custom DNS Server (p. 1094)	Amazon RDS method: <code>rdsadmin.rdsadmin_password_verify.c</code>

- [Database Tasks \(p. 1095\)](#)

Changing the Global Name of a Database (p. 1095)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.rename_global_name</code> Oracle method: <code>alter database rename</code>
Creating and Sizing Tablespaces (p. 1096)	Amazon RDS method: <code>create tablespace</code> Oracle method: <code>alter database</code>
Setting the Default Tablespace (p. 1096)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.alter_default_tablespace</code> Oracle method: <code>alter database default tablespace</code>
Setting the Default Temporary Tablespace (p. 1097)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.alter_default_temp_tablespace</code> Oracle method: <code>alter database default temporary tablespace</code>
Checkpointing a Database (p. 1097)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.checkpoint</code> Oracle method: <code>alter system checkpoint</code>
Setting Distributed Recovery (p. 1097)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.enable_distr_recovery</code> Oracle method: <code>alter system enable distributed recovery</code>
Setting the Database Time Zone (p. 1097)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.alter_db_time_zone</code> Oracle method: <code>alter database set time_zone</code>
Working with Oracle External Tables (p. 1098)	—
Working with Automatic Workload Repository (AWR) (p. 1099)	—
Adjusting Database Links for Use with DB Instances in a VPC (p. 1099)	—

Setting the Default Edition for a DB Instance (p. 1100)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.alter_default_edition</code> Oracle method: <code>alter database default edition</code>
Enabling Auditing for the SYS.AUD\$ Table (p. 1100)	Amazon RDS method: <code>rdsadmin.rdsadmin_master_util.audit_all_sys_aud_table</code> Oracle method: <code>audit</code>
Disabling Auditing for the SYS.AUD\$ Table (p. 1101)	Amazon RDS method: <code>rdsadmin.rdsadmin_master_util.noaudit_all_sys_aud_table</code> Oracle method: <code>noaudit</code>
Cleaning Up Interrupted Online Index Builds (p. 1101)	Amazon RDS method: <code>rdsadmin.rdsadmin_dbms_repair.online_index_clean</code> Oracle method: <code>dbms_repair.online_index_clean</code>
Skipping Corrupt Blocks (p. 1102)	Amazon RDS method: Several <code>rdsadmin.rdsadmin_dbms_repair</code> procedures Oracle method: <code>dbms_repair</code> package
Resizing the Temporary Tablespace in a Read Replica (p. 1104)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.resize_temp_tablespace</code> or <code>rdsadmin.rdsadmin_util.resize_tempfile</code> procedure Oracle method: —

- [Log Tasks \(p. 1105\)](#)

Setting Force Logging (p. 1106)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.force_logging</code> Oracle method: <code>alter database force logging</code>
Setting Supplemental Logging (p. 1106)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.alter_supple</code> Oracle method: <code>alter database add supplemental log</code>
Switching Online Log Files (p. 1107)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.switch_logfi</code> Oracle method: <code>alter system switch logfile</code>
Adding Online Redo Logs (p. 1107)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.add_logfile</code>
Dropping Online Redo Logs (p. 1107)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.drop_logfile</code>

Resizing Online Redo Logs (p. 1108)	—
Retaining Archived Redo Logs (p. 1110)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.set_configur</code>
Accessing Transaction Logs (p. 1111)	Amazon RDS method: <code>rdsadmin.rdsadmin_master_util.creat</code>
	Amazon RDS method: <code>rdsadmin.rdsadmin_master_util.creat</code>

- [RMAN Tasks \(p. 1112\)](#)

Validating DB Instance Files (p. 1115)	Amazon RDS method: <code>rdsadmin_rman_util.procedure</code> Oracle method: <code>RMAN VALIDATE</code>
Enabling and Disabling Block Change Tracking (p. 1118)	Amazon RDS method: <code>rdsadmin_rman_util.procedure</code> Oracle method: <code>ALTER DATABASE</code>
Crosschecking Archived Redo Logs (p. 1119)	Amazon RDS method: <code>rdsadmin_rman_util.crosscheck_archi</code> Oracle method: <code>RMAN BACKUP</code>
Backing Up Archived Redo Logs (p. 1120)	Amazon RDS method: <code>rdsadmin_rman_util.procedure</code> Oracle method: <code>RMAN BACKUP</code>
Performing a Full Database Backup (p. 1125)	Amazon RDS method: <code>rdsadmin_rman_util.backup_database</code> Oracle method: <code>RMAN BACKUP</code>
Performing an Incremental Database Backup (p. 1125)	Amazon RDS method: <code>rdsadmin_rman_util.backup_database</code> Oracle method: <code>RMAN BACKUP</code>
Performing a Tablespace Backup (p. 1127)	Amazon RDS method: <code>rdsadmin_rman_util.backup_database</code> Oracle method: <code>RMAN BACKUP</code>

- [Oracle Scheduler Tasks \(p. 1128\)](#)

Modifying DBMS_SCHEDULER Jobs (p. 1129)	Amazon RDS method: <code>rdsadmin.rdsadmin_dbms_scheduler.set_attribute</code> Oracle method: <code>dbms_scheduler.set_attribute</code>
Disabling SYS-Owned Oracle Scheduler Jobs (p. 1129)	Amazon RDS method: <code>rdsadmin.rdsadmin_dbms_scheduler.disable</code> Oracle method: <code>dbms_scheduler.disable</code>
Enabling SYS-Owned Oracle Scheduler Jobs (p. 1130)	Amazon RDS method: <code>rdsadmin.rdsadmin_dbms_scheduler.enable</code> Oracle method: <code>dbms_scheduler.enable</code>
Modifying the Repeat Interval for Jobs of CALENDAR Type (p. 1130)	Amazon RDS method: <code>rdsadmin.rdsadmin_dbms_scheduler.set_attribute</code> Oracle method: <code>dbms_scheduler.set_attribute</code>
Modifying the Repeat Interval for Jobs of NAMED Type (p. 1130)	Amazon RDS method: <code>rdsadmin.rdsadmin_dbms_scheduler.set_attribute</code> Oracle method: <code>dbms_scheduler.set_attribute</code>

- [Miscellaneous Tasks \(p. 1131\)](#)

Creating New Directories in the Main Data Storage Space (p. 1131)	Amazon RDS method: <code>rdsadmin.rdsadmin_util.create_directory</code> Oracle method: <code>create directory</code>
Listing Files in a DB Instance Directory (p. 1132)	Amazon RDS method: <code>rdsadmin.rds_file_util.listdir</code> Oracle method: —
Reading Files in a DB Instance Directory (p. 1132)	Amazon RDS method: <code>rdsadmin.rds_file_util.read_text_file</code> Oracle method: —

You can also use Amazon RDS procedures for Amazon S3 integration with Oracle and for running OEM Management Agent database tasks. For more information, see [Amazon S3 Integration \(p. 1005\)](#) and [Performing Database Tasks with the Management Agent \(p. 1033\)](#).

Common DBA System Tasks for Oracle DB Instances

This section describes how you can perform common DBA tasks related to the system on your Amazon RDS DB instances running Oracle. To deliver a managed service experience, Amazon RDS doesn't provide shell access to DB instances, and restricts access to certain system procedures and tables that require advanced privileges.

Topics

- [Disconnecting a Session \(p. 1085\)](#)
- [Killing a Session \(p. 1086\)](#)
- [Cancelling a SQL Statement in a Session \(p. 1086\)](#)
- [Enabling and Disabling Restricted Sessions \(p. 1087\)](#)
- [Flushing the Shared Pool \(p. 1088\)](#)
- [Flushing the Buffer Cache \(p. 1088\)](#)
- [Granting SELECT or EXECUTE Privileges to SYS Objects \(p. 1088\)](#)
- [Revoking SELECT or EXECUTE Privileges on SYS Objects \(p. 1089\)](#)
- [Granting Privileges to Non-Master Users \(p. 1090\)](#)
- [Creating Custom Functions to Verify Passwords \(p. 1091\)](#)
- [Setting Up a Custom DNS Server \(p. 1094\)](#)

Disconnecting a Session

To disconnect the current session by ending the dedicated server process, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.disconnect`. The `disconnect` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>sid</code>	number	—	Yes	The session identifier.
<code>serial</code>	number	—	Yes	The serial number of the session.
<code>method</code>	varchar	'IMMEDIATE'	No	Valid values are 'IMMEDIATE' or 'POST_TRANSACTION'.

The following example disconnects a session.

```
begin
    rdsadmin.rdsadmin_util.disconnect(
        sid    => sid,
        serial => serial_number);
end;
/
```

To get the session identifier and the session serial number, query the `V$SESSION` view. The following example gets all sessions for the user `AWSUSER`.

```
select SID, SERIAL#, STATUS from V$SESSION where USERNAME = 'AWSUSER';
```

The database must be open to use this method. For more information about disconnecting a session, see [ALTER SYSTEM](#) in the Oracle documentation.

Killing a Session

To kill a session, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.kill`. The `kill` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>sid</code>	number	—	Yes	The session identifier.
<code>serial</code>	number	—	Yes	The serial number of the session.
<code>method</code>	varchar	null	No	Valid values are ' <code>IMMEDIATE</code> ' or ' <code>PROCESS</code> '.

The following example kills a session.

```
begin
    rdsadmin.rdsadmin_util.kill(
        sid    => sid,
        serial => serial_number);
end;
/
```

To get the session identifier and the session serial number, query the `V$SESSION` view. The following example gets all sessions for the user `AWSUSER`.

```
select SID, SERIAL#, STATUS from V$SESSION where USERNAME = 'AWSUSER';
```

You can specify either `IMMEDIATE` or `PROCESS` as a value for the `method` parameter. By specifying `PROCESS` as the `method` value, you can kill the processes associated with a session. Do this only if killing the session using `IMMEDIATE` as the `method` value was unsuccessful.

Cancelling a SQL Statement in a Session

To cancel a SQL statement in a session, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.cancel`.

Note

This procedure is supported for Oracle version 18.0.0.0 and later.

The `cancel` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>sid</code>	number	—	Yes	The session identifier.
<code>serial</code>	number	—	Yes	The serial number of the session.
<code>sql_id</code>	varchar2	null	No	The SQL identifier of the SQL statement.

The following example cancels a SQL statement in a session.

```
begin
    rdsadmin.rdsadmin_util.cancel(
        sid    => sid,
        serial => serial_number,
        sql_id => sql_id);
end;
/
```

To get the session identifier, the session serial number, and the SQL identifier of a SQL statement, query the V\$SESSION view. The following example gets all sessions and SQL identifiers for the user AWSUSER.

```
select SID, SERIAL#, SQL_ID, STATUS from V$SESSION where USERNAME = 'AWSUSER';
```

Enabling and Disabling Restricted Sessions

To enable and disable restricted sessions, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.restricted_session`. The `restricted_session` procedure has the following parameters.

Parameter Name	Data Type	Default	Yes	Description
<code>p_enable</code>	boolean	true	No	Set to <code>true</code> to enable restricted sessions, <code>false</code> to disable restricted sessions.

The following example shows how to enable and disable restricted sessions.

```
/* Verify that the database is currently unrestricted. */

select LOGINS from V$INSTANCE;

LOGINS
-----
ALLOWED

/* Enable restricted sessions */

exec rdsadmin.rdsadmin_util.restricted_session(p_enable => true);

/* Verify that the database is now restricted. */

select LOGINS from V$INSTANCE;

LOGINS
-----
RESTRICTED

/* Disable restricted sessions */

exec rdsadmin.rdsadmin_util.restricted_session(p_enable => false);
```

```
/* Verify that the database is now unrestricted again. */

select LOGINS from V$INSTANCE;

LOGINS
-----
ALLOWED
```

Flushing the Shared Pool

To flush the shared pool, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.flush_shared_pool`. The `flush_shared_pool` procedure has no parameters.

The following example flushes the shared pool.

```
exec rdsadmin.rdsadmin_util.flush_shared_pool;
```

Flushing the Buffer Cache

To flush the buffer cache, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.flush_buffer_cache`. The `flush_buffer_cache` procedure has no parameters.

The following example flushes the buffer cache.

```
exec rdsadmin.rdsadmin_util.flush_buffer_cache;
```

Granting SELECT or EXECUTE Privileges to SYS Objects

Usually you transfer privileges by using roles, which can contain many objects. To grant privileges to a single object, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.grant_sys_object`. The procedure grants only privileges that the master account already has through a role or direct grant.

The `grant_sys_object` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>p_obj_name</code>	varchar2	—	Yes	The name of the object to grant privileges for. The object can be a directory, function, package, procedure, sequence, table, or view. Object names must be spelled exactly as they appear in <code>DBA_OBJECTS</code> . Most system objects are defined in uppercase, so we recommend that you try that first.
<code>p_grantee</code>	varchar2	—	Yes	The name of the object to grant privileges to. The object can be a schema or a role.

Parameter Name	Data Type	Default	Required	Description
p_privilege	varchar2	null	Yes	—
p_grant_option	boolean	false	No	Set to true to use the with grant option. The p_grant_option parameter is supported for Oracle versions 11.2.0.4.v8 and later, 12.1.0.2.v4 and later, all 12.2.0.1 versions, all 18.0.0.0 versions, and all 19.0.0 versions.

The following example grants select privileges on an object named V_\$SESSION to a user named USER1.

```
begin
    rdsadmin.rdsadmin_util.grant_sys_object(
        p_obj_name  => 'V_$SESSION',
        p_grantee   => 'USER1',
        p_privilege => 'SELECT');
end;
/
```

The following example grants select privileges on an object named V_\$SESSION to a user named USER1 with the grant option.

```
begin
    rdsadmin.rdsadmin_util.grant_sys_object(
        p_obj_name  => 'V_$SESSION',
        p_grantee   => 'USER1',
        p_privilege => 'SELECT',
        p_grant_option => true);
end;
/
```

To be able to grant privileges on an object, your account must have those privileges granted to it directly with the grant option, or via a role granted using with admin option. In the most common case, you may want to grant SELECT on a DBA view that has been granted to the SELECT_CATALOG_ROLE role. If that role isn't already directly granted to your user using with admin option, then you can't transfer the privilege. If you have the DBA privilege, then you can grant the role directly to another user.

The following example grants the SELECT_CATALOG_ROLE and EXECUTE_CATALOG_ROLE to USER1. Since the with admin option is used, USER1 can now grant access to SYS objects that have been granted to SELECT_CATALOG_ROLE.

```
grant SELECT_CATALOG_ROLE to USER1 with admin option;
grant EXECUTE_CATALOG_ROLE to USER1 with admin option;
```

Objects already granted to PUBLIC do not need to be re-granted. If you use the grant_sys_object procedure to re-grant access, the procedure call succeeds.

Revoking SELECT or EXECUTE Privileges on SYS Objects

To revoke privileges on a single object, use the Amazon RDS procedure rdsadmin.rdsadmin_util.revoke_sys_object. The procedure only revokes privileges that the master account already has by using a role or direct grant.

The `revoke_sys_object` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>p_obj_name</code>	varchar2	—	Yes	The name of the object to revoke privileges for. The object can be a directory, function, package, procedure, sequence, table, or view. Object names must be spelled exactly as they appear in <code>DBA_OBJECTS</code> . Most system objects are defined in upper case, so we recommend you try that first.
<code>p_revokee</code>	varchar2	—	Yes	The name of the object to revoke privileges for. The object can be a schema or a role.
<code>p_privilege</code>	varchar2	null	Yes	—

The following example revokes select privileges on an object named `V_$SESSION` from a user named `USER1`.

```
begin
    rdsadmin.rdsadmin_util.revoke_sys_object(
        p_obj_name  => 'V_$SESSION',
        p_revokee   => 'USER1',
        p_privilege => 'SELECT');
end;
/
```

Granting Privileges to Non-Master Users

You can grant select privileges for many objects in the `SYS` schema by using the `SELECT_CATALOG_ROLE` role. The `SELECT_CATALOG_ROLE` role gives users `SELECT` privileges on data dictionary views. The following example grants the role `SELECT_CATALOG_ROLE` to a user named `user1`.

```
grant SELECT_CATALOG_ROLE to user1;
```

You can grant execute privileges for many objects in the `SYS` schema by using the `EXECUTE_CATALOG_ROLE` role. The `EXECUTE_CATALOG_ROLE` role gives users `EXECUTE` privileges for packages and procedures in the data dictionary. The following example grants the role `EXECUTE_CATALOG_ROLE` to a user named `user1`.

```
grant EXECUTE_CATALOG_ROLE to user1;
```

The following example gets the permissions that the roles `SELECT_CATALOG_ROLE` and `EXECUTE_CATALOG_ROLE` allow.

```
select *
```

```
from ROLE_TAB_PRIVS
where ROLE in ('SELECT_CATALOG_ROLE', 'EXECUTE_CATALOG_ROLE')
order by ROLE, TABLE_NAME asc;
```

The following example creates a non-master user named `user1`, grants the `CREATE SESSION` privilege, and grants the `SELECT` privilege on a database named `sh.sales`.

```
create user user1 identified by password;
grant CREATE SESSION to user1;
grant SELECT on sh.sales TO user1;
```

Creating Custom Functions to Verify Passwords

You can create a custom password verification function in two ways. If you want to use standard verification logic, and to store your function in the `SYS` schema, use the `create_verify_function` procedure. If you want to use custom verification logic, or you don't want to store your function in the `SYS` schema, use the `create_passthrough_verify_fcn` procedure.

The `create_verify_function` Procedure

The `create_verify_function` procedure is supported for Oracle version 11.2.0.4.v9 and later, Oracle version 12.1.0.2.v5 and later, all 12.2.0.1 versions, all 18.0.0.0 versions, and all 19.0.0 versions.

You can create a custom function to verify passwords by using the Amazon RDS procedure `rdsadmin.rdsadmin_password_verify.create_verify_function`. The `create_verify_function` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>p_verify_function_name</code>	varchar2	—	Yes	The name for your custom function. This function is created for you in the <code>SYS</code> schema. You assign this function to user profiles.
<code>p_min_length</code>	number	8	No	The minimum number of characters required.
<code>p_max_length</code>	number	256	No	The maximum number of characters allowed.
<code>p_min_letters</code>	number	1	No	The minimum number of letters required.
<code>p_min_uppercase</code>	number	0	No	The minimum number of uppercase letters required.
<code>p_min_lowercase</code>	number	0	No	The minimum number of lowercase letters required.
<code>p_min_digits</code>	number	1	No	The minimum number of digits required.
<code>p_min_special</code>	number	0	No	The minimum number of special characters required.
<code>p_min_different_chars</code>	number	3	No	The minimum number of different characters

Parameter Name	Data Type	Default	Required	Description
				required between the old and new password.
p_disallow_username	boolean	true	No	Set to true to disallow the username in the password.
p_disallow_reverse	boolean	true	No	Set to true to disallow the reverse of the user name in the password.
p_disallow_db_name	boolean	true	No	Set to true to disallow the database or server name in the password.
p_disallow_simple_string	boolean	true	No	Set to true to disallow simple strings as the password.
p_disallow_whitespace	boolean	false	No	Set to true to disallow white space characters in the password.
p_disallow_at_sign	boolean	false	No	Set to true to disallow the @ character in the password.

You can create multiple password verification functions.

There are restrictions on the name of your custom function. Your custom function can't have the same name as an existing system object. The name can be no more than 30 characters long. Also, the name must include one of the following strings: PASSWORD, VERIFY, COMPLEXITY, ENFORCE, or STRENGTH.

The following example creates a function named CUSTOM_PASSWORD_FUNCTION. The function requires that a password has at least 12 characters, 2 uppercase characters, 1 digit, and 1 special character, and that the password disallows the @ character.

```
begin
    rdsadmin.rdsadmin_password_verify.create_verify_function(
        p_verify_function_name => 'CUSTOM_PASSWORD_FUNCTION',
        p_min_length          => 12,
        p_min_uppercase        => 2,
        p_min_digits           => 1,
        p_min_special          => 1,
        p_disallow_at_sign     => true);
end;
/
```

To see the text of your verification function, query DBA_SOURCE. The following example gets the text of a custom password function named CUSTOM_PASSWORD_FUNCTION.

```
col text format a150

select TEXT
  from DBA_SOURCE
 where OWNER = 'SYS' and NAME = 'CUSTOM_PASSWORD_FUNCTION'
order by LINE;
```

To associate your verification function with a user profile, use `alter profile`. The following example associates a verification function with the `DEFAULT` user profile.

```
alter profile DEFAULT limit PASSWORD_VERIFY_FUNCTION CUSTOM_PASSWORD_FUNCTION;
```

To see what user profiles are associated with what verification functions, query `DBA_PROFILES`. The following example gets the profiles that are associated with the custom verification function named `CUSTOM_PASSWORD_FUNCTION`.

```
select * from DBA_PROFILES where RESOURCE_NAME = 'PASSWORD' and LIMIT =
'CUSTOM_PASSWORD_FUNCTION';
```

PROFILE	RESOURCE_NAME	RESOURCE LIMIT
-----	-----	-----
DEFAULT	PASSWORD_VERIFY_FUNCTION	PASSWORD
CUSTOM_PASSWORD_FUNCTION		

The following example gets all profiles and the password verification functions that they are associated with.

```
select * from DBA_PROFILES where RESOURCE_NAME = 'PASSWORD_VERIFY_FUNCTION';
```

PROFILE	RESOURCE_NAME	RESOURCE LIMIT
-----	-----	-----
DEFAULT	PASSWORD_VERIFY_FUNCTION	PASSWORD
CUSTOM_PASSWORD_FUNCTION	PASSWORD_VERIFY_FUNCTION	PASSWORD NULL
RDSADMIN		

The `create_passthrough_verify_fcn` Procedure

The `create_passthrough_verify_fcn` procedure is supported for Oracle version 11.2.0.4.v11 and later, Oracle version 12.1.0.2.v7 and later, all 12.2.0.1 versions, all 18.0.0.0 versions, and all 19.0.0 versions.

You can create a custom function to verify passwords by using the Amazon RDS procedure `rdsadmin.rdsadmin_password_verify.create_passthrough_verify_fcn`. The `create_passthrough_verify_fcn` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>p_verify_function_name</code>	varchar2	—	Yes	The name for your custom verification function. This is a wrapper function that is created for you in the <code>SYS</code> schema, and it doesn't contain any verification logic. You assign this function to user profiles.
<code>p_target_owner</code>	varchar2	—	Yes	The schema owner for your custom verification function.

Parameter Name	Data Type	Default	Required	Description
p_target_function_name	varchar2	—	Yes	The name of your existing custom function that contains the verification logic. Your custom function must return a boolean. Your function should return <code>true</code> if the password is valid and <code>false</code> if the password is invalid.

The following example creates a password verification function that uses the logic from the function named `PASSWORD_LOGIC_EXTRA_STRONG`.

```
begin
    rdsadmin.rdsadmin_password_verify.create_passthrough_verify_fcn(
        p_verify_function_name => 'CUSTOM_PASSWORD_FUNCTION',
        p_target_owner          => 'TEST_USER',
        p_target_function_name  => 'PASSWORD_LOGIC_EXTRA_STRONG');
end;
/
```

To associate the verification function with a user profile, use `alter profile`. The following example associates the verification function with the `DEFAULT` user profile.

```
alter profile DEFAULT limit PASSWORD_VERIFY_FUNCTION CUSTOM_PASSWORD_FUNCTION;
```

Setting Up a Custom DNS Server

Amazon RDS supports outbound network access on your DB instances running Oracle. For more information about outbound network access, including prerequisites, see [Using `utl_http`, `utl_tcp`, and `utl_smtp` with an Oracle DB Instance \(p. 938\)](#).

Amazon RDS Oracle allows Domain Name Service (DNS) resolution from a custom DNS server owned by the customer. You can resolve only fully qualified domain names from your Amazon RDS DB instance through your custom DNS server.

After you set up your custom DNS name server, it takes up to 30 minutes to propagate the changes to your DB instance. After the changes are propagated to your DB instance, all outbound network traffic requiring a DNS lookup queries your DNS server over port 53.

To set up a custom DNS server for your Oracle Amazon RDS DB instance, do the following:

- From the DHCP options set attached to your VPC, set the `domain-name-servers` option to the IP address of your DNS name server. For more information, see [DHCP Options Sets](#).

Note

The `domain-name-servers` option accepts up to four values, but your Amazon RDS DB instance uses only the first value.

- Ensure that your DNS server can resolve all lookup queries, including public DNS names, Amazon EC2 private DNS names, and customer-specific DNS names. If the outbound network traffic contains any DNS lookups that your DNS server can't handle, your DNS server must have appropriate upstream DNS providers configured.
- Configure your DNS server to produce User Datagram Protocol (UDP) responses of 512 bytes or less.

- Configure your DNS server to produce Transmission Control Protocol (TCP) responses of 1024 bytes or less.
- Configure your DNS server to allow inbound traffic from your Amazon RDS DB instances over port 53. If your DNS server is in an Amazon VPC, the VPC must have a security group that contains inbound rules that allow UDP and TCP traffic on port 53. If your DNS server is not in an Amazon VPC, it must have appropriate firewall whitelisting to allow UDP and TCP inbound traffic on port 53.

For more information, see [Security Groups for Your VPC](#) and [Adding and Removing Rules](#).

- Configure the VPC of your Amazon RDS DB instance to allow outbound traffic over port 53. Your VPC must have a security group that contains outbound rules that allow UDP and TCP traffic on port 53.

For more information, see [Security Groups for Your VPC](#) and [Adding and Removing Rules](#).

- The routing path between the Amazon RDS DB instance and the DNS server has to be configured correctly to allow DNS traffic.
 - If the Amazon RDS DB instance and the DNS server are not in the same VPC, a peering connection has to be set up between them. For more information, see [What is VPC Peering?](#)

Common DBA Database Tasks for Oracle DB Instances

This section describes how you can perform common DBA tasks related to databases on your Amazon RDS DB instances running Oracle. To deliver a managed service experience, Amazon RDS doesn't provide shell access to DB instances, and restricts access to certain system procedures and tables that require advanced privileges.

Topics

- [Changing the Global Name of a Database \(p. 1095\)](#)
- [Creating and Sizing Tablespaces \(p. 1096\)](#)
- [Setting the Default Tablespace \(p. 1096\)](#)
- [Setting the Default Temporary Tablespace \(p. 1097\)](#)
- [Checkpointing a Database \(p. 1097\)](#)
- [Setting Distributed Recovery \(p. 1097\)](#)
- [Setting the Database Time Zone \(p. 1097\)](#)
- [Working with Oracle External Tables \(p. 1098\)](#)
- [Working with Automatic Workload Repository \(AWR\) \(p. 1099\)](#)
- [Adjusting Database Links for Use with DB Instances in a VPC \(p. 1099\)](#)
- [Setting the Default Edition for a DB Instance \(p. 1100\)](#)
- [Enabling Auditing for the SYS.AUD\\$ Table \(p. 1100\)](#)
- [Disabling Auditing for the SYS.AUD\\$ Table \(p. 1101\)](#)
- [Cleaning Up Interrupted Online Index Builds \(p. 1101\)](#)
- [Skipping Corrupt Blocks \(p. 1102\)](#)
- [Resizing the Temporary Tablespace in a Read Replica \(p. 1104\)](#)

Changing the Global Name of a Database

To change the global name of a database, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.rename_global_name`. The `rename_global_name` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
p_new_global_name	varchar2	—	Yes	The new global name for the database.

The database must be open for the name change to occur. For more information about changing the global name of a database, see [ALTER DATABASE](#) in the Oracle documentation.

The following example changes the global name of a database to `new_global_name`.

```
exec rdsadmin.rdsadmin_util.rename_global_name(p_new_global_name => 'new_global_name');
```

Creating and Sizing Tablespaces

Amazon RDS only supports Oracle Managed Files (OMF) for data files, log files, and control files. When you create data files and log files, you can't specify the physical file names.

By default, tablespaces are created with auto-extend enabled, and no maximum size. Because of these default settings, tablespaces can grow to consume all allocated storage. We recommend that you specify an appropriate maximum size on permanent and temporary tablespaces, and that you carefully monitor space usage.

The following example creates a tablespace named `users2` with a starting size of 1 gigabyte and a maximum size of 10 gigabytes:

```
create tablespace users2 datafile size 1G autoextend on maxsize 10G;
```

The following example creates temporary tablespace named `temp01`:

```
create temporary tablespace temp01;
```

We recommend that you don't use smallfile tablespaces because you can't resize smallfile tablespaces with Amazon RDS for Oracle. However, you can add a datafile to a smallfile tablespace.

You can resize a bigfile tablespace by using [ALTER TABLESPACE](#). You can specify the size in kilobytes (K), megabytes (M), gigabytes (G), or terabytes (T).

The following example resizes a bigfile tablespace named `users2` to 200 MB.

```
alter tablespace users2 resize 200M;
```

The following example adds an additional datafile to a smallfile tablespace named `users2`.

```
alter tablespace users2 add datafile size 100000M autoextend on next 250M
maxsize UNLIMITED;
```

Setting the Default Tablespace

To set the default tablespace, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.alter_default_tablespace`. The `alter_default_tablespace` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
tablespace_name	varchar	—	Yes	The name of the default tablespace.

The following example sets the default tablespace to `users2`:

```
exec rdsadmin.rdsadmin_util.alter_default_tablespace(tablespace_name => 'users2');
```

Setting the Default Temporary Tablespace

To set the default temporary tablespace, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.alter_default_temp_tablespace`. The `alter_default_temp_tablespace` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
tablespace_name	varchar	—	Yes	The name of the default temporary tablespace.

The following example sets the default temporary tablespace to `temp01`.

```
exec rdsadmin.rdsadmin_util.alter_default_temp_tablespace(tablespace_name => 'temp01');
```

Checkpointing a Database

To checkpoint the database, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.checkpoint`. The `checkpoint` procedure has no parameters.

The following example checkpoints the database.

```
exec rdsadmin.rdsadmin_util.checkpoint;
```

Setting Distributed Recovery

To set distributed recovery, use the Amazon RDS procedures `rdsadmin.rdsadmin_util.enable_distr_recovery` and `disable_distr_recovery`. The procedures have no parameters.

The following example enables distributed recovery.

```
exec rdsadmin.rdsadmin_util.enable_distr_recovery;
```

The following example disables distributed recovery.

```
exec rdsadmin.rdsadmin_util.disable_distr_recovery;
```

Setting the Database Time Zone

There are two different ways that you can set the time zone of your Amazon RDS Oracle database:

- You can use the `Timezone` option.

The `Timezone` option changes the time zone at the host level and impacts all date columns and values such as `SYSDATE`. For more information about the `Timezone` option, see [Oracle Time Zone \(p. 1071\)](#).

- You can use the Amazon RDS procedure `rdsadmin.rdsadmin_util.alter_db_time_zone`.

The `alter_db_time_zone` procedure changes the time zone for only certain data types, and doesn't change `SYSDATE`. There are additional restrictions on setting the time zone listed in the [Oracle documentation](#).

The `alter_db_time_zone` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>p_new_tz</code>	<code>varchar2</code>	—	Yes	The new time zone as a named region or an absolute offset from Coordinated Universal Time (UTC). Valid offsets range from -12:00 to +14:00.

The following example changes the time zone to UTC plus 3 hours.

```
exec rdsadmin.rdsadmin_util.alter_db_time_zone(p_new_tz => '+3:00');
```

The following example changes the time zone to the time zone of the Africa/Algiers region.

```
exec rdsadmin.rdsadmin_util.alter_db_time_zone(p_new_tz => 'Africa/Algiers');
```

After you alter the time zone by using the `alter_db_time_zone` procedure, you must reboot the DB instance for the change to take effect. For more information, see [Rebooting a DB Instance \(p. 166\)](#).

Working with Oracle External Tables

Oracle external tables are tables with data that is not in the database. Instead, the data is in external files that the database can access. By using external tables, you can access data without loading it into the database. For more information about external tables, see [Managing External Tables](#) in the Oracle documentation.

With Amazon RDS, you can store external table files in directory objects. You can create a directory object, or you can use one that is predefined in the Oracle database, such as the `DATA_PUMP_DIR` directory. For information about creating directory objects, see [Creating New Directories in the Main Data Storage Space \(p. 1131\)](#). You can query the `ALL_DIRECTORIES` view to list the directory objects for your Amazon RDS Oracle DB instance.

Note

Directory objects point to the main data storage space (Amazon EBS volume) used by your instance. The space used—along with data files, redo logs, audit, trace, and other files—counts against allocated storage.

You can move an external data file from one Oracle database to another by using the `DBMS_FILE_TRANSFER` package or the `UTL_FILE` package. The external data file is moved from a directory on the source database to the specified directory on the destination database. For information about using `DBMS_FILE_TRANSFER`, see [Importing Using Oracle Data Pump \(p. 987\)](#).

After you move the external data file, you can create an external table with it. The following example creates an external table that uses the `emp_xt_file1.txt` file in the `USER_DIR1` directory.

```
CREATE TABLE emp_xt (
    emp_id      NUMBER,
    first_name  VARCHAR2(50),
    last_name   VARCHAR2(50),
    user_name   VARCHAR2(20)
)
ORGANIZATION EXTERNAL (
    TYPE ORACLE_LOADER
    DEFAULT DIRECTORY USER_DIR1
    ACCESS PARAMETERS (
        RECORDS DELIMITED BY NEWLINE
        FIELDS TERMINATED BY ','
        MISSING FIELD VALUES ARE NULL
        (emp_id,first_name,last_name,user_name)
    )
    LOCATION ('emp_xt_file1.txt')
)
PARALLEL
REJECT LIMIT UNLIMITED;
```

Suppose that you want to move data that is in an Amazon RDS Oracle DB instance into an external data file. In this case, you can populate the external data file by creating an external table and selecting the data from the table in the database. For example, the following SQL statement creates the `orders_xt` external table by querying the `orders` table in the database.

```
CREATE TABLE orders_xt
    ORGANIZATION EXTERNAL
    (
        TYPE ORACLE_DATAPUMP
        DEFAULT DIRECTORY DATA_PUMP_DIR
        LOCATION ('orders_xt.dmp')
    )
    AS SELECT * FROM orders;
```

In this example, the data is populated in the `orders_xt.dmp` file in the `DATA_PUMP_DIR` directory.

Working with Automatic Workload Repository (AWR)

If you use Oracle Database Enterprise Edition and have licensed the Diagnostics and Tuning packs, you can use Automatic Workload Repository (AWR). To enable AWR, change the `CONTROL_MANAGEMENT_PACK_ACCESS` parameter.

AWR reports are typically generated using report generation scripts, such as `awrrpt.sql`, installed on the database host server. You don't have direct access to the host, but you can obtain copies of the scripts from another installation of Oracle Database. Alternatively, you can generate reports using the `DBMS_WORKLOAD_REPOSITORY` package.

Adjusting Database Links for Use with DB Instances in a VPC

To use Oracle database links with Amazon RDS DB instances inside the same VPC or peered VPCs, the two DB instances should have a valid route between them. Verify the valid route between the DB instances by using your VPC routing tables and network access control list (ACL).

The security group of each DB instance must allow ingress to and egress from the other DB instance. The inbound and outbound rules can refer to security groups from the same VPC or a peered VPC. For more information, see [Updating Your Security Groups to Reference Peered VPC Security Groups](#).

If you have configured a custom DNS server using the DHCP Option Sets in your VPC, your custom DNS server must be able to resolve the name of the database link target. For more information, see [Setting Up a Custom DNS Server \(p. 1094\)](#).

For more information about using database links with Oracle Data Pump, see [Importing Using Oracle Data Pump \(p. 987\)](#).

Setting the Default Edition for a DB Instance

You can redefine database objects in a private environment called an edition. You can use edition-based redefinition to upgrade an application's database objects with minimal downtime.

You can set the default edition of an Amazon RDS Oracle DB instance using the Amazon RDS procedure `rdsadmin.rdsadmin_util.alter_default_edition`.

The following example sets the default edition for the Amazon RDS Oracle DB instance to `RELEASE_V1`.

```
exec rdsadmin.rdsadmin_util.alter_default_edition('RELEASE_V1');
```

The following example sets the default edition for the Amazon RDS Oracle DB instance back to the Oracle default.

```
exec rdsadmin.rdsadmin_util.alter_default_edition('ORACLEBASE');
```

For more information about Oracle edition-based redefinition, see [About Editions and Edition-Based Redefinition](#) in the Oracle documentation.

Enabling Auditing for the SYS.AUD\$ Table

To enable auditing on the database audit trail table `SYS.AUD$`, use the Amazon RDS procedure `rdsadmin.rdsadmin_master_util.audit_all_sys_aud_table`. The only supported audit property is `ALL`. You can't audit or not audit individual statements or operations.

Enabling auditing is supported for Oracle DB instances running the following versions:

- 11.2.0.4.v18 and later 11.2 versions
- 12.1.0.2.v14 and later 12.1 versions
- All 12.2.0.1 versions
- All 18.0.0.0 versions
- All 19.0.0.0 versions

The `audit_all_sys_aud_table` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>p_by_access</code>	boolean	true	No	Set to <code>true</code> to audit <code>BY ACCESS</code> . Set to <code>false</code> to audit <code>BY SESSION</code> .

The following query returns the current audit configuration for `SYS.AUD$` for a database.

```
select * from dba_obj_audit_opts where owner='SYS' and object_name='AUD$';
```

The following commands enable audit of ALL on SYS.AUD\$ BY ACCESS.

```
exec rdsadmin.rdsadmin_master_util.audit_all_sys_aud_table;
exec rdsadmin.rdsadmin_master_util.audit_all_sys_aud_table(p_by_access => true);
```

The following command enables audit of ALL on SYS.AUD\$ BY SESSION.

```
exec rdsadmin.rdsadmin_master_util.audit_all_sys_aud_table(p_by_access => false);
```

For more information, see [AUDIT \(Traditional Auditing\)](#) in the Oracle documentation.

Disabling Auditing for the SYS.AUD\$ Table

To disable auditing on the database audit trail table SYS.AUD\$, use the Amazon RDS procedure `rdsadmin.rdsadmin_master_util.noaudit_all_sys_aud_table`. This procedure takes no parameters.

The following query returns the current audit configuration for SYS.AUD\$ for a database:

```
select * from dba_obj_audit_opts where owner='SYS' and object_name='AUD$';
```

The following commands disables audit of ALL on SYS.AUD\$.

```
exec rdsadmin.rdsadmin_master_util.noaudit_all_sys_aud_table;
```

For more information, see [NOAUDIT \(Traditional Auditing\)](#) in the Oracle documentation.

Cleaning Up Interrupted Online Index Builds

To clean up failed online index builds, use the Amazon RDS procedure `rdsadmin.rdsadmin_dbms_repair.online_index_clean`.

The `online_index_clean` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
object_id	binary_integer	ALL_INDEX_ID	No	The object ID of the index. Typically, you can use the object ID from the ORA-08104 error text.

Parameter Name	Data Type	Default	Required	Description
wait_for_lock	binary_integer	rdsadmin.rdsadmin_dbms_repair.lock_nowait	\$admin_dbms	<p>Specify <code>rdsadmin.rdsadmin_dbms_repair.lock_nowait</code> to try to get a lock on the underlying object and retry until an internal limit is reached if the lock fails.</p> <p>Specify <code>rdsadmin.rdsadmin_dbms_repair.lock_noretry</code> to try to get a lock on the underlying object but not retry if the lock fails.</p>

The following example cleans up a failed online index build:

```

declare
    is_clean boolean;
begin
    is_clean := rdsadmin.rdsadmin_dbms_repair.online_index_clean(
        object_id      => 1234567890,
        wait_for_lock => rdsadmin.rdsadmin_dbms_repair.lock_nowait
    );
end;
/

```

For more information, see [ONLINE_INDEX_CLEAN Function](#) in the Oracle documentation.

Skipping Corrupt Blocks

To skip corrupt blocks during index and tables scans, use the `rdsadmin.rdsadmin_dbms_repair` package.

The following procedures wrap the functionality of the `sys.dbms_repair.admin_table` procedure and take no parameters:

- `rdsadmin.rdsadmin_dbms_repair.create_repair_table`
- `rdsadmin.rdsadmin_dbms_repair.create_orphan_keys_table`
- `rdsadmin.rdsadmin_dbms_repair.drop_repair_table`
- `rdsadmin.rdsadmin_dbms_repair.drop_orphan_keys_table`
- `rdsadmin.rdsadmin_dbms_repair.purge_repair_table`
- `rdsadmin.rdsadmin_dbms_repair.purge_orphan_keys_table`

These procedures take no parameters.

The following procedures take the same parameters as their counterparts in the `DBMS_REPAIR` package for Oracle databases:

- `rdsadmin.rdsadmin_dbms_repair.check_object`
- `rdsadmin.rdsadmin_dbms_repair.dump_orphan_keys`

- rdsadmin.rdsadmin_dbms_repair.fix_corrupt_blocks
- rdsadmin.rdsadmin_dbms_repair.rebuild_freelists
- rdsadmin.rdsadmin_dbms_repair.segment_fix_status
- rdsadmin.rdsadmin_dbms_repair.skip_corrupt_blocks

Each procedure takes the same parameters as the corresponding procedure in the DBMS_REPAIR package for Oracle databases. For more information about the parameters for these procedures, see [DBMS_REPAIR](#) in the Oracle documentation.

Complete the following steps to skip corrupt blocks during index and table scans.

1. Run the following procedures to create repair tables if they don't already exist.

```
exec rdsadmin.rdsadmin_dbms_repair.create_repair_table;
exec rdsadmin.rdsadmin_dbms_repair.create_orphan_keys_table;
```

2. Run the following procedures to check for existing records and purge them if appropriate.

```
select count(*) from SYS.REPAIR_TABLE;
select count(*) from SYS.ORPHAN_KEY_TABLE;
select count(*) from SYS.DBA_REPAIR_TABLE;
select count(*) from SYS.DBA_ORPHAN_KEY_TABLE;

exec rdsadmin.rdsadmin_dbms_repair.purge_repair_table;
exec rdsadmin.rdsadmin_dbms_repair.purge_orphan_keys_table;
```

3. Run the following procedure to check for corrupt blocks.

```
set serveroutput on
declare v_num_corrupt int;
begin
    v_num_corrupt := 0;
    rdsadmin.rdsadmin_dbms_repair.check_object (
        schema_name => '&corruptionOwner',
        object_name => '&corruptionTable',
        corrupt_count => v_num_corrupt
    );
    dbms_output.put_line('number corrupt: '||to_char(v_num_corrupt));
end;
/

col corrupt_description format a30
col repair_description format a30
select object_name, block_id, corrupt_type, marked_corrupt, corrupt_description,
       repair_description from sys.repair_table;

select skip_corrupt from dba_tables where owner = '&corruptionOwner' and table_name =
      '&corruptionTable';
```

4. Run the following procedure to enable corruption skipping for affected tables.

```
begin
    rdsadmin.rdsadmin_dbms_repair.skip_corrupt_blocks (
```

```

    schema_name => '&corruptionOwner',
    object_name => '&corruptionTable',
    object_type => rdsadmin.rdsadmin_dbms_repair.table_object,
    flags => rdsadmin.rdsadmin_dbms_repair.skip_flag);
end;
/
select skip_corrupt from dba_tables where owner = '&corruptionOwner' and table_name =
'&corruptionTable';

```

- Run the following procedure to disable corruption skipping.

```

begin
  rdsadmin.rdsadmin_dbms_repair.skip_corrupt_blocks (
    schema_name => '&corruptionOwner',
    object_name => '&corruptionTable',
    object_type => rdsadmin.rdsadmin_dbms_repair.table_object,
    flags => rdsadmin.rdsadmin_dbms_repair.noskip_flag);
end;
/
select skip_corrupt from dba_tables where owner = '&corruptionOwner' and table_name =
'&corruptionTable';

```

- Run the following procedures to drop the repair tables.

```

exec rdsadmin.rdsadmin_dbms_repair.drop_repair_table;
exec rdsadmin.rdsadmin_dbms_repair.drop_orphan_keys_table;

```

Resizing the Temporary Tablespace in a Read Replica

By default, Oracle tablespaces are created with auto-extend enabled and no maximum size. Because of these default settings, tablespaces can grow too large in some cases. We recommend that you specify an appropriate maximum size on permanent and temporary tablespaces, and that you carefully monitor space usage.

To resize the temporary space in a Read Replica for an Oracle DB instance, use either the `rdsadmin.rdsadmin_util.resize_temp_tablespace` or the `rdsadmin.rdsadmin_util.resize_tempfile` Amazon RDS procedure.

The `resize_temp_tablespace` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>temp_tbs</code>	varchar2	—	Yes	The name of the temporary tablespace to resize.
<code>size</code>	varchar2	—	Yes	You can specify the size in bytes (the default), kilobytes (K), megabytes (M), or gigabytes (G).

The `resize_tempfile` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
file_id	binary_integer	—	Yes	The file identifier of the temporary tablespace to resize.
size	varchar2	—	Yes	You can specify the size in bytes (the default), kilobytes (K), megabytes (M), or gigabytes (G).

The following examples resize a temporary tablespace named `TEMP` to the size of 4 gigabytes on a Read Replica.

```
exec rdsadmin.rdsadmin_util.resize_temp_tablespace('TEMP', '4G');
```

```
exec rdsadmin.rdsadmin_util.resize_temp_tablespace('TEMP', '4096000000');
```

The following example resizes a temporary tablespace based on the tempfile with the file identifier 1 to the size of 2 megabytes on a Read Replica.

```
exec rdsadmin.rdsadmin_util.resize_tempfile(1, '2M');
```

For more information about Read Replicas for Oracle DB instances, see [Working with Oracle Read Replicas for Amazon RDS \(p. 998\)](#).

Common DBA Log Tasks for Oracle DB Instances

This section describes how you can perform common DBA tasks related to logging on your Amazon RDS DB instances running Oracle. To deliver a managed service experience, Amazon RDS doesn't provide shell access to DB instances, and restricts access to certain system procedures and tables that require advanced privileges.

For more information, see [Oracle Database Log Files \(p. 413\)](#).

Topics

- [Setting Force Logging \(p. 1106\)](#)
- [Setting Supplemental Logging \(p. 1106\)](#)
- [Switching Online Log Files \(p. 1107\)](#)
- [Adding Online Redo Logs \(p. 1107\)](#)
- [Dropping Online Redo Logs \(p. 1107\)](#)
- [Resizing Online Redo Logs \(p. 1108\)](#)
- [Retaining Archived Redo Logs \(p. 1110\)](#)
- [Accessing Transaction Logs \(p. 1111\)](#)

Setting Force Logging

In force logging mode, Oracle logs all changes to the database except changes in temporary tablespaces and temporary segments (`NOLOGGING` clauses are ignored). For more information, see [Specifying FORCE LOGGING Mode](#) in the Oracle documentation.

To set force logging, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.force_logging`. The `force_logging` procedure has the following parameters.

Parameter Name	Data Type	Default	Yes	Description
<code>p_enable</code>	boolean	true	No	Set to <code>true</code> to put the database in force logging mode, <code>false</code> to remove the database from force logging mode.

The following example puts the database in force logging mode.

```
exec rdsadmin.rdsadmin_util.force_logging(p_enable => true);
```

Setting Supplemental Logging

Supplemental logging ensures that LogMiner and products that use LogMiner technology have sufficient information to support chained rows and storage arrangements such as cluster tables. For more information, see [Supplemental Logging](#) in the Oracle documentation.

Oracle Database doesn't enable supplemental logging by default. To enable and disable supplemental logging, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.alter_supplemental_logging`. For more information about how Amazon RDS manages the retention of archived redo logs for Oracle DB instances, see [Retaining Archived Redo Logs \(p. 1110\)](#).

The `alter_supplemental_logging` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>p_action</code>	varchar2	—	Yes	'ADD' to add supplemental logging, 'DROP' to drop supplemental logging.
<code>p_type</code>	varchar2	null	No	The type of supplemental logging. Valid values are 'ALL', 'FOREIGN KEY', 'PRIMARY KEY', or 'UNIQUE'.

The following example enables supplemental logging.

```
begin
    rdsadmin.rdsadmin_util.alter_supplemental_logging(
        p_action => 'ADD');
end;
/
```

The following example enables supplemental logging for all fixed-length maximum size columns.

```
begin
    rdsadmin.rdsadmin_util.alter_supplemental_logging(
        p_action => 'ADD',
        p_type    => 'ALL');
end;
/
```

The following example enables supplemental logging for primary key columns.

```
begin
    rdsadmin.rdsadmin_util.alter_supplemental_logging(
        p_action => 'ADD',
        p_type    => 'PRIMARY KEY');
end;
/
```

Switching Online Log Files

To switch log files, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.switch_logfile`. The `switch_logfile` procedure has no parameters.

The following example switches log files.

```
exec rdsadmin.rdsadmin_util.switch_logfile;
```

Adding Online Redo Logs

An Amazon RDS DB instance running Oracle starts with four online redo logs, 128 MB each. To add additional redo logs, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.add_logfile`.

The `add_logfile` procedure has the following parameters.

Note

The parameters are mutually exclusive.

Parameter Name	Data Type	Default	Required	Description
bytes	positive	null	No	The size of the log file in bytes.
p_size	varchar2	—	Yes	The size of the log file. You can specify the size in kilobytes (K), megabytes (M), or gigabytes (G).

The following command adds a 100 MB log file.

```
exec rdsadmin.rdsadmin_util.add_logfile(p_size => '100M');
```

Dropping Online Redo Logs

To drop redo logs, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.drop_logfile`. The `drop_logfile` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
grp	positive	—	Yes	The group number of the log.

The following example drops the log with group number 3.

```
exec rdsadmin.rdsadmin_util.drop_logfile(grp => 3);
```

You can only drop logs that have a status of unused or inactive. The following example gets the statuses of the logs.

```
select GROUP#, STATUS from V$LOG;

GROUP#      STATUS
-----
1          CURRENT
2          INACTIVE
3          INACTIVE
4          UNUSED
```

Resizing Online Redo Logs

An Amazon RDS DB instance running Oracle starts with four online redo logs, 128 MB each. The following example shows how you can use Amazon RDS procedures to resize your logs from 128 MB each to 512 MB each.

```
/* Query V$LOG to see the logs.          */
/* You start with 4 logs of 128 MB each. */

select GROUP#, BYTES, STATUS from V$LOG;

GROUP#      BYTES      STATUS
-----
1          134217728  INACTIVE
2          134217728  CURRENT
3          134217728  INACTIVE
4          134217728  INACTIVE

/* Add four new logs that are each 512 MB */

exec rdsadmin.rdsadmin_util.add_logfile(bytes => 536870912);
exec rdsadmin.rdsadmin_util.add_logfile(bytes => 536870912);
exec rdsadmin.rdsadmin_util.add_logfile(bytes => 536870912);
exec rdsadmin.rdsadmin_util.add_logfile(bytes => 536870912);

/* Query V$LOG to see the logs. */
/* Now there are 8 logs.          */

select GROUP#, BYTES, STATUS from V$LOG;

GROUP#      BYTES      STATUS
-----
1          134217728  INACTIVE
2          134217728  CURRENT
3          134217728  INACTIVE
4          134217728  INACTIVE
```

```
5          536870912  UNUSED
6          536870912  UNUSED
7          536870912  UNUSED
8          536870912  UNUSED

/* Drop each inactive log using the group number. */

exec rdsadmin.rdsadmin_util.drop_logfile(grp => 1);
exec rdsadmin.rdsadmin_util.drop_logfile(grp => 3);
exec rdsadmin.rdsadmin_util.drop_logfile(grp => 4);

/* Query V$LOG to see the logs. */
/* Now there are 5 logs. */

select GROUP#, BYTES, STATUS from V$LOG;

GROUP#      BYTES      STATUS
-----
2          134217728  CURRENT
5          536870912  UNUSED
6          536870912  UNUSED
7          536870912  UNUSED
8          536870912  UNUSED

/* Switch logs so that group 2 is no longer current. */

exec rdsadmin.rdsadmin_util.switch_logfile;

/* Query V$LOG to see the logs. */
/* Now one of the new logs is current. */

SQL>select GROUP#, BYTES, STATUS from V$LOG;

GROUP#      BYTES      STATUS
-----
2          134217728  ACTIVE
5          536870912  CURRENT
6          536870912  UNUSED
7          536870912  UNUSED
8          536870912  UNUSED

/* If the status of log 2 is still "ACTIVE", issue a checkpoint to clear it to "INACTIVE".
 */

exec rdsadmin.rdsadmin_util.checkpoint;

/* Query V$LOG to see the logs. */
/* Now the final original log is inactive. */

select GROUP#, BYTES, STATUS from V$LOG;

GROUP#      BYTES      STATUS
-----
2          134217728  INACTIVE
5          536870912  CURRENT
6          536870912  UNUSED
7          536870912  UNUSED
8          536870912  UNUSED
```

```
# Drop the final inactive log.

exec rdsadmin.rdsadmin_util.drop_logfile(grp => 2);

/* Query V$LOG to see the logs.      */
/* Now there are four 512 MB logs. */

select GROUP#, BYTES, STATUS from V$LOG;

GROUP#    BYTES      STATUS
-----  -----
5        536870912 CURRENT
6        536870912 UNUSED
7        536870912 UNUSED
8        536870912 UNUSED
```

Retaining Archived Redo Logs

You can retain archived redo logs locally on your DB instance for use with products like Oracle LogMiner (DBMS_LOGMNR). After you have retained the redo logs, you can use LogMiner to analyze the logs. For more information, see [Using LogMiner to Analyze Redo Log Files](#) in the Oracle documentation.

To retain archived redo logs, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.set_configuration`. The `set_configuration` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>name</code>	varchar	—	Yes	The name of the configuration to update.
<code>value</code>	varchar	—	Yes	The value for the configuration.

The following example retains 24 hours of redo logs.

```
begin
    rdsadmin.rdsadmin_util.set_configuration(
        name  => 'archivelog retention hours',
        value => '24');
end;
/
commit;
```

Note

The commit is required for the change to take effect.

To view how long archived redo logs are kept for your DB instance, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.show_configuration`.

The following example shows the log retention time.

```
set serveroutput on
exec rdsadmin.rdsadmin_util.show_configuration;
```

The output shows the current setting for `archivelog retention hours`. The following output shows that archived redo logs are kept for 48 hours.

```
NAME:archivelog retention hours
VALUE:48
DESCRIPTION:ArchiveLog expiration specifies the duration in hours before archive/redo log
files are automatically deleted.
```

Because the archived redo logs are retained on your DB instance, ensure that your DB instance has enough allocated storage for the retained logs. To determine how much space your DB instance has used in the last X hours, you can run the following query, replacing X with the number of hours.

```
select sum(BLOCKS * BLOCK_SIZE) bytes
  from V$ARCHIVED_LOG
 where FIRST_TIME >= SYSDATE-(X/24) and DEST_ID=1;
```

Archived redo logs are only generated if the backup retention period of your DB instance is greater than zero. By default the backup retention period is greater than zero, so unless you explicitly set yours to zero, archived redo logs are generated for your DB instance. To modify the backup retention period for your DB instance, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

After the archived redo logs are removed from your DB instance, you can't download them again to your DB instance. Amazon RDS retains the archived redo logs outside of your DB instance to support restoring your DB instance to a point in time. Amazon RDS retains the archived redo logs outside of your DB instance based on the backup retention period configured for your DB instance. To modify the backup retention period for your DB instance, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Note

In some cases, you might use JDBC on Linux to download archived redo logs and experience long latency times and connection resets. In such cases, the issues might be caused by the default random number generator setting on your Java client. We recommend setting your JDBC drivers to use a nonblocking random number generator.

Accessing Transaction Logs

Accessing transaction logs is supported for Oracle version 11.2.0.4.v11 and later, Oracle version 12.1.0.2.v7 and later, all 12.2.0.1 versions, all 18.0.0.0 versions, and all 19.0.0 versions.

You might want to access your online and archived redo log files for mining with external tools such as GoldenGate, Attunity, Informatica, and others. If you want to access your online and archived redo log files, you must first create directory objects that provide read-only access to the physical file paths.

The following code creates directories that provide read-only access to your online and archived redo log files:

Important

This code also revokes the `DROP ANY DIRECTORY` privilege.

```
exec rdsadmin.rdsadmin_master_util.create_archivelog_dir;
exec rdsadmin.rdsadmin_master_util.create_onlinelog_dir;
```

After you create directory objects for your online and archived redo log files, you can read the files by using PL/SQL. For more information about reading files from directory objects, see [Listing Files in a DB Instance Directory \(p. 1132\)](#) and [Reading Files in a DB Instance Directory \(p. 1132\)](#).

The following code drops the directories for your online and archived redo log files.

```
exec rdsadmin.rdsadmin_master_util.drop_archivelog_dir;
exec rdsadmin.rdsadmin_master_util.drop_onlinelog_dir;
```

The following code grants and revokes the `DROP ANY DIRECTORY` privilege.

```
exec rdsadmin.rdsadmin_master_util.revoke_drop_any_directory;
exec rdsadmin.rdsadmin_master_util.grant_drop_any_directory;
```

Common DBA Recovery Manager (RMAN) Tasks for Oracle DB Instances

In the following section, you can find how you can perform Oracle Recovery Manager (RMAN) DBA tasks on your Amazon RDS DB instances running Oracle. To deliver a managed service experience, Amazon RDS doesn't provide shell access to DB instances. It also restricts access to certain system procedures and tables that require advanced privileges.

You can use the Amazon RDS package `rdsadmin.rdsadmin_rman_util` to perform RMAN backups of your Amazon RDS for Oracle database to disk. The `rdsadmin.rdsadmin_rman_util` package supports full and incremental database file backups, tablespace backups, and archive log backups.

RMAN backups consume storage space on the Amazon RDS DB instance host. When you perform a backup, you specify an Oracle directory object as a parameter in the procedure call. The backup files are placed in the specified directory. You can use default directories, such as `DATA_PUMP_DIR`, or create a new directory. For more information, see [Creating New Directories in the Main Data Storage Space \(p. 1131\)](#).

After an RMAN backup has finished, you can copy the backup files off the Amazon RDS for Oracle DB instance host. You might do this for the purpose of restoring to a non-RDS host or for long-term storage of backups. For example, you can copy the backup files to an Amazon S3 bucket. For more information, see [using Amazon S3 Integration \(p. 1005\)](#).

The backup files for RMAN backups remain on the Amazon RDS DB instance host until you remove them manually. You can use the `UTL_FILE.FREMOVE` Oracle procedure to remove files from a directory. For more information, see [FREMOVE Procedure](#) in the Oracle documentation.

When backing up archived redo logs or performing a full or incremental backup that includes archived redo logs, redo log retention must be set to a nonzero value. For more information, see [Retaining Archived Redo Logs \(p. 1110\)](#).

Note

For backing up and restoring to another Amazon RDS for Oracle DB instance, you can continue to use the Amazon RDS backup and restore features. For more information, see [Backing Up and Restoring Amazon RDS DB Instances \(p. 253\)](#)

Currently, RMAN restore isn't supported for Amazon RDS for Oracle DB instances.

Topics

- [Common Parameters for RMAN Procedures \(p. 1113\)](#)
- [Validating DB Instance Files \(p. 1115\)](#)
- [Enabling and Disabling Block Change Tracking \(p. 1118\)](#)
- [Crosschecking Archived Redo Logs \(p. 1119\)](#)
- [Backing Up Archived Redo Logs \(p. 1120\)](#)
- [Performing a Full Database Backup \(p. 1125\)](#)
- [Performing an Incremental Database Backup \(p. 1125\)](#)
- [Performing a Tablespace Backup \(p. 1127\)](#)

Common Parameters for RMAN Procedures

You can use procedures in the Amazon RDS package `rdsadmin.rdsadmin_rman_util` to perform tasks with RMAN. Several parameters are common to the procedures in the package. The package has the following common parameters.

Parameter Name	Data Type	Valid Values	Default	Required	Description
<code>p_owner</code>	varchar2	A valid owner of the directory specified in <code>p_directory_name</code> .	—	Yes	The owner of the directory to contain the backup files.
<code>p_directory_name</code>	varchar2	A valid database directory name.	—	Yes	The name of the directory to contain the backup files.
<code>p_label</code>	varchar2	a–z, A–Z, 0–9, '_', '-' , '.'	—	No	A unique string that is included in the backup file names. Note The limit is 30 characters.
<code>p_compress</code>	boolean	TRUE, FALSE	FALSE	No	Specify TRUE to enable BASIC backup compression. Specify FALSE to disable BASIC backup compression.
<code>p_include_archive_log</code>	boolean	TRUE, FALSE	FALSE	No	Specify TRUE to include archived redo logs in the backup. Specify FALSE to exclude archived redo logs from the backup. If you include archived redo logs in the backup, set retention to one hour or greater using the <code>rdsadmin.rdsadmin_util.set_co</code> procedure. Also, call the <code>rdsadmin.rdsadmin_rman_util.co</code> procedure immediately before executing the backup. Otherwise, the backup might

Parameter Name	Data Type	Valid Values	Default	Required	Description
					fail due to missing archived redo logs that have been deleted by Amazon RDS management procedures.
p_include_controlfile	boolean	TRUE, FALSE	FALSE	No	<p>Specify TRUE to include the control file in the backup.</p> <p>Specify FALSE to exclude the control file from the backup.</p>
p_optimize	boolean	TRUE, FALSE	TRUE	No	<p>Specify TRUE to enable backup optimization, if archived redo logs are included, to reduce backup size.</p> <p>Specify FALSE to disable backup optimization.</p>
p_parallel	number	A valid integer between 1 and 254 for Oracle Database Enterprise Edition (EE) 1 for other Oracle Database editions	1	No	Number of channels.
p_rman_to_dbms_output	boolean	TRUE, FALSE	FALSE	No	<p>When TRUE, the RMAN output is sent to the DBMS_OUTPUT package in addition to a file in the BDUMP directory. When using SQL*Plus, execute SET SERVEROUTPUT ON to see the output.</p> <p>When FALSE, the RMAN output is only sent to a file in the BDUMP directory.</p>

Parameter Name	Data Type	Valid Values	Default	Required	Description
p_section_size_mb	number	A valid integer	NULL	No	<p>The section size in megabytes (MB).</p> <p>Validates in parallel by dividing each file into the specified section size.</p> <p>When NULL, the parameter is ignored.</p>
p_validation_type	varchar2	'PHYSICAL' 'PHYSICAL+LOGICAL'	'PHYSICAL'	'No'	<p>The level of corruption detection.</p> <p>Specify 'PHYSICAL' to check for physical corruption. An example of physical corruption is a block with a mismatch in the header and footer.</p> <p>Specify 'PHYSICAL+LOGICAL' to check for logical inconsistencies in addition to physical corruption. An example of logical corruption is a corrupt block.</p>

Validating DB Instance Files

You can use the Amazon RDS package `rdsadmin.rdsadmin_rman_util` to validate Amazon RDS for Oracle DB instance files, such as data files, tablespaces, control files, or server parameter files (SPFILEs).

For more information about RMAN validation, see [Validating Database Files and Backups](#) and [VALIDATE](#) in the Oracle documentation.

Topics

- [Validating a DB Instance \(p. 1115\)](#)
- [Validating a Tablespace \(p. 1116\)](#)
- [Validating a Control File \(p. 1117\)](#)
- [Validating an SPFILE \(p. 1117\)](#)
- [Validating a Data File \(p. 1117\)](#)

Validating a DB Instance

To validate all of the relevant files used by an Amazon RDS Oracle DB instance, use the Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.validate_database`.

This procedure uses the following common parameters for RMAN tasks:

- p_validation_type
- p_parallel
- p_section_size_mb
- p_rman_to_dbms_output

For more information, see [Common Parameters for RMAN Procedures \(p. 1113\)](#).

The following example validates the DB instance using the default values for the parameters.

```
exec rdsadmin.rdsadmin_rman_util.validate_database;
```

The following example validates the DB instance using the specified values for the parameters.

```
BEGIN
    rdsadmin.rdsadmin_rman_util.validate_database(
        p_validation_type      => 'PHYSICAL+LOGICAL',
        p_parallel             => 4,
        p_section_size_mb     => 10,
        p_rman_to_dbms_output => FALSE);
END;
/
```

When the p_rman_to_dbms_output parameter is set to FALSE, the RMAN output is written to a file in the BDUMP directory.

To view the files in the BDUMP directory, run the following SELECT statement.

```
SELECT * FROM table(rdsadmin.rds_file_util.listdir('BDUMP')) order by mtime;
```

To view the contents of a file in the BDUMP directory, run the following SELECT statement.

```
SELECT text FROM table(rdsadmin.rds_file_util.read_text_file('BDUMP','rds-rman-
validate-nnn.txt'));
```

Replace the file name with the name of the file you want to view.

Validating a Tablespace

To validate the files associated with a tablespace, use the Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.validate_tablespace`.

This procedure uses the following common parameters for RMAN tasks:

- p_validation_type
- p_parallel
- p_section_size_mb

- `p_rman_to_dbms_output`

For more information, see [Common Parameters for RMAN Procedures \(p. 1113\)](#).

This procedure also uses the following additional parameter.

Parameter Name	Data Type	Valid Values	Default	Required	Description
<code>p_tablespace_name</code>	<code>varchar2</code>	A valid tablespace name	—	Yes	The name of the tablespace.

Validating a Control File

To validate only the control file used by an Amazon RDS Oracle DB instance, use the Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.validate_current_controlfile`.

This procedure uses the following common parameter for RMAN tasks:

- `p_validation_type`
- `p_rman_to_dbms_output`

For more information, see [Common Parameters for RMAN Procedures \(p. 1113\)](#).

Validating an SPFILE

To validate only the server parameter file (SPFILE) used by an Amazon RDS Oracle DB instance, use the Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.validate_spfile`.

This procedure uses the following common parameter for RMAN tasks:

- `p_validation_type`
- `p_rman_to_dbms_output`

For more information, see [Common Parameters for RMAN Procedures \(p. 1113\)](#).

Validating a Data File

To validate a data file, use the Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.validate_datafile`.

This procedure uses the following common parameters for RMAN tasks:

- `p_validation_type`
- `p_parallel`
- `p_section_size_mb`
- `p_rman_to_dbms_output`

For more information, see [Common Parameters for RMAN Procedures \(p. 1113\)](#).

This procedure also uses the following additional parameters.

Parameter Name	Data Type	Valid Values	Default	Required	Description
p_datafile	varchar2	A valid datafile ID number or a valid datafile name including complete path	—	Yes	The datafile ID number (from v\$datafile.file#) or the full datafile name including the path (from v\$datafile.name).
p_from_block	number	A valid integer	NULL	No	The number of the block where the validation starts within the data file. When this is NULL, 1 is used.
p_to_block	number	A valid integer	NULL	No	The number of the block where the validation ends within the data file. When this is NULL, the maximum block in the data file is used.

Enabling and Disabling Block Change Tracking

You can enable block change tracking for a DB instance using the Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.enable_block_change_tracking`.

You can disable block change tracking for a DB instance using the Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.disable_block_change_tracking`.

These procedures take no parameters. Enabling block change tracking can improve the performance of incremental backups.

These procedures are supported for the following Amazon RDS for Oracle DB engine versions:

- 11.2.0.4.v19 or higher 11.2 versions
- 12.1.0.2.v15 or higher 12.1 versions
- 12.2.0.1.ru-2019-01.rur-2019-01.r1 or higher 12.2 versions
- All 18.0.0.0 versions
- All 19.0.0.0 versions

To determine whether block change tracking is enabled for your DB instance, run the following query.

```
SELECT status, filename FROM V$BLOCK_CHANGE_TRACKING;
```

The following example enables block change tracking for a DB instance.

```
exec rdsadmin.rdsadmin_rman_util.enable_block_change_tracking;
```

The following example disables block change tracking for a DB instance.

```
exec rdsadmin.rdsadmin_rman_util.disable_block_change_tracking;
```

Crosschecking Archived Redo Logs

You can crosscheck archived redo logs using the Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.crosscheck_archivelog`.

You can use this procedure to crosscheck the archived redo logs registered in the control file and optionally delete the expired logs. Each time an RMAN backup is performed, a record is made in the control file. Over time, these records increase the size of the control file. It is best practice to flush old, expired records periodically.

Note

Standard Amazon RDS backups don't use RMAN and therefore don't create records in the control file.

This procedure uses the common parameter `p_rman_to_dbms_output` for RMAN tasks.

For more information, see [Common Parameters for RMAN Procedures \(p. 1113\)](#).

This procedure also uses the following additional parameter.

Parameter Name	Data Type	Valid Values	Default	Required	Description
<code>p_delete_expired</code>	boolean	TRUE, FALSE	TRUE	No	When TRUE, delete expired archived redo logs from the control file. When FALSE, retain the expired archived redo logs in the control file.

This procedure is supported for the following Amazon RDS for Oracle DB engine versions:

- 11.2.0.4.v19 or higher 11.2 versions
- 12.1.0.2.v15 or higher 12.1 versions
- 12.2.0.1.ru-2019-01.rur-2019-01.r1 or higher 12.2 versions
- All 18.0.0.0 versions
- All 19.0.0.0 versions

The following example deletes the expired archived redo logs from the control file.

```
BEGIN
    rdsadmin.rdsadmin_rman_util.crosscheck_archivelog(
        p_delete_expired      => FALSE,
        p_rman_to_dbms_output => FALSE);
END;
/
```

Backing Up Archived Redo Logs

You can use the Amazon RDS package `rdsadmin.rdsadmin_rman_util` to back up archived redo logs for an Amazon RDS Oracle DB instance.

The procedures for backing up archived redo logs are supported for the following Amazon RDS for Oracle DB engine versions:

- 11.2.0.4.v19 or higher 11.2 versions
- 12.1.0.2.v15 or higher 12.1 versions
- 12.2.0.1.ru-2019-01.rur-2019-01.r1 or higher 12.2 versions
- All 18.0.0.0 versions
- All 19.0.0.0 versions

Topics

- [Backing Up All Archived Redo Logs \(p. 1120\)](#)
- [Backing Up an Archived Redo Log from a Date Range \(p. 1121\)](#)
- [Backing Up an Archived Redo Log from an SCN Range \(p. 1122\)](#)
- [Backing Up an Archived Redo Log from a Sequence Number Range \(p. 1123\)](#)

Backing Up All Archived Redo Logs

To back up all of the archived redo logs for an Amazon RDS Oracle DB instance, use the Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.backup_archivelog_all`.

This procedure uses the following common parameters for RMAN tasks:

- `p_owner`
- `p_directory_name`
- `p_label`
- `p_parallel`
- `p_compress`
- `p_rman_to_dbms_output`

For more information, see [Common Parameters for RMAN Procedures \(p. 1113\)](#).

The following example backs up all archived redo logs for the DB instance.

```
BEGIN
    rdsadmin.rdsadmin_rman_util.backup_archivelog_all(
        p_owner      => 'SYS',
```

```

    p_directory_name      => 'MYDIRECTORY',
    p_parallel            => 4,
    p_rman_to_dbms_output => FALSE);
END;
/

```

Backing Up an Archived Redo Log from a Date Range

To back up specific archived redo logs for an Amazon RDS Oracle DB instance by specifying a date range, use the Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.backup_archivelog_date`. The date range specifies which archived redo logs to back up.

This procedure uses the following common parameters for RMAN tasks:

- `p_owner`
- `p_directory_name`
- `p_label`
- `p_parallel`
- `p_compress`
- `p_rman_to_dbms_output`

For more information, see [Common Parameters for RMAN Procedures \(p. 1113\)](#).

This procedure also uses the following additional parameters.

Parameter Name	Data Type	Valid Values	Default	Required	Description
<code>p_from_date</code>	date	A date that is between the <code>start_date</code> and <code>next_date</code> of an archived redo log that exists on disk. The value must be less than or equal to the value specified for <code>p_to_date</code> .	—	Yes	The starting date for the archived log backups.
<code>p_to_date</code>	date	A date that is between the	—	Yes	The ending date for the archived log backups.

Parameter Name	Data Type	Valid Values	Default	Required	Description
		start_date and next_date of an archived redo log that exists on disk. The value must be greater than or equal to the value specified for p_from_date.			

The following example backs up archived redo logs in the date range for the DB instance.

```

BEGIN
    rdsadmin.rdsadmin_rman_util.backup_archivelog_date(
        p_owner          => 'SYS',
        p_directory_name => 'MYDIRECTORY',
        p_from_date      => '03/01/2019 00:00:00',
        p_to_date        => '03/02/2019 00:00:00',
        p_parallel       => 4,
        p_rman_to_dbms_output => FALSE);
END;
/

```

Backing Up an Archived Redo Log from an SCN Range

To back up specific archived redo logs for an Amazon RDS Oracle DB instance by specifying a system change number (SCN) range, use the Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.backup_archivelog_scn`. The SCN range specifies which archived redo logs to back up.

This procedure uses the following common parameters for RMAN tasks:

- `p_owner`
- `p_directory_name`
- `p_label`
- `p_parallel`
- `p_compress`
- `p_rman_to_dbms_output`

For more information, see [Common Parameters for RMAN Procedures \(p. 1113\)](#).

This procedure also uses the following additional parameters.

Parameter Name	Data Type	Valid Values	Default	Required	Description
p_from_scn	number	An SCN of an archived redo log that exists on disk. The value must be less than or equal to the value specified for p_to_scn.	—	Yes	The starting SCN for the archived log backups.
p_to_scn	number	An SCN of an archived redo log that exists on disk. The value must be greater than or equal to the value specified for p_from_scn.	—	Yes	The ending SCN for the archived log backups.

The following example backs up archived redo logs in the SCN range for the DB instance.

```

BEGIN
    rdsadmin.rdsadmin_rman_util.backup_archivelog_scn(
        p_owner          => 'SYS',
        p_directory_name => 'MYDIRECTORY',
        p_from_scn       => 1533835,
        p_to_scn         => 1892447,
        p_parallel       => 4,
        p_rman_to_dbms_output => FALSE);
END;
/

```

Backing Up an Archived Redo Log from a Sequence Number Range

To back up specific archived redo logs for an Amazon RDS Oracle DB instance by specifying a sequence number range, use the Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.backup_archivelog_sequence`. The sequence number range specifies which archived redo logs to back up.

This procedure uses the following common parameters for RMAN tasks:

- `p_owner`
- `p_directory_name`
- `p_label`
- `p_parallel`
- `p_compress`
- `p_rman_to_dbms_output`

For more information, see [Common Parameters for RMAN Procedures \(p. 1113\)](#).

This procedure also uses the following additional parameters.

Parameter Name	Data Type	Valid Values	Default	Required	Description
<code>p_from_sequence</code>	number	A sequence number of an archived redo log that exists on disk. The value must be less than or equal to the value specified for <code>p_to_sequence</code> .	—	Yes	The starting sequence number for the archived log backups.
<code>p_to_sequence</code>	number	A sequence number of an archived redo log that exists on disk. The value must be greater than or equal to the value specified for <code>p_from_sequence</code> .	—	Yes	The ending sequence number for the archived log backups.

The following example backs up archived redo logs in the sequence number range for the DB instance.

```
BEGIN
    rdsadmin.rdsadmin_rman_util.backup_archivelog_sequence(
        p_owner          => 'SYS',
        p_directory_name => 'MYDIRECTORY',
```

```
    p_from_sequence      => 11160,
    p_to_sequence        => 11160,
    p_parallel           => 4,
    p_rman_to_dbms_output => FALSE);
END;
/
```

Performing a Full Database Backup

You can perform a backup of all blocks of data files included in the backup using Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.backup_database_full`.

This procedure uses the following common parameters for RMAN tasks:

- `p_owner`
- `p_directory_name`
- `p_label`
- `p_parallel`
- `p_section_size_mb`
- `p_include_archived_logs`
- `p_optimize`
- `p_compress`
- `p_rman_to_dbms_output`

For more information, see [Common Parameters for RMAN Procedures \(p. 1113\)](#).

This procedure is supported for the following Amazon RDS for Oracle DB engine versions:

- 11.2.0.4.v19 or higher 11.2 versions
- 12.1.0.2.v15 or higher 12.1 versions
- 12.2.0.1.ru-2019-01.rur-2019-01.r1 or higher 12.2 versions
- All 18.0.0.0 versions
- All 19.0.0.0 versions

The following example performs a full backup of the DB instance using the specified values for the parameters.

```
BEGIN
    rdsadmin.rdsadmin_rman_util.backup_database_full(
        p_owner          => 'SYS',
        p_directory_name => 'MYDIRECTORY',
        p_parallel       => 4,
        p_section_size_mb => 10,
        p_rman_to_dbms_output => FALSE);
END;
/
```

Performing an Incremental Database Backup

You can perform an incremental backup of your DB instance using the Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.backup_database_incremental`.

For more information about incremental backups, see [Incremental Backups](#) in the Oracle documentation.

This procedure uses the following common parameters for RMAN tasks:

- p_owner
- p_directory_name
- p_label
- p_parallel
- p_section_size_mb
- p_include_archived_logs
- p_include_controlfile
- p_optimize
- p_compress
- p_rman_to_dbms_output

For more information, see [Common Parameters for RMAN Procedures \(p. 1113\)](#).

This procedure is supported for the following Amazon RDS for Oracle DB engine versions:

- 11.2.0.4.v19 or higher 11.2 versions
- 12.1.0.2.v15 or higher 12.1 versions
- 12.2.0.1.ru-2019-01.rur-2019-01.r1 or higher 12.2 versions
- All 18.0.0.0 versions
- All 19.0.0.0 versions

This procedure also uses the following additional parameter.

Parameter Name	Data Type	Valid Values	Default	Required	Description
p_level	number	0, 1	0	No	<p>Specify 0 to enable a full incremental backup.</p> <p>Specify 1 to enable a non-cumulative incremental backup.</p>

The following example performs an incremental backup of the DB instance using the specified values for the parameters.

```

BEGIN
    rdsadmin.rdsadmin_rman_util.backup_database_incremental(
        p_owner          => 'SYS',
        p_directory_name => 'MYDIRECTORY',
        p_level          => 1,
        p_parallel       => 4,
        p_section_size_mb => 10,
        p_rman_to_dbms_output => FALSE);
END;
/

```

Performing a Tablespace Backup

You can perform a DB instance tablespace using the Amazon RDS procedure `rdsadmin.rdsadmin_rman_util.backup_tablespace`.

This procedure uses the following common parameters for RMAN tasks:

- `p_owner`
- `p_directory_name`
- `p_label`
- `p_parallel`
- `p_section_size_mb`
- `p_include_archived_logs`
- `p_include_controlfile`
- `p_optimize`
- `p_compress`
- `p_rman_to_dbms_output`

For more information, see [Common Parameters for RMAN Procedures \(p. 1113\)](#).

This procedure also uses the following additional parameter.

Parameter Name	Data Type	Valid Values	Default	Required	Description
<code>p_tablespace_name</code>	varchar2	A valid tablespace name.	—	Yes	The name of the tablespace to back up.

This procedure is supported for the following Amazon RDS for Oracle DB engine versions:

- 11.2.0.4.v19 or higher 11.2 versions
- 12.1.0.2.v15 or higher 12.1 versions
- 12.2.0.1.ru-2019-01.rur-2019-01.r1 or higher 12.2 versions
- All 18.0.0.0 versions
- All 19.0.0.0 versions

The following example performs a tablespace backup using the specified values for the parameters.

```

BEGIN
    rdsadmin.rdsadmin_rman_util.backup_tablespace(
        p_owner          => 'SYS',
        p_directory_name => 'MYDIRECTORY',
        p_tablespace_name => MYTABLESPACE,
        p_parallel       => 4,
        p_section_size_mb => 10,
        p_rman_to_dbms_output => FALSE);
END;
/

```

Common DBA Oracle Scheduler Tasks for Oracle DB Instances

Some SYS-owned scheduler jobs can interfere with normal database operations, and Oracle Support recommends they be disabled or the job schedule be modified. You can use the Amazon RDS package `rdsadmin.rdsadmin_dbms_scheduler` to perform tasks for SYS-owned Oracle Scheduler jobs.

These procedures are supported for the following Amazon RDS for Oracle DB engine versions:

- 11.2.0.4.v21 or higher 11.2 versions
- 12.1.0.2.v17 or higher 12.1 versions
- 12.2.0.2.ru-2019-07.rur-2019-07.r1 or higher 12.2 versions

Topics

- [Common Parameters for Oracle Scheduler Procedures \(p. 1128\)](#)
- [Modifying DBMS_SCHEDULER Jobs \(p. 1129\)](#)
- [Disabling SYS-Owned Oracle Scheduler Jobs \(p. 1129\)](#)
- [Enabling SYS-Owned Oracle Scheduler Jobs \(p. 1130\)](#)
- [Modifying the Repeat Interval for Jobs of CALENDAR Type \(p. 1130\)](#)
- [Modifying the Repeat Interval for Jobs of NAMED Type \(p. 1130\)](#)

Common Parameters for Oracle Scheduler Procedures

To perform tasks with Oracle Scheduler, use procedures in the Amazon RDS package `rdsadmin.rdsadmin_dbms_scheduler`. Several parameters are common to the procedures in the package. The package has the following common parameters.

Parameter Name	Data Type	Valid Values	Default	Required	Description
name	varchar2	'SYS.BSLN_MAINTAIN_STATS_JOB'	'SYS.CLEANUP_ONLINE_IND_BUILD'		This name of the job to modify. Note Currently, you can only modify <code>SYS.CLEANUP_ONLINE_IND_BUILD</code> and <code>SYS.BSLN_MAINTAIN_STATS_JOB</code> jobs.
attribute	varchar2	'REPEAT_INTERVAL','SCHEDULE_NAME'	'REPEAT_INTERVAL'		Attribute to modify. To modify the repeat interval for the job, specify 'REPEAT_INTERVAL'. To modify the schedule name for the job, specify 'SCHEDULE_NAME'.

Parameter Name	Data Type	Valid Values	Default	Required	Description
value	varchar2	A valid schedule interval or schedule name, depending on attribute used.	–	Yes	The new value of the attribute.

Modifying DBMS_SCHEDULER Jobs

You can still use the Oracle procedure `dbms_scheduler.set_attribute` to modify certain components of the scheduler. For more information, see [DBMS_SCHEDULER](#) and [SET_ATTRIBUTE Procedure](#) in the Oracle documentation.

When working with Amazon RDS DB instances, prepend the schema name `SYS` to the object name. The following example sets the resource plan attribute for the Monday window object.

```
begin
    dbms_scheduler.set_attribute(
        name      => 'SYS.MONDAY_WINDOW',
        attribute => 'RESOURCE_PLAN',
        value     => 'resource_plan_1');
end;
/
```

Note

Some SYS-owned Oracle Scheduler jobs can interfere with normal database operations. Oracle Support recommends that they be disabled or that the job schedule be modified. Amazon RDS for Oracle doesn't provide the required privileges to modify SYS-owned Oracle Scheduler jobs using the `DBMS_SCHEDULER` package. Instead, you can use the procedures in the Amazon RDS package `rdsadmin.rdsadmin_dbms_scheduler` to perform tasks for SYS-owned Oracle Scheduler jobs. For information about using these procedures, see [Common DBA Oracle Scheduler Tasks for Oracle DB Instances \(p. 1128\)](#).

Disabling SYS-Owned Oracle Scheduler Jobs

To disable a SYS-owned Oracle Scheduler job, use the `rdsadmin.rdsadmin_dbms_scheduler.disable` procedure.

This procedure uses the `name` common parameter for Oracle Scheduler tasks. For more information, see [Common Parameters for Oracle Scheduler Procedures \(p. 1128\)](#).

The following example disables the `SYS.CLEANUP_ONLINE_IND_BUILD` Oracle Scheduler job.

```
BEGIN
    rdsadmin.rdsadmin_dbms_scheduler.disable('SYS.CLEANUP_ONLINE_IND_BUILD');
END;
/
```

Enabling SYS-Owned Oracle Scheduler Jobs

To enable a SYS-owned Oracle Scheduler job, use the `rdsadmin.rdsadmin_dbms_scheduler.disable` procedure.

This procedure uses the `name` common parameter for Oracle Scheduler tasks. For more information, see [Common Parameters for Oracle Scheduler Procedures \(p. 1128\)](#).

The following example enables the `SYS.CLEANUP_ONLINE_IND_BUILD` Oracle Scheduler job.

```
BEGIN
    rdsadmin.rdsadmin_dbms_scheduler.enable('SYS.CLEANUP_ONLINE_IND_BUILD');
END;
/
```

Modifying the Repeat Interval for Jobs of CALENDAR Type

To modify the repeat interval to modify a SYS-owned Oracle Scheduler job of `CALENDAR` type, use the `rdsadmin.rdsadmin_dbms_scheduler.disable` procedure.

This procedure uses the following common parameters for Oracle Scheduler tasks:

- `name`
- `attribute`
- `value`

For more information, see [Common Parameters for Oracle Scheduler Procedures \(p. 1128\)](#).

The following example modifies the repeat interval of the `SYS.CLEANUP_ONLINE_IND_BUILD` Oracle Scheduler job.

```
BEGIN
    rdsadmin.rdsadmin_dbms_scheduler.set_attribute(
        name      => 'SYS.CLEANUP_ONLINE_IND_BUILD',
        attribute => 'repeat_interval',
        value     => 'freq=daily;byday=FRI,SAT;byhour=20;byminute=0;bysecond=0');
END;
/
```

Modifying the Repeat Interval for Jobs of NAMED Type

Some Oracle Scheduler jobs use a schedule name instead of an interval. For these type of jobs, you must create a new named schedule in the master user schema. Use the standard Oracle `sys.dbms_scheduler.create_schedule` procedure to do this. Also, use the `rdsadmin.rdsadmin_dbms_scheduler.set_attribute` procedure to assign the new named schedule to the job.

This procedure uses the following common parameter for Oracle Scheduler tasks:

- `name`
- `attribute`
- `value`

For more information, see [Common Parameters for Oracle Scheduler Procedures \(p. 1128\)](#).

The following example modifies the repeat interval of the `SYS.BSLN_MAINTAIN_STATS_JOB` Oracle Scheduler job.

```

BEGIN
    dbms_scheduler.create_schedule (
        schedule_name    => 'rds_master_user.new_schedule',
        start_date       => SYSTIMESTAMP,
        repeat_interval =>
        'freq=daily;byday=MON,TUE,WED,THU,FRI;byhour=0;byminute=0;bysecond=0',
        end_date         => NULL,
        comments         => 'Repeats daily forever');
END;
/

BEGIN
    rdsadmin.rdsadmin_dbms_scheduler.set_attribute (
        name      => 'SYS.BSLN_MAINTAIN_STATS_JOB',
        attribute => 'schedule_name',
        value     => 'rds_master_user.new_schedule');
END;
/

```

Common DBA Miscellaneous Tasks for Oracle DB Instances

This section describes how you can perform miscellaneous DBA tasks on your Amazon RDS DB instances running Oracle. To deliver a managed service experience, Amazon RDS doesn't provide shell access to DB instances, and restricts access to certain system procedures and tables that require advanced privileges.

Topics

- [Creating New Directories in the Main Data Storage Space \(p. 1131\)](#)
- [Listing Files in a DB Instance Directory \(p. 1132\)](#)
- [Reading Files in a DB Instance Directory \(p. 1132\)](#)

Creating New Directories in the Main Data Storage Space

To create directories, use the Amazon RDS procedure `rdsadmin.rdsadmin_util.create_directory`. You can create up to 10,000 directories, all located in your main data storage space.

The `create_directory` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>p_directory_name</code>	varchar2	—	Yes	The name of the new directory.

The following example creates a new directory named `product_descriptions`.

```
exec rdsadmin.rdsadmin_util.create_directory(p_directory_name => 'product_descriptions');
```

You can list the directories by querying `DBA_DIRECTORIES`. The system chooses the actual host pathname automatically. The following example gets the directory path for the directory named `product_descriptions`:

```
select DIRECTORY_PATH
      from DBA_DIRECTORIES
     where DIRECTORY_NAME='product_descriptions';

DIRECTORY_PATH
-----
/rdsdbdata/userdirs/01
```

The master user name for the DB instance has read and write privileges in the new directory, and can grant access to other users. Execute privileges are not available for directories on a DB instance. Directories are created in your main data storage space and will consume space and I/O bandwidth.

You can drop a directory that you created by using the Oracle `drop_directory` command. Dropping a directory doesn't remove its contents. Because the `create_directory()` method can reuse pathnames, files in dropped directories can appear in a newly created directory. Before you drop a directory, you should use `UTL_FILE.FREMOVE` to remove files from the directory. For more information, see [FREMOVE Procedure](#) in the Oracle documentation.

Listing Files in a DB Instance Directory

To list the files in a directory, use the Amazon RDS procedure `rdsadmin.rds_file_util.listdir`. The `listdir` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>p_directory</code>	varchar2	—	Yes	The name of the directory to list.

The following example lists the files in the directory named `product_descriptions`.

```
select * from table
  (rdsadmin.rds_file_util.listdir(p_directory => 'product_descriptions'));
```

Reading Files in a DB Instance Directory

To read a text file, use the Amazon RDS procedure `rdsadmin.rds_file_util.read_text_file`. The `read_text_file` procedure has the following parameters.

Parameter Name	Data Type	Default	Required	Description
<code>p_directory</code>	varchar2	—	Yes	The name of the directory that contains the file.
<code>p_filename</code>	varchar2	—	Yes	The name of the file to read.

The following example reads the file `rice.txt` from the directory `product_descriptions`.

```
select * from table
```

```
(rdsadmin.rds_file_util.read_text_file(  
    p_directory => 'product_descriptions',  
    p_filename  => 'rice.txt'));
```

Using Kerberos Authentication with Amazon RDS for Oracle

You can use Kerberos authentication to authenticate users when they connect to your Amazon RDS DB instance running Oracle. In this case, your DB instance works with AWS Directory Service for Microsoft Active Directory, also called AWS Managed Microsoft AD, to enable Kerberos authentication. When users authenticate with an Oracle DB instance joined to the trusting domain, authentication requests are forwarded to the directory that you create with AWS Directory Service.

Keeping all of your credentials in the same directory can save you time and effort. You have a centralized place for storing and managing credentials for multiple database instances. Using a directory can also improve your overall security profile.

Amazon RDS supports Kerberos authentication for Oracle DB instances in the following AWS Regions:

- US East (Ohio)
- US East (N. Virginia)
- US West (Oregon)
- Asia Pacific (Seoul)
- Asia Pacific (Singapore)
- Asia Pacific (Sydney)
- Asia Pacific (Tokyo)
- Canada (Central)
- Europe (Frankfurt)
- Europe (Ireland)
- Europe (London)

Note

Kerberos authentication isn't supported for DB instance classes that are deprecated for Oracle DB instances. For more information, see [DB Instance Class Support for Oracle \(p. 917\)](#).

To set up Kerberos authentication for an Oracle DB instance, complete the following general steps, described in more detail later:

1. Use AWS Managed Microsoft AD to create an AWS Managed Microsoft AD directory. You can use the AWS Management Console, the AWS CLI, or the AWS Directory Service API to create the directory.
2. Create an AWS Identity and Access Management (IAM) role that uses the managed IAM policy `AmazonRDSDirectoryServiceAccess`. The role allows Amazon RDS to make calls to your directory.

For the role to allow access, the AWS Security Token Service (AWS STS) endpoint must be activated in the correct AWS Region for your AWS account. AWS STS endpoints are active by default in all AWS Regions, and you can use them without any further actions. For more information, see [Activating and Deactivating AWS STS in an AWS Region](#) in the *IAM User Guide*.

3. Create and configure users in the AWS Managed Microsoft AD directory using the Microsoft Active Directory tools. For more information about creating users in your Microsoft Active Directory, see [Manage Users and Groups in AWS Managed Microsoft AD](#) in the *AWS Directory Service Administration Guide*.
4. Create or modify an Oracle DB instance either from the console, CLI, or RDS API using one of the following methods:
 - [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#)

- [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#)
- [Restoring from a DB Snapshot \(p. 266\)](#)
- [Restoring a DB Instance to a Specified Time \(p. 285\)](#)

When you create or modify the DB instance, provide the domain identifier (`d-* identifier`) that was generated when you created your directory and the name of the role you created. You can locate the DB instance in the same Amazon Virtual Private Cloud (VPC) as the directory or in a different VPC.

5. Use the Amazon RDS master user credentials to connect to the Oracle DB instance. Create the user in Oracle to be identified externally. Externally identified users can log in to the Oracle DB instance using Kerberos authentication.

To get Kerberos authentication using an on-premises or self-hosted Microsoft Active Directory, create a two-way forest trust. For more information on setting up forest trusts using AWS Directory Service, see [When to Create a Trust Relationship](#) in the *AWS Directory Service Administration Guide*.

Topics

- [Setting Up Kerberos Authentication for Oracle DB Instances \(p. 1135\)](#)
- [Managing a DB Instance in a Domain \(p. 1144\)](#)
- [Connecting to Oracle with Kerberos Authentication \(p. 1145\)](#)

Setting Up Kerberos Authentication for Oracle DB Instances

You use AWS Directory Service for Microsoft Active Directory, also called AWS Managed Microsoft AD, to set up Kerberos authentication for an Oracle DB instance. To set up Kerberos authentication, complete the following steps:

- [Step 1: Create a Directory Using the AWS Managed Microsoft AD \(p. 1135\)](#)
- [Step 2: Create a Forest Trust \(p. 1139\)](#)
- [Step 3: Create an IAM Role for Use by Amazon RDS \(p. 1139\)](#)
- [Step 4: Create and Configure Users \(p. 1140\)](#)
- [Step 5: Configure VPC Peering \(p. 1140\)](#)
- [Step 6: Create or Modify an Oracle DB Instance \(p. 1141\)](#)
- [Step 7: Create Kerberos Authentication Oracle Logins \(p. 1142\)](#)
- [Step 8: Configure an Oracle Client \(p. 1143\)](#)

Step 1: Create a Directory Using the AWS Managed Microsoft AD

AWS Directory Service creates a fully managed Microsoft Active Directory in the AWS Cloud. When you create an AWS Managed Microsoft AD directory, AWS Directory Service creates two domain controllers and DNS servers on your behalf. The directory servers are created in different subnets in a VPC. This redundancy helps ensure that your directory remains accessible even if a failure occurs.

When you create an AWS Managed Microsoft AD directory, AWS Directory Service performs the following tasks on your behalf:

- Sets up a Microsoft Active Directory within your VPC.
- Creates a directory administrator account with the user name `Admin` and the specified password. You use this account to manage your directory.

Note

Make sure to save this password. AWS Directory Service doesn't store this password, and it can't be retrieved or reset.

- Creates a security group for the directory controllers.

When you launch an AWS Directory Service for Microsoft Active Directory, AWS creates an Organizational Unit (OU) that contains all of your directory's objects. This OU, which has the NetBIOS name that you entered when you created your directory, is located in the domain root. The domain root is owned and managed by AWS.

The `admin` account that was created with your AWS Managed Microsoft AD directory has permissions for the most common administrative activities for your OU:

- Create, update, or delete users
- Add resources to your domain such as file or print servers, and then assign permissions for those resources to users in your OU
- Create additional OUs and containers
- Delegate authority
- Restore deleted objects from the Active Directory Recycle Bin
- Run AD and DNS Windows PowerShell modules on the Active Directory Web Service

The `admin` account also has rights to perform the following domain-wide activities:

- Manage DNS configurations (add, remove, or update records, zones, and forwarders)
- View DNS event logs
- View security event logs

To create a directory with AWS Managed Microsoft AD

1. In the [AWS Directory Service console](#) navigation pane, choose **Directories**, and then choose **Set up directory**.
2. Choose **AWS Managed Microsoft AD**. AWS Managed Microsoft AD is the only option currently supported for use with Amazon RDS.
3. Choose **Next**.
4. On the **Enter directory information** page, provide the following information:

Edition

Choose the edition that meets your requirements.

Directory DNS name

The fully qualified name for the directory, such as `corp.example.com`.

Note

The fully qualified domain name of the AWS Managed Microsoft AD must not be longer than 61 characters.

Directory NetBIOS name

An optional short name for the directory, such as `CORP`.

Directory description

An optional description for the directory.

Admin password

The password for the directory administrator. The directory creation process creates an administrator account with the user name **Admin** and this password.

The directory administrator password can't include the word "admin." The password is case-sensitive and must be 8–64 characters in length. It must also contain at least one character from three of the following four categories:

- Lowercase letters (a–z)
- Uppercase letters (A–Z)
- Numbers (0–9)
- Nonalphanumeric characters (~!@#\$%^&*_+=`|\{}[];"'<>,.?/)

Confirm password

Retype the administrator password.

5. Choose **Next**.
6. On the **Choose VPC and subnets** page, provide the following information:

VPC

Choose the VPC for the directory. The Oracle DB instance can be created in this same VPC or in a different VPC.

Subnets

Choose the subnets for the directory servers. The two subnets must be in different Availability Zones.

7. Choose **Next**.
8. Review the directory information. If changes are needed, choose **Previous** and make the changes. When the information is correct, choose **Create directory**.

Review & create

Review

Directory type	VPC
Microsoft AD	vpc-8b6b78e9 (172.31.0.0/16)
Directory DNS name	Subnets
corp.example.com	subnet-75128d10 (172.31.64.0/20, us-east-1a) subnet-f51665dd (172.31.32.0/20, us-east-1b)
Directory NetBIOS name	
CORP	
Directory description	
My directory	



Pricing

Edition	Free trial eligible Learn more 30-day limited trial
Standard	
~USD 86.4000/mo (USD 0.1200/hr)*	
* Includes two domain controllers, USD 43.2000/mo for each additional domain controller.	

[Cancel](#) [Previous](#) [Create directory](#)

It takes several minutes for the directory to be created. When it has been successfully created, the **Status** value changes to **Active**.

To see information about your directory, choose the directory ID in the directory listing. Make a note of the **Directory ID** value. You need this value when you create or modify your Oracle DB instance.

Directory details		
Directory type	VPC	Status
Microsoft AD	vpc-8b6b78e9	Creating
Edition	Subnets	Last updated
Standard	subnet-f51665dd subnet-75128d10	Friday, July 26, 2019
Directory ID	Availability zones	Launch time
d-9067004760	us-east-1a, us-east-1b	Friday, July 26, 2019
Directory DNS name	DNS address	
corp.example.com		
Directory NetBIOS name		
CORP		
Description - Edit		
My directory		

Step 2: Create a Forest Trust

If you plan to use AWS Managed Microsoft AD only, move on to [Step 3: Create an IAM Role for Use by Amazon RDS \(p. 1139\)](#).

To get Kerberos authentication using an on-premises or self-hosted Microsoft Active Directory, create a two-way forest trust. For more information about setting up forest trusts using AWS Directory Service, see [When to Create a Trust Relationship](#) in the *AWS Directory Service Administration Guide*.

Step 3: Create an IAM Role for Use by Amazon RDS

You must create an IAM role that uses the managed IAM policy `AmazonRDSDirectoryServiceAccess`. This role allows Amazon RDS to make calls to the AWS Directory Service for you. When you create the role, choose `Directory Service`, and attach the AWS managed policy `AmazonRDSDirectoryServiceAccess` to it.

For more information about creating IAM roles for a service, see [Creating a Role to Delegate Permissions to an AWS Service](#).

Optionally, you can create policies with the required permissions instead of using the managed IAM policy `AmazonRDSDirectoryServiceAccess`. In this case, the role must have the following IAM trust policy.

```
{  
    "Version": "2012-10-17",
```

```
"Statement": [
    {
        "Effect": "Allow",
        "Principal": {
            "Service": [
                "rds.amazonaws.com",
                "directoryservice.rds.amazonaws.com"
            ]
        },
        "Action": "sts:AssumeRole"
    }
]
```

The role must also have the following IAM role policy.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Action": [
                "ds:DescribeDirectories",
                "ds:AuthorizeApplication",
                "ds:UnauthorizeApplication",
                "ds:GetAuthorizedApplicationDetails"
            ],
            "Effect": "Allow",
            "Resource": "*"
        }
    ]
}
```

Step 4: Create and Configure Users

You can create users with the Active Directory Users and Computers tool, which is one of the Active Directory Domain Services and Active Directory Lightweight Directory Services tools. In this case, *users* are individual people or entities that have access to your directory.

To create users in an AWS Directory Service directory, you must be connected to a Windows-based Amazon EC2 instance that is a member of the AWS Directory Service directory. At the same time, you must be logged in as a user that has privileges to create users. For more information, see [Create a User](#) in the *AWS Directory Service Administration Guide*.

Step 5: Configure VPC Peering

If you plan to locate the directory and the DB instance in different VPCs, configure VPC peering by following the instructions in this step. If you plan to locate the directory and the DB instance in the same VPC, skip this step and move on to [Step 6: Create or Modify an Oracle DB Instance \(p. 1141\)](#).

If the same AWS account owns both VPCs, follow the instructions in [What is VPC Peering?](#). Specifically, complete the following steps:

1. Set up appropriate VPC routing rules to ensure the network traffic can flow both ways.
2. Ensure the DB instance's security group can receive ingress traffic from this security group.
3. Ensure that there is no network access control list (ACL) rule to block traffic.

If different AWS accounts own the VPCs, complete the following steps:

1. Configure VPC peering by following the instructions in [What is VPC Peering?](#) in the *Amazon Virtual Private Cloud VPC Peering*. Specifically, complete the following steps:
 - a. Set up appropriate VPC routing rules to ensure the network traffic can flow both ways.
 - b. Ensure that there is no ACL rule to block traffic.
2. Initiate sharing of the directory with the AWS account that the DB instance will be created in by following the instructions in [Tutorial: Sharing Your AWS Managed Microsoft AD Directory for Seamless EC2 Domain-Join](#) in the *AWS Directory Service Administration Guide*.
3. Log in to the AWS Directory Service console using the account for the DB instance, and ensure that the domain has the SHARED status before proceeding.
4. While logged into the AWS Directory Service console using the account for the DB instance, make a note of the **Directory ID** value for the directory.

Step 6: Create or Modify an Oracle DB Instance

Create or modify an Oracle DB instance for use with your directory. You can use the console, CLI, or RDS API to associate a DB instance with a directory. You can do this in one of the following ways:

- Create a new Oracle DB instance using the console, the [create-db-instance](#) CLI command, or the [CreateDBInstance](#) RDS API operation.

For instructions, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).

- Modify an existing Oracle DB instance using the console, the [modify-db-instance](#) CLI command, or the [ModifyDBInstance](#) RDS API operation.

For instructions, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

- Restore an Oracle DB instance from a DB snapshot using the console, the [restore-db-instance-from-db-snapshot](#) CLI command, or the [RestoreDBInstanceFromDBSnapshot](#) RDS API operation.

For instructions, see [Restoring from a DB Snapshot \(p. 266\)](#).

- Restore an Oracle DB instance to a point-in-time using the console, the [restore-db-instance-to-point-in-time](#) CLI command, or the [RestoreDBInstanceToPointInTime](#) RDS API operation.

For instructions, see [Restoring a DB Instance to a Specified Time \(p. 285\)](#).

Kerberos authentication is only supported for Oracle DB instances in a VPC. The DB instance can be in the same VPC as the directory, or in a different VPC. The DB instance must use a security group that allows egress within the directory's VPC so the DB instance can communicate with the directory.

When you use the console, choose **Password and Kerberos authentication** in the **Database authentication** section. Choose the directory or choose **Create a new Directory**.

Database authentication

Database authentication options [Info](#)

Password authentication
Authenticates using database passwords.

Password and Kerberos authentication
Choose a directory in which you want to allow authorized users to authenticate with this DB instance using Kerberos Authentication.

Directory

▼ C

[Create a new directory](#) 

 By choosing a directory and continuing with database instance creation you authorize Amazon RDS to create the IAM role necessary for using Kerberos Authentication

When you use the AWS CLI, the following parameters are required for the DB instance to be able to use the directory that you created:

- For the `--domain` parameter, use the domain identifier ("d-*" identifier) generated when you created the directory.
- For the `--domain-iam-role-name` parameter, use the role you created that uses the managed IAM policy `AmazonRDSDirectoryServiceAccess`.

For example, the following CLI command modifies a DB instance to use a directory.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
    --db-instance-identifier mydbinstance \
    --domain d-ID \
    --domain-iam-role-name role-name
```

For Windows:

```
aws rds modify-db-instance ^
    --db-instance-identifier mydbinstance ^
    --domain d-ID ^
    --domain-iam-role-name role-name
```

Important

If you modify a DB instance to enable Kerberos authentication, reboot the DB instance after making the change.

Step 7: Create Kerberos Authentication Oracle Logins

Use the Amazon RDS master user credentials to connect to the Oracle DB instance as you do any other DB instance. The DB instance is joined to the AWS Managed Microsoft AD domain. Thus, you can provision Oracle logins and users from the Microsoft Active Directory users and groups in your domain. To manage database permissions, you grant and revoke standard Oracle permissions to these logins.

To allow a Microsoft Active Directory user to authenticate with Oracle, use the Amazon RDS master user credentials. You use these credentials to connect to the Oracle DB instance as you do any other DB instance. After you're logged in, create an externally authenticated user in Oracle.

```
CREATE USER "KRBUSER@CORP.EXAMPLE.COM" IDENTIFIED EXTERNALLY;  
GRANT CREATE SESSION TO "KRBUSER@CORP.EXAMPLE.COM";
```

Replace `KRBUSER@CORP.EXAMPLE.COM` with the user name and domain name. Users (both humans and applications) from your domain can now connect to the RDS Oracle instance from a domain joined client machine using Kerberos authentication.

Step 8: Configure an Oracle Client

To configure an Oracle client, meet the following requirements:

- Create a krb5.conf file (or equivalent) to point to the domain. Configure the Oracle client to use this krb5.conf file.
- Verify that traffic can flow between the client host and AWS Directory Service over DNS port 53 and Kerberos ports (88 and 464 for managed AWS Directory Service) over TCP/UDP.
- Verify that traffic can flow between the client host and the DB instance over the database port.

The following is sample krb5.conf content for AWS Managed Microsoft AD:

```
[libdefaults]  
default_realm = EXAMPLE.COM  
default_ccache_name = /tmp/kerbcache  
[realms]  
EXAMPLE.COM = {  
    kdc = example.com  
    admin_server = example.com  
}  
[domain_realm]  
.example.com = EXAMPLE.COM  
example.com = EXAMPLE.COM
```

The following is sample krb5.conf content for on-premise Microsoft AD:

```
[libdefaults]  
default_realm = EXAMPLE.COM  
default_ccache_name = /tmp/kerbcache  
[realms]  
EXAMPLE.COM = {  
    kdc = example.com  
    admin_server = example.com  
}  
ONPREM.COM = {  
    kdc = onprem.com  
    admin_server = onprem.com  
}  
[domain_realm]  
.example.com = EXAMPLE.COM  
example.com = EXAMPLE.COM  
.onprem.com = ONPREM.COM  
onprem.com = ONPREM.COM
```

The following is sample sqlnet.ora content for a SQL*Plus configuration:

```
SOLNET.AUTHENTICATION_SERVICES=(KERBEROS5PRE,KERBEROS5)
SQLNET.KERBEROS5_CONF=path_to_krb5.conf_file
```

For an example of a SQL Developer configuration, see [Document 1609359.1](#) from Oracle Support.

Managing a DB Instance in a Domain

You can use the console, the CLI, or the RDS API to manage your DB instance and its relationship with your Microsoft Active Directory. For example, you can associate a Microsoft Active Directory to enable Kerberos authentication. You can also disassociate a Microsoft Active Directory to disable Kerberos authentication. You can also move a DB instance to be externally authenticated by one Microsoft Active Directory to another.

For example, using the CLI, you can do the following:

- To reattempt enabling Kerberos authentication for a failed membership, use the [modify-db-instance](#) CLI command and specify the current membership's directory ID for the --domain option.
- To disable Kerberos authentication on a DB instance, use the [modify-db-instance](#) CLI command and specify none for the --domain option.
- To move a DB instance from one domain to another, use the [modify-db-instance](#) CLI command and specify the domain identifier of the new domain for the --domain option.

Understanding Domain Membership

After you create or modify your DB instance, the DB instance becomes a member of the domain. You can view the status of the domain membership for the DB instance in the console or by running the [describe-db-instances](#) CLI command. The status of the DB instance can be one of the following:

- **kerberos-enabled** – The DB instance has Kerberos authentication enabled.
- **enabling-kerberos** – AWS is in the process of enabling Kerberos authentication on this DB instance.
- **pending-enable-kerberos** – Enabling Kerberos authentication is pending on this DB instance.
- **pending-maintenance-enable-kerberos** – AWS will attempt to enable Kerberos authentication on the DB instance during the next scheduled maintenance window.
- **pending-disable-kerberos** – Disabling Kerberos authentication is pending on this DB instance.
- **pending-maintenance-disable-kerberos** – AWS will attempt to disable Kerberos authentication on the DB instance during the next scheduled maintenance window.
- **enable-kerberos-failed** – A configuration problem has prevented AWS from enabling Kerberos authentication on the DB instance. Correct the configuration problem before reissuing the command to modify the DB instance.
- **disabling-kerberos** – AWS is in the process of disabling Kerberos authentication on this DB instance.

A request to enable Kerberos authentication can fail because of a network connectivity issue or an incorrect IAM role. If the attempt to enable Kerberos authentication fails when you create or modify a DB instance, first make sure that you are using the correct IAM role. Then modify the DB instance to join the domain.

Note

Only Kerberos authentication with Amazon RDS for Oracle sends traffic to the domain's DNS servers. All other DNS requests are treated as outbound network access on your DB instances running Oracle. For more information about outbound network access with Amazon RDS for Oracle, see [Setting Up a Custom DNS Server \(p. 1094\)](#).

Force-Rotating Kerberos Keys

A secret key is shared between AWS Managed Microsoft AD and Amazon RDS for Oracle DB instance. This key is rotated automatically every 45 days. You can use the following Amazon RDS procedure to force the rotation of this key:

```
SELECT rdsadmin.rdsadmin_kerberos_auth_tasks.rotate_kerberos_keytab AS TASK_ID FROM DUAL;
```

Note

In a Read Replica configuration, this procedure is available only on the source DB instance and not on the Read Replica.

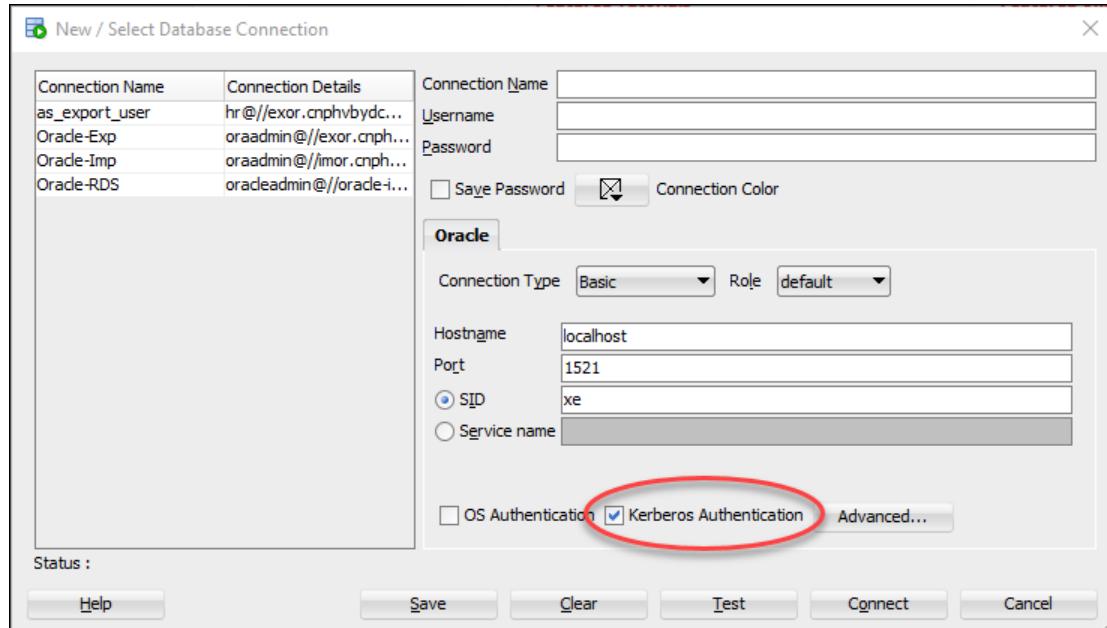
You can view the status of an ongoing task in a bdump file. The bdump files are located in the /rdsdbdata/log/trace directory. Each bdump file name is in the following format.

```
dbtask-task-id.log
```

Replace **task-id** with the identifier of the task that you want to monitor.

Connecting to Oracle with Kerberos Authentication

To connect to Oracle with Kerberos authentication, log in using the Kerberos authentication type. For example, after launching Oracle SQL Developer, choose **Kerberos Authentication** as the authentication type, as shown following.



To connect to Oracle with Kerberos authentication with SQL*Plus, take the following steps:

1. At a command prompt, run the following command:

```
kinit username
```

Replace *username* with the user name and, at the prompt, enter the password stored in the Microsoft Active Directory for the user.

2. Open SQL*Plus and connect using the DNS name and port number for the Oracle DB instance.

For more information about connecting to an Oracle DB instance in SQL*Plus, see [Connecting to Your DB Instance Using SQL*Plus \(p. 959\)](#).

Tools and Third-Party Software for Oracle DB Instances

This section provides information about tools and third-party software for Oracle DB instances on Amazon RDS.

Topics

- [Setting Up Amazon RDS to Host Tools and Third-Party Software for Oracle \(p. 1147\)](#)
- [Using Oracle GoldenGate with Amazon RDS \(p. 1153\)](#)
- [Using the Oracle Repository Creation Utility on Amazon RDS for Oracle \(p. 1165\)](#)
- [Installing a Siebel Database on Oracle on Amazon RDS \(p. 1170\)](#)

Setting Up Amazon RDS to Host Tools and Third-Party Software for Oracle

You can use Amazon RDS to host an Oracle DB instance that supports software and components such as the following:

- Siebel Customer Relationship Management (CRM)
- Oracle Fusion Middleware Metadata — installed by the Repository Creation Utility (RCU)

The following procedures help you create an Oracle DB instance on Amazon RDS that you can use to host additional software and components for Oracle.

Creating a VPC for Use with an Oracle Database

In the following procedure, you create a virtual private cloud (VPC) based on the Amazon VPC service, a private subnet, and a security group. Your Amazon RDS DB instance needs to be available only to your middle-tier components, and not to the public internet. Thus, your Amazon RDS DB instance is hosted in a private subnet, providing greater security.

To create a VPC based on Amazon VPC

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. In the upper-right corner of the AWS Management Console, choose the AWS Region for your VPC. This example uses the US West (Oregon) region.
3. In the upper-left corner, choose **VPC Dashboard**, and then choose **Start VPC Wizard**.
4. On the page **Step 1: Select a VPC Configuration**, choose **VPC with Public and Private Subnets**, and then choose **Select**.
5. On the page **Step 2: VPC with Public and Private Subnets**, shown following, set the following values.

Option	Value
IPv4 CIDR block	10.0.0.0/16 For more information about selecting CIDR blocks for your VPC, see VPC Sizing .

Option	Value
IPv6 CIDR block	No IPv6 CIDR Block
VPC name	The name for your VPC, for example vpc-1 .
Public subnet's IPv4 CIDR	10.0.0.0/24 For more information about subnet sizing, see Subnet Sizing .
Availability Zone	An Availability Zone for your AWS Region.
Public subnet name	The name for your public subnet, for example subnet-public-1 .
Private subnet's IPv4 CIDR	10.0.1.0/24 For more information about subnet sizing, see Subnet Sizing .
Availability Zone	An Availability Zone for your AWS Region.
Private subnet name	The name for your private subnet, for example subnet-private-1 .
Instance type	An instance type for your NAT instance, for example t2.small . Note If you don't see Instance type in the console, choose Use a NAT instance instead .
Key pair name	No key pair
Service endpoints	None
Enable DNS hostnames	Yes
Hardware tenancy	Default

Step 2: VPC with Public and Private Subnets

IPv4 CIDR block: <input type="text" value="10.0.0.0/16"/>	(65531 IP addresses available)
IPv6 CIDR block:	<input checked="" type="radio"/> No IPv6 CIDR Block <input type="radio"/> Amazon provided IPv6 CIDR block
VPC name:	<input type="text" value="vpc-1"/>
Public subnet's IPv4 CIDR: <input type="text" value="10.0.0.0/24"/>	(251 IP addresses available)
Availability Zone: <input type="text" value="us-west-2a"/>	
Public subnet name:	<input type="text" value="subnet-public-1"/>
Private subnet's IPv4 CIDR: <input type="text" value="10.0.1.0/24"/>	(251 IP addresses available)
Availability Zone: <input type="text" value="us-west-2a"/>	
Private subnet name:	<input type="text" value="subnet-private-1"/>
You can add more subnets after AWS creates the VPC.	
Specify the details of your NAT instance (Instance rates apply). Use a NAT gateway instead	
Instance type: <input type="text" value="t2.small"/>	
Key pair name: <input type="text" value="No key pair"/>	
Service endpoints	
<input type="button" value="Add Endpoint"/>	
Enable DNS hostnames: <input checked="" type="radio"/> Yes <input type="radio"/> No	
Hardware tenancy: <input type="text" value="Default"/>	
<input type="button" value="Cancel and Exit"/> <input type="button" value="Back"/> <input type="button" value="Create VPC"/>	

6. Choose **Create VPC**.

An Amazon RDS DB instance in a VPC requires at least two private subnets or at least two public subnets, to support Multi-AZ deployment. For more information about working with multiple Availability Zones, see [Regions and Availability Zones \(p. 132\)](#). Because your database is private, add a second private subnet to your VPC.

To create an additional subnet

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. In the upper-right corner of the AWS Management Console, confirm that you are in the correct AWS Region for your VPC.
3. In the upper-left corner, choose **VPC Dashboard**, choose **Subnets**, and then choose **Create Subnet**.

4. On the **Create Subnet** page, set the following values.

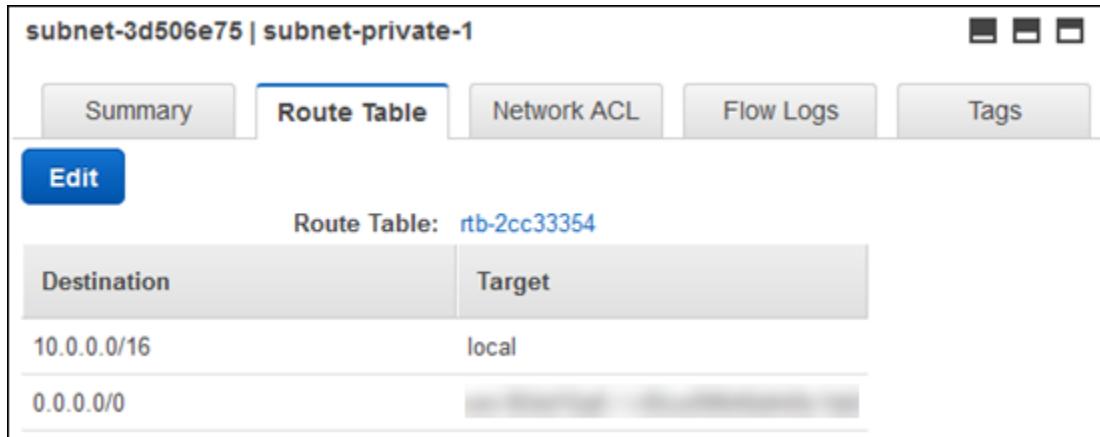
Option	Value
Name tag	The name for your second private subnet, for example subnet-private-2 .
VPC	Your VPC, for example vpc-1 .
Availability Zone	An Availability Zone for your AWS Region. Note Choose an Availability Zone different from the one that you chose for the first private subnet.
CIDR block	10.0.2.0/24

5. Choose **Yes, Create**.

Both private subnets must use the same route table. In the following procedure, you check to make sure the route tables match, and if not you edit one of them.

To ensure the subnets use the same route table.

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. In the upper-right corner of the AWS Management Console, confirm that you are in the correct AWS Region for your VPC.
3. In the upper-left corner, choose **VPC Dashboard**, choose **Subnets**, and then choose your first private subnet, for example **subnet-private-1**.
4. At the bottom of the console, choose the **Route Table** tab, shown following.



5. Make a note of the route table, for example **rtb-0d9fc668**.
6. In the list of subnets, choose the second private subnet, for example **subnet-private-2**.
7. At the bottom of the console, choose the **Route Table** tab.
8. If the route table for the second subnet is not the same as the route table for the first subnet, edit it to match:
 - a. Choose **Edit**.
 - b. For **Change to**, choose the route table that matches your first subnet.
 - c. Choose **Save**.

A security group acts as a virtual firewall for your DB instance to control inbound and outbound traffic. In the following procedure, you create a security group for your DB instance. For more information about security groups, see [Security Groups for Your VPC](#).

To create a VPC security group for a Private Amazon RDS DB Instance

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. In the upper-right corner of the AWS Management Console, confirm that you are in the correct AWS Region for your VPC.
3. In the upper-left corner, choose **VPC Dashboard**, choose **Security Groups**, and then choose **Create Security Group**.
4. On the page **Create Security Group**, set the following values.

Option	Value
Name tag	The name for your security group, for example sgdb-1 .
Group name	The name for your security group, for example sgdb-1 .
Description	A description for your security group.
VPC	Your VPC, for example vpc-1 .

5. Choose **Yes, Create**.

In the following procedure, you add rules to your security group to control inbound traffic to your DB instance. For more information about inbound rules, see [Security Group Rules](#).

To add inbound rules to the security group

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. In the upper-right corner of the AWS Management Console, confirm that you are in the correct AWS Region for your VPC.
3. In the upper-left corner, choose **VPC Dashboard**, choose **Security Groups**, and then choose your security group, for example **sgdb-1**.
4. At the bottom of the console, choose the **Inbound Rules** tab, and then choose **Edit**.
5. Set these values, as shown following.

Option	Value
Type	Oracle (1521)
Protocol	TCP (6)
Port Range	1521
Source	The identifier of your security group. When you choose the box, you see the name of your security group, for example sgdb-1 .



6. Choose **Save**.

Creating an Oracle DB Instance

You can use Amazon RDS to host an Oracle DB instance. When you create the new DB instance, specify the VPC and security group you created previously using the instructions in [Creating a VPC for Use with an Oracle Database \(p. 1147\)](#). Also, choose **No** for **Publicly accessible**.

For information about creating an Oracle DB instance, see [Creating a DB Instance Running the Oracle Database Engine \(p. 942\)](#).

Additional Amazon RDS Interfaces

In the preceding procedures, we use the AWS Management Console to perform tasks. Amazon Web Services also provides the AWS Command Line Interface (AWS CLI), and an application programming interface (API). You can use the AWS CLI or the API to automate many of the tasks for managing Amazon RDS, including tasks to manage an Oracle DB instance with Amazon RDS.

For more information, see [AWS Command Line Interface Reference for Amazon RDS](#) and [Amazon RDS API Reference](#).

Related Topics

- [Setting Up for Amazon RDS \(p. 5\)](#)
- [Using the Oracle Repository Creation Utility on Amazon RDS for Oracle \(p. 1165\)](#)
- [Installing a Siebel Database on Oracle on Amazon RDS \(p. 1170\)](#)
- [Scenarios for Accessing a DB Instance in a VPC \(p. 507\)](#)
- [Connecting to a DB Instance Running the Oracle Database Engine \(p. 955\)](#)

Using Oracle GoldenGate with Amazon RDS

Oracle GoldenGate (GoldenGate) is used to collect, replicate, and manage transactional data between databases. It is a log-based change data capture (CDC) and replication software package used with Oracle databases for online transaction processing (OLTP) systems. GoldenGate creates trail files that contain the most recent changed data from the source database and then pushes these files to the target database. You can use GoldenGate with Amazon RDS for Active-Active database replication, zero-downtime migration and upgrades, disaster recovery, data protection, and in-region and cross-region replication.

The following are important points to know when working with GoldenGate on Amazon RDS:

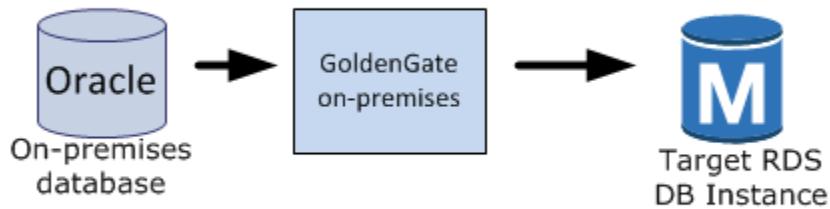
- You are responsible for setting up and managing GoldenGate for use with Amazon RDS.
- You are responsible for managing GoldenGate licensing (bring-your-own-license) for use with Amazon RDS in all AWS regions. For more information, see [Oracle Licensing \(p. 916\)](#).
- Amazon RDS supports GoldenGate for Oracle Database Standard Edition Two (SE2), Standard Edition One (SE1), Standard Edition (SE), and Enterprise Edition (EE).
- Amazon RDS supports GoldenGate for database version 11.2.0.4, 12.1.0.2, 12.2.0.1, 18.0.0.0, and 19.0.0.
- Amazon RDS supports GoldenGate version 11.2.1 and later, including 12.1, 12.2, and 12.3.
- Amazon RDS supports migration and replication across Oracle databases using GoldenGate. We do not support nor prevent customers from migrating or replicating across heterogeneous databases.
- You can use GoldenGate on Amazon RDS Oracle DB instances that use Oracle Transparent Data Encryption (TDE). To maintain the integrity of replicated data, you should configure encryption on the GoldenGate hub using EBS encrypted volumes or trail file encryption. You should also configure encryption for data sent between the GoldenGate hub and the source and target database instances. Amazon RDS Oracle DB instances support encryption with [Oracle Secure Sockets Layer \(p. 1051\)](#) or [Oracle Native Network Encryption \(p. 1049\)](#).
- GoldenGate DDL is supported with GoldenGate version 12.1 and later when using Integrated capture mode.

Overview

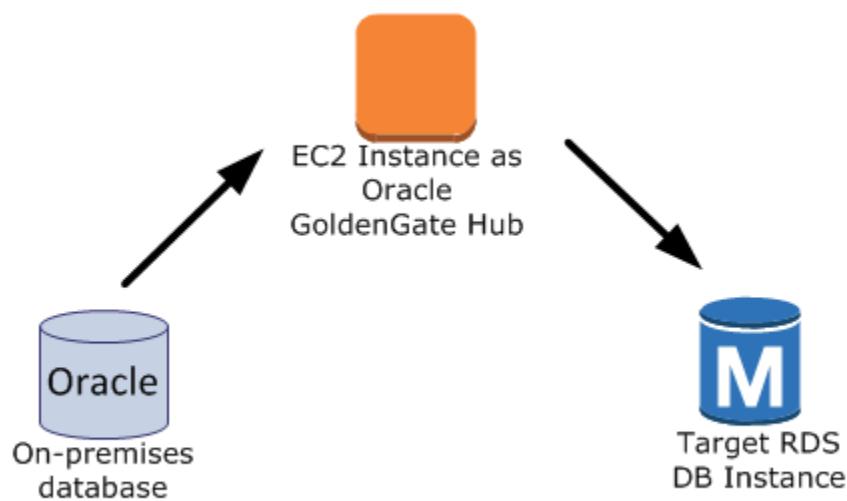
The GoldenGate architecture for use with Amazon RDS consists of three decoupled modules. The source database can be either an on-premises Oracle database, an Oracle database on an Amazon EC2 instance, or an Oracle database on an Amazon RDS DB instance. Next, the GoldenGate hub, which moves transaction information from the source database to the target database, can be either an Amazon EC2 instance with Oracle Database and with GoldenGate installed, or an on-premises Oracle installation. You can have more than one Amazon EC2 hub, and we recommend that you use two hubs if you are using GoldenGate for cross-region replication. Finally, the target database can be either on an Amazon RDS DB instance, on an Amazon EC2 instance, or on an on-premises location.

GoldenGate on Amazon RDS supports the following common scenarios:

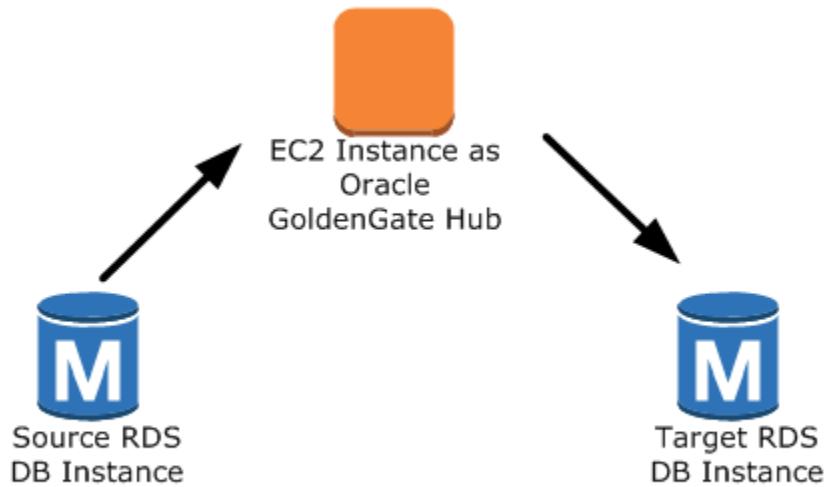
Scenario 1: An on-premises Oracle source database and on-premises GoldenGate hub, that provides data to a target Amazon RDS DB instance.



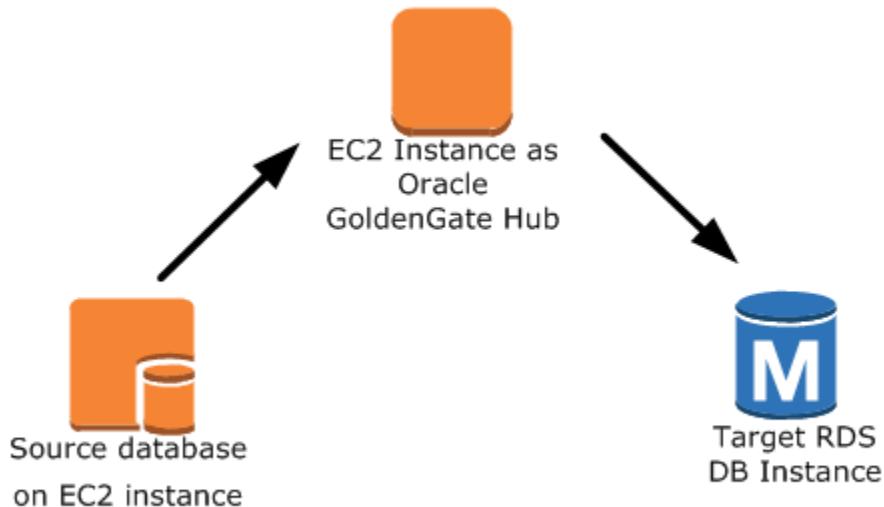
Scenario 2: An on-premises Oracle database that acts as the source database, connected to an Amazon EC2 instance hub that provides data to a target Amazon RDS DB instance.



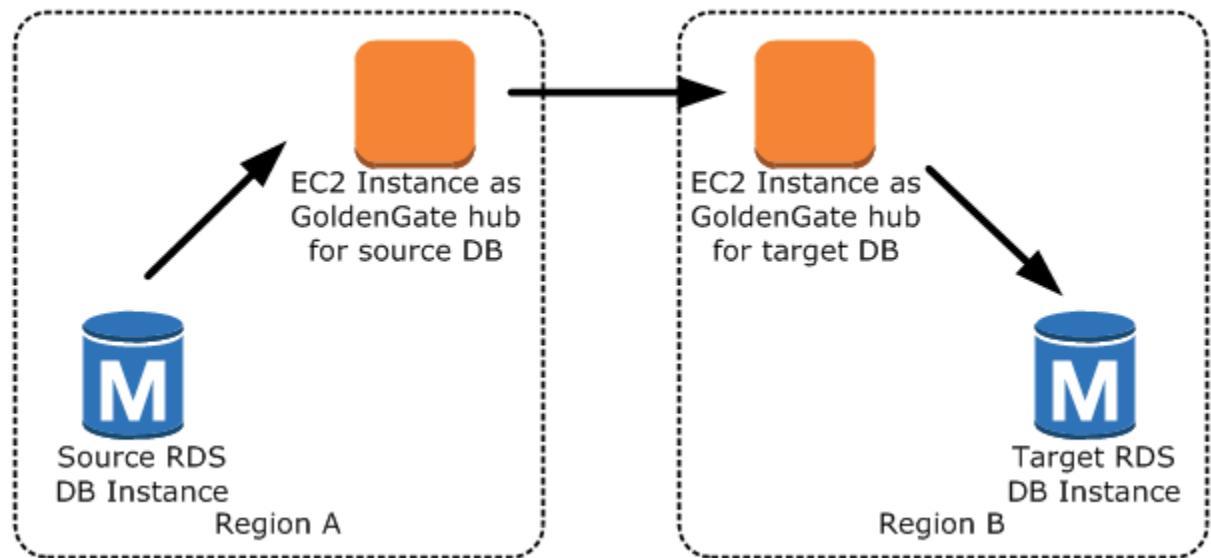
Scenario 3: An Oracle database on an Amazon RDS DB instance that acts as the source database, connected to an Amazon EC2 instance hub that provides data to a target Amazon RDS DB instance.



Scenario 4: An Oracle database on an Amazon EC2 instance that acts as the source database, connected to an Amazon EC2 instance hub that provides data to a target Amazon RDS DB instance.



Scenario 5: An Oracle database on an Amazon RDS DB instance connected to an Amazon EC2 instance hub in the same region, connected to an Amazon EC2 instance hub in a different region that provides data to the target Amazon RDS DB instance in the same region as the second Amazon EC2 instance hub.



Note

Any issues that impact running GoldenGate on an on-premises environment will also impact running GoldenGate on AWS. We strongly recommend that you monitor the GoldenGate hub to ensure that EXTRACT and REPLICAT are resumed if a failover occurs. Since the GoldenGate hub is run on an Amazon EC2 instance, Amazon RDS does not manage the GoldenGate hub and cannot ensure that it is running.

You can use GoldenGate using Amazon RDS to upgrade to major versions of Oracle. For example, you can use GoldenGate using Amazon RDS to upgrade from an Oracle version 8 on-premises database to an Oracle database running version 11.2.0.4 on an Amazon RDS DB instance.

To set up GoldenGate using Amazon RDS, you configure the hub on the Amazon EC2 instance, and then configure the source and target databases. The following steps show how to set up GoldenGate for use with Amazon RDS. Each step is explained in detail in the following sections:

- [Setting Up a GoldenGate Hub on Amazon EC2 \(p. 1156\)](#)
- [Setting Up a Source Database for Use with GoldenGate on Amazon RDS \(p. 1157\)](#)
- [Setting Up a Target Database for Use with GoldenGate on Amazon RDS \(p. 1160\)](#)
- [Working with the EXTRACT and REPLICAT Utilities of GoldenGate \(p. 1161\)](#)

Setting Up a GoldenGate Hub on Amazon EC2

You must complete several steps when you create a GoldenGate hub on an Amazon EC2 instance. First, you create an Amazon EC2 instance with a full client installation of Oracle RDBMS. For Oracle version 11.2.0.4, you must install patch 13328193. The Amazon EC2 instance must also have Oracle GoldenGate software installed. The exact software versions depend on the source and target database versions. For more information about installing GoldenGate, see the [Oracle documentation](#)

Since the Amazon EC2 instance that is serving as the GoldenGate hub stores and processes the transaction information from the source database into trail files, you must have enough allocated storage to store the trail files. You must also ensure that the Amazon EC2 instance has enough processing power to manage the amount of data being processed and enough memory to store the transaction information before it is written to the trail file.

The following tasks set up a GoldenGate hub on an Amazon EC2 instance; each task is explained in detail in this section. The tasks include:

- Create the GoldenGate subdirectories
- Update the GLOBALS parameter file
- Configure the *mgr.prm* file and start the *manager*

Create subdirectories in the GoldenGate directory using the Amazon EC2 command line shell and *ggsci*, the GoldenGate command interpreter. The subdirectories are created under the *gg* directory and include directories for parameter, report, and checkpoint files.

```
prompt$ cd /gg
prompt$ ./ggsci
GGSCI> CREATE SUBDIRS
```

Create a GLOBALS parameter file using the Amazon EC2 command line shell. Parameters that affect all GoldenGate processes are defined in the GLOBALS parameter file. The following example creates the necessary file:

```
$ cd $GGHOME
$ vi GLOBALS
CheckpointTable oggadm1.oggchkpt
```

The last step in setting up and configuring the GoldenGate hub is to configure the *manager*. Add the following lines to the *mgr.prm* file, then start the *manager* using *ggsci*:

```
PORT 8199
PurgeOldExtracts ./dirdat/*, UseCheckpoints, MINKEEPDAYS 5
```

```
GGSCI> start mgr
```

Once you have completed these steps, the GoldenGate hub is ready for use. Next, you set up the source and target databases.

Setting Up a Source Database for Use with GoldenGate on Amazon RDS

When your source database is running version 11.2.0.4 or later, complete the following tasks to set up a source database for use with GoldenGate:

- Set the *compatible* parameter to 11.2.0.4 or later.
- Set the *ENABLE_GOLDENGATE_REPLICATION* parameter to *True*. This parameter turns on supplemental logging for the source database. If your source database is on an Amazon RDS DB instance, you must have a parameter group assigned to the DB instance with the *ENABLE_GOLDENGATE_REPLICATION* parameter set to *true*. For more information about the *ENABLE_GOLDENGATE_REPLICATION* parameter, see the [Oracle documentation](#).
- Set the retention period for archived redo logs for the GoldenGate source database.
- Create a GoldenGate user account on the source database.
- Grant the necessary privileges to the GoldenGate user.
- Add a TNS alias for the source database to *tnsnames.ora* on the GoldenGate hub.

The source database must have the *compatible* parameter set to 11.2.0.4 or later. If you are using an Oracle database on an Amazon RDS DB instance as the source database, you must have a parameter

group with the compatible parameter set to 11.2.0.4 or later associated with the DB instance. If you change the compatible parameter in a parameter group associated with the DB instance, the change requires an instance reboot. You can use the following Amazon RDS CLI commands to create a new parameter group and set the compatible parameter. Note that you must associate the new parameter group with the source DB instance:

For Linux, OS X, or Unix:

```
aws rds create-db-parameter-group \
    --db-parameter-group-name example-goldengate \
    --description "Parameters to allow GoldenGate" \
    --db-parameter-group-family oracle-ee-11.2

aws rds modify-db-parameter-group \
    --db-parameter-group-name example-goldengate \
    --parameters "ParameterName=compatible, ParameterValue=11.2.0.4, ApplyMethod=pending-reboot"

aws rds modify-db-instance \
    --db-instance-identifier example-test \
    --db-parameter-group-name example-goldengate \
    --apply-immediately

aws rds reboot-db-instance \
    --db-instance-identifier example-test
```

For Windows:

```
aws rds create-db-parameter-group ^
    --db-parameter-group-name example-goldengate ^
    --description "Parameters to allow GoldenGate" ^
    --db-parameter-group-family oracle-ee-11.2

aws rds modify-db-parameter-group ^
    --db-parameter-group-name example-goldengate ^
    --parameters "ParameterName=compatible, ParameterValue=11.2.0.4, ApplyMethod=pending-reboot"

aws rds modify-db-instance ^
    --db-instance-identifier example-test ^
    --db-parameter-group-name example-goldengate ^
    --apply-immediately

aws rds reboot-db-instance ^
    --db-instance-identifier example-test
```

Always retain the parameter group with the compatible parameter. If you restore an instance from a DB snapshot, you must modify the restored instance to use the parameter group that has a matching or greater compatible parameter value. This should be done as soon as possible after the restore action and will require a reboot of the instance.

The `ENABLE_GOLDENGATE_REPLICATION` parameter, when set to *True*, turns on supplemental logging for the source database and configures the required GoldenGate permissions. If your source database is on an Amazon RDS DB instance, you must have a parameter group assigned to the DB instance with the `ENABLE_GOLDENGATE_REPLICATION` parameter set to *true*. For more information about the `ENABLE_GOLDENGATE_REPLICATION` parameter, see the [Oracle documentation](#).

The source database must also retain archived redo logs. For example, the following command sets the retention period for archived redo logs to 24 hours:

```
exec rdsadmin.rdsadmin_util.set_configuration('archivelog retention hours',24);
```

The duration for log retention is specified in hours. The duration should exceed any potential downtime of the source instance or any potential communication/networking issues to the source instance, so that GoldenGate can recover logs from the source instance as needed. The absolute minimum value required is one (1) hour of logs retained.

A log retention setting that is too small will result in the following message:

```
ERROR OGG-02028 Failed to attach to logmining server OGG$<extract_name> error 26927 -  
ORA-26927: altering an outbound server with a remote capture is not allowed.
```

Because these logs are retained on your DB instance, you need to ensure that you have enough storage available on your instance to accommodate the log files. To see how much space you have used in the last "X" hours, use the following query, replacing "X" with the number of hours.

```
select sum(blocks * block_size) bytes from v$archived_log  
where next_time>=sysdate-X/24 and dest_id=1;
```

GoldenGate runs as a database user and must have the appropriate database privileges to access the redo and archive logs for the source database, so you must create a GoldenGate user account on the source database. For more information about the permissions for a GoldenGate user account, see the sections 4, section 4.4, and table 4.1 in the [Oracle documentation](#).

The following statements create a user account named *oggadm1*:

```
CREATE tablespace administrator;  
CREATE USER oggadm1 IDENTIFIED BY "XXXXXX"  
default tablespace ADMINISTRATOR temporary tablespace TEMP;
```

Finally, grant the necessary privileges to the GoldenGate user account. The following statements grant privileges to a user named *oggadm1*:

```
grant create session, alter session to oggadm1;  
grant resource to oggadm1;  
grant select any dictionary to oggadm1;  
grant flashback any table to oggadm1;  
grant select any table to oggadm1;  
grant select_catalog_role to <RDS instance master username> with admin option;  
exec RDSADMIN.RDSADMIN_UTIL.GRANT_SYS_OBJECT ('DBA_CLUSTERS', 'OGGADM1');  
grant execute on dbms_flashback to oggadm1;  
grant select on SYS.V_$DATABASE to oggadm1;  
grant alter any table to oggadm1;  
  
EXEC DBMS_GOLDENGATE_AUTH.GRANT_ADMIN_PRIVILEGE (grantee=>'OGGADM1',  
privilege_type=>'capture',  
grant_select_privileges=>true,  
do_grants=>TRUE);
```

Add the following entry to \$ORACLE_HOME/network/admin/tnsnames.ora in the Oracle Home that will be used by the EXTRACT process. For more information on the tnsname.ora file, see the [Oracle documentation](#).

```
OGGSOURCE=  
(DESCRIPTION=  
  (ENABLE=BROKEN)  
  (ADDRESS_LIST=  
    (ADDRESS=(PROTOCOL=TCP)(HOST=goldengate-source.abcdef12345.us-west-2.rds.amazonaws.com)  
    (PORT=8200))
```

```
)  
(CONNECT_DATA=  
  (SID=ORCL)  
)  
)
```

Setting Up a Target Database for Use with GoldenGate on Amazon RDS

The following tasks set up a target DB instance for use with GoldenGate:

- Set the `compatible` parameter to 11.2.0.4 or later
- Set the `ENABLE_GOLDENGATE_REPLICATION` parameter to *True*. If your target database is on an Amazon RDS DB instance, you must have a parameter group assigned to the DB instance with the `ENABLE_GOLDENGATE_REPLICATION` parameter set to *true*. For more information about the `ENABLE_GOLDENGATE_REPLICATION` parameter, see the [Oracle documentation](#).
- Create and manage a GoldenGate user account on the target database
- Grant the necessary privileges to the GoldenGate user
- Add a TNS alias for the target database to `tnsnames.ora` on the GoldenGate hub.

GoldenGate runs as a database user and must have the appropriate database privileges, so you must create a GoldenGate user account on the target database.

The following statements create a user named `oggadm1`:

```
create tablespace administrator;  
create tablespace administrator_idx;  
CREATE USER oggadm1 IDENTIFIED BY "XXXXXX"  
    default tablespace ADMINISTRATOR  
    temporary tablespace TEMP;  
alter user oggadm1 quota unlimited on ADMINISTRATOR;  
alter user oggadm1 quota unlimited on ADMINISTRATOR_IDX;
```

Finally, grant the necessary privileges to the GoldenGate user account. The following statements grant privileges to a user named `oggadm1`:

```
grant create session      to oggadm1;  
grant alter session       to oggadm1;  
grant CREATE CLUSTER     to oggadm1;  
grant CREATE INDEXTYPE   to oggadm1;  
grant CREATE OPERATOR    to oggadm1;  
grant CREATE PROCEDURE   to oggadm1;  
grant CREATE SEQUENCE    to oggadm1;  
grant CREATE TABLE       to oggadm1;  
grant CREATE TRIGGER     to oggadm1;  
grant CREATE TYPE        to oggadm1;  
grant select any dictionary to oggadm1;  
grant create any table   to oggadm1;  
grant alter any table    to oggadm1;  
grant lock any table     to oggadm1;  
grant select any table   to oggadm1;  
grant insert any table   to oggadm1;  
grant update any table   to oggadm1;  
grant delete any table   to oggadm1;  
  
EXEC DBMS_GOLDENGATE_AUTH.GRANT_ADMIN_PRIVILEGE  
  (grantee=>'OGGADM1',privilege_type=>'apply',
```

```
grant_select_privileges=>true, do_grants=>TRUE);
```

Add the following entry to \$ORACLE_HOME/network/admin/tnsnames.ora in the Oracle Home that will be used by the REPLICAT process. For more information on the tnsname.ora file, see the [Oracle documentation](#).

```
OGGTARGET=
(DESCRIPTION=
  (ENABLE=BROKEN)
  (ADDRESS_LIST=
    (ADDRESS=(PROTOCOL=TCP)(HOST=goldengate-target.abcdef12345.us-west-2.rds.amazonaws.com)
     (PORT=8200))
  )
  (CONNECT_DATA=
    (SID=ORCL)
  )
)
```

Working with the EXTRACT and REPLICAT Utilities of GoldenGate

The GoldenGate utilities EXTRACT and REPLICAT work together to keep the source and target databases in sync via incremental transaction replication using trail files. All changes that occur on the source database are automatically detected by EXTRACT, then formatted and transferred to trail files on the GoldenGate on-premises or EC2-instance hub. After initial load is completed, the data is read from these files and replicated to the target database by the REPLICAT utility.

Running GoldenGate's EXTRACT Utility

The EXTRACT utility retrieves, converts, and outputs data from the source database to trail files. EXTRACT queues transaction details to memory or to temporary disk storage. When the transaction is committed to the source database, EXTRACT flushes all of the transaction details to a trail file for routing to the GoldenGate on-premises or EC2-instance hub and then to the target database.

The following tasks enable and start the EXTRACT utility:

- Configure the EXTRACT parameter file on the GoldenGate hub (on-premises or EC2 instance). The following listing shows an example EXTRACT parameter file.

```
EXTRACT EABC
SETENV (ORACLE_SID=ORCL)
SETENV (NLSLANG=AL32UTF8)

USERID oggadm1@OGGSOURCE, PASSWORD XXXXXX
EXTTRAIL /path/to/goldengate/dirdat/ab

IGNOREREPPLICATES
GETAPPLOPS
TRANLOGOPTIONS EXCLUDEUSER OGGADM1

TABLE EXAMPLE.TABLE;
```

- On the GoldenGate hub, launch the GoldenGate command line interface (*ggsci*). Log into the source database. The following example shows the format for logging in:

```
dblogin userid <user>@<db tnsname>
```

- Add a checkpoint table for the database:

```
add checkpointtable
```

- Add transdata to turn on supplemental logging for the database table:

```
add transdata <user>.<table>
```

Alternatively, you can add transdata to turn on supplemental logging for all tables in the database:

```
add transdata <user>.*
```

- Using the `ggsci` command line, enable the `EXTRACT` utility using the following commands:

```
add extract <extract name> tranlog, INTEGRATED tranlog, begin now
add exttrail <path-to-trail-from-the param-file>
    extract <extractname-from-paramfile>,
    MEGABYTES Xm
```

- Register the `EXTRACT` utility with the database so that the archive logs are not deleted. This allows you to recover old, uncommitted transactions if necessary. To register the `EXTRACT` utility with the database, use the following command:

```
register EXTRACT <extract process name>, DATABASE
```

- To start the `EXTRACT` utility, use the following command:

```
start <extract process name>
```

Running GoldenGate's REPLICAT Utility

The `REPLICAT` utility is used to "push" transaction information in the trail files to the target database.

The following tasks enable and start the `REPLICAT` utility:

- Configure the `REPLICAT` parameter file on the GoldenGate hub (on-premises or EC2 instance). The following listing shows an example `REPLICAT` parameter file.

```
REPLICAT RABC
SETENV (ORACLE_SID=ORCL)
SETENV (NLSLANG=AL32UTF8)

USERID oggadm1@OGGTARGET, password XXXXXX

ASSUMETARGETDEFS
MAP EXAMPLE.TABLE, TARGET EXAMPLE.TABLE;
```

- Launch the GoldenGate command line interface (`ggsci`). Log into the target database. The following example shows the format for logging in:

```
dblogin userid <user>@<db tnsname>
```

- Using the `ggsci` command line, add a checkpoint table. Note that the user indicated should be the GoldenGate user account, not the target table schema owner. The following example creates a checkpoint table named `gg_checkpoint`.

```
add checkpointtable <user>.gg_checkpoint
```

- To enable the REPLICAT utility, use the following command:

```
add replicat <replicat name> EXTTRAIL <extract trail file> CHECKPOINTTABLE  
<user>.gg_checkpoint
```

- To start the REPLICAT utility, use the following command:

```
start <replicat name>
```

Troubleshooting Issues When Using GoldenGate with Amazon RDS

This section explains the most common issues when using GoldenGate with Amazon RDS.

Topics

- [Log Retention \(p. 1163\)](#)
- [GoldenGate appears to be properly configured but replication is not working \(p. 1163\)](#)

Log Retention

You must have log retention enabled. If you do not, or if the retention value is too small, you will see the following message:

```
2014-03-06 06:17:27  ERROR    OGG-00446  error 2 (No such file or directory)  
opening redo log /rdsdbdata/db/GGTEST3_A/onlinelog/o1_mf_2_9k4bp1n6_.log  
for sequence 1306Not able to establish initial position for begin time 2014-03-06  
06:16:55.
```

GoldenGate appears to be properly configured but replication is not working

For pre-existing tables, GoldenGate needs to be told which SCN it should work from. Take the following steps to fix this issue:

- Launch the GoldenGate command line interface (ggsci). Log into the source database. The following example shows the format for logging in:

```
dblogin userid <user>@<db tnsname>
```

- Using the ggsci command line, set up the start SCN for the EXTRACT process. The following example sets the SCN to 223274 for the extract:

```
ALTER EXTRACT <extract process name> SCN 223274  
start <extract process name>
```

- Log into the target database. The following example shows the format for logging in:

```
dblogin userid <user>@<db tnsname>
```

- Using the ggsci command line, set up the start SCN for the REPLICAT process. The following example sets the SCN to 223274 for the REPLICAT:

```
start <replicat process name> atcsn 223274
```


Using the Oracle Repository Creation Utility on Amazon RDS for Oracle

You can use Amazon RDS to host an Oracle DB instance that holds the schemas to support your Fusion Middleware components. Before you can use Fusion Middleware components, you must create and populate schemas for them in your database. You create and populate the schemas by using the Oracle Repository Creation Utility (RCU).

You can store the schemas for any Fusion Middleware components in your Amazon RDS DB instance. The following is a list of schemas that have been verified to install correctly:

- Analytics (ACTIVITIES)
- Audit Services (IAU)
- Audit Services Append (IAU_APPEND)
- Audit Services Viewer (IAU_VIEWER)
- Discussions (DISCUSSIONS)
- Metadata Services (MDS)
- Oracle Business Intelligence (BIPLATFORM)
- Oracle Platform Security Services (OPSS)
- Portal and Services (WEBCENTER)
- Portlet Producers (PORTLET)
- Service Table (STB)
- SOA Infrastructure (SOAINFRA)
- User Messaging Service (UCSUMS)
- WebLogic Services (WLS)

Licensing and Versions

Amazon RDS supports Oracle Repository Creation Utility (RCU) version 12c only. You can use the RCU in the following configurations:

- RCU 12c with Oracle database 12.2.0.1
- RCU 12c with Oracle database 12.1.0.2.v4 or later
- RCU 12c with Oracle database 11.2.0.4.v8 or later

Before you can use RCU, you need a license for Oracle Fusion Middleware. You also need to follow the Oracle licensing guidelines for the Oracle database that hosts the repository. For more information, see [Oracle Fusion Middleware Licensing Information User Manual](#) in the Oracle documentation.

Fusion MiddleWare supports repositories on Oracle Database Enterprise Edition and Standard Editions (SE, SE One, or SE Two). Oracle recommends Enterprise Edition for production installations that require partitioning and installations that require online index rebuild.

Before you create your Oracle DB instance, confirm the Oracle database version that you need to support the components that you want to deploy. You can use the Certification Matrix to find the requirements for the Fusion Middleware components and versions you want to deploy. For more information, see [Oracle Fusion Middleware Supported System Configurations](#) in the Oracle documentation.

Amazon RDS supports Oracle database version upgrades as needed. For more information, see [Upgrading a DB Instance Engine Version \(p. 159\)](#).

Before You Begin

Before you begin, you need an Amazon VPC. Because your Amazon RDS DB instance needs to be available only to your Fusion Middleware components, and not to the public Internet, your Amazon RDS DB instance is hosted in a private subnet, providing greater security. For information about how to create an Amazon VPC for use with an Oracle DB instance, see [Creating a VPC for Use with an Oracle Database \(p. 1147\)](#).

Before you begin, you also need an Oracle DB instance. For information about how to create an Oracle DB instance for use with Fusion Middleware metadata, see [Creating an Oracle DB Instance \(p. 1152\)](#).

Recommendations

The following are some recommendations for working with your DB instance in this scenario:

- We recommend that you use Multi-AZ for production workloads. For more information about working with multiple Availability Zones, see [Regions and Availability Zones \(p. 132\)](#).
- For additional security, Oracle recommends that you use Transparent Data Encryption (TDE) to encrypt your data at rest. If you have an Enterprise Edition license that includes the Advanced Security Option, you can enable encryption at rest by using the TDE option. For more information, see [Oracle Transparent Data Encryption \(p. 1074\)](#).

Amazon RDS also provides an encryption at rest option for all database editions. For more information, see [Encrypting Amazon RDS Resources \(p. 429\)](#).

- Configure your VPC Security Groups to allow communication between your application servers and your Amazon RDS DB instance. The application servers that host the Fusion Middleware components can be on Amazon EC2 or on-premises.

Using the Oracle Repository Creation Utility

You use the Oracle Repository Creation Utility (RCU) to create and populate the schemas to support your Fusion Middleware components.

Running RCU Using the Command Line in One Step

If you don't need to edit any of your schemas before populating them, you can run RCU in a single step. Otherwise, see the following section for running RCU in multiple steps.

You can run the RCU in silent mode by using the command-line parameter `-silent`. When you run RCU in silent mode, you can avoid typing passwords on the command line by creating a text file containing the passwords. Create a text file with the password for `dbUser` on the first line, and the password for each component on subsequent lines. You specify the name of the password file as the last parameter to the RCU command.

Example

The following example creates and populates schemas for the SOA Infrastructure component (and its dependencies) in a single step.

For Linux, OS X, or Unix:

```
export ORACLE_HOME=/u01/app/oracle/product/12.2.1.0/fmw
export JAVA_HOME=/usr/java/jdk1.8.0_65
${ORACLE_HOME}/oracle_common/bin/rcu \
-silent \
-createRepository \
-connectString ${dbhost}:${dbport}:${dbname} \
```

```

-dbUser ${dbuser} \
-dbRole Normal \
-honorOMF \
-schemaPrefix ${SCHEMA_PREFIX} \
-component MDS \
-component STB \
-component OPSS \
-component IAU \
-component IAU_APPEND \
-component IAU_VIEWER \
-component UCSUMS \
-component WLS \
-component SOAINFRA \
-f < /tmp/passwordfile.txt

```

For more information, see [Running Repository Creation Utility from the Command Line](#) in the Oracle documentation.

Running RCU Using the Command Line in Multiple Steps

If you need to manually edit your schema scripts, you can run the RCU in multiple steps:

1. Run RCU in **Prepare Scripts for System Load** mode by using the `-generateScript` command-line parameter to create the scripts for your schemas.
2. Manually edit and run the generated script `script_systemLoad.sql`.
3. Run RCU again in **Perform Product Load** mode by using the `-dataLoad` command-line parameter to populate the schemas.
4. Run the generated clean-up script `script_postDataLoad.sql`.

You can run the RCU in silent mode by using the command-line parameter `-silent`. When you run RCU in silent mode, you can avoid typing passwords on the command line by creating a text file containing the passwords. Create a text file with the password for `dbUser` on the first line, and the password for each component on subsequent lines. You specify the name of the password file as the last parameter to the RCU command.

Example

The following example creates schema scripts for the SOA Infrastructure component (and its dependencies).

For Linux, OS X, or Unix:

```

export ORACLE_HOME=/u01/app/oracle/product/12.2.1.0/fmw
export JAVA_HOME=/usr/java/jdk1.8.0_65
${ORACLE_HOME}/oracle_common/bin/rcu \
-silent \
-generateScript \
-connectString ${dbhost}:${dbport}:${dbname} \
-dbUser ${dbuser} \
-dbRole Normal \
-honorOMF \
[-encryptTablespace true] \
-schemaPrefix ${SCHEMA_PREFIX} \
-component MDS \
-component STB \
-component OPSS \
-component IAU \
-component IAU_APPEND \
-component IAU_VIEWER \
-component UCSUMS \

```

```
-component WLS \
-component SOAINFRA \
-scriptLocation /tmp/rcuscripts \
-f < /tmp/passwordfile.txt
```

Now you can edit the generated script, connect to your Oracle DB instance, and run the script. The generated script is named `script_systemLoad.sql`. For information about connecting to your Oracle DB instance, see [Connecting to Your Sample Oracle DB Instance \(p. 53\)](#).

The following example populates the schemas for the SOA Infrastructure component (and its dependencies).

For Linux, OS X, or Unix:

```
export JAVA_HOME=/usr/java/jdk1.8.0_65
${ORACLE_HOME}/oracle_common/bin/rcu \
-silent \
-dataLoad \
-connectString ${dbhost}:${dbport}:${dbname} \
-dbUser ${dbuser} \
-dbRole Normal \
-honorOMF \
-schemaPrefix ${SCHEMA_PREFIX} \
-component MDS \
-component STB \
-component OPSS \
-component IAU \
-component IAU_APPEND \
-component IAU_VIEWER \
-component UCSUMS \
-component WLS \
-component SOAINFRA \
-f < /tmp/passwordfile.txt
```

To finish, you connect to your Oracle DB instance, and run the clean-up script. The script is named `script_postDataLoad.sql`.

For more information, see [Running Repository Creation Utility from the Command Line](#) in the Oracle documentation.

Running RCU in Interactive Mode

To use the RCU graphical user interface, you can run RCU in interactive mode. To run RCU in interactive mode, include the `-interactive` parameter and omit the `-silent` parameter. For more information, see [Understanding Repository Creation Utility Screens](#) in the Oracle documentation.

Example

The following example starts RCU in interactive mode and pre-populates the connection information.

For Linux, OS X, or Unix:

```
export ORACLE_HOME=u01/app/oracle/product/12.2.1.0/fmw
export JAVA_HOME=/usr/java/jdk1.8.0_65
${ORACLE_HOME}/oracle_common/bin/rcu \
-interactive \
-createRepository \
-connectString ${dbhost}:${dbport}:${dbname} \
-dbUser ${dbuser} \
-dbRole Normal
```

Known Issues

The following are some known issues for working with RCU, with some troubleshooting suggestions:

- Oracle Managed Files (OMF) — Amazon RDS uses OMF data files to simplify storage management. You can customize tablespace attributes, such as size and extent management. However, specifying a data file name when you run RCU causes tablespace code to fail with ORA-20900. The RCU can be used with OMF in the following ways:
 - In RCU 12.2.1.0 and later, use the `-honorOMF` command-line parameter.
 - In RCU 12.1.0.3 and later, use multiple steps and edit the generated script. For more information, see [Running RCU Using the Command Line in Multiple Steps \(p. 1167\)](#).
- SYSDBA — Because Amazon RDS is a managed service, you don't have full SYSDBA access to your Oracle DB instance. However, RCU 12c supports users with lower privileges. In most cases, the master user privilege is sufficient to create repositories. In some cases, the RCU might fail with ORA-01031 when attempting to grant SYS object privileges. You can retry and run the `RDSADMIN_UTIL.GRANT_SYS_OBJECT()` stored procedure, or contact AWS Support.
- Dropping Enterprise Scheduler Service — When you use the RCU to drop an Enterprise Scheduler Service repository, the RCU might fail with `Error: Component drop check failed`.

Related Topics

- [Oracle Licensing \(p. 916\)](#)

Installing a Siebel Database on Oracle on Amazon RDS

You can use Amazon RDS to host a Siebel Database on an Oracle DB instance. The Siebel Database is part of the Siebel Customer Relationship Management (CRM) application architecture. For an illustration, see [Generic Architecture of Siebel Business Application](#).

Use the following topic to help set up a Siebel Database on an Oracle DB instance on Amazon RDS. You can also find out how to use Amazon Web Services to support the other components required by the Siebel CRM application architecture.

Note

To install a Siebel Database on Oracle on Amazon RDS, you need to use the master user account. You don't need SYSDBA privilege; master user privilege is sufficient. For more information, see [Master User Account Privileges \(p. 499\)](#).

Licensing and Versions

To install a Siebel Database on Amazon RDS, you must use your own Oracle Database license, and your own Siebel license. You must have the appropriate Oracle Database license (with Software Update License and Support) for the DB instance class and Oracle Database edition. For more information, see [Oracle Licensing \(p. 916\)](#).

Oracle Database Enterprise Edition is the only edition certified by Siebel for this scenario. Amazon RDS supports Siebel CRM version 15.0 or 16.0. Use Oracle 12c, version 12.1.0.2.0. For the procedures following, we use Siebel CRM version 15.0 and Oracle 12.1.0.2 or 12.2.0.1. For more information, see [Oracle 12c with Amazon RDS \(p. 923\)](#).

Amazon RDS supports database version upgrades. For more information, see [Upgrading a DB Instance Engine Version \(p. 159\)](#).

Before You Begin

Before you begin, you need an Amazon VPC. Because your Amazon RDS DB instance needs to be available only to your Siebel Enterprise Server, and not to the public Internet, your Amazon RDS DB instance is hosted in a private subnet, providing greater security. For information about how to create an Amazon VPC for use with Siebel CRM, see [Creating a VPC for Use with an Oracle Database \(p. 1147\)](#).

Before you begin, you also need an Oracle DB instance. For information about how to create an Oracle DB instance for use with Siebel CRM, see [Creating an Oracle DB Instance \(p. 1152\)](#).

Installing and Configuring a Siebel Database

After you create your Oracle DB instance, you can install your Siebel Database. You install the database by creating table owner and administrator accounts, installing stored procedures and functions, and then running the Siebel Database Configuration Wizard. For more information, see [Installing the Siebel Database on the RDBMS](#).

To run the Siebel Database Configuration Wizard, you need to use the master user account. You don't need SYSDBA privilege; master user privilege is sufficient. For more information, see [Master User Account Privileges \(p. 499\)](#).

Using Other Amazon RDS Features with a Siebel Database

After you create your Oracle DB instance, you can use additional Amazon RDS features to help you customize your Siebel Database.

Collecting Statistics with the Oracle Statspack Option

You can add features to your DB instance through the use of options in DB option groups. When you created your Oracle DB instance, you used the default DB option group. If you want to add features to your database, you can create a new option group for your DB instance.

If you want to collect performance statistics on your Siebel Database, you can add the Oracle Statspack feature. For more information, see [Oracle Statspack \(p. 1068\)](#).

Some option changes are applied immediately, and some option changes are applied during the next maintenance window for the DB instance. For more information, see [Working with Option Groups \(p. 196\)](#). After you create a customized option group, modify your DB instance to attach it. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

Performance Tuning with Parameters

You manage your DB engine configuration through the use of parameters in a DB parameter group. When you created your Oracle DB instance, you used the default DB parameter group. If you want to customize your database configuration, you can create a new parameter group for your DB instance.

When you change a parameter, depending on the type of the parameter, the changes are applied either immediately or after you manually reboot the DB instance. For more information, see [Working with DB Parameter Groups \(p. 211\)](#). After you create a customized parameter group, modify your DB instance to attach it. For more information, see [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#).

To optimize your Oracle DB instance for Siebel CRM, you can customize certain parameters. The following table shows some recommended parameter settings. For more information about performance tuning Siebel CRM, see [Siebel CRM Performance Tuning Guide](#).

Parameter Name	Default Value	Guidance for Optimal Siebel CRM Performance
<code>_always_semi_join</code>	<code>CHOOSE</code>	OFF
<code>_b_tree_bitmap_p</code>	<code>TRUE</code>	FALSE
<code>_like_with_bind_</code>	<code>EQUALITY</code>	TRUE
<code>_no_or_expansion</code>	<code>FALSE</code>	FALSE
<code>_optimizer_join_</code>	<code>TRUE</code>	<code>sanity_check</code> TRUE
<code>_optimizer_max_p</code>	<code>2000</code>	100
<code>_optimizer_sortmerge_</code>	<code>join_enabled</code>	<code>FALSE</code>
<code>_partition_view_</code>	<code>ENABLED</code>	FALSE
<code>open_cursors</code>	300	At least 2000 .

Creating Snapshots

After you create your Siebel Database, you can copy the database by using the snapshot features of Amazon RDS. For more information, see [Creating a DB Snapshot \(p. 264\)](#) and [Restoring from a DB Snapshot \(p. 266\)](#).

Support for Other Siebel CRM Components

In addition to your Siebel Database, you can also use Amazon Web Services to support the other components of your Siebel CRM application architecture. You can find more information about the support provided by Amazon AWS for additional Siebel CRM components in the following table.

Siebel CRM Component	Amazon AWS Support
Siebel Enterprise (with one or more Siebel Servers)	You can host your Siebel Servers on Amazon Elastic Compute Cloud (Amazon EC2) instances. You can use Amazon EC2 to launch as many or as few virtual servers as you need. Using Amazon EC2, you can scale up or down easily to handle changes in requirements. For more information, see What Is Amazon EC2? You can put your servers in the same VPC with your DB instance and use the VPC security group to access the database. For more information, see Working with a DB Instance in a VPC (p. 514) .
Web Servers (with Siebel Web Server Extensions)	You can install multiple Web Servers on multiple EC2 instances. You can then use Elastic Load Balancing to distribute incoming traffic among the instances. For more information, see What Is Elastic Load Balancing?
Siebel Gateway Name Server	You can host your Siebel Gateway Name Server on an EC2 instance. You can then put your server in the same VPC with the DB instance and use the VPC security group to access the database. For more information, see Working with a DB Instance in a VPC (p. 514) .

Related Topics

- [Connecting to a DB Instance Running the Oracle Database Engine \(p. 955\)](#)

Oracle Database Engine Release Notes

Updates to your Amazon RDS for Oracle DB instances keep them current. If you apply updates, you can be confident that your DB instance is running a stable, common version of the database software that has been regression-tested by both Oracle and Amazon. We don't support applying one-off patches to individual DB instances.

You can specify any currently supported Oracle version when creating a new DB instance. You can specify the major version (such as Oracle 12.1), and any supported minor version for the specified major version. If no version is specified, Amazon RDS defaults to a supported version, typically the most recent version. If a major version is specified but a minor version is not, Amazon RDS defaults to a recent release of the major version you have specified. To see a list of supported versions, as well as defaults for newly created DB instances, use the [describe-db-engine-versions](#) AWS CLI command.

Oracle Versions 19.0.0, 18.0.0, and 12.2.0.1

For Amazon RDS for Oracle versions 19.0.0.0, 18.0.0.0, and 12.2.0.1, Amazon RDS incorporates bug fixes from Oracle by using Release Updates (RUs) and Release Updates Revisions (RURs). We don't support applying one-off patches to individual DB instances.

To find what RUs and RURs are applied to Amazon RDS for Oracle versions 19.0.0, 18.0.0.0, and 12.2.0.1, see the following table.

RU and RUR	Version 19.0.0.0	Version 18.0.0.0	Version 12.2.0.1
2019 October	—	18.0.0.0.ru-2019-10.12.2019f1020119p.101ru7	2019-10.12.2019f1020119p.101ru7
2019 July	19.0.0.0.ru-2019-07.18.2019f0720119p.071ru10	2019-07.18.2019f0720119p.071ru10	2019-07.18.2019f0720119p.071ru10
2019 April	—	—	12.2.0.1.ru-2019-04.rur-2019-04
2019 January	—	—	12.2.0.1.ru-2019-01.rur-2019-01
2018 October	—	—	12.2.0.1.ru-2018-10.rur-2018-10

Oracle Versions 12.1.0.2 and 11.2.0.4

For Amazon RDS for Oracle versions 12.1.0.2 and 11.2.0.4, Amazon RDS incorporates bug fixes from Oracle via their quarterly Database Patch Set Updates (PSUs). You can be confident that your DB instance is running a stable, common version of the database software that has been regression-tested by both Oracle and Amazon. We don't support applying one-off patches to individual DB instances.

To find what Oracle Patch Set Updates (PSUs) are applied to Amazon RDS for Oracle versions 12.1.0.2 and 11.2.0.4, see the following table.

PSU	Version 12.1.0.2	Version 11.2.0.4
2019 October	12.1.0.2.v18 (p. 1196)	11.2.0.4.v22 (p. 1225)
2019 July	12.1.0.2.v17 (p. 1199)	11.2.0.4.v21 (p. 1228)
2019 April	12.1.0.2.v16 (p. 1201)	11.2.0.4.v20 (p. 1229)
2019 January	12.1.0.2.v15 (p. 1204)	11.2.0.4.v19 (p. 1231)
2018 October	12.1.0.2.v14 (p. 1206)	11.2.0.4.v18 (p. 1232)

PSU	Version 12.1.0.2	Version 11.2.0.4
2018 July	12.1.0.2.v13 (p. 1208)	11.2.0.4.v17 (p. 1234)
2018 April	12.1.0.2.v12 (p. 1210)	11.2.0.4.v16 (p. 1235)
2018 January	12.1.0.2.v11 (p. 1212)	11.2.0.4.v15 (p. 1237)
2017 October	12.1.0.2.v10 (p. 1213)	11.2.0.4.v14 (p. 1238)
2017 July	12.1.0.2.v9 (p. 1215)	11.2.0.4.v13 (p. 1240)
2017 April	12.1.0.2.v8 (p. 1216)	11.2.0.4.v12 (p. 1241)
2017 January	12.1.0.2.v7 (p. 1218)	11.2.0.4.v11 (p. 1242)
2016 October	12.1.0.2.v6 (p. 1219)	11.2.0.4.v10 (p. 1244)
2016 July	12.1.0.2.v5 (p. 1220)	11.2.0.4.v9 (p. 1245)
2016 April	12.1.0.2.v4 (p. 1222)	11.2.0.4.v8 (p. 1246)
2016 January	12.1.0.2.v3 (p. 1222)	11.2.0.4.v7 (p. 1247)
2015 October	12.1.0.2.v2 (p. 1223)	11.2.0.4.v6 (p. 1248) 11.2.0.4.v5 (p. 1249)
2015 April	12.1.0.2.v1 (p. 1224)	11.2.0.4.v4 (p. 1250)
2014 October	—	11.2.0.4.v3 (p. 1251)
2014 July	—	11.2.0.4.v2 (p. 1251) <i>(Deprecated)</i>
2014 January	—	11.2.0.4.v1 (p. 1252)

Topics

- [Database Engine: 19.0.0.0 \(p. 1174\)](#)
- [Database Engine: 18.0.0.0 \(p. 1176\)](#)
- [Database Engine: 12.2.0.1 \(p. 1182\)](#)
- [Database Engine: 12.1.0.2 \(p. 1195\)](#)
- [Database Engine: 11.2.0.4 \(p. 1225\)](#)

Database Engine: 19.0.0.0

The following versions are available for Oracle database engine 19.0.0.0:

- [Version 19.0.0.0.ru-2019-07.rur-2019-07.r1 \(p. 1174\)](#)

Version 19.0.0.0.ru-2019-07.rur-2019-07.r1

Version 19.0.0.0.ru-2019-07.rur-2019-07.r1 adds support for the following:

- Patch 29834717: Database Release Update : 19.4.0.0.190716

- Patch 29774421: OJVM RELEASE UPDATE: 19.4.0.0.190716
- Patch 28852325: RDBMS - DSTV33 UPDATE - TZDATA2018G
- Patch 28852334: DSTV33 UPDATE - TZDATA2018G - NEED OJVM FIX
- Patch 28730253: SUPPORT NEW ERA REIWA FOR JAPANESE IMPERIAL CALENDAR
- Patch 29213893: DBMS_STATS FAILING WITH ERROR ORA-01422 WHEN GATHERING STATS FOR USER \$ TABLE
- PreUpgrade Jar: preupgrade_19_cbuild_4_lf.zip

Oracle Release Update 19.4.0.0.190716, Released July 2019

Bugs fixed:

29850930, 29225758, 29276272, 14735102, 17428816, 19080742, 19697993 20313356, 21374587, 21965541, 23296836, 23606241, 25756945, 25806201 25883179, 25986062, 25997810, 26476244, 26611353, 26739322, 26872233 27126938, 27244999, 27359766, 27369515, 27423500, 27453490, 27666312 27710072, 27846298, 27957203, 28064977, 28072567, 28125947, 28129791 28181021, 28210681, 28271693, 28279456, 28313275, 28326928, 28350595 28371123, 28379065, 28381939, 28390273, 28431445, 28463226, 28475242 28489419, 28502773, 28513333, 28534475, 28547068, 28561704, 28569897 28572533, 28572544, 28572667, 28572834, 28587723, 28593682, 28594086 28597221, 28601957, 28605066, 28606598, 28612239, 28625862, 28627033 28636532, 28643718, 28644549, 28645570, 28646200, 28646939, 28649388 28655209, 28663782, 28673945, 28692275, 28694872, 28696373, 28705231 28710385, 28710734, 28714461, 28715727, 28718469, 28719348, 28720418 28722229, 28730079, 28740708, 28742555, 28749853, 28755011, 28758722 28760206, 28767240, 28772390, 28774416, 28777214, 28781754, 28785531 28789531, 28791852, 28795551, 28802734, 28804517, 28810381, 28811560 28815123, 28815355, 28819640, 28824482, 28827682, 28833912, 28835937 28849776, 28854004, 28862532, 28863432, 28867992, 28873575, 28876253 28876639, 28884931, 28888327, 28889730, 28892794, 28897512, 28899663 28901126, 28905457, 28907196, 28912691, 28915561, 28917080, 28918429 28919145, 28922227, 28922532, 28922608, 28925634, 28925880, 28933158 28936114, 28937717, 28938698, 28940179, 28940281, 28941901, 28942455 28945421, 28945994, 28950868, 28951533, 28952168, 28954762, 28955606 28956908, 28957292, 28957723, 28962775, 28965231, 28966444, 28974083 28977322, 28983095, 28983486, 28986207, 28986231, 28986326, 28986481 28988482, 28989306, 28993295, 28993353, 28994307, 28996376, 29000000 29001305, 29001888, 29002784, 29002927, 29003738, 29006318, 29006621 29007321, 29007775, 29008035, 29008669, 29009513, 29011936, 29012609 29013475, 29014076, 29015118, 29017265, 29018655, 29019121, 29021063 29021352, 29024054, 29024552, 29024732, 29026582, 29026606, 29027456 29027694, 29027940, 29031575, 29031600, 29032234, 29032276, 29032457 29032607, 29033052, 29033145, 29033200, 29033280, 29034587, 29037290 29038528, 29039089, 29039510, 29040739, 29043554, 29043651, 29043725 29044763, 29044954, 29047850, 29048178, 29048289, 29048498, 29048605 29050357, 29050560, 29050765, 29051702, 29051953, 29052726, 29053783 29056024, 29056270, 29056560, 29056767, 29059011, 29061959, 29062692 29062848, 29062860, 29062868, 29110526, 29110783, 29110790, 29110797 29110802, 29110805, 29111598, 29113282, 29113305, 29117526, 29117642 29119077, 29120223, 29122224, 29122254, 29123297, 29123432, 29123482 29124368, 29125036, 29125374, 29125380, 29126345, 29127957, 29128693 29128935, 29129450, 29129497, 29129712, 29130219, 29131539, 29132938 29134447, 29135383, 29135649, 29136111, 29138641, 29141316, 29141341 29141685, 29142609, 29142667, 29144995, 29145214, 29145730, 29149829 29150338, 29151520, 29152357, 29155099, 29157389, 29158680, 29158899 29159909, 29159936, 29160174, 29162095, 29163156, 29163415, 29163437 29163524, 29163567, 29167111, 29167342, 29167374, 29167940, 29168219 29168433, 29169073, 29169215, 29171683, 29171942, 29172618, 29172826 29173140, 29173373, 29173817, 29174004, 29176318, 29177466, 29177543 29177886, 29178385, 29180313, 29180455, 29180559, 29180893, 29181153 29181231, 29181620, 29181743, 29181923, 29182019, 29183912, 29184297 29184666, 29185193, 29186456, 29189302, 29189307, 29189889, 29190235 29190474, 29190740, 29191541, 29192419, 29192468, 29192685, 29193207 29194205, 29194367, 29194493, 29194827, 29194981, 29195279, 29195337 29195758, 29196725, 29199635, 29199733, 29200316, 29200700, 29201494 29201539, 29202104, 29202850, 29203122, 29203166, 29203425, 29203443 29203604, 29205281,

29205323, 29205419, 29205463, 29205767, 29205918 29206109, 29206605, 29207073, 29208260, 29208732, 29211457, 29211724 29212012, 29212433, 29212611, 29213351, 29213775, 29213850, 29213879 29214561, 29214960, 29216746, 29216984, 29217294, 29217472, 29217828 29217848, 29218570, 29219205, 29219273, 29220079, 29221248, 29221891 29222031, 29222784, 29223833, 29223859, 29223967, 29224065, 29224605 29225076, 29225168, 29227602, 29228869, 29229164, 29229754, 29229844 29229955, 29230252, 29230565, 29231133, 29232117, 29232154, 29233415 29234123, 29237538, 29240307, 29241345, 29241651, 29242017, 29242884 29243958, 29245137, 29245160, 29246163, 29247415, 29247712, 29247906 29248495, 29248552, 29248835, 29248858, 29249991, 29250059, 29250317 29251259, 29253184, 29253871, 29254031, 29254930, 29255178, 29255273 29255431, 29255435, 29256426, 29259119, 29259320, 29260452, 29261547 29261906, 29262512, 29262887, 29265448, 29266248, 29266899, 29267292 29268412, 29269171, 29269228, 29270585, 29273539, 29273570, 29273735 29273812, 29273847, 29274428, 29274564, 29274627, 29275461, 29277317 29278218, 29279658, 29279751, 29279854, 29281527, 29281691, 29281796 29282233, 29282898, 29285503, 29285788, 29285956, 29286037, 29287130 29287705, 29292837, 29293072, 29293574, 29297863, 29297915, 29298220 29299049, 29299082, 29299844, 29301463, 29301566, 29302963, 29303918 29304781, 29306226, 29306713, 29311588, 29312310, 29312672, 29312734 29312753, 29313347, 29313417, 29313525, 29314539, 29317756, 29318410 29319441, 29321489, 29323946, 29324568, 29324735, 29325087, 29325105 29325257, 29325765, 29325993, 29327044, 29329042, 29329087, 29329807 29330361, 29331066, 29331209, 29331380, 29331493, 29332292, 29332395 29332771, 29333500, 29336843, 29337310, 29338315, 29338453, 29338780 29338913, 29339101, 29339155, 29341209, 29343086, 29345937, 29346057 29346211, 29346943, 29347620, 29348176, 29350052, 29351386, 29351716 29351735, 29351749, 29352298, 29352724, 29352867, 29352947, 29353271 29353432, 29353960, 29355654, 29356547, 29356704, 29356711, 29356752 29358509, 29358828, 29360285, 29360672, 29360911, 29360950, 29361472 29361801, 29363151, 29364171, 29364177, 29366940, 29367019, 29367561 29368253, 29372541, 29373418, 29373588, 29374179, 29375941, 29376346 29377986, 29378029, 29378834, 29378913, 29379978, 29382784, 29382815 29384781, 29384854, 29384864, 29385429, 29385652, 29386502, 29386635 29386660, 29387073, 29387274, 29388020, 29388072, 29388094, 29388524 29388830, 29389889, 29390011, 29390435, 29390785, 29391030, 29394014 29394140, 29394749, 29395657, 29397954, 29397996, 29398488, 29398863 29399100, 29399121, 29399938, 29402131, 29404483, 29405012, 29405462 29405651, 29405996, 29407804, 29409149, 29410311, 29410834, 29411037 29412066, 29412269, 29417719, 29417884, 29420834, 29421059, 29423826 29424999, 29426241, 29429017, 29429264, 29429566, 29430524, 29431192 29431485, 29434301, 29435474, 29435652, 29436454, 29436514, 29437594 29437712, 29438277, 29438736, 29439522, 29441196, 29443187, 29443250 29444072, 29444282, 29444602, 29446669, 29449477, 29450421, 29451386 29452576, 29452936, 29452953, 29455424, 29457312, 29457370, 29457807 29460252, 29461791, 29462594, 29462767, 29462957, 29464779, 29465177 29467622, 29476473, 29483532, 29483672, 29483685, 29483712, 29486181 29488894, 29489546, 29490256, 29492127, 29492939, 29494245, 29495057 29495684, 29497311, 29500826, 29502561, 29503543, 29503631, 29503827 29504492, 29504682, 29505668, 29507616, 29510278, 29511611, 29514479 29515134, 29515240, 29515766, 29515834, 29516300, 29517168, 29521187 29521748, 29522358, 29522561, 29522662, 29523511, 29525467, 29525886 29527595, 29529147, 29530440, 29530515, 29530812, 29531654, 29541769 29542449, 29543034, 29546817, 29547867, 29548687, 29548722, 29549154 29549730, 29557336, 29558975, 29559187, 29559446, 29559908, 29559981 29565611, 29580983, 29581771, 29584261, 29586143, 29597536, 29597754 29598039, 29598233, 29599300, 29601461, 29604002, 29608000, 29611020 29611991, 29616244, 29616414, 29618074, 29618190, 29622936, 29626732 29628200, 29629650, 29629681, 29631749, 29632095, 29633697, 29635427 29635717, 29637362, 29638285, 29641736, 29653246, 29656819, 29657422 29664087, 29664161, 29670782, 29676089, 29677051, 29677733, 29679856 29687459, 29687763, 29692694, 29695841, 29703932, 29707099, 29742223 29747648, 29751094, 29753244, 29754951, 29756274, 29757651, 29766207 29766503, 29766679, 29775393, 29779196, 29795957, 29838773

Database Engine: 18.0.0.0

The following versions are available for Oracle database engine 18.0.0.0:

- [Version 18.0.0.0.ru-2019-10.rur-2019-10.r1 \(p. 1177\)](#)

- [Version 18.0.0.0.ru-2019-07.rur-2019-07.r1 \(p. 1179\)](#)

Version 18.0.0.0.ru-2019-10.rur-2019-10.r1

Version 18.0.0.0.ru-2019-10.rur-2019-10.r1 adds support for the following:

- Patch 30112122: DATABASE OCT 2019 RELEASE UPDATE 18.0.0.0.191015
- Patch 30133625: OJVM RELEASE UPDATE 12.2.0.1.191015
- Patch 29997937: DSTv34 for RDBMS (TZDATA2019G)
- Patch 29997959: DSTv34 for OJVM (TZDATA2019G)
- PreUpgrade Jar: preupgrade_181_cbuild_8_1f.zip
- Patch 28730253: SUPPORT NEW ERA REIWA FOR JAPANESE IMPERIAL CALENDAR is included in DB PATCH 30138470
- Patch 29213893: DBMS_STATS FAILING WITH ERROR ORA-01422 WHEN GATHERING STATS FOR USER \$ TABLE
- Adds support for [Resizing the Temporary Tablespace in a Read Replica \(p. 1104\)](#).

Oracle Release Update 18.8.0.0.0, Released October 2019

Bugs fixed:

30259008, 30253255, 8932139, 9062315, 13554903, 14221306, 20436508 21095391, 21223848, 21374587, 21547051, 21744603, 21766220, 21806121 21935698, 23003564, 23061453, 23310101, 23761724, 24489904, 24687075 24689376, 24737581, 24841671, 24844841, 24903291, 24925863, 24971597 25035594, 25035599, 25060506, 25287072, 25293659, 25303284, 25309116 25348956, 25405687, 25487146, 25591292, 25607397, 25634405, 25644811 25663488, 25686739, 25696520, 25726981, 25736428, 25743479, 25756945 25824236, 25882883, 25890002, 25908728, 25911069, 25929650, 25943740 25958554, 25986062, 25997810, 26083489, 26115103, 26164661, 26226953 26237338, 26281476, 26284722, 26297826, 26336101, 26362155, 26399691 26422277, 26423085, 26427905, 26440169, 26450454, 26476244, 26521043 26536320, 26595088, 26598422, 26615291, 26617804, 26646549, 26654411 26694735, 26724511, 26731697, 26745002, 26785169, 26790514, 26792891 26798411, 26818960, 26822620, 26843558, 26843664, 26846077, 26871815 26883456, 26894737, 26895149, 26898279, 26910716, 26927998, 26928317 26933599, 26943004, 26956033, 26960097, 26961415, 26966120, 26966916 26970175, 26976568, 26985002, 26986173, 26992964, 27000158, 27005278 27006120, 27006664, 27016033, 27026401, 27028251, 27030974, 27033520 27034688, 27035653, 27036408, 27037839, 27038986, 27041253, 27044575 27047831, 27053044, 27054231, 27058530, 27060167, 27060859, 27061736 27066451, 27066519, 27072923, 27073066, 27075854, 27080874, 27086821 27090765, 27093423, 27101105, 27101527, 27101652, 27105900, 27106301 27110878, 27111780, 27112686, 27115422, 27119621, 27122162, 27125872 27126666, 27128580, 27135647, 27142529, 27143756, 27143882, 27144928 27147979, 27150500, 27151826, 27152892, 27153641, 27155549, 27156355 27160360, 27160922, 27163928, 27164352, 27165231, 27166354, 27169796 27169888, 27170305, 27179264, 27181521, 27181537, 27181897, 27185188 27187440, 27189611, 27190851, 27193810, 27197334, 27199245, 27200959 27202015, 27208795, 27208953, 27210038, 27210263, 27210872, 27214085 27214204, 27215007, 27216046, 27216224, 27220937, 27221900, 27222121 27222626, 27223075, 27224987, 27226913, 27228786, 27229389, 27231051 27232983, 27233563, 27236052, 27236110, 27236722, 27240246, 27240570 27241221, 27241247, 27242226, 27242616, 27244337, 27244785, 27249215 27249531, 27250547, 27251690, 27254851, 27255377, 27256000, 27256488 27256534, 27256584, 27258578, 27259307, 27259386, 27259983, 27262601 27262650, 27262798, 27262945, 27262991, 27263276, 27263996, 27264464 27266245, 27270197, 27274456, 27274536, 27275136, 27275533, 27275776 27276231, 27282707, 27283029, 27283960, 27284375, 27284499, 27285244 27288230, 27288638, 27288894, 27292213, 27293599, 27294480, 27301308 27301568, 27302594, 27302632, 27302681, 27302695, 27302711, 27302714 27302730, 27302777, 27302800, 27302960, 27303287, 27303785, 27303938 27304410, 27304906, 27304936, 27305318, 27307868, 27308088, 27310092 27313687, 27314206, 27314390,

27318117, 27318869, 27320576, 27321179 27321834, 27326204, 27329812, 27330158, 27330161, 27333658, 27333664 27333693, 27333731, 27334316, 27334648, 27335682, 27338912, 27338946 27339115, 27339396, 27339483, 27339495, 27341036, 27345190, 27345231 27345450, 27345498, 27346329, 27346644, 27346709, 27346949, 27347126 27348081, 27348707, 27349393, 27350267, 27351628, 27352600, 27354783 27356373, 27357773, 27358241, 27359178, 27359368, 27360126, 27364854 27364891, 27364916, 27364947, 27365139, 27365702, 27365993, 27367194 27368850, 27369515, 27372756, 27375260, 27375542, 27376871, 27378103 27379233, 27381383, 27381417, 27381498, 27381656, 27383281, 27384222 27386467, 27389352, 27392187, 27393570, 27394086, 27395404, 27395416 27395794, 27396357, 27396365, 27396377, 27396624, 27396666, 27396672 27396813, 27397048, 27398080, 27398660, 27400416, 27400598, 27401637 27404573, 27404668, 27405242, 27405645, 27405696, 27406105, 27410279 27410300, 27410595, 27412805, 27416327, 27416997, 27417186, 27420715 27421101, 27421733, 27422874, 27423251, 27424405, 27425507, 27425622 27426363, 27427805, 27430802, 27432062, 27432338, 27432355, 27432826 27433385, 27433870, 27434050, 27434193, 27434486, 27434974, 27435537 27439835, 27441326, 27441980, 27442041, 27444727, 27445330, 27445462 27445727, 27447452, 27447687, 27448162, 27449814, 27450355, 27450400 27450783, 27451049, 27451182, 27451187, 27451531, 27452046, 27452760 27453225, 27454722, 27457666, 27457891, 27458164, 27459909, 27459948 27460675, 27462994, 27466597, 27467543, 27468303, 27469245, 27469329 27471876, 27472969, 27473800, 27479358, 27480784, 27483974, 27484556 27486253, 27487309, 27487795, 27487919, 27489719, 27493674, 27496224 27496308, 27496424, 27497950, 27498477, 27501327, 27501413, 27501465 27502420, 27504190, 27504770, 27505229, 27505603, 27506774, 27507968 27508985, 27510959, 27511196, 27512439, 27517818, 27518227, 27518310 27520070, 27520900, 27522245, 27523368, 27523800, 27525909, 27526744 27532375, 27533780, 27533819, 27534509, 27537472, 27539757, 27540613 27541286, 27541468, 27542824, 27544030, 27544973, 27545630, 27547732 27550341, 27551855, 27554074, 27555481, 27558557, 27558559, 27558861 27560602, 27560702, 27562488, 27563629, 27563767, 27565906, 27567477 27570318, 27576342, 27576354, 27577758, 27578007, 27579353, 27580996 27585755, 27585800, 27586810, 27586895, 27587672, 27588271, 27591842 27592466, 27593389, 27593501, 27593585, 27595973, 27599689, 27599927 27601118, 27601441, 27602091, 27602488, 27603841, 27604293, 27607563 27607805, 27608669, 27610269, 27613080, 27613247, 27613530, 27615608 27616657, 27617522, 27617978, 27620808, 27623159, 27623844, 27625274 27625620, 27629756, 27631506, 27632114, 27634676, 27634991, 27635508 27644757, 27645940, 27649707, 27652302, 27654521, 27655217, 27658186 27658205, 27662528, 27663370, 27664702, 27666312, 27671633, 27679488 27679664, 27679793, 27679806, 27679961, 27680162, 27680509, 27680669 27682151, 27686599, 27688036, 27688099, 27688692, 27690513, 27690578 27691809, 27691920, 27691939, 27692215, 27693416, 27693713, 27694261 27695063, 27697092, 27698953, 27700466, 27701795, 27704237, 27705761 27707544, 27709046, 27710072, 27718914, 27719187, 27723002, 27723151 27726269, 27726780, 27729678, 27732323, 27733415, 27739006, 27740424 27740844, 27744211, 27745220, 27747869, 27748954, 27751006, 27751755 27753336, 27756900, 27757567, 27757794, 27757888, 27758544, 27758653 27758972, 27759077, 27759457, 27761402, 27766324, 27767081, 27772093 27772815, 27773602, 27774320, 27774539, 27779886, 27780562, 27782339 27783289, 27786772, 27791223, 27793533, 27797290, 27801337, 27803665 27807441, 27810967, 27812560, 27812593, 27813267, 27815347, 27818871 27824540, 27824543, 27825241, 27829295, 27832643, 27833369, 27833672 27834551, 27834569, 27834984, 27835925, 27839353, 27840386, 27843646 27846298, 27846499, 27847259, 27849825, 27850112, 27851757, 27856471 27861226, 27861452, 27861909, 27869075, 27869339, 27873412, 27873643 27876671, 27882176, 27886087, 27892488, 27896443, 27896458, 27898015 27900663, 27902561, 27908644, 27912301, 27918832, 27920184, 27924147 27926113, 27930478, 27931299, 27934468, 27935348, 27935493, 27938736 27940876, 27941110, 27941896, 27945870, 27948050, 27948153, 27950708 27952762, 27959594, 27960021, 27961746, 27964051, 27965400, 27965830 27966472, 27967484, 27970265, 27971503, 27971575, 27972265, 27975778 27977039, 27983174, 27984028, 27986817, 27989556, 27989849, 27991970 27993289, 27994325, 27994333, 27995215, 27995248, 27997875, 27998003 27999073, 27999597, 27999638, 28000269, 28004853, 28006704, 28018962 28019283, 28021205, 28022101, 28022847, 28023081, 28023399, 28023482 28024347, 28026866, 28033429, 28036487, 28043157, 28045903, 28057267 28058612, 28059199, 28067846, 28072130, 28072464, 28072567, 28074713 28085865, 28088762, 28090453, 28091981, 28098865, 28103600, 28103869 28104361, 28104409, 28106402, 28108003, 28108898, 28111583, 28120036 28120951, 28124631, 28125947, 28129791, 28132287, 28135648, 28165439 28165545, 28169711, 28174827, 28174951, 28175445,

28180464, 28181021 28184554, 28184800, 28187706, 28188330, 28189466, 28194173, 28199085, 28201419, 28204443, 28209341, 28210192, 28211734, 28214943, 28215510 28218832, 28220398, 28223871, 28226179, 28227512, 28229360, 28236305 28238264, 28242712, 28256164, 28258608, 28264172, 28271119, 28271693 28279837, 28281094, 28282606, 28285766, 28290434, 28294563, 28302049 28304709, 28305001, 28305362, 28309182, 28312508, 28315031, 28315995 28319114, 28319623, 28320117, 28320399, 28321446, 28323201, 28328895 28329450, 28330714, 28333072, 28338399, 28338999, 28344964, 28350595 28354603, 28357401, 28361083, 28361221, 28361787, 28365111, 28369092 28371123, 28378446, 28379065, 28384353, 28385102, 28390273, 28391210, 28391582, 28392168, 28392251, 28393678, 28396445, 28401116, 28403295 28413955, 28420042, 28420457, 28423598, 28432129, 28434028, 28435825 28445741, 28448314, 28455212, 28468312, 28475164, 28478676, 28481149 28481679, 28483184, 28489150, 28492362, 28493478, 28498976, 28501075 28502403, 28502773, 28503484, 28504545, 28507324, 28508053, 28508296 28508557, 28512761, 28513333, 28514693, 28521330, 28527416, 28528349 28535272, 28542455, 28544633, 28545134, 28545687, 28547068, 28553468 28571483, 28572834, 28578164, 28580528, 28584193, 28584217, 28584444 28585411, 28587723, 28600233, 28612674, 28614072, 28617631, 28617959, 28621470, 28627255, 28627686, 28632559, 28636676, 28644549, 28646200 28670445, 28673203, 28678804, 28679454, 28680029, 28685371, 28692103 28692275, 28695694, 28697806, 28702188, 28708023, 28710827, 28713840 28714058, 28714988, 28728040, 28728272, 28730076, 28742555, 28747182 28749289, 28752599, 28755011, 28758090, 28758722, 28761812, 28767240 28770146, 28774416, 28776811, 28777214, 28781754, 28785022, 28785531 28791725, 28797711, 28803345, 28805612, 28805695, 28809909, 28817449 28819640, 28820669, 28831971, 28849751, 28852691, 28855922, 28856060 28856172, 28867992, 28876639, 28878525, 28881723, 28887509, 28889730, 28891984, 28905457, 28919145, 28925880, 28938924, 28940179, 28951014 28951382, 28956908, 28960211, 28965084, 28986231, 28987454, 28993353 28993590, 29000190, 29002488, 29006527, 29009513, 29015118, 29015706 29024054, 29027694, 29027940, 29032276, 29033896, 29037290, 29051702 29056270, 29056767, 29123482, 29125374, 29136111, 29139591, 29158680 29163567, 29171683, 29177886, 29189889, 29198092, 29200700, 29203604 29205918, 29212433, 29213320, 29213351, 29224605, 29225076, 29230252 29230565, 29233415, 29241345, 29242017, 29247906, 29250317, 29255273 29285503, 29301463, 29312672, 29314539, 29331209, 29331493, 29339155, 29347981, 29351662, 29353821, 29356752, 29364171, 29376346, 29378913 29379978, 29382784, 29386635, 29388020, 29388952, 29391849, 29394749 29395657, 29398488, 29404483, 29405462, 29407804, 29409149, 29412269 29417719, 29429264, 29430524, 29436454, 29437712, 29439522, 29464779 29465177, 29483672, 29483723, 29493122, 29500257, 29504682, 29511611 29515766, 29524599, 29525886, 29530515, 29531541, 29541742, 29542449 29548413, 29557261, 29607136, 29614098, 29616244, 29632265, 29633753 29637526, 29645349, 29668005, 29676089, 29687459, 29707896, 29724041 29782211, 29813494, 29893132, 30034456, 30088912, 30189023

Version 18.0.0.0.ru-2019-07.rur-2019-07.r1

Version 18.0.0.0.ru-2019-07.rur-2019-07.r1 adds support for the following:

- Patch 29757256: Database Release Update : 18.7.0.0.190716
- Patch 29774410: OJVM RELEASE UPDATE: 18.7.0.0.190716
- Patch 27539475: "ORA-3816 - MISSING MESSAGE INFORMATION FOR 3816 ERROR."
- Patch 29213893: "DBMS_STATS FAILING WITH ERROR ORA-01422 WHEN GATHERING STATS FOR USER\$ TABLE"
- Patch 28125601: DSTv32 for RDBMS (TZDATA2018E)
- Patch 28852325: DSTv33 for RDBMS (TZDATA2018G)
- Patch 28852334: DSTv33 for OJVM (TZDATA2018G)
- PreUpgrade Jar: preupgrade_181_cbuild_7_lf.zip
- Patch 28730253: SUPPORT NEW ERA REIWA FOR JAPANESE IMPERIAL CALENDAR

Oracle Release Update 18.7.0.0.190716, Released July 2019

Bugs fixed:

8932139, 9062315, 13554903, 14221306, 20436508, 21095391, 21223848 21547051, 21744603, 21766220, 21806121, 23003564, 23061453, 23310101 23761724, 24489904, 24689376, 24737581, 24841671, 24844841, 24903291 24925863, 24971597, 25035594, 25035599, 25287072, 25293659, 25303284 25309116, 25348956, 25405687, 25487146, 25591292, 25607397, 25634405 25644811, 25663488, 25686739, 25726981, 25736428, 25743479, 25756945 25824236, 25882883, 25890002, 25908728, 25911069, 25929650, 25943740 25958554, 25986062, 25997810, 26083489, 26115103, 26164661, 26226953 26237338, 26281476, 26284722, 26297826, 26336101, 26362155, 26399691 26422277, 26423085, 26427905, 26440169, 26450454, 26476244, 26521043 26536320, 26595088, 26598422, 26615291, 26617804, 26646549, 26654411 26694735, 26731697, 26745002, 26785169, 26790514, 26792891, 26818960 26822620, 26843558, 26843664, 26846077, 26871815, 26883456, 26894737 26895149, 26898279, 26927998, 26928317, 26933599, 26943004, 26956033 26960097, 26961415, 26966120, 26966916, 26970175, 26985002, 26986173 26992964, 27000158, 27005278, 27006120, 27006664, 27016033, 27026401 27028251, 27030974, 27034688, 27035653, 27036408, 27037839, 27038986 27041253, 27044575, 27047831, 27053044, 27054231, 27058530, 27060167 27060859, 27061736, 27066451, 27066519, 27073066, 27075854, 27080874 27086821, 27090765, 27093423, 27101105, 27101527, 27101652, 27105900 27106301, 27110878, 27111780, 27112686, 27115422, 27119621, 27122162 27125872, 27126666, 27128580, 27135647, 27142529, 27143756, 27143882 27147979, 27150500, 27151826, 27152892, 27153641, 27155549, 27156355 27160360, 27160922, 27163928, 27164352, 27165231, 27166354, 27169796 27169888, 27170305, 27179264, 27181521, 27181537, 27181897, 27185188 27187440, 27189611, 27190851, 27193810, 27197334, 27199245, 27200959 27202015, 27208795, 27208953, 27210038, 27210263, 27210872, 27214085 27215007, 27216046, 27216224, 27220937, 27221900, 27222121, 27222626 27223075, 27224987, 27226913, 27228786, 27229389, 27231051, 27232983 27233563, 27236052, 27236110, 27236722, 27240246, 27240570, 27241221 27241247, 27242226, 27242616, 27244337, 27244785, 27249215, 27249531 27250547, 27251690, 27254851, 27255377, 27256000, 27256488, 27256534 27256584, 27258578, 27259307, 27259386, 27259983, 27262601, 27262650 27262798, 27262945, 27262991, 27263276, 27263996, 27266245, 27270197 27274456, 27274536, 27275136, 27275533, 27275776, 27276231, 27282707 27283029, 27283960, 27284375, 27284499, 27285244, 27288230, 27288638 27288894, 27292213, 27293599, 27294480, 27301308, 27301568, 27302594 27302632, 27302681, 27302695, 27302711, 27302714, 27302730, 27302777 27302800, 27302960, 27303287, 27303785, 27303938, 27304410, 27304936 27305318, 27307868, 27308088, 27310092, 27313687, 27314206, 27314390 27318117, 27318869, 27320576, 27321179, 27321834, 27326204, 27329812 27330158, 27330161, 27333658, 27333664, 27333693, 27333731, 27334316 27334648, 27335682, 27338912, 27338946, 27339115, 27339396, 27339483 27339495, 27341036, 27345190, 27345231, 27345450, 27345498, 27346329 27346644, 27346709, 27346949, 27347126, 27348081, 27348707, 27349393 27350267, 27351628, 27352600, 27354783, 27356373, 27357773, 27358241 27359178, 27359368, 27360126, 27364854, 27364891, 27364916, 27364947 27365139, 27365702, 27365993, 27367194, 27368850, 27372756, 27375260 27375542, 27376871, 27378103, 27379233, 27381383, 27381417, 27381498 27381656, 27384222, 27384647, 27389352, 27392187, 27394086, 27395404 27395416, 27395794, 27396357, 27396365, 27396377, 27396624, 27396666 27396672, 27396813, 27397048, 27398080, 27398660, 27400416, 27400598 27401637, 27404668, 27405242, 27405645, 27405696, 27410279, 27410300 27410595, 27412805, 27416327, 27416997, 27417186, 27420715, 27421101 27421733, 27422874, 27423251, 27424405, 27425507, 27425622, 27426363 27427805, 27430802, 27432062, 27432338, 27432355, 27432826, 27433870 27434050, 27434193, 27434486, 27434974, 27435537, 27439835, 27441326 27441980, 27442041, 27444727, 27445330, 27445462, 27445727, 27447452 27447687, 27448162, 27449814, 27450355, 27450400, 27450783, 27451049 27451182, 27451187, 27451531, 27452046, 27452760, 27453225, 27454722 27457666, 27457891, 27458164, 27459909, 27460675, 27462994, 27466597 27467543, 27468303, 27469245, 27469329, 27471876, 27472969, 27473800 27479358, 27480784, 27483974, 27484556, 27486253, 27487309, 27487795 27487919, 27489719, 27493674, 27496224, 27496308, 27496424, 27497950 27498477, 27501327, 27501413, 27501465, 27502420, 27504190, 27504770 27505229, 27505603, 27506774, 27507968, 27508985, 27510959, 27511196 27512439, 27517818, 27518227, 27518310, 27520070, 27520900, 27522245 27523368, 27523800, 27525909, 27526744, 27532375, 27533819, 27534509 27537472,

27539757, 27540613, 27541286, 27541468, 27542824, 27544030 27544973, 27545630, 27547732, 27550341, 27551855, 27554074, 27555481 27558557, 27558559, 27558861, 27560602, 27560702, 27562488, 27563629 27563767, 27565906, 27567477, 27570318, 27576342, 27577758, 27578007, 27579353, 27580996, 27585755, 27585800, 27586810, 27586895, 27587672 27591842, 27592466, 27593389, 27593501, 27593585, 27595973, 27599689 27599927, 27601118, 27602091, 27602488, 27603841, 27604293, 27607563 27607805, 27608669, 27610269, 27613080, 27613247, 27613530, 27615608 27616657, 27617522, 27617978, 27620808, 27623844, 27625274, 27625620 27629756, 27631506, 27634676, 27634991, 27635508, 27644757, 27645940 27649707, 27652302, 27654521, 27655217, 27658186, 27658205, 27662528 27663370, 27664702, 27666312, 27671633, 27679488, 27679664, 27679793 27679806, 27679961, 27680162, 27680509, 27680669, 27682151, 27686599 27688036, 27688099, 27688692, 27690513, 27690578, 27691809, 27691920 27691939, 27692215, 27693416, 27693713, 27694261, 27695063, 27697092 27698953, 27700466, 27701795, 27704237, 27705761, 27707544, 27709046 27710072, 27718914, 27719187, 27723002, 27723151, 27726269, 27726780 27729678, 27732323, 27733415, 27739006, 27740424, 27740844, 27744211 27745220, 27747869, 27748954, 27751006, 27751755, 27753336, 27756900 27757567, 27757794, 27757888, 27758972, 27759077, 27759457, 27761402 27766324, 27767081, 27772093, 27772815, 27773602, 27774320, 27774539 27779886, 27780562, 27782339, 27783289, 27786772, 27791223, 27793533 27797290, 27801337, 27803665, 27807441, 27810967, 27812560, 27812593 27813267, 27815347, 27818871, 27824540, 27824543, 27825241, 27829295 27832643, 27833369, 27833672, 27834551, 27834984, 27835925, 27839353 27840386, 27843646, 27846298, 27846499, 27847259, 27849825, 27851757 27856471, 27861226, 27861452, 27861909, 27869075, 27869339, 27873643 27876671, 27882176, 27892488, 27896443, 27896458, 27898015, 27900663 27908644, 27912301, 27918832, 27920184, 27924147, 27926113, 27930478 27931299, 27934468, 27935348, 27938736, 27940876, 27941110, 27941896 27945870, 27948050, 27948153, 27950708, 27952762, 27959594, 27960021 27961746, 27964051, 27965830, 27966472, 27970265, 27971503, 27971575 27972265, 27975778, 27977039, 27983174, 27984028, 27986817, 27989556 27989849, 27991970, 27993289, 27994325, 27994333, 27995215, 27995248 27997875, 27998003, 27999073, 27999597, 27999638, 28000269, 28004853 28006704, 28018962, 28019283, 28021205, 28022101, 28022847, 28023081 28023399, 28023482, 28024347, 28026866, 28033429, 28036487, 28045903 28057267, 28058612, 28059199, 28067846, 28072130, 28072464, 28072567 28074713, 28085865, 28088762, 28090453, 28091981, 28098865, 28103600 28103869, 28104361, 28106402, 28108003, 28108898, 28111583, 28120036 28120951, 28124631, 28129791, 28132287, 28135648, 28165439, 28165545 28169711, 28174827, 28174951, 28175445, 28180464, 28181021, 28184554 28184800, 28187706, 28188330, 28189466, 28194173, 28199085, 28201419 28204443, 28209341, 28210192, 28211734, 28214943, 28215510, 28218832 28220398, 28223871, 28226179, 28227512, 28229360, 28236305, 28238264 28242712, 28258608, 28264172, 28271119, 28271693, 28279837, 28281094 28282606, 28285766, 28290434, 28302049, 28304709, 28305001, 28305362 28309182, 28312508, 28315031, 28315995, 28319623, 28320117, 28320399 28321446, 28323201, 28328895, 28329450, 28330714, 28333072, 28338399, 28338999, 28344964, 28350595, 28354603, 28357401, 28361083, 28361221 28361787, 28365111, 28369092, 28371123, 28378446, 28379065, 28385102 28390273, 28391582, 28392168, 28392251, 28393678, 28396445, 28403295 28413955, 28420042, 28420457, 28423598, 28432129, 28434028, 28435825 28445741, 28448314, 28455212, 28468312, 28475164, 28478676, 28481149 28481679, 28483184, 28489150, 28492362, 28493478, 28501075, 28502403 28502773, 28503484, 28504545, 28507324, 28508053, 28508296, 28508557 28512761, 28513333, 28514693, 28521330, 28527416, 28528349, 28535272 28544633, 28545134, 28545687, 28547068, 28553468, 28571483, 28572834, 28578164, 28580528, 28584193, 28584217, 28584444, 28587723, 28600233 28612674, 28614072, 28617631, 28617959, 28621470, 28627255, 28632559 28636676, 28644549, 28646200, 28670445, 28673203, 28679454, 28680029 28685371, 28692275, 28695694, 28702188, 28708023, 28710827, 28713840 28714058, 28728040, 28728272, 28730076, 28742555, 28747182, 28749289 28752599, 28755011, 28758090, 28758722, 28761812, 28767240, 28770146 28774416, 28777214, 28781754, 28785022, 28785531, 28803345, 28805612 28805695, 28809909, 28819640, 28849751, 28852691, 28856060, 28856172 28878525, 28881723, 28887509, 28891984, 28919145, 28925880, 28938924, 28940179, 28951014, 28951382, 28956908, 28960211, 28986231, 28987454 28993353, 28993590, 29000190, 29006527, 29015118, 29015706, 29024054 29027694, 29032276, 29033896, 29037290, 29051702, 29056270, 29123482 29125374, 29136111, 29139591, 29158680, 29171683, 29177886, 29189889 29200700, 29203604, 29205918, 29224605, 29230565, 29242017, 29301463 29314539, 29331209, 29331493, 29339155, 29347981, 29356752, 29364171 29376346, 29378913, 29379978,

29382784, 29388020, 29388952, 29394749 29395657, 29405462, 29409149, 29412269, 29429264, 29430524, 29436454 29437712, 29439522, 29504682, 29511611, 29531541, 29542449, 29616244 29676089, 29813494

Database Engine: 12.2.0.1

For Oracle 12c version 12.2.0.1, Oracle changed the way it releases Oracle Database updates. Instead of Patch Set Updates (PSUs), Oracle supplies Release Updates (RUs) and Release Updates Revisions (RURs). RUs contain optimizer changes, feature additions, and security fixes. RURs only contain security fixes for the two preceding quarterly patch cycles. With this new system, you have more control over the features that you install with each update.

The naming conventions have also changed for Oracle 12c version 12.2.0.1 versions. In previous versions, Amazon RDS for Oracle used the PSU naming convention of *oracle-version.vpatch-version*. The *patch-version* corresponded with an Oracle PSU. For example, in Oracle for Amazon RDS version 12.1.0.2.v13, the v13 part of the version number corresponds with an Oracle PSU.

Oracle 12c version 12.2.0.1 naming conventions account for both RU and RUR updates. For example, the first Amazon RDS for Oracle version available is 12.2.0.1.ru-2018-10.rur-2018-10.r1. In this example, 12.2 is the major version, and 0.1 is the minor version. The revision version has the following parts:

- ru-2018-10 – the October RU
- rur-2018-10 – the October RUR for the October RU
- r1 – Internal Amazon RDS revision, which lets Amazon RDS differentiate between emergency patches of pre-existing RU/RURs

For more information about the new Oracle Database versioning system, see the posts [Differences between PSU / BP and RU / RUR at the Upgrade your Database – NOW! blog](#) and [RU and RUR patches for Oracle 12.2 at the Oracle-Help blog](#).

The following versions are available for Oracle database engine 12.2.0.2:

- [Version 12.2.0.1.ru-2019-10.rur-2019-10.r1 \(p. 1182\)](#)
- [Version 12.2.0.1.ru-2019-07.rur-2019-07.r1 \(p. 1185\)](#)
- [Version 12.2.0.1.ru-2019-04.rur-2019-04.r1 \(p. 1188\)](#)
- [Version 12.2.0.1.ru-2019-01.rur-2019-01.r1 \(p. 1191\)](#)
- [Version 12.2.0.1.ru-2018-10.rur-2018-10.r1 \(p. 1193\)](#)

Version 12.2.0.1.ru-2019-10.rur-2019-10.r1

Version 12.2.0.1.ru-2019-10.rur-2019-10.r1 adds support for the following:

- Patch 30138470: DATABASE OCT 2019 RELEASE UPDATE 12.2.0.1.191015
- Patch 30133625: OJVM RELEASE UPDATE 12.2.0.1.191015
- Patch 29997937: DSTv34 for RDBMS (TZDATA2019G)
- Patch 29997959: DSTv34 for OJVM (TZDATA2019G)
- PreUpgrade Jar: preupgrade_12201_cbuild_23_if.zip
- Patch 28730253: SUPPORT NEW ERA REIWA FOR JAPANESE IMPERIAL CALENDAR is included in DB PATCH 30138470
- Adds support for [Resizing the Temporary Tablespace in a Read Replica \(p. 1104\)](#).

Oracle Release Update 12.2.0.1.191015, Released October 2019

Bugs fixed:

29013832, 30244787, 30253255, 8480838, 8932139, 12763598, 13554903 14221306, 14690846, 15931756, 16002385, 16438495, 16727454, 16942578 17027695, 17533661, 17947871, 18308268, 18521691, 18594510, 18774543 18878420, 19072655, 19211433, 19285025, 19327292, 19526548, 19614243 19647894, 19649997, 19702201, 19721304, 20003668, 20087519, 20118035 20120236, 20324049, 20436508, 20532077, 20588486, 20591151, 20617383 20620169, 20736227, 20756305, 20866970, 20976443, 21089435, 21095391 21143725, 21147908, 21159907, 21178363, 21186167, 21197098, 21216226 21320338, 21433452, 21479706, 21517767, 21520266, 21547051, 21638600 21744603, 21882528, 21935698, 21981529, 21985256, 22007324, 22070473 22070853, 22072543, 22087683, 22104866, 22107360, 22179537, 22310426 22347493, 22364044, 22367053, 22379010, 22446455, 22495673, 22503283 22503297, 22504793, 22522515, 22530986, 22564336, 22568728, 22581771 22594071, 22599050, 22628825, 22645009, 22654475, 22700845, 22726044 22729345, 22826067, 22843979, 22845846, 22864303, 22898198, 22921674 22939829, 22950945, 22970869, 22981722, 23019710, 23026585, 23035249 23055900, 23061453, 23065002, 23066146, 23080557, 23104033, 23105538 23110523, 23125560, 23126545, 23127945, 23151677, 23169712, 23177923 23179662, 23184263, 23197730, 23234232, 23237091, 23249829, 23271203 23278750, 23300142, 23306590, 23310101, 23312077, 23328639, 23336559 23342170, 23481673, 23491861, 23499004, 23499160, 23521523, 23527363 23533647, 23548817, 23567857, 23572982, 23581777, 23588722, 23599216 23600861, 23602213, 23645516, 23665623, 23709062, 23715460, 23715518 23730961, 23733981, 23735292, 23741944, 23743596, 23746128, 23749454 23761724, 24010030, 24289874, 24294174, 24303148, 24307571, 24308349 24326444, 24326846, 24328857, 24332831, 24334708, 24336249, 24337882 24341675, 24343905, 24345420, 24346821, 24348685, 24350620, 24352981 24355111, 24357348, 24368004, 24371491, 24373528, 24373756, 24374976 24376875, 24376878, 24383086, 24385983, 24401351, 24403922, 24409977 24415926, 24416451, 24421668, 24423416, 24425056, 24425998, 24435982 24437162, 24440648, 24443539, 24457597, 24460392, 24461826, 24467122 24468470, 24470606, 24471079, 24471473, 24473736, 24485034, 24485161 24485174, 24486059, 24486237, 24509056, 24534401, 24554533, 24555417 24556967, 24560906, 24563422, 24570214, 24570598, 24573817, 24578718 24578797, 24589081, 24589590, 24591506, 24593740, 24595699, 24600330 24609592, 24609996, 24611527, 24616637, 24617969, 24623975, 24624166 24642495, 24654629, 24655717, 24664211, 24668398, 24669189, 24674197 24674955, 24676172, 24677696, 24680959, 24687075, 24689376, 24692973 24693290, 24697323, 24699619, 24710696, 24713381, 24714096, 24717183 24717859, 24718260, 24719799, 24735430, 24737064, 24737403, 24737581 24737954, 24739791, 24744383, 24744686, 24757934, 24759556, 24760407 24764085, 24766309, 24786669, 24792678, 24793511, 24796092, 24797119 24800423, 24801152, 24802934, 24811725, 24812047, 24818566, 24827228 24827654, 24831514, 24835919, 24841671, 24843188, 24844549, 24844841 24845157, 24848746, 24848923, 24850622, 24907917, 24908321, 24911709 24912588, 24922704, 24923080, 24923215, 24923338, 24923790, 24924667 24926999, 24929210, 24938784, 24940060, 24942749, 24953434, 24957555 24960044, 24965426, 24966594, 24966788, 24967993, 24968162, 24976007 24978100, 25022574, 25027852, 25028996, 25029022, 25029423, 25032818 25034396, 25036006, 25036474, 25042823, 25044977, 25045228, 25050160 25051465, 25051628, 25057811, 25058080, 25060506, 25062592, 25063971 25065563, 25072986, 25078611, 25086233, 25087436, 25092777, 25093872 25095982, 25098160, 25099339, 25099497, 25099758, 25100063, 25100579 25103996, 25107662, 25110233, 25114561, 25120284, 25120668, 25120742 25121089, 25123585, 25124363, 25129925, 25130312, 25140197, 25145163 25145215, 25150925, 25159176, 25162645, 25164293, 25166187, 25171041 25171084, 25175723, 25176408, 25178032, 25178101, 25178179, 25179774 25182817, 25184453, 25184555, 25186079, 25191872, 25192044, 25192528 25192729, 25195901, 25199585, 25201454, 25202355, 25203656, 25206864 25207410, 25209912, 25210268, 25210499, 25211628, 25223839, 25224242 25225795, 25226665, 25227381, 25230870, 25230945, 25237577, 25240188 25240590, 25241448, 25241625, 25244807, 25248384, 25251648, 25257085 25259611, 25262869, 25263960, 25265499, 25283790, 25287072, 25293659 25296876, 25299227, 25299807, 25300427, 25303284, 25303756, 25305405 25307368, 25309116, 25313154, 25313411, 25316758, 25317989, 25320555 25323525, 25328518, 25329664, 25335249, 25335360, 25335790, 25337332 25337640, 25348956, 25353983, 25356118, 25357142, 25360661, 25362958 25367588, 25367721, 25382812, 25383204, 25384462, 25386748, 25388896 25392535,

25393714, 25395696, 25397936, 25398306, 25404202, 25405100 25405687, 25405813, 25410017, 25410180, 25410802, 25410877, 25411036 25415713, 25417050, 25417056, 25417958, 25425451, 25425760, 25427662 25429959, 25430120, 25433696, 25435038, 25437699, 25440818, 25442559 25444961, 25445168, 25451531, 25455795, 25457409, 25459958, 25462714 25463844, 25472112, 25476149, 25477657, 25478885, 25479164, 25482971 25489342, 25489367, 25489607, 25492379, 25498930, 25498994, 25516250 25524955, 25528838, 25530080, 25530814, 25535668, 25536819, 25537470 25539063, 25540738, 25546580, 25546608, 25547901, 25551676, 25553616 25554787, 25555252, 25557886, 25558986, 25560487, 25561296, 25569149 25570929, 25575348, 25575369, 25575628, 25579458, 25579761, 25594901 25597525, 25598473, 25599425, 25600342, 25600421, 25602488, 25603923 25606091, 25607726, 25612095, 25614866, 25616268, 25616359, 25616417 25616645, 25631933, 25633101, 25634317, 25634348, 25635149, 25638456 25639019, 25643818, 25643889, 25643931, 25646373, 25647325, 25648731 25653109, 25654459, 25654936, 25655390, 25655966, 25659655, 25660847 25661819, 25662088, 25662101, 25662524, 25663488, 25669791, 25670786 25671354, 25672640, 25674386, 25680221, 25685152, 25686739, 25687460 25691904, 25694206, 25695903, 25696520, 25699321, 25700654, 25709368 25710420, 25715167, 25717371, 25722055, 25722608, 25722720, 25723158 25728085, 25729507, 25734963, 25736747, 25739065, 25741955, 25743479 25747569, 25749273, 25752755, 25754606, 25756945, 25757748, 25760195 25762221, 25764020, 25766822, 25768681, 25772669, 25774077, 25775213 25780343, 25783447, 25784002, 25785331, 25785441, 25788879, 25789041 25789277, 25789579, 25790353, 25792911, 25795865, 25797092, 25797124 25797305, 25800464, 25802510, 25803545, 25807997, 25810263, 25810704 25811650, 25813931, 25818707, 25822410, 25823754, 25825910, 25826740 25830492, 25832935, 25834581, 25835365, 25838361, 25838755, 25852885 25856821, 25858672, 25861398, 25865785, 25870579, 25871177, 25871639 25871753, 25872127, 25872389, 25873336, 25874050, 25874678, 25882264 25883438, 25885148, 25888073, 25888984, 25890056, 25890673, 25890782 25894239, 25895224, 25897615, 25904273, 25904490, 25906117, 25906886 25908728, 25911724, 25914276, 25919622, 25932524, 25932728, 25933494 25941836, 25943271, 25945130, 25947799, 25951571, 25953857, 25954022 25954054, 25957038, 25963024, 25964954, 25967544, 25967985, 25970731 25971286, 25973152, 25975723, 25977302, 25980605, 25980770, 25981498 25982666, 25986062, 25990907, 25995938, 25997810, 26006257, 26007010 26019148, 26024732, 26024784, 26025681, 26029075, 26029777, 26029780 26032573, 26034119, 26036748, 26037215, 26038086, 26039623, 26040483 26045732, 26051656, 26078437, 26078493, 26080410, 26083298, 26088426 26088836, 26090767, 26091640, 26091786, 26095327, 26095405, 26096382 26108080, 26108337, 26110259, 26110632, 26111842, 26112621, 26115103 26121990, 26124078, 26130486, 26137367, 26138085, 26149904, 26153372 26153977, 26168933, 26169341, 26169345, 26170659, 26170715, 26176002 26187943, 26189861, 26198757, 26198926, 26201113, 26203182, 26223039 26237338, 26237431, 26237773, 26238195, 26242031, 26242677, 26243698 26244115, 26245237, 26248143, 26249718, 26256131, 26257953, 26259265 26261327, 26263328, 26263721, 26268756, 26269790, 26271001, 26275023 26275415, 26277439, 26281476, 26285062, 26285933, 26308650, 26309047 26317991, 26318627, 26323308, 26324206, 26324769, 26327624, 26330994 26331743, 26333141, 26336977, 26338953, 26351334, 26353617, 26358670 26359091, 26362155, 26362821, 26366517, 26367012, 26367460, 26371725 26374791, 26375052, 26375250, 26380097, 26385189, 26388538, 26396790 26399626, 26399691, 26406387, 26407408, 26412540, 26418088, 26420561 26421667, 26422277, 26423085, 26426526, 26426967, 26430323, 26430737 26434436, 26434999, 26435073, 26436168, 26438612, 26439748, 26440169 26440749, 26442308, 26444601, 26444887, 26446098, 26452606, 26474662 26474703, 26475419, 26476244, 26478970, 26479173, 26482376, 26486365 26492866, 26493289, 26498354, 26513067, 26513709, 26521043, 26522439 26523432, 26526726, 26526799, 26536320, 26537307, 26542135, 26542236 26544823, 26545688, 26546070, 26546664, 26546754, 26548363, 26556014 26558437, 26569225, 26575788, 26580633, 26582460, 26584641, 26597140 26599395, 26608137, 26608238, 26609492, 26615291, 26615690, 26617804 26623652, 26626879, 26629381, 26633355, 26633558, 26635897, 26637273 26637824, 26639167, 26641610, 26650226, 26654363, 26658759, 26659182 26680105, 26712331, 26714486, 26714910, 26717528, 26724511, 26725687 26727397, 26729494, 26729611, 26740700, 26744595, 26745002, 26751106 26751171, 26755171, 26758193, 26764561, 26765212, 26768025, 26775602 26784509, 26794786, 26797591, 26798411, 26798516, 26802503, 26816582 26820076, 26822620, 26828994, 26829845, 26833932, 26837702, 26840654 26844406, 26844870, 26849779, 26871815, 26875822, 26883456, 26895149 26896659, 26898563, 26907327, 26908788, 26909100, 26909504, 26910716 26911000, 26939314, 26943004, 26944190, 26958896, 26963310, 26966616 26966916, 26967713, 26969321,

26970175, 26970717, 26981902, 26983259 26985002, 26986173, 26992964, 27000158, 27006120, 27006664, 27009164 27013146, 27028251, 27032785, 27033520, 27034890, 27037839, 27038986 27039712, 27044297, 27052607, 27058530, 27060167, 27060859, 27061736 27072923, 27073314, 27079140, 27087426, 27090765, 27092508, 27093423 27097854, 27101105, 27105900, 27106179, 27110878, 27115422, 27117822 27119621, 27122162, 27124624, 27125872, 27133662, 27134734, 27135647 27135993, 27138325, 27142373, 27142529, 27144928, 27151826, 27153641 27160922, 27161071, 27162390, 27162405, 27163928, 27165231, 27169796 27170305, 27181537, 27181897, 27185188, 27199245, 27200959, 27202015 27207110, 27208795, 27213224, 27216046, 27223075, 27229389, 27231051 27236722, 27242226, 27244337, 27248917, 27249531, 27250547, 27251690 27255377, 27256000, 27258578, 27259307, 27262945, 27264464, 27266245 27274456, 27274536, 27275533, 27276231, 27283960, 27284499, 27285244 27288638, 27292213, 27293599, 27302711, 27302730, 27303287, 27304410 27304906, 27305039, 27308088, 27314206, 27314390, 27320576, 27321179 27329612, 27333106, 27334316, 27338912, 27338946, 27339115, 27345231 27346709, 27348081, 27349393, 27350267, 27351628, 27359178, 27364854 27367194, 27369515, 27370965, 27375542, 27381498, 27383281, 27386467 27393570, 27394703, 27395416, 27396624, 27396672, 27396813, 27397048 27400416, 27400598, 27404668, 27405645, 27416997, 27423251, 27424405 27426363, 27432062, 27432826, 27433385, 27433870, 27434193, 27439835 27441326, 27442041, 27445727, 27457891, 27459948, 27466597, 27468303 27493674, 27501373, 27501413, 27502420, 27504770, 27505229, 27508985 27510959, 27525909, 27533780, 27533819, 27534509, 27540613, 27544973 27548131, 27554074, 27555481, 27558861, 27560602, 27562488, 27565906 27567477, 27576342, 27576354, 27588271, 27593501, 27595973, 27601118 27601441, 27607563, 27611612, 27613080, 27613530, 27617978, 27620808 27623159, 27629756, 27632114, 27634676, 27634991, 27658186, 27666312 27671633, 27680669, 27686599, 27687880, 27688036, 27688099, 27688692, 27691920, 27691939, 27693416, 27693713, 27695063, 27698953, 27700466 27704237, 27709046, 27710072, 27726780, 27740424, 27748954, 27751755 27757567, 27757888, 27758544, 27758653, 27758972, 27759077, 27793533 27799032, 27801337, 27824540, 27824543, 27825241, 27829295, 27833672 27834551, 27834569, 27835925, 27839353, 27846298, 27846499, 27847259 27850112, 27855490, 27861226, 27873412, 27882176, 27886087, 27898015 27902561, 27931299, 27935493, 27940876, 27945870, 27951817, 27959048 27959594, 27966472, 27967484, 27983174, 27986817, 27994325, 27995215 27995248, 27997875, 27998003, 27999073, 27999638, 28000269, 28022101 28023081, 28023399, 28023482, 28026866, 28033429, 28040776, 28043157 28045903, 28067846, 28072567, 28074713, 28090453, 28099662, 28108003 28111583, 28120036, 28120951, 28124631, 28125947, 28129791, 28140658 28165439, 28171079, 28174827, 28180464, 28181021, 28184554, 28188330 28194173, 28199085, 28201419, 28215510, 28218832, 28220398, 28223871 28226179, 28229360, 28236305, 28242712, 28256164, 28271119, 28279837 28281094, 28282606, 28290434, 28294563, 28302049, 28305001, 28305362 28319114, 28320399, 28354603, 28357401, 28361221, 28365111, 28369092 28371123, 28378446, 28384353, 28390273, 28391210, 28396445, 28401116 28420042, 28420457, 28423598, 28432129, 28434028, 28435902, 28437315 28454242, 28468312, 28481149, 28483184, 28489150, 28501075, 28502343 28507324, 28508053, 28508557, 28521330, 28522441, 28528349, 28535272 28537715, 28542455, 28545134, 28547068, 28571483, 28572834, 28578164 28585411, 28587723, 28600233, 28612674, 28617631, 28617959, 28621470 28627255, 28636676, 28678804, 28691965, 28692103, 28692275, 28697806 28708023, 28714988, 28728272, 28742555, 28749289, 28749724, 28758722 28774416, 28791725, 28803345, 28817449, 28819640, 28820669, 28831971 28849751, 28852691, 28855922, 28856060, 28856172, 28867992, 28889730 28891984, 28951026, 28951382, 28956908, 28960211, 28965787, 28986231 28987439, 28991884, 28993590, 29002488, 29006527, 29009513, 29024054 29027694, 29032276, 29125374, 29158680, 29163567, 29189889, 29198092 29200700, 29213320, 29224605, 29250230, 29250317, 29301463, 29339155 29347943, 29353821, 29376346, 29378913, 29379978, 29388020, 29405462 29409149, 29436454, 29437712, 29483672, 29483723, 29500257, 29511611 29542449, 29633753, 29637526, 29645349, 29676089, 29690625, 29707896 29724063, 29767177, 29782211, 29813494, 29836659, 29893132, 29902311 30088912, 30189023

Version 12.2.0.1.ru-2019-07.rur-2019-07.r1

Version 12.2.0.1.ru-2019-07.rur-2019-07.r1 adds support for the following:

- Patch 29757449: DATABASE JUL 2019 RELEASE UPDATE 12.2.0.1.190716

- Patch 29774415: OJVM RELEASE UPDATE 12.2.0.1.190716
- Patch 28125601: DSTv33 for RDBMS (TZDATA2018G)
- Patch 28127287: DSTv33 for OJVM (TZDATA2018G)
- PreUpgrade Jar: preupgrade_12201_cbuild_22_lf.zip
- Patch 29213893: DBMS_STATS FAILING WITH ERROR ORA-01422 WHEN GATHERING STATS FOR USER \$ TABLE
- Patch 28730253: SUPPORT NEW ERA REIWA FOR JAPANESE IMPERIAL CALENDAR

Oracle Release Update 12.2.0.1.190716, Released July 2019

Bugs fixed:

8480838, 8932139, 12763598, 13554903, 14221306, 14690846, 15931756 16002385, 16438495, 16727454, 16942578, 17027695, 17533661, 17947871 18308268, 18521691, 18594510, 18774543, 18878420, 19072655, 19211433 19285025, 19327292, 19526548, 19614243, 19647894, 19649997, 19702201 19721304, 20003668, 20087519, 20118035, 20120236, 20324049, 20436508 20532077, 20588486, 20591151, 20617383, 20620169, 20736227, 20756305 20866970, 20976443, 21089435, 21095391, 21143725, 21147908, 21159907 21178363, 21186167, 21197098, 21216226, 21320338, 21433452, 21479706 21517767, 21520266, 21547051, 21638600, 21744603, 21882528, 21981529 21985256, 22007324, 22070473, 22070853, 22072543, 22087683, 22104866 22179537, 22310426, 22347493, 22364044, 22367053, 22379010, 22446455 22495673, 22503283, 22503297, 22504793, 22522515, 22530986, 22564336 22568728, 22581771, 22594071, 22599050, 22628825, 22645009, 22654475 22700845, 22726044, 22729345, 22826067, 22843979, 22845846, 22864303 22898198, 22950945, 22970869, 22981722, 23019710, 23026585, 23035249 23055900, 23061453, 23065002, 23066146, 23080557, 23104033, 23105538 23110523, 23125560, 23126545, 23127945, 23151677, 23169712, 23177923 23179662, 23184263, 23197730, 23234232, 23237091, 23249829, 23271203 23278750, 23300142, 23306590, 23310101, 23312077, 23328639, 23336559 23481673, 23491861, 23499004, 23499160, 23521523, 23527363, 23533647 23548817, 23567857, 23572982, 23581777, 23588722, 23599216, 23600861 23602213, 23645516, 23665623, 23709062, 23715460, 23715518, 23730961 23733981, 23735292, 23741944, 23746128, 23749454, 23761724, 24010030 24289874, 24294174, 24303148, 24307571, 24308349, 24326444, 24326846 24328857, 24332831, 24334708, 24336249, 24337882, 24341675, 24343905 24345420, 24346821, 24348685, 24350620, 24352981, 24355111, 24357348 24368004, 24371491, 24373528, 24373756, 24374976, 24376875, 24376878 24385983, 24401351, 24403922, 24409977, 24415926, 24416451, 24421668 24423416, 24425056, 24425998, 24435982, 24437162, 24443539, 24457597 24460392, 24461826, 24467122, 24468470, 24470606, 24471079, 24471473 24473736, 24485034, 24485161, 24485174, 24486059, 24486237, 24509056 24534401, 24554533, 24555417, 24556967, 24560906, 24563422, 24570214 24570598, 24573817, 24578718, 24578797, 24589081, 24589590, 24591506 24593740, 24595699, 24600330, 24609592, 24609996, 24611527, 24616637 24617969, 24623975, 24624166, 24642495, 24654629, 24655717, 24664211 24668398, 24669189, 24674197, 24674955, 24676172, 24677696, 24680959 24689376, 24692973, 24693290, 24697323, 24699619, 24710696, 24713381 24714096, 24717183, 24717859, 24718260, 24719799, 24735430, 24737064 24737403, 24737581, 24744383, 24744686, 24757934, 24759556, 24760407 24764085, 24766309, 24786669, 24792678, 24793511, 24796092, 24797119 24800423, 24801152, 24802934, 24811725, 24812047, 24827228, 24827654 24831514, 24835919, 24841671, 24843188, 24844549, 24844841, 24845157 24848746, 24848923, 24850622, 24907917, 24908321, 24911709, 24912588 24922704, 24923080, 24923215, 24923338, 24923790, 24924667, 24926999 24929210, 24938784, 24940060, 24942749, 24953434, 24957555, 24960044 24965426, 24966594, 24966788, 24967993, 24968162, 24976007, 24978100 25022574, 25027852, 25028996, 25029022, 25029423, 25032818, 25034396 25036474, 25042823, 25044977, 25045228, 25050160, 25051465, 25051628 25057811, 25058080, 25062592, 25063971, 25065563, 25072986, 25078611 25086233, 25087436, 25092777, 25093872, 25098160, 25099339, 25099497 25099758, 25100063, 25100579, 25103996, 25107662, 25110233, 25114561 25120284, 25120668, 25120742, 25121089, 25123585, 25124363, 25129925 25140197, 25145163, 25145215, 25150925, 25159176, 25162645, 25164293 25166187, 25171041, 25171084, 25175723, 25176408, 25178032, 25178101 25178179, 25179774, 25182817, 25184555, 25186079, 25191872, 25192044 25192528, 25192729,

25199585, 25201454, 25202355, 25203656, 25206864 25207410, 25209912, 25210268, 25210499, 25211628, 25223839, 25224242 25225795, 25226665, 25227381, 25230870, 25230945, 25237577, 25240188 25240590, 25241448, 25241625, 25244807, 25248384, 25251648, 25257085 25259611, 25262869, 25263960, 25265499, 25283790, 25287072, 25293659 25296876, 25299227, 25299807, 25300427, 25303756, 25305405, 25307368 25309116, 25313154, 25313411, 25316758, 25317989, 25320555, 25323525 25328518, 25329664, 25335249, 25335360, 25335790, 25337332, 25337640 25348956, 25353983, 25356118, 25357142, 25360661, 25362958, 25367588 25367721, 25382812, 25383204, 25384462, 25386748, 25388896, 25392535 25395696, 25397936, 25398306, 25404202, 25405100, 25405687, 25405813 25410017, 25410180, 25410802, 25410877, 25411036, 25415713, 25417050 25417056, 25417958, 25425451, 25425760, 25427662, 25429959, 25430120 25433696, 25435038, 25437699, 25440818, 25442559, 25444961, 25445168 25451531, 25455795, 25457409, 25459958, 25462714, 25463844, 25472112 25476149, 25477657, 25478885, 25479164, 25482971, 25489342, 25489367 25489607, 25492379, 25498930, 25498994, 25516250, 25524955, 25528838 25530080, 25530814, 25535668, 25536819, 25537470, 25539063, 25540738 25546580, 25546608, 25547901, 25551676, 25553616, 25554787, 25555252 25557886, 25558986, 25560487, 25561296, 25569149, 25570929, 25575348 25575369, 25575628, 25579458, 25579761, 25594901, 25597525, 25598473 25599425, 25600342, 25600421, 25602488, 25603923, 25606091, 25607726 25612095, 25614866, 25616268, 25616359, 25616417, 25616645, 25631933 25633101, 25634317, 25634348, 25635149, 25638456, 25639019, 25643818 25643889, 25643931, 25646373, 25647325, 25648731, 25653109, 25654459 25654936, 25655390, 25655966, 25659655, 25660847, 25661819, 25662088 25662101, 25662524, 25663488, 25669791, 25670786, 25671354, 25672640 25674386, 25680221, 25685152, 25686739, 25687460, 25691904, 25694206 25695903, 25699321, 25700654, 25709368, 25710420, 25715167, 25717371 25722055, 25722608, 25722720, 25723158, 25728085, 25729507, 25734963 25736747, 25739065, 25741955, 25743479, 25747569, 25749273, 25752755 25754606, 25757748, 25760195, 25762221, 25764020, 25766822, 25768681 25772669, 25774077, 25775213, 25780343, 25783447, 25784002, 25785331 25785441, 25788879, 25789041, 25789277, 25789579, 25790353, 25792911 25795865, 25797092, 25797124, 25797305, 25800464, 25802510, 25803545 25807997, 25810263, 25810704, 25811650, 25813931, 25818707, 25822410 25823754, 25825910, 25826740, 25830492, 25832935, 25834581, 25835365 25838361, 25838755, 25852885, 25856821, 25858672, 25861398, 25865785 25870579, 25871177, 25871639, 25871753, 25872127, 25872389, 25873336 25874050, 25874678, 25882264, 25883438, 25885148, 25888073, 25888984 25890056, 25890673, 25894239, 25895224, 25897615, 25904273, 25904490 25906117, 25906886, 25908728, 25911724, 25914276, 25919622, 25932524 25932728, 25933494, 25941836, 25943271, 25945130, 25947799, 25951571 25953857, 25954022, 25954054, 25957038, 25963024, 25964954, 25967544 25967985, 25970731, 25971286, 25973152, 25975723, 25977302, 25980605 25980770, 25981498, 25982666, 25986062, 25990907, 25995938, 25997810 26006257, 26007010, 26019148, 26024732, 26024784, 26025681, 26029075 26029777, 26029780, 26032573, 26034119, 26036748, 26037215, 26038086 26039623, 26040483, 26045732, 26051656, 26078437, 26078493, 26080410 26083298, 26088426, 26088836, 26090767, 26091640, 26091786, 26095327 26095405, 26096382, 26108080, 26108337, 26110259, 26110632, 26111842 26112621, 26115103, 26121990, 26124078, 26137367, 26138085, 26149904 26153977, 26168933, 26169341, 26169345, 26170659, 26170715, 26176002 26187943, 26189861, 26198757, 26198926, 26201113, 26203182, 26223039 26237338, 26237431, 26237773, 26238195, 26242031, 26243698, 26244115 26245237, 26248143, 26249718, 26256131, 26257953, 26259265, 26261327 26263328, 26263721, 26268756, 26269790, 26271001, 26275023, 26275415 26277439, 26281476, 26285062, 26285933, 26308650, 26309047, 26317991 26318627, 26323308, 26324206, 26324769, 26327624, 26330994, 26331743 26333141, 26336977, 26338953, 26351334, 26353617, 26358670, 26359091 26362155, 26362821, 26366517, 26367012, 26367460, 26371725, 26374791 26375052, 26375250, 26380097, 26385189, 26388538, 26396790, 26399626 26399691, 26406387, 26407408, 26412540, 26418088, 26420561, 26421667 26422277, 26423085, 26426526, 26426967, 26430323, 26430737, 26434436 26434999, 26435073, 26436168, 26438612, 26439748, 26440169, 26440749 26442308, 26444601, 26444887, 26446098, 26452606, 26474662, 26474703 26475419, 26476244, 26478970, 26479173, 26482376, 26486365, 26492866 26493289, 26498354, 26513067, 26513709, 26521043, 26522439, 26523432 26526726, 26526799, 26536320, 26537307, 26542135, 26542236, 26544823 26545688, 26546070, 26546664, 26546754, 26548363, 26556014, 26569225 26575788, 26580633, 26582460, 26584641, 26597140, 26599395, 26608137 26608238, 26609942, 26615291, 26615690, 26617804, 26623652, 26626879 26629381, 26633355, 26633558, 26635897, 26637273, 26637824, 26639167 26641610, 26650226, 26654363, 26658759,

26659182, 26680105, 26712331 26714486, 26714910, 26717528, 26725687, 26727397, 26729494, 26729611 26740700, 26744595, 26745002, 26751106, 26751171, 26755171, 26758193 26764561, 26765212, 26775602, 26784509, 26794786, 26797591, 26798516 26802503, 26816582, 26820076, 26822620, 26828994, 26829845, 26833932 26837702, 26840654, 26844406, 26844870, 26849779, 26875822, 26883456 26895149, 26896659, 26898563, 26907327, 26908788, 26909100, 26909504 26911000, 26939314, 26943004, 26944190, 26958896, 26963310, 26966616 26966916, 26967713, 26969321, 26970175, 26970717, 26981902, 26983259 26985002, 26986173, 26992964, 27006120, 27006664, 27009164, 27013146 27028251, 27032785, 27034890, 27037839, 27038986, 27039712, 27044297 27052607, 27058530, 27060167, 27060859, 27061736, 27073314, 27079140 27087426, 27090765, 27092508, 27093423, 27097854, 27101105, 27105900 27106179, 27110878, 27115422, 27117822, 27119621, 27122162, 27124624 27125872, 27133662, 27134734, 27135647, 27135993, 27138325, 27142373 27142529, 27151826, 27153641, 27161071, 27162390, 27162405, 27163928 27165231, 27169796, 27170305, 27181537, 27181897, 27199245, 27200959 27207110, 27208795, 27213224, 27216046, 27223075, 27229389, 27231051 27236722, 27242226, 27244337, 27248917, 27249531, 27250547, 27251690 27255377, 27256000, 27258578, 27259307, 27262945, 27266245, 27274456 27274536, 27275533, 27276231, 27283960, 27284499, 27285244, 27288638 27292213, 27293599, 27302711, 27302730, 27303287, 27304410, 27305039 27308088, 27314206, 27314390, 27320576, 27321179, 27329612, 27333106 27334316, 27338912, 27338946, 27339115, 27345231, 27346709, 27348081 27349393, 27350267, 27351628, 27359178, 27364854, 27367194, 27370965 27375542, 27381498, 27386467, 27394703, 27395416, 27396624, 27396672 27396813, 27397048, 27400416, 27400598, 27404668, 27405645, 27416997 27423251, 27424405, 27426363, 27432062, 27433870, 27434193, 27439835 27441326, 27442041, 27445727, 27457891, 27466597, 27468303, 27493674 27501373, 27501413, 27502420, 27504770, 27505229, 27508985, 27510959 27525909, 27533819, 27534509, 27540613, 27544973, 27548131, 27554074 27555481, 27558861, 27560602, 27562488, 27565906, 27567477, 27576342 27593501, 27595973, 27607563, 27611612, 27613080, 27613530, 27617978 27620808, 27629756, 27634676, 27634991, 27658186, 27666312, 27671633 27680669, 27686599, 27687880, 27688036, 27688099, 27688692, 27691920 27691939, 27693416, 27693713, 27695063, 27698953, 27700466, 27704237 27709046, 27710072, 27726780, 27740424, 27748954, 27751755, 27757567 27757888, 27758972, 27759077, 27793533, 27799032, 27801337, 27824540 27824543, 27825241, 27829295, 27833672, 27834551, 27835925, 27846298 27846499, 27847259, 27855490, 27861226, 27882176, 27898015, 27931299 27940876, 27945870, 27951817, 27959048, 27959594, 27966472, 27986817 27994325, 27995215, 27995248, 27997875, 27998003, 27999073, 27999638 28000269, 28022101, 28023081, 28023399, 28023482, 28026866, 28033429 28040776, 28045903, 28067846, 28072567, 28074713, 28090453, 28099662 28108003, 28111583, 28120036, 28120951, 28124631, 28129791, 28140658 28165439, 28171079, 28174827, 28180464, 28181021, 28184554, 28188330 28194173, 28199085, 28201419, 28218832, 28220398, 28223871, 28226179 28229360, 28236305, 28271119, 28279837, 28282606, 28290434, 28302049 28305001, 28305362, 28320399, 28354603, 28357401, 28361221, 28365111 28378446, 28390273, 28396445, 28420042, 28420457, 28423598, 28432129 28434028, 28437315, 28454242, 28468312, 28483184, 28489150, 28501075 28502343, 28507324, 28508053, 28508557, 28522441, 28528349, 28535272 28545134, 28547068, 28571483, 28572834, 28578164, 28587723, 28600233 28612674, 28617631, 28617959, 28621470, 28627255, 28636676, 28691965 28692275, 28708023, 28728272, 28742555, 28749289, 28758722, 28774416 28803345, 28819640, 28849751, 28852691, 28856060, 28891984, 28951382 28956908, 28960211, 28965787, 28986231, 28987439, 28991884, 28993590 29006527, 29024054, 29027694, 29032276, 29125374, 29158680, 29198889 29200700, 29224605, 29250230, 29301463, 29339155, 29347943, 29376346 29378913, 29379978, 29388020, 29405462, 29436454, 29437712, 29511611 29542449, 29676089, 29690625, 29813494, 29836659

Version 12.2.0.1.ru-2019-04.rur-2019-04.r1

Version 12.2.0.1.ru-2019-04.rur-2019-04.r1 adds support for the following:

- Patch 29314339: Database Apr 2019 Release Update : 12.2.0.1.190416
- Patch 29249637: OJVM RELEASE UPDATE: 12.2.0.1.190416
- Patch 28852325: DSTv33 for RDBMS (TZDATA2018G)
- Patch 28852334: DSTv33 for OJVM (TZDATA2018G)

- PreUpgrade Jar: preupgrade_12201_cbuild_21_1f.zip
- Patch 28423598: GOLDENGATE AUTH CAUSES ACTIVE DG TO BE UNUSABLE UNTIL BOUNCE
- Patch 28730253: SUPPORT NEW ERA REIWA FOR JAPANESE IMPERIAL CALENDAR

Oracle Release Update 12.2.0.1.190416, Released April 2019

Bugs fixed:

26362155, 28023399, 25741955, 25873336, 26966616, 27097854, 28617631 28742555, 29006527, 8480838, 8932139, 12763598, 13554903, 14221306 14690846, 15931756, 16002385, 16438495, 16727454, 16942578, 17027695 17533661, 17947871, 18308268, 18521691, 18594510, 18774543, 18878420 19072655, 19211433, 19285025, 19327292, 19526548, 19614243, 19647894 19649997, 19702201, 19721304, 20003668, 20087519, 20118035, 20120236 20324049, 20436508, 20532077, 20591151, 20617383, 20620169, 20736227 20756305, 20866970, 20976443, 21089435, 21095391, 21143725, 21147908 21159907, 21178363, 21186167, 21197098, 21216226, 21320338, 21433452 21479706, 21520266, 21547051, 21638600, 21744603, 21882528, 21981529 21985256, 22007324, 22070853, 22072543, 22087683, 22104866, 22179537 22347493, 22364044, 22367053, 22379010, 22446455, 22495673, 22503283 22503297, 22504793, 22530986, 22564336, 22568728, 22581771, 22594071 22599050, 22628825, 22645009, 22654475, 22700845, 22726044, 22729345 22826067, 22843979, 22845846, 22864303, 22898198, 22950945, 22970869 22981722, 23019710, 23026585, 23035249, 23055900, 23061453, 23065002 23066146, 23080557, 23104033, 23105538, 23110523, 23125560, 23126545 23127945, 23151677, 23179662, 23184263, 23197730, 23234232, 23237091 23249829, 23271203, 23278750, 23300142, 23310101, 23312077, 23328639 23336559, 23481673, 23491861, 23499004, 23499160, 23521523, 23527363 23533647, 23548817, 23567857, 23572982, 23581777, 23588722, 23599216 23600861, 23602213, 23645516, 23665623, 23709062, 23715460, 23715518 23730961, 23733981, 23735292, 23741944, 23746128, 23749454, 23761724 24010030, 24289874, 24294174, 24303148, 24307571, 24308349, 24326444 24326846, 24328857, 24332831, 24334708, 24336249, 24337882, 24341675 24343905, 24345420, 24346821, 24348685, 24350620, 24352981, 24368004 24371491, 24373756, 24374976, 24376875, 24376878, 24385983, 24401351 24403922, 24415926, 24421668, 24423416, 24425056, 24425998, 24435982 24437162, 24443539, 24457597, 24460392, 24461826, 24467122, 24468470 24470606, 24471473, 24473736, 24485034, 24485161, 24485174, 24486059 24486237, 24509056, 24534401, 24554533, 24555417, 24556967, 24560906 24563422, 24570214, 24570598, 24573817, 24578718, 24578797, 24589081 24589590, 24593740, 24595699, 24600330, 24609592, 24609996, 24611527 24616637, 24617969, 24623975, 24624166, 24642495, 24654629, 24655717 24664211, 24668398, 24669189, 24674197, 24674955, 24676172, 24677696 24680959, 24689376, 24692973, 24693290, 24697323, 24699619, 24710696 24713381, 24714096, 24717183, 24717859, 24718260, 24719799, 24735430 24737064, 24737403, 24737581, 24744383, 24744686, 24757934, 24759556 24760407, 24764085, 24766309, 24786669, 24792678, 24793511, 24796092 24797119, 24800423, 24801152, 24802934, 24811725, 24812047, 24827228 24827654, 24831514, 24835919, 24843188, 24844549, 24845157, 24848746 24848923, 24850622, 24907917, 24908321, 24911709, 24912588, 24922704 24923080, 24923215, 24923338, 24923790, 24924667, 24926999, 24929210 24938784, 24940060, 24942749, 24953434, 24957555, 24960044, 24966594 24966788, 24968162, 24976007, 24978100, 25022574, 25027852, 25028996 25029022, 25029423, 25032818, 25034396, 25036474, 25042823, 25044977 25045228, 25050160, 25051628, 25057811, 25058080, 25062592, 25063971 25065563, 25072986, 25078611, 25086233, 25087436, 25093872, 25098160 25099339, 25099497, 25099758, 25100063, 25100579, 25103996, 25107662 25110233, 25114561, 25120284, 25120668, 25120742, 25121089, 25123585 25124363, 25129925, 25140197, 25145163, 25145215, 25150925, 25159176 25162645, 25164293, 25166187, 25171084, 25175723, 25176408, 25178032 25178101, 25178179, 25179774, 25182817, 25184555, 25186079, 25191872 25192044, 25192729, 25199585, 25201454, 25202355, 25203656, 25206864 25207410, 25209912, 25210268, 25210499, 25211628, 25223839, 25224242 25225795, 25226665, 25227381, 25230870, 25230945, 25237577, 25240188 25240590, 25241448, 25241625, 25244807, 25248384, 25251648, 25257085 25259611, 25262869, 25263960, 25265499, 25283790, 25287072, 25293659 25296876, 25299227, 25299807, 25300427, 25303756, 25305405, 25307368 25309116, 25313154, 25313411, 25316758, 25317989, 25320555, 25323525 25328518, 25329664, 25335249, 25335360, 25335790, 25337332, 25337640 25348956, 25353983, 25356118, 25357142, 25362958,

25367588, 25367721 25382812, 25383204, 25384462, 25386748, 25388896, 25392535, 25395696, 25397936, 25398306, 25404202, 25405100, 25405687, 25405813, 25410017 25410180, 25410802, 25410877, 25411036, 25417050, 25417056, 25417958 25425451, 25425760, 25427662, 25429959, 25430120, 25433696, 25435038 25437699, 25440818, 25442559, 25444961, 25445168, 25451531, 25455795 25457409, 25459958, 25462714, 25463844, 25472112, 25476149, 25477657 25478885, 25479164, 25489342, 25489367, 25489607, 25492379, 25498930 25498994, 25516250, 25524955, 25528838, 25530080, 25530814, 25535668 25536819, 25537470, 25539063, 25540738, 25546580, 25546608, 25547901 25551676, 25553616, 25554787, 25555252, 25557886, 25558986, 25560487 25561296, 25569149, 25570929, 25575348, 25575628, 25579458, 25579761 25594901, 25597525, 25598473, 25599425, 25600342, 25600421, 25602488 25603923, 25606091, 25607726, 25612095, 25614866, 25616268, 25616359 25616417, 25616645, 25631933, 25633101, 25634317, 25634348, 25635149 25638456, 25639019, 25643818, 25643931, 25646373, 25647325, 25648731 25653109, 25654459, 25654936, 25655390, 25655966, 25659655, 25660847 25661819, 25662088, 25662101, 25662524, 25663488, 25669791, 25670786 25671354, 25672640, 25674386, 25680221, 25685152, 25686739, 25687460 25691904, 25694206, 25695903, 25700654, 25710420, 25715167, 25717371 25722055, 25722608, 25722720, 25728085, 25729507, 25734963, 25736747 25739065, 25743479, 25754606, 25757748, 25760195, 25762221, 25764020 25766822, 25768681, 25772669, 25774077, 25775213, 25780343, 25783447 25784002, 25785331, 25785441, 25788879, 25789041, 25789277, 25789579 25790353, 25795865, 25797092, 25797124, 25797305, 25800464, 25802510 25803545, 25807997, 25810704, 25813931, 25818707, 25822410, 25823754 25825910, 25826740, 25830492, 25832935, 25834581, 25838361, 25852885 25856821, 25858672, 25861398, 25865785, 25870579, 25871177, 25871639 25871753, 25872127, 25872389, 25874050, 25874678, 25882264, 25883438 25885148, 25888073, 25890056, 25890673, 25894239, 25895224, 25897615 25904273, 25904490, 25906117, 25908728, 25911724, 25914276, 25919622 25932524, 25932728, 25933494, 25941836, 25943271, 25945130, 25947799 25953857, 25954022, 25954054, 25957038, 25963024, 25964954, 25967544 25967985, 25970731, 25971286, 25973152, 25975723, 25977302, 25980605 25980770, 25981498, 25982666, 25986062, 25990907, 25995938, 26006257 26007010, 26019148, 26024732, 26025681, 26029075, 26029777, 26029780 26032573, 26034119, 26036748, 26037215, 26038086, 26039623, 26040483 26045732, 26051656, 26078437, 26078493, 26080410, 26083298, 26088426 26088836, 26090767, 26091640, 26091786, 26095327, 26095405, 26096382 26108080, 26108337, 26110259, 26110632, 26111842, 26112621, 26115103 26121990, 26124078, 26137367, 26138085, 26149904, 26153977, 26169341 26169345, 26170715, 26176002, 26187943, 26189861, 26198757, 26198926 26201113, 26203182, 26223039, 26237431, 26237773, 26238195, 26242031 26243698, 26244115, 26245237, 26248143, 26249718, 26256131, 26259265 26261327, 26263328, 26263721, 26269790, 26271001, 26277439, 26285062 26285933, 26308650, 26309047, 26318627, 26323308, 26324206, 26324769 26327624, 26330994, 26331743, 26333141, 26336977, 26338953, 26351334 26353617, 26358670, 26359091, 26362821, 26366517, 26367012, 26367460 26371725, 26374791, 26375250, 26380097, 26385189, 26388538, 26396790 26399626, 26399691, 26406387, 26407408, 26412540, 26418088, 26420561 26421667, 26422277, 26423085, 26426526, 26426967, 26430323, 26430737 26434436, 26434999, 26435073, 26436168, 26438612, 26439748, 26440169 26440749, 26442308, 26444601, 26444887, 26446098, 26452606, 26474703 26475419, 26476244, 26478970, 26479173, 26486365, 26492866, 26493289 26498354, 26513709, 26521043, 26522439, 26523432, 26526726, 26526799 26536320, 26537307, 26542135, 26544823, 26545688, 26546070, 26546664, 26546754, 26548363, 26556014, 26569225, 26575788, 26580633, 26582460 26584641, 26597140, 26599395, 26608137, 26608238, 26609942, 26615291 26615690, 26617804, 26623652, 26626879, 26629381, 26633355, 26633558 26635897, 26637273, 26637824, 26639167, 26641610, 26650226, 26654363 26658759, 26659182, 26680105, 26712331, 26714486, 26714910, 26717528 26725687, 26727397, 26729494, 26729611, 26740700, 26744595, 26745002 26751106, 26751171, 26755171, 26758193, 26764561, 26765212, 26775602 26784509, 26794786, 26797591, 26798516, 26802503, 26816582, 26820076 26822620, 26828994, 26829845, 26833932, 26837702, 26840654, 26844406 26844870, 26849779, 26875822, 26883456, 26895149, 26896659, 26898563 26907327, 26908788, 26909100, 26909504, 26911000, 26939314, 26944190 26958896, 26963310, 26966916, 26967713, 26969321, 26970175, 26970717 26981902, 26983259, 26986173, 26992964, 27006120, 27006664, 27009164 27013146, 27028251, 27034890, 27037839, 27038986, 27039712, 27044297 27052607, 27058530, 27060167, 27060859, 27061736, 27073314, 27079140 27087426, 27090765, 27093423, 27106179, 27110878, 27115422, 27117822 27119621, 27122162, 27124624, 27125872, 27133662, 27135647, 27135993 27138325, 27142373, 27142529, 27151826, 27153641, 27161071, 27162390

27162405, 27163928, 27165231, 27169796, 27170305, 27181537, 27181897 27199245, 27200959, 27207110, 27208795, 27213224, 27216046, 27223075 27229389, 27231051, 27236722, 27242226, 27244337, 27248917, 27250547 27251690, 27255377, 27256000, 27258578, 27259307, 27262945, 27266245 27274456, 27274536, 27275533, 27276231, 27283960, 27284499, 27285244 27288638, 27292213, 27293599, 27302711, 27302730, 27304410, 27305039 27314206, 27314390, 27320576, 27321179, 27329612, 27333106, 27334316 27338912, 27338946, 27339115, 27345231, 27346709, 27348081, 27349393 27350267, 27351628, 27359178, 27364854, 27367194, 27370965, 27375542 27381498, 27386467, 27394703, 27395416, 27396624, 27396672, 27396813 27397048, 27400416, 27400598, 27404668, 27405645, 27416997, 27423251 27424405, 27426363, 27432062, 27433870, 27434193, 27439835, 27441326 27442041, 27445727, 27457891, 27466597, 27493674, 27501373, 27501413 27502420, 27504770, 27505229, 27508985, 27510959, 27525909, 27534509 27540613, 27544973, 27548131, 27554074, 27555481, 27558861, 27560602 27562488, 27565906, 27567477, 27593501, 27595973, 27607563, 27611612 27613080, 27613530, 27617978, 27620808, 27634676, 27634991, 27658186 27666312, 27671633, 27680669, 27686599, 27687880, 27688036, 27688099 27688692, 27691920, 27691939, 27693713, 27695063, 27698953, 27700466 27704237, 27709046, 27726780, 27740424, 27748954, 27751755, 27757567 27757888, 27758972, 27759077, 27793533, 27799032, 27801337, 27824543 27825241, 27829295, 27833672, 27834551, 27835925, 27847259, 27855490 27882176, 27898015, 27931299, 27940876, 27945870, 27951817, 27959048 27966472, 27986817, 27994325, 27995215, 27995248, 27997875, 27998003 27999638, 28000269, 28022101, 28023081, 28023482, 28033429, 28040776 28045903, 28067846, 28072567, 28074713, 28090453, 28099662, 28108003 28111583, 28120036, 28120951, 28140658, 28171079, 28174827, 28180464 28184554, 28188330, 28194173, 28199085, 28218832, 28220398, 28223871 28226179, 28229360, 28279837, 28282606, 28290434, 28302049, 28305001 28305362, 28320399, 28354603, 28361221, 28365111, 28378446, 28390273 28396445, 28420042, 28420457, 28432129, 28434028, 28437315, 28454242 28468312, 28483184, 28489150, 28502343, 28507324, 28508053, 28508557 28522441, 28528349, 28535272, 28545134, 28571483, 28578164, 28587723 28600233, 28617959, 28621470, 28627255, 28636676, 28691965, 28708023 28728272, 28749289, 28803345, 28849751, 28852691, 28856060, 28891984 28951382, 28960211, 28987439, 28991884, 28993590, 29027694, 29189889 29250230

Version 12.2.0.1.ru-2019-01.rur-2019-01.r1

Version 12.2.0.1.ru-2019-01.rur-2019-01.r1 adds support for the following:

- Patch 28822515: Database Jan 2019 Release Update : 12.2.0.1.190115
- Patch 28790651: OJVM RELEASE UPDATE: 12.2.0.1.190115
- Patch 28125601: DSTv32 for RDBMS (TZDATA2018E)
- Patch 28127287: DSTv32 for OJVM (TZDATA2018E)
- PreUpgrade Jar: preupgrade_12201_cbuild_19_lf.zip

Oracle Release Update 12.2.0.1.190115, Released January 2019

Bugs fixed:

26285062, 8480838, 8932139, 13554903, 14221306, 14690846, 15931756 16002385, 16438495, 16727454, 16942578, 17027695, 17533661, 17947871 18308268, 18521691, 18594510, 18774543, 19072655, 19211433, 19285025 19327292, 19526548, 19614243, 19647894, 19649997, 19721304, 20003668 20087519, 20118035, 20120236, 20324049, 20436508, 20532077, 20591151 20620169, 20736227, 20756305, 20866970, 20976443, 21095391, 21143725 21147908, 21159907, 21178363, 21186167, 21197098, 21216226, 21320338 21433452, 21479706, 21520266, 21547051, 21744603, 21882528, 21981529 21985256, 22007324, 22070853, 22072543, 22087683, 22104866, 22179537, 22347493, 22364044, 22367053, 22379010, 22446455, 22495673, 22503283 22503297, 22504793, 22530986, 22564336, 22568728, 22581771, 22594071 22599050, 22628825, 22645009, 22654475, 22700845, 22729345, 22826067 22843979, 22845846, 22864303, 22898198, 22950945, 22970869, 22981722 23019710, 23026585, 23035249, 23055900, 23061453, 23065002, 23066146 23080557, 23104033, 23105538, 23110523, 23125560, 23126545, 23127945 23151677, 23179662, 23184263,

23197730, 23234232, 23249829, 23271203 23300142, 23310101, 23312077, 23481673, 23491861, 23499004, 23499160 23521523, 23527363, 23533647, 23548817, 23567857, 23572982, 23581777 23588722, 23599216, 23600861, 23602213, 23645516, 23665623, 23709062 23715460, 23730961, 23733981, 23735292, 23741944, 23746128, 23749454 23761724, 24010030, 24289874, 24294174, 24303148, 24307571, 24308349 24326444, 24326846, 24332831, 24334708, 24336249, 24337882, 24341675 24343905, 24345420, 24346821, 24348685, 24350620, 24352981, 24368004 24371491, 24373756, 24374976, 24376875, 24376878, 24385983, 24401351 24403922, 24415926, 24421668, 24423416, 24425056, 24425998, 24435982 24437162, 24443539, 24457597, 24460392, 24461826, 24467122, 24468470 24470606, 24471473, 24473736, 24485034, 24485161, 24485174, 24486059 24486237, 24509056, 24534401, 24554533, 24555417, 24556967, 24560906 24563422, 24570214, 24570598, 24573817, 24578718, 24578797, 24589081 24589590, 24593740, 24595699, 24600330, 24609592, 24609996, 24616637 24617969, 24623975, 24624166, 24642495, 24654629, 24655717, 24664211 24668398, 24669189, 24674197, 24674955, 24676172, 24677696, 24680959 24689376, 24692973, 24693290, 24697323, 24699619, 24710696, 24713381 24714096, 24717183, 24717859, 24718260, 24719799, 24735430, 24737064 24737403, 24737581, 24744383, 24744686, 24757934, 24759556, 24760407 24764085, 24766309, 24786669, 24792678, 24793511, 24796092, 24797119 24800423, 24801152, 24802934, 24811725, 24812047, 24827228, 24827654 24831514, 24835919, 24843188, 24844549, 24845157, 24848746, 24848923 24850622, 24907917, 24908321, 24911709, 24912588, 24922704, 24923080 24923215, 24923338, 24923790, 24929210, 24938784, 24940060, 24942749 24953434, 24957555, 24960044, 24966594, 24966788, 24968162, 24976007 24978100, 25022574, 25027852, 25028996, 25029022, 25029423, 25034396 25036474, 25044977, 25045228, 25050160, 25051628, 25057811, 25058080 25062592, 25063971, 25065563, 25072986, 25078611, 25086233, 25087436 25093872, 25098160, 25099339, 25099497, 25099758, 25100063, 25100579 25103996, 25107662, 25110233, 25114561, 25120284, 25120668, 25120742 25121089, 25123585, 25124363, 25129925, 25140197, 25145163, 25145215 25150925, 25159176, 25162645, 25164293, 25166187, 25171084, 25175723 25176408, 25178032, 25178101, 25178179, 25179774, 25182817, 25184555 25186079, 25191872, 25192044, 25192729, 25199585, 25201454, 25202355 25203656, 25206864, 25207410, 25209912, 25210268, 25210499, 25211628 25223839, 25224242, 25225795, 25226665, 25227381, 25230870, 25230945 25237577, 25240188, 25240590, 25241448, 25241625, 25244807, 25248384 25251648, 25257085, 25259611, 25262869, 25263960, 25265499, 25283790 25287072, 25296876, 25299227, 25299807, 25300427, 25305405, 25307368 25309116, 25313154, 25313411, 25316758, 25317989, 25320555, 25323525 25328518, 25329664, 25335249, 25335360, 25335790, 25337332, 25337640 25348956, 25353983, 25356118, 25357142, 25362958, 25382812, 25383204 25384462, 25386748, 25388896, 25392535, 25395696, 25397936, 25405813 25410017, 25410180, 25410802, 25410877, 25411036, 25417050, 25417056 25417958, 25425451, 25425760, 25427662, 25429959, 25430120, 25433696 25435038, 25437699, 25440818, 25444961, 25445168, 25451531, 25455795 25457409, 25459958, 25462714, 25463844, 25472112, 25476149, 25478885 25479164, 25489342, 25489367, 25489607, 25492379, 25498930, 25498994 25516250, 25524955, 25528838, 25530080, 25530814, 25535668, 25536819 25537470, 25539063, 25540738, 25546580, 25546608, 25547901, 25551676 25553616, 25554787, 25555252, 25557886, 25558986, 25560487, 25561296 25569149, 25570929, 25575348, 25575628, 25579458, 25579761, 25594901 25597525, 25598473, 25599425, 25600342, 25600421, 25602488, 25603923 25606091, 25607726, 25612095, 25614866, 25616268, 25616359, 25616417 25616645, 25631933, 25633101, 25634317, 25634348, 25635149, 25638456 25639019, 25643818, 25643931, 25646373, 25647325, 25648731, 25653109 25654459, 25654936, 25655390, 25655966, 25659655, 25660847, 25661819 25662088, 25662101, 25662524, 25669791, 25670786, 25671354, 25672640 25674386, 25680221, 25685152, 25686739, 25687460, 25691904, 25694206 25695903, 25700654, 25710420, 25715167, 25717371, 25722055, 25722608 25722720, 25728085, 25729507, 25734963, 25736747, 25739065, 25754606 25757748, 25760195, 25762221, 25764020, 25766822, 25768681, 25772669 25774077, 25775213, 25780343, 25784002, 25785331, 25785441, 25788879 25789041, 25789277, 25789579, 25790353, 25797092, 25797124, 25797305 25800464, 25802510, 25803545, 25807997, 25810704, 25813931, 25818707 25822410, 25823754, 25825910, 25826740, 25830492, 25832935, 25834581 25838361, 25852885, 25856821, 25858672, 25861398, 25865785, 25870579 25871177, 25871639, 25871753, 25872127, 25872389, 25874050, 25874678 25882264, 25885148, 25888073, 25890056, 25890673, 25894239, 25895224 25897615, 25904273, 25904490, 25906117, 25911724, 25914276, 25919622 25932524, 25932728, 25933494, 25941836, 25943271, 25945130, 25947799 25953857, 25954022, 25954054, 25957038, 25963024, 25964954, 25967544 25967985, 25970731, 25973152, 25975723, 25977302,

25980605, 25980770 25981498, 25982666, 25990907, 25995938, 26006257, 26007010, 26019148, 26024732, 26025681, 26029780, 26032573, 26034119, 26036748, 26037215 26038086, 26039623, 26040483, 26045732, 26078437, 26078493, 26080410 26083298, 26088426, 26088836, 26090767, 26091640, 26091786, 26095327 26095405, 26096382, 26108080, 26108337, 26110632, 26111842, 26112621 26115103, 26121990, 26124078, 26137367, 26138085, 26149904, 26153977 26169341, 26169345, 26170715, 26176002, 26187943, 26189861, 26198757 26198926, 26201113, 26203182, 26223039, 26237431, 26237773, 26238195 26242031, 26243698, 26244115, 26245237, 26248143, 26249718, 26256131 26259265, 26261327, 26263328, 26263721, 26269790, 26271001, 26277439 26285933, 26308650, 26309047, 26318627, 26323308, 26324769, 26327624 26330994, 26331743, 26333141, 26336977, 26338953, 26351334, 26353617 26358670, 26359091, 26362821, 26366517, 26367012, 26371725, 26374791 26375250, 26380097, 26385189, 26388538, 26396790, 26399626, 26399691 26406387, 26407408, 26412540, 26418088, 26420561, 26421667, 26422277 26423085, 26426526, 26426967, 26430737, 26434436, 26434999, 26435073 26436168, 26438612, 26440749, 26442308, 26444601, 26444887, 26446098 26452606, 26474703, 26475419, 26476244, 26478970, 26479173, 26486365 26492866, 26493289, 26498354, 26513709, 26521043, 26522439, 26523432 26526726, 26536320, 26537307, 26542135, 26544823, 26545688, 26546070 26546664, 26546754, 26548363, 26556014, 26569225, 26575788, 26580633 26582460, 26584641, 26597140, 26599395, 26608137, 26609942, 26615291 26615690, 26623652, 26626879, 26629381, 26633355, 26633558, 26635897 26637273, 26637824, 26639167, 26641610, 26650226, 26658759, 26659182 26680105, 26712331, 26714486, 26714910, 26717528, 26725687, 26727397 26729494, 26729611, 26740700, 26744595, 26745002, 26751106, 26751171 26755171, 26758193, 26764561, 26765212, 26775602, 26784509, 26794786 26797591, 26798516, 26802503, 26820076, 26822620, 26828994, 26840654 26844406, 26844870, 26849779, 26875822, 26883456, 26895149, 26896659 26898563, 26907327, 26908788, 26909100, 26909504, 26911000, 26939314 26944190, 26958896, 26963310, 26966916, 26967713, 26969321, 26970717 26981902, 26983259, 26986173, 26992964, 27006664, 27009164, 27013146 27028251, 27034890, 27037839, 27038986, 27039712, 27044297, 27052607 27058530, 27060167, 27060859, 27073314, 27079140, 27087426, 27090765 27093423, 27110878, 27117822, 27119621, 27122162, 27124624, 27125872 27133662, 27135647, 27135993, 27138325, 27142373, 27142529, 27151826 27153641, 27161071, 27162405, 27163928, 27165231, 27169796, 27170305 27181537, 27181897, 27199245, 27200959, 27207110, 27213224, 27223075 27229389, 27236722, 27244337, 27248917, 27250547, 27251690, 27255377 27256000, 27258578, 27259307, 27262945, 27266245, 27274456, 27274536 27276231, 27283960, 27285244, 27292213, 27293599, 27302711, 27302730 27304410, 27305039, 27314206, 27314390, 27320576, 27321179, 27329612 27333106, 27334316, 27338912, 27338946, 27339115, 27345231, 27346709 27348081, 27349393, 27351628, 27359178, 27367194, 27370965, 27375542 27381498, 27386467, 27394703, 27395416, 27396624, 27396813, 27397048 27400416, 27400598, 27404668, 27405645, 27416997, 27433870, 27434193 27439835, 27441326, 27442041, 27445727, 27466597, 27493674, 27501373 27501413, 27502420, 27504770, 27505229, 27508985, 27510959, 27525909 27534509, 27540613, 27544973, 27548131, 27555481, 27558861, 27560602 27567477, 27593501, 27595973, 27607563, 27611612, 27613080, 27613530 27617978, 27620808, 27634676, 27658186, 27687880, 27688036, 27688099 27688692, 27691920, 27691939, 27693713, 27695063, 27698953, 27700466 27709046, 27726780, 27740424, 27748954, 27757888, 27759077, 27793533 27799032, 27801337, 27835925, 27847259, 27855490, 27882176, 27898015 27931299, 27940876, 27945870, 27951817, 27959048, 27994325, 27995248 27997875, 27998003, 27999638, 28000269, 28022101, 28023081, 28023482 28033429, 28040776, 28067846, 28072567, 28074713, 28090453, 28099662 28111583, 28120951, 28140658, 28171079, 28174827, 28180464, 28184554 28188330, 28194173, 28218832, 28220398, 28226179, 28229360, 28282606 28290434, 28302049, 28305001, 28305362, 28320399, 28354603, 28361221 28365111, 28390273, 28396445, 28420042, 28420457, 28434028, 28437315 28454242, 28483184, 28489150, 28502343, 28508053, 28508557, 28522441 28528349, 28535272, 28571483, 28617959, 28621470, 28627255, 28636676 28691965, 28708023, 28728272, 28749289, 28960211, 28993590

Version 12.2.0.1.ru-2018-10.rur-2018-10.r1

Version 12.2.0.1.ru-2018-10.rur-2018-10.r1 adds support for the following:

- October 2018 Release Update: 12.2.0.1.181016 (28662603)

Oracle Release Update 12.2.0.1.181016, Released October 2018

Bugs fixed:

28390273, 28571483, 28483184, 8480838, 13554903, 14221306, 14690846 15931756, 16002385, 16438495, 16727454, 16942578, 17027695, 17533661 17947871, 18308268, 18521691, 18594510, 18774543, 19072655, 19211433 19285025, 19327292, 19526548, 19614243, 19647894, 19649997, 19721304 20003668, 20087519, 20118035, 20120236, 20324049, 20436508, 20532077 20591151, 20620169, 20736227, 20756305, 20866970, 20976443, 21143725 21147908, 21159907, 21178363, 21186167, 21216226, 21320338, 21433452 21479706, 21520266, 21547051, 21744603, 21882528, 21981529, 21985256 22007324, 22070853, 22072543, 22087683, 22104866, 22179537, 22347493 22364044, 22367053, 22379010, 22446455, 22495673, 22503283, 22503297 22504793, 22530986, 22564336, 22568728, 22581771, 22594071, 22599050 22628825, 22645009, 22654475, 22700845, 22729345, 22826067, 22843979 22845846, 22864303, 22898198, 22950945, 22970869, 22981722, 23019710 23026585, 23035249, 23055900, 23061453, 23065002, 23066146, 23080557 23105538, 23110523, 23125560, 23126545, 23127945, 23151677, 23179662 23184263, 23197730, 23234232, 23249829, 23271203, 23300142, 23310101 23312077, 23481673, 23491861, 23499160, 23521523, 23527363, 23533647 23548817, 23567857, 23572982, 23581777, 23588722, 23599216, 23600861 23602213, 23645516, 23665623, 23709062, 23715460, 23730961, 23733981 23735292, 23741944, 23746128, 23749454, 24010030, 24289874, 24294174 24303148, 24307571, 24308349, 24326444, 24326846, 24332831, 24334708 24336249, 24337882, 24341675, 24343905, 24345420, 24346821, 24348685 24350620, 24368004, 24371491, 24373756, 24374976, 24376875, 24376878 24385983, 24401351, 24403922, 24415926, 24421668, 24423416, 24425056 24425998, 24435982, 24437162, 24443539, 24457597, 24461826, 24467122 24468470, 24470606, 24473736, 24485034, 24485161, 24485174, 24486059 24486237, 24509056, 24534401, 24554533, 24555417, 24556967, 24560906 24563422, 24570598, 24573817, 24578718, 24578797, 24589081, 24589590 24593740, 24595699, 24600330, 24609592, 24609996, 24616637, 24617969 24623975, 24624166, 24642495, 24654629, 24655717, 24664211, 24668398 24674197, 24674955, 24676172, 24677696, 24680959, 24689376, 24692973 24693290, 24699619, 24710696, 24713381, 24714096, 24717183, 24717859 24718260, 24719799, 24735430, 24737064, 24737403, 24737581, 24744383 24744686, 24757934, 24759556, 24760407, 24766309, 24786669, 24792678 24793511, 24796092, 24797119, 24800423, 24801152, 24802934, 24811725 24812047, 24827228, 24827654, 24831514, 24835919, 24843188, 24844549 24845157, 24848746, 24848923, 24850622, 24907917, 24908321, 24911709 24912588, 24922704, 24923080, 24923215, 24923338, 24923790, 24929210 24938784, 24940060, 24942749, 24953434, 24957555, 24960044, 24966594 24966788, 24968162, 24976007, 24978100, 25027852, 25029022, 25029423 25034396, 25036474, 25044977, 25045228, 25050160, 25051628, 25057811 25058080, 25062592, 25063971, 25065563, 25072986, 25078611, 25086233 25087436, 25093872, 25098160, 25099339, 25099497, 25099758, 25100063 25100579, 25103996, 25107662, 25110233, 25120284, 25120742, 25121089 25123585, 25124363, 25129925, 25140197, 25145163, 25145215, 25150925 25159176, 25162645, 25164293, 25166187, 25171084, 25175723, 25176408 25178032, 25178101, 25178179, 25179774, 25182817, 25184555, 25186079 25191872, 25192044, 25192729, 25199585, 25201454, 25202355, 25203656 25206864, 25207410, 25209912, 25210268, 25210499, 25211628, 25223839 25224242, 25225795, 25226665, 25227381, 25230945, 25237577, 25240590 25241448, 25241625, 25244807, 25248384, 25251648, 25257085, 25259611 25262869, 25263960, 25265499, 25283790, 25287072, 25296876, 25299227 25299807, 25300427, 25305405, 25307368, 25309116, 25313154, 25313411 25316758, 25317989, 25320555, 25323525, 25328518, 25329664, 25335249 25335360, 25335790, 25337332, 25337640, 25348956, 25353983, 25357142 25362958, 25382812, 25383204, 25384462, 25386748, 25388896, 25392535 25395696, 25397936, 25405813, 25410017, 25410180, 25410802, 25410877 25411036, 25417050, 25417056, 25417958, 25425451, 25425760, 25427662 25429959, 25430120, 25433696, 25435038, 25437699, 25440818, 25444961 25451531, 25455795, 25457409, 25459958, 25462714, 25463844, 25472112 25476149, 25478885, 25479164, 25489342, 25489367, 25489607, 25492379 25498930, 25498994, 25516250, 25524955, 25528838, 25530080, 25530814 25535668, 25536819, 25537470, 25539063, 25540738, 25546580, 25546608 25547901, 25551676, 25553616, 25554787, 25555252, 25557886, 25558986 25560487, 25561296, 25569149, 25570929, 25575348, 25575628, 25579458 25579761, 25594901, 25597525, 25598473, 25600342, 25600421, 25602488 25603923, 25606091, 25607726, 25612095, 25614866, 25616268, 25616359 25616417, 25616645, 25631933, 25633101, 25634317, 25634348, 25635149 25638456,

25639019, 25643818, 25643931, 25646373, 25647325, 25648731 25653109, 25654459, 25654936, 25655390, 25655966, 25659655, 25660847 25661819, 25662088, 25662101, 25662524, 25669791, 25670786, 25672640 25674386, 25680221, 25685152, 25686739, 25687460, 25691904, 25694206, 25695903, 25700654, 25710420, 25715167, 25717371, 25722055, 25722608 25722720, 25728085, 25729507, 25734963, 25736747, 25739065, 25754606 25757748, 25760195, 25762221, 25764020, 25766822, 25768681, 25772669 25774077, 25775213, 25780343, 25784002, 25785331, 25785441, 25788879 25789041, 25789277, 25789579, 25790353, 25797092, 25797124, 25797305 25800464, 25803545, 25807997, 25810704, 25813931, 25818707, 25822410 25823754, 25825910, 25826740, 25830492, 25832935, 25834581, 25838361 25852885, 25856821, 25858672, 25861398, 25865785, 25870579, 25871177 25871639, 25871753, 25872127, 25872389, 25874050, 25874678, 25882264, 25885148, 25888073, 25890056, 25890673, 25894239, 25895224, 25897615 25904273, 25904490, 25906117, 25911724, 25914276, 25919622, 25932524 25932728, 25933494, 25941836, 25943271, 25945130, 25947799, 25953857 25954022, 25954054, 25957038, 25963024, 25964954, 25967544, 25967985 25970731, 25973152, 25975723, 25977302, 25980605, 25980770, 25981498 25982666, 25990907, 25995938, 26006257, 26007010, 26019148, 26024732 26025681, 26029780, 26032573, 26036748, 26037215, 26038086, 26039623 26040483, 26045732, 26078437, 26078493, 26080410, 26083298, 26088426 26088836, 26090767, 26091640, 26091786, 26095327, 26095405, 26096382, 26108080, 26110632, 26111842, 26121990, 26137367, 26138085, 26149904 26153977, 26169341, 26169345, 26170715, 26176002, 26187943, 26189861 26198757, 26198926, 26201113, 26223039, 26237431, 26237773, 26242031 26243698, 26244115, 26245237, 26249718, 26256131, 26259265, 26261327 26263328, 26263721, 26269790, 26271001, 26277439, 26285933, 26308650 26309047, 26318627, 26323308, 26324769, 26327624, 26330994, 26331743 26333141, 26338953, 26351334, 26353617, 26358670, 26362821, 26366517 26367012, 26374791, 26375250, 26380097, 26385189, 26388538, 26396790 26399626, 26399691, 26406387, 26412540, 26418088, 26420561, 26421667, 26422277, 26426526, 26430737, 26434999, 26435073, 26436168, 26438612 26440749, 26442308, 26444601, 26444887, 26446098, 26452606, 26475419 26476244, 26478970, 26479173, 26486365, 26492866, 26493289, 26498354 26513709, 26521043, 26522439, 26523432, 26526726, 26536320, 26537307 26542135, 26544823, 26545688, 26546070, 26546664, 26546754, 26548363 26556014, 26569225, 26575788, 26582460, 26584641, 26597140, 26599395 26608137, 26609942, 26615291, 26615690, 26623652, 26626879, 26629381 26633355, 26633558, 26635897, 26637273, 26637824, 26639167, 26641610 26650226, 26658759, 26659182, 26680105, 26712331, 26714910, 26717528, 26727397, 26729494, 26729611, 26740700, 26744595, 26751106, 26751171 26758193, 26764561, 26765212, 26775602, 26784509, 26794786, 26797591 26802503, 26820076, 26822620, 26828994, 26840654, 26844870, 26849779 26875822, 26883456, 26896659, 26898563, 26907327, 26908788, 26909100 26909504, 26911000, 26939314, 26944190, 26963310, 26966916, 26967713 26969321, 26970717, 26981902, 26983259, 26986173, 26992964, 27006664 27009164, 27013146, 27028251, 27034890, 27038986, 27039712, 27044297 27052607, 27060167, 27060859, 27073314, 27079140, 27087426, 27090765 27093423, 27110878, 27117822, 27119621, 27124624, 27125872, 27133662, 27135647, 27135993, 27138325, 27142373, 27153641, 27161071, 27162405 27163928, 27165231, 27169796, 27170305, 27181537, 27199245, 27207110 27213224, 27229389, 27244337, 27248917, 27250547, 27251690, 27255377 27256000, 27259307, 27262945, 27274536, 27276231, 27285244, 27292213 27293599, 27302711, 27302730, 27304410, 27305039, 27314206, 27314390 27321179, 27329612, 27333106, 27334316, 27338912, 27338946, 27339115 27345231, 27346709, 27348081, 27349393, 27351628, 27359178, 27367194 27370965, 27375542, 27394703, 27395416, 27396624, 27396813, 27400598 27405645, 27416997, 27433870, 27434193, 27439835, 27441326, 27442041, 27466597, 27493674, 27501373, 27501413, 27502420, 27504770, 27505229 27508985, 27510959, 27534509, 27544973, 27548131, 27555481, 27558861 27560602, 27567477, 27595973, 27607563, 27611612, 27613080, 27620808 27687880, 27688036, 27688692, 27691920, 27691939, 27698953, 27700466 27709046, 27726780, 27740424, 27748954, 27757888, 27799032, 27835925 27847259, 27882176, 27898015, 27940876, 27945870, 27951817, 27959048 27994325, 27997875, 27998003, 28000269, 28033429, 28040776, 28074713 28090453, 28099662, 28140658, 28171079, 28174827, 28184554, 28188330 28218832, 28226179, 28290434, 28305001, 28320399, 28354603, 28437315, 28454242, 28508557, 28522441

Database Engine: 12.1.0.2

The following versions are available for database engine 12.1.0.2:

- [Version 12.1.0.2.v18 \(p. 1196\)](#)
- [Version 12.1.0.2.v17 \(p. 1199\)](#)
- [Version 12.1.0.2.v16 \(p. 1201\)](#)
- [Version 12.1.0.2.v15 \(p. 1204\)](#)
- [Version 12.1.0.2.v14 \(p. 1206\)](#)
- [Version 12.1.0.2.v13 \(p. 1208\)](#)
- [Version 12.1.0.2.v12 \(p. 1210\)](#)
- [Version 12.1.0.2.v11 \(p. 1212\)](#)
- [Version 12.1.0.2.v10 \(p. 1213\)](#)
- [Version 12.1.0.2.v9 \(p. 1215\)](#)
- [Version 12.1.0.2.v8 \(p. 1216\)](#)
- [Version 12.1.0.2.v7 \(p. 1218\)](#)
- [Version 12.1.0.2.v6 \(p. 1219\)](#)
- [Version 12.1.0.2.v5 \(p. 1220\)](#)
- [Version 12.1.0.2.v4 \(p. 1222\)](#)
- [Version 12.1.0.2.v3 \(p. 1222\)](#)
- [Version 12.1.0.2.v2 \(p. 1223\)](#)
- [Version 12.1.0.2.v1 \(p. 1224\)](#)

Version 12.1.0.2.v18

Version 12.1.0.2.v18 adds support for the following:

- Patch 29918340: DATABASE PATCH SET UPDATE 12.1.0.2.191015
- Patch 30128197: OJVM PATCH SET UPDATE 12.1.0.2.191015
- Patch 29997937: DSTv34 for RDBMS (TZDATA2019G)
- Patch 29997959: DSTv34 for OJVM (TZDATA2019G)
- Patch 17969866: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 20394750: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 24835919: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 21171382: DBMS_STATS Patch
- Patch 21091901: ONLINE MOVE OF HASH OR REF PARTITION CAN LEAVE LOCAL INDEXES INCONSISTENT
- Patch 30370890: JSON bundle patch
- Patch 20033733: PART :IMC: HIT ORA 600 [KGL-HEAP-SIZE-EXCEEDED]
- Patch 28730253: SUPPORT NEW ERA REIWA FOR JAPANESE IMPERIAL CALENDAR
- PreUpgrade Jar: preupgrade_12.1.0.2.0_18_crlf.zip
- Adds support for [Resizing the Temporary Tablespace in a Read Replica \(p. 1104\)](#).

Oracle patch 29918340, released October 2019

Bugs fixed:

19309466, 19902195, 18250893, 25437699, 19383839, 19781326, 16756406 18456643, 26546664, 22364044, 29483723, 18845653, 18913440, 19915271 20172151, 18417036, 19516448, 24907917, 23713236, 24796092, 23140259 19243521, 19658708, 18272672, 21153266, 19174430, 18900107, 22243719 19548064, 26556014, 20493163, 20688221, 22346829, 21387964, 13542050 22072818, 22250006, 22734547, 22243983, 21623164, 19012119, 19932634 19869255, 22232606, 18681056,

23324000, 25427662, 22068305, 24589081 19439759, 19303936, 22916353, 24835538, 22353346, 19790243, 21106027 26444887, 23088803, 22529728, 26256131, 19134173, 24303148, 21101873 20447445, 27122162, 21188584, 19390567, 26513709, 25780343, 19769480 21097043, 21225209, 26245237, 20677396, 19284031, 19450314, 19016730 23514911, 22205263, 22517782, 20919320, 22075064, 22551446, 29250317 22721409, 18440095, 22496904, 16439813, 18354830, 20596234, 22022760 20936905, 22268833, 23197103, 22515353, 23746128, 21514877, 19809171 21186167, 26111842, 18990023, 13787015, 22492533, 22233505, 20173897 24624166, 17210525, 21260431, 29707896, 25579761, 20181030, 25056052 19370504, 21868720, 23068169, 19124589, 22690648, 21383171, 19402853 19888853, 25107334, 24341675, 17722075, 20882568, 25653109, 23026585 18604692, 20717081, 25546608, 20768076, 27370965, 19081128, 22173980 25722055, 23514710, 19178851, 20951038, 22168163, 25161298, 20569094 24308635, 19791377, 19050649, 20920911, 22475617, 19189525, 19469538 27052607, 20598042, 29633753, 22458049, 18988834, 21159665, 23302839 25307368, 25699321, 21924131, 17409174, 22729345, 22842151, 25051465 19238590, 16941434, 20387265, 29378913, 24397438, 20673810, 23108128 20356733, 22380919, 18436647, 23065323, 20825533, 19124336, 22294260 24790914, 20284155, 23571055, 6194865, 25539063, 17365043, 25914276 20952966, 22961508, 19176223, 21300341, 23237313, 18288842, 27223075 22353199, 200111515, 22083366, 27634991, 25670786, 21419850, 26898563 27986817, 19577410, 22495673, 26248143, 23294548, 23101501, 24737064 19931709, 25423453, 25547060, 23533807, 27726780, 24600330, 28384353 25600421, 18122373, 20043616, 23124895, 18856999, 21450666, 18893947 26633558, 20076781, 26029780, 21196809, 21354456, 22533631, 23725036 20464614, 19562381, 27375542, 24808595, 19189317, 25669791, 18307021 21917884, 19708632, 27213224, 25633101, 29006527, 20711718, 18973548 25982666, 19718981, 20684983, 23567857, 22826718, 25655390, 21773465 20250147, 20144019, 19197175, 26263721, 19597439, 21387128, 22007324 18818069, 21566639, 19180770, 19879746, 21785691, 20424183, 20539050 24285405, 21425496, 26544823, 19957298, 20322560, 22228324, 23172924 22520320, 27751755, 21575362, 25058080, 22365117, 22645009, 25165496 28950969, 18774543, 20124446, 21429602, 29189889, 26153977, 19371175 21863727, 18940497, 19074147, 22923409, 25489342, 21380789, 19154375 25417056, 19044962, 19532017, 19662635, 22374754, 20560611, 25654936 21492036, 18705806, 28420042, 19578247, 22024071, 22238921, 22809871 21184223, 19995869, 23089357, 19404068, 18921743, 19065677, 19018447 19018206, 18308268, 19777862, 29027694, 22223463, 19304354, 22519146 22214989, 19445860, 26654363, 27199245, 22977256, 20890311, 27445727 21142837, 20869721, 24555417, 22258530, 22179537, 21756699, 18801391 20217801, 18819908, 22760595, 19550902, 25483815, 23628685, 19543384 25482971, 23007241, 19593445, 21080143, 27351628, 20582405, 24966594 20031873, 25489367, 29828111, 18618122, 24737581, 21698350, 22501616 26784509, 24739928, 18966843, 19077215, 20704450, 19068970, 20543011 19023822, 24713381, 22836801, 20432873, 21756677, 20328248, 18674047 18849537, 20087383, 25459958, 20315311, 22897344, 29163567, 27534509 26768025, 25178179, 19308965, 18948177, 20764012, 27623159, 19468991 20868862, 21780146, 23315153, 20466628, 21756661, 20397490, 19706965 20302006, 24831514, 23240358, 22178855, 19032777, 20862087, 19329654 18974476, 20603378, 20859910, 19307662, 26203182, 21847223, 20281121 22568797, 19075256, 19076343, 28026866, 29511611, 18866977, 22808310 25635149, 20844426, 20904530, 20441797, 20296619, 18548246, 21442094 25079710, 24674955, 18840932, 18740837, 20294666, 21037923, 25602488 21517440, 22062517, 19180394, 27337759, 19174942, 20671094, 21889720 19347458, 19450116, 18411216, 20117253, 24386767, 20641666, 19931367 25264559, 19930276, 22092979, 25616268, 21625179, 20879709, 23003979 20165574, 28578164, 19272708, 19547370, 22624709, 23084507, 23184263 20228093, 21281532, 25093872, 19805359, 26324206, 19461270, 19434529 18799063, 20354900, 29388020, 20378086, 17008068, 21246723, 20831538 20424899, 20361671, 18674024, 19689979, 24411921, 19873610, 16619249 20562898, 21641414, 21091431, 19440586, 20001168, 22757364, 22175564 22499356, 20725343, 21241052, 19561643, 28199085, 20736227, 19399918 19195895, 20830459, 20017509, 18475439, 25790353, 21828126, 21665897 25555252, 20746251, 22568177, 25764020, 25612095, 25357142, 23096938 19067244, 18043064, 19941352, 21329301, 18885870, 26243698, 26187943 20324049, 19536415, 23709062, 28174827, 20446883, 27314206, 21299490 25313154, 18628388, 21744290, 18254023, 27072923, 25047724, 20591183 27847259, 20459944, 19185876, 18548433, 27207110, 22465352, 24385625 24326444, 20402832, 19627012, 22733141, 29200700, 20468401, 27441326 27620950, 16863642, 19639483, 19315691, 27567477, 21479753, 19174521 23177923, 20401975, 18306996, 18851894, 27034890, 21424824, 20581111 20318889, 20936731, 21060755, 25240188, 26828994, 22256560, 19188927 23328639, 27229389, 20766180, 20229001, 24570598,

25475853, 21172913 17655240, 29379978, 21266085, 19028800, 19035573, 19366375, 24523374, 25599425, 25034396, 19289642, 21502702, 21291274, 18007682, 23521523 20475845, 29408136, 22148226, 22528741, 25417958, 24652769, 26088426 19326908, 19597583, 29500257, 17414008, 23019710, 20897759, 26822620 22046677, 19663176, 20938170, 19891090, 24825843, 26318627, 21960504 24509056, 19054077, 21385422, 26262953, 22657942, 20428621, 21899588 23326313, 19723336, 19835133, 17532734, 17495022, 25300427, 19333670 21842017, 19285025, 21373473, 23260854, 23061453, 19687159, 14643995 22146062, 20977794, 20734332, 17551063, 16938780, 27548131, 21977392 28612674, 24461826, 19676012, 20588502, 23315889, 19520602, 23053606 19841800, 20245930, 19001359, 21476308, 26546754, 19393542, 23533524 21099555, 27995248, 25429959, 19141838, 19644859, 21915719, 19908836 21421886, 19358317, 19524158, 23548817, 25861398, 20803014, 23025340 19335438, 19058490, 23642282, 19207117, 18799993, 25919622, 26569225 25986062, 20835241, 24662775, 19475971, 18967382, 20347562, 20348653 19896336, 24812585, 20048359, 21896069, 20468490, 19524384, 25392535 21147908, 21695575, 20440930, 30295478, 25789277, 19171086, 24718260 17867700, 19791273, 27397048, 21241829, 19591608, 22707244, 18419520 22296366, 22654475, 18914624, 19571367, 28636676, 21522582, 19501299 29893132, 26007010, 19529868, 20425790, 19708342, 27997875, 16870214 18202441, 24415926, 18743542, 19001390, 21157728, 20657411, 19332396 22606521, 21875360, 25091141, 21821302, 28000269, 19149990, 20382309 22855193, 16777441, 19606174, 20848335, 28542455, 25495682, 19382851 20528052, 22762046, 24563422, 23125826, 22503297, 28993590, 25192729 23338911, 22730454, 19354794, 20757079, 19176326, 20298413, 19048007 22018363, 18849970, 21532755, 24300640, 20860659, 22905130, 26121990 21263635, 27710072, 23209741, 22160989, 18499088, 18775971, 22894949 21059919, 18952989, 27348081, 22518784, 25856821, 24457597, 25484507 20794034, 20554364, 21061354, 19468347, 17533661, 19883092, 20657441 24401351, 21285458, 28023399, 18051556, 25330273, 26412540, 24425998 19699191, 24437510, 16875041, 20669434, 18964978, 22972770, 23342170 20828947, 21373076, 25492379, 25551676, 14283239, 25766822, 21967197 22922076, 19601762, 25575628, 26110259, 20368850, 21239530, 20437153 24848928, 20880215, 20798891, 25606091, 19013183, 25042823, 21133343 22695831, 24365589, 25248384, 25634317, 20134113, 19587324, 20273319 28501075, 18542562, 19017309, 26758193, 21063322, 22062026, 20134339 22077517, 22815955, 23854396, 24690216, 22507210, 16354467, 20101006 21795111, 27938623, 23501901, 18797519, 25997810, 23029562, 25879984 21260397, 25029423, 26844406, 19354335, 19730508, 22366558, 19390620 26658759, 25822410, 6599380, 20717359, 24321547, 27097854, 21297872 18964939, 19871910, 29437712, 26366517, 21913183, 25695903, 22366322 20171986, 20603431, 21132297, 25957038, 21542577, 22507234, 23170620 24719736, 25600342, 18868646, 28587723, 29142109, 26637824, 20627866 18110491, 16923858, 24642295, 19518079, 19339555, 20466322, 25823754 25110233, 20169408, 24908321, 20842388, 17274537, 26575788, 20474192 21644640, 28849751, 21794615, 18899974, 20471920, 22806698, 19052488 26198757, 19503821, 23717151, 24350620, 23126410, 20074391, 19157754 22495062, 21220620, 24316947, 19865345, 19065556, 22816287, 25947799 20878790, 23492665, 21322887, 22305887, 20879889, 19617921, 24350831 19578350, 28022101, 26439748, 21893235, 19363645, 21072646, 20898391 19291380, 27060167, 18382302, 27086138, 22536802, 22087683, 21197626 21656630, 20373598, 19248799, 22707866, 28432129, 19155797, 19279273 18886413, 25490238, 20922010, 19990037, 25150925, 20509482, 20778986 27255377, 24717859, 20703000, 22862134, 21526048, 28683167, 24929210 24560906, 20144308, 21620471, 19670108, 19068610, 20267166, 25123585 20476175, 18549238, 19297917, 22950945, 19385656, 20564072, 23528412 19684504, 19330795, 21174504, 28357401, 20899461, 20557786, 21911701 19143550, 20118035, 19024808, 25760195, 20009833, 19604659, 16359751 26039623, 22820579, 19928926, 23314180, 20212067, 24737403, 20480209 18904062, 29030780, 26430737, 20856766, 17258582, 27169796, 21668627 23272045, 20877664, 19487147, 23149541, 24577566, 19430401, 19676905 20925795, 26482376, 21296029, 21629064, 23229229, 22865673, 20708701 25353983, 19280225, 21315084, 20613079, 19375649, 19213447, 19989009 18191823, 27314390, 26336977, 25775213, 24393981, 22568016, 25639019 17319928, 14705949, 19703301, 28390273, 21626377, 20122715, 6418158 23105538, 26198926, 19258504, 21188532, 24792678, 23151677, 17890099 21649497, 26446098, 16887946, 26024732, 18791688, 19721304, 27012701 19490948, 29483672, 19619732, 21164318, 29559723, 23148260, 18090142 21641760, 19818513, 23002524, 22468781, 20139391, 24693382, 19978542 25477657, 23543183, 22165897, 19373893, 22359063, 19409212, 18373438 23035249, 20677974, 18990693, 20470877, 21422580, 21632821, 22351572 20235511, 23220453, 18742258, 18604493, 23008056, 22901797, 18610915 20978259, 20832516, 24801152, 26089440, 20907061, 19523462, 25733479

20505778, 18733351, 19183343, 21787056, 21273804, 22782647, 20544065 25093739, 17835294, 28708023, 24413809, 27846298, 18371441, 26714910 24385983, 20413820, 28986231, 24421668, 25897615, 25643931, 23195445 21281607, 20513399, 18841764, 20558005, 20093776, 18909599, 20618595 23572982, 23104033, 19211433, 20331945, 19512341, 22256431, 19637186 23066146, 19022470, 22686674, 18607546, 26875822, 24573817, 23115139 19649152, 19201867, 21294938, 20898997, 18510194, 21293600, 21842740 22454326, 24683149, 19534363, 25489607, 23061702

Version 12.1.0.2.v17

Version 12.1.0.2.v17 adds support for the following:

- Patch 29494060: DATABASE PATCH SET UPDATE 12.1.0.2.190716
- Patch 29774383: OJVM PATCH SET UPDATE 12.1.0.2.190716
- Patch 28852325: DSTv33 for RDBMS (TZDATA2018G)
- Patch 28852334: DSTv33 for OJVM (TZDATA2018G)
- Patch 17969866: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 20394750: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 24835919: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 21171382: DBMS_STATS Patch
- Patch 21091901: ONLINE MOVE OF HASH OR REF PARTITION CAN LEAVE LOCAL INDEXES INCONSISTENT
- Patch 29958796: JSON bundle patch
- Patch 20033733: PART :IMC:HIT ORA 600 [KGL-HEAP-SIZE-EXCEEDED]
- Patch 28730253: SUPPORT NEW ERA REIWA FOR JAPANESE IMPERIAL CALENDAR

Oracle patch 29494060, released July 2019

Bugs fixed:

19309466, 19902195, 18250893, 25437699, 19383839, 19781326, 16756406 18456643, 26546664, 22364044, 18845653, 19915271, 20172151, 18417036 19516448, 23713236, 24907917, 24796092, 23140259, 19243521, 19658708 18272672, 21153266, 19174430, 22243719, 19548064, 26556014, 20493163 20688221, 21387964, 13542050, 22250006, 22734547, 22243983, 21623164 19012119, 19932634, 19869255, 22232606, 18681056, 23324000, 25427662 22068305, 24589081, 19439759, 19303936, 22916353, 24835538, 22353346 19790243, 21106027, 26444887, 23088803, 22529728, 26256131, 19134173 24303148, 20447445, 21101873, 21188584, 19390567, 26513709, 25780343 19769480, 21097043, 21225209, 26245237, 20677396, 19284031, 19450314 19016730, 22517782, 20919320, 22205263, 22075064, 22551446, 22721409 18440095, 22496904, 16439813, 18354830, 20596234, 22022760, 20936905 23197103, 22515353, 21514877, 19809171, 21186167, 26111842, 18990023 13787015, 22492533, 20173897, 24624166, 17210525, 21260431, 20181030 25056052, 19370504, 21868720, 23068169, 19124589, 21383171, 19402853 22690648, 19888853, 24341675, 17722075, 25107334, 20882568, 25653109 23026585, 18604692, 20717081, 25546608, 27370965, 19081128, 20768076 22173980, 25722055, 23514710, 19178851, 20951038, 22168163, 25161298 20569094, 24308635, 19791377, 19050649, 20920911, 22475617, 19189525 19469538, 27052607, 20598042, 22458049, 18988834, 21159665, 23302839 25307368, 25699321, 17409174, 22729345, 22842151, 19238590, 25051465 16941434, 20387265, 24397438, 29378913, 20673810, 23108128, 20356733 22380919, 18436647, 23065323, 20825533, 19124336, 22294260, 24790914 20284155, 23571055, 25539063, 17365043, 25914276, 20952966, 22961508 19176223, 21300341, 23237313, 18288842, 27223075, 22353199, 20011515 22083366, 27634991, 25670786, 21419850, 26898563, 27986817, 19577410 26248143, 23294548, 23101501, 24737064, 19931709, 25423453, 25547060 23533807, 27726780, 24600330, 25600421, 18122373, 20043616, 23124895 18856999, 21450666, 18893947, 26633558, 20076781, 26029780, 21196809 21354456, 22533631, 23725036, 20464614, 19562381, 27375542, 24808595 19189317, 25669791, 18307021, 21917884, 19708632, 27213224, 25633101 29006527, 20711718, 18973548, 25982666, 19718981, 23567857, 22826718 25655390,

20684983, 21773465, 20250147, 20144019, 19197175, 26263721 19597439, 21387128, 22007324, 19180770, 18818069, 21566639, 19879746 21785691, 20424183, 24285405, 21425496, 26544823, 20322560, 22228324 23172924, 22520320, 27751755, 21575362, 25058080, 22365117, 22645009 25165496, 28950969, 18774543, 20124446, 21429602, 26153977, 29189889 19371175, 21863727, 18940497, 19074147, 22923409, 25489342, 21380789 19154375, 19044962, 25417056, 19532017, 19662635, 22374754, 20560611 25654936, 21492036, 18705806, 28420042, 19578247, 22024071, 22238921 22809871, 21184223, 19995869, 23089357, 19404068, 18921743, 19065677 19018447, 19018206, 18308268, 19777862, 29027694, 22223463, 19304354 22519146, 19445860, 26654363, 27199245, 22977256, 20890311, 27445727 21142837, 20869721, 24555417, 22179537, 21756699, 20217801, 18819908 22760595, 25483815, 23628685, 23007241, 19593445, 21080143, 27351628 20582405, 24966594, 20031873, 25489367, 18618122, 24737581, 21698350 26784509, 24739928, 18966843, 19077215, 20704450, 19068970, 20543011 19023822, 24713381, 20432873, 21756677, 22836801, 20328248, 18674047 18849537, 20087383, 25459958, 20315311, 22897344, 27534509, 25178179 19308965, 18948177, 19468991, 20868862, 21780146, 23315153, 20466628 21756661, 20397490, 19706965, 20302006, 24831514, 23240358, 22178855 19032777, 20862087, 19329654, 18974476, 20603378, 20859910, 19307662 26203182, 21847223, 20281121, 22568797, 19075256, 19076343, 18866977 28026866, 29511611, 22808310, 25635149, 20844426, 20904530, 20441797 21442094, 25079710, 24674955, 18840932, 18740837, 20294666, 25602488 21517440, 22062517, 19180394, 27337759, 19174942, 20671094, 21889720 19450116, 18411216, 20117253, 24386767, 20641666, 19931367, 25264559 19930276, 22092979, 25616268, 21625179, 20879709, 23003979, 20165574 28578164, 19272708, 19547370, 22624709, 23084507, 23184263, 20228093 21281532, 25093872, 19805359, 26324206, 19461270, 19434529, 18799063 20354900, 20378086, 29388020, 17008068, 21246723, 20831538, 20424899 20361671, 18674024, 19689979, 24411921, 19873610, 16619249, 20562898 21641414, 21091431, 19440586, 20001168, 22757364, 22175564, 20725343 21241052, 19561643, 20736227, 19399918, 19195895, 20830459, 20017509 18475439, 25790353, 21828126, 21665897, 25555252, 20746251, 22568177 25764020, 25612095, 25357142, 23096938, 19067244, 18043064, 21329301 18885870, 26243698, 26187943, 20324049, 19536415, 23709062, 28174827 20446883, 27314206, 21299490, 25313154, 18628388, 21744290, 18254023 25047724, 20591183, 27847259, 19185876, 18548433, 27207110, 22465352 24385625, 24326444, 20402832, 19627012, 29200700, 20468401, 27441326 27620950, 16863642, 19639483, 19315691, 27567477, 21479753, 19174521 23177923, 20401975, 18306996, 18851894, 27034890, 20581111, 20318889 20936731, 21060755, 25240188, 26828994, 22256560, 19188927, 23328639 27229389, 20766180, 20229001, 24570598, 25475853, 21172913, 17655240 21266085, 19028800, 29379978, 19035573, 19366375, 24523374, 25599425 25034396, 19289642, 21502702, 21291274, 18007682, 23521523, 20475845 29408136, 22148226, 22528741, 25417958, 24652769, 26088426, 19326908 19597583, 17414008, 23019710, 20897759, 26822620, 22046677, 19663176 20938170, 19891090, 24825843, 26318627, 21960504, 24509056, 19054077 26262953, 22657942, 20428621, 21899588, 23326313, 19723336, 19835133 17532734, 25300427, 19333670, 17495022, 21842017, 19285025, 21373473 23260854, 23061453, 19687159, 14643995, 20977794, 20734332, 17551063 27548131, 21977392, 24461826, 28612674, 19676012, 20588502, 23315889 19520602, 23053606, 19841800, 20245930, 19001359, 21476308, 26546754 19393542, 23533524, 21099555, 27995248, 25429959, 19141838, 19644859 21915719, 19908836, 21421886, 19358317, 19524158, 23548817, 25861398 20803014, 23025340, 19335438, 19058490, 23642282, 19207117, 18799993 25919622, 26569225, 25986062, 20835241, 24662775, 19475971, 18967382 20347562, 20348653, 19896336, 24812585, 20048359, 21896069, 20468490 19524384, 25392535, 21147908, 21695575, 20440930, 25789277, 19171086 24718260, 17867700, 19791273, 27397048, 21241829, 19591608, 22707244 18419520, 22296366, 22654475, 18914624, 19571367, 28636676, 21522582 19501299, 26007010, 19529868, 20425790, 19708342, 27997875, 16870214 18202441, 24415926, 18743542, 19001390, 21157728, 20657411, 19332396 21875360, 22606521, 25091141, 28000269, 19149990, 20382309, 22855193 16777441, 19606174, 20848335, 25495682, 19382851, 20528052, 22762046 24563422, 23125826, 22503297, 28993590, 25192729, 23338911, 22730454 19354794, 20757079, 19176326, 20298413, 19048007, 22018363, 18849970 21532755, 20860659, 22905130, 26121990, 21263635, 27710072, 22160989 23209741, 18499088, 22894949, 21059919, 18952989, 27348081, 22518784 25856821, 24457597, 25484507, 20794034, 21061354, 20554364, 19468347 17533661, 19883092, 20657441, 24401351, 21285458, 28023399, 18051556 25330273, 26412540, 19699191, 24437510, 16875041, 20669434, 18964978 22972770, 20828947, 21373076, 25492379, 25551676, 14283239, 25766822 21967197, 22922076, 25575628, 19601762, 26110259, 20368850, 21239530 20437153, 24848928, 20880215,

20798891, 25606091, 19013183, 25042823 21133343, 22695831, 24365589, 25248384, 25634317, 20134113, 19587324 20273319, 28501075, 18542562, 19017309, 26758193, 21063322, 22062026 20134339, 22077517, 22815955, 23854396, 24690216, 22507210, 16354467 20101006, 21795111, 27938623, 23501901, 18797519, 25997810, 25879984 21260397, 25029423, 19354335, 19730508, 22366558, 26658759, 25822410 6599380, 20717359, 24321547, 27097854, 21297872, 18964939, 19871910 29437712, 26366517, 21913183, 22366322, 20171986, 25695903, 20603431 21132297, 25957038, 21542577, 22507234, 23170620, 24719736, 25600342 18868646, 28587723, 29142109, 26637824, 20627866, 18110491, 16923858 24642295, 19518079, 19339555, 20466322, 25823754, 25110233, 20169408 24908321, 20842388, 17274537, 26575788, 20474192, 21644640, 28849751 21794615, 18899974, 20471920, 22806698, 19052488, 26198757, 19503821 24350620, 20074391, 19157754, 21220620, 22495062, 24316947, 19865345 19065556, 22816287, 25947799, 20878790, 23492665, 21322887, 22305887 20879889, 24350831, 19578350, 28022101, 26439748, 21893235, 19363645 21072646, 20898391, 19291380, 27060167, 18382302, 27086138, 22536802 22087683, 21197626, 21656630, 20373598, 19248799, 22707866, 28432129 19155797, 19279273, 18886413, 25490238, 20922010, 19990037, 25150925 20509482, 27255377, 24717859, 20703000, 22862134, 21526048, 28683167 24929210, 24560906, 20144308, 21620471, 19670108, 19068610, 20267166 25123585, 20476175, 18549238, 19297917, 22950945, 19385656, 23528412 19684504, 19330795, 21174504, 28357401, 20899461, 20557786, 21911701 19143550, 20118035, 19024808, 25760195, 20009833, 19604659, 16359751 26039623, 22820579, 19928926, 23314180, 20212067, 24737403, 20480209 18904062, 29030780, 26430737, 20856766, 17258582, 27169796, 21668627 20877664, 23272045, 19487147, 23149541, 24577566, 19430401, 19676905 20925795, 26482376, 21296029, 21629064, 23229229, 22865673, 20708701 25353983, 19280225, 21315084, 20613079, 19375649, 19213447, 19989009 18191823, 27314390, 26336977, 25775213, 24393981, 25639019, 17319928 14705949, 19703301, 28390273, 21626377, 20122715, 6418158, 23105538 26198926, 19258504, 21188532, 23151677, 24792678, 17890099, 21649497 26446098, 16887946, 26024732, 18791688, 19721304, 27012701, 19490948 19619732, 21164318, 23148260, 18090142, 21641760, 19818513, 23002524 20139391, 24693382, 19978542, 23543183, 22165897, 19373893, 22359063 19409212, 18373438, 23035249, 20677974, 18990693, 20470877, 21422580 21632821, 22351572, 20235511, 23220453, 18742258, 18604493, 23008056 22901797, 18610915, 20978259, 20832516, 24801152, 26089440, 20907061 19523462, 20505778, 19183343, 21787056, 21273804, 22782647, 25093739 17835294, 28708023, 24413809, 27846298, 18371441, 26714910, 24385983 20413820, 24421668, 28986231, 25897615, 25643931, 23195445, 21281607 20513399, 20558005, 20093776, 18909599, 20618595, 23572982, 19211433 20331945, 19512341, 22256431, 19637186, 19022470, 18607546, 26875822 24573817, 23115139, 19649152, 19201867, 21294938, 20898997, 18510194 21842740, 22454326, 24683149, 19534363, 25489607, 23061702

Version 12.1.0.2.v16

Version 12.1.0.2.v16 adds support for the following:

- Patch 29141015: Database Patch Set Update : 12.1.0.2.190416
- Patch 29251241: OJVM PATCH SET UPDATE 12.1.0.2.190416
- Patch 28852325: DSTv33 for RDBMS (TZDATA2018G)
- Patch 28852334: DSTv33 for OJVM (TZDATA2018G)
- Patch 17969866: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 20394750: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 24835919: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 21171382: DBMS_STATS Patch
- Patch 21091901: ONLINE MOVE OF HASH OR REF PARTITION CAN LEAVE LOCAL INDEXES INCONSISTENT
- Patch 29600862: JSON bundle patch
- Patch 20033733: PART :IMC:HIT ORA 600 [KGL-HEAP-SIZE-EXCEEDED]
- Patch 28730253: SUPPORT NEW ERA REIWA FOR JAPANESE IMPERIAL CALENDAR

Oracle patch 22785785, released April 2019

Bugs fixed:

19309466, 19902195, 18250893, 25437699, 19383839, 19781326, 16756406 18456643, 26546664, 22364044, 18845653, 19915271, 20172151, 18417036 19516448, 23713236, 24796092, 23140259, 19243521, 19658708, 18272672 21153266, 19174430, 22243719, 26556014, 20493163, 20688221, 21387964 13542050, 22734547, 22243983, 21623164, 19012119, 19932634, 19869255 22232606, 18681056, 23324000, 25427662, 22068305, 24589081, 19439759 19303936, 22916353, 24835538, 22353346, 21106027, 19790243, 26444887 23088803, 22529728, 26256131, 19134173, 20447445, 21188584, 19390567 26513709, 25780343, 19769480, 21097043, 21225209, 26245237, 20677396 19284031, 19450314, 19016730, 20919320, 22517782, 22075064, 22551446 22721409, 18440095, 22496904, 16439813, 18354830, 20596234, 22022760 20936905, 23197103, 21514877, 21186167, 26111842, 18990023, 13787015 22492533, 20173897, 24624166, 17210525, 21260431, 20181030, 25056052 19370504, 21868720, 23068169, 19124589, 21383171, 19402853, 19888853 24341675, 17722075, 20882568, 25653109, 23026585, 18604692, 20717081 25546608, 27370965, 19081128, 22173980, 23514710, 25722055, 19178851 20951038, 22168163, 25161298, 20569094, 24308635, 19791377, 19050649 20920911, 22475617, 19189525, 19469538, 27052607, 20598042, 22458049 18988834, 23302839, 25307368, 17409174, 22729345, 22842151, 19238590 16941434, 20387265, 24397438, 20673810, 23108128, 20356733, 22380919 18436647, 23065323, 20825533, 19124336, 22294260, 24790914, 20284155 23571055, 25539063, 17365043, 25914276, 20952966, 22961508, 19176223 21300341, 23237313, 18288842, 27223075, 22353199, 20011515, 22083366 25670786, 27634991, 21419850, 26898563, 19577410, 27986817, 26248143 23294548, 24737064, 23101501, 19931709, 25423453, 25547060, 23533807 27726780, 24600330, 25600421, 18122373, 20043616, 23124895, 18856999 21450666, 18893947, 26633558, 20076781, 26029780, 21196809, 21354456 23725036, 20464614, 22533631, 19562381, 27375542, 24808595, 19189317 25669791, 18307021, 21917884, 19708632, 27213224, 25633101, 20711718 29006527, 18973548, 25982666, 19718981, 23567857, 22826718, 25655390 21773465, 20250147, 20144019, 19197175, 26263721, 19597439, 21387128 22007324, 19180770, 19879746, 21785691, 20424183, 24285405, 26544823 20322560, 22228324, 23172924, 22520320, 21575362, 27751755, 25058080 22365117, 22645009, 25165496, 28950969, 18774543, 20124446, 21429602 26153977, 19371175, 21863727, 18940497, 19074147, 22923409, 25489342 21380789, 19154375, 19044962, 19532017, 19662635, 22374754, 20560611, 25654936, 21492036, 18705806, 28420042, 19578247, 22024071, 22238921 22809871, 21184223, 19995869, 23089357, 19404068, 18921743, 19065677 19018447, 19018206, 18308268, 19777862, 22223463, 19304354, 29027694 22519146, 27199245, 19445860, 26654363, 22977256, 20890311, 27445727 21142837, 20869721, 24555417, 22179537, 21756699, 20217801, 18819908 22760595, 25483815, 23628685, 23007241, 19593445, 21080143, 27351628 20582405, 20031873, 25489367, 18618122, 24737581, 26784509, 24739928 18966843, 19077215, 20704450, 19068970, 20543011, 19023822, 24713381 20432873, 21756677, 20328248, 18674047, 18849537, 20087383, 25459958 20315311, 22897344, 27534509, 25178179, 19308965, 18948177, 19468991 20868862, 21780146, 23315153, 20466628, 21756661, 20397490, 19706965 20302006, 24831514, 23240358, 22178855, 19032777, 20862087, 19329654 18974476, 20603378, 20859910, 19307662, 21847223, 20281121, 22568797 19075256, 19076343, 18866977, 22808310, 25635149, 20844426, 20904530 20441797, 21442094, 25079710, 24674955, 18840932, 18740837, 20294666 25602488, 21517440, 22062517, 27337759, 19174942, 19180394, 20671094 21889720, 18411216, 20117253, 24386767, 20641666, 25264559, 22092979 25616268, 21625179, 20879709, 23003979, 20165574, 28578164, 19272708 19547370, 22624709, 23084507, 23184263, 20228093, 21281532, 25093872 19805359, 19461270, 26324206, 19434529, 18799063, 20378086, 17008068 21246723, 20831538, 20424899, 20361671, 18674024, 19689979, 24411921 19873610, 16619249, 20562898, 21641414, 21091431, 19440586, 22757364 20001168, 22175564, 20725343, 21241052, 19561643, 20736227, 19399918 19195895, 20830459, 20017509, 18475439, 25790353, 21828126, 21665897 25555252, 20746251, 22568177, 25764020, 25612095, 25357142, 23096938 19067244, 18043064, 21329301, 18885870, 26243698, 26187943, 20324049 19536415, 23709062, 28174827, 20446883, 27314206, 21299490, 25313154 18628388, 21744290, 18254023, 25047724, 20591183, 27847259, 19185876 22207110, 22465352, 24385625, 24326444, 20402832, 19627012, 20468401 27441326, 27620950, 16863642, 19639483, 19315691, 27567477, 21479753 19174521, 20401975, 18306996, 18851894, 27034890, 20581111, 20318889 20936731, 21060755, 25240188, 26828994, 22256560, 19188927, 27229389 20766180,

20229001, 24570598, 25475853, 21172913, 17655240, 21266085 19028800, 19035573, 19366375, 24523374, 25599425, 25034396, 19289642 21502702, 21291274, 18007682, 23521523, 20475845, 22148226, 22528741 29408136, 25417958, 24652769, 26088426, 19326908, 19597583, 17414008 23019710, 20897759, 26822620, 22046677, 19663176, 20938170, 19891090 24825843, 26318627, 21960504, 24509056, 19054077, 26262953, 22657942 20428621, 21899588, 23326313, 19723336, 19835133, 17532734, 19333670 25300427, 21842017, 19285025, 21373473, 23260854, 23061453, 19687159 14643995, 20977794, 20734332, 17551063, 27548131, 21977392, 24461826 19676012, 20588502, 23315889, 19520602, 23053606, 19841800, 20245930 19001359, 21476308, 26546754, 19393542, 23533524, 21099555, 27995248 25429959, 19141838, 19644859, 21915719, 19908836, 21421886, 19358317 19524158, 23548817, 25861398, 20803014, 23025340, 19335438, 19058490 23642282, 19207117, 18799993, 25919622, 26569225, 20835241, 25986062 24662775, 19475971, 18967382, 20347562, 20348653, 19896336, 24812585 20048359, 21896069, 20468490, 19524384, 25392535, 21147908, 20440930 25789277, 19171086, 24718260, 17867700, 19791273, 27397048, 21241829 19591608, 22707244, 18419520, 22296366, 22654475, 18914624, 19571367 28636676, 21522582, 19501299, 19529868, 20425790, 26007010, 19708342 27997875, 16870214, 18202441, 24415926, 18743542, 19001390, 21157728 19332396, 21875360, 25091141, 28000269, 19149990, 20382309, 22855193 16777441, 19606174, 20848335, 25495682, 19382851, 20528052, 22762046 24563422, 23125826, 22503297, 28993590, 25192729, 23338911, 22730454 19176326, 20298413, 19048007, 18849970, 21532755, 20860659, 22905130 26121990, 21263635, 22160989, 18499088, 22894949, 21059919, 18952989 22518784, 27348081, 25856821, 24457597, 25484507, 20794034, 21061354 19468347, 17533661, 19883092, 20657441, 24401351, 21285458, 18051556 25330273, 28023399, 26412540, 19699191, 24437510, 20669434, 16875041 18964978, 22972770, 20828947, 21373076, 25492379, 25551676, 14283239 25766822, 21967197, 22922076, 25575628, 26110259, 20368850, 21239530 20437153, 24848928, 20880215, 20798891, 25606091, 19013183, 21133343 22695831, 24365589, 25634317, 19587324, 20134113, 20273319, 18542562 26758193, 19017309, 21063322, 22062026, 20134339, 22077517, 22815955 23854396, 24690216, 22507210, 16354467, 20101006, 21795111, 27938623 23501901, 18797519, 25879984, 21260397, 25029423, 19354335, 19730508 22366558, 26658759, 25822410, 6599380, 20717359, 24321547, 27097854 21297872, 18964939, 26366517, 21913183, 22366322, 20171986, 20603431 21132297, 25957038, 21542577, 22507234, 23170620, 24719736, 25600342 18868646, 26637824, 20627866, 28587723, 29142109, 18110491, 16923858 24642295, 19518079, 20466322, 19339555, 25823754, 25110233, 24908321 20842388, 17274537, 26575788, 20474192, 21644640, 28849751, 21794615 18899974, 20471920, 22806698, 19052488, 19503821, 24350620, 20074391 19157754, 21220620, 24316947, 19865345, 19065556, 22816287, 25947799 20878790, 23492665, 21322887, 22305887, 20879889, 24350831, 19578350 28022101, 26439748, 21893235, 19363645, 21072646, 20898391, 19291380 27060167, 27086138, 18382302, 22536802, 22087683, 21197626, 21656630 20373598, 19248799, 22707866, 19155797, 19279273, 28432129, 18886413 25490238, 20922010, 19990037, 25150925, 20509482, 27255377, 24717859 20703000, 22862134, 21526048, 28683167, 24929210, 24560906, 20144308 21620471, 19670108, 19068610, 20267166, 25123585, 20476175, 18549238 22950945, 19385656, 23528412, 19684504, 21174504, 20899461, 20557786 21911701, 19143550, 20118035, 19024808, 25760195, 20009833, 19604659 16359751, 26039623, 19928926, 23314180, 20212067, 24737403, 20480209 18904062, 26430737, 29030780, 20856766, 27169796, 21668627, 17258582 20877664, 19487147, 23149541, 24577566, 19430401, 19676905, 20925795 21296029, 21629064, 23229229, 22865673, 20708701, 25353983, 19280225 21315084, 20613079, 19375649, 19213447, 19989009, 18191823, 27314390 26336977, 25775213, 24393981, 25639019, 17319928, 14705949, 19703301 21626377, 20122715, 6418158, 23105538, 26198926, 19258504, 21188532 23151677, 17890099, 21649497, 26446098, 16887946, 26024732, 18791688 19721304, 27012701, 19490948, 19619732, 21164318, 23148260, 18090142 21641760, 19818513, 23002524, 20139391, 24693382, 19978542, 23543183 22165897, 22359063, 19373893, 19409212, 18373438, 23035249, 18990693 20470877, 21422580, 21632821, 22351572, 20235511, 23220453, 18742258 18604493, 23008056, 22901797, 18610915, 20978259, 20832516, 24801152 26089440, 20907061, 19523462, 20505778, 19183343, 21787056, 21273804 22782647, 25093739, 17835294, 28708023, 24413809, 18371441, 26714910 24385983, 20413820, 24421668, 25897615, 25643931, 23195445, 21281607 20513399, 20558005, 20093776, 18909599, 20618595, 23572982, 19211433 20331945, 19512341, 22256431, 19637186, 19022470, 18607546, 26875822 24573817, 23115139, 19649152, 19201867, 21294938, 20898997, 18510194 21842740, 22454326, 24683149, 19534363, 25489607

Version 12.1.0.2.v15

Version 12.1.0.2.v15 adds support for the following:

- Patch 28729169: Oracle Database Patch Set Update 12.1.0.2.190115
- Patch 28790654: Oracle JVM Patch Set Update 12.1.0.2.190115
- Patch 28125601: DSTv32 for RDBMS (TZDATA2018E)
- Patch 28127287: DSTv32 for OJVM (TZDATA2018E)
- Patch 17969866: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 20394750: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 21171382: DBMS_STATS Patch
- Patch 29125200: JSON bundle patch
- Patch 20033733: KGL heap size patch

Oracle patch 28729169, released January 2019

Bugs fixed:

19309466, 19902195, 18250893, 25437699, 19383839, 16756406, 18456643 26546664, 22364044, 18845653, 19915271, 20172151, 18417036, 23713236 24796092, 23140259, 19243521, 19658708, 18272672, 21153266, 19174430 22243719, 20493163, 20688221, 21387964, 13542050, 22734547, 21623164 19012119, 19932634, 19869255, 22232606, 18681056, 23324000, 25427662 22068305, 24589081, 19439759, 19303936, 22916353, 24835538, 22353346 21106027, 26444887, 23088803, 22529728, 26256131, 19134173, 20447445 21188584, 19390567, 26513709, 25780343, 19769480, 21097043, 21225209 26245237, 20677396, 19284031, 19450314, 19016730, 20919320, 22075064 22551446, 22721409, 18440095, 22496904, 16439813, 18354830, 20596234 22022760, 20936905, 23197103, 21514877, 26111842, 18990023, 13787015 22492533, 20173897, 24624166, 17210525, 21260431, 20181030, 25056052 19370504, 21868720, 23068169, 19124589, 19402853, 21383171, 19888853 24341675, 17722075, 20882568, 25653109, 23026585, 18604692, 20717081 25546608, 27370965, 19081128, 22173980, 23514710, 19178851, 20951038 22168163, 25161298, 20569094, 24308635, 19791377, 19050649, 20920911 19189525, 22475617, 19469538, 27052607, 20598042, 22458049, 18988834 23302839, 25307368, 17409174, 22729345, 22842151, 19238590, 16941434 20387265, 24397438, 20673810, 23108128, 20356733, 22380919, 18436647 23065323, 20825533, 19124336, 22294260, 24790914, 20284155, 23571055 25539063, 17365043, 25914276, 20952966, 22961508, 19176223, 21300341 23237313, 18288842, 22353199, 27223075, 22083366, 25670786, 21419850 26898563, 19577410, 26248143, 23294548, 24737064, 19931709, 25423453 25547060, 23533807, 27726780, 24600330, 25600421, 18122373, 20043616 23124895, 18856999, 21450666, 18893947, 26633558, 20076781, 26029780 21196809, 21354456, 23725036, 20464614, 19562381, 27375542, 24808595 19189317, 25669791, 18307021, 21917884, 19708632, 27213224, 25633101 20711718, 18973548, 25982666, 19718981, 23567857, 22826718, 25655390 21773465, 20250147, 20144019, 19197175, 26263721, 19597439, 21387128 22007324, 19180770, 19879746, 21785691, 20424183, 24285405, 26544823 20322560, 22228324, 23172924, 22520320, 21575362, 25058080, 22365117 22645009, 25165496, 28950969, 18774543, 20124446, 21429602, 26153977 19371175, 21863727, 18940497, 19074147, 22923409, 25489342, 21380789 19154375, 19044962, 19532017, 19662635, 22374754, 20560611, 25654936 21492036, 18705806, 19578247, 22024071, 22238921, 22809871, 21184223 23089357, 19404068, 18921743, 19065677, 19018447, 19018206, 18308268 19777862, 22223463, 19304354, 22519146, 27199245, 22977256, 20890311 27445727, 21142837, 20869721, 24555417, 22179537, 21756699, 20217801 18819908, 22760595, 25483815, 23628685, 23007241, 19593445, 21080143 27351628, 20031873, 25489367, 18618122, 24737581, 26784509, 24739928 18966843, 19077215, 20704450, 19068970, 20543011, 19023822, 24713381 20432873, 21756677, 20328248, 18674047, 18849537, 20087383, 25459958 20315311, 22897344, 27534509, 25178179, 19308965, 18948177, 19468991 20868862, 21780146, 23315153, 20466628, 21756661, 20397490, 19706965 20302006, 24831514, 23240358, 22178855, 19032777, 20862087, 19329654 18974476, 20603378, 20859910, 19307662, 21847223, 20281121, 19075256 22568797, 19076343,

18866977, 22808310, 25635149, 20844426, 20904530 20441797, 21442094, 25079710, 24674955, 18840932, 18740837, 20294666 25602488, 21517440, 22062517, 27337759, 19174942, 20671094, 21889720 18411216, 20117253, 24386767, 20641666, 25264559, 22092979, 25616268 21625179, 20879709, 23003979, 20165574, 19272708, 19547370, 22624709 23084507, 20228093, 21281532, 25093872, 19805359, 19461270, 19434529 18799063, 20378086, 17008068, 21246723, 20831538, 20424899, 20361671 18674024, 19689979, 24411921, 19873610, 16619249, 20562898, 21641414 21091431, 19440586, 22757364, 22175564, 20725343, 21241052, 19561643 20736227, 19399918, 19195895, 20830459, 20017509, 18475439, 25790353 21828126, 21665897, 25555252, 20746251, 25764020, 25612095, 22568177 25357142, 23096938, 19067244, 18043064, 21329301, 18885870, 26243698 26187943, 20324049, 19536415, 23709062, 28174827, 20446883, 27314206 21299490, 25313154, 21744290, 18254023, 20591183, 25047724, 27847259 19185876, 27207110, 22465352, 24326444, 20402832, 19627012, 20468401 27441326, 27620950, 16863642, 19639483, 19315691, 27567477, 21479753 19174521, 20401975, 18306996, 18851894, 27034890, 20581111, 20318889 20936731, 21060755, 25240188, 26828994, 22256560, 19188927, 27229389 24570598, 20229001, 25475853, 21172913, 17655240, 21266085, 19028800 19035573, 19366375, 24523374, 25034396, 25599425, 19289642, 21291274 18007682, 23521523, 20475845, 22148226, 22528741, 25417958, 24652769 26088426, 19326908, 19597583, 17414008, 23019710, 20897759, 26822620 22046677, 20938170, 19891090, 24825843, 26318627, 21960504, 24509056 19054077, 26262953, 22657942, 20428621, 21899588, 23326313, 19723336 19835133, 17532734, 19333670, 21842017, 19285025, 21373473, 23260854 23061453, 19687159, 14643995, 20977794, 20734332, 17551063, 27548131 21977392, 24461826, 19676012, 20588502, 23315889, 19520602, 23053606 19841800, 20245930, 19001359, 21476308, 26546754, 19393542, 23533524 21099555, 27995248, 25429959, 19141838, 19644859, 21915719, 19908836 21421886, 19358317, 19524158, 23548817, 25861398, 20803014, 23025340 19335438, 19058490, 23642282, 19207117, 18799993, 25919622, 26569225 20835241, 24662775, 19475971, 18967382, 20347562, 20348653, 19896336 24812585, 20048359, 21896069, 19524384, 20468490, 25392535, 21147908 20440930, 25789277, 19171086, 24718260, 17867700, 19791273, 21241829 27397048, 19591608, 22707244, 18419520, 22296366, 22654475, 18914624 19571367, 28636676, 21522582, 19501299, 20425790, 19529868, 19708342 27997875, 16870214, 18202441, 24415926, 18743542, 19001390, 21875360 25091141, 28000269, 19149990, 20382309, 22855193, 16777441, 19606174 20848335, 25495682, 19382851, 20528052, 22762046, 24563422, 23125826 22503297, 28993590, 25192729, 23338911, 22730454, 19176326, 19048007 18849970, 21532755, 20860659, 22905130, 26121990, 21263635, 22160989 18499088, 22894949, 21059919, 18952989, 22518784, 25856821, 25484507 20794034, 19468347, 17533661, 19883092, 20657441, 24401351, 21285458 18051556, 25330273, 26412540, 19699191, 24437510, 20669434, 18964978 22972770, 20828947, 21373076, 25492379, 25551676, 14283239, 25766822 21967197, 22922076, 25575628, 26110259, 20368850, 21239530, 20437153 24848928, 20880215, 20798891, 25606091, 19013183, 21133343, 22695831 24365589, 25634317, 19587324, 20273319, 18542562, 26758193, 21063322 22062026, 20134339, 22077517, 22815955, 24690216, 22507210, 16354467 20101006, 21795111, 27938623, 23501901, 18797519, 25879984, 21260397 25029423, 19354335, 19730508, 22366558, 26658759, 25822410, 6599380 20717359, 24321547, 21297872, 27097854, 18964939, 26366517, 21913183 22366322, 20171986, 20603431, 21132297, 25957038, 21542577, 22507234 23170620, 24719736, 25600342, 18868646, 26637824, 20627866, 18110491 16923858, 24642295, 19518079, 20466322, 25823754, 25110233, 24908321 20842388, 17274537, 26575788, 20474192, 21644640, 21794615, 18899974 20471920, 22806698, 19052488, 19503821, 24350620, 20074391, 19157754 21220620, 24316947, 19865345, 19065556, 22816287, 25947799, 20878790 23492665, 21322887, 22305887, 20879889, 24350831, 19578350, 28022101 19363645, 21072646, 20898391, 19291380, 27060167, 27086138, 22536802 22087683, 21656630, 20373598, 19248799, 22707866, 19155797, 19279273 18886413, 25490238, 20922010, 19990037, 25150925, 20509482, 24717859 20703000, 22862134, 21526048, 24929210, 24560906, 28683167, 20144308 21620471, 19670108, 19068610, 20267166, 25123585, 20476175, 18549238 22950945, 19385656, 23528412, 19684504, 21174504, 20899461, 20557786 21911701, 19143550, 20118035, 19024808, 25760195, 20009833, 19604659 16359751, 26039623, 19928926, 23314180, 20212067, 24737403, 20480209 26430737, 20856766, 27169796, 21668627, 20877664, 19487147, 23149541 24577566, 19430401, 19676905, 20925795, 21296029, 21629064, 23229229 22865673, 20708701, 25353983, 19280225, 21315084, 20613079, 19375649 19213447, 19989009, 18191823, 27314390, 25775213, 26336977, 24393981 25639019, 17319928, 19703301, 21626377, 20122715, 6418158, 23105538 26198926, 19258504, 21188532, 23151677, 17890099, 21649497, 26446098 16887946, 26024732, 18791688, 19721304,

27012701, 19490948, 19619732 21164318, 23148260, 18090142, 21641760, 19818513, 20139391, 24693382 19978542, 23543183, 22165897, 22359063, 19409212, 23035249, 18990693 20470877, 21422580, 21632821, 22351572, 20235511, 23220453, 18604493 18742258, 23008056, 22901797, 18610915, 20978259, 20832516, 24801152 26089440, 20907061, 20505778, 19183343, 21787056, 21273804, 25093739 17835294, 24413809, 28708023, 18371441, 26714910, 24385983, 20413820 24421668, 25897615, 25643931, 23195445, 21281607, 20513399, 20558005 20093776, 18909599, 20618595, 23572982, 19211433, 20331945, 19512341 22256431, 19637186, 19022470, 18607546, 24573817, 23115139, 19649152 19201867, 21294938, 20898997, 18510194, 21842740, 22454326, 24683149 19534363, 25489607

Version 12.1.0.2.v14

Version 12.1.0.2.v14 adds support for the following:

- Patch 28259833: Oracle Database Patch Set Update 12.1.0.2.181016
- Patch 28440711: Oracle JVM Patch Set Update 12.1.0.2.181016
- Patch 28125601: DSTv32 for RDBMS (TZDATA2018E)
- Patch 28127287: DSTv32 for OJVM (TZDATA2018E)
- Patch 17969866: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 20394750: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 21171382: DBMS_STATS Patch
- Patch 28697469: JSON Database Patch
- Patch 20033733: KGL heap size patch

Oracle patch 28259833, released October 2018

Bugs fixed:

19309466, 19902195, 18250893, 25437699, 19383839, 16756406, 18456643 26546664, 22364044, 18845653, 19915271, 20172151, 18417036, 23713236 24796092, 23140259, 19243521, 19658708, 18272672, 21153266, 19174430 22243719, 20688221, 20493163, 21387964, 13542050, 22734547, 21623164 19012119, 19932634, 19869255, 22232606, 18681056, 23324000, 25427662 22068305, 24589081, 19439759, 19303936, 22916353, 24835538, 22353346 21106027, 26444887, 23088803, 22529728, 26256131, 19134173, 20447445 21188584, 19390567, 26513709, 25780343, 19769480, 21097043, 21225209 26245237, 20677396, 19284031, 19450314, 19016730, 20919320, 22075064 22551446, 22721409, 18440095, 22496904, 16439813, 18354830, 20596234 22022760, 20936905, 23197103, 21514877, 26111842, 18990023, 22492533 20173897, 24624166, 17210525, 21260431, 20181030, 25056052, 19370504 21868720, 23068169, 19124589, 19402853, 19888853, 24341675, 17722075 20882568, 25653109, 23026585, 18604692, 20717081, 25546608, 27370965 19081128, 22173980, 23514710, 19178851, 20951038, 22168163, 25161298 20569094, 24308635, 19791377, 19050649, 20920911, 19189525, 19469538 27052607, 20598042, 22458049, 18988834, 23302839, 25307368, 17409174 22729345, 22842151, 19238590, 16941434, 20387265, 24397438, 20673810 23108128, 20356733, 22380919, 18436647, 23065323, 20825533, 19124336 22294260, 24790914, 20284155, 25539063, 17365043, 25914276, 20952966 22961508, 19176223, 21300341, 23237313, 18288842, 22353199, 22083366 25670786, 21419850, 26898563, 19577410, 23294548, 24737064, 19931709 25423453, 25547060, 23533807, 27726780, 24600330, 25600421, 18122373 20043616, 23124895, 18856999, 21450666, 18893947, 20076781, 26633558 26029780, 21196809, 21354456, 23725036, 20464614, 19562381, 24808595 27375542, 19189317, 25669791, 18307021, 21917884, 19708632, 27213224 25633101, 20711718, 18973548, 25982666, 19718981, 22826718, 25655390 23567857, 21773465, 20250147, 19197175, 26263721, 19597439, 21387128 22007324, 19180770, 19879746, 21785691, 20424183, 24285405, 26544823 20322560, 22228324, 23172924, 22520320, 21575362, 25058080, 22365117 22645009, 25165496, 18774543, 20124446, 21429602, 26153977, 19371175 21863727, 18940497, 19074147, 22923409, 25489342, 21380789, 19154375 19044962, 19532017, 19662635, 22374754, 20560611, 25654936, 21492036 18705806, 19578247, 22024071, 22238921, 22809871, 21184223, 23089357 19404068, 18921743, 19065677, 19018447, 19018206,

18308268, 19777862 22223463, 19304354, 22519146, 27199245, 20890311, 22977256, 21142837, 20869721, 24555417, 22179537, 21756699, 20217801, 18819908, 22760595 25483815, 23007241, 19593445, 21080143, 27351628, 20031873, 18618122 24737581, 26784509, 24739928, 18966843, 19077215, 20704450, 19068970 20543011, 19023822, 24713381, 20432873, 21756677, 20328248, 18674047 18849537, 25459958, 20315311, 22897344, 27534509, 25178179, 19308965 18948177, 19468991, 20868862, 21780146, 20466628, 21756661, 20397490 23315153, 19706965, 20302006, 24831514, 23240358, 22178855, 19032777 20862087, 19329654, 18974476, 20603378, 20859910, 19307662, 21847223 20281121, 19075256, 19076343, 18866977, 22808310, 25635149, 20844426 20904530, 20441797, 21442094, 25079710, 24674955, 18840932, 18740837 20294666, 25602488, 21517440, 22062517, 27337759, 19174942, 20671094 21889720, 18411216, 20117253, 24386767, 20641666, 25264559, 22092979 21625179, 20879709, 23003979, 20165574, 19272708, 19547370, 22624709 23084507, 20228093, 21281532, 19805359, 19461270, 19434529, 18799063 20378086, 17008068, 21246723, 20831538, 20424899, 20361671, 18674024 19689979, 24411921, 19873610, 16619249, 20562898, 21641414, 21091431 19440586, 22757364, 22175564, 21241052, 20725343, 19561643, 20736227 19399918, 19195895, 20830459, 20017509, 25790353, 21828126, 21665897 25555252, 20746251, 25764020, 25612095, 25357142, 23096938, 19067244 18043064, 21329301, 18885870, 26243698, 26187943, 20324049, 19536415 23709062, 28174827, 20446883, 27314206, 21299490, 25313154, 21744290 18254023, 20591183, 27847259, 19185876, 27207110, 22465352, 24326444 20402832, 19627012, 27441326, 27620950, 16863642, 19639483, 19315691 21479753, 19174521, 20401975, 18306996, 18851894, 27034890, 20581111 20318889, 20936731, 21060755, 26828994, 22256560, 19188927, 27229389 24570598, 25475853, 21172913, 17655240, 21266085, 19028800, 19035573 19366375, 24523374, 25034396, 19289642, 21291274, 18007682, 23521523 20475845, 22148226, 22528741, 25417958, 24652769, 26088426, 19326908 19597583, 17414008, 23019710, 20897759, 26822620, 22046677, 20938170 24825843, 19891090, 21960504, 26318627, 24509056, 19054077, 26262953 22657942, 20428621, 21899588, 19723336, 19835133, 17532734, 19333670 21842017, 19285025, 21373473, 23260854, 19687159, 23061453, 14643995 20977794, 20734332, 17551063, 27548131, 21977392, 24461826, 19676012 20588502, 23315889, 19520602, 23053606, 19841800, 20245930, 19001359 21476308, 26546754, 19393542, 23533524, 21099555, 25429959, 19141838 19644859, 21915719, 19908836, 21421886, 19358317, 19524158, 23548817 25861398, 20803014, 23025340, 19335438, 19058490, 19207117, 23642282 18799993, 25919622, 26569225, 20835241, 24662775, 19475971, 18967382 20347562, 20348653, 19896336, 24812585, 20048359, 21896069, 19524384 25392535, 21147908, 20440930, 25789277, 19171086, 24718260, 17867700 19791273, 21241829, 19591608, 22707244, 18419520, 22296366, 18914624 19571367, 22654475, 21522582, 19501299, 20425790, 19708342, 27997875 16870214, 18202441, 24415926, 18743542, 19001390, 21875360, 25091141 28000269, 19149990, 20382309, 22855193, 16777441, 19606174, 20848335 25495682, 19382851, 20528052, 22762046, 24563422, 23125826, 22503297 25192729, 23338911, 22730454, 19176326, 19048007, 18849970, 21532755 20860659, 22905130, 21263635, 22160989, 18499088, 21059919, 18952989 22894949, 22518784, 25856821, 25484507, 20794034, 19468347, 17533661 19883092, 20657441, 24401351, 21285458, 18051556, 25330273, 19699191 24437510, 20669434, 18964978, 22972770, 20828947, 21373076, 25551676 25492379, 14283239, 25766822, 22922076, 25575628, 20368850, 21239530 20437153, 24848928, 20880215, 20798891, 25606091, 19013183, 21133343 22695831, 24365589, 25634317, 19587324, 20273319, 18542562, 26758193 21063322, 22062026, 20134339, 22077517, 22815955, 24690216, 22507210 20101006, 16354467, 21795111, 27938623, 23501901, 18797519, 25879984 21260397, 25029423, 19354335, 19730508, 22366558, 26658759, 6599380 20717359, 24321547, 21297872, 18964939, 26366517, 21913183, 22366322 20171986, 20603431, 21132297, 25957038, 21542577, 22507234, 23170620 24719736, 25600342, 18868646, 20627866, 26637824, 18110491, 16923858 24642295, 19518079, 20466322, 25823754, 25110233, 24908321, 20842388 17274537, 26575788, 20474192, 21644640, 21794615, 18899974, 20471920 22806698, 19052488, 19503821, 24350620, 20074391, 19157754, 21220620 24316947, 19865345, 19065556, 22816287, 25947799, 20878790, 23492665 21322887, 22305887, 20879889, 24350831, 19578350, 19363645, 21072646 20898391, 19291380, 27060167, 27086138, 22536802, 22087683, 21656630 20373598, 19248799, 22707866, 19155797, 19279273, 18886413, 25490238 20922010, 19990037, 25150925, 20509482, 24717859, 20703000, 22862134 21526048, 24929210, 24560906, 20144308, 21620471, 19670108, 19068610 20267166, 25123585, 20476175, 18549238, 22950945, 19385656, 23528412 19684504, 21174504, 20899461, 20557786, 21911701, 19143550, 20118035 19024808, 25760195, 20009833, 19604659, 16359751, 26039623, 19928926 23314180, 20212067, 24737403, 20480209, 26430737, 20856766, 27169796

21668627, 20877664, 19487147, 23149541, 24577566, 19430401, 19676905 20925795, 21296029, 21629064, 23229229, 22865673, 20708701, 25353983 19280225, 21315084, 19213447, 19989009, 18191823, 27314390, 25775213 24393981, 25639019, 17319928, 19703301, 21626377, 20122715, 6418158 23105538, 26198926, 19258504, 21188532, 23151677, 17890099, 21649497 26446098, 16887946, 26024732, 18791688, 19721304, 19490948, 27012701 19619732, 21164318, 18090142, 21641760, 19818513, 20139391, 24693382 19978542, 23543183, 22165897, 22359063, 19409212, 23035249, 18990693 20470877, 21422580, 21632821, 22351572, 20235511, 23220453, 18604493 23008056, 22901797, 18610915, 20832516, 24801152, 26089440, 20907061 20505778, 19183343, 21787056, 21273804, 25093739, 17835294, 24413809 18371441, 26714910, 24385983, 20413820, 24421668, 25897615, 25643931 23195445, 21281607, 20513399, 20558005, 20093776, 18909599, 20618595 23572982, 19211433, 20331945, 19512341, 22256431, 19637186, 19022470 18607546, 24573817, 19649152, 23115139, 19201867, 21294938, 20898997 18510194, 21842740, 22454326, 24683149, 19534363, 25489607

Version 12.1.0.2.v13

Version 12.1.0.2.v13 adds support for the following:

- Patch 27547329: Oracle Database Patch Set Update 12.1.0.2.180717
- Patch 27923320: Oracle JVM Patch Set Update 12.1.0.2.180717
- Patch 28125601: DSTv32 for RDBMS (TZDATA2018E)
- Patch 28127287: DSTv32 for OJVM (TZDATA2018E)
- Patch 17969866: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 20394750: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 21171382: DBMS_STATS Patch
- Patch 28307069: JSON Database Patch
- Patch 20033733: KGL heap size patch

Oracle patch 27547329, released July 2018

Bugs fixed:

19309466, 19902195, 18250893, 25437699, 19383839, 16756406, 18456643 26546664, 18845653, 19915271, 20172151, 18417036, 23713236, 24796092 19243521, 19658708, 21153266, 19174430, 22243719, 20688221, 21387964 13542050, 22734547, 21623164, 19012119, 19932634, 19869255, 22232606 18681056, 23324000, 25427662, 22068305, 24589081, 19439759, 19303936 22916353, 24835538, 22353346, 21106027, 26444887, 23088803, 22529728 26256131, 19134173, 20447445, 21188584, 19390567, 26513709, 19769480 21097043, 21225209, 20677396, 19284031, 26245237, 19450314, 19016730 20919320, 22075064, 22551446, 22721409, 18440095, 22496904, 16439813 18354830, 20596234, 22022760, 20936905, 23197103, 21514877, 26111842 18990023, 22492533, 20173897, 24624166, 17210525, 21260431, 20181030 25056052, 19370504, 21868720, 23068169, 19124589, 19402853, 19888853 24341675, 17722075, 20882568, 23026585, 25653109, 20717081, 25546608 19081128, 27370965, 22173980, 19178851, 20951038, 22168163, 25161298 20569094, 24308635, 19791377, 19050649, 20920911, 19189525, 19469538 20598042, 22458049, 18988834, 17409174, 22729345, 22842151, 19238590 16941434, 20387265, 24397438, 20673810, 23108128, 20356733, 22380919 18436647, 23065323, 20825533, 19124336, 22294260, 24790914, 20284155 25539063, 17365043, 20952966, 22961508, 19176223, 21300341, 23237313 18288842, 22353199, 22083366, 21419850, 26898563, 19577410, 23294548 19931709, 25423453, 25547060, 23533807, 24600330, 25600421, 18122373 20043616, 23124895, 18856999, 21450666, 18893947, 20076781, 26029780 21196809, 21354456, 20464614, 23725036, 19562381, 24808595, 19189317 18307021, 25669791, 21917884, 19708632, 27213224, 25633101, 20711718 18973548, 25982666, 22826718, 25655390, 21773465, 20250147, 19197175 19597439, 26263721, 21387128, 19180770, 19879746, 21785691, 20424183 24285405, 26544823, 20322560, 22228324, 22520320, 23172924, 21575362 22365117, 22645009, 25165496, 18774543, 20124446, 21429602, 19371175 21863727, 18940497, 19074147, 22923409, 21380789, 19154375, 19044962 19532017, 19662635, 22374754, 20560611,

25654936, 21492036, 18705806 19578247, 22024071, 22238921, 22809871, 21184223, 23089357, 19404068 18921743, 19065677, 19018447, 19018206, 18308268, 19777862, 22223463 19304354, 22519146, 27199245, 20890311, 21142837, 20869721, 24555417 22179537, 21756699, 20217801, 18819908, 22760595, 25483815, 23007241 19593445, 21080143, 20031873, 18618122, 26784509, 24739928, 18966843 19077215, 20704450, 19068970, 20543011, 19023822, 24713381, 20432873 21756677, 20328248, 18674047, 18849537, 25459958, 20315311, 22897344 27534509, 25178179, 19308965, 18948177, 19468991, 20868862, 21780146 20466628, 21756661, 20397490, 19706965, 24831514, 23240358, 22178855 20302006, 19032777, 20862087, 19329654, 18974476, 20603378, 20859910 19307662, 21847223, 20281121, 19075256, 19076343, 18866977, 20844426 20904530, 20441797, 21442094, 25079710, 24674955, 18840932, 18740837 20294666, 25602488, 21517440, 22062517, 27337759, 19174942, 20671094 21889720, 18411216, 20117253, 24386767, 20641666, 25264559, 22092979 21625179, 20879709, 23003979, 20165574, 19272708, 19547370, 22624709 23084507, 20228093, 21281532, 19805359, 19461270, 19434529, 18799063 20378086, 17008068, 21246723, 20831538, 20424899, 20361671, 18674024 19689979, 24411921, 19873610, 16619249, 20562898, 21091431, 21641414 19440586, 22757364, 22175564, 21241052, 19561643, 19399918, 19195895 20830459, 20017509, 25790353, 21828126, 21665897, 20746251, 25764020 25612095, 25357142, 23096938, 19067244, 18043064, 21329301, 18885870 26187943, 20324049, 19536415, 20446883, 21299490, 27314206, 25313154 21744290, 18254023, 20591183, 27847259, 19185876, 22465352, 27207110 20402832, 19627012, 27441326, 27620950, 16863642, 19639483, 19315691 21479753, 19174521, 20401975, 18306996, 18851894, 27034890, 20581111 20318889, 20936731, 21060755, 22256560, 19188927, 24570598, 25475853 21172913, 17655240, 21266085, 19028800, 19035573, 19366375, 24523374 25034396, 19289642, 21291274, 18007682, 23521523, 20475845, 22148226 22528741, 25417958, 24652769, 26088426, 19326908, 19597583, 17414008 23019710, 20897759, 22046677, 20938170, 24825843, 21960504, 24509056 19054077, 22657942, 26262953, 20428621, 21899588, 19723336, 19835133 17532734, 19333670, 21842017, 19285025, 21373473, 23260854, 19687159 14643995, 20977794, 20734332, 17551063, 27548131, 21977392, 24461826 19676012, 20588502, 23315889, 19520602, 23053606, 19841800, 20245930 19001359, 21476308, 26546754, 19393542, 23533524, 21099555, 25429959 19141838, 19644859, 21915719, 19908836, 21421886, 19358317, 19524158 23548817, 25861398, 20803014, 23025340, 19335438, 19058490, 19207117 18799993, 26569225, 25919622, 20835241, 24662775, 19475971, 18967382 20347562, 20348653, 19896336, 24812585, 20048359, 21896069, 19524384 25392535, 20440930, 25789277, 19171086, 24718260, 17867700, 19791273 21241829, 19591608, 22707244, 18419520, 22296366, 18914624, 19571367 19501299, 20425790, 19708342, 27997875, 16870214, 18202441, 24415926 18743542, 19001390, 21875360, 25091141, 19149990, 20382309, 22855193 16777441, 19606174, 20848335, 25495682, 19382851, 20528052, 22762046 24563422, 23125826, 22503297, 25192729, 23338911, 22730454, 19176326 19048007, 18849970, 21532755, 20860659, 22905130, 21263635, 22160989 18499088, 21059919, 18952989, 22518784, 25856821, 25484507, 20794034 19468347, 17533661, 19883092, 20657441, 24401351, 21285458, 18051556 25330273, 19699191, 24437510, 20669434, 18964978, 20828947, 21373076 25551676, 14283239, 25766822, 22922076, 25575628, 20368850, 21239530 20437153, 20880215, 20798891, 25606091, 19013183, 21133343, 22695831 24365589, 19587324, 18542562, 26758193, 22062026, 20134339, 22077517 22815955, 24690216, 22507210, 20101006, 21795111, 27938623, 23501901 18797519, 21260397, 25029423, 19354335, 19730508, 22366558, 26658759 6599380, 20717359, 24321547, 21297872, 18964939, 26366517, 21913183 22366322, 20171986, 20603431, 21132297, 25957038, 21542577, 22507234 23170620, 24719736, 25600342, 18868646, 20627866, 18110491, 16923858 24642295, 19518079, 20466322, 25823754, 25110233, 24908321, 20842388 17274537, 26575788, 20474192, 21644640, 21794615, 18899974, 20471920 22806698, 19052488, 19503821, 24350620, 20074391, 19157754, 21220620 24316947, 19865345, 19065556, 22816287, 25947799, 20878790, 23492665 21322887, 20879889, 24350831, 19578350, 19363645, 21072646, 20898391 19291380, 27060167, 27086138, 22536802, 22087683, 20373598, 19248799 22707866, 19155797, 19279273, 18886413, 25490238, 20922010, 19990037 25150925, 20509482, 24717859, 20703000, 22862134, 21526048, 24929210 24560906, 20144308, 21620471, 19670108, 19068610, 20267166, 25123585 20476175, 18549238, 22950945, 19385656, 23528412, 19684504, 21174504 20899461, 20557786, 21911701, 19143550, 19024808, 20118035, 20009833 25760195, 19604659, 16359751, 26039623, 19928926, 23314180, 20212067 24737403, 20480209, 26430737, 27169796, 21668627, 20877664, 19487147 23149541, 24577566, 19430401, 19676905, 20925795, 21296029, 21629064 23229229, 22865673, 20708701, 19280225, 25353983, 21315084, 19213447 19989009, 18191823, 24393981, 25639019, 17319928, 19703301,

21626377 20122715, 6418158, 23105538, 26198926, 19258504, 21188532, 17890099 21649497, 26446098, 16887946, 26024732, 18791688, 19721304, 19490948 19619732, 21164318, 18090142, 21641760, 19818513, 20139391, 24693382 19978542, 23543183, 22165897, 22359063, 19409212, 23035249, 18990693 20470877, 21422580, 21632821, 22351572, 20235511, 23220453, 18604493 23008056, 18610915, 20832516, 24801152, 26089440, 20907061, 20505778 19183343, 21787056, 21273804, 25093739, 17835294, 24413809, 18371441 24385983, 20413820, 26714910, 24421668, 25897615, 25643931, 21281607 20513399, 23195445, 20558005, 20093776, 18909599, 20618595, 23572982 19211433, 20331945, 19512341, 22256431, 19637186, 19022470, 18607546 24573817, 19649152, 19201867, 21294938, 20898997, 18510194, 22454326 19534363, 24683149, 25489607

Version 12.1.0.2.v12

Version 12.1.0.2.v12 adds support for the following:

- Patch 27338041: DATABASE PATCH SET UPDATE 12.1.0.2.180417
- Patch 27475603: OJVM PATCH SET UPDATE 12.1.0.2.180417
- Patch 27015449: RDBMS - PROACTIVE DSTV31 UPDATE - TZDATA2017C
- Patch 27015468: PROACTIVE DSTV31 UPDATE - TZDATA2017C - NEED OJVM FIX
- Patch 17969866: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 20394750: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 21171382: AUTO DOP COMPUTES A HIGH DOP UNNECESSARILY
- Patch 27666699: JSON Database Patch
- Patch 20033733: PART :IMC: HIT ORA 600 [KGL-HEAP-SIZE-EXCEEDED]

Oracle patch 27338041, released April 2018

Bugs fixed:

19309466, 24570598, 25475853, 21172913, 19902195, 18250893, 17655240 25437699, 19383839, 21266085, 19028800, 19035573, 16756406, 19366375 18456643, 26546664, 24523374, 25034396, 19289642, 18845653, 19915271 21291274, 18007682, 20172151, 18417036, 23713236, 24796092, 23521523 20475845, 22148226, 22528741, 19243521, 19658708, 21153266, 24652769 26088426, 19326908, 19597583, 17414008, 20897759, 23019710, 19174430 22046677, 22243719, 20938170, 24825843, 21960504, 24509056, 19054077 22657942, 20688221, 20428621, 21899588, 21387964, 13542050, 19723336 19835133, 17532734, 19333670, 21842017, 19285025, 21373473, 22734547 23260854, 19687159, 14643995, 21623164, 20977794, 20734332, 19012119 19869255, 19932634, 17551063, 18681056, 22232606, 27548131, 21977392 23324000, 24461826, 19676012, 20588502, 25427662, 22068305, 23315889 19520602, 23053606, 19841800, 19439759, 20245930, 19303936, 19001359 21476308, 26546754, 22916353, 19393542, 23533524, 21099555, 24835538 22353346, 25429959, 19141838, 19644859, 21106027, 21915719, 26444887 23088803, 19908836, 21421886, 22529728, 26256131, 19358317, 19134173 19524158, 20447445, 23548817, 25861398, 20803014, 23025340, 21188584 19335438, 19390567, 19058490, 19207117, 26513709, 18799993, 26569225 20835241, 24662775, 19769480, 19475971, 21097043, 21225209, 20677396 19284031, 19450314, 19016730, 18967382, 20919320, 22075064, 20347562 20348653, 22551446, 19896336, 22721409, 24812585, 20048359, 21896069 18440095, 22496904, 19524384, 25392535, 16439813, 18354830, 20596234 20440930, 22022760, 20936905, 19171086, 23197103, 24718260, 17867700 19791273, 21514877, 26111842, 18990023, 21241829, 19591608, 22707244 18419520, 22492533, 22296366, 20173897, 24624166, 17210525, 18914624 19571367, 21260431, 19501299, 20181030, 25056052, 20425790, 19708342 19370504, 21868720, 23068169, 19124589, 19402853, 19888853, 16870214 24341675, 17722075, 18202441, 24415926, 18743542, 19001390, 20882568 23026585, 20717081, 25546608, 19081128, 22173980, 21875360, 25091141 19178851, 19149990, 20382309, 20951038, 22855193, 22168163, 16777441 25161298, 19606174, 20569094, 24308635, 20848335, 19791377, 19050649 19382851, 20920911, 20528052, 22762046, 19189525, 24563422, 23125826 22503297, 19469538, 25192729, 23338911, 20598042, 22458049, 18988834 22730454, 19176326, 19048007, 17409174, 22729345, 18849970, 21532755 20860659, 22842151, 22905130, 19238590, 16941434,

20387265, 21263635 24397438, 20673810, 23108128, 22160989, 20356733, 22380919, 18499088, 18436647, 23065323, 21059919, 20825533, 18952989, 22518784, 19124336 25856821, 22294260, 25484507, 20794034, 19468347, 20284155, 17533661 19883092, 20657441, 24401351, 25539063, 17365043, 21285458, 20952966 22961508, 18051556, 25330273, 19176223, 21300341, 23237313, 18288842 19699191, 22353199, 24437510, 22083366, 21419850, 20669434, 18964978 26898563, 19577410, 23294548, 20828947, 21373076, 25551676, 14283239 25766822, 19931709, 22922076, 25423453, 25547060, 25575628, 23533807 20368850, 21239530, 20437153, 20880215, 25600421, 20798891, 25606091 18122373, 20043616, 23124895, 19013183, 18856999, 21450666, 21133343 22695831, 18893947, 24365589, 20076781, 21196809, 21354456, 19587324 20464614, 19562381, 18542562, 26758193, 24808595, 22062026, 19189317 18307021, 21917884, 19708632, 27213224, 25633101, 20711718, 20134339 22077517, 22815955, 24690216, 18973548, 25982666, 22507210, 22826718 25655390, 21773465, 20250147, 20101006, 21795111, 19197175, 23501901 18797519, 19597439, 21387128, 19180770, 19879746, 19354335, 21785691 19730508, 20424183, 22366558, 26658759, 24285405, 6599380, 20717359 26544823, 21297872, 20322560, 18964939, 22520320, 21575362, 26366517 21913183, 22366322, 20171986, 22365117, 22645009, 25165496, 20603431 21132297, 25957038, 21542577, 22507234, 18774543, 23170620, 24719736 25600342, 20627866, 20124446, 18110491, 21429602, 16923858, 24642295 19518079, 19371175, 20466322, 21863727, 18940497, 19074147, 22923409 25823754, 25110233, 24908321, 20842388, 17274537, 21380789, 26575788 19154375, 20474192, 19044962, 19532017, 21644640, 19662635, 22374754 20560611, 25654936, 21794615, 18899974, 21492036, 18705806, 20471920 22806698, 19052488, 22024071, 22238921, 19503821, 24350620, 22809871 20074391, 21184223, 23089357, 19157754, 21220620, 19404068, 24316947 18921743, 19865345, 19065677, 19065556, 22816287, 19018447, 19018206 19777862, 25947799, 22223463, 19304354, 20878790, 22519146, 23492665 21322887, 20879889, 24350831, 20890311, 19578350, 21142837, 20869721 24555417, 22179537, 21756699, 20217801, 18819908, 19363645, 25483815 21072646, 20898391, 19291380, 27060167, 27086138, 23007241, 19593445 21080143, 22536802, 22087683, 20373598, 19248799, 20031873, 22707866 19155797, 19279273, 18886413, 18618122, 25490238, 20922010, 19990037 25150925, 20509482, 24739928, 20703000, 18966843, 19077215, 22862134 21526048, 24929210, 24560906, 20704450, 20144308, 19068970, 20543011 21620471, 19023822, 19670108, 19068610, 20267166, 24713381, 20432873 21756677, 20476175, 25123585, 18549238, 20328248, 18674047, 22950945 19385656, 18849537, 23528412, 19684504, 25459958, 20315311, 22897344 20899461, 25178179, 20557786, 21911701, 19308965, 19143550, 19024808 18948177, 19468991, 20009833, 20868862, 21780146, 20466628, 21756661 20397490, 19706965, 24831514, 23240358, 22178855, 19604659, 16359751 19032777, 20862087, 19329654, 19928926, 18974476, 23314180, 20212067 20603378, 24737403, 20480209, 20859910, 26430737, 19307662, 21847223 21668627, 20281121, 27169796, 19075256, 20877664, 19487147, 19076343 23149541, 18866977, 24577566, 19430401, 19676905, 20844426, 20904530 20925795, 20441797, 21296029, 21629064, 21442094, 23229229, 25079710 22865673, 20708701, 19280225, 21315084, 24674955, 19213447, 18840932 18740837, 20294666, 19989009, 25602488, 18191823, 21517440, 22062517 19174942, 27337759, 17319928, 20671094, 21889720, 19703301, 21626377 20122715, 23105538, 18411216, 6418158, 26198926, 20117253, 19258504 21188532, 24386767, 17890099, 21649497, 26446098, 16887946, 26024732 25264559, 18791688, 19721304, 22092979, 19490948, 19619732, 21164318 21625179, 20879709, 23003979, 20165574, 18090142, 19272708, 21641760 19818513, 19547370, 22624709, 20139391, 23084507, 24693382, 20228093 21281532, 19978542, 23543183, 22165897, 22359063, 19409212, 19805359 19461270, 23035249, 19434529, 18799063, 18990693, 20470877, 20378086 17008068, 21246723, 21422580, 21632821, 20831538, 22351572, 20424899 20361671, 18674024, 19689979, 20235511, 23220453, 24411921, 19873610 16619249, 18604493, 20562898, 21091431, 19440586, 22757364, 18610915 22175564, 21241052, 19561643, 19399918, 19195895, 20832516, 20830459 20017509, 24801152, 21828126, 20907061, 21665897, 20746251, 20505778 19183343, 25764020, 25612095, 25357142, 23096938, 21787056, 21273804 19067244, 18043064, 21329301, 18885870, 20324049, 26187943, 19536415 25093739, 17835294, 20446883, 21299490, 25313154, 24413809, 21744290 18254023, 20591183, 18371441, 24385983, 20413820, 24421668, 25897615 19185876, 25643931, 21281607, 20513399, 22465352, 20558005, 20402832 19627012, 20093776, 18909599, 20618595, 27441326, 27620950, 23572982 16863642, 19639483, 19315691, 19211433, 20331945, 19512341, 22256431 21479753, 19637186, 19174521, 19022470, 18607546, 20401975, 18306996 24573817, 18851894, 19649152, 27034890, 20581111, 19201867, 20318889 20936731, 21060755, 21294938, 20898997, 18510194, 22256560, 22454326 19534363, 25489607, 19188927

Version 12.1.0.2.v11

Version 12.1.0.2.v11 adds support for the following:

- Patch 26925311: DATABASE PATCH SET UPDATE 12.1.0.2.180116
- Patch 27001733: OJVM PATCH SET UPDATE 12.1.0.2.180116
- Patch 27015449: RDBMS - PROACTIVE DSTV31 UPDATE - TZDATA2017C
- Patch 27015468: PROACTIVE DSTV31 UPDATE - TZDATA2017C - NEED OJVM FIX
- Patch 17969866: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 20394750: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 21171382: AUTO DOP COMPUTES A HIGH DOP UNNECESSARILY
- Patch 27315904: JSON Database Patch
- Patch 20033733: ORA 600 [KGL-HEAP-SIZE-EXCEEDED]

Oracle patch 26925311, released January 2018

Bugs fixed:

21099555, 22175564, 19141838, 22083366, 20842388, 19865345, 20117253 20830459, 19791273, 20671094, 21542577, 23105538, 19243521, 20951038 22165897, 19238590, 21281532, 17008068, 19908836, 24401351, 24577566 21184223, 25427662, 20717359, 19134173, 20569094, 20031873, 20387265 20322560, 21575362, 19149990, 21263635, 18886413, 17551063, 24719736 22160989, 22519146, 21623164, 22507210, 19703301, 23338911, 19366375 18007682, 19001390, 18202441, 24285405, 25655390, 20267166, 19358317 19706965, 19068970, 24739928, 18549238, 22148226, 18797519, 26544823 20825533, 23521523, 21196809, 18940497, 19670108, 19649152, 18866977 18948177, 19404068, 22496904, 22826718, 18964978, 19176326, 19035573 20413820, 20717081, 19176223, 21106027, 20904530, 20134339, 19074147 20868862, 18411216, 23035249, 25475853, 21072646, 21322887, 22507234 20425790, 20862087, 18966843, 25861398, 24929210, 24624166, 21329301 20562898, 19333670, 19468991, 20124446, 19883092, 23543183, 20878790 22855193, 18510194, 19658708, 19591608, 19402853, 23149541, 24796092 20618595, 22238921, 21795111, 21787056, 22380919, 19469538, 21266085 17835294, 19721304, 19068610, 19791377, 22178855, 16777441, 22173980 20746251, 20048359, 21896069, 19185876, 20898391, 20281121, 20907061 22950945, 21281607, 6599380, 19577410, 22092979, 19001359, 20603378 23089357, 23572982, 19490948, 21387964, 22294260, 20832516, 17532734 22351572, 18849970, 19309466, 19081128, 20627866, 20844426, 24908321 21188532, 18791688, 21442094, 20890311, 20596234, 20368850, 26366517 18973548, 19303936, 21296029, 22536802, 20882568, 21479753, 19461270 20235511, 20936905, 22077517, 21220620, 18964939, 19430401, 22806698 22296366, 21153266, 19409212, 20703000, 22657942, 20657441, 19879746 20557786, 26758193, 23237313, 26198926, 19684504, 26088426, 21294938 19024808, 24693382, 20528052, 20977794, 18799993, 20466322, 24642295 18740837, 19662635, 18440095, 21794615, 20382309, 20228093, 19065556 20212067, 25547060, 21868720, 22905130, 20938170, 19524384, 25459958 24350831, 17722075, 20446883, 20144308, 25056052, 18952989, 24523374 16870214, 21773465, 19928926, 19835133, 21629064, 21354456, 20466628 23007241, 24386767, 25490238, 19931709, 19730508, 18819908, 20250147 23124895, 25643931, 23220453, 19188927, 20074391, 18307021, 23533807 20356733, 14643995, 26430737, 18090142, 19065677, 19547370, 26024732 21225209, 21960504, 18371441, 20397490, 26575788, 23315889, 20172151 18967382, 22729345, 19174430, 22068305, 25654936, 18419520, 21241829 19536415, 26546664, 19171086, 21889720, 21132297, 20470877, 22465352 22168163, 19335438, 24397438, 20076781, 20447445, 18856999, 20471920 19869255, 21620471, 18990693, 23096938, 17890099, 19124336, 24812585 18990023, 20101006, 21300341, 20848335, 21744290, 21241052, 20897759 21668627, 19304354, 19052488, 20543011, 20794034, 23025340, 25606091 23260854, 18681056, 19562381, 24570598, 20952966, 19896336, 20828947 25539063, 18618122, 20328248, 24365589, 20440930, 18456643, 19699191 23065323, 22865673, 19201867, 22816287, 21514877, 22022760, 18743542 20798891, 20347562, 25161298, 23294548, 19777862, 24560906, 22551446

19687159, 21373076, 19174942, 20424899, 24461826, 21641760, 21899588 22862134, 18899974, 21476308, 20598042, 21297872, 24308635, 19058490 19032777, 20171986, 22815955, 25150925, 19399918, 24718260, 19434529 22492533, 19018447, 21273804, 18051556, 22757364, 18851894, 23125826 20424183, 21842017, 19022470, 19284031, 18043064, 26898563, 20173897 23713236, 22062026, 20475845, 17274537, 19440586, 16887946, 22374754 18974476, 22961508, 24825843, 17319928, 20401975, 20708701, 22062517 24674955, 17655240, 22809871, 19805359, 16439813, 19155797, 20859910 19393542, 17210525, 22024071, 19189525, 21847223, 21649497, 19075256 25079710, 25823754, 19370504, 20315311, 22762046, 22075064, 20936731 20437153, 25165496, 18845653, 19280225, 19248799, 20560611, 18988834 21756699, 22256431, 18921743, 20245930, 21532755, 18799063, 22454326 20373598, 20476175, 19571367, 20925795, 19018206, 25264559, 24385983 20509482, 20711718, 24509056, 20588502, 20181030, 21911701, 18849537 23501901, 25034396, 19183343, 22842151, 21917884, 21142837, 20603431 19189317, 23003979, 19644859, 19390567, 19279273, 26546754, 20669434 16863642, 22528741, 22707244, 25546608, 19619732, 20348653, 18607546 19315691, 19676905, 20165574, 17867700, 23528412, 20558005, 20734332 19532017, 20922010, 19818513, 19450314, 22353346, 16941434, 20361671 25423453, 20009833, 22366558, 20294666, 23197103, 18191823, 20860659 22707866, 19195895, 19371175, 19307662, 19154375, 20043616, 20324049 21977392, 18914624, 22529728, 22256560, 25330273, 19708342, 20139391 19593445, 21291274, 19382851, 19520602, 19174521, 21875360, 19676012 19326908, 20217801, 20093776, 18840932, 21097043, 21246723, 20803014 21665897, 19143550, 23026585, 20428621, 19627012, 24415926, 22087683 23548817, 14283239, 21422580, 19213447, 19518079, 26446098, 18610915 23492665, 18674024, 24831514, 21863727, 24413809, 18306996, 19915271 21626377, 19524158, 20122715, 20513399, 18110491, 22366322, 20284155 25091141, 21080143, 20017509, 22359063, 19363645, 19597439, 21239530 23108128, 19888853, 19383839, 20880215, 21756677, 22458049, 19534363 19354335, 19044962, 19639483, 25982666, 19475971, 22353199, 21060755 22243719, 22916353, 20378086, 21260431, 21756661, 24808595, 22923409 19028800, 20877664, 22518784, 21059919, 20879889, 21380789, 19723336 19077215, 21421886, 19604659, 21285458, 23533524, 26569225, 23170620 22365117, 18288842, 19048007, 19308965, 19689979, 17409174, 19503821 23068169, 24662775, 21526048, 25429959, 19197175, 19180770, 24555417 24573817, 19902195, 26444887, 25313154, 24835538, 23324000, 20318889 21492036, 19013183, 20591183, 19012119, 20464614, 22645009, 21625179 19067244, 25178179, 23053606, 21632821, 19841800, 19512341, 19211433 22695831, 20331945, 19587324, 24316947, 19578350, 19637186, 19054077 18674047, 19708632, 20898997, 21091431, 19285025, 19289642, 25947799 21133343, 20835241, 20869721, 21172913, 25602488, 19258504, 17365043 21419850, 21644640, 19468347, 21373473, 25093739, 22721409, 16359751 24421668, 21164318, 25484507, 25489607, 22520320, 19769480, 19439759 19272708, 23088803, 19978542, 19329654, 20402832, 19873610, 23229229 21517440, 13542050, 25897615, 19291380, 21915719, 25600342, 25192729 20879709, 20677396, 19076343, 19561643, 19990037, 18909599, 19487147 22897344, 20831538, 25600421, 19016730, 18250893, 23240358, 22179537 16619249, 18354830, 24411921, 25764020, 18254023, 16756406, 21188584 19989009, 25766822, 17414008, 20688221, 20441797, 20704450, 21780146 25612095, 25957038, 24652769, 25483815, 19157754, 19207117, 24437510 18885870, 21785691, 20673810, 24341675, 21450666, 18893947, 18705806 22223463, 18417036, 16923858, 23084507, 23314180, 20919320, 22503297 20474192, 22046677, 21299490, 19501299, 19385656, 20432873, 18542562 20920911, 20899461, 21315084, 21429602, 21387128, 18122373, 20581111 22624709, 26111842, 19606174, 24690216, 18436647, 19023822, 25110233 19124589, 19178851, 19597583, 20480209, 18499088, 19050649

Version 12.1.0.2.v10

Version 12.1.0.2.v10 adds support for the following:

- Oracle October 2017 PSU, a combination of database PSU (patch 26713565) + OJVM component PSU (patch 26635845)
- Oracle recommended RDBMS patches for Oracle GoldenGate (patch 17969866)
- DBMS_STATS AUTO DOP COMPUTES A HIGH DOP UNNECESSARILY (patch 21171382)
- JSON bundle patch (patch 26750145)
- KGL heap size patch (patch 20033733)

- Timezone file DSTv30 (patch 25881255, OJVM patch 25881271)

Oracle patch 26713565, released October 2017

Bugs fixed:

21099555, 22175564, 19141838, 22083366, 20842388, 19865345, 20117253 20830459, 19791273, 20671094, 21542577, 19243521, 20951038, 22165897 19238590, 21281532, 17008068, 19908836, 24577566, 21184223, 25427662 19134173, 20569094, 20031873, 20387265, 20322560, 21575362, 19149990 21263635, 17551063, 18886413, 24719736, 22160989, 22519146, 21623164 22507210, 23338911, 19703301, 19366375, 18007682, 19001390, 18202441 24285405, 25655390, 20267166, 19358317, 19706965, 19068970, 24739928 18549238, 22148226, 18797519, 26544823, 20825533, 23521523, 21196809 18940497, 19670108, 19649152, 18866977, 18948177, 22496904, 19404068 18964978, 19176326, 19035573, 20413820, 20717081, 19176223, 21106027 20904530, 20134339, 19074147, 20868862, 23035249, 18411216, 21072646 25475853, 21322887, 22507234, 20425790, 20862087, 18966843, 25861398 21329301, 20562898, 19333670, 19468991, 20124446, 19883092, 22855193 20878790, 18510194, 19658708, 19591608, 19402853, 23149541, 20618595 22238921, 21795111, 21787056, 22380919, 19469538, 21266085, 17835294 19721304, 19068610, 19791377, 22178855, 16777441, 22173980, 20746251 20048359, 21896069, 19185876, 20898391, 20281121, 20907061, 22950945 6599380, 19577410, 22092979, 19001359, 20603378, 23089357, 21387964 19490948, 22294260, 20832516, 17532734, 22351572, 19309466, 19081128 20627866, 20844426, 24908321, 21188532, 18791688, 21442094, 20890311 20596234, 20368850, 18973548, 19303936, 21296029, 20882568, 21479753 19461270, 20235511, 22077517, 20936905, 21220620, 18964939, 19430401 22806698, 22296366, 21153266, 19409212, 22657942, 20703000, 20657441 19879746, 20557786, 26198926, 26088426, 19684504, 21294938, 19024808 24693382, 20528052, 20977794, 18799993, 20466322, 24642295, 18740837 19662635, 18440095, 21794615, 20228093, 19065556, 20212067, 25547060 21868720, 20938170, 22905130, 19524384, 25459958, 24350831, 17722075 20144308, 20446883, 25056052, 18952989, 24523374, 16870214, 19928926 19835133, 21629064, 21354456, 20466628, 24386767, 25490238, 19931709 19730508, 18819908, 20250147, 23124895, 25643931, 23220453, 19188927 20074391, 18307021, 23533807, 20356733, 26430737, 14643995, 18090142 19065677, 19547370, 21225209, 21960504, 18371441, 20397490, 26575788 23315889, 20172151, 18967382, 19174430, 22068305, 25654936, 21241829 19536415, 19171086, 26546664, 21132297, 21889720, 22465352, 22168163 19335438, 24397438, 20076781, 20447445, 18856999, 20471920, 19869255 21620471, 18990693, 23096938, 19124336, 17890099, 24812585, 18990023 21300341, 20101006, 20848335, 21744290, 21241052, 20897759, 21668627 19304354, 19052488, 20543011, 20794034, 23025340, 25606091, 23260854 18681056, 19562381, 20952966, 19896336, 20828947, 25539063, 18618122 20328248, 20440930, 18456643, 19699191, 22865673, 19201867, 22816287 22022760, 21514877, 18743542, 20798891, 20347562, 25161298, 23294548 24560906, 22551446, 19777862, 19687159, 21373076, 19174942, 20424899 21899588, 22862134, 18899974, 21476308, 20598042, 24308635, 21297872 19058490, 19032777, 20171986, 22815955, 19399918, 19434529, 19018447 18051556, 21273804, 22757364, 18851894, 23125826, 20424183, 21842017 19022470, 19284031, 18043064, 23713236, 20173897, 22062026, 20475845 17274537, 19440586, 22961508, 24825843, 18974476, 22374754, 16887946 17319928, 20401975, 20708701, 24674955, 22062517, 22809871, 17655240 19805359, 16439813, 19155797, 20859910, 19393542, 17210525, 22024071 19189525, 21847223, 21649497, 19075256, 25823754, 25079710, 20315311 22762046, 22075064, 20936731, 20437153, 18845653, 19280225, 19248799 20560611, 18988834, 21756699, 22256431, 21532755, 18921743, 20245930 22454326, 18799063, 20373598, 20476175, 19571367, 20925795, 19018206 25264559, 20711718, 20509482, 20181030, 20588502, 21911701, 18849537 23501901, 25034396, 19183343, 22842151, 21917884, 21142837, 20603431 19189317, 23003979, 19644859, 19390567, 19279273, 26546754, 20669434 16863642, 22528741, 22707244, 25546608, 19619732, 20348653, 18607546 19315691, 19676905, 20165574, 17867700, 20558005, 20734332, 19532017 20922010, 19818513, 19450314, 22353346, 16941434, 20361671, 25423453 20009833, 22366558, 20294666, 23197103, 18191823, 20860659, 19195895 19371175, 19307662, 19154375, 20043616, 21977392, 18914624, 22529728 19708342, 20139391, 25330273, 19593445, 21291274, 19382851, 19520602 19174521, 21875360, 19676012, 19326908, 20217801, 20093776, 18840932 21097043, 21246723, 20803014, 21665897, 19143550, 23026585, 20428621 19627012, 22087683, 23548817, 14283239, 21422580, 19213447, 26446098 19518079, 23492665, 18610915, 18674024,

21863727, 24413809, 18306996 19915271, 21626377, 19524158, 20122715, 20513399, 18110491, 20284155 25091141, 21080143, 20017509, 22359063, 19363645, 19597439, 21239530 23108128, 19383839, 20880215, 21756677, 19888853, 22458049, 19534363 19354335, 19044962, 19639483, 25982666, 19475971, 22353199, 21060755 22243719, 22916353, 20378086, 24808595, 21756661, 21260431, 22923409 19028800, 20877664, 21059919, 20879889, 21380789, 19723336, 19077215 21421886, 19604659, 21285458, 23533524, 23170620, 22365117, 18288842 19048007, 19308965, 19689979, 17409174, 23068169, 19503821, 24662775 25429959, 21526048, 19197175, 19180770, 24555417, 24573817, 19902195 26444887, 24835538, 23324000, 20318889, 21492036, 19013183, 20591183 19012119, 20464614, 21625179, 19067244, 23053606, 21632821, 19841800 19512341, 22695831, 20331945, 19587324, 24316947, 19578350, 19637186 19054077, 18674047, 19708632, 20898997, 19285025, 21091431, 19289642 25947799, 21133343, 20835241, 20869721, 21172913, 25602488, 19258504 17365043, 21419850, 21644640, 19468347, 21373473, 25093739, 16359751 24421668, 21164318, 25489607, 25484507, 22520320, 19769480, 19439759 19272708, 19978542, 19329654, 20402832, 19873610, 23229229, 13542050 21517440, 25897615, 19291380, 21915719, 25600342, 20879709, 20677396 19076343, 19561643, 19990037, 22897344, 18909599, 19487147, 25600421 20831538, 19016730, 18250893, 23240358, 22179537, 16619249, 18354830 24411921, 18254023, 16756406, 21188584, 19989009, 25766822, 17414008 20688221, 20441797, 20704450, 21780146, 25612095, 25957038, 24652769 25483815, 19157754, 19207117, 24437510, 18885870, 21785691, 20673810 24341675, 21450666, 18893947, 18705806, 22223463, 18417036, 16923858 23084507, 23314180, 20919320, 22503297, 20474192, 22046677, 21299490 19501299, 19385656, 20432873, 18542562, 20920911, 20899461, 21429602 21387128, 21315084, 18122373, 20581111, 26111842, 22624709, 19606174 24690216, 18436647, 19023822, 25110233, 19124589, 19178851, 19597583 18499088, 19050649

Version 12.1.0.2.v9

Version 12.1.0.2.v9 adds support for the following:

- Oracle July 2017 PSU, a combination of database PSU (patch 26609783) + OJVM component PSU (patch 26027162)
- Oracle recommended RDBMS patches for Oracle GoldenGate (patch 17969866)
- DBMS_STATS AUTO DOP COMPUTES A HIGH DOP UNNECESSARILY (patch 21171382)
- JSON bundle patch (patch 26083365)
- KGL heap size patch (patch 20033733 for 12.1.0.2)
- Timezone file DSTv30 (patch 25881255, OJVM patch 25881271)
- Adds support for [Validating DB Instance Files \(p. 1115\)](#) with the RMAN logical validation utility
- Adds support for [Setting the Default Edition for a DB Instance \(p. 1100\)](#)

Oracle patch 26609783, released July 2017

Bugs fixed:

21099555, 22175564, 19141838, 22083366, 20842388, 19865345, 20117253 19791273, 20671094, 21542577, 20951038, 19243521, 22165897, 19238590 21281532, 17008068, 19908836, 24577566, 21184223, 25427662, 19134173 20569094, 20031873, 20387265, 20322560, 21575362, 19149990, 21263635 17551063, 18886413, 22160989, 22507210, 19703301, 19366375, 18007682 19001390, 18202441, 24285405, 25655390, 20267166, 19358317, 19706965 19068970, 24739928, 18549238, 22148226, 18797519, 26544823, 20825533 21196809, 18940497, 19670108, 19649152, 18866977, 18948177, 22496904 19404068, 18964978, 19176326, 19035573, 20413820, 20717081, 19176223 21106027, 20904530, 20134339, 19074147, 20868862, 18411216, 21072646 25475853, 21322887, 22507234, 20425790, 20862087, 18966843, 21329301 20562898, 19333670, 19468991, 20124446, 19883092, 20878790, 18510194 19658708, 19591608, 19402853, 20618595, 21787056, 22380919, 21266085 19469538, 17835294, 19721304, 19068610, 19791377, 22178855, 16777441 22173980, 20746251, 20048359, 21896069, 19185876, 20898391, 20281121 20907061, 6599380, 19577410, 22092979, 19001359, 20603378, 23089357 21387964, 19490948, 22294260, 20832516, 17532734,

22351572, 19309466 19081128, 20627866, 20844426, 24908321, 21188532, 18791688, 21442094, 20890311, 20596234, 20368850, 18973548, 19303936, 21296029, 20882568 21479753, 19461270, 20235511, 22077517, 20936905, 21220620, 18964939 19430401, 22296366, 21153266, 19409212, 22657942, 20703000, 20657441 19879746, 20557786, 19684504, 21294938, 19024808, 24693382, 20528052 20977794, 18799993, 20466322, 18740837, 19662635, 18440095, 20228093 19065556, 20212067, 25547060, 21868720, 22905130, 19524384, 25459958 24350831, 17722075, 20446883, 25056052, 18952989, 24523374, 16870214 19928926, 19835133, 21629064, 21354456, 20466628, 24386767, 25490238 19931709, 19730508, 18819908, 20250147, 23124895, 25643931, 23220453 19188927, 20074391, 18307021, 23533807, 20356733, 14643995, 18090142 19065677, 19547370, 21225209, 21960504, 26575788, 20397490, 20172151 18967382, 19174430, 21241829, 19536415, 26546664, 19171086, 21132297 21889720, 22465352, 22168163, 19335438, 24397438, 20076781, 20447445 18856999, 20471920, 19869255, 21620471, 18990693, 23096938, 19124336 17890099, 24812585, 18990023, 21300341, 20101006, 20848335, 21744290 20897759, 21668627, 19304354, 19052488, 20543011, 20794034, 23025340 25606091, 23260854, 18681056, 19562381, 20952966, 19896336, 20828947 25539063, 18618122, 20328248, 20440930, 18456643, 19699191, 22865673 19201867, 22022760, 21514877, 18743542, 20798891, 20347562, 25161298 23294548, 24560906, 22551446, 19777862, 19687159, 21373076, 19174942 20424899, 21899588, 18899974, 21476308, 20598042, 24308635, 21297872 19058490, 19032777, 20171986, 22815955, 19399918, 19434529, 19018447 18051556, 21273804, 22757364, 18851894, 19022470, 19284031, 18043064 20173897, 22062026, 20475845, 17274537, 19440586, 24825843, 18974476 22374754, 16887946, 17319928, 20401975, 20708701, 24674955, 22062517 22809871, 17655240, 19805359, 16439813, 19155797, 20859910, 19393542 17210525, 22024071, 19189525, 21847223, 21649497, 19075256, 25823754 25079710, 20315311, 22762046, 22075064, 20936731, 20437153, 18845653 19280225, 19248799, 20560611, 18988834, 21756699, 18921743, 20245930 18799063, 20373598, 20476175, 19571367, 20925795, 19018206, 25264559 20711718, 20509482, 20181030, 20588502, 21911701, 18849537, 23501901 19183343, 21917884, 21142837, 20603431, 19189317, 19644859, 19390567 26546754, 19279273, 20669434, 16863642, 22528741, 25546608, 19619732 20348653, 18607546, 19315691, 19676905, 20165574, 17867700, 20558005 20734332, 19532017, 20922010, 19818513, 19450314, 22353346, 16941434 20361671, 25423453, 20009833, 22366558, 20294666, 23197103, 18191823 19195895, 19371175, 19307662, 19154375, 20043616, 21977392, 18914624 22529728, 20139391, 25330273, 19593445, 21291274, 19382851, 19520602 19174521, 21875360, 19676012, 19326908, 20217801, 20093776, 18840932 21097043, 21246723, 20803014, 21665897, 19143550, 23026585, 20428621 19627012, 14283239, 21422580, 19213447, 19518079, 18610915, 18674024 24413809, 18306996, 19915271, 21626377, 19524158, 20122715, 20513399 20284155, 25091141, 21080143, 20017509, 22359063, 19363645, 19597439 21239530, 19383839, 2080215, 21756677, 19888853, 22458049, 19534363 19354335, 19044962, 19639483, 25982666, 19475971, 22353199, 21060755 22243719, 22916353, 20378086, 24808595, 21756661, 21260431, 22923409 19028800, 20877664, 21059919, 20879889, 21380789, 19723336, 19077215 21421886, 19604659, 21285458, 23533524, 23170620, 22365117, 18288842 19048007, 19308965, 19689979, 17409174, 19503821, 21526048, 19197175 19180770, 24573817, 19902195, 24835538, 23324000, 20318889, 19013183 20591183, 19012119, 20464614, 19067244, 21632821, 19841800, 19512341 22695831, 20331945, 19587324, 24316947, 19578350, 19637186, 19054077 18674047, 19708632, 20898997, 21091431, 19289642, 21133343, 20835241 20869721, 21172913, 19258504, 17365043, 21419850, 21644640, 19468347 21373473, 25093739, 16359751, 21164318, 25484507, 22520320, 19769480 19439759, 19272708, 19978542, 19329654, 20402832, 19873610, 23229229 13542050, 21517440, 19291380, 21915719, 25600342, 20879709, 20677396 19076343, 19561643, 19990037, 18909599, 19487147, 25600421, 20831538 19016730, 18250893, 16619249, 18354830, 24411921, 16756406, 18254023 21188584, 19989009, 25766822, 17414008, 20688221, 20441797, 20704450 21780146, 25612095, 25957038, 25483815, 19157754, 19207117, 24437510 18885870, 21785691, 20673810, 21450666, 18893947, 18705806, 22223463 18417036, 16923858, 23314180, 20919320, 20474192, 22046677, 21299490 19501299, 19385656, 20432873, 20920911, 20899461, 21387128, 21315084 18122373, 20581111, 22624709, 19606174, 24690216, 18436647, 19023822 25110233, 19124589, 19178851, 19597583, 18499088, 19050649

Version 12.1.0.2.v8

Version 12.1.0.2.v8 adds support for the following:

- Oracle patch 25433980, a combination of database PSU (patch 25171037) + OJVM component PSU (patch 25437695)
- Oracle recommended RDBMS patches for Oracle GoldenGate (patch 17969866 for 12.1.0.2)
- Oracle Forms patch 18307021 for 12.1.0.2
- DBMS_STATS Patch (patch 21171382 for 12.1.0.2)
- JSON bundle patch (patch 25531469 for 12.1.0.2)
- KGL heap size patch (patch 20033733 for 12.1.0.2)
- Fixed a bug that affected PSU apply after upgrade to 12.1.0.2.v5, v6, and v7
- Timezone file DSTv28 (patch 24701840)
- Adds support for the DBMS_CHANGE_NOTIFICATION package
- Adds support for XSTREAM packages and views (may require additional licensing)

Oracle patch 25171037, released April 2017

Bugs fixed:

21099555, 22175564, 19141838, 22083366, 20842388, 20117253, 19865345 19791273, 21542577, 20951038, 19243521, 22165897, 17008068, 19908836 21281532, 19238590, 24577566, 21184223, 19134173, 20569094, 20031873 20322560, 20387265, 21575362, 19149990, 21263635, 17551063, 18886413 22160989, 22507210, 19366375, 19703301, 19001390, 24285405, 18202441 20267166, 19358317, 19706965, 19068970, 18549238, 24739928, 18797519 22148226, 20825533, 21196809, 19649152, 19670108, 18940497, 18948177 22496904, 18964978, 19176326, 19035573, 20413820, 19176223, 21106027 20904530, 20134339, 19074147, 20868862, 18411216, 25475853, 21322887 21072646, 22507234, 20425790, 20862087, 18966843, 21329301, 20562898 19333670, 20124446, 19468991, 19883092, 20878790, 18510194, 19658708 19591608, 19402853, 20618595, 21787056, 22380919, 19469538, 21266085 17835294, 19721304, 19068610, 19791377, 22178855, 16777441, 22173980 20048359, 20746251, 21896069, 19185876, 20898391, 20907061, 20281121 6599380, 19577410, 22092979, 19001359, 20603378, 23089357, 21387964 19490948, 22294260, 17532734, 20832516, 22351572, 19309466, 20627866 19081128, 20844426, 21188532, 18791688, 20890311, 21442094, 20596234 20368850, 18973548, 19303936, 21296029, 20882568, 19461270, 21479753 22077517, 20936905, 20235511, 21220620, 18964939, 19430401, 22296366 21153266, 19409212, 20703000, 22657942, 19879746, 20657441, 21294938 19684504, 19024808, 20528052, 24693382, 20977794, 18799993, 20466322 18740837, 19662635, 18440095, 20228093, 19065556, 20212067, 21868720 22905130, 19524384, 24350831, 17722075, 20446883, 25056052, 18952989 24523374, 16870214, 19928926, 19835133, 21629064, 21354456, 20466628 24386767, 25490238, 19931709, 19730508, 18819908, 20250147, 23124895 23220453, 19188927, 20074391, 18307021, 20356733, 14643995, 19065677 19547370, 21960504, 21225209, 20397490, 18967382, 19174430, 21241829 19536415, 19171086, 21889720, 22465352, 22168163, 19335438, 24397438 20447445, 18856999, 19869255, 20471920, 21620471, 23096938, 18990693 19124336, 17890099, 24812585, 18990023, 21300341, 20101006, 20848335 21744290, 20897759, 21668627, 19304354, 20543011, 19052488, 20794034 23025340, 23260854, 18681056, 20952966, 19896336, 25539063, 18618122 20328248, 20440930, 18456643, 19699191, 19201867, 22865673, 22022760 20798891, 18743542, 25161298, 20347562, 22551446, 19777862, 19687159 21373076, 19174942, 20424899, 21899588, 18899974, 21476308, 20598042 21297872, 24308635, 20171986, 19058490, 19032777, 22815955, 19399918 19434529, 21273804, 19018447, 22757364, 18851894, 19022470, 19284031 18043064, 20173897, 22062026, 20475845, 17274537, 19440586, 18974476 24825843, 22374754, 16887946, 17319928, 20401975, 20708701, 22062517 22809871, 17655240, 16439813, 19805359, 19155797, 20859910, 19393542 22024071, 17210525, 19189525, 21847223, 21649497, 25079710, 19075256 20315311, 22762046, 22075064, 20936731, 18845653, 19280225, 19248799 20560611, 18988834, 21756699, 18921743, 20245930, 18799063, 20373598 19571367, 20476175, 20925795, 19018206, 25264559, 20711718, 20509482 20181030, 20588502, 21911701, 18849537, 23501901, 19183343, 21917884 21142837, 19189317, 19644859, 19390567, 19279273, 20669434, 16863642 22528741, 25546608, 19619732, 18607546, 20348653, 19315691, 19676905 20165574, 17867700, 20558005, 20734332, 19532017, 20922010, 19818513 19450314, 22353346, 16941434, 20361671, 20009833, 22366558,

20294666 18191823, 23197103, 19195895, 19371175, 19307662, 19154375, 20043616 21977392, 18914624, 22529728, 25330273, 20139391, 19593445, 21291274 19382851, 19520602, 19174521, 21875360, 19676012, 19326908, 20217801 20093776, 18840932, 21097043, 21246723, 20803014, 21665897, 19143550 20428621, 19627012, 14283239, 21422580, 19213447, 19518079, 18610915 18674024, 24413809, 18306996, 19915271, 19524158, 20122715, 20284155 20017509, 22359063, 19363645, 19597439, 21239530, 19383839, 20880215 21756677, 19888853, 22458049, 19534363, 19354335, 19044962, 19639483 19475971, 22353199, 22243719, 21060755, 22916353, 20378086, 24808595 21756661, 21260431, 22923409, 19028800, 20877664, 21059919, 20879889 21380789, 19723336, 19077215, 19604659, 21421886, 21285458, 23533524 23170620, 22365117, 18288842, 19048007, 19308965, 19689979, 19503821 21526048, 19197175, 19180770, 19902195, 23324000, 20318889, 19013183 20591183, 19012119, 20464614, 19067244, 21632821, 19841800, 19512341 22695831, 20331945, 19587324, 24316947, 19578350, 19637186, 19054077 18674047, 19708632, 20898997, 21091431, 19289642, 21133343, 20869721 21172913, 19258504, 17365043, 21419850, 19468347, 21373473, 25093739 16359751, 21164318, 22520320, 19769480, 19439759, 19272708, 19978542 19329654, 20402832, 19873610, 23229229, 13542050, 21517440, 19291380 21915719, 20879709, 20677396, 19076343, 19561643, 19990037, 19487147 18909599, 20831538, 19016730, 18250893, 16619249, 18354830, 24411921 16756406, 18254023, 21188584, 19989009, 17414008, 20688221, 20704450 20441797, 25483815, 19157754, 24437510, 18885870, 21785691, 20673810 21450666, 18893947, 18705806, 22223463, 16923858, 18417036, 23314180 20919320, 20474192, 22046677, 21299490, 19501299, 19385656, 20920911 2089461, 21387128, 21315084, 18122373, 20581111, 19606174, 24690216 18436647, 19023822, 19124589, 19178851, 19597583, 18499088, 19050649

Version 12.1.0.2.v7

Version 12.1.0.2.v7 adds support for the following:

- Oracle patch 24917069, a combination of database PSU (patch 24732082) + OJVM component PSU (patch 24917972)
- Oracle recommended RDBMS patches for Oracle GoldenGate (patch 17969866 for 12.1.0.2)
- Oracle Forms patch 18307021 for 12.1.0.2
- DBMS_STATS Patch (patch 21171382 for 12.1.0.2)
- JSON bundle patch (patch 25089615 for 12.1.0.2)
- KGL heap size patch (patch 20033733 for 12.1.0.2)

Oracle patch 24917069, released January 2017

Bugs fixed:

24917972, 25067795, 24534298, 25076732, 25076756, 24315824, 21659726 24448240, 24448282, 23177536, 22675136, 23265914, 23265965, 23727148 22674709, 22670413, 22670385, 21188537, 22139226, 22118835, 22118851 21555660, 21811517, 19623450, 21566993, 21566944, 19176885, 21068507 21047803, 21047766, 20415564, 20408829, 20408866, 19877336, 19855285 19909862, 19895362, 19895326, 19153980, 19231857, 19223010, 19245191, 19699946, 21099555, 22175564, 19141838, 22083366, 20842388, 20117253, 19865345 19791273, 21542577, 20951038, 19243521, 22165897, 19908836, 21281532 19238590, 24577566, 21184223, 19134173, 20031873, 20387265, 21575362 19149990, 21263635, 17551063, 18886413, 22160989, 22507210, 19366375 19703301, 19001390, 24285405, 18202441, 20267166, 19358317, 19706965 24739928, 19068970, 18549238, 18797519, 22148226, 20825533, 21196809 19649152, 19670108, 18940497, 18948177, 22496904, 18964978, 19035573 19176326, 20413820, 19176223, 21106027, 20904530, 20134339, 19074147 20868862, 18411216, 21072646, 21322887, 22507234, 20425790, 18966843 21329301, 20562898, 19333670, 20124446, 19468991, 19883092, 18510194 19658708, 19591608, 19402853, 20618595, 21787056, 22380919, 19469538 21266085, 17835294, 19721304, 19791377, 19068610, 22178855, 16777441 22173980, 20048359, 20746251, 21896069, 20898391, 19185876, 20907061 20281121, 6599380, 19577410, 22092979, 19001359, 20603378, 23089357 19490948, 21387964, 22294260,

20832516, 17532734, 19309466, 20627866 19081128, 20844426, 21188532, 18791688, 20890311, 21442094, 20596234 18973548, 21296029, 19303936, 20882568, 19461270, 21479753, 22077517 20936905, 20235511, 21220620, 18964939, 19430401, 22296366, 21153266 19409212, 22657942, 19879746, 20657441, 21294938, 19684504, 24693382 20528052, 19024808, 20977794, 18799993, 20466322, 18740837, 19662635 20228093, 20212067, 19065556, 19524384, 17722075, 20446883, 25056052 24523374, 18952989, 16870214, 19928926, 19835133, 21629064, 21354456 20466628, 24386767, 19931709, 19730508, 18819908, 23124895, 23220453 19188927, 20074391, 18307021, 20356733, 14643995, 19547370, 19065677 21960504, 21225209, 20397490, 18967382, 19174430, 21241829, 19536415 19171086, 22465352, 22168163, 19335438, 24397438, 20447445, 18856999 19869255, 20471920, 21620471, 18990693, 17890099, 24812585, 18990023 21300341, 20101006, 20848335, 21744290, 20897759, 21668627, 19304354 19052488, 20794034, 23025340, 23260854, 18681056, 20952966, 19896336 20328248, 18618122, 20440930, 18456643, 19699191, 19201867, 22865673 22022760, 20798891, 18743542, 25161298, 20347562, 19777862, 22551446 19687159, 21373076, 19174942, 20424899, 21899588, 18899974, 21476308 20598042, 24308635, 19032777, 19058490, 22815955, 19399918, 19434529 21273804, 19018447, 22757364, 18851894, 19022470, 19284031, 18043064 20173897, 22062026, 20475845, 17274537, 19440586, 24825843, 18974476 22374754, 16887946, 17319928, 20401975, 20708701, 22809871, 17655240 16439813, 19805359, 19155797, 20859910, 19393542, 17210525, 22024071 21847223, 19189525, 21649497, 19075256, 20315311, 22762046, 22075064 20936731, 19280225, 18845653, 20560611, 19248799, 21756699, 18988834 20245930, 18921743, 18799063, 20373598, 19571367, 20476175, 20925795 25264559, 19018206, 20711718, 20509482, 20181030, 20588502, 18849537 23501901, 19183343, 21917884, 19189317, 19644859, 19390567, 19279273 20669434, 22528741, 16863642, 19619732, 18607546, 20348653, 19315691 19676905, 20165574, 17867700, 20558005, 20734332, 19532017, 20922010 19818513, 19450314, 22353346, 20361671, 20009833, 22366558, 20294666 23197103, 18191823, 19195895, 19307662, 19371175, 20043616, 19154375 18914624, 22529728, 20139391, 21291274, 19382851, 19520602, 19174521 21875360, 19676012, 19326908, 20217801, 20093776, 18840932, 21097043 21246723, 20803014, 21665897, 19143550, 20428621, 19627012, 14283239 19518079, 18610915, 18674024, 24413809, 18306996, 19524158, 19915271 20122715, 20284155, 20017509, 22359063, 19363645, 19597439, 21239530 19888853, 21756677, 20880215, 22458049, 19534363, 19354335, 19044962 19639483, 19475971, 22353199, 21060755, 22243719, 22916353, 20378086 24808595, 21260431, 21756661, 22923409, 20877664, 19028800, 21059919 20879889, 21380789, 19723336, 19077215, 19604659, 21421886, 21285458 23533524, 23170620, 22365117, 18288842, 19308965, 19048007, 19689979 21526048, 19197175, 19180770, 19902195, 23324000, 20318889, 19013183 20591183, 19012119, 20464614, 19067244, 21632821, 19512341, 19841800 22695831, 20331945, 19587324, 24316947, 19578350, 19637186, 18674047 19054077, 20898997, 19708632, 21091431, 19289642, 21133343, 20869721 21172913, 19258504, 17365043, 19468347, 21373473, 16359751, 19769480 19439759, 19272708, 19978542, 20402832, 19329654, 19873610, 23229229 21517440, 13542050, 19291380, 21915719, 20879709, 20677396, 19076343 19561643, 19990037, 19487147, 18909599, 20831538, 18250893, 19016730 16619249, 18354830, 18254023, 21188584, 19989009, 17414008, 20688221 20704450, 20441797, 19157754, 24437510, 18885870, 21785691, 18893947 21450666, 18705806, 22223463, 16923858, 18417036, 23314180, 20919320 20474192, 22046677, 19385656, 19501299, 20920911, 20899461, 21315084 21387128, 18122373, 20581111, 19606174, 24690216, 18436647, 19023822 19178851, 19124589, 19597583, 18499088, 19050649

Version 12.1.0.2.v6

Version 12.1.0.2.v6 adds support for the following:

- Oracle patch 24433133, a combination of database PSU (patch 24006101) + OJVM component PSU (patch 24315824)
- Oracle recommended RDBMS patches for Oracle GoldenGate (patch 17969866 for 12.1.0.2)
- Oracle Forms patch 18307021 for 12.1.0.2
- DBMS_STATS Patch (patch 21171382 for 12.1.0.2)
- JSON bundle patch (patch 24568656 for 12.1.0.2)
- Fixed a bug that caused 12c upgrade scripts to drop customer directories
- Made DIAG log directory available to customers

Baseline: Oracle Database Patch Set Update 12.1.0.2.161018 (patch 24006101, released October 2016)

Bugs fixed:

21099555, 22175564, 19141838, 22083366, 20842388, 20117253, 19865345 19791273, 19243521, 20951038, 19908836, 21281532, 19238590, 24577566 21184223, 19134173, 20387265, 19149990, 21263635, 18886413, 17551063 22160989, 22507210, 19703301, 19366375, 19001390, 18202441, 20267166 19358317, 19706965, 18549238, 19068970, 18797519, 22148226, 20825533 19649152, 19670108, 18940497, 18948177, 18964978, 19035573, 19176326 20413820, 19176223, 20904530, 20134339, 19074147, 20868862, 18411216 21322887, 22507234, 20425790, 18966843, 21329301, 19333670, 19468991 20124446, 19883092, 19658708, 19591608, 19402853, 20618595, 21787056 22380919, 21266085, 17835294, 19721304, 19791377, 19068610, 22178855 22173980, 20746251, 20048359, 20898391, 19185876, 20281121, 20907061 6599380, 19577410, 22092979, 20603378, 19001359, 19490948, 21387964 20832516, 17532734, 19309466, 19081128, 20627866, 20844426, 21188532 18791688, 21442094, 20890311, 20596234, 18973548, 21296029, 19303936 19461270, 21479753, 20936905, 20235511, 21220620, 18964939, 19430401 22296366, 21153266, 19409212, 22657942, 20657441, 19879746, 19684504 20528052, 19024808, 20977794, 18799993, 20466322, 18740837, 19662635 20228093, 19065556, 20212067, 19524384, 17722075, 20446883, 18952989 16870214, 19928926, 19835133, 21629064, 20466628, 24386767, 19931709 19730508, 18819908, 23124895, 19188927, 20074391, 20356733, 14643995 19547370, 19065677, 21960504, 21225209, 20397490, 18967382, 19174430 21241829, 19536415, 19171086, 22465352, 22168163, 19335438, 20447445 18856999, 20471920, 19869255, 21620471, 18990693, 17890099, 18990023 20101006, 21300341, 20848335, 21744290, 20897759, 21668627, 19304354 19052488, 20794034, 23260854, 18681056, 20952966, 19896336, 18618122 20328248, 20440930, 18456643, 19699191, 19201867, 22865673, 18743542 20798891, 20347562, 22551446, 19777862, 19687159, 21373076, 19174942 20424899, 21899588, 18899974, 20598042, 19032777, 19058490, 22815955 19399918, 19434529, 21273804, 19018447, 22757364, 18851894, 19284031 19022470, 18043064, 20173897, 22062026, 20475845, 17274537, 19440586 16887946, 22374754, 17319928, 20708701, 17655240, 16439813, 19805359 19155797, 20859910, 19393542, 22024071, 17210525, 21847223, 19189525 21649497, 19075256, 22762046, 22075064, 19280225, 18845653, 20560611 19248799, 21756699, 18988834, 20245930, 18921743, 18799063, 20373598 20476175, 19571367, 20925795, 19018206, 20509482, 20711718, 20588502 18849537, 19183343, 21917884, 19189317, 19644859, 19390567, 19279273 20669434, 16863642, 22528741, 19619732, 18607546, 20348653, 19315691 19676905, 20165574, 17867700, 20558005, 20734332, 19532017, 20922010 19450314, 22353346, 20361671, 20009833, 22366558, 20294666, 18191823 19307662, 19371175, 19195895, 20043616, 19154375, 18914624, 20139391 21291274, 19174521, 19520602, 19382851, 21875360, 19676012, 19326908 20217801, 20093776, 21097043, 21246723, 21665897, 19143550, 20428621 19627012, 14283239, 19518079, 18610915, 18674024, 18306996, 19524158 19915271, 20122715, 20284155, 20017509, 19363645, 19597439, 21239530 19888853, 20880215, 21756677, 19534363, 19354335, 19044962, 19639483 22353199, 22243719, 22916353, 20378086, 21756661, 21260431, 22923409 20877664, 19028800, 20879889, 19723336, 19077215, 21421886, 19604659 19308965, 19048007, 18288842, 19689979, 21526048, 19180770, 19197175 19902195, 20318889, 19013183, 19012119, 20464614, 19067244, 21632821 19512341, 19841800, 20331945, 19587324, 24316947, 19578350, 19637186 18674047, 19054077, 20898997, 19708632, 21091431, 19289642, 20869721 19258504, 17365043, 19468347, 21373473, 16359751, 19439759, 19769480 19272708, 19978542, 20402832, 19329654, 19873610, 23229229, 21517440 13542050, 19291380, 21915719, 19076343, 19561643, 19990037, 19487147 18909599, 20831538, 18250893, 19016730, 16619249, 18354830, 21188584 19989009, 17414008, 20688221, 20704450, 20441797, 19157754, 18885870 21785691, 21450666, 18893947, 18705806, 22223463, 16923858, 18417036 20919320, 20474192, 22046677, 19385656, 19501299, 20920911, 20899461 21387128, 21315084, 18122373, 20581111, 19606174, 18436647, 19023822 19178851, 19124589, 19597583, 18499088, 19050649

Version 12.1.0.2.v5

Version 12.1.0.2.v5 adds support for the following:

- Oracle patch 23615289, a combination of database PSU (patch 23054246) + OJVM component PSU (patch 23177536)
- Timezone file DSTv26 (patch 22873635 for 12.1.0.2)
- Oracle recommended RDBMS patches for Oracle GoldenGate (patch 17969866 for 12.1.0.2)
- Oracle Forms patch 18307021 for 12.1.0.2
- Added the ability to create custom password verify functions. For more information, see [Creating Custom Functions to Verify Passwords \(p. 1091\)](#).
- Fixed a bug that prevented implicit recompilation of views owned by SYS

Baseline: Oracle Database Patch Set Update 12.1.0.2.160719 (patch 23054246, released July 2016)

Bugs fixed:

19189525, 21847223, 21099555, 21649497, 19075256, 19141838, 22762046 22075064, 20117253, 19865345, 19791273, 18845653, 19280225, 19248799 19243521, 20951038, 18988834, 21756699, 21281532, 19238590, 21184223 18921743, 20245930, 18799063, 19134173, 20373598, 19571367, 20476175 20925795, 19018206, 20509482, 20711718, 20387265, 20588502, 19149990 21263635, 18849537, 18886413, 17551063, 22507210, 19183343, 19366375 19703301, 21917884, 19001390, 18202441, 19189317, 20267166, 19644859 19390567, 19358317, 19279273, 19706965, 18549238, 16863642, 19068970 22528741, 18797519, 20825533, 19619732, 18607546, 20348653, 19649152 19670108, 18940497, 18948177, 19315691, 19676905, 18964978, 19176326 20165574, 19035573, 20413820, 17867700, 20558005, 19176223, 19532017 20904530, 20134339, 19450314, 19074147, 22353346, 20868862, 18411216 22507234, 20361671, 20425790, 18966843, 20009833, 22366558, 21329301 20294666, 18191823, 19333670, 19195895, 19371175, 19307662, 19154375 20043616, 20124446, 18914624, 19468991, 19883092, 21291274, 19382851 19520602, 19174521, 21875360, 19676012, 19326908, 19658708, 19591608 19402853, 20093776, 20618595, 21787056, 22380919, 21246723, 17835294 19721304, 19068610, 19791377, 21665897, 22178855, 22173980, 20048359 20746251, 19143550, 20898391, 19185876, 19627012, 20281121, 19577410 22092979, 19001359, 14283239, 19518079, 18610915, 19490948, 17532734 18674024, 18306996, 19309466, 19081128, 19524158, 19915271, 20122715 21188532, 18791688, 20284155, 20890311, 21442094, 20596234, 18973548 21296029, 19303936, 19597439, 20936905, 20235511, 21220620, 20880215 18964939, 21756677, 19888853, 19534363, 19430401, 19354335, 19044962 19639483, 22296366, 22353199, 21153266, 19409212, 19879746, 20657441 19684504, 20528052, 19024808, 20977794, 20378086, 18799993, 21756661 21260431, 18740837, 22923409, 19028800, 20877664, 20228093, 20879889 19065556, 19723336, 19077215, 19604659, 21421886, 19524384, 17722075 19308965, 18288842, 19048007, 19689979, 20446883, 18952989, 16870214 19928926, 19835133, 21629064, 21526048, 19197175, 19180770, 20466628 19902195, 19931709, 20318889, 19013183, 19730508, 19012119, 19067244 20074391, 20356733, 14643995, 19512341, 19841800, 20331945, 19587324 19065677, 19547370, 19578350, 21225209, 19637186, 20397490, 18967382 19174430, 21241829, 19054077, 18674047, 20898997, 19708632, 19536415 21091431, 19289642, 20869721, 22168163, 19335438, 19258504, 20447445 17365043, 18856999, 19468347, 19869255, 20471920, 21373473, 21620471 16359751, 18990693, 17890099, 19769480, 19439759, 19272708, 18990023 19978542, 19329654, 20101006, 21300341, 20402832, 19873610, 20848335 23229229, 21744290, 21668627, 21517440, 13542050, 19304354, 19052488 20794034, 19291380, 21915719, 23260854, 18681056, 20952966, 19896336 19076343, 19561643, 18618122, 19990037, 20440930, 18456643, 19699191 19201867, 19487147, 18909599, 20831538, 19016730, 18250893, 20798891 18743542, 20347562, 16619249, 18354830, 22551446, 19777862, 19687159 21373076, 19174942, 20424899, 21188584, 19989009, 17414008, 20688221 21899588, 20441797, 19157754, 19058490, 19032777, 22815955, 19399918 18885870, 19434529, 21273804, 19018447, 21450666, 18893947, 18851894 16923858, 18417036, 20919320, 19022470, 19284031, 20474192, 20173897 22046677, 22062026, 19501299, 19385656, 20920911, 17274537, 20899461 21315084, 19440586, 16887946, 22374754, 17319928, 19606174, 20708701 18436647, 17655240, 19023822, 19124589, 19178851, 16439813, 19805359 19597583, 18499088, 19155797, 19050649, 19393542

Version 12.1.0.2.v4

Version 12.1.0.2.v4 adds support for the following:

- Oracle PSU 12.1.0.2.160419 (22291127)
- Timezone file DSTv25 (patch 22037014)
- Oracle recommended RDBMS patches for Oracle GoldenGate (patch 17969866)
- Adds the ability for the master user to grant the EM_EXPRESS_BASIC and EM_EXPRESS_ALL roles
- Adds the ability for the master user to grant privileges on SYS objects with the grant option using the RDSADMIN.RDSADMIN_UTIL.GRANT_SYS_OBJECT procedure
- Adds master user privileges to support most common schemas created by the Oracle Fusion Middleware Repository Creation Utility (RCU)

Baseline: Oracle Database Patch Set Update 12.1.0.2.160419 (patch 22291127, released April 2016)

Bugs fixed:

21847223, 19189525, 19075256, 19141838, 22762046, 20117253, 19865345 19791273, 19280225, 18845653, 19248799, 20951038, 19243521, 21756699 18988834, 21281532, 19238590, 18921743, 20245930, 18799063, 19134173 20373598, 19571367, 20476175, 20925795, 19018206, 20711718, 20387265 20509482, 20588502, 19149990, 18849537, 17551063, 18886413, 19183343 19703301, 21917884, 19001390, 18202441, 19189317, 19644859, 19358317 19390567, 19279273, 19706965, 22528741, 19068970, 20825533, 19619732 18607546, 20348653, 19649152, 19670108, 18940497, 18948177, 19315691 19676905, 18964978, 19035573, 20165574, 19176326, 20413820, 20558005 19176223, 19532017, 20904530, 20134339, 19450314, 22353346, 19074147 18411216, 20361671, 20425790, 18966843, 21329301, 20294666, 19333670 19195895, 19307662, 19371175, 20043616, 19154375, 20124446, 18914624 19468991, 19883092, 19382851, 19520602, 19174521, 21875360, 19676012 19326908, 19658708, 19591608, 20093776, 20618595, 21787056, 17835294 19721304, 19791377, 19068610, 22173980, 20746251, 20048359, 19143550 19185876, 19627012, 20281121, 19577410, 22092979, 19001359, 19518079 18610915, 19490948, 18674024, 18306996, 19309466, 19081128, 19915271 20122715, 21188532, 18791688, 20284155, 20890311, 21442094, 20596234 18973548, 19303936, 19597439, 20936905, 20235511, 19888853, 21756677 18964939, 19354335, 19430401, 19044962, 19639483, 21153266, 22353199 19409212, 20657441, 19879746, 19684504, 19024808, 21260431, 21756661 18799993, 20877664, 19028800, 20879889, 19065556, 19723336, 19077215 19604659, 21421886, 19524384, 18288842, 19048007, 19689979, 20446883 18952989, 16870214, 19928926, 19835133, 21526048, 20466628, 19197175 19180770, 19902195, 20318889, 19730508, 19012119, 19067244, 20074391 20356733, 14643995, 19512341, 19841800, 20331945, 19587324, 19547370 19065677, 21225209, 19637186, 20397490, 18967382, 19174430, 19054077 18674047, 19536415, 19708632, 21091431, 19289642, 22168163, 20869721 19335438, 19258504, 20447445, 17365043, 18856999, 19468347, 20471920 19869255, 21620471, 16359751, 18990693, 17890099, 19769480, 19439759 19272708, 18990023, 19978542, 20402832, 20101006, 21300341, 19329654 19873610, 21744290, 13542050, 21517440, 21668627, 19304354, 19052488 20794034, 19291380, 21915719, 18681056, 20952966, 19896336, 19076343 19561643, 19990037, 18618122, 20440930, 18456643, 19699191, 19487147 18909599, 20831538, 18250893, 19016730, 18743542, 20347562, 16619249 18354830, 19777862, 19687159, 19174942, 20424899, 19989009, 20688221 21899588, 20441797, 19157754, 19032777, 19058490, 19399918, 18885870 19434529, 21273804, 19018447, 18893947, 16923858, 18417036, 20919320 19022470, 19284031, 20474192, 22046677, 20173897, 22062026, 19385656 19501299, 17274537, 20899461, 21315084, 19440586, 22374754, 16887946 19606174, 18436647, 17655240, 19023822, 19178851, 19124589, 16439813 19805359, 19597583, 18499088, 19155797, 19050649, 19393542

Version 12.1.0.2.v3

Version 12.1.0.2.v3 adds support for the following:

- Oracle PSU 12.1.0.2.160119 (21948354).
- Timezone file DSTv25 (patch 22037014 for 12.1.0.2). 12.1.0.1 includes DSTv24, patch 20875898 (unchanged from 12.1.0.1.v3), because a backport of DSTv25 was unavailable at build time.
- Fixed an issue that prevented customers from creating more than 10 Directory objects in the database.
- Fixed an issue that prevented customers from re-granting read privileges on the ADUMP and BDUMP Directory objects.

Baseline: Oracle Database Patch Set Update 12.1.0.2.160119 (patch 21948354, released January 2016)

Bugs fixed:

19189525, 19075256, 19141838, 19865345, 19791273, 19280225, 18845653 20951038, 19243521, 19248799, 21756699, 18988834, 19238590, 21281532 20245930, 18921743, 18799063, 19134173, 19571367, 20476175, 20925795 19018206, 20509482, 20387265, 20588502, 19149990, 18849537, 18886413 17551063, 19183343, 19703301, 19001390, 18202441, 19189317, 19644859 19358317, 19390567, 19279273, 19706965, 19068970, 19619732, 20348653 18607546, 18940497, 19670108, 19649152, 18948177, 19315691, 19676905 18964978, 19035573, 20165574, 19176326, 20413820, 20558005, 19176223 19532017, 20134339, 19074147, 18411216, 20361671, 20425790, 18966843 20294666, 19307662, 19371175, 19195895, 19154375, 19468991, 19174521 19520602, 19382851, 21875360, 19326908, 19658708, 20093776, 20618595 21787056, 17835294, 19791377, 19068610, 20048359, 20746251, 19143550 19185876, 19627012, 20281121, 19577410, 22092979, 19001359, 19518079 18610915, 19490948, 18674024, 18306996, 19309466, 19081128, 19915271 20122715, 21188532, 20284155, 18791688, 20890311, 21442094, 18973548 19303936, 19597439, 20235511, 18964939, 19430401, 19044962, 19409212 19879746, 20657441, 19684504, 19024808, 18799993, 20877664, 19028800 19065556, 19723336, 19077215, 19604659, 21421886, 19524384, 19048007 18288842, 19689979, 20446883, 18952989, 16870214, 19928926, 21526048 19180770, 19197175, 19902195, 20318889, 19730508, 19012119, 19067244 20074391, 19512341, 19841800, 14643995, 20331945, 19587324, 19547370 19065677, 19637186, 21225209, 20397490, 18967382, 19174430, 18674047 19054077, 19536415, 19708632, 19289642, 20869721, 19335438, 17365043 18856999, 19869255, 20471920, 19468347, 21620471, 16359751, 18990693 17890099, 19439759, 19769480, 19272708, 19978542, 20101006, 21300341 20402832, 19329654, 19873610, 21668627, 21517440, 19304354, 19052488 20794034, 19291380, 18681056, 19896336, 19076343, 19561643, 18618122 20440930, 18456643, 19699191, 18909599, 19487147, 18250893, 19016730 18743542, 20347562, 16619249, 18354830, 19687159, 19174942, 20424899 19989009, 20688221, 20441797, 19157754, 19032777, 19058490, 19399918 18885870, 19434529, 19018447, 18417036, 20919320, 19022470, 19284031 20474192, 20173897, 22062026, 19385656, 19501299, 17274537, 20899461 19440586, 16887946, 19606174, 18436647, 17655240, 19023822, 19178851 19124589, 19805359, 19597583, 19155797, 19393542, 19050649

Version 12.1.0.2.v2

Version 12.1.0.2.v2 adds support for the following:

- Oracle PSU 12.1.0.2.5 (21359755)
- Includes the Daylight Saving Time Patch, patch 20875898: DST-24, that came out after the April 2015 PSU.

Baseline: Oracle Database Patch Set Update 12.1.0.2.5 (patch 21359755, released October 2015)

Bugs fixed:

19189525, 19075256, 19865345, 19791273, 19280225, 18845653, 19248799 19243521, 18988834, 19238590, 21281532, 18921743, 20245930, 19134173 19571367, 20476175, 20925795, 19018206,

20387265, 19149990, 18849537 19183343, 19703301, 19001390, 18202441, 19189317, 19644859, 19390567 19358317, 19279273, 19706965, 19068970, 19619732, 18607546, 20348653 18940497, 19670108, 19649152, 18948177, 19315691, 19676905, 18964978 20165574, 19035573, 19176326, 20413820, 20558005, 19176223, 19532017 20134339, 19074147, 18411216, 20361671, 20425790, 18966843, 20294666 19371175, 19307662, 19195895, 19154375, 19468991, 19174521, 19520602 19382851, 19658708, 20093776, 17835294, 19068610, 19791377, 20746251 20048359, 19143550, 19185876, 19627012, 20281121, 19577410, 19001359 19518079, 18610915, 18674024, 18306996, 19309466, 19081128, 19915271 20122715, 20284155, 18791688, 21442094, 19303936, 19597439, 20235511 18964939, 19430401, 19044962, 19409212, 20657441, 19684504, 19024808 19028800, 19065556, 19723336, 19077215, 21421886, 19524384, 19048007 18288842, 18952989, 16870214, 19928926, 19180770, 19197175, 19730508 19012119, 19067244, 20074391, 19841800, 19512341, 14643995, 20331945 19587324, 19065677, 19547370, 19637186, 21225209, 20397490, 18967382 19174430, 18674047, 19054077, 19708632, 19536415, 19289642, 19335438 17365043, 18856999, 20471920, 19468347, 21620471, 16359751, 18990693 19439759, 19769480, 19272708, 19978542, 19329654, 20402832, 19873610 19304354, 19052488, 19291380, 18681056, 19896336, 19076343, 19561643 18618122, 20440930, 18456643, 19699191, 18909599, 19487147, 18250893 19016730, 18743542, 20347562, 16619249, 18354830, 19687159, 19174942 20424899, 19989009, 20688221, 20441797, 19157754, 19058490, 19032777 19399918, 18885870, 19434529, 19018447, 18417036, 20919320, 19284031 19022470, 20474192, 22062026, 19385656, 19501299, 17274537, 20899461 19440586, 19606174, 18436647, 19023822, 19178851, 19124589, 19805359 19597583, 19155797, 19393542, 19050649

Version 12.1.0.2.v1

Version 12.1.0.2.v1 adds support for the following:

- Oracle PSU 12.1.0.2.3 (20299023)
- The In-Memory option allows storing a subset of data in an in-memory column format optimized for performance.
- Installs additional Oracle Text knowledge bases from Oracle Database. Examples media (English and French)
- Provides access to DBMS_REPAIR through RDSADMIN.RDSADMIN_DBMS_REPAIR
- Grants ALTER DATABASE LINK, ALTER PUBLIC DATABASE LINK, EXEMPT ACCESS POLICY, EXEMPT IDENTITY POLICY, and EXEMPT REDACTION POLICY to master user

Note

Version 12.1.0.2.v1 supports Enterprise Edition only.

Baseline: Oracle Database Patch Set Update 12.1.0.2.3 (patch 20299023, released April 2015)

Bugs fixed:

19189525, 19065556, 19075256, 19723336, 19077215, 19865345, 18845653 19280225, 19524384, 19248799, 18988834, 19048007, 18288842, 19238590 18921743, 18952989, 16870214, 19928926, 19134173, 19180770, 19018206 19197175, 19149990, 18849537, 19730508, 19183343, 19012119, 19001390 18202441, 19067244, 19189317, 19644859, 19358317, 19390567, 20074391 19279273, 19706965, 19068970, 19841800, 19512341, 14643995, 19619732 20348653, 18607546, 18940497, 19670108, 19649152, 19065677, 19547370 18948177, 19315691, 19637186, 19676905, 18964978, 19035573, 19176326 18967382, 19174430, 19176223, 19532017, 18674047, 19074147, 19054077 19536415, 19708632, 19289642, 20425790, 19335438, 18856999, 19371175 19468347, 19195895, 19154375, 16359751, 18990693, 19439759, 19769480 19272708, 19978542, 19329654, 19873610, 19174521, 19520602, 19382851 19658708, 19304354, 19052488, 19291380, 18681056, 19896336, 17835294 19076343, 19791377, 19068610, 19561643, 18618122, 20440930, 18456643 18909599, 19487147, 19143550, 19185876, 19016730, 20347562 19627012, 16619249, 18354830,

19577410, 19687159, 19001359, 19174942 19518079, 18610915, 18674024, 18306996, 19309466, 19081128, 19915271 19157754, 19058490, 20284155, 18791688, 18885870, 19303936, 19434529 19018447, 18417036, 19597439, 20235511, 19022470, 18964939, 19430401 19044962, 19385656, 19501299, 17274537, 19409212, 19440586, 19606174 18436647, 19023822, 19684504, 19178851, 19124589, 19805359, 19024808 19597583, 19155797, 19393542, 19050649, 19028800

Related Topics

- [Upgrading the Oracle DB Engine \(p. 979\)](#)
- [Oracle on Amazon RDS \(p. 914\)](#)

Database Engine: 11.2.0.4

The following versions are available for database engine 11.2.0.4:

- [Version 11.2.0.4.v22 \(p. 1225\)](#)
- [Version 11.2.0.4.v21 \(p. 1228\)](#)
- [Version 11.2.0.4.v20 \(p. 1229\)](#)
- [Version 11.2.0.4.v19 \(p. 1231\)](#)
- [Version 11.2.0.4.v18 \(p. 1232\)](#)
- [Version 11.2.0.4.v17 \(p. 1234\)](#)
- [Version 11.2.0.4.v16 \(p. 1235\)](#)
- [Version 11.2.0.4.v15 \(p. 1237\)](#)
- [Version 11.2.0.4.v14 \(p. 1238\)](#)
- [Version 11.2.0.4.v13 \(p. 1240\)](#)
- [Version 11.2.0.4.v12 \(p. 1241\)](#)
- [Version 11.2.0.4.v11 \(p. 1242\)](#)
- [Version 11.2.0.4.v10 \(p. 1244\)](#)
- [Version 11.2.0.4.v9 \(p. 1245\)](#)
- [Version 11.2.0.4.v8 \(p. 1246\)](#)
- [Version 11.2.0.4.v7 \(p. 1247\)](#)
- [Version 11.2.0.4.v6 \(p. 1248\)](#)
- [Version 11.2.0.4.v5 \(p. 1249\)](#)
- [Version 11.2.0.4.v4 \(p. 1250\)](#)
- [Version 11.2.0.4.v3 \(p. 1251\)](#)
- [Version 11.2.0.4.v2 \(Deprecated\) \(p. 1251\)](#)
- [Version 11.2.0.4.v1 \(p. 1252\)](#)

Version 11.2.0.4.v22

Version 11.2.0.4.v22 adds support for the following:

- Patch 29913194: DATABASE PATCH SET UPDATE 11.2.0.4.191015
- Patch 30132974: OJVM PATCH SET UPDATE 11.2.0.4.191015
- Patch 29997937: DSTV34 for RDBMS (TZDATA2019G)
- Patch 29997959: DSTV34 OJVM (TZDATA2019B)
- Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 19440386: FAILED TO RAISE ORA-1 FOR PK UPDATE WHEN CONSTRAINT=IMMEDIATE

- Patch 19277336: INTEGRATED REPLICAT INVALIDATES DEPENDENT PACKAGES RESULTING IN AN ORA-4068
- Patch 24286409: MERGE REQUEST ON TOP OF DATABASE PSU 11.2.0.4.6 FOR BUGS 20647412 21534893
- Patch 24010393: MERGE REQUEST ON TOP OF DATABASE PSU 11.2.0.4.6 FOR BUGS 12897813 21281961
- Patch 19306797: HEARTBEAT REDO IS NOT GENERATED NON-RAC HOSTS WHEN SUPPLIMENTAL LOGGING ENABLED
- Patch 19563715: LOGMINER DOES NOT MAKE PROGRESS WHEN 4GB OR MORE MEMORY IS USED IN GOLDENGATE
- Patch 20425790: LOGMINER PATCHES SHOULD TRANSPARENTLY FUNCTION IN A NON-PARTITIONING ENABLED DB
- Patch 17031322: 46719: OCIXMLDBREWRITEXML RETURNED BIND VARIABLES ARE NOT WHITESPACE PRESERVING
- Patch 30303921: MERGE REQUEST ON TOP OF DATABASE PSU 11.2.0.4.190416 FOR BUGS 29879564 14312810
- Patch 30293609: MERGE REQUEST ON TOP OF DATABASE PSU 11.2.0.4.190416 FOR BUGS 29600521 23262847
- Patch 26744595: LGSB:APPLY ABORTS W/ ORA-26786 (ROW-EXISTS) COLLISION WITH HCC(PR)-NO HCC(SB)
- Patch 28730253: SUPPORT NEW ERA REIWA FOR JAPANESE IMPERIAL CALENDAR

Oracle patch 29913194, released October 2019

Bugs fixed:

17184721, 21174504, 20169408, 21538558, 16091637, 18092127, 17381384 15979965, 20671094, 16731148, 16314254, 18441944, 13837378, 17835048 17291347, 23105538, 28254374, 13558557, 21656630, 21842740, 17008068 18382302, 17201159, 17853498, 25427662, 21197626, 17246576, 20717359 18356166, 18681862, 18440047, 20569094, 20031873, 16875449, 20387265 19788842, 17296856, 21330264, 14010183, 17648596, 17025461, 18886413 17551063, 24719736, 17258582, 17267114, 21063322, 22507210, 17912217 17889583, 18202441, 17040764, 17478145, 16524926, 25655390, 19358317 22730454, 22148226, 18747196, 26544823, 18641419, 17036973, 18948177 17811789, 16542886, 14285317, 18009564, 17359610, 16618694, 8322815 16832076, 18247991, 16692232, 22507234, 28022101, 17570240, 13871092 24624166, 26631046, 24348685, 19429927, 17848897, 17441661, 14034426 17465741, 20273319, 19207156, 16596890, 17437634, 20506706, 18510194 21343897, 28849751, 18339044, 21453153, 17951233, 21795111, 22321741 18430495, 21787056, 22380919, 20506715, 19692824, 19469538, 17811429 17903598, 19721304, 11786053, 29511611, 18230522, 19554106, 19458377 21281607, 17612828, 6599380, 18029658, 22092979, 19516448, 17040527 22321756, 17811438, 18641461, 18682983, 14657740, 25635149, 21502702 13364795, 19490948, 21387964, 17346671, 17588480, 22351572, 18235390 26474853, 18849970, 17982832, 17889549, 19309466, 16472716, 23008056 20627866, 24908321, 20134113, 25775213, 20596234, 18331850, 18641451 17019356, 20882568, 17344412, 19461270, 21179898, 17546761, 24842886 17231779, 14521849, 18203835, 18203838, 18964939, 18203837, 17313525 22195457, 18139690, 16837842, 14106803, 22296366, 17842825, 22657942 21352646, 16360112, 22594718, 20657441, 22195441, 17389192, 26198926 14565184, 19781326, 17019345, 17205719, 18740837, 18440095, 14764829 14354737, 22195448, 17019086, 13944971, 16571443, 21868720, 17186905 17080436, 18673342, 28501075, 22905130, 17027426, 19972569, 27374796 19972568, 16833845, 19972566, 20144308, 17282229, 19972564, 16870214 16410570, 21629064, 19615136, 21354456, 26039623, 19871910, 17390431 18762750, 23007241, 25248384, 16613964, 18098207, 17957017, 17484762 18471685, 19730508, 18264060, 21538485, 17323222, 17754782, 17600719 18317531, 17852463, 17596908, 17655634, 18166013, 16228604, 20074391 27053456, 24790914, 19972570, 26482376, 20856766, 18090142, 19891090 18996843, 16042673, 19854503, 17835627, 22901797, 20334344, 17393683 20861693, 18000422, 17551709, 26575788, 23315889, 20506699, 19006849

18277454, 18456514, 17258090, 19174430, 20657411, 17174582, 25654936 17242746, 27097854, 16399083, 17824637, 21132297, 17762296, 22465352 22168163, 28612674, 18604692, 17397545, 16450169, 12364061, 20067212 19373893, 18856999, 19211724, 19463893, 19463897, 27351628, 21343775 17853456, 18373438, 18673304, 20004021, 28000269, 26030218, 21668627 16194160, 17477958, 23140259, 16538760, 12982566, 24570598, 20828947 27255377, 18259031, 20296213, 21425496, 28855981, 18293054, 17610798 19699191, 23065323, 17311728, 18135678, 18774543, 23294548, 16785708 10136473, 22551446, 19777862, 24560906, 17786518, 18315328, 25879984 18334586, 12747740, 22250006, 18096714, 19032867, 21641760, 17390160 18899974, 17232014, 20598042, 16354467, 26245237, 26679352, 17484731 18673325, 16422541, 18155762, 19827973, 14015842, 22683225, 17726838 18554871, 23177648, 18051556, 20803583, 18282562, 17922254, 15990359 21972320, 16855292, 16668584, 21343838, 20299015, 29483672, 17446237 18043064, 18093615, 17694209, 23713236, 17288409, 18308268, 20475845 17274537, 13955826, 16934803, 18841764, 17634921, 17501491, 16315398 23725036, 22683212, 17006183, 13829543, 18191164, 26746894, 22809871 17655240, 28819280, 19393542, 18384391, 29633753, 21538567, 17695685 16198143, 21847223, 28199085, 25823754, 17892268, 20142975, 19584068 17165204, 25165496, 27072923, 18604493, 18508861, 21756699, 18554763 16901385, 21532755, 18189036, 17443671, 17385178, 14829250, 17936109 20476175, 20925795, 20509482, 17478514, 27441326, 16850630, 13951456 16595641, 14054676, 15861775, 21142837, 16912439, 17299889, 17297939 23003979, 16833527, 18619917, 17798953, 17630484, 19697993, 17816865 25914276, 18607546, 17571306, 21286665, 16898135, 17341326, 16819962 26910644, 17851160, 17586955, 20558005, 19049453, 21051840, 17587063 16956380, 18328509, 25042823, 14735792, 25423453, 14133975, 29033139 19718981, 18061914, 18522509, 16233738, 17518652, 21051833, 18765602 20294666, 23194294, 20860659, 18272672, 20324049, 18199537, 17332800 13609098, 22502493, 18384537, 14338435, 17945983, 27710072, 16392068 21067387, 17752995, 21097043, 21051862, 16863422, 17237521, 25505382 29483723, 18244962, 19544839, 28357401, 19433930, 24433711, 24717859 17156148, 18973907, 23026585, 17449815, 17877323, 18180390, 17088068 17037130, 20004087, 21422580, 19466309, 11733603, 25505371, 18610915 21051858, 18084625, 29027694, 18674024, 26243698, 21051852, 18091059 18306996, 16306373, 25369547, 19930276, 17787259, 19915271, 18193833 20631274, 20513399, 16344544, 26439748, 25879656, 14692762, 18614015 22782647, 17346091, 18413820, 19297917, 18228645, 17721717, 13960236 18685892, 18436307, 11883252, 19888853, 21756677, 17891943, 19475971 22353199, 16384983, 19121551, 27825893, 25634317, 12816846, 17982555 17761775, 17227073, 13936038, 22243719, 17265217, 25505394, 17071721 16721594, 18262334, 21756661, 17891946, 15913355, 17672719, 17602269 17239687, 17042658, 25555252, 17238511, 21059919, 17811456, 17284817 17752121, 20879889, 28806384, 21380789, 19601762, 17394950, 17011832 28305362, 16579084, 22195465, 16875230, 14602788, 28790634, 18325460 27567477, 30275351, 26569225, 24476265, 24476274, 12611721, 18674465 16903536, 17006570, 19689979, 28076295, 16043574, 18783224, 22836801 14705949, 24662775, 16494615, 21526048, 17392698, 19197175, 16069901 17811447, 29200700, 27870645, 28876684, 17308789, 24835538, 22195477 17865671, 17343514, 19013183, 17325413, 18316692, 16180763, 30275359 17348614, 14368995, 21983325, 17393915, 16285691, 19788303, 19211433 20331945, 17883081, 17705023, 24316947, 17614227, 23571055, 19578350 22195485, 14084247, 24975421, 26078387, 23115139, 13645875, 23328639 16777840, 19727057, 21698350, 14852021, 18744139, 18674047, 17716305 19285025, 18482502, 17622427, 19289642, 27534509, 25947799, 22195492 14458214, 20869721, 21172913, 17767676, 18723434, 25505407, 17786278 19258504, 17082983, 17365043, 21351877, 13498382, 18331812, 16065166 25489607, 16685417, 21566639, 18031668, 22893153, 17551674, 16943711 19272701, 21517440, 25897615, 17649265, 13866822, 18094246, 24528741 17783588, 14245531, 17082359, 18280813, 26007010, 20448824, 23330119 16268425, 19487147, 25600421, 18018515, 17302277, 17215560, 24411921 19271443, 25764020, 14176370, 17016369, 20777150, 16756406, 23330124 22977256, 20441797, 19769489, 28100487, 17545847, 25093656, 18260550 13853126, 17227277, 17551699, 23536835, 30237239, 25957038, 24652769 20725343, 19207117, 9756271, 17495022, 18868646, 17614134, 26667023 17546973, 19680952, 18704244, 26667015, 17050888, 18273830, 18828868 17360606, 16992075, 24563422, 17375354, 12905058, 18362222, 21429602 17254374, 27086138, 28364007, 26667032, 17571039, 17468141, 18436647 17570606, 17235750, 21168487, 17279227, 16220077, 16929165

Version 11.2.0.4.v21

Version 11.2.0.4.v21 adds support for the following:

- Patch 29497421: DATABASE PATCH SET UPDATE 11.2.0.4.190716
- Patch 29610422: OJVM PATCH SET UPDATE 11.2.0.4.190716
- Patch 28852325: DSTv33 for RDBMS (TZDATA2018G)
- Patch 28852334: DSTv33 for OJVM (TZDATA2018G)
- Patch 30018733: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 28730253: SUPPORT NEW ERA REIWA FOR JAPANESE IMPERIAL CALENDAR

Oracle patch 29497421, released July 2019

Bugs fixed:

17184721, 21174504, 20169408, 21538558, 16091637, 18092127, 17381384 15979965, 20671094, 16731148, 16314254, 18441944, 13837378, 17835048 23105538, 17291347, 28254374, 13558557, 21842740, 21656630, 17008068 18382302, 17201159, 25427662, 17853498, 21197626, 20717359, 17246576 18356166, 18681862, 18440047, 20569094, 20031873, 16875449, 20387265 19788842, 17296856, 21330264, 14010183, 17648596, 17025461, 18886413 17551063, 17258582, 24719736, 17267114, 21063322, 22507210, 17912217 17889583, 18202441, 17040764, 17478145, 16524926, 25655390, 19358317 22730454, 22148226, 18747196, 26544823, 18641419, 17036973, 18948177 17811789, 16542886, 14285317, 18009564, 17359610, 16618694, 8322815 16832076, 18247991, 16692232, 22507234, 28022101, 17570240, 13871092 24624166, 26631046, 24348685, 19429927, 17848897, 17441661, 14034426 17465741, 20273319, 19207156, 16596890, 17437634, 18510194, 21343897 20506706, 28849751, 18339044, 21453153, 17951233, 21795111, 22321741 18430495, 21787056, 22380919, 19692824, 19469538, 20506715, 17811429 17903598, 19721304, 11786053, 29511611, 18230522, 19554106, 19458377 21281607, 17612828, 6599380, 18029658, 22092979, 17040527, 22321756 17811438, 18641461, 18682983, 21502702, 14657740, 25635149, 13364795 19490948, 21387964, 17346671, 17588480, 22351572, 18235390, 18849970 26474853, 17889549, 19309466, 16472716, 23008056, 20627866, 24908321 20134113, 25775213, 20596234, 18331850, 18641451, 20882568, 17019356 17344412, 19461270, 21179898, 17546761, 24842886, 17231779, 14521849 18203835, 18203838, 18964939, 18203837, 17313525, 22195457, 18139690 16837842, 14106803, 22296366, 17842825, 22657942, 21352646, 16360112 22594718, 20657441, 22195441, 17389192, 26198926, 14565184, 17019345 19781326, 17205719, 18740837, 18440095, 14764829, 14354737, 22195448 17019086, 13944971, 16571443, 21868720, 17186905, 17080436, 18673342 28501075, 22905130, 17027426, 19972569, 27374796, 19972568, 16833845 19972566, 20144308, 17282229, 19972564, 16870214, 16410570, 21629064 19615136, 19871910, 21354456, 26039623, 17390431, 18762750, 23007241 25248384, 16613964, 18098207, 17957017, 17484762, 18471685, 19730508 18264060, 21538485, 17323222, 17754782, 17600719, 18317531, 17852463 17596908, 17655634, 18166013, 16228604, 20074391, 27053456, 24790914 19972570, 26482376, 20856766, 18090142, 19891090, 18996843, 16042673 19854503, 17835627, 22901797, 20334344, 17393683, 20861693, 18000422 17551709, 26575788, 23315889, 20506699, 19006849, 18277454, 18456514 17258090, 19174430, 20657411, 17174582, 25654936, 17242746, 27097854 16399083, 17824637, 21132297, 22465352, 17762296, 22168163, 28612674 18604692, 17397545, 16450169, 12364061, 20067212, 18856999, 19211724 19463893, 19463897, 27351628, 21343775, 17853456, 18373438, 18673304 20004021, 26030218, 28000269, 21668627, 16194160, 17477958, 23140259 16538760, 12982566, 24570598, 20828947, 27255377, 18259031, 20296213 28855981, 21425496, 18293054, 17610798, 19699191, 23065323, 17311728 18135678, 18774543, 23294548, 16785708, 10136473, 19777862, 22551446 24560906, 17786518, 18315328, 25879984, 18334586, 22250006, 12747740 18096714, 19032867, 21641760, 17390160, 18899974, 17232014, 20598042 26679352, 26245237, 16354467, 17484731, 18673325, 16422541, 18155762 19827973, 14015842, 22683225, 17726838, 18554871, 23177648, 18051556 20803583, 18282562, 17922254, 15990359, 21972320, 16855292, 16668584 21343838, 20299015, 17446237, 18043064, 18093615, 23713236, 17694209 17288409, 20475845, 18308268, 17274537, 13955826, 16934803, 17634921 17501491, 16315398, 23725036,

22683212, 17006183, 13829543, 18191164 26746894, 22809871, 17655240, 28819280, 18384391, 19393542, 21538567 17695685, 16198143, 21847223, 25823754, 17892268, 20142975, 19584068 17165204, 25165496, 18604493, 18508861, 21756699, 18554763, 16901385 21532755, 18189036, 17443671, 17385178, 14829250, 17936109, 20925795 20509482, 17478514, 27441326, 16850630, 13951456, 16595641, 14054676 15861775, 21142837, 16912439, 17299889, 17297939, 23003979, 16833527 18619917, 17798953, 19697993, 17816865, 25914276, 18607546, 17571306 21286665, 16898135, 17341326, 16819962, 26910644, 17851160, 17586955 20558005, 19049453, 21051840, 17587063, 16956380, 25042823, 14735792 18328509, 25423453, 14133975, 29033139, 19718981, 18061914, 16233738 18522509, 17518652, 21051833, 20294666, 18765602, 23194294, 20860659 18272672, 20324049, 18199537, 17332800, 13609098, 22502493, 18384537 14338435, 27710072, 17945983, 16392068, 21067387, 17752995, 21097043 21051862, 16863422, 17237521, 25505382, 18244962, 28357401, 19544839 19433930, 24433711, 24717859, 17156148, 18973907, 23026585, 17449815 17877323, 18180390, 17088068, 17037130, 20004087, 21422580, 19466309 11733603, 25505371, 18610915, 21051858, 18084625, 29027694, 18674024 26243698, 21051852, 18091059, 18306996, 25369547, 16306373, 19930276 17787259, 18193833, 19915271, 20513399, 20631274, 16344544, 26439748 25879656, 14692762, 18614015, 22782647, 17346091, 18413820, 19297917 18228645, 17721717, 13960236, 18685892, 18436307, 11883252, 19888853 21756677, 17891943, 19475971, 22353199, 16384983, 19121551, 27825893 25634317, 12816846, 17982555, 17761775, 13936038, 17227073, 22243719 17265217, 25505394, 17071721, 16721594, 18262334, 21756661, 17891946 15913355, 17672719, 17602269, 17239687, 17042658, 25555252, 17238511 21059919, 17811456, 17284817, 17752121, 20879889, 28806384, 19601762 21380789, 17394950, 28305362, 17011832, 16579084, 22195465, 16875230 14602788, 28790634, 18325460, 27567477, 26569225, 24476265, 24476274 12611721, 18674465, 16903536, 17006570, 19689979, 28076295, 16043574 18783224, 22836801, 24662775, 16494615, 21526048, 17392698, 19197175 16069901, 29200700, 17811447, 27870645, 28876684, 17308789, 24835538 22195477, 17865671, 17343514, 19013183, 17325413, 18316692, 16180763 17348614, 14368995, 21983325, 17393915, 16285691, 19211433, 20331945 17883081, 17705023, 24316947, 23571055, 17614227, 19578350, 22195485 14084247, 24975421, 26078387, 23115139, 13645875, 23328639, 16777840 21698350, 19727057, 14852021, 18744139, 18674047, 17716305, 19285025 18482502, 17622427, 19289642, 27534509, 25947799, 22195492, 14458214 20869721, 21172913, 17767676, 18723434, 25505407, 17786278, 19258504 17082983, 17365043, 21351877, 13498382, 18331812, 16065166, 25489607 16685417, 21566639, 18031668, 22893153, 17551674, 16943711, 19272701 21517440, 25897615, 17649265, 13866822, 18094246, 24528741, 17783588 14245531, 17082359, 18280813, 26007010, 20448824, 23330119, 16268425 19487147, 25600421, 18018515, 17302277, 17215560, 24411921, 19271443 25764020, 14176370, 17016369, 20777150, 16756406, 23330124, 22977256 20441797, 19769489, 28100487, 17545847, 25093656, 18260550, 13853126 17227277, 17551699, 23536835, 25957038, 24652769, 20725343, 19207117 9756271, 17495022, 18868646, 17614134, 26667023, 17546973, 19680952 18704244, 26667015, 17050888, 18273830, 18828868, 17360606, 16992075 24563422, 17375354, 12905058, 18362222, 21429602, 17254374, 27086138 26667032, 28364007, 17571039, 17468141, 18436647, 17570606, 17235750 21168487, 17279227, 16220077, 16929165

Version 11.2.0.4.v20

Version 11.2.0.4.v20 adds support for the following:

- Patch 29141056: DATABASE PATCH SET UPDATE 11.2.0.4.190416
- Patch 29251270: OJVM PATCH SET UPDATE 11.2.0.4.190416
- Patch 28852325: DSTv33 for RDBMS (TZDATA2018G)
- Patch 28852334: DSTv33 for OJVM (TZDATA2018G)
- Patch 29638593: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 28730253: SUPPORT NEW ERA REIWA FOR JAPANESE IMPERIAL CALENDAR

Oracle patch 22768427, released April 2019

Bugs fixed:

17184721, 21174504, 21538558, 16091637, 18092127, 17381384, 15979965 20671094, 16731148, 16314254, 13837378, 18441944, 17291347, 23105538 17835048, 28254374, 13558557, 21656630, 21842740, 17008068, 18382302 17201159, 25427662, 17853498, 21197626, 20717359, 17246576, 18356166 18681862, 18440047, 20569094, 20031873, 16875449, 20387265, 19788842 17296856, 21330264, 14010183, 17648596, 18886413, 17025461, 17551063 24719736, 17267114, 21063322, 22507210, 17912217, 17889583, 18202441 17040764, 17478145, 16524926, 25655390, 19358317, 22730454, 22148226 18747196, 26544823, 18641419, 17036973, 18948177, 17811789, 16542886 14285317, 18009564, 17359610, 16618694, 8322815, 16832076, 18247991 16692232, 22507234, 17570240, 13871092, 24624166, 26631046, 19429927 24348685, 17848897, 17441661, 14034426, 17465741, 20273319, 19207156 16596890, 18510194, 17437634, 21343897, 20506706, 28849751, 18339044 21453153, 17951233, 21795111, 22321741, 18430495, 21787056, 22380919 19469538, 19692824, 20506715, 17811429, 17903598, 19721304, 11786053 18230522, 19554106, 19458377, 21281607, 17612828, 6599380, 22092979 17040527, 22321756, 17811438, 18641461, 18682983, 14657740, 25635149 13364795, 19490948, 21387964, 17346671, 17588480, 22351572, 18235390 18849970, 26474853, 17889549, 19309466, 20627866, 23008056, 16472716 24908321, 20134113, 25775213, 20596234, 18331850, 18641451, 20882568 17019356, 17344412, 19461270, 21179898, 17546761, 24842886, 17231779 14521849, 18203835, 18203838, 18964939, 18203837, 17313525, 22195457 18139690, 16837842, 14106803, 22296366, 17842825, 22657942, 21352646 16360112, 22594718, 20657441, 22195441, 17389192, 26198926, 14565184 19781326, 17019345, 17205719, 18740837, 18440095, 14764829, 14354737 22195448, 17019086, 13944971, 16571443, 21868720, 17186905, 17080436 18673342, 22905130, 17027426, 27374796, 19972569, 19972568, 20144308 19972566, 17282229, 19972564, 16870214, 16410570, 21629064, 19615136 21354456, 26039623, 17390431, 18762750, 23007241, 16613964, 17957017 18098207, 17484762, 18471685, 19730508, 18264060, 21538485, 17323222 17754782, 17600719, 18317531, 17852463, 17596908, 17655634, 18166013 16228604, 20074391, 27053456, 24790914, 19972570, 20856766, 18090142 19891090, 18996843, 16042673, 19854503, 17835627, 22901797, 20334344 17393683, 20861693, 18000422, 17551709, 26575788, 23315889, 20506699 19006849, 18277454, 18456514, 19174430, 17258090, 17174582, 25654936 17242746, 27097854, 16399083, 17824637, 21132297, 22465352, 17762296 22168163, 18604692, 17397545, 16450169, 12364061, 20067212, 18856999 19211724, 19463893, 27351628, 19463897, 21343775, 17853456, 18373438 18673304, 20004021, 28000269, 26030218, 21668627, 16194160, 17477958 23140259, 16538760, 12982566, 24570598, 20828947, 27255377, 18259031 20296213, 21425496, 28855981, 18293054, 17610798, 19699191, 23065323 17311728, 18135678, 18774543, 23294548, 16785708, 10136473, 24560906 19777862, 22551446, 17786518, 25879984, 18315328, 18334586, 12747740 18096714, 19032867, 21641760, 17390160, 18899974, 17232014, 20598042 26679352, 26245237, 16354467, 17484731, 18673325, 16422541, 18155762 19827973, 14015842, 22683225, 17726838, 18554871, 23177648, 18051556 20803583, 18282562, 21972320, 15990359, 17922254, 16855292, 16668584 21343838, 20299015, 17446237, 18043064, 18093615, 23713236, 17694209 17288409, 20475845, 18308268, 17274537, 13955826, 16934803, 17634921 17501491, 16315398, 23725036, 22683212, 17006183, 13829543, 18191164 26746894, 22809871, 17655240, 28819280, 18384391, 19393542, 21538567 16198143, 21847223, 25823754, 17892268, 20142975, 19584068, 17165204 25165496, 18604493, 21756699, 18508861, 18554763, 16901385, 21532755 18189036, 17443671, 17385178, 14829250, 17936109, 20925795, 20509482 17478514, 27441326, 16850630, 13951456, 16595641, 14054676, 15861775 21142837, 16912439, 17299889, 17297939, 23003979, 16833527, 18619917 17798953, 19697993, 17816865, 25914276, 18607546, 17571306, 21286665 16898135, 16819962, 17341326, 26910644, 17851160, 17586955, 20558005 19049453, 21051840, 17587063, 16956380, 18328509, 25423453, 14133975 29033139, 19718981, 18061914, 18522509, 17518652, 21051833, 20294666 18765602, 20860659, 18272672, 20324049, 18199537, 17332800, 13609098 22502493, 18384537, 14338435, 17945983, 16392068, 21067387, 17752995 21051862, 16863422, 17237521, 25505382, 18244962, 19544839, 24433711 24717859, 17156148, 18973907, 23026585, 17449815, 17877323, 18180390 17088068, 17037130, 20004087, 21422580, 19466309, 11733603, 25505371 18610915, 21051858, 18084625, 29027694, 18674024, 26243698, 21051852 18091059, 18306996, 25369547, 16306373, 17787259, 18193833, 19915271 20513399, 20631274, 26439748, 16344544, 25879656, 14692762, 18614015 22782647, 17346091, 18228645, 17721717, 13960236, 18685892, 18436307 11883252, 19888853, 21756677, 17891943, 19475971, 22353199, 16384983 19121551, 25634317, 27825893, 12816846, 17982555, 17761775, 17227073 22243719, 17265217, 25505394, 17071721, 16721594, 18262334, 21756661 17891946, 15913355, 17672719, 17602269, 17239687, 17042658, 25555252 17238511, 21059919,

17811456, 17284817, 17752121, 20879889, 28806384 21380789, 17394950, 17011832, 16579084, 22195465, 14602788, 28790634 18325460, 27567477, 26569225, 24476265, 24476274, 12611721, 18674465 16903536, 17006570, 19689979, 28076295, 16043574, 18783224, 24662775 16494615, 21526048, 17392698, 19197175, 16069901, 17811447, 27870645 28876684, 17308789, 24835538, 22195477, 17865671, 17343514, 19013183 17325413, 18316692, 16180763, 17348614, 14368995, 21983325, 17393915 16285691, 19211433, 20331945, 17883081, 17705023, 24316947, 17614227 19578350, 22195485, 14084247, 24975421, 26078387, 23115139, 13645875 16777840, 19727057, 14852021, 18744139, 18674047, 17716305, 19285025 18482502, 17622427, 19289642, 27534509, 25947799, 22195492, 14458214 20869721, 21172913, 17767676, 18723434, 25505407, 17786278, 19258504 17082983, 17365043, 21351877, 13498382, 18331812, 16065166, 25489607 16685417, 18031668, 22893153, 17551674, 16943711, 19272701, 21517440 25897615, 17649265, 13866822, 18094246, 24528741, 17783588, 14245531 17082359, 26007010, 18280813, 20448824, 23330119, 16268425, 19487147 25600421, 18018515, 17302277, 17215560, 24411921, 19271443, 25764020 14176370, 17016369, 20777150, 16756406, 23330124, 22977256, 20441797 19769489, 28100487, 17545847, 25093656, 18260550, 13853126, 17227277 17551699, 23536835, 25957038, 24652769, 20725343, 19207117, 9756271 18868646, 17614134, 26667023, 17546973, 19680952, 18704244, 26667015 17050888, 18273830, 18828868, 17360606, 16992075, 24563422, 17375354 12905058, 18362222, 21429602, 17254374, 26667032, 28364007, 27086138 17571039, 17468141, 18436647, 17570606, 17235750, 21168487, 17279227 16220077, 16929165

Version 11.2.0.4.v19

Version 11.2.0.4.v19 adds support for the following:

- Patch 28729262: Oracle Database Patch Set Update 11.2.0.4.190115
- Patch 28790660: Oracle JVM Patch Set Update 11.2.0.4.190115
- Patch 28125601: DSTv32 for RDBMS (TZDATA2018E)
- Patch 27015468: DSTv32 for OJVM (TZDATA2018E)
- Patch 27216420: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches

Oracle patch 28729262, released January 2019

Bugs fixed:

17288409, 21051852, 24316947, 17811429, 17205719, 18607546, 25654936 17484762, 17816865, 20506699, 24835538, 25957038, 19692824, 23330119 17922254, 17754782, 13364795, 16934803, 17311728, 18604692, 20387265 26679352, 17284817, 17441661, 20671094, 24560906, 16992075, 25635149 17446237, 14015842, 19972569, 21756677, 17375354, 17449815, 20925795 21538558, 17019086, 19463897, 26575788, 13866822, 17235750, 17982555 17478514, 18317531, 14338435, 18235390, 19461270, 20803583, 13944971 19475971, 20142975, 17811789, 16929165, 18704244, 24662775, 20506706 17546973, 21422580, 17359610, 20334344, 14054676, 25489607, 17088068 17570606, 18264060, 17346091, 17343514, 21538567, 19680952, 18471685 19211724, 21132297, 23105538, 13951456, 25775213, 16315398, 21847223 18744139, 16850630, 23177648, 19049453, 18673304, 18090142, 17883081 19915271, 18641419, 18262334, 25600421, 17006183, 16065166, 18277454 18685892, 16833527, 10136473, 18051556, 17865671, 25879984, 18554871 17852463, 18774543, 17853498, 18334586, 19487147, 20879889, 17551709 17588480, 19827973, 17344412, 17842825, 18828868, 20509482, 17025461 13609098, 19429927, 26039623, 11883252, 17239687, 16410570, 23007241 17602269, 19197175, 22195457, 18316692, 17313525, 12611721, 21174504 19544839, 18964939, 20294666, 17600719, 26667015, 18191164, 17571306 19393542, 18482502, 20777150, 27086138, 19466309, 22243719, 17165204 17040527, 18098207, 16785708, 24790914, 19891090, 17465741, 16180763 17174582, 12982566, 16777840, 27097854, 19463893, 22195465, 16875449 22148226, 12816846, 17237521, 6599380, 19358317, 17811438, 25505394 17811447, 21983325, 17945983, 18762750, 16912439, 17184721, 18061914 20598042, 26631046, 21380789, 17282229, 18948177, 18331850, 21142837 18202441, 17082359, 18723434, 19554106, 21532755, 21972320, 25505371 20273319, 14034426, 18339044, 19458377, 17752995, 20448824, 17891943 17767676, 17258090, 16668584, 18384391, 21063322, 17040764, 17381384 15913355, 18356166,

14084247, 20596234, 21641760, 20506715, 13853126 21756661, 18203837, 18610915, 14245531, 16043574, 21756699, 22195441 17848897, 17877323, 26667032, 28790634, 19272701, 21453153, 20569094 17468141, 17786518, 20861693, 17912217, 17037130, 16956380, 18155762 17478145, 17394950, 18189036, 18641461, 17551674, 18619917, 17027426 17019356, 21352646, 16268425, 24476274, 22195492, 19584068, 26544823 18436307, 22507210, 17265217, 13498382, 17634921, 19469538, 21526048 19258504, 23003979, 16354467, 18043064, 19174430, 20004087, 17443671 22195485, 18000422, 22321756, 20004021, 17571039, 27053456, 25897615 16832076, 21067387, 22905130, 16344544, 21429602, 18009564, 14354737 18135678, 21286665, 18614015, 14521849, 20441797, 28876684, 18362222 25655390, 16472716, 17835048, 17050888, 17936109, 14010183, 17325413 18747196, 19207156, 17231779, 21842740, 17761775, 16721594, 17082983 20067212, 21179898, 17279227, 17302277, 18084625, 20717359, 24624166 15990359, 24842886, 26746894, 18203835, 23026585, 17297939, 17811456 16731148, 22380919, 21168487, 14133975, 17215560, 13829543, 18740837 17694209, 17385178, 18091059, 8322815, 18259031, 17586955, 19689979 25165496, 28254374, 17201159, 17655634, 18331812, 19730508, 17551699 17648596, 18868646, 16220077, 16069901, 17393915, 17348614, 17957017 17274537, 18096714, 17308789, 18436647, 14285317, 19289642, 14764829 17622427, 18328509, 23115139, 16943711, 22195477, 22502493, 14368995 17346671, 18996843, 17783588, 18604493, 21343838, 16618694, 17672719 18856999, 18783224, 17851160, 17546761, 22168163, 17798953, 18273830 22092979, 16596890, 19972566, 13871092, 20828947, 26667023, 17726838 16384983, 22296366, 17360606, 13645875, 22321741, 16542886, 18199537 25879656, 25634317, 21787056, 23140259, 17889549, 21172913, 14565184 26245237, 20475845, 27825893, 17071721, 21281607, 17610798, 18308268 20299015, 21343897, 22893153, 22594718, 28076295, 20657441, 17397545 18230522, 16360112, 19769489, 12905058, 18641451, 12747740, 18430495 25423453, 17016369, 17042658, 14602788, 17551063, 26243698, 19972568 21517440, 23725036, 19788842, 18508861, 14657740, 17332800, 13837378 17186905, 19972564, 17019345, 19699191, 18315328, 27441326, 17437634 24570598, 22353199, 18093615, 19006849, 28806384, 17392698, 19013183 17296856, 18674024, 26569225, 17232014, 16855292, 21051840, 14692762 17762296, 17705023, 23294548, 22351572, 22507234, 19121551, 20324049 21330264, 26198926, 19854503, 23315889, 26030218, 26910644, 21868720 19309466, 27567477, 25764020, 18681862, 17365043, 17390160, 18554763 20031873, 20558005, 24717859, 21795111, 18456514, 13955826, 16306373 18139690, 17501491, 17752121, 17299889, 21668627, 23713236, 24652769 17889583, 18673325, 22551446, 17242746, 18293054, 18674465, 19721304 19211433, 19888853, 25914276, 24563422, 17951233, 18094246, 17649265 19615136, 17011832, 17477958, 16870214, 18522509, 20631274, 16091637 17323222, 16595641, 16524926, 18228645, 17484731, 18282562, 17596908 18272672, 18031668, 17156148, 16494615, 22683225, 20869721, 17545847 25093656, 28819280, 18682983, 17655240, 24528741, 17614134, 13558557 25427662, 17341326, 22465352, 29033139, 17891946, 17716305, 22657942 27374796, 16392068, 18440095, 19271443, 21351877, 20513399, 18092127 17614227, 18440047, 18849970, 14106803, 16903536, 20725343, 18973907 18673342, 17389192, 25505382, 22809871, 19032867, 17612828, 17006570 16194160, 25369547, 25505407, 16685417, 17721717, 21354456, 17390431 17570240, 16863422, 13960236, 28100487, 18325460, 17008068, 19727057 28855981, 16422541, 17267114, 19972570, 18244962, 21538485, 18203838 18765602, 16198143, 17246576, 14829250, 28364007, 17835627, 20860659 21629064, 18247991, 14458214, 21051862, 17786278, 16692232, 17227277 24348685, 16042673, 24476265, 24975421, 22901797, 16314254, 19285025 16228604, 16756406, 14176370, 16837842, 20144308, 17393683, 23536835 25823754, 18899974, 17787259, 24719736, 20331945, 26078387, 20074391 19490948, 15861775, 16399083, 25947799, 25555252, 18018515, 22683212 18260550, 21051858, 17080436, 16613964, 17036973, 16579084, 24433711 18384537, 27870645, 18280813, 20296213, 16901385, 15979965, 17518652 23330124, 20856766, 18441944, 16450169, 9756271, 27534509, 22730454 19718981, 17291347, 17892268, 11733603, 16285691, 17587063, 21343775 18180390, 16538760, 26474853, 18193833, 21387964, 21051833, 17238511 19777862, 17824637, 23065323, 21656630, 19697993, 17903598, 16571443 18306996, 18166013, 19578350, 14852021, 18674047, 17853456, 12364061 24411921, 19207117, 22195448

Version 11.2.0.4.v18

Version 11.2.0.4.v18 adds support for the following:

- Patch 28204707: Oracle Database Patch Set Update 11.2.0.4.181016

- Patch 28440700: Oracle JVM Patch Set Update 11.2.0.4.181016
- Patch 28125601: DSTv32 for RDBMS (TZDATA2018E)
- Patch 27015468: DSTv32 for OJVM (TZDATA2018E)
- Patch 27216420: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patches 27659043 and 19692824 are now included in the Database Patch Set Update

Oracle patch 28204707, released October 2018

Bugs fixed:

17288409, 21051852, 24316947, 17811429, 17205719, 18607546, 25654936 17484762, 17816865, 20506699, 24835538, 25957038, 19692824, 23330119 17922254, 17754782, 13364795, 16934803, 17311728, 18604692, 26679352 20387265, 17284817, 17441661, 20671094, 24560906, 25635149, 16992075 17446237, 14015842, 19972569, 21756677, 17375354, 21538558, 20925795 17449815, 17019086, 19463897, 26575788, 13866822, 17235750, 17982555 17478514, 18317531, 14338435, 18235390, 19461270, 20803583, 13944971 19475971, 20142975, 17811789, 16929165, 18704244, 24662775, 20506706 17359610, 17546973, 21422580, 20334344, 14054676, 25489607, 17570606 17088068, 17346091, 18264060, 17343514, 21538567, 19680952, 18471685 19211724, 21132297, 25775213, 13951456, 16315398, 21847223, 18744139 16850630, 23177648, 19049453, 18090142, 18673304, 17883081, 19915271 18641419, 18262334, 25600421, 17006183, 16065166, 18277454, 18685892 16833527, 10136473, 18051556, 17865671, 25879984, 18554871, 17852463 18774543, 17853498, 18334586, 19487147, 20879889, 17551709, 17588480 19827973, 17344412, 17842825, 18828868, 20509482, 17025461, 26039623 19429927, 13609098, 11883252, 16410570, 17239687, 23007241, 17602269 19197175, 22195457, 18316692, 17313525, 12611721, 21174504, 19544839 20294666, 18964939, 17600719, 26667015, 18191164, 17571306, 19393542 20777150, 18482502, 27086138, 19466309, 22243719, 17165204, 17040527 18098207, 24790914, 16785708, 19891090, 17465741, 16180763, 17174582 12982566, 16777840, 19463893, 22195465, 16875449, 22148226, 12816846 17237521, 6599380, 19358317, 17811438, 25505394, 17811447, 21983325 17945983, 18762750, 16912439, 17184721, 18061914, 20598042, 26631046 21380789, 17282229, 18948177, 18331850, 21142837, 18202441, 17082359 18723434, 21972320, 21532755, 19554106, 25505371, 20273319, 14034426 18339044, 19458377, 17752995, 20448824, 17891943, 17767676, 17258090 16668584, 18384391, 21063322, 17040764, 17381384, 15913355, 18356166 14084247, 20596234, 21641760, 20506715, 13853126, 21756661, 18610915 18203837, 14245531, 16043574, 21756699, 22195441, 17848897, 17877323 26667032, 21453153, 19272701, 20569094, 17468141, 17786518, 20861693 17912217, 17037130, 16956380, 18155762, 17478145, 17394950, 18641461 18189036, 17551674, 18619917, 17019356, 17027426, 21352646, 16268425 24476274, 22195492, 19584068, 26544823, 18436307, 22507210, 17265217 13498382, 17634921, 19469538, 21526048, 19258504, 23003979, 16354467 18043064, 19174430, 20004087, 17443671, 22195485, 18000422, 22321756 20004021, 17571039, 25897615, 27053456, 16832076, 21067387, 22905130 16344544, 21429602, 18009564, 14354737, 21286665, 18135678, 14521849 18614015, 20441797, 18362222, 25655390, 16472716, 17835048, 17050888 17936109, 14010183, 17325413, 18747196, 19207156, 17231779, 21842740 17761775, 16721594, 17082983, 20067212, 21179898, 17279227, 17302277 18084625, 20717359, 24624166, 15990359, 24842886, 26746894, 18203835 23026585, 17297939, 17811456, 16731148, 22380919, 21168487, 14133975 13829543, 17215560, 18740837, 17694209, 17385178, 18091059, 8322815 18259031, 28254374, 19689979, 25165496, 17586955, 17201159, 17655634 18331812, 17551699, 19730508, 17648596, 18868646, 16220077, 16069901 17393915, 17348614, 17957017, 17274537, 18096714, 17308789, 18436647 14285317, 19289642, 14764829, 17622427, 18328509, 23115139, 16943711 22195477, 22502493, 14368995, 17346671, 18996843, 17783588, 18604493 21343838, 16618694, 17672719, 18856999, 18783224, 17851160, 17546761 22168163, 17798953, 18273830, 22092979, 16596890, 19972566, 13871092 20828947, 26667023, 17726838, 16384983, 22296366, 17360606, 13645875 22321741, 25634317, 16542886, 18199537, 25879656, 21787056, 23140259 17889549, 21172913, 26245237, 14565184, 27825893, 20475845, 17071721 21281607, 17610798, 18308268, 20299015, 21343897, 22893153, 22594718 20657441, 17397545, 18230522, 16360112, 19769489, 12905058, 18641451 12747740, 18430495, 25423453, 17016369, 17042658, 14602788, 17551063 26243698, 19972568, 21517440, 23725036, 19788842, 18508861, 14657740 17332800, 13837378, 17186905, 19972564, 17019345, 19699191, 18315328 27441326,

17437634, 24570598, 22353199, 18093615, 19006849, 17392698 19013183, 17296856, 18674024, 26569225, 17232014, 16855292, 21051840 14692762, 17762296, 17705023, 23294548, 22351572, 22507234, 19121551 20324049, 21330264, 26198926, 19854503, 23315889, 26910644, 26030218 21868720, 19309466, 25764020, 18681862, 17365043, 17390160, 20031873 20558005, 18554763, 24717859, 21795111, 18456514, 13955826, 16306373 18139690, 17501491, 17752121, 17299889, 21668627, 23713236, 24652769 17889583, 18673325, 22551446, 18674465, 17242746, 19721304, 18293054 19211433, 19888853, 25914276, 24563422, 17951233, 18094246, 17649265 19615136, 17011832, 17477958, 16870214, 18522509, 20631274, 16091637 17323222, 16595641, 16524926, 17484731, 18228645, 18282562, 17596908 18272672, 18031668, 17156148, 16494615, 22683225, 20869721, 17545847 25093656, 18682983, 17655240, 24528741, 17614134, 25427662, 13558557 17341326, 22465352, 17891946, 17716305, 22657942, 27374796, 16392068 18440095, 19271443, 21351877, 20513399, 18092127, 17614227, 18440047 18849970, 16903536, 14106803, 20725343, 18973907, 18673342, 17389192 19032867, 25505382, 22809871, 17612828, 17006570, 16194160, 25369547 25505407, 16685417, 17721717, 21354456, 17390431, 17570240, 13960236 16863422, 28100487, 18325460, 17008068, 19727057, 16422541, 17267114 19972570, 18244962, 21538485, 18203838, 18765602, 16198143, 17246576 14829250, 28364007, 17835627, 20860659, 21629064, 18247991, 14458214 21051862, 17786278, 16692232, 17227277, 24348685, 24476265, 16042673 24975421, 22901797, 16314254, 19285025, 16228604, 16756406, 14176370 16837842, 20144308, 17393683, 23536835, 25823754, 18899974, 17787259 24719736, 20331945, 26078387, 19490948, 20074391, 15861775, 16399083 25555252, 25947799, 18018515, 22683212, 18260550, 21051858, 17080436 16613964, 17036973, 16579084, 24433711, 18384537, 27870645, 18280813 20296213, 16901385, 15979965, 17518652, 23330124, 20856766, 18441944 16450169, 9756271, 27534509, 22730454, 19718981, 17291347, 17892268 11733603, 16285691, 17587063, 21343775, 18180390, 26474853, 16538760 18193833, 21387964, 21051833, 17238511, 19777862, 17824637, 23065323 21656630, 17903598, 16571443, 18166013, 18306996, 19578350, 14852021 17853456, 18674047, 12364061, 24411921, 19207117, 22195448

Version 11.2.0.4.v17

Version 11.2.0.4.v17 adds support for the following:

- Patch 27734982: Oracle Database Patch Set Update 11.2.0.4.180717
- Patch 27923163: Oracle JVM Patch Set Update 11.2.0.4.180717
- Patch 28125601: DSTv32 for RDBMS (TZDATA2018E)
- Patch 27015468: DSTv32 for OJVM (TZDATA2018E)
- Patch 27216420: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 27659043: MES Bundle 405
- Patch 19692824: DBCONTROL is not coming up on OEL 7

Oracle patch 27734982, released July 2018

Bugs fixed:

17288409, 21051852, 24316947, 17811429, 17205719, 18607546, 25654936 17816865, 20506699, 24835538, 25957038, 23330119, 17922254, 17754782 13364795, 16934803, 17311728, 20387265, 17284817, 17441661, 20671094 24560906, 16992075, 17446237, 14015842, 19972569, 21756677, 17375354 21538558, 20925795, 17449815, 19463897, 26575788, 13866822, 17235750 17982555, 17478514, 18317531, 14338435, 18235390, 19461270, 20803583 13944971, 19475971, 20142975, 17811789, 16929165, 18704244, 24662775 20506706, 17546973, 21422580, 20334344, 14054676, 25489607, 17088068 17346091, 18264060, 17343514, 21538567, 19680952, 18471685, 19211724 21132297, 13951456, 16315398, 21847223, 18744139, 16850630, 23177648 19049453, 18090142, 18673304, 17883081, 19915271, 18641419, 18262334 25600421, 17006183, 16065166, 18277454, 16833527, 10136473, 18051556 17865671, 18554871, 17852463, 18774543, 17853498, 18334586, 19487147 20879889, 17551709, 17588480, 19827973, 17344412, 17842825, 18828868 20509482, 17025461, 13609098, 11883252, 17239687, 23007241, 17602269 19197175, 22195457, 18316692,

17313525, 12611721, 21174504, 19544839 20294666, 18964939, 17600719, 26667015, 18191164, 17571306, 19393542 20777150, 18482502, 27086138, 19466309, 22243719, 17165204, 17040527 18098207, 16785708, 17465741, 16180763, 17174582, 12982566, 16777840 19463893, 22195465, 16875449, 22148226, 12816846, 17237521, 6599380 19358317, 17811438, 25505394, 17811447, 21983325, 17945983, 18762750 16912439, 17184721, 18061914, 20598042, 21380789, 17282229, 18948177 18331850, 21142837, 18202441, 17082359, 18723434, 21972320, 21532755 19554106, 25505371, 14034426, 18339044, 19458377, 17752995, 20448824 17891943, 17767676, 17258090, 16668584, 18384391, 17040764, 17381384 15913355, 18356166, 14084247, 20596234, 21641760, 20506715, 13853126 21756661, 18203837, 14245531, 16043574, 21756699, 22195441, 17848897 17877323, 21453153, 19272701, 20569094, 17468141, 17786518, 20861693 17912217, 17037130, 16956380, 18155762, 17478145, 17394950, 18641461 18189036, 18619917, 17027426, 21352646, 16268425, 24476274, 22195492 19584068, 26544823, 18436307, 22507210, 17265217, 13498382, 17634921 19469538, 21526048, 19258504, 23003979, 18043064, 19174430, 20004087 17443671, 22195485, 18000422, 20004021, 22321756, 17571039, 27053456 25897615, 21067387, 16832076, 22905130, 16344544, 21429602, 18009564 14354737, 21286665, 18135678, 14521849, 18614015, 20441797, 18362222 25655390, 16472716, 17835048, 17050888, 17936109, 14010183, 17325413 18747196, 17761775, 16721594, 17082983, 20067212, 21179898, 17302277 18084625, 20717359, 24624166, 15990359, 24842886, 26746894, 18203835 23026585, 17297939, 17811456, 16731148, 22380919, 21168487, 14133975 13829543, 17215560, 17694209, 17385178, 18091059, 8322815, 18259031 19689979, 25165496, 17586955, 17201159, 17655634, 18331812, 19730508 17648596, 18868646, 16220077, 16069901, 17393915, 17348614, 17957017 17274537, 18096714, 17308789, 18436647, 14285317, 19289642, 14764829 17622427, 18328509, 16943711, 22195477, 22502493, 14368995, 17346671 18996843, 17783588, 18604493, 21343838, 16618694, 17672719, 18856999 18783224, 17851160, 17546761, 22168163, 17798953, 18273830, 22092979 16596890, 19972566, 20828947, 13871092, 26667023, 17726838, 16384983 22296366, 17360606, 13645875, 22321741, 16542886, 18199537, 25879656 21787056, 17889549, 21172913, 14565184, 27825893, 20475845, 17071721 21281607, 18308268, 17610798, 20299015, 21343897, 22893153, 20657441 17397545, 18230522, 16360112, 19769489, 12905058, 18641451, 12747740 18430495, 25423453, 17016369, 17042658, 14602788, 17551063, 19972568 21517440, 23725036, 19788842, 18508861, 14657740, 17332800, 13837378 17186905, 19972564, 19699191, 18315328, 27441326, 17437634, 24570598 22353199, 18093615, 19006849, 17392698, 19013183, 17296856, 18674024 26569225, 17232014, 16855292, 21051840, 14692762, 17762296, 17705023 23294548, 22351572, 22507234, 19121551, 20324049, 21330264, 26198926 19854503, 23315889, 26910644, 26030218, 21868720, 19309466, 25764020 18681862, 17365043, 17390160, 20031873, 20558005, 18554763, 24717859 21795111, 18456514, 16306373, 13955826, 18139690, 17501491, 17752121 17299889, 21668627, 23713236, 24652769, 17889583, 18673325, 22551446 17242746, 19721304, 18293054, 19211433, 19888853, 24563422, 17951233 18094246, 17649265, 19615136, 17011832, 17477958, 16870214, 18522509 20631274, 16091637, 17323222, 16595641, 16524926, 18228645, 18282562 17596908, 18031668, 17156148, 16494615, 22683225, 20869721, 17545847 25093656, 17655240, 24528741, 17614134, 25427662, 13558557, 17341326 22465352, 17891946, 17716305, 22657942, 27374796, 16392068, 18440095 19271443, 21351877, 20513399, 18092127, 17614227, 18440047, 18849970 16903536, 14106803, 18973907, 18673342, 17389192, 19032867, 25505382 22809871, 17612828, 17006570, 16194160, 25369547, 25505407, 16685417 17721717, 21354456, 17390431, 17570240, 16863422, 28100487, 18325460 17008068, 19727057, 16422541, 17267114, 19972570, 18244962, 21538485 18203838, 18765602, 16198143, 17246576, 14829250, 17835627, 20860659 21629064, 18247991, 14458214, 21051862, 17786278, 16692232, 24348685 17227277, 24476265, 16042673, 16314254, 19285025, 16228604, 16756406 16837842, 20144308, 17393683, 23536835, 25823754, 18899974, 17787259 24719736, 20331945, 19490948, 20074391, 15861775, 16399083, 25947799 18018515, 22683212, 18260550, 21051858, 17080436, 16613964, 17036973 16579084, 24433711, 18384537, 27870645, 18280813, 20296213, 16901385 15979965, 23330124, 18441944, 16450169, 27534509, 9756271, 17892268 11733603, 16285691, 17587063, 21343775, 18180390, 26474853, 16538760 18193833, 21387964, 21051833, 17238511, 19777862, 17824637, 23065323 17903598, 16571443, 18306996, 19578350, 14852021, 17853456, 18674047 12364061, 24411921, 19207117, 22195448

Version 11.2.0.4.v16

Version 11.2.0.4.v16 adds support for the following:

- Patch 27338049: DATABASE PATCH SET UPDATE 11.2.0.4.180417
- Patch 27475598: OJVM PATCH SET UPDATE 11.2.0.4.180417
- Patch 27015449: RDBMS - PROACTIVE DSTV31 UPDATE - TZDATA2017C
- Patch 27015468: PROACTIVE DSTV31 UPDATE - TZDATA2017C - NEED OJVM FIX
- Patch 27216420: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 27659043: MES 405 BUNDLE ON TOP OF RDBMS 11.2.0.4.180116 PSU
- Patch 19692824: DBCONTROL is not coming up on OEL 7
- Adds support for the `DBMS_ADVANCED_REWRITE` package
- Fixed a bug where `DBA_LOCKS` and associated views available in new DB instances of 11.2.0.4.v15 were not created in upgrades to 11.2.0.4.v15. Views are now created in new and upgraded DB instances of 11.2.0.4.v16 and later.

Oracle patch 27338049, released April 2018

Bugs fixed:

21174504, 17184721, 21538558, 16091637, 18092127, 17381384, 15979965 20671094, 16731148, 16314254, 13837378, 18441944, 17835048, 13558557 17008068, 17201159, 25427662, 17853498, 20717359, 17246576, 18356166 18681862, 18440047, 20569094, 20031873, 16875449, 20387265, 19788842 17296856, 21330264, 14010183, 17648596, 17551063, 17025461, 24719736 17267114, 22507210, 17912217, 17889583, 18202441, 17040764, 17478145 16524926, 25655390, 19358317, 22148226, 18747196, 26544823, 18641419 17036973, 18948177, 17811789, 16542886, 14285317, 18009564, 16618694 8322815, 16832076, 18247991, 16692232, 22507234, 17570240, 13871092 24624166, 17848897, 17441661, 14034426, 17465741, 16596890, 17437634 21343897, 20506706, 21453153, 18339044, 22321741, 21795111, 17951233 18430495, 21787056, 22380919, 19469538, 20506715, 17811429, 19721304 17903598, 18230522, 19554106, 19458377, 21281607, 17612828, 6599380 22092979, 22321756, 17040527, 17811438, 18641461, 14657740, 13364795 21387964, 19490948, 22351572, 17346671, 17588480, 18235390, 26474853 18849970, 17889549, 19309466, 16472716, 20596234, 18331850, 18641451 17344412, 21179898, 19461270, 17546761, 24842886, 14521849, 18203835 18203838, 18964939, 18203837, 17313525, 22195457, 18139690, 16837842 22296366, 14106803, 17842825, 21352646, 22657942, 16360112, 20657441 22195441, 17389192, 26198926, 14565184, 17205719, 18440095, 14764829 22195448, 14354737, 13944971, 16571443, 21868720, 17186905, 17080436 18673342, 22905130, 17027426, 27374796, 19972569, 19972568, 20144308 19972566, 17282229, 19972564, 16870214, 21629064, 19615136, 21354456 17390431, 18762750, 23007241, 16613964, 17957017, 18098207, 18471685 19730508, 21538485, 18264060, 17323222, 17754782, 17600719, 18317531 17852463, 17596908, 17655634, 16228604, 27053456, 20074391, 19972570 18090142, 18996843, 19854503, 16042673, 17835627, 20334344, 17393683 20861693, 18000422, 17551709, 26575788, 23315889, 20506699, 19006849 18277454, 18456514, 19174430, 17258090, 17174582, 25654936, 17242746 16399083, 17824637, 21132297, 22465352, 17762296, 22168163, 17397545 16450169, 12364061, 20067212, 18856999, 19211724, 19463893, 19463897 21343775, 17853456, 18673304, 20004021, 26030218, 21668627, 16194160 17477958, 16538760, 12982566, 24570598, 20828947, 18259031, 20296213 18293054, 17610798, 19699191, 23065323, 17311728, 18135678, 18774543 23294548, 16785708, 10136473, 24560906, 22551446, 19777862, 17786518 18315328, 18334586, 12747740, 18096714, 19032867, 21641760, 18899974 17390160, 17232014, 20598042, 18673325, 16422541, 18155762, 14015842 19827973, 22683225, 17726838, 18554871, 23177648, 18051556, 20803583 21972320, 15990359, 17922254, 18282562, 16855292, 16668584, 21343838 20299015, 17446237, 18093615, 18043064, 23713236, 17694209, 17288409 20475845, 17274537, 13955826, 16934803, 17634921, 17501491, 16315398 22683212, 17006183, 13829543, 18191164, 17655240, 26746894, 22809871 18384391, 19393542, 21538567, 16198143, 21847223, 25823754, 17892268 20142975, 19584068, 17165204, 25165496, 18604493, 21756699, 18508861 16901385, 18554763, 21532755, 18189036, 17443671, 17385178, 14829250 17936109, 20925795, 20509482, 17478514, 27441326, 16850630, 13951456 16595641, 14054676, 15861775, 21142837, 16912439, 17299889, 17297939 23003979, 18619917, 16833527, 17798953, 17816865, 18607546, 17571306 21286665, 17341326, 26910644, 17851160, 20558005, 17586955, 19049453 21051840, 17587063, 16956380, 18328509, 25423453, 14133975, 18061914 18522509,

21051833, 18765602, 20860659, 20324049, 18199537, 17332800 13609098, 22502493, 18384537, 14338435, 17945983, 16392068, 21067387 17752995, 21051862, 16863422, 25505382, 17237521, 18244962, 19544839 24433711, 24717859, 17156148, 18973907, 23026585, 17877323, 17449815 18180390, 17088068, 17037130, 20004087, 21422580, 19466309, 11733603 25505371, 21051858, 18084625, 18674024, 21051852, 18091059, 25369547 16306373, 18306996, 18193833, 19915271, 17787259, 20513399, 20631274 25879656, 16344544, 14692762, 18614015, 17346091, 18228645, 17721717 18436307, 21756677, 19888853, 11883252, 17891943, 19475971, 22353199 16384983, 19121551, 12816846, 17982555, 17761775, 22243719, 17265217 25505394, 17071721, 16721594, 21756661, 18262334, 17891946, 15913355 17672719, 17602269, 17239687, 17042658, 17238511, 17811456, 17284817 17752121, 20879889, 21380789, 17394950, 17011832, 16579084, 22195465 14602788, 18325460, 24476265, 26569225, 24476274, 12611721, 16903536 17006570, 19689979, 16043574, 18783224, 24662775, 16494615, 21526048 17392698, 19197175, 16069901, 17811447, 17308789, 22195477, 24835538 17865671, 17343514, 19013183, 17325413, 18316692, 16180763, 17348614 14368995, 21983325, 17393915, 16285691, 19211433, 20331945, 17883081 17705023, 24316947, 17614227, 19578350, 22195485, 14084247, 13645875 16777840, 19727057, 14852021, 18744139, 18674047, 17716305, 19285025 18482502, 17622427, 19289642, 22195492, 25947799, 14458214, 20869721 21172913, 17767676, 18723434, 25505407, 17786278, 19258504, 17082983 21351877, 17365043, 13498382, 18331812, 16065166, 25489607, 16685417 18031668, 22893153, 16943711, 19272701, 21517440, 25897615, 17649265 13866822, 18094246, 24528741, 17783588, 14245531, 17082359, 18280813 20448824, 23330119, 16268425, 19487147, 25600421, 18018515, 17302277 17215560, 24411921, 19271443, 25764020, 17016369, 20777150, 23330124 16756406, 20441797, 19769489, 17545847, 25093656, 18260550, 13853126 17227277, 23536835, 25957038, 24652769, 19207117, 9756271, 18868646 17614134, 26667023, 17546973, 18704244, 19680952, 26667015, 17050888 18828868, 18273830, 17360606, 24563422, 16992075, 17375354, 12905058 18362222, 21429602, 27086138, 17571039, 17468141, 18436647, 17235750 21168487, 16220077, 16929165

Version 11.2.0.4.v15

Version 11.2.0.4.v15 adds support for the following:

- Patch 26925576: DATABASE PATCH SET UPDATE 11.2.0.4.180116
- Patch 26925532: OJVM PATCH SET UPDATE 11.2.0.4.180116
- Patch 27015449: RDBMS - PROACTIVE DSTV31 UPDATE - TZDATA2017C
- Patch 27015468: PROACTIVE DSTV31 UPDATE - TZDATA2017C - NEED OJVM FIX
- Patch 27216420: Oracle GoldenGate – Oracle RDBMS Server Recommended Patches
- Patch 27244661: MES 405 BUNDLE ON TOP OF RDBMS 11.2.0.4.180116 PSU
- Patch 19692824: DBCONTROL is not coming up on OEL 7
- Adds support for DBA_LOCKS and associated views

Oracle patch 26925576, released January 2018

Bugs fixed:

17288409, 21051852, 24316947, 17811429, 17205719, 18607546, 25654936 17816865, 20506699, 24835538, 25957038, 23330119, 17922254, 17754782 13364795, 16934803, 17311728, 20387265, 17284817, 17441661, 20671094 24560906, 16992075, 17446237, 14015842, 19972569, 21756677, 17375354 21538558, 20925795, 17449815, 26575788, 19463897, 13866822, 17235750 17982555, 17478514, 18317531, 14338435, 18235390, 20803583, 19461270 19475971, 13944971, 20142975, 17811789, 16929165, 18704244, 24662775 20506706, 21422580, 17546973, 20334344, 14054676, 25489607, 17088068 17346091, 18264060, 17343514, 21538567, 19680952, 18471685, 19211724 21132297, 13951456, 16315398, 21847223, 18744139, 16850630, 23177648 19049453, 18090142, 18673304, 17883081, 19915271, 18641419, 18262334 25600421, 17006183, 16065166, 18277454, 16833527, 10136473, 18051556 17865671, 18554871, 17852463, 17853498, 18334586, 20879889, 17551709 17588480, 19827973, 17344412, 17842825, 18828868, 20509482, 17025461 13609098,

11883252, 17239687, 23007241, 17602269, 19197175, 18316692 22195457, 17313525, 12611721, 21174504, 19544839, 18964939, 17600719 26667015, 18191164, 17571306, 19393542, 20777150, 18482502, 19466309 22243719, 17165204, 17040527, 18098207, 16785708, 17465741, 16180763 17174582, 12982566, 16777840, 19463893, 22195465, 16875449, 22148226 12816846, 17237521, 6599380, 19358317, 17811438, 25505394, 17811447 21983325, 17945983, 18762750, 16912439, 17184721, 20598042, 18061914 21380789, 17282229, 18948177, 18331850, 21142837, 18202441, 17082359 18723434, 21972320, 21532755, 19554106, 25505371, 14034426, 18339044 19458377, 17752995, 20448824, 17891943, 17767676, 17258090, 16668584 18384391, 17040764, 17381384, 15913355, 18356166, 14084247, 20596234 21641760, 20506715, 13853126, 21756661, 18203837, 14245531, 16043574 21756699, 22195441, 17848897, 17877323, 21453153, 19272701, 20569094 17468141, 17786518, 20861693, 17912217, 17037130, 16956380, 18155762 17478145, 17394950, 18641461, 18189036, 18619917, 17027426, 21352646 16268425, 24476274, 22195492, 19584068, 26544823, 18436307, 22507210 17265217, 13498382, 17634921, 19469538, 21526048, 19258504, 23003979 19174430, 18043064, 20004087, 17443671, 22195485, 18000422, 20004021 22321756, 17571039, 25897615, 27053456, 21067387, 16832076, 22905130 16344544, 21429602, 18009564, 14354737, 21286665, 18135678, 14521849 18614015, 20441797, 18362222, 25655390, 16472716, 17835048, 17050888 17936109, 14010183, 17325413, 18747196, 17761775, 16721594, 17082983 20067212, 21179898, 17302277, 18084625, 20717359, 24624166, 15990359 26746894, 24842886, 18203835, 23026585, 17297939, 17811456, 16731148 22380919, 21168487, 14133975, 13829543, 17215560, 17694209, 17385178 18091059, 8322815, 18259031, 25165496, 19689979, 17586955, 17201159 17655634, 18331812, 19730508, 18868646, 17648596, 16220077, 16069901 17393915, 17348614, 17957017, 17274537, 18096714, 17308789, 18436647 14285317, 19289642, 14764829, 17622427, 18328509, 16943711, 22195477 22502493, 14368995, 17346671, 18996843, 17783588, 21343838, 16618694 17672719, 18856999, 18783224, 17851160, 17546761, 22168163, 17798953 18273830, 22092979, 16596890, 19972566, 20828947, 13871092, 26667023 17726838, 16384983, 22296366, 17360606, 13645875, 22321741, 16542886 25879656, 18199537, 21787056, 17889549, 21172913, 14565184, 20475845 17071721, 21281607, 17610798, 20299015, 21343897, 22893153, 20657441 17397545, 18230522, 16360112, 19769489, 12905058, 18641451, 12747740 18430495, 25423453, 17016369, 17042658, 14602788, 17551063, 19972568 21517440, 19788842, 18508861, 14657740, 17332800, 13837378, 17186905 19972564, 19699191, 18315328, 17437634, 24570598, 22353199, 18093615 19006849, 19013183, 17296856, 18674024, 26569225, 17232014, 16855292 21051840, 14692762, 17762296, 17705023, 23294548, 22507234, 19121551 20324049, 21330264, 26198926, 19854503, 23315889, 26910644, 26030218 21868720, 19309466, 25764020, 18681862, 17365043, 20031873, 20558005 18554763, 17390160, 24717859, 21795111, 18456514, 16306373, 13955826 18139690, 17501491, 17752121, 21668627, 17299889, 23713236, 24652769 17889583, 18673325, 22551446, 19721304, 18293054, 17242746, 19211433 19888853, 17951233, 18094246, 17649265, 19615136, 17011832, 17477958 16870214, 18522509, 20631274, 16091637, 17323222, 16595641, 16524926 18228645, 18282562, 17596908, 18031668, 17156148, 16494615, 22683225 20869721, 17545847, 25093656, 17655240, 24528741, 17614134, 25427662 13558557, 22465352, 17341326, 17891946, 17716305, 22657942, 16392068 18440095, 19271443, 21351877, 20513399, 18092127, 17614227, 18440047 18849970, 16903536, 14106803, 18973907, 18673342, 22809871, 17389192 19032867, 25505382, 17612828, 17006570, 16194160, 25369547, 16685417 25505407, 17721717, 21354456, 17390431, 17570240, 16863422, 18325460 17008068, 19727057, 16422541, 19972570, 17267114, 18244962, 21538485 18203838, 18765602, 16198143, 17246576, 14829250, 17835627, 20860659 21629064, 18247991, 14458214, 21051862, 17786278, 16692232, 17227277 24476265, 16042673, 16314254, 19285025, 16228604, 16756406, 16837842 20144308, 17393683, 23536835, 25823754, 18899974, 17787259, 24719736 20331945, 19490948, 20074391, 15861775, 16399083, 25947799, 18018515 22683212, 21051858, 18260550, 17080436, 16613964, 17036973, 16579084 24433711, 18384537, 18280813, 20296213, 16901385, 15979965, 23330124 18441944, 16450169, 9756271, 17892268, 11733603, 16285691, 17587063 21343775, 18180390, 26474853, 16538760, 18193833, 21387964, 21051833 17238511, 19777862, 23065323, 17824637, 16571443, 17903598, 18306996 19578350, 14852021, 17853456, 18674047, 12364061, 19207117, 24411921, 22195448

Version 11.2.0.4.v14

Version 11.2.0.4.v14 adds support for the following:

- Oracle October 2017 PSU, a combination of database PSU (patch 26392168) + OJVM component PSU (patch 26635834)
- Oracle recommended RDBMS patches for Oracle GoldenGate (patch 26950781)
- RSA Micro-Edition Suite Bundle (patch 26963526)
- Timezone file DSTv30 (patch 25881255, OJVM patch 25881271)

Oracle patch 26392168, released October 2017

Bugs fixed:

17288409, 21051852, 24316947, 17811429, 17205719, 18607546, 25654936 20506699, 17816865, 25957038, 23330119, 17922254, 17754782, 13364795 16934803, 17311728, 20387265, 17284817, 17441661, 24560906, 16992075 17446237, 14015842, 19972569, 21756677, 17375354, 21538558, 20925795 17449815, 26575788, 19463897, 13866822, 17235750, 17982555, 17478514 18317531, 14338435, 18235390, 20803583, 19461270, 13944971, 20142975 17811789, 16929165, 18704244, 24662775, 20506706, 17546973, 20334344 25489607, 14054676, 17088068, 17346091, 18264060, 17343514, 21538567 19680952, 18471685, 19211724, 21132297, 13951456, 21847223, 16315398 18744139, 16850630, 23177648, 19049453, 18673304, 17883081, 19915271 18641419, 18262334, 25600421, 17006183, 16065166, 18277454, 16833527 10136473, 18051556, 17865671, 17852463, 18554871, 17853498, 18334586 20879889, 17551709, 17588480, 19827973, 17344412, 17842825, 18828868 20509482, 17025461, 11883252, 13609098, 17239687, 17602269, 19197175 18316692, 22195457, 17313525, 12611721, 19544839, 18964939, 26667015 17600719, 18191164, 19393542, 17571306, 20777150, 18482502, 19466309 22243719, 17040527, 17165204, 18098207, 16785708, 17465741, 16180763 17174582, 12982566, 16777840, 19463893, 22195465, 16875449, 22148226 12816846, 17237521, 6599380, 19358317, 17811438, 25505394, 17811447 17945983, 21983325, 18762750, 16912439, 17184721, 18061914, 17282229 18331850, 18202441, 17082359, 18723434, 21532755, 21972320, 19554106 25505371, 14034426, 18339044, 19458377, 17752995, 20448824, 17891943 17258090, 17767676, 16668584, 18384391, 17040764, 17381384, 15913355 18356166, 14084247, 20596234, 20506715, 21756661, 13853126, 18203837 14245531, 16043574, 21756699, 22195441, 17848897, 17877323, 19272701 21453153, 20569094, 17468141, 20861693, 17786518, 17912217, 17037130 16956380, 18155762, 17478145, 17394950, 18641461, 18189036, 18619917 17027426, 21352646, 16268425, 24476274, 22195492, 19584068, 26544823 18436307, 22507210, 17265217, 17634921, 13498382, 19469538, 21526048 19258504, 18043064, 20004087, 17443671, 22195485, 18000422, 20004021 22321756, 17571039, 21067387, 16832076, 22905130, 16344544, 21429602 18009564, 14354737, 21286665, 18135678, 14521849, 18614015, 20441797 18362222, 25655390, 16472716, 17835048, 17050888, 17936109, 14010183 17325413, 18747196, 17761775, 16721594, 17082983, 20067212, 21179898 17302277, 18084625, 24624166, 15990359, 26746894, 24842886, 23026585 18203835, 17297939, 17811456, 16731148, 22380919, 21168487, 14133975 13829543, 17215560, 17694209, 17385178, 18091059, 8322815, 18259031 19689979, 17586955, 17201159, 17655634, 18331812, 19730508, 18868646 17648596, 16220077, 16069901, 17348614, 17393915, 17957017, 17274537 18096714, 17308789, 18436647, 14285317, 19289642, 14764829, 17622427 18328509, 16943711, 22195477, 14368995, 22502493, 17346671, 18996843 17783588, 21343838, 16618694, 17672719, 18856999, 18783224, 17851160 17546761, 22168163, 17798953, 18273830, 22092979, 16596890, 19972566 20828947, 13871092, 26667023, 17726838, 16384983, 22296366, 17360606 22321741, 13645875, 25879656, 18199537, 16542886, 21787056, 17889549 14565184, 20475845, 21281607, 17071721, 17610798, 20299015, 21343897 22893153, 20657441, 17397545, 18230522, 16360112, 19769489, 12905058 18641451, 12747740, 18430495, 25423453, 17016369, 17042658, 14602788 17551063, 19972568, 21517440, 19788842, 18508861, 14657740, 17332800 13837378, 17186905, 19972564, 19699191, 18315328, 17437634, 22353199 18093615, 19006849, 19013183, 17296856, 18674024, 17232014, 16855292 17762296, 14692762, 21051840, 17705023, 23294548, 22507234, 19121551 21330264, 26198926, 19854503, 23315889, 26030218, 21868720, 19309466 18681862, 17365043, 20558005, 18554763, 17390160, 18456514, 16306373 13955826, 18139690, 17501491, 17752121, 21668627, 17299889, 23713236 24652769, 17889583, 18673325, 22551446, 19721304, 18293054, 17242746 19211433, 19888853, 17951233, 18094246, 17649265, 19615136, 17011832 16870214, 17477958, 18522509, 20631274, 16091637, 17323222, 16595641 16524926, 18228645, 18282562, 17596908, 18031668, 17156148, 16494615 22683225,

20869721, 17545847, 25093656, 17655240, 24528741, 17614134 25427662, 13558557, 17341326, 17891946, 17716305, 22657942, 18440095 16392068, 19271443, 21351877, 18092127, 17614227, 18440047, 18849970 16903536, 14106803, 18973907, 18673342, 17389192, 25505382, 19032867 17612828, 16194160, 17006570, 25369547, 25505407, 16685417, 17721717 17390431, 17570240, 16863422, 18325460, 17008068, 19727057, 16422541 19972570, 17267114, 18244962, 21538485, 18203838, 18765602, 16198143 17246576, 14829250, 17835627, 18247991, 14458214, 21051862, 17786278 16692232, 17227277, 24476265, 16042673, 16314254, 19285025, 16228604 16837842, 20144308, 17393683, 23536835, 25823754, 18899974, 17787259 24719736, 20331945, 19490948, 20074391, 15861775, 16399083, 25947799 18018515, 22683212, 21051858, 18260550, 17080436, 16613964, 17036973 16579084, 24433711, 18384537, 18280813, 20296213, 16901385, 15979965 23330124, 18441944, 16450169, 9756271, 17892268, 11733603, 16285691 17587063, 21343775, 26474853, 18180390, 16538760, 18193833, 21387964 21051833, 17238511, 19777862, 23065323, 17824637, 17903598, 16571443 18306996, 19578350, 14852021, 17853456, 18674047, 12364061, 24411921 19207117, 22195448

Version 11.2.0.4.v13

Version 11.2.0.4.v13 adds support for the following:

- Oracle July 2017 PSU, a combination of database PSU (patch 26609445) + OJVM component PSU (patch 26027154)
- Oracle recommended RDBMS patches for Oracle GoldenGate (patch 26554712)
- RSA Micro-Edition Suite Bundle (patch 26770426)
- Timezone file DSTv30 (patch 25881255, OJVM patch 25881271)
- Adds support for [Validating DB Instance Files \(p. 1115\)](#) with the `RMAN` logical validation utility
- Adds support for [Setting the Default Edition for a DB Instance \(p. 1100\)](#)

Oracle patch 26609445, released July 2017

Bugs fixed:

17288409, 21051852, 24316947, 17811429, 17205719, 18607546, 20506699 17816865, 25957038, 23330119, 17922254, 17754782, 13364795, 16934803 17311728, 20387265, 17284817, 17441661, 24560906, 16992075, 17446237 14015842, 19972569, 21756677, 17375354, 21538558, 20925795, 17449815 26575788, 19463897, 13866822, 17235750, 17982555, 17478514, 18317531 14338435, 18235390, 20803583, 19461270, 13944971, 20142975, 17811789 16929165, 18704244, 20506706, 17546973, 20334344, 14054676, 17088068 17346091, 18264060, 17343514, 21538567, 19680952, 18471685, 19211724 13951456, 21847223, 16315398, 18744139, 16850630, 23177648, 19049453 18673304, 17883081, 19915271, 18641419, 18262334, 25600421, 17006183 16065166, 18277454, 16833527, 10136473, 18051556, 17865671, 17852463 18554871, 17853498, 18334586, 20879889, 17551709, 17588480, 19827973 17344412, 17842825, 18828868, 20509482, 17025461, 11883252, 13609098 17239687, 17602269, 19197175, 18316692, 22195457, 17313525, 12611721 19544839, 18964939, 17600719, 18191164, 19393542, 17571306, 20777150 18482502, 19466309, 22243719, 17040527, 17165204, 18098207, 16785708 17465741, 16180763, 17174582, 12982566, 16777840, 19463893, 22195465 16875449, 22148226, 12816846, 17237521, 6599380, 19358317, 17811438 25505394, 17811447, 17945983, 21983325, 18762750, 16912439, 17184721 18061914, 17282229, 18331850, 18202441, 17082359, 18723434, 21972320 19554106, 25505371, 14034426, 18339044, 19458377, 17752995, 20448824 17891943, 17258090, 17767676, 16668584, 18384391, 17040764, 17381384 15913355, 18356166, 14084247, 20596234, 20506715, 21756661, 13853126 18203837, 14245531, 16043574, 21756699, 22195441, 17848897, 17877323 21453153, 17468141, 20861693, 17786518, 17912217, 17037130, 16956380 18155762, 17478145, 17394950, 18641461, 18189036, 18619917, 17027426 21352646, 16268425, 24476274, 22195492, 19584068, 26544823, 18436307 22507210, 17265217, 17634921, 13498382, 19469538, 21526048, 19258504 18043064, 20004087, 17443671, 22195485, 18000422, 20004021, 22321756 17571039, 21067387, 16832076, 22905130,

16344544, 18009564, 14354737 21286665, 18135678, 14521849, 18614015, 20441797, 18362222, 25655390 16472716, 17835048, 17050888, 17936109, 14010183, 17325413, 18747196 17761775, 16721594, 17082983, 20067212, 21179898, 17302277, 18084625 15990359, 24842886, 18203835, 17297939, 17811456, 16731148, 22380919 21168487, 14133975, 13829543, 17215560, 17694209, 17385178, 18091059 8322815, 18259031, 19689979, 17586955, 17201159, 17655634, 18331812 19730508, 18868646, 17648596, 16220077, 16069901, 17348614, 17393915 17957017, 17274537, 18096714, 17308789, 18436647, 14285317, 19289642 14764829, 17622427, 18328509, 16943711, 22195477, 14368995, 22502493 17346671, 18996843, 17783588, 21343838, 16618694, 17672719, 18856999 18783224, 17851160, 17546761, 22168163, 17798953, 18273830, 22092979 16596890, 19972566, 13871092, 17726838, 16384983, 22296366, 17360606 22321741, 13645875, 25879656, 18199537, 16542886, 21787056, 17889549 14565184, 17071721, 17610798, 20299015, 21343897, 22893153, 20657441 17397545, 18230522, 16360112, 19769489, 12905058, 18641451, 12747740 18430495, 25423453, 17016369, 17042658, 14602788, 17551063, 19972568 21517440, 19788842, 18508861, 14657740, 17332800, 13837378, 17186905 19972564, 19699191, 18315328, 17437634, 22353199, 18093615, 19006849 19013183, 17296856, 18674024, 17232014, 16855292, 17762296, 14692762 21051840, 17705023, 22507234, 19121551, 21330264, 19854503, 26030218 21868720, 19309466, 18681862, 17365043, 20558005, 18554763, 17390160 18456514, 16306373, 13955826, 18139690, 17501491, 17752121, 21668627 17299889, 17889583, 18673325, 19721304, 18293054, 17242746, 19888853 17951233, 18094246, 17649265, 19615136, 17011832, 16870214, 17477958 18522509, 20631274, 16091637, 17323222, 16595641, 16524926, 18228645 18282562, 17596908, 18031668, 17156148, 16494615, 22683225, 17545847 25093656, 17655240, 24528741, 17614134, 25427662, 13558557, 17341326 17891946, 17716305, 22657942, 18440095, 16392068, 19271443, 21351877 18092127, 17614227, 18440047, 16903536, 14106803, 18973907, 18673342 17389192, 25505382, 19032867, 17612828, 16194160, 17006570, 25369547 25505407, 16685417, 17721717, 17390431, 17570240, 16863422, 18325460 19727057, 16422541, 19972570, 17267114, 18244962, 21538485, 18203838 18765602, 16198143, 17246576, 14829250, 17835627, 18247991, 14458214 21051862, 16692232, 17786278, 17227277, 24476265, 16042673, 16314254 16228604, 16837842, 17393683, 23536835, 25823754, 18899974, 17787259 20331945, 20074391, 15861775, 16399083, 18018515, 22683212, 21051858 18260550, 17080436, 16613964, 17036973, 16579084, 24433711, 18384537 18280813, 20296213, 16901385, 15979965, 23330124, 18441944, 16450169 9756271, 17892268, 11733603, 16285691, 17587063, 21343775, 18180390 16538760, 18193833, 21387964, 21051833, 17238511, 19777862, 17824637 16571443, 18306996, 19578350, 14852021, 17853456, 18674047, 12364061 24411921, 19207117, 22195448

Version 11.2.0.4.v12

Version 11.2.0.4.v12 adds support for the following:

- Oracle patch 25440428, a combination of database PSU (patch 24732075) + OJVM component PSU (patch 25434033)
- Oracle recommended RDBMS patches for Oracle GoldenGate (patch 25734992)
- MES Bundle (patch 24975421 for 11.2.0.4)
- Timezone file DSTv28 (patch 24701840)
- Adds support for the DBMS_CHANGE_NOTIFICATION package
- Adds support for XSTREAM packages and views (may require additional licensing)

Oracle patch 24732075, released April 2017

Bugs fixed:

17288409, 21051852, 24316947, 17811429, 17205719, 18607546, 20506699 17816865, 17922254, 23330119, 17754782, 16934803, 13364795, 17311728 17284817, 17441661, 24560906, 16992075, 17446237, 14015842, 19972569 21756677, 17375354, 20925795, 21538558, 17449815, 19463897, 13866822 17235750, 17982555, 17478514, 18317531, 14338435, 18235390, 20803583 13944971,

20142975, 17811789, 16929165, 18704244, 20506706, 17546973 20334344, 14054676, 17088068, 17346091, 18264060, 17343514, 21538567 19680952, 18471685, 19211724, 13951456, 21847223, 16315398, 18744139 16850630, 23177648, 19049453, 18673304, 17883081, 19915271, 18641419 18262334, 17006183, 16065166, 18277454, 16833527, 10136473, 18051556 17865671, 17852463, 18554871, 17853498, 18334586, 17551709, 17588480 19827973, 17344412, 17842825, 18828868, 17025461, 11883252, 13609098 17239687, 17602269, 19197175, 18316692, 22195457, 17313525, 12611721 19544839, 18964939, 17600719, 18191164, 19393542, 17571306, 20777150 18482502, 19466309, 22243719, 17040527, 17165204, 18098207, 16785708 17465741, 17174582, 16180763, 12982566, 16777840, 19463893, 22195465 16875449, 12816846, 22148226, 17237521, 6599380, 19358317, 25505394 17811438, 17811447, 17945983, 21983325, 18762750, 16912439, 17184721 18061914, 17282229, 18331850, 18202441, 17082359, 18723434, 21972320 19554106, 25505371, 14034426, 18339044, 19458377, 17752995, 20448824 17891943, 17258090, 17767676, 16668584, 18384391, 17040764, 17381384 15913355, 18356166, 14084247, 20596234, 20506715, 21756661, 13853126 18203837, 14245531, 16043574, 21756699, 22195441, 17848897, 17877323 21453153, 17468141, 20861693, 17786518, 17912217, 17037130, 16956380 18155762, 17478145, 17394950, 18641461, 18189036, 18619917, 17027426 21352646, 16268425, 24476274, 22195492, 19584068, 18436307, 22507210 17265217, 17634921, 13498382, 21526048, 19258504, 20004087, 17443671 22195485, 18000422, 22321756, 20004021, 17571039, 21067387, 22905130 16344544, 18009564, 14354737, 21286665, 18135678, 18614015, 20441797 18362222, 17835048, 16472716, 17936109, 17050888, 14010183, 17325413 18747196, 17761775, 16721594, 17082983, 20067212, 21179898, 17302277 18084625, 15990359, 24842886, 18203835, 17297939, 17811456, 22380919 16731148, 21168487, 14133975, 13829543, 17215560, 17694209, 17385178 18091059, 8322815, 17586955, 17201159, 17655634, 18331812, 19730508 18868646, 17648596, 16220077, 16069901, 17348614, 17393915, 17274537 17957017, 18096714, 17308789, 18436647, 14285317, 19289642, 14764829 17622427, 18328509, 16943711, 22195477, 14368995, 22502493, 17346671 18996843, 17783588, 21343838, 16618694, 17672719, 18856999, 18783224 17851160, 17546761, 17798953, 18273830, 22092979, 16596890, 19972566 16384983, 17726838, 22296366, 17360606, 22321741, 13645875, 18199537 16542886, 21787056, 17889549, 14565184, 17071721, 17610798, 20299015 21343897, 22893153, 20657441, 17397545, 18230522, 16360112, 19769489 12905058, 18641451, 12747740, 18430495, 17016369, 17042658, 14602788 17551063, 19972568, 21517440, 18508861, 19788842, 14657740, 17332800 13837378, 19972564, 17186905, 18315328, 19699191, 17437634, 22353199 18093615, 19006849, 19013183, 17296856, 18674024, 17232014, 16855292 17762296, 14692762, 21051840, 17705023, 22507234, 19121551, 21330264 19854503, 21868720, 19309466, 18681862, 20558005, 18554763, 17390160 18456514, 16306373, 13955826, 18139690, 17501491, 17752121, 21668627 17299889, 17889583, 18673325, 19721304, 18293054, 17242746, 17951233 18094246, 17649265, 19615136, 17011832, 16870214, 17477958, 18522509 20631274, 16091637, 17323222, 16595641, 16524926, 18228645, 18282562 17596908, 18031668, 17156148, 16494615, 22683225, 17545847, 25093656 17655240, 24528741, 17614134, 13558557, 17341326, 17891946, 17716305 22657942, 18440095, 16392068, 19271443, 21351877, 18092127, 17614227 18440047, 16903536, 14106803, 18973907, 18673342, 25505382, 19032867 17389192, 17612828, 16194160, 17006570, 25369547, 25505407, 17721717 17390431, 17570240, 16863422, 18325460, 19727057, 16422541, 19972570 17267114, 18244962, 21538485, 18765602, 18203838, 16198143, 17246576 14829250, 17835627, 18247991, 14458214, 21051862, 16692232, 17786278 17227277, 24476265, 16042673, 16314254, 16228604, 16837842, 17393683 23536835, 17787259, 20331945, 20074391, 15861775, 16399083, 18018515 22683212, 18260550, 21051858, 17080436, 16613964, 17036973, 16579084 24433711, 18384537, 18280813, 20296213, 16901385, 15979965, 23330124 18441944, 16450169, 9756271, 17892268, 11733603, 16285691, 17587063 21343775, 18180390, 16538760, 18193833, 21387964, 21051833, 17238511 17824637, 16571443, 18306996, 14852021, 17853456, 18674047, 12364061 24411921, 22195448

Version 11.2.0.4.v11

Version 11.2.0.4.v11 adds support for the following:

- Oracle patch 24918033, a combination of database PSU (patch 24006111) + OJVM component PSU (patch 24917954)
- Oracle recommended RDBMS patches for Oracle GoldenGate (patch 24491261)

- MES Bundle (patch 24975421 for 11.2.0.4)

Oracle patch 24918033, released January 2017

Bugs fixed:

18933818, 19176885, 17201047, 25067795, 14774730, 19153980, 21911849 23727132, 18166577, 24448240, 17056813, 21811517, 19909862, 22675136 24534298, 19895326, 22253904, 17804361, 19231857, 17528315, 19058059 19554117, 19007266, 17285560, 22670385, 18458318, 19187988, 23265914 19006757, 19374518, 19223010, 25076732, 22118835, 19852360, 20408829 21047766, 21566944, 17288409, 21051852, 24316947, 17811429, 18607546, 17205719, 20506699 17816865, 17922254, 23330119, 17754782, 16934803, 13364795, 17311728 17441661, 17284817, 16992075, 17446237, 14015842, 19972569, 21756677 17375354, 20925795, 21538558, 17449815, 19463897, 13866822, 17235750 17982555, 17478514, 18317531, 14338435, 18235390, 20803583, 13944971 20142975, 17811789, 16929165, 18704244, 20506706, 17546973, 20334344 14054676, 17088068, 17346091, 18264060, 17343514, 21538567, 19680952 18471685, 19211724, 13951456, 21847223, 16315398, 18744139, 16850630 23177648, 19049453, 18673304, 17883081, 19915271, 18641419, 18262334 17006183, 16065166, 18277454, 16833527, 10136473, 18051556, 17865671 17852463, 18554871, 17853498, 18334586, 17551709, 17588480, 19827973 17344412, 17842825, 18828868, 17025461, 11883252, 13609098, 17239687 17602269, 19197175, 22195457, 18316692, 17313525, 12611721, 19544839 18964939, 17600719, 18191164, 19393542, 17571306, 20777150, 18482502 19466309, 22243719, 17040527, 17165204, 18098207, 16785708, 17465741 17174582, 16180763, 16777840, 12982566, 19463893, 22195465, 22148226 16875449, 12816846, 17237521, 6599380, 19358317, 17811438, 17811447 17945983, 21983325, 18762750, 16912439, 17184721, 18061914, 17282229 18331850, 18202441, 17082359, 18723434, 21972320, 19554106, 14034426 18339044, 19458377, 17752995, 20448824, 17891943, 17258090, 17767676 16668584, 18384391, 17040764, 17381384, 15913355, 18356166, 14084247 20596234, 20506715, 21756661, 13853126, 18203837, 14245531, 16043574 21756699, 22195441, 17848897, 17877323, 21453153, 17468141, 20861693 17786518, 17912217, 17037130, 16956380, 18155762, 17478145, 17394950 18641461, 18189036, 18619917, 17027426, 21352646, 16268425, 24476274 22195492, 19584068, 18436307, 22507210, 17265217, 17634921, 13498382 21526048, 19258504, 20004087, 17443671, 22195485, 18000422, 22321756 20004021, 17571039, 21067387, 16344544, 18009564, 14354737, 21286665 18135678, 18614015, 20441797, 18362222, 17835048, 16472716, 17936109 17050888, 17325413, 14010183, 18747196, 17761775, 16721594, 17082983 20067212, 21179898, 17302277, 18084625, 15990359, 18203835, 17297939 17811456, 22380919, 16731148, 21168487, 14133975, 13829543, 17215560 17694209, 17385178, 18091059, 8322815, 17586955, 17201159, 17655634 18331812, 19730508, 18868646, 17648596, 16220077, 16069901, 17348614 17393915, 17274537, 17957017, 18096714, 17308789, 18436647, 14285317 19289642, 14764829, 18328509, 17622427, 16943711, 22195477, 14368995 22502493, 17346671, 18996843, 17783588, 21343838, 16618694, 17627219 18856999, 18783224, 17851160, 17546761, 17798953, 18273830, 22092979 16596890, 19972566, 16384983, 17726838, 22296366, 17360606, 22321741 13645875, 18199537, 16542886, 21787056, 17889549, 14565184, 17071721 17610798, 20299015, 21343897, 22893153, 20657441, 17397545, 18230522 16360112, 19769489, 12905058, 18641451, 12747740, 18430495, 17016369 17042658, 14602788, 17551063, 19972568, 21517440, 18508861, 19788842 14657740, 17332800, 13837378, 19972564, 17186905, 18315328, 19699191 17437634, 22353199, 18093615, 19006849, 19013183, 17296856, 18674024 17232014, 16855292, 17762296, 14692762, 21051840, 17705023, 22507234 19121551, 21330264, 19854503, 21868720, 19309466, 18681862, 20558005 18554763, 17390160, 18456514, 16306373, 13955826, 18139690, 17501491 17752121, 21668627, 17299889, 17889583, 18673325, 19721304, 18293054 17242746, 17951233, 18094246, 17649265, 19615136, 17011832, 16870214 17477958, 18522509, 20631274, 16091637, 17323222, 16595641, 16524926 18228645, 18282562, 17596908, 18031668, 17156148, 16494615, 22683225 17545847, 17655240, 24528741, 17614134, 13558557, 17341326, 17891946 17716305, 22657942, 16392068, 19271443, 21351877, 18092127, 17614227 18440047, 16903536, 14106803, 18973907, 18673342, 19032867, 17389192 17612828, 16194160, 17006570, 17721717, 17390431, 17570240, 16863422 18325460, 19727057, 16422541, 19972570, 17267114, 18244962, 21538485 18765602, 18203838, 16198143, 17246576, 14829250, 17835627, 18247991 14458214, 21051862, 16692232, 17786278, 17227277, 24476265, 16042673, 16314254, 16228604, 16837842, 17393683, 23536835, 17787259, 20331945 20074391, 15861775,

16399083, 18018515, 22683212, 18260550, 21051858 17080436, 16613964, 17036973, 16579084, 24433711, 18384537, 18280813 20296213, 16901385, 15979965, 23330124, 18441944, 16450169, 9756271 17892268, 11733603, 16285691, 17587063, 21343775, 18180390, 16538760 18193833, 21387964, 21051833, 17238511, 17824637, 16571443, 18306996 14852021, 17853456, 18674047, 12364061, 22195448

Version 11.2.0.4.v10

Version 11.2.0.4.v10 adds support for the following:

- Oracle patch 24436313, a combination of database PSU (patch 24006111) + OJVM component PSU (patch 24315821)
- Oracle recommended RDBMS patches for Oracle GoldenGate (patch 24491261)
- MES Bundle (patch 24975421 for 11.2.0.4)

Baseline: Oracle Database Patch Set Update 11.2.0.4.161018 (patch 24006111, released October 2016)

Bugs fixed:

17288409, 21051852, 24316947, 17811429, 18607546, 17205719, 20506699 17816865, 17922254, 23330119, 17754782, 16934803, 13364795, 17311728 17441661, 17284817, 16992075, 17446237, 14015842, 19972569, 21756677 17375354, 20925795, 21538558, 17449815, 19463897, 13866822, 17235750 17982555, 17478514, 18317531, 14338435, 18235390, 20803583, 13944971 20142975, 17811789, 16929165, 18704244, 20506706, 17546973, 20334344 14054676, 17088068, 17346091, 18264060, 17343514, 21538567, 19680952 18471685, 19211724, 13951456, 21847223, 16315398, 18744139, 16850630 23177648, 19049453, 18673304, 17883081, 19915271, 18641419, 18262334 17006183, 16065166, 18277454, 16833527, 10136473, 18051556, 17865671 17852463, 18554871, 17853498, 18334586, 17551709, 17588480, 19827973 17344412, 17842825, 18828868, 17025461, 11883252, 13609098, 17239687 17602269, 19197175, 22195457, 18316692, 17313525, 12611721, 19544839 18964939, 17600719, 18191164, 19393542, 17571306, 20777150, 18482502 19466309, 22243719, 17040527, 17165204, 18098207, 16785708, 17465741 17174582, 16180763, 16777840, 12982566, 19463893, 22195465, 22148226 16875449, 12816846, 17237521, 6599380, 19358317, 17811438, 17811447 17945983, 21983325, 18762750, 16912439, 17184721, 18061914, 17282229 18331850, 18202441, 17082359, 18723434, 21972320, 19554106, 14034426 18339044, 19458377, 17752995, 20448824, 17891943, 17258090, 17767676 16668584, 18384391, 17040764, 17381384, 15913355, 18356166, 14084247 20596234, 20506715, 21756661, 13853126, 18203837, 14245531, 16043574 21756699, 22195441, 17848897, 17877323, 21453153, 17468141, 20861693 17786518, 17912217, 17037130, 16956380, 18155762, 17478145, 17394950 18641461, 18189036, 18619917, 17027426, 21352646, 16268425, 24476274 22195492, 19584068, 18436307, 22507210, 17265217, 17634921, 13498382 21526048, 19258504, 20004087, 17443671, 22195485, 18000422, 22321756 20004021, 17571039, 21067387, 16344544, 18009564, 14354737, 21286665 18135678, 18614015, 20441797, 18362222, 17835048, 16472716, 17936109 17050888, 17325413, 14010183, 18747196, 17761775, 16721594, 17082983 20067212, 21179898, 17302277, 18084625, 15990359, 18203835, 17297939 17811456, 22380919, 16731148, 21168487, 14133975, 13829543, 17215560 17694209, 17385178, 18091059, 8322815, 17586955, 17201159, 17655634 18331812, 19730508, 18868646, 17648596, 16220077, 16069901, 17348614 17393915, 17274537, 17957017, 18096714, 17308789, 18436647, 14285317 19289642, 14764829, 18328509, 17622427, 16943711, 22195477, 14368995 22502493, 17346671, 18996843, 17783588, 21343838, 16618694, 17672719 18856999, 18783224, 17851160, 17546761, 17798953, 18273830, 22092979 16596890, 19972566, 16384983, 17726838, 22296366, 17360606, 22321741 13645875, 18199537, 16542886, 21787056, 17889549, 14565184, 17071721 17610798, 20299015, 21343897, 22893153, 20657441, 17397545, 18230522 16360112, 19769489, 12905058, 18641451, 12747740, 18430495, 17016369 17042658, 14602788, 17551063, 19972568, 21517440, 18508861, 19788842 14657740, 17332800, 13837378, 19972564, 17186905, 18315328, 19699191 17437634, 22353199, 18093615, 19006849, 19013183, 17296856, 18674024 17232014, 16855292, 17762296, 14692762, 21051840, 17705023, 22507234 19121551, 21330264,

19854503, 21868720, 19309466, 18681862, 20558005 18554763, 17390160, 18456514, 16306373, 13955826, 18139690, 17501491 17752121, 21668627, 17299889, 17889583, 18673325, 19721304, 18293054 17242746, 17951233, 18094246, 17649265, 19615136, 17011832, 16870214 17477958, 18522509, 20631274, 16091637, 17323222, 16595641, 16524926 18228645, 18282562, 17596908, 18031668, 17156148, 16494615, 22683225 17545847, 17655240, 24528741, 17614134, 13558557, 17341326, 17891946 17716305, 22657942, 16392068, 19271443, 21351877, 18092127, 17614227 18440047, 16903536, 14106803, 18973907, 18673342, 19032867, 17389192 17612828, 16194160, 17006570, 17721717, 17390431, 17570240, 16863422 18325460, 19727057, 16422541, 19972570, 17267114, 18244962, 21538485 18765602, 18203838, 16198143, 17246576, 14829250, 17835627, 18247991 14458214, 21051862, 16692232, 17786278, 17227277, 24476265, 16042673 16314254, 16228604, 16837842, 17393683, 23536835, 17787259, 20331945 20074391, 15861775, 16399083, 18018515, 22683212, 18260550, 21051858 17080436, 16613964, 17036973, 16579084, 24433711, 18384537, 18280813 20296213, 16901385, 15979965, 23330124, 18441944, 16450169, 9756271 17892268, 11733603, 16285691, 17587063, 21343775, 18180390, 16538760 18193833, 21387964, 21051833, 17238511, 17824637, 16571443, 18306996 14852021, 17853456, 18674047, 12364061, 22195448

Version 11.2.0.4.v9

Version 11.2.0.4.v9 adds support for the following:

- Oracle patch 23615392, a combination of database PSU (patch 23054359) + OJVM component PSU (patch 23177551)
- Timezone file DSTv26 (patch 22873635 for 11.2.0.4)
- Oracle recommended RDBMS patches for Oracle GoldenGate (patch 24320398 for 11.2.0.4.160719)
- MES Bundle (patch 22695784 for 11.2.0.4)
- Added the ability to create custom password verify functions. For more information, see [Creating Custom Functions to Verify Passwords \(p. 1091\)](#).
- Fixed a bug that prevented implicit recompilation of views owned by SYS

Baseline: Oracle Database Patch Set Update 11.2.0.4.160719 (patch 23054359, released July 2016)

Bugs fixed:

17288409, 21051852, 17811429, 18607546, 17205719, 20506699, 17816865 23330119, 17922254, 17754782, 16934803, 13364795, 17311728, 17441661 17284817, 16992075, 17446237, 14015842, 19972569, 21756677, 17375354 21538558, 20925795, 17449815, 19463897, 13866822, 17982555, 17235750 17478514, 18317531, 14338435, 18235390, 20803583, 13944971, 20142975 17811789, 16929165, 18704244, 20506706, 17546973, 20334344, 14054676 17088068, 17346091, 18264060, 17343514, 21538567, 19680952, 18471685 19211724, 13951456, 21847223, 16315398, 18744139, 16850630, 23177648 19049453, 18673304, 17883081, 19915271, 18641419, 18262334, 17006183 16065166, 18277454, 16833527, 10136473, 18051556, 17865671, 17852463 18554871, 17853498, 18334586, 17551709, 17588480, 19827973, 17344412 17842825, 18828868, 17025461, 11883252, 13609098, 17239687, 17602269 19197175, 22195457, 18316692, 17313525, 12611721, 19544839, 18964939 17600719, 18191164, 19393542, 17571306, 18482502, 20777150, 19466309 17040527, 17165204, 18098207, 16785708, 17465741, 17174582, 16180763 16777840, 12982566, 19463893, 22195465, 16875449, 12816846, 17237521 19358317, 17811438, 17811447, 17945983, 21983325, 18762750, 16912439 17184721, 18061914, 17282229, 18331850, 18202441, 17082359, 18723434 21972320, 19554106, 14034426, 18339044, 19458377, 17752995, 20448824 17891943, 17258090, 17767676, 16668584, 18384391, 17040764, 17381384 15913355, 18356166, 14084247, 20596234, 20506715, 21756661, 13853126 18203837, 14245531, 16043574, 21756699, 22195441, 17848897, 17877323 21453153, 17468141, 20861693, 17786518, 17912217, 17037130, 16956380 18155762, 17478145, 17394950, 18641461, 18189036, 18619917, 17027426 21352646, 16268425, 22195492, 19584068, 18436307, 22507210, 17265217 17634921, 13498382, 21526048, 19258504, 20004087,

17443671, 22195485 18000422, 22321756, 20004021, 17571039, 21067387, 16344544, 18009564, 14354737, 21286665, 18135678, 18614015, 20441797, 18362222, 17835048 16472716, 17936109, 17050888, 17325413, 14010183, 18747196, 17761775 16721594, 17082983, 20067212, 21179898, 17302277, 18084625, 15990359 18203835, 17297939, 22380919, 17811456, 16731148, 21168487, 13829543 17215560, 14133975, 17694209, 17385178, 18091059, 8322815, 17586955 17201159, 17655634, 18331812, 19730508, 18868646, 17648596, 16220077 16069901, 17348614, 17393915, 17274537, 17957017, 18096714, 17308789 18436647, 14285317, 19289642, 14764829, 18328509, 17622427, 16943711 22195477, 14368995, 22502493, 17346671, 18996843, 17783588, 21343838 16618694, 17672719, 18856999, 18783224, 17851160, 17546761, 17798953 18273830, 22092979, 16596890, 19972566, 16384983, 17726838, 22296366 17360606, 22321741, 13645875, 18199537, 16542886, 21787056, 17889549 14565184, 17071721, 17610798, 20299015, 21343897, 22893153, 20657441 17397545, 18230522, 16360112, 19769489, 12905058, 18641451, 12747740 18430495, 17016369, 17042658, 14602788, 17551063, 19972568, 21517440 18508861, 19788842, 14657740, 17332800, 13837378, 19972564, 17186905 18315328, 19699191, 17437634, 22353199, 18093615, 19006849, 19013183 17296856, 18674024, 17232014, 16855292, 17762296, 14692762, 21051840 17705023, 22507234, 19121551, 21330264, 19854503, 21868720, 19309466 18681862, 18554763, 20558005, 17390160, 18456514, 16306373, 13955826 18139690, 17501491, 17752121, 21668627, 17299889, 17889583, 18673325 19721304, 18293054, 17242746, 17951233, 18094246, 17649265, 19615136 17011832, 16870214, 17477958, 18522509, 20631274, 16091637, 17323222 16595641, 16524926, 18228645, 18282562, 17596908, 18031668, 17156148 16494615, 22683225, 17545847, 17655240, 17614134, 13558557, 17341326 17891946, 17716305, 16392068, 19271443, 21351877, 18092127, 17614227 18440047, 16903536, 14106803, 18973907, 18673342, 19032867, 17389192 17612828, 16194160, 17006570, 17721717, 17390431, 17570240, 16863422 18325460, 19727057, 16422541, 19972570, 17267114, 18244962, 21538485 18765602, 18203838, 16198143, 17246576, 14829250, 17835627, 18247991 14458214, 21051862, 16692232, 17786278, 17227277, 16042673, 16314254 16228604, 16837842, 17393683, 23536835, 17787259, 20331945, 20074391 15861775, 16399083, 18018515, 22683212, 18260550, 21051858, 17080436 16613964, 17036973, 16579084, 18384537, 18280813, 20296213, 16901385 15979965, 23330124, 18441944, 16450169, 9756271, 17892268, 11733603 16285691, 17587063, 21343775, 16538760, 18180390, 18193833, 21387964 21051833, 17238511, 17824637, 16571443, 18306996, 14852021, 17853456 18674047, 12364061, 22195448

Version 11.2.0.4.v8

Version 11.2.0.4.v8 adds support for the following:

- Oracle PSU 11.2.0.4.160419 (22502456)
- Timezone file DSTv25 (patch 22037014)
- Oracle recommended RDBMS patches for Oracle GoldenGate (patch 22576728)
- MES Bundle (patch 22695784 for 11.2.0.4)
- Adds the ability for the master user to grant privileges on SYS objects with the grant option using the RDSADMIN.RDSADMIN_UTIL.GRANT_SYS_OBJECT procedure
- Adds master user privileges to support most common schemas created by the Oracle Fusion Middleware Repository Creation Utility (RCU)

Baseline: Oracle Database Patch Set Update 11.2.0.4.160419 (patch 22502456, released April 2016)

Bugs fixed:

17288409, 21051852, 17811429, 18607546, 17205719, 20506699, 17816865 17922254, 17754782, 16934803, 13364795, 17311728, 17441661, 17284817 16992075, 17446237, 14015842, 19972569, 21756677, 21538558, 20925795 17449815, 17375354, 19463897, 13866822, 17982555, 17235750, 17478514 18317531, 14338435, 18235390, 20803583, 13944971, 20142975, 17811789 16929165, 18704244, 20506706, 17546973, 20334344, 14054676, 17088068 17346091, 18264060, 17343514,

21538567, 19680952, 18471685, 19211724 13951456, 21847223, 16315398, 18744139, 16850630, 19049453, 18673304 17883081, 19915271, 18641419, 18262334, 17006183, 16065166, 18277454 16833527, 10136473, 18051556, 17865671, 17852463, 18554871, 17853498 18334586, 17551709, 17588480, 19827973, 17344412, 17842825, 18828868 17025461, 11883252, 13609098, 17239687, 17602269, 19197175, 22195457 18316692, 17313525, 12611721, 19544839, 18964939, 17600719, 18191164 19393542, 17571306, 18482502, 20777150, 19466309, 17040527, 17165204 18098207, 16785708, 17465741, 17174582, 16180763, 16777840, 12982566 19463893, 22195465, 16875449, 12816846, 17237521, 19358317, 17811438 17811447, 21983325, 17945983, 18762750, 16912439, 17184721, 18061914 17282229, 18331850, 18202441, 17082359, 18723434, 21972320, 19554106 14034426, 18339044, 19458377, 17752995, 20448824, 17891943, 17258090 17767676, 16668584, 18384391, 17040764, 17381384, 15913355, 18356166 14084247, 20596234, 20506715, 21756661, 13853126, 18203837, 14245531 21756699, 16043574, 22195441, 17848897, 17877323, 21453153, 17468141 20861693, 17786518, 17912217, 17037130, 18155762, 16956380, 17478145 17394950, 18641461, 18189036, 18619917, 17027426, 21352646, 16268425 22195492, 19584068, 18436307, 17265217, 17634921, 13498382, 21526048 19258504, 20004087, 17443671, 22195485, 18000422, 20004021, 22321756 17571039, 21067387, 16344544, 18009564, 14354737, 21286665, 18135678 18614015, 20441797, 18362222, 17835048, 16472716, 17936109, 17050888 17325413, 14010183, 18747196, 17761775, 16721594, 17082983, 20067212 21179898, 17302277, 18084625, 15990359, 18203835, 17297939, 17811456 16731148, 21168487, 13829543, 17215560, 14133975, 17694209, 17385178 18091059, 8322815, 17586955, 17201159, 17655634, 18331812, 19730508 18868646, 17648596, 16220077, 16069901, 17348614, 17393915, 17274537 17957017, 18096714, 17308789, 18436647, 14285317, 19289642, 14764829 18328509, 17622427, 22195477, 16943711, 22502493, 14368995, 17346671 18996843, 17783588, 21343838, 16618694, 17672719, 18856999, 18783224 17851160, 17546761, 17798953, 18273830, 22092979, 16596890, 19972566 16384983, 17726838, 17360606, 22321741, 13645875, 18199537, 16542886 21787056, 17889549, 14565184, 17071721, 17610798, 20299015, 21343897 22893153, 20657441, 17397545, 18230522, 16360112, 19769489, 12905058 18641451, 12747740, 18430495, 17016369, 17042658, 14602788, 17551063 19972568, 21517440, 18508861, 19788842, 14657740, 17332800, 13837378 19972564, 17186905, 18315328, 19699191, 17437634, 22353199, 18093615 19006849, 19013183, 17296856, 18674024, 17232014, 16855292, 17762296 14692762, 21051840, 17705023, 19121551, 21330264, 19854503, 21868720 19309466, 18681862, 18554763, 20558005, 17390160, 18456514, 16306373 13955826, 18139690, 17501491, 17752121, 21668627, 17299889, 17889583 18673325, 19721304, 18293054, 17242746, 17951233, 17649265, 18094246 19615136, 17011832, 16870214, 17477958, 18522509, 20631274, 16091637 17323222, 16595641, 16524926, 18228645, 18282562, 17596908, 17156148 18031668, 16494615, 22683225, 17545847, 17655240, 17614134, 13558557 17341326, 17891946, 17716305, 16392068, 19271443, 21351877, 18092127 18440047, 17614227, 14106803, 16903536, 18973907, 18673342, 19032867 17389192, 17612828, 16194160, 17006570, 17721717, 17390431, 17570240 16863422, 18325460, 19727057, 16422541, 19972570, 17267114, 18244962 21538485, 18765602, 18203838, 16198143, 17246576, 14829250, 17835627 18247991, 14458214, 21051862, 16692232, 17786278, 17227277, 16042673 16314254, 16228604, 16837842, 17393683, 17787259, 20331945, 20074391 15861775, 16399083, 18018515, 22683212, 18260550, 21051858, 17036973 16613964, 17080436, 16579084, 18384537, 18280813, 20296213, 16901385 15979965, 18441944, 16450169, 9756271, 17892268, 11733603, 16285691 17587063, 21343775, 16538760, 18180390, 18193833, 21387964, 21051833 17238511, 17824637, 16571443, 18306996, 14852021, 18674047, 17853456 12364061, 22195448

Version 11.2.0.4.v7

Version 11.2.0.4.v7 adds support for the following:

- Oracle PSU 11.2.0.4.160119 (21948347)
- Timezone file DSTv25 - patch 22037014 for 11.2.0.4 and 12.1.0.2 (12.1.0.1 includes DSTv24, patch 20875898 (unchanged from 12.1.0.1.v3), as a backport of DSTv25 was unavailable at build time)
- Fixed an issue that prevented customers from creating more than 10 Directory objects in the database
- Fixed an issue that prevented customers from re-granting read privileges on the ADUMP and BDUMP Directory objects

Baseline: Oracle Database Patch Set Update 11.2.0.4.160119 (patch 21948347, released January 2016)

Bugs fixed:

17288409, 21051852, 18607546, 17205719, 17811429, 17816865, 20506699 17922254, 17754782, 16934803, 13364795, 17311728, 17441661, 17284817 16992075, 17446237, 14015842, 19972569, 17449815, 21538558, 20925795 17375354, 19463897, 17982555, 17235750, 13866822, 17478514, 18317531 18235390, 14338435, 20803583, 13944971, 20142975, 17811789, 16929165 18704244, 20506706, 17546973, 20334344, 14054676, 17088068, 18264060 17346091, 17343514, 21538567, 19680952, 18471685, 19211724, 13951456 21847223, 16315398, 18744139, 16850630, 19049453, 18673304, 17883081 19915271, 18641419, 18262334, 17006183, 16065166, 18277454, 16833527 10136473, 18051556, 17865671, 17852463, 18554871, 17853498, 18334586 17588480, 17551709, 19827973, 17842825, 17344412, 18828868, 17025461 11883252, 13609098, 17239687, 17602269, 19197175, 22195457, 18316692 17313525, 12611721, 19544839, 18964939, 17600719, 18191164, 19393542 17571306, 18482502, 20777150, 19466309, 17040527, 17165204, 18098207 16785708, 17174582, 16180763, 17465741, 16777840, 12982566, 19463893 22195465, 12816846, 16875449, 17237521, 19358317, 17811438, 17811447 17945983, 18762750, 17184721, 16912439, 18061914, 17282229, 18331850 18202441, 17082359, 18723434, 21972320, 19554106, 14034426, 18339044 19458377, 17752995, 20448824, 17891943, 17258090, 17767676, 16668584 18384391, 17040764, 17381384, 15913355, 18356166, 14084247, 20506715 13853126, 18203837, 14245531, 21756699, 16043574, 22195441, 17848897 17877323, 21453153, 17468141, 20861693, 17786518, 17912217, 17037130 18155762, 16956380, 17478145, 17394950, 18189036, 18641461, 18619917 17027426, 21352646, 16268425, 22195492, 19584068, 18436307, 17265217 17634921, 13498382, 21526048, 20004087, 22195485, 17443671, 18000422 22321756, 20004021, 17571039, 21067387, 16344544, 18009564, 14354737 18135678, 18614015, 20441797, 18362222, 17835048, 16472716, 17936109 17050888, 17325413, 14010183, 18747196, 17761775, 16721594, 17082983 20067212, 21179898, 17302277, 18084625, 15990359, 18203835, 17297939 17811456, 16731148, 21168487, 17215560, 13829543, 14133975, 17694209 18091059, 17385178, 8322815, 17586955, 17201159, 17655634, 18331812 19730508, 18868646, 17648596, 16220077, 16069901, 17348614, 17393915 17274537, 17957017, 18096714, 17308789, 18436647, 14285317, 19289642 14764829, 18328509, 17622427, 22195477, 16943711, 14368995, 17346671 18996843, 17783588, 21343838, 16618694, 17672719, 18856999, 18783224 17851160, 17546761, 17798953, 18273830, 22092979, 19972566, 16384983 17726838, 17360606, 22321741, 13645875, 18199537, 16542886, 21787056 17889549, 14565184, 17071721, 17610798, 20299015, 21343897, 20657441 17397545, 18230522, 16360112, 19769489, 12905058, 18641451, 12747740 18430495, 17042658, 17016369, 14602788, 17551063, 19972568, 21517440 18508861, 19788842, 14657740, 17332800, 13837378, 19972564, 17186905 18315328, 19699191, 17437634, 19006849, 19013183, 17296856, 18674024 17232014, 16855292, 21051840, 14692762, 17762296, 17705023, 19121551 21330264, 19854503, 19309466, 18681862, 18554763, 20558005, 17390160 18456514, 16306373, 13955826, 18139690, 17501491, 21668627, 17299889 17752121, 17889583, 18673325, 18293054, 17242746, 17951233, 17649265 18094246, 19615136, 17011832, 16870214, 17477958, 18522509, 20631274 16091637, 17323222, 16595641, 16524926, 18228645, 18282562, 17596908 17156148, 18031668, 16494615, 17545847, 17655240, 17614134, 13558557 17341326, 17891946, 17716305, 16392068, 19271443, 21351877, 18092127 18440047, 17614227, 14106803, 16903536, 18973907, 18673342, 19032867 17389192, 17612828, 16194160, 17006570, 17721717, 17570240, 17390431 16863422, 18325460, 19727057, 16422541, 19972570, 17267114, 18244962 21538485, 18765602, 18203838, 16198143, 17246576, 14829250, 17835627 18247991, 14458214, 21051862, 16692232, 17786278, 17227277, 16042673 16314254, 16228604, 16837842, 17393683, 17787259, 20331945, 20074391 15861775, 16399083, 18018515, 21051858, 18260550, 17036973, 16613964 17080436, 16579084, 18384537, 18280813, 20296213, 16901385, 15979965 18441944, 16450169, 9756271, 17892268, 11733603, 16285691, 17587063 21343775, 16538760, 18180390, 18193833, 21051833, 17238511, 17824637 16571443, 18306996, 14852021, 18674047, 17853456, 12364061, 22195448

Version 11.2.0.4.v6

Version 11.2.0.4.v6 adds support for the following:

- Enable SSL encryption for Standard Edition and Standard Edition One

Version 11.2.0.4.v5

Version 11.2.0.4.v5 adds support for the following:

- Oracle PSU 11.2.0.4.8 (21352635)
- Includes the Daylight Saving Time Patch, patch 20875898: DST-24, that came out after the April 2015 PSU.

Baseline: Oracle Database Patch Set Update 11.2.0.4.8 (patch 21352635, released October 2015)

Bugs fixed:

17288409, 21051852, 18607546, 17205719, 17811429, 17816865, 20506699 17922254, 17754782, 16934803, 13364795, 17311728, 17441661, 17284817 16992075, 17446237, 14015842, 19972569, 21538558, 20925795, 17449815 17375354, 19463897, 17982555, 17235750, 13866822, 18317531, 17478514 18235390, 14338435, 20803583, 13944971, 20142975, 17811789, 16929165 18704244, 20506706, 17546973, 20334344, 14054676, 17088068, 18264060 17346091, 17343514, 21538567, 19680952, 18471685, 19211724, 13951456 16315398, 18744139, 16850630, 19049453, 18673304, 17883081, 19915271 18641419, 18262334, 17006183, 16065166, 18277454, 16833527, 10136473 18051556, 17865671, 17852463, 18554871, 17853498, 18334586, 17588480 17551709, 19827973, 17842825, 17344412, 18828868, 17025461, 11883252 13609098, 17239687, 17602269, 19197175, 18316692, 17313525, 12611721 19544839, 18964939, 17600719, 18191164, 19393542, 17571306, 18482502 20777150, 19466309, 17040527, 17165204, 18098207, 16785708, 17174582 16180763, 17465741, 16777840, 12982566, 19463893, 12816846, 16875449 17237521, 19358317, 17811438, 17811447, 17945983, 18762750, 17184721 16912439, 18061914, 17282229, 18331850, 18202441, 17082359, 18723434 19554106, 14034426, 18339044, 19458377, 17752995, 20448824, 17891943 17258090, 17767676, 16668584, 18384391, 17040764, 17381384, 15913355 18356166, 14084247, 20506715, 13853126, 18203837, 14245531, 16043574 17848897, 17877323, 17468141, 17786518, 17912217, 17037130, 18155762 16956380, 17478145, 17394950, 18189036, 18641461, 18619917, 17027426 21352646, 16268425, 19584068, 18436307, 17265217, 17634921, 13498382 20004087, 17443671, 18000422, 20004021, 17571039, 21067387, 16344544 18009564, 14354737, 18135678, 18614015, 20441797, 18362222, 17835048 16472716, 17936109, 17050888, 17325413, 14010183, 18747196, 17761775 16721594, 17082983, 20067212, 21179898, 17302277, 18084625, 15990359 18203835, 17297939, 17811456, 16731148, 17215560, 13829543, 14133975 17694209, 18091059, 17385178, 8322815, 17586955, 17201159, 17655634 18331812, 19730508, 18868646, 17648596, 16220077, 16069901, 17348614 17393915, 17274537, 17957017, 18096714, 17308789, 18436647, 14285317 19289642, 14764829, 18328509, 17622427, 16943711, 14368995, 17346671 18996843, 17783588, 16618694, 17672719, 18856999, 18783224, 17851160 17546761, 17798953, 18273830, 19972566, 16384983, 17726838, 17360606 13645875, 18199537, 16542886, 17889549, 14565184, 17071721, 20299015 17610798, 20657441, 17397545, 18230522, 16360112, 19769489, 12905058 18641451, 12747740, 18430495, 17042658, 17016369, 14602788, 19972568 18508861, 19788842, 14657740, 17332800, 13837378, 19972564, 17186905 18315328, 19699191, 17437634, 19006849, 19013183, 17296856, 18674024 17232014, 16855292, 21051840, 14692762, 17762296, 17705023, 19121551 19854503, 19309466, 18681862, 18554763, 20558005, 17390160, 18456514 16306373, 13955826, 18139690, 17501491, 17299889, 17752121, 17889583 18673325, 18293054, 17242746, 17951233, 17649265, 18094246, 19615136 17011832, 16870214, 174777958, 18522509, 20631274, 16091637, 17323222 16595641, 16524926, 18228645, 18282562, 17596908, 17156148, 18031668 16494615, 17545847, 17614134, 13558557, 17341326, 17891946, 17716305 16392068, 19271443, 18092127, 18440047, 17614227, 14106803, 16903536 18973907, 18673342, 17389192, 16194160, 17006570, 17612828, 17721717 17570240, 17390431, 16863422, 18325460, 19727057, 16422541, 19972570 17267114, 18244962, 21538485, 18765602, 18203838, 16198143, 17246576 14829250, 17835627, 18247991, 14458214, 21051862, 16692232, 17786278 17227277, 16042673, 16314254,

16228604, 16837842, 17393683, 17787259 20331945, 20074391, 15861775, 16399083, 18018515, 18260550, 21051858 17036973, 16613964, 17080436, 16579084, 18384537, 18280813, 20296213 16901385, 15979965, 18441944, 16450169, 9756271, 17892268, 11733603 16285691, 17587063, 16538760, 18180390, 18193833, 21051833, 17238511 17824637, 16571443, 18306996, 14852021, 18674047, 17853456, 12364061

Version 11.2.0.4.v4

Version 11.2.0.4.v4 adds support for the following:

- Oracle PSU 11.2.0.4.6 (20299013)
- Installs additional Oracle Text knowledge bases from Oracle Database. Examples media (English and French)
- Provides access to DBMS_REPAIR through RDSADMIN.RDSADMIN_DBMS_REPAIR
- Grants ALTER DATABASE LINK, ALTER PUBLIC DATABASE LINK, EXEMPT ACCESS POLICY, EXEMPT IDENTITY POLICY, and EXEMPT REDACTION POLICY to master user

Baseline: Oracle Database Patch Set Update 11.2.0.4.6 (patch 20299013, released April 2015)

Bugs fixed:

17288409, 17798953, 18273830, 18607546, 17811429, 17205719, 20506699 17816865, 19972566, 17922254, 17754782, 16384983, 17726838, 13364795 16934803, 17311728, 17284817, 17441661, 17360606, 13645875, 18199537 16992075, 16542886, 17446237, 14015842, 17889549, 14565184, 19972569 17071721, 20299015, 17610798, 17375354, 17449815, 17397545, 19463897 18230522, 13866822, 17235750, 17982555, 16360112, 18317531, 17478514 19769489, 12905058, 14338435, 18235390, 13944971, 18641451, 20142975 17811789, 16929165, 18704244, 12747740, 18430495, 20506706, 17546973 14054676, 17088068, 17346091, 18264060, 17016369, 17042658, 17343514 14602788, 19972568, 19680952, 18471685, 19788842, 18508861, 14657740 17332800, 19211724, 13837378, 13951456, 16315398, 17186905, 18744139 19972564, 16850630, 18315328, 17437634, 19049453, 18673304, 17883081 19006849, 19915271, 19013183, 18641419, 17296856, 18674024, 18262334 17006183, 18277454, 16833527, 17232014, 16855292, 10136473, 17762296 14692762, 17705023, 18051556, 17865671, 17852463, 18554871, 17853498 19121551, 18334586, 19854503, 17551709, 19309466, 17588480, 19827973 17344412, 17842825, 18828868, 18681862, 18554763, 17390160, 18456514 16306373, 17025461, 13955826, 18139690, 11883252, 13609098, 17501491 17239687, 17752121, 17299889, 17602269, 19197175, 17889583, 18316692 17313525, 18673325, 12611721, 19544839, 18293054, 17242746, 18964939 17600719, 18191164, 19393542, 17571306, 18482502, 19466309, 17951233 17649265, 18094246, 19615136, 17040527, 17011832, 17165204, 18098207 16785708, 16870214, 17465741, 16180763, 17174582, 17477958, 12982566 16777840, 18522509, 20631274, 16091637, 17323222, 19463893, 16595641 16875449, 12816846, 16524926, 17237521, 18228645, 18282562, 17596908 19358317, 17811438, 17811447, 17945983, 18762750, 17156148, 18031668 16912439, 17184721, 16494615, 18061914, 17282229, 17545847, 18331850 18202441, 17082359, 18723434, 19554106, 17614134, 13558557, 17341326 14034426, 17891946, 18339044, 17716305, 19458377, 17752995, 16392068 19271443, 17891943, 18092127, 17258090, 17767676, 16668584, 18384391 17614227, 17040764, 16903536, 17381384, 14106803, 15913355, 18973907 18356166, 18673342, 17389192, 14084247, 16194160, 17612828, 17006570 20506715, 17721717, 13853126, 17390431, 18203837, 17570240, 14245531 16043574, 16863422, 17848897, 17877323, 18325460, 19727057, 17468141 17786518, 17912217, 16422541, 19972570, 17267114, 17037130, 18244962 18765602, 18203838, 18155762, 16956380, 16198143, 17246576, 17478145 17394950, 14829250, 18189036, 18641461, 18619917, 17835627, 17027426 16268425, 18247991, 19584068, 14458214, 18436307, 17265217, 17634921 13498382, 16692232, 17786278, 17227277, 16042673, 16314254, 17443671 18000422, 16228604, 16837842, 17571039, 17393683, 16344544, 17787259 18009564, 20074391, 14354737, 15861775, 18135678, 18614015, 16399083 18362222, 18018515, 16472716, 17835048, 17050888, 17936109, 14010183 17325413, 18747196, 17080436,

16613964, 17036973, 17761775, 16579084 16721594, 17082983, 18384537, 18280813, 20296213, 17302277, 16901385 18084625, 15979965, 15990359, 18203835, 17297939, 17811456, 16731148 13829543, 14133975, 17215560, 17694209, 18091059, 17385178, 8322815 17586955, 18441944, 17201159, 16450169, 9756271, 17655634, 19730508 17892268, 18868646, 17648596, 16220077, 16069901, 11733603, 16285691 17587063, 18180390, 16538760, 18193833, 17348614, 17393915, 17957017 17274537, 18096714, 17308789, 17238511, 18436647, 17824637, 14285317 19289642, 14764829, 17622427, 18328509, 16571443, 16943711, 14368995 18306996, 17346671, 14852021, 18996843, 17783588, 16618694, 17853456 18674047, 17672719, 18856999, 12364061, 18783224, 17851160, 17546761

Version 11.2.0.4.v3

Version 11.2.0.4.v3 adds support for the following:

- Oracle PSU 11.2.0.4.4 (19121551)
- Latest DST file (DSTv23 – patch 19396455, released Oct 2014). This patch is incorporated by default in new instances only.

Baseline: Oracle Database Patch Set Update 11.2.0.4.4 (patch 19121551, released October 2014)

Bugs fixed:

19396455, 18759211, 17432124, 16799735, 17288409, 17205719, 17811429, 17754782, 17726838, 13364795, 17311728 17284817, 17441661, 13645875, 18199537, 16992075, 16542886, 17446237 14565184, 17071721, 17610798, 17375354, 17449815, 17397545, 19463897 18230522, 17235750, 16360112, 13866822, 17982555, 17478514, 12905058 14338435, 13944971, 16929165, 12747740, 17546973, 14054676, 17088068 18264060, 17343514, 17016369, 17042658, 14602788, 14657740, 17332800 19211724, 13951456, 16315398, 17186905, 18744139, 16850630, 17437634 19049453, 18673304, 17883081, 18641419, 17296856, 18262334, 17006183 18277454, 17232014, 16855292, 10136473, 17705023, 17865671, 18554871 19121551, 17588480, 17551709, 17344412, 17842825, 18681862, 17390160 13955826, 13609098, 18139690, 17501491, 17239687, 17752121, 17299889 17602269, 18673325, 17313525, 17242746, 19544839, 17600719, 18191164 17571306, 19466309, 17951233, 18094246, 17165204, 17011832, 17040527 16785708, 16180763, 17477958, 17174582, 17465741, 18522509, 17323222 19463893, 16875449, 16524926, 17237521, 17596908, 17811438, 17811447 18031668, 16912439, 16494615, 18061914, 17545847, 17082359, 19554106 17614134, 17341326, 17891946, 19458377, 17716305, 17752995, 16392068 19271443, 17767676, 17614227, 17040764, 17381384, 18973907, 18673342 14084247, 17389192, 17006570, 17612828, 17721717, 13853126, 18203837 17390431, 17570240, 14245531, 16043574, 16863422, 19727057, 17468141 17786518, 17037130, 17267114, 18203838, 16198143, 16956380, 17478145 14829250, 17394950, 17027426, 16268425, 18247991, 19584068, 14458214 18436307, 17265217, 13498382, 16692232, 17786278, 17227277, 16042673 16314254, 17443671, 16228604, 16837842, 17393683, 17787259, 18009564 15861775, 16399083, 18018515, 16472716, 17050888, 14010183, 17325413 16613964, 17080436, 17036973, 17761775, 16721594, 18280813, 15979965 18203835, 17297939, 16731148, 17811456, 14133975, 17385178, 17586955 16450169, 17655634, 9756271, 17892268, 17648596, 16220077, 16069901 11733603, 16285691, 17587063, 18180390, 17393915, 18096714, 17238511 17824637, 14285317, 19289642, 14764829, 18328509, 17622427, 16943711 17346671, 18996843, 14852021, 17783588, 16618694, 17672719, 17546761

Version 11.2.0.4.v2 (Deprecated)

Version 11.2.0.4.v2 adds support for the following:

- Oracle PSU 11.2.0.4.3 (18522509)
- User access to DBMS_TRANSACTION package to clean-up failed distributed transactions

- Latest DST file (DSTv22 – patch 18759211, released June 2014). This patch is incorporated by default only in new Oracle DB instances.
- Grants DBMS_REPUTIL to DBA role (upgrade to 11.2.0.4 revokes it from public)
- Privileges granted on DBMS_TRANSACTION, v\$pending_xatrans\$, and v\$xatrans\$
- Resolves a problem with DDL commands when user objects have "SYSTEM" in their names
- Installs schema objects to support XA Transactions, allowing transactions to be managed by an external transaction manager
- Permits truncation of temporary SYS and SYSTEM objects, allowing tools like LogMiner to function correctly

Baseline: Oracle Database Patch Set Update 11.2.0.4.3 (patch 18522509, released July 2014)

Bugs fixed:

17432124, 18759211, 18522509, 18031668, 17478514, 17752995, 17288409, 16392068, 17205719, 17811429, 17767676, 17614227 17040764, 17381384, 17754782, 17726838, 13364795, 17311728, 17389192 17006570, 17612828, 17284817, 17441661, 13853126, 17721717, 13645875 18203837, 17390431, 16542886, 16992075, 16043574, 17446237, 16863422 14565184, 17071721, 17610798, 17468141, 17786518, 17375354, 17397545 18203838, 16956380, 17478145, 16360112, 17235750, 17394950, 13866822 17478514, 17027426, 12905058, 14338435, 16268425, 13944971, 18247991 14458214, 16929165, 17265217, 13498382, 17786278, 17227277, 17546973 14054676, 17088068, 16314254, 17016369, 14602788, 17443671, 16228604 16837842, 17332800, 17393683, 13951456, 16315398, 18744139, 17186905 16850630, 17437634, 19049453, 17883081, 15861775, 17296856, 18277454 16399083, 16855292, 18018515, 10136473, 16472716, 17050888, 17865671 17325413, 14010183, 18554871, 17080436, 16613964, 17761775, 16721594 17588480, 17551709, 17344412, 18681862, 15979965, 13609098, 18139690 17501491, 17239687, 17752121, 17602269, 18203835, 17297939, 17313525 16731148, 17811456, 14133975, 17600719, 17385178, 17571306, 16450169 17655634, 18094246, 17892268, 17165204, 17011832, 17648596, 16785708 17477958, 16180763, 16220077, 17465741, 17174582, 18522509, 16069901 16285691, 17323222, 18180390, 17393915, 16875449, 18096714, 17238511

Version 11.2.0.4.v1

Version 11.2.0.4.v1 adds support for the following:

- Oracle PSU 11.2.0.4.1
- [Creating New Directories in the Main Data Storage Space \(p. 1131\)](#)

Baseline: Oracle Database Patch Set Update 11.2.0.4.1 (released January 2014)

Bugs fixed:

17432124, 16850630, 17551709, 13944971, 17811447, 13866822, 17811429, 16069901 16721594, 17443671, 17478514, 17612828, 17610798, 17239687, 17501491 17446237, 16450169, 17811438, 17288409, 17811456, 12905058, 17088068 16285691, 17332800

Related Topics

- [Upgrading the Oracle DB Engine \(p. 979\)](#)
- [Oracle on Amazon RDS \(p. 914\)](#)

PostgreSQL on Amazon RDS

Amazon RDS supports DB instances running several versions of PostgreSQL. You can create DB instances and DB snapshots, point-in-time restores and backups. DB instances running PostgreSQL support Multi-AZ deployments, Read Replicas (version 9.3.5 and later), Provisioned IOPS, and can be created inside a VPC. You can also use Secure Socket Layer (SSL) to connect to a DB instance running PostgreSQL.

Before creating a DB instance, you should complete the steps in the [Setting Up for Amazon RDS \(p. 5\)](#) section of this guide.

You can use any standard SQL client application to run commands for the instance from your client computer. Such applications include *pgAdmin*, a popular Open Source administration and development tool for PostgreSQL, or *psql*, a command line utility that is part of a PostgreSQL installation. To deliver a managed service experience, Amazon RDS doesn't provide host access to DB instances, and it restricts access to certain system procedures and tables that require advanced privileges. Amazon RDS supports access to databases on a DB instance using any standard SQL client application. Amazon RDS doesn't allow direct host access to a DB instance by using Telnet or Secure Shell (SSH).

Amazon RDS for PostgreSQL is compliant with many industry standards. For example, you can use Amazon RDS for PostgreSQL databases to build HIPAA-compliant applications and to store healthcare-related information, including protected health information (PHI) under an executed Business Associate Agreement (BAA) with AWS. Amazon RDS for PostgreSQL also meets Federal Risk and Authorization Management Program (FedRAMP) security requirements. Amazon RDS for PostgreSQL has received a FedRAMP Joint Authorization Board (JAB) Provisional Authority to Operate (P-ATO) at the FedRAMP HIGH Baseline within the AWS GovCloud (US-West) Region. For more information on supported compliance standards, see [AWS Cloud Compliance](#).

To import PostgreSQL data into a DB instance, follow the information in the [Importing Data into PostgreSQL on Amazon RDS \(p. 1302\)](#) section.

Topics

- [Common Management Tasks for PostgreSQL on Amazon RDS \(p. 1254\)](#)
- [Creating a DB Instance Running the PostgreSQL Database Engine \(p. 1257\)](#)
- [Connecting to a DB Instance Running the PostgreSQL Database Engine \(p. 1270\)](#)
- [Updating Applications to Connect to PostgreSQL DB Instances Using New SSL/TLS Certificates \(p. 1275\)](#)
- [Modifying a DB Instance Running the PostgreSQL Database Engine \(p. 1279\)](#)
- [Upgrading the PostgreSQL DB Engine for Amazon RDS \(p. 1289\)](#)
- [Upgrading a PostgreSQL DB Snapshot \(p. 1296\)](#)
- [Working with PostgreSQL Read Replicas \(p. 1298\)](#)
- [Importing Data into PostgreSQL on Amazon RDS \(p. 1302\)](#)
- [Common DBA Tasks for PostgreSQL \(p. 1322\)](#)
- [Using Kerberos Authentication with Amazon RDS for PostgreSQL \(p. 1351\)](#)
- [Working with the Database Preview Environment \(p. 1362\)](#)
- [Amazon RDS for PostgreSQL Versions and Extensions \(p. 1364\)](#)

Common Management Tasks for PostgreSQL on Amazon RDS

The following are the common management tasks you perform with an Amazon RDS for PostgreSQL DB instance, with links to relevant documentation for each task.

Task Area	Relevant Documentation
Setting up Amazon RDS for first-time use There are prerequisites you must complete before you create your DB instance. For example, DB instances are created by default with a firewall that prevents access to it. You therefore must create a security group with the correct IP addresses and network configuration to access the DB instance.	Setting Up for Amazon RDS (p. 5)
Understanding Amazon RDS DB instances If you are creating a DB instance for production purposes, you should understand how instance classes, storage types, and Provisioned IOPS work in Amazon RDS.	Choosing the DB Instance Class (p. 107) Amazon RDS Storage Types (p. 135) Provisioned IOPS SSD Storage (p. 137)
Finding supported PostgreSQL versions Amazon RDS supports several versions of PostgreSQL.	Supported PostgreSQL Database Versions (p. 1365)
Setting up high availability and failover support A production DB instance should use Multi-AZ deployments. Multi-AZ deployments provide increased availability, data durability, and fault tolerance for DB instances.	High Availability (Multi-AZ) for Amazon RDS (p. 144)
Understanding the Amazon Virtual Private Cloud (VPC) network If your AWS account has a default VPC, then your DB instance is automatically created inside the default VPC. In some cases, your account might not have a default VPC, and you might want the DB instance in a VPC. In these cases, create the VPC and subnet groups before you create the DB instance.	Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform (p. 506) Working with a DB Instance in a VPC (p. 514)
Importing data into Amazon RDS PostgreSQL You can use several different tools to import data into your PostgreSQL DB instance on Amazon RDS.	Importing Data into PostgreSQL on Amazon RDS (p. 1302)
Setting up read-only Read Replicas (masters and standbys) PostgreSQL on Amazon RDS supports Read Replicas in both the same AWS Region and in a different AWS Region from the master instance.	Working with Read Replicas (p. 182) Working with PostgreSQL Read Replicas (p. 1298) Creating a Read Replica in a Different AWS Region (p. 188)

Task Area	Relevant Documentation
<p>Understanding security groups</p> <p>By default, DB instances are created with a firewall that prevents access to them. You therefore must create a security group with the correct IP addresses and network configuration to access the DB instance.</p> <p>In general, if your DB instance is on the EC2-Classic platform, you need to create a DB security group. If your DB instance is on the EC2-VPC platform, you need to create a VPC security group.</p>	<p>Determining Whether You Are Using the EC2-VPC or EC2-Classic Platform (p. 506)</p> <p>Controlling Access with Security Groups (p. 486)</p>
<p>Setting up parameter groups and features</p> <p>If your DB instance is going to require specific database parameters, you should create a parameter group before you create the DB instance.</p>	<p>Working with DB Parameter Groups (p. 211)</p>
<p>Performing common DBA tasks for PostgreSQL</p> <p>Some of the more common tasks for PostgreSQL DBAs include:</p> <ul style="list-style-type: none"> • Creating Roles (p. 1322) • Managing PostgreSQL Database Access (p. 1323) • Working with PostgreSQL Parameters (p. 1323) • Working with PostgreSQL Autovacuum on Amazon RDS (p. 1332) • Audit Logging for a PostgreSQL DB Instance (p. 1340) • Working with PostGIS (p. 1343) • Using pgBadger for Log Analysis with PostgreSQL (p. 1345) • Using a Custom DNS Server for Outbound Network Access (p. 1348) 	<p>Common DBA Tasks for PostgreSQL (p. 1322)</p>
<p>Connecting to your PostgreSQL DB instance</p> <p>After creating a security group and associating it to a DB instance, you can connect to the DB instance using any standard SQL client application such as pgadmin III.</p>	<p>Connecting to a DB Instance</p> <p>Running the PostgreSQL Database Engine (p. 1270)</p> <p>Using SSL with a PostgreSQL DB Instance (p. 1411)</p>
<p>Backing up and restoring your DB instance</p> <p>You can configure your DB instance to take automated backups, or take manual snapshots, and then restore instances from the backups or snapshots.</p>	<p>Backing Up and Restoring Amazon RDS DB Instances (p. 253)</p>
<p>Monitoring the activity and performance of your DB instance</p> <p>You can monitor a PostgreSQL DB instance by using CloudWatch Amazon RDS metrics, events, and enhanced monitoring.</p>	<p>Viewing DB Instance Metrics (p. 304)</p> <p>Viewing Amazon RDS Events (p. 388)</p>
<p>Upgrading the PostgreSQL database version</p> <p>You can do both major and minor version upgrades for your PostgreSQL DB instance.</p>	<p>Upgrading a PostgreSQL DB Instance (p. 1411)</p> <p>Major Version Upgrades for PostgreSQL (p. 1290)</p>

Task Area	Relevant Documentation
Working with log files You can access the log files for your PostgreSQL DB instance.	PostgreSQL Database Log Files (p. 419)
Understanding the best practices for PostgreSQL DB instances Find some of the best practices for working with PostgreSQL on Amazon RDS.	Best Practices for Working with PostgreSQL (p. 101)

Creating a DB Instance Running the PostgreSQL Database Engine

The basic building block of Amazon RDS is the DB instance. This is the environment in which you run your PostgreSQL databases.

Important

Before you can create or connect to a DB instance, you must complete the tasks in [Setting Up for Amazon RDS \(p. 5\)](#).

Note

If you are using the console, a new console interface is available for database creation. Choose either the **New Console** or the **Original Console** instructions based on the console that you are using. The **New Console** instructions are open by default.

New Console

To launch a PostgreSQL DB instance

You can create a DB instance running PostgreSQL with the AWS Management Console with **Easy Create** enabled or disabled. With **Easy Create** enabled, you specify only the DB engine type, DB instance size, and DB instance identifier. **Easy Create** uses the default setting for other configuration options. With **Easy Create** not enabled, you specify more configuration options when you create a database, including ones for availability, security, backups, and maintenance.

Note

For this example, **Standard Create** is enabled, and **Easy Create** isn't enabled. For information about creating a PostgreSQL DB instance with **Easy Create** enabled, see [Creating a PostgreSQL DB Instance and Connecting to a Database on a PostgreSQL DB Instance \(p. 55\)](#).

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the AWS Management Console, choose the AWS Region where you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
4. Choose **Create database**.
5. In **Choose a database creation method**, choose **Standard Create**.
6. In **Engine options**, choose **PostgreSQL**.

Create database

Choose a database creation method [Info](#)

Standard Create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy Create

Use recommended best-practice configurations. Some configuration options can be changed after the database is created.

Engine options

Engine type [Info](#)

Amazon Aurora



MySQL



MariaDB



PostgreSQL



Oracle



Microsoft SQL Server



Version [Info](#)

PostgreSQL 12

7. In **Templates**, choose the template that matches your use case. If you choose **Production**, the following are also chosen in a later step:

- **Multi-AZ failover option**
- **Provisioned IOPS storage option**
- **Enable deletion protection option**

We recommend these features for any production environment.

8. To enter your master password, do the following:

- a. In the **Settings** section, open **Credential Settings**.
- b. Clear the **Auto generate a password** check box.
- c. (Optional) Change the **Master username** value and enter the same password in **Master password** and **Confirm password**.

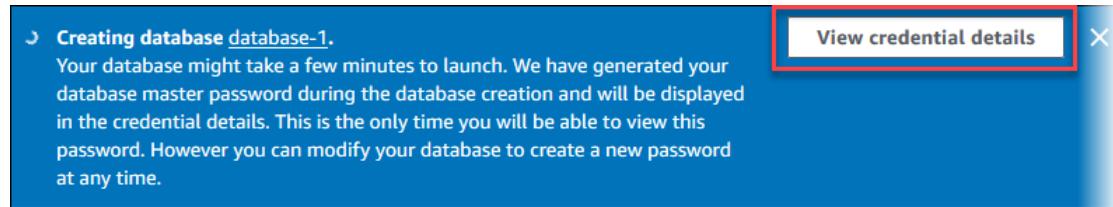
By default, the new DB instance uses an automatically generated password for the master user.

9. For the remaining sections, specify your DB instance settings. For information about each setting, see [Settings for PostgreSQL DB Instances \(p. 1265\)](#).

10. Choose **Create database**.

If you chose to use an automatically generated password, the **View credential details** button appears on the **Databases** page.

To view the master user name and password for the DB instance, choose **View credential details**.



To connect to the DB instance as the master user, use the user name and password that appear.

Important

You can't view the master user password again. If you don't record it, you might have to change it. If you need to change the master user password after the DB instance is available, you can modify the DB instance to do so. For more information about modifying a DB instance, see [Modifying a DB Instance Running the PostgreSQL Database Engine \(p. 1279\)](#).

11. For **Databases**, choose the name of the new PostgreSQL DB instance.

On the RDS console, the details for the new DB instance appear. The DB instance has a status of **creating** until the DB instance is created and ready for use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and storage allocated, it can take several minutes for the new instance to be available.

A screenshot of the Amazon RDS 'Summary' page for a DB instance named 'database-1'. The instance is currently in the 'Creating' state, which is highlighted with a red circle. Other details shown include the DB identifier 'database-1', CPU type 'CPU', engine 'PostgreSQL', and region 'Region -'. Below the summary, there are tabs for 'Connectivity & security', 'Monitoring', 'Logs & events', 'Configuration', and 'Maintenance & b...'. The 'Connectivity & security' tab is active.

DB identifier	CPU	Info	Class
database-1		Creating	db.t2...
Role	Current activity	Engine	Region
Instance		PostgreSQL	-

Original Console

To launch a PostgreSQL DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the top-right corner of the AWS Management Console, choose the AWS Region where you want to create the DB instance.

3. In the navigation pane, choose **Databases**.

If the navigation pane is closed, choose the menu icon at the top left to open it.

4. Choose **Create database** to start open on the **Select engine** page.

Create database

Choose a database creation method Info

Standard Create
You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy Create
Use recommended best-practice configurations. Some configuration options can be changed after the database is created.

Engine options

Engine type Info

Amazon Aurora 

MySQL 

MariaDB 

PostgreSQL 

Oracle 

Microsoft SQL Server 

Version Info

5. On the **Select engine** page, choose the PostgreSQL icon, and then choose **Next**.
6. Next, the **Use case** page asks if you are planning to use the DB instance you are creating for production. If you are, choose **Production**. If you choose this option, the following are also chosen in a later step:
 - **Multi-AZ failover option**
 - **Provisioned IOPS storage option**
 - **Enable deletion protection option**

Choose **Next** when you are finished.

7. On the **Specify DB Details** page, specify your DB instance information. Choose **Next** when you are finished.

For This Parameter	Do This
License Model	PostgreSQL has only one license model. Choose postgresql-license to use the general license agreement for PostgreSQL.
DB Engine Version	Choose the version of PostgreSQL you want to use.
DB Instance Class	Choose db.t2.small for a configuration that equates to 2 GiB memory, 1 ECU (1 virtual core with 1 ECU), 64-bit platform, and moderate I/O capacity. For more information about all the DB instance class options, see Choosing the DB Instance Class (p. 107) .
Multi-AZ Deployment	<p>Choose Yes to have a standby replica of your DB instance created in another Availability Zone for failover support. We recommend Multi-AZ for production workloads to maintain high availability. For development and testing, you can choose No.</p> <p>For more information, see High Availability (Multi-AZ) for Amazon RDS (p. 144).</p>
Storage autoscaling	<p>Enable storage autoscaling to enable Amazon RDS to automatically increase storage when needed to avoid having your DB instance run out of storage space.</p> <p>Maximum storage threshold to set the upper limit for Amazon RDS to automatically increase storage to for your DB instance.</p> <p>For more information, see Amazon RDS DB Instance Storage (p. 135).</p>
Storage Type	Choose the storage type General Purpose (SSD) . For more information about storage, see Amazon RDS DB Instance Storage (p. 135) .
Allocated Storage	Type 20 to allocate 20 GiB of storage for your database. In some cases, allocating a higher amount of storage for your DB instance than the size of your database can improve I/O performance. For more information about storage allocation, see Amazon Relational Database Service Features .
DB Instance Identifier	Type a name for the DB instance that is unique for your account in the AWS Region you chose. You can add some intelligence to the name, such as including the AWS Region and DB engine you chose, for example postgresql-test .
Master Username	Type a name using alphanumeric characters to use as the master user name to log on to your DB instance. For information on the default privileges granted to the master user name, see Amazon RDS for PostgreSQL Versions and Extensions (p. 1364)

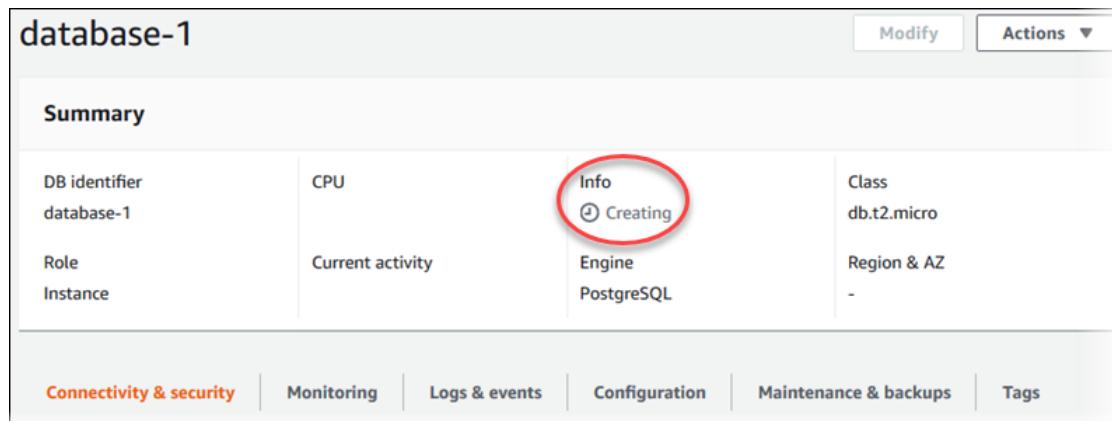
For This Parameter	Do This
Master Password and Confirm Password	Type a password that contains from 8 to 128 printable ASCII characters (excluding /," and @) for your master password, then type the password again in the Confirm Password box.

8. On the **Configure Advanced Settings** page, provide additional information that RDS needs to launch the PostgreSQL DB instance. The table shows settings for an example DB instance. Specify your DB instance information, then choose **Create database**.

For This Parameter	Do This
VPC	This setting depends on the platform you are on. If you are a new customer to AWS, choose the default VPC shown. If you are creating a DB instance on the previous E2-Classic platform that does not use a VPC, choose Not in VPC . For more information about VPC, see Amazon Virtual Private Cloud VPCs and Amazon RDS (p. 505) .
Subnet Group	This setting depends on the platform you are on. If you are a new customer to AWS, choose default , which is the default DB subnet group that was created for your account. If you are creating a DB instance on the previous E2-Classic platform and you want your DB instance in a specific VPC, choose the DB subnet group you created for that VPC. For more information about VPC, see Amazon Virtual Private Cloud VPCs and Amazon RDS (p. 505) .
Publicly Accessible	Choose Yes to give the DB instance a public IP address, meaning that it is accessible outside the VPC; otherwise, choose No , so the DB instance is only accessible from inside the VPC. For more information about hiding DB instances from public access, see Hiding a DB Instance in a VPC from the Internet (p. 515) .
Availability Zone	Use the default value of No Preference unless you want to specify an Availability Zone.
VPC Security Group	If you are a new customer to AWS, choose the default VPC. If you created a VPC security group, choose the VPC security group you previously created. When you choose Create new VPC security group in the RDS console, a new security group is created with an inbound rule that allows access to the DB instance from the IP address detected in your browser.
Database Name	Type a name for your database of up to 63 alpha-numeric characters. If you do not provide a name, the default "postgres" database is created. To create additional databases, connect to the DB instance and use the SQL command <code>CREATE DATABASE</code> . For more information about connecting to the DB instance, see Connecting to a DB Instance Running the PostgreSQL Database Engine (p. 1270) .

For This Parameter	Do This
Database Port	Specify a port you want to use to access the database. PostgreSQL installations default to port 5432.
DB Parameter Group	Use the default value unless you have created your own parameter group.
Option Group	Use the default value unless you have created your own option group.
Copy Tags To Snapshots	Choose this option to have any DB instance tags copied to a DB snapshot when you create a snapshot. For more information, see Tagging Amazon RDS Resources (p. 177) .
Enable Encryption	Choose Yes to enable encryption at rest for this DB instance. For more information, see Encrypting Amazon RDS Resources (p. 429) .
Backup Retention Period	Set the number of days you want automatic backups of your database to be retained. For testing purposes, you can set this value to 1 .
Backup Window	Unless you have a specific time that you want to have your database backup, use the default of No Preference .
Enable Enhanced Monitoring	Choose Yes to enable real-time OS monitoring. Amazon RDS provides metrics in real time for the operating system (OS) that your DB instance runs on. You are only charged for Enhanced Monitoring that exceeds the free tier provided by Amazon CloudWatch Logs.
Monitoring Role	Choose Default to use the default IAM role.
Granularity	Choose 60 to monitor the instance every minute.
Auto minor version upgrade	Choose Enable auto minor version upgrade to enable your DB instance to receive preferred minor DB engine version upgrades automatically when they become available. Amazon RDS performs automatic minor version upgrades in the maintenance window.
Maintenance Window	Choose the 30-minute window in which pending modifications to your DB instance are applied. If the time period doesn't matter, choose No Preference .
Enable deletion protection	Enable deletion protection to prevent your DB instance from being deleted. If you create a production DB instance with the AWS Management Console, deletion protection is enabled by default. For more information, see Deleting a DB Instance (p. 174) .

9. On the final page, choose **Create database**.
10. On the Amazon RDS console, the new DB instance appears in the list of DB instances. The DB instance will have a status of **creating** until the DB instance is created and ready for use. When the state changes to **available**, you can connect to the DB instance. Depending on the DB instance class and store allocated, it could take several minutes for the new instance to be available.



AWS CLI

To create a PostgreSQL DB instance, use the AWS CLI `create-db-instance` command with the following parameters:

- `--db-instance-identifier`
- `--allocated-storage`
- `--db-instance-class`
- `--engine`
- `--master-username`
- `--master-user-password`

Example

For Linux, OS X, or Unix:

```
aws rds create-db-instance
  --db-instance-identifier pgdbinstance \
  --allocated-storage 20 \
  --db-instance-class db.t2.small \
  --engine postgres \
  --master-username masterawsuser \
  --master-user-password masteruserpassword
```

For Windows:

```
aws rds create-db-instance
  --db-instance-identifier pgdbinstance ^
  --allocated-storage 20 ^
  --db-instance-class db.t2.small ^
  --engine postgres ^
  --master-username masterawsuser ^
  --master-user-password masteruserpassword
```

This command should produce output similar to the following:

```
DBINSTANCE pgdbinstance db.t2.small postgres 20 sa creating 3 **** n 9.3
SECGROUP default active
PARAMGRP default.PostgreSQL9.3 in-sync
```

RDS API

To create a PostgreSQL DB instance, use the Amazon RDS API[CreateDBInstance](#) command with the following parameters:

- Engine = *postgres*
- DBInstanceIdentifier = *pgdbinstance*
- DBInstanceClass = *db.t2.small*
- AllocatedStorage = *20*
- BackupRetentionPeriod = *3*
- MasterUsername = *masterawsuser*
- MasterUserPassword = *masteruserpassword*

Example

```
https://rds.amazonaws.com/
    ?Action/CreateDBInstance
    &AllocatedStorage=20
    &BackupRetentionPeriod=3
    &DBInstanceClass=db.t2.small
    &DBInstanceIdentifier=pgdbinstance
    &DBName=mydatabase
    &DBSecurityGroups.member.1=mysecuritygroup
    &DBSubnetGroup=mydbsubnetgroup
    &Engine=postgres
    &MasterUserPassword=<masteruserpassword>
    &MasterUsername=<masterawsuser>
    &SignatureMethod=HmacSHA256
    &SignatureVersion=4
    &Version=2013-09-09
    &X-Amz-Algorithm=AWS4-HMAC-SHA256
    &X-Amz-Credential=AKIADQKE4SARGYLE/20140212/us-west-2/rds/aws4_request
    &X-Amz-Date=20140212T190137Z
    &X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
    &X-Amz-Signature=60d520ca0576c191b9eac8dbfe5617ebb6a6a9f3994d96437a102c0c2c80f88d
```

Settings for PostgreSQL DB Instances

The following table contains details about settings that you choose when you create a PostgreSQL DB instance.

Setting	Setting Description
Allocated storage	The amount of storage to allocate for your DB instance (in gigabytes). In some cases, allocating a higher amount of storage for your DB instance than the size of your database can improve I/O performance. For more information, see Amazon RDS DB Instance Storage (p. 135) .
Auto minor version upgrade	Enable auto minor version upgrade to enable your DB instance to receive preferred minor DB engine version upgrades automatically when they become available. Amazon RDS performs automatic minor version upgrades in the maintenance window.

Setting	Setting Description
Availability zone	The Availability Zone for your DB instance. Use the default value of No Preference unless you want to specify an Availability Zone. For more information, see Regions and Availability Zones (p. 132) .
Backup retention period	The number of days that you want automatic backups of your DB instance to be retained. For any nontrivial DB instance, set this value to 1 or greater. For more information, see Working With Backups (p. 254) .
Backup window	The time period during which Amazon RDS automatically takes a backup of your DB instance. Unless you have a specific time that you want to have your database backed up, use the default of No Preference . For more information, see Working With Backups (p. 254) .
Copy tags To snapshots	This option copies any DB instance tags to a DB snapshot when you create a snapshot. For more information, see Tagging Amazon RDS Resources (p. 177) .
Database authentication	The database authentication option you want to use. Choose Password authentication to authenticate database users with database passwords only. Choose Password and IAM DB authentication to authenticate database users with database passwords and user credentials through IAM users and roles. For more information, see IAM Database Authentication for MySQL and PostgreSQL (p. 461) .
Database name	The name for the database on your DB instance. The name must contain 1 to 64 alphanumeric characters. If you don't provide a name, Amazon RDS doesn't create a database on the DB instance you are creating.
Database port	The port that you want to access the DB instance through. PostgreSQL installations default to port 3306. If you use a DB security group with your DB instance, this port value must be the same one you provided when creating the DB security group. The firewalls at some companies block connections to the default PostgreSQL port. If your company firewall blocks the default port, enter another port for your DB instance.
DB engine version	The version of PostgreSQL that you want to use.

Setting	Setting Description
DB instance class	<p>The configuration for your DB instance.</p> <p>If possible, choose an instance class large enough that a typical query working set can be held in memory. When working sets are held in memory, the system can avoid writing to disk, and this improves performance.</p> <p>For more information, see Choosing the DB Instance Class (p. 107).</p>
DB instance identifier	<p>The name for your DB instance. Your DB instance identifier can contain up to 63 alphanumeric characters, and must be unique for your account in the AWS Region you chose. You can add some intelligence to the name, such as including the AWS Region you chose, for example <code>postgresql-instance1</code>.</p>
DB parameter group	<p>A parameter group for your DB instance. You can choose the default parameter group or you can create a custom parameter group.</p> <p>For more information, see Working with DB Parameter Groups (p. 211).</p>
Deletion protection	<p>Enable deletion protection to prevent your DB instance from being deleted. If you create a production DB instance with the AWS Management Console, deletion protection is enabled by default.</p> <p>For more information, see Deleting a DB Instance (p. 174).</p>
Encryption	<p>Enable Encryption to enable encryption at rest for this DB instance.</p> <p>For more information, see Encrypting Amazon RDS Resources (p. 429).</p>
Enhanced Monitoring	<p>Enable enhanced monitoring to enable gathering metrics in real time for the operating system that your DB instance runs on.</p> <p>For more information, see Enhanced Monitoring (p. 307).</p>
License model	<p>PostgreSQL has only one license model, general-public-license, the general license agreement for PostgreSQL.</p>
Log exports	<p>The types of PostgreSQL database log files to generate.</p> <p>For more information, see PostgreSQL Database Log Files (p. 419).</p>
Maintenance window	<p>The 30-minute window in which pending modifications to your DB instance are applied. If the time period doesn't matter, choose No Preference.</p> <p>For more information, see The Amazon RDS Maintenance Window (p. 156).</p>

Setting	Setting Description
Master password	The password for your master user account. The password must contain from 8 to 41 printable ASCII characters (excluding /, ", a space, and @).
Master username	The name that you use as the master user name to log on to your DB instance.
Multi-AZ deployment	<p>Create a standby instance to create a passive secondary replica of your DB instance in another Availability Zone for failover support. We recommend Multi-AZ for production workloads to maintain high availability.</p> <p>For development and testing, you can choose Do not create a standby instance.</p> <p>For more information, see High Availability (Multi-AZ) for Amazon RDS (p. 144).</p>
Option group	<p>An option group for your DB instance. You can choose the default option group or you can create a custom option group.</p> <p>For more information, see Working with Option Groups (p. 196).</p>
Performance Insights	<p>Enable Performance Insights to monitor your DB instance load so that you can analyze and troubleshoot your database performance.</p> <p>Choose a retention period to determine how much rolling data history to keep. The default of seven days is in the free tier. Long-term retention (two years) is priced per vCPU per month.</p> <p>Choose a master key to use to protect the key used to encrypt this database volume. Choose from the master keys in your account, or enter the key from a different account.</p> <p>For more information, see Using Amazon RDS Performance Insights (p. 319).</p>
Publicly accessible	<p>Yes to give the DB instance a public IP address, meaning that it's accessible outside the VPC. To be publicly accessible, the DB instance also has to be in a public subnet in the VPC.</p> <p>No to make the DB instance accessible only from inside the VPC.</p> <p>For more information, see Hiding a DB Instance in a VPC from the Internet (p. 515).</p>

Setting	Setting Description
Storage autoscaling	<p>Enable storage autoscaling to enable Amazon RDS to automatically increase storage when needed to avoid having your DB instance run out of storage space.</p> <p>Use Maximum storage threshold to set the upper limit for Amazon RDS to automatically increase storage for your DB instance. The default is 1,000 GiB.</p> <p>For more information, see Managing Capacity Automatically with Amazon RDS Storage Autoscaling (p. 232).</p>
Storage type	<p>The storage type for your DB instance.</p> <p>For more information, see Amazon RDS Storage Types (p. 135).</p>
Subnet group	<p>This setting depends on the platform that you are on. If you are a new customer to AWS, keep the subnet group set to default, which is the default DB subnet group that was created for your account.</p> <p>If you are creating a DB instance on the earlier E2-Classic platform, you might want your DB instance in a specific VPC. In this case, choose the DB subnet group that you created for that VPC.</p>
Virtual Private Cloud (VPC)	<p>This setting depends on the platform that you are on. If you are a new customer to AWS, choose the default VPC shown. If you are creating a DB instance on the earlier E2-Classic platform that doesn't use a VPC, choose Not in VPC.</p> <p>For more information, see Amazon Virtual Private Cloud VPCs and Amazon RDS (p. 505).</p>
VPC security groups	<p>If you are a new customer to AWS, Create new to create a new VPC security group. Otherwise, Choose existing, then choose from security groups that you previously created.</p> <p>When you choose Create new in the RDS console, a new security group is created. This new security group has an inbound rule that allows access to the DB instance from the IP address detected in your browser.</p> <p>For more information, see Working with DB Security Groups (EC2-Classic Platform) (p. 491).</p>

Connecting to a DB Instance Running the PostgreSQL Database Engine

After Amazon RDS provisions your DB instance, you can use any standard SQL client application to connect to the instance. To list the details of an Amazon RDS DB instance, you can use the AWS Management Console, the AWS CLI [describe-db-instances](#) command, or the Amazon RDS API [DescribeDBInstances](#) operation. You need the following information to connect:

- The host or host name for the DB instance, for example:

```
myinstance.123456789012.us-east-1.rds.amazonaws.com
```

- The port on which the DB instance is listening. For example, the default PostgreSQL port is 5432.
- The user name and password for the DB instance.

Following are two ways to connect to a PostgreSQL DB instance. The first example uses pgAdmin, a popular open-source administration and development tool for PostgreSQL. The second example uses psql, a command line utility that is part of a PostgreSQL installation.

Using pgAdmin to Connect to a PostgreSQL DB Instance

You can use the open-source tool pgAdmin to connect to a PostgreSQL DB instance.

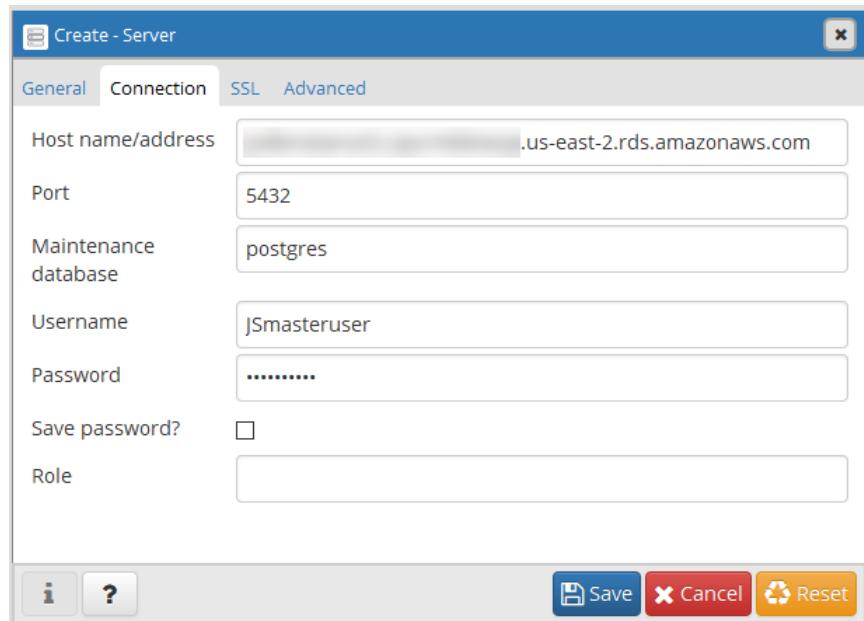
To connect to a PostgreSQL DB instance using pgAdmin

1. Find the endpoint (DNS name) and port number for your DB Instance.
 - a. Open the RDS console and then choose **Databases** to display a list of your DB instances.
 - b. Choose the PostgreSQL DB instance name to display its details.
 - c. On the **Connectivity & security** tab, copy the endpoint. Also, note the port number. You need both the endpoint and the port number to connect to the DB instance.

The screenshot shows the AWS RDS 'Summary' page for a database named 'database-1'. The 'Connectivity & security' tab is selected. The 'Endpoint' field is highlighted with a red oval.

DB identifier	database-1
Role	Instance
Connectivity & security	
Endpoint	database-1.us-west-1.rds.amazonaws.com
Port	5432

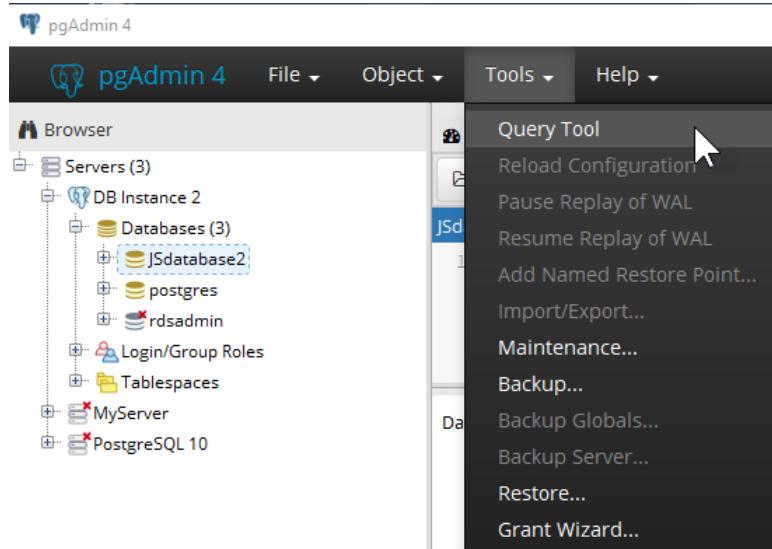
2. Install pgAdmin from <http://www.pgadmin.org/>. You can download and use pgAdmin without having a local instance of PostgreSQL on your client computer.
3. Launch the pgAdmin application on your client computer.
4. On the **Dashboard** tab, choose **Add New Server**.
5. In the **Create - Server** dialog box, type a name on the **General** tab to identify the server in pgAdmin.
6. On the **Connection** tab, type the following information from your DB instance:
 - For **Host**, type the endpoint, for example `mypostgresql.c6c8dntfzzhgv0.us-east-2.rds.amazonaws.com`.
 - For **Port**, type the assigned port.
 - For **Username**, type the user name that you entered when you created the DB instance.
 - For **Password**, type the password that you entered when you created the DB instance.



7. Choose **Save**.

If you have any problems connecting, see [Troubleshooting Connection Issues \(p. 1273\)](#).

8. To access a database in the pgAdmin browser, expand **Servers**, the DB instance, and **Databases**. Choose the DB instance's database name.



9. To open a panel where you can enter SQL commands, choose **Tools, Query Tool**.

Using psql to Connect to a PostgreSQL DB Instance

You can use a local instance of the psql command line utility to connect to a PostgreSQL DB instance. You need either PostgreSQL or the psql client installed on your client computer. To connect to your PostgreSQL DB instance using psql, you need to provide host information and access credentials.

Use one of the following formats to connect to a PostgreSQL DB instance on Amazon RDS. When you connect, you're prompted for a password. For batch jobs or scripts, use the `--no-password` option.

If this is the first time you are connecting to this DB instance, try using the default database name `postgres` for the `--dbname` option.

For Unix, use the following format.

```
psql \
--host=<DB instance endpoint> \
--port=<port> \
--username=<master user name> \
--password \
--dbname=<database name>
```

For Windows, use the following format.

```
psql ^
--host=<DB instance endpoint> ^
--port=<port> ^
--username=<master user name> ^
--password ^
--dbname=<database name>
```

For example, the following command connects to a database called `mypgdb` on a PostgreSQL DB instance called `mypostgresql` using fictitious credentials.

```
psql --host=mypostgresql.c6c8mwvfdgv0.us-west-2.rds.amazonaws.com --port=5432 --
username=awsuser --password --dbname=mypgdb
```

Troubleshooting Connection Issues

If you can't connect to the DB instance, the most common error is `Could not connect to server: Connection timed out`. If you receive this error, do the following:

- Check that the host name used is the DB instance endpoint and that the port number used is correct.
- Make sure that the DB instance's public accessibility is set to **Yes**.
- Check that the security group assigned to the DB instance has rules to allow access through any firewall your connection might go through. For example, if the DB instance was created using the default port of 5432, your company might have firewall rules blocking connections to that port from company devices.

To fix this failure, modify the DB instance to use a different port. Also, make sure that the security group applied to the DB instance allows connections to the new port.

- Check whether the DB instance was created using a security group that doesn't authorize connections from the device or Amazon EC2 instance where the application is running. For the connection to work, the security group you assigned to the DB instance at its creation must allow access to the DB instance. For example, if the DB instance was created in a VPC, it must have a VPC security group that authorizes connections.

You can add or edit an inbound rule in the security group. For **Source**, choose **My IP**. This allows access to the DB instance from the IP address detected in your browser. For more information, see [Amazon Virtual Private Cloud VPCs and Amazon RDS \(p. 505\)](#).

Alternatively, if the DB instance was created outside of a VPC, it must have a database security group that authorizes those connections.

By far the most common connection problem is with the security group's access rules assigned to the DB instance. If you used the default DB security group when you created the DB instance, the security group likely didn't have access rules that allow you to access the instance. For more information about Amazon RDS security groups, see [Controlling Access with Security Groups \(p. 486\)](#).

If you receive an error like `FATAL: database some-name does not exist` when connecting, try using the default database name `postgres` for the `--dbname` option.

Updating Applications to Connect to PostgreSQL DB Instances Using New SSL/TLS Certificates

As of September 19, 2019, Amazon RDS has published new Certificate Authority (CA) certificates for connecting to your RDS DB instances using Secure Socket Layer or Transport Layer Security (SSL/TLS). The previous CA certificates expire on March 5, 2020. Following, you can find information about updating your applications to use the new certificates. If your application connects to an RDS DB instance using SSL/TLS, you must take the following steps before **March 5, 2020**. Doing this means you can avoid interruption of connectivity between your applications and your RDS DB instances.

This topic can help you to determine whether any client applications use SSL/TLS to connect to your DB instances. If they do, you can further check whether those applications require certificate verification to connect.

Note

Some applications are configured to connect to PostgreSQL DB instances only if they can successfully verify the certificate on the server.

For such applications, you must update your client application trust stores to include the new CA certificates.

After you update your CA certificates in the client application trust stores, you can rotate the certificates on your DB instances. We strongly recommend testing these procedures in a development or staging environment before implementing them in your production environments.

For more information about certificate rotation, see [Rotating Your SSL/TLS Certificate \(p. 433\)](#). For information about using SSL/TLS with PostgreSQL DB instances, see [Using SSL with a PostgreSQL DB Instance \(p. 1411\)](#).

Topics

- [Determining Whether Applications Are Connecting to PostgreSQL DB Instances Using SSL \(p. 1275\)](#)
- [Determining Whether a Client Requires Certificate Verification in Order to Connect \(p. 1276\)](#)
- [Updating Your Application Trust Store \(p. 1277\)](#)
- [Using SSL/TLS Connections for Different Types of Applications \(p. 1278\)](#)

Determining Whether Applications Are Connecting to PostgreSQL DB Instances Using SSL

Check the DB instance configuration for the value of the `rds.force_ssl` parameter. By default, the `rds.force_ssl` parameter is set to 0 (off). If the `rds.force_ssl` parameter is set to 1 (on), clients are required to use SSL/TLS for connections. For more information about parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#).

If you are using RDS PostgreSQL version 9.5 or later major version and `rds.force_ssl` is not set to 1 (on), query the `pg_stat_ssl` view to check connections using SSL. For example, the following query returns only SSL connections and information about the clients using SSL.

```
select datname, usename, ssl, client_addr from pg_stat_ssl inner join pg_stat_activity on pg_stat_ssl.pid = pg_stat_activity.pid where ssl is true and usename<>'rdsadmin';
```

Only rows using SSL/TLS connections are displayed with information about the connection. The following is sample output.

```
datname | username | ssl | client_addr
-----+-----+-----+
benchdb | pgadmin | t   | 53.95.6.13
postgres | pgadmin | t   | 53.95.6.13
(2 rows)
```

This query displays only the current connections at the time of the query. The absence of results doesn't indicate that no applications are using SSL connections. Other SSL connections might be established at a different time.

If you are using RDS PostgreSQL version 9.3 or 9.4 and `rds.force_ssl` is not set to 1 (on), then to determine whether your applications are connecting with SSL, you must enable the `log_connections` parameter. This parameter logs SSL connection information when a client connects. The following is an example of an SSL connection entry in the error log.

```
2018-04-19 10:19:20 UTC:123.45.67.8(6789):[unknown]@[unknown]:[17196]:LOG: connection
received: host=123.45.67.8 port=12345
2018-04-19 10:19:20 UTC:123.45.67.8(6789):test_user@Test_DB:[17196]:LOG: connection
authorized: user=test_user database=Test_DB SSL enabled (protocol=xxxx, cipher=xxxx,
compression=off)
```

For more information about viewing the error log, see [Amazon RDS Database Log Files \(p. 390\)](#).

Determining Whether a Client Requires Certificate Verification in Order to Connect

When a client, such as `psql` or JDBC, is configured with SSL support, the client first tries to connect to the database with SSL by default. If the client can't connect with SSL, it reverts to connecting without SSL. The default `sslmode` mode used is different between libpq-based clients (such as `psql`) and JDBC. The libpq-based clients default to `prefer`, where JDBC clients default to `verify-full`. The certificate on the server is verified only when `sslrootcert` is provided with `sslmode` set to `require`, `verify-ca`, or `verify-full`. An error is thrown if the certificate is invalid.

Use `PGSSLROOTCERT` to verify the certificate with the `PGSSLMODE` environment variable, with `PGSSLMODE` set to `require`, `verify-ca`, or `verify-full`.

```
PGSSLMODE=require PGSSLROOTCERT=/fullpath/rds-ca-2019-root.pem psql -h
pgdbidentifier.cxxxxxxxxx.us-east-2.rds.amazonaws.com -U masteruser -d postgres
```

Use the `sslrootcert` argument to verify the certificate with `sslmode` in connection string format, with `sslmode` set to `require`, `verify-ca`, or `verify-full` to verify the certificate.

```
psql "host=pgdbidentifier.cxxxxxxxxx.us-east-2.rds.amazonaws.com sslmode=require
sslrootcert=/full/path/rds-ca-2019-root.pem user=masteruser dbname=postgres"
```

For example, in the preceding case, if you are using an invalid root certificate, then you see an error similar to the following on your client.

```
psql: SSL error: certificate verify failed
```

Updating Your Application Trust Store

For information about updating the trust store for PostgreSQL applications, see [Secure TCP/IP Connections with SSL](#) in the PostgreSQL documentation.

Note

When you update the trust store, you can retain older certificates in addition to adding the new certificates.

Updating Your Application Trust Store for JDBC

You can update the trust store for applications that use JDBC for SSL/TLS connections.

To update the trust store for JDBC applications

1. Download the 2019 root certificate that works for all AWS Regions and put the file in your trust store directory.

For information about downloading the root certificate, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).

2. Convert the certificate to .der format using the following command.

```
openssl x509 -outform der -in rds-ca-2019-root.pem -out rds-ca-2019-root.der
```

Replace the file name with the one that you downloaded.

3. Import the certificate into the key store using the following command.

```
keytool -import -alias rds-root -keystore clientkeystore -file rds-ca-2019-root.der
```

4. Confirm that the key store was updated successfully.

```
keytool -list -v -keystore clientkeystore.jks
```

Enter the key store password when you are prompted for it.

Your output should contain the following:

```
rds-root,date, trustedCertEntry,  
Certificate fingerprint (SHA1):  
D4:0D:DB:29:E3:75:0D:FF:A6:71:C3:14:0B:BF:5F:47:8D:1C:80:96  
# This fingerprint should match the output from the below command  
openssl x509 -fingerprint -in rds-ca-2019-root.pem -noout
```

Using SSL/TLS Connections for Different Types of Applications

The following provides information about using SSL/TLS connections for different types of applications:

- **psql**

The client is invoked from the command line by specifying options either as a connection string or as environment variables. For SSL/TLS connections, the relevant options are `sslmode` (environment variable `PGSSLMODE`), `sslrootcert` (environment variable `PGSSLROOTCERT`).

For the complete list of options, see [Parameter Key Words](#) in the PostgreSQL documentation. For the complete list of environment variables, see [Environment Variables](#) in the PostgreSQL documentation.

- **pgAdmin**

This browser-based client is a more user-friendly interface for connecting to a PostgreSQL database.

For information about configuring connections, see the [pgAdmin documentation](#).

- **JDBC**

JDBC enables database connections with Java applications.

For general information about connecting to a PostgreSQL database with JDBC, see [Connecting to the Database](#) in the PostgreSQL documentation. For information about connecting with SSL/TLS, see [Configuring the Client](#) in the PostgreSQL documentation.

- **Python**

A popular Python library for connecting to PostgreSQL databases is `psycopg2`.

For information about using `psycopg2`, see the [psycopg2 documentation](#). For a short tutorial on how to connect to a PostgreSQL database, see [Psycopg2 Tutorial](#). You can find information about the options the `connect` command accepts in [The psycopg2 module content](#).

Important

After you have determined that your database connections use SSL/TLS and have updated your application trust store, you can update your database to use the rds-ca-2019 certificates. For instructions, see step 3 in [Updating Your CA Certificate by Modifying Your DB Instance \(p. 434\)](#).

Modifying a DB Instance Running the PostgreSQL Database Engine

You can change the settings of a DB instance to accomplish tasks such as adding additional storage or changing the DB instance class. This topic guides you through modifying an Amazon RDS PostgreSQL DB instance, and describes the settings for PostgreSQL instances. For information about additional tasks, such as renaming, rebooting, deleting, tagging, or upgrading an Amazon RDS DB instance, see [Amazon RDS DB Instance Lifecycle \(p. 147\)](#). We recommend that you test any changes on a test instance before modifying a production instance so you better understand the impact of a change. This is especially important when upgrading database versions.

You can have the changes apply immediately or have them applied during the DB instance's next maintenance window. Applying changes immediately can cause an outage in some cases; for more information on the impact of the **Apply Immediately** option when modifying a DB instance, see [Modifying an Amazon RDS DB Instance \(p. 150\)](#).

Console

To modify a PostgreSQL DB instance

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Databases**, and then choose the DB instance that you want to modify.
3. Choose **Modify**. The **Modify DB Instance** page appears.
4. Change any of the settings that you want. For information about each setting, see [Settings for PostgreSQL DB Instances \(p. 1280\)](#).
5. When all the changes are as you want them, choose **Continue**.
6. To apply the changes immediately, choose **Apply Immediately**. Choosing this option can cause an outage in some cases. For more information, see [Using the Apply Immediately Parameter \(p. 150\)](#).
7. On the confirmation page, review your changes. If they are correct, choose **Modify DB Instance** to save your changes.

Alternatively, choose **Back** to edit your changes, or choose **Cancel** to cancel your changes.

AWS CLI

To modify a PostgreSQL DB instance, use the AWS CLI command `modify-db-instance`.

Example

The following code modifies pgdbinstance by setting the backup retention period to 1 week (7 days) and disabling automatic minor version upgrades. These changes are applied during the next maintenance window.

Parameters

- `--db-instance-identifier`—the name of the DB instance
- `--backup-retention-period`—the number of days to retain automatic backups.
- `--auto-minor-version-upgrade`—allow automatic minor version upgrades. To disallow automatic minor version upgrades, use `--no-auto-minor-version-upgrade`.
- `--no-apply-immediately`—apply changes during the next maintenance window. To apply changes immediately, use `--apply-immediately`.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier pgdbinstance \
--backup-retention-period 7 \
--auto-minor-version-upgrade \
--no-apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier pgdbinstance ^
--backup-retention-period 7 ^
--auto-minor-version-upgrade ^
--no-apply-immediately
```

API

To modify a PostgreSQL DB instance, use the [ModifyDBInstance](#) operation.

Example

The following code modifies pgdbinstance by setting the backup retention period to 1 week (7 days) and disabling automatic minor version upgrades. These changes are applied during the next maintenance window.

Parameters

- **DBInstanceIdentifier**—the name of the DB instance
- **BackupRetentionPeriod**—the number of days to retain automatic backups.
- **AutoMinorVersionUpgrade=true**—allow automatic minor version upgrades. To disallow automatic minor version upgrades, set the value to **false**.
- **ApplyImmediately=false**—apply changes during the next maintenance window. To apply changes immediately, set the value to **true**.

```
https://rds.us-east-1.amazonaws.com/
?Action=ModifyDBInstance
&ApplyImmediately=false
&AutoMinorVersionUpgrade=true
&BackupRetentionPeriod=7
&DBInstanceIdentifier=mydbinstance
&SignatureMethod=HmacSHA256
&SignatureVersion=4
&Version=2013-09-09
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20131016/us-east-1/rds/aws4_request
&X-Amz-Date=20131016T233051Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=087a8eb41cb1ab0fc9ec1575f23e73757fffc6a1e42d7d2b30b9cc0be988cff97
```

Settings for PostgreSQL DB Instances

The following table contains details about which settings you can modify, which settings you can't modify, when the changes can be applied, and whether the changes cause downtime for the DB instance.

Setting	Setting Description	When the Change Occurs	Downtime Notes
Allocated storage	<p>The storage, in gigabytes, that you want to allocate for your DB instance. You can only increase the allocated storage, you can't reduce the allocated storage.</p> <p>You can't modify allocated storage if the DB instance status is storage-optimization or if the allocated storage for the DB instance has been modified in the last six hours.</p> <p>The maximum storage allowed depends on the storage type. For more information, see Amazon RDS DB Instance Storage (p. 135).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	No downtime. Performance might be degraded during the change.
Auto minor version upgrade	Enable auto minor version upgrade to enable your DB instance to receive preferred minor DB engine version upgrades automatically when they become available. Amazon RDS performs automatic minor version upgrades in the maintenance window.	–	–
Backup retention period	<p>The number of days that automatic backups are retained. To disable automatic backups, set the backup retention period to 0.</p> <p>For more information, see Working With Backups (p. 254).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false and you change the setting from a nonzero value to another nonzero value, the change is applied asynchronously, as soon as possible. Otherwise, the change occurs during the next maintenance window.</p>	An outage occurs if you change from 0 to a nonzero value, or from a nonzero value to 0.
Backup window	<p>The time range during which automated backups of your databases occur. The backup window is a start time in Universal Coordinated Time (UTC), and a duration in hours.</p> <p>For more information, see Working With Backups (p. 254).</p>	The change is applied asynchronously, as soon as possible.	–
Certificate authority	The certificate that you want to use.	If Apply Immediately is set to true, the change occurs immediately.	An outage occurs during this change.

Setting	Setting Description	When the Change Occurs	Downtime Notes
		If Apply Immediately is set to false, the change occurs during the next maintenance window.	
Copy tags to snapshots	<p>If you have any DB instance tags, this option copies them when you create a DB snapshot.</p> <p>For more information, see Tagging Amazon RDS Resources (p. 177).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–
Database authentication	<p>The database authentication option you want to use.</p> <p>Choose Password authentication to authenticate database users with database passwords only.</p> <p>Choose Password and IAM DB authentication to authenticate database users with database passwords and user credentials through IAM users and roles. For more information, see IAM Database Authentication for MySQL and PostgreSQL (p. 461).</p>	<p>If Apply immediately is set to true, the change occurs immediately.</p> <p>If Apply immediately is set to false, the change occurs during the next maintenance window.</p>	–
Database port	<p>The port that you want to use to access the database.</p> <p>The port value must not match any of the port values specified for options in the option group for the DB instance.</p>	The change occurs immediately. This setting ignores the Apply Immediately setting.	The DB instance is rebooted immediately.
DB engine version	The version of the PostgreSQL database engine that you want to use. Before you upgrade your production DB instances, we recommend that you test the upgrade process on a test instance to verify its duration and to validate your applications.	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change.
DB instance class	<p>The DB instance class that you want to use.</p> <p>For more information, see Choosing the DB Instance Class (p. 107)</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change.

Setting	Setting Description	When the Change Occurs	Downtime Notes
DB instance identifier	<p>The DB instance identifier. This value is stored as a lowercase string.</p> <p>For more information about the effects of renaming a DB instance, see Renaming a DB Instance (p. 163).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change. The DB instance is rebooted.
DB parameter group	<p>The parameter group that you want associated with the DB instance.</p> <p>For more information, see Working with DB Parameter Groups (p. 211).</p>	The parameter group change occurs immediately.	<p>An outage doesn't occur during this change. When you change the parameter group, changes to some parameters are applied to the DB instance immediately without a reboot. Changes to other parameters are applied only after the DB instance is rebooted.</p> <p>For more information, see Rebooting a DB Instance (p. 166).</p>
Deletion protection	<p>Enable deletion protection to prevent your DB instance from being deleted.</p> <p>For more information, see Deleting a DB Instance (p. 174).</p>	–	–
Enhanced monitoring	<p>Enable enhanced monitoring to enable gathering metrics in real time for the operating system that your DB instance runs on.</p> <p>For more information, see Enhanced Monitoring (p. 307).</p>	–	–
License model	You can't change the license model because PostgreSQL has only one license model.	–	–
Log exports	<p>The types of PostgreSQL database log files to publish to Amazon CloudWatch Logs.</p> <p>For more information, see PostgreSQL Database Log Files (p. 419).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Maintenance window	<p>The time range during which system maintenance occurs. System maintenance includes upgrades, if applicable. The maintenance window is a start time in Universal Coordinated Time (UTC), and a duration in hours.</p> <p>If you set the window to the current time, there must be at least 30 minutes between the current time and end of the window to ensure any pending changes are applied.</p> <p>For more information, see The Amazon RDS Maintenance Window (p. 156).</p>	The change occurs immediately. This setting ignores the Apply Immediately setting.	If there are one or more pending actions that cause an outage, and the maintenance window is changed to include the current time, then those pending actions are applied immediately, and an outage occurs.
Multi-AZ deployment	<p>Yes to deploy your DB instance in multiple Availability Zones. Otherwise, No.</p> <p>For more information, see Regions and Availability Zones (p. 132).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	–
New master password	The password for your master user. The password must contain from 8 to 30 alphanumeric characters.	The change is applied asynchronously, as soon as possible. This setting ignores the Apply Immediately setting.	–
Option group	No options are available for PostgreSQL DB instances.	–	–
	For more information, see Working with Option Groups (p. 196) .		

Setting	Setting Description	When the Change Occurs	Downtime Notes
Performance Insights	<p>Enable Performance Insights to monitor your DB instance load so that you can analyze and troubleshoot your database performance.</p> <p>Choose a retention period to determine how much rolling data history to retain. The default of seven days is in the free tier. Long-term retention (two years) is priced per vCPU per month.</p> <p>You can't change the master key after the database is created.</p> <p>For more information, see Using Amazon RDS Performance Insights (p. 319).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–
Public accessibility	<p>Yes to give the DB instance a public IP address, meaning that it's accessible outside the VPC. To be publicly accessible, the DB instance also has to be in a public subnet in the VPC.</p> <p>No to make the DB instance accessible only from inside the VPC.</p> <p>For more information, see Hiding a DB Instance in a VPC from the Internet (p. 515).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–
Security group	<p>The security group that you want associated with the DB instance.</p> <p>For more information, see Working with DB Security Groups (EC2-Classic Platform) (p. 491).</p>	The change is applied asynchronously, as soon as possible. This setting ignores the Apply Immediately setting.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Storage autoscaling	<p>Enable storage autoscaling to enable Amazon RDS to automatically increase storage when needed to avoid having your DB instance run out of storage space.</p> <p>Use Maximum storage threshold to set the upper limit for Amazon RDS to automatically increase storage for your DB instance. The default is 1,000 GiB.</p> <p>For more information, see Managing Capacity Automatically with Amazon RDS Storage Autoscaling (p. 232).</p>	The change occurs immediately. This setting ignores the Apply immediately setting.	–

Setting	Setting Description	When the Change Occurs	Downtime Notes
Storage type	<p>The storage type that you want to use.</p> <p>For more information, see Amazon RDS Storage Types (p. 135).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	<p>The following changes all result in a brief outage while the process starts. After that, you can use your database normally while the change takes place.</p> <ul style="list-style-type: none"> • From General Purpose (SSD) to Magnetic. • From General Purpose (SSD) to Provisioned IOPS (SSD). The outage only happens if the DB instance is single-AZ and you are using a custom parameter group. There is no outage for a multi-AZ DB instance. • From Magnetic to General Purpose (SSD). • From Magnetic to Provisioned IOPS (SSD). • From Provisioned IOPS (SSD) to Magnetic. • From Provisioned IOPS (SSD) to General Purpose (SSD). The outage only happens if the DB instance is single-AZ and you are using a custom parameter group. There is no outage for a multi-AZ DB instance.

Setting	Setting Description	When the Change Occurs	Downtime Notes
Subnet group	<p>The subnet group for the DB instance. You can use this setting to move your DB instance to a different VPC. If your DB instance is not in a VPC, you can use this setting to move your DB instance into a VPC.</p> <p>For more information, see Moving a DB Instance Not in a VPC into a VPC (p. 520).</p>	<p>If Apply Immediately is set to true, the change occurs immediately.</p> <p>If Apply Immediately is set to false, the change occurs during the next maintenance window.</p>	An outage occurs during this change. The DB instance is rebooted.

Upgrading the PostgreSQL DB Engine for Amazon RDS

When Amazon RDS supports a new version of a database engine, you can upgrade your DB instances to the new version. There are two kinds of upgrades for PostgreSQL DB instances: major version upgrades and minor version upgrades.

Major version upgrades can contain database changes that are not backward-compatible with existing applications. As a result, you must manually perform major version upgrades of your DB instances. You can initiate a major version upgrade by modifying your DB instance. However, before you perform a major version upgrade, we recommend that you follow the steps described in [Major Version Upgrades for PostgreSQL \(p. 1290\)](#).

In contrast, *minor version upgrades* include only changes that are backward-compatible with existing applications. You can initiate a minor version upgrade manually by modifying your DB instance. Alternatively, you can enable the **Auto minor version upgrade** option when creating or modifying a DB instance. Doing so means that your DB instance is automatically upgraded after Amazon RDS tests and approves the new version. For more details, see [Automatic Minor Version Upgrades for PostgreSQL \(p. 1294\)](#). For information about manually performing a minor version upgrade, see [Manually Upgrading the Engine Version \(p. 159\)](#).

If your PostgreSQL DB instance is using read replicas, you must upgrade all of the read replicas before upgrading the source instance. If the DB instance is in a Multi-AZ deployment, both the writer and standby replicas are upgraded, and the instance might not be available until the upgrade is complete.

Topics

- [Overview of Upgrading PostgreSQL \(p. 1289\)](#)
- [Major Version Upgrades for PostgreSQL \(p. 1290\)](#)
- [Automatic Minor Version Upgrades for PostgreSQL \(p. 1294\)](#)

Overview of Upgrading PostgreSQL

To safely upgrade your DB instances, Amazon RDS uses the `pg_upgrade` utility described in the [PostgreSQL documentation](#).

Amazon RDS takes two DB snapshots during the upgrade process if your backup retention period is greater than 0. The first DB snapshot is of the DB instance before any upgrade changes have been made. If the upgrade doesn't work for your databases, you can restore this snapshot to create a DB instance running the old version. The second DB snapshot is taken after the upgrade completes.

Note

Amazon RDS takes DB snapshots during the upgrade process only if you have set the backup retention period for your DB instance to a number greater than 0. To change your backup retention period, see [Modifying a DB Instance Running the PostgreSQL Database Engine \(p. 1279\)](#).

If your DB instance is in a Multi-AZ deployment, both the primary writer DB instance and standby DB instances are upgraded. The writer and standby DB instances are upgraded at the same time, and you experience an outage until the upgrade is complete.

After an upgrade is complete, you can't revert to the previous version of the database engine. If you want to return to the previous version, restore the DB snapshot that was taken before the upgrade to create a new DB instance.

Major Version Upgrades for PostgreSQL

Major version upgrades can contain database changes that are not backward-compatible with previous versions of the database. This functionality can cause your existing applications to stop working correctly.

As a result, Amazon RDS doesn't apply major version upgrades automatically. To perform a major version upgrade, you modify your DB instance manually. Make sure that you thoroughly test any upgrade to verify that your applications work correctly before applying the upgrade to your production DB instances. When you do a PostgreSQL major version upgrade, we recommend that you follow the steps described in [How to Perform a Major Version Upgrade \(p. 1291\)](#).

Choosing from Multiple Major Versions

You can upgrade a PostgreSQL database to its next major version. From some PostgreSQL database versions, you can skip to a higher major version when upgrading. The following table lists the source PostgreSQL database versions and the associated target major versions available for upgrading.

Note

- Upgrade targets are enabled to a higher version released at the same time as the source minor version or later.
- If a database uses the PostGIS extension, you cannot skip major versions for some source to target combinations. For these circumstances, upgrade in steps to the next major version one at a time until you reach the desired target version.
- The pgRouting extension isn't supported for an upgrade that skips a major version to versions 11.x. A major version is skipped when the upgrade goes from versions 9.4.x, 9.5.x, or 9.6.x to versions 11.x. It's safe to drop the pgRouting extension and then add it again after an upgrade.
- The tsearch2 and chkpass extensions aren't supported in PostgreSQL 11 or later. If you are upgrading to version 11.x, drop these extensions before the upgrade.

Current Source Version	Major Version Targets	Additional Major Version Targets (without PostGIS Extension)
9.3.x	9.4.x	
9.3.23	9.4.x, 9.5.13	9.6.9
9.3.24	9.4.x, 9.5.14	9.6.10
9.3.25	9.4.x, 9.5.15	9.6.11
9.4.x	9.5.x	
9.4.20	9.5.x	11.1
9.4.21	9.5.x	10.7, 11.2
9.4.23	9.5.x	10.9, 11.4
9.4.24	9.5.x, 9.6.15, 10.10, 11.5	
9.5.x	9.6.x	

Current Source Version	Major Version Targets	Additional Major Version Targets (without PostGIS Extension)
9.5.15	9.6.x	11.1
9.5.16	9.6.x	10.7, 11.2
9.5.18	9.6.x	10.9, 11.4
9.5.19	9.6.x, 10.10, 11.5	
9.6.x	10.x	
9.6.11	10.x	11.1
9.6.12	10.x	11.2
9.6.14	10.x	11.4
9.6.15	10.x, 11.5	
10.x	11.x	

How to Perform a Major Version Upgrade

We recommend the following process when upgrading an Amazon RDS PostgreSQL DB instance:

1. **Have a version-compatible parameter group ready** – If you are using a custom parameter group, you have two options. You can specify a default parameter group for the new DB engine version. Or you can create your own custom parameter group for the new DB engine version.

If you associate a new parameter group with a DB instance, reboot the database after the upgrade completes. If the instance needs to be rebooted to apply the parameter group changes, the instance's parameter group status shows pending-reboot. You can view an instance's parameter group status in the console or by using a describe command, such as [describe-db-instances](#).

2. **Check for unsupported usage:**

- **Prepared transactions** – Commit or roll back all open prepared transactions before attempting an upgrade.

You can use the following query to verify that there are no open prepared transactions on your instance.

```
SELECT count(*) FROM pg_catalog.pg_prepared_xacts;
```

- **The line data type** – If you are upgrading an RDS PostgreSQL 9.3 instance, remove all uses of the line data type before attempting an upgrade. The line data type wasn't fully implemented in PostgreSQL until version 9.4.

To verify that there are no uses of the line data type, use the following query on each database to be upgraded.

```
SELECT count(*) FROM pg_catalog.pg_class c, pg_catalog.pg_namespace n,
pg_catalog.pg_attribute a
WHERE c.oid = a.attrelid
AND NOT a.attisdropped
AND a.atttypid = 'pg_catalog.line'::pg_catalog.regtype
```

```
AND c.relnamespace = n.oid
AND n.nspname !~ '^pg_temp_'
AND n.nspname !~ '^pg_toast_temp_'
AND n.nspname NOT IN ('pg_catalog', 'information_schema');
```

Note

To list all databases for an instance, use the following query.

```
SELECT d.datname FROM pg_catalog.pg_database d WHERE d.datallowconn = true;
```

- **Reg* data types** – Remove all uses of the *reg** data types before attempting an upgrade. Except for `regtype` and `regclass`, you can't upgrade the *reg** data types. The `pg_upgrade` utility can't persist this data type, which is used by Amazon RDS to do the upgrade.

To verify that there are no uses of unsupported *reg** data types, use the following query for each database.

```
SELECT count(*) FROM pg_catalog.pg_class c, pg_catalog.pg_namespace n,
pg_catalog.pg_attribute a
WHERE c.oid = a.attrelid
    AND NOT a.attisdropped
    AND a.atttypid IN ('pg_catalog.regproc'::pg_catalog.regtype,
'pg_catalog.regprocedure'::pg_catalog.regtype,
'pg_catalog.regoper'::pg_catalog.regtype,
'pg_catalog.regoperator'::pg_catalog.regtype,
'pg_catalog.regconfig'::pg_catalog.regtype,
'pg_catalog.regdictionary'::pg_catalog.regtype)
    AND c.relnamespace = n.oid
    AND n.nspname NOT IN ('pg_catalog', 'information_schema');
```

3. **Handle read replicas** – A read replica can't undergo a major version upgrade but the read replica's source instance can. If a read replica's source instance undergoes a major version upgrade, all read replicas for that source instance remain with the previous engine version. In this case, the read replicas can no longer replicate changes performed on the source instance.

We recommend that you either promote your read replicas, or delete and recreate them after the source instance has upgraded to a different major version. For more information, see [Working with PostgreSQL Read Replicas \(p. 1298\)](#).

4. **Perform a backup** – We recommend that you perform a backup before performing the major version upgrade so that you have a known restore point for your database. If your backup retention period is greater than 0, the upgrade process creates DB snapshots of your DB instance before and after upgrading. To change your backup retention period, see [Modifying a DB Instance Running the PostgreSQL Database Engine \(p. 1279\)](#). To perform a backup manually, see [Creating a DB Snapshot \(p. 264\)](#).

5. **Update certain extensions before the major version upgrade** – If you plan to skip a major version with the upgrade, you need to update certain extensions *before* performing the major version upgrade. Upgrading from versions 9.4.x, 9.5.x, or 9.6.x to versions 11.x skip a major version. The extensions to be updated include:

- `address_standardizer`
- `address_standardizer_data_us`
- `postGIS`
- `postgis_tiger_geocoder`
- `postgis_topology`

Run the following command for each extension you are using.

```
ALTER EXTENSION PostgreSQL-extension UPDATE TO 'new-version'
```

For the extension versions you can update to, see [PostgreSQL Extensions and Modules Supported on Amazon RDS \(p. 1383\)](#).

6. **Drop certain extensions before the major version upgrade** – An upgrade that skips a major version to version 11.x doesn't support updating the pgRouting extension. Upgrading from versions 9.4.x, 9.5.x, or 9.6.x to versions 11.x skip a major version. It's safe to drop the pgRouting extension and then reinstall it to a compatible version after the upgrade. For the extension versions you can update to, see [PostgreSQL Extensions and Modules Supported on Amazon RDS \(p. 1383\)](#).

The tsearch2 and chkpass extensions are no longer supported for PostgreSQL versions 11 or later. If you are upgrading to version 11.x, drop the tsearch2, and chkpass extensions before the upgrade.

7. **Perform an upgrade dry run** – We highly recommend testing a major version upgrade on a duplicate of your production database before attempting the upgrade on your production database. To create a duplicate test instance, you can either restore your database from a recent snapshot or do a point-in-time restore of your database to its latest restorable time. For more information, see [Restoring from a Snapshot \(p. 267\)](#) or [Restoring a DB Instance to a Specified Time \(p. 285\)](#). For details on performing the upgrade, see [Manually Upgrading the Engine Version \(p. 159\)](#).

During the major version upgrade, the public and template1 databases and the public schema in every database on the instance are temporarily renamed. These objects appear in the logs with their original name and a random string appended. The string is appended so that custom settings such as locale and owner are preserved during the major version upgrade. After the upgrade completes, the objects are renamed back to their original names.

Note

During the major version upgrade process, you can't do a point-in-time restore of your instance. After Amazon RDS performs the upgrade, it takes an automatic backup of the instance. You can perform a point-in-time restore to times before the upgrade began and after the automatic backup of your instance has completed.

8. **If an upgrade fails with precheck procedure errors, resolve the issues** – During the major version upgrade process, Amazon RDS for PostgreSQL first runs a precheck procedure to identify any issues that might cause the upgrade to fail. The precheck procedure checks all potential incompatible conditions across all databases in the instance.

If the precheck encounters an issue, it creates a log event indicating the upgrade precheck failed. The precheck process details are in an upgrade log named pg_upgrade_precheck.log for all the databases of a DB instance. Amazon RDS appends a timestamp to the file name. For more information about viewing logs, see [Amazon RDS Database Log Files \(p. 390\)](#).

Resolve all of the issues identified in the precheck log and then retry the major version upgrade. The following is an example of a precheck log.

```
-----  
Upgrade could not be run on Wed Apr 4 18:30:52 2018  
-----  
The instance could not be upgraded from 9.6.11 to 10.6 for the following reasons.  
Please take appropriate action on databases that have usage incompatible with the  
requested major engine version upgrade and try the upgrade again.  
  
* There are uncommitted prepared transactions. Please commit or rollback all prepared  
transactions.* One or more role names start with 'pg_'. Rename all role names that start  
with 'pg_'.
```

```
* The following issues in the database 'my"million$"db' need to be corrected before upgrading:  
** The ["line","reg*"] data types are used in user tables. Remove all usage of these data types.  
** The database name contains characters that are not supported by RDS PostgreSQL. Rename the database.  
** The database has extensions installed that are not supported on the target database version. Drop the following extensions from your database: ["tsearch2"].  
  
* The following issues in the database 'mydb' need to be corrected before upgrading:  
The database has views or materialized views that depend on 'pg_stat_activity'. Drop the views.
```

9. **Upgrade your production instance** – When the dry-run major version upgrade is successful, you should be able to upgrade your production database with confidence. For more information, see [Manually Upgrading the Engine Version \(p. 159\)](#).

After the major version upgrade is complete, we recommend the following:

- Run the `ANALYZE` operation to refresh the `pg_statistic` table.
- A PostgreSQL upgrade doesn't upgrade any PostgreSQL extensions. To upgrade an extension, use the `ALTER EXTENSION UPDATE` command.

Note

If you are running the PostGIS extension in your Amazon RDS PostgreSQL DB instance, make sure that you follow the [PostGIS upgrade instructions](#) in the PostGIS documentation before you update the extension.

To update an extension, run the following command.

```
ALTER EXTENSION extension_name UPDATE TO 'new_version'
```

For the list of supported versions of PostgreSQL extensions, see [PostgreSQL Extensions and Modules Supported on Amazon RDS \(p. 1383\)](#).

- Optionally, use Amazon RDS to view two logs that the `pg_upgrade` utility produces. These are `pg_upgrade_internal.log` and `pg_upgrade_server.log`. Amazon RDS appends a timestamp to the file name for these logs. You can view these logs as you can any other log. For more information, see [Amazon RDS Database Log Files \(p. 390\)](#).

You can also upload the upgrade logs to Amazon CloudWatch Logs. For more information, see [Publishing PostgreSQL Logs to CloudWatch Logs \(p. 420\)](#).

- To verify that everything works as expected, test your application on the upgraded database with a similar workload. After the upgrade is verified, you can delete this test instance.

Automatic Minor Version Upgrades for PostgreSQL

If you enable the **Auto minor version upgrade** option when creating or modifying a DB instance, you can have your DB instance automatically upgraded.

For each RDS for PostgreSQL major version, one minor version is designated by RDS as the automatic upgrade version. After a minor version has been tested and approved by Amazon RDS, the minor version upgrade occurs automatically during your maintenance window. RDS doesn't automatically set newer released minor versions as the automatic upgrade version. Before RDS designates a newer automatic upgrade version, several criteria are considered, such as the following:

- Known security issues
- Bugs in the PostgreSQL community version

- Overall fleet stability since the minor version was released

You can use the following AWS CLI command and script to determine the current automatic upgrade minor versions.

```
aws rds describe-db-engine-versions --engine postgres | grep -A 1 AutoUpgrade| grep -A 2 true |grep PostgreSQL | sort --unique | sed -e 's/"Description": "//g'
```

A PostgreSQL DB instance is automatically upgraded during your maintenance window if the following criteria are met:

- The DB instance has the **Auto minor version upgrade** option enabled.
- The DB instance is running a minor DB engine version that is less than the current automatic upgrade minor version.

For more information, see [Automatically Upgrading the Minor Engine Version \(p. 160\)](#).

A PostgreSQL upgrade doesn't upgrade any PostgreSQL extensions. To update an extension after a minor version upgrade, use the `ALTER EXTENSION UPDATE` command.

Note

If you are running the PostGIS extension in your Amazon RDS PostgreSQL DB instance, make sure that you follow the [PostGIS upgrade instructions](#) in the PostGIS documentation before you update the extension.

For example, to update an extension, run the following command.

```
ALTER EXTENSION extension_name UPDATE TO 'new_version'
```

For the list of supported versions of PostgreSQL extensions, see [PostgreSQL Extensions and Modules Supported on Amazon RDS \(p. 1383\)](#).

Upgrading a PostgreSQL DB Snapshot

With Amazon RDS, you can create a storage volume DB snapshot of your PostgreSQL DB instance. When you create a DB snapshot, the snapshot is based on the engine version used by your Amazon RDS instance. In addition to upgrading the DB engine version of your DB instance, you can also upgrade the engine version for your DB snapshots.

After restoring a DB snapshot upgraded to a new engine version, make sure to test that the upgrade was successful. For more information about a major version upgrade, see [Upgrading the PostgreSQL DB Engine for Amazon RDS \(p. 1289\)](#). To learn how to restore a DB snapshot, see [Restoring from a DB Snapshot \(p. 266\)](#).

You can upgrade manual DB snapshots that are either encrypted or not encrypted.

For the list of engine versions that are available for upgrading a DB snapshot, see [Upgrading the PostgreSQL DB Engine for Amazon RDS](#).

Note

- The DB snapshot must be from the same AWS Region as the account.
- You can't upgrade DB snapshots that are copied within region, copied across regions, or shared across accounts.
- You can't upgrade automated DB snapshots that are created during the automated backup process.

Console

To upgrade a DB snapshot

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the navigation pane, choose **Snapshots**.
3. Choose the snapshot that you want to upgrade.
4. For **Actions**, choose **Upgrade snapshot**. The **Upgrade snapshot** page appears.
5. Choose the **New engine version** to upgrade to.
6. Choose **Save changes** to upgrade the snapshot.

During the upgrade process, all snapshot actions are disabled for this DB snapshot. Also, the DB snapshot status changes from **available** to **upgrading**, and then changes to **active** upon completion. If the DB snapshot can't be upgraded because of snapshot corruption issues, the status changes to **unavailable**. You can't recover the snapshot from this state.

Note

If the DB snapshot upgrade fails, the snapshot is rolled back to the original state with the original version.

AWS CLI

To upgrade a DB snapshot to a new database engine version, use the AWS CLI [modify-db-snapshot](#) command.

Parameters

- **--db-snapshot-identifier** – The identifier of the DB snapshot to upgrade. The identifier must be a unique Amazon Resource Name (ARN). For more information, see [Working with Amazon Resource Names \(ARNs\) in Amazon RDS \(p. 224\)](#).
- **--engine-version** – The engine version to upgrade the DB snapshot to.

Example

For Linux, OS X, or Unix:

```
aws rds modify-db-snapshot \
--db-snapshot-identifier my_db_snapshot \
--engine-version new_version
```

For Windows:

```
aws rds modify-db-snapshot ^
--db-snapshot-identifier my_db_snapshot ^
--engine-version new_version
```

RDS API

To upgrade a DB snapshot to a new database engine version, call the Amazon RDS API [ModifyDBSnapshot](#) operation.

- **DBSnapshotIdentifier** – The identifier of the DB snapshot to upgrade. The identifier must be a unique Amazon Resource Name (ARN). For more information, see [Working with Amazon Resource Names \(ARNs\) in Amazon RDS \(p. 224\)](#).
- **EngineVersion** – The engine version to upgrade the DB snapshot to.

Example

```
https://rds.us-west-2.amazonaws.com/
?Action=ModifyDBSnapshot
&DBSnapshotIdentifier=mydbsnapshot
&EngineVersion=newversion
&SignatureVersion=4
&Version=2014-10-31
&X-Amz-Algorithm=AWS4-HMAC-SHA256
&X-Amz-Credential=AKIADQKE4SARGYLE/20161222/us-west-1/rds/aws4_request
&X-Amz-Date=20161222T233051Z
&X-Amz-SignedHeaders=content-type;host;user-agent;x-amz-content-sha256;x-amz-date
&X-Amz-Signature=8052a76dfb18469393c5f0182cdab0ebc224a9c7c5c949155376c1c250fc7ec3
```

Working with PostgreSQL Read Replicas

You usually use Read Replicas to configure replication between Amazon RDS DB instances. For general information about Read Replicas, see [Working with Read Replicas \(p. 182\)](#).

This section contains specific information about working with Read Replicas on PostgreSQL.

Topics

- [Read Replica Configuration with PostgreSQL \(p. 1298\)](#)
- [Monitoring PostgreSQL Read Replicas \(p. 1299\)](#)
- [Read Replica Limitations with PostgreSQL \(p. 1299\)](#)
- [Replication Interruptions with PostgreSQL Read Replicas \(p. 1299\)](#)
- [Troubleshooting a PostgreSQL Read Replica Problem \(p. 1299\)](#)

Read Replica Configuration with PostgreSQL

Amazon RDS PostgreSQL 9.3.5 and later uses PostgreSQL native streaming replication to create a read-only copy of a source (a "master" in PostgreSQL terms) DB instance. This Read Replica (a "standby" in PostgreSQL terms) DB instance is an asynchronously created physical replication of the master DB instance. It's created by a special connection that transmits write ahead log (WAL) data between the source DB instance and the Read Replica where PostgreSQL asynchronously streams database changes as they are made.

PostgreSQL uses a "replication" role to perform streaming replication. The role is privileged, but can't be used to modify any data. PostgreSQL uses a single process for handling replication.

Before a DB instance can serve as a source DB instance, you must enable automatic backups on the source DB instance by setting the backup retention period to a value other than 0.

Creating a PostgreSQL Read Replica doesn't require an outage for the master DB instance. Amazon RDS sets the necessary parameters and permissions for the source DB instance and the Read Replica without any service interruption. A snapshot is taken of the source DB instance, and this snapshot becomes the Read Replica. No outage occurs when you delete a Read Replica.

You can create up to five Read Replicas from one source DB instance. For replication to operate effectively, each Read Replica should have the same amount of compute and storage resources as the source DB instance. If you scale the source DB instance, you should also scale the Read Replicas.

Amazon RDS overrides any incompatible parameters on a Read Replica if it prevents the Read Replica from starting. For example, suppose that the `max_connections` parameter value is higher on the source DB instance than on the Read Replica. In that case, Amazon RDS updates the parameter on the Read Replica to be the same value as that on the source DB instance.

PostgreSQL DB instances use a secure connection that you can encrypt by setting the `ssl` parameter to 1 for both the source and the Read Replica instances.

You can create a Read Replica from either single-AZ or Multi-AZ DB instance deployments. You use Multi-AZ deployments to improve the durability and availability of critical data, but you can't use the Multi-AZ secondary to serve read-only queries. Instead, you can create Read Replicas from high-traffic Multi-AZ DB instances to offload read-only queries. If the source instance of a Multi-AZ deployment fails over to the secondary, any associated Read Replicas automatically switch to use the secondary (now primary) as their replication source. For more information, see [High Availability \(Multi-AZ\) for Amazon RDS \(p. 144\)](#).

You can create a Read Replica as a Multi-AZ DB instance. Amazon RDS creates a standby of your replica in another Availability Zone for failover support for the replica. Creating your Read Replica as a Multi-AZ DB instance is independent of whether the source database is a Multi-AZ DB instance.

If you use the [postgres_fdw](#) extension to access data from a remote server, the Read Replica will also have access to the remote server. For more information about using `postgres_fdw`, see [Accessing External Data with the postgres_fdw Extension \(p. 1347\)](#).

Monitoring PostgreSQL Read Replicas

For PostgreSQL Read Replicas, you can monitor replication lag in Amazon CloudWatch by viewing the Amazon RDS `ReplicaLag` metric. The `ReplicaLag` metric reports the value of `SELECT extract(epoch from now() - pg_last_xact_replay_timestamp()) AS slave_lag.`

Read Replica Limitations with PostgreSQL

The following are limitations for PostgreSQL Read Replicas:

- Each PostgreSQL Read Replica is read-only and can't be made a writable Read Replica.
- You can't create a Read Replica from another Read Replica (that is, you can't create cascading Read Replicas).
- You can promote a PostgreSQL Read Replica to be a new source DB instance. However, the Read Replica doesn't become the new source DB instance automatically. The Read Replica, when promoted, stops receiving WAL communications and is no longer a read-only instance. You must set up any replication you intend to have going forward because the promoted Read Replica is now a new source DB instance.
- A PostgreSQL Read Replica reports a replication lag of up to five minutes if there are no user transactions occurring on the source DB instance.

Replication Interruptions with PostgreSQL Read Replicas

In several situations, a PostgreSQL source DB instance can unintentionally break replication with a Read Replica. These situations include the following:

- The `max_wal_senders` parameter is set too low to provide enough data to the number of Read Replicas. This situation causes replication to stop.
- The PostgreSQL parameter `wal_keep_segments` dictates how many WAL files are kept to provide data to the Read Replicas. The parameter value specifies the number of logs to keep. If you set the parameter value too low, you can cause a Read Replica to fall so far behind that streaming replication stops. In this case, Amazon RDS reports a replication error and begins recovery on the Read Replica by replaying the source DB instance's archived WAL logs. This recovery process continues until the Read Replica has caught up enough to continue streaming replication. For more information, see [Troubleshooting a PostgreSQL Read Replica Problem \(p. 1299\)](#).
- A PostgreSQL Read Replica requires a reboot if the source DB instance endpoint changes.

When the WAL stream that provides data to a Read Replica is broken, PostgreSQL switches into recovery mode to restore the Read Replica by using archived WAL files. When this process is complete, PostgreSQL attempts to re-establish streaming replication.

Troubleshooting a PostgreSQL Read Replica Problem

PostgreSQL uses replication slots for cross-region replication, so the process for troubleshooting same-region replication problems and cross-region replication problems is different.

Troubleshooting PostgreSQL Read Replica Problems Within an AWS Region

The PostgreSQL parameter, `wal_keep_segments`, dictates how many Write Ahead Log (WAL) files are kept to provide data to the Read Replicas. The parameter value specifies the number of logs to keep. If you set the parameter value too low, you can cause a Read Replica to fall so far behind that streaming replication stops. In this case, Amazon RDS reports a replication error and begins recovery on the Read Replica by replaying the source DB instance's archived WAL logs. This recovery process continues until the Read Replica has caught up enough to continue streaming replication.

The PostgreSQL log on the Read Replica shows when Amazon RDS is recovering a Read Replica that is this state by replaying archived WAL files.

```
2014-11-07 19:01:10 UTC:@:[23180]:DEBUG: switched WAL source from archive to stream
after
failure 2014-11-07 19:01:10 UTC:@:[11575]:LOG: started streaming WAL from primary
at
1A/D3000000 on timeline 1 2014-11-07 19:01:10 UTC:@:[11575]:FATAL: could not
receive
data from WAL stream: ERROR: requested WAL segment 00000001000001A00000D3 has
already been
removed 2014-11-07 19:01:10 UTC:@:[23180]:DEBUG: could not restore file
"00000002.history" from archive: return code 0 2014-11-07 19:01:15
UTC:@:[23180]:DEBUG: switched WAL source from stream to archive after failure
recovering 00000001000001A00000D3 2014-11-07 19:01:16 UTC:@:[23180]:LOG: restored
log file "00000001000001A00000D3"
from archive
```

After a certain amount of time, Amazon RDS replays enough archived WAL files on the replica to catch up and allow the Read Replica to begin streaming again. At this point, PostgreSQL resumes streaming and writes a similar line to the following to the log file.

```
2014-11-07 19:41:36 UTC:@:[24714]:LOG: started streaming WAL from primary at 1B/
B6000000
on timeline 1
```

You can determine how many WAL files you should keep by looking at the checkpoint information in the log. The PostgreSQL log shows the following information at each checkpoint. By looking at the "# recycled" transaction log files of these log statements, you can understand how many transaction files will be recycled during a time range and use this information to tune the `wal_keep_segments` parameter.

```
2014-11-07 19:59:35 UTC:@:[26820]:LOG: checkpoint complete: wrote 376 buffers (0.2%); 0
transaction log file(s) added, 0 removed, 1 recycled; write=35.681 s, sync=0.013 s,
total=35.703 s; sync files=10, longest=0.013 s, average=0.001 s
```

For example, suppose that the PostgreSQL log shows that 35 files are recycled from the "checkpoint completed" log statements within a 5-minute time frame. In that case, we know that with this usage pattern a Read Replica relies on 35 transaction files in five minutes. A Read Replica can't survive five minutes in a nonstreaming state if the source DB instance is set to the default `wal_keep_segments` parameter value of 32.

Troubleshooting PostgreSQL Read Replica Problems Across AWS Regions

PostgreSQL (versions 9.4.7 and 9.5.2 exclusively) uses physical replication slots to manage Write Ahead Log (WAL) retention on the source DB instance. For each cross-region Read Replica instance, Amazon

RDS creates and associates a physical replication slot. You can use two Amazon CloudWatch metrics, Oldest Replication Slot Lag and Transaction Logs Disk Usage, to see how far behind the most lagging replica is in terms of WAL data received and to see how much storage is being used for WAL data. The Transaction Logs Disk Usage value can substantially increase when a cross-region Read Replica is lagging significantly.

If the workload on your DB instance generates a large amount of WAL data, you might need to change the DB instance class of your source DB instance and Read Replica. In that case, you change it to one with high (10 Gbps) network performance for the replica to keep up. The Amazon CloudWatch metric `Transaction Logs Generation` can help you understand the rate at which your workload is generating WAL data.

To determine the status of a cross-region Read Replica, you can query `pg_replication_slots` on the source instance, as in the following example:

```

postgres=# select * from pg_replication_slots;
          slot_name           | plugin | slot_type | datoid | database |
active | active_pid | xmin | catalog_xmin | restart_lsn
-----+-----+-----+-----+-----+
rds_us_east_1_db_uzwlh0lddgpbblksc6hgw4nkte |       | physical |       |       | t
| 12598 |       |           | 4E/95000060
(1 row)

```

Importing Data into PostgreSQL on Amazon RDS

Suppose that you have an existing PostgreSQL deployment that you want to move to Amazon RDS. The complexity of your task depends on the size of your database and the types of database objects that you're transferring. For example, consider a database that contains datasets on the order of gigabytes, along with stored procedures and triggers. Such a database is going to be more complicated than a simple database with only a few megabytes of test data and no triggers or stored procedures.

We recommend that you use native PostgreSQL database migration tools under the following conditions:

- You have a homogeneous migration, where you are migrating from a database with the same database engine as the target database.
- You are migrating an entire database.
- The native tools allow you to migrate your system with minimal downtime.

In most other cases, performing a database migration using AWS Database Migration Service (AWS DMS) is the best approach. AWS DMS can migrate databases without downtime and, for many database engines, continue ongoing replication until you are ready to switch over to the target database. You can migrate to either the same database engine or a different database engine using AWS DMS. If you are migrating to a different database engine than your source database, you can use the AWS Schema Conversion Tool (AWS SCT). You use AWS SCT to migrate schema objects that are not migrated by AWS DMS. For more information about AWS DMS, see [What Is AWS Database Migration Service?](#)

Modify your DB parameter group to include the following settings *for your import only*. You should test the parameter settings to find the most efficient settings for your DB instance size. You also need to revert back to production values for these parameters after your import completes.

Modify your DB instance settings to the following:

- Disable DB instance backups (set backup_retention to 0).
- Disable Multi-AZ.

Modify your DB parameter group to include the following settings. You should only use these settings when importing data. You should test the parameter settings to find the most efficient settings for your DB instance size. You also need to revert back to production values for these parameters after your import completes.

Parameter	Recommended Value When Importing	Description
<code>maintenance_work_mem</code>	524288, 1048576, 2097152 or 4194304 (in KB). These settings are comparable to 512 MB, 1 GB, 2 GB, and 4 GB.	The value for this setting depends on the size of your host. This parameter is used during CREATE INDEX statements and each parallel command can use this much memory. Calculate the best value so that you don't set this value so high that you run out of memory.
<code>checkpoint_segments</code>	256	The value for this setting consumes more disk space, but gives you less contention on your WAL logs. For PostgreSQL versions 9.5.x and 9.6.x, this value would be <code>max_wal_size</code> .
<code>checkpoint_timeout</code>	1800	The value for this setting allows for less frequent WAL rotation.

Parameter	Recommended Value When Importing	Description
<code>synchronous_commit</code>	Off	Disable this setting to speed up writes. Turning this parameter off can increase the risk of data loss in the event of a server crash (do not turn off FSYNC)
<code>wal_buffers</code>	8192	This value is in 8 KB units. This again helps your WAL generation speed
<code>autovacuum</code>	Off	Disable the PostgreSQL auto vacuum parameter while you are loading data so that it doesn't use resources

Use the `pg_dump -Fc` (compressed) or `pg_restore -j` (parallel) commands with these settings.

Note

The PostgreSQL command `pg_dumpall` requires `super_user` permissions that are not granted when you create a DB instance, so it cannot be used for importing data.

Topics

- [Importing a PostgreSQL Database from an Amazon EC2 Instance \(p. 1303\)](#)
- [Using the \copy Command to Import Data to a Table on a PostgreSQL DB Instance \(p. 1305\)](#)
- [Importing Amazon S3 Data into an RDS for PostgreSQL DB Instance \(p. 1305\)](#)
- [Transporting PostgreSQL Databases Between DB Instances \(p. 1316\)](#)

Importing a PostgreSQL Database from an Amazon EC2 Instance

If you have data in a PostgreSQL server on an Amazon EC2 instance and want to move it to a PostgreSQL DB instance, you can use the following process. The following list shows the steps to take. Each step is discussed in more detail in the following sections.

1. Create a file using `pg_dump` that contains the data to be loaded
2. Create the target DB instance
3. Use `psql` to create the database on the DB instance and load the data
4. Create a DB snapshot of the DB instance

Step 1: Create a File Using `pg_dump` That Contains the Data to Load

The `pg_dump` utility uses the `COPY` command to create a schema and data dump of a PostgreSQL database. The dump script generated by `pg_dump` loads data into a database with the same name and recreates the tables, indexes, and foreign keys. You can use the `pg_restore` command and the `-d` parameter to restore the data to a database with a different name.

Before you create the data dump, you should query the tables to be dumped to get a row count so you can confirm the count on the target DB instance.

The following command creates a dump file called `mydb2dump.sql` for a database called `mydb2`.

```
prompt>pg_dump dbname=mydb2 -f mydb2dump.sql
```

Step 2: Create the Target DB Instance

Create the target PostgreSQL DB instance using either the Amazon RDS console, AWS CLI, or API. Create the instance with the backup retention setting set to 0 and disable Multi-AZ. Doing so allows faster data import. You must create a database on the instance before you can dump the data. The database can have the same name as the database that is contained the dumped data. Alternatively, you can create a database with a different name. In this case, you use the `pg_restore` command and the `-d` parameter to restore the data into the newly named database.

For example, the following commands can be used to dump, restore, and rename a database.

```
pg_dump -Fc -v -h [endpoint of instance] -U [master username] [database] > [database].dump
createdb [new database name]
pg_restore -v -h [endpoint of instance] -U [master username] -d [new database
name] [database].dump
```

Step 3: Use psql to Create the Database on the DB Instance and Load Data

You can use the same connection you used to execute the `pg_dump` command to connect to the target DB instance and recreate the database. Using `psql`, you can use the master user name and master password to create the database on the DB instance

The following example uses `psql` and a dump file named `mydb2dump.sql` to create a database called `mydb2` on a PostgreSQL DB instance called `mypginstance`:

For Linux, OS X, or Unix:

```
psql \
-f mydb2dump.sql \
--host mypginstance.c6c8mntzhgv0.us-west-2.rds.amazonaws.com \
--port 8199 \
--username myawsuser \
--password password \
--dbname mydb2
```

For Windows:

```
psql ^
-f mydb2dump.sql ^
--host mypginstance.c6c8mntzhgv0.us-west-2.rds.amazonaws.com ^
--port 8199 ^
--username myawsuser ^
--password password ^
--dbname mydb2
```

Step 4: Create a DB Snapshot of the DB Instance

Once you have verified that the data was loaded into your DB instance, we recommend that you create a DB snapshot of the target PostgreSQL DB instance. DB snapshots are complete backups of your DB instance that can be used to restore your DB instance to a known state. A DB snapshot taken immediately after the load protects you from having to load the data again in case of a mishap. You can also use such a snapshot to seed new DB instances. For information about creating a DB snapshot, see [Creating a DB Snapshot \(p. 264\)](#).

Using the \copy Command to Import Data to a Table on a PostgreSQL DB Instance

You can run the `\copy` command from the `psql` prompt to import data into a table on a PostgreSQL DB instance. The table must already exist on the DB instance. For more information on the `\copy` command, see the [PostgreSQL documentation](#).

Note

The `\copy` command doesn't provide confirmation of actions, such as a count of rows inserted. PostgreSQL does provide error messages if the copy command fails due to an error.

Create a .csv file from the data in the source table, log on to the target database on the PostgreSQL instance using `psql`, and then run the following command. This example uses `source-table` as the source table name, `source-table.csv` as the .csv file, and `target-db` as the target database:

```
target-db=> \copy source-table from 'source-table.csv' with DELIMITER ',';
```

You can also run the following command from your client computer command prompt. This example uses `source-table` as the source table name, `source-table.csv` as the .csv file, and `target-db` as the target database:

For Linux, OS X, or Unix:

```
$psql target-db \
-U <admin user> \
-p <port> \
-h <DB instance name> \
-c "\copy source-table from 'source-table.csv' with DELIMITER ','"
```

For Windows:

```
$psql target-db ^
-U <admin user> ^
-p <port> ^
-h <DB instance name> ^
-c "\copy source-table from 'source-table.csv' with DELIMITER ','"
```

Importing Amazon S3 Data into an RDS for PostgreSQL DB Instance

You can import data from Amazon S3 into a table belonging to an RDS for PostgreSQL DB instance. To do this, you use the `aws_s3` PostgreSQL extension that Amazon RDS provides.

Note

To import from Amazon S3 into RDS for PostgreSQL, your database must be running PostgreSQL version 10.7 or later.

For more information on storing data with Amazon S3, see [Create a Bucket](#) in the *Amazon Simple Storage Service Getting Started Guide*. For instructions on how to upload a file to an Amazon S3 bucket, see [Add an Object to a Bucket](#) in the *Amazon Simple Storage Service Getting Started Guide*.

Topics

- [Overview of Importing Amazon S3 Data \(p. 1306\)](#)
- [Setting Up Access to an Amazon S3 Bucket \(p. 1307\)](#)
- [Using the aws_s3.table_import_from_s3 Function to Import Amazon S3 Data \(p. 1311\)](#)

- [Function Reference \(p. 1313\)](#)

Overview of Importing Amazon S3 Data

To import data stored in an Amazon S3 bucket to a PostgreSQL database table, follow these steps.

To import S3 data into Amazon RDS

1. Install the required PostgreSQL extensions. These include the `aws_s3` and `aws_commons` extensions. To do so, start `psql` and use the following command.

```
psql=> CREATE EXTENSION aws_s3 CASCADE;
NOTICE: installing required extension "aws_commons"
```

The `aws_s3` extension provides the [aws_s3.table_import_from_s3 \(p. 1313\)](#) function that you use to import Amazon S3 data. The `aws_commons` extension provides additional helper functions.

2. Identify the database table and Amazon S3 file to use.

The [aws_s3.table_import_from_s3 \(p. 1313\)](#) function requires the name of the PostgreSQL database table that you want to import data into. The function also requires that you identify the Amazon S3 file to import. To provide this information, take the following steps.

- a. Identify the PostgreSQL database table to put the data in. For example, the following is a sample `t1` database table used in the examples for this topic.

```
psql=> CREATE TABLE t1 (col1 varchar(80), col2 varchar(80), col3 varchar(80));
```

- b. Get the following information to identify the Amazon S3 file that you want to import:

- Bucket name – A *bucket* is a container for Amazon S3 objects or files.
- File path – The file path locates the file in the Amazon S3 bucket.
- AWS Region – The AWS Region is the location of the Amazon S3 bucket. For example, if the S3 bucket is in the US East (N. Virginia) Region, use `us-east-1`. For a listing of AWS Region names and associated values, see [Regions and Availability Zones \(p. 132\)](#).

To find how to get this information, see [View an Object](#) in the *Amazon Simple Storage Service Getting Started Guide*. You can confirm the information by using the AWS CLI command `aws s3 cp`. If the information is correct, this command downloads a copy of the Amazon S3 file.

```
aws s3 cp s3://sample_s3_bucket/sample_file_path ./
```

- c. Use the [aws_commons.create_s3_uri \(p. 1315\)](#) function to create an `aws_commons._s3_uri_1` structure to hold the Amazon S3 file information. You provide this `aws_commons._s3_uri_1` structure as a parameter in the call to the [aws_s3.table_import_from_s3 \(p. 1313\)](#) function.

For a `psql` example, see the following.

```
psql=> SELECT aws_commons.create_s3_uri(
    'sample_s3_bucket',
    'sample.csv',
    'us-east-1'
) AS s3_uri \gset
```

3. Provide permission to access the Amazon S3 file.

To import data from an Amazon S3 file, you need to give the RDS for PostgreSQL DB instance permission to access the Amazon S3 bucket the file is in. To do this, you use either an AWS Identity and Access Management (IAM) role or security credentials. For more information, see [Setting Up Access to an Amazon S3 Bucket \(p. 1307\)](#).

4. Import the Amazon S3 data by calling the `aws_s3.table_import_from_s3` function.

After you complete the previous preparation tasks, use the [aws_s3.table_import_from_s3 \(p. 1313\)](#) function to import the Amazon S3 data. For more information, see [Using the aws_s3.table_import_from_s3 Function to Import Amazon S3 Data \(p. 1311\)](#).

Setting Up Access to an Amazon S3 Bucket

To import data from an Amazon S3 file, you need to give the RDS for PostgreSQL DB instance permission to access the Amazon S3 bucket the file is in. You provide access to an Amazon S3 bucket in one of two ways, as described in the following topics.

Topics

- [Using an IAM Role to Access an Amazon S3 Bucket \(p. 1307\)](#)
- [Using Security Credentials to Access an Amazon S3 Bucket \(p. 1310\)](#)
- [Troubleshooting Access to Amazon S3 \(p. 1311\)](#)

Using an IAM Role to Access an Amazon S3 Bucket

Before you load data from an Amazon S3 file, give your RDS for PostgreSQL DB instance permission to access the Amazon S3 bucket the file is in. This way, you don't have to manage additional credential information or provide it in the [aws_s3.table_import_from_s3 \(p. 1313\)](#) function call.

To do this, create an IAM policy that provides access to the Amazon S3 bucket. Create an IAM role and attach the policy to the role. Then assign the IAM role to your DB instance.

To give an RDS for PostgreSQL DB instance access to Amazon S3 through an IAM role

1. Create an IAM policy. This policy provides the bucket and object permissions that allow your RDS for PostgreSQL DB instance to access Amazon S3.

Include in the policy the following required actions to allow the transfer of files from an Amazon S3 bucket to Amazon RDS:

- `s3:GetObject`
- `s3>ListBucket`

Include in the policy the following resources to identify the Amazon S3 bucket and objects in the bucket. This shows the Amazon Resource Name (ARN) format for accessing Amazon S3.

- `arn:aws:s3:::your-s3-bucket`
- `arn:aws:s3:::your-s3-bucket/*`

For more information on creating an IAM policy for Amazon RDS for PostgreSQL, see [Creating and Using an IAM Policy for IAM Database Access \(p. 464\)](#). See also [Tutorial: Create and Attach Your First Customer Managed Policy](#) in the *IAM User Guide*.

The following AWS CLI command creates an IAM policy named `rds-s3-import-policy` with these options. It grants access to a bucket named `your-s3-bucket`.

Note

After you create the policy, note the Amazon Resource Name (ARN) of the policy. You need the ARN for a subsequent step when you attach the policy to an IAM role.

Example

For Linux, OS X, or Unix:

```
aws iam create-policy \
--policy-name rds-s3-import-policy \
--policy-document '{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "s3import",
            "Action": [
                "s3:GetObject",
                "s3>ListBucket"
            ],
            "Effect": "Allow",
            "Resource": [
                "arn:aws:s3:::your-s3-bucket",
                "arn:aws:s3:::your-s3-bucket/*"
            ]
        }
    ]
}'
```

For Windows:

```
aws iam create-policy ^
--policy-name rds-s3-import-policy ^
--policy-document '{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "s3import",
            "Action": [
                "s3:GetObject",
                "s3>ListBucket"
            ],
            "Effect": "Allow",
            "Resource": [
                "arn:aws:s3:::your-s3-bucket",
                "arn:aws:s3:::your-s3-bucket/*"
            ]
        }
    ]
}'
```

2. Create an IAM role. You do this so Amazon RDS can assume this IAM role on your behalf to access your Amazon S3 buckets. For more information, see [Creating a Role to Delegate Permissions to an IAM User](#) in the *IAM User Guide*.

The following example shows using the AWS CLI command to create a role named `rds-s3-import-role`.

Example

For Linux, OS X, or Unix:

```
aws iam create-role \
--role-name rds-s3-import-role \
--assume-role-policy-document '{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Principal": {
                "Service": "rds.amazonaws.com"
            },
            "Action": "sts:AssumeRole"
        }
    ]
}'
```

For Windows:

```
aws iam create-role ^
--role-name rds-s3-import-role ^
--assume-role-policy-document '{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Principal": {
                "Service": "rds.amazonaws.com"
            },
            "Action": "sts:AssumeRole"
        }
    ]
}'
```

3. Attach the IAM policy that you created to the IAM role that you created.

The following AWS CLI command attaches the policy created earlier to the role named `rds-s3-import-role`. Replace `your-policy-arn` with the policy ARN that you noted in an earlier step.

Example

For Linux, OS X, or Unix:

```
aws iam attach-role-policy \
--policy-arn your-policy-arn \
--role-name rds-s3-import-role
```

For Windows:

```
aws iam attach-role-policy ^
--policy-arn your-policy-arn ^
--role-name rds-s3-import-role
```

4. Add the IAM role to the DB instance. You do so by using the AWS Management Console or AWS CLI, as described following.

Console

To add an IAM role for a PostgreSQL DB instance using the console

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. Choose the PostgreSQL DB instance name to display its details.
3. On the **Connectivity & security** tab, in the **Manage IAM roles** section, choose the role to add under **Add IAM roles to this instance**.
4. Under **Feature**, choose **s3Import**.
5. Choose **Add role**.

AWS CLI

To add an IAM role for a PostgreSQL DB instance using the CLI

- Use the following command to add the role to the PostgreSQL DB instance named `my-db-instance`. Replace `your-role-arn` with the role ARN that you noted in a previous step. Use `s3Import` for the value of the `--feature-name` option.

Example

For Linux, OS X, or Unix:

```
aws rds add-role-to-db-instance \
--db-instance-identifier my-db-instance \
--feature-name s3Import \
--role-arn your-role-arn \
--region your-region
```

For Windows:

```
aws rds add-role-to-db-instance ^
--db-instance-identifier my-db-instance ^
--feature-name s3Import ^
--role-arn your-role-arn ^
--region your-region
```

Using Security Credentials to Access an Amazon S3 Bucket

If you prefer, you can use security credentials to provide access to an Amazon S3 bucket instead of providing access with an IAM role. To do this, use the `credentials` parameter in the [aws_s3.table_import_from_s3 \(p. 1313\)](#) function call.

The `credentials` parameter is a structure of type `aws_commons._aws_credentials_1`, which contains AWS credentials. Use the [aws_commons.create_aws_credentials \(p. 1316\)](#) function to set the access key and secret key in an `aws_commons._aws_credentials_1` structure, as shown following.

```
psql=> SELECT aws_commons.create_aws_credentials(
  'sample_access_key', 'sample_secret_key', '')
AS creds \gset
```

After creating the `aws_commons._aws_credentials_1` structure, use the [aws_s3.table_import_from_s3 \(p. 1313\)](#) function with the `credentials` parameter to import the data, as shown following.

```
psql=> SELECT aws_s3.table_import_from_s3(
  't', '',
  '(format csv)',
  :'s3_uri',
  :'creds'
);
```

Or you can include the [aws_commons.create_aws_credentials \(p. 1316\)](#) function call inline within the `aws_s3.table_import_from_s3` function call.

```
psql=> SELECT aws_s3.table_import_from_s3(
  't', '',
  '(format csv)',
  :'s3_uri',
  aws_commons.create_aws_credentials('sample_access_key', 'sample_secret_key', '')
);
```

Troubleshooting Access to Amazon S3

If you encounter connection problems when attempting to import Amazon S3 file data, see the following for recommendations:

- [Troubleshooting Amazon RDS Identity and Access \(p. 478\)](#)
- [Troubleshooting Amazon S3](#)
- [Troubleshooting Amazon S3 and IAM](#)

Using the `aws_s3.table_import_from_s3` Function to Import Amazon S3 Data

Import your Amazon S3 data by calling the [aws_s3.table_import_from_s3 \(p. 1313\)](#) function.

Note

The following examples use the IAM role method for providing access to the Amazon S3 bucket. Thus, there are no credential parameters in the `aws_s3.table_import_from_s3` function calls.

The following shows a typical PostgreSQL example using `psql`.

```
psql=> SELECT aws_s3.table_import_from_s3(
  't1',
  '',
  '(format csv)',
  :'s3_uri'
);
```

The parameters are the following:

- `t1` – The name for the table in the PostgreSQL DB instance to copy the data into.
- `''` – An optional list of columns in the database table. You can use this parameter to indicate which columns of the S3 data go in which table columns. If no columns are specified, all the columns are copied to the table. For an example of using a column list, see [Importing an Amazon S3 File That Uses a Custom Delimiter \(p. 1312\)](#).
- `(format csv)` – PostgreSQL COPY arguments. The copy process uses the arguments and format of the [PostgreSQL COPY](#) command. In the preceding example, the COPY command uses the comma-separated value (CSV) file format to copy the data.

- **s3_uri** – A structure that contains the information identifying the Amazon S3 file. For an example of using the [aws_commons.create_s3_uri \(p. 1315\)](#) function to create an **s3_uri** structure, see [Overview of Importing Amazon S3 Data \(p. 1306\)](#).

For the full reference of this function, see [aws_s3.table_import_from_s3 \(p. 1313\)](#).

The following examples show how to specify different kinds of files when importing Amazon S3 data.

Topics

- [Importing an Amazon S3 File That Uses a Custom Delimiter \(p. 1312\)](#)
- [Importing an Amazon S3 Compressed \(gzip\) File \(p. 1313\)](#)
- [Importing an Encoded Amazon S3 File \(p. 1313\)](#)

Importing an Amazon S3 File That Uses a Custom Delimiter

The following example shows how to import a file that uses a custom delimiter. It also shows how to control where to put the data in the database table using the `column_list` parameter of the [aws_s3.table_import_from_s3 \(p. 1313\)](#) function.

For this example, assume that the following information is organized into pipe-delimited columns in the Amazon S3 file.

```
1|foo1|bar1|elephant1
2|foo2|bar2|elephant2
3|foo3|bar3|elephant3
4|foo4|bar4|elephant4
...
```

To import a file that uses a custom delimiter

1. Create a table in the database for the imported data.

```
psql=> CREATE TABLE test (a text, b text, c text, d text, e text);
CREATE TABLE
```

2. Use the following form of the [aws_s3.table_import_from_s3 \(p. 1313\)](#) function to import data from the Amazon S3 file.

You can include the [aws_commons.create_s3_uri \(p. 1315\)](#) function call inline within the `aws_s3.table_import_from_s3` function call to specify the file.

```
psql=> SELECT aws_s3.table_import_from_s3(
    'test',
    'a,b,d,e',
    'DELIMITER ''|''',
    aws_commons.create_s3_uri('sampleBucket', 'pipeDelimitedSampleFile', 'us-east-2')
);
```

The data is now in the table in the following columns.

```
psql=> SELECT * FROM test;
a | b | c | d | e
---+---+---+---+---+
1 | foo1 | bar1 | elephant1
2 | foo2 | bar2 | elephant2
```

```
3 | foo3 | | bar3 | elephant3
4 | foo4 | | bar4 | elephant4
```

Importing an Amazon S3 Compressed (gzip) File

The following example shows how to import a file from Amazon S3 that is compressed with gzip.

Ensure that the file contains the following Amazon S3 metadata:

- Key: Content-Encoding
- Value: gzip

For more about adding these values to Amazon S3 metadata, see [How Do I Add Metadata to an S3 Object?](#) in the *Amazon Simple Storage Service Console User Guide*.

Import the gzip file into your RDS for PostgreSQL DB instance as shown following.

```
psql=> CREATE TABLE test_gzip(id int, a text, b text, c text, d text);
CREATE TABLE
psql=> SELECT aws_s3.table_import_from_s3(
  'test_gzip', '', '(format csv)',
  'myS3Bucket', 'test-data.gz', 'us-east-2'
);
```

Importing an Encoded Amazon S3 File

The following example shows how to import a file from Amazon S3 that has Windows-1252 encoding.

```
psql=> SELECT aws_s3.table_import_from_s3(
  'test_table', '', 'encoding ''WIN1252'''',
  aws_commons.create_s3_uri('sampleBucket', 'SampleFile', 'us-east-2')
);
```

Function Reference

Functions

- [aws_s3.table_import_from_s3 \(p. 1313\)](#)
- [aws_commons.create_s3_uri \(p. 1315\)](#)
- [aws_commons.create_aws_credentials \(p. 1316\)](#)

[aws_s3.table_import_from_s3](#)

Imports Amazon S3 data into an Amazon RDS table. The `aws_s3` extension provides the `aws_s3.table_import_from_s3` function.

The three required parameters are `table_name`, `column_list` and `options`. These identify the database table and specify how the data is copied into the table.

You can also use these parameters:

- The `s3_info` parameter specifies the Amazon S3 file to import. When you use this parameter, access to Amazon S3 is provided by an IAM role for the PostgreSQL DB instance.

```
aws_s3.table_import_from_s3 (
```

```

    table_name text,
    column_list text,
    options text,
    s3_info aws_commons._s3_uri_1
)

```

- The `credentials` parameter specifies the credentials to access Amazon S3. When you use this parameter, you don't use an IAM role.

```

aws_s3.table_import_from_s3 (
    table_name text,
    column_list text,
    options text,
    s3_info aws_commons._s3_uri_1,
    credentials aws_commons._aws_credentials_1
)

```

The `aws_s3.table_import_from_s3` parameters are described in the following table.

Parameter	Description
<code>table_name</code>	A required text string containing the name of the PostgreSQL database table to import the data into.
<code>column_list</code>	A required text string containing an optional list of the PostgreSQL database table columns in which to copy the data. If the string is empty, all columns of the table are used. For an example, see Importing an Amazon S3 File That Uses a Custom Delimiter (p. 1312) .
<code>options</code>	A required text string containing arguments for the PostgreSQL <code>COPY</code> command. These arguments specify how the data is to be copied into the PostgreSQL table. For more details, see the PostgreSQL COPY documentation .
<code>s3_info</code>	An <code>aws_commons._s3_uri_1</code> composite type containing the following information about the S3 object: <ul style="list-style-type: none"> <code>bucket</code> – The Amazon S3 bucket name containing the file. <code>file_path</code> – The Amazon S3 path of the file. <code>region</code> – The AWS Region that the file is in. For a listing of AWS Region names and associated region values, see Regions and Availability Zones (p. 132). To create an <code>aws_commons._s3_uri_1</code> composite structure, see aws_commons.create_s3_uri (p. 1315) .
<code>credentials</code>	An <code>aws_commons._aws_credentials_1</code> composite type containing the following credentials to use for the import operation: <ul style="list-style-type: none"> Access key Secret key Session token To create an <code>aws_commons._aws_credentials_1</code> composite structure, see aws_commons.create_aws_credentials (p. 1316) .

Alternate Parameters

To help with testing, you can use an expanded set of parameters instead of the `s3_info` and `credentials` parameters. Following are additional syntax variations for the `aws_s3.table_import_from_s3` function.

- Instead of using the `s3_info` parameter to identify an Amazon S3 file, use the combination of the `bucket`, `file_path`, and `region` parameters. With this form of the function, access to Amazon S3 is provided by an IAM role on the PostgreSQL DB instance.

```
aws_s3.table_import_from_s3 (
    table_name text,
    column_list text,
    options text,
    bucket text,
    file_path text,
    region text
)
```

- Instead of using the `credentials` parameter to specify Amazon S3 access, use the combination of the `access_key`, `session_key`, and `session_token` parameters.

```
aws_s3.table_import_from_s3 (
    table_name text,
    column_list text,
    options text,
    bucket text,
    file_path text,
    region text,
    access_key text,
    secret_key text,
    session_token text
)
```

Find descriptions for these alternate parameters in the following table.

Parameter	Description
<code>bucket</code>	A text string containing the name of the Amazon S3 bucket that contains the file.
<code>file_path</code>	A text string containing the Amazon S3 path of the file.
<code>region</code>	A text string containing the AWS Region that the file is in. For a listing of AWS Region names and associated values, see Regions and Availability Zones (p. 132) .
<code>access_key</code>	A text string containing the access key to use for the import operation. The default is NULL.
<code>secret_key</code>	A text string containing the secret key to use for the import operation. The default is NULL.
<code>session_token</code>	(Optional) A text string containing the session key to use for the import operation. The default is NULL.

[aws_commons.create_s3_uri](#)

Creates an `aws_commons._s3_uri_1` structure to hold Amazon S3 file information. You use the results of the `aws_commons.create_s3_uri` function in the `s3_info` parameter of the `aws_s3.table_import_from_s3` ([p. 1313](#)) function. The function syntax is as follows.

```
aws_commons.create_s3_uri(
    bucket text,
    file_path text,
    region text
)
```

The `aws_commons.create_s3_uri` function parameters are described in the following table.

Parameter	Description
<code>bucket</code>	A required text string containing the Amazon S3 bucket name for the file.
<code>file_path</code>	A required text string containing the Amazon S3 path of the file.
<code>region</code>	A required text string containing the AWS Region the file is in. For a listing of AWS Region names and associated region values, see Regions and Availability Zones (p. 132) .

[aws_commons.create_aws_credentials](#)

Sets an access key and secret key in an `aws_commons._aws_credentials_1` structure. Use the results of the `aws_commons.create_aws_credentials` function in the `credentials` parameter of the [aws_s3.table_import_from_s3 \(p. 1313\)](#) function. The function syntax is as follows.

```
aws_commons.create_aws_credentials(
    access_key text,
    secret_key text,
    session_token text
)
```

The `aws_commons.create_aws_credentials` function parameters are described in the following table.

Parameter	Description
<code>access_key</code>	A required text string containing the access key to use for importing an Amazon S3 file. The default is NULL.
<code>secret_key</code>	A required text string containing the secret key to use for importing an Amazon S3 file. The default is NULL.
<code>session_token</code>	An optional text string containing the session token to use for importing an Amazon S3 file. The default is NULL. Note, if you provide an optional <code>session_token</code> , you can use temporary credentials.

Transporting PostgreSQL Databases Between DB Instances

By using PostgreSQL Transportable Databases for Amazon RDS, you can transport a PostgreSQL database between two DB instances. This provides an extremely fast method of migrating large databases between separate DB instances. To transport databases using this method, your DB instances must both run the same major version of PostgreSQL.

To use transportable databases, install the `pg_transport` extension. This extension provides a physical transport mechanism to move each database. By streaming the database files with minimal processing, physical transport moves data much faster than traditional dump and load processes and takes minimal downtime. PostgreSQL transportable databases use a pull model where the destination DB instance imports the database from the source DB instance.

Note

PostgreSQL transportable databases are available in RDS for PostgreSQL versions 10.10 and later, and 11.5 and later.

Topics

- [Limitations for Using PostgreSQL Transportable Databases \(p. 1317\)](#)
- [Setting Up to Transport PostgreSQL Databases \(p. 1317\)](#)
- [Transporting a PostgreSQL Database Using the `transport.import_from_server` Function \(p. 1318\)](#)
- [What Happens During Database Transport \(p. 1318\)](#)
- [`transport.import_from_server` Function Reference \(p. 1319\)](#)
- [Configuration Parameters for the `pg_transport` Extension \(p. 1320\)](#)

Limitations for Using PostgreSQL Transportable Databases

Transportable databases have the following limitations:

- **Read Replicas** – You can't transportable databases on read replicas or parent instances of read replicas.
- **Unsupported Column Types** – You can't use the `reg` data types in any database tables that you plan to transport with this method. These types depend on system catalog object IDs (OIDs), which often change during transport.
- **Tablespaces** – All source database objects must be in the default `pg_default` tablespace.
- **Compatibility** – Both the source and destination DB instances must run the same major version of PostgreSQL.

Before transport begins, the `transport.import_from_server` function compares the source and destination DB instances to ensure database compatibility. This includes verifying PostgreSQL major version compatibility. Also, the function verifies that the destination DB instance likely has enough space to receive the source database. The function performs several additional checks to make sure that the transport is smooth.

- **Extensions** – The only extension that you can install on the source DB instance during transport is `pg_transport`.
- **Roles and ACLs** – The source database's access privileges and ownership information aren't carried over to the destination database. All database objects are created and owned by the local destination user of the transport.
- **Concurrent Transports** – You can run up to 32 total transports at the same time on a DB instance, including both imports and exports. To define the worker processes used for each transport, use the `pg_transport.work_mem` and `pg_transport.num_workers` parameters. To accommodate concurrent transports, you might need to increase the `max_worker_processes` parameter quite a bit. For more information, see [Configuration Parameters for the `pg_transport` Extension \(p. 1320\)](#).

Setting Up to Transport PostgreSQL Databases

To prepare to transport a PostgreSQL database from one DB instance to another, take the following steps.

To set up for transporting a PostgreSQL database

1. Make sure that the source DB instance's security group allows inbound traffic from the destination DB instance. This is required because the destination DB instance starts the database transport with an import call to the source DB instance. For information about how to use security groups, see [Controlling Access with Security Groups \(p. 486\)](#).
2. For both the source and destination DB instances, add `pg_transport` to the `shared_preload_libraries` parameter for each parameter group. The `shared_preload_libraries` parameter is static and requires a database restart for changes to take effect. For information about parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#).
3. For both the source and destination DB instances, install the required `pg_transport` PostgreSQL extension.

To do so, start `psql` as a user with the `rds_superuser` role for each DB instance, and then run the following command.

```
psql=> CREATE EXTENSION pg_transport;
```

Transporting a PostgreSQL Database Using the `transport.import_from_server` Function

After you complete the process described in [Setting Up to Transport PostgreSQL Databases \(p. 1317\)](#), you can start the transport. To do so, run the `transport.import_from_server` (p. 1319) function on the destination DB instance.

Note

Both the destination user for transport and the source user for the connection must be members of the `rds_superuser` role.

The destination DB instance can't already contain a database with the same name as the source database to be transported, or the transport fails.

The following shows an example transport.

```
SELECT transport.import_from_server(
    'source-db-instance-endpoint',
    'source-db-instance-port',
    'source-db-instance-user',
    'source-user-password',
    'source-database-name',
    'destination-user-password',
    false);
```

This function requires that you provide database user passwords. Thus, we recommend that you change the passwords of the user roles you used after transport is complete. Or, you can use SQL bind variables to create temporary user roles. Use these temporary roles for the transport and then discard the roles afterwards.

For details of the `transport.import_from_server` function and its parameters, see [transport.import_from_server Function Reference \(p. 1319\)](#).

What Happens During Database Transport

The `transport.import_from_server` function creates the in-transit database on the destination DB instance. The in-transit database is inaccessible on the destination DB instance for the duration of the transport.

When transport begins, all current sessions on the source database are ended. Any databases other than the source database on the source DB instance aren't affected by the transport.

The source database is put into a special read-only mode. While it's in this mode, you can connect to the source database and run read-only queries. However, write-enabled queries and some other types of commands are blocked. Only the specific source database that is being transported is affected by these restrictions.

During transport, you can't restore the destination DB instance to a point in time. This is because the transport isn't transactional and doesn't use the PostgreSQL write-ahead log to record changes. If the destination DB instance has automatic backups enabled, a backup is automatically taken after transport completes. Point-in-time restores are available for times after the backup finishes.

If the transport fails, the `pg_transport` extension attempts to undo all changes to the source and destination DB instances. This includes removing the destination's partially transported database. Depending on the type of failure, the source database might continue to reject write-enabled queries. If this happens, use the following command to allow write-enabled queries.

```
ALTER DATABASE my-database SET default_transaction_read_only = false;
```

transport.import_from_server Function Reference

The `transport.import_from_server` function transports a PostgreSQL database by importing it from a source DB instance to a destination DB instance. It does this by using a physical database connection transport mechanism.

Syntax

```
transport.import_from_server(
    host text,
    port int,
    username text,
    password text,
    database text,
    local_password text,
    dry_run bool
)
```

Return Value

None.

Parameters

You can find descriptions of the `transport.import_from_server` function parameters in the following table.

Parameter	Description
<code>host</code>	The endpoint of the source DB instance.
<code>port</code>	An integer representing the port of the source DB instance. PostgreSQL DB instances often use port 5432.
<code>username</code>	The user of the source DB instance. This user must be a member of the <code>rds_superuser</code> role.

Parameter	Description
password	The user password of the source DB instance.
database	The name of the database in the source DB instance to transport.
local_password	The local password of the current user for the destination DB instance. This user must be a member of the <code>rds_superuser</code> role.
dry_run	An optional Boolean value specifying whether to perform a dry run. The default is <code>false</code> , which means the transport proceeds. To confirm compatibility between the source and destination DB instances without performing the actual transport, set <code>dry_run</code> to <code>true</code> .

Example

For an example, see [Transporting a PostgreSQL Database Using the `transport.import_from_server` Function \(p. 1318\)](#).

Configuration Parameters for the pg_transport Extension

Use the following parameters to configure the pg_transport extension behavior.

```
SET pg_transport.num_workers = integer;
SET pg_transport.work_mem = kilobytes;
SET pg_transport.timing = Boolean;
```

You can find descriptions of these parameters in the following table.

Parameter	Description
<code>pg_transport.num_workers</code>	<p>The number of workers to use for a physical transport. The default is 3. Valid values are 1–32. Even large transports typically reach their maximum throughput with fewer than 8 workers.</p> <p>During transport, the <code>pg_transport.num_workers</code> setting on the destination DB instance is used on both the destination and source DB instances.</p> <p>A related parameter is the PostgreSQL <code>max_worker_processes</code> parameter. The transport process creates several background worker processes. Thus, your setting for the <code>pg_transport.num_workers</code> parameter might require you to set the <code>max_worker_processes</code> parameter significantly higher on both the source and destination DB instances.</p> <p>We recommend that you set <code>max_worker_processes</code> on both the source and destination DB instances to at least three times the destination DB instance's setting for the <code>pg_transport.num_workers</code> parameter. Add a few more to provide nontransport background worker processes.</p> <p>For more information about the <code>max_worker_processes</code> parameter, see the PostgreSQL documentation about Asynchronous Behavior.</p>
<code>pg_transport.timing</code>	A Boolean value that specifies whether to report timing information during the transport. The default is <code>true</code> . Valid values are <code>true</code> to report timing information and <code>false</code> to disable the reporting of timing information.

Parameter	Description
	We don't recommend that you set this parameter to <code>false</code> . Disabling <code>pg_transport.timing</code> significantly reduces your ability to track the progress of transports.
<code>pg_transport.work_mem</code>	The maximum amount of memory to allocate for each worker. The default is 131,072 kilobytes (KB). The minimum value is 64 megabytes (65,536 KB). Valid values are in kilobytes (KBs) as binary base-2 units, where 1 KB = 1,024 bytes. The transport might use less memory than is specified in this parameter. Even large transports typically reach their maximum throughput with less than 256 MB (262,144 KB) of memory per worker.

Common DBA Tasks for PostgreSQL

This section describes the Amazon RDS implementations of some common DBA tasks for DB instances running the PostgreSQL database engine. To deliver a managed service experience, Amazon RDS doesn't provide shell access to DB instances, and it restricts access to certain system procedures and tables that require advanced privileges.

For information about working with PostgreSQL log files on Amazon RDS, see [PostgreSQL Database Log Files \(p. 419\)](#).

Topics

- [Creating Roles \(p. 1322\)](#)
- [Managing PostgreSQL Database Access \(p. 1323\)](#)
- [Working with PostgreSQL Parameters \(p. 1323\)](#)
- [Working with PostgreSQL Autovacuum on Amazon RDS \(p. 1332\)](#)
- [Audit Logging for a PostgreSQL DB Instance \(p. 1340\)](#)
- [Working with the pgaudit Extension \(p. 1341\)](#)
- [Working with the pg_repack Extension \(p. 1342\)](#)
- [Working with PostGIS \(p. 1343\)](#)
- [Using pgBadger for Log Analysis with PostgreSQL \(p. 1345\)](#)
- [Viewing the Contents of pg_config \(p. 1346\)](#)
- [Working with the orafce Extension \(p. 1346\)](#)
- [Accessing External Data with the postgres_fdw Extension \(p. 1347\)](#)
- [Using a Custom DNS Server for Outbound Network Access \(p. 1348\)](#)
- [Restricting Password Management \(p. 1349\)](#)

Creating Roles

When you create a DB instance, the master user system account that you create is assigned to the `rds_superuser` role. The `rds_superuser` role is a predefined Amazon RDS role similar to the PostgreSQL superuser role (customarily named `postgres` in local instances), but with some restrictions. As with the PostgreSQL superuser role, the `rds_superuser` role has the most privileges for your DB instance. You should not assign this role to users unless they need the most access to the DB instance.

The `rds_superuser` role can do the following:

- Add extensions that are available for use with Amazon RDS. For more information, see [Supported PostgreSQL Features \(p. 1406\)](#) and the [PostgreSQL documentation](#).
- Manage tablespaces, including creating and deleting them. For more information, see the [Tablespaces](#) section in the PostgreSQL documentation.
- View all users not assigned the `rds_superuser` role using the `pg_stat_activity` command and kill their connections using the `pg_terminate_backend` and `pg_cancel_backend` commands.
- Grant and revoke the `rds_replication` role for all roles that are not the `rds_superuser` role. For more information, see the [GRANT](#) section in the PostgreSQL documentation.

The following example shows how to create a user and then grant the user the `rds_superuser` role. User-defined roles, such as `rds_superuser`, have to be granted.

```
create role testuser with password 'testuser' login;
CREATE ROLE
grant rds_superuser to testuser;
```

GRANT ROLE

Managing PostgreSQL Database Access

In Amazon RDS for PostgreSQL, you can manage which users have privileges to connect to which databases. In other PostgreSQL environments, you sometimes perform this kind of management by modifying the `pg_hba.conf` file. In Amazon RDS, you can use database grants instead.

New databases in PostgreSQL are always created with a default set of privileges. The default privileges allow `PUBLIC` (all users) to connect to the database and to create temporary tables while connected.

To control which users are allowed to connect to a given database in Amazon RDS, first revoke the default `PUBLIC` privileges. Then grant back the privileges on a more granular basis. The following example code shows how.

```
psql> revoke all on database <database-name> from public;
psql> grant connect, temporary on database <database-name> to <user/role name>;
```

For more information about privileges in PostgreSQL databases, see the [GRANT](#) command in the PostgreSQL documentation.

Working with PostgreSQL Parameters

PostgreSQL parameters that you set for a local PostgreSQL instance in the `postgresql.conf` file are maintained in the DB parameter group for your DB instance. If you create a DB instance using the default parameter group, the parameter settings are in the parameter group called `default.postgres9.6`.

When you create a DB instance, the parameters in the associated DB parameter group are loaded. You can modify parameter values by changing values in the parameter group. You can also change parameter values, if you have the security privileges to do so, by using the `ALTER DATABASE`, `ALTER ROLE`, and `SET` commands. You can't use the command line `postgres` command or the `env PGOPTIONS` command, because you have no access to the host.

Keeping track of PostgreSQL parameter settings can occasionally be difficult. Use the following command to list current parameter settings and the default value.

```
select name, setting, boot_val, reset_val, unit
from pg_settings
order by name;
```

For an explanation of the output values, see the [pg_settings](#) topic in the PostgreSQL documentation.

If you set the memory settings too large for `max_connections`, `shared_buffers`, or `effective_cache_size`, you will prevent the PostgreSQL instance from starting up. Some parameters use units that you might not be familiar with; for example, `shared_buffers` sets the number of 8-KB shared memory buffers used by the server.

The following error is written to the `postgres.log` file when the instance is attempting to start up, but incorrect parameter settings are preventing it from starting.

```
2013-09-18 21:13:15 UTC::@[8097]:FATAL:  could not map anonymous shared
memory: Cannot allocate memory
2013-09-18 21:13:15 UTC::@[8097]:HINT:  This error usually means that
PostgreSQL's request for a shared memory segment exceeded available memory or
swap space. To reduce the request size (currently 3514134274048 bytes), reduce
PostgreSQL's shared memory usage, perhaps by reducing shared_buffers or
max_connections.
```

There are two types of PostgreSQL parameters, static and dynamic. Static parameters require that the DB instance be rebooted before they are applied. Dynamic parameters can be applied immediately. The following table shows parameters that you can modify for a PostgreSQL DB instance and each parameter's type.

Parameter Name	Apply_Type	Description
application_name	Dynamic	Sets the application name to be reported in statistics and logs.
array_nulls	Dynamic	Enables input of NULL elements in arrays.
authentication_timeout	Dynamic	Sets the maximum allowed time to complete client authentication.
autovacuum	Dynamic	Starts the autovacuum subprocess.
autovacuum_analyze_scale_factor	Dynamic	Number of tuple inserts, updates, or deletes before analyze as a fraction of reltuples.
autovacuum_analyze_threshold	Dynamic	Minimum number of tuple inserts, updates, or deletes before analyze.
autovacuum_naptime	Dynamic	Time to sleep between autovacuum runs.
autovacuum_vacuum_cost_delay	Dynamic	Vacuum cost delay, in milliseconds, for autovacuum.
autovacuum_vacuum_cost_limit	Dynamic	Vacuum cost amount available before napping, for autovacuum.
autovacuum_vacuum_scale_factor	Dynamic	Number of tuple updates or deletes before vacuum as a fraction of reltuples.
autovacuum_vacuum_threshold	Dynamic	Minimum number of tuple updates or deletes before vacuum.
backslash_quote	Dynamic	Sets whether a backslash (\) is allowed in string literals.
bgwriter_delay	Dynamic	Background writer sleep time between rounds.
bgwriter_lru_maxpages	Dynamic	Background writer maximum number of LRU pages to flush per round.
bgwriter_lru_multiplier	Dynamic	Multiple of the average buffer usage to free per round.
bytea_output	Dynamic	Sets the output format for bytes.
check_function_bodies	Dynamic	Checks function bodies during CREATE FUNCTION.
checkpoint_completion_target	Dynamic	Time spent flushing dirty buffers during checkpoint, as a fraction of the checkpoint interval.
checkpoint_segments	Dynamic	Sets the maximum distance in log segments between automatic write-ahead log (WAL) checkpoints.
checkpoint_timeout	Dynamic	Sets the maximum time between automatic WAL checkpoints.

Parameter Name	Apply_Type	Description
checkpoint_warning	Dynamic	Enables warnings if checkpoint segments are filled more frequently than this.
client_encoding	Dynamic	Sets the client's character set encoding.
client_min_messages	Dynamic	Sets the message levels that are sent to the client.
commit_delay	Dynamic	Sets the delay in microseconds between transaction commit and flushing WAL to disk.
commit_siblings	Dynamic	Sets the minimum concurrent open transactions before performing commit_delay.
constraint_exclusion	Dynamic	Enables the planner to use constraints to optimize queries.
cpu_index_tuple_cost	Dynamic	Sets the planner's estimate of the cost of processing each index entry during an index scan.
cpu_operator_cost	Dynamic	Sets the planner's estimate of the cost of processing each operator or function call.
cpu_tuple_cost	Dynamic	Sets the planner's estimate of the cost of processing each tuple (row).
cursor_tuple_fraction	Dynamic	Sets the planner's estimate of the fraction of a cursor's rows that will be retrieved.
datestyle	Dynamic	Sets the display format for date and time values.
deadlock_timeout	Dynamic	Sets the time to wait on a lock before checking for deadlock.
debug_pretty_print	Dynamic	Indents parse and plan tree displays.
debug_print_parse	Dynamic	Logs each query's parse tree.
debug_print_plan	Dynamic	Logs each query's execution plan.
debug_print_rewritten	Dynamic	Logs each query's rewritten parse tree.
default_statistics_target	Dynamic	Sets the default statistics target.
default_tablespace	Dynamic	Sets the default tablespace to create tables and indexes in.
default_transaction_deferrable	Dynamic	Sets the default deferrable status of new transactions.
default_transaction_isolation	Dynamic	Sets the transaction isolation level of each new transaction.
default_transaction_read_only	Dynamic	Sets the default read-only status of new transactions.
default_with_oids	Dynamic	Creates new tables with OIDs by default.
effective_cache_size	Dynamic	Sets the planner's assumption about the size of the disk cache.

Parameter Name	Apply_Type	Description
<code>effective_io_concurrency</code>	Dynamic	Number of simultaneous requests that can be handled efficiently by the disk subsystem.
<code>enable_bitmapscan</code>	Dynamic	Enables the planner's use of bitmap-scan plans.
<code>enable_hashagg</code>	Dynamic	Enables the planner's use of hashed aggregation plans.
<code>enable_hashjoin</code>	Dynamic	Enables the planner's use of hash join plans.
<code>enable_indexscan</code>	Dynamic	Enables the planner's use of index-scan plans.
<code>enable_material</code>	Dynamic	Enables the planner's use of materialization.
<code>enable_mergejoin</code>	Dynamic	Enables the planner's use of merge join plans.
<code>enable_nestloop</code>	Dynamic	Enables the planner's use of nested-loop join plans.
<code>enable_seqscan</code>	Dynamic	Enables the planner's use of sequential-scan plans.
<code>enable_sort</code>	Dynamic	Enables the planner's use of explicit sort steps.
<code>enable_tidscan</code>	Dynamic	Enables the planner's use of TID scan plans.
<code>escape_string_warning</code>	Dynamic	Warns about backslash (\) escapes in ordinary string literals.
<code>extra_float_digits</code>	Dynamic	Sets the number of digits displayed for floating-point values.
<code>fromCollapse_limit</code>	Dynamic	Sets the FROM-list size beyond which subqueries are not collapsed.
<code>fsync</code>	Dynamic	Forces synchronization of updates to disk.
<code>full_page_writes</code>	Dynamic	Writes full pages to WAL when first modified after a checkpoint.
<code>geqo</code>	Dynamic	Enables genetic query optimization.
<code>geqo_effort</code>	Dynamic	GEQO: effort is used to set the default for other GEQO parameters.
<code>geqo_generations</code>	Dynamic	GEQO: number of iterations of the algorithm.
<code>geqo_pool_size</code>	Dynamic	GEQO: number of individuals in the population.
<code>geqo_seed</code>	Dynamic	GEQO: seed for random path selection.
<code>geqo_selection_bias</code>	Dynamic	GEQO: selective pressure within the population.
<code>geqo_threshold</code>	Dynamic	Sets the threshold of FROM items beyond which GEQO is used.
<code>gin_fuzzy_search_limit</code>	Dynamic	Sets the maximum allowed result for exact search by GIN.
<code>hot_standby_feedback</code>	Dynamic	Determines whether a hot standby sends feedback messages to the primary or upstream standby.

Parameter Name	Apply_Type	Description
intervalstyle	Dynamic	Sets the display format for interval values.
joinCollapse_limit	Dynamic	Sets the FROM-list size beyond which JOIN constructs are not flattened.
lc_messages	Dynamic	Sets the language in which messages are displayed.
lc_monetary	Dynamic	Sets the locale for formatting monetary amounts.
lc_numeric	Dynamic	Sets the locale for formatting numbers.
lc_time	Dynamic	Sets the locale for formatting date and time values.
log_autovacuum_min_duration	Dynamic	Sets the minimum execution time above which autovacuum actions will be logged.
log_checkpoints	Dynamic	Logs each checkpoint.
log_connections	Dynamic	Logs each successful connection.
log_disconnections	Dynamic	Logs end of a session, including duration.
log_duration	Dynamic	Logs the duration of each completed SQL statement.
log_error_verbosity	Dynamic	Sets the verbosity of logged messages.
log_executor_stats	Dynamic	Writes executor performance statistics to the server log.
log_filename	Dynamic	Sets the file name pattern for log files.
log_hostname	Dynamic	Logs the host name in the connection logs.
log_lock_waits	Dynamic	Logs long lock waits.
log_min_duration_statement	Dynamic	Sets the minimum execution time above which statements will be logged.
log_min_error_statement	Dynamic	Causes all statements generating an error at or above this level to be logged.
log_min_messages	Dynamic	Sets the message levels that are logged.
log_parser_stats	Dynamic	Writes parser performance statistics to the server log.
log_planner_stats	Dynamic	Writes planner performance statistics to the server log.
log_rotation_age	Dynamic	Automatic log file rotation will occur after N minutes.
log_rotation_size	Dynamic	Automatic log file rotation will occur after N kilobytes.
log_statement	Dynamic	Sets the type of statements logged.
log_statement_stats	Dynamic	Writes cumulative performance statistics to the server log.
log_temp_files	Dynamic	Logs the use of temporary files larger than this number of kilobytes.

Parameter Name	Apply_Type	Description
<code>maintenance_work_mem</code>	Dynamic	Sets the maximum memory to be used for maintenance operations.
<code>max_stack_depth</code>	Dynamic	Sets the maximum stack depth, in kilobytes.
<code>max_standby_archive_delay</code>	Dynamic	Sets the maximum delay before canceling queries when a hot standby server is processing archived WAL data.
<code>max_standby_streaming_delay</code>	Dynamic	Sets the maximum delay before canceling queries when a hot standby server is processing streamed WAL data.
<code>max_wal_size</code>	Static	Sets the WAL size that triggers the checkpoint. For PostgreSQL version 9.6 and earlier, <code>max_wal_size</code> is in units of 16 MB. For PostgreSQL version 10 and later, <code>max_wal_size</code> is in units of 1 MB.
<code>min_wal_size</code>	Static	Sets the minimum size to shrink the WAL to. For PostgreSQL version 9.6 and earlier, <code>min_wal_size</code> is in units of 16 MB. For PostgreSQL version 10 and later, <code>min_wal_size</code> is in units of 1 MB.
<code>quote_all_identifiers</code>	Dynamic	Adds quotes ("") to all identifiers when generating SQL fragments.
<code>random_page_cost</code>	Dynamic	Sets the planner's estimate of the cost of a non-sequentially fetched disk page.
<code>rds.adaptive_autovacuum</code>	Dynamic	Automatically tunes the autovacuum parameters whenever the transaction ID thresholds are exceeded.
<code>rds.log_retention_period</code>	Dynamic	Sets log retention such that Amazon RDS deletes PostgreSQL logs that are older than N minutes.
<code>rds.restrict_password_commands</code>	Static	Restricts who can manage passwords to users with the <code>rds_password</code> role. Set this parameter to 1 to enable password restriction. The default is 0.
<code>search_path</code>	Dynamic	Sets the schema search order for names that are not schema-qualified.
<code>seq_page_cost</code>	Dynamic	Sets the planner's estimate of the cost of a sequentially fetched disk page.
<code>session_replication_role</code>	Dynamic	Sets the sessions behavior for triggers and rewrite rules.
<code>sql_inheritance</code>	Dynamic	Causes subtables to be included by default in various commands.
<code>ssl_renegotiation_limit</code>	Dynamic	Sets the amount of traffic to send and receive before renegotiating the encryption keys.
<code>standard_conforming_strings</code>	Dynamic	Causes ... strings to treat backslashes literally.
<code>statement_timeout</code>	Dynamic	Sets the maximum allowed duration of any statement.

Parameter Name	Apply_Type	Description
<code>synchronize_seqscans</code>	Dynamic	Enables synchronized sequential scans.
<code>synchronous_commit</code>	Dynamic	Sets the current transactions synchronization level.
<code>tcp_keepalives_count</code>	Dynamic	Maximum number of TCP keepalive retransmits.
<code>tcp_keepalives_idle</code>	Dynamic	Time between issuing TCP keepalives.
<code>tcp_keepalives_interval</code>	Dynamic	Time between TCP keepalive retransmits.
<code>temp_buffers</code>	Dynamic	Sets the maximum number of temporary buffers used by each session.
<code>temp_tablespaces</code>	Dynamic	Sets the tablespaces to use for temporary tables and sort files.
<code>timezone</code>	Dynamic	Sets the time zone for displaying and interpreting time stamps.
<code>track_activities</code>	Dynamic	Collects information about executing commands.
<code>track_counts</code>	Dynamic	Collects statistics on database activity.
<code>track_functions</code>	Dynamic	Collects function-level statistics on database activity.
<code>track_io_timing</code>	Dynamic	Collects timing statistics on database I/O activity.
<code>transaction_deferrable</code>	Dynamic	Indicates whether to defer a read-only serializable transaction until it can be executed with no possible serialization failures.
<code>transaction_isolation</code>	Dynamic	Sets the current transactions isolation level.
<code>transaction_read_only</code>	Dynamic	Sets the current transactions read-only status.
<code>transform_null_equals</code>	Dynamic	Treats <code>expr=NULL</code> as <code>expr IS NULL</code> .
<code>update_process_title</code>	Dynamic	Updates the process title to show the active SQL command.
<code>vacuum_cost_delay</code>	Dynamic	Vacuum cost delay in milliseconds.
<code>vacuum_cost_limit</code>	Dynamic	Vacuum cost amount available before napping.
<code>vacuum_cost_page_dirty</code>	Dynamic	Vacuum cost for a page dirtied by vacuum.
<code>vacuum_cost_page_hit</code>	Dynamic	Vacuum cost for a page found in the buffer cache.
<code>vacuum_cost_page_miss</code>	Dynamic	Vacuum cost for a page not found in the buffer cache.
<code>vacuum_defer_cleanup_age</code>	Dynamic	Number of transactions by which vacuum and hot cleanup should be deferred, if any.
<code>vacuum_freeze_min_age</code>	Dynamic	Minimum age at which vacuum should freeze a table row.
<code>vacuum_freeze_table_age</code>	Dynamic	Age at which vacuum should scan a whole table to freeze tuples.
<code>wal_writer_delay</code>	Dynamic	WAL writer sleep time between WAL flushes.

Parameter Name	Apply_Type	Description
<code>work_mem</code>	Dynamic	Sets the maximum memory to be used for query workspaces.
<code>xmlbinary</code>	Dynamic	Sets how binary values are to be encoded in XML.
<code>xmloption</code>	Dynamic	Sets whether XML data in implicit parsing and serialization operations is to be considered as documents or content fragments.
<code>autovacuum_freeze_max_age</code>	Static	Age at which to autovacuum a table to prevent transaction ID wraparound.
<code>autovacuum_max_workers</code>	Static	Sets the maximum number of simultaneously running autovacuum worker processes.
<code>max_connections</code>	Static	Sets the maximum number of concurrent connections.
<code>max_files_per_process</code>	Static	Sets the maximum number of simultaneously open files for each server process.
<code>max_locks_per_transaction</code>	Static	Sets the maximum number of locks per transaction.
<code>max_pred_locks_per_transaction</code>	Static	Sets the maximum number of predicate locks per transaction.
<code>max_prepared_transactions</code>	Static	Sets the maximum number of simultaneously prepared transactions.
<code>shared_buffers</code>	Static	Sets the number of shared memory buffers used by the server.
<code>ssl</code>	Static	Enables SSL connections.
<code>temp_file_limit</code>	Static	Sets the maximum size in KB to which the temporary files can grow.
<code>track_activity_query_size</code>	Static	Sets the size reserved for <code>pg_stat_activity.current_query</code> , in bytes.
<code>wal_buffers</code>	Static	Sets the number of disk-page buffers in shared memory for WAL.

Amazon RDS uses the default PostgreSQL units for all parameters. The following table shows the PostgreSQL default unit and value for each parameter.

Parameter Name	Unit
<code>effective_cache_size</code>	8 KB
<code>segment_size</code>	8 KB
<code>shared_buffers</code>	8 KB
<code>temp_buffers</code>	8 KB
<code>wal_buffers</code>	8 KB

Parameter Name	Unit
wal_segment_size	8 KB
log_rotation_size	KB
log_temp_files	KB
maintenance_work_mem	KB
max_stack_depth	KB
ssl_renegotiation_limit	KB
temp_file_limit	KB
work_mem	KB
log_rotation_age	minutes
autovacuum_vacuum_cost_delay	ms
bgwriter_delay	ms
deadlock_timeout	ms
lock_timeout	ms
log_autovacuum_min_duration	ms
log_min_duration_statement	ms
max_standby_archive_delay	ms
max_standby_streaming_delay	ms
statement_timeout	ms
vacuum_cost_delay	ms
wal_receiver_timeout	ms
wal_sender_timeout	ms
wal_writer_delay	ms
archive_timeout	s
authentication_timeout	s
autovacuum_naptime	s
checkpoint_timeout	s
checkpoint_warning	s
post_auth_delay	s
pre_auth_delay	s
tcp_keepalives_idle	s
tcp_keepalives_interval	s

Parameter Name	Unit
wal_receiver_status_interval	s

Working with PostgreSQL Autovacuum on Amazon RDS

We strongly recommend that you use the autovacuum feature for PostgreSQL databases to maintain the health of your PostgreSQL DB instance. Autovacuum automates the execution of the VACUUM and the ANALYZE commands. Autovacuum checks for tables that have had a large number of inserted, updated, or deleted tuples. Autovacuum then reclaims storage by removing obsolete data or tuples from the PostgreSQL database.

Autovacuum is enabled by default for all new Amazon RDS PostgreSQL DB instances, and the related autovacuum configuration parameters are appropriately set by default. Because our defaults are somewhat generic, you can benefit from tuning parameters to your specific workload. The following section can help you perform the needed autovacuum tuning.

Topics

- [Allocating Memory for Autovacuum \(p. 1332\)](#)
- [Reducing the Likelihood of Transaction ID Wraparound \(p. 1333\)](#)
- [Determining if the Tables in Your Database Need Vacuuming \(p. 1333\)](#)
- [Determining Which Tables Are Currently Eligible for Autovacuum \(p. 1334\)](#)
- [Determining if Autovacuum Is Currently Running and For How Long \(p. 1335\)](#)
- [Performing a Manual Vacuum Freeze \(p. 1337\)](#)
- [Reindexing a Table When Autovacuum Is Running \(p. 1338\)](#)
- [Other Parameters That Affect Autovacuum \(p. 1339\)](#)
- [Setting Table-Level Autovacuum Parameters \(p. 1340\)](#)
- [Autovacuum Logging \(p. 1340\)](#)

Allocating Memory for Autovacuum

One of the most important parameters influencing autovacuum performance is the `maintenance_work_mem` parameter. This parameter determines how much memory that you allocate for autovacuum to use to scan a database table and to hold all the row IDs that are going to be vacuumed. If you set the value of the `maintenance_work_mem` parameter too low, the vacuum process might have to scan the table multiple times to complete its work. Such multiple scans can have a negative impact on performance.

When doing calculations to determine the `maintenance_work_mem` parameter value, keep in mind two things:

- The default unit is kilobytes (KB) for this parameter.
- The `maintenance_work_mem` parameter works in conjunction with the `autovacuum_max_workers` parameter. If you have many small tables, allocate more `autovacuum_max_workers` and less `maintenance_work_mem`. If you have large tables (say, larger than 100 GB), allocate more memory and fewer worker processes. You need to have enough memory allocated to succeed on your biggest table. Each `autovacuum_max_workers` can use the memory you allocate. Thus, you should make sure the combination of worker processes and memory equal the total memory that you want to allocate.

In general terms, for large hosts set the `maintenance_work_mem` parameter to a value between one and two gigabytes (between 1,048,576 and 2,097,152 KB). For extremely large hosts, set the parameter to a value between two and four gigabytes (between 2,097,152 and 4,194,304 KB). The value you set for this parameter should depend on the workload. Amazon RDS has updated its default for this parameter to be kilobytes calculated as follows:

```
GREATEST({DBInstanceClassMemory/63963136*1024}, 65536).
```

Reducing the Likelihood of Transaction ID Wraparound

In some cases, parameter group settings related to autovacuum might not be aggressive enough to prevent transaction ID wraparound. To address this, Amazon RDS for PostgreSQL provides a mechanism that adapts the autovacuum parameter values automatically. *Adaptive autovacuum parameter tuning* is a feature for RDS for PostgreSQL versions 9.4 and newer. A detailed explanation of [TransactionID wraparound](#) is found in the PostgreSQL documentation.

Adaptive autovacuum parameter tuning is enabled by default for RDS PostgreSQL instances with the dynamic parameter `rds.adaptive_autovacuum` set to ON. We strongly recommend that you keep this enabled. However, to turn off adaptive autovacuum parameter tuning, set the `rds.adaptive_autovacuum` parameter to 0 or OFF.

Transaction ID wraparound is still possible even when RDS tunes the autovacuum parameters. We encourage you to implement an Amazon CloudWatch alarm for transaction ID wraparound. For more information, see the blog post [Implement an Early Warning System for Transaction ID Wraparound in Amazon RDS for PostgreSQL](#).

With adaptive autovacuum parameter tuning enabled, RDS will begin adjusting autovacuum parameters when the CloudWatch metric `MaximumUsedTransactionIDs` reaches the greater of the following values:

- 750,000,000
- The autovacuum parameter `autovacuum_freeze_max_age`

RDS continues to adjust parameters for autovacuum if a table continues to trend toward transaction ID wraparound. Each of these adjustments dedicates more resources to autovacuum to avoid wraparound. RDS updates the following autovacuum-related parameters:

- `autovacuum_vacuum_cost_delay`
- `autovacuum_vacuum_cost_limit`
- `autovacuum_work_mem`
- `autovacuum_naptime`

RDS modifies these parameters only if the new value makes autovacuum more aggressive. The parameters are modified in memory on the DB instance. The values in the parameter group aren't changed. To view the current in-memory settings, use the PostgreSQL `SHOW` SQL command.

Whenever RDS modifies any of these autovacuum parameters, it generates an event for the affected DB instance that is visible on the AWS Management Console (<https://console.aws.amazon.com/rds/>) and through the RDS API. After the `MaximumUsedTransactionIDs` CloudWatch metric returns below the threshold, RDS resets the autovacuum related parameters in memory back to the values specified in the parameter group and generates another event corresponding to this change.

Determining if the Tables in Your Database Need Vacuuming

You can use the following query to show the number of unvacuumed transactions in a database. The `datfrozenxid` column of a database's `pg_database` row is a lower bound on the normal transaction

IDs appearing in that database. This column is the minimum of the per-table `realfrozenxid` values within the database.

```
SELECT datname, age(datfrozenxid) FROM pg_database ORDER BY age(datfrozenxid) desc limit 20;
```

For example, the results of running the preceding query might be the following.

datname	age
mydb	1771757888
template0	1721757888
template1	1721757888
rdsadmin	1694008527
postgres	1693881061

(5 rows)

When the age of a database reaches 2 billion transaction IDs, transaction ID (XID) wraparound occurs and the database becomes read-only. This query can be used to produce a metric and run a few times a day. By default, autovacuum is set to keep the age of transactions to no more than 200,000,000 ([autovacuum_freeze_max_age](#)).

A sample monitoring strategy might look like this:

- Set the `autovacuum_freeze_max_age` value to 200 million transactions.
- If a table reaches 500 million unvacuumed transactions, that triggers a low-severity alarm. This isn't an unreasonable value, but it can indicate that autovacuum isn't keeping up.
- If a table ages to 1 billion, this should be treated as an alarm to take action on. In general, you want to keep ages closer to `autovacuum_freeze_max_age` for performance reasons. We recommend you investigate using the recommendations that follow.
- If a table reaches 1.5 billion unvacuumed transactions, that triggers a high-severity alarm. Depending on how quickly your database uses transaction IDs, this alarm can indicate that the system is running out of time to run autovacuum. In this case, we recommend you resolve this immediately.

If a table is constantly breaching these thresholds, you need to modify your autovacuum parameters further. By default, using VACUUM manually (which has cost-based delays disabled) is more aggressive than using the default autovacuum, but it is also more intrusive to the system as a whole.

We recommend the following:

- Be aware and enable a monitoring mechanism so that you are aware of the age of your oldest transactions.

For information on creating a process that warns you about transaction ID wraparound, see the AWS Database Blog post [Implement an Early Warning System for Transaction ID Wraparound in Amazon RDS for PostgreSQL](#).

- For busier tables, perform a manual vacuum freeze regularly during a maintenance window, in addition to relying on autovacuum. For information on performing a manual vacuum freeze, see [Performing a Manual Vacuum Freeze \(p. 1337\)](#).

Determining Which Tables Are Currently Eligible for Autovacuum

Often, it is one or two tables in need of vacuuming. Tables whose `realfrozenxid` value is greater than the number of transactions in `autovacuum_freeze_max_age` are always targeted by autovacuum.

Otherwise, if the number of tuples made obsolete since the last VACUUM exceeds the "vacuum threshold", the table is vacuumed.

The [autovacuum threshold](#) is defined as:

```
Vacuum-threshold = vacuum-base-threshold + vacuum-scale-factor * number-of-tuples
```

While you are connected to your database, run the following query to see a list of tables that autovacuum sees as eligible for vacuuming:

```

WITH vbt AS (SELECT setting AS autovacuum_vacuum_threshold FROM
pg_settings WHERE name = 'autovacuum_vacuum_threshold')
     , vsf AS (SELECT setting AS autovacuum_vacuum_scale_factor FROM
pg_settings WHERE name = 'autovacuum_vacuum_scale_factor')
     , fma AS (SELECT setting AS autovacuum_freeze_max_age FROM
pg_settings WHERE name = 'autovacuum_freeze_max_age')
     , sto AS (select opt_oid, split_part(setting, '=', 1) as param,
split_part(setting, '=', 2) as value from (select oid opt_oid,
unnest(reloptions) setting from pg_class) opt)
SELECT
      '''||ns.nspname||'.'||c.relname||''' as relation
     , pg_size.pretty(pg_table_size(c.oid)) as table_size
     , age(relfrozenxid) as xid_age
     , coalesce(cfma.value::float, autovacuum_freeze_max_age::float)
autovacuum_freeze_max_age
     , (coalesce(cvbt.value::float, autovacuum_vacuum_threshold::float)
+ coalesce(cvsf.value::float, autovacuum_vacuum_scale_factor::float) *
c.reltuples) as autovacuum_vacuum_tuples
     , n_dead_tup as dead_tuples
FROM pg_class c join pg_namespace ns on ns.oid = c.relnamespace
join pg_stat_all_tables stat on stat.relid = c.oid
join vbt on (1=1) join vsf on (1=1) join fma on (1=1)
left join sto cvbt on cvbt.param = 'autovacuum_vacuum_threshold' and
c.oid = cvbt.opt_oid
left join sto cvsf on cvsf.param = 'autovacuum_vacuum_scale_factor' and
c.oid = cvsf.opt_oid
left join sto cfma on cfma.param = 'autovacuum_freeze_max_age' and
c.oid = cfma.opt_oid
WHERE c.relkind = 'r' and nspname <> 'pg_catalog'
and (
      age(relfrozenxid) >= coalesce(cfma.value::float,
autovacuum_freeze_max_age::float)
      or
      coalesce(cvbt.value::float, autovacuum_vacuum_threshold::float) +
coalesce(cvsf.value::float, autovacuum_vacuum_scale_factor::float) *
c.reltuples <= n_dead_tup
      -- or 1 = 1
)
ORDER BY age(relfrozenxid) DESC LIMIT 50;
```

Determining if Autovacuum Is Currently Running and For How Long

If you need to manually vacuum a table, you need to determine if autovacuum is currently running. If it is, you might need to adjust parameters to make it run more efficiently, or terminate autovacuum so you can manually run VACUUM.

Use the following query to determine if autovacuum is running, how long it has been running, and if it is waiting on another session.

If you are using Amazon RDS PostgreSQL 9.6+ or higher, use this query:

```
SELECT datname, usename, pid, state, wait_event, current_timestamp - xact_start AS
xact_runtime, query
FROM pg_stat_activity
WHERE upper(query) LIKE '%VACUUM%'
ORDER BY xact_start;
```

After running the query, you should see output similar to the following.

datname	usename	pid	state	wait_event	xact_runtime	query
mydb	rdsadmin	16473	active		33 days 16:32:11.600656	autovacuum: VACUUM ANALYZE public.mytable1 (to prevent wraparound)
mydb	rdsadmin	22553	active		14 days 09:15:34.073141	autovacuum: VACUUM ANALYZE public.mytable2 (to prevent wraparound)
mydb	rdsadmin	41909	active		3 days 02:43:54.203349	autovacuum: VACUUM ANALYZE public.mytable3
mydb	rdsadmin	618	active		00:00:00	SELECT datname, usename, pid, state, wait_event, current_timestamp - xact_start AS xact_runtime, query+ FROM pg_stat_activity WHERE query like '%VACUUM%' ORDER BY xact_start;

If you are using a version less than Amazon RDS PostgreSQL 9.6, but, 9.3.12 or later, 9.4.7 or later, or 9.5.2+, use this query:

```
SELECT datname, usename, pid, waiting, current_timestamp - xact_start AS xact_runtime,
query
FROM pg_stat_activity
WHERE upper(query) LIKE '%VACUUM%'
ORDER BY xact_start;
```

After running the query, you should see output similar to the following.

datname	usename	pid	waiting	xact_runtime	query
mydb	rdsadmin	16473	f	33 days 16:32:11.600656	autovacuum: VACUUM ANALYZE public.mytable1 (to prevent wraparound)
mydb	rdsadmin	22553	f	14 days 09:15:34.073141	autovacuum: VACUUM ANALYZE public.mytable2 (to prevent wraparound)
mydb	rdsadmin	41909	f	3 days 02:43:54.203349	autovacuum: VACUUM ANALYZE public.mytable3
mydb	rdsadmin	618	f	00:00:00	SELECT datname, usename, pid, waiting, current_timestamp - xact_start AS xact_runtime, query+ FROM pg_stat_activity WHERE query like '%VACUUM%'

```
|           |           |           |           | ORDER BY xact_start;  
+-----+
```

Several issues can cause a long-running autovacuum session (that is, multiple days long). The most common issue is that your `maintenance_work_mem` parameter value is set too low for the size of the table or rate of updates.

We recommend that you use the following formula to set the `maintenance_work_mem` parameter value.

```
GREATEST({DBInstanceClassMemory/63963136*1024},65536)
```

Short running autovacuum sessions can also indicate problems:

- It can indicate that there aren't enough `autovacuum_max_workers` for your workload. In this case, you need to indicate the number of workers.
- It can indicate that there is an index corruption (autovacuum crashes and restart on the same relation but make no progress). In this case, run a manual vacuum freeze verbose `__table__` to see the exact cause.

Performing a Manual Vacuum Freeze

You might want to perform a manual vacuum on a table that has a vacuum process already running. This is useful if you have identified a table with an age approaching 2 billion transactions (or above any threshold you are monitoring).

The following steps are a guideline, and there are several variations to the process. For example, during testing, suppose that you find that the `maintenance_work_mem` parameter value was set too small and that you need to take immediate action on a table. However, perhaps you don't want to bounce the instance at the moment. Using the queries in previous sections, you determine which table is the problem and notice a long running autovacuum session. You know that you need to change the `maintenance_work_mem` parameter setting, but you also need to take immediate action and vacuum the table in question. The following procedure shows what to do in this situation:

To manually perform a vacuum freeze

1. Open two sessions to the database containing the table you want to vacuum. For the second session, use "screen" or another utility that maintains the session if your connection is dropped.
2. In session one, get the PID of the autovacuum session running on the table. This action requires that you are running Amazon RDS PostgreSQL 9.3.12 or later, 9.4.7 or later, or 9.5.2 or later to have full visibility into the running rdsadmin processes.

Run the following query to get the PID of the autovacuum session.

```
SELECT datname, username, pid, waiting, current_timestamp - xact_start  
AS xact_runtime, query  
FROM pg_stat_activity WHERE upper(query) LIKE '%VACUUM%' ORDER BY  
xact_start;
```

3. In session two, calculate the amount of memory you need for this operation. In this example, we determine that we can afford to use up to 2 GB of memory for this operation, so we set `maintenance_work_mem` for the current session to 2 GB.

```
set maintenance_work_mem='2 GB';  
SET
```

4. In session two, issue a `vacuum freeze verbose` command for the table. The verbose setting is useful because, although there is no progress report for this in PostgreSQL currently, you can see activity.

```
\timing on
Timing is on.
vacuum freeze verbose pgbench_branches;
```

```
INFO: vacuuming "public.pgbench_branches"
INFO: index "pgbench_branches_pkey" now contains 50 row versions in 2 pages
DETAIL: 0 index row versions were removed.
0 index pages have been deleted, 0 are currently reusable.
CPU 0.00s/0.00u sec elapsed 0.00 sec.
INFO: index "pgbench_branches_test_index" now contains 50 row versions in 2 pages
DETAIL: 0 index row versions were removed.
0 index pages have been deleted, 0 are currently reusable.
CPU 0.00s/0.00u sec elapsed 0.00 sec.
INFO: "pgbench_branches": found 0 removable, 50 nonremovable row versions
      in 43 out of 43 pages
DETAIL: 0 dead row versions cannot be removed yet.
There were 9347 unused item pointers.
0 pages are entirely empty.
CPU 0.00s/0.00u sec elapsed 0.00 sec.
VACUUM
Time: 2.765 ms
```

5. In session one, if autovacuum was blocking, you see in `pg_stat_activity` that waiting is "T" for your vacuum session. In this case, you need to terminate the autovacuum process as follows.

```
SELECT pg_terminate_backend('the_pid');
```

6. At this point, your session begins. It's important to note that autovacuum restarts immediately because this table is probably the highest on its list of work. Initiate your `vacuum freeze verbose` command in session 2 and then terminate the autovacuum process in session 1.

Reindexing a Table When Autovacuum Is Running

If an index has become corrupt, autovacuum continues to process the table and fails. If you attempt a manual vacuum in this situation, you will receive an error message similar to the following:

```
mydb=# vacuum freeze pgbench_branches;
ERROR: index "pgbench_branches_test_index" contains unexpected
       zero page at block 30521
HINT: Please REINDEX it.
```

When the index is corrupted and autovacuum is attempting to run against the table, you contend with an already running autovacuum session. When you issue a "`REINDEX`" command, you take out an exclusive lock on the table. Write operations are blocked, and also reads that use that specific index.

To reindex a table when autovacuum is running on the table

1. Open two sessions to the database containing the table you want to vacuum. For the second session, use "screen" or another utility that maintains the session if your connection is dropped.
2. In session one, get the PID of the autovacuum session running on the table. This action requires that you are running Amazon RDS PostgreSQL 9.3.12 or later, 9.4.7 or later, or 9.5.2 or later to have full visibility into the running rdsadmin processes.

Run the following query to get the PID of the autovacuum session:

```
SELECT datname, usename, pid, waiting, current_timestamp - xact_start
AS xact_runtime, query
FROM pg_stat_activity WHERE upper(query) like '%VACUUM%' ORDER BY
xact_start;
```

3. In session two, issue the reindex command.

```
\timing on
Timing is on.
reindex index pgbench_branches_test_index;
REINDEX
Time: 9.966 ms
```

4. In session one, if autovacuum was blocking, you see in pg_stat_activity that waiting is "T" for your vacuum session. In this case, you will need to terminate the autovacuum process.

```
select pg_terminate_backend('the_pid');
```

5. At this point, your session begins. It's important to note that autovacuum restarts immediately because this table is probably the highest on its list of work. Initiate your command in session 2 and then terminate the autovacuum process in session 1.

Other Parameters That Affect Autovacuum

The following query shows the values of some of the parameters that directly affect autovacuum and its behavior. The [autovacuum parameters](#) are described fully in the PostgreSQL documentation.

```
SELECT name, setting, unit, short_desc
FROM pg_settings
WHERE name IN (
'autovacuum_max_workers',
'autovacuum_analyze_scale_factor',
'autovacuum_naptime',
'autovacuum_analyze_threshold',
'autovacuum_analyze_scale_factor',
'autovacuum_vacuum_threshold',
'autovacuum_vacuum_scale_factor',
'autovacuum_vacuum_threshold',
'autovacuum_vacuum_cost_delay',
'autovacuum_vacuum_cost_limit',
'vacuum_cost_limit',
'autovacuum_freeze_max_age',
'maintenance_work_mem',
'vacuum_freeze_min_age');
```

While these all affect autovacuum, some of the most important ones are:

- [maintenance_mork_mem](#)
- [autovacuum_freeze_max_age](#)
- [autovacuum_max_workers](#)
- [autovacuum_vacuum_cost_delay](#)
- [Autovacuum_vacuum_cost_limit](#)

Setting Table-Level Autovacuum Parameters

Autovacuum-related [storage parameters](#) can be set at a table level, which can be better than altering the behavior of the entire database. For large tables, you might need to set aggressive settings and you might not want to make autovacuum behave that way for all tables.

The following query shows which tables currently have table-level options in place.

```
SELECT relname, reloptions
FROM pg_class
WHERE reloptions IS NOT null;
```

An example where this might be useful is on tables that are much larger than the rest of your tables. Suppose that you have one 300-GB table and 30 other tables less than 1 GB. In this case, you might set some specific parameters for your large table so you don't alter the behavior of your entire system.

```
ALTER TABLE mytable set (autovacuum_vacuum_cost_delay=0);
```

Doing this disables the cost-based autovacuum delay for this table at the expense of more resource usage on your system. Normally, autovacuum pauses for autovacuum_vacuum_cost_delay each time autovacuum_cost_limit is reached. You can find more details in the PostgreSQL documentation about [cost-based vacuuming](#).

Autovacuum Logging

By default, the *postgresql.log* doesn't contain information about the autovacuum process. If you are using PostgreSQL 9.4.5 or later, you can see output in the PostgreSQL error log from the autovacuum worker operations by setting the `rds.force_autovacuum_logging_level` parameter. Allowed values are `disabled`, `debug5`, `debug4`, `debug3`, `debug2`, `debug1`, `info`, `notice`, `warning`, `error`, `log`, `fatal`, and `panic`. The default value is `disabled` because the other allowable values can add significant amount of information to your logs.

We recommend that you set the value of the `rds.force_autovacuum_logging_level` parameter to `log` and that you set the `log_autovacuum_min_duration` parameter to a value from 1,000 to 5,000 milliseconds. If you set this value to 5,000, Amazon RDS writes any activity to the log that takes more than five seconds. It also shows "vacuum skipped" messages when application locking is causing autovacuum to intentionally skip tables. If you are troubleshooting a problem and need more detail, you can use a different logging level value, such as `debug1` or `debug3`. Use these debug parameters for a short period of time because these settings produce extremely verbose content written to the error log file. For more information about these debug settings, see the [PostgreSQL documentation](#).

Note

PostgreSQL version 9.4.7 and later includes improved visibility of autovacuum sessions by allowing the `rds_superuser` account to view autovacuum sessions in `pg_stat_activity`. For example, you can identify and terminate an autovacuum session that is blocking a command from running, or executing slower than a manually issued vacuum command.

Audit Logging for a PostgreSQL DB Instance

There are several parameters you can set to log activity that occurs on your PostgreSQL DB instance. These parameters include the following:

- The `log_statement` parameter can be used to log user activity in your PostgreSQL database. For more information, see [PostgreSQL Database Log Files \(p. 419\)](#).
- The `rds.force_admin_logging_level` parameter logs actions by the RDS internal user (`rdsadmin`) in the databases on the DB instance, and writes the output to the PostgreSQL error log. Allowed values

are disabled, debug5, debug4, debug3, debug2, debug1, info, notice, warning, error, log, fatal, and panic. The default value is disabled.

- The `rds.force_autovacuum_logging_level` parameter logs autovacuum worker operations in all databases on the DB instance, and writes the output to the PostgreSQL error log. Allowed values are disabled, debug5, debug4, debug3, debug2, debug1, info, notice, warning, error, log, fatal, and panic. The default value is disabled. The Amazon RDS recommended setting for `rds.force_autovacuum_logging_level` is LOG. Set `log_autovacuum_min_duration` to a value from 1000 or 5000. Setting this value to 5,000 writes activity to the log that takes more than 5 seconds and shows "vacuum skipped" messages. For more information on this parameter, see [Best Practices for Working with PostgreSQL \(p. 101\)](#).

Working with the pgaudit Extension

The `pgaudit` extension provides detailed session and object audit logging for Amazon RDS for PostgreSQL version 9.6.3 and later and version 9.5.7 version and later. You can enable session auditing or object auditing using this extension.

With session auditing, you can log audit events from various sources and includes the fully qualified command text when available. For example, you can use session auditing to log all READ statements that connect to a database by setting `pgaudit.log` to 'READ'.

With object auditing, you can refine the audit logging to work with specific commands. For example, you can specify that you want audit logging for READ operations on a specific number of tables.

To use object based logging with the `pgaudit` extension

1. Create a specific database role called `rds_pgaudit`. Use the following command to create the role.

```
CREATE ROLE rds_pgaudit;
CREATE ROLE
```

2. Modify the parameter group that is associated with your DB instance to use the shared preload libraries that contain `pgaudit` and set the parameter `pgaudit.role`. The `pgaudit.role` must be set to the role `rds_pgaudit`.

The following command modifies a custom parameter group.

```
aws rds modify-db-parameter-group
  --db-parameter-group-name rds-parameter-group-96
  --parameters
  "ParameterName=pgaudit.role,ParameterValue=rds_pgaudit,ApplyMethod=pending-reboot"
  --parameters
  "ParameterName=shared_preload_libraries,ParameterValue=pgaudit,ApplyMethod=pending-reboot"
  --region us-west-2
```

3. Reboot the instance so that the DB instance picks up the changes to the parameter group. The following command reboots a DB instance.

```
aws rds reboot-db-instance --db-instance-identifier rds-test-instance --region us-west-2
```

4. Run the following command to confirm that pgaudit has been initialized.

```
show shared_preload_libraries;
shared_preload_libraries
-----
rdsutils,pgaudit
(1 row)
```

5. Run the following command to create the pgaudit extension.

```
CREATE EXTENSION pgaudit;
CREATE EXTENSION
```

6. Run the following command to confirm pgaudit.role is set to *rds_pgaudit*.

```
show pgaudit.role;
pgaudit.role
-----
rds_pgaudit
```

To test the audit logging, run several commands that you have chosen to audit. For example, you might run the following commands.

```
CREATE TABLE t1 (id int);
CREATE TABLE
GRANT SELECT ON t1 TO rds_pgaudit;
GRANT
select * from t1;
id
-----
(0 rows)
```

The database logs should contain an entry similar to the following.

```
...
2017-06-12 19:09:49 UTC::rds_test@postgres:[11701]:LOG: AUDIT:
OBJECT,1,1,READ,SELECT,TABLE,public.t1,select * from t1;
...
```

For information on viewing the logs, see [Amazon RDS Database Log Files \(p. 390\)](#).

Working with the pg_repack Extension

You can use the pg_repack extension to remove bloat from tables and indexes. This extension is supported on Amazon RDS for PostgreSQL versions 9.6.3 and later. For more information on the pg_repack extension, see the [GitHub project documentation](#).

To use the pg_repack extension

1. Install the pg_repack extension on your Amazon RDS for PostgreSQL DB instance by running the following command.

```
CREATE EXTENSION pg_repack;
```

2. Use the pg_repack client utility to connect to a database. Use a database role that has *rds_superuser* privileges to connect to the database. In the following connection example, the *rds_test* role has *rds_superuser* privileges, and the database endpoint used is *rds-test-instance.cw7jjfgdr4on8.us-west-2.rds.amazonaws.com*.

```
pg_repack -h rds-test-instance.cw7jjfgdr4on8.us-west-2.rds.amazonaws.com -U rds_test -k postgres
```

Connect using the *-k* option. The *-a* option is not supported.

3. The response from the pg_repack client provides information on the tables on the DB instance that are repacked.

```
INFO: repacking table "pgbench_tellers"  
INFO: repacking table "pgbench_accounts"  
INFO: repacking table "pgbench_branches"
```

Working with PostGIS

PostGIS is an extension to PostgreSQL for storing and managing spatial information. If you are not familiar with PostGIS, you can get a good general overview at [PostGIS Introduction](#).

You need to perform a bit of setup before you can use the PostGIS extension. The following list shows what you need to do; each step is described in greater detail later in this section.

- Connect to the DB instance using the master user name used to create the DB instance.
- Load the PostGIS extensions.
- Transfer ownership of the extensions to the *rds_superuser* role.
- Transfer ownership of the objects to the *rds_superuser* role.
- Test the extensions.

Step 1: Connect to the DB Instance Using the Master User Name Used to Create the DB Instance

First, you connect to the DB instance using the master user name that was used to create the DB instance. That name is automatically assigned the *rds_superuser* role. You need the *rds_superuser* role that is needed to do the remaining steps.

The following example uses SELECT to show you the current user; in this case, the current user should be the master username you chose when creating the DB instance.

```
select current_user;
current_user
-----
myawsuser
(1 row)
```

Step 2: Load the PostGIS Extensions

Use the CREATE EXTENSION statements to load the PostGIS extensions. You must also load the extension. You can then use the \dn *psql* command to list the owners of the PostGIS schemas.

```
create extension postgis;
CREATE EXTENSION
create extension fuzzystrmatch;
CREATE EXTENSION
create extension postgis_tiger_geocoder;
CREATE EXTENSION
create extension postgis_topology;
CREATE EXTENSION
\dn
List of schemas
 Name      | Owner
-----+-----
 public    | myawsuser
 tiger     | rdsadmin
 tiger_data | rdsadmin
 topology   | rdsadmin
(4 rows)
```

Step 3: Transfer Ownership of the Extensions to the *rds_superuser* Role

Use the ALTER SCHEMA statements to transfer ownership of the schemas to the *rds_superuser* role.

```
alter schema tiger owner to rds_superuser;
ALTER SCHEMA
alter schema tiger_data owner to rds_superuser;
ALTER SCHEMA
alter schema topology owner to rds_superuser;
ALTER SCHEMA
\dn
List of schemas
 Name      | Owner
-----+-----
 public    | myawsuser
 tiger     | rds_superuser
 tiger_data | rds_superuser
 topology   | rds_superuser
(4 rows)
```

Step 4: Transfer Ownership of the Objects to the *rds_superuser* Role

Use the following function to transfer ownership of the PostGIS objects to the *rds_superuser* role. Run the following statement from the *psql* prompt to create the function.

```
CREATE FUNCTION exec(text) returns text language plpgsql volatile AS $$ BEGIN EXECUTE $1;
RETURN $1; END; $$;
```

Next, run this query to run the exec function that in turn executes the statements and alters the permissions.

```
SELECT exec('ALTER TABLE ' || quote_ident(s.nspname) || '.' || quote_ident(s.relname) || '
OWNER TO rds_superuser;');
FROM (
  SELECT nspname, relname
  FROM pg_class c JOIN pg_namespace n ON (c.relnamespace = n.oid)
  WHERE nspname in ('tiger','topology') AND
    relkind IN ('r','S','v') ORDER BY relkind = 'S')
s;
```

Step 5: Test the Extensions

Add `tiger` to your search path using the following command.

```
SET search_path=public,tiger;
```

Test `tiger` by using the following SELECT statement.

```
select na.address, na.streetname, na.streettypeabbrev, na.zip
from normalize_address('1 Devonshire Place, Boston, MA 02109') as na;
address | streetname | streettypeabbrev | zip
-----+-----+-----+-----+
1 | Devonshire | Pl | 02109
(1 row)
```

Test `topology` by using the following SELECT statement.

```
select topology.createtopology('my_new_topo',26986,0.5);
createtopology
-----
1
(1 row)
```

Using pgBadger for Log Analysis with PostgreSQL

You can use a log analyzer such as [pgbadger](#) to analyze PostgreSQL logs. The *pgbadger* documentation states that the `%l` pattern (log line for session/process) should be a part of the prefix. However, if you provide the current rds log_line_prefix as a parameter to *pgbadger* it should still produce a report.

For example, the following command correctly formats an Amazon RDS PostgreSQL log file dated 2014-02-04 using *pgbadger*.

```
./pgbadger -p '%t:%r:%u@%d:[%p]:' postgresql.log.2014-02-04-00
```

Viewing the Contents of pg_config

In PostgreSQL version 9.6.1, you can see the compile-time configuration parameters of the currently installed version of PostgreSQL using the new view pg_config. You can use the view by calling the pg_config function as shown in the following sample.

```
select * from pg_config();
   name      |           setting
-----
-----+
-----+
BINDIR          | /rdsdbbin/postgres-9.6.1.R1/bin
DOCDIR          | /rdsdbbin/postgres-9.6.1.R1/share/doc
HTMLDIR         | /rdsdbbin/postgres-9.6.1.R1/share/doc
INCLUDEDIR       | /rdsdbbin/postgres-9.6.1.R1/include
PKGINCLUDEDIR   | /rdsdbbin/postgres-9.6.1.R1/include
INCLUDEDIR-SERVER | /rdsdbbin/postgres-9.6.1.R1/include/server
LIBDIR          | /rdsdbbin/postgres-9.6.1.R1/lib
PKGLIBDIR       | /rdsdbbin/postgres-9.6.1.R1/lib
LOCALEDIR       | /rdsdbbin/postgres-9.6.1.R1/share/locale
MANDIR          | /rdsdbbin/postgres-9.6.1.R1/share/man
SHAREDIR         | /rdsdbbin/postgres-9.6.1.R1/share
SYSCONFDIR      | /rdsdbbin/postgres-9.6.1.R1/etc
PGXS            | /rdsdbbin/postgres-9.6.1.R1/lib/pgxs/src/makefiles/pgxs.mk
CONFIGURE        | '--prefix=/rdsdbbin/postgres-9.6.1.R1' '--with-openssl' '--with-perl'
                  '--with-tcl' '--with-uuid' '--with-libxml' '--with-libraries=/rdsdbbin
                  /postgres-9.6.1.R1/lib' '--with-includes=/rdsdbbin/postgres-9.6.1.R1/include' '--enable-
debug'
CC               | gcc
CPPFLAGS         | -D_GNU_SOURCE -I/usr/include/libxml2 -I/rdsdbbin/postgres-9.6.1.R1/
include
CFLAGS           | -Wall -Wmissing-prototypes -Wpointer-arith -Wdeclaration-after-
statement
                  -Wendif-labels -Wmissing-format-attribute -Wformat-security -fno-strict-
aliasing -fwrapv -fexcess-precision=standard -g -O2
CFLAGS_SL        | -fpic
LDFLAGS          | -L../../src/common -L/rdsdbbin/postgres-9.6.1.R1/lib -Wl,--as-needed -
Wl,
                  -rpath,'/rdsdbbin/postgres-9.6.1.R1/lib',--enable-new-dtags
LDFLAGS_EX       |
LDFLAGS_SL       |
LIBS             | -lpgcommon -lpgport -lxmll2 -lssl -lcrypto -lz -lreadline -lrt -lcrypt
                  -ldl -lm
VERSION          | PostgreSQL 9.6.1
(23 rows)
```

If you attempt to access the view directly, the request fails.

```
select * from pg_config;
ERROR: permission denied for relation pg_config
```

Working with the orafce Extension

The orafce extension provides functions that are common in commercial databases, and can make it easier for you to port a commercial database to PostgreSQL. Amazon RDS for PostgreSQL versions 9.6.6 and later support this extension. For more information about orafce, see the [orafce project on GitHub](#).

Note

Amazon RDS for PostgreSQL doesn't support the `utl_file` package that is part of the `orafce` extension. This is because the `utl_file` schema functions provide read and write operations on operating-system text files, which requires superuser access to the underlying host.

To use the orafce extension

1. Connect to the DB instance with the master user name that you used to create the DB instance.

Note

If you want to enable `orafce` on a different database in the same instance, use the `/c dbname psql` command to change from the master database after initiating the connection.

2. Enable the `orafce` extension with the `CREATE EXTENSION` statement.

```
CREATE EXTENSION orafce;
```

3. Transfer ownership of the oracle schema to the `rds_superuser` role with the `ALTER SCHEMA` statement.

```
ALTER SCHEMA oracle OWNER TO rds_superuser;
```

Note

If you want to see the list of owners for the oracle schema, use the `\dn` psql command.

Accessing External Data with the `postgres_fdw` Extension

You can access data in a table on a remote database server with the `postgres_fdw` extension. If you set up a remote connection from your PostgreSQL DB instance, access is also available to your Read Replica.

To use `postgres_fdw` to access a remote database server

1. Install the `postgres_fdw` extension.

```
CREATE EXTENSION postgres_fdw;
```

2. Create a foreign data server using `CREATE SERVER`.

```
CREATE SERVER foreign_server
FOREIGN DATA WRAPPER postgres_fdw
OPTIONS (host 'xxx.xx.xxx.xx', port '5432', dbname 'foreign_db');
```

3. Create a user mapping to identify the role to be used on the remote server.

```
CREATE USER MAPPING FOR local_user
SERVER foreign_server
OPTIONS (user 'foreign_user', password 'password');
```

4. Create a table that maps to the table on the remote server.

```
CREATE FOREIGN TABLE foreign_table (
    id integer NOT NULL,
    data text)
SERVER foreign_server
OPTIONS (schema_name 'some_schema', table_name 'some_table');
```

Using a Custom DNS Server for Outbound Network Access

Amazon RDS for PostgreSQL supports outbound network access on your DB instances and allows Domain Name Service (DNS) resolution from a custom DNS server owned by the customer. You can resolve only fully qualified domain names from your Amazon RDS DB instance through your custom DNS server.

Topics

- [Enabling Custom DNS Resolution \(p. 1348\)](#)
- [Disabling Custom DNS Resolution \(p. 1348\)](#)
- [Setting Up a Custom DNS Server \(p. 1348\)](#)

Enabling Custom DNS Resolution

To enable the DNS resolution in your customer VPC, you need to associate a custom DB parameter group to your RDS PostgreSQL instance, turn on the parameter **rds.custom_dns_resolution** by setting it to 1, and restart the DB instance for the changes to take place.

Disabling Custom DNS Resolution

In order to disable the DNS resolution in your customer VPC, you need to turn off the parameter **rds.custom_dns_resolution** of your custom DB parameter group by setting it to 0, then restart the DB instance for the changes to take place.

Setting Up a Custom DNS Server

After you set up your custom DNS name server, it takes up to 30 minutes to propagate the changes to your DB instance. After the changes are propagated to your DB instance, all outbound network traffic requiring a DNS lookup queries your DNS server over port 53.

To set up a custom DNS server for your Amazon RDS PostgreSQL DB instance, do the following:

1. From the DHCP options set attached to your VPC, set the **domain-name-servers** option to the IP address of your DNS name server. For more information, see [DHCP Options Sets](#).

Note

The **domain-name-servers** option accepts up to four values, but your Amazon RDS DB instance uses only the first value.

2. Ensure that your DNS server can resolve all lookup queries, including public DNS names, Amazon EC2 private DNS names, and customer-specific DNS names. If the outbound network traffic contains any DNS lookups that your DNS server can't handle, your DNS server must have appropriate upstream DNS providers configured.

3. Configure your DNS server to produce User Datagram Protocol (UDP) responses of 512 bytes or less.
4. Configure your DNS server to produce Transmission Control Protocol (TCP) responses of 1024 bytes or less.
5. Configure your DNS server to allow inbound traffic from your Amazon RDS DB instances over port 53. If your DNS server is in an Amazon VPC, the VPC must have a security group that contains inbound rules that allow UDP and TCP traffic on port 53. If your DNS server is not in an Amazon VPC, it must have appropriate firewall whitelisting to allow UDP and TCP inbound traffic on port 53.

For more information, see [Security Groups for Your VPC](#) and [Adding and Removing Rules](#).

6. Configure the VPC of your Amazon RDS DB instance to allow outbound traffic over port 53. Your VPC must have a security group that contains outbound rules that allow UDP and TCP traffic on port 53.

For more information, see [Security Groups for Your VPC](#) and [Adding and Removing Rules](#).

7. The routing path between the Amazon RDS DB instance and the DNS server has to be configured correctly to allow DNS traffic.

If the Amazon RDS DB instance and the DNS server are not in the same VPC, a peering connection has to be set up between them. For more information, see [What is VPC Peering?](#)

Restricting Password Management

You can restrict who can manage database user passwords to a special role. By doing this, you can have more control over password management on the client side.

You enable restricted password management with the static parameter `rds.restrict_password_commands` and use a role called `rds_password`. When the parameter `rds.restrict_password_commands` is set to 1, only users that are members of the `rds_password` role can run certain SQL commands. The restricted SQL commands are commands that modify database user passwords and password expiration time.

To use restricted password management, your DB instance must be running Amazon RDS for PostgreSQL 10.6 or higher. Because the `rds.restrict_password_commands` parameter is static, changing this parameter requires a database restart.

When a database has restricted password management enabled, if you try to run restricted SQL commands you get the following error: ERROR: must be a member of `rds_password` to alter passwords.

Following are some examples of SQL commands that are restricted when restricted password management is enabled.

```
postgres=> CREATE ROLE myrole WITH PASSWORD 'mypassword';
postgres=> CREATE ROLE myrole WITH PASSWORD 'mypassword' VALID UNTIL '2020-01-01';
postgres=> ALTER ROLE myrole WITH PASSWORD 'mypassword' VALID UNTIL '2020-01-01';
postgres=> ALTER ROLE myrole WITH PASSWORD 'mypassword';
postgres=> ALTER ROLE myrole VALID UNTIL '2020-01-01';
postgres=> ALTER ROLE myrole RENAME TO myrole2;
```

Some `ALTER ROLE` commands that include `RENAME TO` might also be restricted. They might be restricted because renaming a PostgreSQL role that has an MD5 password clears the password.

The `rds_superuser` role has membership for the `rds_password` role by default, and you can't change this. You can give other roles membership for the `rds_password` role by using the `GRANT` SQL command. We recommend that you give membership to `rds_password` to only a few roles that you use solely for password management. These roles require the `CREATEROLE` attribute to modify other roles.

Make sure that you verify password requirements such as expiration and needed complexity on the client side. We recommend that you restrict password-related changes by using your own client-side

utility. This utility should have a role that is a member of `rds_password` and has the `CREATEROLE` role attribute.

Using Kerberos Authentication with Amazon RDS for PostgreSQL

You can use Kerberos authentication to authenticate users when they connect to your Amazon RDS DB instance running PostgreSQL. In this case, your DB instance works with AWS Directory Service for Microsoft Active Directory, also called AWS Managed Microsoft AD, to enable Kerberos authentication. When users authenticate with a PostgreSQL DB instance joined to the trusting domain, authentication requests are forwarded to the directory that you create with AWS Directory Service.

Keeping all of your credentials in the same directory can save you time and effort. You have a centralized place for storing and managing credentials for multiple DB instances. Using a directory can also improve your overall security profile.

Amazon RDS supports Kerberos authentication for PostgreSQL DB instances in the following AWS Regions:

- US East (Ohio)
- US East (N. Virginia)
- US West (N. California)
- US West (Oregon)
- Asia Pacific (Mumbai)
- Asia Pacific (Seoul)
- Asia Pacific (Singapore)
- Asia Pacific (Sydney)
- Asia Pacific (Tokyo)
- Canada (Central)
- Europe (Frankfurt)
- Europe (Ireland)
- Europe (London)
- South America (São Paulo)

Topics

- [Overview of Kerberos Authentication for PostgreSQL DB Instances \(p. 1351\)](#)
- [Setting Up Kerberos Authentication for PostgreSQL DB Instances \(p. 1352\)](#)
- [Managing a DB Instance in a Domain \(p. 1360\)](#)
- [Connecting to PostgreSQL with Kerberos Authentication \(p. 1361\)](#)

Overview of Kerberos Authentication for PostgreSQL DB Instances

To set up Kerberos authentication for a PostgreSQL DB instance, take the following steps, described in more detail later:

1. Use AWS Managed Microsoft AD to create an AWS Managed Microsoft AD directory. You can use the AWS Management Console, the AWS CLI, or the AWS Directory Service API to create the directory.
2. Create an AWS Identity and Access Management (IAM) role that uses the managed IAM policy `AmazonRDSDirectoryServiceAccess`. The role allows Amazon RDS to make calls to your directory.

For the IAM role to allow access, the AWS Security Token Service (AWS STS) endpoint must be activated in the correct AWS Region for your AWS account. AWS STS endpoints are active by default in all AWS Regions, and you can use them without any further actions. For more information, see [Activating and Deactivating AWS STS in an AWS Region](#) in the *IAM User Guide*.

3. Create and configure users in the AWS Managed Microsoft AD directory using the Microsoft Active Directory tools. For more information about creating users in your Microsoft Active Directory, see [Manage Users and Groups in AWS Managed Microsoft AD](#) in the *AWS Directory Service Administration Guide*.
4. If you plan to locate the directory and the DB instance in different virtual private clouds (VPCs), configure VPC peering. For more information, see [What is VPC Peering?](#) in the *Amazon Virtual Private Cloud VPC Peering*.
5. Create or modify a PostgreSQL DB instance either from the console, CLI, or RDS API using one of the following methods:
 - [Creating a DB Instance Running the PostgreSQL Database Engine](#)
 - [Modifying a DB Instance Running the PostgreSQL Database Engine](#)
 - [Restoring from a DB Snapshot](#)
 - [Restoring a DB Instance to a Specified Time](#)

When you create or modify the DB instance, provide the domain identifier (d-* identifier) that was generated when you created your directory. Also provide the name of the IAM role that you created. You can locate the DB instance in the same VPC as the directory or in a different VPC.

6. Use the RDS master user credentials to connect to the PostgreSQL DB instance. Create the user in PostgreSQL to be identified externally. Externally identified users can log in to the PostgreSQL DB instance using Kerberos authentication.

To get Kerberos authentication using an on-premises or self-hosted Microsoft Active Directory, create a forest trust. For more information on setting up forest trusts using AWS Directory Service, see [When to Create a Trust Relationship](#) in the *AWS Directory Service Administration Guide*.

Setting Up Kerberos Authentication for PostgreSQL DB Instances

You use AWS Directory Service for Microsoft Active Directory (AWS Managed Microsoft AD) to set up Kerberos authentication for a PostgreSQL DB instance. To set up Kerberos authentication, take the following steps.

Topics

- [Step 1: Create a Directory Using the AWS Managed Microsoft AD \(p. 1352\)](#)
- [Step 2: Create an IAM Role for Use by Amazon RDS \(p. 1356\)](#)
- [Step 3: Create and Configure Users \(p. 1357\)](#)
- [Step 4: Configure VPC Peering \(p. 1357\)](#)
- [Step 5: Create or Modify a PostgreSQL DB Instance \(p. 1358\)](#)
- [Step 6: Create Kerberos Authentication PostgreSQL Logins \(p. 1359\)](#)
- [Step 7: Configure a PostgreSQL Client \(p. 1359\)](#)

Step 1: Create a Directory Using the AWS Managed Microsoft AD

AWS Directory Service creates a fully managed Microsoft Active Directory in the AWS Cloud. When you create an AWS Managed Microsoft AD directory, AWS Directory Service creates two domain controllers

and DNS servers for you. The directory servers are created in different subnets in a VPC. This redundancy helps make sure that your directory remains accessible even if a failure occurs.

When you create an AWS Managed Microsoft AD directory, AWS Directory Service performs the following tasks on your behalf:

- Sets up a Microsoft Active Directory within your VPC.
- Creates a directory administrator account with the user name `Admin` and the specified password. You use this account to manage your directory.

Important

Make sure to save this password. AWS Directory Service doesn't store this password, and it can't be retrieved or reset.

- Creates a security group for the directory controllers.

When you launch AWS Directory Service for Microsoft Active Directory, AWS creates an Organizational Unit (OU) that contains all of your directory's objects. This OU, which has the NetBIOS name that you entered when you created your directory, is located in the domain root. The domain root is owned and managed by AWS.

The `Admin` account that was created with your AWS Managed Microsoft AD directory has permissions for the most common administrative activities for your OU:

- Create, update, or delete users
- Add resources to your domain such as file or print servers, and then assign permissions for those resources to users in your OU
- Create additional OUs and containers
- Delegate authority
- Restore deleted objects from the Active Directory Recycle Bin
- Run Active Directory and Domain Name Service (DNS) modules for Windows PowerShell on the Active Directory Web Service

The `Admin` account also has rights to perform the following domain-wide activities:

- Manage DNS configurations (add, remove, or update records, zones, and forwarders)
- View DNS event logs
- View security event logs

To create a directory with AWS Managed Microsoft AD

1. In the [AWS Directory Service console](#) navigation pane, choose **Directories**, and then choose **Set up directory**.
2. Choose **AWS Managed Microsoft AD**. AWS Managed Microsoft AD is the only option currently supported for use with Amazon RDS.
3. Choose **Next**.
4. On the **Enter directory information** page, provide the following information:

Edition

Choose the edition that meets your requirements.

Directory DNS name

The fully qualified name for the directory, such as `corp.example.com`.

Directory NetBIOS name

An optional short name for the directory, such as CORP.

Directory description

An optional description for the directory.

Admin password

The password for the directory administrator. The directory creation process creates an administrator account with the user name Admin and this password.

The directory administrator password can't include the word "admin." The password is case-sensitive and must be 8–64 characters in length. It must also contain at least one character from three of the following four categories:

- Lowercase letters (a–z)
- Uppercase letters (A–Z)
- Numbers (0–9)
- Nonalphanumeric characters (~!@#\$%^&*_+=`|\{}{};:"'<>,./?)

Confirm password

Retype the administrator password.

Important

Make sure that you save this password. AWS Directory Service doesn't store this password, and it can't be retrieved or reset.

5. Choose **Next**.
6. On the **Choose VPC and subnets** page, provide the following information:

VPC

Choose the VPC for the directory. You can create the PostgreSQL DB instance in this same VPC or in a different VPC.

Subnets

Choose the subnets for the directory servers. The two subnets must be in different Availability Zones.

7. Choose **Next**.
8. Review the directory information. If changes are needed, choose **Previous** and make the changes. When the information is correct, choose **Create directory**.

Review & create

Review

Directory type	VPC
Microsoft AD	vpc-8b6b78e9 (172.31.0.0/16)
Directory DNS name	Subnets
corp.example.com	subnet-75128d10 (172.31.64.0/20, us-east-1a) subnet-f51665dd (172.31.32.0/20, us-east-1b)
Directory NetBIOS name	
CORP	
Directory description	
My directory	



Pricing

Edition	Free trial eligible Learn more 30-day limited trial
Standard	
~USD 86.4000/mo (USD 0.1200/hr)*	
* Includes two domain controllers, USD 43.2000/mo for each additional domain controller.	

[Cancel](#) [Previous](#) [Create directory](#)

It takes several minutes for the directory to be created. When it has been successfully created, the **Status** value changes to **Active**.

To see information about your directory, choose the directory ID in the directory listing. Make a note of the **Directory ID** value. You need this value when you create or modify your PostgreSQL DB instance.

Directory details		Reset user password	G
Directory type	VPC	Status	
Microsoft AD	vpc-8b6b78e9 Edit	Creating	
Edition	Subnets	Last updated	
Standard	subnet-f51665dd Edit subnet-75128d10 Edit	Friday, July 26, 2019	
Directory ID	Availability zones	Launch time	
d-9067004760	us-east-1a, us-east-1b	Friday, July 26, 2019	
Directory DNS name	DNS address		
corp.example.com			
Directory NetBIOS name			
CORP			
Description - Edit			
My directory			

Step 2: Create an IAM Role for Use by Amazon RDS

For RDS to call AWS Directory Service for you, you must create an IAM role that uses the managed IAM policy `AmazonRDSDirectoryServiceAccess`. This role allows Amazon RDS to make calls to the AWS Directory Service. When you create this IAM role, choose `Directory Service`, and attach the AWS managed policy `AmazonRDSDirectoryServiceAccess` to it.

For more information about creating IAM roles for a service, see [Creating a Role to Delegate Permissions to an AWS Service](#).

Optionally, you can create policies with the required permissions instead of using the managed IAM policy `AmazonRDSDirectoryServiceAccess`. In this case, the IAM role must have the following IAM trust policy.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Service": [
          "rds.amazonaws.com",
          "directoryservice.rds.amazonaws.com"
        ]
      },
      "Action": "sts:AssumeRole"
    }
  ]
}
```

```
}
```

The role must also have the following IAM role policy.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Action": [
        "ds:DescribeDirectories",
        "ds:AuthorizeApplication",
        "ds:UnauthorizeApplication",
        "ds:GetAuthorizedApplicationDetails"
      ],
      "Effect": "Allow",
      "Resource": "*"
    }
  ]
}
```

Step 3: Create and Configure Users

You can create users by using the Active Directory Users and Computers tool, which is one of the Active Directory Domain Services and Active Directory Lightweight Directory Services tools. In this case, *users* are individual people or entities that have access to your directory.

To create users in an AWS Directory Service directory, you must be connected to a Windows-based Amazon EC2 instance. Also, this EC2 instance must be a member of the AWS Directory Service directory. At the same time, you must be logged in as a user that has privileges to create users. For more information, see [Create a User](#) in the *AWS Directory Service Administration Guide*.

Step 4: Configure VPC Peering

If you plan to locate the directory and the DB instance in the same VPC, skip this step and move on to [Step 5: Create or Modify a PostgreSQL DB Instance \(p. 1358\)](#). If you plan to locate the directory and the DB instance in different VPCs, configure VPC peering by following the instructions in this step.

If the same AWS account owns both VPCs, follow the instructions in [What is VPC Peering?](#) in the *Amazon VPC VPC Peering* guide. Specifically, complete the following steps:

1. Set up appropriate VPC routing rules to make sure that the network traffic can flow both ways.
2. Ensure the DB instance's security group can receive ingress traffic from this security group.
3. Ensure that there is no network access control list (ACL) rule to block traffic.

If different AWS accounts own the VPCs, complete the following steps:

1. Configure VPC peering by following the instructions in [What is VPC Peering?](#) Specifically, complete the following steps:
 - a. Set up appropriate VPC routing rules to make sure that the network traffic can flow both ways.
 - b. Make sure that there is no ACL rule to block traffic.
2. Start sharing the directory with the AWS account where you plan to create the DB instance. To do this, follow the instructions in [Tutorial: Sharing Your AWS Managed Microsoft AD Directory for Seamless EC2 Domain-Join](#) in the *AWS Directory Service Administration Guide*.

3. Sign in to the AWS Directory Service console using the account for the DB instance, and make sure that the domain has the **SHARED** status before proceeding.
4. While logged into the AWS Directory Service console using the account for the DB instance, make a note of the **Directory ID** value for the directory.

Step 5: Create or Modify a PostgreSQL DB Instance

Next, create or modify a PostgreSQL DB instance for use with your directory. You can use the console, CLI, or RDS API to associate a DB instance with a directory. You can do this in one of the following ways:

- Create a new PostgreSQL DB instance using the console, the `create-db-instance` CLI command, or the [CreateDBInstance](#) RDS API operation.

For instructions, see [Creating a DB Instance Running the PostgreSQL Database Engine](#).

- Modify an existing PostgreSQL DB instance using the console, the `modify-db-instance` CLI command, or the [ModifyDBInstance](#) RDS API operation.

For instructions, see [Modifying a DB Instance Running the PostgreSQL Database Engine](#).

- Restore a PostgreSQL DB instance from a DB snapshot using the console, the `restore-db-instance-from-db-snapshot` CLI command, or the [RestoreDBInstanceFromDBSnapshot](#) RDS API operation.

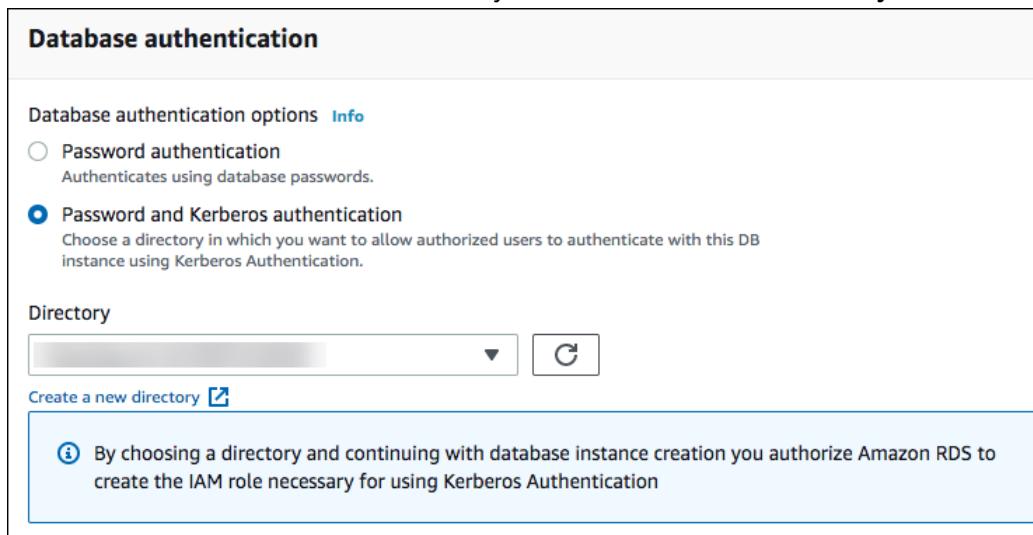
For instructions, see [Restoring from a DB Snapshot](#).

- Restore a PostgreSQL DB instance to a point-in-time using the console, the `restore-db-instance-to-point-in-time` CLI command, or the [RestoreDBInstanceToPointInTime](#) RDS API operation.

For instructions, see [Restoring a DB Instance to a Specified Time](#).

Kerberos authentication is only supported for PostgreSQL DB instances in a VPC. The DB instance can be in the same VPC as the directory, or in a different VPC. The DB instance must use a security group that allows egress within the directory's VPC so the DB instance can communicate with the directory.

When you use the console, choose **Password and Kerberos authentication** in the **Database authentication** section. Choose the directory or choose **Create a new Directory**.



When you use the AWS CLI, the following parameters are required for the DB instance to be able to use the directory that you created:

- For the --domain parameter, use the domain identifier ("d-*" identifier) generated when you created the directory.
- For the --domain-iam-role-name parameter, use the role you created that uses the managed IAM policy [AmazonRDSdirectoryServiceAccess](#).

For example, the following CLI command modifies a DB instance to use a directory.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
    --db-instance-identifier mydbinstance \
    --domain d-Directory-ID \
    --domain-iam-role-name role-name
```

For Windows:

```
aws rds modify-db-instance ^
    --db-instance-identifier mydbinstance ^
    --domain d-Directory-ID ^
    --domain-iam-role-name role-name
```

Important

If you modify a DB instance to enable Kerberos authentication, reboot the DB instance after making the change.

Step 6: Create Kerberos Authentication PostgreSQL Logins

Next, use the RDS master user credentials to connect to the PostgreSQL DB instance as you do with any other DB instance. The DB instance is joined to the AWS Managed Microsoft AD domain. Thus, you can provision PostgreSQL logins and users from the Microsoft Active Directory users and groups in your domain. To manage database permissions, you grant and revoke standard PostgreSQL permissions to these logins.

To allow a Microsoft Active Directory user to authenticate with PostgreSQL, use the RDS master user credentials. You use these credentials to connect to the PostgreSQL DB instance as you do with any other DB instance. After you're logged in, create an externally authenticated user in PostgreSQL and grant the `rds_ad` role to this user.

```
CREATE USER "username@CORP.EXAMPLE.COM" WITH LOGIN;
GRANT rds_ad TO "username@CORP.EXAMPLE.COM";
```

Replace `username` with the user name and include the domain name in uppercase. Users (both humans and applications) from your domain can now connect to the RDS PostgreSQL instance from a domain-joined client machine using Kerberos authentication.

Step 7: Configure a PostgreSQL Client

To configure a PostgreSQL client, meet the following requirements:

- Create a `krb5.conf` file (or equivalent) to point to the domain.
- Verify that traffic can flow between the client host and AWS Directory Service. Use a network utility such as Netcat for the following:

- Verify traffic over DNS for port 53.
- Verify traffic over TCP/UDP for port 53 and for Kerberos, which includes ports 88 and 464 for AWS Directory Service.
- Verify that traffic can flow between the client host and the DB instance over the database port. For example, use `psql` to connect and access the database.

Managing a DB Instance in a Domain

You can use the console, the CLI, or the RDS API to manage your DB instance and its relationship with your Microsoft Active Directory. For example, you can associate a Microsoft Active Directory to enable Kerberos authentication. You can also remove the association for a Microsoft Active Directory to disable Kerberos authentication. You can also move a DB instance to be externally authenticated by one Microsoft Active Directory to another.

For example, using the CLI, you can do the following:

- To reattempt enabling Kerberos authentication for a failed membership, use the [modify-db-instance](#) CLI command. Specify the current membership's directory ID for the `--domain` option.
- To disable Kerberos authentication on a DB instance, use the [modify-db-instance](#) CLI command. Specify `none` for the `--domain` option.
- To move a DB instance from one domain to another, use the [modify-db-instance](#) CLI command. Specify the domain identifier of the new domain for the `--domain` option.

Understanding Domain Membership

After you create or modify your DB instance, the DB instance becomes a member of the domain. You can view the status of the domain membership for the DB instance in the console or by running the [describe-db-instances](#) CLI command. The status of the DB instance can be one of the following:

- `kerberos-enabled` – The DB instance has Kerberos authentication enabled.
- `enabling-kerberos` – AWS is in the process of enabling Kerberos authentication on this DB instance.
- `pending-enable-kerberos` – Enabling Kerberos authentication is pending on this DB instance.
- `pending-maintenance-enable-kerberos` – AWS will attempt to enable Kerberos authentication on the DB instance during the next scheduled maintenance window.
- `pending-disable-kerberos` – Disabling Kerberos authentication is pending on this DB instance.
- `pending-maintenance-disable-kerberos` – AWS will attempt to disable Kerberos authentication on the DB instance during the next scheduled maintenance window.
- `enable-kerberos-failed` – A configuration problem prevented AWS from enabling Kerberos authentication on the DB instance. Correct the configuration problem before reissuing the command to modify the DB instance.
- `disabling-kerberos` – AWS is in the process of disabling Kerberos authentication on this DB instance.

A request to enable Kerberos authentication can fail because of a network connectivity issue or an incorrect IAM role. In some cases, the attempt to enable Kerberos authentication might fail when you create or modify a DB instance. If so, make sure that you are using the correct IAM role, then modify the DB instance to join the domain.

Note

Only Kerberos authentication with RDS for PostgreSQL sends traffic to the domain's DNS servers. All other DNS requests are treated as outbound network access on your DB instances

running PostgreSQL. For more information about outbound network access with Amazon RDS for PostgreSQL, see [Using a Custom DNS Server for Outbound Network Access](#).

Connecting to PostgreSQL with Kerberos Authentication

You can connect to PostgreSQL with Kerberos authentication with the pgAdmin interface or with a command line interface such as psql. For more information about connecting, see [Connecting to a DB Instance Running the PostgreSQL Database Engine](#).

pgAdmin

To use pgAdmin to connect to PostgreSQL with Kerberos authentication, take the following steps:

1. Launch the pgAdmin application on your client computer.
2. On the **Dashboard** tab, choose **Add New Server**.
3. In the **Create - Server** dialog box, type a name on the **General** tab to identify the server in pgAdmin.
4. On the **Connection** tab, type the following information from your DB instance:
 - For **Host**, type the endpoint. For example, use the format *PostgreSQL-instance-endpoint.AWS-Region.rds.amazonaws.com*.
 - For **Port**, type the assigned port.
 - For **Maintenance database**, type the name of the initial database to which the client will connect.
 - For **Username**, type the user name that you entered when you created the DB instance.
 - For **Password**, type the password that you entered when you created the DB instance.
5. Choose **Save**.

psql

To use psql to connect to PostgreSQL with Kerberos authentication, take the following steps:

1. At a command prompt, run the following command.

```
kinit username
```

Replace *username* with the user name. At the prompt, enter the password stored in the Microsoft Active Directory for the user.

2. If the PostgreSQL DB instance is using a publicly accessible VPC, put a private IP address for your DB instance endpoint in your /etc/hosts file on the EC2 client. For example, the following commands obtain the private IP address and then put it in the /etc/hosts file.

```
% dig +short PostgreSQL-instance-endpoint.AWS-Region.rds.amazonaws.com
;; Truncated, retrying in TCP mode.
ec2-34-210-197-118.AWS-Region.compute.amazonaws.com.
34.210.197.118

% echo " 34.210.197.118 PostgreSQL-instance-endpoint.AWS-Region.rds.amazonaws.com" >> /etc/hosts
```

3. Use the following psql command to log in to a PostgreSQL DB instance that is integrated with Microsoft Active Directory.

```
psql postgres -U username@CORP.EXAMPLE.COM -p 5432 postgres -h PostgreSQL-instance-endpoint.AWS-Region.rds.amazonaws.com
```

Working with the Database Preview Environment

When you create a DB instance in Amazon RDS, you know that the PostgreSQL version it's based on has been tested and is fully supported by Amazon. The PostgreSQL community releases new versions and new extensions continuously. You can try out new PostgreSQL versions and extensions before they are fully supported. To do that, you can create a new DB instance in the Database Preview Environment.

DB instances in the Database Preview Environment are similar to DB instances in a production environment. However, keep in mind several important factors:

- All DB instances are deleted 60 days after you create them, along with any backups and snapshots.
- You can only create a DB instance in a virtual private cloud (VPC) based on the Amazon VPC service.
- You can only create M4, T2, and R4 instance types. For more information about RDS instance classes, see [Choosing the DB Instance Class \(p. 107\)](#).
- You can't get help from AWS Support with DB instances. You can post your questions in the [RDS Database Preview Environment Forum](#).
- You can only use General Purpose SSD and Provisioned IOPS SSD storage.
- You can't copy a snapshot of a DB instance to a production environment.
- Some Amazon RDS features aren't available in the preview environment, as described following.

Topics

- [Features Not Supported in the Preview Environment \(p. 1362\)](#)
- [PostgreSQL Extensions Supported in the Preview Environment \(p. 1362\)](#)
- [Creating a New DB Instance in the Preview Environment \(p. 1364\)](#)

Features Not Supported in the Preview Environment

The following features are not available in the preview environment:

- Cross-region snapshot copy
- Cross-region Read Replicas
- Extensions not in the following table of supported extensions

PostgreSQL Extensions Supported in the Preview Environment

The PostgreSQL extensions supported in the Database Preview Environment are listed in the following table.

Extension	Version
amcheck	1.1
bloom	1.0
btree_gin	1.3
btree_gist	1.5
citext	1.5

Extension	Version
cube	1.4
dblink	1.2
dict_int	1.0
dict_xsyn	1.0
earthdistance	1.1
fuzzystrmatch	1.1
hstore	1.5
hstore_plper	1.0
intagg	1.1
antarray	1.2
isn	1.2
log_fdw	1.0
ltree	1.1
pg_buffercache	1.3
pg_freespacemap	1.2
pg_prewarm	1.2
pg_stat_statements	1.5
pg_trgm	1.4
pg_visibility	1.2
pgcrypto	1.3
pgrowlocks	1.2
pgstattuple	1.5
plperl	1.0
plpgsql	1.0
pltcl	1.0
postgres_fdw	1.0
sslinfo	1.2
tablefunc	1.0
test_parser	1.0
tsm_system_rows	1.0
tsm_system_time	1.0

Extension	Version
unaccent	1.1
uuid-ossp	1.1

Creating a New DB Instance in the Preview Environment

Use the following procedure to create a DB instance in the preview environment.

To create a DB instance in the preview environment

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. Choose **Dashboard** from the navigation pane.
3. Choose **Switch to database preview environment**.

Database Preview Environment

Get early access to new DB engine versions, before they're generally available. The RDS database preview environment lets you work with upcoming beta, release candidate, and early production versions of PostgreSQL engines. Preview environment instances are fully functional, so you can easily test new features and functionality with your applications.

[Info](#)

[Preview PostgreSQL in US EAST \(Ohio\)](#)

You also can navigate directly to the [Database Preview Environment](#).

Note

If you want to create an instance in the Database Preview Environment with the API or CLI the endpoint is `rds-preview.us-east-2.amazonaws.com`.

4. Continue with the procedure as described in [New Console \(p. 1257\)](#).

Amazon RDS for PostgreSQL Versions and Extensions

Amazon RDS supports DB instances running several editions of PostgreSQL. Use this section to see how to work with PostgreSQL on Amazon RDS. You should also be aware of the limits for PostgreSQL DB instances.

You can specify any currently supported PostgreSQL version when creating a new DB instance. You can specify the major version (such as PostgreSQL 10), and any supported minor version for the specified major version. If no version is specified, Amazon RDS defaults to a supported version, typically the most recent version. If a major version is specified but a minor version is not, Amazon RDS defaults to a recent release of the major version you have specified. To see a list of supported versions, as well as defaults for newly created DB instances, use the [describe-db-engine-versions](#) AWS CLI command.

For information about importing PostgreSQL data into a DB instance, see [Importing Data into PostgreSQL on Amazon RDS \(p. 1302\)](#).

Topics

- [Supported PostgreSQL Database Versions \(p. 1365\)](#)
- [Supported PostgreSQL Features and Extensions \(p. 1383\)](#)

Supported PostgreSQL Database Versions

Amazon RDS supports the following PostgreSQL versions.

Topics

- [PostgreSQL Version 12.0 in the Database Preview Environment \(p. 1366\)](#)
- [PostgreSQL Version 12 Beta 3 on Amazon RDS in the Database Preview Environment \(p. 1367\)](#)
- [PostgreSQL Version 12 Beta 2 on Amazon RDS in the Database Preview Environment \(p. 1367\)](#)
- [PostgreSQL Version 11.5 on Amazon RDS \(p. 1367\)](#)
- [PostgreSQL Version 11.4 on Amazon RDS \(p. 1367\)](#)
- [PostgreSQL Version 11.2 on Amazon RDS \(p. 1368\)](#)
- [PostgreSQL Version 11.1 on Amazon RDS \(p. 1368\)](#)
- [PostgreSQL Version 11 on Amazon RDS in the Database Preview Environment \(p. 1369\)](#)
- [PostgreSQL Version 10.10 on Amazon RDS \(p. 1369\)](#)
- [PostgreSQL Version 10.9 on Amazon RDS \(p. 1369\)](#)
- [PostgreSQL Version 10.7 on Amazon RDS \(p. 1369\)](#)
- [PostgreSQL Version 10.6 on Amazon RDS \(p. 1370\)](#)
- [PostgreSQL Version 10.5 on Amazon RDS \(p. 1370\)](#)
- [PostgreSQL Version 10.4 on Amazon RDS \(p. 1371\)](#)
- [PostgreSQL Version 10.3 on Amazon RDS \(p. 1371\)](#)
- [PostgreSQL Version 10.1 on Amazon RDS \(p. 1372\)](#)
- [PostgreSQL Version 9.6.15 on Amazon RDS \(p. 1372\)](#)
- [PostgreSQL Version 9.6.14 on Amazon RDS \(p. 1372\)](#)
- [PostgreSQL Version 9.6.12 on Amazon RDS \(p. 1372\)](#)
- [PostgreSQL Version 9.6.11 on Amazon RDS \(p. 1373\)](#)
- [PostgreSQL Version 9.6.10 on Amazon RDS \(p. 1373\)](#)
- [PostgreSQL Version 9.6.9 on Amazon RDS \(p. 1373\)](#)
- [PostgreSQL Version 9.6.8 on Amazon RDS \(p. 1374\)](#)
- [PostgreSQL Version 9.6.6 on Amazon RDS \(p. 1374\)](#)
- [PostgreSQL Version 9.6.5 on Amazon RDS \(p. 1374\)](#)
- [PostgreSQL Version 9.6.3 on Amazon RDS \(p. 1374\)](#)
- [PostgreSQL Version 9.6.2 on Amazon RDS \(p. 1375\)](#)
- [PostgreSQL Version 9.6.1 on Amazon RDS \(p. 1375\)](#)
- [PostgreSQL Version 9.5.19 on Amazon RDS \(p. 1376\)](#)

- [PostgreSQL Version 9.5.18 on Amazon RDS \(p. 1376\)](#)
- [PostgreSQL Version 9.5.16 on Amazon RDS \(p. 1376\)](#)
- [PostgreSQL Version 9.5.15 on Amazon RDS \(p. 1376\)](#)
- [PostgreSQL Version 9.5.14 on Amazon RDS \(p. 1376\)](#)
- [PostgreSQL Version 9.5.13 on Amazon RDS \(p. 1377\)](#)
- [PostgreSQL Version 9.5.12 on Amazon RDS \(p. 1377\)](#)
- [PostgreSQL Version 9.5.10 on Amazon RDS \(p. 1377\)](#)
- [PostgreSQL Version 9.5.9 on Amazon RDS \(p. 1377\)](#)
- [PostgreSQL Version 9.5.7 on Amazon RDS \(p. 1377\)](#)
- [PostgreSQL Version 9.5.6 on Amazon RDS \(p. 1378\)](#)
- [PostgreSQL Version 9.5.4 on Amazon RDS \(p. 1378\)](#)
- [PostgreSQL Version 9.5.2 on Amazon RDS \(p. 1378\)](#)
- [PostgreSQL Version 9.4.24 on Amazon RDS \(p. 1379\)](#)
- [PostgreSQL Version 9.4.23 on Amazon RDS \(p. 1379\)](#)
- [PostgreSQL Version 9.4.21 on Amazon RDS \(p. 1379\)](#)
- [PostgreSQL Version 9.4.20 on Amazon RDS \(p. 1380\)](#)
- [PostgreSQL Version 9.4.19 on Amazon RDS \(p. 1380\)](#)
- [PostgreSQL Version 9.4.18 on Amazon RDS \(p. 1380\)](#)
- [PostgreSQL Version 9.4.17 on Amazon RDS \(p. 1380\)](#)
- [PostgreSQL Version 9.4.15 on Amazon RDS \(p. 1380\)](#)
- [PostgreSQL Version 9.4.14 on Amazon RDS \(p. 1380\)](#)
- [PostgreSQL Version 9.4.12 on Amazon RDS \(p. 1380\)](#)
- [PostgreSQL Version 9.4.11 on Amazon RDS \(p. 1381\)](#)
- [PostgreSQL Version 9.4.9 on Amazon RDS \(p. 1381\)](#)
- [PostgreSQL Version 9.4.7 on Amazon RDS \(p. 1381\)](#)
- [PostgreSQL Version 9.3.25 on Amazon RDS \(p. 1381\)](#)
- [PostgreSQL Version 9.3.24 on Amazon RDS \(p. 1382\)](#)
- [PostgreSQL Version 9.3.23 on Amazon RDS \(p. 1382\)](#)
- [PostgreSQL Version 9.3.22 on Amazon RDS \(p. 1382\)](#)
- [PostgreSQL Version 9.3.20 on Amazon RDS \(p. 1382\)](#)
- [PostgreSQL Version 9.3.19 on Amazon RDS \(p. 1382\)](#)
- [PostgreSQL Version 9.3.17 on Amazon RDS \(p. 1382\)](#)
- [PostgreSQL Version 9.3.16 on Amazon RDS \(p. 1383\)](#)
- [PostgreSQL Version 9.3.14 on Amazon RDS \(p. 1383\)](#)
- [PostgreSQL Version 9.3.12 on Amazon RDS \(p. 1383\)](#)

PostgreSQL Version 12.0 in the Database Preview Environment

PostgreSQL version 12.0 is now available in the Amazon RDS Database Preview Environment. PostgreSQL version 12.0 contains several improvements that are described in the following PostgreSQL documentation:

- [PostgreSQL 12 Released!](#)
- [PostgreSQL 12 Beta 3 Released!](#)
- [PostgreSQL 12 Beta 2 Released!](#)
- [PostgreSQL 12 Beta 1 Released!](#)

For information on the Database Preview Environment, see [the section called “Working with the Database Preview Environment” \(p. 1362\)](#). To access the Preview Environment from the console, select <https://console.aws.amazon.com/rds-preview/>.

PostgreSQL version 12.0 supports the `aws_s3` extension and you can now set `pg_prewarm` as `shared_preload_values`. For information on extensions supported in this version, see [PostgreSQL Version 12.x Extensions and Modules Supported on Amazon RDS in the Database Preview Environment \(p. 1384\)](#).

PostgreSQL Version 12 Beta 3 on Amazon RDS in the Database Preview Environment

PostgreSQL version 12 Beta 3 contains several improvements that are described in [PostgreSQL 12 Beta 3 Released!](#)

For information on the Database Preview Environment, see [the section called “Working with the Database Preview Environment” \(p. 1362\)](#). To access the Preview Environment from the console, select <https://console.aws.amazon.com/rds-preview/>. For information on extensions supported in this version, see [PostgreSQL Version 12.x Extensions and Modules Supported on Amazon RDS in the Database Preview Environment \(p. 1384\)](#).

PostgreSQL Version 12 Beta 2 on Amazon RDS in the Database Preview Environment

PostgreSQL version 12 Beta 2 contains several improvements that are described in the following links:

- [PostgreSQL 12 Beta 1 Released!](#)
- [PostgreSQL 12 Beta 2 Released!](#)

For information on the Database Preview Environment, see [the section called “Working with the Database Preview Environment” \(p. 1362\)](#). To access the Preview Environment from the console, select <https://console.aws.amazon.com/rds-preview/>.

PostgreSQL Version 11.5 on Amazon RDS

PostgreSQL version 11.5 contains several bug fixes for issues in release 11.4. For more information on the fixes in PostgreSQL 11.5, see the [PostgreSQL documentation](#).

This version also includes the following changes:

- A new extension `pg_transport` is added.
- The extension `aws_s3` has been updated to support virtual-hosted style requests. For more information, see [Amazon S3 Path Deprecation Plan – The Rest of the Story](#).
- The `PostGIS` extension is updated to version 2.5.2.

PostgreSQL Version 11.4 on Amazon RDS

This release contains an important security fix and also bug fixes and improvements done by the PostgreSQL community. For more information on the security fix, see the [PostgreSQL community announcement](#) and the security fix CVE-2019-10164.

With this release, the `pg_hint_plan` extension has been updated to version 1.3.4.

For more information on the fixes in PostgreSQL 11.4, see the [PostgreSQL documentation](#).

PostgreSQL Version 11.2 on Amazon RDS

PostgreSQL version 11.2 contains several bug fixes for issues in release 11.1. For more information on the fixes in PostgreSQL 11.2, see the [PostgreSQL documentation](#).

This version also includes the following changes:

- A new [pgTAP](#) extension version 1.0.
- Support for Amazon S3 import. For more information, see [Importing Amazon S3 Data into an RDS for PostgreSQL DB Instance \(p. 1305\)](#).
- Multiple major version upgrade is available to PostgreSQL 11.2 from certain previous PostgreSQL versions. For more information, see [Major Version Upgrades for PostgreSQL \(p. 1290\)](#).

For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 11.1 on Amazon RDS

PostgreSQL version 11.1 contains several improvements that were announced in [PostgreSQL 11.1 Released!](#) This version includes SQL stored procedures that enable embedded transactions within a procedure. This version also includes major improvements to partitioning and parallelism and many useful performance improvements. For example, by using a non-null constant for a column default, you can now use an ALTER TABLE command to add a column without causing a table rewrite.

PostgreSQL version 11.1 contains several bug fixes for issues in release 11. For complete details, see the [PostgreSQL Release 11.1 documentation](#). Some changes in this version include the following:

- Partitioning – Partitioning improvements include support for hash partitioning, enabling creation of a default partition, and dynamic row movement to another partition based on the key column update.
- Performance – Performance improvements include parallelism while creating indexes, materialized views, hash joins, and sequential scans to make the operations perform better.
- Stored procedures – SQL stored procedures now added support embedded transactions.
- Support for Just-In-Time (JIT) capability – RDS PostgreSQL 11 instances are created with JIT capability, speeding evaluation of expressions. To enable this feature, set jit to ON.
- Segment size – The write-ahead logging (WAL) segment size has been changed from 16 MB to 64 MB.
- Autovacuum improvements – To provide valuable logging, the parameter `rds.force_autovacuum_logging` is ON by default in conjunction with the `log_autovacuum_min_duration` parameter set to 10 seconds. To increase autovacuum effectiveness, the values for the `autovacuum_max_workers` and `autovacuum_vacuum_cost_limit` parameters are computed based on host memory capacity to provide larger default values.
- Improved transaction timeout – The parameter `idle_in_transaction_session_timeout` is set to 12 hours. Any session that has been idle more than 12 hours is terminated.
- Performance metrics – The `pg_stat_statements` module is included in `shared_preload_libraries` by default. This avoids having to reboot the instance immediately after creation. However, this functionality still requires you to run the statement `CREATE EXTENSION pg_stat_statements;`. Also, `track_io_timing` is enabled by default to add more granular data to `pg_stat_statements`.
- The `tsearch2` module is no longer supported – If your application uses `tsearch2` functions, update it to use the equivalent functions provided by the core PostgreSQL engine. For more information about the `tsearch2` module, see [PostgreSQL tsearch2](#).

- The `chkpass` module is no longer supported – For more information about the `chkpass` module, see [PostgreSQL `chkpass`](#).
- Extension updates for RDS PostgreSQL 11.1 include the following:
 - `pgaudit` is updated to 1.3.0.
 - `pg_hint_plan` is updated to 1.3.2.
 - `pglogical` is updated to 2.2.1.
 - `plcoffee` is updated to 2.3.8.
 - `plv8` is updated to 2.3.8.
 - `PostGIS` is updated to 2.5.1.
 - `prefix` is updated to 1.2.8.
 - `wal2json` is updated to hash 9e962bad.

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 11 on Amazon RDS in the Database Preview Environment

Note

PostgreSQL Version 11 on Amazon RDS has been released in the production environment. It is no longer supported in the Database Preview Environment.

PostgreSQL version 11 contains several improvements that are described in [PostgreSQL 11 Released!](#)

For information on the Database Preview Environment, see [the section called “Working with the Database Preview Environment” \(p. 1362\)](#). To access the Preview Environment from the console, select <https://console.aws.amazon.com/rds-preview/>.

PostgreSQL Version 10.10 on Amazon RDS

PostgreSQL version 10.10 contains several bug fixes for issues in release 10.9. For more information on the fixes in PostgreSQL 10.10, see the [PostgreSQL documentation](#). Changes in this version include the following:

1. The `aws_s3` extension is updated to support virtual-hosted style requests. For more information, see [Amazon S3 Path Deprecation Plan – The Rest of the Story](#).
2. The `PostGIS` extension is updated to version 2.5.2.

PostgreSQL Version 10.9 on Amazon RDS

This release contains an important security fix and also bug fixes and improvements done by the PostgreSQL community. For more information on the security fix, see the [PostgreSQL community announcement](#) and [security fix CVE-2019-10164](#).

With this release, the `pg_hint_plan` extension has been updated to version 1.3.3.

For more information on the fixes in PostgreSQL 10.9, see the [PostgreSQL documentation](#).

PostgreSQL Version 10.7 on Amazon RDS

PostgreSQL version 10.7 contains several bug fixes for issues in release 10.6. For more information on the fixes in 10.7, see the [PostgreSQL documentation](#).

This version also includes the following changes:

- Support for Amazon S3 import. For more information, see [Importing Amazon S3 Data into an RDS for PostgreSQL DB Instance \(p. 1305\)](#).
- Multiple major version upgrade is available to PostgreSQL 10.7 from certain previous PostgreSQL versions. For more information, see [Major Version Upgrades for PostgreSQL \(p. 1290\)](#).

For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL Version 10.6 on Amazon RDS

PostgreSQL version 10.6 contains several bug fixes for issues in release 10.5. For more information on the fixes in PostgreSQL 10.6, see the [PostgreSQL documentation](#).

This version also includes the following changes:

- A new `rds.restrict_password_commands` parameter and a new `rds_password` role have been introduced. When the `rds.restrict_password_commands` parameter is enabled, only users who have the `rds_password` role can make user password and password expiration changes. By restricting password-related operations to a limited set of roles, you can implement policies such as password complexity requirements from the client side. The `rds.restrict_password_commands` parameter is static, so it requires a database restart to change it. For more information, see [Restricting Password Management \(p. 1349\)](#).
- The logical decoding plugin `wal2json` has been updated to commit `9e962ba`.

For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

Note

Amazon RDS for PostgreSQL has announced the removal of the `tsearch2` extension in the next major release. We encourage customers still using pre-8.3 text search to migrate to the equivalent built-in features. For more information about migrating, see the [PostgreSQL documentation](#).

PostgreSQL Version 10.5 on Amazon RDS

PostgreSQL version 10.5 contains several bug fixes for issues in release 10.4. For more information on the fixes in 10.5, see the [PostgreSQL documentation](#).

This version also includes the following changes:

- Support for the `pglogical` extension version 2.2.0. Prerequisites for using this extension are the same as the prerequisites for using logical replication for PostgreSQL as described in [Logical Replication for PostgreSQL on Amazon RDS \(p. 1406\)](#).
- Support for the `pg_similarity` extension version 1.0.
- Support for the `pageinspect` extension version 1.6.
- Support for the `libprotobuf` extension version 1.3.0 for the PostGIS component.
- An update for the `pg_hint_plan` extension to version 1.3.1.
- An update for the `wal2json` extension to version `01c5c1e`.

For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 10.4 on Amazon RDS

PostgreSQL version 10.4 contains several bug fixes for issues in release 10.3. For more information on the fixes in 10.4, see the [PostgreSQL documentation](#).

This version also includes the following changes:

- Support for PostgreSQL 10 Logical Replication using the native publication and subscription framework. RDS PostgreSQL databases can function as both publishers and subscribers. You can specify replication to other PostgreSQL databases at the database-level or at the table-level. With logical replication, the publisher and subscriber databases need not be physically identical (block-to-block) to each other. This allows for use cases such as data consolidation, data distribution, and data replication across different database versions for 10.4 and above. For more details, refer to [Logical Replication for PostgreSQL on Amazon RDS \(p. 1406\)](#).
- The temporary file size limitation is user-configurable. You require the **rds_superuser** role to modify the `temp_file_limit` parameter.
- Update of the GDAL library, which is used by the PostGIS extension. See [Working with PostGIS \(p. 1343\)](#).
- Update of the `ip4r` extension to version 2.1.1.
- Update of the `pg_repack` extension to version 1.4.3. See [Working with the pg_repack Extension \(p. 1342\)](#).
- Update of the `plv8` extension to version 2.1.2.

For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

Note

The `tsearch2` extension is to be removed in the next major release. We encourage customers still using pre-8.3 text search to migrate to the equivalent built-in features. For more information about migrating, see the [PostgreSQL documentation](#).

PostgreSQL Version 10.3 on Amazon RDS

PostgreSQL version 10.3 contains several bug fixes for issues in release 10. For more information on the fixes in 10.3, see the [PostgreSQL documentation](#).

Version 2.1.0 of PL/v8 is now available. If you use PL/v8 and upgrade PostgreSQL to a new PL/v8 version, you immediately take advantage of the new extension but the catalog metadata doesn't reflect this fact. For the steps to synchronize your catalog metadata with the new version of PL/v8, see [Upgrade PL/v8 \(p. 1404\)](#).

For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 10.1 on Amazon RDS

PostgreSQL version 10.1 contains several bug fixes for issues in release 10. For more information on the fixes in 10.1, see the [PostgreSQL documentation](#) and the [PostgreSQL 10 community announcement](#).

For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL version 10.1 includes the following changes:

- **Declarative table partitioning** – PostgreSQL 10 adds table partitioning to SQL syntax and native tuple routing.
- **Parallel queries** – When you create a new PostgreSQL 10.1 instance, parallel queries are enabled for the `default.postgres10` parameter group. The parameter `max_parallel_workers_per_gather` is set to 2 by default, but you can modify it to support your specific workload requirements.
- **Support for the International Components for Unicode (ICU)** – You can use the ICU library to provide explicitly versioned collations. Amazon RDS for PostgreSQL 10.1 is compiled with ICU version 60.2. For more information about ICU implementation in PostgreSQL, see [Collation Support](#).
- **Huge pages** – Huge pages is a feature of the Linux kernel that uses multiple page size capabilities of modern hardware architectures. Amazon RDS for PostgreSQL supports huge pages with a global configuration parameter. When you create a new PostgreSQL 10.1 instance with RDS, the `huge_pages` parameter is set to "on" for the `default.postgres10` parameter group. You can modify this setting to support your specific workload requirements.
- **PL/v8 update** – PL/v8 is a procedural language that allows you to write functions in JavaScript that you can then call from SQL. This release of PostgreSQL supports version 2.1.0 of PL/v8.
- **Renaming of xlog and location** – In PostgreSQL version 10 the abbreviation "xlog" has changed to "wal", and the term "location" has changed to "lsn". For more information, see <https://www.postgresql.org/docs/10/static/release-10.html#id-1.11.6.8.4>.
- **tsearch2 module** – Amazon RDS continues to provide the `tsearch2` module in PostgreSQL version 10, but is to remove it in the next major version release. If your application uses `tsearch2` functions update it to use the equivalent functions the core engine provides. For more information about using `tsearch2`, see [tsearch2 module](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.6.15 on Amazon RDS

PostgreSQL version 9.6.15 contains several bug fixes for issues in release 9.6.14. For more information on the fixes in PostgreSQL 9.6.15, see the [PostgreSQL documentation](#).

The PostGIS extension is updated to version 2.5.2.

PostgreSQL Version 9.6.14 on Amazon RDS

This release contains bug fixes and improvements done by the PostgreSQL community.

With this release, the `pg_hint_plan` extension has been updated to version 1.2.5.

For more information on the fixes in PostgreSQL 9.6.14, see the [PostgreSQL documentation](#).

PostgreSQL Version 9.6.12 on Amazon RDS

PostgreSQL version 9.6.12 contains several bug fixes for issues in release 9.6.11. For more information on the fixes in 9.6.12, see the [PostgreSQL documentation](#).

For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL Version 9.6.11 on Amazon RDS

PostgreSQL version 9.6.11 contains several bug fixes for issues in release 9.6.10. For more information on the fixes in PostgreSQL 9.6.11, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

With this version, the logical decoding plugin `wal2json` has been updated to commit `9e962ba`.

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.6.10 on Amazon RDS

PostgreSQL version 9.6.10 contains several bug fixes for issues in release 9.6.9. For more information on the fixes in 9.6.10, see the [PostgreSQL documentation](#).

This version includes the following changes:

- Support for the `pglogical` extension version 2.2.0. Prerequisites for using this extension are the same as the prerequisites for using logical replication for PostgreSQL as described in [Logical Replication for PostgreSQL on Amazon RDS \(p. 1406\)](#).
- Support for the `pg_similarity` extension version 2.2.0.
- An update for the `wal2json` extension to version `01c5c1e`.
- An update for the `pg_hint_plan` extension to version 1.2.3.

For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.6.9 on Amazon RDS

PostgreSQL version 9.6.9 contains several bug fixes for issues in release 9.6.8. For more information on the fixes in 9.6.9, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

This version includes the following changes:

- The temporary file size limitation is user-configurable. You require the `rds_superuser` role to modify the `temp_file_limit` parameter.
- Update of the `GDAL` library, which is used by the PostGIS extension. See [Working with PostGIS \(p. 1343\)](#).
- Update of the `ip4r` extension to version 2.1.1.
- Update of the `pgaudit` extension to version 1.1.1. See [Working with the pgaudit Extension \(p. 1341\)](#).

Update of the `pg_repack` extension to version 1.4.3. See [Working with the pg_repack Extension \(p. 1342\)](#).

- Update of the `plv8` extension to version 2.1.2.

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.6.8 on Amazon RDS

PostgreSQL version 9.6.8 contains several bug fixes for issues in release 9.6.6. For more information on the fixes in 9.6.8, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.6.6 on Amazon RDS

PostgreSQL version 9.6.6 contains several bug fixes for issues in release 9.6.5. For more information on the fixes in 9.6.6, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

This version includes the following features:

- Supports the `orafce` extension, version 3.6.1. This extension contains functions that are native to commercial databases, and can be helpful if you are porting a commercial database to PostgreSQL. For more information about using `orafce` with Amazon RDS, see [Working with the orafce Extension \(p. 1346\)](#).
- Supports the `prefix` extension, version 1.2.6. This extension provides an operator for text prefix searches. For more information about `prefix`, see the [prefix project on GitHub](#).
- Supports version 2.3.4 of PostGIS, version 2.4.2 of pgrouting, and an updated version of wal2json.

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.6.5 on Amazon RDS

PostgreSQL version 9.6.5 contains several bug fixes for issues in release 9.6.4. For more information on the fixes in 9.6.5, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

This version also includes support for the `pgrouting` and `postgresql-hll` extensions, and the `decoder_raw` optional module.

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.6.3 on Amazon RDS

PostgreSQL version 9.6.3 contains several new features and bug fixes. This version includes the following features:

- Supports the extension `pg_repack` version 1.4.0. You can use this extension to remove bloat from tables and indexes. For more information on using `pg_repack` with Amazon RDS, see [Working with the pg_repack Extension \(p. 1342\)](#).
- Supports the extension `pgaudit` version 1.1.0. This extension provides detailed session and object audit logging. For more information on using `pgaudit` with Amazon RDS, see [Working with the pgaudit Extension \(p. 1341\)](#).
- Supports `wal2json`, an output plugin for logical decoding.
- Supports the `auto_explain` module. You can use this module to log execution plans of slow statements automatically. The following example shows how to use `auto_explain` from within an Amazon RDS PostgreSQL session:

```
LOAD '$libdir/plugins/auto_explain';
```

For more information on using `auto_explain`, see the [PostgreSQL documentation](#).

PostgreSQL Version 9.6.2 on Amazon RDS

PostgreSQL version 9.6.2 contains several new features and bug fixes. The new version also includes the following extension versions:

- PostGIS version 2.3.2
- `pg_freespacemap` version 1.1—Provides a way to examine the free space map (FSM). This extension provides an overloaded function called `pg_freespace`. The functions show the value recorded in the free space map for a given page, or for all pages in the relation.
- `pg_hint_plan` version 1.1.3— Provides control of execution plans by using hinting phrases at the beginning of SQL statements.
- `log_fdw` version 1.0—Using this extension from Amazon RDS, you can load and query your database engine log from within the database. For more information, see [Using the log_fdw Extension \(p. 1402\)](#).
- With this version release, you can now edit the `max_worker_processes` parameter in a DB parameter group.

PostgreSQL version 9.6.2 on Amazon RDS also supports altering enum values. For more information, see [ALTER ENUM for PostgreSQL \(p. 1410\)](#).

For more information on the fixes in 9.6.2, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL Version 9.6.1 on Amazon RDS

PostgreSQL version 9.6.1 contains several new features and improvements. For more information about the fixes and improvements in PostgreSQL 9.6.1, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#). For information about performing parallel queries and phrase searching using Amazon RDS for PostgreSQL 9.6.1, see the [AWS Database Blog](#).

PostgreSQL version 9.6.1 includes the following changes:

- **Parallel query execution:** Supports parallel execution of large read-only queries, allowing sequential scans, hash joins, nested loops, and aggregates to be run in parallel. By default, parallel query execution is not enabled. To enable parallel query execution, set the parameter `max_parallel_workers_per_gather` to a value larger than zero.
- **Updated postgres_fdw extension:** Supports remote JOINs, SORTs, UPDATEs, and DELETE operations.
- **PL/v8 update:** Provides version 1.5.3 of the PL/v8 language.
- **PostGIS version update:** Supports POSTGIS="2.3.0 r15146" GEOS="3.5.0-CAPI-1.9.0 r4084" PROJ="Rel. 4.9.2, 08 September 2015" GDAL="GDAL 2.1.1, released 2016/07/07" LIBXML="2.9.1" LIBJSON="0.12" RASTER
- **Vacuum improvement:** Avoids scanning pages unnecessarily during vacuum freeze operations.

- **Full-text search support for phrases:** Supports the ability to specify a phrase-search query in tsquery input using the new operators <-> and <N>.
- **Two new extensions are supported:**
 - bloom, an index access method based on [Bloom filters](#)
 - pg_visibility, which provides a means for examining the visibility map and page-level visibility information of a table.
- With the release of version 9.6.2, you can now edit the max_worker_processes parameter in a PostgreSQL version 9.6.1 DB parameter group.

You can create a new PostgreSQL 9.6.1 database instance using the AWS Management Console, AWS CLI, or RDS API. You can also upgrade an existing PostgreSQL 9.5 instance to version 9.6.1 using major version upgrade. If you want to upgrade a DB instance from version 9.3 or 9.4 to 9.6, you must perform a point-and-click upgrade to the next major version first. Each upgrade operation involves a short period of unavailability for your DB instance.

PostgreSQL Version 9.5.19 on Amazon RDS

PostgreSQL version 9.5.19 contains several bug fixes for issues in release 9.6.18. For more information on the fixes in PostgreSQL 9.5.19, see the [PostgreSQL documentation](#).

The PostGIS extension is updated to version 2.5.2.

PostgreSQL Version 9.5.18 on Amazon RDS

This release contains bug fixes and improvements done by the PostgreSQL community.

With this release, the pg_hint_plan extension has been updated to version 1.1.8.

For more information on the fixes in PostgreSQL 9.5.18, see the [PostgreSQL documentation](#).

PostgreSQL Version 9.5.16 on Amazon RDS

PostgreSQL version 9.5.16 contains several bug fixes for issues in release 9.5.15. For more information on the fixes in 9.5.16, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.5.15 on Amazon RDS

PostgreSQL version 9.5.15 contains several bug fixes for issues in release 9.5.14. For more information on the fixes in 9.5.15, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.5.14 on Amazon RDS

PostgreSQL version 9.5.14 contains several bug fixes for issues in release 9.5.13. For more information on the fixes in 9.5.14, see the [PostgreSQL documentation](#).

For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.5.13 on Amazon RDS

PostgreSQL version 9.5.13 contains several bug fixes for issues in release 9.5.12. For more information on the fixes in 9.5.13, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

This version includes the following extension updates:

- Update of the `pgaudit` extension to version 1.0.6. See [Working with the pgaudit Extension \(p. 1341\)](#).
- Update of the `pg_hint_plan` extension to version 1.1.5.
- Update of the `plv8` extension to version 2.1.2.

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.5.12 on Amazon RDS

PostgreSQL version 9.5.12 contains several bug fixes for issues in release 9.5.10. For more information on the fixes in 9.5.12, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.5.10 on Amazon RDS

PostgreSQL version 9.5.10 contains several bug fixes for issues in version 9.5.9. For more information on the fixes in 9.5.10, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL Version 9.5.9 on Amazon RDS

PostgreSQL version 9.5.9 contains several bug fixes for issues in version 9.5.8. For more information on the fixes in 9.5.9, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL Version 9.5.7 on Amazon RDS

PostgreSQL version 9.5.7 contains several new features and bug fixes. This version includes the following features:

- Supports the extension `pgaudit` version 1.0.5. This extension provides detailed session and object audit logging. For more information on using `pgaudit` with Amazon RDS, see [Working with the pgaudit Extension \(p. 1341\)](#).
- Supports `wal2json`, an output plugin for logical decoding.
- Supports the `auto_explain` module. You can use this module to log execution plans of slow statements automatically. The following example shows how to use `auto_explain` from within an Amazon RDS PostgreSQL session.

```
LOAD '$libdir/plugins/auto_explain';
```

For more information on using `auto_explain`, see the [PostgreSQL documentation](#).

PostgreSQL Version 9.5.6 on Amazon RDS

PostgreSQL version 9.5.6 contains several new features and bug fixes. The new version also includes the following extension versions:

- PostGIS version 2.2.5
- [pg_freespacemap](#) version 1.1–Provides a way to examine the free space map (FSM). This extension provides an overloaded function called pg_freespace. This function shows the value recorded in the free space map for a given page, or for all pages in the relation.
- [pg_hint_plan](#) version 1.1.3– Provides control of execution plans by using hinting phrases at the beginning of SQL statements.

PostgreSQL version 9.5.6 on Amazon RDS also supports altering enum values. For more information, see [ALTER ENUM for PostgreSQL \(p. 1410\)](#).

For more information on the fixes in 9.5.6, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL Version 9.5.4 on Amazon RDS

PostgreSQL version 9.5.4 contains several fixes to issue found in previous versions. For more information on the fixes in 9.5.4, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

Beginning with PostgreSQL version 9.4, PostgreSQL supports the streaming of WAL changes using logical replication decoding. Amazon RDS supports logical replication for PostgreSQL version 9.4.9 and higher and 9.5.4 and higher. For more information about PostgreSQL logical replication on Amazon RDS, see [Logical Replication for PostgreSQL on Amazon RDS \(p. 1406\)](#).

Beginning with PostgreSQL version 9.5.4 for Amazon RDS, the command ALTER USER WITH BYPASSRLS is supported.

PostgreSQL versions 9.4.9 and later and version 9.5.4 and later support event triggers, and Amazon RDS supports event triggers for these versions. You can use the master user account can be used to create, modify, rename, and delete event triggers. Event triggers are at the DB instance level, so they can apply to all databases on an instance. For more information about PostgreSQL event triggers on Amazon RDS, see [Event Triggers for PostgreSQL on Amazon RDS \(p. 1408\)](#).

PostgreSQL Version 9.5.2 on Amazon RDS

PostgreSQL version 9.5.2 contains several fixes to issues found in previous versions. For more information on the features in 9.5.2, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL version 9.5.2 doesn't support the db.m1 or db.m2 DB instance classes. If you need to upgrade a DB instance running PostgreSQL version 9.4 to version 9.5.2 to one of these instance classes, you need to scale compute. To do that, you need a comparable db.t2 or db.m3 DB instance class before you can upgrade a DB instance running PostgreSQL version 9.4 to version 9.5.2. For more information on DB instance classes, see [Choosing the DB Instance Class \(p. 107\)](#).

Native PostgreSQL version 9.5.2 introduced the command ALTER USER WITH BYPASSRLS.

This release includes updates from previous versions, including the following:

- **CVE-2016-2193:** Fixes an issue where a query plan might be reused for more than one ROLE in the same session. Reusing a query plan can cause the query to use the wrong set of Row Level Security (RLS) policies.

- **CVE-2016-3065:** Fixes a server crash bug triggered by using `pageinspect` with BRIN index pages. Because an attacker might be able to expose a few bytes of server memory, this crash is being treated as a security issue.

Major enhancements in RDS PostgreSQL 9.5 include the following:

- UPSERT: Allow INSERTs that would generate constraint conflicts to be turned into UPDATEs or ignored
- Add the GROUP BY analysis features GROUPING SETS, CUBE, and ROLLUP
- Add row-level security control
- Create mechanisms for tracking the progress of replication, including methods for identifying the origin of individual changes during logical replication
- Add Block Range Indexes (BRIN)
- Add substantial performance improvements for sorting
- Add substantial performance improvements for multi-CPU machines
- PostGIS 2.2.2 - To use this latest version of PostGIS, use the ALTER EXTENSION UPDATE statement to update after you upgrade to version 9.5.2. Example:
`ALTER EXTENSION POSTGIS UPDATE TO '2.2.2'`
- Improved visibility of autovacuum sessions by allowing the `rds_superuser` account to view autovacuum sessions in `pg_stat_activity`. For example, you can identify and terminate an autovacuum session that is blocking a command from running, or executing slower than a manually issued vacuum command.

RDS PostgreSQL version 9.5.2 includes the following new extensions:

- **address_standardizer** – A single-line address parser that takes an input address and normalizes it based on a set of rules stored in a table, helper lex, and gaz tables.
- **hstore_plperl** – Provides transforms for the `hstore` type for PL/Perl.
- **tsm_system_rows** – Provides the table sampling method `SYSTEM_ROWS`, which can be used in the `TABLESAMPLE` clause of a `SELECT` command.
- **tsm_system_time** – Provides the table sampling method `SYSTEM_TIME`, which can be used in the `TABLESAMPLE` clause of a `SELECT` command.

PostgreSQL Version 9.4.24 on Amazon RDS

PostgreSQL version 9.4.24 contains several bug fixes for issues in release 9.4.23. For more information on the fixes in PostgreSQL 9.4.24, see the [PostgreSQL documentation](#).

The PostGIS extension is updated to version 2.5.2.

PostgreSQL Version 9.4.23 on Amazon RDS

This release contains bug fixes and improvements done by the PostgreSQL community. For more information on the fixes in PostgreSQL 9.4.23, see the [PostgreSQL documentation](#).

PostgreSQL Version 9.4.21 on Amazon RDS

PostgreSQL version 9.4.21 contains several bug fixes for issues in release 9.4.20. For more information on the fixes in 9.4.21, the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.4.20 on Amazon RDS

PostgreSQL version 9.4.20 contains several bug fixes for issues in release 9.4.19. For more information on the fixes in 9.4.20, the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.4.19 on Amazon RDS

PostgreSQL version 9.4.19 contains several bug fixes for issues in release 9.4.18. For more information on the fixes in 9.4.19, see the [PostgreSQL documentation](#).

For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.4.18 on Amazon RDS

PostgreSQL version 9.4.18 contains several bug fixes for issues in release 9.4.17. For more information on the fixes in 9.4.18, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

The plv8 extension has been updated to version 2.1.2. For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.4.17 on Amazon RDS

PostgreSQL version 9.4.17 contains several bug fixes for issues in release 9.4.15. For more information on the fixes in 9.4.17, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.4.15 on Amazon RDS

PostgreSQL version 9.4.15 contains several bug fixes for issues in release 9.4.14. For more information on the fixes in 9.4.15, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL Version 9.4.14 on Amazon RDS

PostgreSQL version 9.4.14 contains several bug fixes for issues in release 9.4.12. For more information on the fixes in 9.4.14, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL Version 9.4.12 on Amazon RDS

PostgreSQL version 9.4.12 contains several fixes to issue found in previous versions.

For more information on the fixes in 9.4.12, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL Version 9.4.11 on Amazon RDS

PostgreSQL version 9.4.11 contains several fixes to issue found in previous versions.

For more information on the fixes in 9.4.11, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

Beginning with PostgreSQL version 9.4, PostgreSQL supports the streaming of WAL changes using logical replication decoding. Amazon RDS supports logical replication for PostgreSQL version 9.4.9 and higher and 9.5.4 and higher. For more information about PostgreSQL logical replication on Amazon RDS, see [Logical Replication for PostgreSQL on Amazon RDS \(p. 1406\)](#).

PostgreSQL versions 9.4.9 and later and version 9.5.4 and later support event triggers, and Amazon RDS supports event triggers for these versions. The master user account can be used to create, modify, rename, and delete event triggers. Event triggers are at the DB instance level, so they can apply to all databases on an instance. For more information about PostgreSQL event triggers on Amazon RDS, see [Event Triggers for PostgreSQL on Amazon RDS \(p. 1408\)](#).

PostgreSQL Version 9.4.9 on Amazon RDS

PostgreSQL version 9.4.9 contains several fixes to issue found in previous versions. For more information on the fixes in 9.4.9, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

Beginning with PostgreSQL version 9.4, PostgreSQL supports the streaming of WAL changes using logical replication decoding. Amazon RDS supports logical replication for PostgreSQL version 9.4.9 and higher and 9.5.4 and higher. For more information about PostgreSQL logical replication on Amazon RDS, see [Logical Replication for PostgreSQL on Amazon RDS \(p. 1406\)](#).

PostgreSQL versions 9.4.9 and later and version 9.5.4 and later support event triggers, and Amazon RDS supports event triggers for these versions. The master user account can be used to create, modify, rename, and delete event triggers. Event triggers are at the DB instance level, so they can apply to all databases on an instance. For more information about PostgreSQL event triggers on Amazon RDS, see [Event Triggers for PostgreSQL on Amazon RDS \(p. 1408\)](#).

PostgreSQL Version 9.4.7 on Amazon RDS

PostgreSQL version 9.4.7 contains several fixes to issue found in previous versions. For more information on the fixes in 9.4.7, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL version 9.4.7 includes improved visibility of autovacuum sessions by allowing the rds_superuser account to view autovacuum sessions in pg_stat_activity. For example, you can identify and terminate an autovacuum session that is blocking a command from running, or executing slower than a manually issued vacuum command.

PostgreSQL Version 9.3.25 on Amazon RDS

Note

Amazon RDS for PostgreSQL announced retirement of PostgreSQL 9.3 in September 2018 and has stopped support for PostgreSQL version 9.3. We encourage you to upgrade to PostgreSQL 9.4 or a higher version as early as possible. For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL version 9.3.25 contains several bug fixes for issues in release 9.3.24. For more information on the fixes in 9.3.25, see the [PostgreSQL documentation](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.3.24 on Amazon RDS

PostgreSQL version 9.3.24 contains several bug fixes for issues in release 9.3.23. For more information on the fixes in 9.3.24, see the [PostgreSQL documentation](#).

Note

Amazon RDS for PostgreSQL has deprecated PostgreSQL version 9.3.x.

For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.3.23 on Amazon RDS

PostgreSQL version 9.3.23 contains several bug fixes for issues in release 9.3.22. For more information on the fixes in 9.3.23, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

The plv8 extension has been updated to version 2.1.2. For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

Note

Amazon RDS for PostgreSQL has announced version 9.3.x is scheduled for retirement in September, 2018. We encourage you to upgrade your 9.3.x databases to the latest version at your earliest convenience.

PostgreSQL Version 9.3.22 on Amazon RDS

PostgreSQL version 9.3.22 contains several bug fixes for issues in release 9.3.20. For more information on the fixes in 9.3.22, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

For the complete list of extensions supported by Amazon RDS for PostgreSQL, see [Supported PostgreSQL Features and Extensions \(p. 1383\)](#).

PostgreSQL Version 9.3.20 on Amazon RDS

PostgreSQL version 9.3.20 contains several bug fixes for issues in version 9.3.19. For more information on the fixes in 9.3.20, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL Version 9.3.19 on Amazon RDS

PostgreSQL version 9.3.19 contains several bug fixes for issues in version 9.3.18. For more information on the fixes in 9.3.19, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL Version 9.3.17 on Amazon RDS

PostgreSQL version 9.3.17 contains several fixes for bugs found in previous versions. This version contains the same extension components as version 9.3.16. For a list of fixes in version 9.3.17, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL Version 9.3.16 on Amazon RDS

PostgreSQL version 9.3.16 contains several fixes for bugs found in previous versions. This version contains the same extension components as version 9.3.14. For a list of fixes in version 9.3.16, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL Version 9.3.14 on Amazon RDS

PostgreSQL version 9.3.14 contains several fixes for bugs found in previous versions. For a list of fixes in version 9.3.14, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL Version 9.3.12 on Amazon RDS

PostgreSQL version 9.3.12 contains several fixes for bugs found in previous versions. For a list of fixes in version 9.3.12, see the [PostgreSQL documentation](#). For information on upgrading the engine version for your PostgreSQL DB instance, see [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#).

PostgreSQL version 9.3.12 includes improved visibility of autovacuum sessions by allowing the `rds_superuser` account to view autovacuum sessions in `pg_stat_activity`. For example, you can identify and terminate an autovacuum session that is blocking a command from running, or executing slower than a manually issued `vacuum` command.

Supported PostgreSQL Features and Extensions

Amazon RDS supports many of the most common PostgreSQL extensions and features.

Topics

- [PostgreSQL Extensions and Modules Supported on Amazon RDS \(p. 1383\)](#)
- [Upgrade PL/v8 \(p. 1404\)](#)
- [Supported PostgreSQL Features \(p. 1406\)](#)
- [Limits for PostgreSQL DB Instances \(p. 1410\)](#)
- [Upgrading a PostgreSQL DB Instance \(p. 1411\)](#)
- [Using SSL with a PostgreSQL DB Instance \(p. 1411\)](#)

PostgreSQL Extensions and Modules Supported on Amazon RDS

PostgreSQL supports many PostgreSQL extensions and modules. Extensions and modules expand on the functionality provided by the PostgreSQL engine. The following sections show the extensions and modules supported by Amazon RDS for the major PostgreSQL versions.

Topics

- [PostgreSQL Version 12.x Extensions and Modules Supported on Amazon RDS in the Database Preview Environment \(p. 1384\)](#)
- [PostgreSQL Version 11.x Extensions and Modules Supported on Amazon RDS \(p. 1386\)](#)
- [PostgreSQL Version 10.x Extensions and Modules Supported on Amazon RDS \(p. 1388\)](#)
- [PostgreSQL Version 9.6.x Extensions and Modules Supported on Amazon RDS \(p. 1391\)](#)
- [PostgreSQL Version 9.5.x Extensions Supported on Amazon RDS \(p. 1394\)](#)
- [PostgreSQL Version 9.4.x Extensions and Modules Supported on Amazon RDS \(p. 1396\)](#)
- [PostgreSQL Version 9.3.x Extensions Supported on Amazon RDS \(p. 1398\)](#)
- [PostgreSQL Extension Support for PostGIS on Amazon RDS \(p. 1400\)](#)

- [Using the log_fdw Extension \(p. 1402\)](#)

You can find a list of extensions supported by Amazon RDS in the default DB parameter group for that PostgreSQL version. You can also see the current extensions list using `psql` by showing the `rds.extensions` parameter as in the following example.

```
SHOW rds.extensions;
```

Note

Parameters added in a minor version release might display inaccurately when using the `rds.extensions` parameter in `psql`.

PostgreSQL Version 12.x Extensions and Modules Supported on Amazon RDS in the Database Preview Environment

The following table shows PostgreSQL extensions and modules for PostgreSQL version 12.x that are currently supported by PostgreSQL in the Amazon RDS Preview Environment. "N/A" indicates that the extension or module is not available for that PostgreSQL version. For more information on PostgreSQL extensions, see [Packaging Related Objects into an Extension](#).

Extension	Version 12 Beta 3	Version 12.0
address_standardizer	3.0.0alpha3	3.0.0alpha3
address_standardizer_data_us	3.0.0alpha3	3.0.0alpha3
amcheck	1.2	1.2
aws_commons —see Importing Data into PostgreSQL on Amazon RDS	1.0	1.0
aws_s3 —see Importing Data into PostgreSQL on Amazon RDS	N/A	1.0
bloom	1.0	1.0
btree_gin	1.3	1.3
btree_gist	1.5	1.5
citext	1.6	1.6
cube	1.4	1.4
dblink	1.2	1.2
dict_int	1.0	1.0
dict_xsyn	1.0	1.0
earthdistance	1.1	1.1
fuzzystrmatch	1.1	1.1
hstore	1.6	1.6
hstore_plperl	1.0	1.0

Extension	Version 12 Beta 3	Version 12.0
intagg	1.1	1.1
intarray	1.2	1.2
ip4r	2.2	2.2
isn	1.2	1.2
jsonb_plperl	1.0	1.0
log_fdw—see Using the log_fdw Extension (p. 1402)	1.1	1.1
ltree	1.1	1.1
pageinspect	1.7	1.7
pg_buffercache	1.3	1.3
pg_freespacemap	1.2	1.2
pg_prewarm	1.2	1.2
pg_similarity	1.0	1.0
pg_stat_statements	1.7	1.7
pg_trgm	1.4	1.4
pg_visibility	1.2	1.2
pgcrypto	1.3	1.3
pgrouting	2.6.1	2.6.1
pgrowlocks	1.2	1.2
pgstattuple	1.5	1.5
pgTAP	1.0.0	1.0.0
plperl	1.0	1.0
plpgsql	1.0	1.0
pltcl	1.0	1.0
PostGIS	3.0.0alpha3	3.0.0alpha3
postgis_tiger_geocoder	3.0.0alpha3	3.0.0alpha3
postgis_topology	3.0.0alpha3	3.0.0alpha3
postgres_fdw	1.0	1.0
prefix	1.2.0	1.2.0
sslinfo	1.2	1.2
tablefunc	1.0	1.0

Extension	Version 12 Beta 3	Version 12.0
test_parser	1.0	1.0
tsm_system_rows	1.0	1.0
tsm_system_time	1.0	1.0
unaccent	1.1	1.1
uuid-ossp	1.1	1.1

PostgreSQL Version 11.x Extensions and Modules Supported on Amazon RDS

The following tables show PostgreSQL extensions and modules for PostgreSQL version 11.x that are currently supported by PostgreSQL on Amazon RDS. "N/A" indicates that the extension or module is not available for that PostgreSQL version. For more information on PostgreSQL extensions, see [Packaging Related Objects into an Extension](#).

Extension	Version 11.1	Version 11.2	Version 11.4	Version 11.5
address_standardizer	2.4.2	2.4.2	2.4.2	2.4.2
address_standardizer_data_us	2.4.2	2.4.2	2.4.2	2.4.2
bloom	1.0	1.0	1.0	1.0
btree_gin	1.2	1.2	1.2	1.2
btree_gist	1.5	1.5	1.5	1.5
citext	1.4	1.4	1.4	1.4
cube	1.2	1.2	1.2	1.2
dblink	1.2	1.2	1.2	1.2
dict_int	1.0	1.0	1.0	1.0
dict_xsyn	1.0	1.0	1.0	1.0
earthdistance	1.1	1.1	1.1	1.1
fuzzystrmatch	1.1	1.1	1.1	1.1
hstore	1.4	1.4	1.4	1.4
hstore_plperl	1.0	1.0	1.0	1.0
intagg	1.1	1.1	1.1	1.1
intarray	1.2	1.2	1.2	1.2
ip4r	2.3	2.3	2.3	2.3
isn	1.1	1.1	1.1	1.1

Extension	Version 11.1	Version 11.2	Version 11.4	Version 11.5
log_fdw —see Using the log_fdw Extension (p. 1402)	1.0	1.0	1.0	1.0
libprotobuf	1.3.0	1.3.0	1.3.0	1.3.0
ltree	1.1	1.1	1.1	1.1
orafce	3.6.1	3.6.1	3.6.1	3.6.1
pgaudit	1.3.0	1.3.0	1.3.0	1.3.0
pg_buffercache	1.3	1.3	1.3	1.3
pg_freespacemap	1.2	1.2	1.2	1.2
pg_hint_plan	1.3.2	1.3.2	1.3.4	1.3.4
pg_prewarm	1.1	1.1	1.1	1.1
pg_repack	1.4.4	1.4.4	1.4.4	1.4.4
pg_similarity	1.0	1.0	1.0	1.0
pg_stat_statements	1.5	1.5	1.5	1.5
pg_transport	N/A	N/A	N/A	1.0
pg_trgm	1.3	1.3	1.3	1.3
pg_visibility	1.2	1.2	1.2	1.2
pgcrypto	1.3	1.3	1.3	1.3
pageinspect	1.6	1.6	1.6	1.6
pglogical	2.2.1	2.2.1	2.2.1	2.2.1
pgrowlocks	1.2	1.2	1.2	1.2
pgrouting	2.6.1	2.6.1	2.6.1	2.6.1
pgstattuple	1.5	1.5	1.5	1.5
pgTAP	N/A	1.0	1.0	1.0
plcoffee	2.3.8	2.3.8	2.3.8	2.3.8
plls	2.3.8	2.3.8	2.3.8	2.3.8
plperl	1.0	1.0	1.0	1.0
plpgsql	1.0	1.0	1.0	1.0
pltcl	1.0	1.0	1.0	1.0
plv8	2.3.8	2.3.8	2.3.8	2.3.8
PostGIS	2.5.1	2.5.1	2.5.1	2.5.2
postgis_tiger_geocoder	2.5.1	2.5.1	2.5.1	2.5.1

Extension	Version 11.1	Version 11.2	Version 11.4	Version 11.5
postgis_topology	2.5.0	2.5.0	2.5.0	2.5.0
postgres_fdw	1.0	1.0	1.0	1.0
postgresql-hll	2.11	2.11	2.11	2.11
prefix	1.2.8	1.2.8	1.2.8	1.2.8
sslinfo	1.2	1.2	1.2	1.2
tablefunc	1.0	1.0	1.0	1.0
test_parser	1.0	1.0	1.0	1.0
tsm_system_rows	1.0	1.0	1.0	1.0
tsm_system_time	1.0	1.0	1.0	1.0
unaccent	1.1	1.1	1.1	1.1
uuid-ossp	1.1	1.1	1.1	1.1

The following modules are supported as shown for PostgreSQL version 11.x.

Module	Version 11.1	Version 11.2	Version 11.4	Version 11.5
amcheck	Supported	Supported	Supported	Supported
auto_explain	Supported	Supported	Supported	Supported
decoder_ras	Supported	Supported	Supported	Supported
ICU	Version 60.2 supported	Version 60.2 supported	Version 60.2 supported	Version 60.2 supported
test_decode	Supported	Supported	Supported	Supported
wal2json	Commit hash 9e962bad	Commit hash 9e962bad	Commit hash 9e962bad	Commit hash 9e962bad

PostgreSQL Version 10.x Extensions and Modules Supported on Amazon RDS

The following tables show PostgreSQL extensions and modules for PostgreSQL version 10 that are currently supported by PostgreSQL on Amazon RDS. "N/A" indicates that the extension or module is not available for that PostgreSQL version. For more information on PostgreSQL extensions, see [Packaging Related Objects into an Extension](#).

Extension	10.1	10.3	10.4	10.5	10.6	10.7	10.9	10.10
address_standard	214.22.4.22.4.22.4.22.4.22.4.22.4.2	214.22.4.22.4.22.4.22.4.22.4.22.4.2	214.22.4.22.4.22.4.22.4.22.4.22.4.2	214.22.4.22.4.22.4.22.4.22.4.22.4.2	214.22.4.22.4.22.4.22.4.22.4.22.4.2	214.22.4.22.4.22.4.22.4.22.4.22.4.2	N/A	N/A

Extension	10.1	10.3	10.4	10.5	10.6	10.7	10.9	10.10
address_standard	2.1.4	2.2.4	2.2.4	2.2.4	2.2.4	2.2.4	2.4.2	2.4.2
bloom	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
btree_gin	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2
btree_gist	1.5	1.5	1.5	1.5	1.5	1.5	1.5	1.5
chkpass	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
citext	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4
cube	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2
dblink	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2
dict_int	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
dict_xsyn	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
earthdistance	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
fuzzystrmatch	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
hstore	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4
hstore_plperl	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
intagg	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
intarray	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2
ip4r	2.0	2.0	2.1.	12.1.12.1.12.1.12.1.12.1.1	2.1.1			
isn	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
log_fdw— see Using the log_fdw Extension (p. 1402)	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
libprotobuf	N/A	N/A	N/A	1.3.0	1.3.0	1.3.0	1.3.0	1.3.0
ltree	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
orafce	3.6.1	3.6.1	3.6.1	3.6.1	3.6.1	3.6.1	3.6.1	3.6.1
pgaudit	1.2.0	1.2.0	1.2.0	1.2.0	1.2.0	1.2.0	1.2.0	1.2.0
pg_buffercache	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3
pg_freespacemap	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2
pg_hint_plan	1.3.0	1.3.0	1.3.0	1.3.1	1.3.1	1.3.1	1.3.3	1.3.3
pg_prewarm	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
pg_repack	1.4.2	1.4.2	1.4.3	1.4.3	1.4.3	1.4.3	1.4.3	1.4.3

Extension	10.1	10.3	10.4	10.5	10.6	10.7	10.9	10.10
pg_similarity	N/A	N/A	N/A	1.0	1.0	1.0	1.0	1.0
pg_stat_statements	1.5	1.5	1.5	1.5	1.5	1.5	1.5	1.5
pg_trgm	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3
pg_visibility	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2
pgcrypto	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3
pageinspect	N/A	N/A	N/A	1.6	1.6	1.6	1.6	1.6
pglogical	N/A	N/A	N/A	2.2.0	2.0	2.0	2.0	2.2.0
pgrowlocks	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2
pgrouting	2.5.22.5.22.5.22.5.22.5.22.5.22.5.22.5.2	2.5.22.5.22.5.22.5.22.5.22.5.22.5.22.5.2	2.5.22.5.22.5.22.5.22.5.22.5.22.5.22.5.2	2.5.22.5.22.5.22.5.22.5.22.5.22.5.22.5.2	2.5.22.5.22.5.22.5.22.5.22.5.22.5.22.5.2	2.5.22.5.22.5.22.5.22.5.22.5.22.5.22.5.2	2.5.22.5.22.5.22.5.22.5.22.5.22.5.22.5.2	2.5.22.5.22.5.22.5.22.5.22.5.22.5.22.5.2
pgstattuple	1.5	1.5	1.5	1.5	1.5	1.5	1.5	1.5
plcoffee	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2
plls	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2
plperl	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
plpgsql	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
pltcl	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
plv8	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2	2.1.02.1.02.1.22.1.22.1.22.1.22.1.22.1.2
PostGIS	2.4.22.4.22.4.42.4.42.4.42.4.42.4.42.4.4	2.4.22.4.22.4.42.4.42.4.42.4.42.4.42.4.4	2.4.22.4.22.4.42.4.42.4.42.4.42.4.42.4.4	2.4.22.4.22.4.42.4.42.4.42.4.42.4.42.4.4	2.4.22.4.22.4.42.4.42.4.42.4.42.4.42.4.4	2.4.22.4.22.4.42.4.42.4.42.4.42.4.42.4.4	2.4.22.4.22.4.42.4.42.4.42.4.42.4.42.4.4	2.4.22.4.22.4.42.4.42.4.42.4.42.4.42.4.4
postgis_tiger_geometry	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2
postgis_topology	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2	2.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.22.4.2
postgres_fdw	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
postgresql-hll	2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2	2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2	2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2	2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2	2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2	2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2	2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2	2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2.10.2
prefix	1.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.0	1.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.0	1.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.0	1.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.0	1.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.0	1.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.0	1.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.0	1.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.01.2.0
sslinfo	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2
tablefunc	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
test_parser	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
tsearch2 (deprecated in version 10)	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
tsm_system_rows	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
tsm_system_time	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0

Extension	10.1	10.3	10.4	10.5	10.6	10.7	10.9	10.10
unaccent	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
uuid-ossp	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1

The `tsearch2` extension is deprecated in version 10. The PostgreSQL team plans to remove `tsearch2` from the next major release of PostgreSQL.

The following modules are supported as shown for versions of PostgreSQL 10.

Module	Versio 10.1	10.3	10.4	10.5	10.6	10.7	10.9	10.10
amcheck	Not supported	Supported	Supported	Supported	Supported	Supported	Supported	Supported
auto_explain	Supported	Supported						
decoder_raw	Supported	Supported						
ICU	Version 60.2 supported	Version 60.2 supported						
test_decode	Supported	Supported						
wal2json	Commit hash 5352cc4352535204c59e962ba9e962ba9e962ba	Commit hash 9e962ba						

PostgreSQL Version 9.6.x Extensions and Modules Supported on Amazon RDS

The following tables show PostgreSQL extensions and modules for PostgreSQL version 9.6.x that are currently supported by PostgreSQL on Amazon RDS. "N/A" indicates that the extension or module is not available for that PostgreSQL version. For more information on PostgreSQL extensions, see [Packaging Related Objects into an Extension](#).

Extension	9.6	9.6.1	9.6.15									
address_standard	2.3.0	2.3.2	2.3.2	2.3.2	2.3.2	2.3.2	2.3.2	2.3.2	2.3.2	2.3.2	2.3.4	
address_standard	2.3.0	2.3.2	2.3.2	2.3.2	2.3.2	2.3.2	2.3.2	2.3.2	2.3.2	2.3.2	2.3.4	
bloom	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
btree_gin	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
btree_gist	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2
chkpass	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
citext	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3
cube	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2
dblink	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2

Extension	9.6	9.6	9.6	9.6	9.6	9.6	9.6	9.6	9.6	9.6	9.6	9.6.1	9.6.15
dict_int	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
dict_xsyn	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
earthdistance	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
fuzzystrmatch	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
hstore	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4
hstore_plperl	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
intagg	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
intarray	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2
ip4r	2.0	2.0	2.0	2.0	2.0	2.0	2.1.	12.1.	12.1.	12.1.	12.1.	2.1.1	
isn	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
log_fdw— see Using the log_fdw Extension (p. 1402)	N/ A	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
ltree	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
orafce	N/ A	N/ A	N/ A	N/ A	3.6.	13.6.	13.6.	13.6.	13.6.	13.6.	13.6.	13.6.1	3.6.1
pgaudit	N/ A	N/ A	1.1	1.1	1.1	1.1	1.1.	11.1.	11.1.	11.1.	11.1.	1.1.1	
pg_buffercache	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2
pg_freespacem	N/ A	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
pg_hint_plan	N/ A	1.1.	1.3.	1.3.	1.31.	1.31.	2.2.	21.2.	21.2.	31.2.	31.2.	2.5	1.2.5
pg_prewarm	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
pg_repack	N/ A	N/ A	1.4.	0.4.	4.11.	4.21.	4.21.	4.31.	4.31.	4.31.	4.31.	4.3	1.4.3
pg_similarity	N/ A	1.0	1.0	1.0	1.0	1.0	1.0						
pg_stat_statement	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4
pg_trgm	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3
pg_visibility	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
pgcrypto	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3
pglogical	N/ A	2.2.	02.2.	02.2.	02.2.	2.2.0							

The following modules are supported as shown for versions of PostgreSQL 9.6.

PostgreSQL Version 9.5.x Extensions Supported on Amazon RDS

The following tables show PostgreSQL extensions and modules for PostgreSQL version 9.5.x that are currently supported by PostgreSQL on Amazon RDS. "N/A" indicates that the extension or module is not available for that PostgreSQL version. For more information on PostgreSQL extensions, see [Packaging Related Objects into an Extension](#).

Extension	9.5	9.5	9.5	9.5	9.5	9.5	9.5	9.5	9.5	9.5	9.5	9.5	9.5	9.5.19
pgaudit	N/ A	N/ A	N/ A	1.0	0.5	0.5	0.5	0.5	1.0	51.0.61	0.61	0.61	0.61	0.61
pg_buffercache	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
pg_freespacemap	N/ A	N/ A	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
pg_hint_plan	N/ A	N/ A	1.1	3.1.3	3.1.3	3.1.3	31.1.31	1.51	1.51	1.51	1.51	1.51	1.81	1.8
pg_prewarm	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
pg_stat_statements	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3
pg_trgm	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
pg_visibility	N/ A	N/ A	N/ A	N/ A	N/ A	N/ A	N/ A	N/ A	N/ A	N/ A	N/ A	N/ A	N/ A	N/A
pgcrypto	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2
pgrowlocks	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
pgstattuple	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3	1.3
plcoffee	1.4.4	4.4.4	4.4.4	4.4.4	4.4.4	4.4.4	42.1.02	1.02	1.02	1.02	1.02	1.02	1.02	1.0
plls	1.4.4	4.4.4	4.4.4	4.4.4	4.4.4	4.4.4	42.1.02	1.02	1.02	1.02	1.02	1.02	1.02	1.0
plperl	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
plpgsql	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
pltcl	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
plv8	1.4.4	4.4.4	4.4.4	4.4.4	4.4.4	4.4.4	42.1.02	1.02	1.02	1.02	1.02	1.02	1.02	1.02
PostGIS	2.2.2	2.2.2	2.2.2	2.2.2	2.2.2	2.2.2	2.52	2.52	2.52	2.52	2.52	2.52	2.52.5.2	
postgis_tiger	2.2.2	2.2.2	2.2.2	2.2.2	2.2.2	2.2.2	2.52	2.52	2.52	2.52	2.52	2.52	2.52	2.5
postgis_topology	2.2.2	2.2.2	2.2.2	2.2.2	2.2.2	2.2.2	2.52	2.52	2.52	2.52	2.52	2.52	2.52	2.5
postgres_fdw	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
sslinfo	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
tablefunc	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
test_parser	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
tsearch2	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
tsm_system_row	N/ A	N/ A	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
tsm_system_time	N/ A	N/ A	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
unaccent	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0

The following modules are supported as shown for versions of PostgreSQL 9.5.

PostgreSQL Version 9.4.x Extensions and Modules Supported on Amazon RDS

The following tables show the PostgreSQL extensions and modules for PostgreSQL version 9.4.x that are currently supported by PostgreSQL on Amazon RDS. "N/A" indicates that the extension or module is not available for that PostgreSQL version. For more information on PostgreSQL extensions, see [Packaging Related Objects into an Extension](#).

Extension	9.1	9.2	9.3	9.4	9.4	9.4	9.4	9.4	9.4	9.4	9.4	9.4	9.4	9.4.24
hstore_plperl	N/A													
intagg	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
intarray	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
ip4r	2.0	2.0	2.0	2.0	2.0	2.0	2.0	2.0	2.0	2.0	2.0	2.0	2.0	2.0
isn	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
log_fdw—see Using the log_fdw Extension (p. 1402)	N/A													
ltree	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
pg_buffercache	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
pg_freespacemap	N/A													
pg_hint_plan	N/A													
pg_prewarm	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
pg_stat_statements	1.2	2.1	2.1	2.2	2.2	2.2	2.2	2.2	2.2	2.2	2.2	2.2	2.2	2.2
pg_trgm	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
pg_visibility	N/A													
pgcrypto	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
pgrowlocks	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
pgstattuple	1.2	2.1	2.1	2.2	2.2	2.2	2.2	2.2	2.2	2.2	2.2	2.2	2.2	2.2
plcoffee	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4
plls	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4
plperl	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
plpgsql	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
pltcl	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
plv8	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4	1.4
PostGIS	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1
postgis_tiger_geocoder	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1
postgis_topology	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1	2.1
postgres_fdw	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
sslinfo	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0

Extension	9.4	9.4	9.4	9.4	9.4	9.4	9.4	9.4	9.4	9.4	9.4	9.4	9.4.24
tablefunc	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
test_parser	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
tsearch2	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
tsm_system_rows	N/A												
tsm_system_time	N/A												
unaccent	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
uuid-ossp	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0

The following modules are supported as shown for versions of PostgreSQL 9.4.

Module	9.4.7	9.4.9	9.4.11	9.4.12	9.4.14	9.4.17
test_decoding	N/A	N/A	N/A	Supported	Supported	Supported

PostgreSQL Version 9.3.x Extensions Supported on Amazon RDS

Note

Amazon RDS for PostgreSQL has deprecated PostgreSQL version 9.3.x. We strongly recommend that you upgrade to a major version, preferably version 9.6.x or 10.x. See [Upgrading the PostgreSQL DB Engine for Amazon RDS \(p. 1289\)](#).

The following table shows PostgreSQL extensions for PostgreSQL version 9.3.x that are currently supported by PostgreSQL on Amazon RDS. "N/A" indicates that the extension is not available for that PostgreSQL version. For more information on PostgreSQL extensions, see [Packaging Related Objects into an Extension](#).

Extension	9.3.1	9.3.1	9.3.1	9.3.1	9.3.1	9.3.1	9.3.1	9.3.1	9.3.1	9.3.25
address_standard	N/A									
address_standard_us	N/A									
bloom	N/A									
btree_gin	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
btree_gist	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
chkpass	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
citext	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
cube	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
dblink	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1

Extension	9.3.1	9.3.1	9.3.1	9.3.1	9.3.1	9.3.1	9.3.1	9.3.1	9.3.2	9.3.25
dict_int	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
dict_xsyn	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
earthdistance	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
fuzzystrmatch	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
hstore	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2	1.2
hstore_plperl	N/A									
intagg	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
intarray	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
ip4r	N/A									
isn	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
log_fdw— see Using the log_fdw Extension (p. 1402)	N/A									
ltree	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
pg_buffercache	N/A									
pg_freespace	N/A									
pg_hint_plan	N/A									
pg_prewarm	N/A									
pg_stat_statements	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
pg_trgm	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
pg_visibility	N/A									
pgcrypto	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
pgrowlocks	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1	1.1
pgstattuple	N/A									
plcoffee	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4
plls	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4
plperl	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0

Extension	9.3.1	9.3.1	9.3.1	9.3.1	9.3.1	9.3.1	9.3.1	9.3.1	9.3.2	9.3.25
plpgsql	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
pltcl	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
plv8	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4	1.4.4	2.1.02	2.1.22	2.1.2
PostGIS	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8
postgis_tiger_geod	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8
postgis_topology	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8	2.1.8
postgres_fdw	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
sslinfo	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
tablefunc	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
test_parser	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
tsearch2	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
tsm_system_row	N/A	N/A	N/A							
tsm_system_time	N/A	N/A	N/A							
unaccent	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
uuid-ossp	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0

PostgreSQL Extension Support for PostGIS on Amazon RDS

Before you can use the PostGIS extension, you must create it by running the following command.

```
CREATE EXTENSION POSTGIS;
```

The following table shows the PostGIS component versions that ship with the Amazon RDS for PostgreSQL versions.

PostgreSQL	PostGIS	GEOS	GDAL	PROJ
9.3.12	2.1.8 r13780	3.5.0-C API-1.9.0 r4084	1.11.4, released 2016/01/25	Rel. 4.9.2, September 8th, 2015
9.3.14	2.1.8 r13780	3.5.0-C API-1.9.0 r4084	1.11.5, released 2016/07/01	Rel. 4.9.2, September 8th, 2015
9.3.16	2.1.8 r13780	3.5.0-C API-1.9.0 r4084	1.11.5, released 2016/07/01	Rel. 4.9.2, September 8th, 2015

PostgreSQL	PostGIS	GEOS	GDAL	PROJ
9.3.17	2.1.8 r13780	3.5.0-CAPI-1.9.0 r4084	1.11.5, released 2016/07/01	Rel. 4.9.2, September 8th, 2015
9.4.7	2.1.8 r13780	3.5.0-CAPI-1.9.0 r4084	1.11.4, released 2016/01/25	Rel. 4.9.2, September 8th, 2015
9.4.9	2.1.8 r13780	3.5.0-CAPI-1.9.0 r4084	1.11.5, released 2016/07/01	Rel. 4.9.2, September 8th, 2015
9.4.11	2.1.8 r13780	3.5.0-CAPI-1.9.0 r4084	1.11.5, released 2016/07/01	Rel. 4.9.2, September 8th, 2015
9.4.12	2.1.8 r13780	3.5.0-CAPI-1.9.0 r4084	1.11.5, released 2016/07/01	Rel. 4.9.2, September 8th, 2015
9.5.2	2.2.2 r14797	3.5.0-CAPI-1.9.0 r4084	2.0.2, released 2016/01/26	Rel. 4.9.2, September 8th, 2015
9.5.4	2.2.2 r14797	3.5.0-CAPI-1.9.0 r4084	2.0.3, released 2016/07/01	Rel. 4.9.2, September 8th, 2015
9.5.6	2.2.5 r15298	3.5.1-CAPI-1.9.1 r4246	2.0.3, released 2016/07/01	Rel. 4.9.3, September 15th, 2016
9.5.7	2.2.5 r15298	3.5.1-CAPI-1.9.1 r4246	2.0.3, released 2016/07/01	Rel. 4.9.3, September 15th, 2016
9.6.1	2.3.0 r15146	3.5.0-CAPI-1.9.0 r4084	2.1.1, released 2016/07/07	Rel. 4.9.2, September 8th, 2016
9.6.2	2.3.2 r15302	3.5.1-CAPI-1.9.1 r4246	2.1.3, released 2017/20/01	Rel. 4.9.3, September 15th, 2016
9.6.3	2.3.2 r15302	3.5.1-CAPI-1.9.1 r4246	2.1.3, released 2017/20/01	Rel. 4.9.3, September 15th, 2016
9.6.6	2.3.4 r16009	3.6.2-CAPI-1.10.2 4d2925d6	2.1.3, released 2017/20/01	Rel. 4.9.3, September 15th, 2016
9.6.8	2.3.4 r16009	3.6.2-CAPI-1.10.2 4d2925d6	2.1.3, released 2017/20/01	Rel. 4.9.3, September 15th, 2016

PostgreSQL	PostGIS	GEOS	GDAL	PROJ
9.6.9	2.3.7 r16523	3.6.2-CAPI-1.10.2 4d2925d6	2.1.4, released 2017/06/23	Rel. 4.9.3, September 15th, 2016
9.6.10	2.3.7 r16523	3.6.2-CAPI-1.10.2 4d2925d6	2.1.4, released 2017/06/23	Rel. 4.9.3, September 15th, 2016
9.6.11	2.3.7 r16523	3.6.2-CAPI-1.10.2 4d2925d6	2.1.4, released 2017/06/23	Rel. 4.9.3, September 15th, 2016
10.1	2.4.2	3.6.2-CAPI-1.10.2 4d2925d6	2.1.3, released 2017/20/01	Rel. 4.9.3, September 15th, 2016
10.3	2.4.2	3.6.2-CAPI-1.10.2 4d2925d6	2.1.3, released 2017/20/01	Rel. 4.9.3, September 15th, 2016
10.4	2.4.4 r16526	3.6.2-CAPI-1.10.2 4d2925d6	2.1.4, released 2017/06/23	Rel. 4.9.3, September 15th, 2016
10.5	2.4.4 r16526	3.6.2-CAPI-1.10.2 4d2925d6	2.1.4, released 2017/06/23	Rel. 4.9.3, September 15th, 2016
10.6	2.4.4 r16526	3.6.2-CAPI-1.10.2 4d2925d6	2.1.4, released 2017/06/23	Rel. 4.9.3, September 15th, 2016
11.1	2.5.1 r17027	3.7.0-CAPI-1.11.0 673b9939	2.3.1, released 2018/06/22	Rel. 5.2.0, September 15th, 2018

Note

PostgreSQL 10.5 added support for the `libprotobuf` extension version 1.3.0 to the PostGIS component.

Using the log_fdw Extension

The `log_fdw` extension is new for Amazon RDS for PostgreSQL version 9.6.2 and later. Using this extension, you can access your database engine log using a SQL interface. In addition to viewing the `stderr` log files that are generated by default on RDS, you can view CSV logs (set the `log_destination` parameter to `csvlog`) and build foreign tables with the data neatly split into several columns.

This extension introduces two new functions that make it easy to create foreign tables for database logs:

- `list_postgres_log_files()` – Lists the files in the database log directory and the file size in bytes.
- `create_foreign_table_for_log_file(table_name text, server_name text, log_file_name text)` – Builds a foreign table for the specified file in the current database.

All functions created by `log_fdw` are owned by `rds_superuser`. Members of the `rds_superuser` role can grant access to these functions to other database users.

The following example shows how to use the `log_fdw` extension.

To use the `log_fdw` extension

1. Get the `log_fdw` extension.

```
postgres=> CREATE EXTENSION log_fdw;  
CREATE EXTENSION
```

2. Create the log server as a foreign data wrapper.

```
postgres=> CREATE SERVER log_server FOREIGN DATA WRAPPER log_fdw;  
CREATE SERVER
```

3. Select all from a list of log files.

```
postgres=> SELECT * from list_postgres_log_files() order by 1;
```

A sample response is as follows.

file_name	file_size_bytes
postgresql.log.2016-08-09-22.csv	1111
postgresql.log.2016-08-09-23.csv	1172
postgresql.log.2016-08-10-00.csv	1744
postgresql.log.2016-08-10-01.csv	1102

(4 rows)

4. Create a table with a single 'log_entry' column for non-CSV files.

```
postgres=> SELECT create_foreign_table_for_log_file('my_postgres_error_log',  
        'log_server', 'postgresql.log.2016-08-09-22.csv');
```

A sample response is as follows.

```
-----  
(1 row)
```

5. Select a sample of the log file. The following code retrieves the log time and error message description.

```
postgres=> SELECT log_time, message from my_postgres_error_log order by 1;
```

A sample response is as follows.

log_time		message

Tue Aug 09 15:45:18.172 2016 PDT		ending log output to stderr
Tue Aug 09 15:45:18.175 2016 PDT		database system was interrupted; last known up at 2016-08-09 22:43:34 UTC
Tue Aug 09 15:45:18.223 2016 PDT		checkpoint record is at 0/90002E0
Tue Aug 09 15:45:18.223 2016 PDT		redo record is at 0/90002A8; shutdown FALSE
Tue Aug 09 15:45:18.223 2016 PDT		next transaction ID: 0/1879; next OID: 24578
Tue Aug 09 15:45:18.223 2016 PDT		next MultiXactId: 1; next MultiXactOffset: 0
Tue Aug 09 15:45:18.223 2016 PDT		oldest unfrozen transaction ID: 1822, in database 1 (7 rows)

Upgrade PL/v8

If you use PL/v8 and upgrade PostgreSQL to a new PL/v8 version, you immediately take advantage of the new extension but the catalog metadata doesn't reflect this fact. The following steps synchronize your catalog metadata with the new version of PL/v8. These steps are optional but we highly recommended you complete them to avoid metadata mismatch warnings.

1. Verify that you need to update.

Run the following command while connected to your instance.

```
select * from pg_available_extensions where name in ('plv8','plls','plcoffee');
```

If your results contain values for an installed version that is a lower number than the default version, you should continue with this procedure to update your extensions.

For example, the following result set indicates you should update:

name	default_version	installed_version	comment

plls	2.1.0	1.5.3	PL/LiveScript (v8) trusted procedural language
plcoffee	2.1.0	1.5.3	PL/CoffeeScript (v8) trusted procedural language
plv8	2.1.0	1.5.3	PL/JavaScript (v8) trusted procedural language
(3 rows)			

2. Take a snapshot of your instance.

The upgrade drops all your PL/v8 functions. Take a snapshot of your instance as a precaution. You can continue with the following steps while the snapshot is being created.

For steps to create a snapshot see, [Creating a DB Snapshot \(p. 264\)](#)

3. Get a count of the functions you need to drop and recreate.

Obtain the count of the number of PL/v8 functions in your instance so you can validate that they are all in place after the upgrade.

The following code returns the number of functions written in PL/v8, plcoffee, or pll:

```
select proname, nspname, lanname
from pg_proc p, pg_language l, pg_namespace n
where p.prolang = l.oid
and n.oid = p.pronamespace
and lanname in ('plv8','plcoffee','pll');
```

4. Use pg_dump to create a schema-only dump file.

The following code creates a file on your client machine in the /tmp directory.

```
./pg_dump -Fc --schema-only -U master postgres > /tmp/test.dmp
```

This example uses the following flags:

- -Fc "format custom"
- --schema-only "will only dump commands necessary to create schema (functions in our case)"
- -U "rds master username"
- database "the database name in our instance"

For more information on pg_dump see, [pg_dump](#).

5. Extract the "CREATE FUNCTION" DDL statement that is present in the dump file.

The following code extracts the DDL statement needed to create the functions. You use this in subsequent steps to recreate the functions. The code uses the grep command to extract the statements to a file.

```
./pg_restore -l /tmp/test.dmp | grep FUNCTION > /tmp/function_list/
```

For more information on pg_restore see, [pg_restore](#).

6. Drop the functions and extensions.

The following code drops any PL/v8 based objects. The cascade option ensures that any dependent are dropped.

```
drop extension plv8 cascade;
```

If your PostgreSQL instance contains objects based on plcoffee or pll, repeat this step for those extensions.

7. Create the extensions.

The following code creates the PL/v8, plcoffee, and pll extensions:

```
create extension plv8;
create extension plcoffee;
create extension pll;
```

8. Create the functions using the dump file and "driver" file.

The following code re-creates the functions that you extracted previously.

```
./pg_restore -U master -d postgres -Fc -L /tmp/function_list /tmp/test.dmp
```

9. Verify your functions count.

Validate that your functions have all been re-creating by re-running the following code:

```
select * from pg_available_extensions where name in
('plv8','plls','plcoffee');
```

Note

PL/v8 version 2 adds the following extra row to your result set:

proname	nspname	lanname
plv8_version	pg_catalog	plv8

Supported PostgreSQL Features

Amazon RDS supports many of the most common PostgreSQL features. These include:

Topics

- [Logical Replication for PostgreSQL on Amazon RDS \(p. 1406\)](#)
- [Event Triggers for PostgreSQL on Amazon RDS \(p. 1408\)](#)
- [Huge Pages for Amazon RDS for PostgreSQL \(p. 1409\)](#)
- [Tablespaces for PostgreSQL on Amazon RDS \(p. 1409\)](#)
- [Autovacuum for PostgreSQL on Amazon RDS \(p. 1409\)](#)
- [RAM Disk for the stats_temp_directory \(p. 1409\)](#)
- [ALTER ENUM for PostgreSQL \(p. 1410\)](#)

Logical Replication for PostgreSQL on Amazon RDS

Beginning with PostgreSQL version 10.4, RDS supports the publication and subscription SQL Syntax for PostgreSQL 10 Logical Replication.

To enable logical replication for an Amazon RDS for PostgreSQL DB instance

1. The AWS user account requires the `rds_superuser` role to perform logical replication for the PostgreSQL database on Amazon RDS.
2. Set the `rds.logical_replication` static parameter to 1.
3. Modify the inbound rules of the security group for the publisher instance (production) to allow the subscriber instance (replica) to connect. This is usually done by including the IP address of the subscriber in the security group.
4. Restart the DB instance for the changes to the static parameter `rds.logical_replication` to take effect.

For more information on PostgreSQL logical replication, see the [PostgreSQL documentation](#).

Logical Decoding and Logical Replication

Beginning with PostgreSQL version 9.4, RDS supports the streaming of WAL changes using logical replication slots. Amazon RDS supports logical decoding for a PostgreSQL DB instance version 9.4.9 and higher and 9.5.4 and higher. You can set up logical replication slots on your instance and stream database changes through these slots to a client such as `pg_recvlogical`. Logical replication slots are created at the database level and support replication connections to a single database.

The most common clients for PostgreSQL logical replication are the AWS Database Migration Service or a custom-managed host on an AWS EC2 instance. The logical replication slot knows nothing about

the receiver of the stream, and there is no requirement that the target be a replica database. If you set up a logical replication slot and don't read from the slot, data can be written and quickly fill up your DB instance's storage.

PostgreSQL logical replication and logical decoding on Amazon RDS are enabled with a parameter, a replication connection type, and a security role. The client for logical decoding can be any client that is capable of establishing a replication connection to a database on a PostgreSQL DB instance.

To enable logical decoding for an Amazon RDS for PostgreSQL DB instance

1. The user account requires the **rds_superuser** role to enable logical replication. The user account also requires the **rds_replication** role to grant permissions to manage logical slots and to stream data using logical slots.
2. Set the **rds.logical_replication** static parameter to 1. As part of applying this parameter, we also set the parameters **wal_level**, **max_wal_senders**, **max_replication_slots**, and **max_connections**. These parameter changes can increase WAL generation, so you should only set the **rds.logical_replication** parameter when you are using logical slots.
3. Reboot the DB instance for the static **rds.logical_replication** parameter to take effect.
4. Create a logical replication slot as explained in the next section. This process requires that you specify a decoding plugin. Currently we support the **test_decoding** output plugin that ships with PostgreSQL.

For more information on PostgreSQL logical decoding, see the [PostgreSQL documentation](#).

Working with Logical Replication Slots

You can use SQL commands to work with logical slots. For example, the following command creates a logical slot named `test_slot` using the default PostgreSQL output plugin `test_decoding`.

```
SELECT * FROM pg_create_logical_replication_slot('test_slot', 'test_decoding');
```

The output should be similar to the following.

```
slot_name      | xlog_position
-----+-----
regression_slot | 0/16B1970
(1 row)
```

To list logical slots, use the following command.

```
SELECT * FROM pg_replication_slots;
```

To drop a logical slot, use the following command.

```
SELECT pg_drop_replication_slot('test_slot');
```

The output should be similar to the following.

```
pg_drop_replication_slot
-----
(1 row)
```

For more examples on working with logical replication slots, see [Logical Decoding Examples](#) in the PostgreSQL documentation.

Once you create the logical replication slot, you can start streaming. The following example shows how logical decoding is controlled over the streaming replication protocol, using the program pg_recvlogical included in the PostgreSQL distribution. This requires that client authentication is set up to allow replication connections.

```
pg_recvlogical -d postgres --slot test_slot -U master
--host sg-postgresql1.c6c8mresaghv0.us-west-2.rds.amazonaws.com
-f - --start
```

Event Triggers for PostgreSQL on Amazon RDS

PostgreSQL versions 9.4.9 and later and version 9.5.4 and later support event triggers, and Amazon RDS supports event triggers for these versions. The master user account can be used to create, modify, rename, and delete event triggers. Event triggers are at the DB instance level, so they can apply to all databases on an instance.

For example, the following code creates an event trigger that prints the current user at the end of every DDL command.

```
CREATE OR REPLACE FUNCTION raise_notice_func()
RETURNS event_trigger
LANGUAGE plpgsql AS
$$
BEGIN
    RAISE NOTICE 'In trigger function: %', current_user;
END;
$$;

CREATE EVENT TRIGGER event_trigger_1
    ON ddl_command_end
EXECUTE PROCEDURE raise_notice_func();
```

For more information about PostgreSQL event triggers, see [Event Triggers](#) in the PostgreSQL documentation.

There are several limitations to using PostgreSQL event triggers on Amazon RDS. These include:

- You cannot create event triggers on read replicas. You can, however, create event triggers on a read replica master. The event triggers are then copied to the read replica. The event triggers on the read replica don't fire on the read replica when changes are pushed from the master. However, if the read replica is promoted, the existing event triggers fire when database operations occur.
- To perform a major version upgrade to a PostgreSQL DB instance that uses event triggers, you must delete the event triggers before you upgrade the instance.

Huge Pages for Amazon RDS for PostgreSQL

Amazon RDS for PostgreSQL supports multiple page sizes for PostgreSQL versions 9.4.11 and later, 9.5.6 and later, and 9.6.2 and later. This support includes 4 K and 2 MB page sizes.

Huge pages reduce overhead when using large contiguous chunks of memory. You allocate huge pages for your application by using calls to *mmap* or SYSV shared memory. You enable huge pages on an Amazon RDS for PostgreSQL database by using the `huge_pages` parameter. Set this parameter to "on" to enable huge pages.

For PostgreSQL versions 10 and above, huge pages are enabled for all instance classes. For PostgreSQL versions below 10, huge pages are enabled by default for db.r4.*, db.m4.16xlarge, and db.m5.* instance classes. For other instance classes, huge pages are disabled by default.

When you set the `huge_pages` parameter to "on," Amazon RDS uses huge pages based on the available shared memory. If the DB instance is unable to use huge pages due to shared memory constraints, Amazon RDS prevents the instance from starting and sets the status of the DB instance to an incompatible parameters state. In this case, you can set the `huge_pages` parameter to "off" to allow Amazon RDS to start the DB instance.

The `shared_buffers` parameter is key to setting the shared memory pool that is required for using huge pages. The default value for the `shared_buffers` parameter is set to a percentage of the total 8K pages available for that instance's memory. When you use huge pages, those pages are allocated in the huge pages collocated together. Amazon RDS puts a DB instance into an incompatible parameters state if the shared memory parameters are set to require more than 90 percent of the DB instance memory. For more information about setting shared memory for PostgreSQL, see the [PostgreSQL documentation](#).

Note

Huge pages are not supported for the db.m1, db.m2, and db.m3 DB instance classes.

Tablespaces for PostgreSQL on Amazon RDS

Tablespaces are supported in PostgreSQL on Amazon RDS for compatibility; since all storage is on a single logical volume, tablespaces cannot be used for IO splitting or isolation. We have benchmarks and practical experience that shows that a single logical volume is the best setup for most use cases.

Autovacuum for PostgreSQL on Amazon RDS

The PostgreSQL autovacuum feature is turned on by default for new PostgreSQL DB instances. Autovacuum is optional, but we highly recommend that you do not turn autovacuum off. For more information on using autovacuum with Amazon RDS for PostgreSQL, see [Working with PostgreSQL Autovacuum on Amazon RDS \(p. 1332\)](#).

RAM Disk for the stats_temp_directory

The Amazon RDS for PostgreSQL parameter, `rds.pg_stat_ramdisk_size`, can be used to specify the system memory allocated to a RAM disk for storing the PostgreSQL `stats_temp_directory`. The RAM disk parameter is available for all PostgreSQL versions on Amazon RDS.

Under certain workloads, setting this parameter can improve performance and decrease IO requirements. For more information about the `stats_temp_directory`, see [the PostgreSQL documentation](#).

To enable a RAM disk for your `stats_temp_directory`, set the `rds.pg_stat_ramdisk_size` parameter to a non-zero value in the parameter group used by your DB instance. The parameter value is in MB. You must reboot the DB instance before the change takes effect.

For example, the following AWS CLI command sets the RAM disk parameter to 256 MB.

```
postgres=>aws rds modify-db-parameter-group \
    --db-parameter-group-name pg-95-ramdisk-testing \
    --parameters "ParameterName=rds.pg_stat_ramdisk_size, ParameterValue=256,
    ApplyMethod=pending-reboot"
```

After you reboot, run the following command to see the status of the `stats_temp_directory`:

```
postgres=>show stats_temp_directory;
```

The command should return the following:

```
stats_temp_directory
-----
/rdsdbramdisk/pg_stat_tmp
(1 row)
```

ALTER ENUM for PostgreSQL

Amazon RDS for PostgreSQL versions 9.6.2 and 9.5.6 and later support the ability to alter enumerations. This feature is not available in other versions on Amazon RDS.

The following code shows an example of altering an enum value.

```
postgres=> CREATE TYPE rainbow AS ENUM ('red', 'orange', 'yellow', 'green', 'blue',
    'purple');
CREATE TYPE
postgres=> CREATE TABLE t1 (colors rainbow);
CREATE TABLE
postgres=> INSERT INTO t1 VALUES ('red'), ('orange');
INSERT 0 2
postgres=> SELECT * from t1;
colors
-----
red
orange
(2 rows)
postgres=> ALTER TYPE rainbow RENAME VALUE 'red' TO 'crimson';
ALTER TYPE
postgres=> SELECT * from t1;
colors
-----
crimson
orange
(2 rows)
```

Limits for PostgreSQL DB Instances

The following is a list of limitations for PostgreSQL on Amazon RDS:

- You can have up to 40 PostgreSQL DB instances.
- For storage limits, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

- Amazon RDS reserves up to 3 connections for system maintenance. If you specify a value for the user connections parameter, you need to add 3 to the number of connections that you expect to use.

Upgrading a PostgreSQL DB Instance

There are two types of upgrades you can manage for your PostgreSQL DB instance:

- OS Updates – Occasionally, Amazon RDS might need to update the underlying operating system of your DB instance to apply security fixes or OS changes. You can decide when Amazon RDS applies OS updates by using the RDS console, AWS Command Line Interface (AWS CLI), or RDS API.

For more information about OS updates, see [Applying Updates for a DB Instance \(p. 154\)](#).

- Database Engine Upgrades – When Amazon RDS supports a new version of a database engine, you can upgrade your DB instances to the new version. There are two kinds of upgrades: major version upgrades and minor version upgrades. Amazon RDS supports both major and minor version upgrades for PostgreSQL DB instances.

For more information about PostgreSQL DB engine upgrades, see [Upgrading the PostgreSQL DB Engine for Amazon RDS \(p. 1289\)](#).

Using SSL with a PostgreSQL DB Instance

Amazon RDS supports Secure Socket Layer (SSL) encryption for PostgreSQL DB instances. Using SSL, you can encrypt a PostgreSQL connection between your applications and your PostgreSQL DB instances. You can also force all connections to your PostgreSQL DB instance to use SSL.

For general information about SSL support and PostgreSQL databases, see [SSL Support](#) in the PostgreSQL documentation. For information about using an SSL connection over JDBC, see [Configuring the Client](#) in the PostgreSQL documentation.

Topics

- [Requiring an SSL Connection to a PostgreSQL DB Instance \(p. 1412\)](#)
- [Determining the SSL Connection Status \(p. 1412\)](#)

SSL support is available in all AWS regions for PostgreSQL. Amazon RDS creates an SSL certificate for your PostgreSQL DB instance when the instance is created. If you enable SSL certificate verification, then the SSL certificate includes the DB instance endpoint as the Common Name (CN) for the SSL certificate to guard against spoofing attacks.

To connect to a PostgreSQL DB instance over SSL

1. Download the certificate.

For information about downloading certificates, see [Using SSL/TLS to Encrypt a Connection to a DB Instance \(p. 432\)](#).

2. Import the certificate into your operating system.
3. Connect to your PostgreSQL DB instance over SSL.

When you connect using SSL, your client can choose whether to verify the certificate chain. If your connection parameters specify `sslmode=verify-ca` or `sslmode=verify-full`, then your client requires the RDS CA certificates to be in their trust store or referenced in the connection URL. This requirement is to verify the certificate chain that signs your database certificate.

When a client, such as `psql` or JDBC, is configured with SSL support, the client first tries to connect to the database with SSL by default. If the client can't connect with SSL, it reverts to connecting

without SSL. The default `sslmode` mode used is different between libpq-based clients (such as `psql`) and JDBC. The libpq-based clients default to `prefer`, and JDBC clients default to `verify-full`.

Use the `sslrootcert` parameter to reference the certificate, for example `sslrootcert=rds-ssl-ca-cert.pem`.

The following is an example of using `psql` to connect to a PostgreSQL DB instance.

```
$ psql -h testpg.cdhmuqifdpib.us-east-1.rds.amazonaws.com -p 5432 \
  "dbname=testpg user=testuser sslrootcert=rds-ca-2015-root.pem sslmode=verify-full"
```

Requiring an SSL Connection to a PostgreSQL DB Instance

You can require that connections to your PostgreSQL DB instance use SSL by using the `rds.force_ssl` parameter. By default, the `rds.force_ssl` parameter is set to 0 (off). You can set the `rds.force_ssl` parameter to 1 (on) to require SSL for connections to your DB instance. Updating the `rds.force_ssl` parameter also sets the PostgreSQL `ssl` parameter to 1 (on) and modifies your DB instance's `pg_hba.conf` file to support the new SSL configuration.

You can set the `rds.force_ssl` parameter value by updating the parameter group for your DB instance. If the parameter group for your DB instance isn't the default one, and the `ssl` parameter is already set to 1 when you set `rds.force_ssl` to 1, you don't need to reboot your DB instance. Otherwise, you must reboot your DB instance for the change to take effect. For more information on parameter groups, see [Working with DB Parameter Groups \(p. 211\)](#).

When the `rds.force_ssl` parameter is set to 1 for a DB instance, you see output similar to the following when you connect, indicating that SSL is now required:

```
$ psql postgres -h SOMEHOST.amazonaws.com -p 8192 -U someuser
psql (9.3.12, server 9.4.4)
WARNING: psql major version 9.3, server major version 9.4.
Some psql features might not work.
SSL connection (cipher: DHE-RSA-AES256-SHA, bits: 256)
Type "help" for help.

postgres=>
```

Determining the SSL Connection Status

The encrypted status of your connection is shown in the logon banner when you connect to the DB instance:

```
Password for user master:
psql (9.3.12)
SSL connection (cipher: DHE-RSA-AES256-SHA, bits: 256)
Type "help" for help.

postgres=>
```

You can also load the `sslinfo` extension and then call the `ssl_is_used()` function to determine if SSL is being used. The function returns `t` if the connection is using SSL, otherwise it returns `f`.

```
postgres=> create extension sslinfo;
CREATE EXTENSION
```

```
postgres=> select ssl_is_used();
ssl_is_used
-----
t
(1 row)
```

You can use the `select ssl_cipher()` command to determine the SSL cipher:

```
postgres=> select ssl_cipher();
ssl_cipher
-----
DHE-RSA-AES256-SHA
(1 row)
```

If you enable `set rds.force_ssl` and restart your instance, non-SSL connections are refused with the following message:

```
$ export PGSSLMODE=disable
$ psql postgres -h SOMEHOST.amazonaws.com -p 8192 -U someuser
psql: FATAL: no pg_hba.conf entry for host "host.ip", user "someuser", database "postgres",
      SSL off
$
```

For information about the `sslmode` option, see [Database Connection Control Functions](#) in the PostgreSQL documentation.

Limits for Amazon RDS

This topic describes the resource limits and naming constraints for Amazon RDS.

Topics

- [Limits in Amazon RDS \(p. 1414\)](#)
- [Naming Constraints in Amazon RDS \(p. 1415\)](#)
- [File Size Limits in Amazon RDS \(p. 1417\)](#)

Limits in Amazon RDS

Each AWS account has limits, for each AWS Region, on the number of Amazon RDS resources that can be created. Once a limit for a resource has been reached, additional calls to create that resource fail with an exception.

The following table lists the resources and their limits per region.

Resource	Default Limit
Cross-region snapshots copy requests	5
DB Instances	40
Event subscriptions	20
Manual snapshots	100
Option groups	20
Parameter groups	50
Read replicas per master	5
Reserved instances	40
Rules per DB security group	20
Rules per VPC security group	50 inbound 50 outbound
DB Security groups	25
VPC Security groups	5
Subnet groups	50
Subnets per subnet group	20
Tags per resource	50
Total storage for all DB instances	100 TiB

Note

By default, you can have up to a total of 40 DB instances. Of those 40, you can have up to 10 instances of each SQL Server DB edition (Enterprise, Standard, Web, and Express) under

the "license-included" model. All 40 can be MySQL, MariaDB, or PostgreSQL. All 40 can be Oracle under the "bring-your-own-license" (BYOL) model. If your application requires more DB instances, you can request additional DB instances using this request form [Request RDS DB instance limit](#).

Backups managed by AWS Backup are considered manual snapshots for the manual snapshot limit. For information about AWS Backup, see the [AWS Backup Developer Guide](#).

Naming Constraints in Amazon RDS

The following table describes naming constraints in Amazon RDS.

DB instance identifier	<ul style="list-style-type: none"> Must contain 1 to 63 alphanumeric characters or hyphens. First character must be a letter. Cannot end with a hyphen or contain two consecutive hyphens. Must be unique for all DB instances per AWS account, per region.
Database name	<p>Database name constraints differ for each database engine.</p> <p>MySQL and MariaDB</p> <ul style="list-style-type: none"> Must contain 1 to 64 alphanumeric characters. Cannot be a word reserved by the database engine. <p>Oracle</p> <ul style="list-style-type: none"> Cannot be longer than 8 characters. <p>PostgreSQL</p> <ul style="list-style-type: none"> Must contain 1 to 63 alphanumeric characters. Must begin with a letter or an underscore. Subsequent characters can be letters, underscores, or digits (0-9). Cannot be a word reserved by the database engine. <p>SQL Server</p> <ul style="list-style-type: none"> Not applicable. For SQL Server, you create your databases after you create your DB instance. Database names follow the usual SQL Server naming rules with the following differences. Database names in SQL Server: <ul style="list-style-type: none"> Can't start with "rdsadmin" Can't start or end with a space or a tab Can't contain any of the characters that create a new line Can't contain a single quote
Master user name	<p>Master user name constraints differ for each database engine.</p> <p>MariaDB</p>

	<ul style="list-style-type: none">• Must contain 1 to 16 alphanumeric characters.• Cannot be a word reserved by the database engine.
	<p>MySQL</p> <ul style="list-style-type: none">• Must contain 1 to 16 alphanumeric characters.• First character must be a letter.• Cannot be a word reserved by the database engine.
	<p>Oracle</p> <ul style="list-style-type: none">• Must contain 1 to 30 alphanumeric characters.• First character must be a letter.• Cannot be a word reserved by the database engine.
	<p>PostgreSQL</p> <ul style="list-style-type: none">• Must contain 1 to 63 alphanumeric characters.• First character must be a letter.• Cannot be a word reserved by the database engine.
	<p>SQL Server</p> <ul style="list-style-type: none">• Must contain 1 to 64 alphanumeric characters.• First character must be a letter.• Cannot be a word reserved by the database engine.
Master password	<p>The password for the master database user can be any printable ASCII character except "/", "", or "@". Master password constraints differ for each database engine.</p> <p>MySQL and MariaDB</p> <ul style="list-style-type: none">• Must contain 8 to 41 characters. <p>Oracle</p> <ul style="list-style-type: none">• Must contain 8 to 30 characters. <p>PostgreSQL</p> <ul style="list-style-type: none">• Must contain 8 to 128 characters. <p>SQL Server</p> <ul style="list-style-type: none">• Must contain 8 to 128 characters.
DB parameter group name	<ul style="list-style-type: none">• Must contain from 1 to 255 alphanumeric characters.• First character must be a letter.• Hyphens are allowed, but the name cannot end with a hyphen or contain two consecutive hyphens.

DB subnet group name	<ul style="list-style-type: none">Must contain from 1 to 255 characters.Alphanumeric characters, spaces, hyphens, underscores, and periods are allowed.
----------------------	--

File Size Limits in Amazon RDS

File size limits apply to Amazon RDS DB instances.

MySQL File Size Limits in Amazon RDS

For Amazon RDS MySQL DB instances, the maximum provisioned storage limit constrains the size of a table to a maximum size of 16 TB when using InnoDB file-per-table tablespaces. This limit also constrains the system tablespace to a maximum size of 16 TB. InnoDB file-per-table tablespaces (with tables each in their own tablespace) are set by default for Amazon RDS MySQL DB instances. For more information, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

Note

Some existing DB instances have a lower limit. For example, MySQL DB instances created prior to April 2014 have a file and table size limit of 2 TB. This 2-TB file size limit also applies to DB instances or Read Replicas created from DB snapshots taken before April 2014, regardless of when the DB instance was created.

There are advantages and disadvantages to using InnoDB file-per-table tablespaces, depending on your application. To determine the best approach for your application, go to [InnoDB File-Per-Table Mode](#) in the MySQL documentation.

We don't recommend allowing tables to grow to the maximum file size. In general, a better practice is to partition data into smaller tables, which can improve performance and recovery times.

One option that you can use for breaking a large table up into smaller tables is partitioning. Partitioning distributes portions of your large table into separate files based on rules that you specify. For example, if you store transactions by date, you can create partitioning rules that distribute older transactions into separate files using partitioning. Then periodically, you can archive the historical transaction data that doesn't need to be readily available to your application. For more information, see [Partitioning](#) in the MySQL documentation.

To determine the file size of a table

Use the following SQL command to determine if any of your tables are too large and are candidates for partitioning. To update table statistics, issue an `ANALYZE TABLE` command on each table. For more information, see [ANALYZE TABLE](#) in the MySQL documentation.

```
SELECT TABLE_SCHEMA, TABLE_NAME,
       round(((DATA_LENGTH + INDEX_LENGTH) / 1024 / 1024), 2) AS "Approximate size (MB)",
       DATA_FREE
  FROM information_schema.TABLES
 WHERE TABLE_SCHEMA NOT IN ('mysql', 'information_schema', 'performance_schema');
```

To enable InnoDB file-per-table tablespaces

- To enable InnoDB file-per-table tablespaces, set the `innodb_file_per_table` parameter to 1 in the parameter group for the DB instance.

To disable InnoDB file-per-table tablespaces

- To disable InnoDB file-per-table tablespaces, set the *innodb_file_per_table* parameter to 0 in the parameter group for the DB instance.

For information on updating a parameter group, see [Working with DB Parameter Groups \(p. 211\)](#).

When you have enabled or disabled InnoDB file-per-table tablespaces, you can issue an `ALTER TABLE` command. You can use this command to move a table from the global tablespace to its own tablespace, or from its own tablespace to the global tablespace as shown in the following example.

```
ALTER TABLE table_name ENGINE=InnoDB, ALGORITHM=COPY;
```

MariaDB File Size Limits in Amazon RDS

For Amazon RDS MariaDB DB instances, the maximum provisioned storage limit constrains the size of a table to a maximum size of 16 TB when using InnoDB file-per-table tablespaces. This limit also constrains the system tablespace to a maximum size of 16 TB. InnoDB file-per-table tablespaces (with tables each in their own tablespace) is set by default for Amazon RDS MariaDB DB instances. For more information, see [Amazon RDS DB Instance Storage \(p. 135\)](#).

There are advantages and disadvantages to using InnoDB file-per-table tablespaces, depending on your application. To determine the best approach for your application, go to [InnoDB File-Per-Table Mode](#) in the MySQL documentation.

We don't recommend allowing tables to grow to the maximum file size. In general, a better practice is to partition data into smaller tables, which can improve performance and recovery times.

One option that you can use for breaking a large table up into smaller tables is partitioning. Partitioning distributes portions of your large table into separate files based on rules that you specify. For example, if you store transactions by date, you can create partitioning rules that distribute older transactions into separate files using partitioning. Then periodically, you can archive the historical transaction data that doesn't need to be readily available to your application. For more information, go to <https://dev.mysql.com/doc/refman/5.6/en/partitioning.html> in the MySQL documentation.

To determine the file size of a table

Use the following SQL command to determine if any of your tables are too large and are candidates for partitioning. To update table statistics, issue an `ANALYZE TABLE` command on each table. For more information, see [ANALYZE TABLE](#) in the MySQL documentation.

```
SELECT TABLE_SCHEMA, TABLE_NAME,
       round(((DATA_LENGTH + INDEX_LENGTH) / 1024 / 1024), 2) As "Approximate size (MB)",
       DATA_FREE
  FROM information_schema.TABLES
 WHERE TABLE_SCHEMA NOT IN ('mysql', 'information_schema', 'performance_schema');
```

To enable InnoDB file-per-table tablespaces

- To enable InnoDB file-per-table tablespaces, set the *innodb_file_per_table* parameter to 1 in the parameter group for the DB instance.

To disable InnoDB file-per-table tablespaces

- To disable InnoDB file-per-table tablespaces, set the *innodb_file_per_table* parameter to 0 in the parameter group for the DB instance.

For information on updating a parameter group, see [Working with DB Parameter Groups \(p. 211\)](#).

When you have enabled or disabled InnoDB file-per-table tablespaces, you can issue an `ALTER TABLE` command. You can use this command to move a table from the global tablespace to its own tablespace, or from its own tablespace to the global tablespace as shown in the following example.

```
ALTER TABLE table_name ENGINE=InnoDB, ALGORITHM=COPY;
```

Troubleshooting for Amazon RDS

Use the following sections to help troubleshoot problems you have with DB instances in Amazon RDS and Aurora.

Topics

- [Cannot Connect to Amazon RDS DB Instance \(p. 1420\)](#)
- [Amazon RDS Security Issues \(p. 1421\)](#)
- [Resetting the DB Instance Owner Role Password \(p. 1422\)](#)
- [Amazon RDS DB Instance Outage or Reboot \(p. 1422\)](#)
- [Amazon RDS DB Parameter Changes Not Taking Effect \(p. 1423\)](#)
- [Amazon RDS DB Instance Running Out of Storage \(p. 1423\)](#)
- [Amazon RDS Insufficient DB Instance Capacity \(p. 1424\)](#)
- [Amazon RDS MySQL and MariaDB Issues \(p. 1425\)](#)
- [Amazon RDS Oracle GoldenGate Issues \(p. 1430\)](#)
- [Cannot Connect to Amazon RDS SQL Server DB Instance \(p. 1431\)](#)
- [Cannot Connect to Amazon RDS PostgreSQL DB Instance \(p. 1431\)](#)
- [Cannot Set Backup Retention Period to 0 \(p. 1432\)](#)

For information about debugging problems using the Amazon RDS API, see [Troubleshooting Applications on Amazon RDS \(p. 1434\)](#).

Cannot Connect to Amazon RDS DB Instance

When you cannot connect to a DB instance, the following are common causes:

- The access rules enforced by your local firewall and the ingress IP addresses that you authorized to access your DB instance in the instance's security group aren't in sync. The problem is most likely the ingress rules in your security group.

By default, DB instances don't allow access; access is granted through a security group. To grant access, you must create your own security group with specific ingress and egress rules for your situation. If necessary, add rules to the security group associated with the VPC that allow traffic related to the source in and out of the DB instance. You can specify an IP address, a range of IP addresses, or another VPC security group.

For more information about setting up a security group, see [Provide Access to Your DB Instance in Your VPC by Creating a Security Group \(p. 8\)](#).

- The port you specified when you created the DB instance cannot be used to send or receive communications due to your local firewall restrictions. In this case, check with your network administrator to determine if your network allows the specified port to be used for inbound and outbound communication.
- Your DB instance is still being created and is not yet available. Depending on the size of your DB instance, it can take up to 20 minutes before an instance is available.

Testing a Connection to an Amazon RDS DB Instance

You can test your connection to a DB instance using common Linux or Windows tools.

From a Linux or Unix terminal, you can test the connection by typing the following (replace <DB-instance-endpoint> with the endpoint and <port> with the port of your DB instance):

```
nc -zv <DB-instance-endpoint> <port>
```

For example, the following shows a sample command and the return value:

```
nc -zv postgresql1.c6c8mn7tsdgv0.us-west-2.rds.amazonaws.com 8299
Connection to postgresql1.c6c8mn7tsdgv0.us-west-2.rds.amazonaws.com 8299 port [tcp/vvr-data] succeeded!
```

Windows users can use Telnet to test the connection to a DB instance. Note that Telnet actions are not supported other than for testing the connection. If a connection is successful, the action returns no message. If a connection is not successful, you receive an error message such as the following:

```
C:\>telnet sg-postgresql1.c6c8mntzhgv0.us-west-2.rds.amazonaws.com 819
Connecting To sg-postgresql1.c6c8mntzhgv0.us-west-2.rds.amazonaws.com...Could not open
connection to the host, on port 819: Connect failed
```

If Telnet actions return success, your security group is properly configured.

Note

Amazon RDS does not accept internet control message protocol (ICMP) traffic, including ping.

Troubleshooting Connection Authentication

If you can connect to your DB instance but you get authentication errors, you might want to reset the master user password for the DB instance. You can do this by modifying the RDS instance.

For more information about modifying a DB instance, see one of the following topics:

- [Modifying a DB Instance Running the MySQL Database Engine \(p. 802\)](#)
- [Modifying a DB Instance Running the Oracle Database Engine \(p. 966\)](#)
- [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#)
- [Modifying a DB Instance Running the PostgreSQL Database Engine \(p. 1279\)](#)

Amazon RDS Security Issues

To avoid security issues, never use your master AWS user name and password for a user account. Best practice is to use your master AWS account to create IAM users and assign those to DB user accounts. You can also use your master account to create other user accounts, if necessary.

For more information on creating IAM users, see [Create an IAM User \(p. 5\)](#).

Error Message "Failed to retrieve account attributes, certain console functions may be impaired."

There are several reasons you would get this error; it could be because your account is missing permissions, or your account has not been properly set up. If your account is new, you may not have

waited for the account to be ready. If this is an existing account, you could lack permissions in your access policies to perform certain actions such as creating a DB instance. To fix the issue, your IAM administrator needs to provide the necessary roles to your account. For more information, see [the IAM documentation](#).

Resetting the DB Instance Owner Role Password

You can reset the assigned permissions for your DB instance by resetting the master password. For example, if you lock yourself out of the `db_owner` role on your SQL Server database, you can reset the `db_owner` role password by modifying the DB instance master password. By changing the DB instance password, you can regain access to the DB instance, access databases using the modified password for the `db_owner`, and restore privileges for the `db_owner` role that may have been accidentally revoked. You can change the DB instance password by using the Amazon RDS console, the AWS CLI command [modify-db-instance](#), or by using the [ModifyDBInstance](#) operation.

For more information about modifying a SQL Server DB instance, see [Modifying a DB Instance Running the Microsoft SQL Server Database Engine \(p. 643\)](#).

Amazon RDS DB Instance Outage or Reboot

A DB instance outage can occur when a DB instance is rebooted, when the DB instance is put into a state that prevents access to it, and when the database is restarted. A reboot can occur when you manually reboot your DB instance or when you change a DB instance setting that requires a reboot before it can take effect.

When you modify a setting for a DB instance, you can determine when the change is applied by using the **Apply Immediately** setting.

To see a table that shows DB instance actions and the effect that setting the **Apply Immediately** value has, see [Modifying an Amazon RDS DB Instance \(p. 150\)](#).

A DB instance reboot only occurs when you change a setting that requires a reboot, or when you manually cause a reboot. A reboot can occur immediately if you change a setting and request that the change take effect immediately or it can occur during the DB instance's maintenance window.

A DB instance reboot occurs immediately when one of the following occurs:

- You change the backup retention period for a DB instance from 0 to a nonzero value or from a nonzero value to 0 and set **Apply Immediately** to *true*.
- You change the DB instance class, and **Apply Immediately** is set to *true*.
- You change the storage type from **Magnetic (Standard)** to **General Purpose (SSD)** or **Provisioned IOPS (SSD)**, or from **Provisioned IOPS (SSD)** or **General Purpose (SSD)** to **Magnetic (Standard)**. from standard to PIOPS.

A DB instance reboot occurs during the maintenance window when one of the following occurs:

- You change the backup retention period for a DB instance from 0 to a nonzero value or from a nonzero value to 0, and **Apply Immediately** is set to *false*.
- You change the DB instance class, and **Apply Immediately** is set to *false*.

When you change a static parameter in a DB parameter group, the change will not take effect until the DB instance associated with the parameter group is rebooted. The change requires a manual reboot; the DB instance will not automatically be rebooted during the maintenance window.

Amazon RDS DB Parameter Changes Not Taking Effect

If you change a parameter in a DB parameter group but you don't see the changes take effect, you most likely need to reboot the DB instance associated with the DB parameter group. When you change a dynamic parameter, the change takes effect immediately; when you change a static parameter, the change won't take effect until you reboot the DB instance associated with the parameter group.

You can reboot a DB instance using the RDS console or explicitly calling the [RebootDBInstance](#) API operation (without failover, if the DB instance is in a Multi-AZ deployment). The requirement to reboot the associated DB instance after a static parameter change helps mitigate the risk of a parameter misconfiguration affecting an API call, such as calling [ModifyDBInstance](#) to change DB instance class. For more information, see [Modifying Parameters in a DB Parameter Group \(p. 213\)](#).

Amazon RDS DB Instance Running Out of Storage

If your DB instance runs out of storage space, it might no longer be available. We highly recommend that you constantly monitor the `FreeStorageSpace` metric published in CloudWatch to ensure that your DB instance has enough free storage space.

If your database instance runs out of storage, its status will change to *storage-full*. For example, a call to the `DescribeDBInstances` action for a DB instance that has used up its storage will output the following:

```
aws rds describe-db-instances --db-instance-identifier mydbinstance

DBINSTANCE mydbinstance 2009-12-22T23:06:11.915Z db.m3.large mysql5.6 50 sa
storage-full mydbinstance.clla4j4jgyp.us-east-1.rds.amazonaws.com 3306
us-east-1b 3
SECGROUP default active
PARAMGRP default.mysql5.6 in-sync
```

To recover from this scenario, add more storage space to your instance using the `ModifyDBInstance` action or the following AWS CLI command:

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier mydbinstance \
--allocated-storage 60 \
--apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier mydbinstance ^
--allocated-storage 60 ^
--apply-immediately
```

```
DBINSTANCE mydbinstance 2009-12-22T23:06:11.915Z db.m3.large mysql5.6 50 sa
storage-full mydbinstance.clla4j4jgyp.us-east-1.rds.amazonaws.com 3306
us-east-1b 3 60
```

```
SECGRP default active
PARAMGRP default.mysql5.6 in-sync
```

Now, when you describe your DB instance, you will see that your DB instance will have *modifying* status, which indicates the storage is being scaled.

```
aws rds describe-db-instances --db-instance-identifier mydbinstance
```

```
DBINSTANCE mydbinstance 2009-12-22T23:06:11.915Z db.m3.large mysql5.6 50 sa
modifying mydbinstance.clla4j4jgyp.us-east-1.rds.amazonaws.com
3306 us-east-1b 3 60
SECGRP default active
PARAMGRP default.mysql5.6 in-sync
```

Once storage scaling is complete, your DB instance status will change to *available*.

```
aws rds describe-db-instances --db-instance-identifier mydbinstance
```

```
DBINSTANCE mydbinstance 2009-12-22T23:06:11.915Z db.m3.large mysql5.6 60 sa
available mydbinstance.clla4j4jgyp.us-east-1.rds.amazonaws.com 3306
us-east-1b 3
SECGRP default active
PARAMGRP default.mysql5.6 in-sync
```

Note that you can receive notifications when your storage space is exhausted using the `DescribeEvents` action. For example, in this scenario, if you do a `DescribeEvents` call after these operations you will see the following output:

```
aws rds describe-events --source-type db-instance --source-identifier mydbinstance
```

```
2009-12-22T23:44:14.374Z mydbinstance Allocated storage has been exhausted db-instance
2009-12-23T00:14:02.737Z mydbinstance Applying modification to allocated storage db-
instance
2009-12-23T00:31:54.764Z mydbinstance Finished applying modification to allocated storage
```

Amazon RDS Insufficient DB Instance Capacity

The `InsufficientDBInstanceCapacity` error is returned when you try to create or modify a DB instance, or when you try to restore a DB instance from a DB snapshot. When this error is returned, the following are common causes:

- The specific DB instance class is not available in the requested Availability Zone. You can try one of the following to solve the problem:
 - Retry the request with a different DB instance class.
 - Retry the request with a different Availability Zone.
 - Retry the request without specifying an explicit Availability Zone.

For information about troubleshooting instance capacity issues for Amazon EC2, see [Insufficient Instance Capacity](#) in the *Amazon Elastic Compute Cloud User Guide*.

- The DB instance is on the EC2-Classic platform and therefore is not in a VPC. Some DB instance classes require a VPC. For example, if you are on the EC2-Classic platform and try to increase capacity by

switching to a DB instance class that requires a VPC, this error results. For information about Amazon EC2 instance types that are only available in a VPC, see [Instance Types Available in EC2-Classic](#) in the *Amazon Elastic Compute Cloud User Guide*. To correct the problem, you can move the DB instance into a VPC. For more information, see [Moving a DB Instance Not in a VPC into a VPC \(p. 520\)](#).

For information about modifying a DB instance, see [Modifying an Amazon RDS DB Instance \(p. 150\)](#).

Amazon RDS MySQL and MariaDB Issues

You can diagnose and correct problems with MySQL and MariaDB DB instances.

Index Merge Optimization Returns Wrong Results

This issue applies only to MySQL DB instances.

Queries that use index merge optimization might return wrong results due to a bug in the MySQL query optimizer that was introduced in MySQL 5.5.37. When you issue a query against a table with multiple indexes the optimizer scans ranges of rows based on the multiple indexes, but does not merge the results together correctly. For more information on the query optimizer bug, go to <http://bugs.mysql.com/bug.php?id=72745> and <http://bugs.mysql.com/bug.php?id=68194> in the MySQL bug database.

For example, consider a query on a table with two indexes where the search arguments reference the indexed columns.

```
SELECT * FROM table1
  WHERE indexed_col1 = 'value1' AND indexed_col2 = 'value2';
```

In this case, the search engine searches both indexes. However, due to the bug, the merged results are incorrect.

To resolve this issue, you can do one of the following:

- Set the `optimizer_switch` parameter to `index_merge=off` in the DB parameter group for your MySQL DB instance. For information on setting DB parameter group parameters, see [Working with DB Parameter Groups \(p. 211\)](#).
- Upgrade your MySQL DB instance to MySQL version 5.6, 5.7, or 8.0. For more information, see [Upgrading a MySQL DB Snapshot \(p. 818\)](#).
- If you cannot upgrade your instance or change the `optimizer_switch` parameter, you can work around the bug by explicitly identifying an index for the query, for example:

```
SELECT * FROM table1
  USE INDEX covering_index
  WHERE indexed_col1 = 'value1' AND indexed_col2 = 'value2';
```

For more information, go to [Index Merge Optimization](#).

Diagnosing and Resolving Lag Between Read Replicas

After you create a MySQL or MariaDB Read Replica and the Read Replica is available, Amazon RDS first replicates the changes made to the source DB instance from the time the create Read Replica operation

was initiated. During this phase, the replication lag time for the Read Replica will be greater than 0. You can monitor this lag time in Amazon CloudWatch by viewing the Amazon RDS `ReplicaLag` metric.

The `ReplicaLag` metric reports the value of the `Seconds_Behind_Master` field of the MySQL or MariaDB `SHOW SLAVE STATUS` command. For more information, see [SHOW SLAVE STATUS](#). When the `ReplicaLag` metric reaches 0, the replica has caught up to the source DB instance. If the `ReplicaLag` metric returns -1, replication might not be active. To troubleshoot a replication error, see [Diagnosing and Resolving a MySQL or MariaDB Read Replication Failure \(p. 1427\)](#). A `ReplicaLag` value of -1 can also mean that the `Seconds_Behind_Master` value cannot be determined or is `NULL`.

The `ReplicaLag` metric returns -1 during a network outage or when a patch is applied during the maintenance window. In this case, wait for network connectivity to be restored or for the maintenance window to end before you check the `ReplicaLag` metric again.

Because the MySQL and MariaDB read replication technology is asynchronous, you can expect occasional increases for the `BinLogDiskUsage` metric on the source DB instance and for the `ReplicaLag` metric on the Read Replica. For example, a high volume of write operations to the source DB instance can occur in parallel, while write operations to the Read Replica are serialized using a single I/O thread, can lead to a lag between the source instance and Read Replica. For more information about Read Replicas and MySQL, go to [Replication Implementation Details](#) in the MySQL documentation. For more information about Read Replicas and MariaDB, go to [Replication Overview](#) in the MariaDB documentation.

You can reduce the lag between updates to a source DB instance and the subsequent updates to the Read Replica by doing the following:

- Set the DB instance class of the Read Replica to have a storage size comparable to that of the source DB instance.
- Ensure that parameter settings in the DB parameter groups used by the source DB instance and the Read Replica are compatible. For more information and an example, see the discussion of the `max_allowed_packet` parameter in the next section.
- Disable the query cache. For tables that are modified often, using the query cache can increase replica lag because the cache is locked and refreshed often. If this is the case, you might see less replica lag if you disable the query cache. You can disable the query cache by setting the `query_cache_type` parameter to 0 in the DB parameter group for the DB instance. For more information on the query cache, see [Query Cache Configuration](#).
- Warm the buffer pool on the Read Replica for InnoDB for MySQL, InnoDB for MariaDB 10.2 or higher, or XtraDB for MariaDB 10.1 or lower. If you have a small set of tables that are being updated often, and you are using the InnoDB or XtraDB table schema, then dump those tables on the Read Replica. Doing this causes the database engine to scan through the rows of those tables from the disk and then cache them in the buffer pool, which can reduce replica lag. The following shows an example.

For Linux, OS X, or Unix:

```
PROMPT> mysqldump \
-h <endpoint> \
--port=<port> \
-u=<username> \
-p <password> \
database_name table1 table2 > /dev/null
```

For Windows:

```
PROMPT> mysqldump ^
-h <endpoint> ^
--port=<port> ^
-u=<username> ^
-p <password> ^
```

```
database_name table1 table2 > /dev/null
```

Diagnosing and Resolving a MySQL or MariaDB Read Replication Failure

Amazon RDS monitors the replication status of your Read Replicas and updates the **Replication State** field of the Read Replica instance to **Error** if replication stops for any reason. You can review the details of the associated error thrown by the MySQL or MariaDB engines by viewing the **Replication Error** field. Events that indicate the status of the Read Replica are also generated, including [RDS-EVENT-0045 \(p. 375\)](#), [RDS-EVENT-0046 \(p. 375\)](#), and [RDS-EVENT-0047 \(p. 373\)](#). For more information about events and subscribing to events, see [Using Amazon RDS Event Notification \(p. 370\)](#). If a MySQL error message is returned, review the error in the [MySQL error message documentation](#). If a MariaDB error message is returned, review the error in the [MariaDB error message documentation](#).

Common situations that can cause replication errors include the following:

- The value for the `max_allowed_packet` parameter for a Read Replica is less than the `max_allowed_packet` parameter for the source DB instance.

The `max_allowed_packet` parameter is a custom parameter that you can set in a DB parameter group that is used to specify the maximum size of data manipulation language (DML) that can be executed on the database. If the `max_allowed_packet` parameter value for the source DB instance is smaller than the `max_allowed_packet` parameter value for the Read Replica, the replication process can throw an error and stop replication. The most common error is `packet bigger than 'max_allowed_packet' bytes`. You can fix the error by having the source and Read Replica use DB parameter groups with the same `max_allowed_packet` parameter values.

- Writing to tables on a Read Replica. If you are creating indexes on a Read Replica, you need to have the `read_only` parameter set to `0` to create the indexes. If you are writing to tables on the Read Replica, it can break replication.
- Using a non-transactional storage engine such as MyISAM. Read replicas require a transactional storage engine. Replication is only supported for the following storage engines: InnoDB for MySQL, InnoDB for MariaDB 10.2 or higher, or XtraDB for MariaDB 10.1 or lower.

You can convert a MyISAM table to InnoDB with the following command:

```
alter table <schema>.<table_name> engine=innodb;
```

- Using unsafe non-deterministic queries such as `SYSDATE()`. For more information, see [Determination of Safe and Unsafe Statements in Binary Logging](#).

The following steps can help resolve your replication error:

- If you encounter a logical error and you can safely skip the error, follow the steps described in [Skipping the Current Replication Error \(p. 881\)](#). Your MySQL or MariaDB DB instance must be running a version that includes the `mysql_rds_skip_repl_error` procedure. For more information, see [mysql.rds_skip_repl_error \(p. 904\)](#).
- If you encounter a binlog position issue, you can change the slave replay position with the `mysql_rds_next_master_log` command. Your MySQL or MariaDB DB instance must be running a version that supports the `mysql_rds_next_master_log` command in order to change the slave replay position. For version information, see [mysql.rds_next_master_log \(p. 905\)](#).
- If you encounter a temporary performance issue due to high DML load, you can set the `innodb_flush_log_at_trx_commit` parameter to 2 in the DB parameter group on the Read Replica. Doing this can help the Read Replica catch up, though it temporarily reduces atomicity, consistency, isolation, and durability (ACID).

- You can delete the Read Replica and create an instance using the same DB instance identifier so that the endpoint remains the same as that of your old Read Replica.

If a replication error is fixed, the **Replication State** changes to **replicating**. For more information, see [Troubleshooting a MySQL Read Replica Problem \(p. 856\)](#).

Creating Triggers with Binary Logging Enabled Requires SUPER Privilege

When trying to create triggers in an RDS MySQL or MariaDB DB instance, you might receive the following error:

```
"You do not have the SUPER privilege and binary logging is enabled"
```

To use triggers when binary logging is enabled requires the SUPER privilege, which is restricted for RDS MySQL and MariaDB DB instances. You can create triggers when binary logging is enabled without the SUPER privilege by setting the `log_bin_trust_function_creators` parameter to true. To set the `log_bin_trust_function_creators` to true, create a new DB parameter group or modify an existing DB parameter group.

To create a new DB parameter group that allows you to create triggers in your RDS MySQL or MariaDB DB instance with binary logging enabled, use the following CLI commands. To modify an existing parameter group, start with step 2.

To create a new parameter group to allow triggers with binary logging enabled using the CLI

1. Create a new parameter group.

For Linux, OS X, or Unix:

```
aws rds create-db-parameter-group \
--db-parameter-group-name allow-triggers \
--db-parameter-group-family mysql5.5 \
--description "parameter group allowing triggers"
```

For Windows:

```
aws rds create-db-parameter-group ^
--db-parameter-group-name allow-triggers ^
--db-parameter-group-family mysql5.5 ^
--description "parameter group allowing triggers"
```

2. Modify the DB parameter group to allow triggers.

For Linux, OS X, or Unix:

```
aws rds modify-db-parameter-group \
--db-parameter-group-name allow-triggers \
--parameters "name=log_bin_trust_function_creators,value=true, method=pending-reboot"
```

For Windows:

```
aws rds modify-db-parameter-group ^
```

```
--db-parameter-group-name allow-triggers ^
--parameters "name=log_bin_trust_function_creators,value=true, method=pending-
reboot"
```

3. Modify your DB instance to use the new DB parameter group.

For Linux, OS X, or Unix:

```
aws rds modify-db-instance \
--db-instance-identifier mydbinstance \
--db-parameter-group-name allow-triggers \
--apply-immediately
```

For Windows:

```
aws rds modify-db-instance ^
--db-instance-identifier mydbinstance ^
--db-parameter-group-name allow-triggers ^
--apply-immediately
```

4. In order for the changes to take effect, manually reboot the DB instance.

```
aws rds reboot-db-instance mydbinstance
```

Diagnosing and Resolving Point-In-Time Restore Failures

Restoring a DB Instance That Includes Temporary Tables

When attempting a Point-In-Time Restore (PITR) of your MySQL or MariaDB DB instance, you might encounter the following error:

```
Database instance could not be restored because there has been incompatible database
activity for restore
functionality. Common examples of incompatible activity include using temporary tables, in-
memory tables,
or using MyISAM tables. In this case, use of Temporary table was detected.
```

PITR relies on both backup snapshots and binlogs from MySQL or MariaDB to restore your DB instance to a particular time. Temporary table information can be unreliable in binlogs and can cause a PITR failure. If you use temporary tables in your MySQL or MariaDB DB instance, you can minimize the possibility of a PITR failure by performing more frequent backups. A PITR failure is most probable in the time between a temporary table's creation and the next backup snapshot.

Restoring a DB Instance That Includes In-Memory Tables

You might encounter a problem when restoring a database that has in-memory tables. In-memory tables are purged during a restart. As a result, your in-memory tables might be empty after a reboot. We recommend that when you use in-memory tables, you architect your solution to handle empty tables in the event of a restart. If you are using in-memory tables with replicated DB instances, you might need to recreate the Read Replicas after a restart if a Read Replica reboots and is unable to restore data from an empty in-memory table.

For more information about backups and PITR, see [Working With Backups \(p. 254\)](#) and [Restoring a DB Instance to a Specified Time \(p. 285\)](#).

Slave Down or Disabled Error

When you call the `mysql.rds_skip_repl_error` command, you might receive the following error message: `Slave is down or disabled.`

This error message appears because replication has stopped and could not be restarted.

If you need to skip a large number of errors, the replication lag can increase beyond the default retention period for binary log files. In this case, you might encounter a fatal error due to binary log files being purged before they have been replayed on the replica. This purge causes replication to stop, and you can no longer call the `mysql.rds_skip_repl_error` command to skip replication errors.

You can mitigate this issue by increasing the number of hours that binary log files are retained on your replication master. After you have increased the binlog retention time, you can restart replication and call the `mysql.rds_skip_repl_error` command as needed.

To set the binlog retention time, use the [mysql.rds_set_configuration \(p. 908\)](#) procedure and specify a configuration parameter of 'binlog retention hours' along with the number of hours to retain binlog files on the DB cluster, up to 720 (30 days). The following example sets the retention period for binlog files to 48 hours:

```
CALL mysql.rds_set_configuration('binlog retention hours', 48);
```

Read Replica Create Fails or Replication Breaks With Fatal Error 1236

After changing default parameter values for a MySQL or MariaDB DB instance, you might encounter one of the following problems:

- You are unable to create a Read Replica for the DB instance.
- Replication fails with `fatal error 1236`.

Some default parameter values for MySQL or MariaDB DB instances help to ensure the database is ACID compliant and Read Replicas are crash-safe by making sure that each commit is fully synchronized by writing the transaction to the binary log before it is committed. Changing these parameters from their default values to improve performance can cause replication to fail when a transaction has not been written to the binary log.

To resolve this issue, set the following parameter values:

- `sync-binlog = 1`
- `innodb_support_xa = 1`
- `innodb_flush_log_at_trx_commit = 1`

Amazon RDS Oracle GoldenGate Issues

Retaining Logs for Sufficient Time

The source database must retain archived redo logs. The duration for log retention is specified in hours. The duration should exceed any potential downtime of the source instance or any potential period of communication or networking issues for the source instance, so that Oracle GoldenGate can recover

logs from the source instance as needed. The absolute minimum value required is one (1) hour of logs retained. If you don't have log retention enabled, or if the retention value is too small, you will receive the following message:

```
2014-03-06 06:17:27  ERROR  OGG-00446  error 2 (No such file or directory)
opening redo log /rdsdbdata/db/GGTEST3_A/onlinelog/o1_mf_2_9k4bp1n6_.log
for sequence 1306Not able to establish initial position for begin time 2014-03-06
06:16:55.
```

Cannot Connect to Amazon RDS SQL Server DB Instance

When you have problems connecting to a DB instance using SQL Server Management Studio, the following are some common causes:

- The access rules enforced by your local firewall and the IP addresses you authorized to access your DB instance in the instance's security group are not in sync. If you use your DB instance's endpoint and port with Microsoft SQL Server Management Studio and cannot connect, the problem is most likely the inbound or outbound rules on your firewall.

To grant access, you can create your own security group with specific inbound and outbound rules for your situation. You can also modify an existing security group. A common issue is incorrect IP address ranges in inbound rules. Add or edit an inbound rule in the security group: for **Source**, choose **My IP**. This allows access to the DB instance from the IP address detected in your browser.

For more information about security groups, see [Controlling Access with Security Groups \(p. 486\)](#).

- The port you specified when you created the DB instance cannot be used to send or receive communications due to your local firewall restrictions. In this case, check with your network administrator to determine if your network allows the specified port to be used for inbound and outbound communication.
- Your DB instance is still being created and is not yet available. Depending on the size of your DB instance, it can take up to 20 minutes before an instance is available.

If you can send and receive communications through the port you specified, check for the following SQL Server errors:

- **Could not open a connection to SQL Server - Microsoft SQL Server, Error: 53** – You must include the port number when you specify the server name when using Microsoft SQL Server Management Studio. For example, the server name for a DB instance (including the port number) might be: **sqlsvr-pdz.c6c8mdfntzgv0.region.rds.amazonaws.com,1433**.
- **No connection could be made because the target machine actively refused it - Microsoft SQL Server, Error: 10061** – In this case, you reached the DB instance but the connection was refused. This error is often caused by an incorrect user name or password.

Cannot Connect to Amazon RDS PostgreSQL DB Instance

The most common problem when attempting to connect to a PostgreSQL DB instance is that the security group assigned to the DB instance has incorrect access rules. By default, DB instances do not allow access; access is granted through a security group. To grant access, you must create your own security

group with specific ingress and egress rules for your situation. For more information about creating a security group for your DB instance, see [Provide Access to Your DB Instance in Your VPC by Creating a Security Group \(p. 8\)](#).

The most common error is `could not connect to server: Connection timed out`. If you receive this error, check that the host name is the DB instance endpoint and that the port number is correct. Check that the security group assigned to the DB instance has the necessary rules to allow access through your local firewall.

Cannot Set Backup Retention Period to 0

There are several reasons why you may need to set the backup retention period to 0. For example, you can disable automatic backups immediately by setting the retention period to 0. If you set the value to 0 and receive a message saying that the retention period must be between 1 and 35, check to make sure you haven't setup a read replica for the instance. Read replicas require backups for managing read replica logs, thus, you can't set the retention period of 0.

Amazon RDS Application Programming Interface (API) Reference

In addition to the AWS Management Console, and the AWS Command Line Interface (AWS CLI), Amazon Relational Database Service (Amazon RDS) also provides an application programming interface (API). You can use the API to automate tasks for managing your DB instances and other objects in Amazon RDS.

- For an alphabetical list of API operations, see [Actions](#).
- For an alphabetical list of data types, see [Data Types](#).
- For a list of common query parameters, see [Common Parameters](#).
- For descriptions of the error codes, see [Common Errors](#).

For more information about the AWS CLI, see [AWS Command Line Interface Reference for Amazon RDS](#).

Topics

- [Using the Query API \(p. 1433\)](#)
- [Troubleshooting Applications on Amazon RDS \(p. 1434\)](#)

Using the Query API

The following sections discuss the parameters and request authentication used with the Query API.

Query Parameters

HTTP Query-based requests are HTTP requests that use the HTTP verb GET or POST and a Query parameter named `Action`.

Each Query request must include some common parameters to handle authentication and selection of an action.

Some operations take lists of parameters. These lists are specified using the `param.n` notation. Values of *n* are integers starting from 1.

For information about Amazon RDS regions and endpoints, go to [Amazon Relational Database Service \(RDS\)](#) in the Regions and Endpoints section of the *Amazon Web Services General Reference*.

Query Request Authentication

You can only send Query requests over HTTPS, and you must include a signature in every Query request. You must use either AWS signature version 4 or signature version 2. For more information, see [Signature Version 4 Signing Process](#) and [Signature Version 2 Signing Process](#).

Troubleshooting Applications on Amazon RDS

Amazon RDS provides specific and descriptive errors to help you troubleshoot problems while interacting with the Amazon RDS API.

Topics

- [Retrieving Errors \(p. 1434\)](#)
- [Troubleshooting Tips \(p. 1434\)](#)

For information about troubleshooting for Amazon RDS DB instances, see [Troubleshooting for Amazon RDS \(p. 1420\)](#).

Retrieving Errors

Typically, you want your application to check whether a request generated an error before you spend any time processing results. The easiest way to find out if an error occurred is to look for an `Error` node in the response from the Amazon RDS API.

XPath syntax provides a simple way to search for the presence of an `Error` node, as well as an easy way to retrieve the error code and message. The following code snippet uses Perl and the `XML::XPath` module to determine if an error occurred during a request. If an error occurred, the code prints the first error code and message in the response.

```
use XML::XPath;
my $xp = XML::XPath->new(xml =>$response);
if ( $xp->find("//Error") )
{print "There was an error processing your request:\n", " Error code: ",
$xp->findvalue("//Error[1]/Code"), "\n", " ",
$xp->findvalue("//Error[1]/Message"), "\n\n"; }
```

Troubleshooting Tips

We recommend the following processes to diagnose and resolve problems with the Amazon RDS API.

- Verify that Amazon RDS is operating normally in the AWS Region you are targeting by visiting <http://status.aws.amazon.com>.
- Check the structure of your request

Each Amazon RDS operation has a reference page in the *Amazon RDS API Reference*. Double-check that you are using parameters correctly. In order to give you ideas regarding what might be wrong, look at the sample requests or user scenarios to see if those examples are doing similar operations.

- Check the forum

Amazon RDS has a development community forum where you can search for solutions to problems others have experienced along the way. To view the forum, go to

<https://forums.aws.amazon.com/>

Document History

Current API version: 2014-10-31

The following table describes important changes in each release of the *Amazon RDS User Guide* after May 2018. For notification about updates to this documentation, you can subscribe to an RSS feed.

update-history-change	update-history-description	update-history-date
Amazon RDS Performance Insights supports viewing more SQL text for Amazon RDS for Microsoft SQL Server (p. 1435)	Amazon RDS Performance Insights now supports viewing more SQL text in the Performance Insights dashboard for Amazon RDS for Microsoft SQL Server DB instances. For more information, see Viewing More SQL Text in the Performance Insights Dashboard .	December 17, 2019
Amazon RDS on AWS Outposts (Preview) (p. 1435)	With Amazon RDS on AWS Outposts, you can create AWS-managed relational databases in your on-premises data centers. RDS on Outposts enables you to run RDS databases on AWS Outposts. For more information, see Amazon RDS on AWS Outposts (Preview) .	December 3, 2019
Amazon RDS Proxy (p. 1435)	You can reduce the overhead of connection management on your cluster, and reduce the chance of "too many connections" errors, by using the Amazon RDS Proxy. You associate each proxy with an RDS DB instance or Aurora DB cluster. Then you use the proxy endpoint in the connection string for your application. The Amazon RDS Proxy is currently in a public preview state. It supports the RDS MySQL database engine. For more information, see Managing Connections with Amazon RDS Proxy (Preview) .	December 3, 2019
Amazon RDS for Oracle supports cross-region Read Replicas (p. 1435)	Amazon RDS for Oracle now supports cross-region Read Replicas with Active Data Guard. For more information, see Working with Read Replicas and Working with Oracle Read Replicas .	November 26, 2019

Performance Insights supports analyzing statistics of running Oracle queries (p. 1435)	You can now analyze statistics of running queries with Performance Insights for Oracle DB instances. For more information, see Analyzing Statistics of Running Queries .	November 25, 2019
Amazon RDS for Microsoft SQL Server supports publishing logs to CloudWatch Logs (p. 1435)	You can configure your Amazon RDS SQL Server DB instance to publish log events directly to Amazon CloudWatch Logs. For more information, see Publishing SQL Server Logs to Amazon CloudWatch Logs .	November 25, 2019
Amazon RDS for Microsoft SQL Server supports new DB instance classes (p. 1435)	You can now create Amazon RDS DB instances running SQL Server that use the db.x1e and db.x1 DB instance classes. For more information, see DB Instance Class Support for Microsoft SQL Server .	November 25, 2019
Amazon RDS for Microsoft SQL Server supports differential and log restores (p. 1435)	You can restore differential backups and logs using SQL Server native backup and restore. For more information, see Using Native Backup and Restore .	November 25, 2019
Multi-AZ supported on Amazon RDS for Microsoft SQL Server in new regions (p. 1435)	Multi-AZ on SQL Server is now available in China, Middle East (Bahrain), and Europe (Stockholm). For more information, see Multi-AZ Deployments for Microsoft SQL Server .	November 22, 2019
Amazon RDS for Microsoft SQL Server now supports bulk insert and S3 integration (p. 1435)	You can transfer files between a SQL Server DB instance and an Amazon S3 bucket. Then you can use Amazon S3 with SQL Server features such as bulk insert. For more information, see Integrating an Amazon RDS for SQL Server DB Instance with Amazon S3 .	November 21, 2019
Amazon RDS for Oracle October 2019 RU, RUR, and PSU (p. 1435)	Amazon RDS for Oracle has released database engine updates for October 2019. For more information, see Oracle Database Engine Release Notes .	November 19, 2019

Performance Insights counters for Amazon RDS for Microsoft SQL Server (p. 1435)	You can now add performance counters to your Performance Insights charts for Microsoft SQL Server DB instances. For more information, see Performance Insights Counters for Amazon RDS for Microsoft SQL Server .	November 12, 2019
Amazon RDS for Microsoft SQL Server supports new DB instance class sizes (p. 1435)	You can now create Amazon RDS DB instances running SQL Server that use the 8xlarge and 16xlarge instance sizes for the db.m5 and db.r5 DB instance classes. Instance sizes ranging from small to 2xlarge are now available for the db.t3 instance class. For more information, see DB Instance Class Support for Microsoft SQL Server .	November 11, 2019
Support for PostgreSQL snapshot upgrades (p. 1435)	If you have existing manual DB snapshots of your Amazon RDS PostgreSQL DB instances, you can now upgrade them to a later version of the PostgreSQL database engine. For more information, see Upgrading a PostgreSQL DB Snapshot .	November 7, 2019
Amazon RDS for Oracle supports a new major version (p. 1435)	You can now create Amazon RDS DB instances running Oracle version 19.0. For more information, see Oracle 19c with Amazon RDS .	November 7, 2019
Amazon RDS for PostgreSQL version 12.0 in the Database Preview Environment (p. 1435)	Amazon RDS for PostgreSQL now supports PostgreSQL Version 12.0 in the Database Preview Environment. For more information, see PostgreSQL Version 12.0 in the Database Preview Environment .	November 1, 2019
Amazon RDS for PostgreSQL supports Kerberos authentication (p. 1435)	You can now use Kerberos authentication to authenticate users when they connect to your Amazon RDS DB instance running PostgreSQL. For more information, see Using Kerberos Authentication with Amazon RDS for PostgreSQL .	October 28, 2019

OEM Management Agent database tasks for Oracle DB instances (p. 1435)	Amazon RDS for Oracle DB instances now support procedures to invoke certain EMCTL commands on the Management Agent. For more information, see OEM Agent Database Tasks .	October 24, 2019
Amazon RDS for PostgreSQL versions 11.5, 10.10, 9.6.15, 9.5.19, and 9.4.24 (p. 1435)	Amazon RDS for PostgreSQL supports versions 11.5, 10.10, 9.6.15, 9.5.19, and 9.4.24. For more information, see Supported PostgreSQL Database Versions .	October 8, 2019
Amazon RDS for PostgreSQL supports PostgreSQL Transportable Databases (p. 1435)	PostgreSQL Transportable Databases provide an extremely fast method of migrating an RDS PostgreSQL database between two DB instances. For more information, see Transporting PostgreSQL Databases Between DB Instances .	October 8, 2019
Amazon RDS for Oracle supports Kerberos authentication (p. 1435)	You can now use Kerberos authentication to authenticate users when they connect to your Amazon RDS DB instance running Oracle. For more information, see Using Kerberos Authentication with Amazon RDS for Oracle .	September 30, 2019
Amazon RDS for PostgreSQL version 12 Beta 3 in the Database Preview Environment (p. 1435)	Amazon RDS for PostgreSQL now supports PostgreSQL Version 12 Beta 3 in the Database Preview Environment. For more information, see PostgreSQL Version 12 Beta 3 on Amazon RDS in the Database Preview Environment .	August 28, 2019
Support for MySQL 8.0.16 (p. 1435)	You can now create Amazon RDS DB instances running MySQL version 8.0.16. For more information, see MySQL on Amazon RDS Versions .	August 19, 2019
Amazon RDS for Oracle supports a new major version (p. 1435)	You can now create Amazon RDS DB instances running Oracle version 18.0. For more information, see Oracle 18c with Amazon RDS .	August 15, 2019

Amazon RDS for Oracle July 2019 RU, RUR, and PSU (p. 1435)	Amazon RDS for Oracle has released database engine version 12.2.0.1.ru-2019-07.rur-2019-07.r1 to support the July 2019 Release Update (RU) and Release Update Revision (RUR). Amazon RDS for Oracle has also released database engine versions 12.1.0.2.v17 and 11.2.0.4.v21 to support the July 2019 Oracle Database Patch Set Update (PSU). For more information, see Oracle Database Engine Release Notes .	August 8, 2019
Management Agent for OEM 13c Release 3 (p. 1435)	Amazon RDS for Oracle DB instances now support the Management Agent for Oracle Enterprise Manager (OEM) Cloud Control 13c Release 3. For more information, see Oracle Management Agent for Enterprise Manager Cloud Control .	August 7, 2019
Amazon RDS for PostgreSQL version 12 Beta 2 in the Database Preview Environment (p. 1435)	Amazon RDS for PostgreSQL now supports PostgreSQL Version 12 Beta 2 in the Database Preview Environment. For more information, see PostgreSQL Version 12 Beta 2 on Amazon RDS in the Database Preview Environment .	August 6, 2019
Amazon RDS supports server collations for SQL Server (p. 1435)	Amazon RDS for SQL Server supports a selection of collations for new DB instances. For more information, see Collations and Character Sets for Microsoft SQL Server .	July 29, 2019
Amazon RDS for PostgreSQL versions 11.4, 10.9, 9.6.14, 9.5.18, and 9.4.23 (p. 1435)	Amazon RDS for PostgreSQL supports minor versions 11.4, 10.9, 9.6.14, 9.5.18, and 9.4.23. For more information, see Supported PostgreSQL Database Versions .	July 3, 2019
Amazon RDS for Oracle supports Oracle APEX version 19.1.v1 (p. 1435)	Amazon RDS for Oracle now supports Oracle Application Express (APEX) version 19.1.v1. For more information, see Oracle Application Express .	June 28, 2019

Amazon RDS storage autoscaling (p. 1435)	Storage autoscaling for Amazon RDS DB instances enables Amazon RDS to automatically expand the storage associated with a DB instance to reduce the chance of out-of-space conditions. For information about storage autoscaling, see Working with Storage for Amazon RDS DB Instances .	June 20, 2019
Amazon RDS for Oracle supports db.z1d DB instance classes (p. 1435)	You can now create Amazon RDS DB instances running Oracle that use the db.z1d DB instance classes. For more information, see DB Instance Class .	June 13, 2019
Amazon RDS Performance Insights supports viewing more SQL text for Amazon RDS for Oracle (p. 1435)	Amazon RDS Performance Insights now supports viewing more SQL text in the Performance Insights dashboard for Amazon RDS for Oracle DB instances. For more information, see Viewing More SQL Text in the Performance Insights Dashboard .	June 10, 2019
Amazon RDS adds support native restores of SQL Server databases up to 16 TB (p. 1435)	You can now do native restores of up to 16 TB from SQL Server to Amazon RDS. For more information, see Amazon RDS for SQL Server: Limitations and Recommendations .	June 4, 2019
Amazon RDS adds support for Microsoft SQL Server Audit (p. 1435)	Using Amazon RDS for Microsoft SQL Server, you can audit server and database level events using SQL Server Audit, and view the results on your DB instance or send the audit log files directly to Amazon S3. For more information, see SQL Server Audit .	May 23, 2019
Improvements to Amazon RDS recommendations (p. 1435)	Amazon RDS has improved its automated recommendations for database resources. For example, Amazon RDS now provides recommendations for database parameters. For more information, see Using Amazon RDS Recommendations .	May 22, 2019

Support for more databases per DB instance for Amazon RDS SQL Server (p. 1435)	You can create up to 30 databases on each of your DB instances running Microsoft SQL Server. For more information, see Limits for Microsoft SQL Server DB Instances .	May 21, 2019
Support for 64 TiB and 80k IOPS of storage for Amazon RDS for MariaDB, MySQL and PostgreSQL (p. 1435)	You can now create Amazon RDS DB instances for MariaDB, MySQL and PostgreSQL with up to 64 TiB of storage and up to 80,000 provisioned IOPS. For more information, see DB Instance Storage .	May 20, 2019
Amazon RDS for MySQL supports upgrade prechecks (p. 1435)	When you upgrade a DB instance from MySQL 5.7 to MySQL 8.0, Amazon RDS performs prechecks for incompatibilities. For more information, see Prechecks for Upgrades from MySQL 5.7 to 8.0 .	May 17, 2019
Support for the MySQL password validation plugin (p. 1435)	You can now use the MySQL validate_password plugin for improved security of Amazon RDS for MySQL DB instances. For more information, see Using the Password Validation Plugin .	May 16, 2019
Amazon RDS for Oracle April 2019 RU, RUR, and PSU (p. 1435)	Amazon RDS for Oracle has released database engine version 12.2.0.1.ru-2019-04.rur-2019-04.r1 to support the April 2019 Release Update (RU) and Release Update Revision (RUR). Amazon RDS for Oracle has also released database engine versions 12.1.0.2.v16 and 11.2.0.4.v20 to support the April 2019 Oracle Database Patch Set Update (PSU). For more information, see Oracle Database Engine Release Notes .	May 16, 2019
Performance Insights counters for Amazon RDS for Oracle (p. 1435)	You can now add performance counters to your Performance Insights charts for Oracle DB instances. For more information, see Performance Insights Counters for Amazon RDS for Oracle .	May 8, 2019

Amazon RDS for PostgreSQL versions 11.2, 10.7, 9.6.12, 9.5.16, and 9.4.21 (p. 1435)	Amazon RDS for PostgreSQL supports minor versions 11.2, 10.7, 9.6.12, 9.5.16, and 9.4.21. For more information, see Supported PostgreSQL Database Versions .	May 1, 2019
Support for per-second billing (p. 1435)	Amazon RDS is now billed in 1-second increments in all AWS Regions except AWS GovCloud (US) for on-demand instances. For more information, see DB Instance Billing for Amazon RDS .	April 25, 2019
Support for importing data from Amazon S3 for Amazon RDS for PostgreSQL (p. 1435)	You can now import data from Amazon S3 file into a table in an RDS PostgreSQL DB instance. For more information, see Importing Amazon S3 Data into an RDS PostgreSQL DB Instance .	April 24, 2019
Support for restoring 5.7 backups from Amazon S3 (p. 1435)	You can now create a backup of your MySQL version 5.7 database, store it on Amazon S3, and then restore the backup file onto a new Amazon RDS DB instance running MySQL. For more information, see Restoring a Backup into an Amazon RDS MySQL DB Instance .	April 17, 2019
Support for multiple major version upgrades for Amazon RDS for PostgreSQL (p. 1435)	With Amazon RDS for PostgreSQL, you can now choose from multiple major versions when you upgrade the DB engine. This feature enables you to skip ahead to a newer major version when you upgrade select PostgreSQL engine versions. For more information, see Upgrading the PostgreSQL DB Engine .	April 16, 2019
Support for 64 TiB of storage for Amazon RDS for Oracle (p. 1435)	You can now create Amazon RDS DB instances for Oracle with up to 64 TiB of storage and up to 80,000 provisioned IOPS. For more information, see DB Instance Storage .	April 4, 2019
Support for MySQL 8.0.15 (p. 1435)	You can now create Amazon RDS DB instances running MySQL version 8.0.15. For more information, see MySQL on Amazon RDS Versions .	April 3, 2019

Support for MariaDB 10.3.13 (p. 1435)	You can now create Amazon RDS DB instances running MariaDB version 10.3.13. For more information, see MariaDB on Amazon RDS Versions .	April 3, 2019
Microsoft SQL Server 2008 R2 deprecated on Amazon RDS (p. 1435)	Support for Microsoft SQL Server 2008 R2 is deprecated, coinciding with the Microsoft plan to terminate extended support for this version on July 9, 2019. Any existing Microsoft SQL Server 2008 R2 snapshots are to be automatically upgraded to the latest minor version of Microsoft SQL Server 2012 starting on June 1, 2019. For more information, see Microsoft SQL Server 2008 R2 Support on Amazon RDS .	April 2, 2019
Always On Availability Groups supported in Microsoft SQL Server 2017 (p. 1435)	You can now use Always On Availability Groups in SQL Server 2017 Enterprise Edition 14.00.3049.1 or later. For more information, see Multi-AZ Deployments for Microsoft SQL Server .	March 29, 2019
View volume metrics (p. 1435)	You can now view metrics for the Amazon Elastic Block Store (Amazon EBS) volumes, which are the physical devices used for database and log storage. For more information, see Viewing Enhanced Monitoring .	March 20, 2019
Support for MySQL 5.7.25 (p. 1435)	You can now create Amazon RDS DB instances running MySQL version 5.7.25. For more information, see MySQL on Amazon RDS Versions .	March 19, 2019
Amazon RDS for Oracle supports RMAN DBA tasks (p. 1435)	Amazon RDS for Oracle now supports Oracle Recovery Manager (RMAN) DBA tasks, including RMAN backups. For more information, see Common DBA Recovery Manager (RMAN) Tasks for Oracle DB Instances .	March 14, 2019
Amazon RDS for PostgreSQL supports version 11.1 (p. 1435)	You can now create Amazon RDS DB instances running PostgreSQL version 11.1. For more information, see PostgreSQL Version 11.1 on Amazon RDS .	March 12, 2019

Multiple-file restore is available in Amazon RDS for SQL Server (p. 1435)	You can now restore from multiple files with Amazon RDS for SQL Server. For more information, see Restoring a Database .	March 11, 2019
MariaDB 10.2.21 (p. 1435)	You can now create Amazon RDS DB instances running MariaDB version 10.2.21. For more information, see MariaDB on Amazon RDS Versions .	March 11, 2019
Amazon RDS for Oracle supports Read Replicas (p. 1435)	Amazon RDS for Oracle now supports Read Replicas with Active Data Guard. For more information, see Working with Read Replicas and Working with Oracle Read Replicas .	March 11, 2019
Amazon RDS Performance Insights is available for Amazon RDS for MariaDB (p. 1435)	Amazon RDS Performance Insights is now available for Amazon RDS for MariaDB. For more information, see Using Amazon RDS Performance Insights .	March 11, 2019
MySQL 8.0.13 and 5.7.24 (p. 1435)	You can now create Amazon RDS DB instances running MySQL versions 8.0.13 and 5.7.24. For more information, see MySQL on Amazon RDS Versions .	March 8, 2019
Amazon RDS Performance Insights is available for Amazon RDS for SQL Server (p. 1435)	Amazon RDS Performance Insights is now available for Amazon RDS for SQL Server. For more information, see Using Amazon RDS Performance Insights .	March 4, 2019
Amazon RDS for Oracle supports Amazon S3 integration (p. 1435)	You can now transfer files between an Amazon RDS for Oracle DB instance and an Amazon S3 bucket. For more information, see Integrating Amazon RDS for Oracle and Amazon S3 .	February 26, 2019
Amazon RDS for MySQL and Amazon RDS for MariaDB support db.t3 DB instance classes (p. 1435)	You can now create Amazon RDS DB instances running MySQL or MariaDB that use the db.t3 DB instance classes. For more information, see DB Instance Class .	February 20, 2019

Amazon RDS for MySQL and Amazon RDS for MariaDB support db.r5 DB instance classes (p. 1435)	You can now create Amazon RDS DB instances running MySQL or MariaDB that use the db.r5 DB instance classes. For more information, see DB Instance Class .	February 20, 2019
Performance Insights counters for Amazon RDS MySQL and PostgreSQL (p. 1435)	You can now add performance counters to your Performance Insights charts for Amazon RDS MySQL and PostgreSQL DB instances. For more information, see Performance Insights Dashboard Components .	February 19, 2019
Amazon RDS for PostgreSQL now supports adaptive autovacuum parameter tuning (p. 1435)	Adaptive autovacuum parameter tuning with Amazon RDS for PostgreSQL helps prevent transaction ID wraparound by adjusting autovacuum parameter values automatically. For more information, see Reducing the Likelihood of Transaction ID Wraparound .	February 12, 2019
Amazon RDS for Oracle supports Oracle APEX versions 18.1.v1 and 18.2.v1 (p. 1435)	Amazon RDS for Oracle now supports Oracle Application Express (APEX) versions 18.1.v1 and 18.2.v1. For more information, see Oracle Application Express .	February 11, 2019
Amazon RDS for Oracle January 2019 RU, RUR, and PSU (p. 1435)	Amazon RDS for Oracle has released database engine version 12.2.0.1.ru-2019-01.rur-2019-01.r1 to support the January 2019 Release Update (RU) and Release Update Revision (RUR). Amazon RDS for Oracle has also released database engine versions 12.1.0.2.v15 and 11.2.0.4.v19 to support the January 2019 Oracle Database Patch Set Update (PSU). For more information, see Oracle Database Engine Release Notes .	February 8, 2019
Amazon RDS Performance Insights supports viewing more SQL text for Amazon RDS MySQL (p. 1435)	Amazon RDS Performance Insights now supports viewing more SQL text in the Performance Insights dashboard for Amazon RDS MySQL DB instances. For more information, see Viewing More SQL Text in the Performance Insights Dashboard .	February 6, 2019

Amazon RDS for PostgreSQL supports db.t3 DB instance classes (p. 1435)	You can now create Amazon RDS DB instances running PostgreSQL that use the db.t3 DB instance classes. For more information, see DB Instance Class .	January 25, 2019
Amazon RDS for Oracle supports db.t3 DB instance classes (p. 1435)	You can now create Amazon RDS DB instances running Oracle that use the db.t3 DB instance classes. For more information, see DB Instance Class .	January 25, 2019
Amazon RDS Performance Insights supports viewing more SQL text for Amazon RDS PostgreSQL (p. 1435)	Amazon RDS Performance Insights now supports viewing more SQL text in the Performance Insights dashboard for Amazon RDS PostgreSQL DB instances. For more information, see Viewing More SQL Text in the Performance Insights Dashboard .	January 24, 2019
Amazon RDS for Oracle supports a new version of SQLT (p. 1435)	Amazon RDS for Oracle now supports SQLT version 12.2.180725. For more information, see Oracle SQLT .	January 22, 2019
Amazon RDS for PostgreSQL supports new minor versions (p. 1435)	Amazon RDS for PostgreSQL now supports the following new minor versions: 10.6, 9.6.11, 9.5.15, 9.4.20, and 9.3.25. For more information, see Amazon RDS for PostgreSQL Versions and Extensions .	December 19, 2018
Amazon RDS for PostgreSQL supports db.r5 DB instance classes (p. 1435)	You can now create Amazon RDS DB instances running PostgreSQL that use the db.r5 DB instance classes. For more information, see DB Instance Class .	December 19, 2018
Amazon RDS for PostgreSQL now supports restricted password management (p. 1435)	Amazon RDS for PostgreSQL enables you to restrict who can manage user passwords and password expiration changes by using the parameter <code>rds.restrict_password_commands</code> and the role <code>rds_password</code> . For more information, see Restricting Password Management .	December 19, 2018

Amazon RDS for PostgreSQL supports uploading database logs to Amazon CloudWatch Logs (p. 1435)	Amazon RDS for PostgreSQL supports uploading database logs to CloudWatch Logs. For more information, see Publishing PostgreSQL Logs to CloudWatch Logs .	December 10, 2018
Amazon RDS for Oracle supports db.r5 DB instance classes (p. 1435)	You can now create Amazon RDS DB instances running Oracle that use the db.r5 DB instance classes. For more information, see DB Instance Class .	November 20, 2018
Retain backups when deleting a DB instance (p. 1435)	Amazon RDS supports retaining automated backups when you delete a DB instance. For more information, see Working with Backups .	November 15, 2018
Amazon RDS for PostgreSQL supports db.m5 DB instance classes (p. 1435)	You can now create Amazon RDS DB instances running PostgreSQL that use the db.m5 DB instance classes. For more information, see DB Instance Class .	November 15, 2018
Amazon RDS for Oracle supports a new major version (p. 1435)	You can now create Amazon RDS DB instances running Oracle version 12.2. For more information, see Oracle 12c Version 12.2.0.1 with Amazon RDS .	November 13, 2018
Amazon RDS for Oracle October 2018 PSU (p. 1435)	Amazon RDS for Oracle has released database engine versions 12.1.0.2.v14 and 11.2.0.4.v18 to support the October 2018 Oracle Database Patch Set Update (PSU). For more information, see Oracle Database Engine Release Notes .	November 13, 2018
Amazon RDS for SQL Server supports Always On (p. 1435)	Amazon RDS for SQL Server supports Always On Availability Groups. For more information, see Multi-AZ Deployments for Microsoft SQL Server .	November 8, 2018
Amazon RDS for PostgreSQL supports outbound network access using custom DNS servers (p. 1435)	Amazon RDS for PostgreSQL supports outbound network access using custom DNS servers. For more information, see Using a Custom DNS Server for Outbound Network Access .	November 8, 2018

Amazon RDS for MariaDB, MySQL, and PostgreSQL supports 32 TiB of storage (p. 1435)	You can now create Amazon RDS DB instances with up to 32 TiB of storage for MySQL, MariaDB, and PostgreSQL. For more information, see DB Instance storage .	November 7, 2018
Amazon RDS for Oracle supports extended data types (p. 1435)	You can now enable extended data types on Amazon RDS DB instances running Oracle. With extended data types, the maximum size is 32,767 bytes for the VARCHAR2, NVARCHAR2, and RAW data types. For more information, see Using Extended Data Types .	November 6, 2018
Amazon RDS for Oracle supports db.m5 DB instance classes (p. 1435)	You can now create Amazon RDS DB instances running Oracle that use the db.m5 DB instance classes. For more information, see DB Instance Class .	November 2, 2018
Amazon RDS for Oracle migration from SE, SE1, or SE2 to EE (p. 1435)	You can now migrate from any Oracle Database Standard Edition (SE, SE1, or SE2) to Oracle Database Enterprise Edition (EE). For more information, see Migrating Between Oracle Editions .	October 31, 2018
Amazon RDS can now stop Multi-AZ instances (p. 1435)	Amazon RDS can now stop a DB instance that is part of a Multi-AZ deployment. Formerly, the stop instance feature had a limitation for multi-AZ instances. For more information, see Stopping an Amazon RDS DB Instance Temporarily .	October 29, 2018
Amazon RDS Performance Insights is available for Amazon RDS Oracle (p. 1435)	Amazon RDS Performance Insights is now available for Amazon RDS Oracle. For more information, see Using Amazon RDS Performance Insights .	October 29, 2018
Amazon RDS for PostgreSQL supports PostgreSQL version 11 in the Database Preview Environment (p. 1435)	Amazon RDS for PostgreSQL now supports PostgreSQL version 11 in the Database Preview Environment. For more information, see PostgreSQL Version 11 on Amazon RDS in the Database Preview Environment .	October 25, 2018

MySQL supports a new major version (p. 1435)	You can now create Amazon RDS DB instances running MySQL version 8.0. For more information, see MySQL on Amazon RDS Versions .	October 23, 2018
MariaDB supports a new major version (p. 1435)	You can now create Amazon RDS DB instances running MariaDB version 10.3. For more information, see MariaDB on Amazon RDS Versions .	October 23, 2018
Amazon RDS for Oracle supports Oracle JVM (p. 1435)	Amazon RDS for Oracle now supports the Oracle Java Virtual Machine (JVM) option. For more information, see Oracle Java Virtual Machine .	October 16, 2018
Custom parameter group for restore and point in time recovery (p. 1435)	You can now specify a custom parameter group when you restore a snapshot or perform a point in time recovery operation. For more information, see Restoring from a DB Snapshot and Restoring a DB Instance to a Specified Time .	October 15, 2018
Amazon RDS for Oracle supports 32 TiB storage (p. 1435)	You can now create Oracle RDS DB instances with up to 32 TiB of storage. For more information, see DB instance storage .	October 15, 2018
Amazon RDS for MySQL supports GTIDs (p. 1435)	Amazon RDS for MySQL now supports global transaction identifiers (GTIDs), which are unique across all DB instances and in a replication configuration. For more information, see Using GTID-Based Replication for Amazon RDS MySQL .	October 10, 2018
MySQL 5.7.23, 5.6.41, and 5.5.61 (p. 1435)	You can now create Amazon RDS DB instances running MySQL versions 5.7.23, 5.6.41, and 5.5.61. For more information, see MySQL on Amazon RDS Versions .	October 8, 2018
Amazon RDS for PostgreSQL supports new minor versions (p. 1435)	Amazon RDS for PostgreSQL now supports the following new minor versions: 10.5, 9.6.10, 9.5.14, 9.4.19, and 9.3.24. For more information, see Amazon RDS for PostgreSQL Versions and Extensions .	October 4, 2018

Amazon RDS for Oracle supports a new version of SQLT (p. 1435)	Amazon RDS for Oracle now supports SQLT version 12.2.180331. For more information, see Oracle SQLT .	October 4, 2018
Amazon RDS for Oracle July 2018 PSU (p. 1435)	Amazon RDS for Oracle has released database engine versions 12.1.0.2.v13 and 11.2.0.4.v17 to support the July 2018 Oracle Database Patch Set Update (PSU). For more information, see Oracle Database Engine Release Notes .	October 3, 2018
Amazon RDS for PostgreSQL now supports IAM authentication (p. 1435)	Amazon RDS for PostgreSQL now supports IAM authentication. For more information see IAM Database Authentication for MySQL and PostgreSQL .	September 27, 2018
You can enable deletion protection for your Amazon RDS DB instances (p. 1435)	When you enable deletion protection for a DB instance, the database cannot be deleted by any user. For more information, see Deleting a DB Instance .	September 26, 2018
Amazon RDS for MySQL and Amazon RDS for MariaDB support db.m5 DB instance classes (p. 1435)	You can now create Amazon RDS DB instances running MySQL or MariaDB that use the db.m5 DB instance classes. For more information, see DB Instance Class .	September 18, 2018
Amazon RDS now supports upgrades to SQL Server 2017 (p. 1435)	You can upgrade your existing DB instance to SQL Server 2017 from any version except SQL Server 2008. To upgrade from SQL Server 2008, first upgrade to one of the other versions first. For information, see Upgrading the Microsoft SQL Server DB Engine .	September 11, 2018
Amazon RDS for PostgreSQL now supports PostgreSQL Version 11 Beta 3 in the Database Preview Environment (p. 1435)	In this release, the Write-Ahead Log (WAL) segment size (<code>wal_segment_size</code>) is now set to 64MB. For more about PostgreSQL version 11 Beta 3, see PostgreSQL 11 Beta 3 Released . For information on the Database Preview Environment, see Working with the Database Preview Environment .	September 7, 2018

Amazon Aurora User Guide (p. 1435)	The <i>Amazon Aurora User Guide</i> describes all Amazon Aurora concepts and provides instructions on using the various features with both the console and the command line interface. The <i>Amazon RDS User Guide</i> now covers non-Aurora database engines.	August 31, 2018
Amazon RDS Performance Insights is available for Amazon RDS MySQL (p. 1435)	Amazon RDS Performance Insights is now available for Amazon RDS MySQL. For more information, see Using Amazon RDS Performance Insights .	August 28, 2018
Aurora with PostgreSQL compatibility now supports Aurora Auto Scaling (p. 1435)	Auto Scaling of Aurora replicas is now available for Aurora with PostgreSQL compatibility. For more information, see Using Amazon Aurora Auto Scaling with Aurora Replicas .	August 16, 2018
Aurora Serverless for Aurora MySQL (p. 1435)	Aurora Serverless is an on-demand, autoscaling configuration for Amazon Aurora. For more information, see Using Amazon Aurora Serverless .	August 9, 2018
MySQL 5.7.22 and 5.6.40 (p. 1435)	You can now create Amazon RDS DB instances running MySQL versions 5.7.22 and 5.6.40. For more information, see MySQL on Amazon RDS Versions .	August 6, 2018
Aurora is now available in the China (Ningxia) region (p. 1435)	Aurora MySQL and Aurora PostgreSQL are now available in the China (Ningxia) region. For more information, see Availability for Amazon Aurora MySQL and Availability for Amazon Aurora PostgreSQL .	August 6, 2018
Amazon RDS for MySQL supports delayed replication (p. 1435)	Amazon RDS for MySQL now supports delayed replication as a strategy for disaster recovery. For more information, see Configuring Delayed Replication with MySQL .	August 6, 2018
Amazon RDS Performance Insights is available for Aurora MySQL (p. 1435)	Amazon RDS Performance Insights is now available for Aurora MySQL. For more information, see Using Amazon RDS Performance Insights .	August 6, 2018

Amazon RDS Performance Insights integration with Amazon CloudWatch (p. 1435)	Amazon RDS Performance Insights automatically publishes metrics to Amazon CloudWatch. For more information, see Performance Insights Metrics Published to CloudWatch .	August 6, 2018
Amazon RDS recommendations (p. 1435)	Amazon RDS now provides automated recommendations for database resources. For more information, see Using Amazon RDS Recommendations .	July 25, 2018
Amazon RDS for PostgreSQL supports new minor versions (p. 1435)	Amazon RDS for PostgreSQL now supports the following new minor versions: 10.4, 9.6.9, 9.5.13, 9.4.18, and 9.3.23. For more information, see Amazon RDS for PostgreSQL Versions and Extensions .	July 25, 2018
Incremental snapshot copies across AWS Regions (p. 1435)	Amazon RDS supports incremental snapshot copies across AWS Regions for both unencrypted and encrypted instances. For more information, see Copying Snapshots Across AWS Regions .	July 24, 2018
Amazon RDS Performance Insights is available for Amazon RDS for PostgreSQL (p. 1435)	Amazon RDS Performance Insights is now available for Amazon RDS for PostgreSQL. For more information, see Using Amazon RDS Performance Insights .	July 18, 2018
Amazon RDS for Oracle supports Oracle APEX version 5.1.4.v1 (p. 1435)	Amazon RDS for Oracle now supports Oracle Application Express (APEX) version 5.1.4.v1. For more information, see Oracle Application Express .	July 10, 2018
Amazon RDS for Oracle supports publishing logs to Amazon CloudWatch Logs (p. 1435)	Amazon RDS for Oracle now supports publishing alert, audit, trace, and listener log data to a log group in CloudWatch Logs. For more information, see Publishing Oracle Logs to Amazon CloudWatch Logs .	July 9, 2018
MariaDB 10.2.15, 10.1.34, and 10.0.35 (p. 1435)	You can now create Amazon RDS DB instances running MariaDB versions 10.2.15, 10.1.34, and 10.0.35. For more information, see MariaDB on Amazon RDS Versions .	July 5, 2018

Aurora PostgreSQL 1.2 is available and compatible with PostgreSQL 9.6.8 (p. 1435)	Aurora PostgreSQL 1.2 is now available and is compatible with PostgreSQL 9.6.8. For more information, see Version 1.2 .	June 27, 2018
Read Replicas for Amazon RDS PostgreSQL support Multi-AZ deployments (p. 1435)	RDS Read Replicas in Amazon RDS PostgreSQL now support multiple Availability Zones. For more information, see Working with PostgreSQL Read Replicas .	June 25, 2018
Performance Insights available for Aurora PostgreSQL (p. 1435)	Performance Insights is generally available for Aurora PostgreSQL, with support for extended retention of performance data. For more information, see Using Amazon RDS Performance Insights .	June 21, 2018
Aurora PostgreSQL available in western US (Northern California) region (p. 1435)	Aurora PostgreSQL is now available in the western United States (Northern California) region. For more information, see Availability for Amazon Aurora PostgreSQL .	June 11, 2018
Amazon RDS for Oracle now supports CPU configuration (p. 1435)	Amazon RDS for Oracle supports configuring the number of CPU cores and the number of threads for each core for the processor of a DB instance class. For more information, see Configuring the Processor of the DB Instance Class .	June 5, 2018
Amazon RDS for Oracle April 2018 PSU (p. 1435)	Amazon RDS for Oracle has released database engine versions 12.1.0.2.v12 and 11.2.0.4.v16 to support the April 2018 Oracle Database Patch Set Update (PSU). For more information, see Oracle Database Engine Release Notes .	June 1, 2018

Earlier Updates

The following table describes the important changes in each release of the *Amazon RDS User Guide* before June 2018.

Change	Description	Date Changed
Amazon RDS for PostgreSQL now supports PostgreSQL Version	PostgreSQL version 11 Beta 1 contains several improvements that are described in PostgreSQL 11 Beta 1 Released!	May 31, 2018

Change	Description	Date Changed
11 Beta 1 in the Database Preview Environment	For information on the Database Preview Environment, see Working with the Database Preview Environment (p. 1362) .	
Amazon RDS for Oracle now supports TLS versions 1.0 and 1.2	Amazon RDS for Oracle supports Transport Layer Security (TLS) versions 1.0 and 1.2. For more information, see TLS Versions for the Oracle SSL Option (p. 1052) .	May 30, 2018
Aurora MySQL supports publishing logs to Amazon CloudWatch Logs	Aurora MySQL now supports publishing general, slow, audit, and error log data to a log group in CloudWatch Logs. For more information, see Publishing Aurora MySQL to CloudWatch Logs .	May 23, 2018
Database Preview Environment for Amazon RDS PostgreSQL	You can now launch a new instance of Amazon RDS PostgreSQL in a preview mode. For more information about the Database Preview Environment see, Working with the Database Preview Environment (p. 1362) .	May 22, 2018
Amazon RDS for Oracle DB instances support new DB instance classes	Oracle DB instances now support the db.x1e and db.x1 DB instance classes. For more information, see Choosing the DB Instance Class (p. 107) and DB Instance Class Support for Oracle (p. 917) .	May 22, 2018
Amazon RDS PostgreSQL now supports postgres_fdw on a Read Replica.	You can now use postgres_fdw to connect to a remote server from a Read Replica. For more information see, Accessing External Data with the postgres_fdw Extension (p. 1347) .	May 17, 2018
Amazon RDS for Oracle now supports setting sqlnet.ora parameters	You can now set sqlnet.ora parameters with Amazon RDS for Oracle. For more information, see Modifying Oracle sqlnet.ora Parameters (p. 976) .	May 10, 2018
Aurora PostgreSQL available in Asia Pacific (Seoul) region.	Aurora PostgreSQL is now available in the Asia Pacific (Seoul) region. For more information, see Availability for Amazon Aurora PostgreSQL .	May 9, 2018
Aurora MySQL supports backtracking	Aurora MySQL now supports "rewinding" a DB cluster to a specific time, without restoring data from a backup. For more information, see Backtracking an Aurora DB Cluster .	May 9, 2018
Aurora MySQL supports encrypted migration and replication from external MySQL	Aurora MySQL now supports encrypted migration and replication from an external MySQL database. For more information, see Migrating Data from an External MySQL Database to an Amazon Aurora MySQL DB Cluster and Replication Between Aurora and MySQL or Between Aurora and Another Aurora DB Cluster .	April 25, 2018

Change	Description	Date Changed
Aurora with PostgreSQL compatibility support for the Copy-on-Write protocol.	You can now clone databases in an Aurora PostgreSQL database cluster. For more information see, Cloning Databases in an Aurora DB Cluster .	April 10, 2018
MariaDB 10.2.12, 10.1.31, and 10.0.34	You can now create Amazon RDS DB instances running MariaDB versions 10.2.12, 10.1.31, and 10.0.34. For more information, see MariaDB on Amazon RDS Versions (p. 534) .	March 21, 2018
Aurora PostgreSQL Support for new regions	Aurora PostgreSQL is now available in the EU (London) and Asia Pacific (Singapore) regions. For more information, see Availability for Amazon Aurora PostgreSQL .	March 13, 2018
MySQL 5.7.21, 5.6.39, and 5.5.59	You can now create Amazon RDS DB instances running MySQL versions 5.7.21, 5.6.39, and 5.5.59. For more information, see MySQL on Amazon RDS Versions (p. 742) .	March 9, 2018
Amazon RDS for Oracle now supports Oracle REST Data Services	Amazon RDS for Oracle supports Oracle REST Data Services as part of the APEX option. For more information, see Oracle Application Express (p. 1016) .	March 9, 2018
Amazon Aurora with MySQL compatibility available in new AWS Region	Aurora MySQL is now available in the Asia Pacific (Singapore) region. For the complete list of AWS Regions for Aurora MySQL, see Availability for Amazon Aurora MySQL .	March 6, 2018
Support for PostgreSQL 10.1	Amazon RDS now supports version 10.1 of PostgreSQL. For more information, see PostgreSQL Version 10.1 on Amazon RDS (p. 1372)	February 27, 2018
Oracle January 2018 PSU	Amazon RDS for Oracle has released database engine versions 12.1.0.2.v11 and 11.2.0.4.v15 to support the January 2018 Oracle Database Patch Set Update (PSU). For more information, see Oracle Database Engine Release Notes (p. 1173) .	February 22, 2018
Amazon RDS DB instances running Microsoft SQL Server support change data capture (CDC)	DB instances running Amazon RDS for Microsoft SQL Server now support change data capture (CDC). For more information, see Change Data Capture Support for Microsoft SQL Server DB Instances (p. 609) .	February 6, 2018
Aurora MySQL supports a new major version	You can now create Aurora MySQL DB clusters running MySQL version 5.7. For more information, see Amazon Aurora MySQL Database Engine Updates 2018-02-06 .	February 6, 2018

Change	Description	Date Changed
Support for PostgreSQL 9.6.6	Amazon RDS PostgreSQL now supports version 9.6.6. This release also includes support for the prefix and orafce extensions. For more information, see PostgreSQL Version 9.6.6 on Amazon RDS (p. 1374) .	January 19, 2018
Publish MySQL and MariaDB logs to Amazon CloudWatch Logs	You can now publish MySQL and MariaDB log data to CloudWatch Logs. For more information, see Publishing MySQL Logs to CloudWatch Logs (p. 406) and Publishing MariaDB Logs to Amazon CloudWatch Logs (p. 395) .	January 17, 2018
Multi-AZ support for Read Replicas	You can now create a Read Replica as a Multi-AZ DB instance. Amazon RDS creates a standby of your replica in another Availability Zone for failover support for the replica. Creating your Read Replica as a Multi-AZ DB instance is independent of whether the source database is a Multi-AZ DB instance. For more information, see Working with Read Replicas (p. 182) .	January 11, 2018
Amazon RDS for MariaDB supports a new major version	You can now create Amazon RDS DB instances running MariaDB version 10.2. For more information, see MariaDB 10.2 Support on Amazon RDS (p. 535) .	January 3, 2018
Amazon Aurora with PostgreSQL compatibility available in new AWS Region	Aurora PostgreSQL is now available in the EU (Paris) region. For the complete list of AWS Regions for Aurora PostgreSQL, see Availability for Amazon Aurora PostgreSQL .	December 22, 2017
Aurora PostgreSQL supports new instance types	Aurora PostgreSQL now supports new instance types. For the complete list of instance types, see Choosing the DB Instance Class .	December 20, 2017
Oracle October 2017 PSU	Amazon RDS for Oracle has released database engine versions 12.1.0.2.v10 and 11.2.0.4.v14 to support the October 2017 Oracle Database Patch Set Update (PSU). For more information, see Oracle Database Engine Release Notes (p. 1173) .	December 19, 2017
Amazon Aurora with MySQL compatibility available in new AWS Region	Aurora MySQL is now available in the EU (Paris) region. For the complete list of AWS Regions for Aurora MySQL, see Availability for Amazon Aurora MySQL .	December 18, 2017
Aurora MySQL supports hash joins	This feature can improve query performance when you need to join a large amount of data by using an equijoin. For more information, see Working with Hash Joins in Aurora MySQL .	December 11, 2017
Aurora MySQL supports native functions to invoke AWS Lambda functions	You can call the native functions <code>lambda_sync</code> and <code>lambda_async</code> when you use Aurora MySQL. For more information, see Invoking a Lambda Function from an Amazon Aurora MySQL DB Cluster .	December 11, 2017

Change	Description	Date Changed
Added Aurora PostgreSQL HIPAA compliance	Aurora PostgreSQL now supports building HIPAA compliant applications, see Working with Amazon Aurora PostgreSQL .	December 6, 2017
Additional AWS Regions available for Amazon Aurora with PostgreSQL compatibility	Amazon Aurora with PostgreSQL compatibility is now available in four new AWS Regions. For more information, see Availability for Amazon Aurora PostgreSQL .	November 22, 2017
Modify storage for Amazon RDS DB instances running Microsoft SQL Server	You can now modify the storage of your Amazon RDS DB instances running SQL Server. For more information, see Modifying a DB Instance Running the Microsoft SQL Server Database Engine (p. 643) .	November 21, 2017
Amazon RDS supports 16 TiB storage for Linux-based engines	You can now create MySQL, MariaDB, PostgreSQL, and Oracle RDS DB instances with up to 16 TiB of storage. For more information, see Amazon RDS DB Instance Storage (p. 135) .	November 21, 2017
Amazon RDS supports fast scale up of storage	You can now add storage to MySQL, MariaDB, PostgreSQL, and Oracle RDS DB instances in a few minutes. For more information, see Amazon RDS DB Instance Storage (p. 135) .	November 21, 2017
Amazon RDS supports MariaDB versions 10.1.26 and 10.0.32	You can now create Amazon RDS DB instances running MariaDB versions 10.1.26 and 10.0.32. For more information, see MariaDB on Amazon RDS Versions (p. 534) .	November 20, 2017
Amazon RDS for Microsoft SQL Server now supports new DB instance classes	You can now create Amazon RDS DB instances running SQL Server that use the db.r4 and db.m4.16xlarge DB instance classes. For more information, see DB Instance Class Support for Microsoft SQL Server (p. 602) .	November 20, 2017
Amazon RDS for MySQL and MariaDB now supports new DB instance classes	You can now create Amazon RDS DB instances running MySQL and MariaDB that use the db.r4, db.m4.16xlarge, db.t2.xlarge, and db.t2.2xlarge DB instance classes. For more information, see Choosing the DB Instance Class (p. 107) .	November 20, 2017
SQL Server 2017	You can now create Amazon RDS DB instances running Microsoft SQL Server 2017. You can also create DB instances running SQL Server 2016 SP1 CU5. For more information, see Microsoft SQL Server on Amazon RDS (p. 598) .	November 17, 2017
Restore MySQL backups from Amazon S3	You can now create a backup of your on-premises database, store it on Amazon S3, and then restore the backup file onto a new Amazon RDS DB instance running MySQL. For more information, see Restoring a Backup into an Amazon RDS MySQL DB Instance (p. 826) .	November 17, 2017

Change	Description	Date Changed
Auto Scaling with Aurora Replicas	Amazon Aurora MySQL now supports Aurora Auto Scaling. Aurora Auto Scaling dynamically adjusts the number of Aurora Replicas based on increases or decreases in connectivity or workload. For more information, see Using Amazon Aurora Auto Scaling with Aurora Replicas .	November 17, 2017
Oracle default edition support	Amazon RDS for Oracle DB instances now supports setting the default edition for the DB instance. For more information, see Setting the Default Edition for a DB Instance (p. 1100) .	November 3, 2017
Oracle DB instance file validation	Amazon RDS for Oracle DB instances now supports validating DB instance files with the Oracle Recovery Manager (RMAN) logical validation utility. For more information, see Validating DB Instance Files (p. 1115) .	November 3, 2017
Oracle July 2017 PSU	Amazon RDS for Oracle has released database engine versions 12.1.0.2.v9 and 11.2.0.4.v13 to support the July 2017 Oracle Database Patch Set Update (PSU). For more information, see Oracle Database Engine Release Notes (p. 1173) .	November 3, 2017
Management Agent for OEM 13c	Amazon RDS Oracle DB instances now support the Management Agent for Oracle Enterprise Manager (OEM) Cloud Control 13c. For more information, see Oracle Management Agent for Enterprise Manager Cloud Control (p. 1028) .	November 1, 2017
PostgreSQL 9.6.5, 9.5.9, 9.4.14, and 9.3.19	You can now create Amazon RDS DB instances running PostgreSQL versions 9.6.5., 9.5.9, 9.4.14, and 9.3.19. For more information, see Supported PostgreSQL Database Versions (p. 1365) .	November 1, 2017
Storage reconfiguration for Microsoft SQL Server snapshots	You can now reconfigure the storage when you restore a snapshot to an Amazon RDS DB instance running Microsoft SQL Server. For more information, see Restoring from a DB Snapshot (p. 266) .	October 26, 2017
Asynchronous key prefetch for Aurora with MySQL compatibility	Asynchronous key prefetch (AKP) improves the performance of noncached index joins, by prefetching keys in memory ahead of when they are needed. For more information, see Working with Asynchronous Key Prefetch in Amazon Aurora .	October 26, 2017
MySQL 5.7.19, 5.6.37, and 5.5.57	You can now create Amazon RDS DB instances running MySQL versions 5.7.19, 5.6.37, and 5.5.57. For more information, see MySQL on Amazon RDS Versions (p. 742) .	October 25, 2017
General availability of Amazon Aurora with PostgreSQL compatibility	Amazon Aurora with PostgreSQL compatibility makes it simple and cost-effective to set up, operate, and scale your new and existing PostgreSQL deployments, thus freeing you to focus on your business and applications. For more information, see Working with Amazon Aurora PostgreSQL .	October 24, 2017

Change	Description	Date Changed
Amazon RDS for Oracle DB instances support new DB instance classes	Amazon RDS Oracle DB instances now support Memory Optimized Next Generation (db.r4) instance classes. Amazon RDS Oracle DB instances also now support the following new current generation instance classes: db.m4.16xlarge, db.t2.xlarge, and db.t2.2xlarge. For more information, see Choosing the DB Instance Class (p. 107) and DB Instance Class Support for Oracle (p. 917) .	October 23, 2017
New feature	Your new and existing Reserved Instances can now cover multiple sizes in the same DB instance class. Size-flexible reserved instances are available for DB instances with the same AWS Region, database engine, and instance family, and across AZ configuration. Size-flexible reserved instances are available for the following database engines: Amazon Aurora, MariaDB, MySQL, Oracle (Bring Your Own License), PostgreSQL. For more information, see Size-Flexible Reserved DB Instances (p. 241) .	October 11, 2017
New feature	You can now use the Oracle SQLT option to tune a SQL statement for optimal performance. For more information, see Oracle SQLT (p. 1063) .	September 22, 2017
New feature	If you have existing manual DB snapshots of your Amazon RDS Oracle DB instances, you can now upgrade them to a later version of the Oracle database engine. For more information, see Upgrading an Oracle DB Snapshot (p. 984) .	September 20, 2017
New feature	You can now use Oracle Spatial to store, retrieve, update, and query spatial data in your Amazon RDS DB instances running Oracle. For more information, see Oracle Spatial (p. 1060) .	September 15, 2017
New feature	You can now use Oracle Locator to support internet and wireless service-based applications and partner-based GIS solutions with your Amazon RDS DB instances running Oracle. For more information, see Oracle Locator (p. 1043) .	September 15, 2017
New feature	You can now use Oracle Multimedia to store, manage, and retrieve images, audio, video, and other heterogeneous media data in your Amazon RDS DB instances running Oracle. For more information, see Oracle Multimedia (p. 1046) .	September 15, 2017
New feature	You can now export audit logs from your Amazon Aurora MySQL DB clusters to Amazon CloudWatch Logs. For more information, see Publishing Aurora MySQL Logs to Amazon CloudWatch Logs .	September 14, 2017
New feature	Amazon RDS now supports multiple versions of Oracle Application Express (APEX) for your DB instances running Oracle. For more information, see Oracle Application Express (p. 1016) .	September 13, 2017

Change	Description	Date Changed
New feature	You can now use Amazon Aurora to migrate an unencrypted or encrypted DB snapshot or Amazon RDS MySQL DB instance to an encrypted Aurora MySQL DB cluster. For more information, see Migrating an RDS MySQL Snapshot to Aurora and Migrating Data from a MySQL DB Instance to an Amazon Aurora MySQL DB Cluster by Using an Aurora Read Replica .	September 5, 2017
New feature	You can use Amazon RDS for Microsoft SQL Server databases to build HIPAA-compliant applications. For more information, see Compliance Program Support for Microsoft SQL Server DB Instances (p. 604) .	August 31, 2017
New feature	You can now use Amazon RDS for MariaDB databases to build HIPAA-compliant applications. For more information, see MariaDB on Amazon RDS (p. 532) .	August 31, 2017
New feature	You can now create Amazon RDS DB instances running Microsoft SQL Server with allocated storage up to 16 TiB, and Provisioned IOPS to storage ranges of 1:1–50:1. For more information, see Amazon RDS DB Instance Storage (p. 135) .	August 22, 2017
New feature	You can now use Multi-AZ deployments for DB instances running Microsoft SQL Server in the EU (Frankfurt) region. For more information, see Multi-AZ Deployments for Microsoft SQL Server (p. 681) .	August 3, 2017
New feature	You can now create Amazon RDS DB instances running MariaDB versions 10.1.23 and 10.0.31. For more information, see MariaDB on Amazon RDS Versions (p. 534) .	July 17, 2017
New feature	Amazon RDS now supports Microsoft SQL Server Enterprise Edition with the License Included model in all AWS Regions. For more information, see Licensing Microsoft SQL Server on Amazon RDS (p. 615) .	July 13, 2017
New feature	Amazon RDS for Oracle now supports Linux kernel huge pages for increased database scalability. The use of huge pages results in smaller page tables and less CPU time spent on memory management, increasing the performance of large database instances. You can use huge pages with your Amazon RDS DB instances running all editions of Oracle versions 12.1.0.2 and 11.2.0.4. For more information, see Using Huge Pages with an Oracle DB Instance (p. 935) .	July 7, 2017
New feature	Updated to support encryption at rest (EAR) for db.t2.small and db.t2.medium DB instance classes for all non-Aurora DB engines. For more information, see Availability of Amazon RDS Encryption (p. 430) .	June 27, 2017
New feature	Updated to support Amazon Aurora in the Europe (Frankfurt) region. For more information, see Availability for Amazon Aurora MySQL .	June 16, 2017

Change	Description	Date Changed
New feature	You can now specify an option group when you copy a DB snapshot across AWS regions. For more information, see Option Group Considerations (p. 270) .	June 12, 2017
New feature	You can now copy DB snapshots created from specialized DB instances across AWS regions. You can copy snapshots from DB instances that use Oracle TDE, Microsoft SQL Server TDE, and Microsoft SQL Server Multi-AZ with Mirroring. For more information, see Copying a DB Snapshot (p. 271) .	June 12, 2017
New feature	Amazon Aurora now allows you to quickly and cost-effectively copy all of your databases in an Amazon Aurora DB cluster. For more information, see Cloning Databases in an Aurora DB Cluster .	June 12, 2017
New feature	Amazon RDS now supports Microsoft SQL Server 2016 SP1 CU2. For more information, see Microsoft SQL Server on Amazon RDS (p. 598) .	June 7, 2017
New feature	Amazon RDS for Oracle has released database engine versions 12.1.0.2.v8 and 11.2.0.4.v12 to support the April 2017 Oracle Database Patch Set Update (PSU). For more information, see Oracle Database Engine Release Notes (p. 1173) .	May 23, 2017
New Feature	Amazon RDS now supports PostgreSQL versions 9.6.2, 9.5.6, 9.4.11, and 9.3.16. For more information, see Supported PostgreSQL Database Versions (p. 1365)	May 3, 2017
Preview	Public preview of Amazon Aurora with PostgreSQL Compatibility. For more information, see Working with Amazon Aurora PostgreSQL .	April 19, 2017
New feature	Amazon Aurora now allows you to execute an <code>ALTER TABLE <i>tbl_name</i> ADD COLUMN <i>col_name</i> <i>column_definition</i></code> operation nearly instantaneously. The operation completes without requiring the table to be copied and without materially impacting other DML statements. For more information, see Altering Tables in Amazon Aurora Using Fast DDL .	April 5, 2017
New feature	We have added a new monitoring command, <code>SHOW VOLUME STATUS</code> , to display the number of nodes and disks in a volume. For more information, see Displaying Volume Status for an Aurora DB Cluster .	April 5, 2017
New feature	Amazon RDS for Oracle now includes the January 2017 Oracle Database Patch Set Update (PSU). This adds support for database engine versions 12.1.0.2.v7 and 11.2.0.4.v11. For more information, see Oracle Database Engine Release Notes (p. 1173) .	March 21, 2017
New feature	You can now use your own custom logic in your custom password verification functions for Oracle on Amazon RDS. For more information, see Creating Custom Functions to Verify Passwords (p. 1091) .	March 21, 2017

Change	Description	Date Changed
New feature	You can now access your online and archived redo log files on your Oracle DB instances on Amazon RDS. For more information, see Accessing Transaction Logs (p. 1111) .	March 21, 2017
New feature	You can now copy both encrypted and unencrypted DB cluster snapshots between accounts in the same region. For more information, see Copying a DB Cluster Snapshot Across Accounts .	March 7, 2017
New feature	You can now share encrypted DB cluster snapshots between accounts in the same region. For more information, see Sharing a DB Cluster Snapshot .	March 7, 2017
New feature	You can now replicate encrypted Amazon Aurora MySQL DB clusters to create cross-region Aurora Replicas. For more information, see Replicating Aurora MySQL DB Clusters Across AWS Regions .	March 7, 2017
New feature	You can now require that all connections to your DB instance running Microsoft SQL Server use Secure Sockets Layer (SSL). For more information, see Using SSL with a Microsoft SQL Server DB Instance (p. 687) .	February 27, 2017
New feature	You can now set your local time zone to one of 15 additional time zones. For more information, see Supported Time Zones (p. 611) .	February 27, 2017
New feature	You can now use the Amazon RDS procedure <code>msdb.dbo.rds_shrink_tempdbfile</code> to shrink the tempdb database on your DB instances running Microsoft SQL Server. For more information, see Shrinking the tempdb Database (p. 724) .	February 17, 2017
New feature	You can now compress your backup file when you export your Enterprise and Standard Edition Microsoft SQL Server database from an Amazon RDS DB instance to Amazon S3. For more information, see Compressing Backup Files (p. 669) .	February 17, 2017
New feature	Amazon RDS now supports custom DNS servers to resolve DNS names used in outbound network access on your DB instances running Oracle. For more information, see Setting Up a Custom DNS Server (p. 1094) .	January 26, 2017
New feature	Amazon RDS now supports creating an encrypted Read Replica in another region. For more information, see Creating a Read Replica in a Different AWS Region (p. 188) and CreateDBInstanceReadReplica .	January 23, 2017
New feature	Amazon RDS now supports upgrading a MySQL DB snapshot from MySQL 5.1 to MySQL 5.5. For more information, see Upgrading a MySQL DB Snapshot (p. 818) and ModifyDBSnapshot .	January 20, 2017

Change	Description	Date Changed
New feature	Amazon RDS now supports copying an encrypted DB snapshot to another region for the MariaDB, MySQL, Oracle, PostgreSQL, and Microsoft SQL Server database engines. For more information, see Copying a DB Snapshot (p. 271) and CopyDBSnapshot .	December 20, 2016
New feature	Amazon RDS now supports migrating an Amazon RDS MySQL 5.6 DB snapshot to a new DB instance running MariaDB 10.1. For more information, see Migrating Data from a MySQL DB Snapshot to a MariaDB DB Instance (p. 575) .	December 20, 2016
New feature	Amazon Aurora MySQL now supports spatial indexing. Spatial indexing improves query performance on large datasets for queries that use spatial data. For more information, see Amazon Aurora MySQL and Spatial Data .	December 14, 2016
New feature	Amazon RDS for Oracle now includes the October 2016 Oracle Database Patch Set Update (PSU). This adds support for Oracle database engine versions 12.1.0.2.v6 and 11.2.0.4.v10. For more information, see Oracle Database Engine Release Notes (p. 1173) .	December 12, 2016
New feature	Amazon RDS now supports outbound network access on your DB instances running Oracle. You can use <code>utl_http</code> , <code>utl_tcp</code> , and <code>utl_smtp</code> to connect from your DB instance to the network. For more information, see Using utl_http, utl_tcp, and utl_smtp with an Oracle DB Instance (p. 938) .	December 5, 2016
New feature	Amazon RDS has retired support for MySQL version 5.1. However, you can restore existing MySQL 5.1 snapshots to a MySQL 5.5 instance. For more information, see Supported Storage Engines for MySQL on Amazon RDS (p. 744) .	November 15, 2016
New feature	Amazon RDS now supports PostgreSQL version 9.6.1. For more information, see PostgreSQL Version 9.6.1 on Amazon RDS (p. 1375) .	November 11, 2016
New feature	Amazon RDS now supports Microsoft SQL Server 2016 RTM CU2. For more information, see Microsoft SQL Server on Amazon RDS (p. 598) .	November 4, 2016
New feature	Amazon RDS now supports major version upgrades for DB instances running Oracle. You can now upgrade your Oracle DB instances from 11g to 12c. For more information, see Upgrading the Oracle DB Engine (p. 979) .	November 2, 2016
New feature	You can now create DB instances running Microsoft SQL Server 2014 Enterprise Edition. Amazon RDS now supports SQL Server 2014 SP2 for all editions and all regions. For more information, see Microsoft SQL Server on Amazon RDS (p. 598) .	October 25, 2016

Change	Description	Date Changed
New feature	Amazon Aurora MySQL now integrates with other AWS services: You can load text or XML data into a table from an Amazon S3 bucket, or invoke an AWS Lambda function from database code. For more information, see Integrating Aurora MySQL with Other AWS Services .	October 18, 2016
New feature	You can now access the tempdb database on your Amazon RDS DB instances running Microsoft SQL Server. You can access the tempdb database by using Transact-SQL through Microsoft SQL Server Management Studio (SSMS), or any other standard SQL client application. For more information, see Accessing the tempdb Database on Microsoft SQL Server DB Instances on Amazon RDS (p. 724) .	September 29, 2016
New feature	You can now use the UTL_MAIL package with your Amazon RDS DB instances running Oracle. For more information, see Oracle UTL_MAIL (p. 1076) .	September 20, 2016
New feature	Amazon RDS for Oracle now includes the July 2016 Oracle Database Patch Set Update (PSU). This adds support for Oracle database engine versions 12.1.0.2.v5, 12.1.0.1.v6, and 11.2.0.4.v9. For more information, see Oracle Database Engine Release Notes (p. 1173) .	September 20, 2016
New features	You can now set the time zone of your new Microsoft SQL Server DB instances to a local time zone, to match the time zone of your applications. For more information, see Local Time Zone for Microsoft SQL Server DB Instances (p. 610) .	September 19, 2016
New features	Added support for new PostgreSQL versions 9.5.4, 9.4.9, and 9.3.14. Also added support for PostgreSQL logical replication, PostgreSQL event triggers, and RAM disk for the PostgreSQL stats_temp_directory. For more information, see Supported PostgreSQL Database Versions (p. 1365) , Logical Replication for PostgreSQL on Amazon RDS (p. 1406) , Event Triggers for PostgreSQL on Amazon RDS (p. 1408) , and RAM Disk for the stats_temp_directory (p. 1409) .	September 14, 2016
New feature	You can now use the Oracle Label Security option to control access to individual table rows in your Amazon RDS DB instances running Oracle 12c. With Oracle Label Security, you can enforce regulatory compliance with a policy-based administration model, and ensure that an access to sensitive data is restricted to only users with the appropriate clearance level. For more information, see Oracle Label Security (p. 1040) .	September 8, 2016

Change	Description	Date Changed
New feature	You can now connect to an Amazon Aurora DB cluster using the reader endpoint, which load-balances connections across the Aurora Replicas that are available in the DB cluster. As clients request new connections to the reader endpoint, Aurora distributes the connection requests among the Aurora Replicas in the DB cluster. This functionality can help balance your read workload across multiple Aurora Replicas in your DB cluster. For more information, see Amazon Aurora Endpoints .	September 8, 2016
New feature	You can now support the Oracle Enterprise Manager Cloud Control on your Amazon RDS DB instances running Oracle. You can enable the Management Agent on your DB instances, and share data with your Oracle Management Service (OMS). For more information, see Oracle Management Agent for Enterprise Manager Cloud Control (p. 1028) .	September 1, 2016
New feature	This release adds support to get an ARN for a resource. For more information, see Getting an Existing ARN (p. 227) .	August 23, 2016
New feature	You can now assign up to 50 tags for each Amazon RDS resource, for managing your resources and tracking costs. For more information, see Tagging Amazon RDS Resources (p. 177) .	August 19, 2016
New feature	Amazon RDS now supports the License Included model for Oracle Standard Edition Two. For more information, see Creating a DB Instance Running the Oracle Database Engine (p. 942) . You can now change the license model of your Amazon RDS DB instances running Microsoft SQL Server and Oracle. For more information, see Licensing Microsoft SQL Server on Amazon RDS (p. 615) and Oracle Licensing (p. 916) .	August 5, 2016
New feature	You can now use the AWS Management Console to easily move your DB instance to a different VPC, or to a different subnet group in the same VPC. For more information, see Updating the VPC for a DB Instance (p. 519) . If your DB instance is not in a VPC, you can now use the AWS Management Console to easily move your DB instance into a VPC. For more information, see Moving a DB Instance Not in a VPC into a VPC (p. 520) .	August 4, 2016

Change	Description	Date Changed
New feature	Amazon RDS now supports native backup and restore for Microsoft SQL Server databases using full backup files (.bak files). You can now easily migrate SQL Server databases to Amazon RDS, and import and export databases in a single, easily-portable file, using Amazon S3 for storage, and AWS KMS for encryption. For more information, see Importing and Exporting SQL Server Databases (p. 657) .	July 27, 2016
New feature	You can now copy the source files from a MySQL database to an Amazon Simple Storage Service (Amazon S3) bucket, and then restore an Amazon Aurora DB cluster from those files. This option can be considerably faster than migrating data using <code>mysqldump</code> . For more information, see Migrating Data from an External MySQL Database to an Aurora MySQL DB Cluster .	July 20, 2016
New feature	You can now restore an unencrypted Amazon Aurora DB cluster snapshot to create an encrypted Amazon Aurora DB cluster by including an AWS Key Management Service (AWS KMS) encryption key during the restore operation. For more information, see Encrypting Amazon RDS Resources .	June 30, 2016
New feature	Amazon RDS for Oracle now includes the April 2016 Oracle Database Patch Set Update (PSU). This PSU adds support for Oracle database engine versions 12.1.0.2.v4, 12.1.0.1.v5, and 11.2.0.4.v8. For more information, see Oracle Database Engine Release Notes (p. 1173) .	June 17, 2016
New feature	You can use the Oracle Repository Creation Utility (RCU) to create a repository on Amazon RDS for Oracle. For more information, see Using the Oracle Repository Creation Utility on Amazon RDS for Oracle (p. 1165) .	June 17, 2016
New feature	Adds support for PostgreSQL cross-region Read Replicas. For more information, see Creating a Read Replica in a Different AWS Region (p. 188) .	June 16, 2016
New feature	You can now use the AWS Management Console to easily add Multi-AZ with Mirroring to a Microsoft SQL Server DB instance. For more information, see Adding Multi-AZ to a Microsoft SQL Server DB Instance (p. 681) .	June 9, 2016
New feature	You can now use Multi-AZ Deployments Using SQL Server Mirroring in the following additional regions: Asia Pacific (Sydney), Asia Pacific (Tokyo), and South America (Sao Paulo). For more information, see Multi-AZ Deployments for Microsoft SQL Server (p. 681) .	June 9, 2016
New feature	Updated to support MariaDB version 10.1. For more information, see MariaDB on Amazon RDS (p. 532) .	June 1, 2016

Change	Description	Date Changed
New feature	Updated to support Amazon Aurora cross-region DB clusters that are Read Replicas. For more information, see Replicating Aurora MySQL DB Clusters Across AWS Regions .	June 1, 2016
New feature	Enhanced Monitoring is now available for Oracle DB instances. For more information, see Enhanced Monitoring (p. 307) and Modifying a DB Instance Running the Oracle Database Engine (p. 966) .	May 27, 2016
New feature	Updated to support manual snapshot sharing for Amazon Aurora DB cluster snapshots. For more information, see Sharing a DB Cluster Snapshot .	May 18, 2016
New feature	You can now use the MariaDB Audit Plugin to log database activity on MariaDB and MySQL database instances. For more information, see Options for MariaDB Database Engine (p. 587) and Options for MySQL DB Instances (p. 873) .	April 27, 2016
New feature	In-place, major version upgrades are now available for upgrading from MySQL version 5.6 to version 5.7. For more information, see Upgrading the MySQL DB Engine (p. 812) .	April 26, 2016
New feature	Enhanced Monitoring is now available for Microsoft SQL Server DB instances. For more information, see Enhanced Monitoring (p. 307) .	April 22, 2016
New feature	Added support for PostgreSQL versions 9.5.2, 9.4.7, and 9.3.12. For more information, see Supported PostgreSQL Database Versions (p. 1365) .	April 8, 2016
New feature	Updated to support Oracle database versions 11.2.0.4.v7, 12.1.0.1.v4, and 12.1.0.2.v3 with the January 2016 Oracle Patch Set Updates (PSU). For more information, see Oracle Database Engine Release Notes (p. 1173) .	April 1, 2016
New feature	Updated to provide an Amazon Aurora Clusters view in the Amazon RDS console. For more information, see Viewing an Aurora DB Cluster .	April 1, 2016
New feature	Updated to support SQL Server Multi-AZ with mirroring in the Asia Pacific (Seoul) region. For more information, see Multi-AZ Deployments for Microsoft SQL Server (p. 681) .	March 31, 2016
New feature	Updated to support Amazon Aurora Multi-AZ with mirroring in the Asia Pacific (Seoul) region. For more information, see Availability for Amazon Aurora MySQL .	March 31, 2016
New feature	PostgreSQL DB instances have the ability to require connections to use SSL. For more information, see Using SSL with a PostgreSQL DB Instance (p. 1411) .	March 25, 2016

Change	Description	Date Changed
New feature	Enhanced Monitoring is now available for PostgreSQL DB instances. For more information, see Enhanced Monitoring (p. 307) .	March 25, 2016
New feature	Microsoft SQL Server DB instances can now use Windows Authentication for user authentication. For more information, see Using Windows Authentication with a Microsoft SQL Server DB Instance (p. 691) .	March 23, 2016
New feature	Enhanced Monitoring is now available in the Asia Pacific (Seoul) region. For more information, see Enhanced Monitoring (p. 307) .	March 16, 2016
New feature	You can now customize the order in which Aurora Replicas are promoted to primary instance during a failover. For more information, see Fault Tolerance for an Aurora DB Cluster .	March 14, 2016
New feature	Updated to support encryption when migrating to an Aurora DB cluster. For more information, see Migrating Data to an Aurora DB Cluster .	March 2, 2016
New feature	Updated to support local time zone for Aurora DB clusters. For more information, see Local Time Zone for Aurora DB Clusters .	March 1, 2016
New feature	Updated to add support for MySQL version 5.7 for current generation Amazon RDS DB instance classes.	February 22, 2016
New feature	Updated to support db.r3 and db.t2 DB instance classes in the AWS GovCloud (US-West) region.	February 11, 2016
New feature	Updated to support encrypting copies of DB snapshots and sharing encrypted DB snapshots. For more information, see Copying a Snapshot (p. 269) and Sharing a DB Snapshot (p. 278) .	February 11, 2016
New feature	Updated to support Amazon Aurora in the Asia Pacific (Sydney) region. For more information, see Availability for Amazon Aurora MySQL .	February 11, 2016
New feature	Updated to support SSL for Oracle DB Instances. For more information, see Using SSL with an Oracle DB Instance (p. 921) .	February 9, 2016
New feature	Updated to support local time zone for MySQL and MariaDB DB instances. For more information, see Local Time Zone for MySQL DB Instances (p. 749) and Local Time Zone for MariaDB DB Instances (p. 541) .	December 21, 2015
New feature	Updated to support Enhanced Monitoring of OS metrics for MySQL and MariaDB instances and Aurora DB clusters. For more information, see Viewing DB Instance Metrics (p. 304) .	December 18, 2015

Change	Description	Date Changed
New feature	Updated to support Oracle Standard Edition Two with Bring-Your- Own-License licensing. Also added support for Oracle versions 11.2.0.4.v5, 12.1.0.1.v3, and 12.1.0.2.v2. For more information, see Oracle Database Engine Release Notes (p. 1173) .	December 14, 2015
New feature	Updated to support db.t2, db.r3, and db.m4 DB instance classes for MySQL version 5.5. For more information, see Choosing the DB Instance Class (p. 107) .	December 4, 2015
New feature	Updated to support modifying the database port for an existing DB instance.	December 3, 2015
New feature	Updated to support three new extensions for PostgreSQL versions 9.3.10 and 9.4.5 DB instances. For more information, see Supported PostgreSQL Database Versions (p. 1365) .	December 1, 2015
New feature	Updated to support PostgreSQL versions 9.3.10 and 9.4.5 DB instances. For more information, see Supported PostgreSQL Database Versions (p. 1365) .	November 27, 2015
New feature	Updated to support major version upgrades of the database engine for PostgreSQL instances. For more information, see Upgrading the PostgreSQL DB Engine for Amazon RDS (p. 1289) .	November 19, 2015
New feature	Updated to support modifying the public accessibility of an existing DB instance. Updated to support db.m4 standard DB instance classes.	November 11, 2015
New feature	Updated to support manual DB snapshot sharing. For more information, see Sharing a DB Snapshot (p. 278) .	October 28, 2015
New feature	Updated to support Microsoft SQL Server 2014 for the Web, Express, and Standard editions.	October 26, 2015
New feature	Updated to support the MySQL-based MariaDB database engine. For more information, see MariaDB on Amazon RDS (p. 532) .	October 7, 2015
New feature	Updated to support Amazon Aurora in the Asia Pacific (Tokyo) region. For more information, see Availability for Amazon Aurora MySQL .	October 7, 2015
New feature	Updated to support db.t2 burst-capable DB instance classes for all DB engines and the addition of the db.t2.large DB instance class. For more information, see Choosing the DB Instance Class (p. 107) .	September 25, 2015
New feature	Updated to support Oracle DB instances on R3 and T2 DB instance classes. For more information, see Choosing the DB Instance Class (p. 107) .	August 5, 2015
New feature	Updated to support PostgreSQL versions 9.4.4 and 9.3.9. For more information, see Supported PostgreSQL Database Versions (p. 1365) .	July 30, 2015

Change	Description	Date Changed
New feature	Microsoft SQL Server Enterprise Edition is now available with the License Included service model. For more information, see Licensing Microsoft SQL Server on Amazon RDS (p. 615) .	July 29, 2015
New feature	Amazon Aurora has officially released. The Amazon Aurora DB engine supports multiple DB instances in a DB cluster. For detailed information, see What Is Amazon Aurora? .	July 27, 2015
New feature	Updated to support copying tags to DB snapshots.	July 20, 2015
New feature	Updated to support Oracle 12c database version "12.1.0.2", including the In-Memory option, Oracle 11g April PSU patches, and improved integration with AWS CloudHSM.	July 20, 2015
New feature	Updated to support increases in storage size for all DB engines and an increase in Provisioned IOPS for SQL Server.	June 18, 2015
New feature	Updated options for reserved DB instances.	June 15, 2015
New feature	Updated to support Oracle version 12c.	April 2, 2015
New feature	Updated to support PostgreSQL versions 9.3.6 and 9.4.1.	March 18, 2015
New feature	Updated to support using Amazon CloudHSM with Oracle DB instances using TDE.	January 8, 2015
New feature	Updated to support encrypting data at rest and new API version 2014-10-31.	January 6, 2015
New feature	Updated to support Oracle version 11.2.0.4.v3 that includes the PSU released in October 2014.	November 20, 2014
New feature	Updated to include the new Amazon DB engine: Aurora. The Amazon Aurora DB engine supports multiple DB instances in a DB cluster. Amazon Aurora is currently in preview release and is subject to change. For detailed information, see What Is Amazon Aurora? .	November 12, 2014
New feature	Updated to support PostgreSQL Read Replicas.	November 10, 2014
New features	Updated to support Oracle 11.2.0.4v2.	October 16, 2014
New API and features	Updated to support the GP2 storage type and new API version 2014-09-01. Updated to support the ability to copy an existing option or parameter group to create a new option or parameter group.	October 7, 2014
New feature	Updated to support InnoDB Cache Warming for DB instances running MySQL version 5.6.19 and later.	September 3, 2014
New feature	Updated to support SSL certificate verification when connecting to MySQL version 5.6, SQL Server, and PostgreSQL database engines.	August 5, 2014

Change	Description	Date Changed
New feature	Updated to support the db.t2 burst-capable DB instance classes.	August 4, 2014
New feature	Updated to support the db.r3 memory-optimized DB instance classes for use with the MySQL (version 5.6), SQL Server, and PostgreSQL database engines.	May 28, 2014
New feature	Updated to support SQL Server Multi-AZ deployments using SQL Server Mirroring.	May 19, 2014
New feature	Updated to support upgrades from MySQL version 5.5 to version 5.6.	April 23, 2014
New feature	Updated to support Oracle 11.2.0.4.	April 23, 2014
New feature	Updated to support Oracle GoldenGate.	April 3, 2014
New feature	Updated to support the M3 DB instance classes.	February 20, 2014
New feature	Updated to support the Oracle Timezone option.	January 13, 2014
New feature	Updated to support replication between Amazon RDS MySQL DB instances in different regions.	November 26, 2013
New feature	Updated to support the PostgreSQL DB engine.	November 14, 2013
New feature	Updated to support SQL Server transparent data encryption (TDE).	November 7, 2013
New API and new feature	Updated to support cross region DB snapshot copies; new API version, 2013-09-09.	October 31, 2013
New features	Updated to support Oracle Statspack.	September 26, 2013
New features	Updated to support using replication to import or export data between instances of MySQL running in Amazon RDS and instances of MySQL running on-premises or on Amazon EC2.	September 5, 2013
New features	Updated to support the db.cr1.8xlarge DB instance class for MySQL 5.6.	September 4, 2013
New feature	Updated to support replication of Read Replicas.	August 28, 2013
New feature	Updated to support parallel Read Replica creation.	July 22, 2013
New feature	Updated to support fine-grained permissions and tagging for all Amazon RDS resources.	July 8, 2013
New feature	Updated to support MySQL 5.6 for new instances, including support for the MySQL 5.6 memcached interface and binary log access.	July 1, 2013
New feature	Updated to support major version upgrades from MySQL 5.1 to MySQL 5.5.	June 20, 2013
New feature	Updated DB parameter groups to allow expressions for parameter values.	June 20, 2013

Change	Description	Date Changed
New API and new feature	Updated to support Read Replica status; new API version, 2013-05-15.	May 23, 2013
New features	Updated to support Oracle Advanced Security features for native network encryption and Oracle Transparent Data Encryption.	April 18, 2013
New features	Updated to support major version upgrades for SQL Server and additional functionality for Provisioned IOPS.	March 13, 2013
New feature	Updated to support VPC By Default for RDS.	March 11, 2013
New API and feature	Updated to support log access; new API version 2013-02-12	March 4, 2013
New feature	Updated to support RDS event notification subscriptions.	February 4, 2013
New API and feature	Updated to support DB instance renaming and the migration of DB security group members in a VPC to a VPC security group.	January 14, 2013
New feature	Updated for AWS GovCloud (US-West) support.	December 17, 2012
New feature	Updated to support m1.medium and m1.xlarge DB Instance classes.	November 6, 2012
New feature	Updated to support Read Replica promotion.	October 11, 2012
New feature	Updated to support SSL in Microsoft SQL Server DB Instances.	October 10, 2012
New feature	Updated to support Oracle micro DB Instances.	September 27, 2012
New feature	Updated to support SQL Server 2012.	September 26, 2012
New API and feature	Updated to support provisioned IOPS. API version 2012-09-17.	September 25, 2012
New features	Updated for SQL Server support for DB Instances in VPC and Oracle support for Data Pump.	September 13, 2012
New feature	Updated for support for SQL Server Agent.	August 22, 2012
New feature	Updated for support for tagging of DB Instances.	August 21, 2012
New features	Updated for support for Oracle APEX and XML DB, Oracle time zones, and Oracle DB Instances in a VPC.	August 16, 2012
New features	Updated for support for SQL Server Database Engine Tuning Advisor and Oracle DB Instances in VPC.	July 18, 2012
New feature	Updated for support for option groups and first option, Oracle Enterprise Manager Database Control.	May 29, 2012
New feature	Updated for support for Read Replicas in Amazon Virtual Private Cloud.	May 17, 2012
New feature	Updated for Microsoft SQL Server support.	May 8, 2012

Change	Description	Date Changed
New features	Updated for support for forced failover, Multi-AZ deployment of Oracle DB Instances, and nondefault character sets for Oracle DB Instances.	May 2, 2012
New feature	Updated for Amazon Virtual Private Cloud (VPC) Support.	February 13, 2012
Updated content	Updated for new Reserved Instance types.	December 19, 2011
New feature	Updated for Oracle engine support.	May 23, 2011
Updated content	Console updates.	May 13, 2011
Updated content	Edited content for shortened backup and maintenance windows.	February 28, 2011
New feature	Added support for MySQL 5.5.	January 31, 2011
New feature	Added support for Read Replicas.	October 4, 2010
New feature	Added support for AWS Identity and Access Management (IAM).	September 2, 2010
New feature	Added DB Engine Version Management.	August 16, 2010
New feature	Added Reserved DB Instances.	August 16, 2010
New Feature	Amazon RDS now supports SSL connections to your DB Instances.	June 28, 2010
New Guide	This is the first release of the Amazon RDS User Guide.	June 7, 2010

AWS Glossary

For the latest AWS terminology, see the [AWS Glossary](#) in the *AWS General Reference*.